

# CloudNets: Combining Clouds with Virtual Networking



Stefan Schmid

December, 2013

# CloudNets: Combining Clouds with Virtual Networking



Stefan Schmid

December, 2013

# Vision: Virtual Networking Cloud Resources.



**Cloud computing is a big success!  
But what is the point of clouds if they cannot be accessed?**

A screenshot of a SPIEGEL ONLINE NETZWERK article titled "Cloud Computing". The page features a large image of a smartphone displaying a grid of icons. Text on the page discusses the concept of cloud computing, mentioning Google and Amazon. The URL in the address bar is "http://www.spiegel.de/netzwelt/reviews/review/0,1518,720111,00.html".

A screenshot of The New York Times Technology section. The top navigation bar includes links for HOME PAGE, TODAY'S PAPER, VIDEO, MOST POPULAR, and TIMES TOPICS. Below the navigation is a banner for "Business Day Technology". The main headline is "Companies Slowly Join Cloud-Computing". Sub-headlines include "SAN FRANCISCO — This year, Netflix made what looked like a peculiar choice: the DVD-by-mail company decided that over the next two years, it would move most of its Web technology — customer movie queues, search tools and the like — over to the computer servers of one of its chief rivals, Amazon.com." A sidebar on the right features a video thumbnail for "3 Tips to Lose your Belly Fat" and a link to "Get DealBook by E-Mail".



# Next Natural Step for Virtualization!

## Success of Node Virtualization

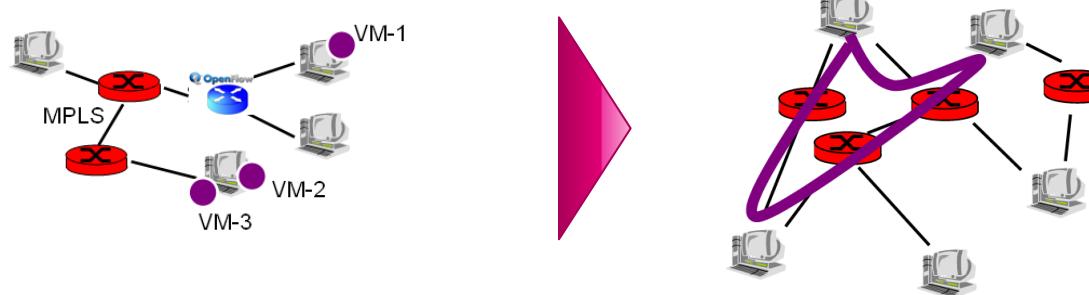
- a.k.a. end-host virtualization
- VMWare revamped server business
- OpenStack
- VM = flexible allocation, migration..
- «**Elastic computing**»

## Trend of Link Virtualization

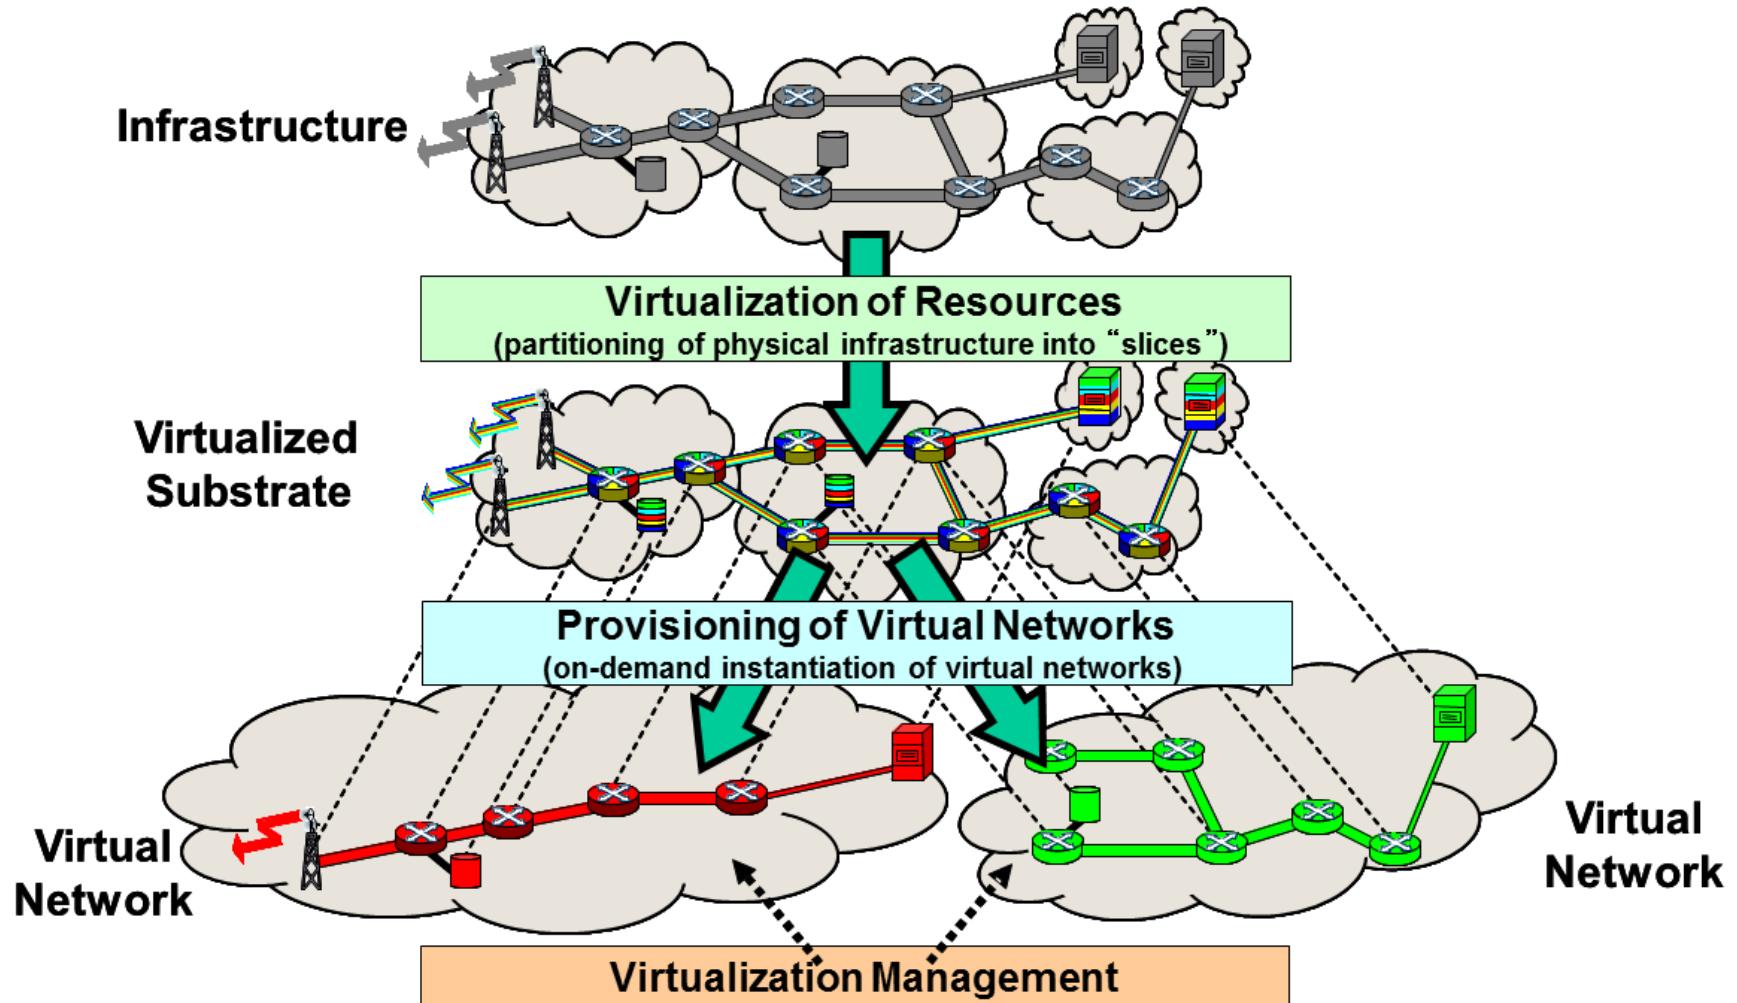
- MPLS, VPN networks, VLANs
- Software Defined Networks (SDN), OpenFlow, ...
- «The VMWare of the net»
- «**Elastic networking**»

## Unified, fully virtualized networks: **CloudNets**

„Combine **networking** with heterogeneous **cloud resources** (e.g., storage, CPU, ...)!"

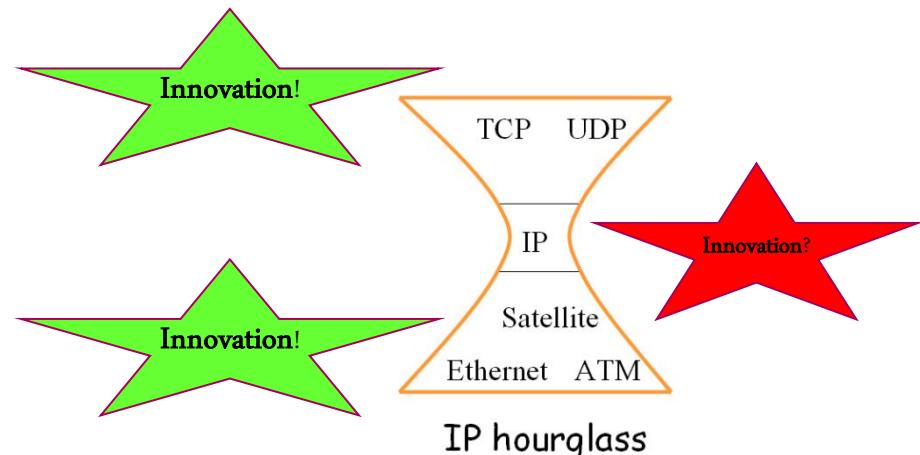
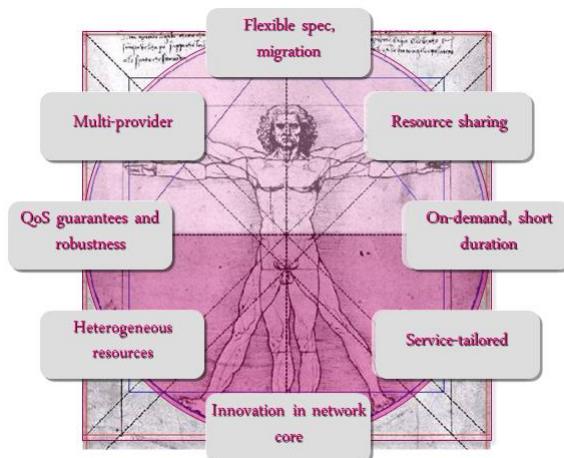


# The Vision: Sharing Resources.



Today: Internet is one solution for everything!

**CloudNets:** Flexibly specifiable virtual networks, executing different **protocol stacks**, cohabiting the same substrate.



- Vision: **facilitate innovation** in network core
  - Make **network core programmable** (e.g., own intrusion detection system)
  - **Service-tailored** networks (for social networks, bulk data transfer, live streaming, etc.)
  - Co-existing **virtual networks** with QoS guarantees
  - No dependencies on IPv4, BGP, ...

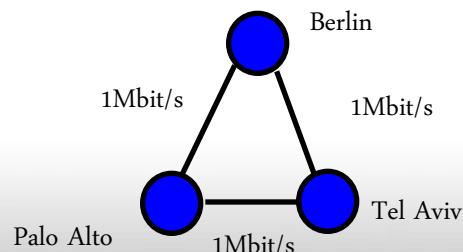


# Requests with Flexible Specification.

Optimization and Migration?

## „VPN++“

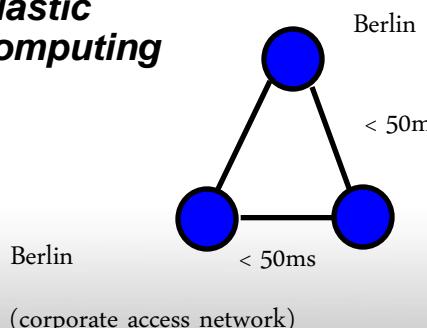
Goal: Fully specified CloudNet mapping constraints (e.g., end-points for a **telco**), but with **QoS guarantees** (e.g., bandwidth) along links



„November 22,  
1pm-2pm!“

## Spillover/Out-Sourcing

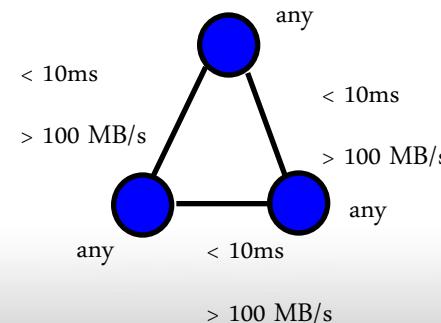
**Elastic computing**



„50 TB storage, 10 Tflops computation!“

„any European cloud provider  
(e.g. due to legal issues?)“

## Datacenters



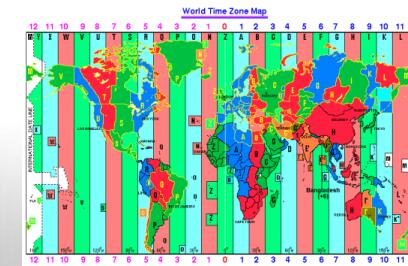
„Guaranteed resources, job deadlines met, no overhead!“

„Network may delay execution:  
costly for per hour priced VM!“

See, e.g., Octopus system (SIGCOMM 2011)

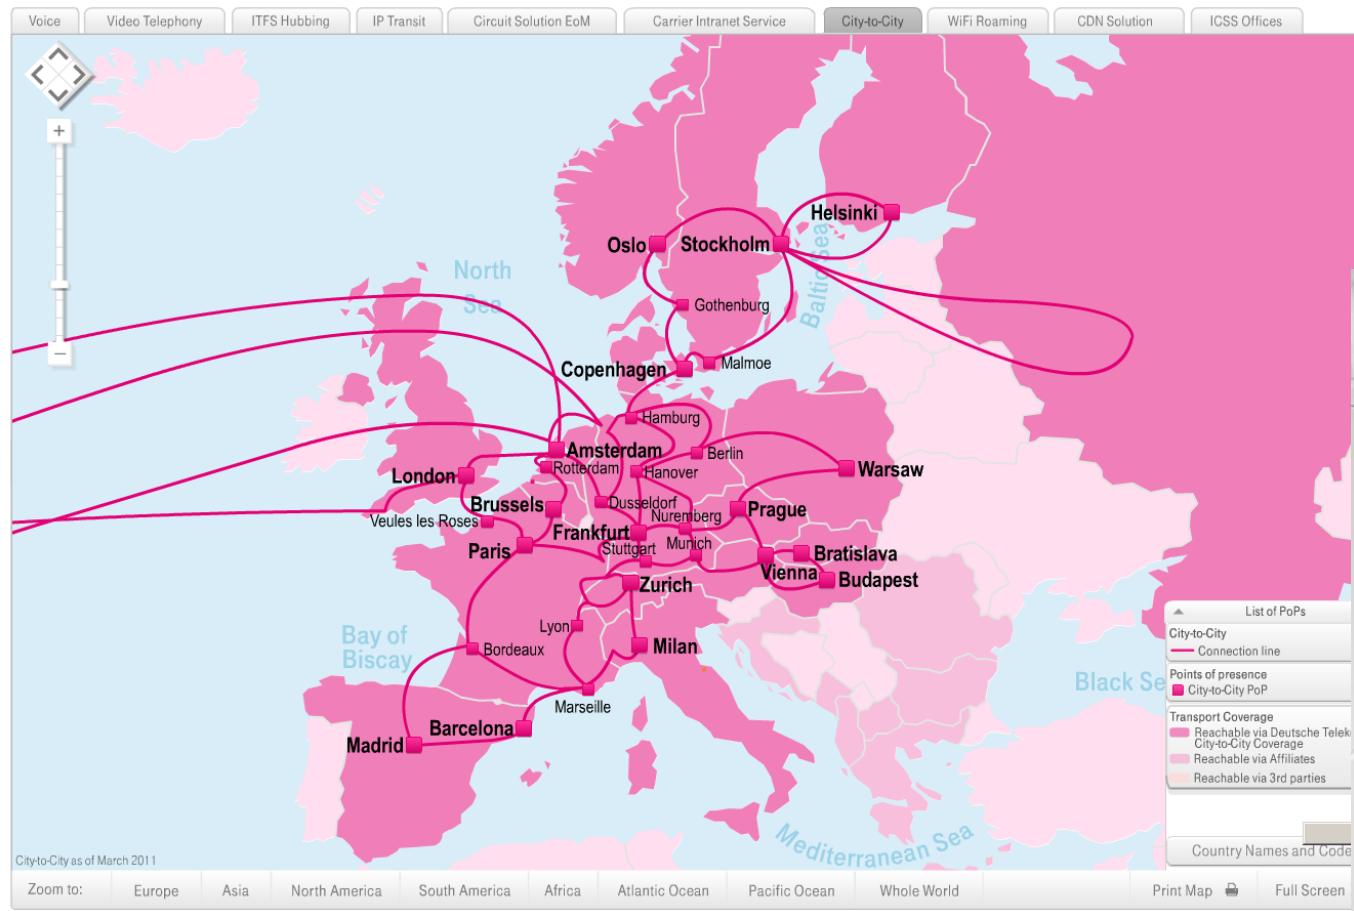
## Migration / Service Deployment

Goal: Move with the sun, with the commuters, (QoS) allow for **maintenance**, avoid roaming costs...: e.g., SAP/game/translator server, small CDN server...



# Vision: Not only in data centers, but WAN

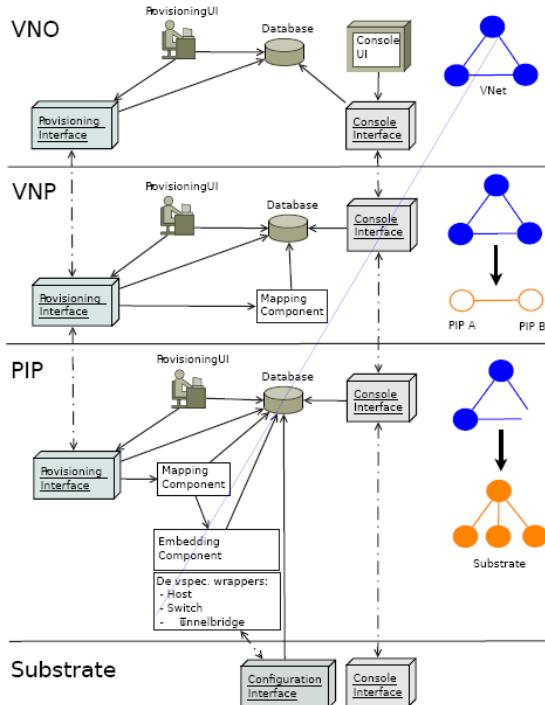
ISP network with resources at Points-of-Presence («nano datacenters»):  
For service deployment!



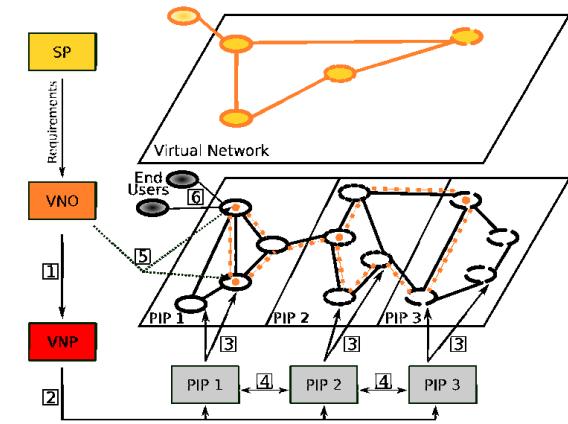
# Opens New Business Roles.

Focus of our work architecture (unlike, e.g., single authority in GINI **testbed**), on **multitude of players**, providers, ...!  
 (Of course, cross-layer infos if all same company...)

## Roles



- Service Provider (SP): uses CloudNets to offer its services (streaming, OSN, CDN, ...): e.g., **value-added application** CloudNet, or transport CloudNet
- Virtual Network Operator (VNO): Installs and **operates** CloudNet over topology provided by VNP, offers **tailored** connectivity service, triggers cross-provider migration (by setting requirements), ...
- Virtual Network Provider (VNP): „Broker“/reseller that assembles virtual resources from different PIPs to provide **virtual topology** (no need to know PIP, can be **recursive**, ...)
- Physical Infrastructure Provider (PIP): Owns and manages physical **infrastructure**



QoS from PIP up to VNO or service provider: accounting via complete set of contracts! (unlike „sending party pays“)



# Federated CloudNet Architecture.

SIGCOMM VISA 2009

New Business Roles.



As in Internet today:  
Netflix, Google, World  
of Warcraft...

As in Internet today:  
Telekom, AT&T, ...

+ resource control interface  
(bootstrapping etc.)

## Roles in CloudNet Arch.

### Service Provider (SP)

(offers services over the top)

### Virtual Network Operator (VNO)

(operates CloudNet, Layer 3+, triggers migration)

### Virtual Network Provider (VNP)

(resource broker, compiles resources)

### Physical Infrastructure Provider (PIP)

(resource and bit pipe provider)



# Federated CloudNet Architecture.

SIGCOMM VISA 2009

New Business Roles.



As in Internet today:  
Netflix, Google, World  
of Warcraft...

As in Internet today:  
Telekom, AT&T, ...  
+ resource control interface  
(bootstrapping etc.)

## Roles in CloudNet Arch.

knows  
application  
(offers services at the top)

### Virtual Network Operator (VNO)

(operates CloudNet, Layer 3+, triggers migration)

### Virtual Network Provider (VNP)

(resource broker, allocates resources)

### Physical Infrastructure Provider (PIP)

knows network  
(uses resources at  
PoPs!)

# Federated CloudNet Architecture.

SIGCOMM VISA 2009

New Business Roles.

## Roles in CloudNet Arch.

### Service Provider (SP)

(offers services over the top)

### Virtual Network Operator (VNO)

(operates CloudNet, Layer 3+, triggers migration)

### Virtual Network Provider (VNP)

(resource broker, compiles resources)

### Physical Infrastructure Provider (PIP)

(resource and bit pipe provider)

Provide layer 2: assembles CloudNets,  
resource and management interfaces,  
provides indirection layer, across PIPs!

View: PIP graph

Resource broker... (recursive)

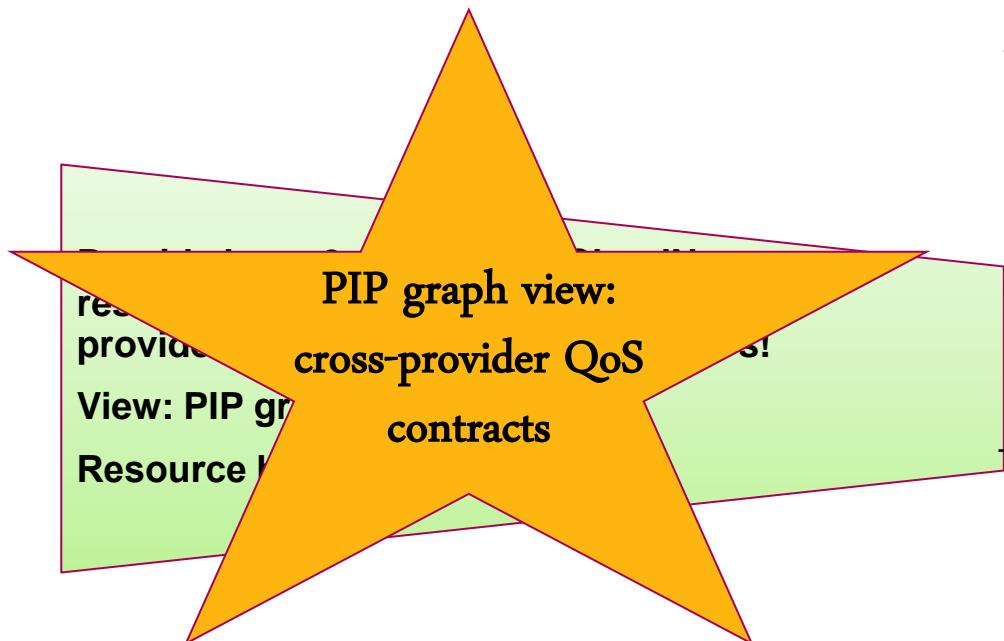
A virtualized infrastructure opens new roles for the allocation of  
resources and the operation of networks!



# Federated CloudNet Architecture.

SIGCOMM VISA 2009

New Business Roles.



## Roles in CloudNet Arch.

### Service Provider (SP)

(offers services over the top)

### Virtual Network Operator (VNO)

(operates CloudNet, Layer 3+, triggers migration)

### Virtual Network Provider (VNP)

(resource broker, compiles resources)

### Physical Infrastructure Provider (PIP)

(resource and bit pipe provider)

A virtualized infrastructure opens new roles for the allocation of resources and the operation of networks!



# Federated CloudNet Architecture.

SIGCOMM VISA 2009

New Business Roles.

## Roles in CloudNet Arch.

### Service Provider (SP)

(offers services over the top)

### Virtual Network Operator (VNO)

(operates CloudNet, Layer 3+, triggers migration)

### Virtual Network Provider (VNP)

(resource broker, compiles resources)

### Physical Infrastructure Provider (PIP)

(resource and bit pipe provider)

Build upon layer 2 (op on virt IDs): clean slate!  
(OSN, ...)

Routing, addressing, multi-path/redundancy...  
(view inside CloudNet!)

Trigger migration (use provisioning interface of  
SP or live)

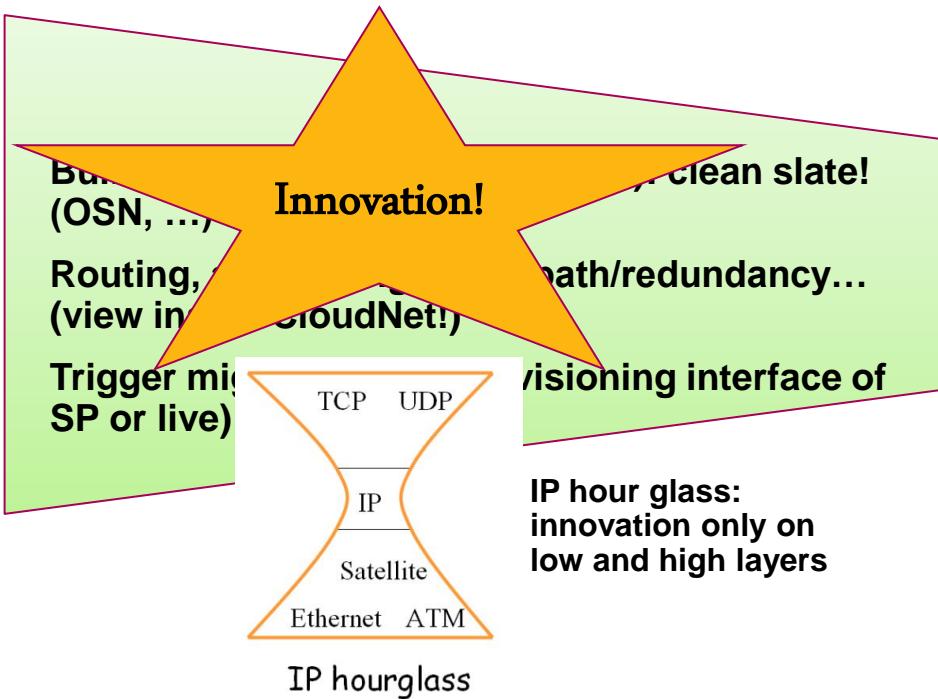
A virtualized infrastructure opens new roles for the allocation of  
resources and the operation of networks!



# Federated CloudNet Architecture.

SIGCOMM VISA 2009

New Business Roles.



## Roles in CloudNet Arch.

### Service Provider (SP)

(offers services over the top)

### Virtual Network Operator (VNO)

(operates CloudNet, Layer 3+, triggers migration)

### Virtual Network Provider (VNP)

(resource broker, compiles resources)

### Physical Infrastructure Provider (PIP)

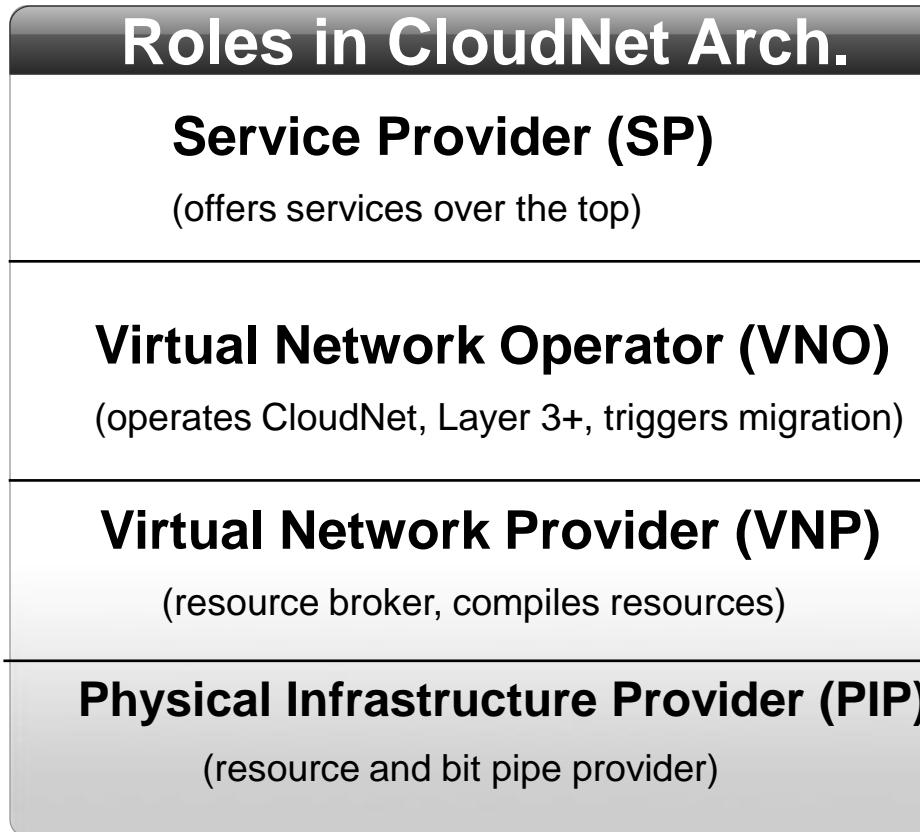
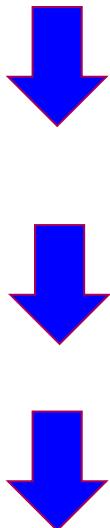
(resource and bit pipe provider)

A virtualized infrastructure opens new roles for the allocation of resources and the operation of networks!



New Business Roles.

**Contract based:**  
Communicate requirements  
/specs down the hierarchy



APIs: e.g., provisioning  
interfaces (migration)



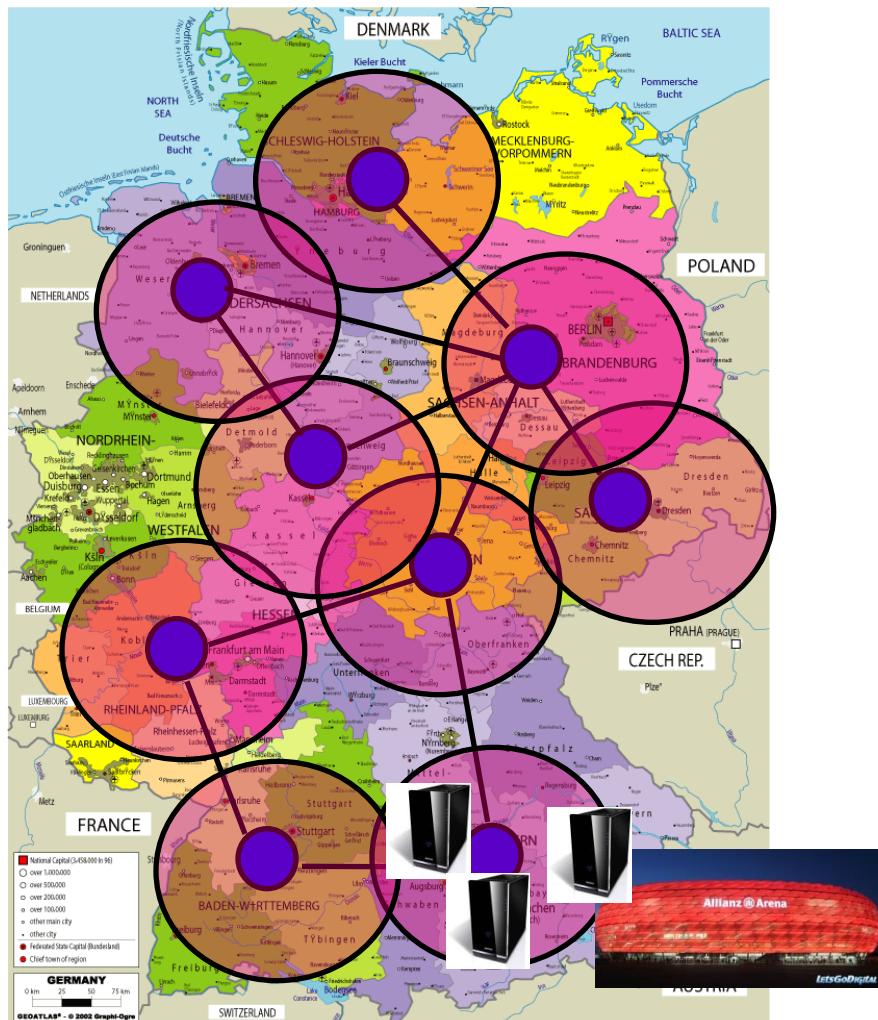
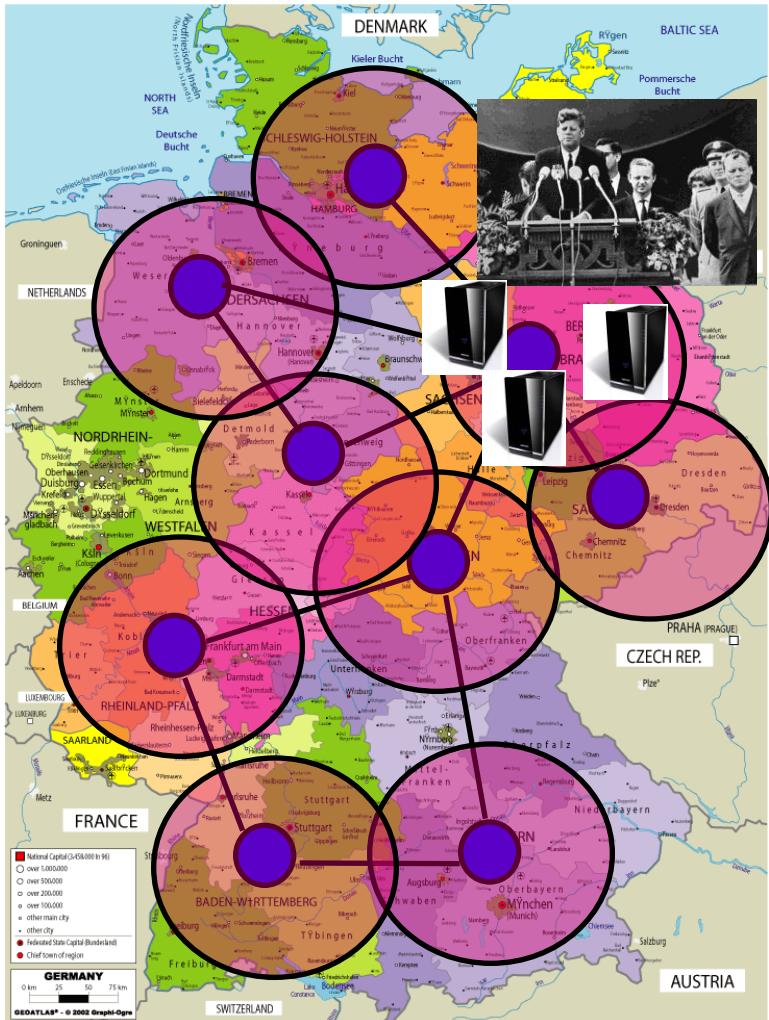
# Roles: Zoom In.

- **Service Provider (SP):**
  - Netflix, Akamai, startup, ..
  - Specifies resources abstractly (latency, push content to customer...)
  - If services migratable: offer provisioning interface (API that VNO can call)
- **Virtual Network Operator (VNO):**
  - Implements layer 3 on top of VNP layer 2: addressing, routing, ...
  - Decides how to realize SP specification: use redundant physical paths => additional specs
  - To fulfill specs, calls API of SP to migrate (if application migratable), or makes live migration itself (not to violate specs)
  - Clean slate architecture possible!
- **Virtual Network Provider (VNP):**
  - Builds layer 2
  - „Broker“/reseller that assembles virtual resources from different PIPs to provide **virtual topology** (no need to know PIP, can be **recursive**, ...)
- **Physical Infrastructure Provider (PIP):**
  - Bit-pipe provider, owns and manages physical **infrastructure**



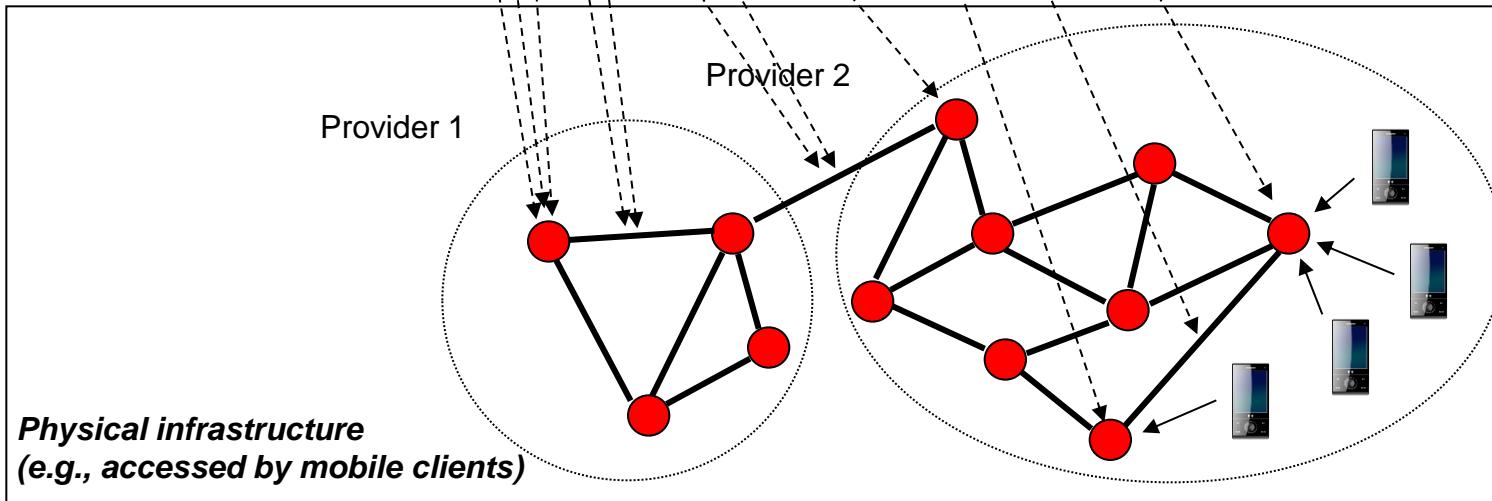
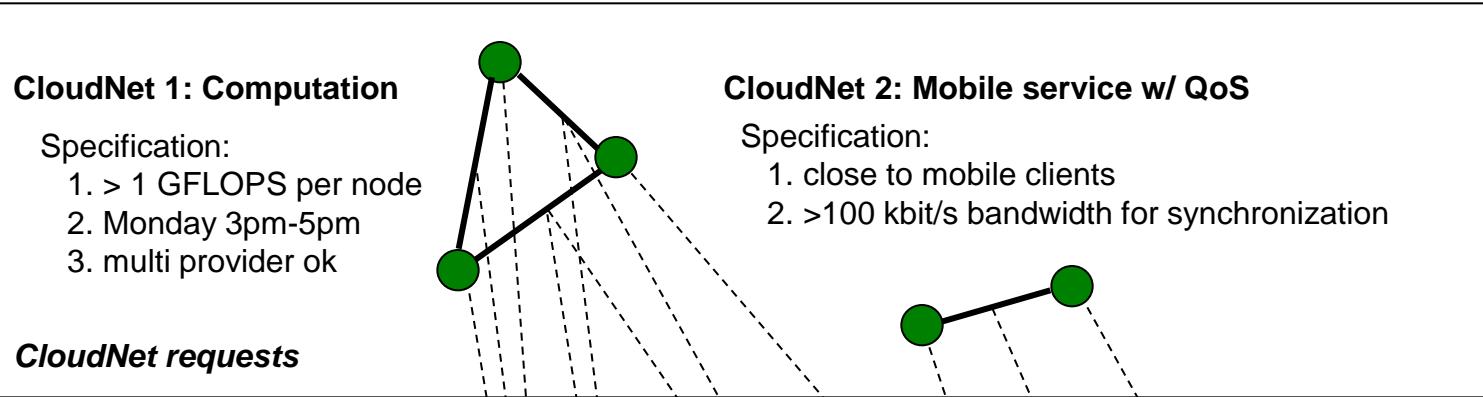
# Use Cases (1): Migrate Resources.

Resource allocation and migration where needed (energy saving otherwise!).



# Use Cases (2).

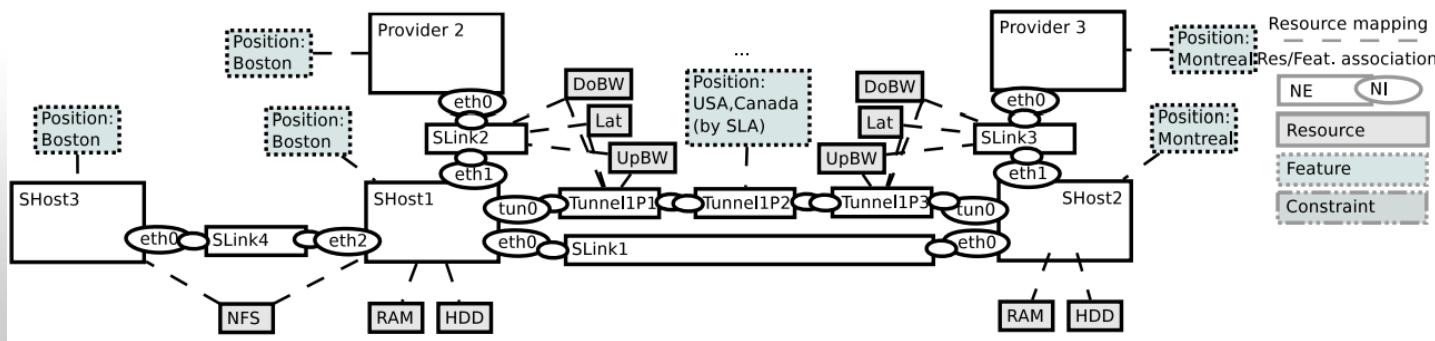
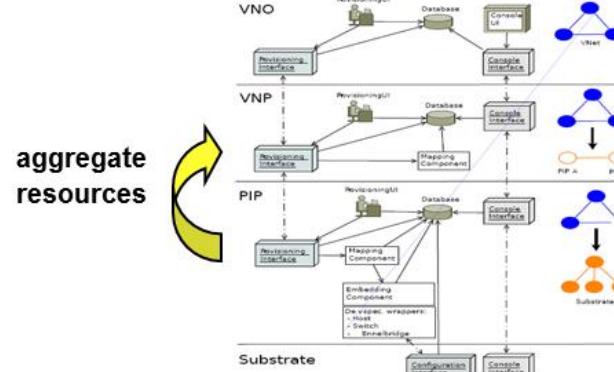
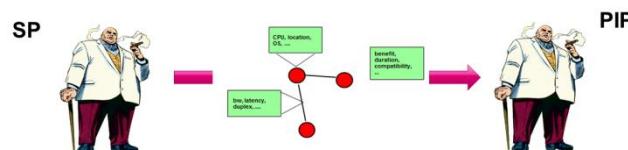
Connecting Providers (Geographic Footprint).



## Flexible Specification

### How to store and communicate CloudNets?

- Specify without losing flexibility
- Communicate non-topological requirements: consistency across multiple roles?
- Allow for aggregation and abstraction

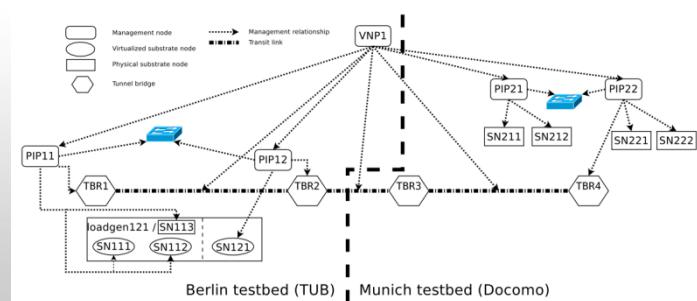
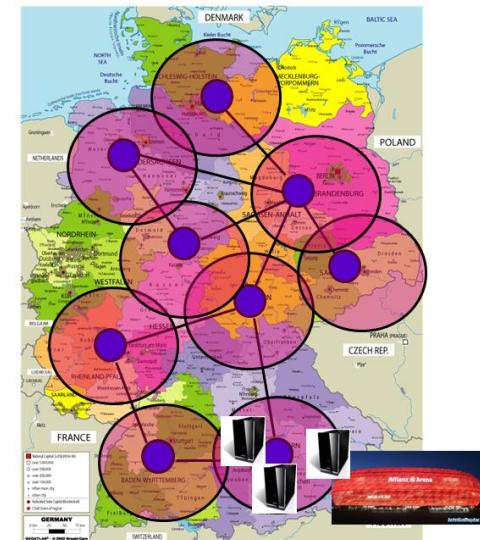
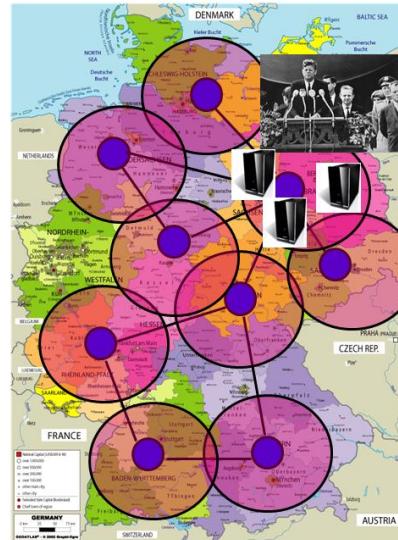
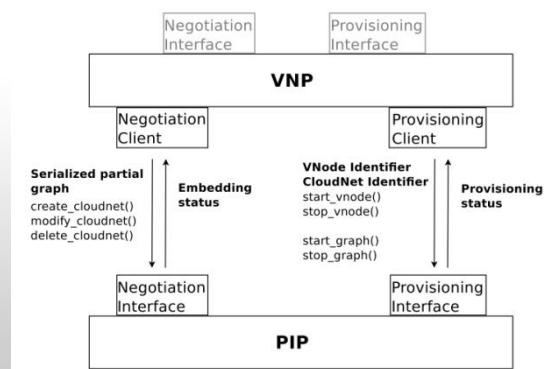


# Research Overview.

# Prototype

## Plugin architecture

- Own cloud operating system
  - Currently: VLAN based, but OpenFlow plugin started
  - Provisioning interfaces, negotiation interfaces
  - PIP and VNP role implemented
  - Seamless migration, e.g., of streaming service
  - Wide-area: OpenVPN tunnels
  - Wide-area testbed: Munich (NTT Docomo) and Berlin

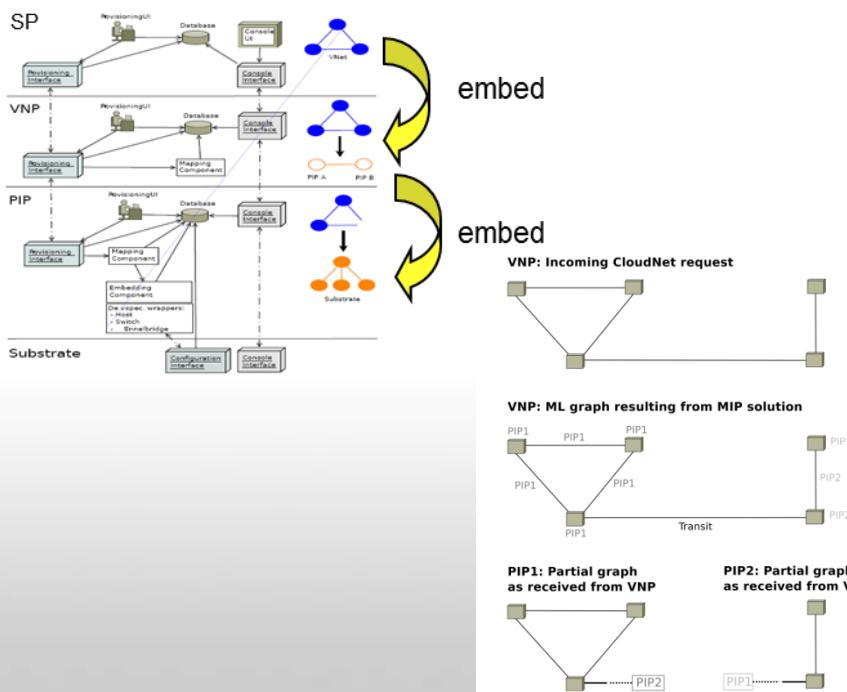


# Research Overview.

## Embedding

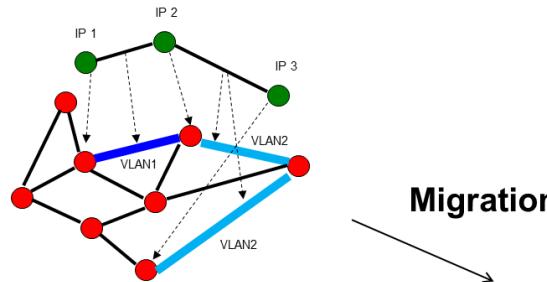
### Two steps:

- Quick heuristic (spec => greedy)
- Optimizing “long-lived” or “heavy hitter” CloudNets only (mixed integer program)



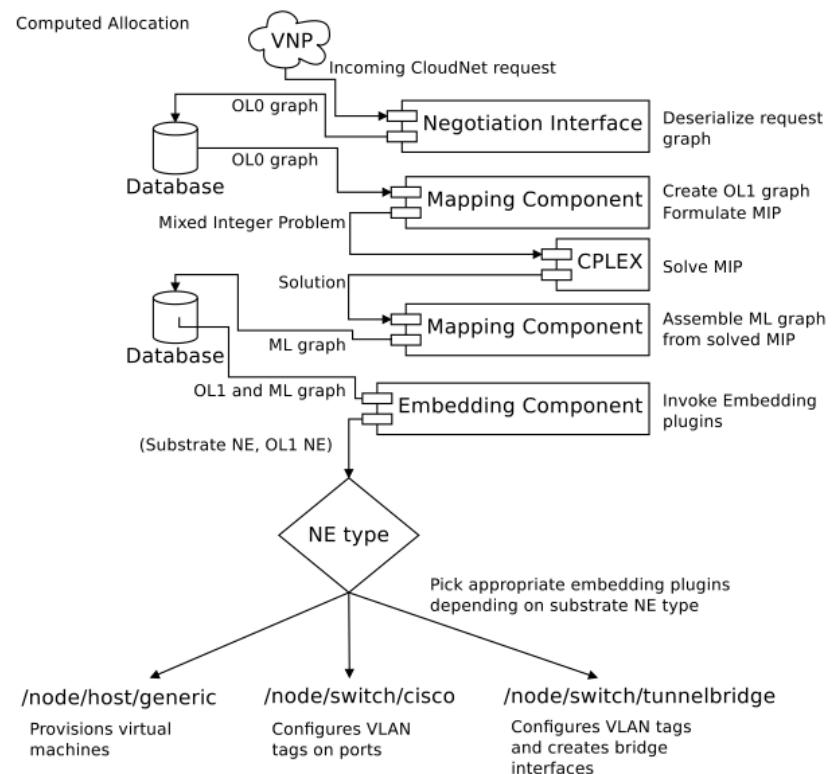
## Migration

### E.g., move VM by reconfiguring VLANs



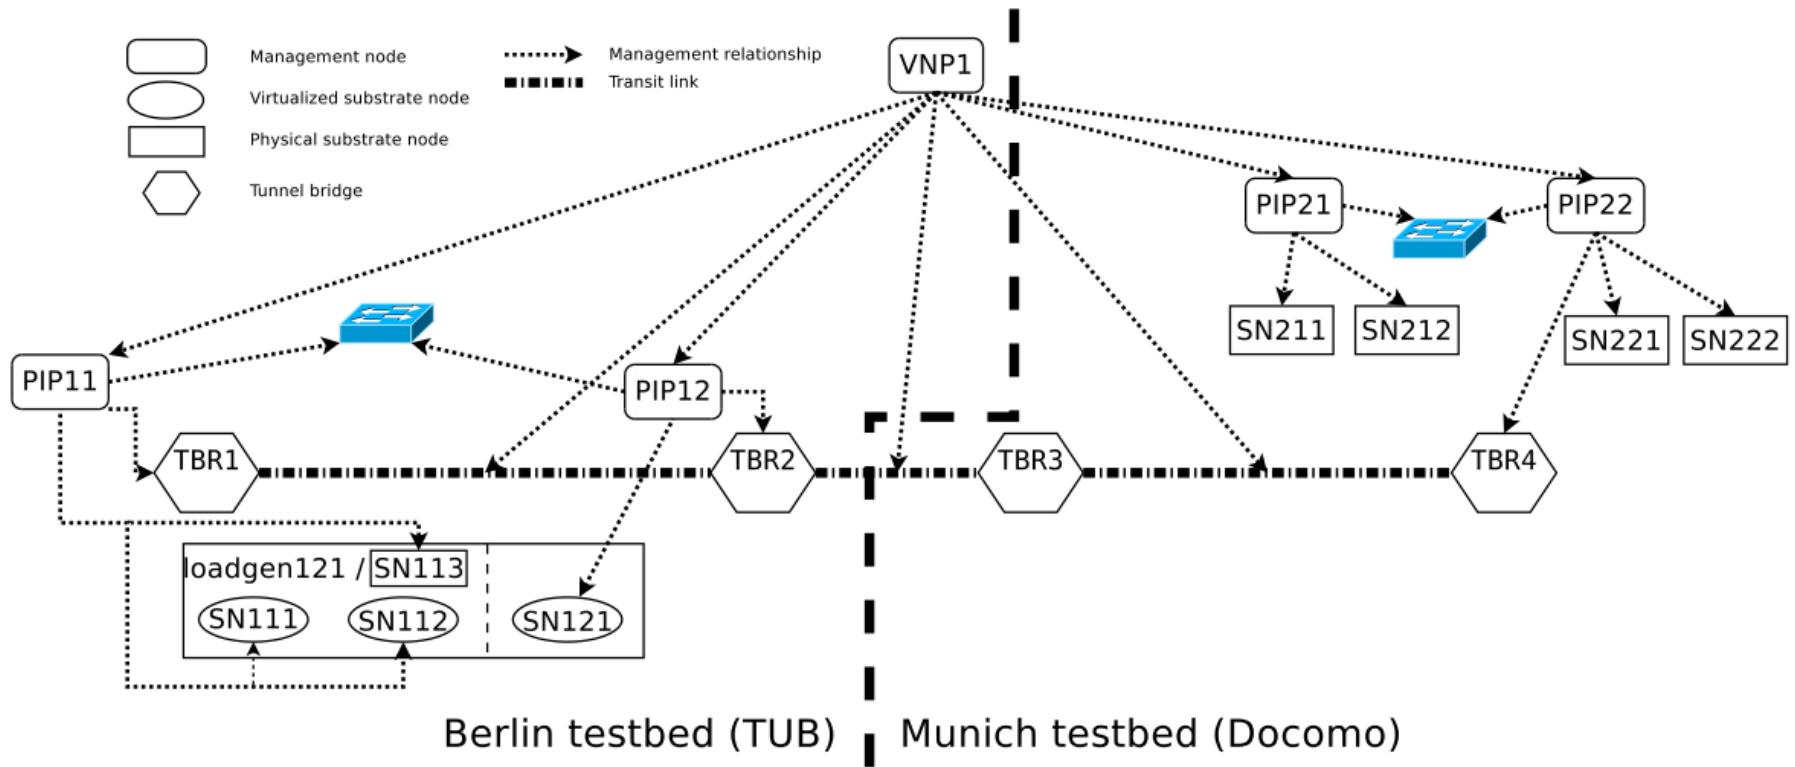
# Architecture / Anatomy of Prototype.

- **Plugin based**
  - VLink technology (VLAN, OpenFlow, MPLS, ...)
  - Embedding algorithm (two stage!)
  - Cloud operating system
- **Prototype: proof-of-concept (flexibility&generality rather than speed)**
  - VLink plugin: **VLANs** (VPN tunnel for WAN)
  - Embedding: MIP
  - Cloud OS: own implementation
- **Plan: other plugins, OpenStack, ...**



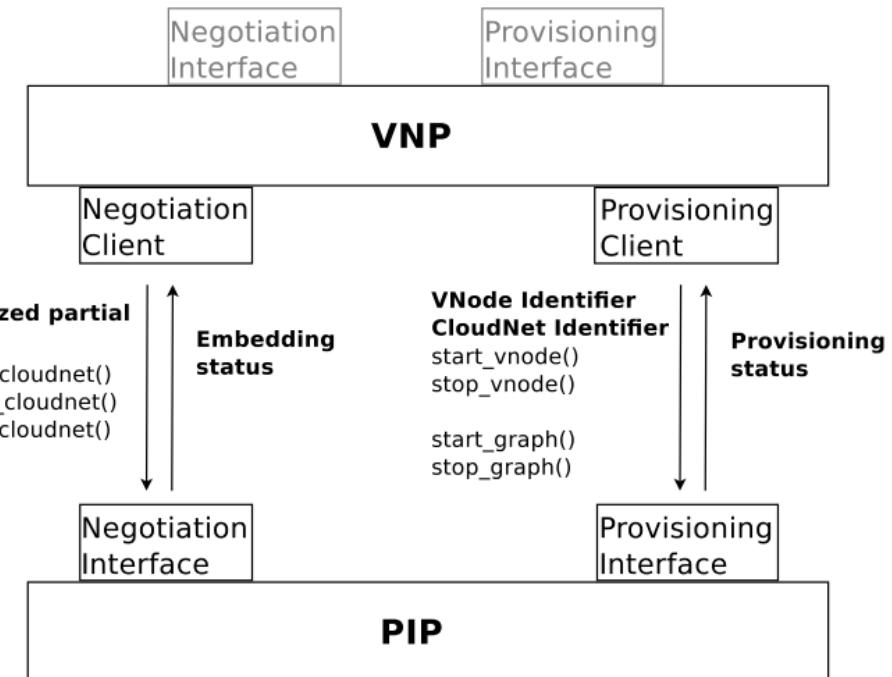
# Testbed.

Two sites: TU Berlin and NTT Docomo Eurolabs Munich

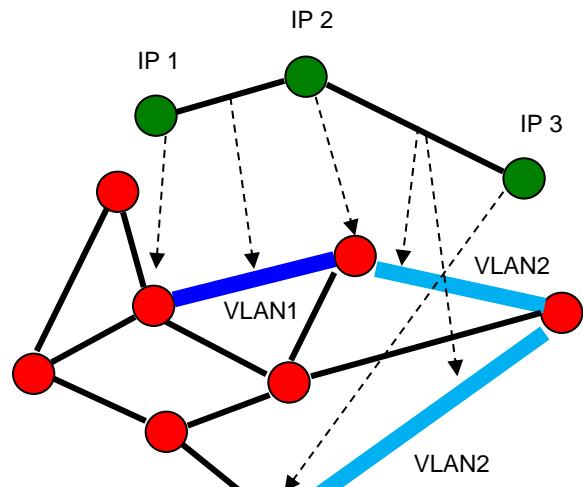


## State of the Art

- Prototype based on KVM+Xen (nodes) and VLANs (links)
- Layer 2 conform: no need for **end-to-end** / routing (routing inside CloudNet only)
- Different VNets may have same internal virtual node addresses
  - VLAN ensures isolation
- PIP and VNP implemented

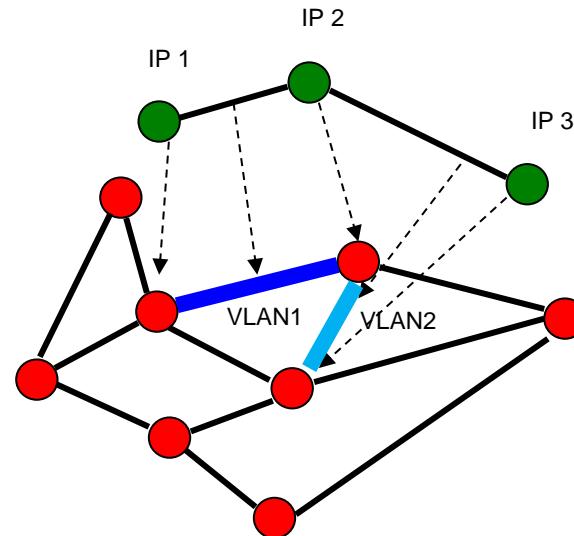


# The Prototype (2).



**Migration**

- Each virtual link is a VLAN (broadcast domain)
- Migration: reconfigure VLANs, not addresses of virtual nodes!
- Transparent for users...



- Open vSwitch supports VLAN bridging
- Demultiplexing eth1 plus VLAN-tag to port in kernel
- To VM looks like Ethernet (no VLAN)

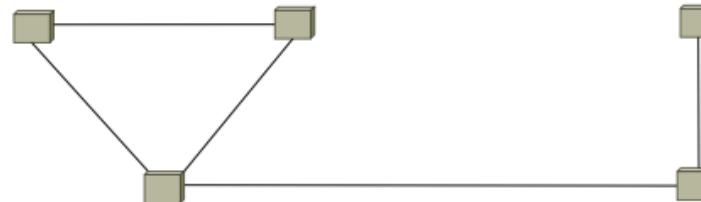


# The Prototype (3).

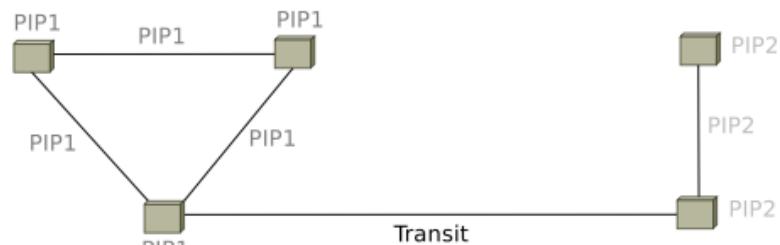
## Life of a CloudNet request:

- Topology broken down for PIPs
- Transit link (tunnel bridge and OpenVPN)
  - One VPN tunnel for control plane
  - One VPN tunnel for data plane
- To VM looks like Ethernet (no VLAN)

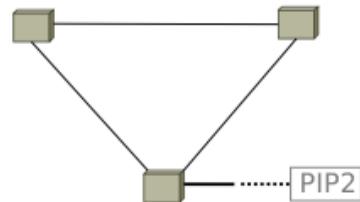
**VNP: Incoming CloudNet request**



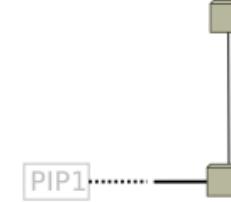
**VNP: ML graph resulting from MIP solution**



**PIP1: Partial graph as received from VNP**



**PIP2: Partial graph as received from VNP**

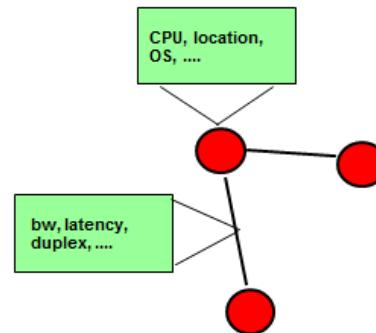


# Need for a Language.

ICCCN 2012

- Communicate CloudNets, substrate resources and embeddings to business partners or customers:

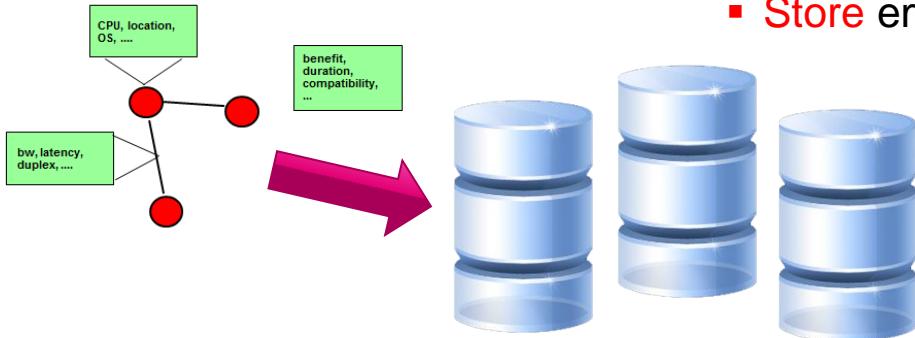
SP



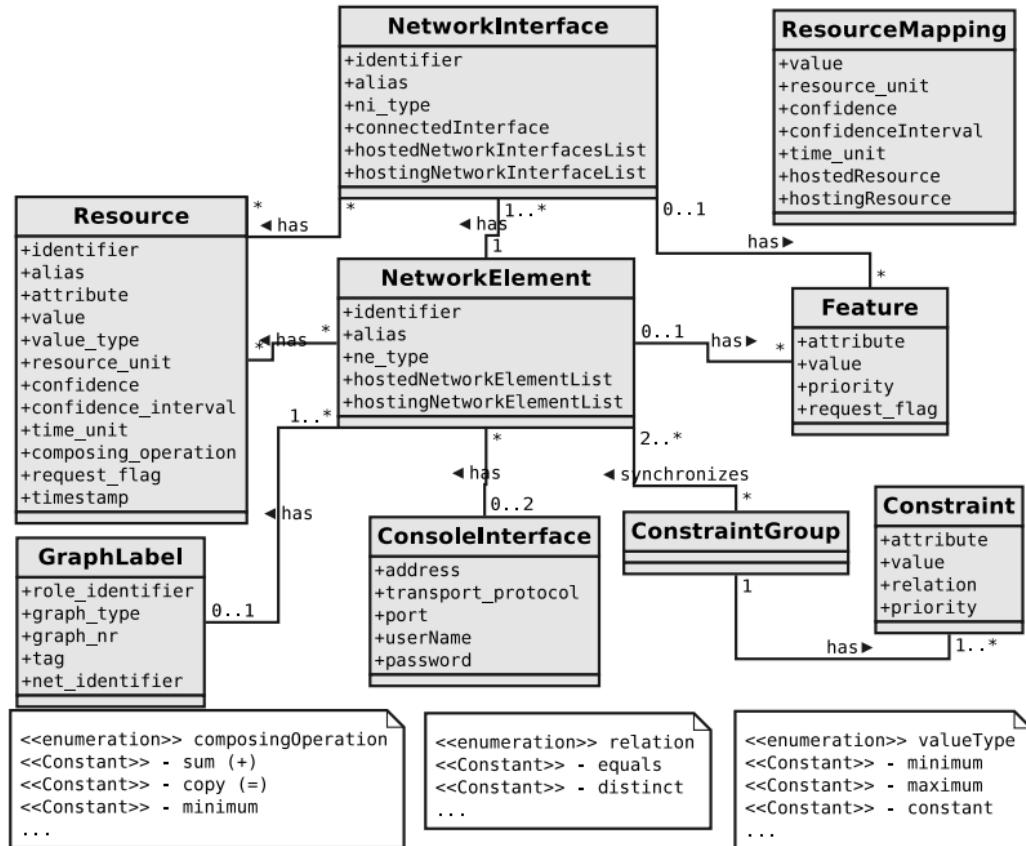
PIP



- Store embedding state internally:



# Exploiting Flexibilities: Resource Description Language.

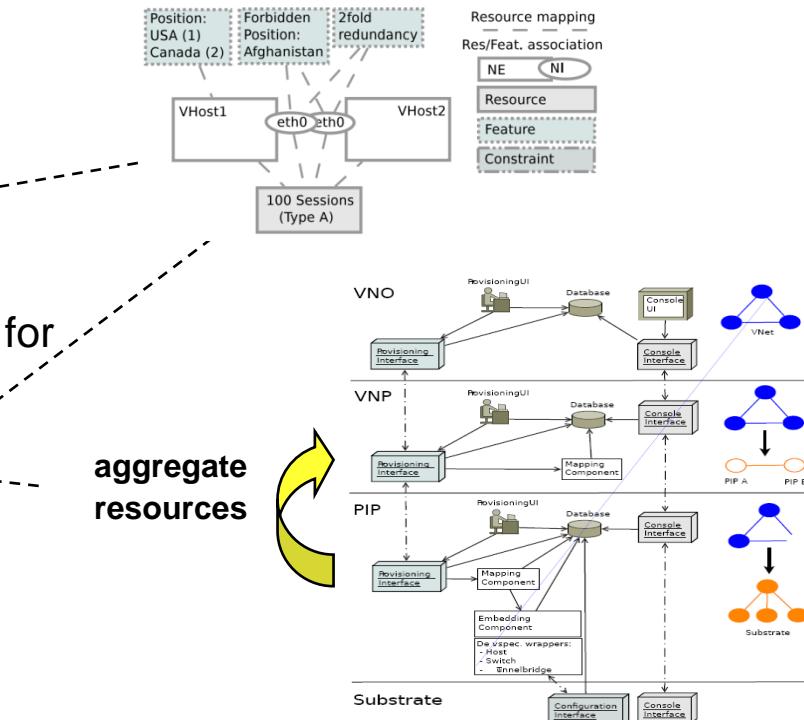
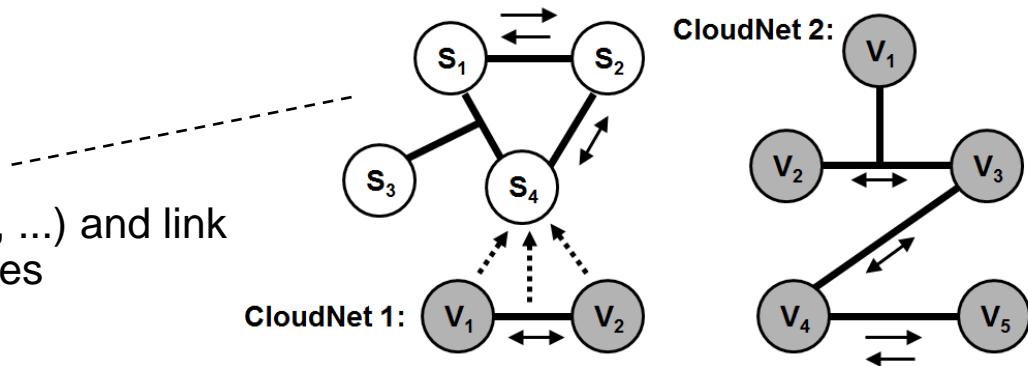


- Network Elements = nodes or links!
  - Connected via Network Interfaces
- Support for omission
- Support for multicast links
- Support for white and black lists...



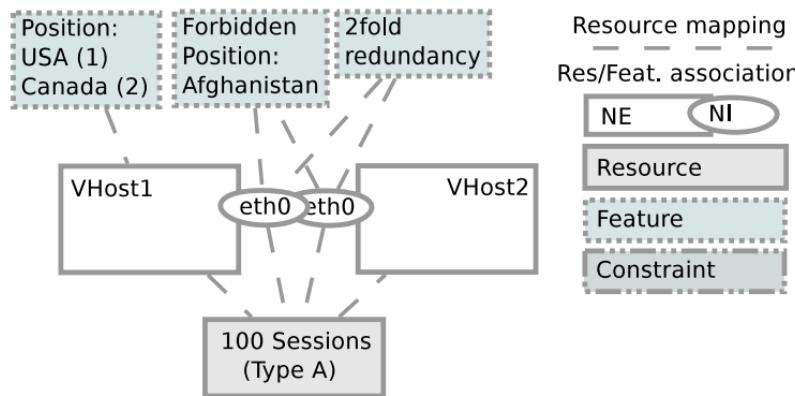
# FleRD Requirements.

- Support all kinds of node (storage, computation, ...) and link (latency, bw, full-duplex/asymmetric, ...) resources (**heterogeneity**)
- Extensible, allow for **syntactic changes** over time, no need for **global agreement** on semantic values
- Facilitate resource leasing and allow PIPs to open **abstract views** on their substrate
- Allow for **vagueness and omission**: customers are unlikely to specify each CloudNet detail (e.g., KVM or Xen is fine, outsource to any European cloud provider): this opens ways for optimization (**exploiting flexibilities**)!
- Allow for **aggregation** of resources (**business secret?**)
- Non-topological requirements (e.g., wordsize compatibility)



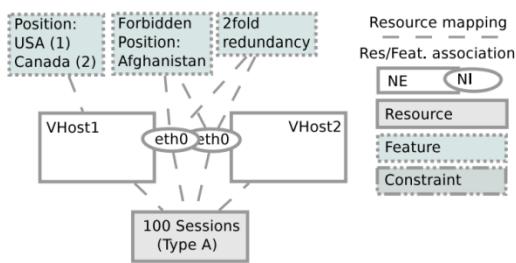
# Exploiting Flexibilities: Resource Description Language.

## Use Case: Web Service

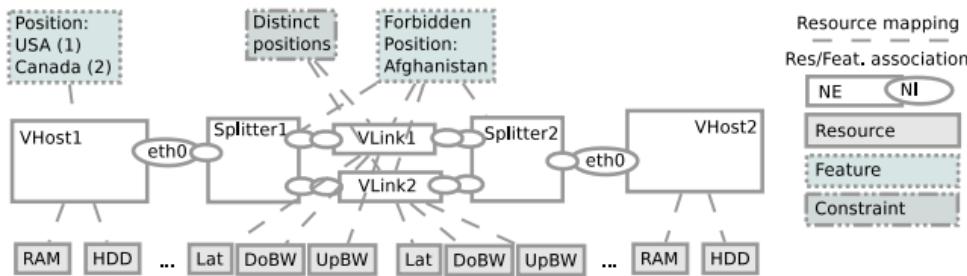


# Exploiting Flexibilities: Resource Description Language.

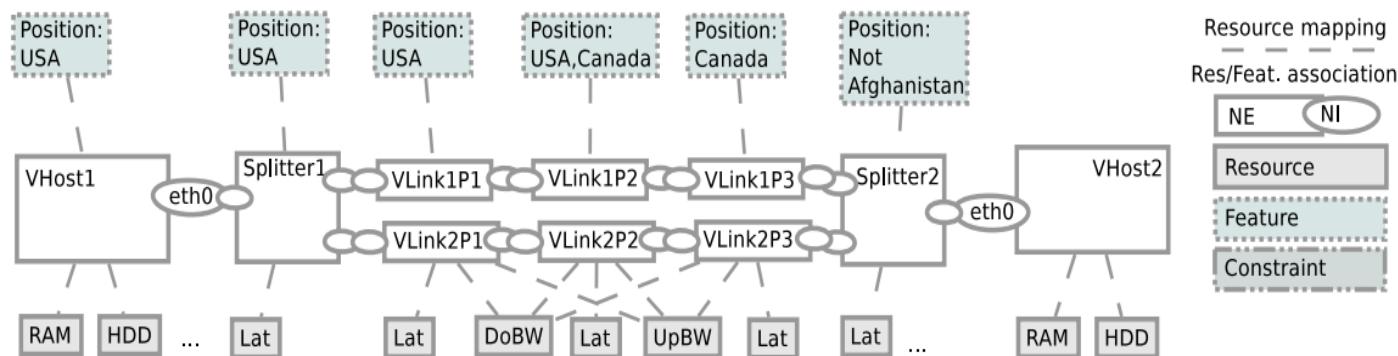
Web Service / Overlay 0:



Overlay 1: with splitters, two virtual links...

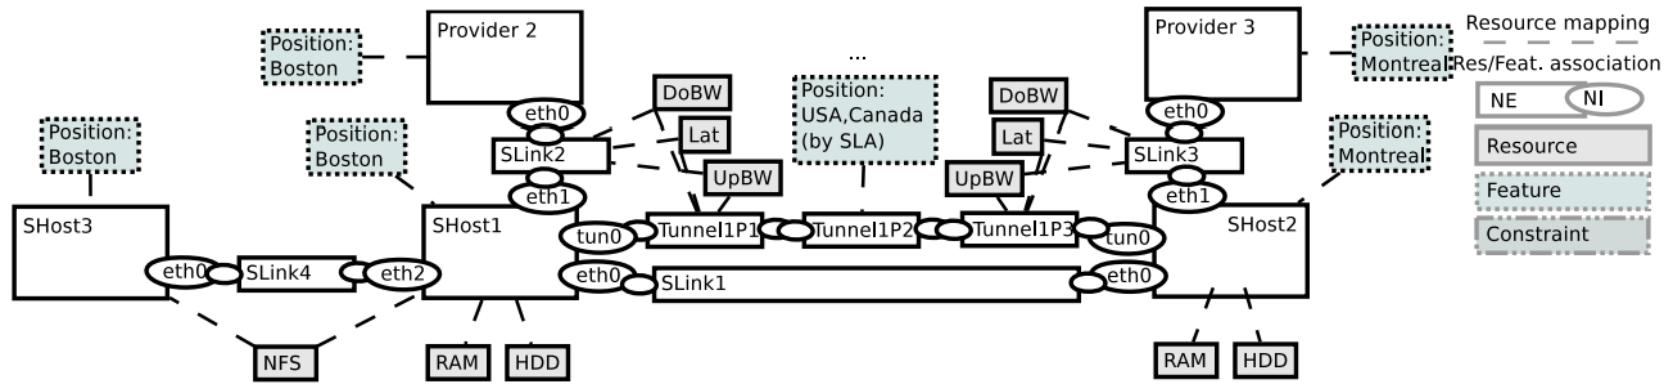


Mapping Layer: an virtual element per substrate element (n:m mapping)

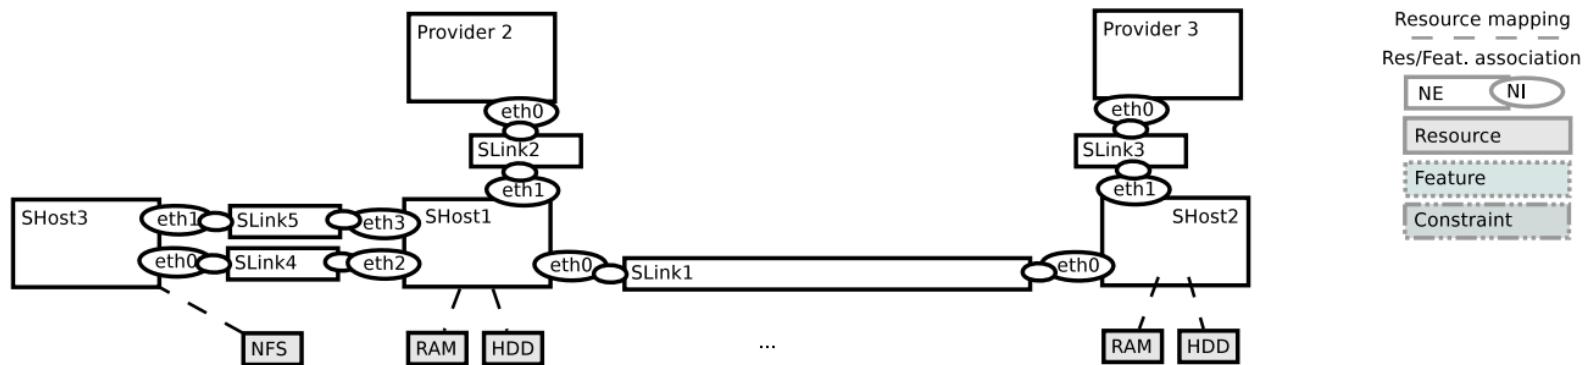


# Exploiting Flexibilities: Resource Description Language.

Underlay 1: all-provider view (splitter once collocated and once separate)



Underlay 0: provider 1 view (NFS at SHost3)

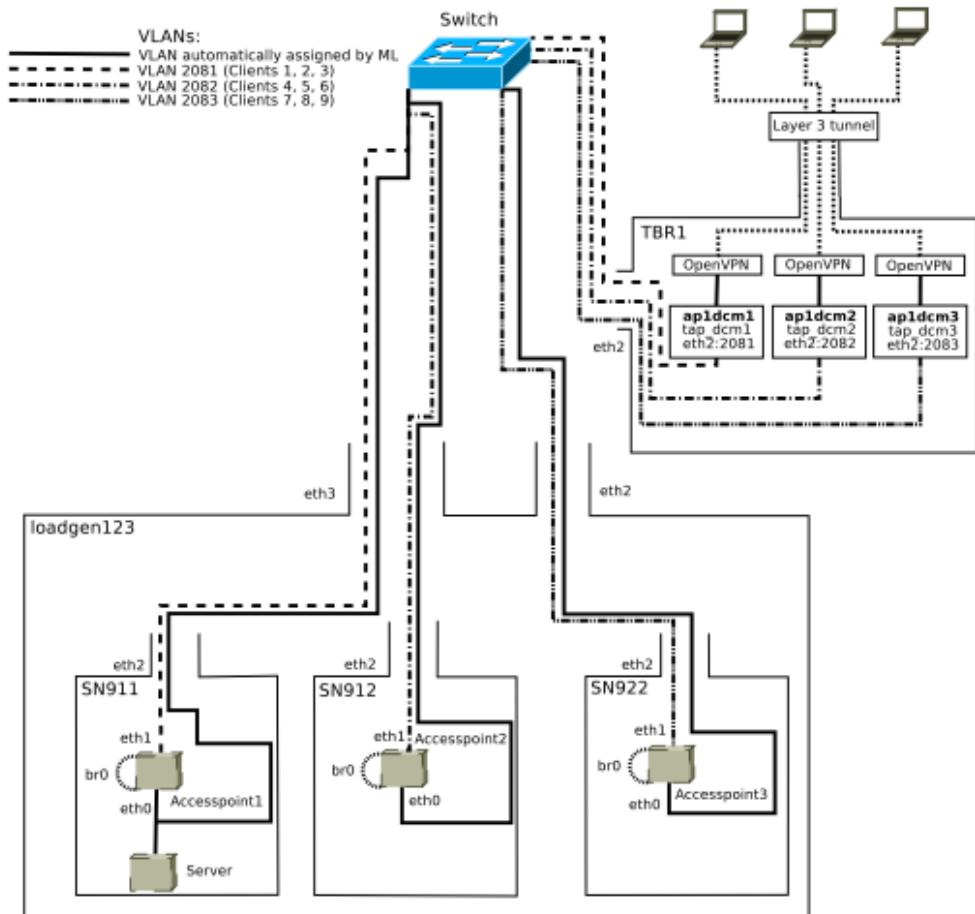


# Demo.

## YouTube Migration Demo

The video player displays a YouTube video titled "network virtualization" by schmiste78. The video content shows a diagram of a testbed virtual network (TESTBED VNET) with three switches (SN A, SN B, SN C) and three access points (AP1, AP2, AP3). It also shows a real-world setup with multiple monitors displaying network traffic. Below the video, a caption reads "clients join stream at three access points". The video player interface includes standard controls like play/pause, volume, and progress bar (0:13 / 3:04), along with video settings (360p) and sharing options (82 Aufrufe).

<http://www.youtube.com/watch?v=llJce0F1zHQ>



# (Theoretical) Research Overview.

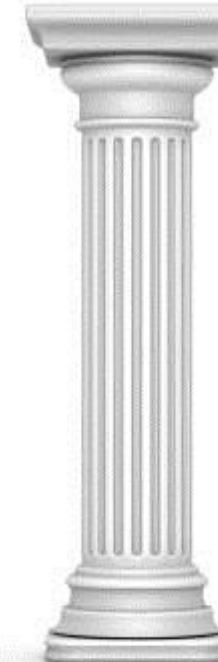
**Access Control  
and Embedding**



**Service Migration**

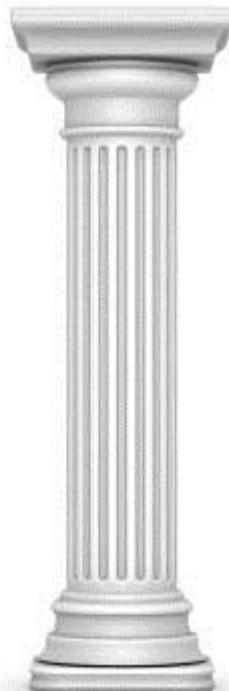


**Security Issues**

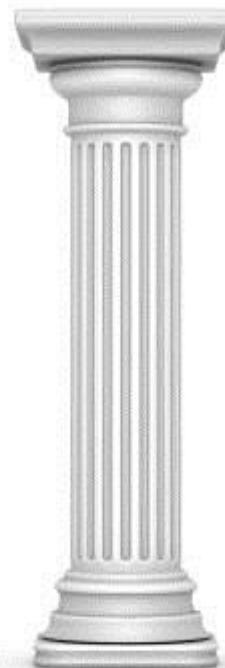


# Offline Embedding.

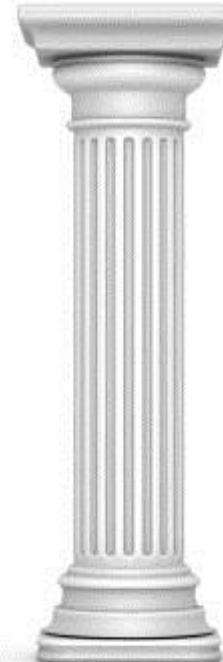
**Access Control  
and Embedding**



**Service Migration**



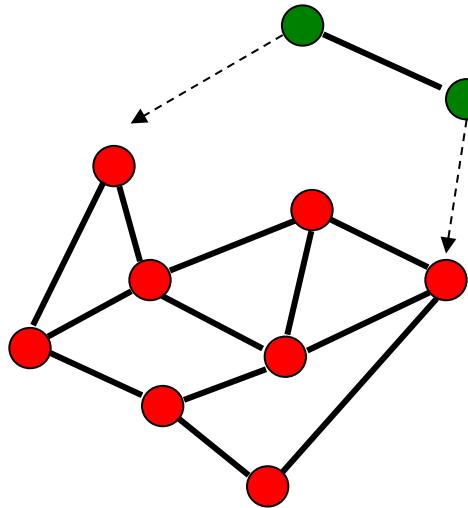
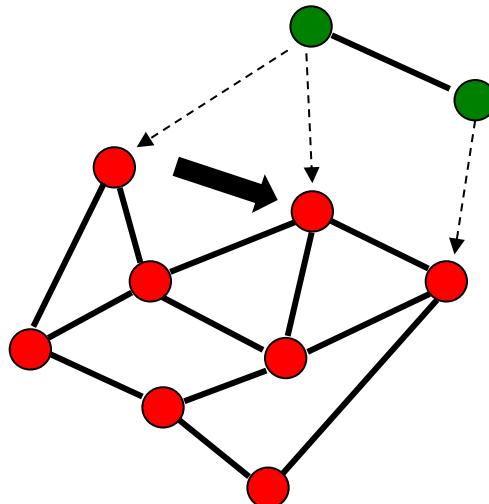
**Security Issues**



# How to Embed CloudNets Efficiently?

Computationally hard...  
Our 2-stage approach:

Stage 1: Map **quickly** and heuristically  
(dedicated resources)



Stage 2: **Migrate** long-lived CloudNets  
to «better» locations (min max load, max  
free resources, ...)  
Typically: **heavy-tailed** durations, so old  
CloudNets will stay longer!



## General Mathematical Program (MIP)

### Nodes:

map_node:	$\sum_{v \in NE_S} new(u, v) = 1$	$\forall u \in NE_VN$
set_new:	$alloc_{rs}(u, v, rv) \leq cap_{rs}(v) new(u, v)$	$\forall u \in NE_VN, v \in NE_S, rv \in R_V, rs \in R_S$
req_min:	$alloc_{rv}(u, v) \geq new(u, v) req(u, rv, s)$	$\forall u \in NE_VN, rv \in R_V, rs \in R_S, s = \text{minimum}$
req_max:	$alloc_{rv}(u, v) \leq new(u, v) req(u, rv, s)$	$\forall u \in NE_VN, rv \in R_V, rs \in R_S, s = \text{maximum}$
req_const:	$alloc_{rv}(u, v) = new(u, v) req(u, rv, s)$	$\forall u \in NE_VN, rv \in R_V, rs \in R_S, s = \text{constant}$

### Mapping:

relate_v:	$alloc_{rv}(u, v) \geq min\_alloc_{rv} \cdot new(u, v)$	$\forall u \in NE_V, v \in NE_S, rv \in R_V$
allowed:	$suit(u, v) \geq new(u, v)$	$\forall u \in NE_V, v \in NE_S$
ne_capacity:	$\sum_{u \in NE_V} \sum_{rv \in R_V} alloc_{rs}(u, v, rv) \leq cap_{rs}(v)$	$\forall v \in NE_S, rs \in R_S$
capacity:	$\sum_{v \in NE_S} \sum_{u \in NE_V} \sum_{rv \in R_V} alloc_{rs}(u, v, rv) \leq cap(rs)$	$\forall rs \in R_S$
load:	$weight_{rs}/cap(rs) \cdot \sum_{v \in NE_S} \sum_{u \in NE_V} \sum_{rv \in R_V} alloc_{rs}(u, v, rv) \leq load(rs)$	$\forall rs \in R_S$
max_load:	$load(rs) \leq max\_load$	$\forall rs \in R_S$

### Resource-Variable Relation:

resource:	$\sum_{rs \in R_S} prop(rv, rs) alloc_{rs}(u, v, rv) = alloc_{rv}(u, v)$	$\forall u \in NE_V, v \in NE_S, rv \in R_V$
flow_res:	$\sum_{rs \in R_S} prop(rv, rs) flow_{rs}(f, v, w, rv) = flow_{rv}(f, v, w)$	$\forall f \in Fl(u), (v, w) \in NE_S^2, rv \in R_f, \forall u \in NE_VL$

### Links:

map_link:	$\sum_{v \in NE_S} new(u, v) \geq 1$	$\forall u \in NE_VL$
map_flow:	$new(f, v) \leq new(u, v)$	$\forall f \in Fl(u), v \in NE_S, \forall u \in NE_VL$
map_src:	$new(f, v) \geq new(q_f, v)$	$\forall f \in Fl(u), v \in NE_S, q_f \text{ source of } f; \forall u \in NE_VL$
map_sink:	$new(f, v) \geq new(d_f, v)$	$\forall f \in Fl(u), v \in NE_S, d_f \text{ sink of } f; \forall u \in NE_VL$
req_min:	$\sum_{w \in NE_S} (flow_{rv}(f, v, w) - flow_{rv}(f, w, v)) \geq new(q_f, v) req(u, rv, s) - new(d_f, v) \infty$	$\forall f \in Fl(u), v \in NE_S, rv \in R_f; \forall u \in NE_VL, s = \text{minimum}$
req_max:	$\sum_{w \in NE_S} (flow_{rv}(f, v, w) - flow_{rv}(f, w, v)) \leq new(q_f, v) req(u, rv, s) + new(d_f, v) \infty$	$\forall f \in Fl(u), v \in NE_S, rv \in R_f; \forall u \in NE_VL, s = \text{maximum}$
req_const:	$\sum_{w \in NE_S} (flow_{rv}(f, v, w) - flow_{rv}(f, w, v)) = new(q_f, v) req(u, rv, s) - new(d_f, v) req(u, rv, s)$	$\forall f \in Fl(u), v \in NE_S, rv \in R_f; \forall u \in NE_VL, s = \text{constant}$

### Link Allocation:

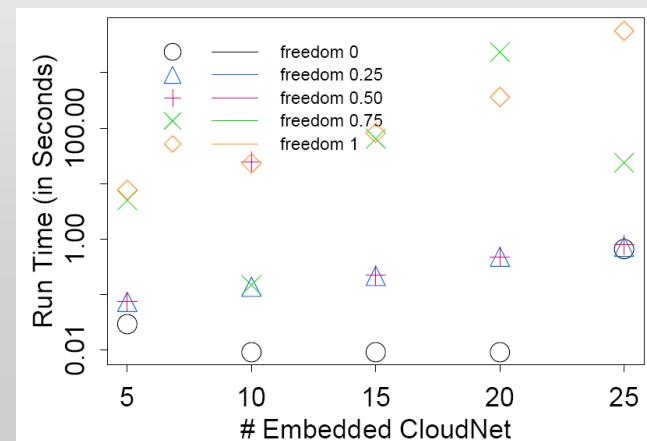
exp_out:	$\sum_{w \in NE_S} flow_{rs}(f, v, w, rv) \leq alloc_{rs}(u, v, rv)$	$\forall f \in Fl(u), v \in NE_S, rv \in R_f, rs \in R_S, \forall u \in NE_VL$
exp_in:	$\sum_{w \in NE_S} flow_{rs}(f, w, v, rv) \leq alloc_{rs}(u, v, rv)$	$\forall f \in Fl(u), v \in NE_S, rv \in R_f, rs \in R_S, \forall u \in NE_VL$
direction:	$flow_{rs}(f, v, w, rv) \leq new(u, v) cap_{rs}(v, w)$	$\forall f \in Fl(u), (v, w) \in NE_S^2, rv \in R_f, rs \in R_S, \forall u \in NE_VL$
relate_f:	$\sum_{w \in NE_S} flow_{rs}(f, v, w, rv) + flow_{rs}(f, w, v, rv) \geq new(f, v)$	$\forall f \in Fl(u), \forall u \in NE_VL, v \in NE_S, rv \in R_f, rs \in R_S$

### Migration:

new:	$\sum_{v \in NE_S} old(u, v) \geq mig(u)$	$\forall u \in NE_V$
migrated:	$old(u, v) - new(u, v) \leq mig(u)$	$\forall u \in NE_V, v \in NE_S$

### Advantages:

1. Generic (backbone vs datacenter) and allows for migration
2. Allows for different objective functions
3. Optimal embedding: for background optimization of heavy-tailed (i.e., long-lived) or «heavy hitter» CloudNets, quick placement e.g., by clustering  
*But: slow...*



## General Mathematical Program (MIP)

Nodes:  
map\_node  
set\_new:  
req\_min:  
req\_max:  
req\_con:

Mapping:  
relate\_V

allowed:  
ne\_capac

capacity:  
load:  
max\_load

Resource:  
resource:  
flow\_res

Links:  
map\_link

map\_flow

map\_src:  
map\_sink

req\_min:  
req\_max:  
req\_cons

Link Alloc

exp\_out:  
exp\_in:  
direction:  
relate\_f

Migration

new:  
migrated

### Advantages:

## Advantages of MIP:

- Very general
- Supports easy replacement of objective functions
- Can use standard, optimized software tools such as CPLEX, Gorubi, etc.

(inter)

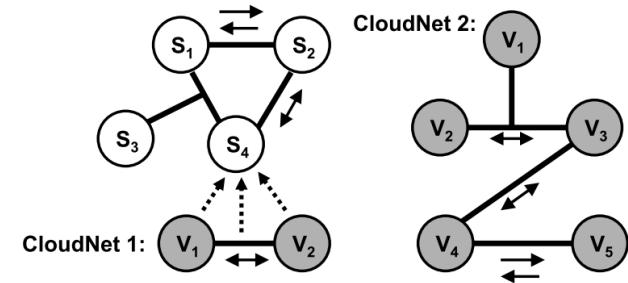
ound  
long-  
nt e.g.,



25



# Generality of the MIP.



## Objective functions:

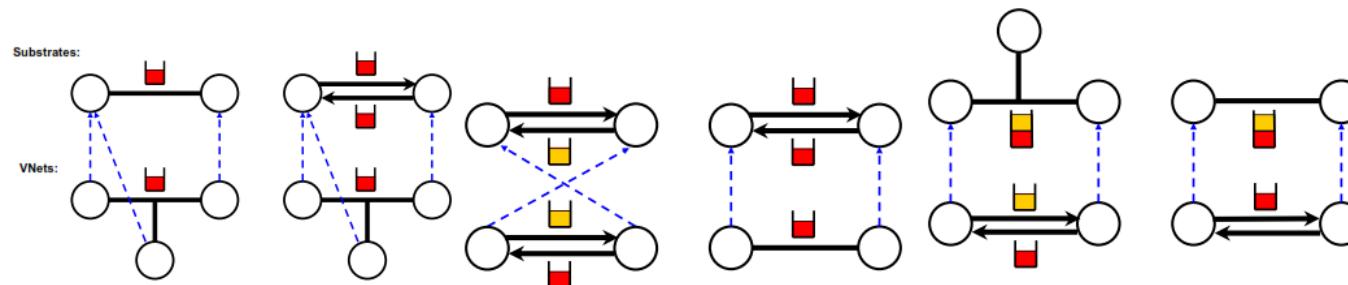
- minimize maximum **load** (= load balance)
- maximize **free resources** (= compress as much as possible), ...

## Migration support:

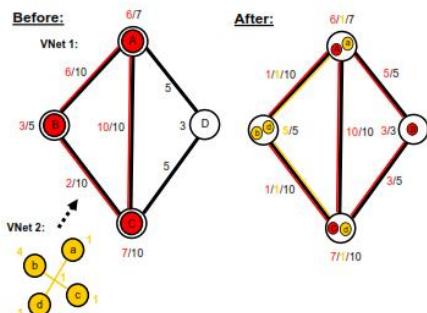
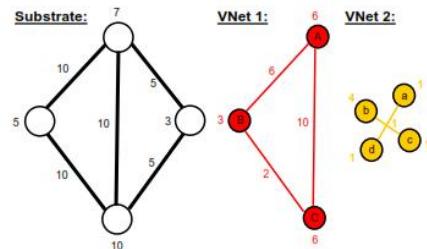
- costs for **migration**: per element, may depend on destination, etc.
- answer questions such as «**what is cost/benefit if I migrate now?**»

## Embedding:

- embedding full-duplex on full-duplex links
- full-duplex on half-duplex links
- or even **multiple endpoint links** (e.g., wireless) supported!



# On the Use of Migration.



Res.	w/o	w/ Link	w/ Link&Node	Opt
1	3	3	4	<b>4</b>
2	5	5	9	<b>9</b>
3	7	8	13	<b>13</b>
4	1	1	17	<b>17</b>
5	17	22	24	<b>24</b>
6	2	2	27	<b>27</b>
7	31	32	32	<b>32</b>
8	37	37	37	<b>37</b>

**Migration: Useful to increase the number of embeddable CloudNets, especially in under-provisioned scenarios**

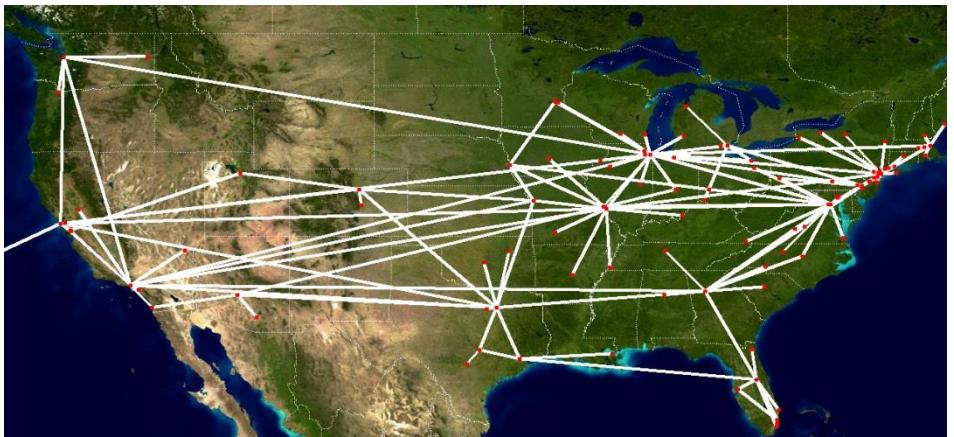
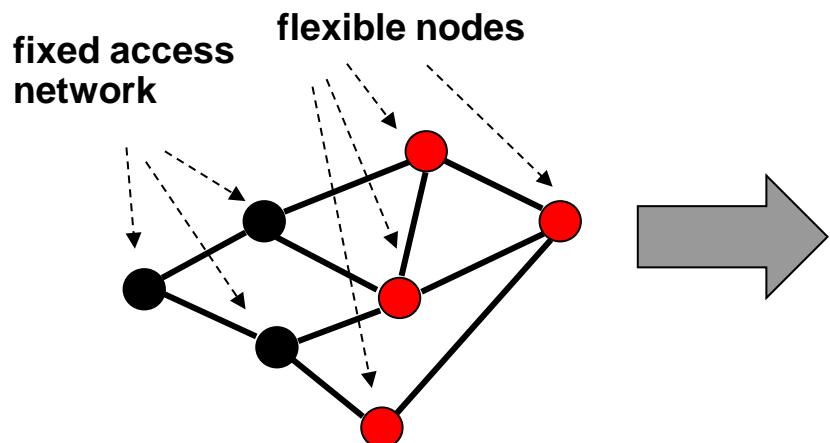


# Performance of the MIP: Setup.

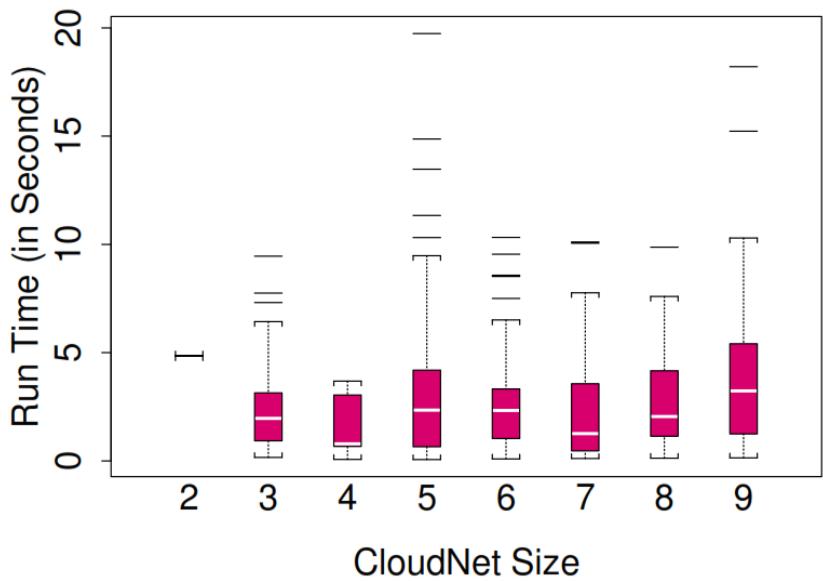
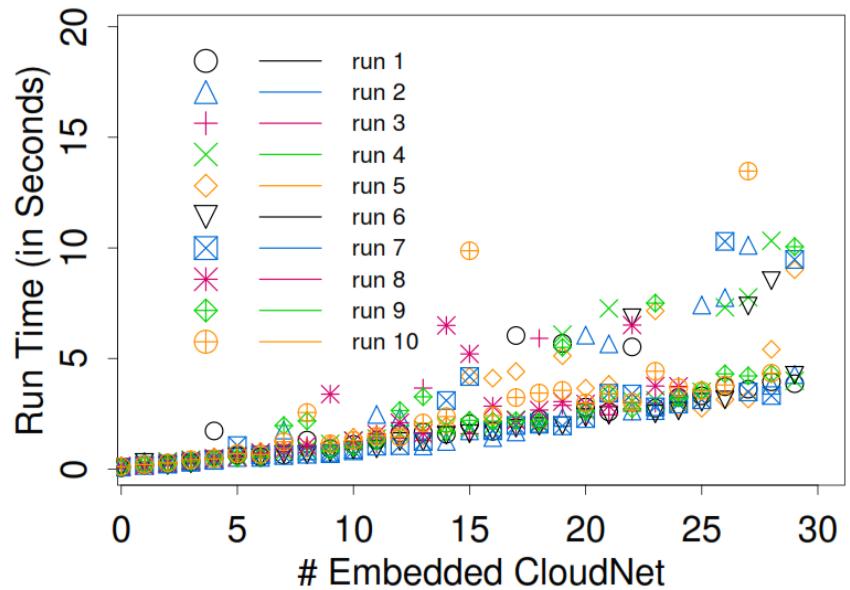
**Substrate:** Rocketfuel ISP topologies (with 25 nodes)

**CloudNets:** Out-sourcing scenario, CloudNets with up to ten nodes, subset of nodes fixed (access points) and subset flexible (cloud resources)

**Solver:** CPLEX on 8-core Xeon (2.5GHz)



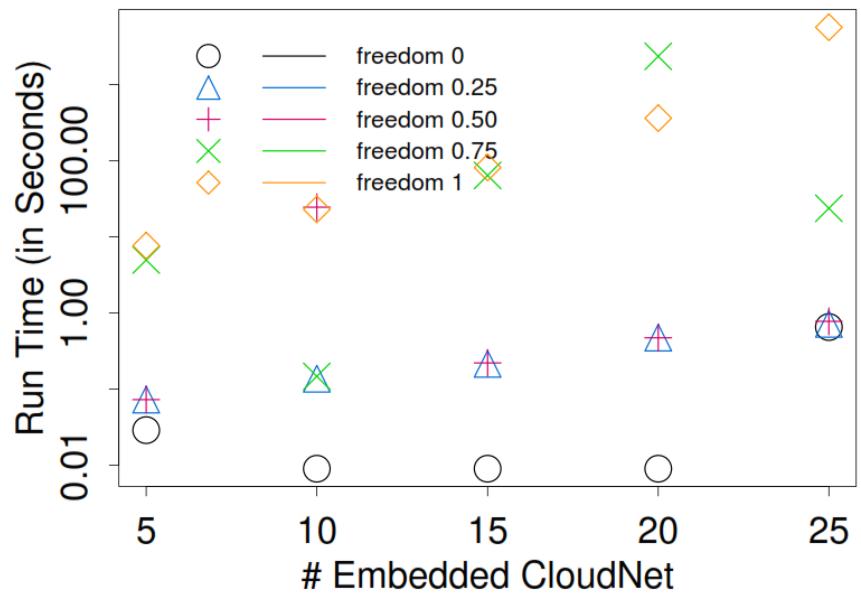
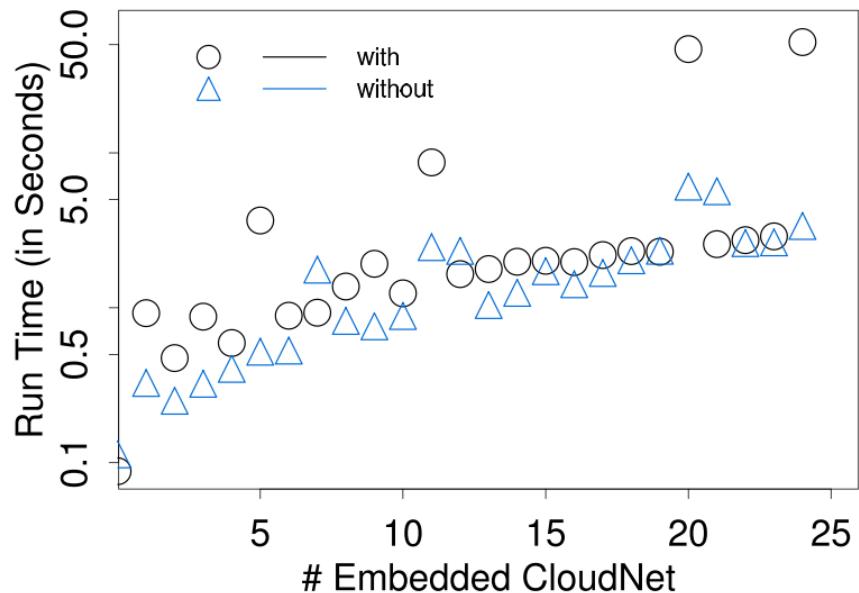
# Performance of the MIP.



- Runtime below 1 minute per CloudNet, slightly increasing under load
- Impact of CloudNet size relatively small



# Performance of the MIP.

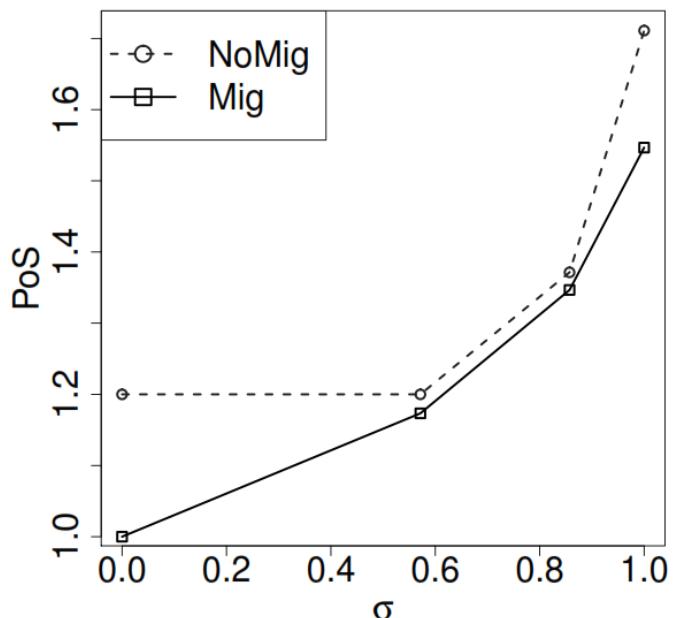
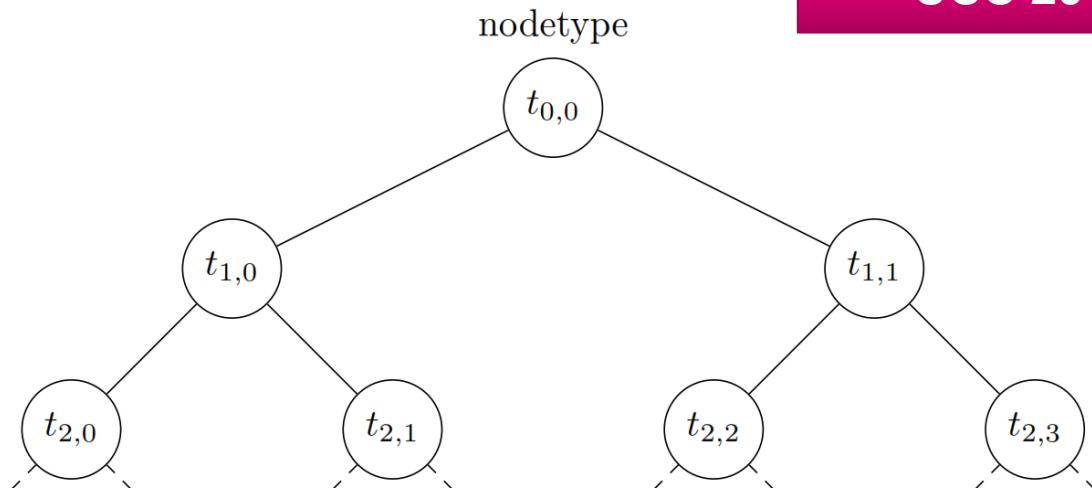


- Enabling option to migrate can increase execution time significantly (log scale!)
- Also number of flexible CloudNet components is important

# Use of Flexibility.

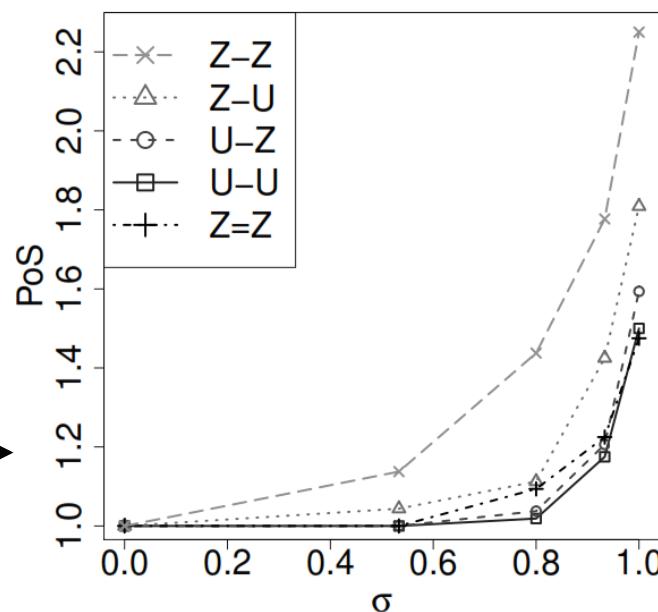
## PoS

How much link resources are needed to embed a CloudNet with specificity s%?



Up to 60%, even a little bit more if no migrations are possible!

Skewed (Zipf) distributions worst when not matching.



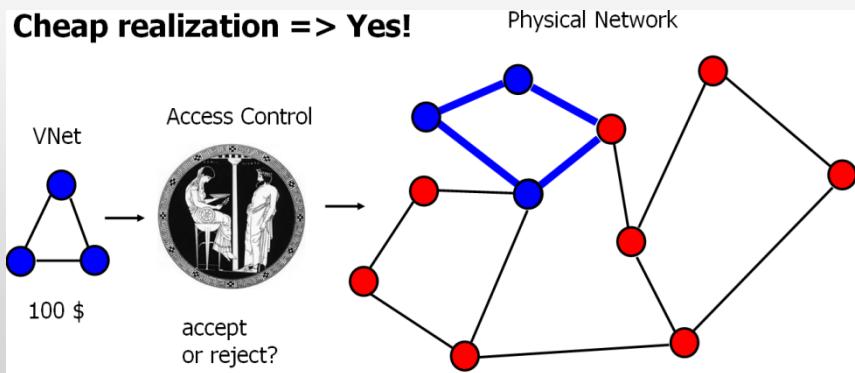
# CloudNet Embedding.

## Online Access Control

### Goal:

Decide online which VNet requests to accept, such that **profit** is maximized

### Cheap realization => Yes!

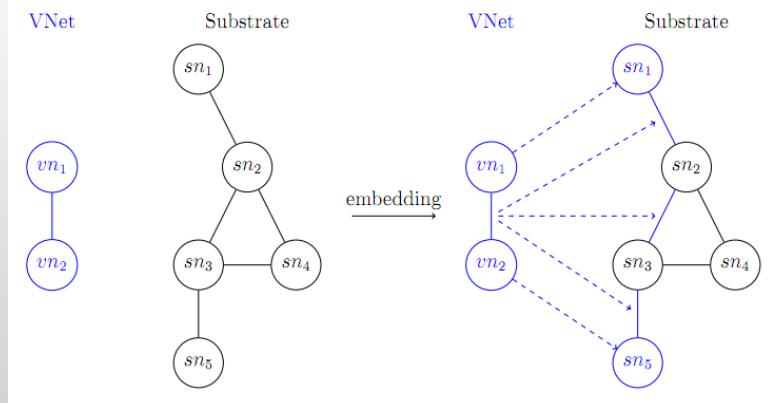


## Mapping and Allocation

### Goal:

Where to realize CloudNet such that spec is met? Objective, e.g.: minimize allocation **resources**, minimize **max load**, save energy,

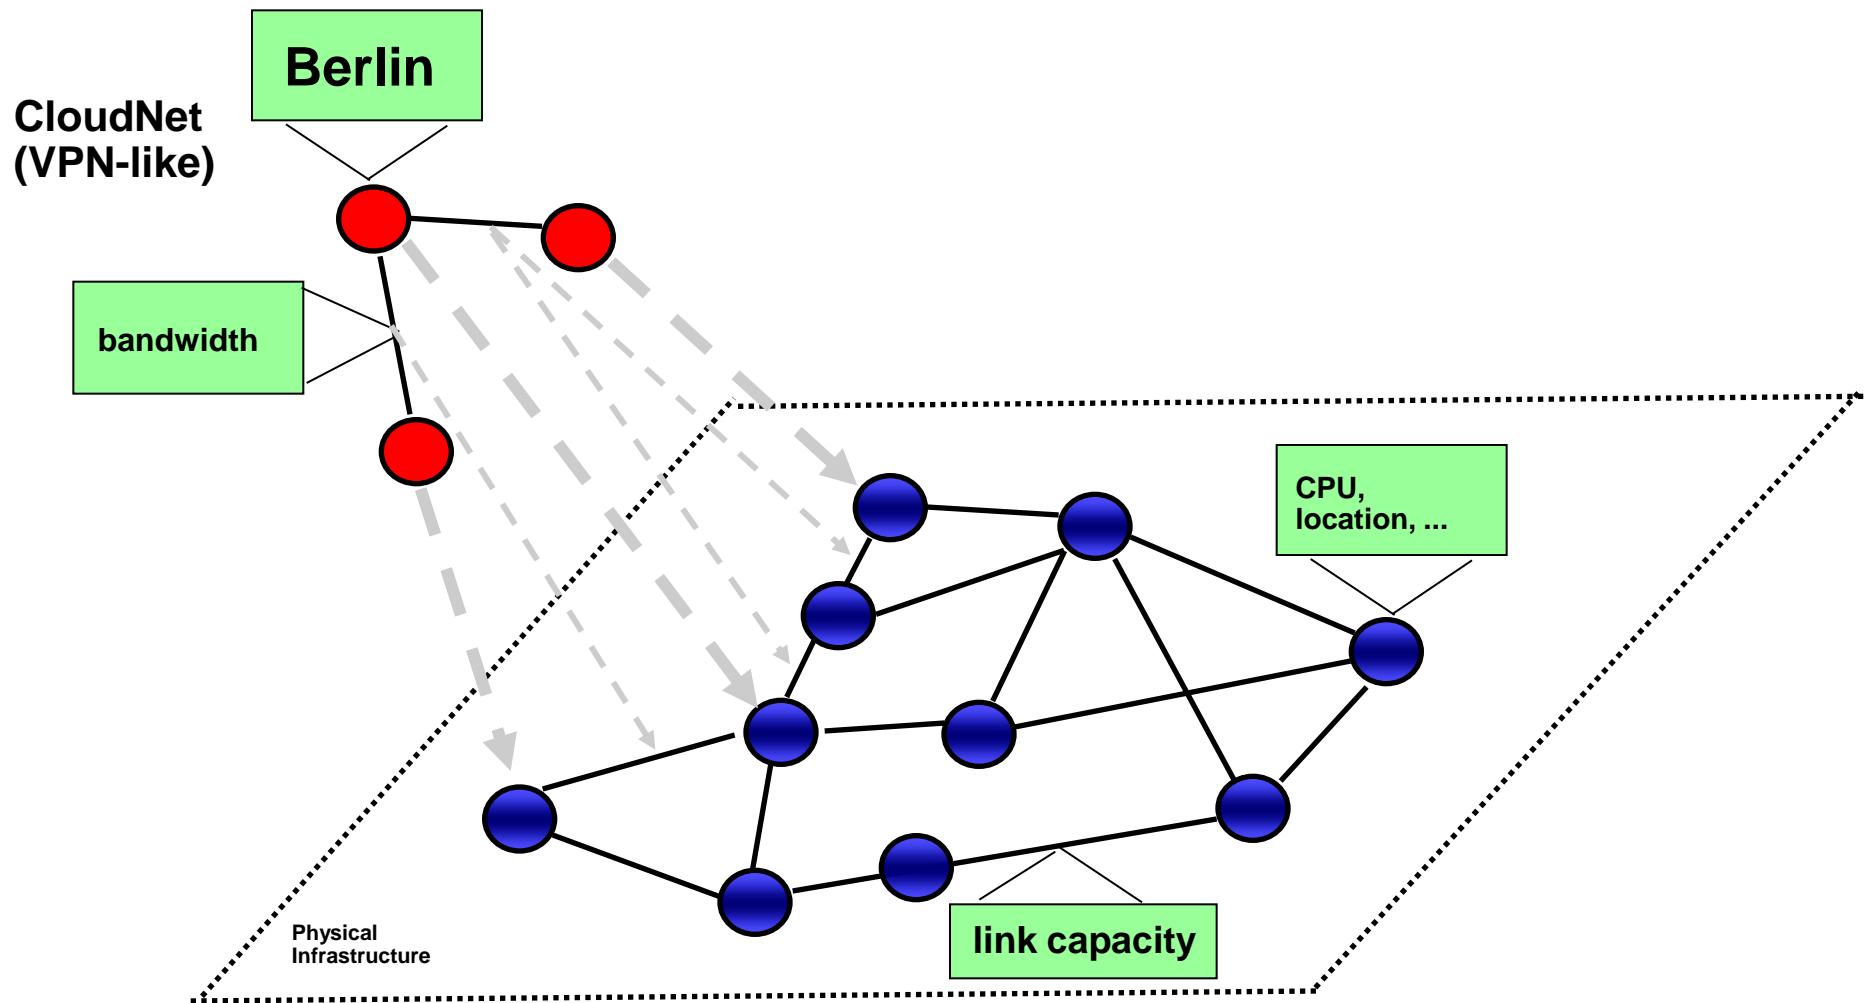
...



Currently focus on optimizing existing CloudNets (**heavy-tailed lifetime** assumption): but we are also working on quick embeddings (**clustering**, iterative, ...)



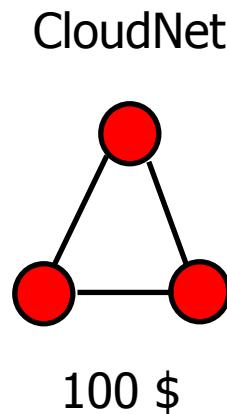
# Competitive Access Control: Model (1)



# Competitive Access Control: Model (2)

Specification of CloudNet request:

- terminal **locations** to be connected
- **benefit** if CloudNet accepted (all-or-nothing, no preemption)
- desired **bandwidth** and allowed **traffic** patterns
- a **routing** model
- **duration** (from when until when?)

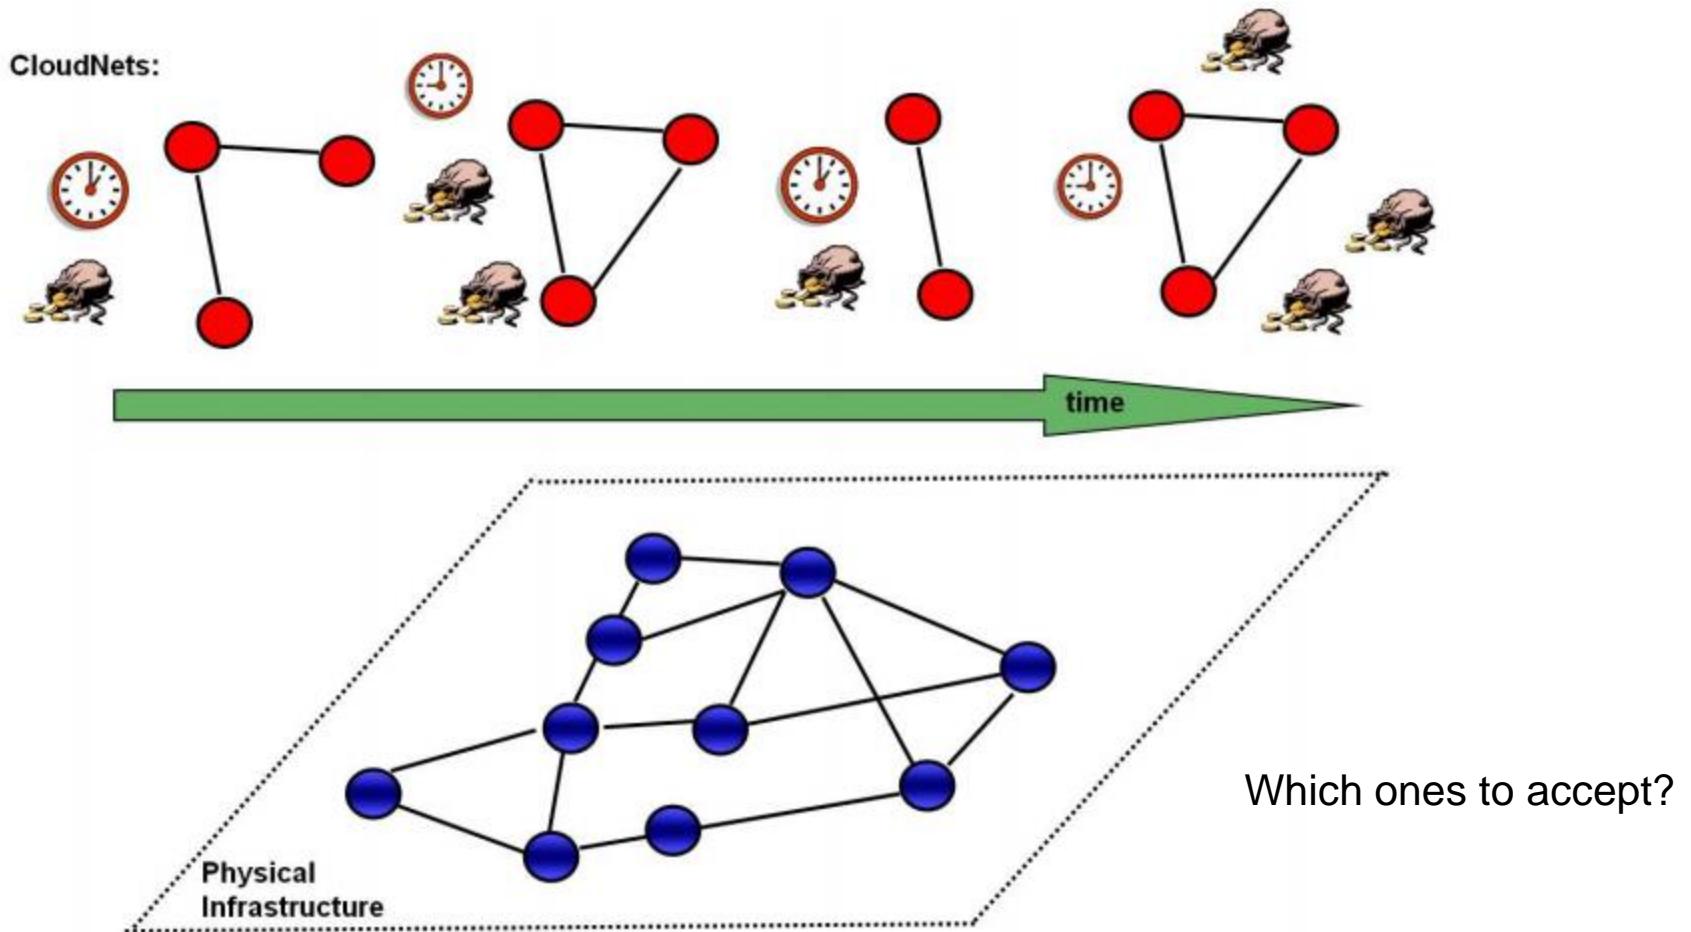


If CloudNets with these specifications arrive over time,  
which ones to accept online?

Objective: maximize **sum of benefits** of accepted CloudNets



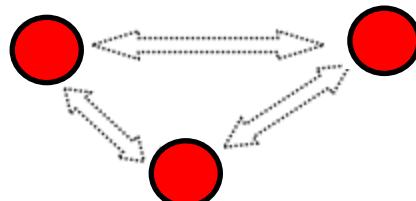
# Competitive Access Control: Model (3)



# CloudNet Specifications (1): Traffic Model.

## Customer Pipe

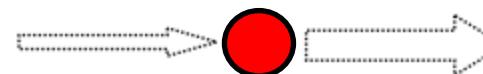
Every pair  $(u,v)$  of nodes requires a certain bandwidth.



Detailed constraints, only this **traffic matrix** needs to be fulfilled!

## Hose Model

Each node  $v$  has **max ingress** and **max egress bandwidth**: each traffic matrix fulfilling them must be served.

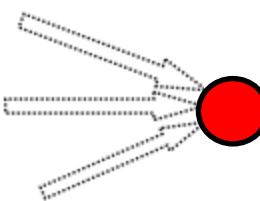


More flexible, must support many traffic matrices!

«virtual switch»

## Aggregate Ingress Model

Sum of ingress bandwidths must be at most a parameter  $I$ .



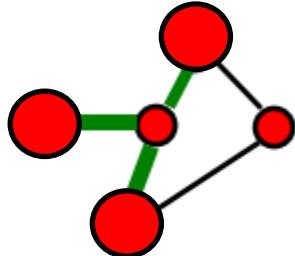
Simple and flexible! Good for **multicasts** etc.: no overhead, duplicate packets for output links, not input links already!



# CloudNet Specifications (2): Routing Model.

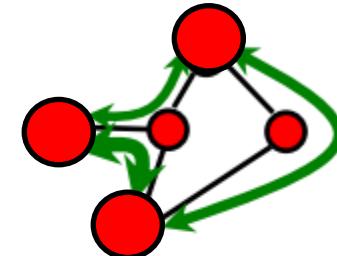
## Tree

VNet is embedded as **Steiner tree**:



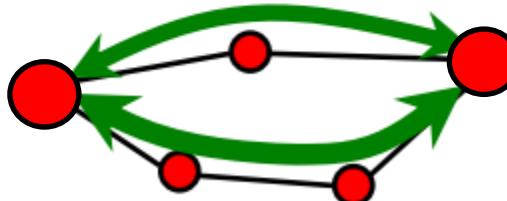
## Single Path

Each pair of nodes communicates along a single path.



## Multi Path

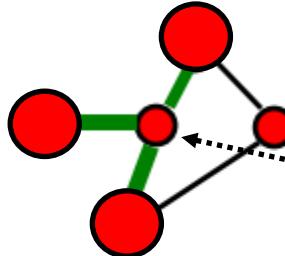
A **linear combination** specifies split of traffic between two nodes.



# CloudNet Specifications (2): Routing Model.

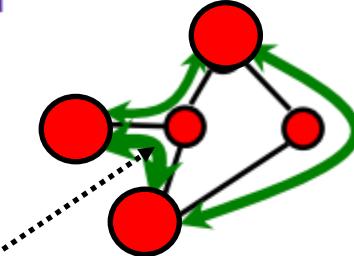
## Tree

VNet is embedded as **Steiner tree**:



## Single Path

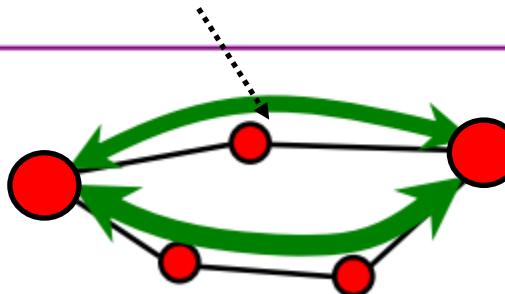
Each pair of nodes communicates along a single path.



Relay nodes may add to embedding costs!  
(resources depend, e.g., on packet rate)

## Multi Path

A **linear combination** specifies split of traffic between two nodes.



# Competitive Embeddings.

Competitive analysis framework:

## Online Algorithm

Online algorithms make decisions at time  $t$  without any knowledge of inputs / requests at times  $t' > t$ .

## Competitive Analysis

An  $r$ -competitive online algorithm ALG gives a **worst-case performance guarantee**: the performance is at most a factor  $r$  worse than an optimal offline algorithm OPT!

## Competitive Ratio

Competitive ratio  $r$ ,

$$r = \text{Cost(ALG)} / \text{cost(OPT)}$$

The **price of not knowing the future!**

No need for complex predictions but still good!



# Buchbinder&Naor: Primal-Dual Approach.

Algorithm design and analysis follows online primal-dual approach recently invented by Buchbinder&Naor!

(Application to general VNet embeddings, traffic&routing models, router loads, duration, approx oracles, ...)

## 1. Formulate dynamic primal and dual LP

$\begin{aligned} \min Z_j^T \cdot \mathbf{1} + X^T \cdot C & \text{ s.t.} \\ Z_j^T \cdot D_j + X^T \cdot A_j & \geq B_j^T \\ X, Z_j & \geq \mathbf{0} \end{aligned}$	$\begin{aligned} \max B_j^T \cdot Y_j & \text{ s.t.} \\ A_j \cdot Y_j & \leq C \\ D_j \cdot Y_j & \leq 1 \\ Y_j & \geq \mathbf{0} \end{aligned}$
(I)	(II)

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.

## 2. Derive GIPO algorithm which always produces feasible primal solutions and where Primal $\geq 2^* \text{Dual}$

**Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the  $j$ th round:

1.  $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$  (oracle procedure)
2. If  $\gamma(j, \ell) < b_j$  then, (accept)
  - (a)  $y_{j,\ell} \leftarrow 1$ .
  - (b) For each row  $e$  : If  $A_{e,(j,\ell)} \neq 0$  do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

- (c)  $z_j \leftarrow b_j - \gamma(j, \ell)$ .
3. Else, (reject)
  - (a)  $z_j \leftarrow 0$ .



# Result.

## Theorem

The presented online algorithm GIPO is **log-competitive** in the amount of resources in the physical network!  
If capacities can be exceeded by a log factor, it is even **constant competitive**.

However, competitive ratio also depends on max benefit!



# Algorithm and Proof Sketch (1).

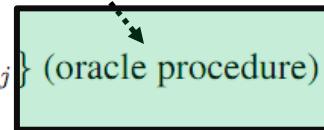
**Embedding oracle:** GIPO invokes an oracle procedure to determine cost of CloudNet embedding!

---

**Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the  $j$ th round:

1.  $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$  (oracle procedure)
2. If  $\gamma(j, \ell) < b_j$  then, (accept)
  - (a)  $y_{j,\ell} \leftarrow 1$ .
  - (b) For each row  $e$  : If  $A_{e,(j,\ell)} \neq 0$  do



$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)}/c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)}/c_e} - 1).$$

- (c)  $z_j \leftarrow b_j - \gamma(j, \ell)$ .
3. Else, (reject)
  - (a)  $z_j \leftarrow 0$ .



# Algorithm and Proof Sketch (1).

If resource cost lower than benefit: accept!

---

**Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the  $j$ th round:

1.  $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$  (oracle procedure)
2. If  $\gamma(j, \ell) < b_j$  then, (accept)
  - (a)  $y_{j,\ell} \leftarrow 1$ .
  - (b) For each row  $e$  : If  $A_{e,(j,\ell)} \neq 0$  do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

- (c)  $z_j \leftarrow b_j - \gamma(j, \ell)$ .
  3. Else, (reject)
    - (a)  $z_j \leftarrow 0$ .
- 



# Algorithm and Proof Sketch (1).

---

**Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the  $j$ th round:

1.  $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$  (oracle procedure)
2. If  $\gamma(j, \ell) < b_j$  then, (accept)

- (a)  $y_{j,\ell} \leftarrow 1$ .
- (b) For each row  $e$  : If  $A_{e,(j,\ell)} \neq 0$  do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)}/c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)}/c_e} - 1).$$

- (c)  $z_j \leftarrow b_j - \gamma(j, \ell)$ .

3. Else, (reject)
    - (a)  $z_j \leftarrow 0$ .
- 

update allocations for  
accepted CloudNet...



# Algorithm and Proof Sketch (1).

---

**Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the  $j$ th round:

1.  $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$  (oracle procedure)
2. If  $\gamma(j, \ell) < b_j$  then, (accept)
  - (a)  $y_{j,\ell} \leftarrow 1$ .
  - (b) For each row  $e$  : If  $A_{e,(j,\ell)} \neq 0$  do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)}/c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)}/c_e} - 1).$$

- (c)  $z_j \leftarrow b_j - \gamma(j, \ell)$ .

3. Else, (reject)
    - (a)  $z_j \leftarrow 0$ .

**otherwise reject  
(no change in substrate)**



# Algorithm and Proof Sketch (1).

---

**Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the  $j$ th round:

1.  $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$  (oracle procedure)
2. If  $\gamma(j, \ell) < b_j$  then, (accept)
  - (a)  $y_{j,\ell} \leftarrow 1$ .
  - (b) For each row  $e$  : If  $A_{e,(j,\ell)} \neq 0$  do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)}/c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)}/c_e} - 1).$$

- (c)  $z_j \leftarrow b_j - \gamma(j, \ell)$ .

3. Else, (reject)
    - (a)  $z_j \leftarrow 0$ .

**otherwise reject  
(no change in substrate)**

---

**Algorithm efficient... except for oracle (static, optimal embedding)!  
What if we only use a suboptimal embedding here?**



# Algorithm and Proof Sketch (2).

Problem: computation of optimal embeddings NP-hard!  
Thus: use approximate embeddings! (E.g., Steiner tree)

**GIPO:**

**Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).  
Upon the  $j$ th round:

1.  $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$  (oracle procedure)
2. If  $\gamma(j, \ell) < b_j$  then, (accept)
  - (a)  $y_{j,\ell} \leftarrow 1$ .
  - (b) For each row  $e$  : If  $A_{e,(j,\ell)} \neq 0$  do
$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$
  - (c)  $z_j \leftarrow b_j - \gamma(j, \ell)$ .
3. Else, (reject)
  - (a)  $z_j \leftarrow 0$ .

Competitive ratio  $\rho$

**Embedding approx.:**

<insert your favorite  
approx algo>

Approx ratio  $r$

## Lemmas

The approximation does not reduce the overall competitive ratio by much: we get  $\rho^*r$  ratio!



# Proof Sketch (1): Simplified LP.

$$\begin{aligned}
 & \min \sum_{e \in E} x_e \cdot c(e) + \sum_{v \in V} x_v \cdot c(v) + \sum_i z_i \cdot d_i \quad s.t. \\
 & (\text{Covering Const.}) \forall i \forall \Delta \in \Delta_i z_i + \alpha(i, \Delta) \geq b_i \\
 & \forall i \forall \Delta \in \Delta_i x_e, x_v, z_i \geq 0
 \end{aligned}$$

maximize  
benefit!

(I)

realization of i-th  
request (will be integer,  
accept fully or not at all)

$$\begin{aligned}
 & \max \sum_i b_i \cdot \sum_{\Delta_{ij} \in \Delta_i} f_{ij} \quad s.t. \\
 & (\text{Vertex Capacity Const.}) \quad \forall v \in V \text{ } flow(v) \leq c(v) \\
 & (\text{Edge Capacity Const.}) \quad \forall e \in E \text{ } flow(e) \leq c(e) \\
 & (\text{Demand Const.}) \quad \forall i \sum_{\Delta_{ij} \in \Delta_i} f_{ij} \leq d_i \\
 & f \geq 0
 \end{aligned}$$

(II)

Fig. 1: (I) The Primal linear embedding program. (II)  
The Dual linear embedding program.



# Proof Sketch (2): Simplified LP.

essentially, exponential load...

$$\begin{aligned} \min & \sum_{e \in E} x_e \cdot c(e) + \sum_{v \in V} x_v \cdot c(v) + \sum_i z_i \cdot d_i \quad s.t. \\ (\text{Covering Const.}) \quad & \forall i \forall \Delta \in \Delta_i z_i + \alpha(i, \Delta) \geq b_i \\ & \forall i \forall \Delta \in \Delta_i x_e, x_v, z_i \geq 0 \end{aligned}$$

(I)

$$\begin{aligned} \max & \sum_i b_i \cdot \sum_{\Delta_{ij} \in \Delta_i} f_{ij} \quad s.t. \\ (\text{Vertex Capacity Const.}) \quad & \forall v \in V \text{ } flow(v) \leq c(v) \\ (\text{Edge Capacity Const.}) \quad & \forall e \in E \text{ } flow(e) \leq c(e) \\ (\text{Demand Const.}) \quad & \forall i \sum_{\Delta_{ij} \in \Delta_i} f_{ij} \leq d_i \\ & f \geq 0 \end{aligned}$$

(II)

Fig. 1: (I) The Primal linear embedding program. (II)  
The Dual linear embedding program.



# Proof Sketch (3): Simplified LP.

---

**Algorithm 1** The ISTP Algorithm.

---

Input:  $G = (V, E)$  (possibly infinite), sequence of requests  $\{r_i\}_{i=1}^{\infty}$  where  $r_i \triangleq (U_i, c_i, d_i, b_i)$ .

Upon arrival of request  $r_i$ :

- 1)  $j \leftarrow \operatorname{argmin}\{\alpha(i, j) : \Delta_{ij} \in \Delta_i\}$  (find a lightest realization over the terminal set  $U_i$  using an oracle).

- 2) If  $\alpha(i, j) < b_i$  then, (accept  $r_i$ )

- a)  $f_{ij} \leftarrow d_i$ .

- b) For each  $e \in E(\Delta_{ij})$  do

$$x_e \leftarrow x_e \cdot 2^{d_i/c(e)} + \frac{1}{|V(\Delta_{ij})|} \cdot (2^{d_i/c(e)} - 1).$$

oracle  
(triangle only)

update primal  
variables if accepted

- c) For each  $v \in V(\Delta_{ij})$  do

$$x_v \leftarrow x_v \cdot 2^{c_i/c(v)} + \frac{d_i/c_i}{|V(\Delta_{ij})|} \cdot (2^{c_i/c(v)} - 1).$$

- d)  $z_i \leftarrow b_i - \alpha(i, j)$ .

- 3) Else, (reject  $r_i$ )

- a)  $z_i \leftarrow 0$ .



# Proof Sketch (4): Simplified LP.

Step (2b) increases the cost  $\sum_e x_e \cdot c(e)$  as follows  
 (change  $\Delta(x_e) = \sum_e (x_e^t - x_e^{t-1}) \cdot c(e)$ ):

$$\begin{aligned}
 \Delta(x_e) &\leq \sum_{e \in \Delta} \left[ x_e \cdot (2^{d_i/c(e)} - 1) + \frac{1}{|V(\Delta_{ij})|} \cdot (2^{d_i/c(e)} - 1) \right] \cdot c(e) \\
 &= \sum_{e \in \Delta} \left( x_e + \frac{1}{|V(\Delta_{ij})|} \right) \cdot (2^{d_i/c(e)} - 1) \cdot c(e) \\
 &\leq c_{\min}(e) \cdot (2^{d_i/c_{\min}(e)} - 1) \sum_{e \in \Delta} \left( x_e + \frac{1}{|V(\Delta_{ij})|} \right) \\
 &\leq d_i \cdot (2^1 - 1) \sum_{e \in \Delta} \left( x_e + \frac{1}{|V(\Delta_{ij})|} \right) \\
 &\leq d_i \cdot \sum_{e \in \Delta} x_e + d_i \cdot \sum_{e \in \Delta} \frac{1}{|V(\Delta_{ij})|} \\
 &\leq d_i \cdot \sum_{e \in \Delta} x_e + d_i. \tag{1}
 \end{aligned}$$

after each request,  
 primal variables  
 constitute feasible  
 solutions...

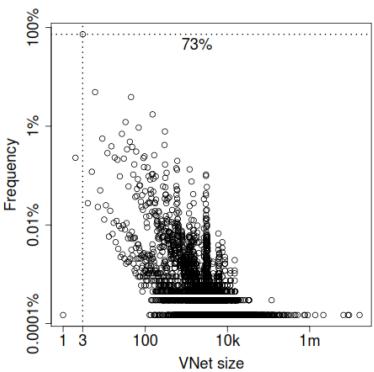
Step (2c) increases the cost  $\sum_v x_v \cdot c(v)$  as follows  
 (change  $\Delta(x_v) = \sum_v (x_v^t - x_v^{t-1}) \cdot c(v)$ ):

$$\begin{aligned}
 \delta(x_v) &\leq \sum_{v \in \Delta} \left[ x_v \cdot (2^{c_i/c(v)} - 1) + \frac{d_i/c_i}{|V(\Delta_{ij})|} \cdot (2^{c_i/c(v)} - 1) \right] \cdot c(v) \\
 &= \sum_{v \in \Delta} \left( x_v + \frac{d_i/c_i}{|V(\Delta_{ij})|} \right) \cdot (2^{c_i/c(v)} - 1) \cdot c(v) \\
 &\leq c_{\min}(v) \cdot (2^{c_i/c_{\min}(v)} - 1) \sum_{v \in \Delta} \left( x_v + \frac{d_i/c_i}{|V(\Delta_{ij})|} \right) \\
 &\leq c_i \cdot (2^1 - 1) \sum_{v \in \Delta} \left( x_v + \frac{d_i/c_i}{|V(\Delta_{ij})|} \right) \\
 &\leq c_i \cdot \sum_{v \in \Delta} x_v + c_i \cdot \sum_{v \in \Delta} \frac{d_i/c_i}{|V(\Delta_{ij})|} \\
 &\leq c_i \cdot \sum_{v \in \Delta} x_v + d_i. \tag{2}
 \end{aligned}$$



# On the Benefit of Collocation.

Google cluster: many small networks, over 90% allow for collocation



Greedy vs SecondNet vs ViNE:  
Greedy collocation algorithm beats them all...!

---

### Algorithm 1 The LoCo Algorithm

---

**Require:** VNet  $G = (V, E)$ ,  $M = \{s\}$  for some  $s \in V(G)$ ,  
 $P = (\Gamma(s))$

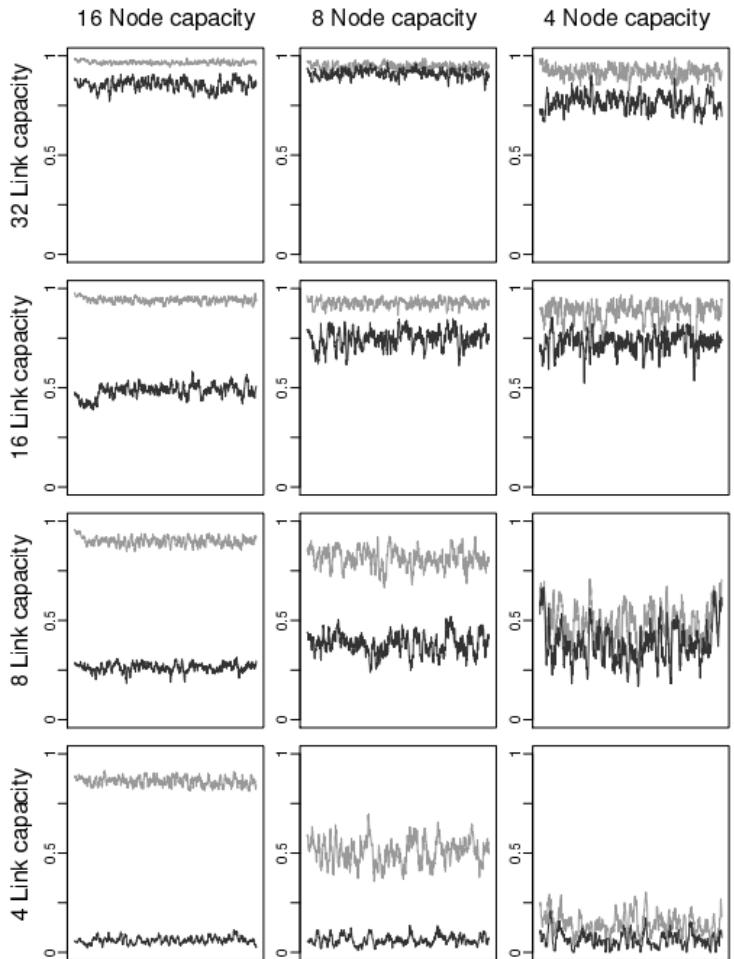
**while**  $|P| > 0$  **do**

- sort**  $P$  (\* decreasing link capacities \*)
- choose**  $u = P[0]$  (\* next node to map \*)
- map**  $u$  (\* forward checking \*)
- map**  $\{u, v\} \quad \forall v \in M$ , **where**  $\{u, v\} \in E(G)$
- $M = M \cup \{u\}$  **and**  $P = P \setminus \{u\}$

**end while**

**if** (embedding failed), **backtrack** on  $s$

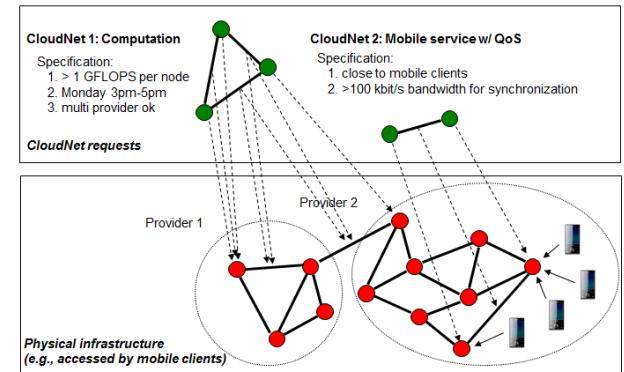
---



# Mixed Integer Programs

# Problem 1: Classic VNet Embedding (VNEP).

- Map virtual nodes to substrate nodes
  - Collocation possible
  - But not splitting of virtual nodes
- Map virtual links
  - One
  - Linear combination
  - Hose
- Mixed Integer Model
  - VINO ☺
- Open Problems
  - Everything ☺

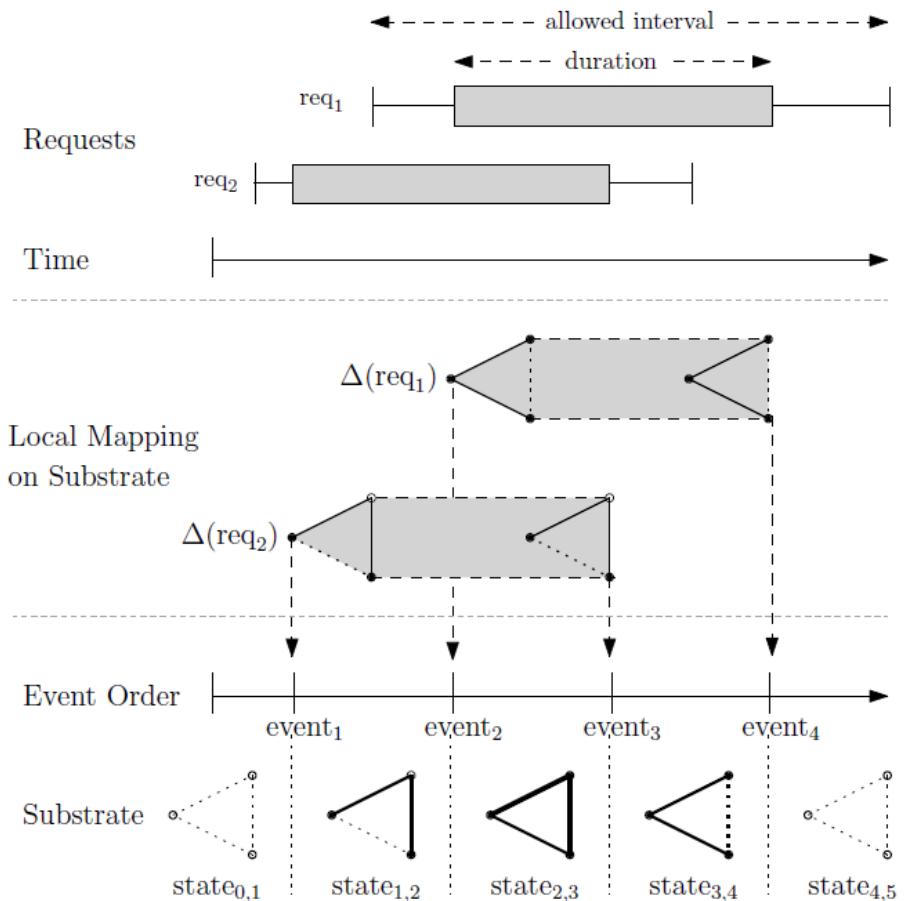


The Informal Guide to the  
Virtual Network Embedding MIP Creator  
TU Berlin, Germany  
Matthias Rost  
mrost@inet.tu-berlin.de



# Problem 2: Embedding with Time Flexibilities (TVNEP).

- VNets come with time flexibilities
- Example: delay-tolerant computations, bulk data transfers, etc.
- **Where** to embed and **when** to schedule VNets?



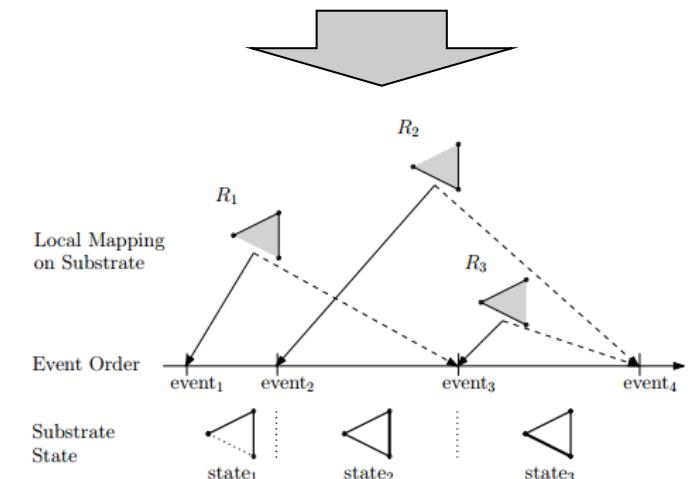
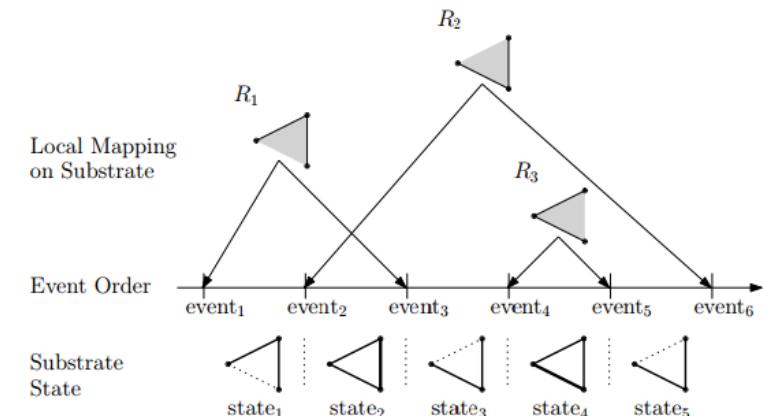
# Problem 2: Embedding with Time Flexibilities (TVNEP).

- Continuous time (less binary variables): **state model** (explicit states) and **delta model** (only differences)

- Delta model yields bad **relaxations**
- State model has **more variables**, but still better

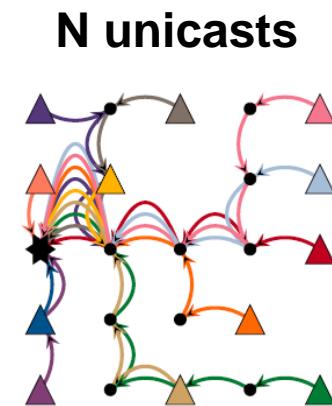
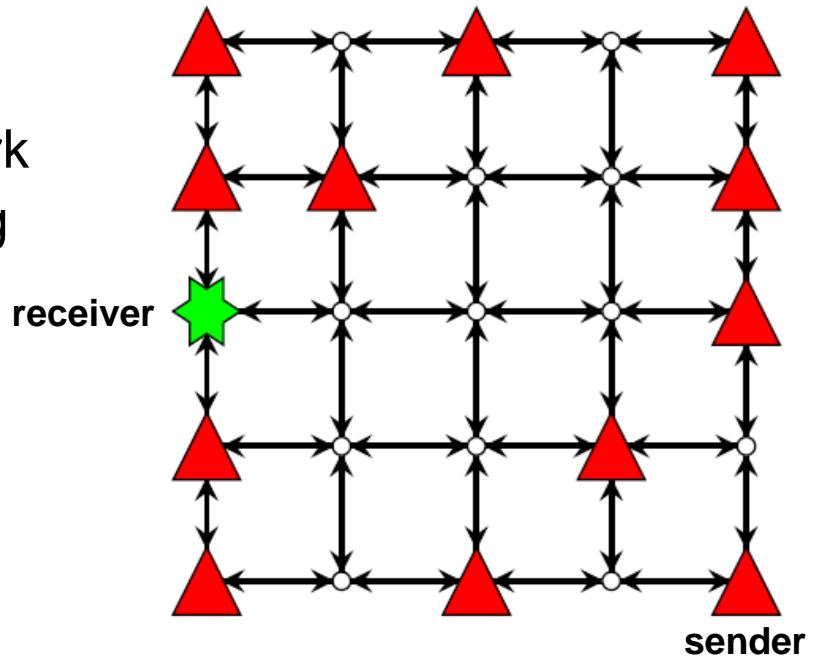
- Compact variant:

- **State reduction**: feasibility check only at start of request sufficient (when finishes less resources)
- Minimize smear-out: Distribute start and end to as few event points as possible
- «Merge» **multiple endpoints** with start points
- Compute **temporal dependencies** graph cuts

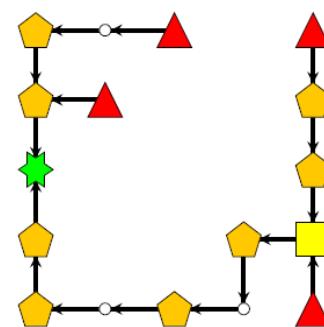


# Problem 3: VirtuCast / In-Network Processing (CVSAP).

- Network Function Virtualization
  - Can aggregate / split streams in network
  - E.g., streaming or wide-area monitoring
- Universal nodes need to be activated:  
Joint optimization of processing and communication?
  - Note: DAG may not be optimal!



vs.



Communication costs!

Processing costs!

**Generalization?**

Rost et al.:

OPODIS 2013

# Problem 3: VirtuCast / In-Network Processing (CVSAP).

- Multi-Commodity flow bad: for 200 Steiner nodes and 6800 edges, 1.3 mio binary variables!
- Single-Commodity approach: solvable!
- Idea: single commodity and then path decomposition
- Open question: can we even relax path variables and optimally round it afterwards?

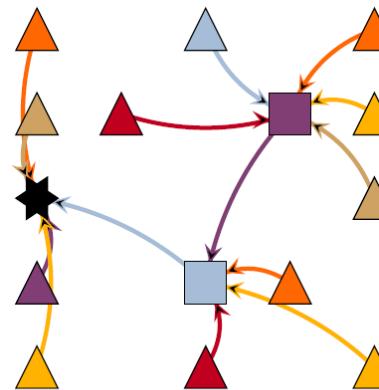


Figure: Virtual Arborescence

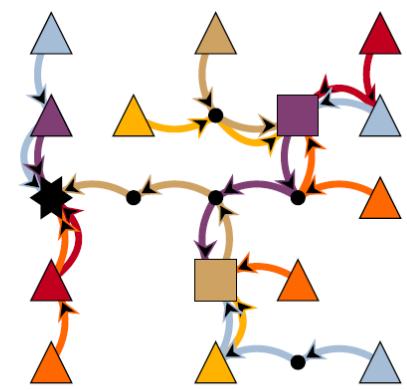
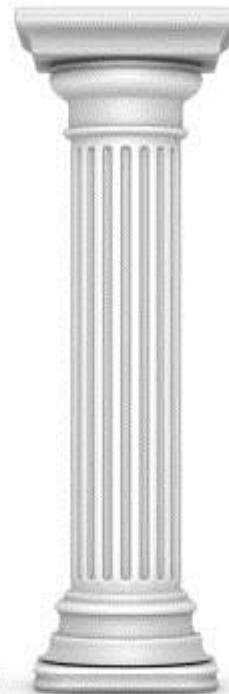


Figure: Flow in original graph

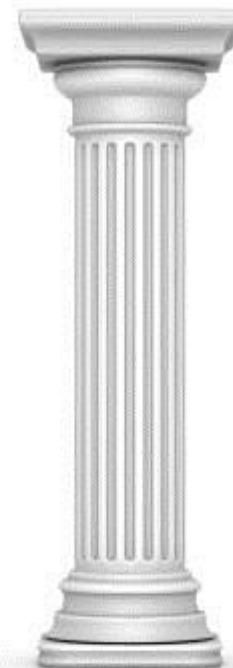


# Migration.

**Access Control  
and Embedding**



**Service Migration**



**Security Issues**

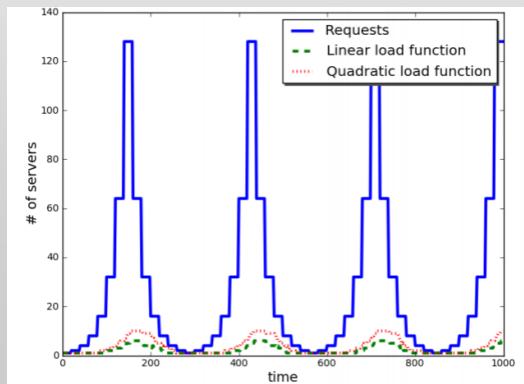
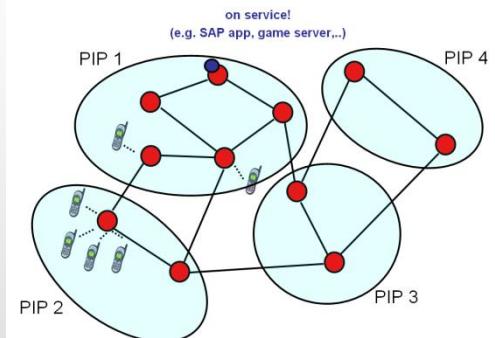


# Online Service Migration.

## Online Migration and Allocation

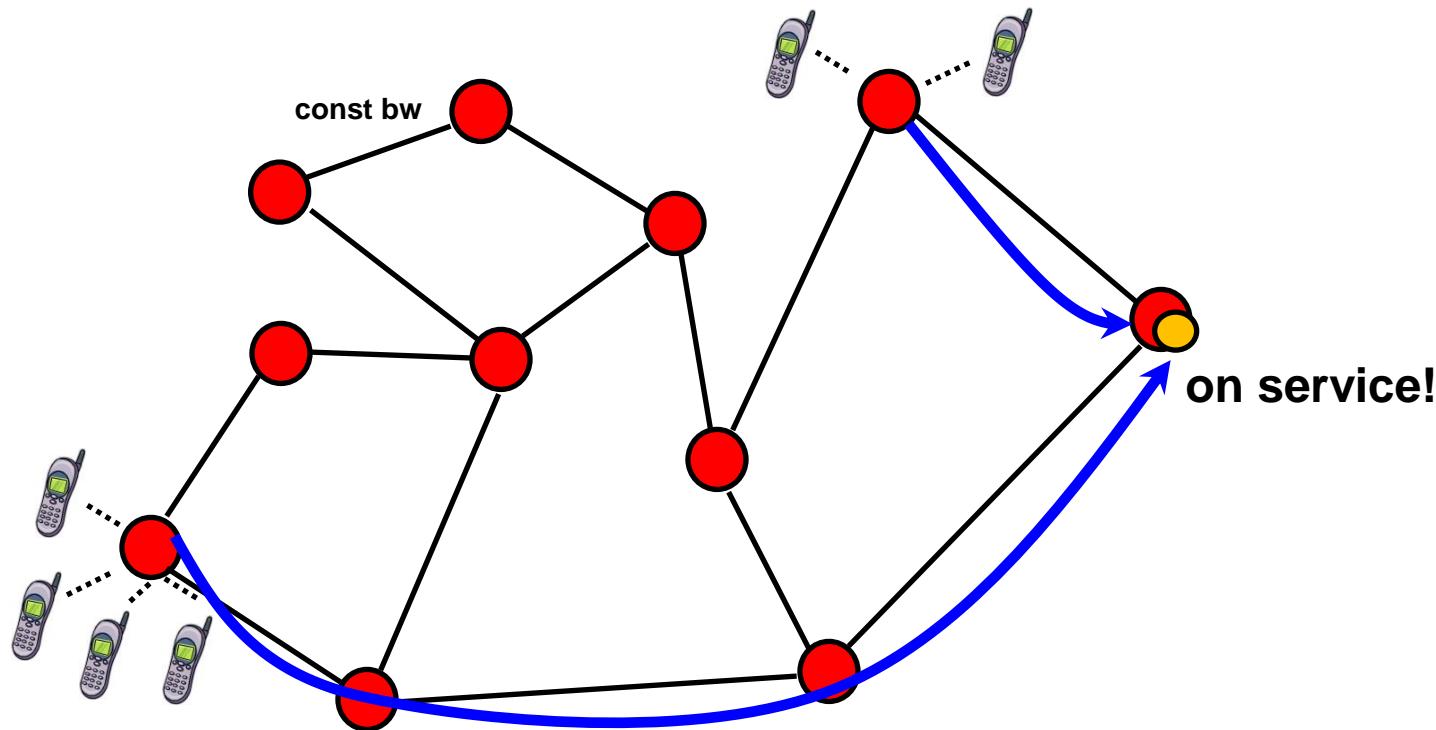
### Goal:

E.g., **QoS** (=“move with the sun”, or with  
commuters); or for **maintenance** or to turn off  
resources (**energy** conservation).



# The Virtual Service Migration Problem.

Bienkowski et al.:  
SIGCOMM VISA 2010



**Given a virtual network with guaranteed bandwidth: where to migrate service?**

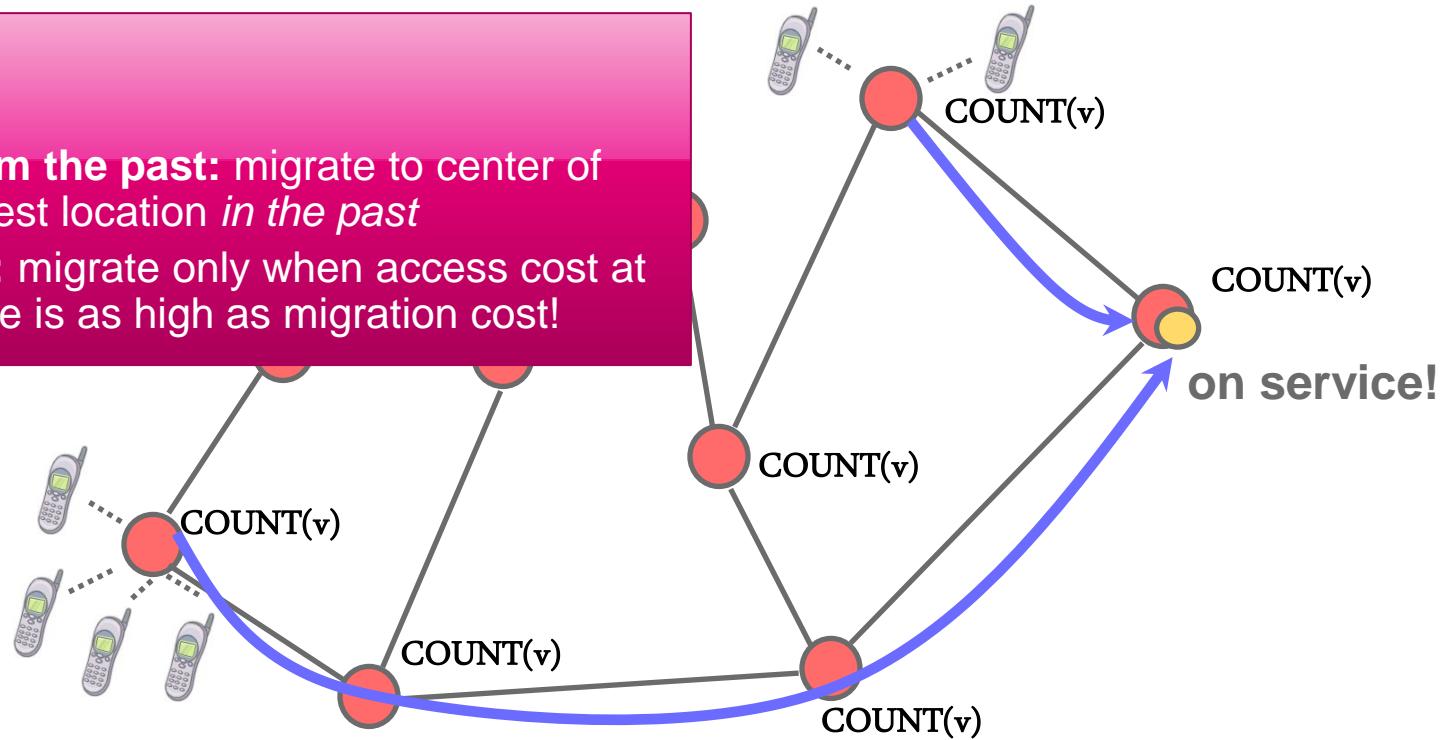
Simple model: **one service**, constant migration cost (interruption), access along graph.



# The Virtual Service Migration Problem.

## Idea:

- **Learn from the past:** migrate to center of gravity of best location *in the past*
- **Amortize:** migrate only when access cost at current node is as high as migration cost!



Given a virtual network with guaranteed ban

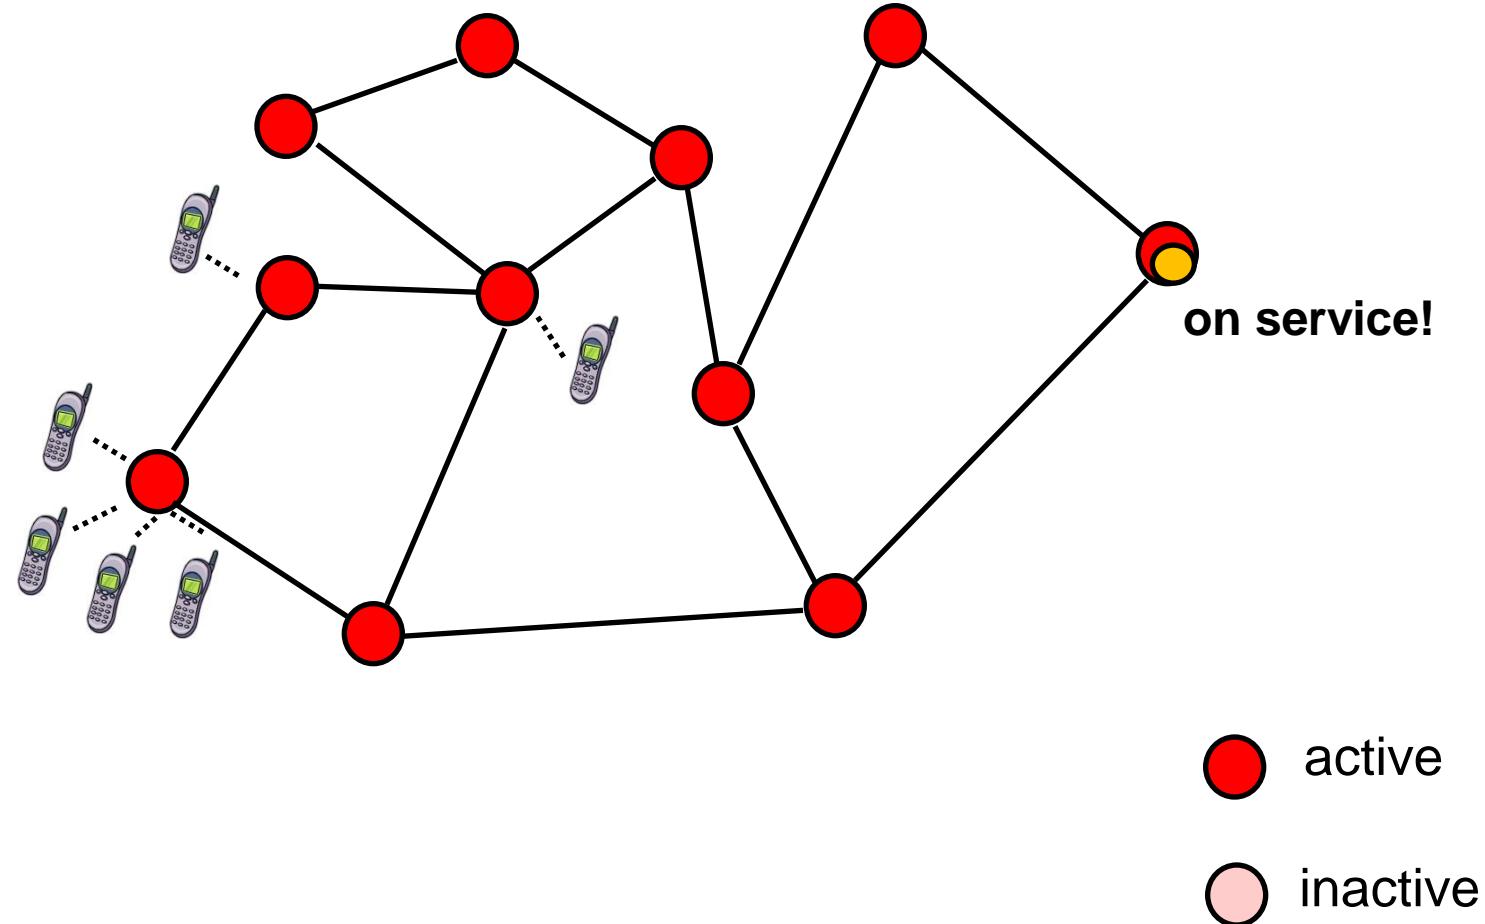
Simple model: **one service**, constant migration cost (i.e.

### Center of Gravity Migration

1. Each node  $v$ :  $COUNT(v)$  = access cost **epoch**
2. Call nodes  $v$  with  $COUNT(v) < m$  **active**.
3. If service is at node  $w$ , a **phase** ends when  $COUNT(w) \geq m$
4. The service is migrated to the **center of gravity** of the remaining active nodes
5. If no such node is left, the epoch ends.

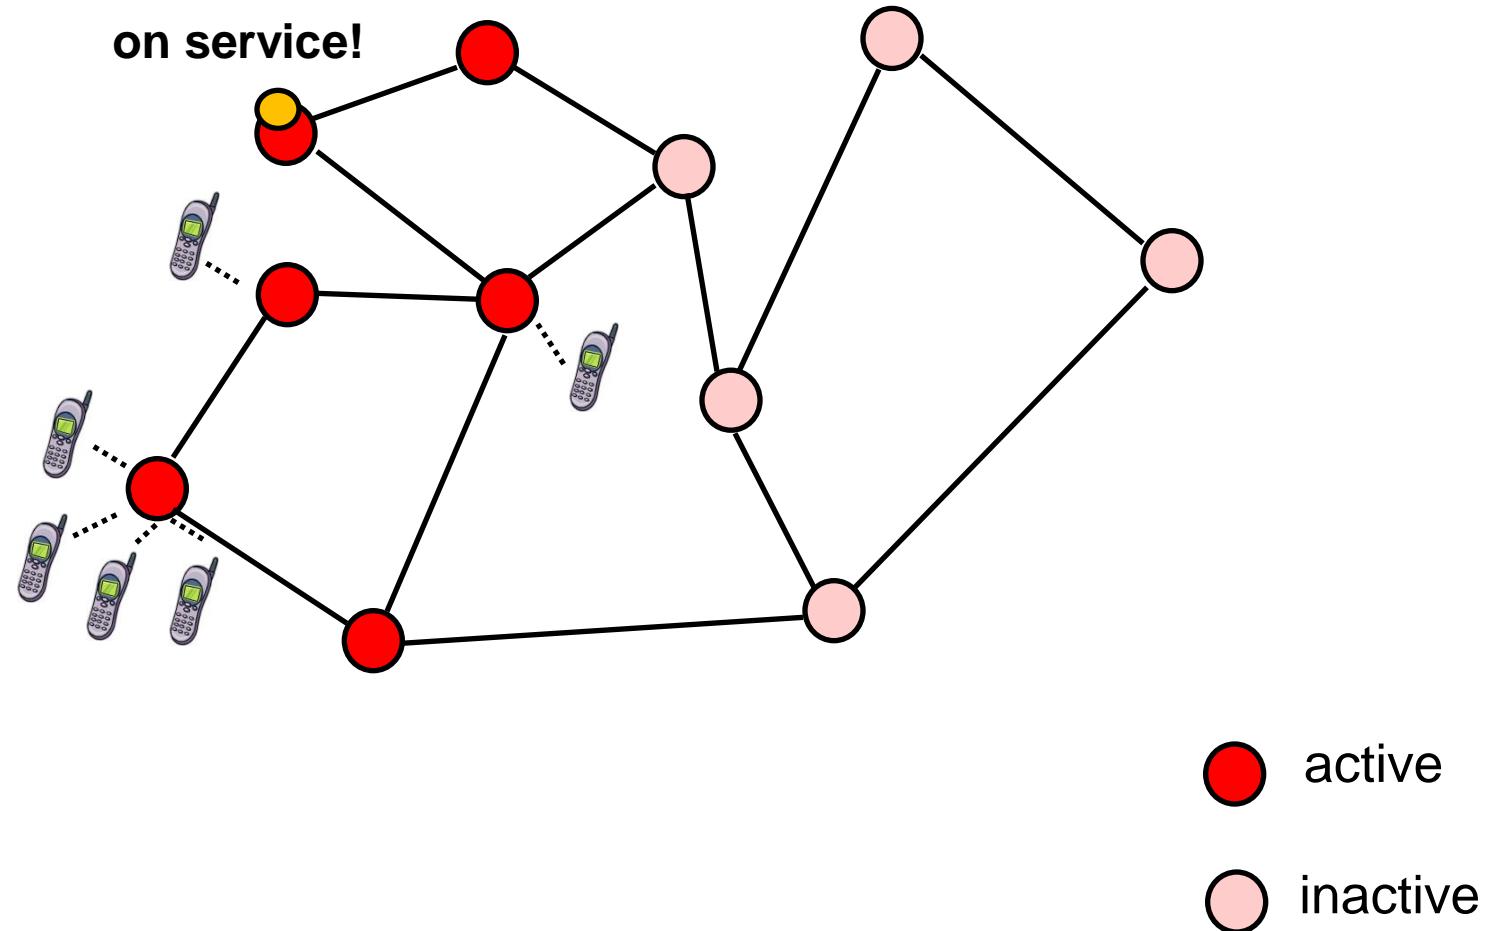
# Center-of-Gravity Algo: Example.

Before phase 1:



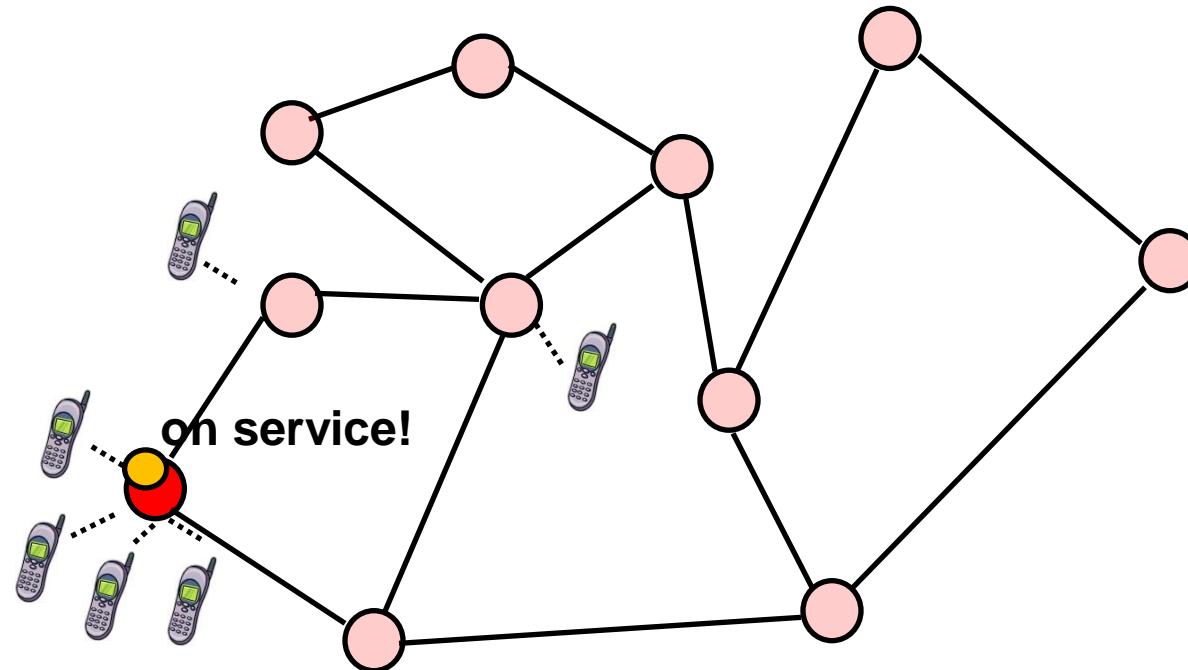
# Center-of-Gravity Algo: Example.

Before phase 2:



# Center-of-Gravity Algo: Example.

End of epoch:



Of course, not converging if demand is dynamic!  
(Simplified example.)

- active
- inactive



# Center-of-Gravity Algo: Result.

Competitive analysis? Assume constant bandwidths!

$$r = \text{ALG} / \text{OPT} ?$$

Lower bound cost of OPT:

In an **epoch**, each node has **at least** access cost  **$m$** , or there was a migration of cost  **$m$** .

Upper bound cost of ALG:

We can show that each **phase** has cost **at most  $2m$**  (access plus migration), and there are at most  **$\log(n)$**  many phases per **epoch**!

## Theorem

**ALG is  $\log(n)$  competitive!**

A special uniform metrical task system (graph metric for access)!



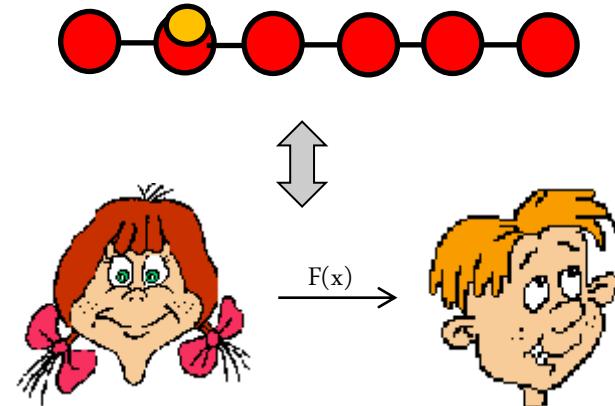
## Theorem

«Center of Gravity» algorithm is  
 $\log(n)$  competitive!

$\log(n)/\log\log(n)$  lower bound  
follows from online function  
tracking reduction!

Also a much simpler randomized  
algorithm achieves this!

on service!



## Theorem

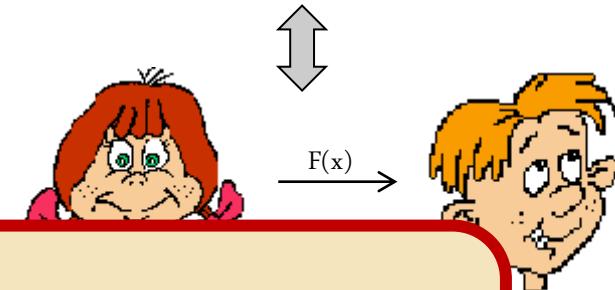
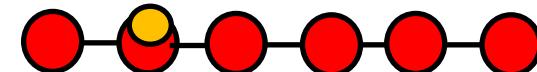
«Center of Gravity» algorithm is  
 $\log(n)$  competitive!

$\log(n)/\log\log(n)$  lower bound  
follows from online function  
tracking reduction!

There is an asymptotically  
optimal called FOLLOWER!

Also a much simpler randomized  
algorithm achieves this!

on service!



# The Online Algorithm FOLLOWER.

## Concepts:

- **Learn from the past:** migrate to **center of gravity** of best location *in the past*
- **Amortize:** migrate only when access cost at current node is as high as migration cost!

## Simplified Follower

1.  $F_i$  are requests handled while service at  $f_i$
2. to compute  $f_{i+1}$  (new pos), Follower only takes into account requests during  $f_i$ :  $F_i$
3. migrate to center of gravity of  $F_i$ , as soon as migration costs there are amortized (and «reset counters» immediately)!

### Algorithm Follower

```
1:  $i := 0; k_0 := 0 \forall j: F_j = \{ \}$  {The server starts at an arbitrary node  $f_0$ }
Upon a new request  $r$  do:
2: Serve request  $r$  with server at  $f_i$ 
3:  $F_i := F_i \cup r$ 
4:  $f' :=$  arbitrary  $u \in CG(F_i)$ 
5:  $x' := d(f_i, f')$  {for co.di., and  $x' := 1$  for co.nb.m.}
6: if  $C(f_i, F_i) \geq g(x'|k_i)$  then
7:    $f_{i+1} := f'; x_i := x'$ 
8:    $y(w) := d(f_i, w) + d(w, f_{i+1})$  {for co.di., and for co.nb.m.  $y(w) := 2$  for  $w \neq f_{i+1}$  and  $y(w) := 1$  otherwise }
9:    $slack(w \in V) := g(y(w)|k_i) - C(f_i, F_i)$ 
10:   $w_i :=$  Node  $w$  with minimum  $slack(w)$  such that  $slack(w) \geq 0$ 
11:  Move server to  $w_i$  and if  $w_i \neq f_{i+1}$  onto  $f_{i+1}$ 
12:   $k_{i+1} := k_i + y(w_i)$ 
13:   $i := i + 1$ 
14: end if
```



# The Online Algorithm FOLLOWER.

## Concepts:

- **Learn from the past:** migrate to **center of gravity** of best location *in the past*
- **Amortize:** migrate only when access cost at current node is as high as migration cost!

## Simplified Follower

1.  $F_i$  are requests handled while service at  $f_i$
2. to compute  $f_{i+1}$  (new pos), Follower only takes into account requests during  $f_i$ :  $F_i$
3. migrate to center of gravity of  $F_i$ , as soon as migration costs there are amortized (and «reset counters» immediately)!

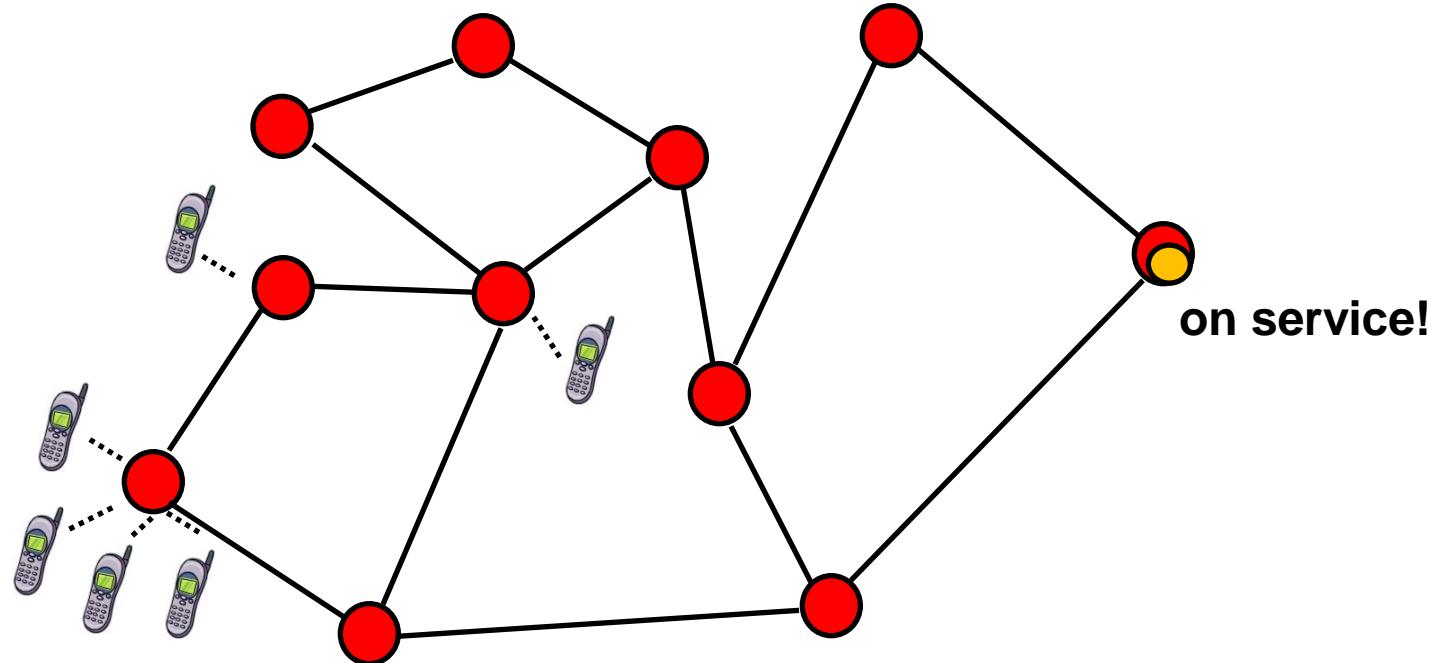
### Algorithm Follower

```
1:  $i := 0; k_0 := 0 \forall j: F_j = \{ \}$  {The server starts at an arbitrary node  $f_0$ }
Upon a new request  $r$  do:
2: Serve request  $r$  with server at  $f_i$ 
3:  $F_i := F_i \cup r$ 
4:  $f' :=$  arbitrary  $u \in CG(F_i)$ 
5:  $x' := d(f_i, f')$  {for co.di., and  $x' := 1$  for co.nb.m.}
6: if  $C(f_i, F_i) \geq g(x'|k_i)$  then
7:    $f_{i+1} := f'; x_i := x'$ 
8:    $y(w) := d(f_i, w) + d(w, f_{i+1})$  {for co.di., and for co.nb.m.  $y(w) := 2$  for  $w \neq f_{i+1}$  and  $y(w) := 1$  otherwise }
9:    $slack(w \in V) := g(y(w)|k_i) - C(f_i, F_i)$ 
10:   $w_i :=$  Node  $w$  with minimum  $slack(w)$  such that  $slack(w) \geq 0$ 
11:  Move server to  $w_i$  and if  $w_i \neq f_{i+1}$  onto  $f_{i+1}$ 
12:   $k_{i+1} := k_i + y(w_i)$ 
13:   $i := i + 1$ 
14: end if
```

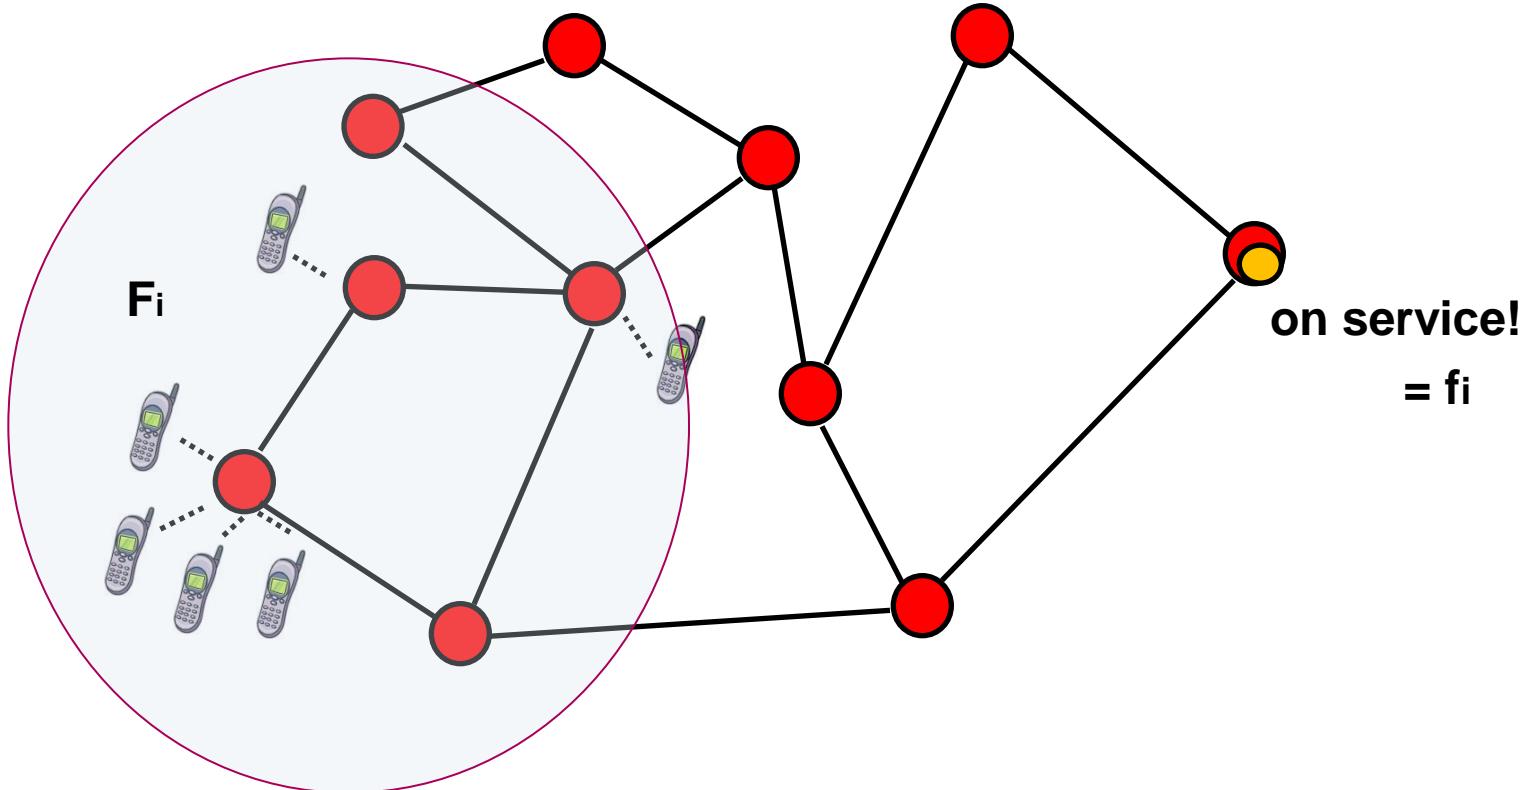
Also works for migrations with discount!  
(Reseller/broker gives discount!)



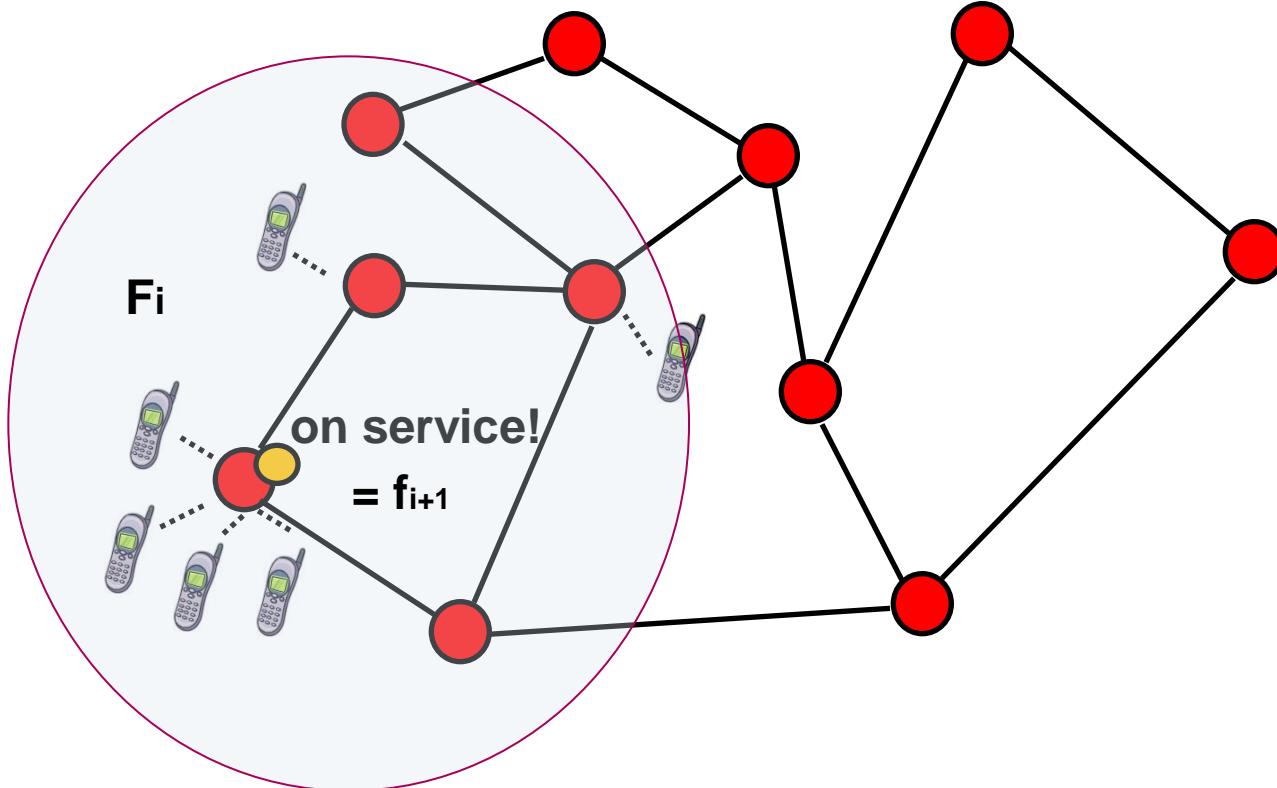
# Intuition.



# Intuition.



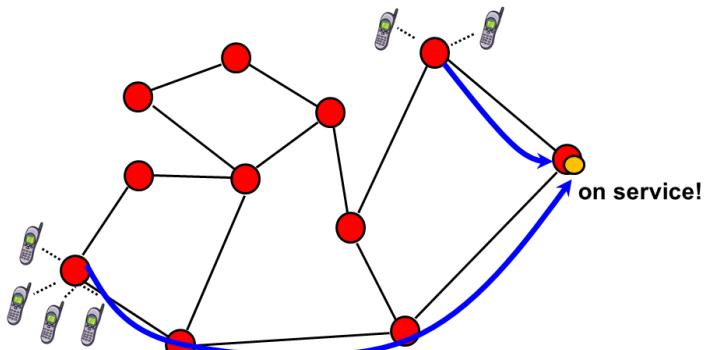
# Intuition.



# Modeling Access and Migration Costs.

## Access Costs

Latency along shortest path in graph.  
(Graph distances, and in particular: metric!)



## Migration Costs

Generalized models:

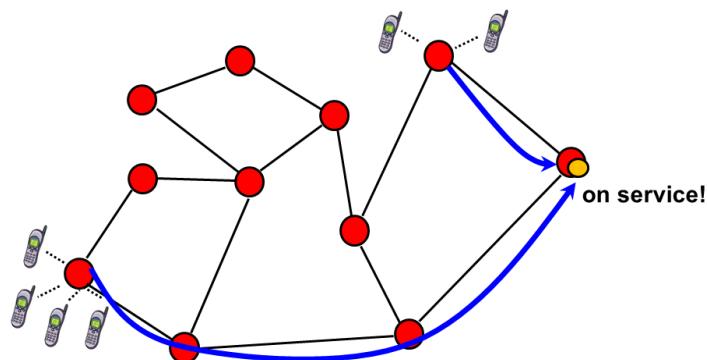
- E.g., depends on bandwidth along path (duration of service interruption)
- E.g., depends on distance travelled (latency)
- Discount: e.g., VNP (number of migrations, distance travelled, ...)



# Modeling Access and Migration Costs.

## Access Costs

Latency along shortest path in graph.  
(Graph distances, and in particular: metric!)



## Migration Costs

Generalized model:

- E.g.  $g(x|y)$ : cost of migrating

(distance  $x$  given already travelled  $y$ )

General cost function  $g(x|y)$ : cost of migrating  
distance  $x$  given already travelled  $y$

Or  $g(1|y)$ : cost of migration given we already  
migrated  $y$  times

Or  $g(0|y)$ : cost of migration given we have  
travelled, ...)



# Competitive Ratio of FOLLOWER.

Competitive analysis? FOLLOWER / OPT?

## Theorem

If no discounts are given,  
Follower is  $\log(n)/\log\log(n)$  competitive!

Simple model with *migration costs* = *bandwidth*, and *homogeneous*

Page migration model with  
*migration costs* = *distance*,  
but discounts

## Theorem

If migration costs depend on travelled  
distance (page migration), competitive  
ratio is  $O(1)$ , even with discounts.



# Related Work.

- Metrical Task Systems:
  - Classical online problem where server at certain location («state») serves requests at certain costs; state transitions also come at certain costs («migration»)
  - Depending on migration cost function more general (we have **graph access costs**) and less general (we allow for **migration discounts**)
  - E.g., **uniform space metrical task system**: migration costs constant, but access costs more general than graph distances! Lower bound of  **$\log(n)$**  vs  **$\log(n)/\log\log(n)$**  upper bound in our case.
- Online Page Migration
  - Classical online problem from the 80ies; we generalize cost function to distance discounts, while keeping O(1)-competitive

**Our work lies between!**



# Simulation.

## Commuter Scenario

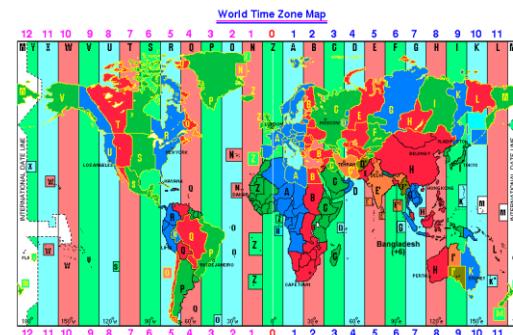
Dynamics due to mobility: requests cycle through a 24h pattern: in the morning, requests distributed widely (people in suburbs), then focus in city centers; in the evening, reverse.



Predictable scenarios,  
but we do not exploit that.  
Reality less predictable!

## Time Zone Scenario

Dynamics due to time zone effects: request originate in China first, then more requests come from European countries, and finally from the U.S.

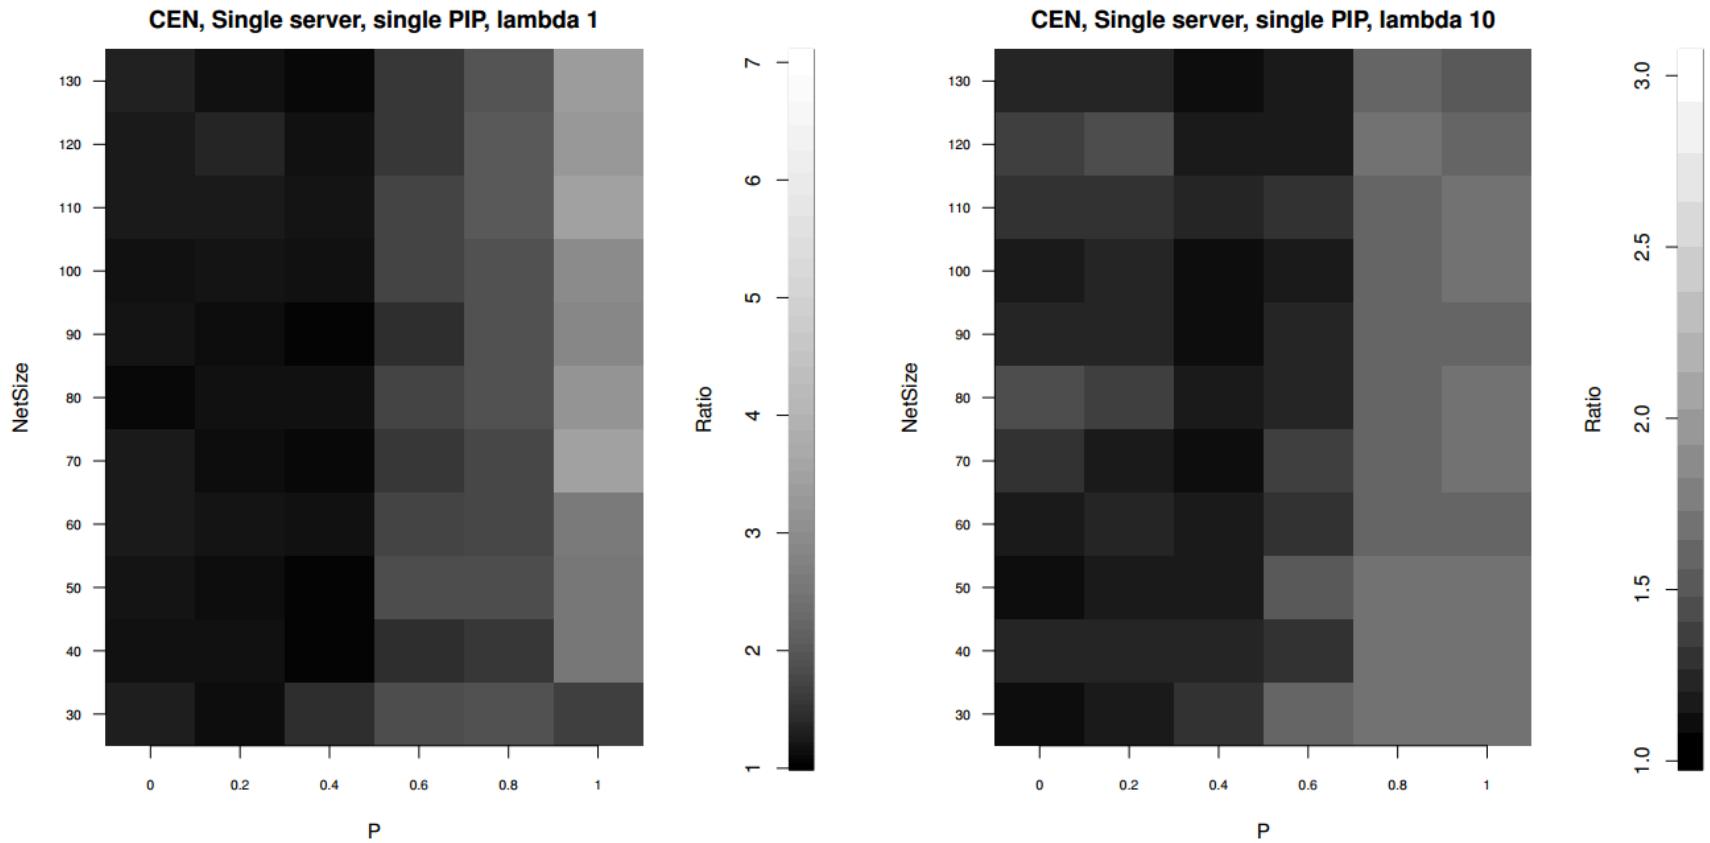


## Static Algorithm

Algorithm which uses optimal static server placements for a given request seq.



# Results.



Competitive ratio generally relatively low. Increases for more correlated requests and more dynamics.



# Related Work.

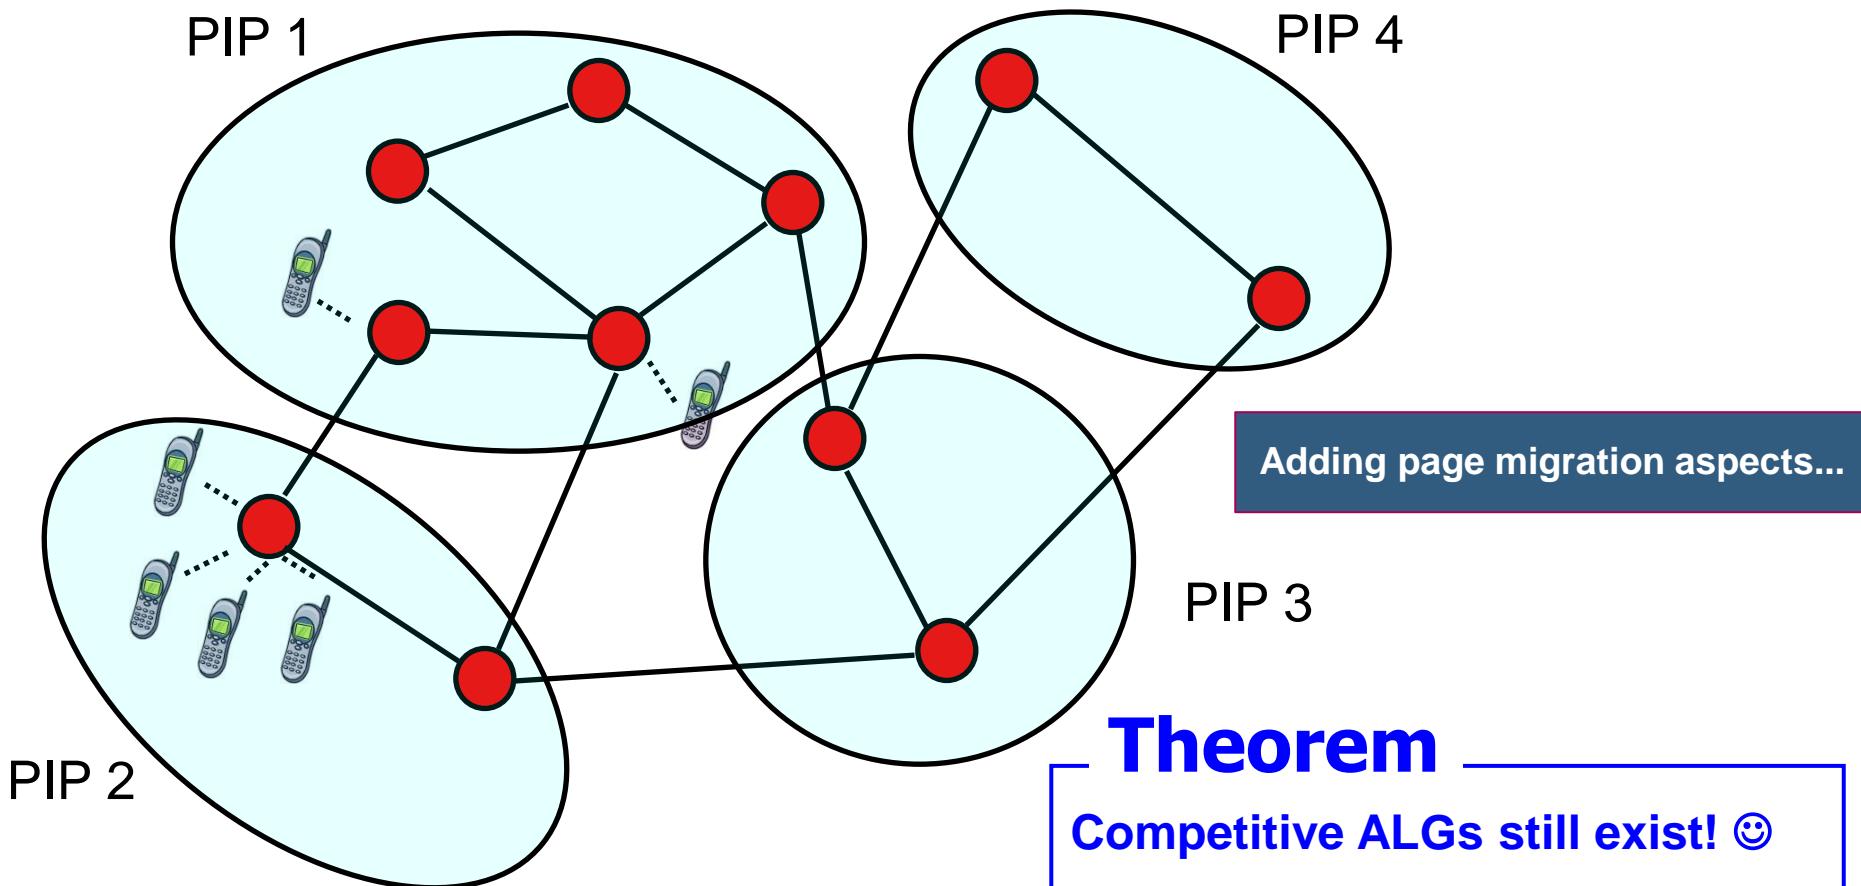
- Metrical Task Systems:
  - Classical online problem where server at certain location («state») serves requests at certain costs; state transitions also come at certain costs («migration»)
  - Depending on migration cost function more general (we have **graph access costs**) and less general (we allow for **migration discounts**)
  - E.g., **uniform space metrical task system**: migration costs constant, but access costs more general than graph distances! Lower bound of  **$\log(n)$**  vs  **$\log(n)/\log\log(n)$**  upper bound in our case.
- Online Page Migration
  - Classical online problem from the 80ies; we generalize cost function to distance discounts, while keeping O(1)-competitive

**Our work lies between!**

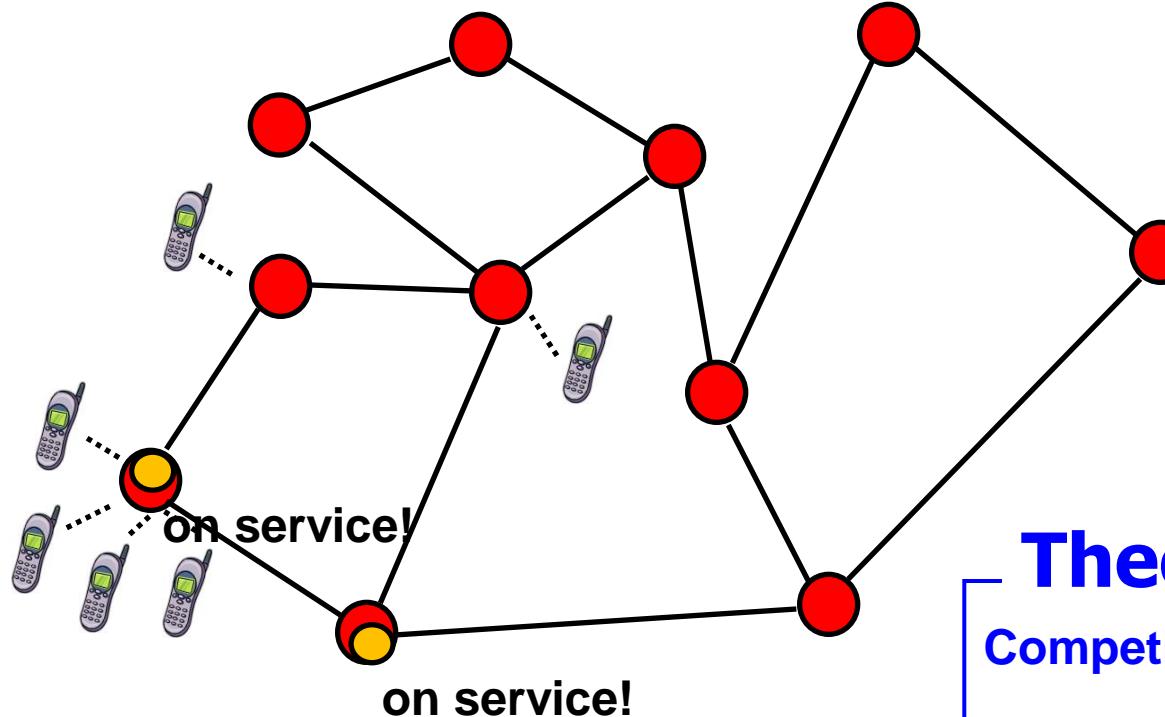


# Extension: Inter-Provider Migration.

Migration across provider boundary costs **transit/roaming costs** (# transit providers), detailed topology not known, etc.



Multiple servers allocated and migrated dynamically depending on demand and load, servers have running costs, etc.



**Theorem**  
Competitive ALGs still exist! ☺

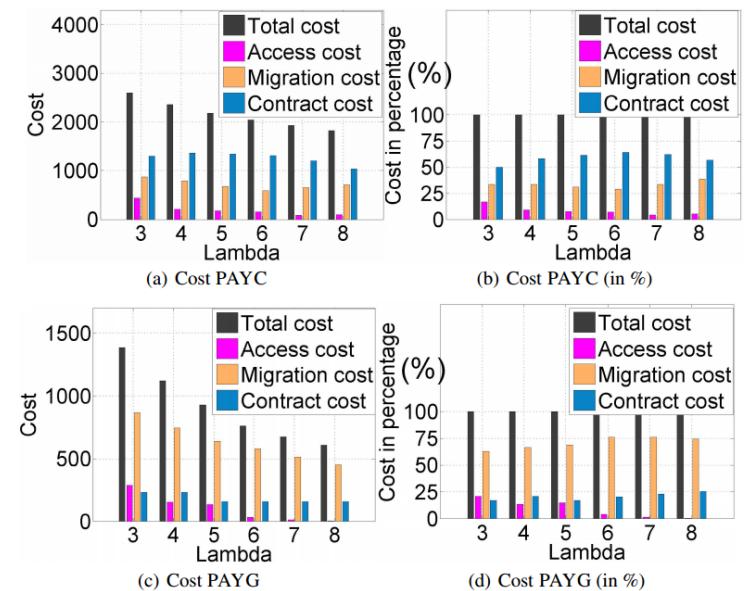
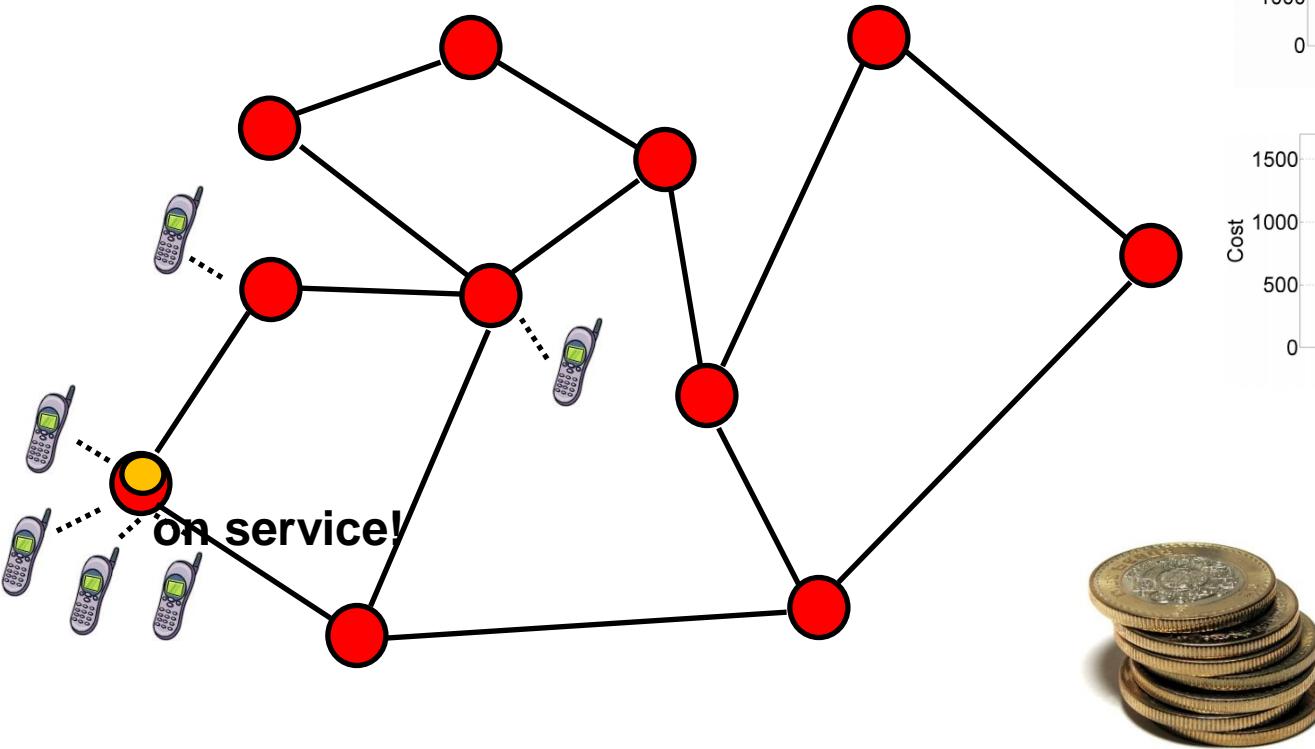


# Extension: Economical Aspects.

Hu et al.: ICDCN 2013

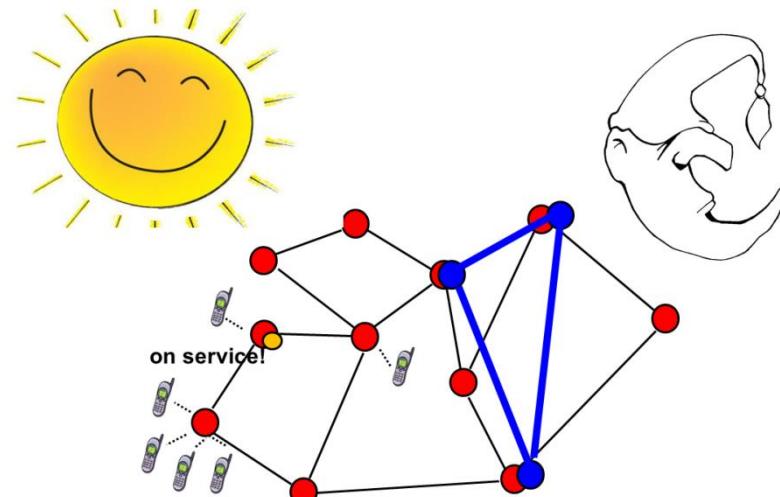
How to price resources?

Pay-as-you-go vs Pay-as-you-come?



# Migration of Entire CloudNets.

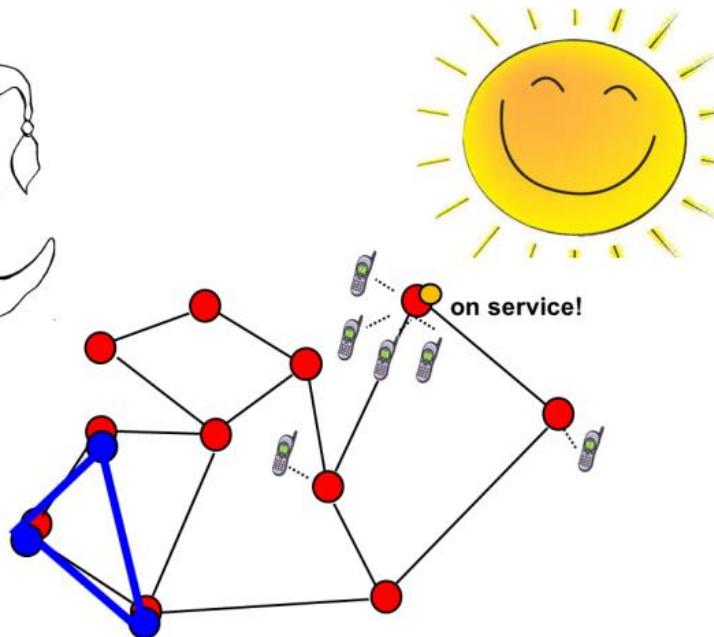
INFOCOM Demo 2013



2 pm in Europe



**Latency-resource  
tradeoffs: move-with-the-  
sun vs move-with-the-  
moon?**



2 pm in Japan

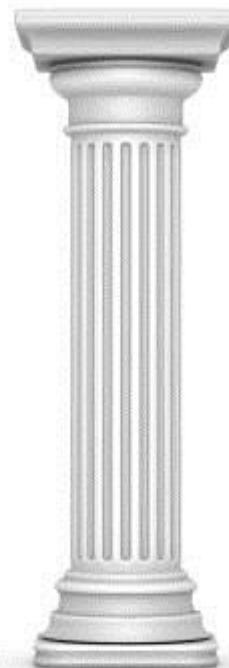


# Security of Embedding.

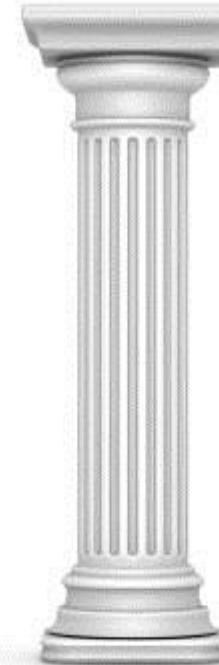
**Access Control  
and Embedding**



**Service Migration**

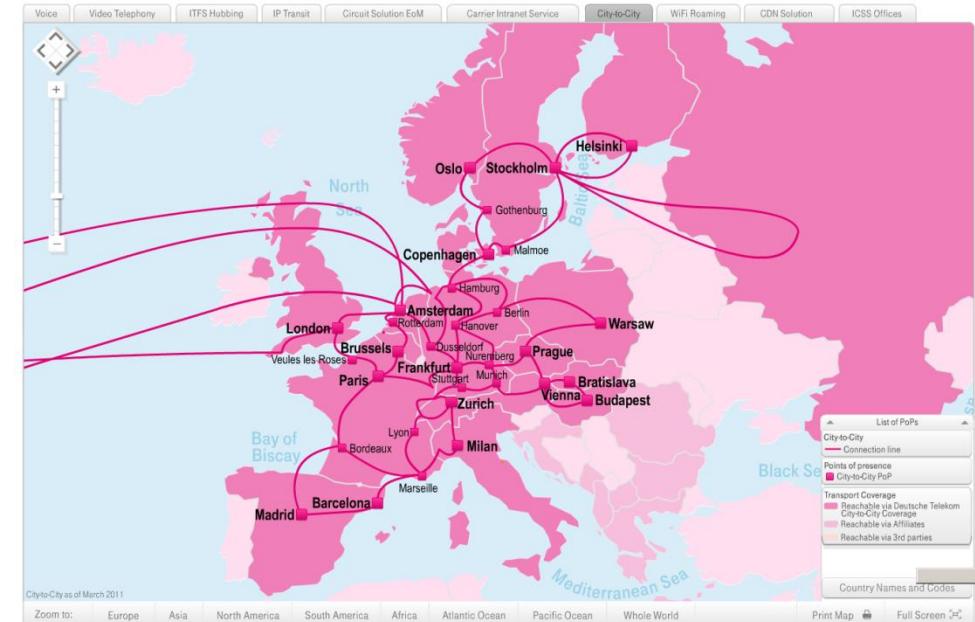
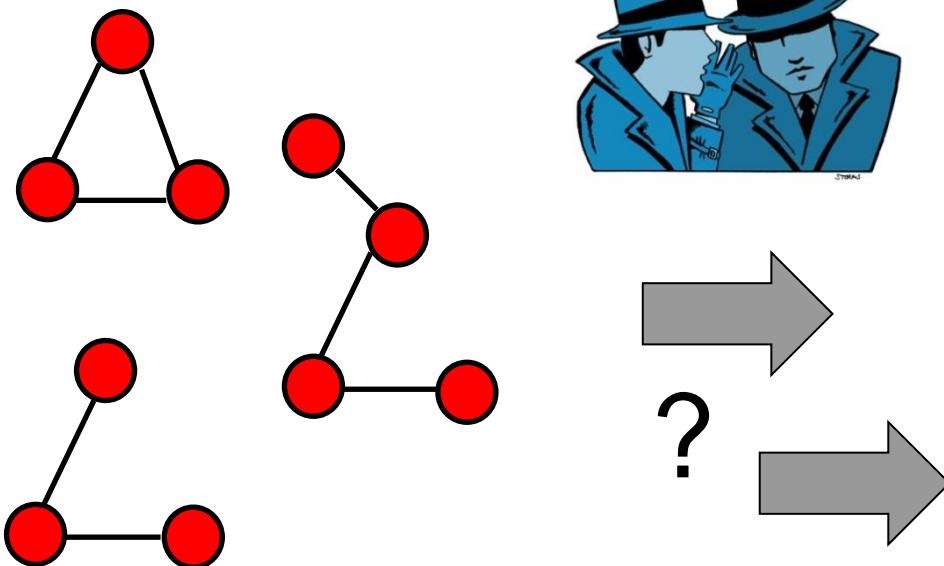


**Security Issues**



# Security Issues.

- Are CloudNet embeddings a threat for ISPs?
- Do embeddings leak information about infrastructure?

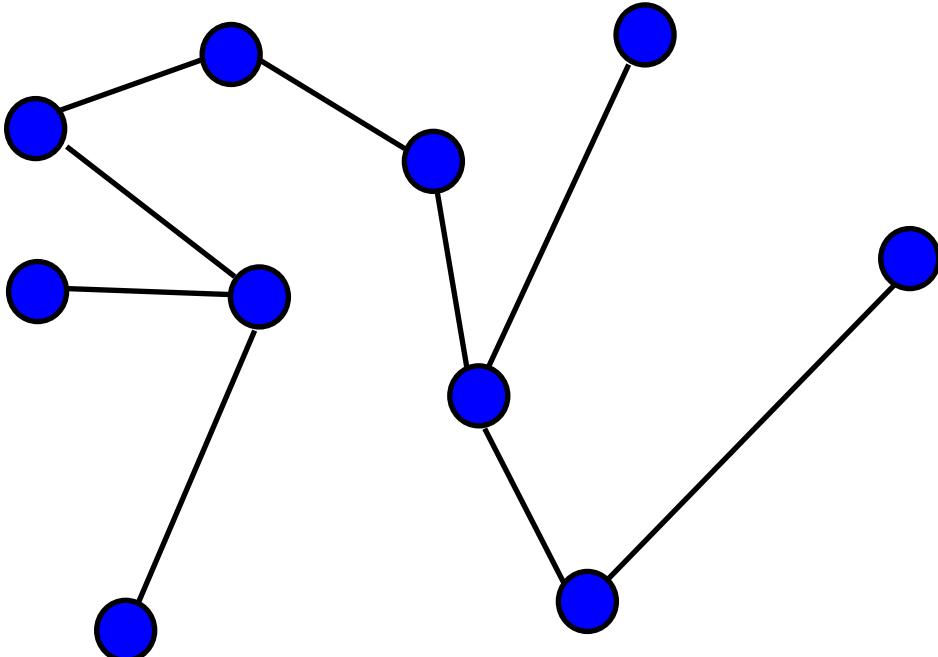
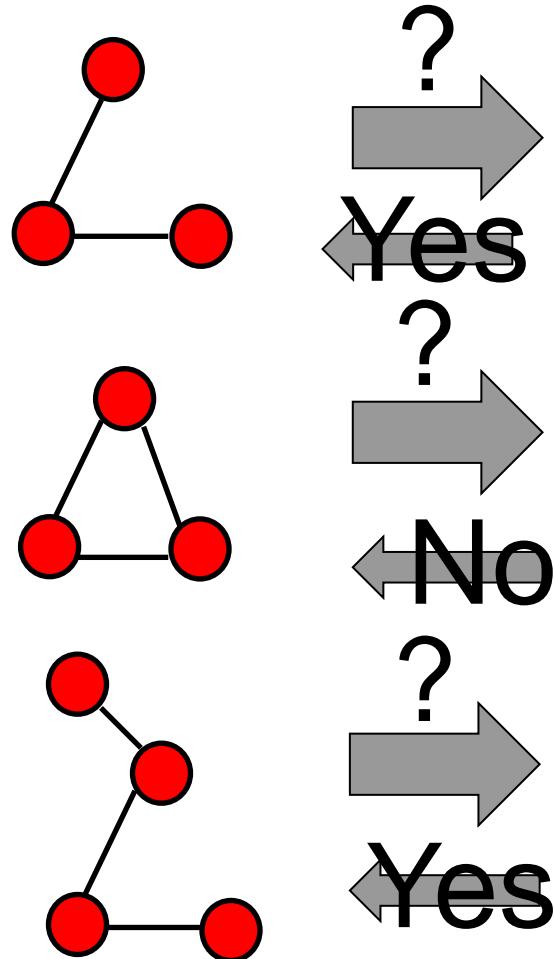


# Request Complexity.

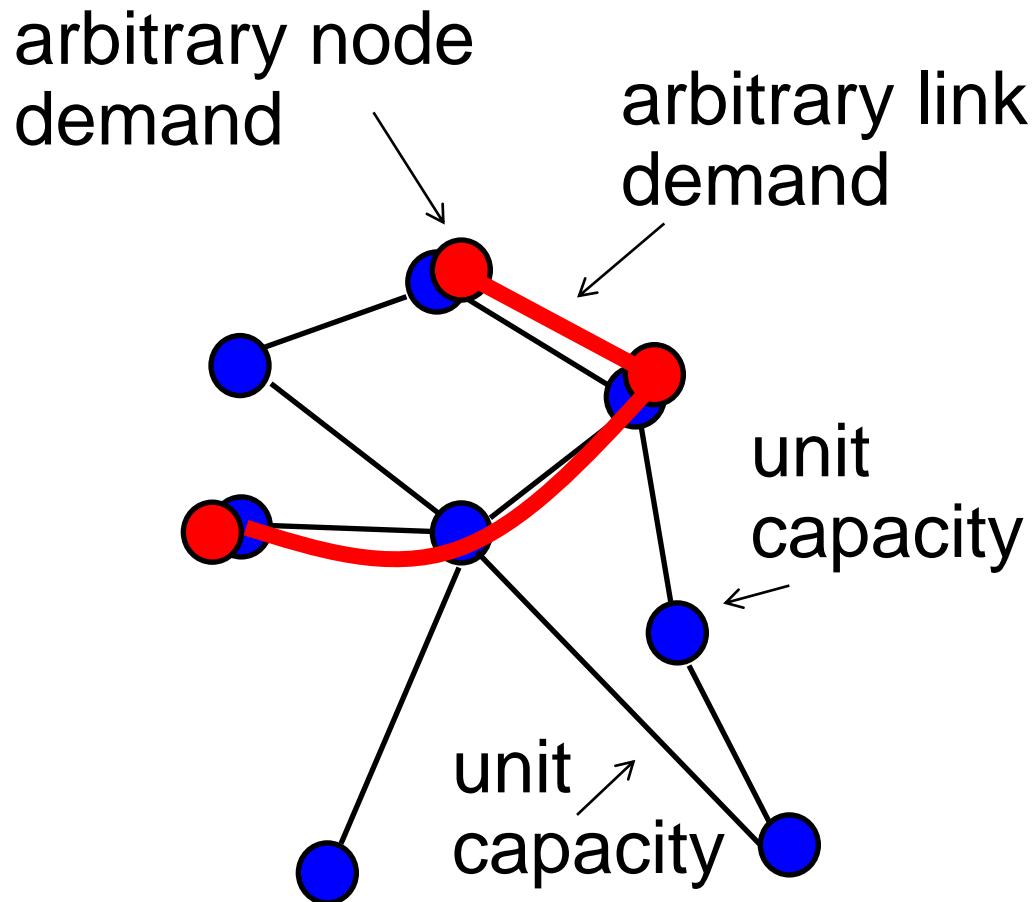
Are CloudNet embeddings  
a threat for ISPs?

## Request Complexity

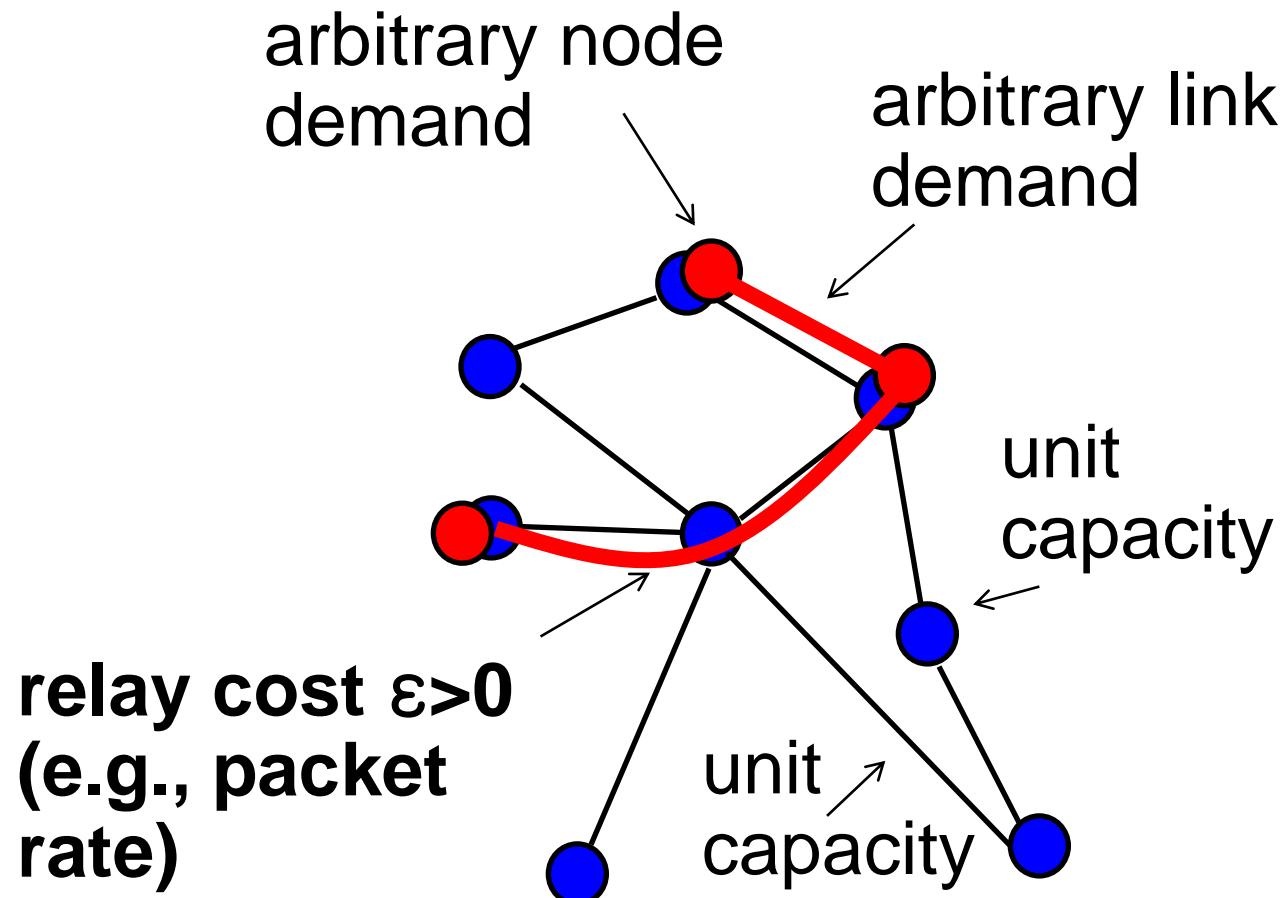
How many embeddings needed to  
fully reveal topology?



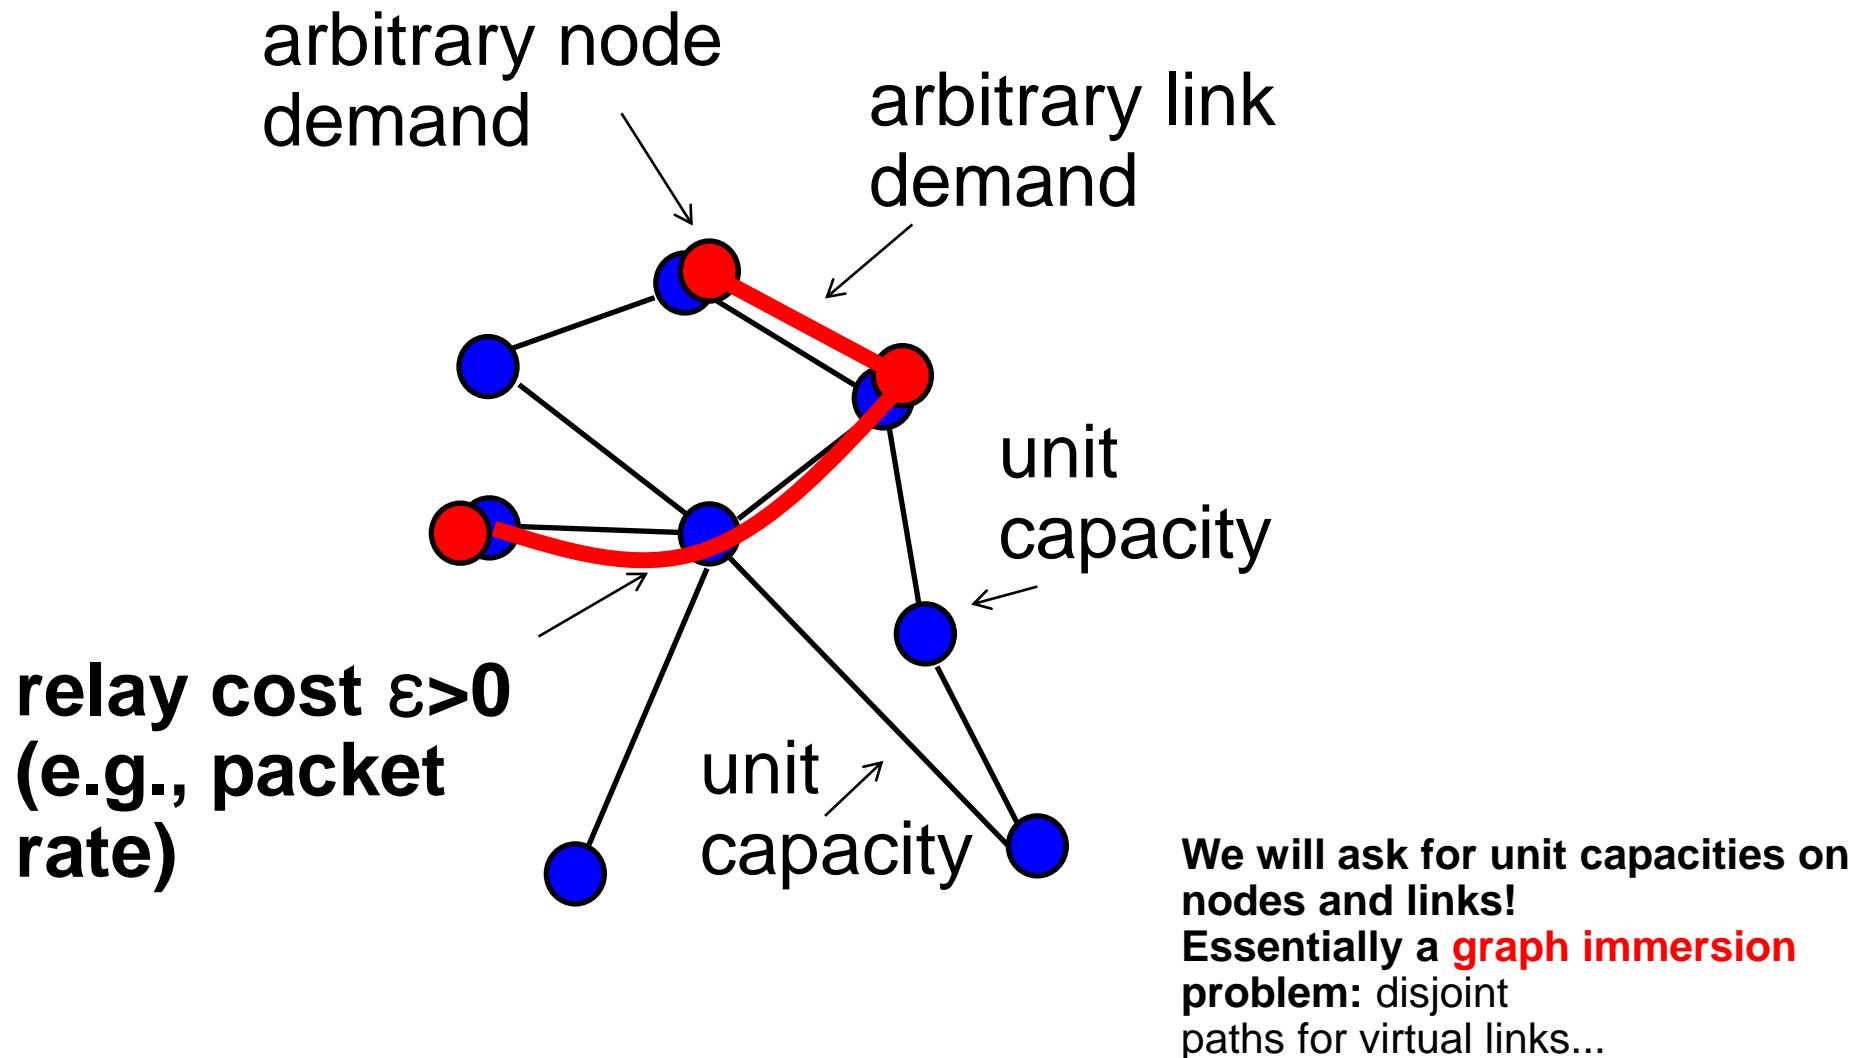
# Embedding Model.



# Embedding Model.

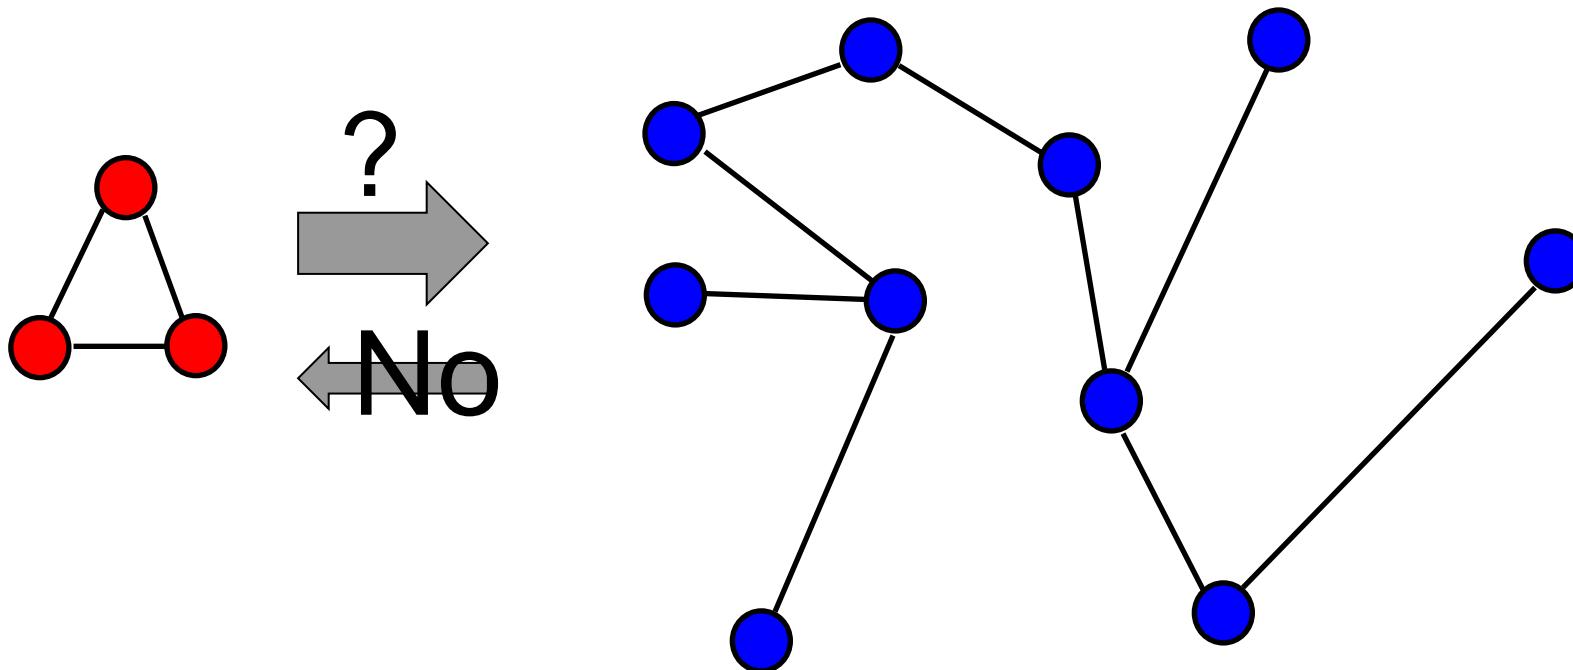


# Embedding Model.



# Some Properties Simple...

«Is the network 2-connected?»

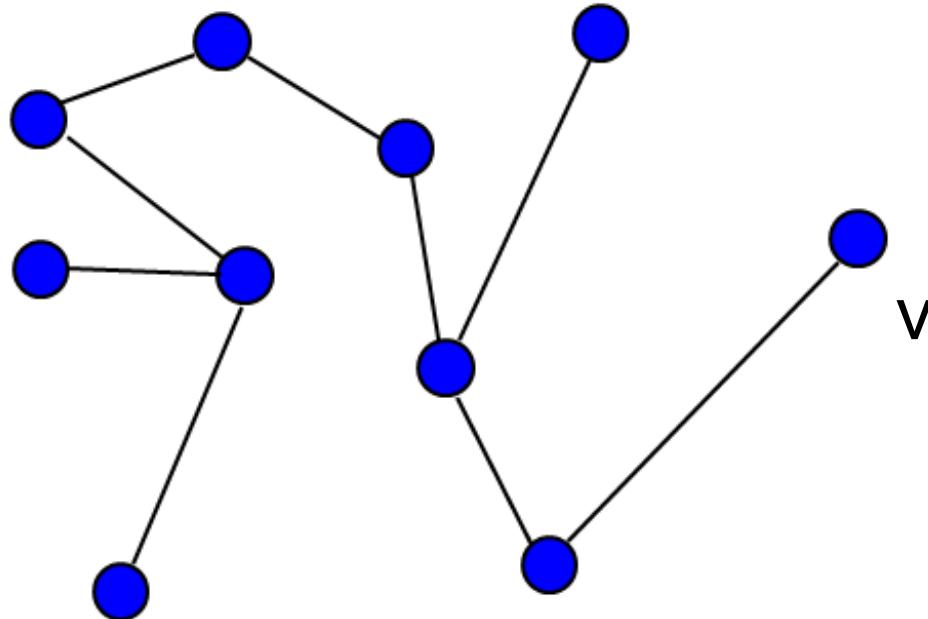


# Example: Tree.

## How to discover a tree?

Graph growing:

1. Test whether triangle fits? (loop-free)
2. Try to add neighbors to node as long as possible, then continue with other node

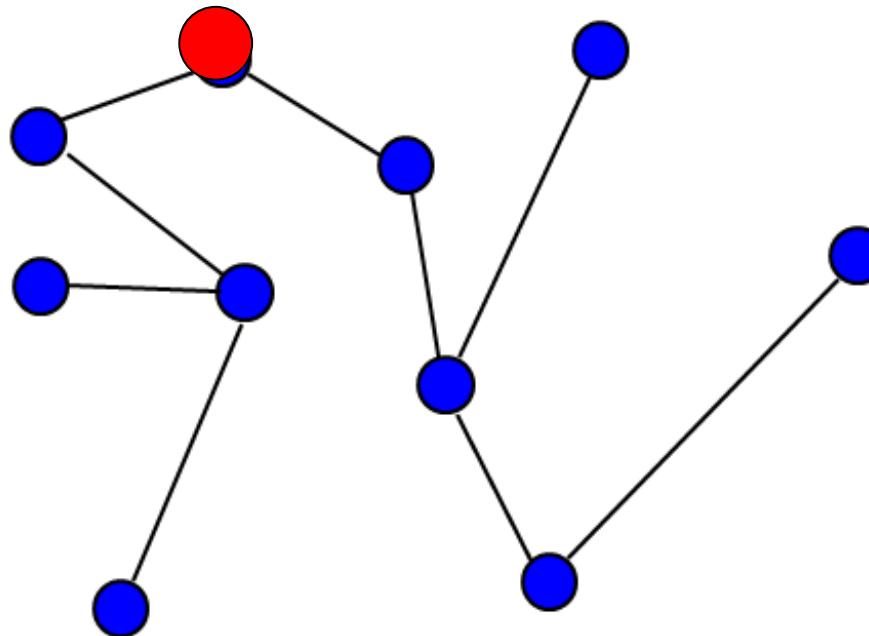


# Example: Tree.

## How to discover a tree?

Graph growing:

1. Test whether triangle fits? (loop-free)
2. Try to add neighbors to node as long as possible, then continue with other node

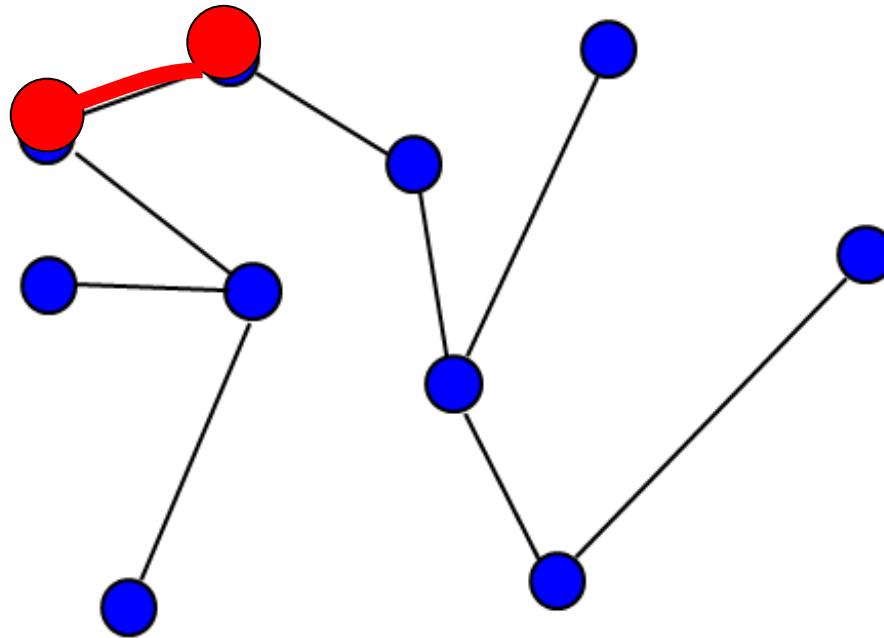


## Example: Tree.

# How to discover a tree?

## Graph growing:

1. Test whether triangle fits? (loop-free)
  2. Try to add neighbors to node as long as possible, then continue with other node

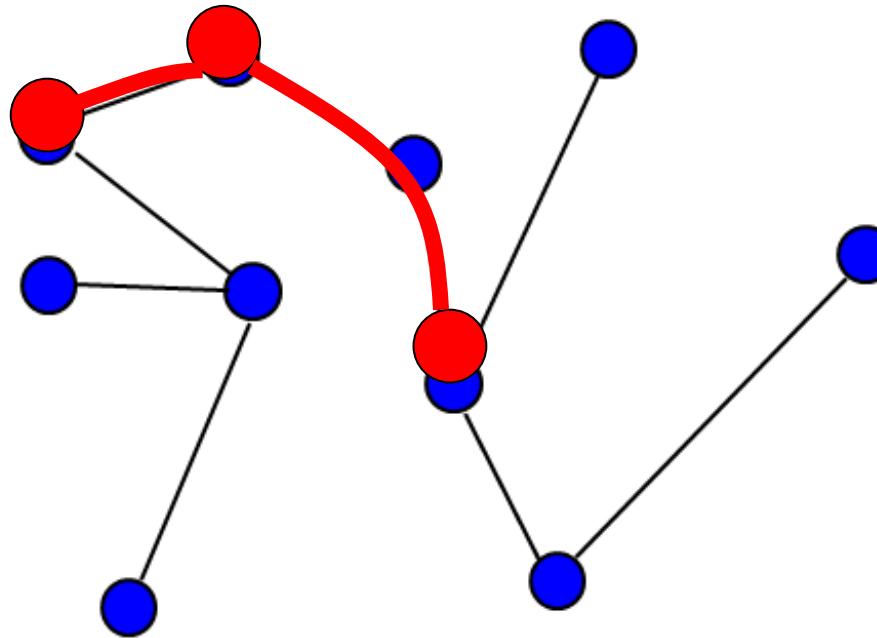


# Example: Tree.

## How to discover a tree?

Graph growing:

1. Test whether triangle fits? (loop-free)
2. Try to add neighbors to node as long as possible, then continue with other node

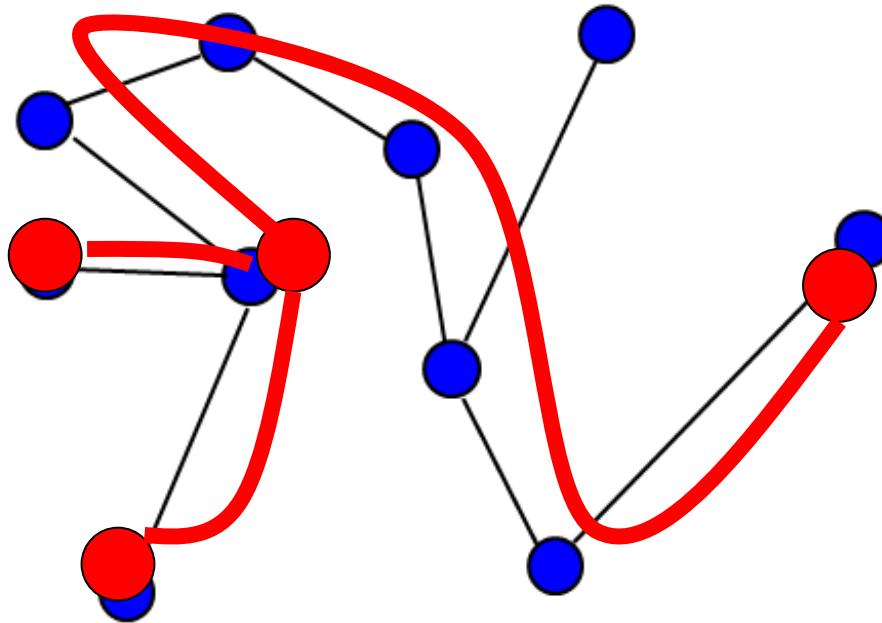


## Example: Tree.

# How to discover a tree?

## Graph growing:

1. Test whether triangle fits? (loop-free)
  2. Try to add neighbors to node as long as possible, then continue with other node

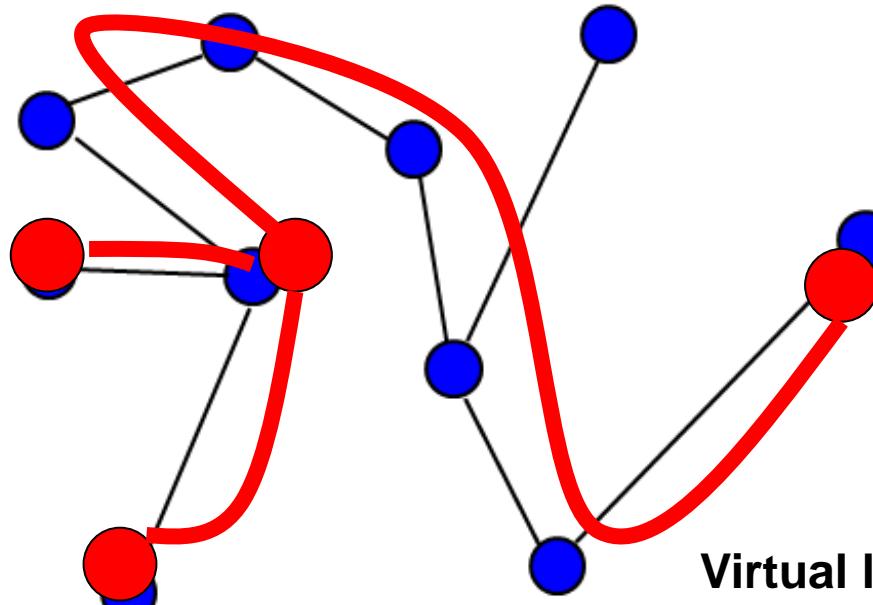


# Example: Tree.

How to discover a tree?

Graph growing:

1. Test whether triangle fits? (loop-free)
2. Try to add neighbors to node as long as possible, then continue with other node



Virtual links may be embedded over multiple physical links!



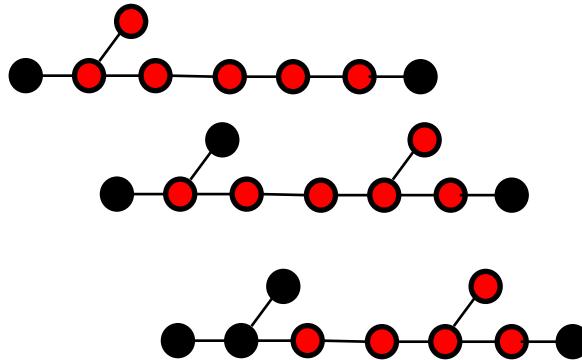
# Tree Solution: Graph Growing.

## TREE ALGORITHM: line strategy

1. Binary search on longest path («anchor»):

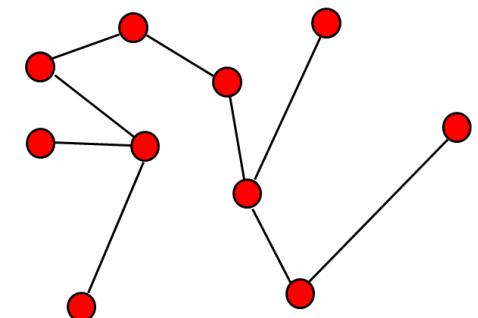


2. Last and first node *explored*, explore «branches» at pending nodes



### Analysis: Amortized analysis on links:

Per discovered physical link at most one query, plus at most one per physical node (no incident links).



---

#### Algorithm 1 Tree Discovery: TREE

---

```
1:  $G := \{\{v\}, \emptyset\}$  /* current request graph */
2:  $\mathcal{P} := \{v\}$  /* pending set of unexplored nodes*/
3: while  $\mathcal{P} \neq \emptyset$  do
4:   choose  $v \in \mathcal{P}$ ,  $S := \text{exploreSequence}(v)$ 
5:   if  $S \neq \emptyset$  then
6:      $G := G \cup S$ , add all nodes of  $S$  to  $\mathcal{P}$ 
7:   else
8:     remove  $v$  from  $\mathcal{P}$ 
```

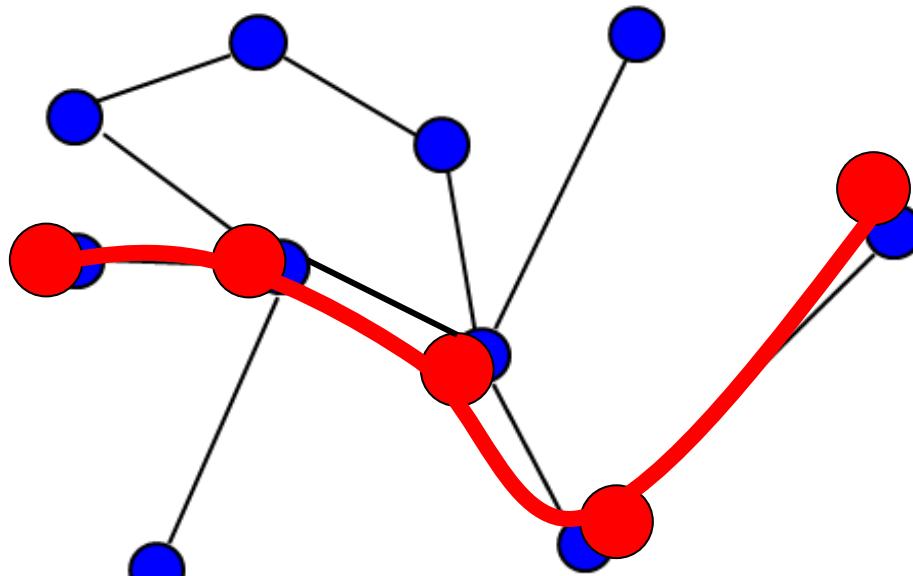
#### *exploreSequence(v)*

```
1:  $S := \emptyset$ 
2: if  $\text{request}(GvC, H)$  then
3:   find max  $j$  s.t.  $GvC^j \mapsto H$  (binary search)
4:    $S := C^j$ 
5: return  $S$ 
```

---



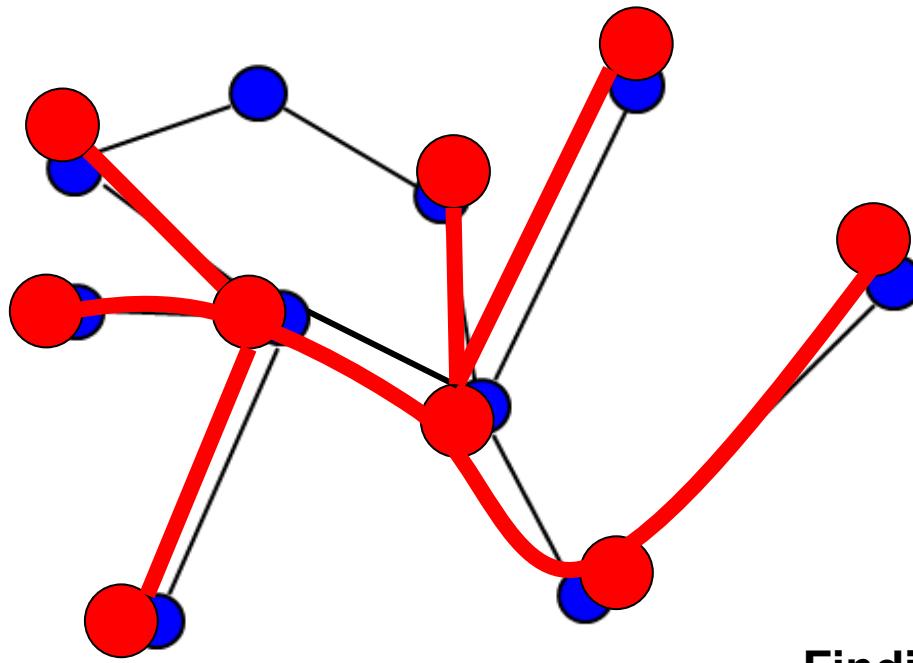
# Greedy Graph Growing on General Graphs? (1)



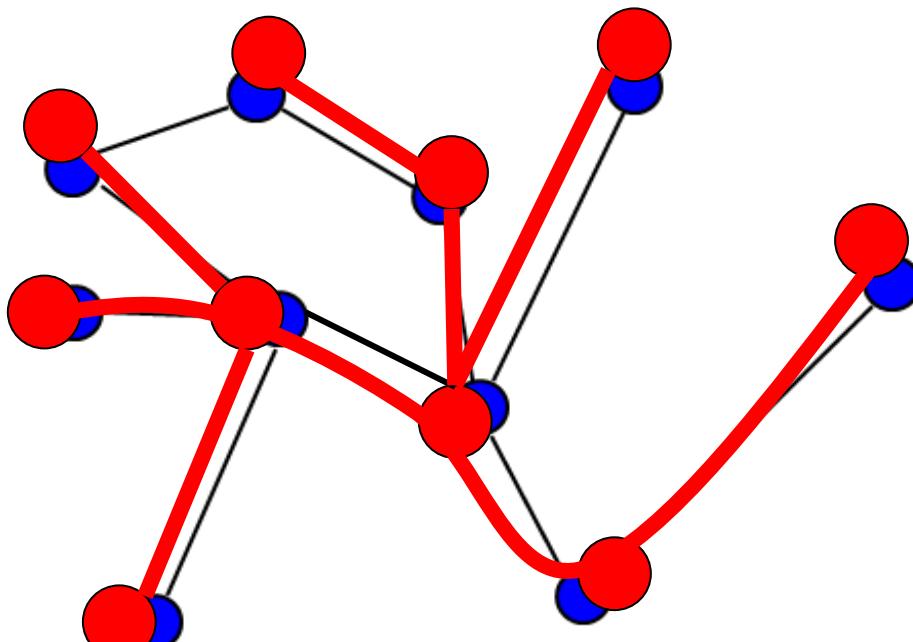
Finding path...



# Greedy Graph Growing on General Graphs? (1)



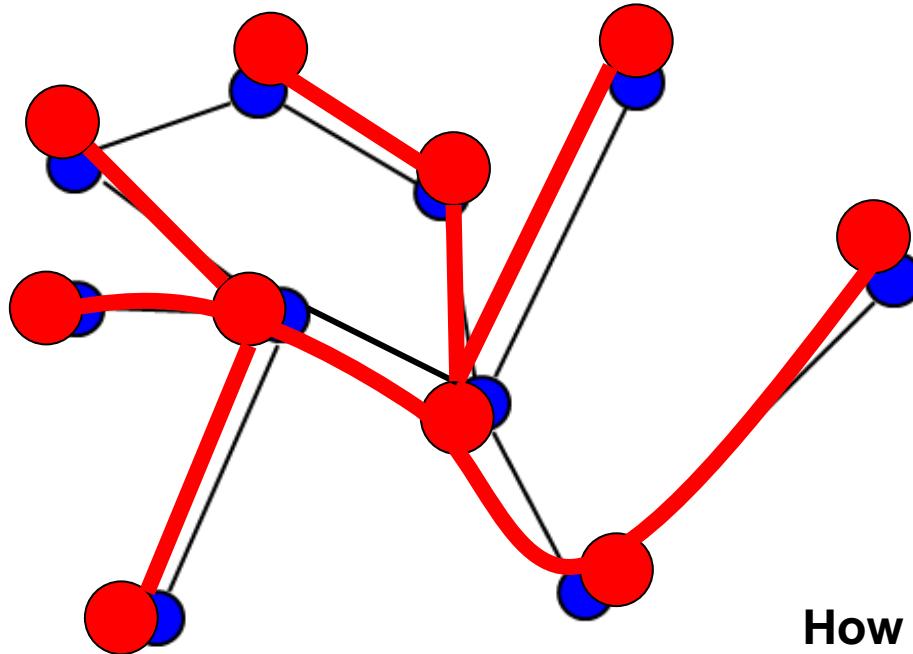
## Greedy Graph Growing on General Graphs? (1)



## Finding more neighbors...



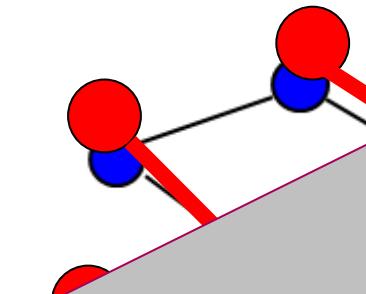
# Greedy Graph Growing on General Graphs? (1)



**How to close the gap?  
Adding connections between  
existing CloudNet nodes is  
expensive: try all pairs!**



# Greedy Graph Growing on General Graphs? (1)



Take-aways:

(1) Allocate resources on all links of highly connected components first: finding these links later is expensive.

(2) In particular, if graph X can be embedded on Y, try to embed Y first!

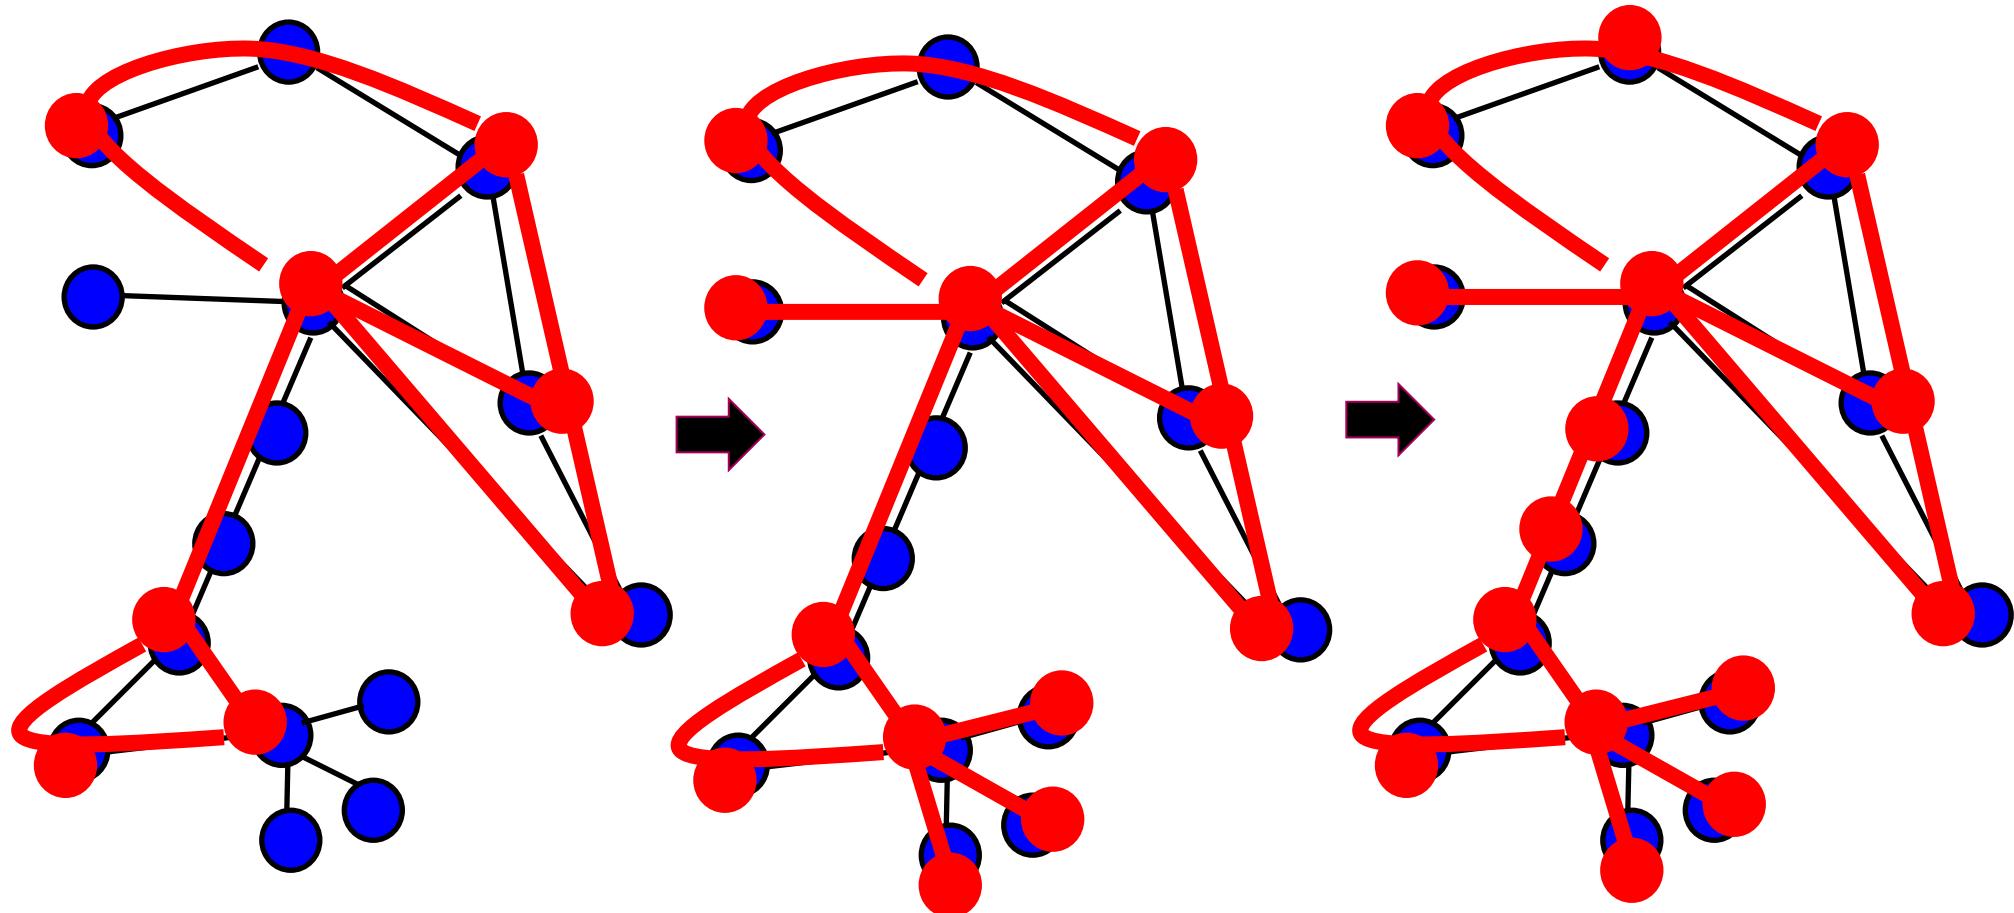
**Close the gap?  
Establishing connections between  
existing CloudNet nodes is  
expensive: try all pairs!**



# Greedy Graph Growing on General Graphs? (2)

**Simple solution:** First try to find the «knitting»!

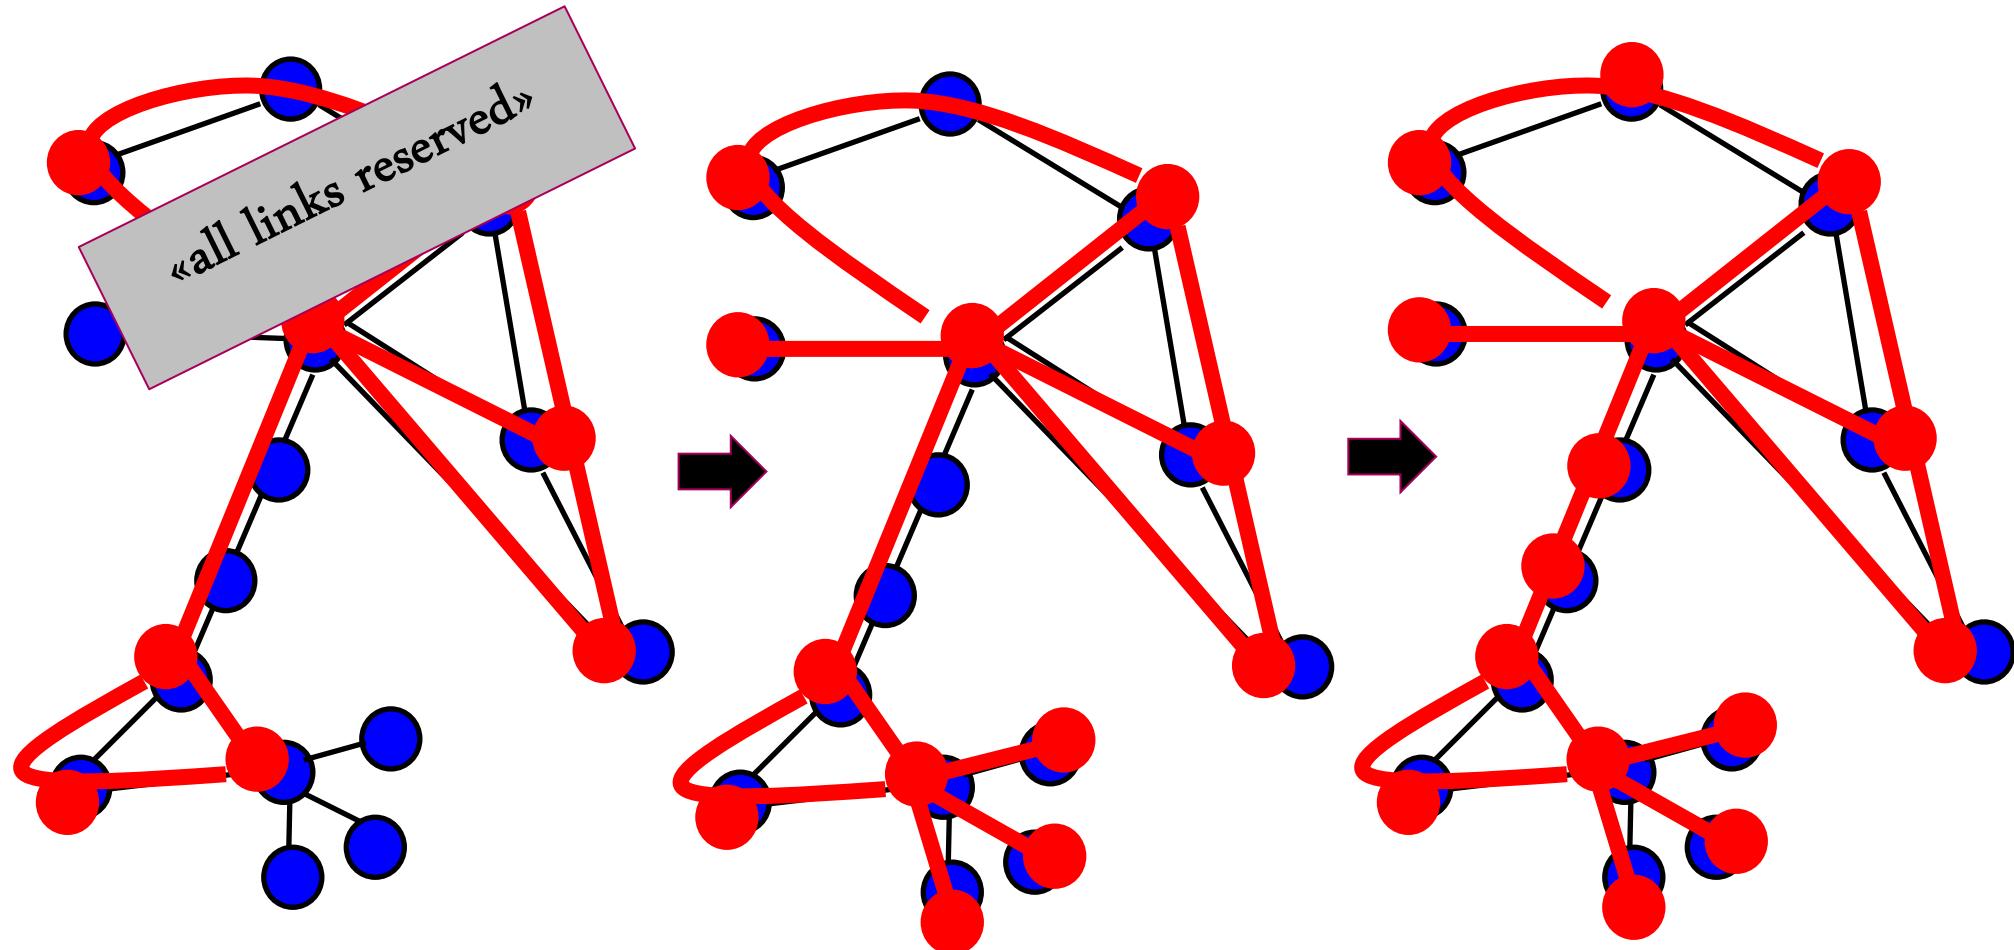
- The «two-or-more» connected components
- Later «expand nodes» and «expand edges»



# Greedy Graph Growing on General Graphs? (2)

**Simple solution:** First try to find the «knitting»!

- The «two-or-more» connected components
- Later «expand nodes» and «expand edges»

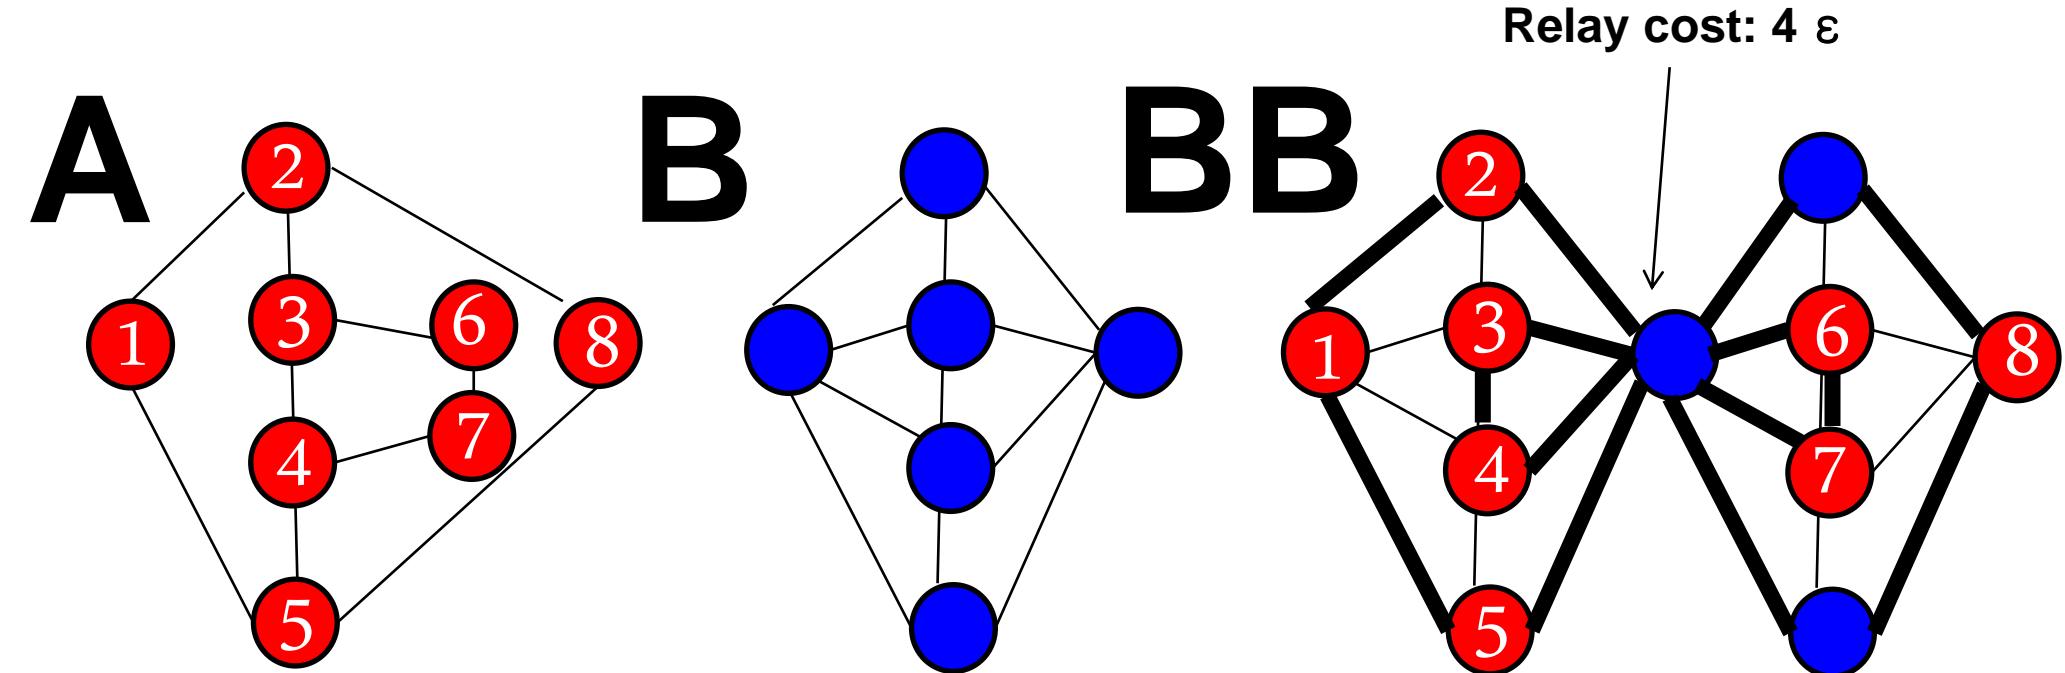


# Greedy Graph Growing on General Graphs? (3)

Idea: Ask graph «motif» only if it's guaranteed that it cannot be embedded over a more highly connected subgraph! (And connectivity has to be added later.)

Careful: What goes first also depends on entire motif sequences!

- A cannot be embedded into B and
- B cannot be embedded into A
- But A can be embedded into BB!



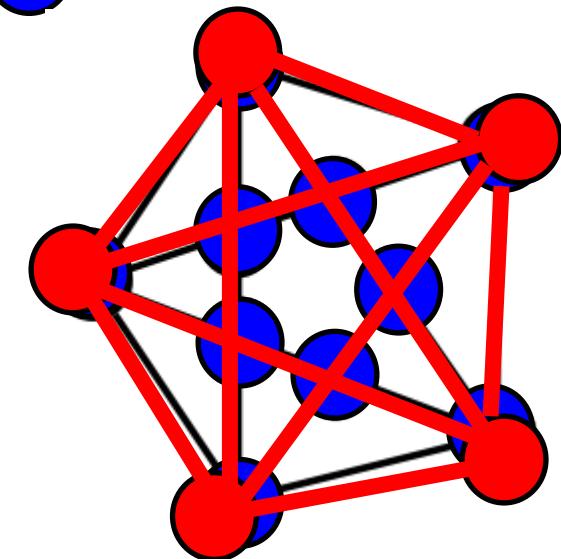
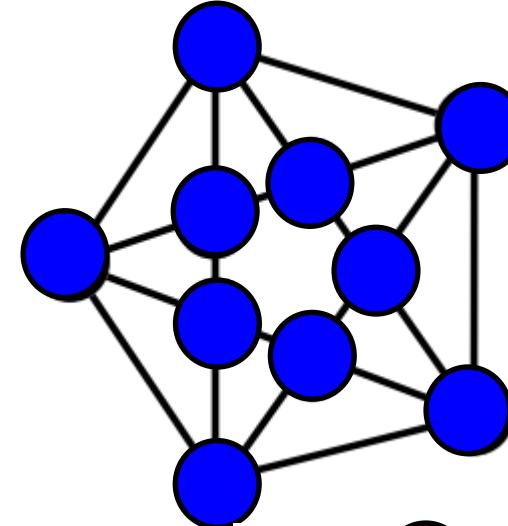
# Remark.

## Minor vs embedding:

Even with unit link capacity, for small epsilon, graph A may be embeddable ( $\rightarrow$ ) into graph B although A is not a minor of B!

### Graph Minor

Graph A is a minor of B if A can be obtained from B by (1) deleting nodes, (2) deleting edges, or (3) contracting two nodes along edges.



Planar graph (and hence K5-minor free):  
But K5 can be embedded here!



# Dictionary Attack: Expansion Framework.

## Motif

Basic “knittings” of the graph.

## Dictionary

Define an **order** on motif sequences:  
Constraints on which sequence to ask first  
in order not to overlook a part of the  
topology. (E.g., by embedding links across  
multiple hops.)

## Poset

Poset = partially ordered set

- (1) Reflexive:  $G \rightarrow G$
- (2) Transitive:  $G \rightarrow G'$  and  $G' \rightarrow G''$ , then  $G \rightarrow G''$
- (3) Antisymmetric:  $G \rightarrow G'$  and  $G' \rightarrow G$  implies  $G=G'$  (isomorphic)

## Framework

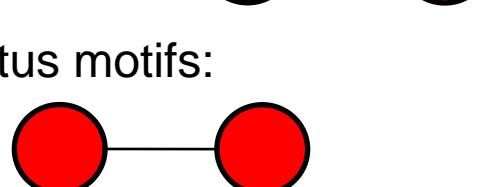
**Explore branches according  
to dictionary order,  
exploiting poset property.**

## Examples

Tree motifs:

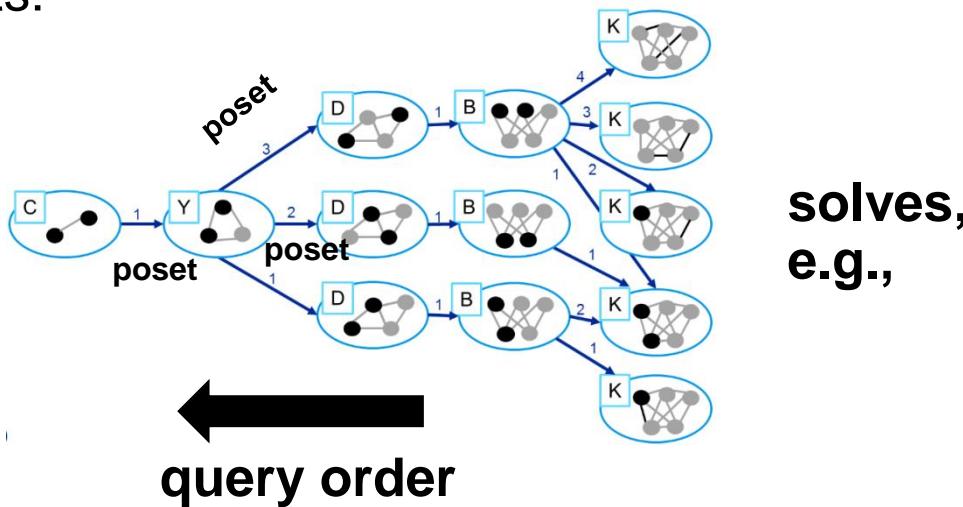


Cactus motifs:

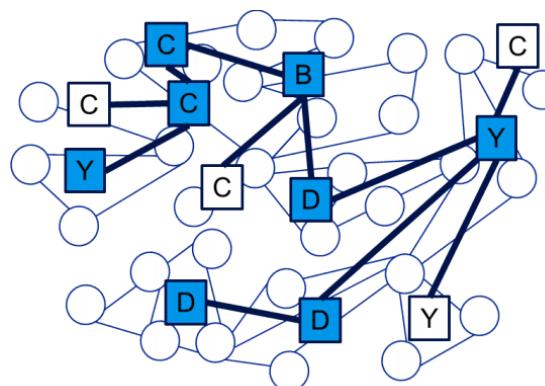
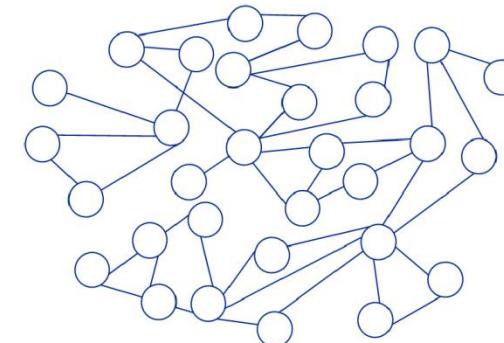


# Dictionary Attack: Expansion Framework.

**Dictionary dag** (for chain C, cycle Y, diamond D, ...) with attachment points:



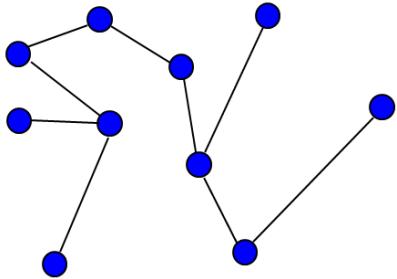
solves,  
e.g.,



**Complexity:**  
Depends on dictionary depth and number of attachment points



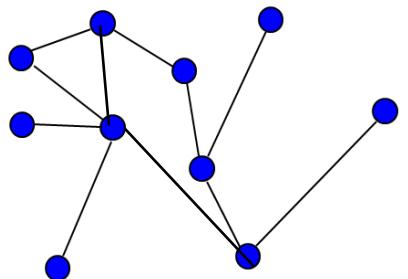
# Overview of Results.



## Tree

Can be explored in  $O(n)$  requests. This is optimal!

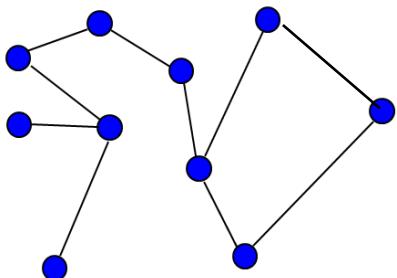
Lower bound: via number of possible trees and **binary** information.



## General Graph

Can be explored in  $O(n^2)$  requests. This is optimal!

Idea: Make **spanning tree** and then try all edges.  
(Edges directly does not work!)



## Cactus Graph

Can be explored in  $O(n)$  requests. This is optimal!

Via «graph motifs»!  
A general framework  
exploiting **poset relation**.



# Overview of Results.

---

**Algorithm 2** Cactus Discovery: CAC
 

---

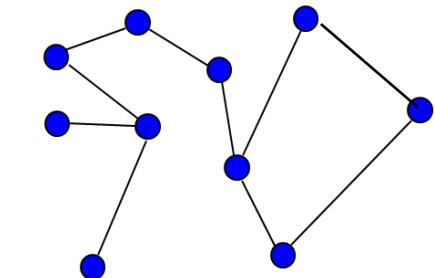
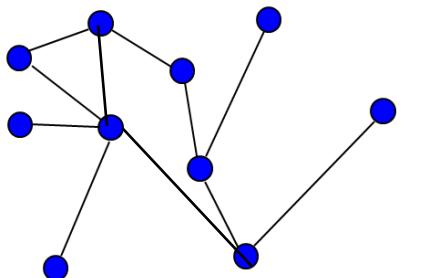
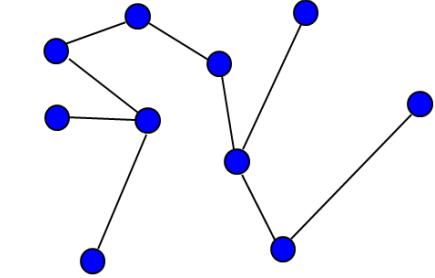
```

1:  $G := \{\{v\}, \emptyset\}$  /* current request graph */
2:  $\mathcal{P} := \{v\}$  /* pending set of unexplored nodes*/
3: while  $\mathcal{P} \neq \emptyset$  do
4:   choose  $v \in \mathcal{P}$ ,  $S := exploreSequence(v)$ 
5:   if  $S \neq \emptyset$  then
6:      $G := G \cup S$ , add all nodes of  $S$  to  $\mathcal{P}$ 
7:     for all  $e \in S$  do  $edgeExpansion(e)$ 
8:   else
9:     remove  $v$  from  $\mathcal{P}$ 

exploreSequence( $v$ )
1:  $S := \emptyset$ 
2: if  $GvYCY \mapsto H$  then
3:   find max  $j$  s.t.  $GvY^jCY \mapsto H$ 
4:    $S := Y^jCY$ ,  $P' := \{C\}$ 
5:   while  $P' \neq \emptyset$  do
6:     for all  $C_i \in P'$  do
7:        $A := prefix(C_i, S)$ ,  $B := postfix(C_i, S)$ ;
8:       if  $GvACYCB \mapsto H$  then
9:         find max  $j, k$  s.t.  $GvAC(Y^jC)^kB \mapsto H$ 
10:        for  $l := 1, \dots, k$  do
11:           $P'' := P'' \cup \{C_l\}$ 
12:           $S := AC(Y^jC)^kB$ 
13:         $P' := P'', P'' := \emptyset$ 
14: if request( $GvSY, H$ ) then
15:   find max  $j$  s.t.  $GvSY^j \mapsto H$ 
16:    $S := SY^j$ 
17: if request( $GvSC, H$ ) then
18:    $S := SC$ 
19: return  $S$ 

edgeExpansion( $e$ )
1: let  $u, v$  be the endpoints of edge  $e$ , remove  $e$  from  $G$ 
2: find max  $j$  s.t.  $GvC^j u \mapsto H$ 
3:  $G := GvC^j u$ , add newly discovered nodes to  $\mathcal{P}$ 
  
```

---



lower bound: via number of possible trees and **binary** information.

Idea: Make **spanning tree** and then try all edges.  
(Edges directly does not work!)

Via «graph motifs»!  
A general framework exploiting **poset relation**.



# Dictionary Attacks: Expand Framework.

## Motif

Basic “knitting”

## Poset

Partially ordered relation fulfills  
symmetry, transitivity

## Framework

Explore branch-and-bound  
to dictionary  
exploiting poset

### Algorithm 5 Motif Graph Discovery DICT

```
1:  $H' := \{\{v\}, \emptyset\}$  /* current request graph */
2:  $\mathcal{P} := \{v\}$  /* pending set of unexplored nodes*/
3: while  $\mathcal{P} \neq \emptyset$  do
4:   choose  $v \in \mathcal{P}$ ,  $T := \text{find\_motif\_sequence}(v, \emptyset, \emptyset)$ 
5:   if  $T \neq \emptyset$  then
6:      $H' := H' v T$ , add all nodes of  $T$  to  $\mathcal{P}$ 
7:     for all  $e \in T$  do  $\text{edgeExpansion}(e)$ 
8:   else
9:     remove  $v$  from  $\mathcal{P}$ 
```

#### $\text{find\_motif\_sequence}(v, T^<, T^>)$

```
1: find max  $i, j, \text{BF}, \text{AF}$  s.t.  

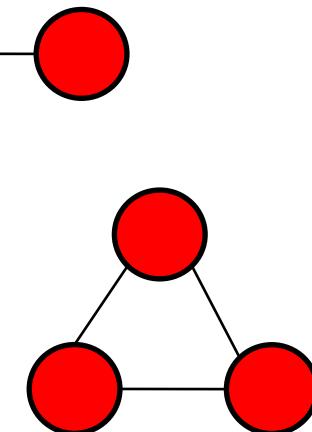
 $H' v (T^<) \text{BF } (D[i])^j \text{ AF } (T^>) \mapsto H$  where  

 $\text{BF}, \text{AF} \in \{\emptyset, C\}^2$  /*issue requests*/
2: if  $(i, j, \text{BF}, \text{AF}) = (0, 0, C, \emptyset)$  then
3:   return  $T^< C T^>$ 
4: if  $\text{BF} = C$  then
5:    $\text{BF} = \text{find\_motif\_sequence}(v, T^<, (D[i])^j \text{ AF } T^>)$ 
6: if  $\text{AF} = C$  then
7:    $\text{AF} = \text{find\_motif\_sequence}(v, T^< \text{BF } (D[i])^j, T^>)$ 
8: return  $\text{BF } (D[i])^j \text{ AF}$ 
```

#### $\text{edge\_expansion}(e)$

```
1: let  $u, v$  be the endpoints of edge  $e$ , remove  $e$  from  $H'$ 
2: find max  $j$  s.t.  $H' v C^j u \mapsto H$  /*issue requests*/
3:  $H' := H' v C^j u$ , add newly discovered nodes to  $\mathcal{P}$ 
```

equences:  
ence to ask first  
part of the  
ing links across



# Application: BigFoot EU Project.

- “BigData is the answer: what was the question?”
  - How to **interact** with fast growing data volumes?
  - Which technology to obtain **insights**?
- BigFoot in three points:
  - **Automatic and self-tuned deployments** through virtualization
  - Cross-layer **optimization**
  - Data **interaction made easy**
- Applications:
  - **SMART-GRID DATA**
    - Billing & revenue assurance
    - Customer segmentation for service personalization
    - Pattern analysis for infrastructure provisioning
  - **ICT SECURITY DATA**
    - Attack attribution
    - Multi-feature classification



<http://www.bigfootproject.eu>



# Application: BigFoot EU Project.

- “BigData is the answer: what was the question?”
  - How to **interact** with fast growing data volumes?
  - Which technology to obtain **insights**?
  
- BigFoot in three points:
  - **Automatic and self-tuned deployments** through virtualization
  - Cross-layer **optimization**
  - Data **interaction made easy**
  
- Applications:
  - **SMART-GRID DATA**
    - Billing & revenue assurance
    - Customer segmentation for service personalization
    - Pattern analysis for infrastructure provisioning
  - **ICT SECURITY DATA**
    - Attack attribution
    - Multi-feature classification

Our focus!

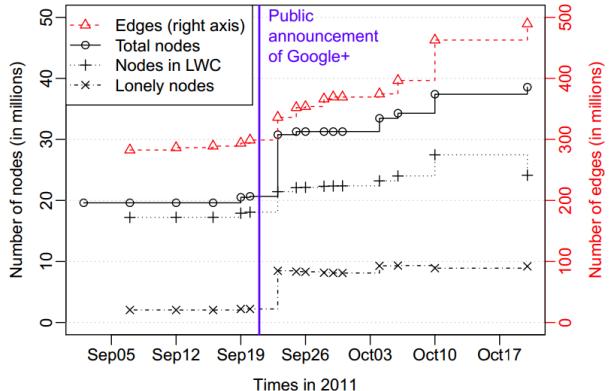


<http://www.bigfootproject.eu>

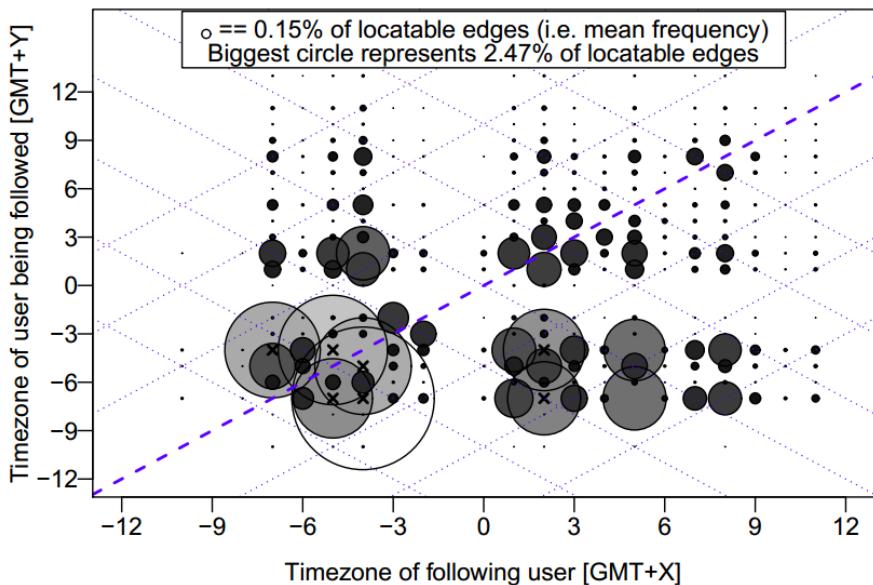
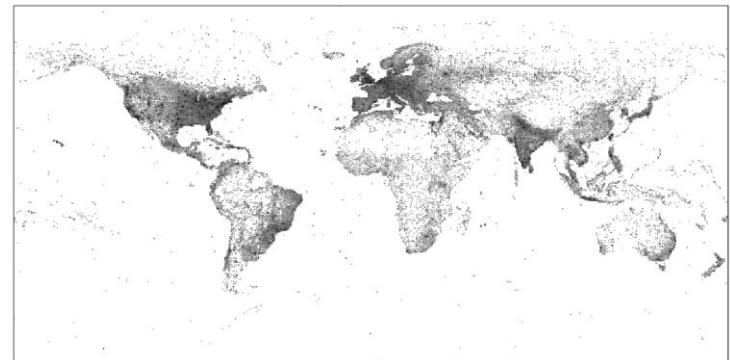


# Big Data: OSN Analysis (Google+).

Schiöberg et al.:  
ACM WebSci 2012



Now 100M+ users...  
... from all over the world!



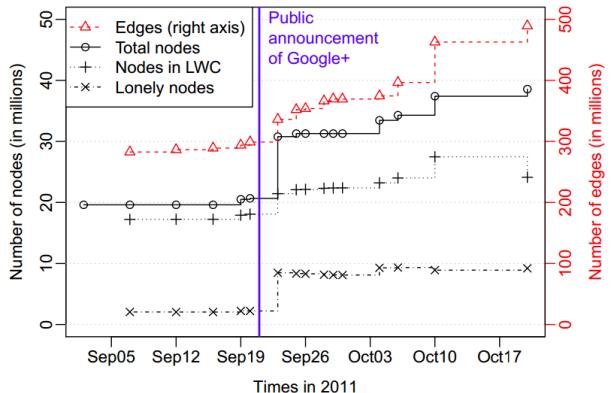
## Research:

- # of «followers» (easy)
- Which time zones follow which time zones? (easy)
- K-cores, «stars», ... (easy)
- Diameter??
- Evolution of centralities??

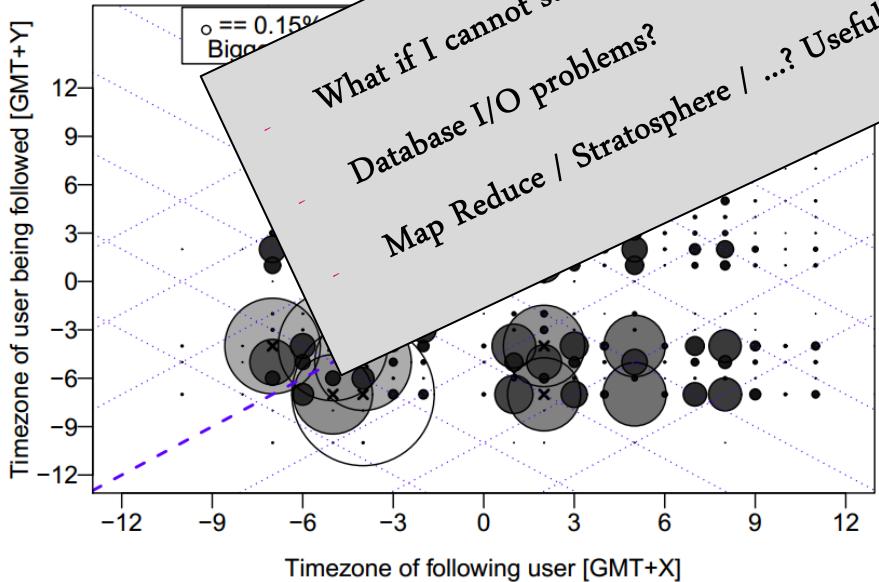
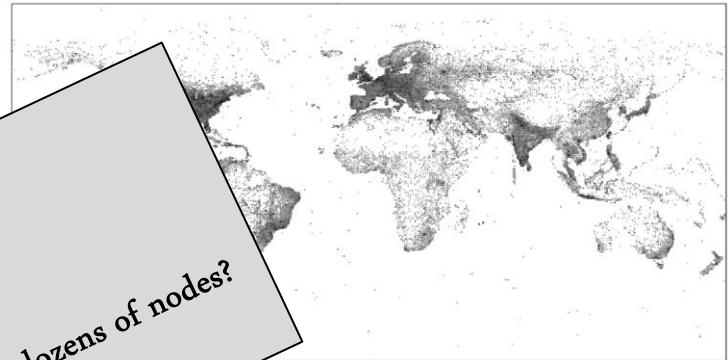


# Big Data: OSN Analysis (Google+).

Schiöberg et al.:  
ACM WebSci 2012



Now 100M+ users...  
... from all over the world



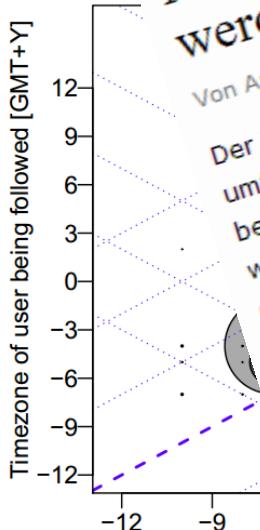
## Research:

- # of «followers» (easy)
- Which time zones follow which time zones? (easy)
- K-cores, «stars», ... (easy)
- Diameter??
- Evolution of centralities??



# Big Data: OSN Analysis (Google)

Schiöberg et al.:  
ACM WebSci 2012

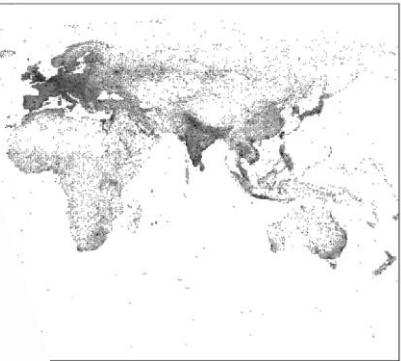


Rüstungskonzern entwickelt Software, mit der Social-Media-Daten zusammengeführt und durchsucht werden können, für die 'nationale Sicherheit'

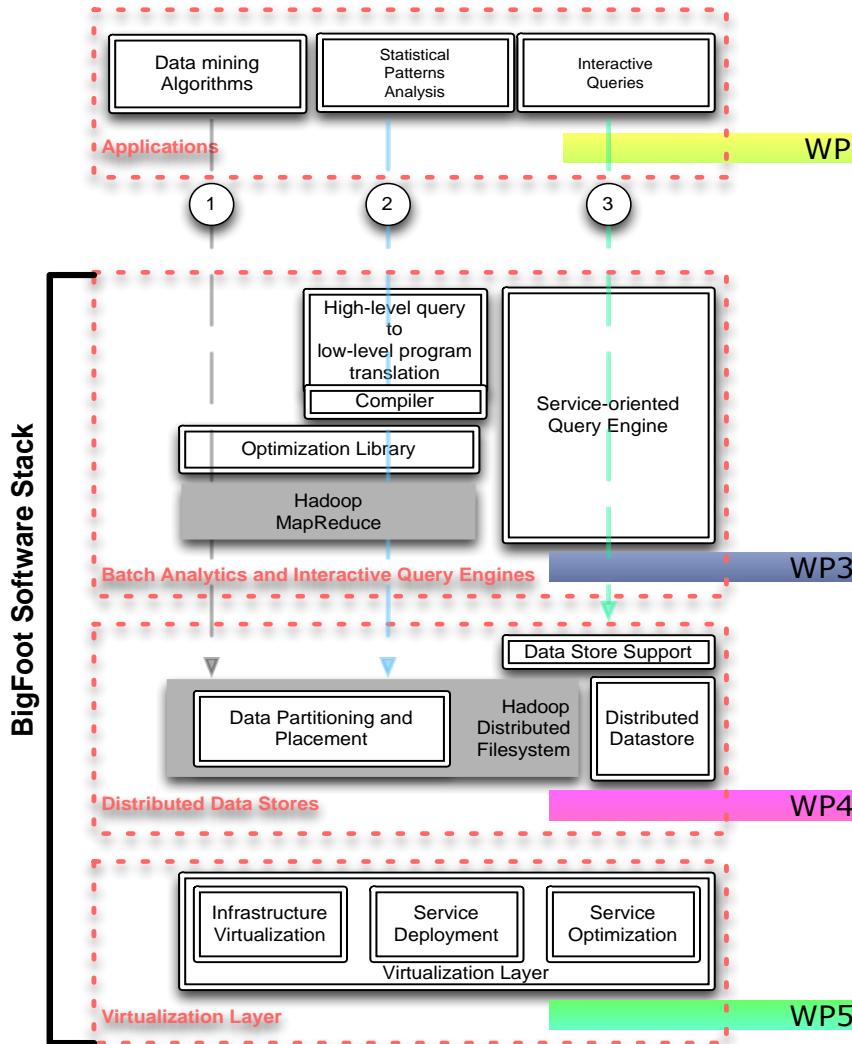
Von Andrea Jonjic | Veröffentlicht am: 11.02.2013 um 15:11h | 4 Responses

Der US-Rüstungs- und Elektronikkonzern Raytheon arbeitet laut *Guardian* an einer umfangreichen Software zur Durchleuchtung von Internetnutzern, indem Daten von beispielsweise Facebook, Twitter und Foursquare zusammengeführt und ausgewertet werden. RIOT – Rapid Information Overlay Technology – heißt die Suchmaschine, mit deren Hilfe aus Social-Media-Daten ermittelt werden soll, ob eine Person ein Risiko für die nationale Sicherheit darstellt. Der *Guardian* zeigt in einem Firmenvideo von Raytheon, wie der Konzern an einem Mitarbeiter die Fähigkeiten von RIOT demonstriert. Aus dessen Foursquare Check-ins wird eine Grafik der 10 Orte erstellt, an denen er am häufigsten eincheckt und schnell lässt sich erkennen, dass er meist montags um 6:00 Uhr im Fitnessstudio ist. Wer ihn oder seinen Laptop also mal erwischen wolle, könne dann ja dort vorbeischauen, so der 'principal investigator' von Raytheon, Brian Urch.

Es gibt ebenfalls eine "Personensuche", die Kontakte auf Twitter und anderen Plattformen in einer Netzwerkgrafik anzeigt und so Vernetzungen von Personen sichtbar macht. Bei all diesen Daten handelt es sich zwar um öffentlich einsehbare Informationen – problematisch ist allerdings die einfache Möglichkeit verschiedene Quellen *...abilities??*



# BigFoot Architecture.



## Testbed:

- **Hadoop (maybe Stratosphere)**
- **OpenStack resource management (+ own cloud operating system)**
- **OpenFlow (link virtualization)**



# Conclusion.

- CloudNets:
  - Elastic computing and networking
  - Federated architecture
- Competitive analysis: a framework to design and prove performance of online algorithms
- Good when:
  - No reliable prediction models exist
  - No data available
  - **Worst case** guarantees matter
- Examples: online embedding and service migration
- Fully incorporated in prototype

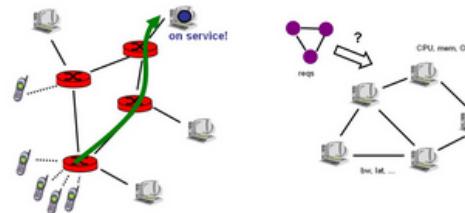


# The Project Website.

## Combining Clouds with Virtual Networking

### The CloudNet Project

Internet Network Architectures (INET)  
TU Berlin / Telekom Innovation Labs (T-Labs)  
Contact: [Stefan Schmid](#)



News  
Overview  
People  
Magazines  
Publications  
Demo  
Talks/Posters

#### News

- Watch on YouTube: migration demonstrator [video!](#)
- We are looking for students and interns with good algorithmic background to contribute to Virtu! [Contact us](#) for more details or have a look at some [open topics](#).

#### Project Overview

CloudNets are virtual networks (VNets) connecting cloud resources. The [network virtualization](#) paradigm allows to run multiple CloudNets on top of a shared physical infrastructure. These CloudNets can have different properties (provide different security or QoS guarantees, run different protocols, etc.) and can be managed independently of each other. Moreover, (parts of) a CloudNet can be migrated dynamically to locations where the service is most useful or most cost efficient (e.g., in terms of energy conservation). Depending on the circumstances and the technology, these migrations can be done live and without interrupting ongoing sessions. The flexibility of the paradigm and the decoupling of the services from the underlying resource networks has many advantages; for example, it facilitates a more efficient use of the given resources, it promises faster innovations by overcoming the ossification of today's Internet architecture, it simplifies the network management, and it can improve service performance.

We are currently developing a [prototype system](#) for this paradigm (currently based on VLANs), which raises many scientific challenges. For example, we address the problem of where to embed CloudNet requests (e.g., see [1] for online CloudNet embeddings and [2] for a general mathematical embedding program), or devise algorithms to migrate CloudNets to new locations (e.g., due to user mobility) taking into account the



# Collaborators and Publications.

## ▪ People

- **T-Labs / TU Berlin:** Anja Feldmann, Carlo Fürst, Johannes Grassler, Arne Ludwig, Matthias Rost, Gregor Schaffrath, Stefan Schmid
- **Uni Wroclaw:** Marcin Bienkowski
- **Uni Tel Aviv:** Guy Even, Moti Medina
- **NTT DoCoMo Eurolabs:** Group around Wolfgang Kellerer
- **LAAS:** Gilles Tredan
- **ABB:** Yvonne Anne Pignolet
- **IBM Research:** Johannes Schneider
- **Arizona State Uni:** Xinhui Hu, Andrea Richa

## ▪ Publications

- **Prototype:** J. Information Technology 2013, ICCCN 2012, ERCIM News 2012, SIGCOMM VISA 2009
- **Migration:** ToN 2013, INFOCOM 2013, ICDCN 2013 + Elsevier TCS Journal, Hot-ICE 2011, IPTComm 2011, SIGCOMM VISA 2010
- **Embedding:** IPDPS 2014, OPODIS 2013, CLOUDNETS 2013 (*Best Paper*), INFOCOM 2013 (Mini-Conference), CLOUDNETS 2012, 2 x ACM UCC 2012, DISC 2012, ICDCN 2012 (*Best Paper Distributed Computing Track*)



# Contact.



**Dr. Stefan Schmid**

Telekom Innovation Laboratories  
Ernst-Reuter-Platz 7, D-10587 Berlin  
E-mail: [stefan@net.t-labs.tu-berlin.de](mailto:stefan@net.t-labs.tu-berlin.de)

Project website:  
<http://www.net.t-labs.tu-berlin.de/~stefan/virtu.shtml>



Telekom **Innovation** Laboratories

Stefan Schmid