

Disconnected cooperation in resilient networks and the algorithmic challenges of local fast re-routing

Stefan Schmid @ Workshop on Distributed Algorithms on Realistic Network Models (PODC'21)



Disconnected cooperation in resilient networks and the algorithmic challenges of local fast re-routing

Stefan Schmid @ Workshop on Distributed Algorithms on Realistic Network Models (PODC'21)



Kudos to Marco Chiesa for some slides!

Communication Networks

Critical infrastructure of digital society

- Popularity of **datacentric applications**: health, business, entertainment, social networking, AI/ML, etc.
- Evident during ongoing **pandemic**: online learning, online conferences, etc.
- Much traffic especially to, from, and inside **datacenters**



Facebook datacenter

Increasingly stringent dependability requirements!

Requirements vs Reality

Entire countries disconnected...

Data Centre ▶ Networks

Google routing blunder sent Japan's Internet dark on Friday

Another big BGP blunder

By Richard Chirgwin 27 Aug 2017 at 22:35

40 □ SHARE ▾

Last Friday, someone in Google fat-thumbed a border gateway protocol (BGP) advertisement and sent Japanese Internet traffic into a black hole.

The trouble began when The Chocolate Factory "leaked" a big route table to Verizon, the result of which was traffic from Japanese giants like NTT and KDDI was sent to Google on the expectation it would be treated as transit.

... 1000s passengers stranded...

British Airways' latest Total Inability To Support Upwardness of Planes* caused by Amadeus system outage

Stuck on the ground awaiting a load sheet? Here's why

By Gareth Corfield 19 Jul 2018 at 11:16

109 □ SHARE ▾



BA flights around the world were grounded as a result of the Amadeus outage

... even 911 services affected!

Officials: Human error to blame in Minn. 911 outage

According to a press release, CenturyLink told department of public safety that human error by an employee of a third party vendor was to blame for the outage

Aug 16, 2018

Duluth News Tribune

SAINT PAUL, Minn. — The Minnesota Department of Public Safety Emergency Communication Networks division was told by its 911 provider that an Aug. 1 outage was caused by human error.

Many outages due to human error!
(Misconfigurations, not attacks...)

Even Tech-Savvy Companies Struggle



We discovered a misconfiguration on this pair of switches that caused what's called a “bridge loop” in the network.

A network change was [...] executed incorrectly [...] more “stuck” volumes and added more requests to the re-mirroring storm.



Service outage was due to a series of internal network events that corrupted router data tables.

Experienced a network connectivity issue [...] interrupted the airline's flight departures, airport processing and reservations systems

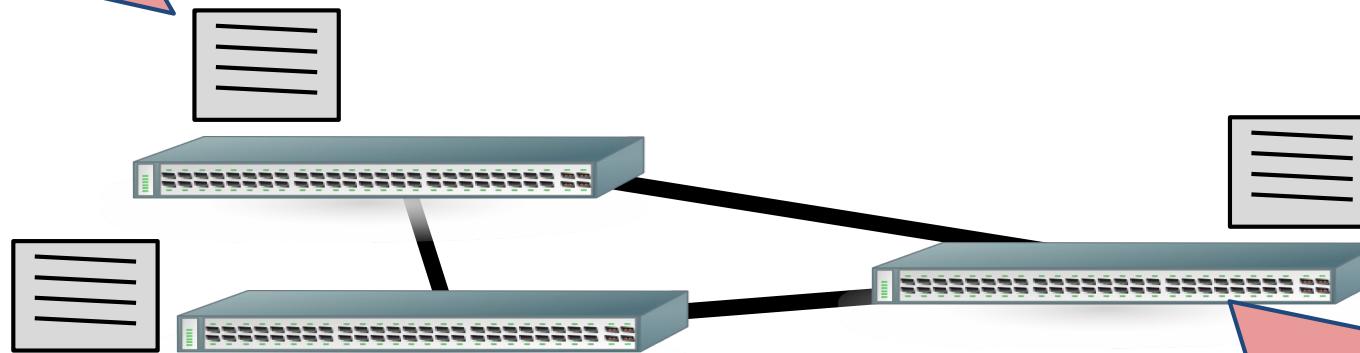


Also here: due to human errors.

No Surprise: Networks Are Complex

Manual, device-centric
network configurations
(CLI, LANmanager)

Un-evolved best practices
(tcpdump, traceroute - from the 1990s)

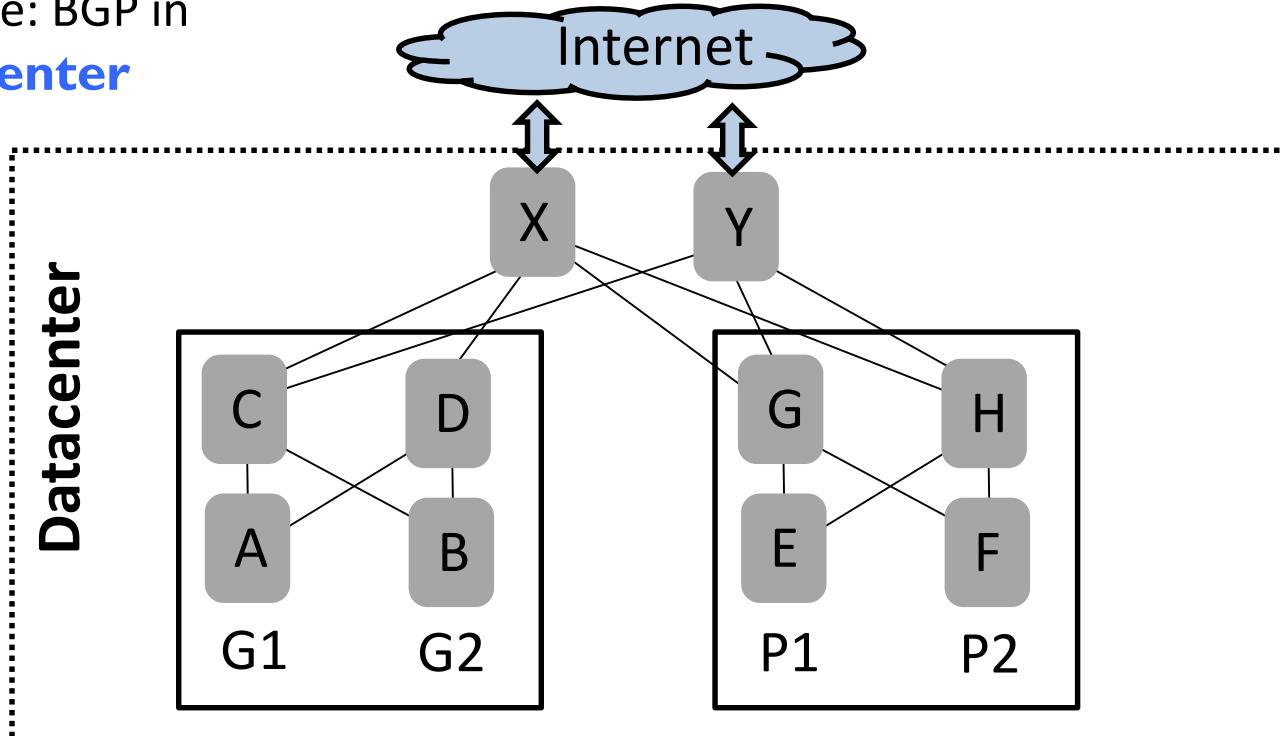


Complex, leaky, low-level interfaces
(VLANs, Spanning Tree, Routing)

500-router network: typically
>1 million lines of configuration

Particularly Challenging for Humans: Reasoning about Policy-Compliance under Failures

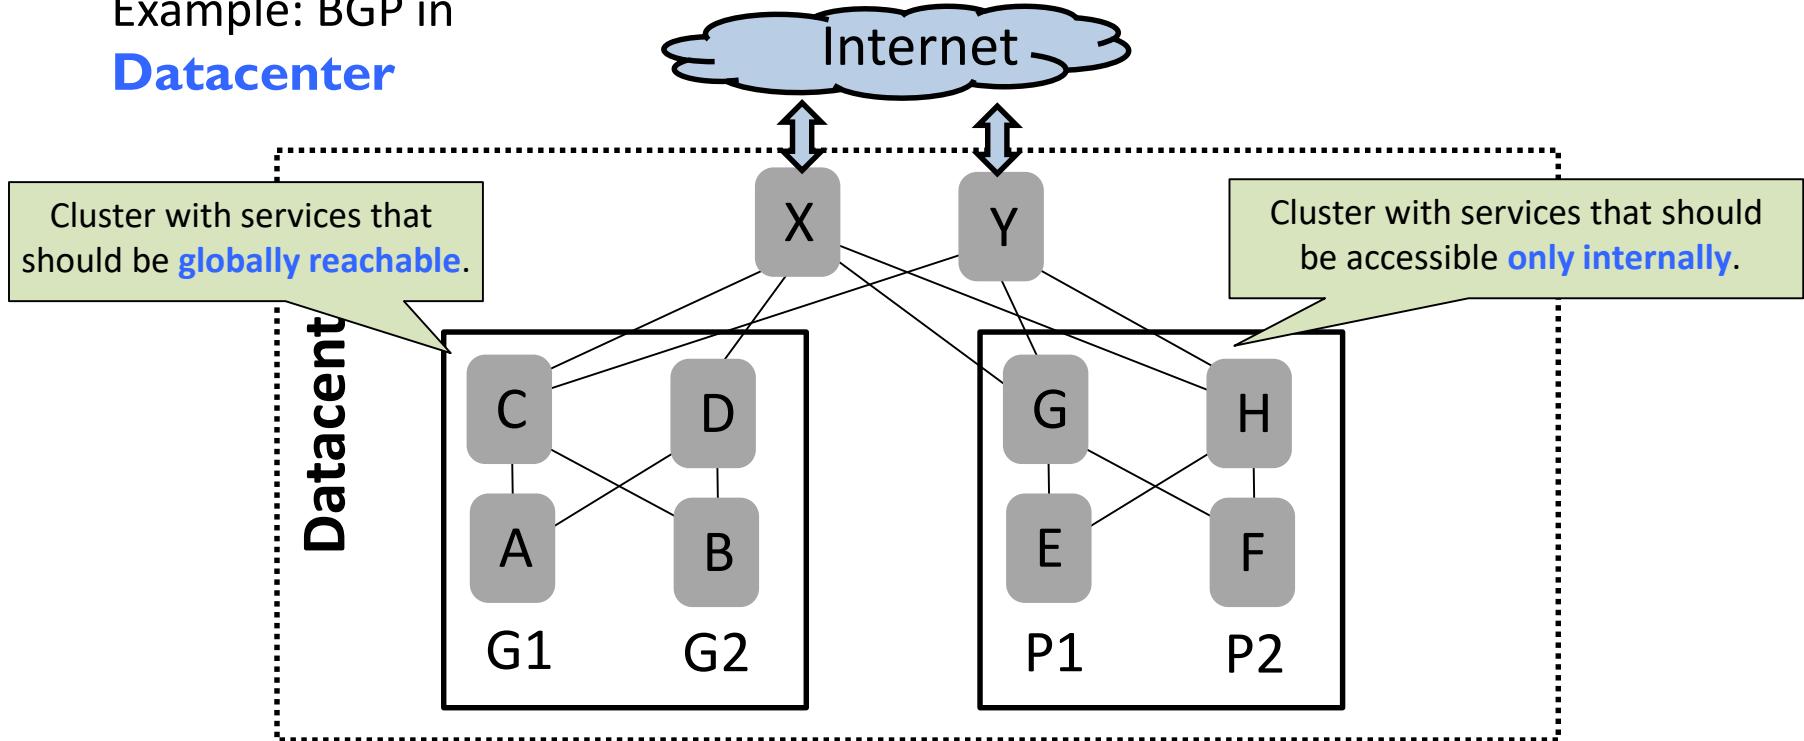
Example: BGP in
Datacenter



Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

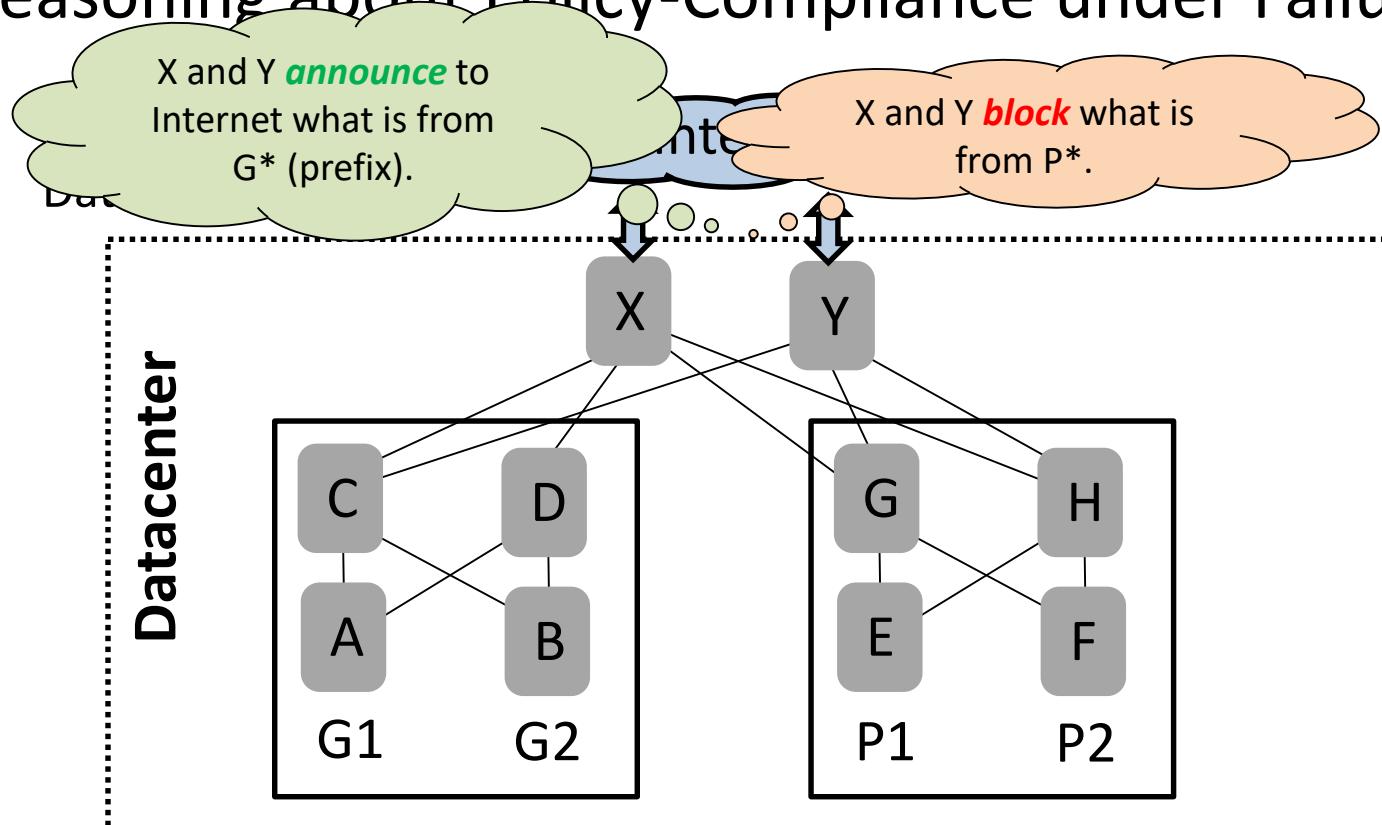
Particularly Challenging for Humans: Reasoning about Policy-Compliance under Failures

Example: BGP in
Datacenter



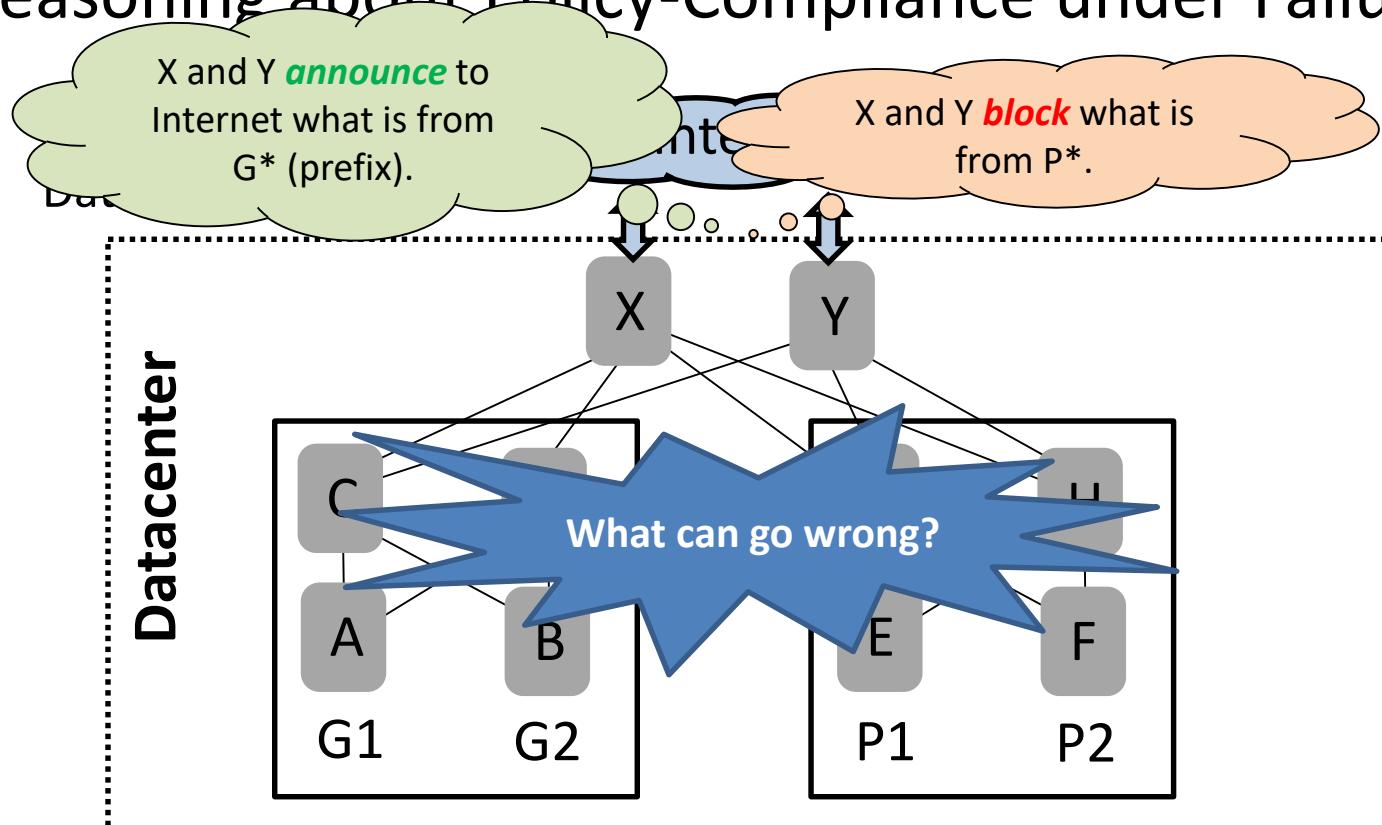
Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

Particularly Challenging for Humans: Reasoning about Policy-Compliance under Failures



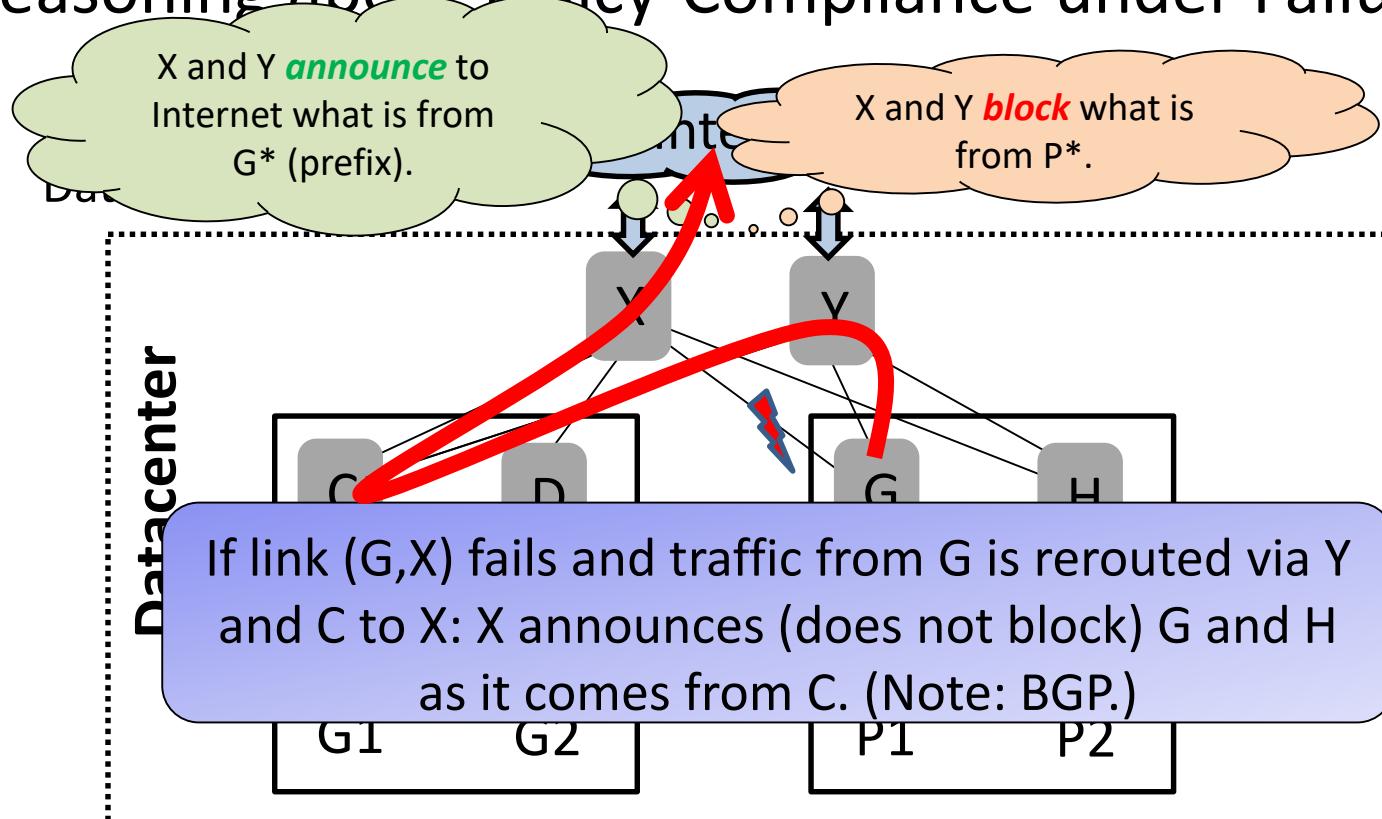
Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

Particularly Challenging for Humans: Reasoning about Policy-Compliance under Failures



Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

Particularly Challenging for Humans: Reasoning about Policy-Compliance under Failures



Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

We're Falling Behind the Curve: Increasing Complexity, Software from the 90s

- Anecdote **Wall Street bank**: outage of a datacenter
 - Lost revenue measured in **1 mio\$/min**
- Quickly, an emergency team was assembled with experts in compute, storage and networking:
 - **The compute team:** *reams of logs*, written experiments to reproduce and *isolate the error*
 - **The storage team:** *system logs* were affected, *workaround programs*.
 - “All the **networking team** had were *two tools invented over twenty years ago* to merely test end-to-end connectivity. Neither tool could reveal *problems with the switches*, the *congestion* experienced.”



Source: «The world's fastest and most programmable networks»
White Paper Barefoot Networks

Roadmap

- A Brief History of Resilient Networking
- Algorithms for Local Fast Re-Routing (FRR)
- Accounting for Congestion
- Accounting for Network Policy

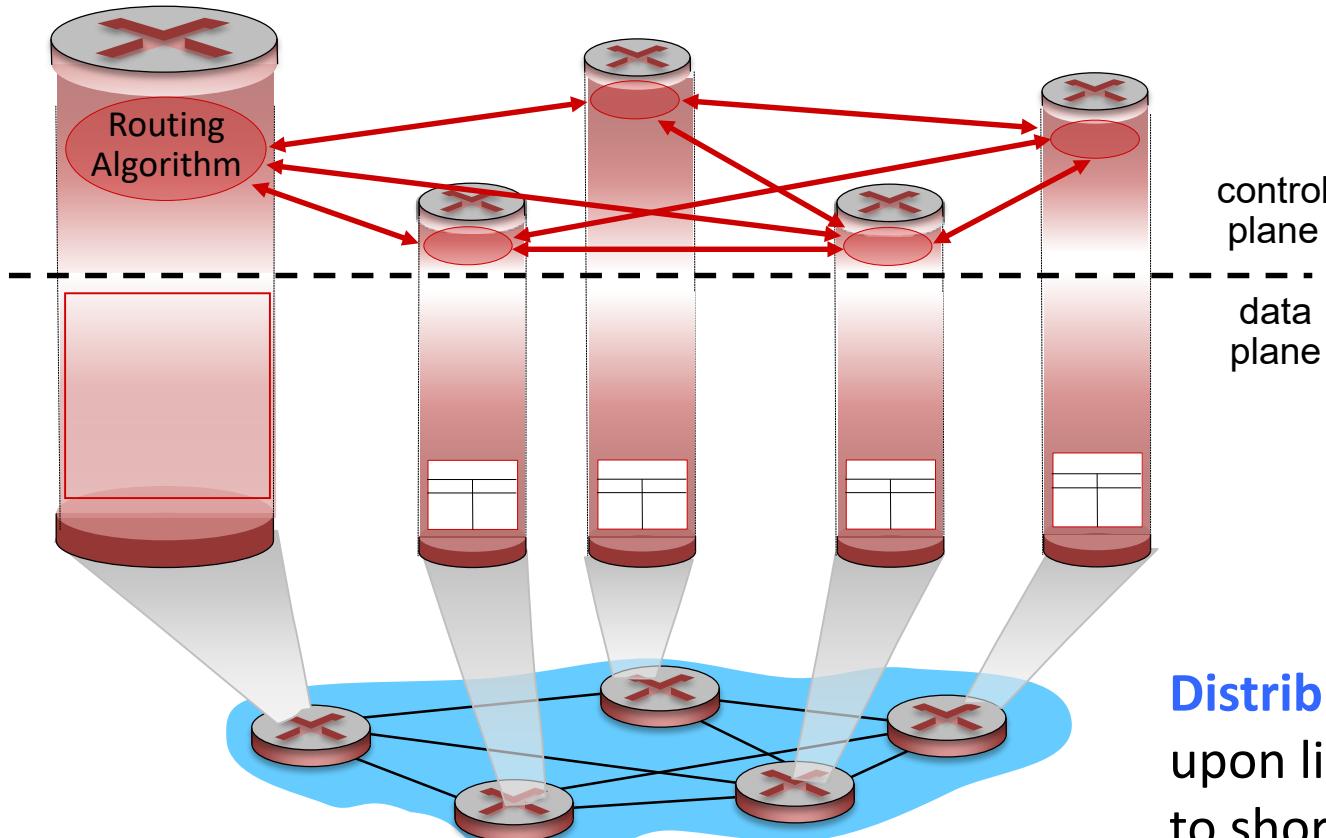


Roadmap

- **A Brief History of Resilient Networking**
- Algorithms for Local Fast Re-Routing (FRR)
- Accounting for Congestion
- Accounting for Network Policy

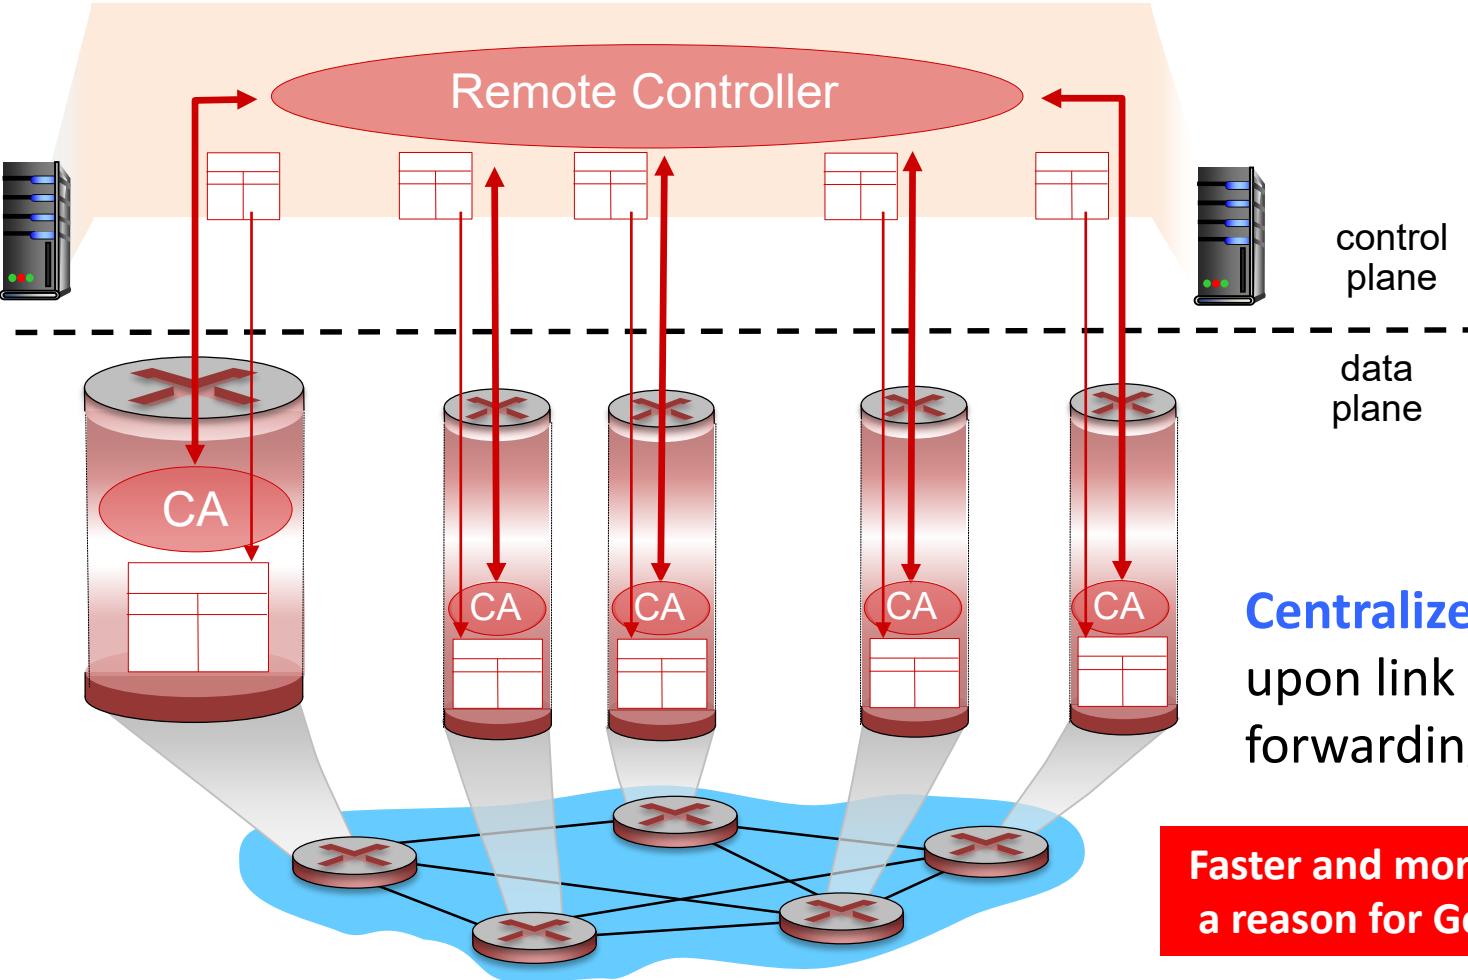


Traditional Networks



Distributed algorithms:
upon link failure, *reconverge*
to shortest paths

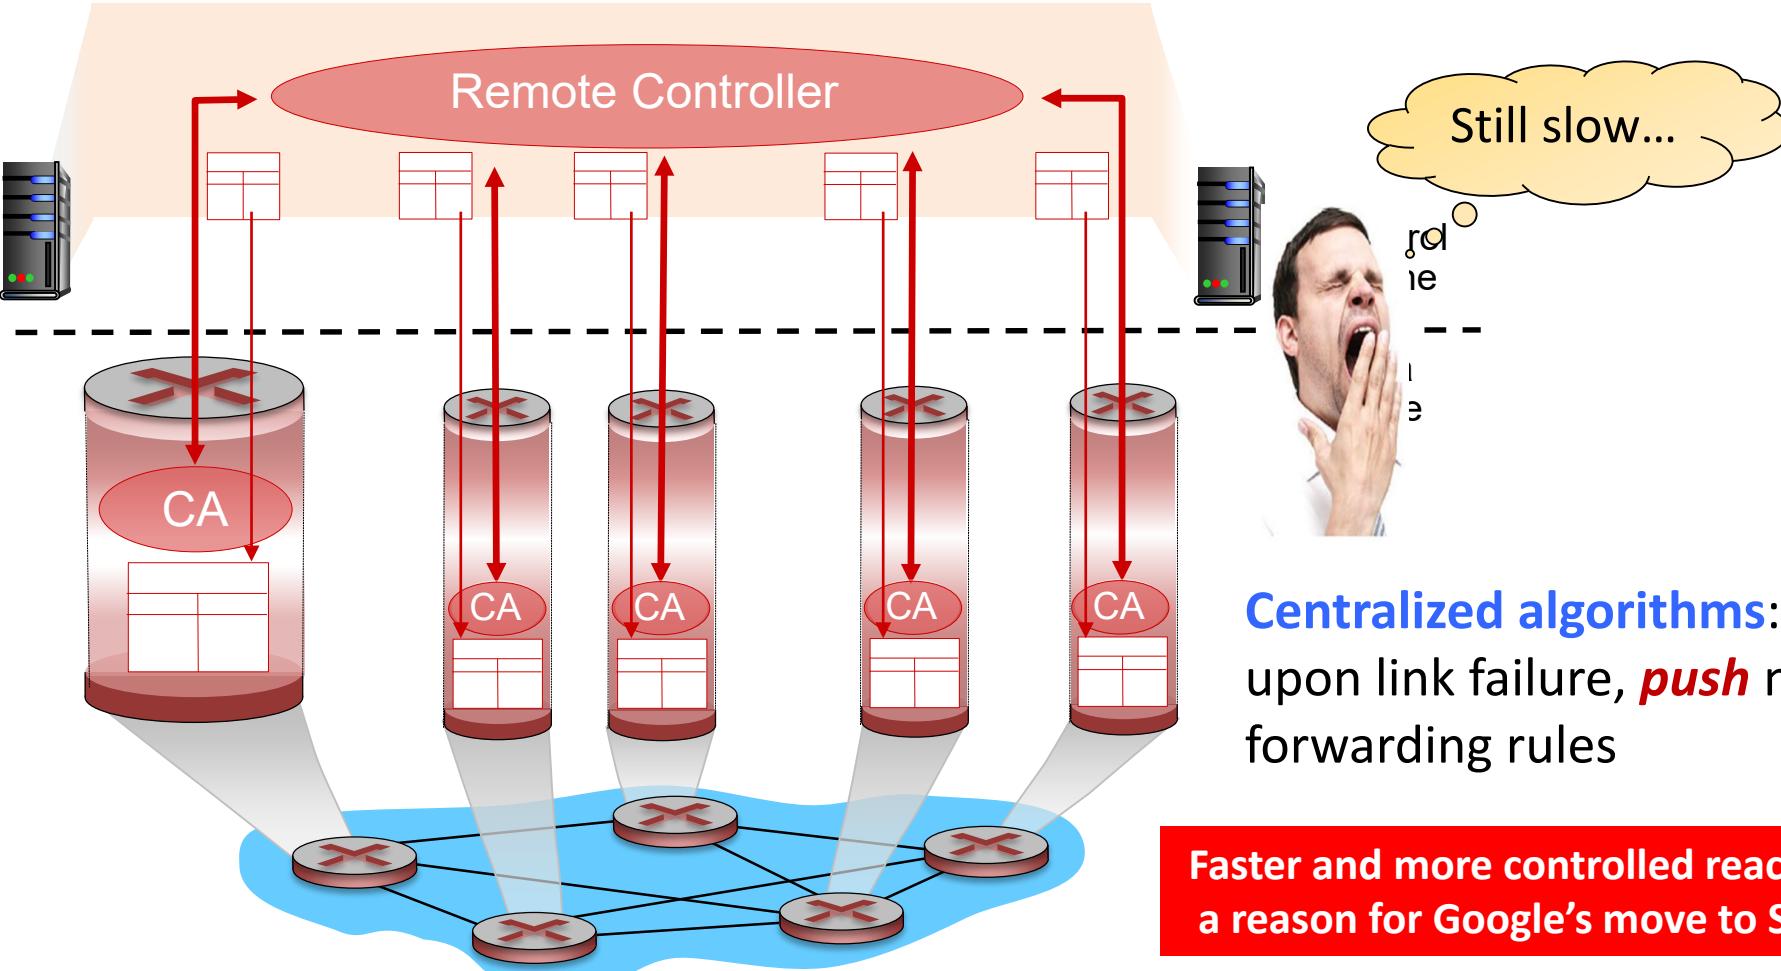
Software-Defined Networks (SDN)



Centralized algorithms:
upon link failure, **push** new
forwarding rules

Faster and more controlled reaction:
a reason for Google's move to SDN!

Software-Defined Networks (SDN)



Restoration in control plane takes time -> packet drops!



Failover: Control Plane vs Data Plane

- Slower reaction in the **control plane** than in the **data plane**



Minister of Education

VS



Teacher in the Classroom

Approaches for Failover

In Control Plane

- Distributed recomputation of shortest paths (“**re-convergence**”)
- Centralized recomputation of paths (SDN)
- **Link-reversal** algorithms (e.g., Gafni et al.)

vs

In Data Plane

- Static forwarding table
- Rules pre-installed **before** failures are known

Approaches for Failover

In Control Plane

Slow but “globally informed”.

Distributed recomputation of shortest paths (“**reconvergence**”)

- Centralized recomputation of paths (SDN)
- **Link-reversal** algorithms (e.g., Gafni et al.)

In Data Plane

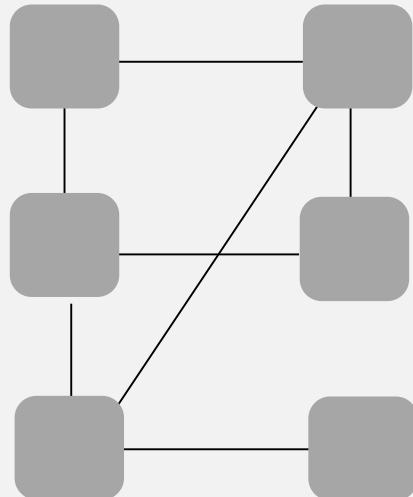
Fast but “local knowledge”.

- Static forwarding table
- Rules pre-installed **before** failures are known

vs

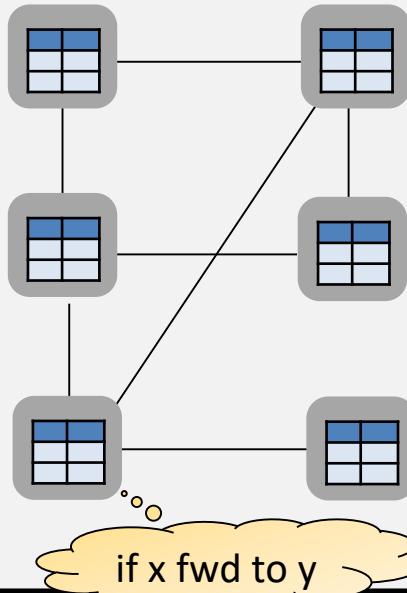
The FRR Problem

Phase 1: Rule installation



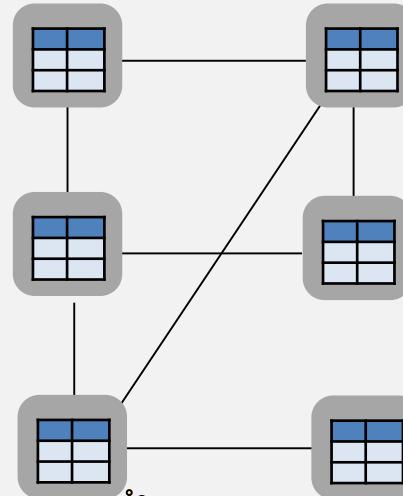
The FRR Problem

Phase 1: Rule installation

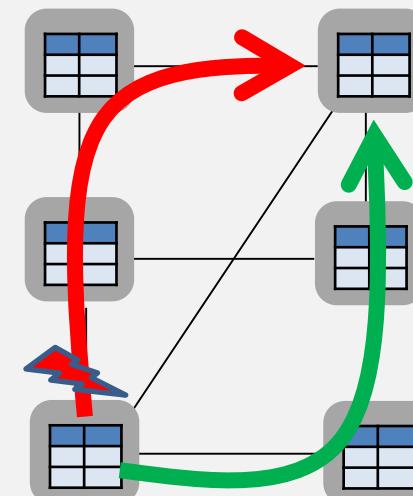


The FRR Problem

Phase 1: Rule installation

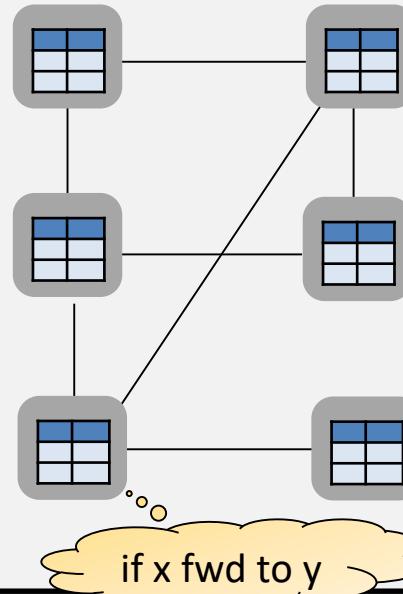


Phase 2: Failures and routing

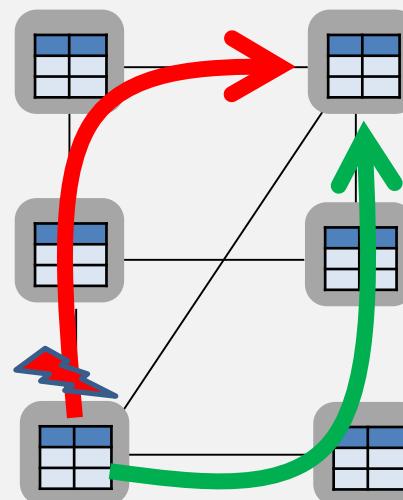


The FRR Problem

Phase 1: Rule installation



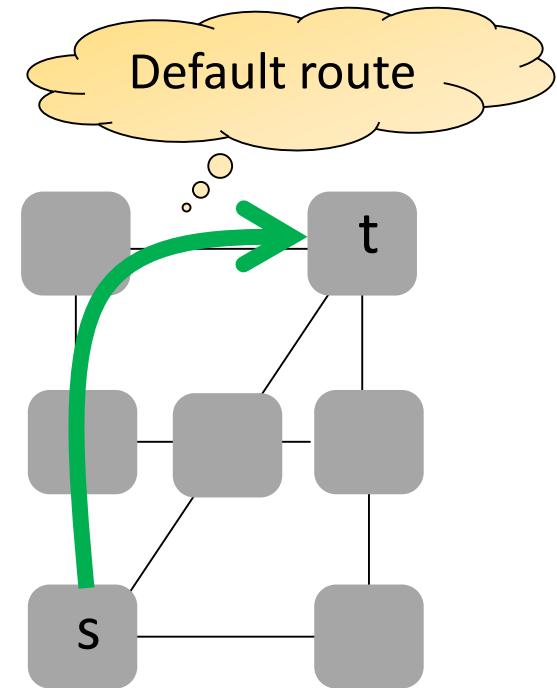
Phase 2: Failures and routing



The FRR Problem

- **Pre-installed** local-fast failover rules
 - *Can depend on local failures and, e.g., destination, import, source*
- **At runtime**, rules are just "*executed*"

Advantage: no need to wait
for reconvergence.

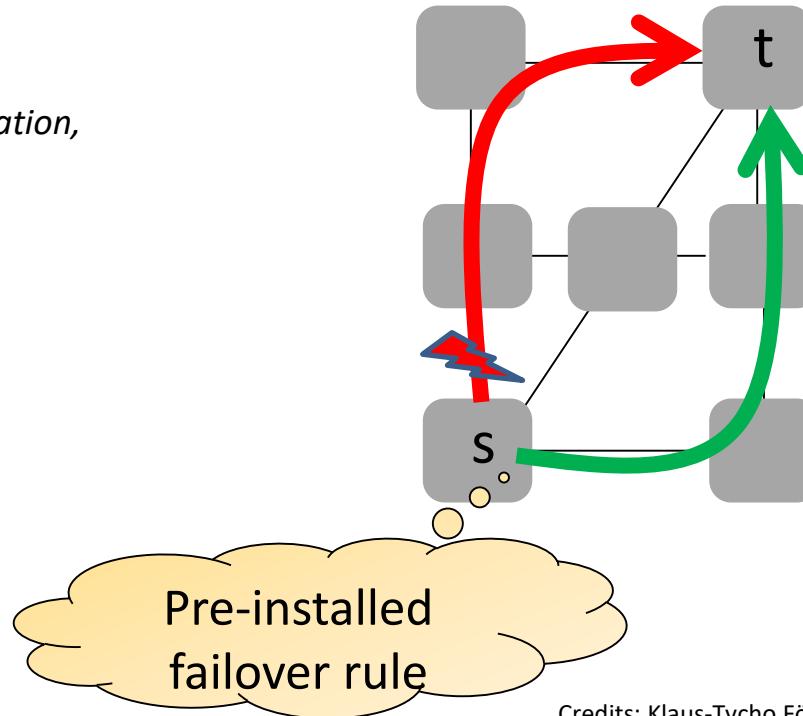


The FRR Problem

Good alternative
under 1 failure!

- **Pre-installed** local-fast failover rules
 - *Can depend on local failures and, e.g., destination, import, source*
- **At runtime**, rules are just “executed”

Advantage: no need to wait
for reconvergence.

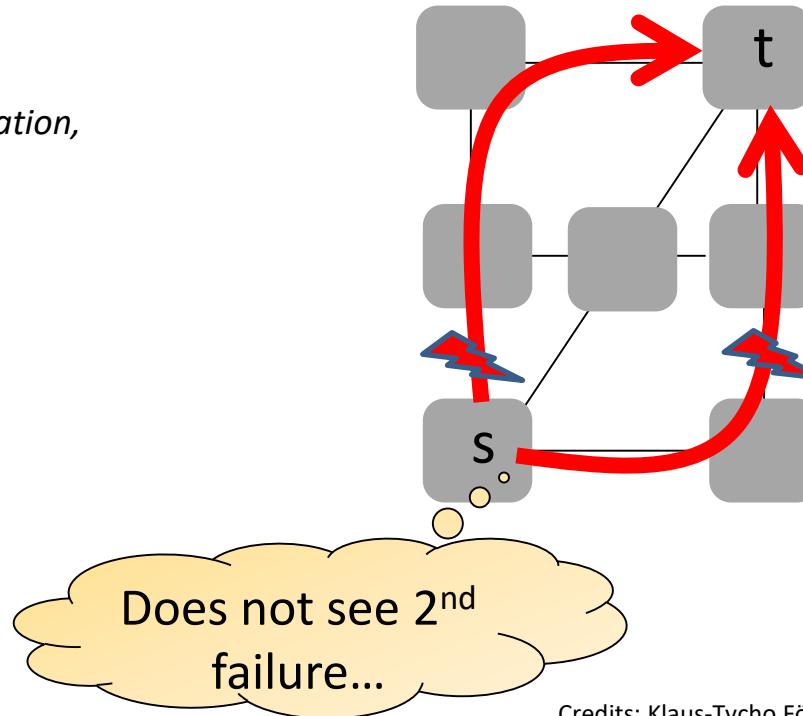


The FRR Problem

Good alternative under 1 failure!

- Pre-installed local-fast failover rules
 - Can depend on local failures and, e.g., destination, import, source
 - At runtime, rules are just "executed"

Advantage: no need to wait for reconvergence.

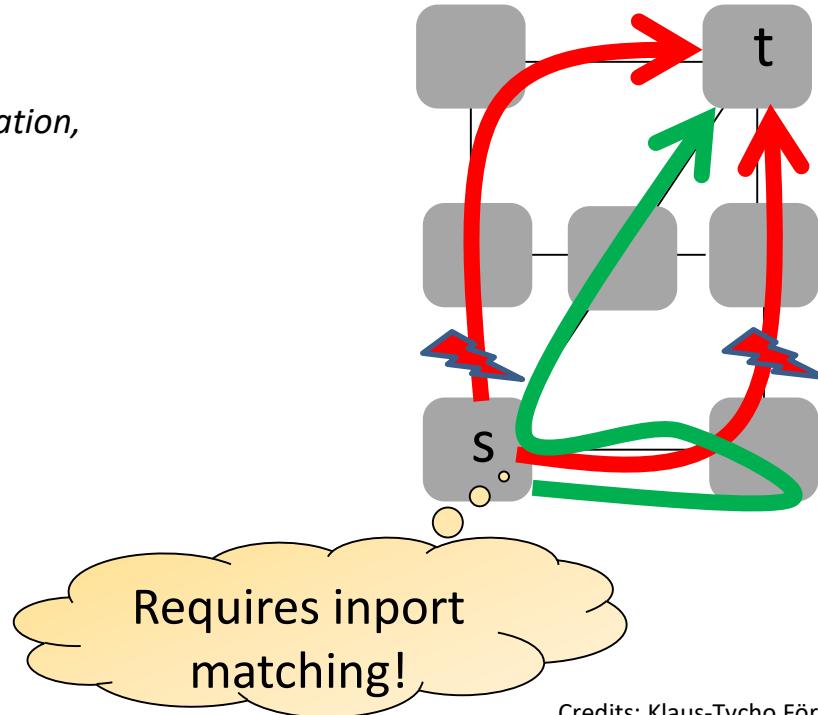


The FRR Problem

Can get complex under multiple failures..

- **Pre-installed** local-fast failover rules
 - *Can depend on local failures and, e.g., destination, import, source*
- **At runtime**, rules are just “executed”

Advantage: no need to wait for reconvergence.

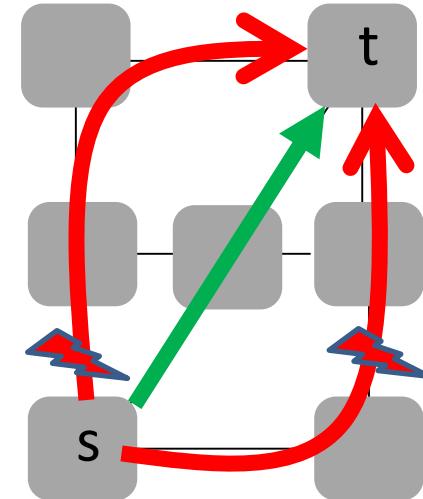


The FRR Problem

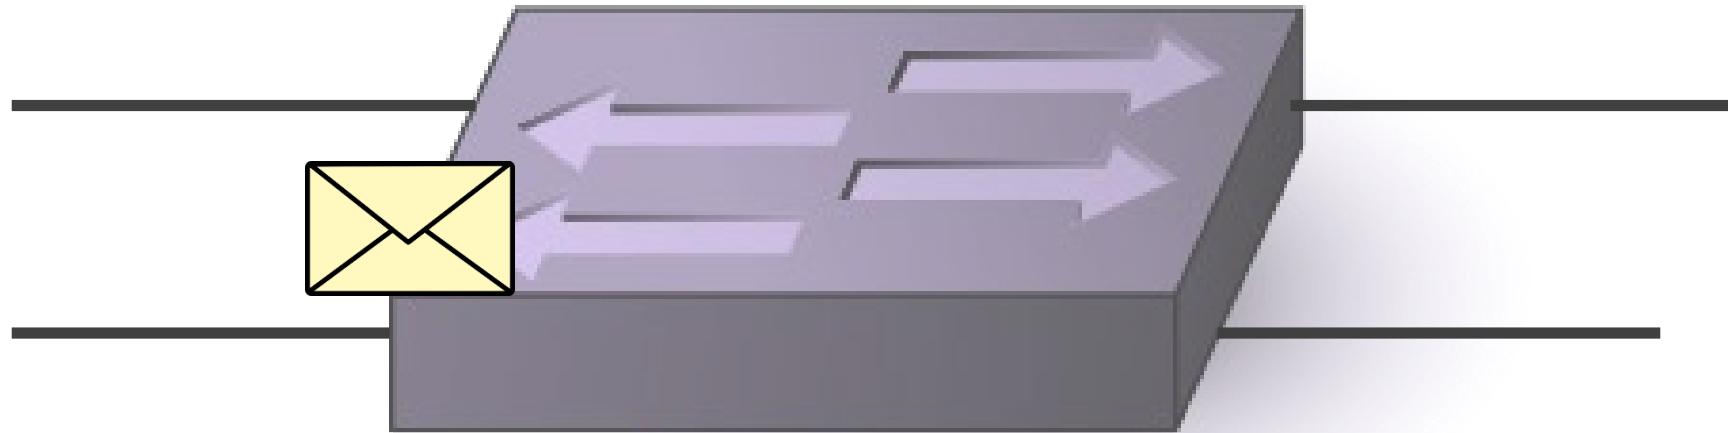
With global knowledge: simpler!

- **Pre-installed** local-fast failover rules
 - *Can depend on local failures and, e.g., destination, import, source*
- **At runtime**, rules are just “executed”

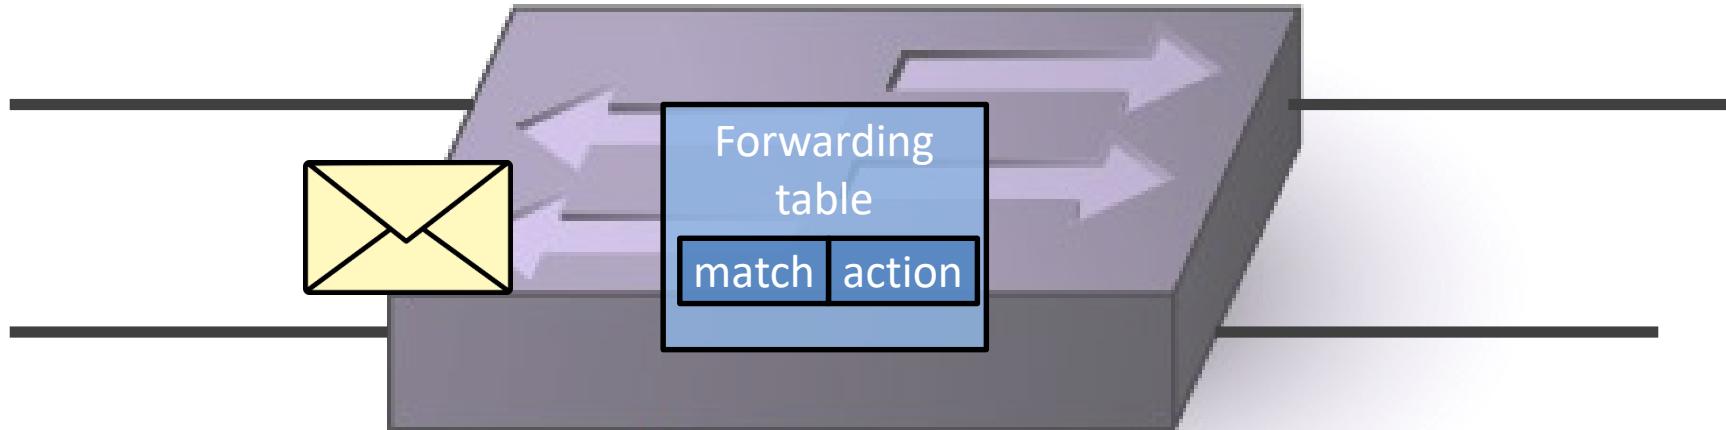
Advantage: no need to wait
for reconvergence.



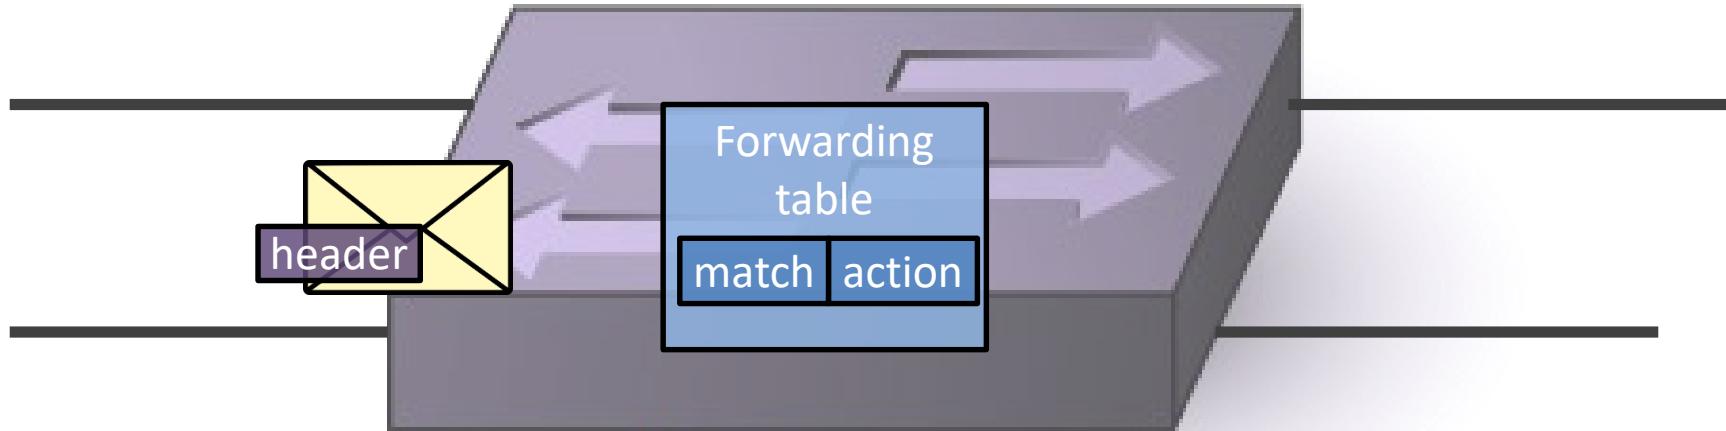
What information is **locally** available in a switch for handling a packet?



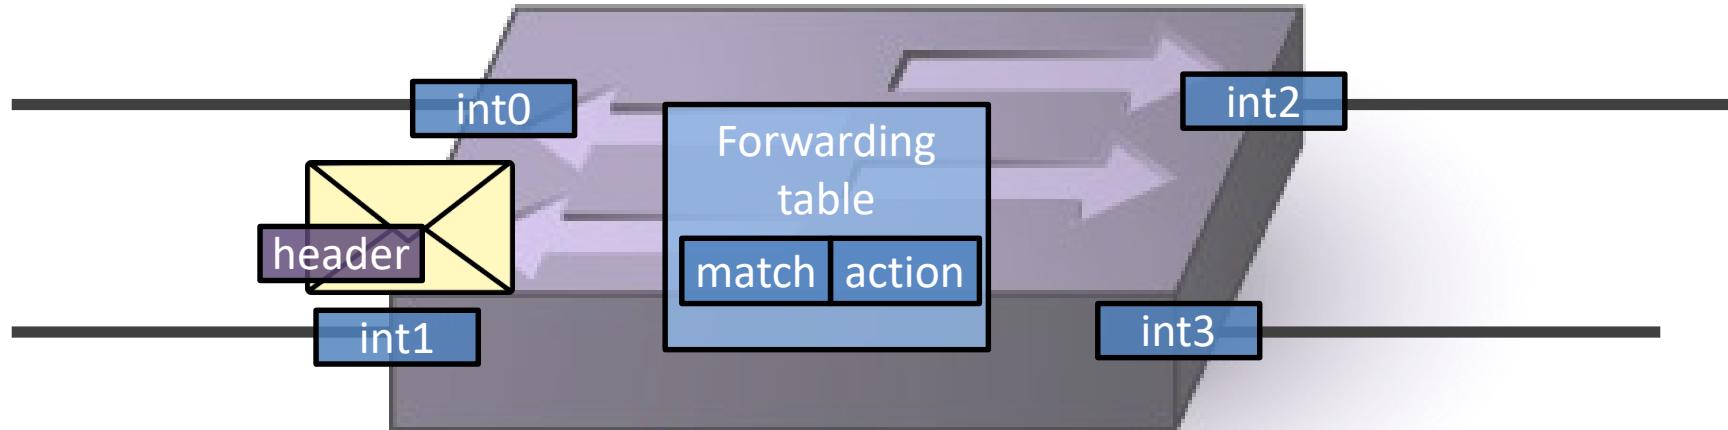
Locally Available Information: The Forwarding Table: Match -> Action



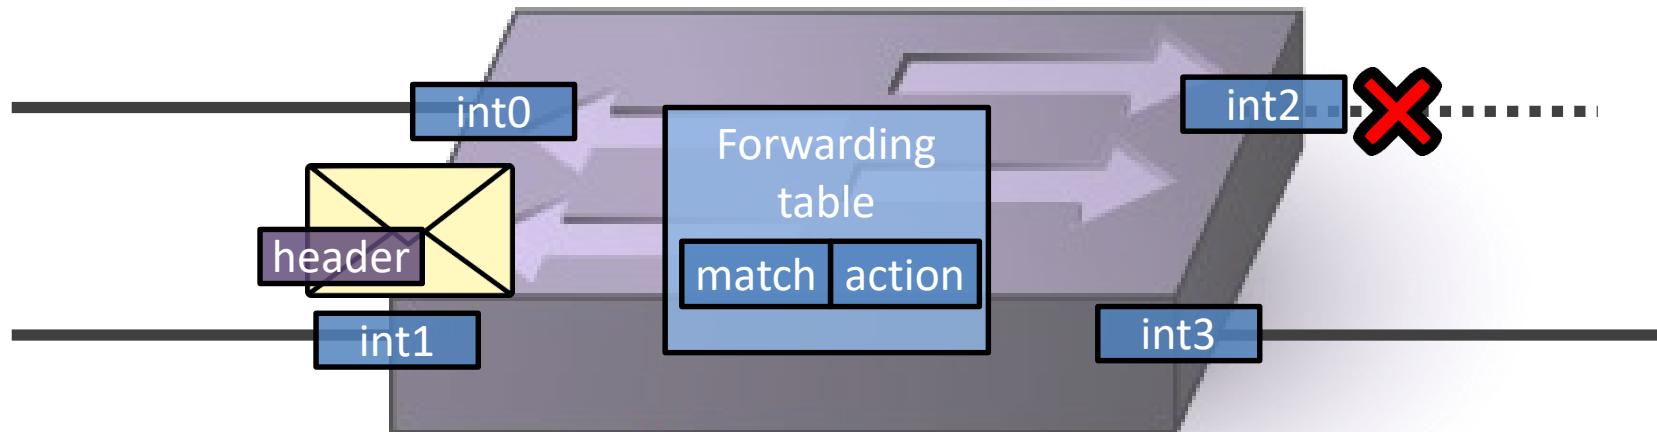
Locally Available Information: The Packet Header



Locally Available Information: The Import of the Received Packet



Locally Available Information: The Outgoing Port Depends on Failed Links



Raises an Interesting Question

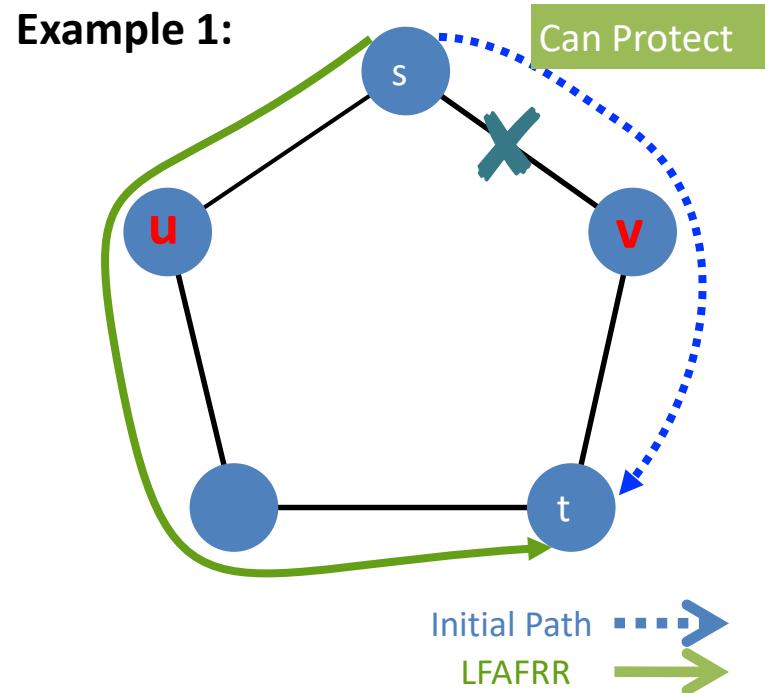
Can we pre-install local fast failover rules which ensure reachability under multiple failures? *In particular:* **How many failures** can be tolerated by static forwarding tables?

Remark: Traditional Approach LFA

- Traditionally: forwarding along **shortest paths**
- Loop-Free Alternative (LFA)**: failover to alternative neighbor, from there shortest path

Example 1:

- If (s,v) fails, s can *failover to u*
- u has shortest path to t that does not go through (s,v) again
- WORKS**: can protect (s,v)

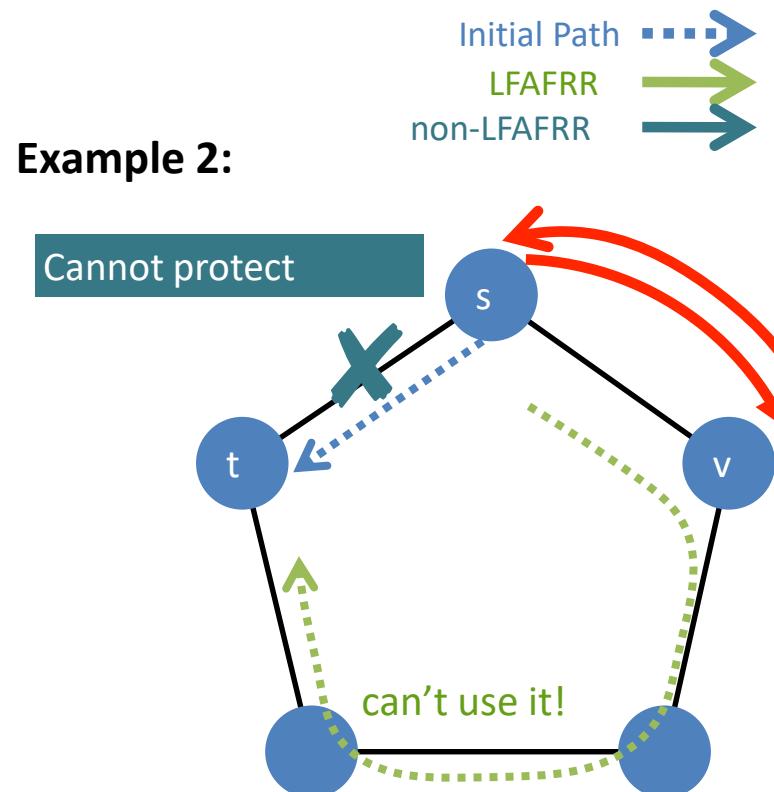


Remark: Traditional Approach LFA

- Traditionally: forwarding along **shortest paths**
- Loop-Free Alternative (LFA)**: failover to alternative neighbor, from there shortest path

Example 2:

- If (s,t) fails, s can only try to failover to v
- However, when v 's shortest route to t goes along s again: loop
- DOES NOT WORK**: Cannot protect (s,t)



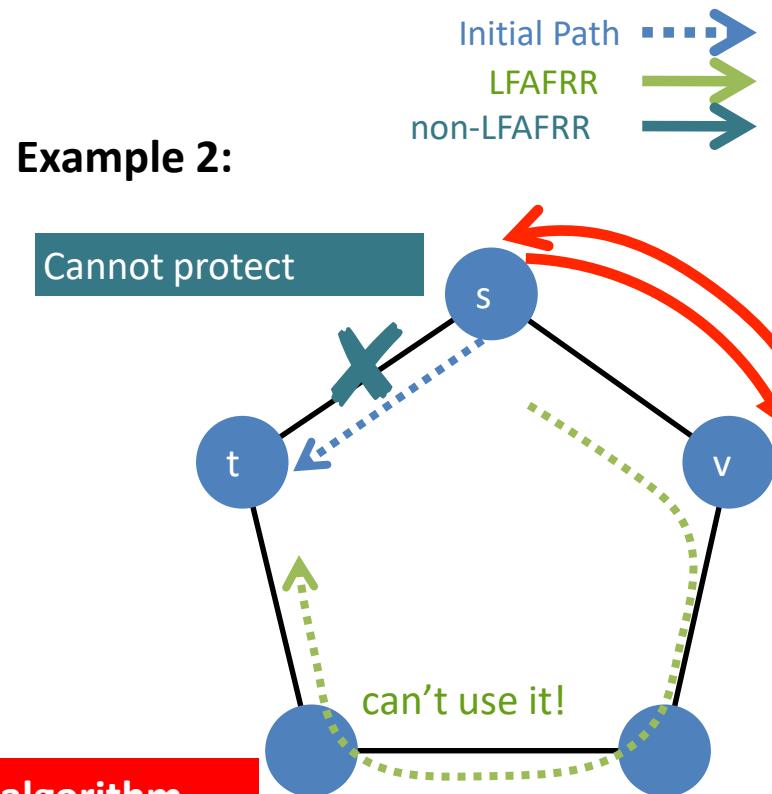
Remark: Traditional Approach LFA

- Traditionally: forwarding along **shortest paths**
- Loop-Free Alternative (LFA)**: failover to alternative neighbor, from there shortest path

Example 2:

- If (s,t) fails, s can only try to failover to v
- However, when v 's shortest route to t goes along s again: loop
- DOES NOT WORK**: Cannot protect (s,t)

Even though loop-free alternative path exists, an LFA algorithm cannot use it. Protection ratio of LFA depends on topology.

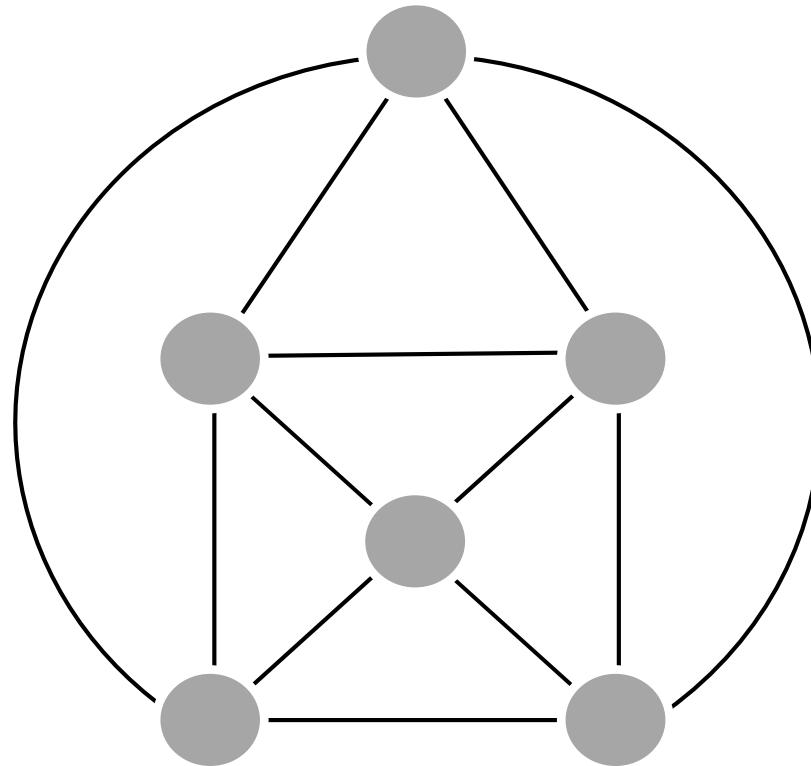


Roadmap

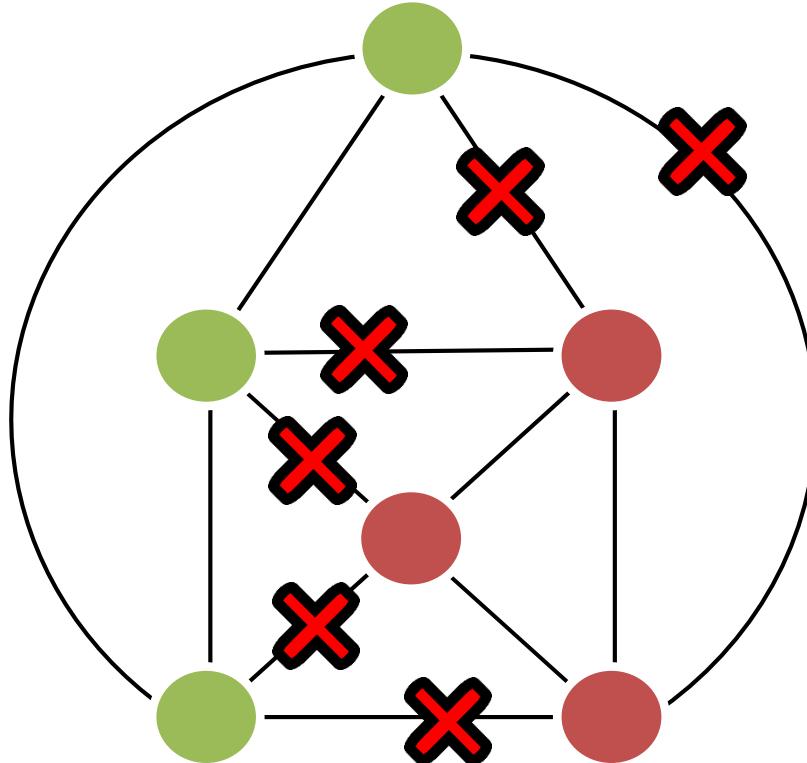
- A Brief History of Resilient Networking
- **Algorithms for Local Fast Re-Routing (FRR)**
- Accounting for Congestion
- Accounting for Network Policy



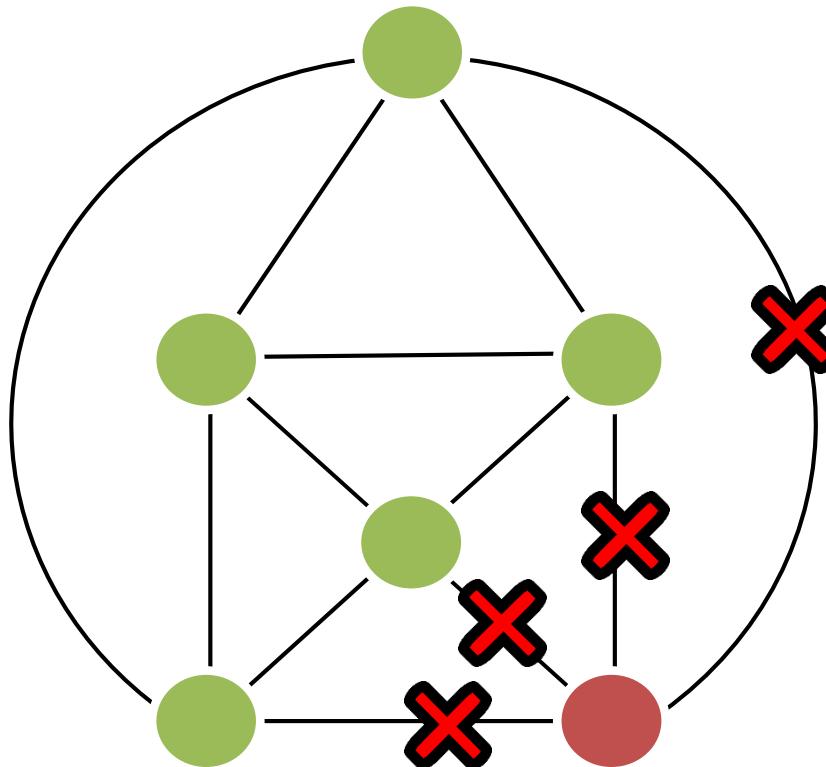
So: How many failures can be tolerated by static forwarding tables?



If we partition the network,
there is not much to do



The connectivity k of a network N : the minimum number of link deletions that partitions N



The connectivity of this network is *four*

Resilience Criteria

Ideal resilience

Given a k -connected graphs, we can tolerate *any $k-1$ link failures*.

Perfect resilience

Any source s can always reach any destination t as long as the underlying network is *physically connected*.

Can this be achieved? Assume undirected link failures.

Resilience Criteria

Ideal resilience

Given a k -connected graphs, we can tolerate *any $k-1$ link failures*.

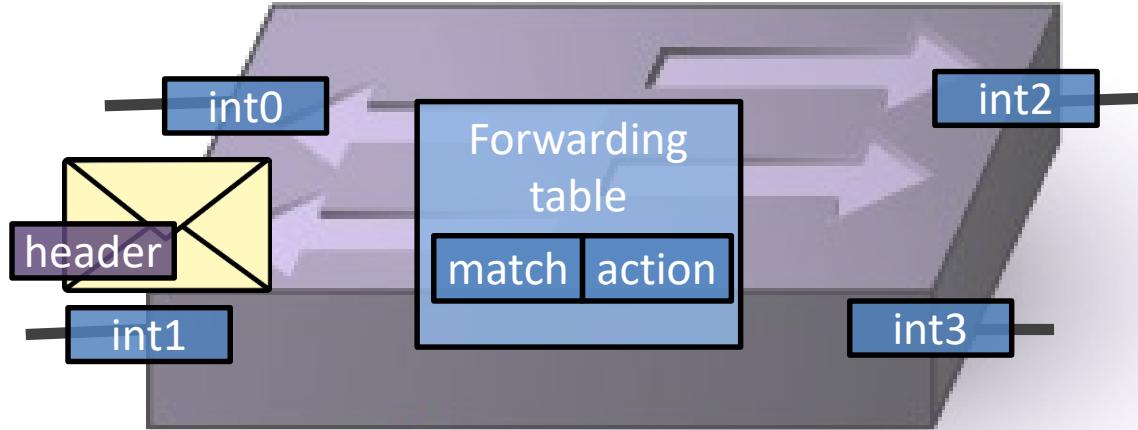
Perfect resilience

Any source s can always reach any destination t as long as the underlying network is *physically connected*.

Can this be achieved? Assume undirected link failures.

Spectrum of Models

Recall our switch model:

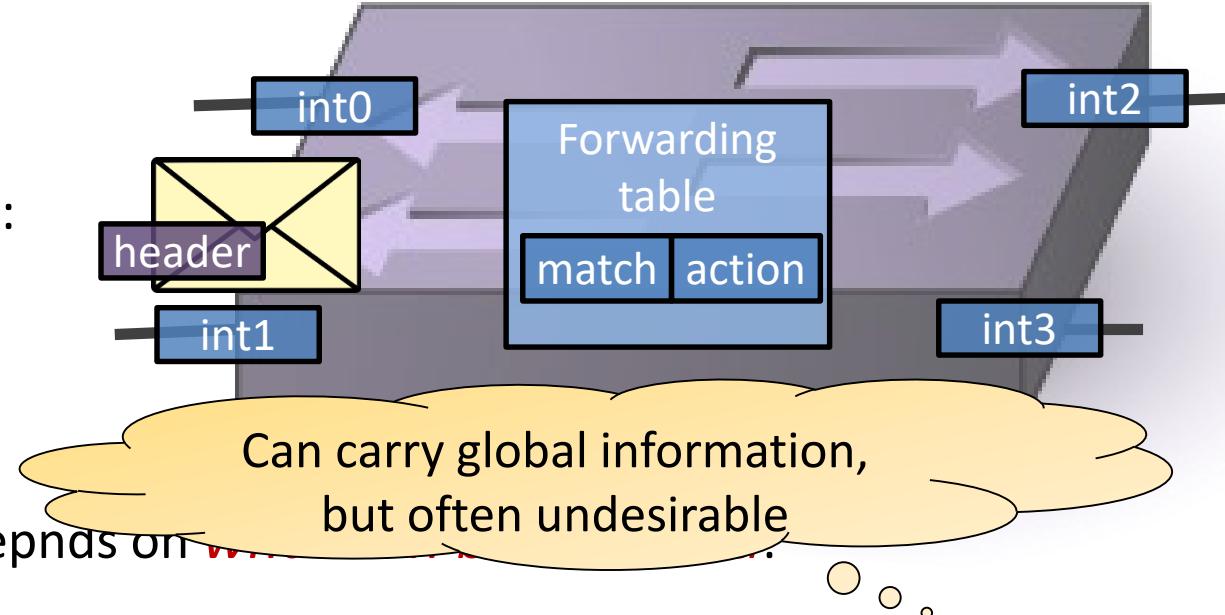


Achievable resilience depnds on *what can be matched*:

Per-destination	Per source	Incoming port	Probabilistic forwarding	Packet header rewriting
-----------------	------------	---------------	--------------------------	-------------------------

Spectrum of Models

Recall our switch model:

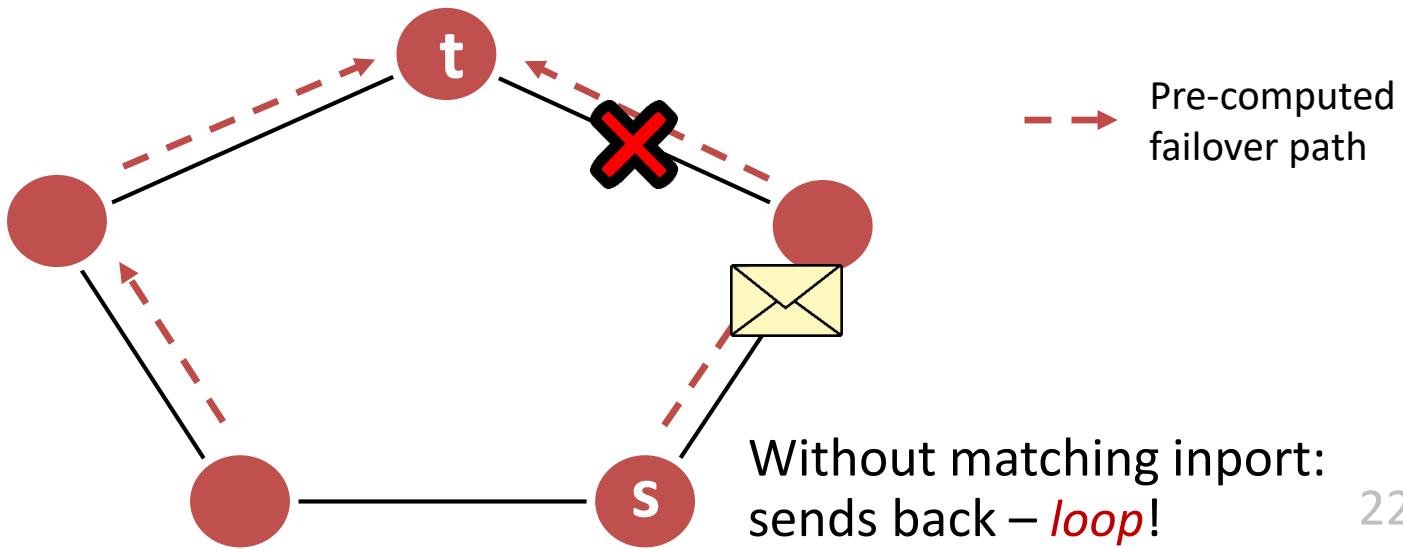


Achievable resilience depends on

Per-destination	Per source	Incoming port	Probabilistic forwarding	Packet header rewriting
-----------------	------------	---------------	--------------------------	-------------------------

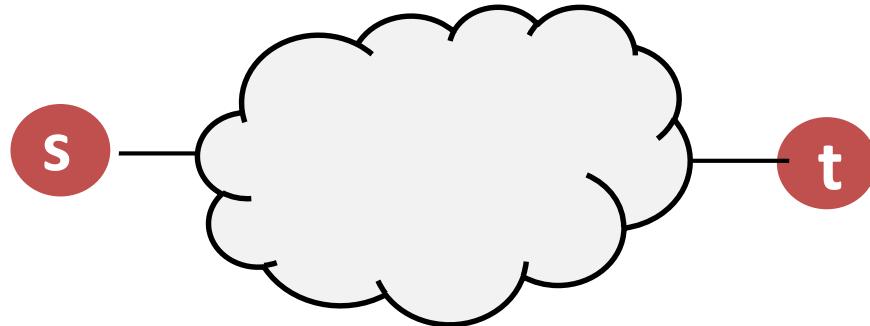
Per-destination routing *cannot cope* with *even one* link failure

Per-destination	Per source	Incoming port	Probabilistic forwarding	Packet header rewriting	Resiliency
X					0



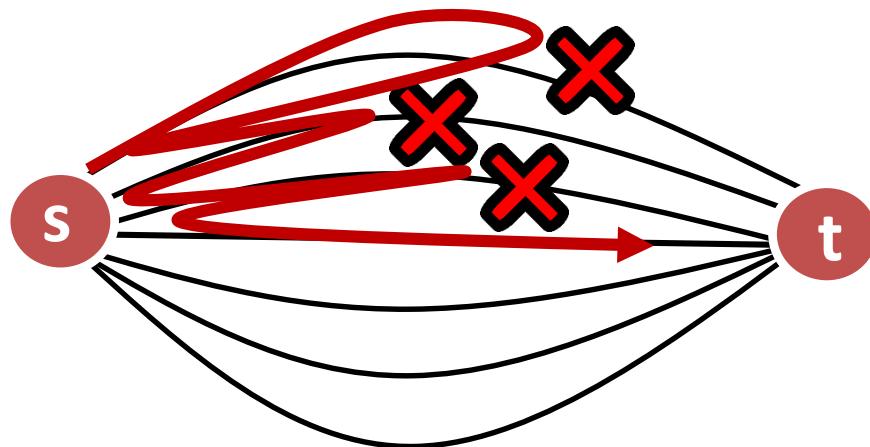
Can we achieve $k - 1$ resiliency in k -connected graph here?

Per-destination	Per source	Incoming port	Probabilistic forwarding	Packet header rewriting	Resiliency
X	X	X			?



Can we achieve $k - 1$ resiliency in k -connected graph here?

Per-destination	Per source	Incoming port	Probabilistic forwarding	Packet header rewriting	Resiliency
X	X	X			Yes



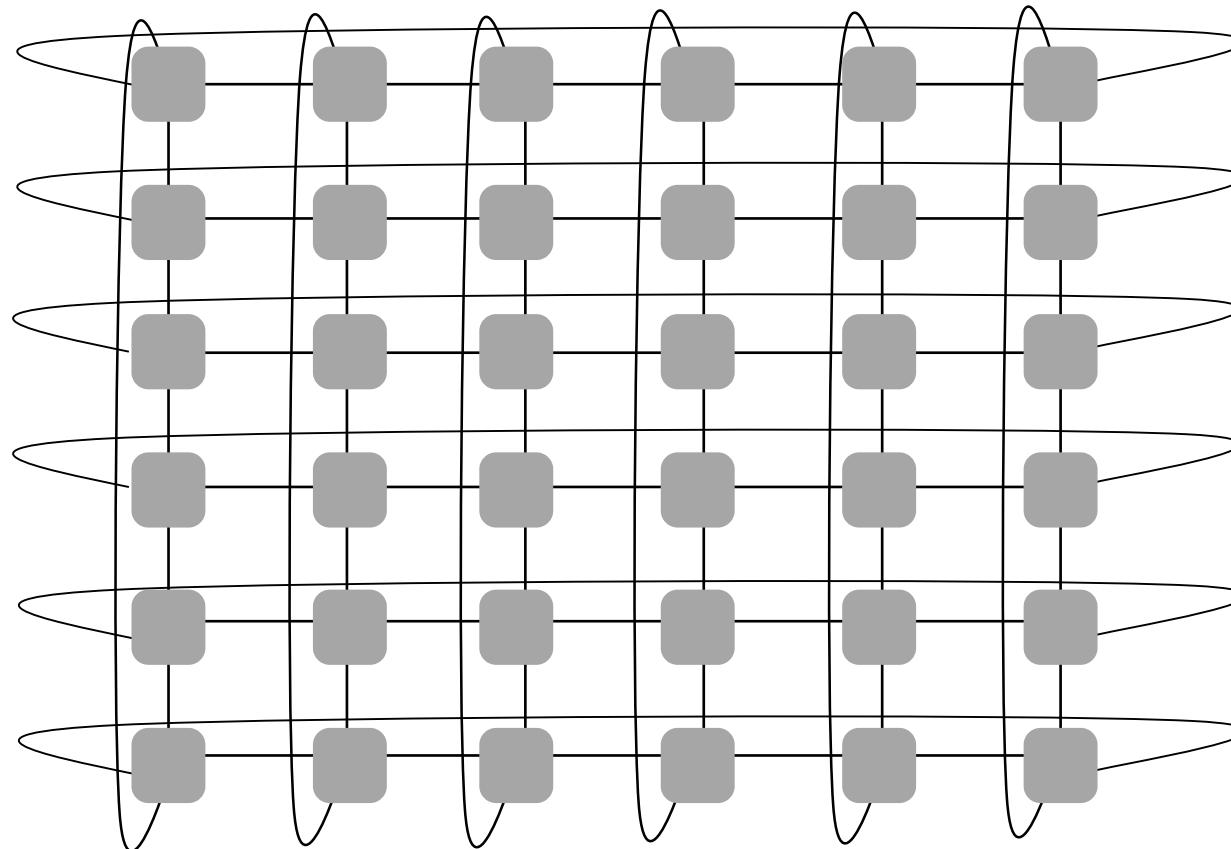
k disjoint paths: try one after the other, routing *back to source* each time.

Can we achieve $k - 1$ resiliency in k -connected graph here?

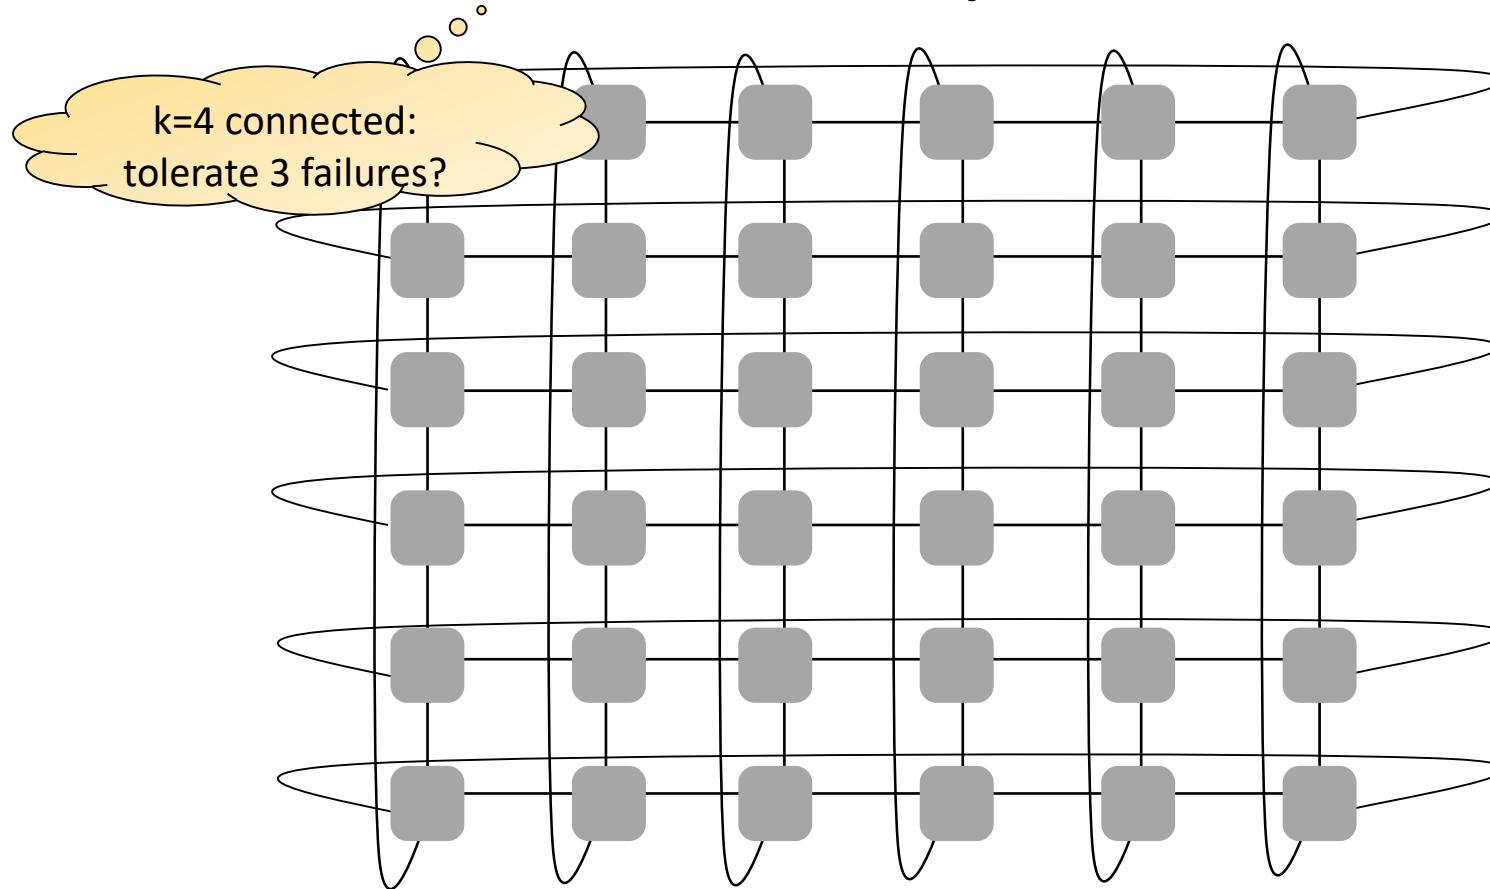
Per-destination	Per source	Incoming port	Probabilistic forwarding	Packet header rewriting	Resiliency
X		X			?

What about this scenario?
Practically important. From now
on called “ideal resilience”.

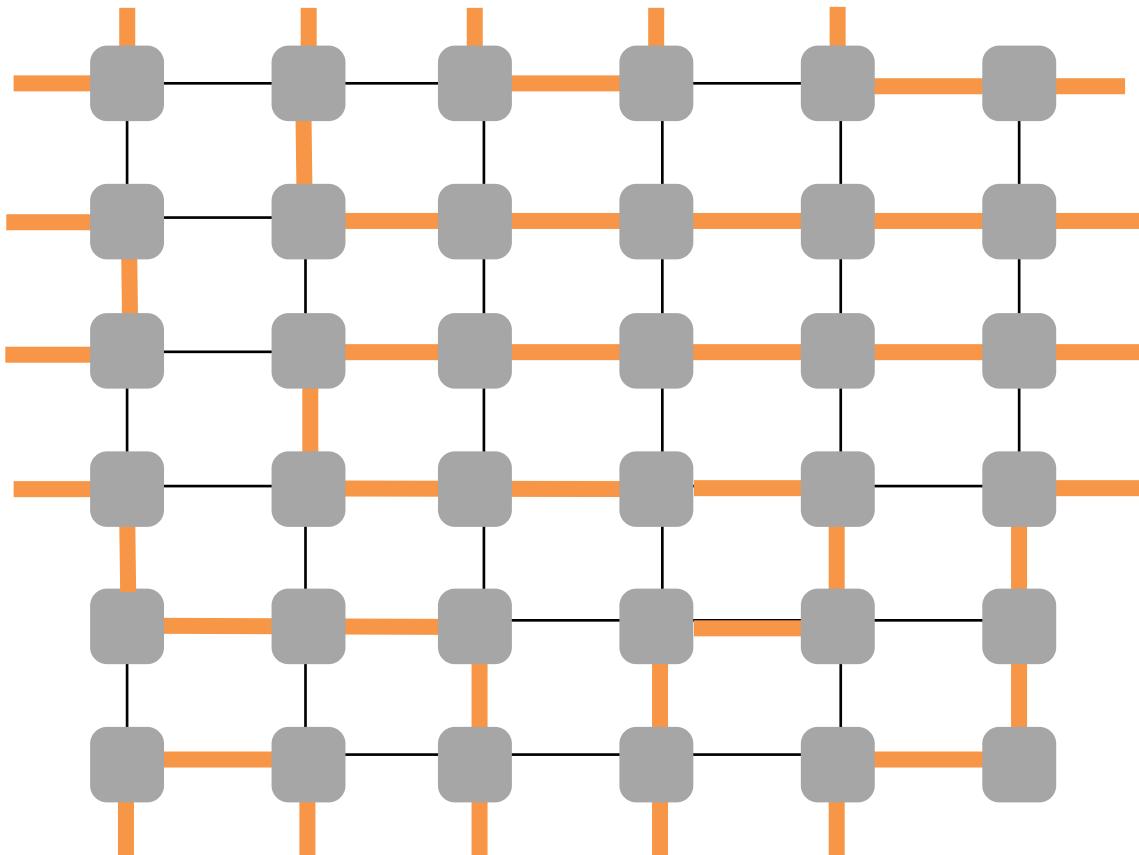
Ideal Resilience: Example 2-dim Torus?



Ideal Resilience: Example 2-dim Torus?



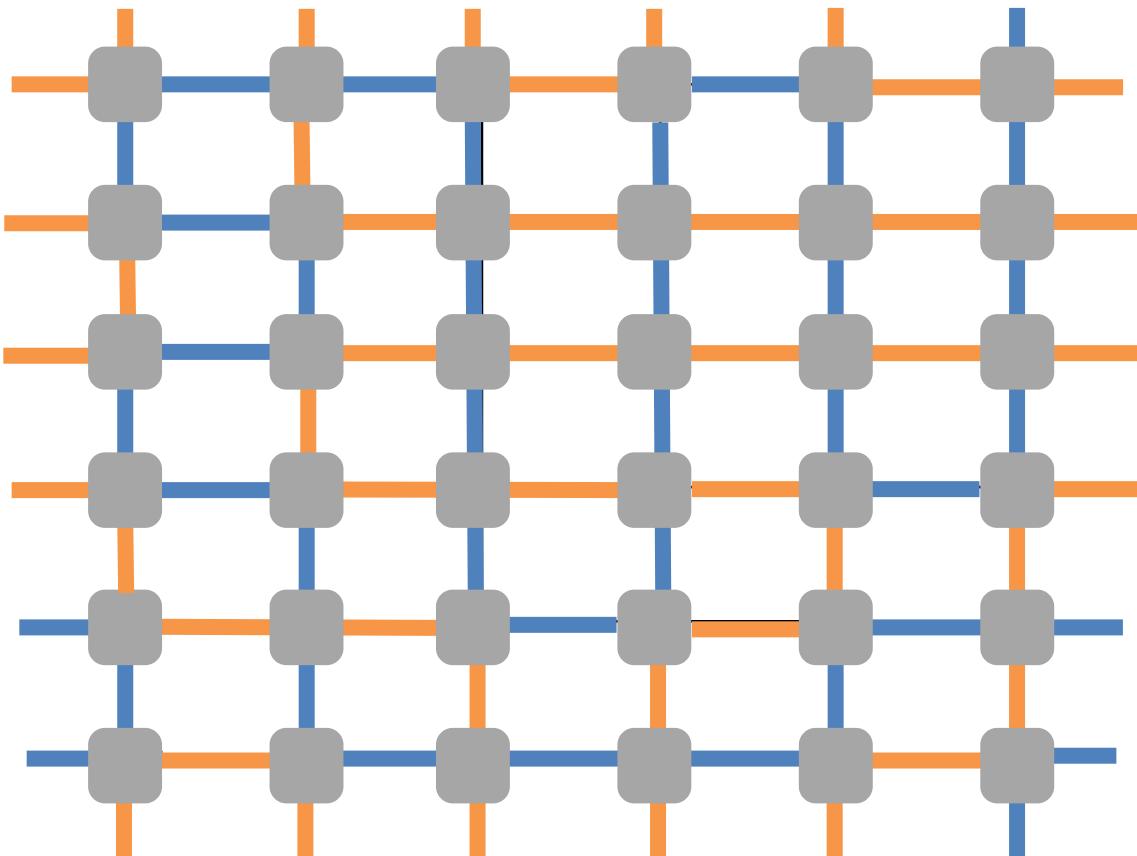
Idea: Decomposition into Hamilton Cycles



- Decompose torus into 2-edge-disjoint Hamilton Cycles (HC)

— *1st Hamilton cycle*

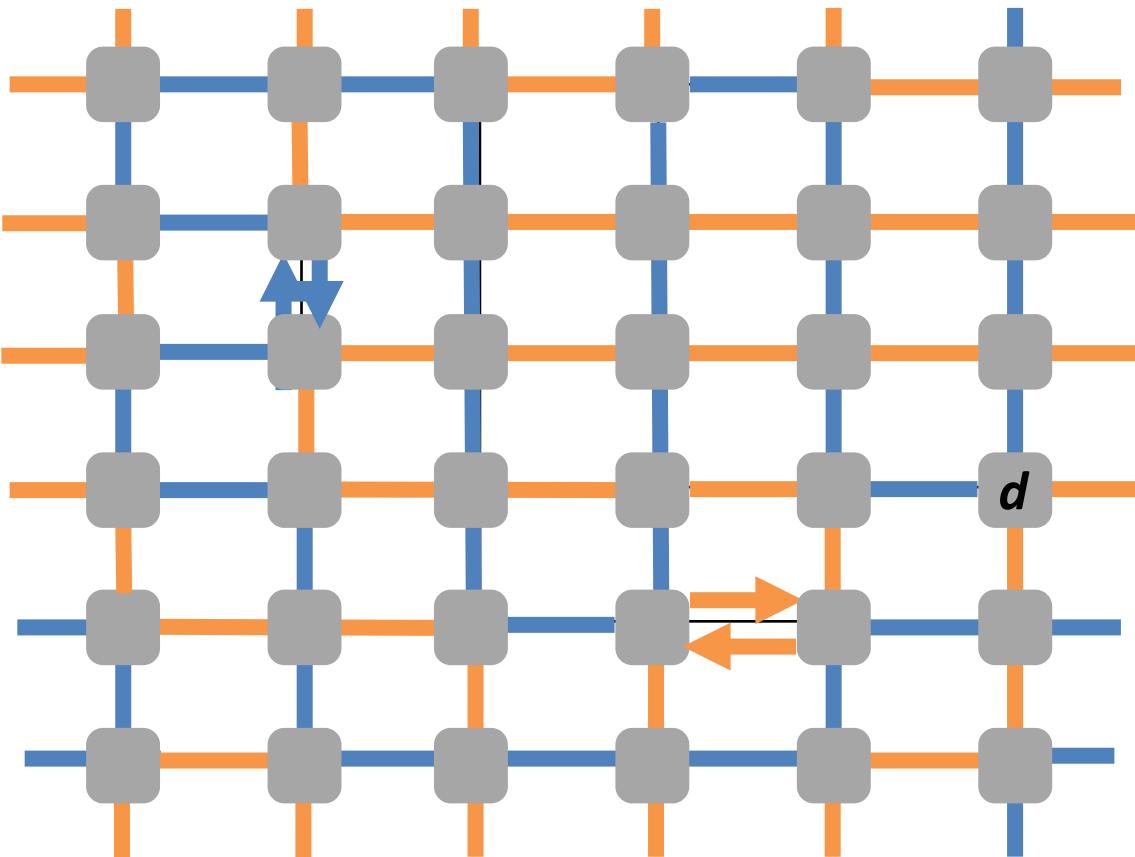
Idea: Decomposition into Hamilton Cycles



- Decompose torus into 2-edge-disjoint Hamilton Cycles (HC)

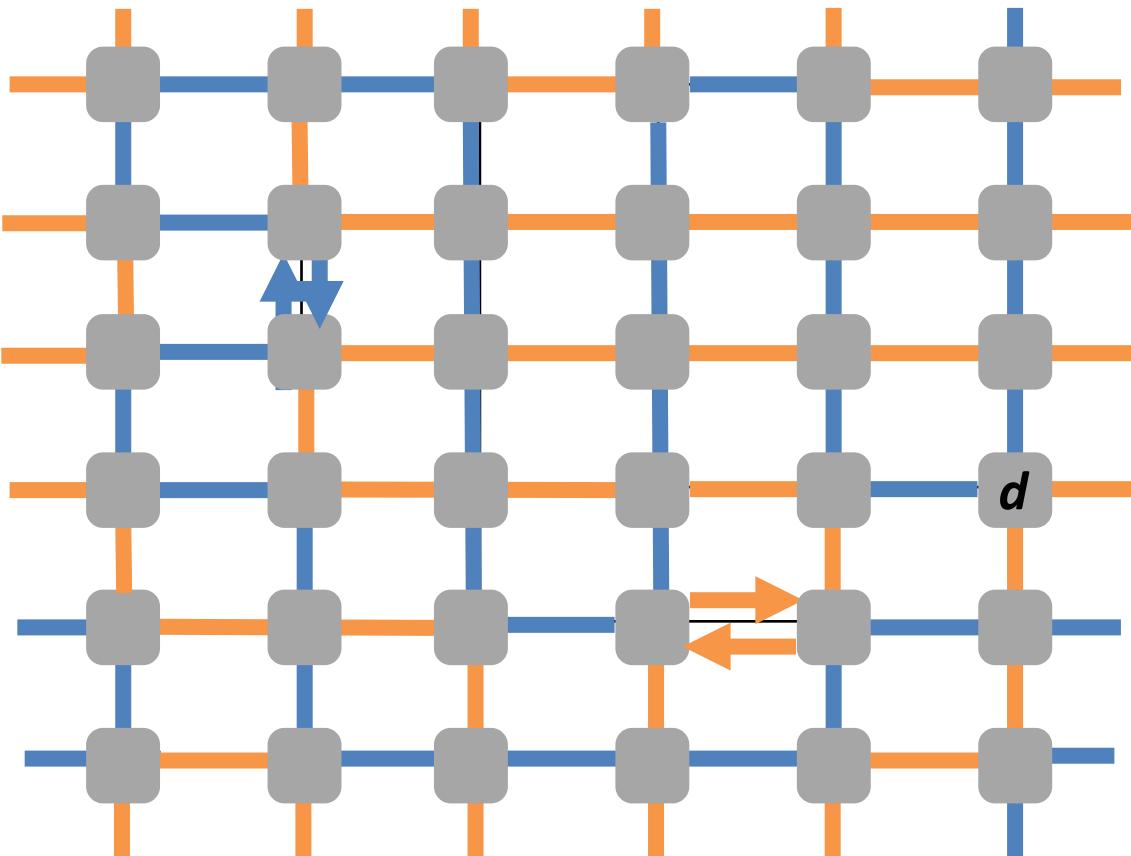
1st Hamilton cycle
2nd Hamilton cycle

Idea: Decomposition into Hamilton Cycles



- Decompose torus into 2-edge-disjoint Hamilton Cycles (HC)
- Can route in both directions:
4-arc-disjoint HCs

Idea: Decomposition into Hamilton Cycles



- Decompose torus into 2-edge-disjoint Hamilton Cycles (HC)
- Can route in both directions:
4-arc-disjoint HCs

3-resilient routing to destination d:

- go along *1st directed HC*, if hit failure, *reverse direction*
- if again failure switch to *2nd HC*, if again failure *reverse direction*
- No more failures possible!

Ideal Resilience with Hamilton Cycles

Chiesa et al.: if k -connected graph has k arc disjoint Hamilton Cycles, $k-1$ resilient routing can be constructed!

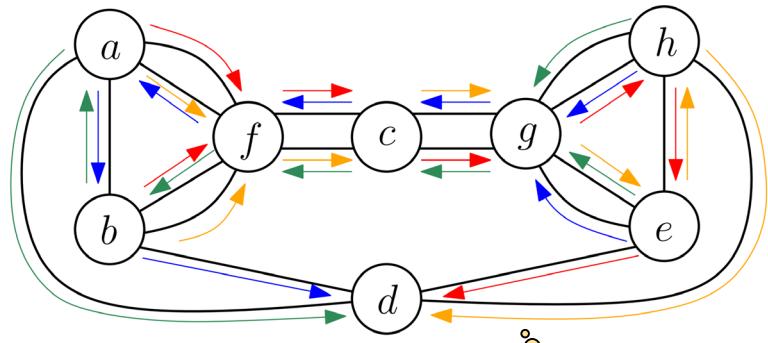
What about graphs which cannot be decomposed into Hamilton cycles?

Ideal Resilience in General k-Connected Graphs

- Use directed trees (i.e. *arborescences*) instead of Hamilton cycles
 - *Arc-disjoint*, spanning, and *rooted* at destination
- Classic result: k-connectivity guarantees k-arborescence decomposition

Basic idea:

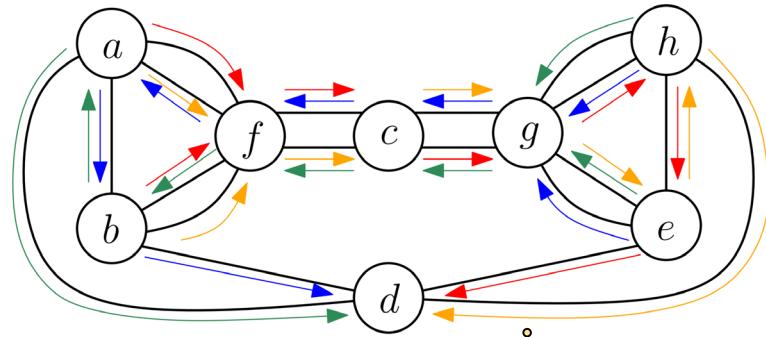
- Idea: route towards root on one arborescence
- After failure: change arborescence (e.g. in circular fashion)
- Incoming port defines current arborescence
- After $k-1$ failures: At least one arborescence intact



4-connected,
4 arborescences

Ideal Resilience in General k-Connected Graphs

- Use directed trees (i.e. *arborescences*) instead of Hamilton cycles
 - *Arc-disjoint*, spanning, and *rooted* at destination
- Classic result: k-connectivity guarantees k-arborescence decomposition



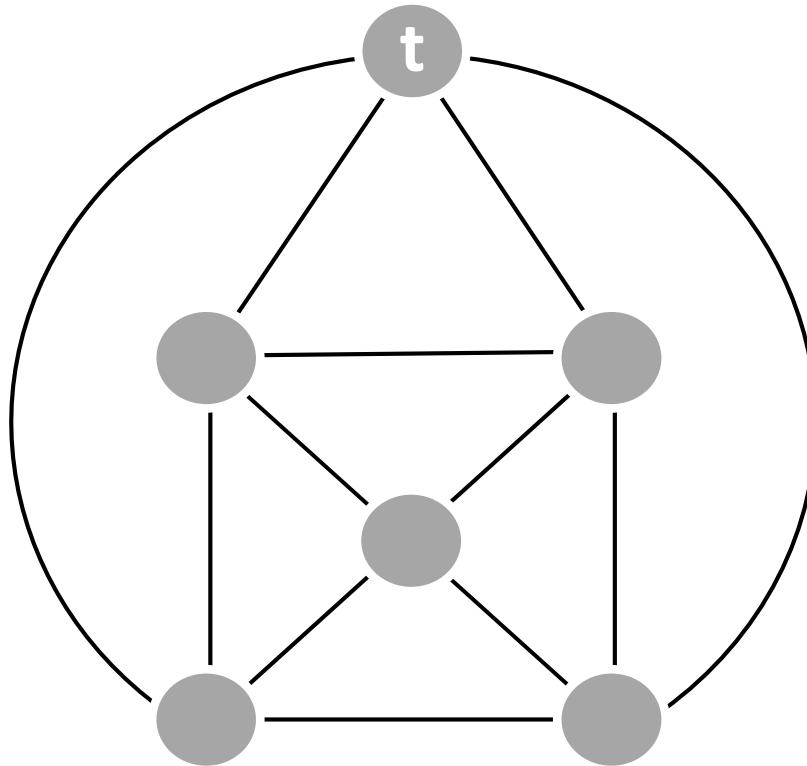
Basic idea:

- Idea: route towards root on one arborescence
- After failure: change arborescence (e.g. in circular fashion)
- Incoming port defines current arborescence
- After $k-1$ failures: At least one arborescence intact

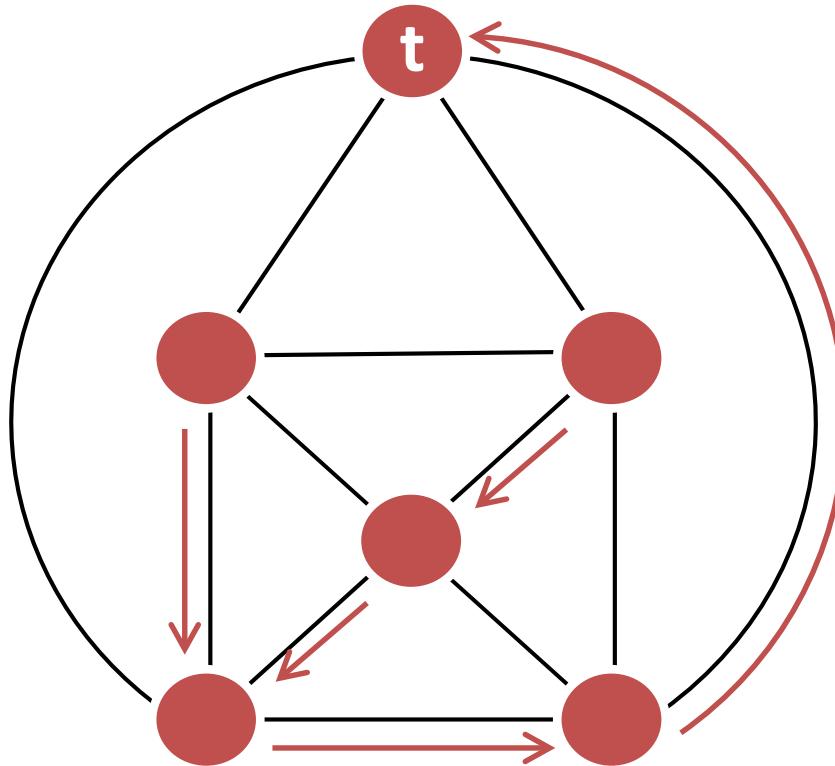
The challenge: how to avoid earlier tree?

4-connected,
4 arborescences

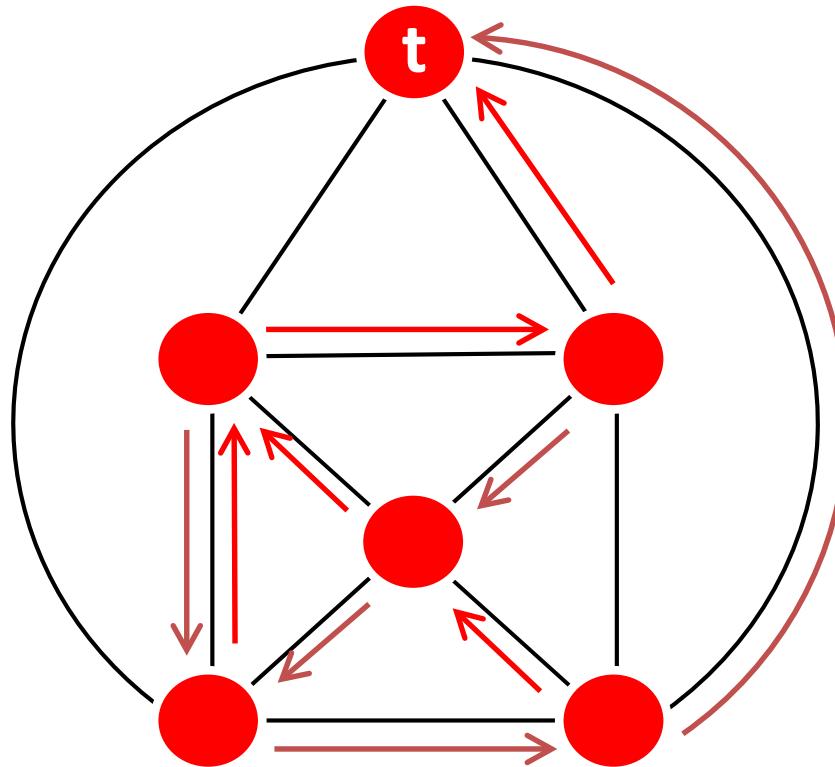
A k -connected network contains
 k arc-disjoint spanning arborescences [Edmonds, 1972]



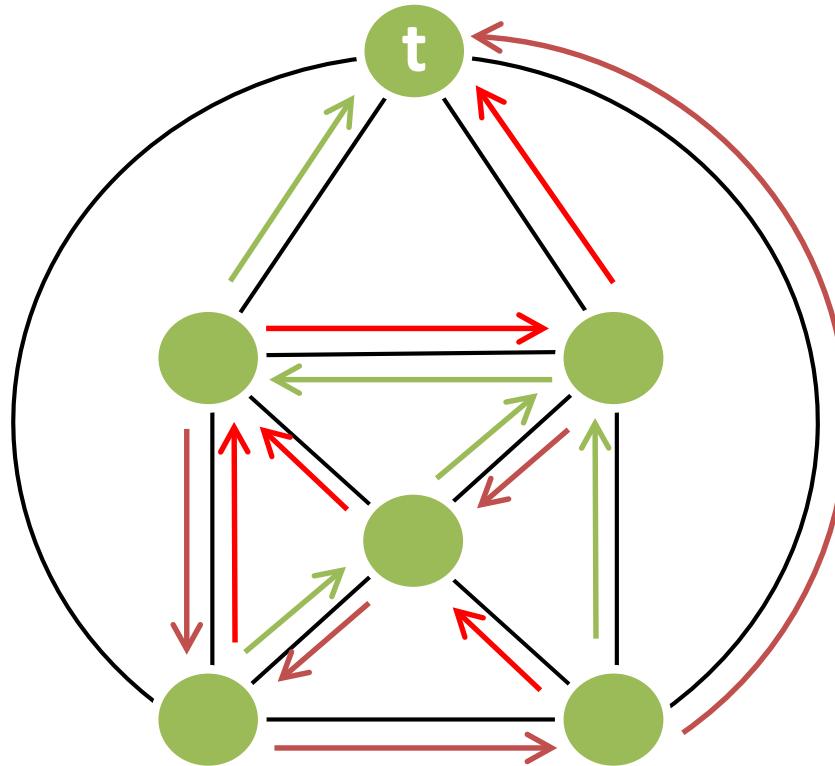
A k -connected network contains
 k arc-disjoint spanning arborescences [Edmonds, 1972]



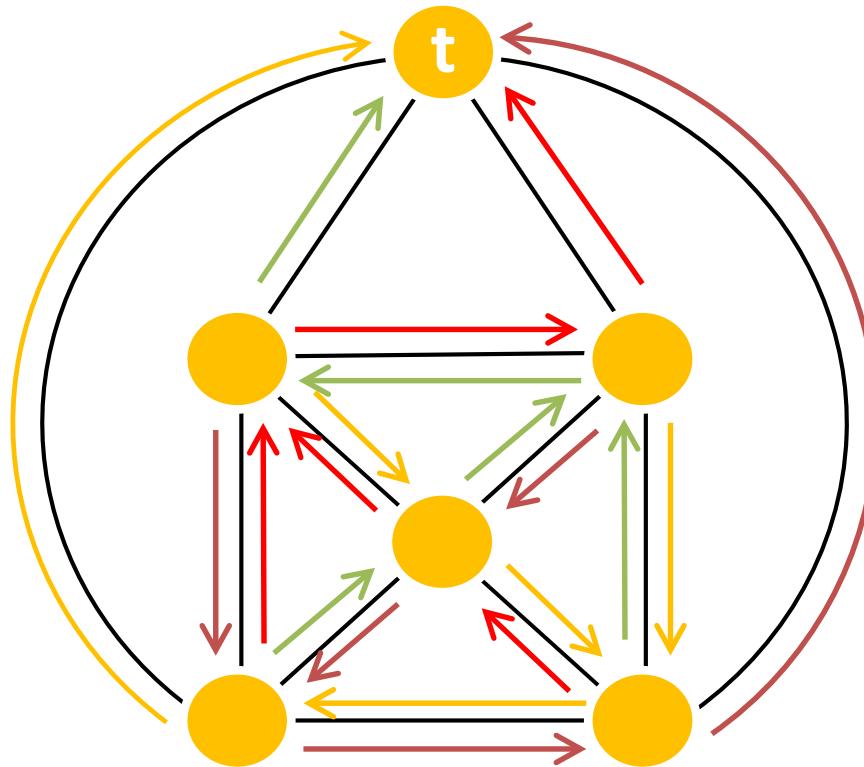
A k -connected network contains
 k arc-disjoint spanning arborescences [Edmonds, 1972]



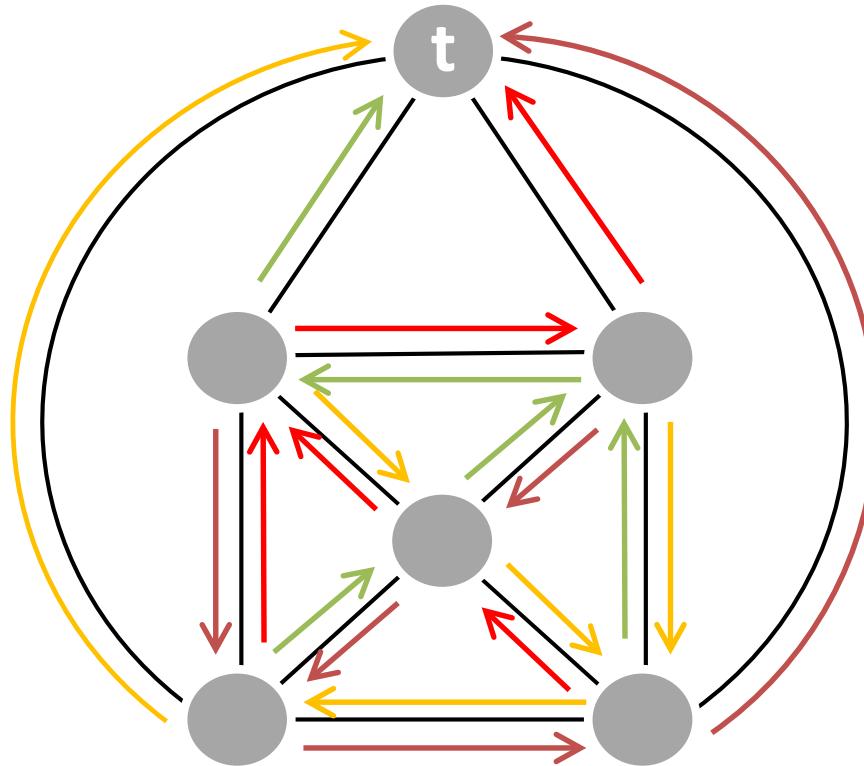
A k -connected network contains
 k arc-disjoint spanning arborescences [Edmonds, 1972]



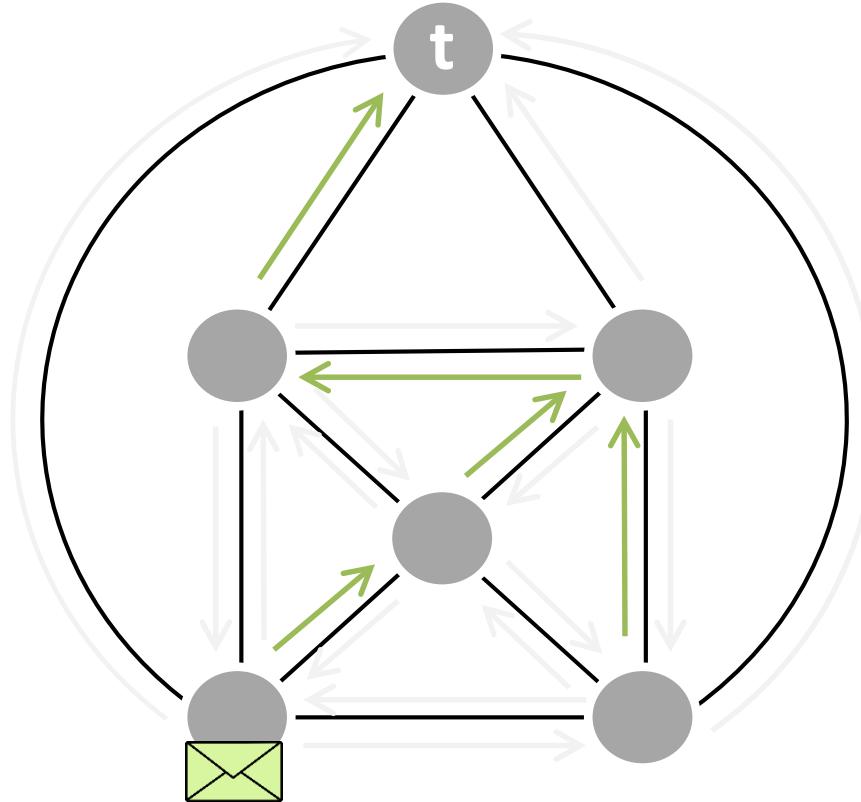
A k -connected network contains
 k arc-disjoint spanning arborescences [Edmonds, 1972]



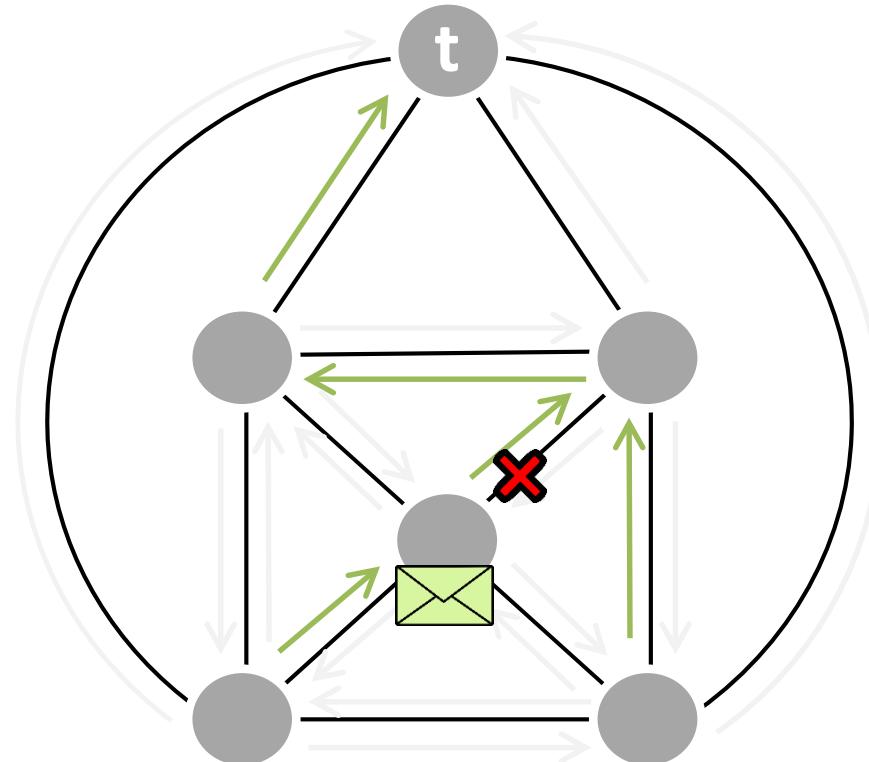
A k -connected network contains
 k arc-disjoint spanning arborescences [Edmonds, 1972]



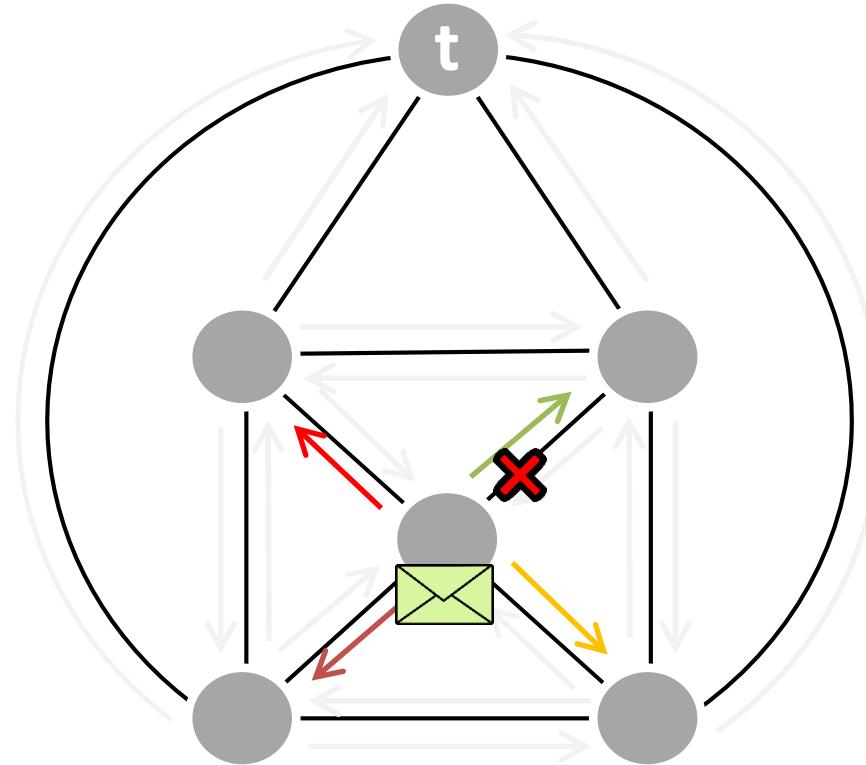
General technique: routing along the same tree



When a failed link is hit...



... how do we choose the next arborescence?



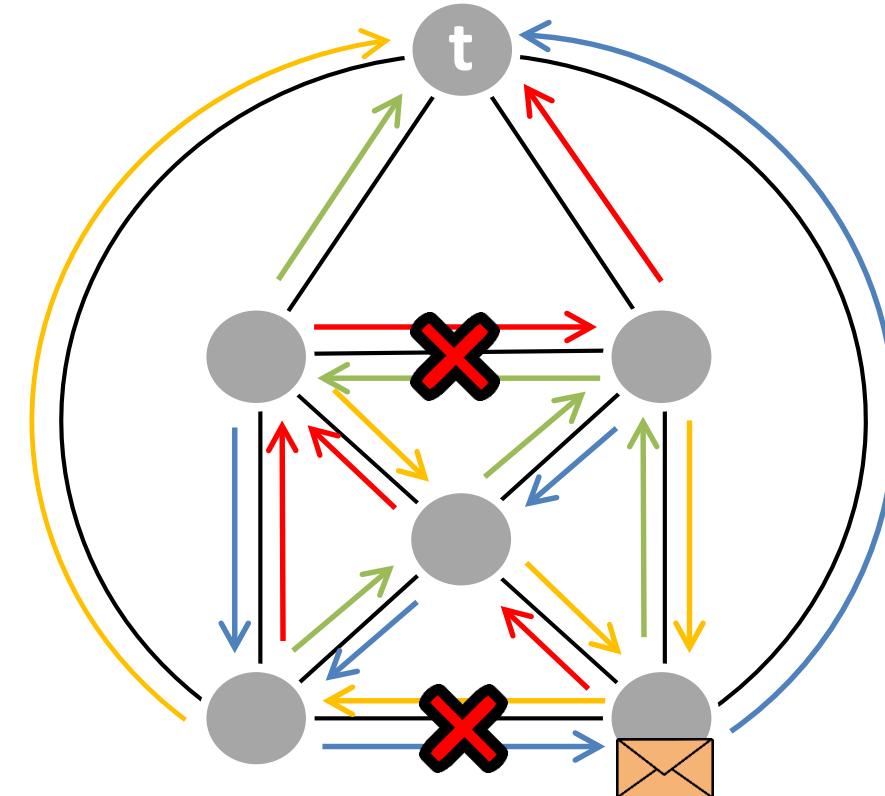
But how do we choose the next arborescence?

Circular-arborescence routing:

- compute an order of the arborescences
- switch to the next arborescence when hitting a failed link

Circular arborescence-routing is $(k/2-1)$ -resilient

Arborescence order



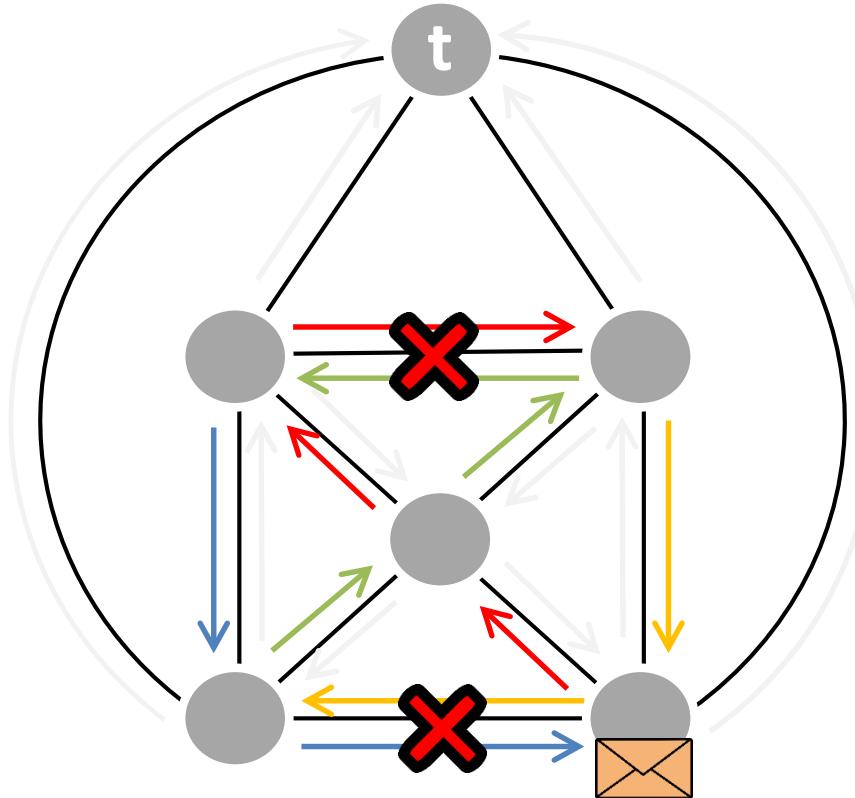
Intuition: each single failure may affect two arborescences

Circular arborescence-routing is $(k/2-1)$ -resilient

Arborescence order



*Go along arborescence 1
to destination...*



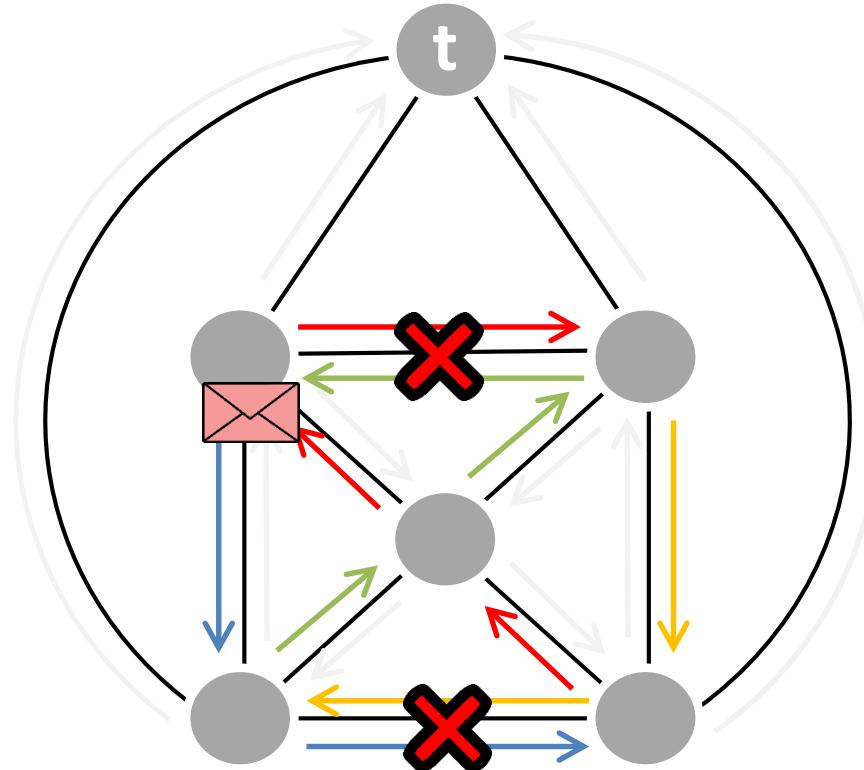
Intuition: each single failure may affect two arborescences

Circular arborescence-routing is $(k/2-1)$ -resilient

Arborescence order



Go along arborescence 2 to destination...



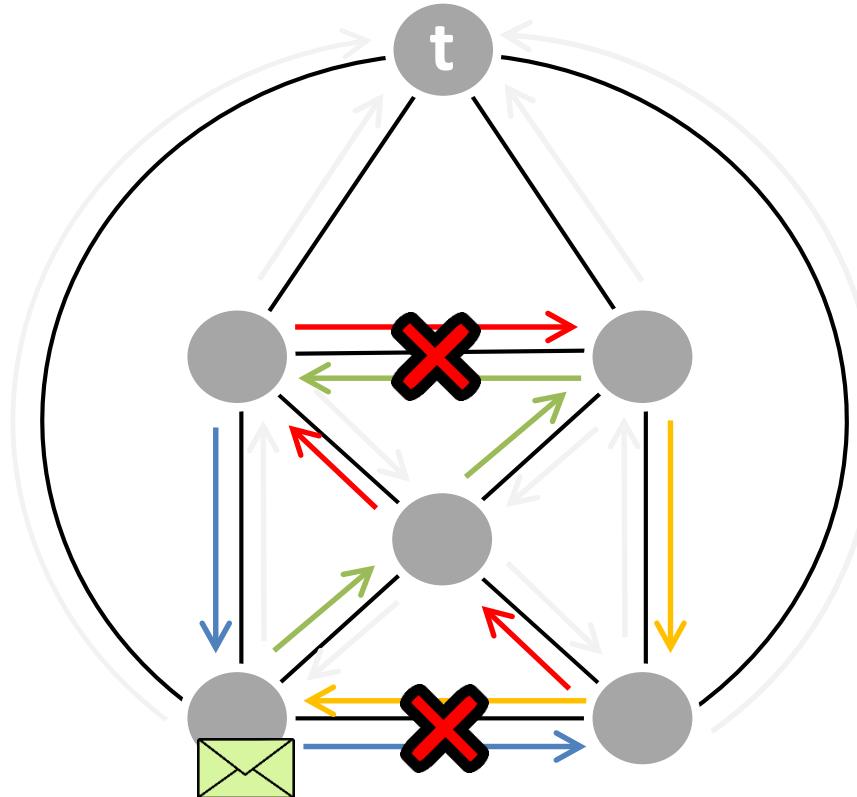
Intuition: each single failure may affect two arborescences

Circular arborescence-routing is $(k/2-1)$ -resilient

Arborescence order



Go along arborescence 3 to destination...



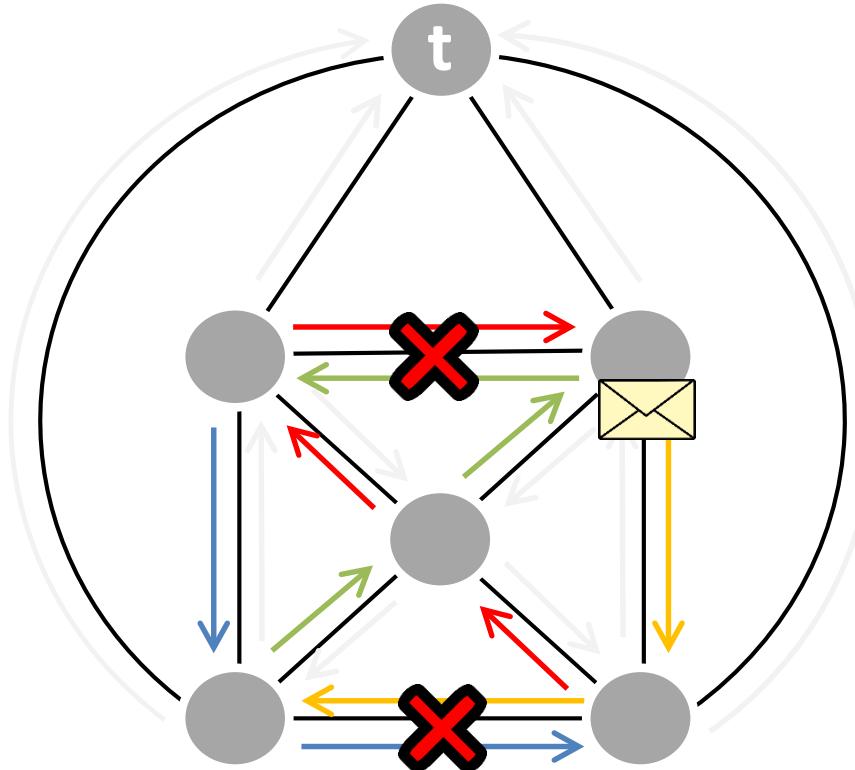
Intuition: each single failure may affect two arborescences

Circular arborescence-routing is $(k/2-1)$ -resilient

Arborescence order



Go along arborescence 4 to destination...

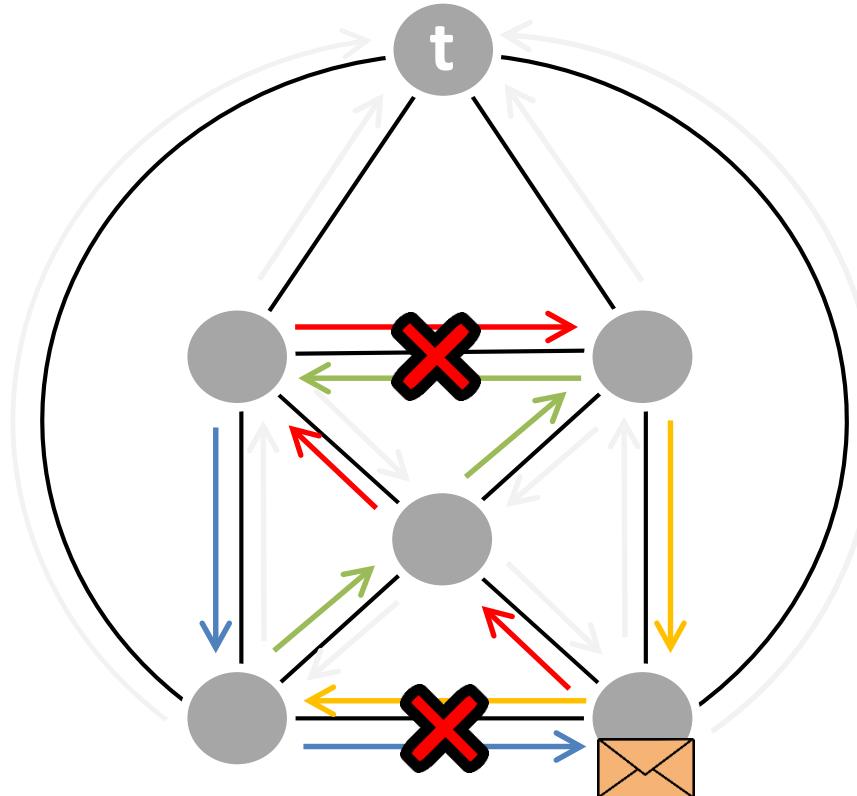


Intuition: each single failure may affect two arborescences

Circular arborescence-routing is $(k/2-1)$ -resilient

Arborescence order

1	2	3	4
---	---	---	---



Intuition: each single failure may affect two arborescences

All $k=4$ arborescences used
(2 failures disconnected affected all four):
LOOP!

An Alternative Algorithm: Bouncing Arborescence

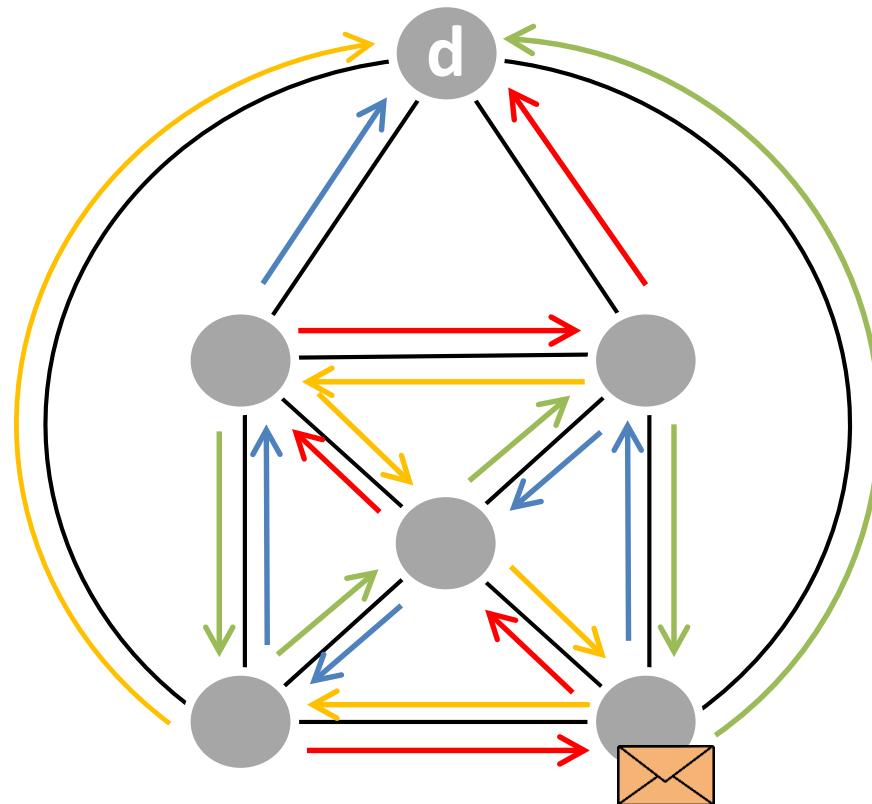
Bouncing-arborescence algorithm:

- Reroute on the tree that shares the failed link

This algorithm is ***1-resilient***.

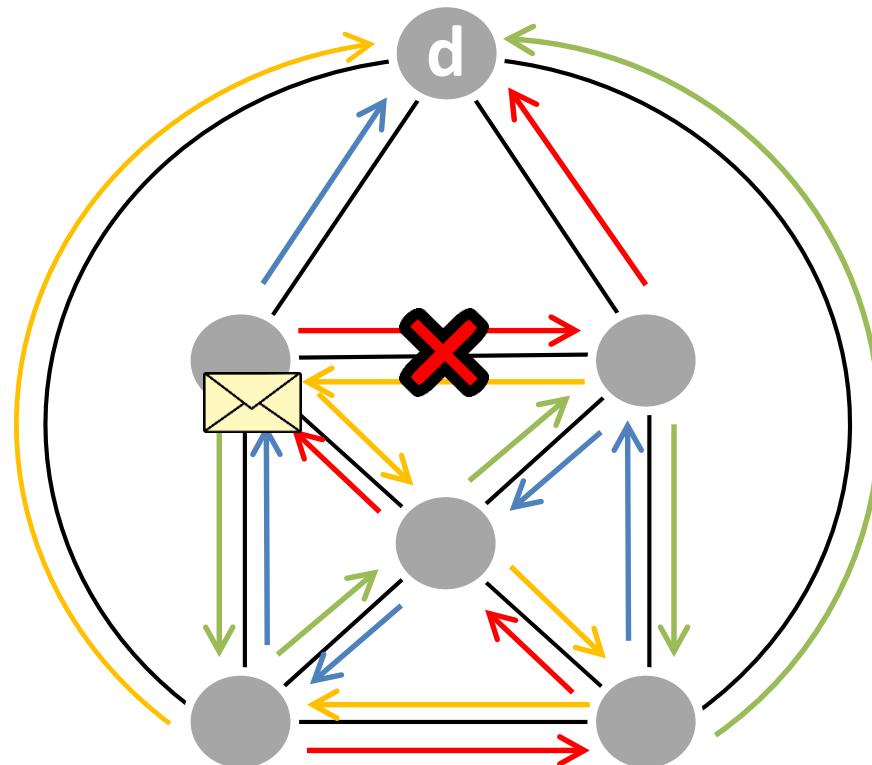
Bouncing-Arborescence is 1-Resilient

Start with red...



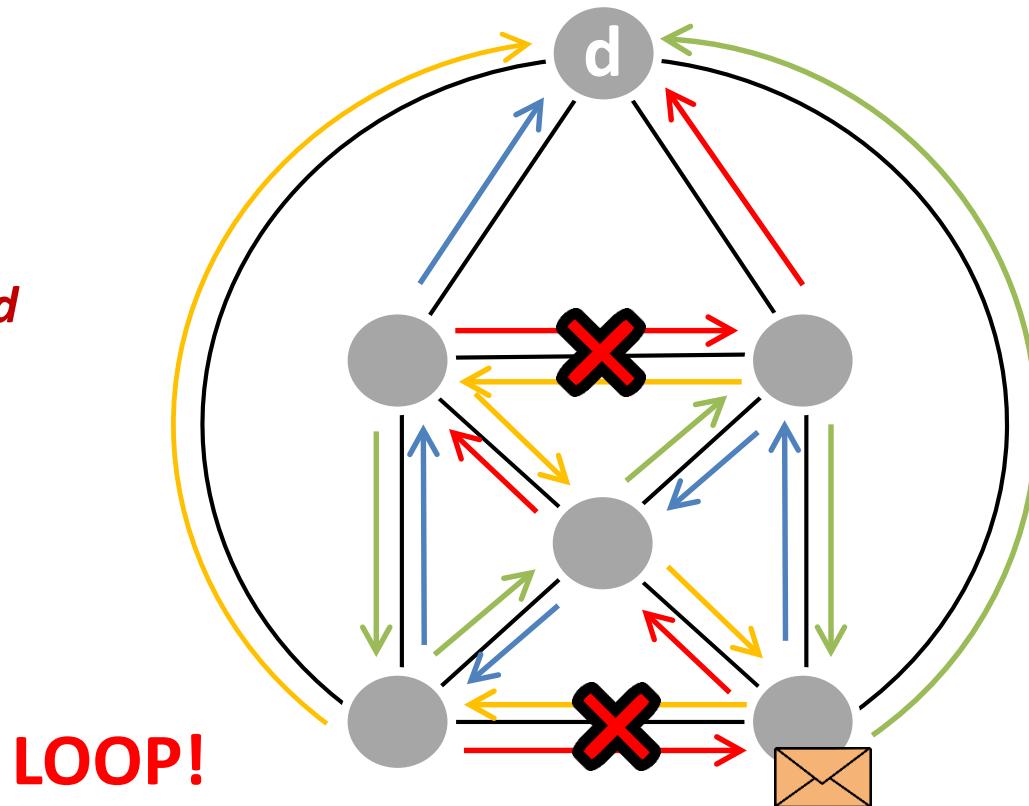
Bouncing-Arborescence is 1-Resilient

... bounce to yellow...



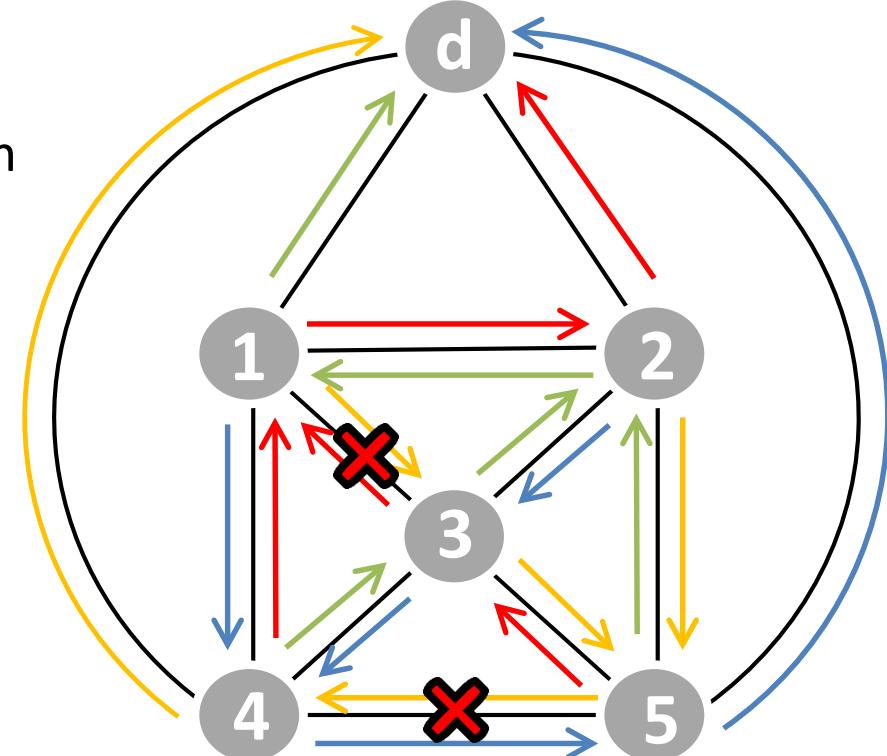
Bouncing-Arborescence is 1-Resilient

*... bounce to red
(again!)...*



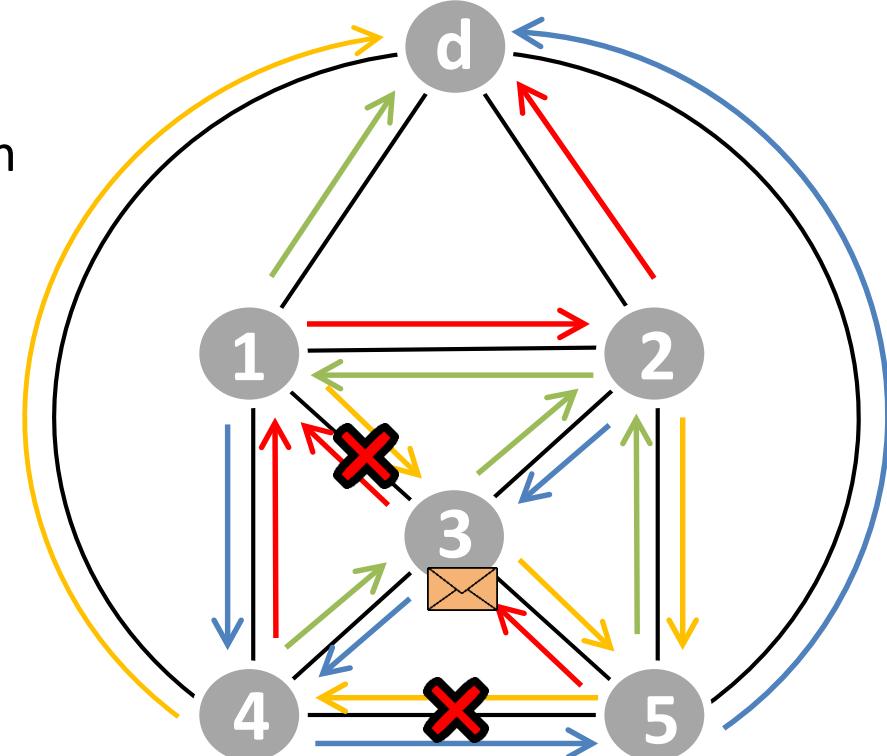
Idea: Bounce on „Good Arborescences“

- Define **well-bouncing arc**:
 - When bounce get to the destination
 - Without hitting any other failures



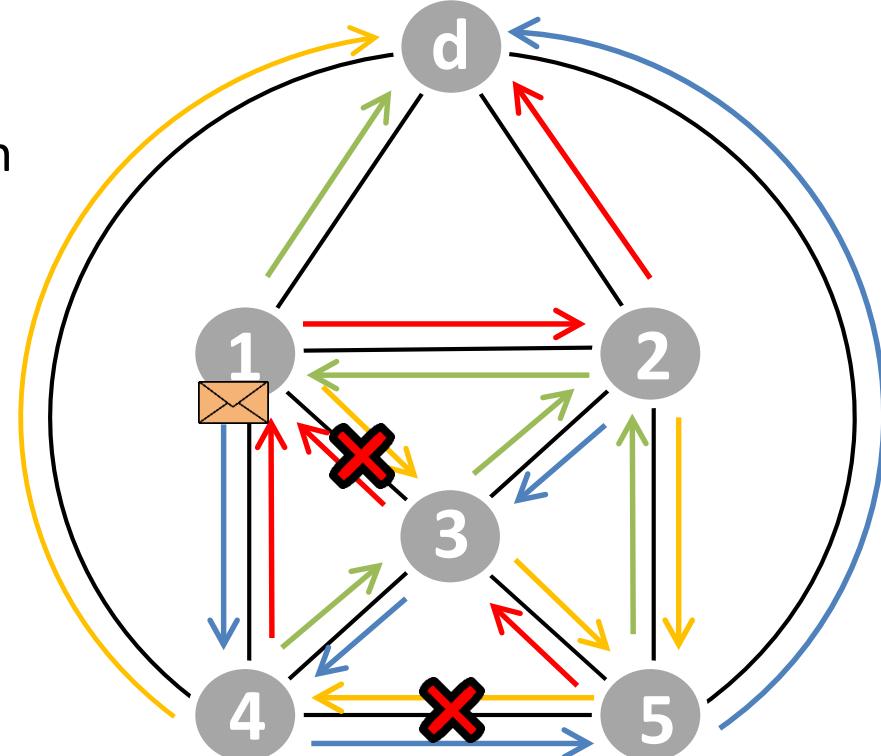
Idea: Bounce on „Good Arborescences“

- Define **well-bouncing arc**:
 - When bounce get to the destination
 - Without hitting any other failures
 - (3,1) is not well-bouncing



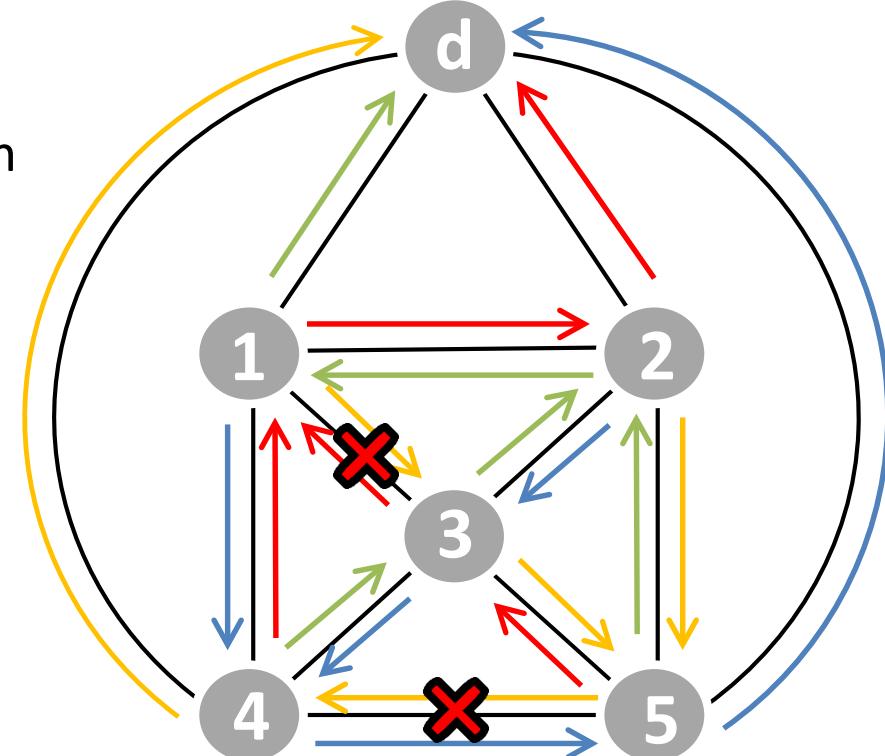
Idea: Bounce on „Good Arborescences“

- Define **well-bouncing arc**:
 - When bounce get to the destination
 - Without hitting any other failures
 - (3,1) is not well-bouncing
 - (1,3) is well-bouncing



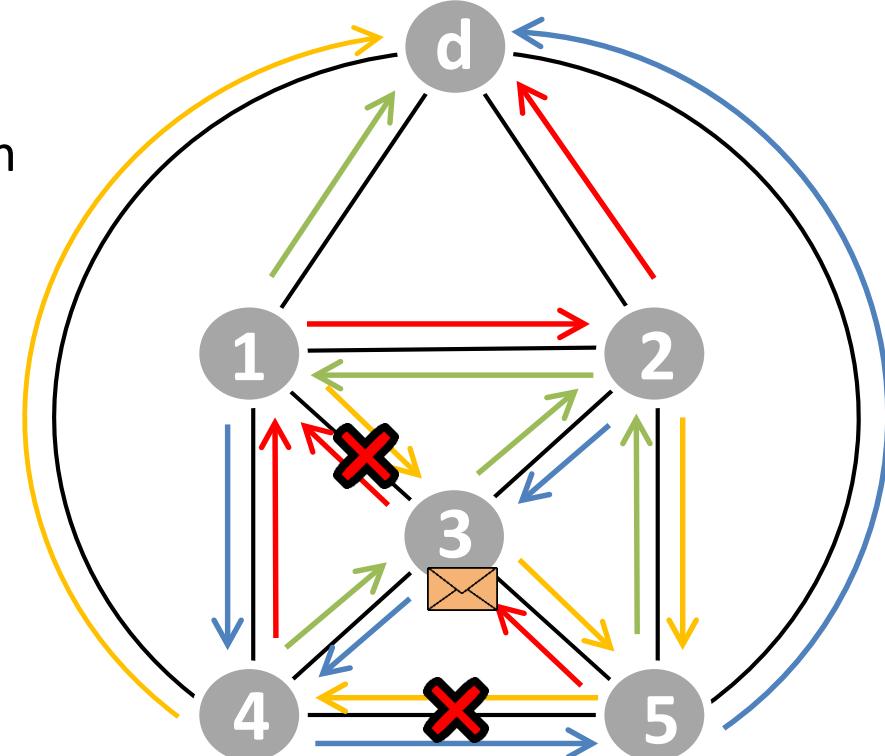
Idea: Bounce on „Good Arborescences“

- Define **well-bouncing arc**:
 - When bounce get to the destination
 - Without hitting any other failures
 - (3,1) is not well-bouncing
 - (1,3) is well-bouncing
- Define **good arborescence**:
 - every failed arc is well-bouncing



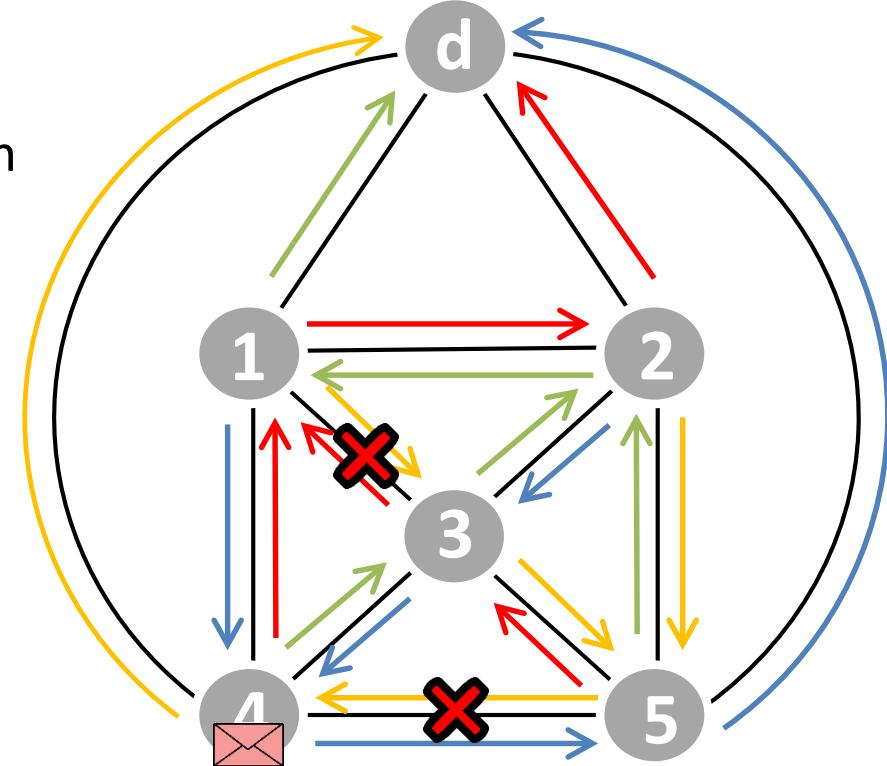
Idea: Bounce on „Good Arborescences“

- Define **well-bouncing arc**:
 - When bounce get to the destination
 - Without hitting any other failures
 - (3,1) is not well-bouncing
 - (1,3) is well-bouncing
- Define **good arborescence**:
 - every failed arc is well-bouncing
 - Red is not a good arborescence



Idea: Bounce on „Good Arborescences“

- Define **well-bouncing arc**:
 - When bounce get to the destination
 - Without hitting any other failures
 - (3,1) is not well-bouncing
 - (1,3) is well-bouncing
- Define **good arborescence**:
 - every failed arc is well-bouncing
 - Red is not a good arborescence
 - Blue is a good arborescence



Ideas

- One can show that there is always a good arborescence
- An tempting idea:
 - route on an arborescence X until a failed link is hit:
 - if X is a good arborescence, bounce!
 - otherwise, route circular
- Too good to be true:
 - The “goodness” of an arborescence depends on the actual set of failed links!
 - How do we know a arborescence is good?

Resilience Criteria

Ideal resilience

Given a k -connected graphs, we can tolerate *any $k-1$ link failures*.

Perfect resilience

Any source s can always reach any destination t as long as the underlying network is *physically connected*.

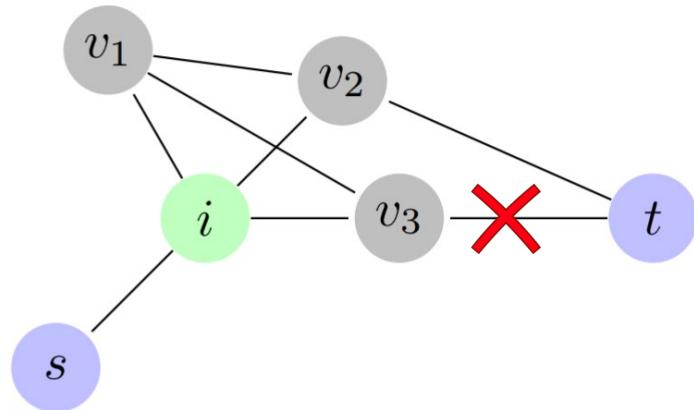
Can this be achieved? Assume undirected link failures.

Resilience Criteria

Perfect resilience is impossible to achieve in general.

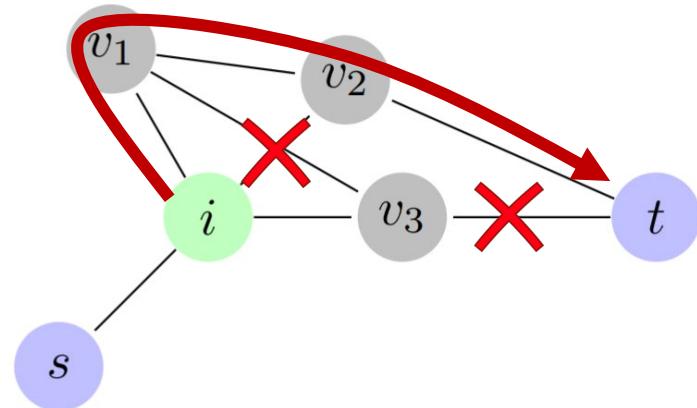
Relevant Neighbors

- Routing table of node i : matches in-ports of i to out-ports of i
 - ... depending on the incident failures
- But not all neighbors are **relevant**: only if potentially required to reach destination!
 - *Without local failures*: just v_2, v_3 for i , since v_1 does not give extra connectivity



Relevant Neighbors

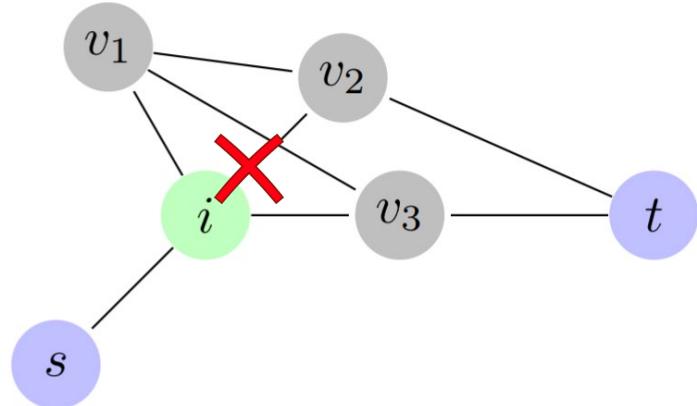
- Routing table of node i : matches in-ports of i to out-ports of i
 - ... depending on the incident failures
- But not all neighbors are **relevant**: only if potentially required to reach destination!
 - *Without local failures*: just v_2, v_3 for i , since v_1 does not give extra connectivity
 - *With additional failures* v_1 becomes relevant, since v_1 might be only choice to reach destination t
 - Note: v_1 is unaware of these non-incident failures!



High-level definition of **relevant**: From the local view-point of the node i , a relevant neighbor might be only neighbor to reach destination (without taking a detour over a current neighbor).

How to Achieve Perfect Resilience?

- Necessary: need to *try all relevant* neighbors
 - Here, if local link to v_2 broken:
 v_1 and v_3
- That is, if packet
 - comes from v_3 : eventually try v_1
 - comes from v_1 : eventually try v_3



Impossibility: On Planar Graphs

Some observations:

- Additional failures only ***add relevant neighbors*** to nodes
- Any node of ***degree 2*** of G after failures must forward packets with incoming port p to port p'
- If all neighbors are relevant, the forwarding function of a node must be a ***cyclic permutation***

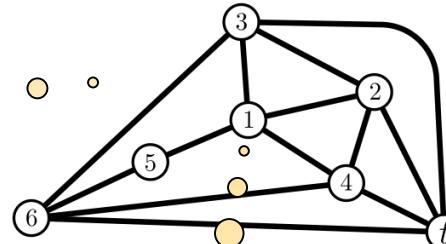
Impossibility: On Planar Graphs

Some observations:

- Additional failures only **add relevant neighbors** to nodes
- Any node of **degree 2** of G after failures must forward packets with incoming port p to port p'
- If all neighbors are relevant, the forwarding function of a node must be a **cyclic permutation**

Idea of the counter example:

All neighbors of all nodes are relevant (even without failures).



So we must fix a
permutation for node 1.

Considered node 1 will not
see any local failures.

Impossibility: On Planar Graphs

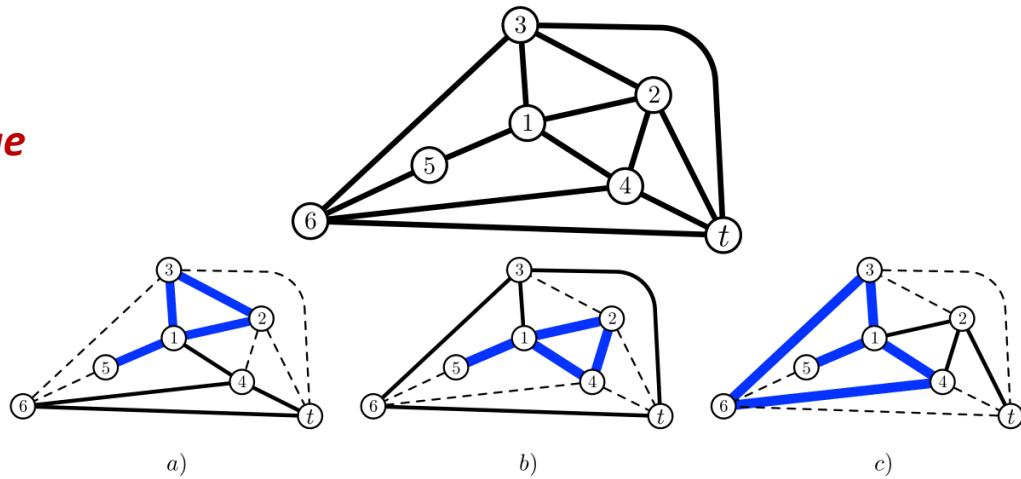
Some observations:

- Additional failures only **add relevant neighbors** to nodes
- Any node of **degree 2** of G after failures must forward packets with incoming port p to port p'
- If all neighbors are relevant, the forwarding function of a node must be a **cyclic permutation**

Proof idea, with three cases:

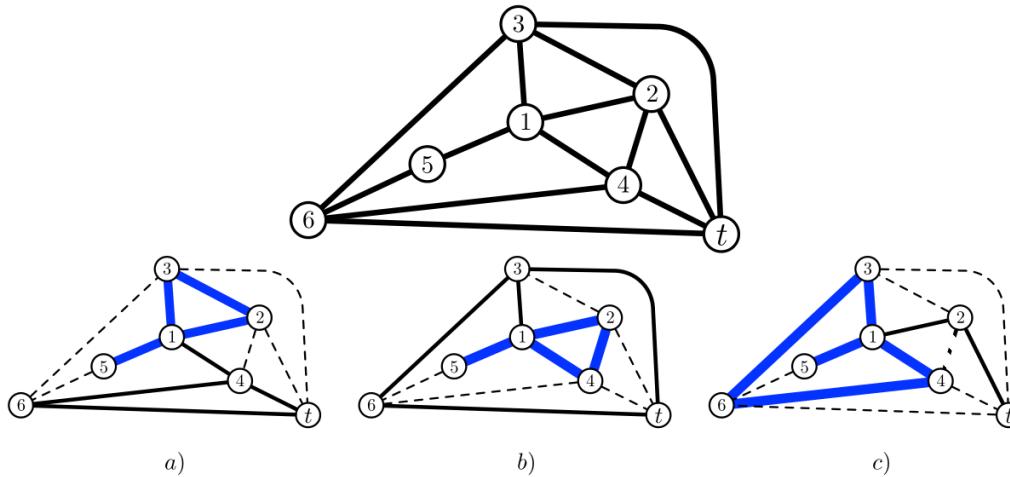
- If the **dashed** links fail (**non-local** to node 1), in any forwarding pattern, packets will be stuck in one of the **blue loops**...
- ... even though there is at least one **remaining path** to the target

Go through all possible permutations @1 and give counter example.



Impossibility: On Planar Graphs

Arriving on
input 5,
forwarded
to 2.

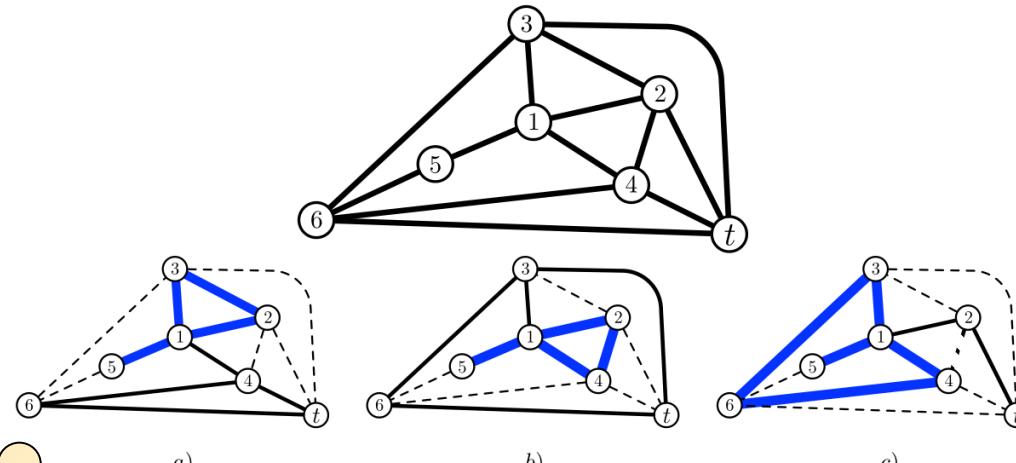


node 1:
5 → 2 implies
(5,2,3,4) (b)
(5,2,4,3) (a)

Possible cyclic permutations: when a packet arrives from 2, due to cyclic permutation, it can only be forwarded to either 3 or 4. Leads to **loops** in scenarios (b) (4 goes to 5, 2 can only go to 4) and (a) (3 goes to 5, 2 can only go to 3), respectively.

Impossibility: On Planar Graphs

Arriving on
input 5,
forwarded
to 3.

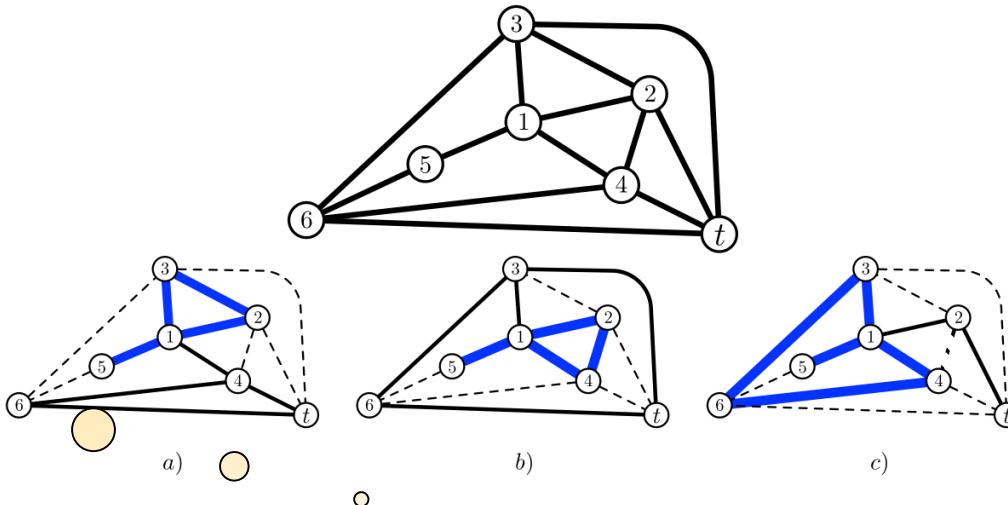


For node 1:
5->2 implies
(5,2,3,4) (b)
(5,2,4,3) (a)

For node 1:
5->3 implies
(5,3,4,2) (a)
(5,3,2,4) (c)

Possible cyclic permutations: when a packet then arrives on port 4, it can only be forwarded to either 2 or 5. Leads to *loops* in scenarios (a) (2 will go to 5, 5 can only go to 1 and 3 only to 2) and (c) (5 goes to 3, 4 goes to 5, rest degree-2), respectively.

Impossibility: On Planar Graphs



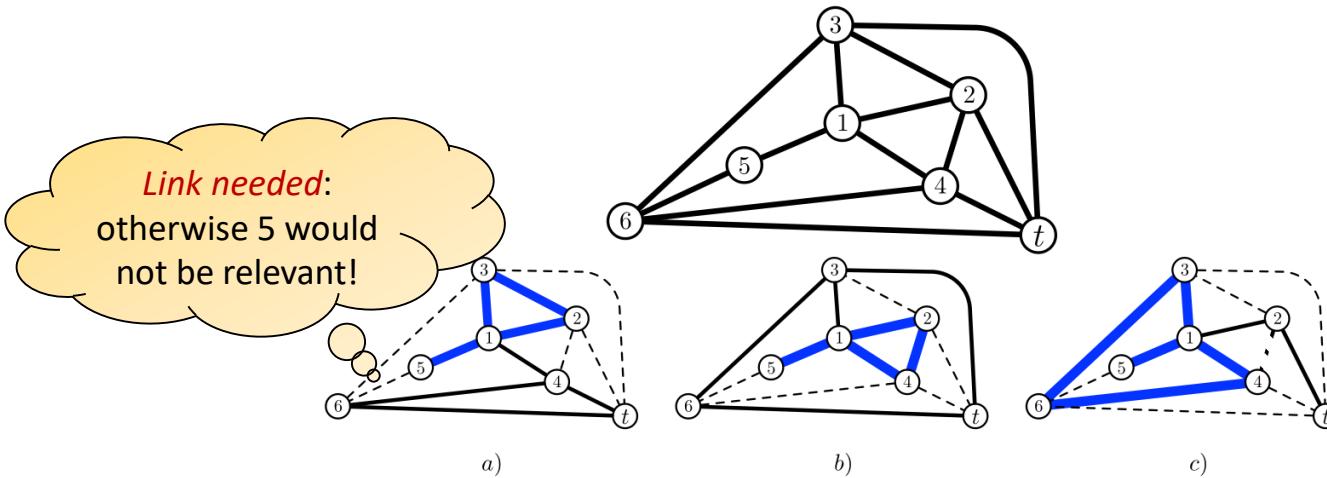
For node 1:
5->2 implies
(5,2,3,4) (b)
(5,2,4,3) (a)

For node 1:
5->3 implies
(5,3,4,2) (a)
(5,3,2,4) (c)

For node 1:
5->4 implies
(5,4,2,3) (c)
(5,4,3,2) (b)

Possible cyclic permutations: packet arriving on port 3 can only be forwarded to either 5 or 2. Leads to *loops* in scenarios (c) and (b), respectively.

Impossibility: On Planar Graphs



For node 1:
5->2 implies
(5,2,3,4) (b)
(5,2,4,3) (a)

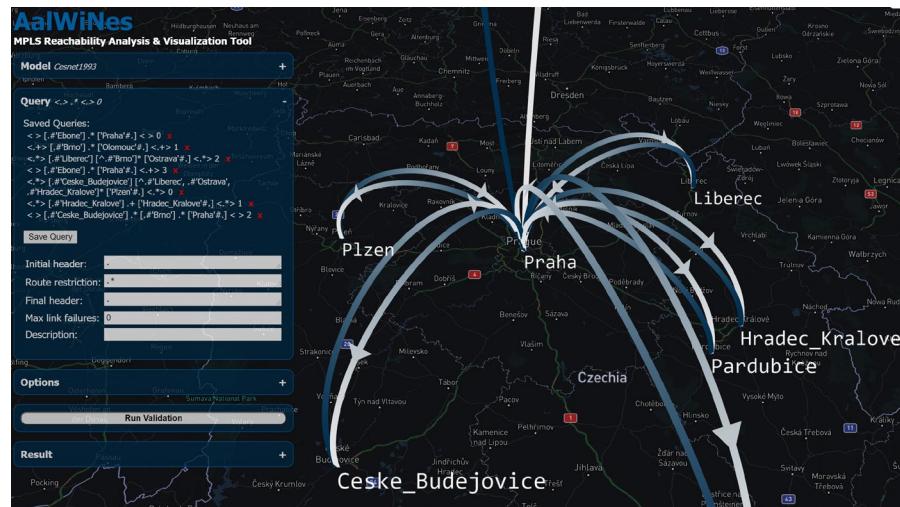
For node 1:
5->3 implies
(5,3,4,2) (a)
(5,3,2,4) (c)

For node 1:
5->4 implies
(5,4,2,3) (c)
(5,4,3,2) (b)

Possible cyclic permutations: packet arriving on port 3 can only be forwarded to either 5 or 2. Leads to *loops* in scenarios (c) and (b), respectively.

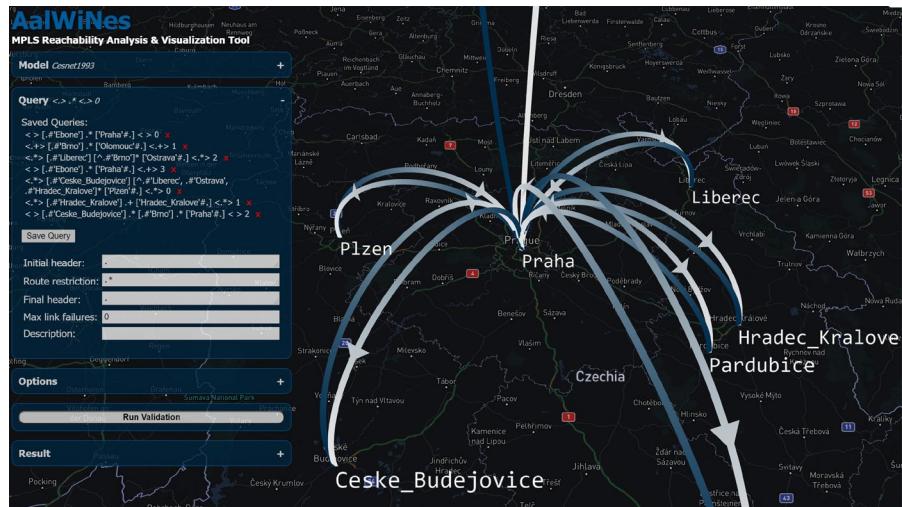
A Pity: Planar Graphs Are Important

- Internet Topology Zoo and Rocketfuel topologies
 - 88% of the graphs are *planar*



A Pity: Planar Graphs Are Important

- Internet Topology Zoo and Rocketfuel topologies
 - 88% of the graphs are **planar**
 - However:
 - Almost a third (32%) belong to the family of **cactus** graphs
 - Roughly half of the graphs (49%) are **outerplanar**
 - ... and they work 😊



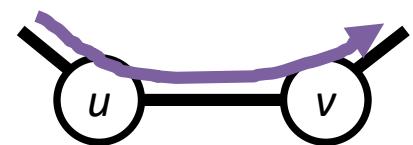
Where Can Perfect Resilience Be Achieved?

For example on **outerplanar graphs**:

- Via *geometric routing*, well studied in sensor networks etc.
- Embed graph in the plane s.t. all nodes are on the outer face
 - Note: If a link l belongs to the outer face of a planar graph G , it also belongs to the outer face for all subgraphs of G
- Apply *right-hand rule* to forwarding (skipping failures)
 - Ensures packets use only the links of the outer face and do not change the direction despite failures
- Strategy traverses all nodes on the outer face
- Also works for any graph which is *outerplanar without the source* (e.g., K4)

Some Observations

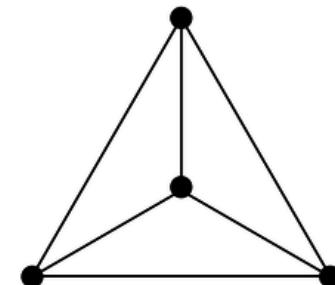
- $K_5, K_{3,3}$: *no perfect resilience*
- Perfect resiliency on graph $G \rightarrow$ any *subgraph* G' of G also allows for perfect resiliency
 - Idea: Take routing on G , fail edges to create G' , routing must still work
- **Contraction** works as well, by a simulation argument
 - A bit technical
- Combined: Perfect resilience on graph $G \rightarrow$ any minor G' of G as well
 - But since $K_5, K_{3,3}$ not: *non-planar graphs not perfectly resilient*



What we know about perfect resilience

Possible:

- On all outerplanar graphs [right-hand rule]
- On every graph that is outerplanar without the destination (e.g. non-outerplanar planar K_4)



Impossible:

- On some planar graphs
- Every non-planar graph
- Perfect resilience must hold on minors



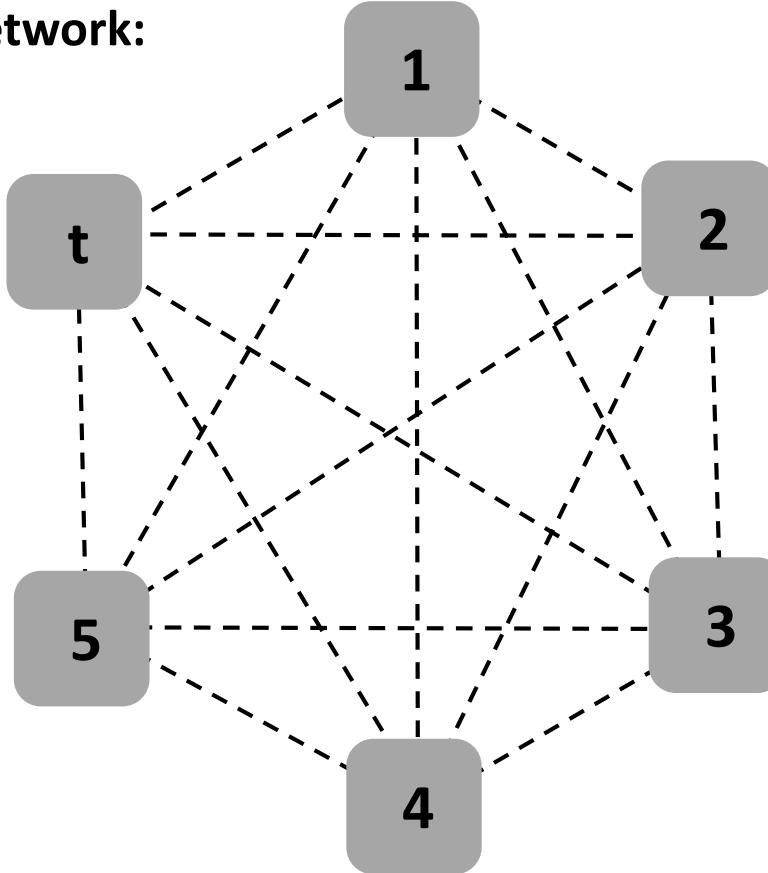
Roadmap

- A Brief History of Resilient Networking
- Algorithms for Local Fast Re-Routing (FRR)
- **Accounting for Congestion**
- Accounting for Network Policy



Congestion-Aware FRR

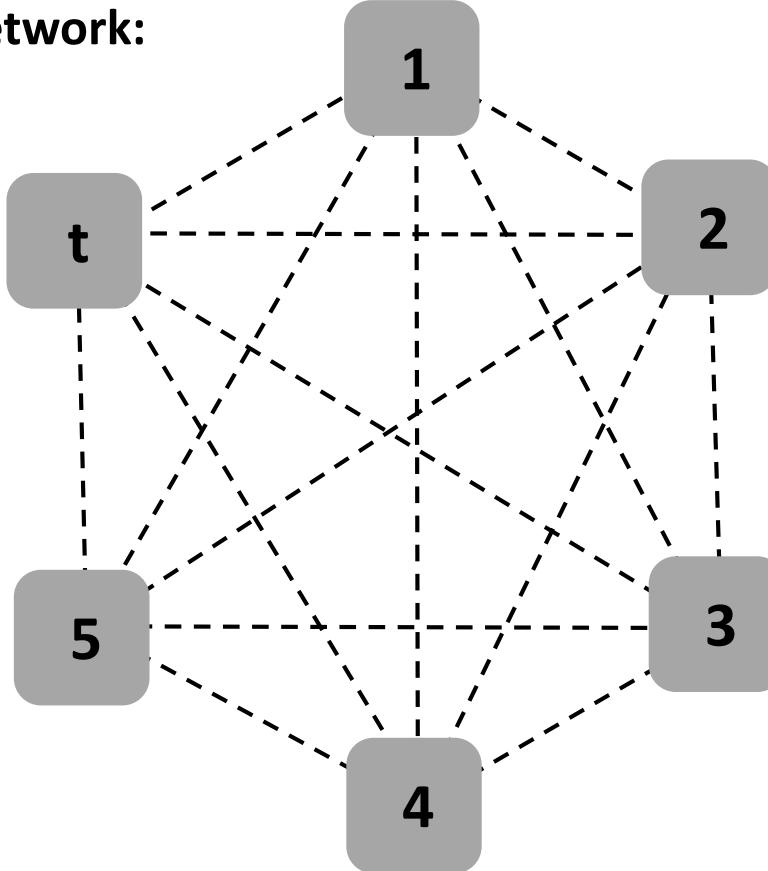
A most simple network:
the clique



Congestion-Aware FRR

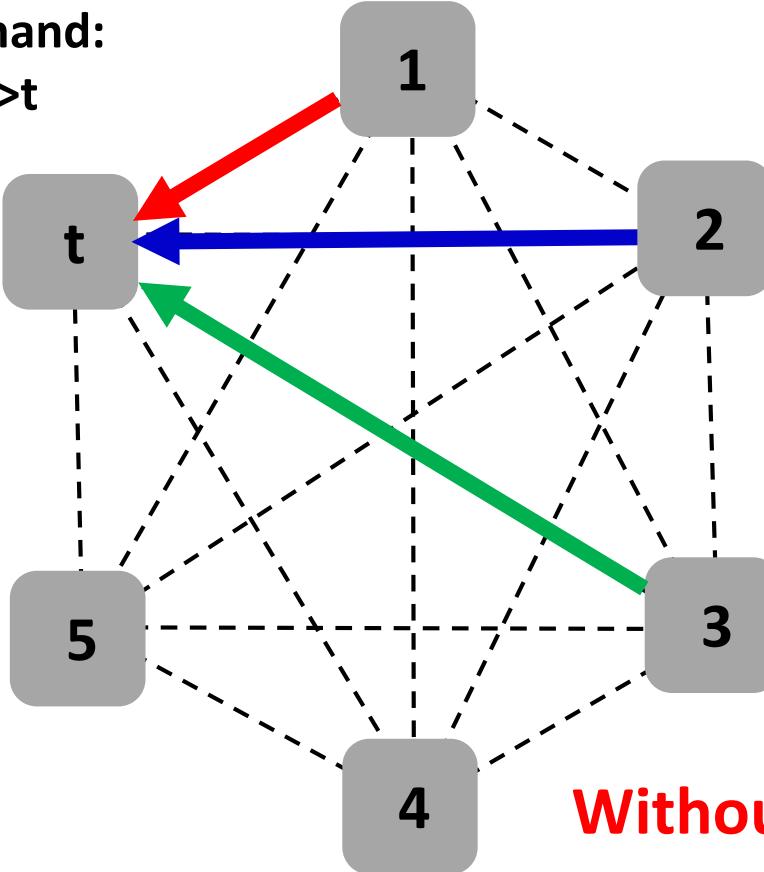
A most simple network:
the clique

Assume we can
match source.



Congestion-Aware FRR

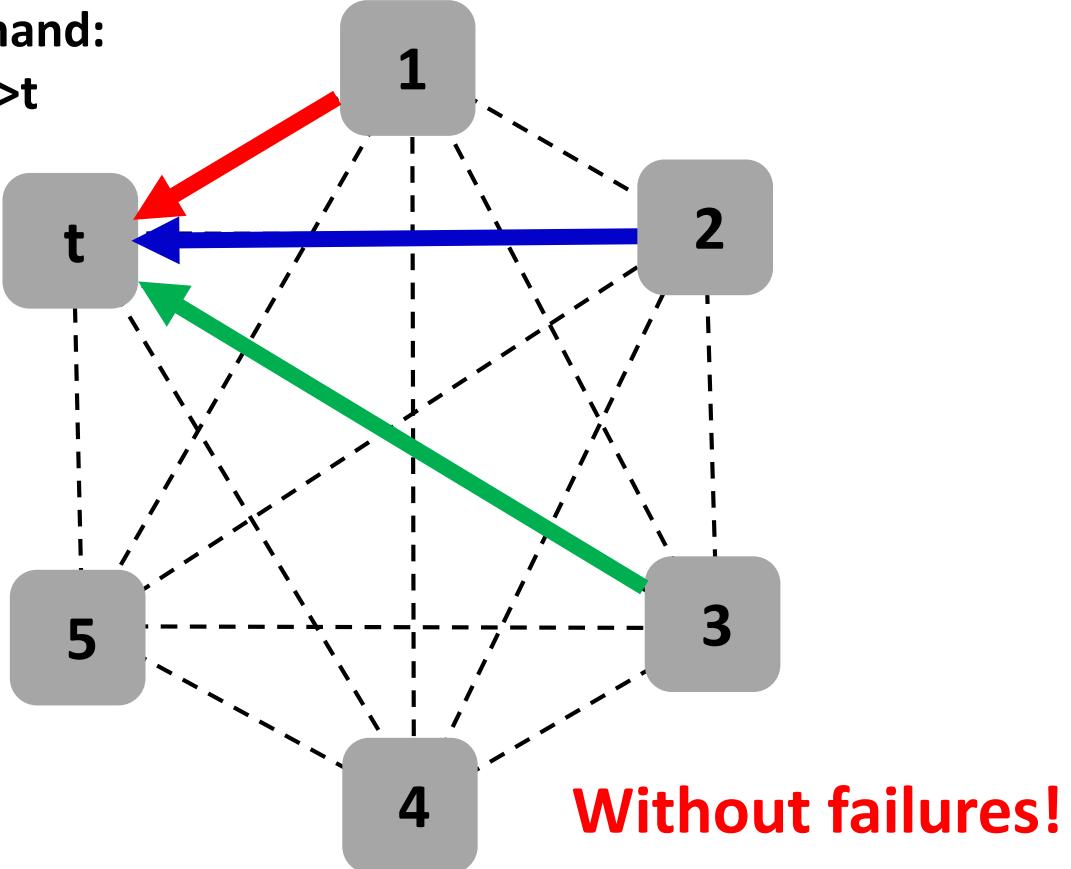
Traffic demand:
 $\{1,2,3\} \rightarrow t$



Congestion-Aware FRR

Traffic demand:
 $\{1,2,3\} \rightarrow t$

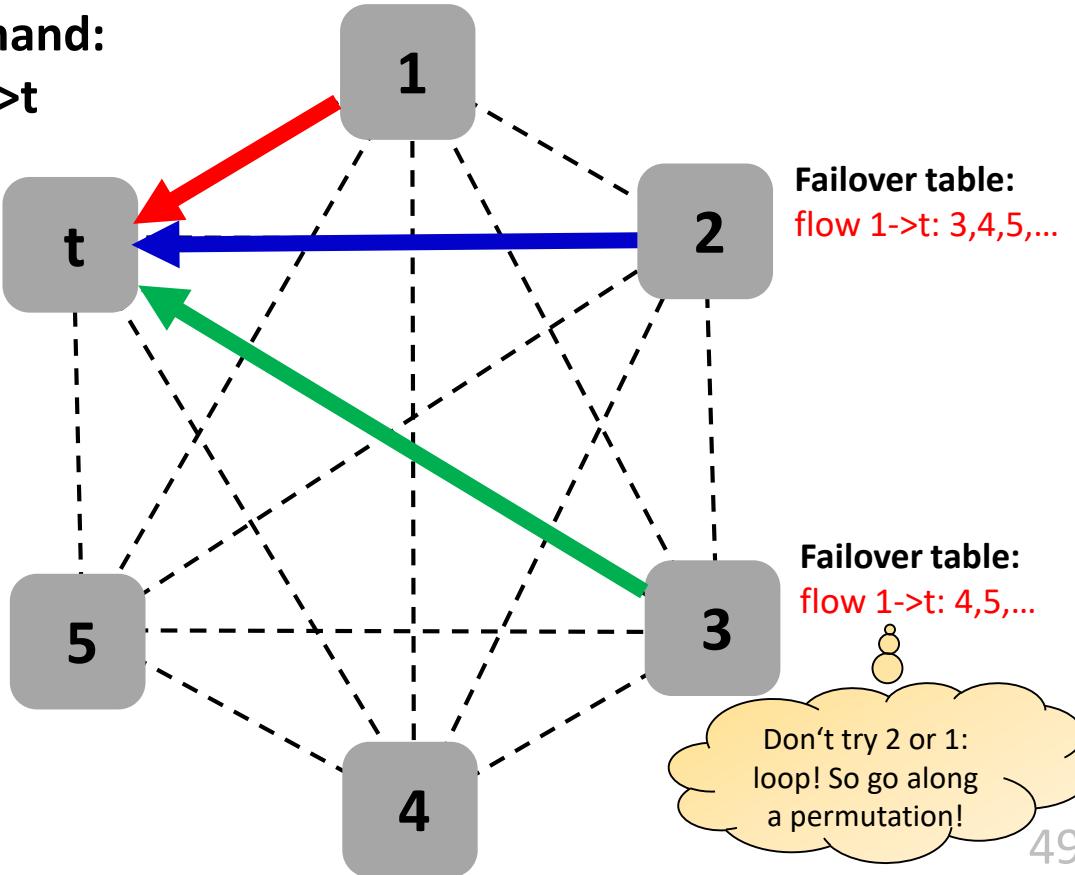
Assume single destination
(incast scenario).



Congestion-Aware FRR

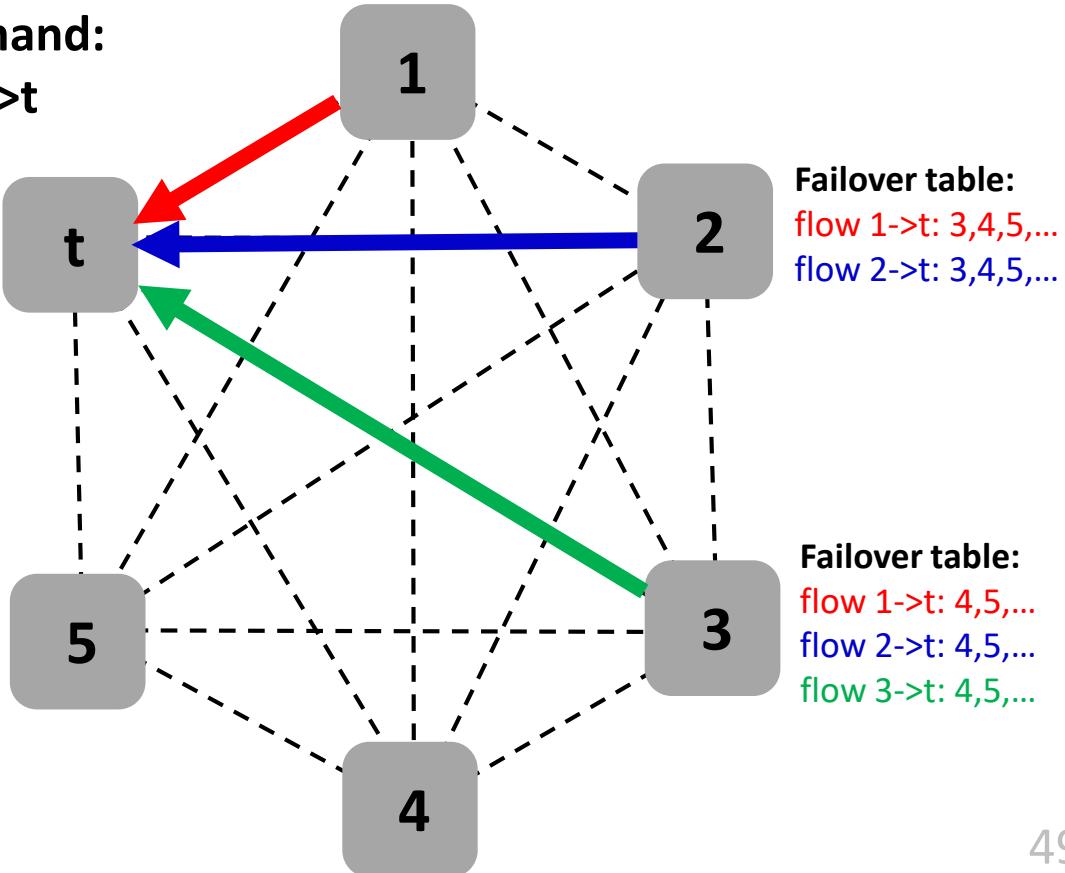
Preinstalled failover rules
for **red** flow

Traffic demand:
 $\{1,2,3\} \rightarrow t$



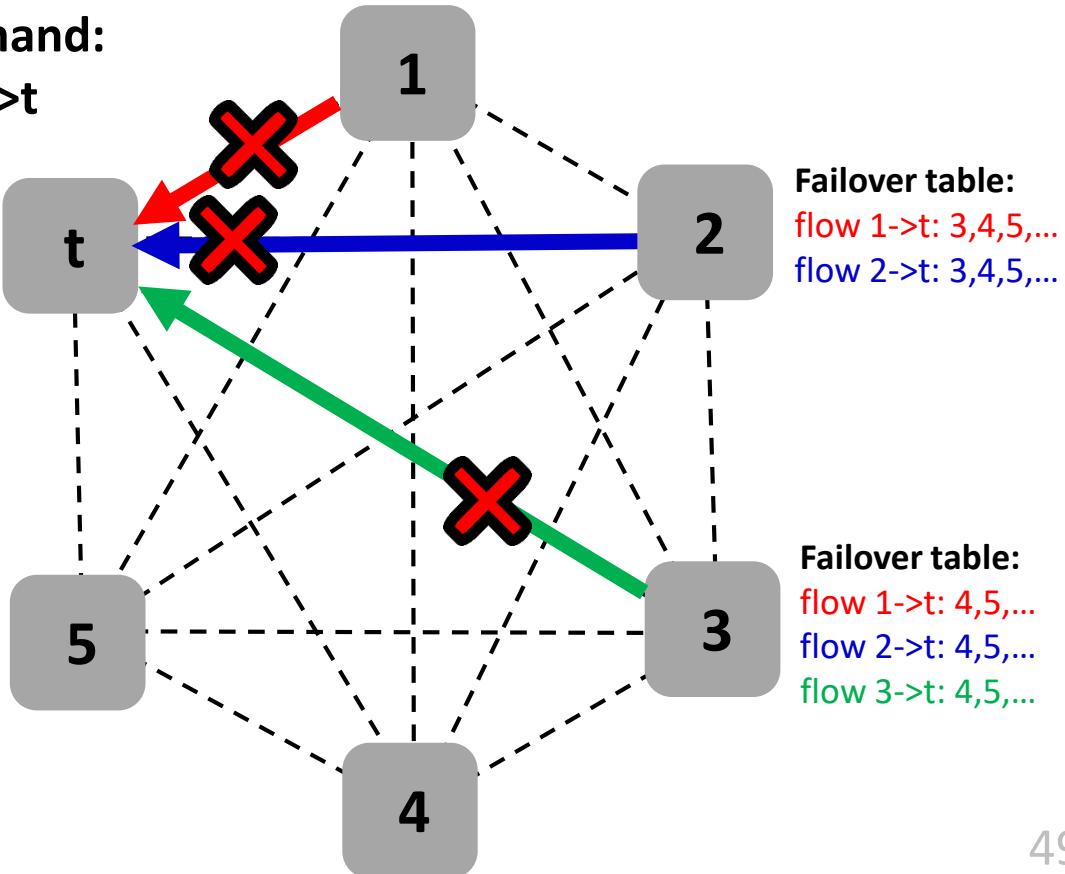
Congestion-Aware FRR

Traffic demand:
 $\{1,2,3\} \rightarrow t$



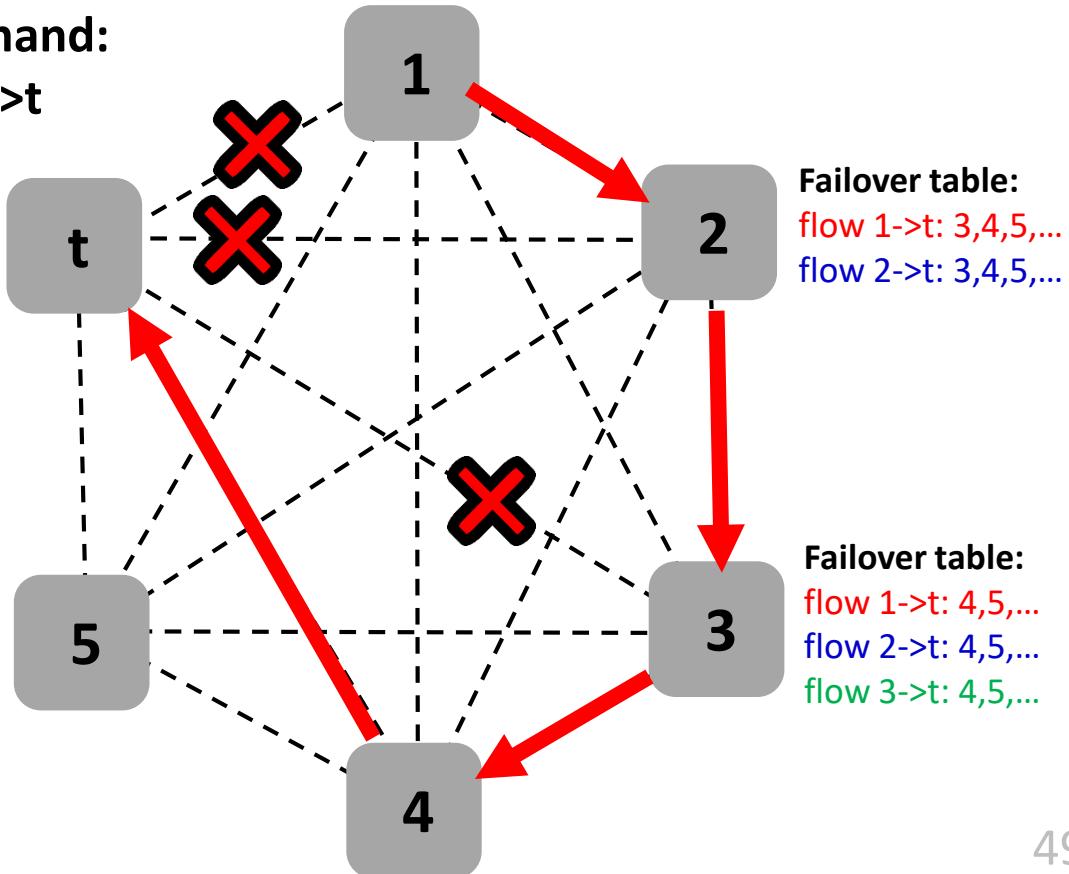
Congestion-Aware FRR

Traffic demand:
 $\{1,2,3\} \rightarrow t$



Congestion-Aware FRR

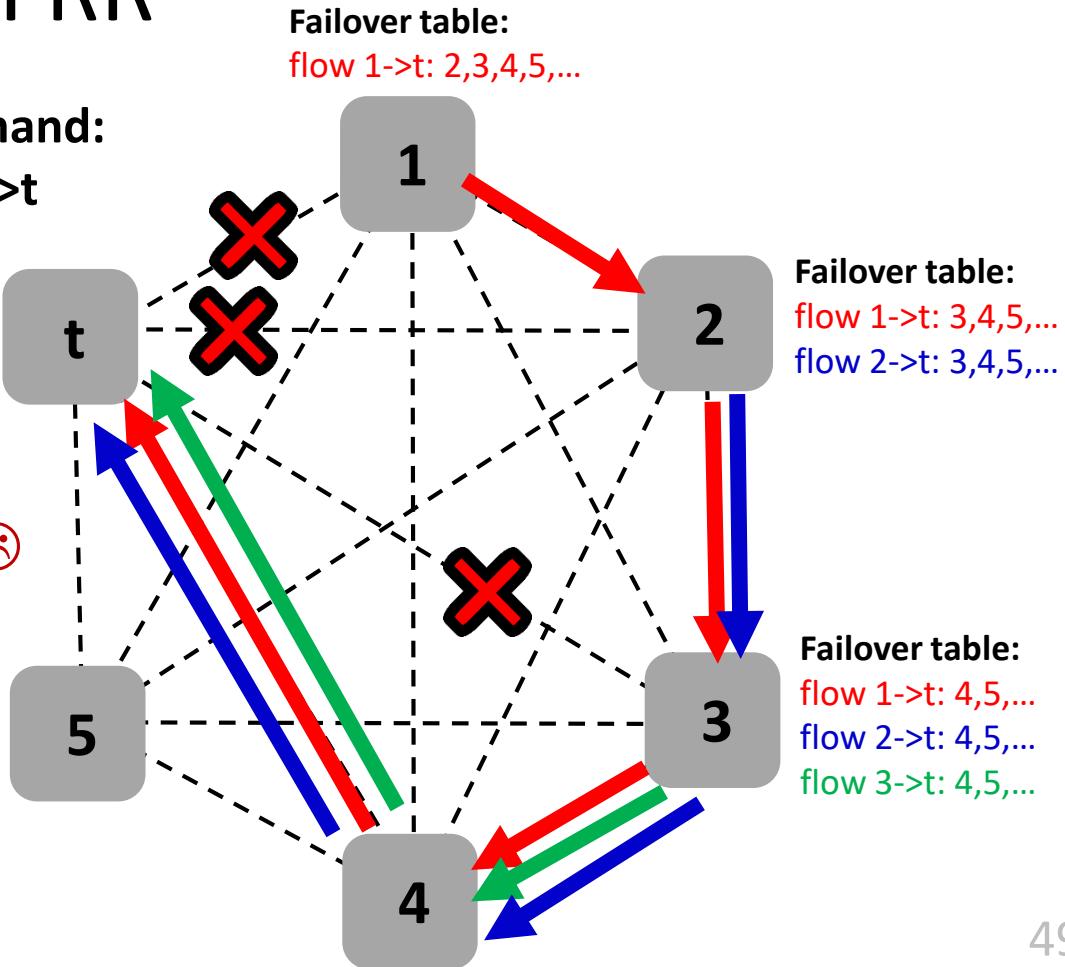
Traffic demand:
 $\{1,2,3\} \rightarrow t$



Congestion-Aware FRR

Traffic demand:
 $\{1,2,3\} \rightarrow t$

Max load is 3 😞

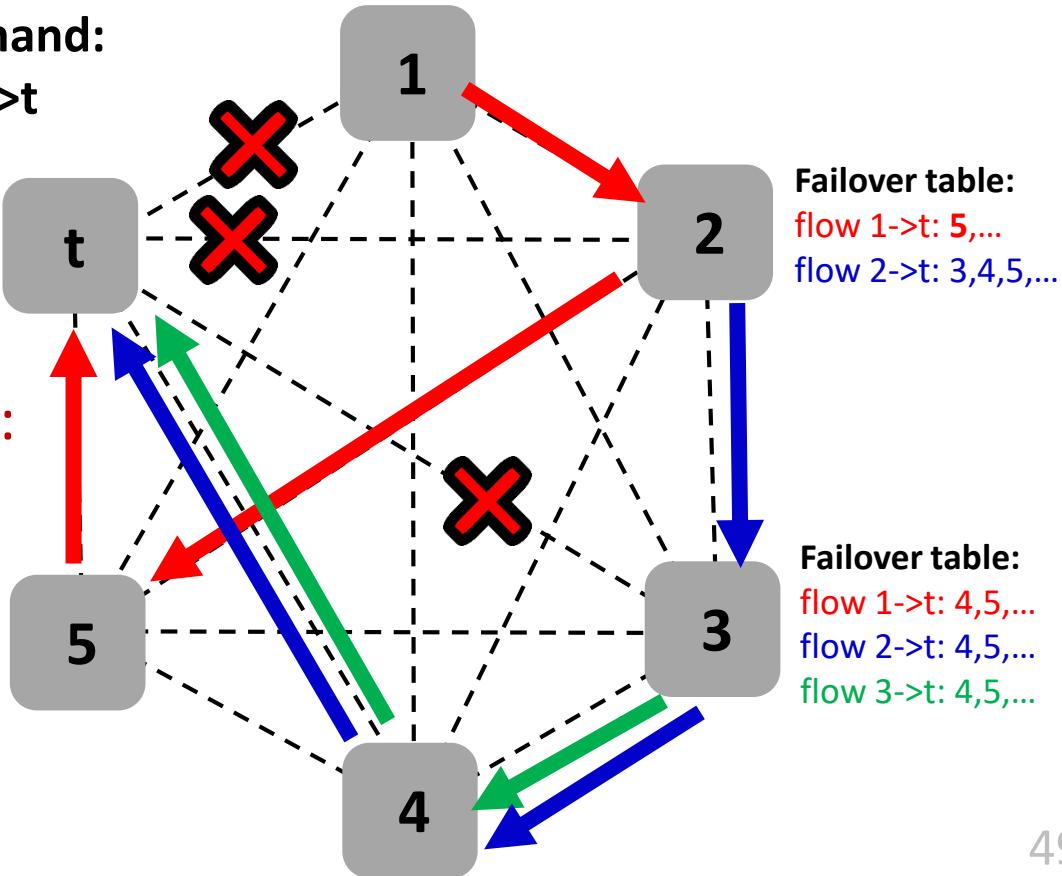


Congestion-Aware FRR

A better solution: load 2 ☺

Traffic demand:

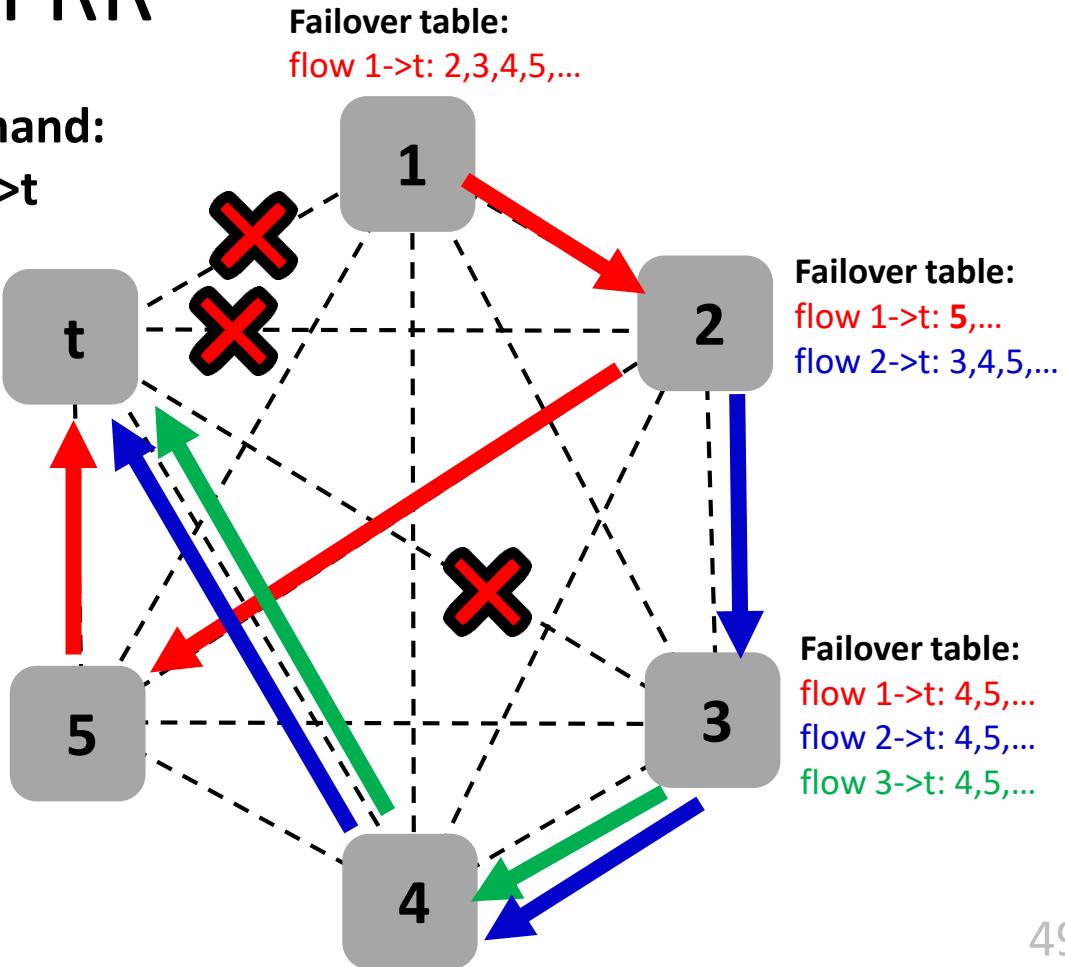
$\{1,2,3\} \rightarrow t$



Congestion-Aware FRR

Observation: we can represent failover tables as a matrix.
To load balance:
prefixes of rows should be different!

Traffic demand:
 $\{1,2,3\} \rightarrow t$



Failover Matrix Representation

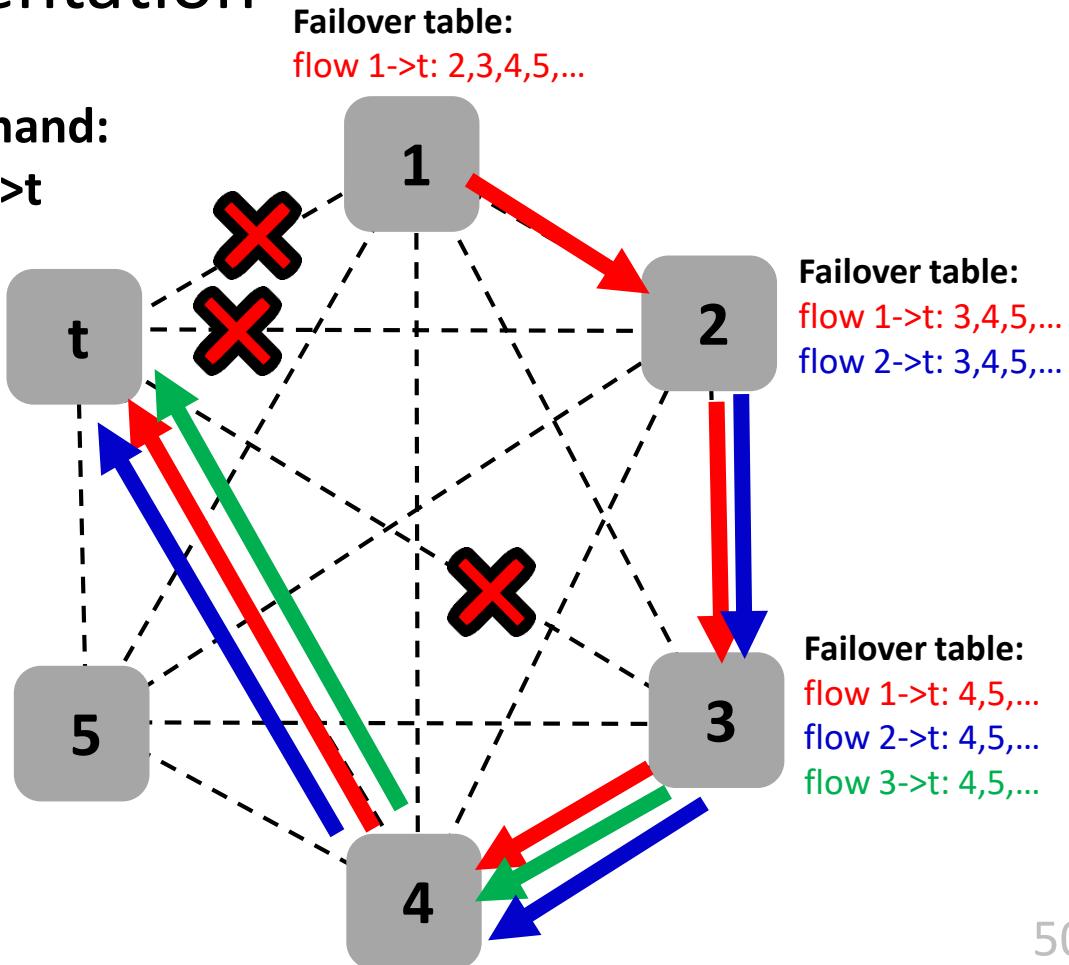
Matrix:

source 1: 2,3,4,5

source 2: 3,4,5,1

source 3: 4,5,1,2

Traffic demand:
 $\{1,2,3\} \rightarrow t$



Failover Matrix Representation

Matrix:

source 1: 2,3,4,5

source 2: 3,4,5,1

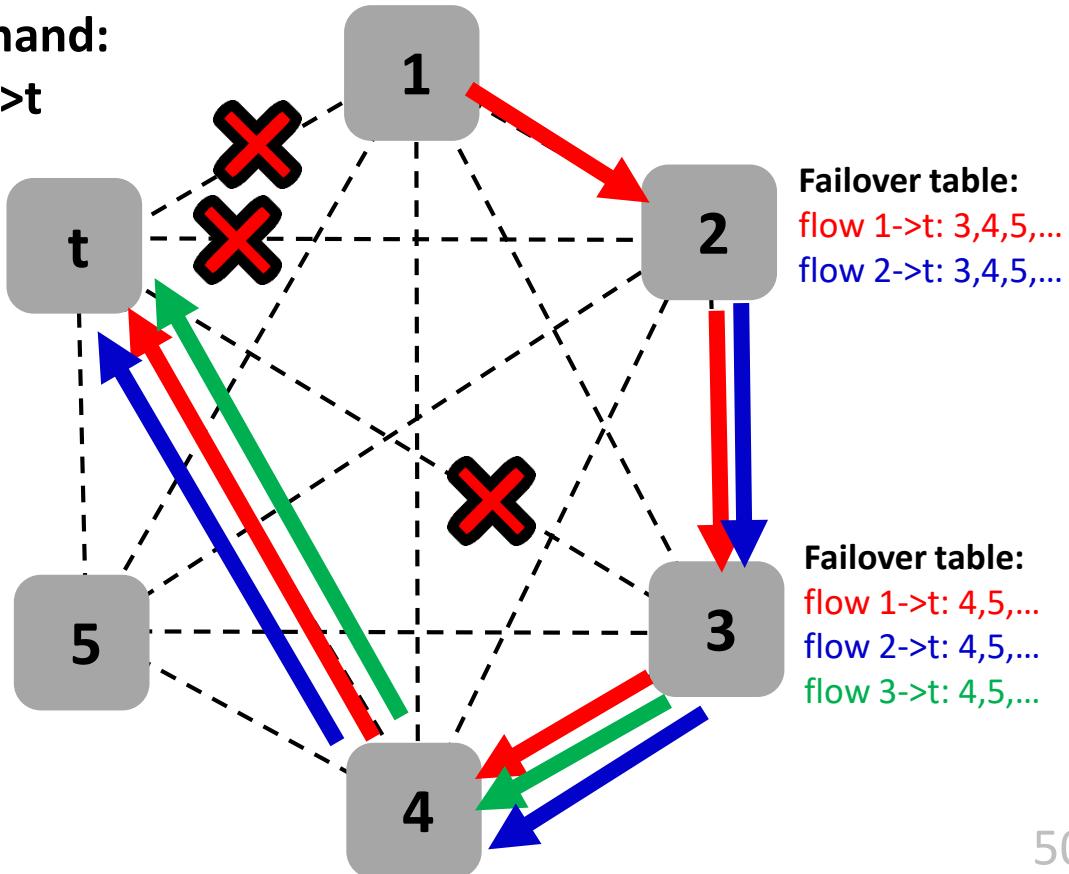
source 3: 4,5,1,2

Problem: failing link $(3,t)$ will affect all three rerouted flows...

In general: easy to create high load on node 4, as failures can be „reused”.

Traffic demand:
 $\{1,2,3\} \rightarrow t$

Failover table:
flow 1 $\rightarrow t$: 2,3,4,5,...



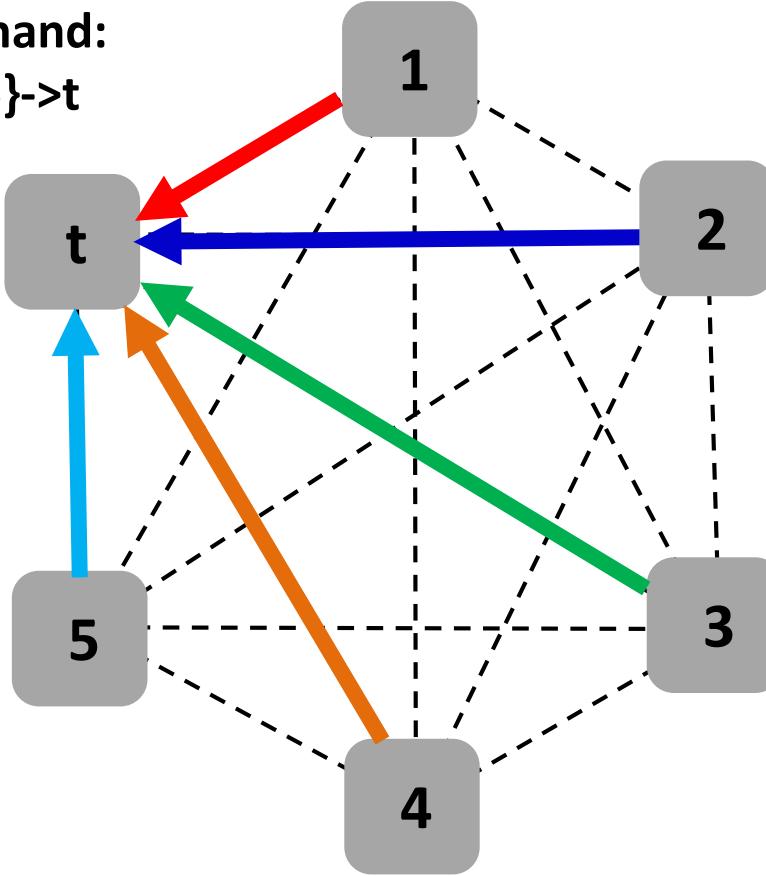
What Are Good Failover Matrices?

- The matrices should be **Latin squares**: each node appears exactly once on each row and each column. No repetitions implies **loop-freedom**.
- Latin squares property gives high **resilience**, but is not sufficient for minimizing **load**.

Challenging Example: Incast

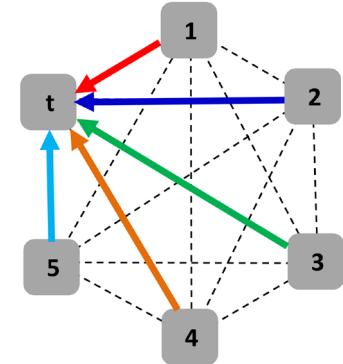
Traffic demand:
 $\{1,2,3,4,5\} \rightarrow t$

In the following, consider
all-to-one demand pattern.



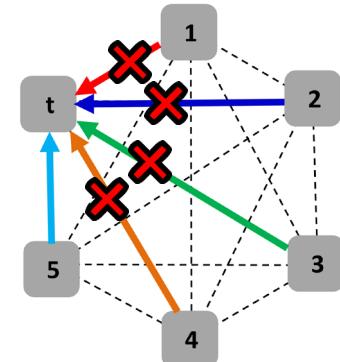
A Bad Matrix for Load

Src 1:	2	3	4	5
Src 2:	3	4	5	1
Src 3:	4	5	1	2
Src 4:	5	1	2	3
Src 5:	1	2	3	4



A Bad Matrix for Load

Src 1:	2	3	4	5
Src 2:	3	4	5	1
Src 3:	4	5	1	2
Src 4:	5	1	2	3
Src 5:	1	2	3	4

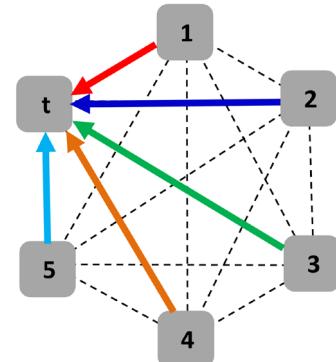


Failing $(1,t)$, $(2,t)$, $(3,t)$, $(4,t)$, gives **load 4** on node 5 / link $(5,t)$.

If the adversary fails the l first links to destination d (that is, $\{(v_i, t), i = 1, \dots, l\}$), then l sources will route through (v_{i+1}, t) . **Load l for l failures.** Can we do better?

Good Failover Matrices?

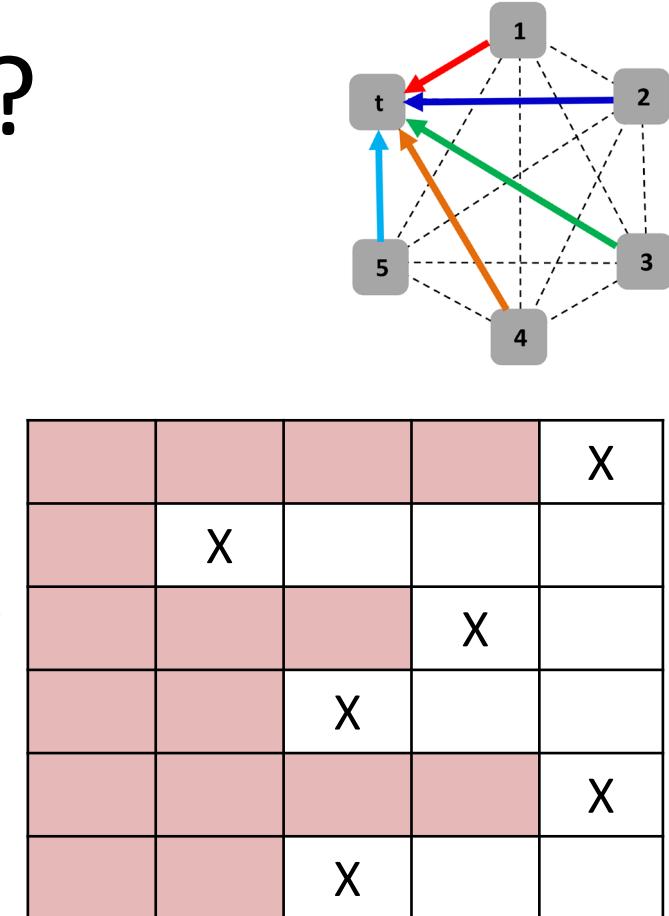
- To bring the flow from a source i to a node X , need to fail *all links* in corresponding *row*
 - Worst case: all *to destination*
- The same for each other flow/row which should reach X



				X
		X		
i				X
			X	
				X
			X	
				X

Good Failover Matrices?

- To bring the flow from a source i to a node X , need to fail *all links* in corresponding *row*
 - Worst case: all *to destination*
- The same for each other flow/row which should reach X
- Adversary will try to *reuse link* failures: **good matrices** have *prefixes with little overlap* (resp. large number of unique nodes)



Connection to Block Designs

- A closely related problem: generating **block designs**
 - and its geometric counterpart, generating **projective planes** of high order
- Using *symmetric balanced incomplete block designs (BIBDs)*
- Gives a latin failover matrix M with intersection properties representing a failover scheme that is ***optimal up to a constant factor***
- Also used in the context **disconnected cooperation**, e.g.:
 - G. Malewicz, A. Russell, and A. A. Shvartsman. Distributed Scheduling for Disconnected Cooperation. *Distributed Computing*, 18(6), 2005.

Overview of Results

Good news: Theory of local algorithms without communication: symmetric block design theory.

Bad news (counting argument): High load unavoidable even in well-connected residual networks: a price of locality.
Given L failures, load at least \sqrt{L} , although network still highly connected ($n-L$ connected). E.g., $L=n/2$, load could be 2 still, but due to locality at least \sqrt{n} .

Randomized Failover

- Recall: deterministic lower bound of \sqrt{L} for L failures, although load could be $O(1)$ for $L < L/2$. A large *price of locality*.
- So what about *randomized* approaches?



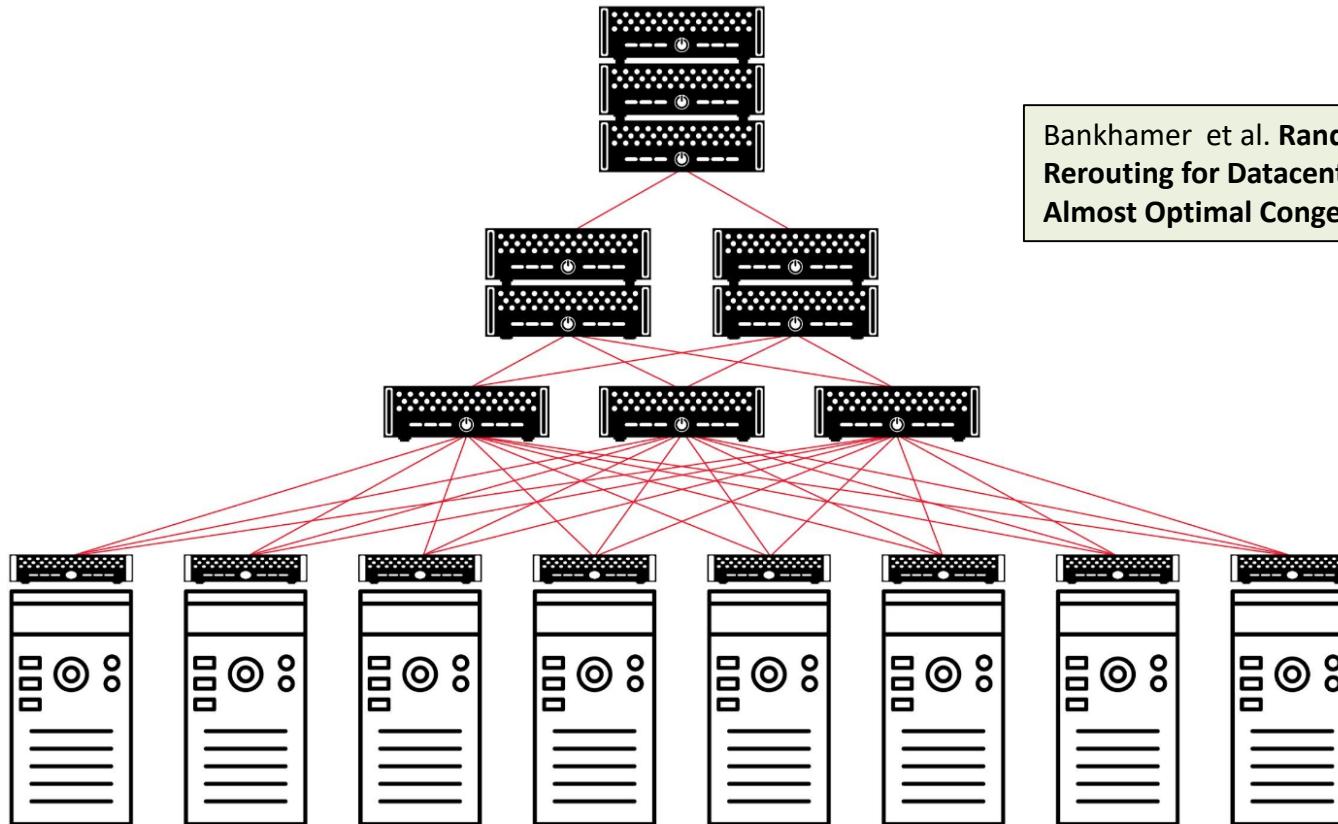
The Power of Randomization

	<i>3-Permutations</i>	<i>Intervals</i>	<i>Shared-Permutations</i>
Rule Set Resilience	Destination + Hop $\Theta(n)$	Destination $\Theta(n/\log n)$	Destination + Hop $\Theta(n)$
Congestion	$\mathcal{O}(\log^2 n \cdot \log \log n)$	$\mathcal{O}(\log n \cdot \log \log n)$	$\mathcal{O}(\sqrt{\log n})$

- While deterministic algorithms can at best achieve a **polynomial** load, randomized algorithms can achieve a **polylogarithmic load**.
- Even when just matching the destination.
 - Losing a $\log n$ factor in resilience.
 - Matching also the hop count can overcome this.

Bankhamer et al. **Local Fast Rerouting with Low Congestion: A Randomized Approach**.
27th IEEE International Conference on Network Protocols (ICNP), 2019.

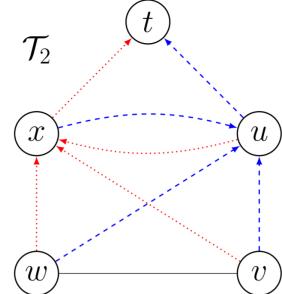
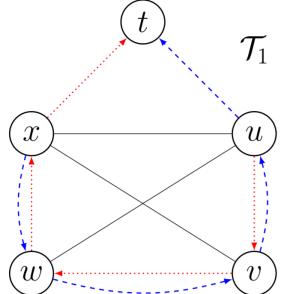
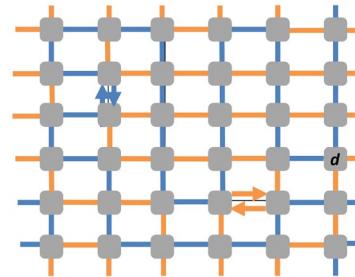
Benefits in Datacenter Networks



Bankhamer et al. Randomized Local Fast
Rerouting for Datacenter Networks with
Almost Optimal Congestion. DISC, 2021.

What About Path Length and Stretch?

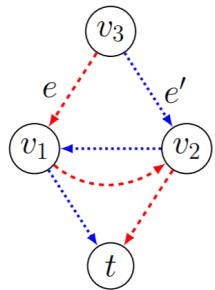
- So far: ignored the length of the failover routes
 - Hamilton cycles are particularly bad
 - The heights of general arborescences may be lower



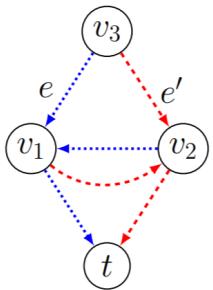
- Idea (so far heuristic):
 - Postprocess the arborescences to lower their heights
 - Two different t -rooted arc-disjoint spanning arborescence decompositions, T_1 and T_2
 - The mean path length of T_1 is higher than that of T_2

Swapping Operations Which Maintain Decomposition

1

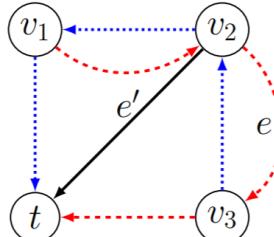


Before swapping

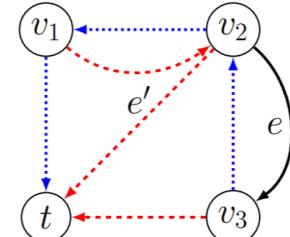


After swapping

2



Before swapping



After swapping

Roadmap

- A Brief History of Resilient Networking
- Algorithms for Local Fast Re-Routing (FRR)
- Accounting for Congestion
- **Accounting for Network Policy**



Roadmap

- A Brief History of Resilient Networking
- Algorithms for Local Fast Re-Routing (FRR)
- Accounting for Congestion
- **Accounting for Network Policy**

An example with
header rewriting.

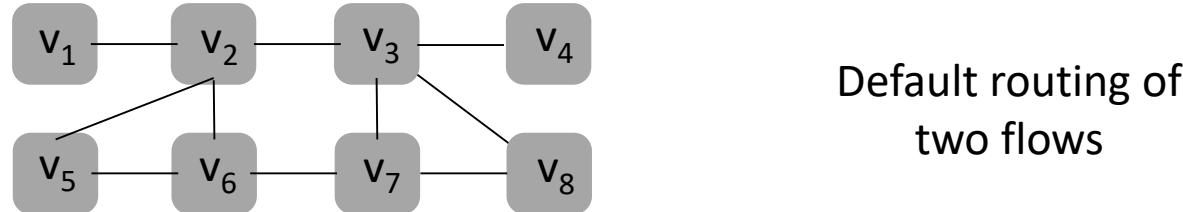


Case Study: MPLS Networks

- Widely deployed networks by Internet Service Providers (**ISPs**)
- Often used for **traffic engineering**
 - Avoid congestion by going non-shortest paths
- Allows for ***header re-writing*** upon failures
 - Header based on **stack of labels**

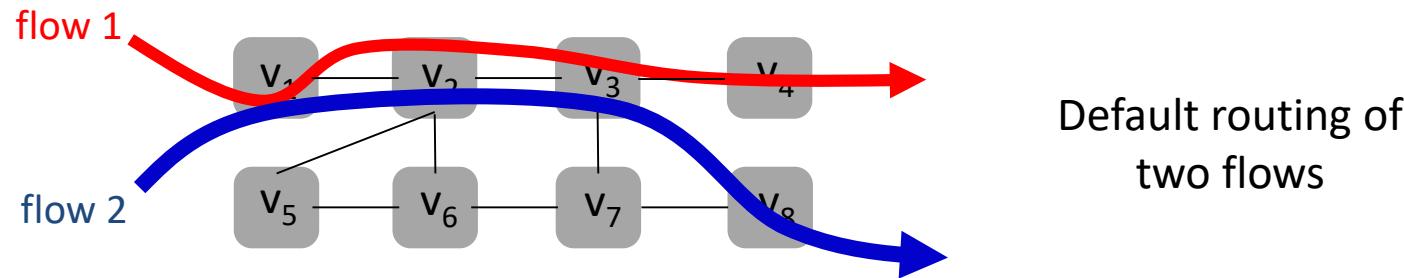
How (MPLS) Networks Work

- Forwarding based on **top label** of label **stack**



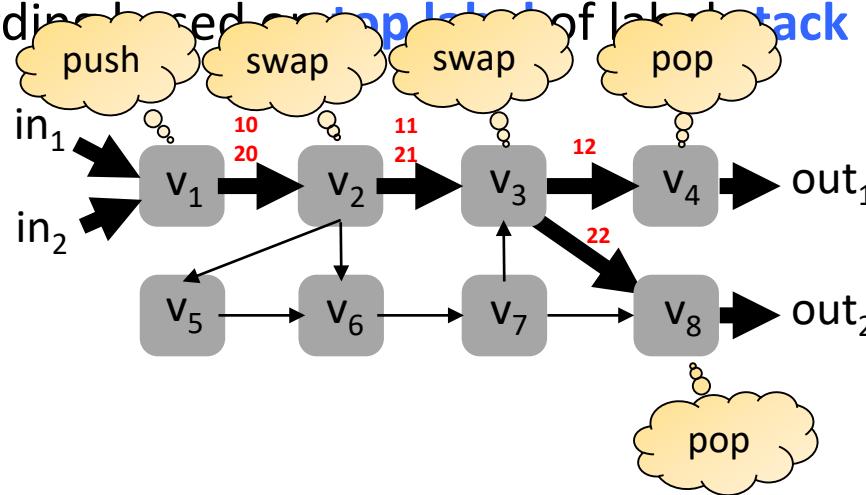
How (MPLS) Networks Work

- Forwarding based on **top label** of label **stack**



How (MPLS) Networks Work

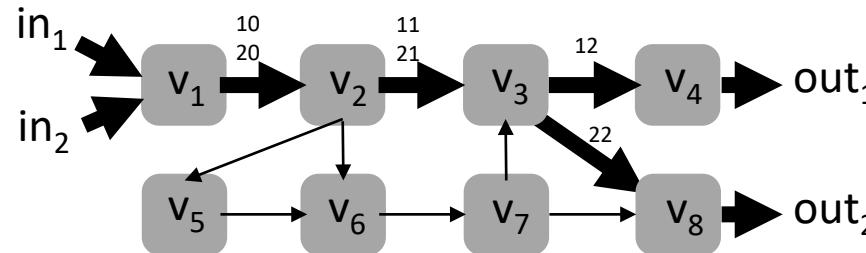
- Forwarding based on top **left** of label stack



Default routing of
two flows

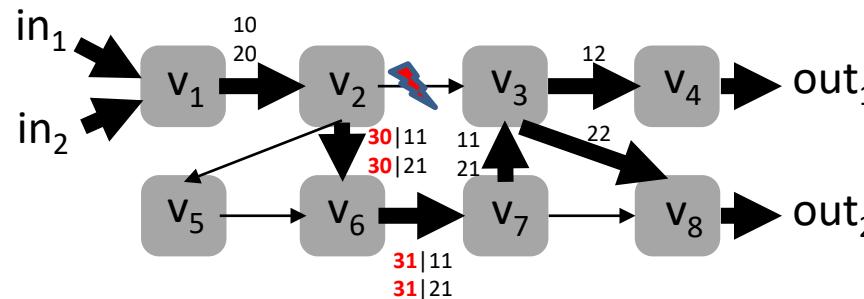
Fast Reroute Around 1 Failure

- Forwarding based on **top label** of label **stack** (in packet header)



Default routing of
two flows

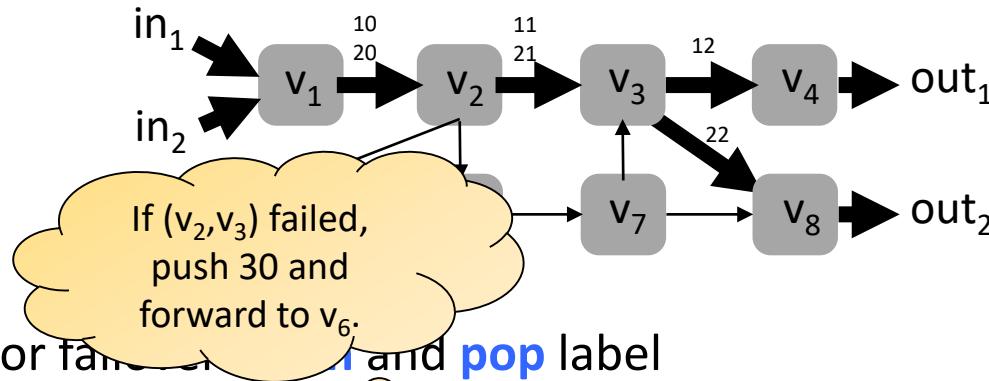
- For failover: **push** and **pop** label



One failure: **push** 30:
route around (v₂, v₃)

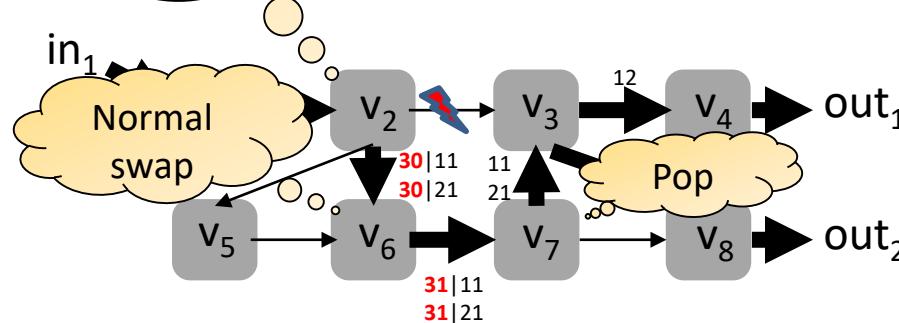
Fast Reroute Around 1 Failure

- Forwarding based on **top label** of label **stack** (in packet header)



Default routing of
two flows

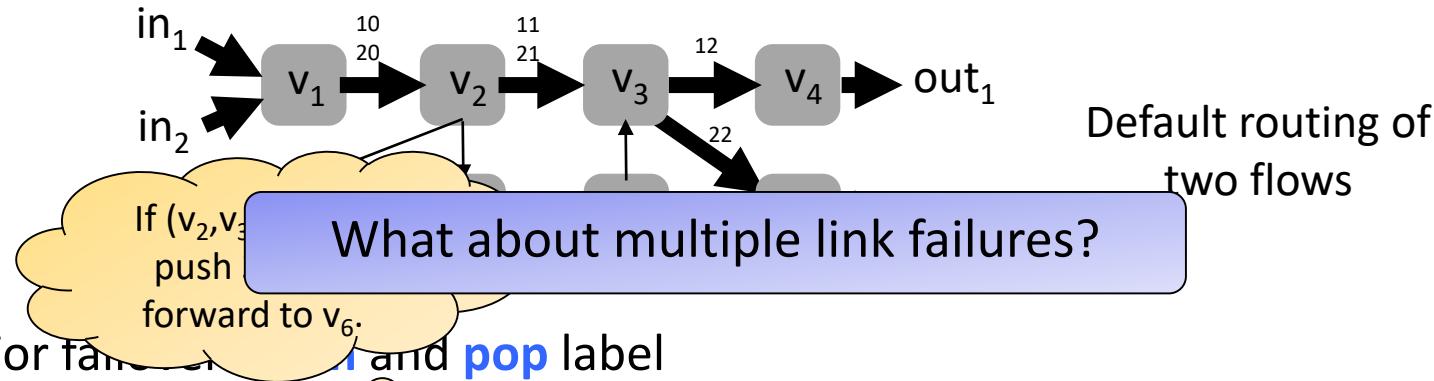
- For failure, **swap** and **pop** label



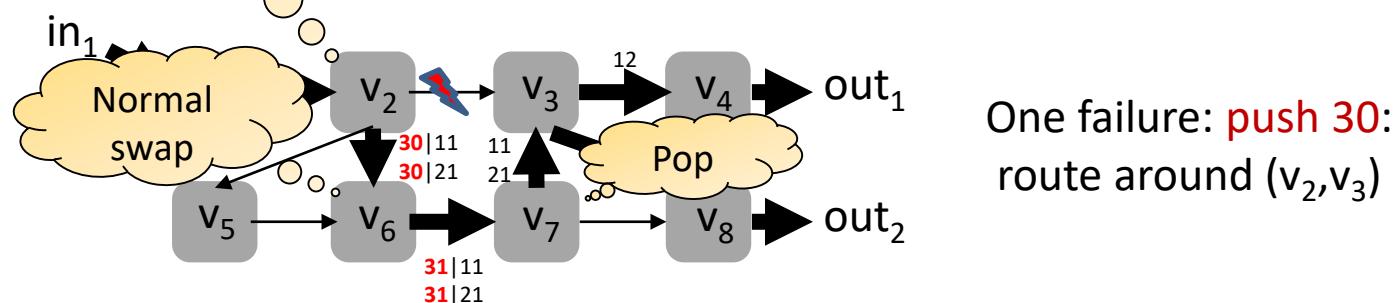
One failure: **push 30:**
route around (v_2, v_3)

Fast Reroute Around 1 Failure

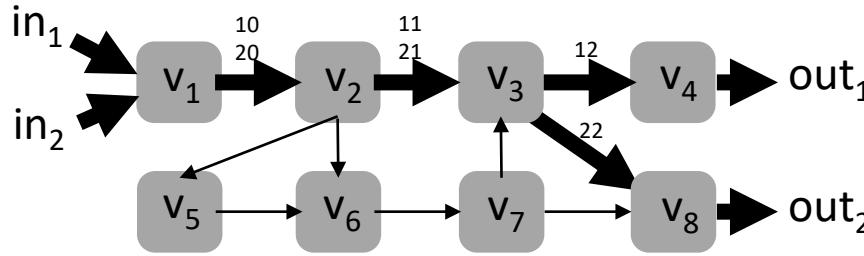
- Forwarding based on **top label** of label **stack** (in packet header)



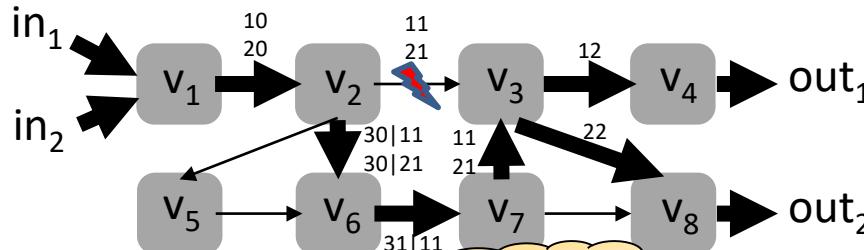
- For failure, **swap** and **pop** label



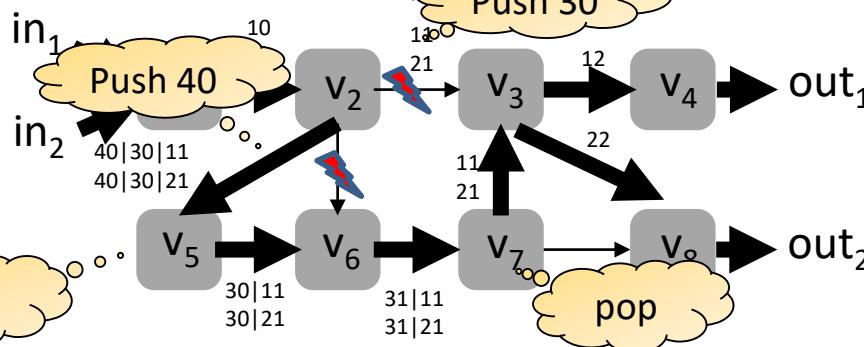
2 Failures: Push *Recursively*



Original Routing



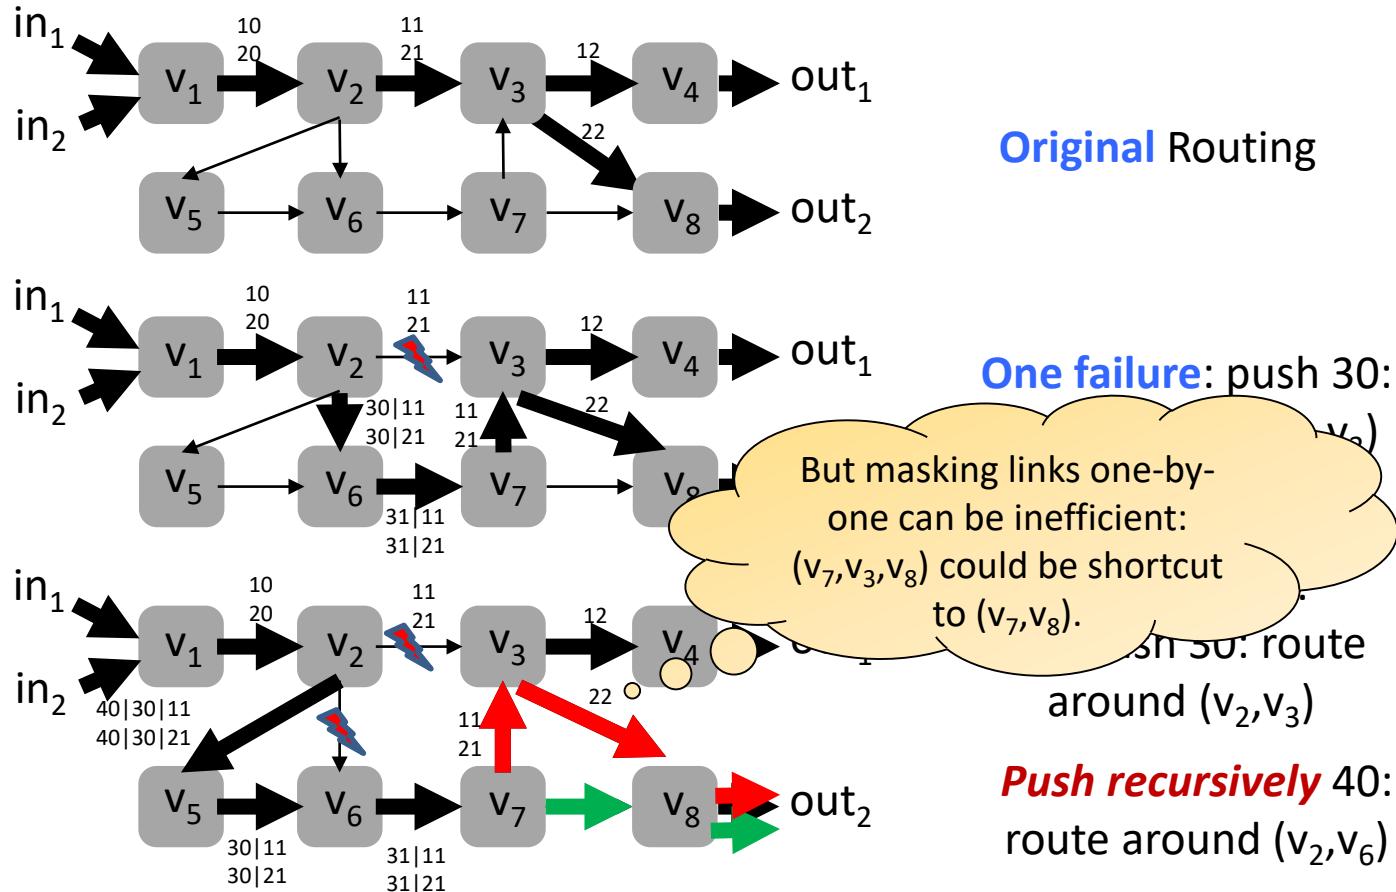
One failure: push 30:
route around (v_2, v_3)



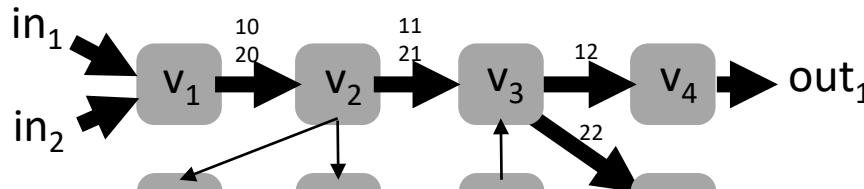
Two failures:
first push 30: route
around (v_2, v_3)

Push recursively 40:
route around (v_2, v_6)

2 Failures: Push *Recursively*

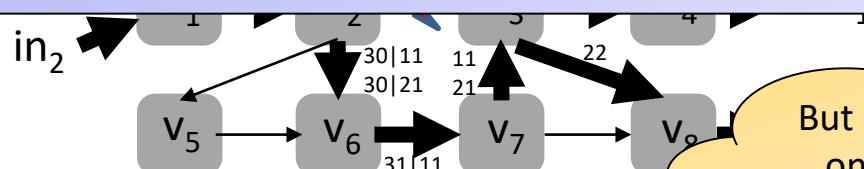


2 Failures: Push *Recursively*



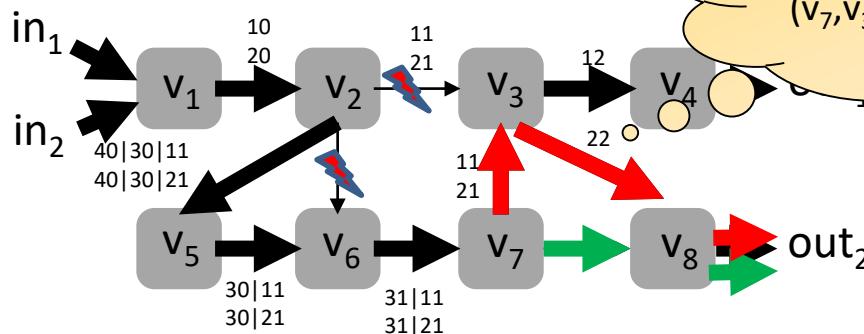
Original Routing

More efficient but also more complex:
Cisco does **not recommend** using this option!



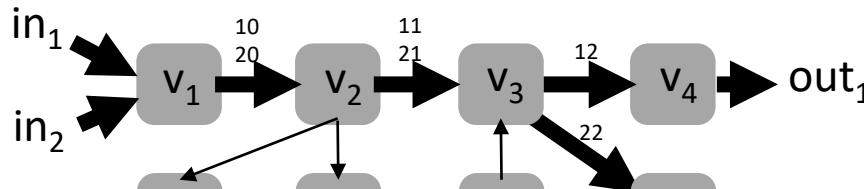
One failure: push 30:

But masking links one-by-one can be inefficient:
 (v_7, v_3, v_8) could be shortcut
to (v_7, v_8) .



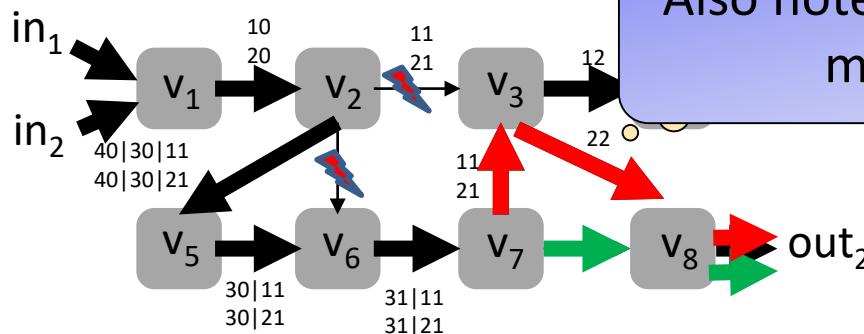
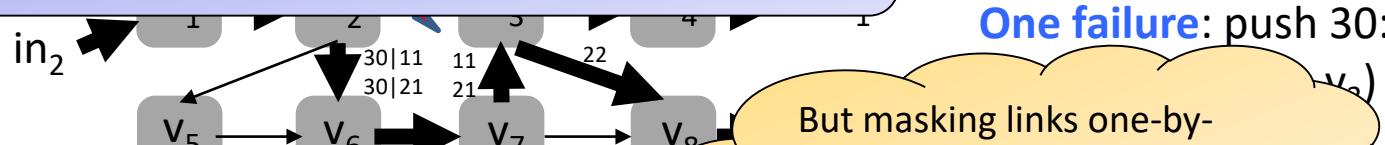
Push recursively 40:
route around (v_2, v_6)

2 Failures: Push *Recursively*



Original Routing

More efficient but also more complex:
Cisco does **not recommend** using this option!



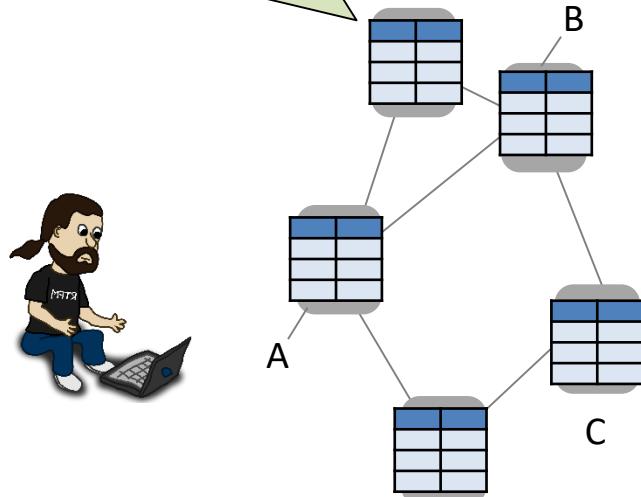
Also note: due to push, **header size** may grow arbitrarily!

around (v_2, v_6)

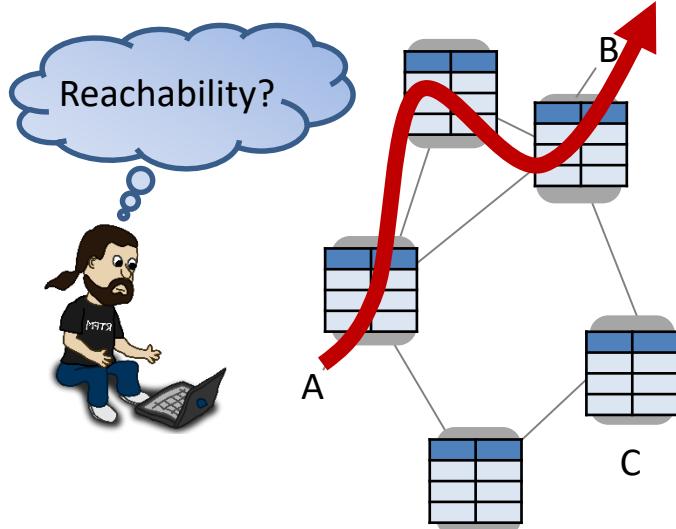
Push recursively 40:
route around (v_2, v_6)

Responsibilities of a Sysadmin

Routers and switches store list of **forwarding rules**, and conditional **failover rules**.



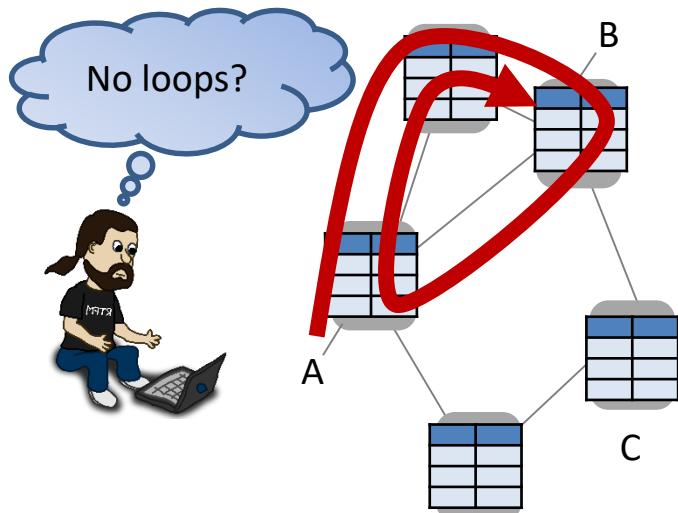
Responsibilities of a Sysadmin



Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?

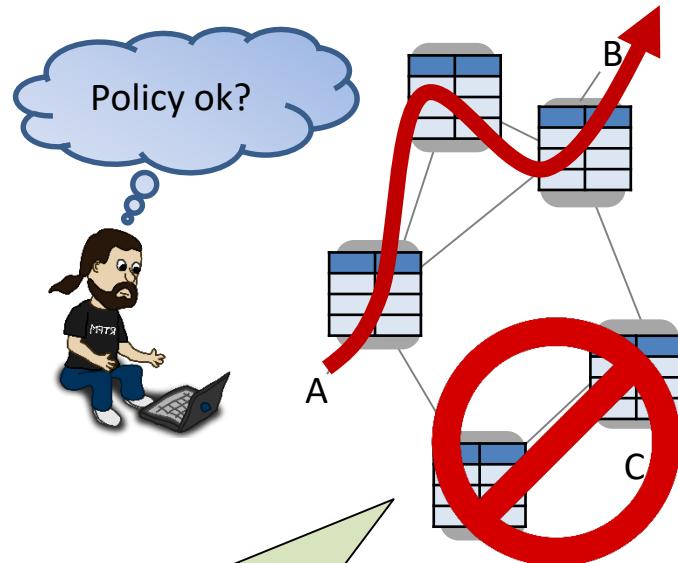
Responsibilities of a Sysadmin



Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?

Responsibilities of a Sysadmin

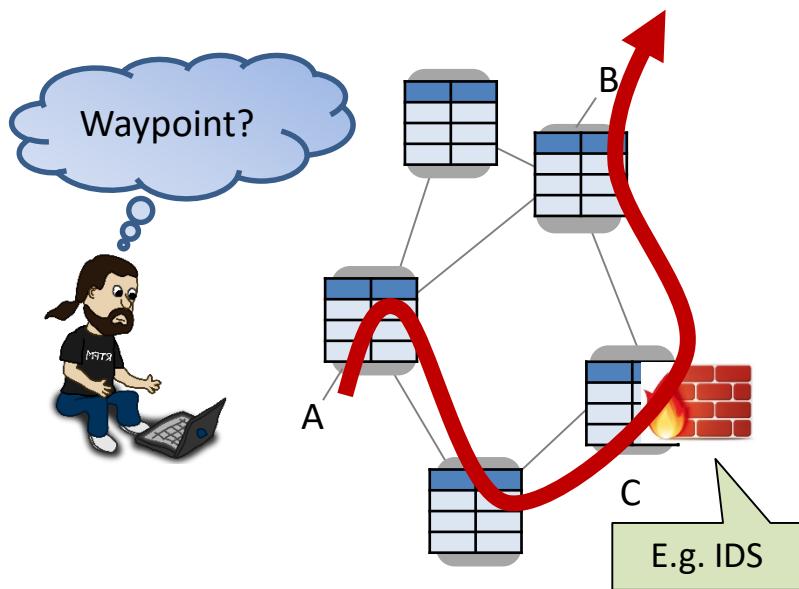


E.g. **NORDUnet**: no traffic via Iceland (expensive!).

Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?
- **Policy:** Is it ensured that traffic from A to B never goes via C?

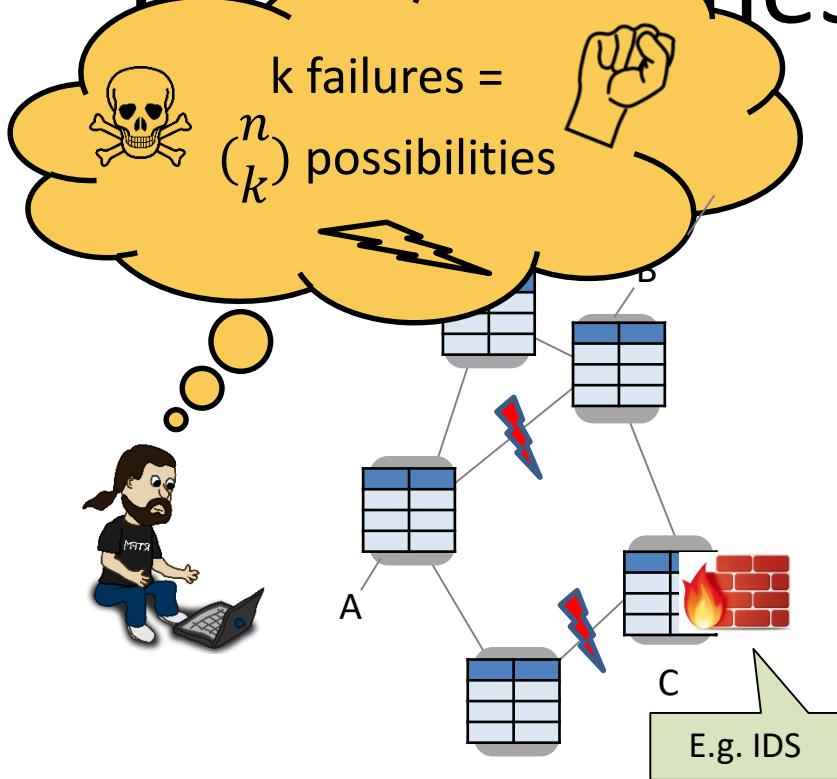
Responsibilities of a Sysadmin



Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?
- **Policy:** Is it ensured that traffic from A to B never goes via C?
- **Waypoint enforcement:** Is it ensured that traffic from A to B is always routed via a node C (e.g., intrusion detection system or a firewall)?

Responsibilities of a Sysadmin

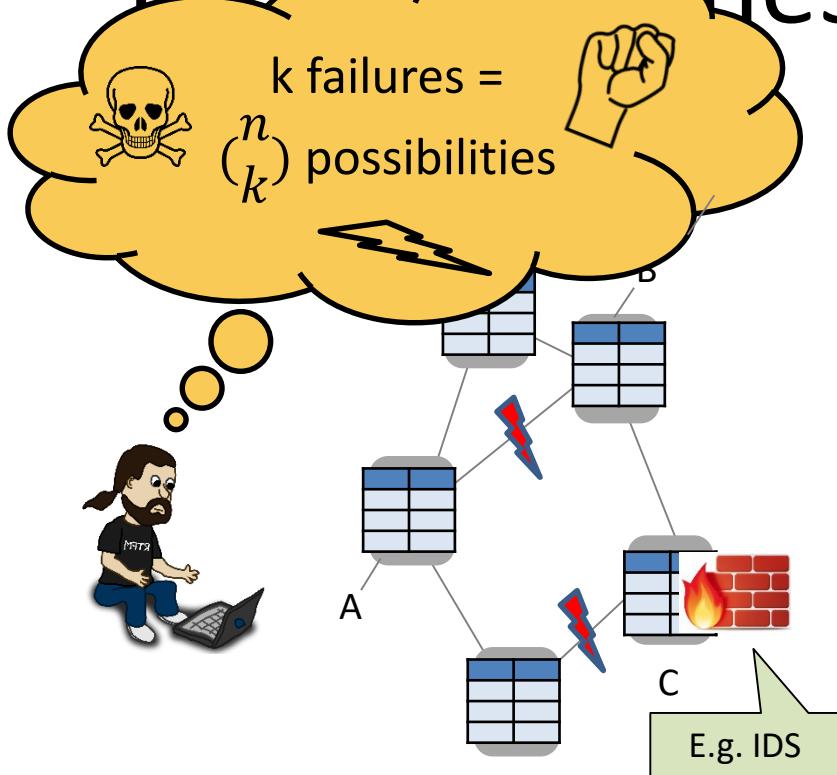


Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?
- **Policy:** Is it ensured that traffic from A to B never goes via C?
- **Waypoint enforcement:** Is it ensured that traffic from A to B is always routed via a node C (e.g., intrusion detection system or a firewall)?

... and everything even under multiple failures?!

Responsibilities of a Sysadmin



Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?
- **Policy:** Is it ensured that traffic from A to B never goes via C?
- **Waypoint enforcement:** Is it ensured that traffic from A to B is always routed via a node C (e.g., intrusion detection system or a firewall)?

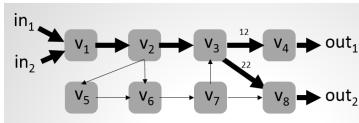
... and everything even under multiple failures?!

Generalization: service chaining!

Approach: Automation and Formal Methods

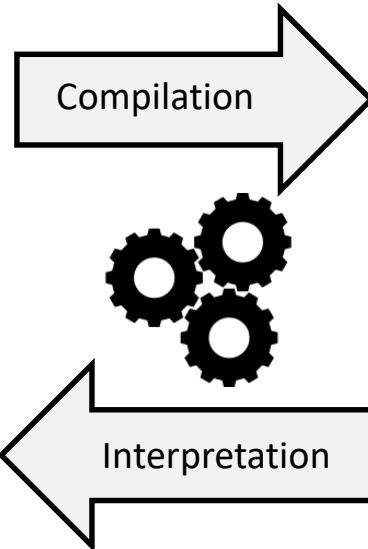


FT	In-I	In-Label	Out-I	op
τ_{v_1}	in_1	\perp	(v_1, v_2)	$push(10)$
τ_{v_2}	(v_1, v_2)	10	(v_2, v_3)	$push(11)$
τ_{v_3}	(v_1, v_2)	20	(v_2, v_3)	$swap(21)$
τ_{v_4}	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
τ_{v_5}	(v_2, v_3)	21	(v_3, v_4)	$swap(22)$
τ_{v_6}	(v_3, v_4)	11	(v_3, v_5)	$swap(12)$
τ_{v_7}	(v_3, v_4)	21	(v_3, v_5)	$swap(22)$
τ_{v_8}	(v_2, v_5)	40	(v_5, v_6)	pop
τ_{v_9}	(v_2, v_6)	30	(v_6, v_7)	$swap(31)$
$\tau_{v_{10}}$	(v_5, v_6)	30	(v_6, v_7)	$swap(31)$
$\tau_{v_{11}}$	(v_5, v_6)	61	(v_6, v_7)	$swap(62)$
$\tau_{v_{12}}$	(v_5, v_6)	71	(v_6, v_7)	$swap(72)$
$\tau_{v_{13}}$	(v_6, v_7)	31	(v_6, v_8)	pop
$\tau_{v_{14}}$	(v_6, v_7)	62	(v_7, v_8)	$swap(11)$
$\tau_{v_{15}}$	(v_6, v_7)	72	(v_7, v_8)	$swap(22)$
$\tau_{v_{16}}$	(v_3, v_8)	22	out_1	pop
$\tau_{v_{17}}$	(v_7, v_8)	22	out_2	pop



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_6)	30	(v_2, v_5)	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
τ'_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$

Router configurations
(Cisco, Juniper, etc.)



$pX \Rightarrow qXX$

$pX \Rightarrow qYX$

$qY \Rightarrow rYY$

$rY \Rightarrow r$

$rX \Rightarrow pX$

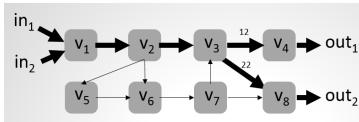
Pushdown Automaton and
Prefix Rewriting Systems

Approach: Automation Methods

Use cases: Sysadmin **issues queries** to test certain properties, or do it on a **regular basis** automatically!



FT	In-I	In-Label	Out-I	op
τ_{v_1}	in_1	\perp	(v_1, v_2)	$push(10)$
	in_2	\perp	(v_1, v_2)	$push(20)$
τ_{v_2}	(v_1, v_2)	10	(v_2, v_3)	$swap(11)$
	(v_1, v_2)	20	(v_2, v_3)	$swap(21)$
τ_{v_3}	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_2, v_3)	21	(v_3, v_4)	$swap(22)$
τ_{v_4}	(v_3, v_4)	11	(v_4, v_5)	$swap(12)$
	(v_3, v_4)	21	(v_3, v_5)	$swap(22)$
τ_{v_5}	(v_4, v_5)	12	out_1	pop
	(v_2, v_5)	40	(v_5, v_6)	pop
τ_{v_6}	(v_2, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	61	(v_6, v_7)	$swap(62)$
τ_{v_7}	(v_5, v_7)	71	(v_7, v_8)	$swap(72)$
	(v_6, v_7)	31	(v_7, v_8)	pop
	(v_6, v_7)	62	(v_7, v_8)	$swap(11)$
	(v_6, v_7)	72	(v_7, v_8)	$swap(22)$
τ_{v_8}	(v_3, v_8)	22	out_2	pop
	(v_7, v_8)	22	out_2	pop



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_6)	30	(v_2, v_5)	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
τ'_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$

Router configurations
(Cisco, Juniper, etc.)

Compilation



$pX \Rightarrow qXX$

$pX \Rightarrow qYX$

$qY \Rightarrow rYY$

$rY \Rightarrow r$

$rX \Rightarrow pX$

Interpretation

Pushdown Automaton and
Prefix Rewriting Systems

AalWiNes Tool

AalWiNes Indian Ocean
MPLS Reachability Analysis & Visualization Tool

Model Aarnet

Query <ip> [#Sydney1] .* [Brisbane2#] <ip> 0

Examples:
<ip> [#Sydney1] .* [Brisbane2#] <ip> 0
<smpls ip> [#Sydney1] .* [Brisbane2#] <mpls* smpls ip> 1

Initial header: ip
Route restriction: [#Sydney1] .* [Brisbane2#]
Final header: ip
Max link failures: 0

Options

Run Validation

Result Satisfied

Query: <ip> [#Sydney1] .* [Brisbane2#] <ip> 0

```
<ip6> : [#Sydney1]
push($43)
<s43,ip6> : [Sydney1#Brisbane1]
swap($44)
<s44,ip6> : [Brisbane1#Brisbane2]
pop()
<ip6> : [Brisbane2#]
```

About AalWiNes

A tool for MPLS reachability analysis and visualization from:

- Aalborg University
Department of Computer Science
- University of Vienna
Communication Technologies Group

Have a look at the [Tool Website](#) & [Tool and query language documentation](#)

Query:
regular
expression

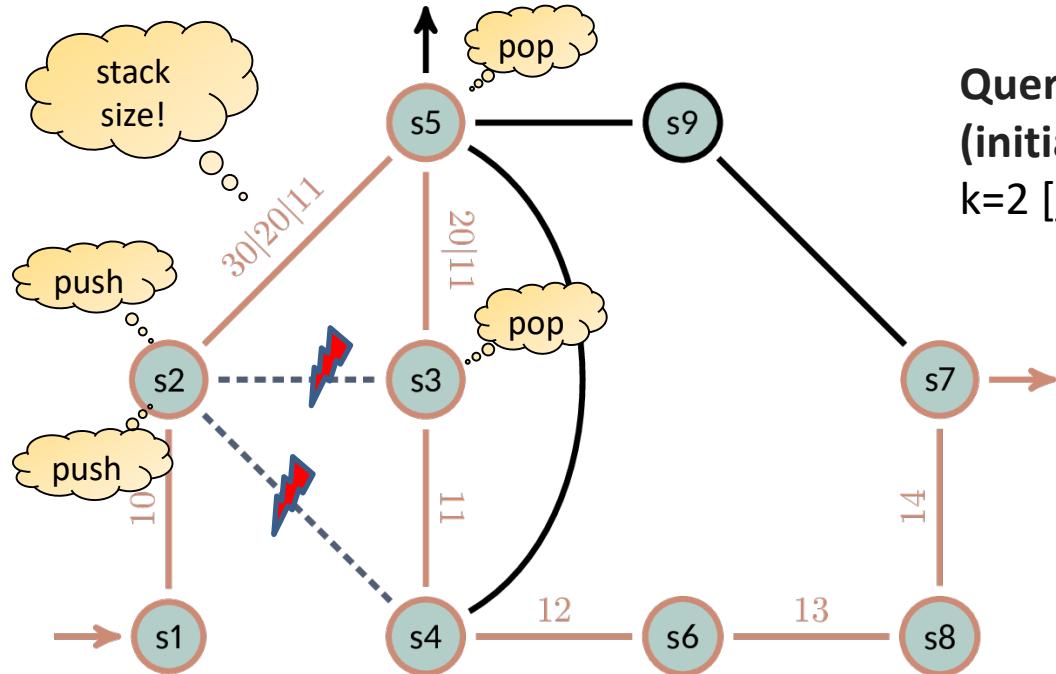
Witness

Dozens of
networks

Online demo: <https://demo.aalwines.cs.aau.dk/>
Source code: <https://github.com/DEIS-Tools/AalWiNes>

Example

Can traffic starting with [] go through s5, under up to k=2 failures?



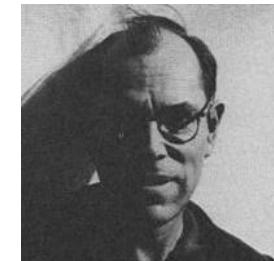
Query: 3 regular expressions
(initial and final header, route)
k=2 [] s1 >> s5 >> s7 []

2 failures

YES
(Polynomial time!)

Why AalWiNes is Fast (Polytime): Automata Theory

- For fast verification, we can use the result by **Büchi**: the set of all reachable configurations of a pushdown automaton a is **regular set**
- We hence simply use **Nondeterministic Finite Automata (NFAs)** when reasoning about the pushdown automata
- The resulting **regular operations** are all **polynomial time**



Julius Richard Büchi
1924-1984
Swiss logician

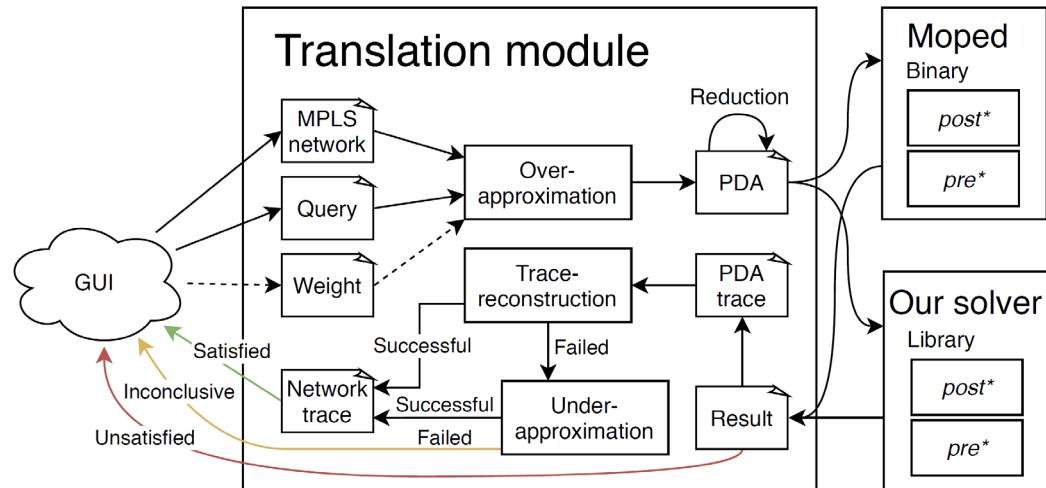
AalWiNes

Part 1: Parses query
and constructs Push-
Down System (PDS)

- In Python 3

Part 2: Reachability
analysis of
constructed PDS

- Using **Moped** tool



Resp. our new weighted extension and
much faster implementation in C++.

Network Model

- Network: a 7-tuple

$$N = (V, E, I_v^{in}, I_v^{out}, \lambda_v, L, \delta_v^F)$$

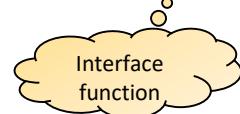
The diagram illustrates the components of a network model tuple. It consists of five yellow thought bubbles connected by dotted lines to the corresponding elements in the tuple definition:

- A bubble labeled "Links" connects to the element I_v^{in} .
- A bubble labeled "Outgoing interfaces" connects to the element I_v^{out} .
- A bubble labeled "Nodes" connects to the element V .
- A bubble labeled "Incoming interfaces" connects to the element E .
- A bubble labeled "Set of labels in packet header" connects to the element δ_v^F .

Network Model

- Network: a 7-tuple

$$N = (V, E, I_v^{in}, I_v^{out}, \lambda_v, L, \delta_v^F)$$



Interface function: maps outgoing interface to next hop node and incoming interface to previous hop node

$$\lambda_v : I_v^{in} \cup I_v^{out} \rightarrow V$$

That is: $(\lambda_v(in), v) \in E$ and $(v, \lambda_v(out)) \in E$

Network Model

- Network: a 7-tuple

$$N = (V, E, I_v^{in}, I_v^{out}, \lambda_v, L, \delta_v^F)$$



Routing function: for each set of **failed links** $F \subseteq E$, the routing function

$$\delta_v^F : I_v^{in} \times L^* \rightarrow 2^{(I_v^{out} \times L^*)}$$

defines, for all **incoming interfaces** and packet **headers**, **outgoing interfaces** together with **modified headers**.

Routing

Packet routing sequence can be represented using **sequence of tuples**:

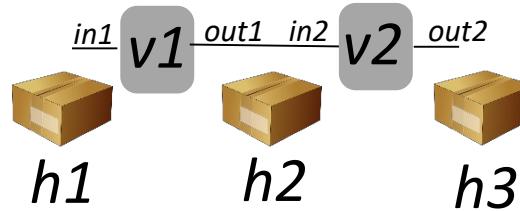


- Example: **routing** (in)finite sequence of tuples

$$(v_1, in_1, h_1, out_1, h_2, F_1),$$

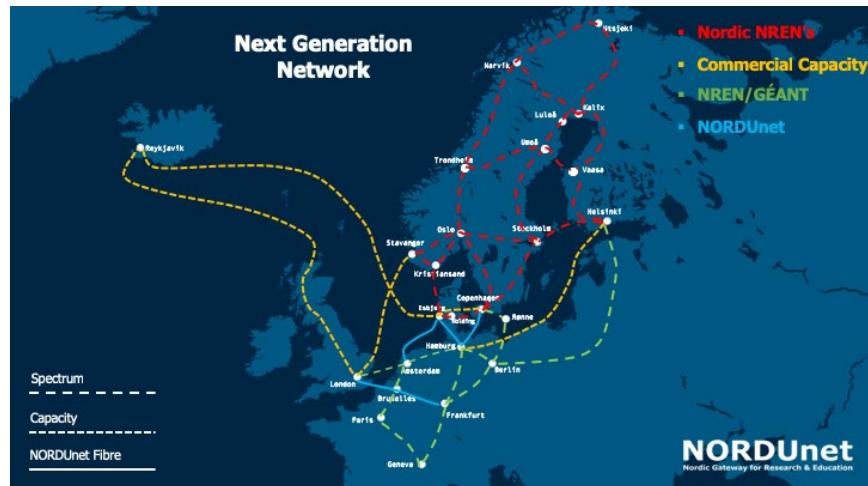
$$(v_2, in_2, h_2, out_2, h_3, F_2),$$

...



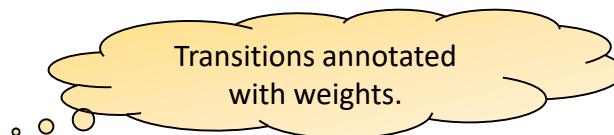
Case Study: NORDUnet

- Regional service provider
 - **24 MPLS routers** geographically distributed across several countries
 - Running **Juniper** operating system
 - More than 30,000 labels
 - Ca. **1 million** forwarding rules in our model
 - For most queries of operators:
answer *within seconds*



Generalizes to Quantitative Properties

- AalWiNes can also be used to test **quantitative properties**
- If query is satisfied, find trace that minimizes:
 - **Hops**
 - Latency (based on a latency value per link)
 - Tunnels



- Approach: **weighted** pushdown automata
 - Fast *poly-time algorithms* exist also for weighted pushdown automata (area of dataflow analysis)
 - Indeed, experiments show: *acceptable overhead* of weighted (quantitative) analysis

Conclusion

- Fast rerouting requires *local decision making*
- Different fault-tolerance *metrics*: ideal resilience, perfect resilience
- What can be achieved depends on *what can be matched* locally
- Locally *balancing load* under failures is hard, but randomization helps

What About The Control Plane?

Resilient Capacity-Aware Routing

Stefan Schmid¹, Nicolas Schnepf², and Jiří Srba²

¹ Faculty of Computer Science, University of Vienna

² Department of Computer Science, Aalborg University

Still many open questions
too, see e.g., **TACAS 2021**

Abstract. To ensure a high availability, most modern communication networks provide resilient routing mechanisms that quickly change routes upon failures. However, a fundamental algorithmic question underlying such mechanisms is hardly understood: how to efficiently verify whether a given network reroutes flows along *feasible* paths, without violating capacity constraints, for up to k link failures? We chart the algorithmic complexity landscape of resilient routing under link failures, considering shortest path routing based on link weights (e.g., the widely deployed ECMP protocol). We study two models: a *pessimistic* model where flows interfere in a worst-case manner along equal-cost shortest paths, and an *optimistic* model where flows are routed in a best-case manner and we present a complete picture of the algorithmic complexities for these models. We further propose a strategic search algorithm that checks only the critical failure scenarios while still providing correctness guarantees. Our experimental evaluation on a large benchmark of Internet and datacenter topologies confirms an improved performance of our strategic search algorithm by several orders of magnitude.

What About Segment Routing?

TI-MFA: Keep Calm and Reroute Segments Fast

Klaus-Tycho Foerster¹ Mahmoud Parham¹ Marco Chiesa² Stefan Schmid¹
¹ University of Vienna, Austria ² KTH Royal Institute of Technology, Sweden

Abstract—Segment Routing (SR) promises to provide scalable and fine-grained traffic engineering. However, little is known today on how to implement resilient routing in SR, i.e., routes which tolerate one or even multiple failures. This paper initiates the theoretical study of static fast failover mechanisms which do not depend on reconvergence and hence support a very fast reaction to failures. We introduce formal models and identify fundamental tradeoffs on what can and cannot be achieved in terms of static resilient routing. In particular, we identify an inherent price in terms of performance if routing paths need to be resilient, even in the absence of failures. Our main contribution is a first algorithm which is resilient even to multiple failures and which comes with provable resiliency and performance guarantees. We complement our formal analysis with simulations on real topologies, which show the benefits of our approach over existing algorithms.

I. INTRODUCTION

Segment Routing (SR) [1], [2], [3] emerged to address operational issues of MPLS-based traffic engineering solutions. SR provides a fine-grained flow management and traffic

- 2) The next node (the node on the top of the stack) does not have to be directly adjacent to the current node. In this case, to transport packets to the next node, segment routing relies on *shortest path routing* (a “segment”). If the next element is a link, that link has to be directly adjacent to the current node.

The definition of intermediate points in the label stack allows to increase the path diversity beyond shortest paths. However, intermediate points can also be used for Fast Rerouting (FRR): if a link along the current shortest path from s to t failed, say at some node u , an alternative intermediate destination (or waypoint) w can be defined to reroute around the failure. In this case, node v receives a label stack “ t ” (only the destination node t is on the stack), pushes the intermediate destination w (a waypoint), and hence sends out a packet with label stack “ wt ”. Once the packet reaches node w , this node is popped from the stack and the packet (with stack “ t ”) routed to t . If additional failures occur, additional intermediate destinations can be defined recursively.

See e.g., *GI 2018*
and *OPODIS 2020*

¹ Maximally Resilient Replacement Paths for a Family of Product Graphs

Mahmoud Parham 

³ University of Vienna, Faculty of Computer Science, Vienna, Austria
⁴ mahmoud.parham@univie.ac.at

Klaus-Tycho Foerster 

⁶ University of Vienna, Faculty of Computer Science, Vienna, Austria
⁸ klaus-tycho.foerster@univie.ac.at

Petar Kosić

⁹ University of Vienna, Faculty of Computer Science, Vienna, Austria
¹¹ petar.kosic@univie.ac.at

Stefan Schmid 

¹² University of Vienna, Faculty of Computer Science, Vienna, Austria
¹⁴ stefan_schmid@univie.ac.at

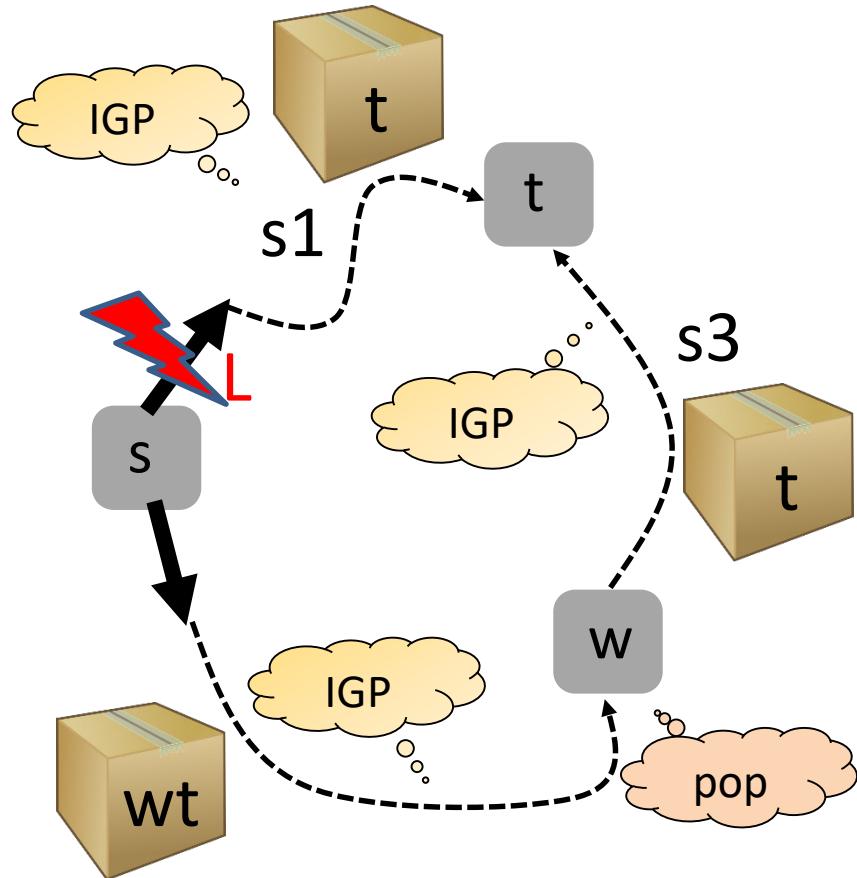
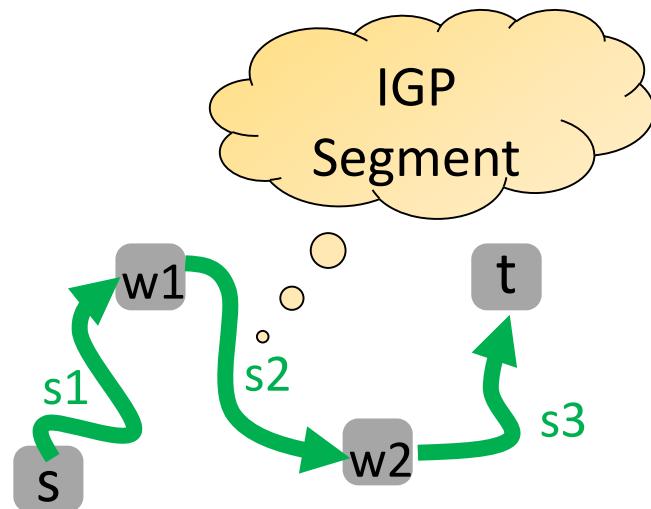
— Abstract —

Modern communication networks support fast path restoration mechanisms which allow to reroute traffic in case of (possibly multiple) link failures, in a completely *decentralized* manner and without requiring global route reconvergence. However, devising resilient path restoration algorithms is challenging as these algorithms need to be inherently *local*. Furthermore, the resulting failover paths often have to fulfill additional requirements related to the policy and function implemented by the network, such as the traversal of certain waypoints (e.g., a firewall).

This paper presents local algorithms which ensure a maximally resilient path restoration for a large family of product graphs, including the widely used tori and generalized hypercube topologies. Our algorithms provably ensure that even under multiple link failures, traffic is rerouted to the other

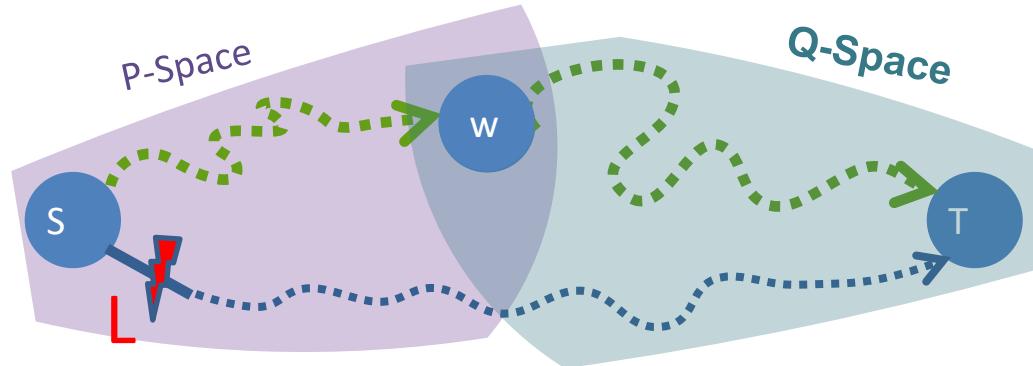
endpoints of every failed link whenever possible (*i.e.* *detouring failed links*) *enforcing waypoints and*

What About Segment Routing?



What About Segment Routing?

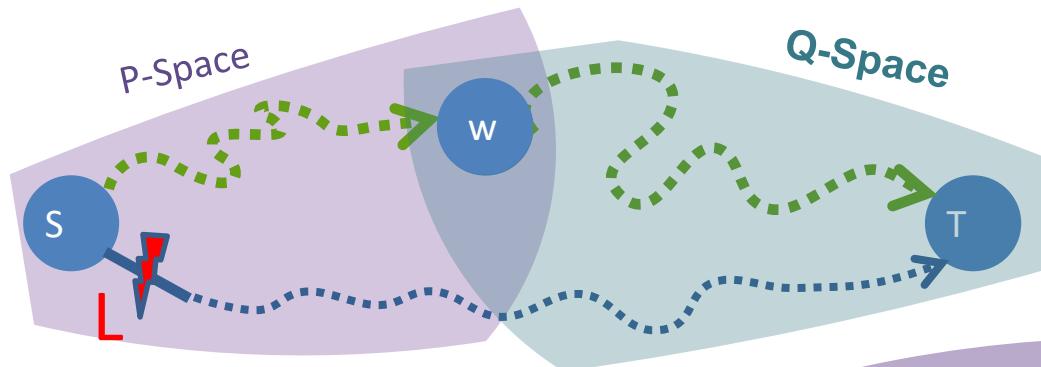
- We need two definitions:
 - **P-Space**: the nodes whose shortest path from S **does not use** L
 - **Q-Space**: the nodes whose shortest path to T **does not use** L



Idea: choose segment endpoint **w** at intersection!

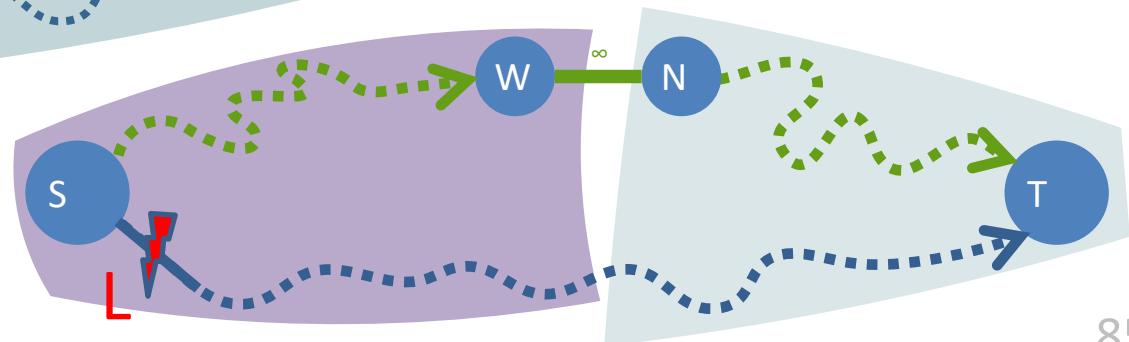
Two Cases

P-Space and Q-Space: Are **connected** subgraphs, **cover** all nodes, **overlap** or are **adjacent**



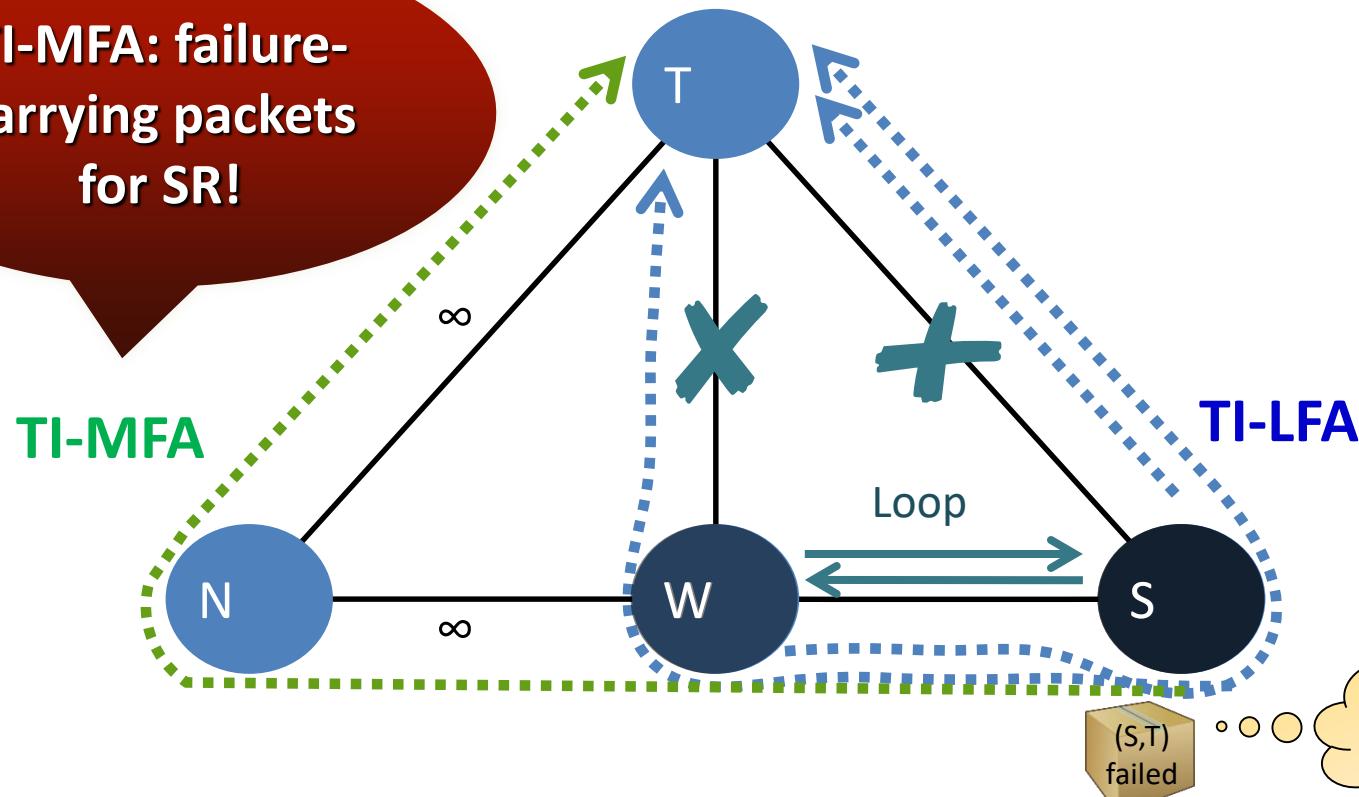
Case 1: S can simply **push** W

Case 2: S pushes W and (W,N), forces packet to enter Q-space



TI-LFA Under Double Failure

TI-MFA: failure-carrying packets for SR!



Efficient Implementation of FRR?

PURR: A Primitive for Reconfigurable Fast Reroute

(hope for the best and program for the worst)

Marco Chiesa

KTH Royal Institute of Technology

Roshan Sedar

Universitat Politècnica de Catalunya

Gianni Antichi

Queen Mary University of London

Michael Borokhovich

Independent Researcher

Andrzej Kamisiński

AGH University of Science and
Technology in Kraków

Georgios Nikolaidis

Barefoot Networks

Stefan Schmid

Faculty of Computer Science
University of Vienna

ACM Reference Format:

Marco Chiesa, Roshan Sedar, Gianni Antichi, Michael Borokhovich, Andrzej Kamisiński, Georgios Nikolaidis, and Stefan Schmid. 2019. PURR: A Primitive for Reconfigurable Fast Reroute: (hope for the best and program for the worst). In *The 15th International Conference on emerging Networking EXperiments and Technologies (CoNEXT '19), December 9–12, 2019, Orlando, FL, USA*. ACM, New York, NY, USA, 14 pages. <https://doi.org/10.1145/3359989.3365410>

See e.g.,
CoNEXT 2019

ABSTRACT

Highly dependable communication networks usually rely on some kind of Fast Re-Route (FRR) mechanism which allows to quickly re-route traffic upon failures, entirely in the data plane. This paper studies the design of FRR mechanisms for emerging reconfigurable switches.

Our main contribution is an FRR primitive for *programmable* data planes, PURR, which provides low failover latency and high switch throughput, by *avoiding packet recirculation*. PURR tolerates multiple concurrent failures and comes with minimal memory requirements, ensuring *compact* forwarding tables, by unveiling an intriguing connection to classic “string theory” (*i.e.*, stringology), and in particular, the shortest common supersequence problem. PURR is well-suited for high-speed match-action forwarding architectures (*e.g.*, PISA) and supports the implementation of arbitrary network-wide FRR mechanisms. Our simulations and prototype implementation (on an FPGA and Tofino) show that PURR improves

1 INTRODUCTION

Emerging applications, *e.g.*, in the context of business [21] and entertainment [57], pose stringent requirements on the dependability and performance of the underlying communication networks, which have become a critical infrastructure of our digital society. In order to meet such requirements, many communication networks provide *Fast Re-Route* (FRR) mechanisms [5, 39, 64] which allow to quickly reroute traffic upon unexpected failures, entirely in the

A Recent Survey

A Survey of Fast-Recovery Mechanisms in Packet-Switched Networks

Marco Chiesa, Andrzej Kaminski, Jacek Rak, Gabor Retvari, and Stefan Schmid.
IEEE Communications Surveys and Tutorials (**COMST**), 2021.

References

On the Feasibility of Perfect Resilience with Local Fast Failover

Klaus-Tycho Foerster, Juho Hirvonen, Yvonne-Anne Pignolet, Stefan Schmid, and Gilles Tredan.

SIAM Symposium on Algorithmic Principles of Computer Systems (**APOCS**), Alexandria, Virginia, USA, January 2021.

Brief Announcement: What Can(not) Be Perfectly Rerouted Locally

Klaus-Tycho Foerster, Juho Hirvonen, Yvonne-Anne Pignolet, Stefan Schmid, and Gilles Tredan.

International Symposium on Distributed Computing (**DISC**), Freiburg, Germany, October 2020.

Improved Fast Rerouting Using Postprocessing

Klaus-Tycho Foerster, Andrzej Kaminski, Yvonne-Anne Pignolet, Stefan Schmid, and Gilles Tredan.

IEEE Transactions on Dependable and Secure Computing (**TDSC**), 2020.

Resilient Capacity-Aware Routing

Stefan Schmid, Nicolas Schnepf and Jiri Srba.

27th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (**TACAS**), Virtual Conference, March 2021.

AalWiNes: A Fast and Quantitative What-If Analysis Tool for MPLS Networks

Peter Gjøl Jensen, Morten Konggaard, Dan Kristiansen, Stefan Schmid, Bernhard Clemens Schrenk, and Jiri Srba.

16th ACM International Conference on emerging Networking EXperiments and Technologies (**CoNEXT**), Barcelona, Spain, December 2020.

P-Rex: Fast Verification of MPLS Networks with Multiple Link Failures

Jesper Stenbjerg Jensen, Troels Beck Krogh, Jonas Sand Madsen, Stefan Schmid, Jiri Srba, and Marc Tom Thorgersen.

14th ACM International Conference on emerging Networking EXperiments and Technologies (**CoNEXT**), Heraklion/Crete, Greece, December 2018.

Polynomial-Time What-If Analysis for Prefix-Manipulating MPLS Networks

Stefan Schmid and Jiri Srba.

37th IEEE Conference on Computer Communications (**INFOCOM**), Honolulu, Hawaii, USA, April 2018.

More References

[Randomized Local Fast Rerouting for Datacenter Networks with Almost Optimal Congestion](#)

Gregor Bankhamer, Robert Elsässer, and Stefan Schmid..

International Symposium on Distributed Computing (**DISC**), Freiburg, Germany, October 2021.

[Bonsai: Efficient Fast Failover Routing Using Small Arborescences](#)

Klaus-Tycho Foerster, Andrzej Kamisinski, Yvonne-Anne Pignolet, Stefan Schmid, and Gilles Tredan.

49th IEEE/IFIP International Conference on Dependable Systems and Networks (**DSN**), Portland, Oregon, USA, June 2019.

[CASA: Congestion and Stretch Aware Static Fast Rerouting](#)

Klaus-Tycho Foerster, Yvonne-Anne Pignolet, Stefan Schmid, and Gilles Tredan

38th IEEE Conference on Computer Communications (**INFOCOM**), Paris, France, April 2019.

[Load-Optimal Local Fast Rerouting for Dense Networks](#)

Michael Borokhovich, Yvonne-Anne Pignolet, Gilles Tredan, and Stefan Schmid.

IEEE/ACM Transactions on Networking (**TON**), 2018.

[PURR: A Primitive for Reconfigurable Fast Reroute](#)

Marco Chiesa, Roshan Sedar, Gianni Antichi, Michael Borokhovich, Andrzej Kamisinski, Georgios Nikolaidis, and Stefan Schmid.

15th ACM International Conference on emerging Networking EXperiments and Technologies (**CoNEXT**), Orlando, Florida, USA, December 2019.

Artefact Evaluation: Available, Functional, Reusable.

[On the Resiliency of Static Forwarding Tables](#)

In IEEE/ACM Transactions on Networking (**ToN**), 2017

M. Chiesa, I. Nikolaevskiy, S. Mitrovic, A. Gurkov, A. Madry, M. Schapira, S. Shenker



Questions?