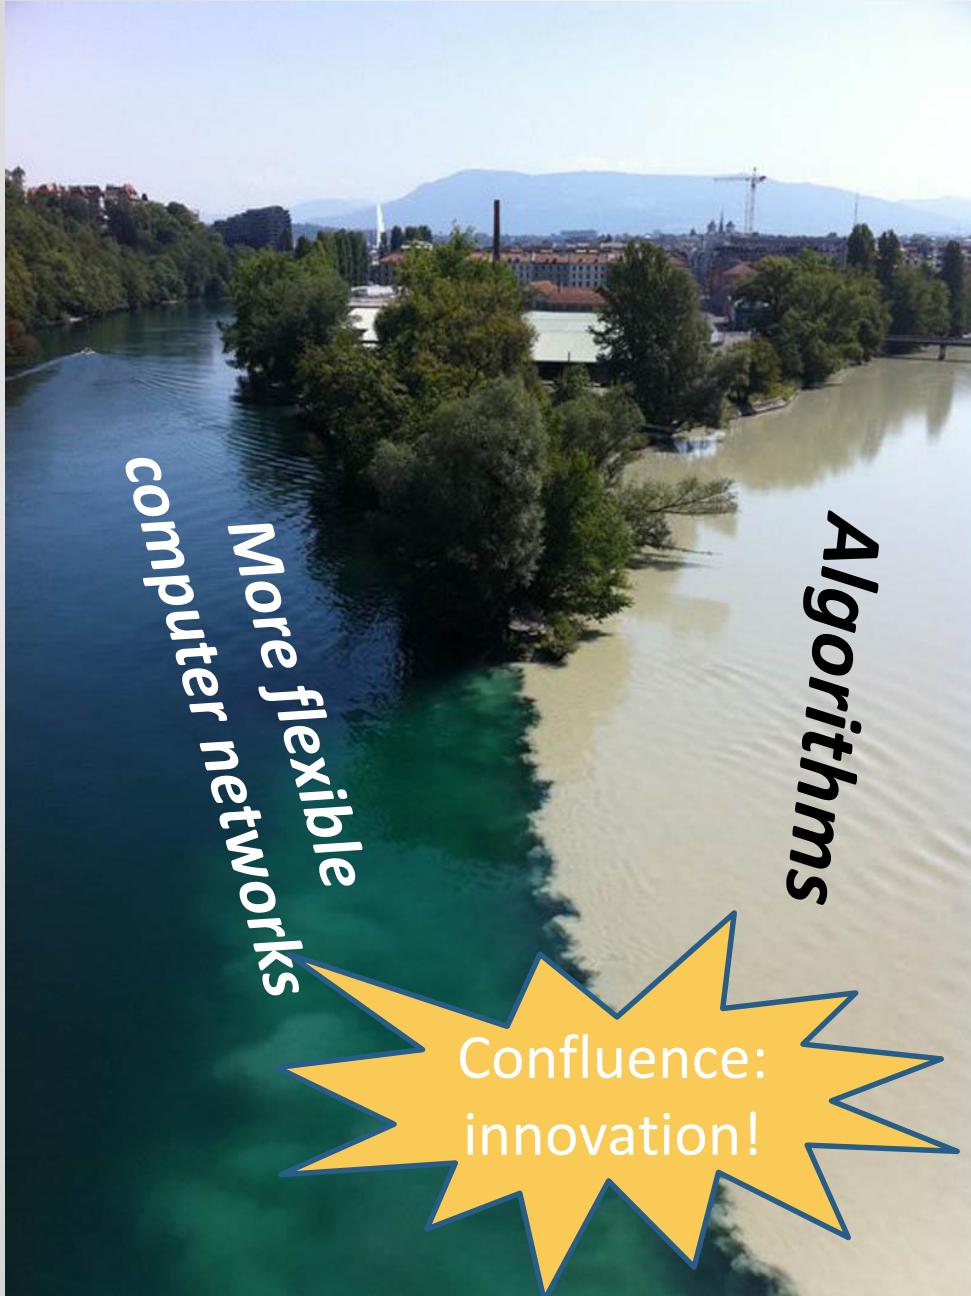


Tutorial: Algorithmic Issues in Network Resource Reservation Problems

Stefan Schmid

Aalborg University, Denmark & TU Berlin, Germany

A rehash: It's a great time to be a scientist!



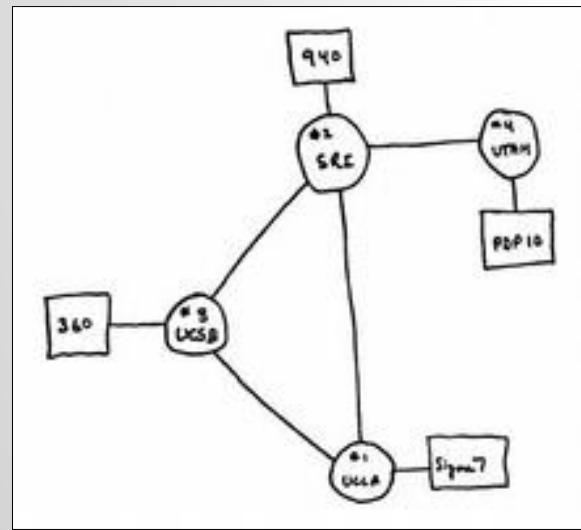
How to exploit
these flexibilities?
How not to shoot
in our feet?
Can be challenging!

"We are at an interesting inflection point!"
Keynote by George Varghese
at SIGCOMM 2014



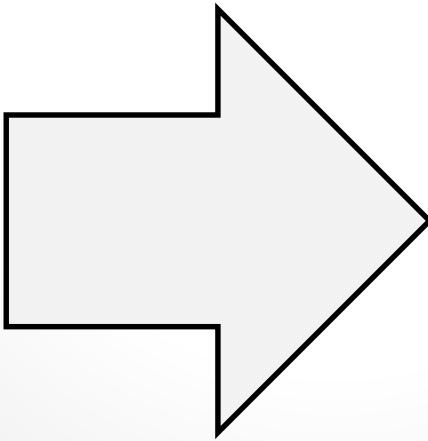
New Flexibilities: It's About Time!

- Datacenter networks, enterprise networks, Internet: a **critical infrastructure** of the information society
- We have seen a huge **shift in scale and applications...**
- ... but many Internet protocols **hardly changed!**



Applications: file transfer, email

Goal: **connectivity** between researchers



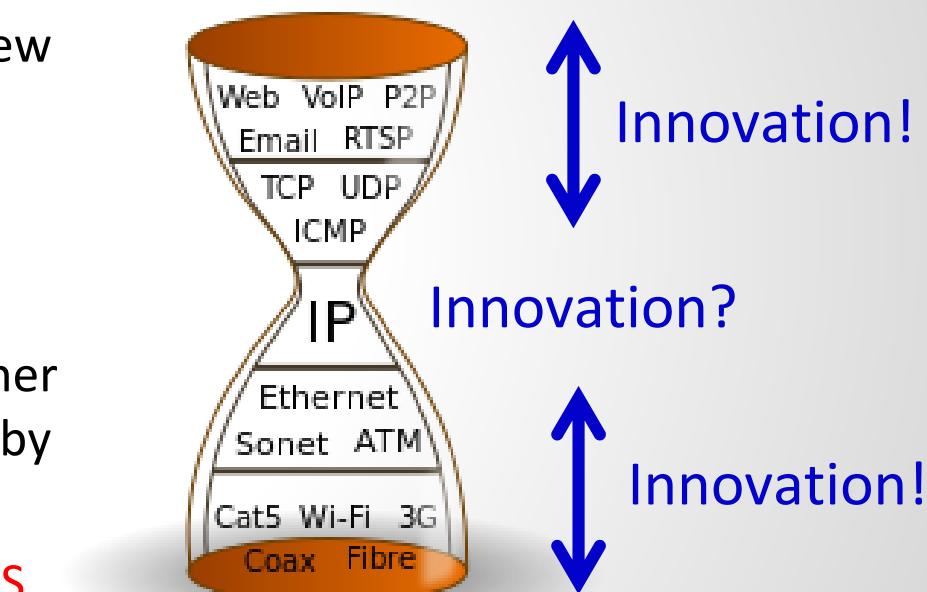
Applications: live streaming, IoT, etc.

Goal: **quality-of-service**, predictable performance, low latency, ...

Opportunity 1 of Network Virtualization: Overcoming Ossification

- ❑ Recent concern: Ossification in the network core
 - ❑ Are computer networks **future-proof**?
 - ❑ Meet the **new requirements** of new applications?

- ❑ Example Internet-of-Things:
 - ❑ IPv4: ~**4.3** billion addresses, Gartner study: **20+** billion “smart things” by 2020
 - ❑ New security threats: recent **DDoS attack** based on IoT (almost 1TB/s, coming from webcams, babyphones, etc.)



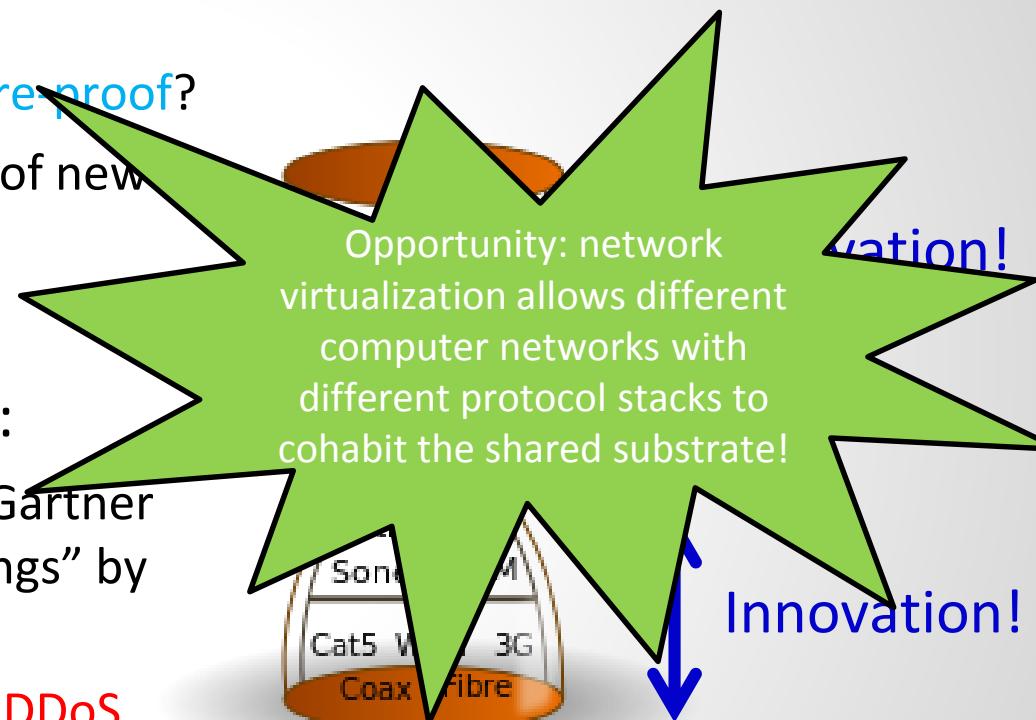
Opportunity 1 of Network Virtualization: Overcoming Ossification

- Recent concern: Ossification in the network core

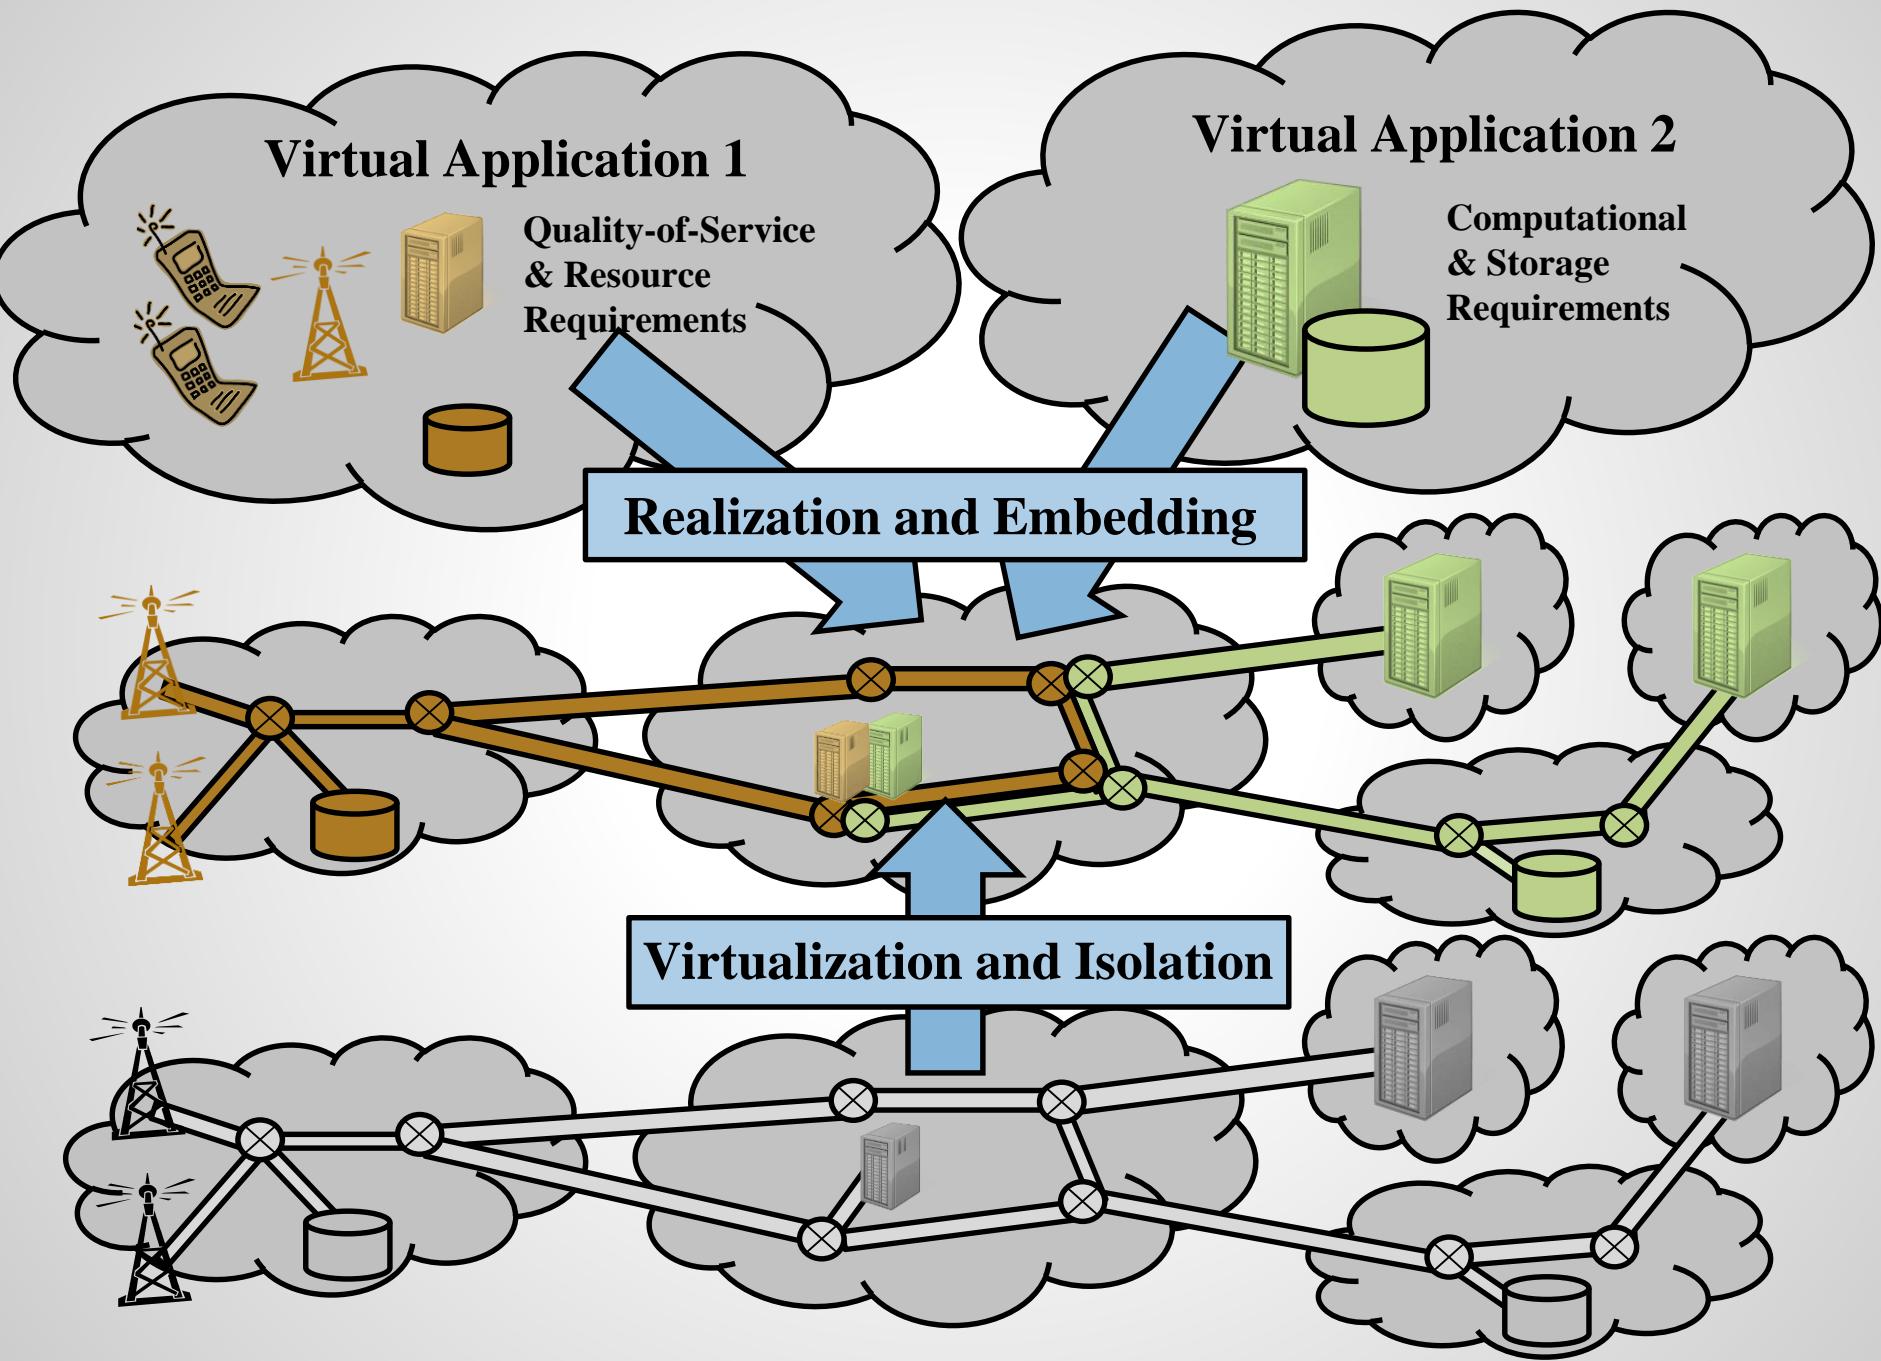
- Are computer networks **future proof?**
- Meet the **new requirements** of new applications?

- Example Internet-of-Things:

- IPv4: ~**4.3** billion addresses, Gartner study: **20+** billion “smart things” by 2020
- New security threats: recent **DDoS attack** based on IoT (almost 1TB/s, coming from webcams, babyphones, etc.)

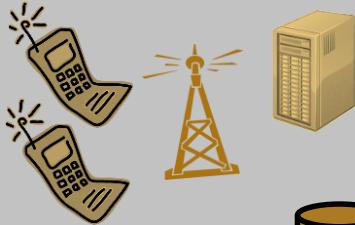


Opportunity 2: Enable Resource Sharing for Improved Utilization



Opportunity 2: Enable Resource Sharing for Improved Utilization

Virtual Application 1



Quality-of-Service
& Resource Requirements

Virtual Application 2



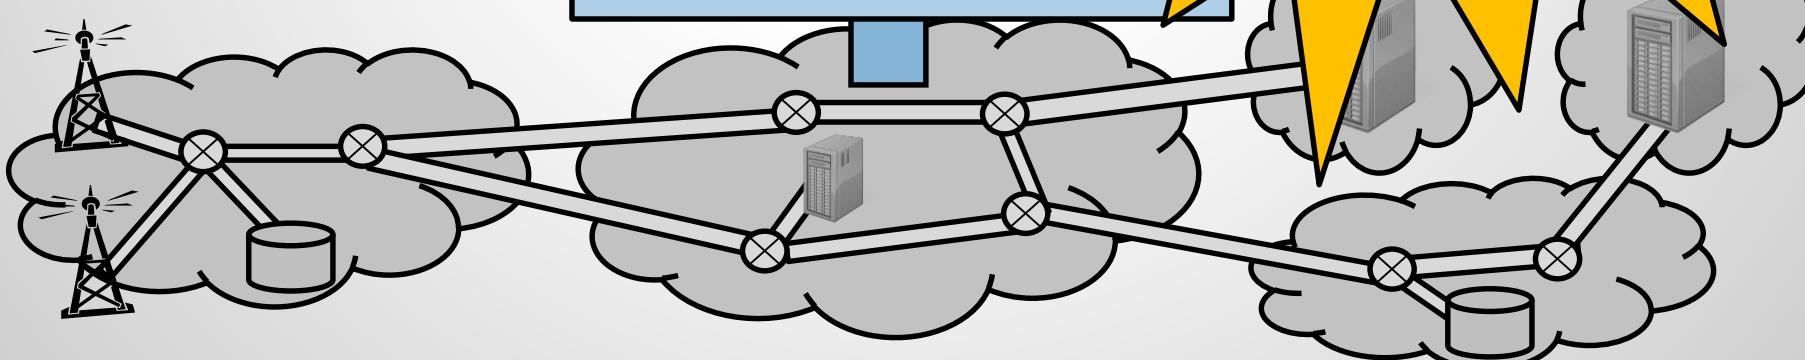
Computational
& Storage
Requirements

Realization and Embedding

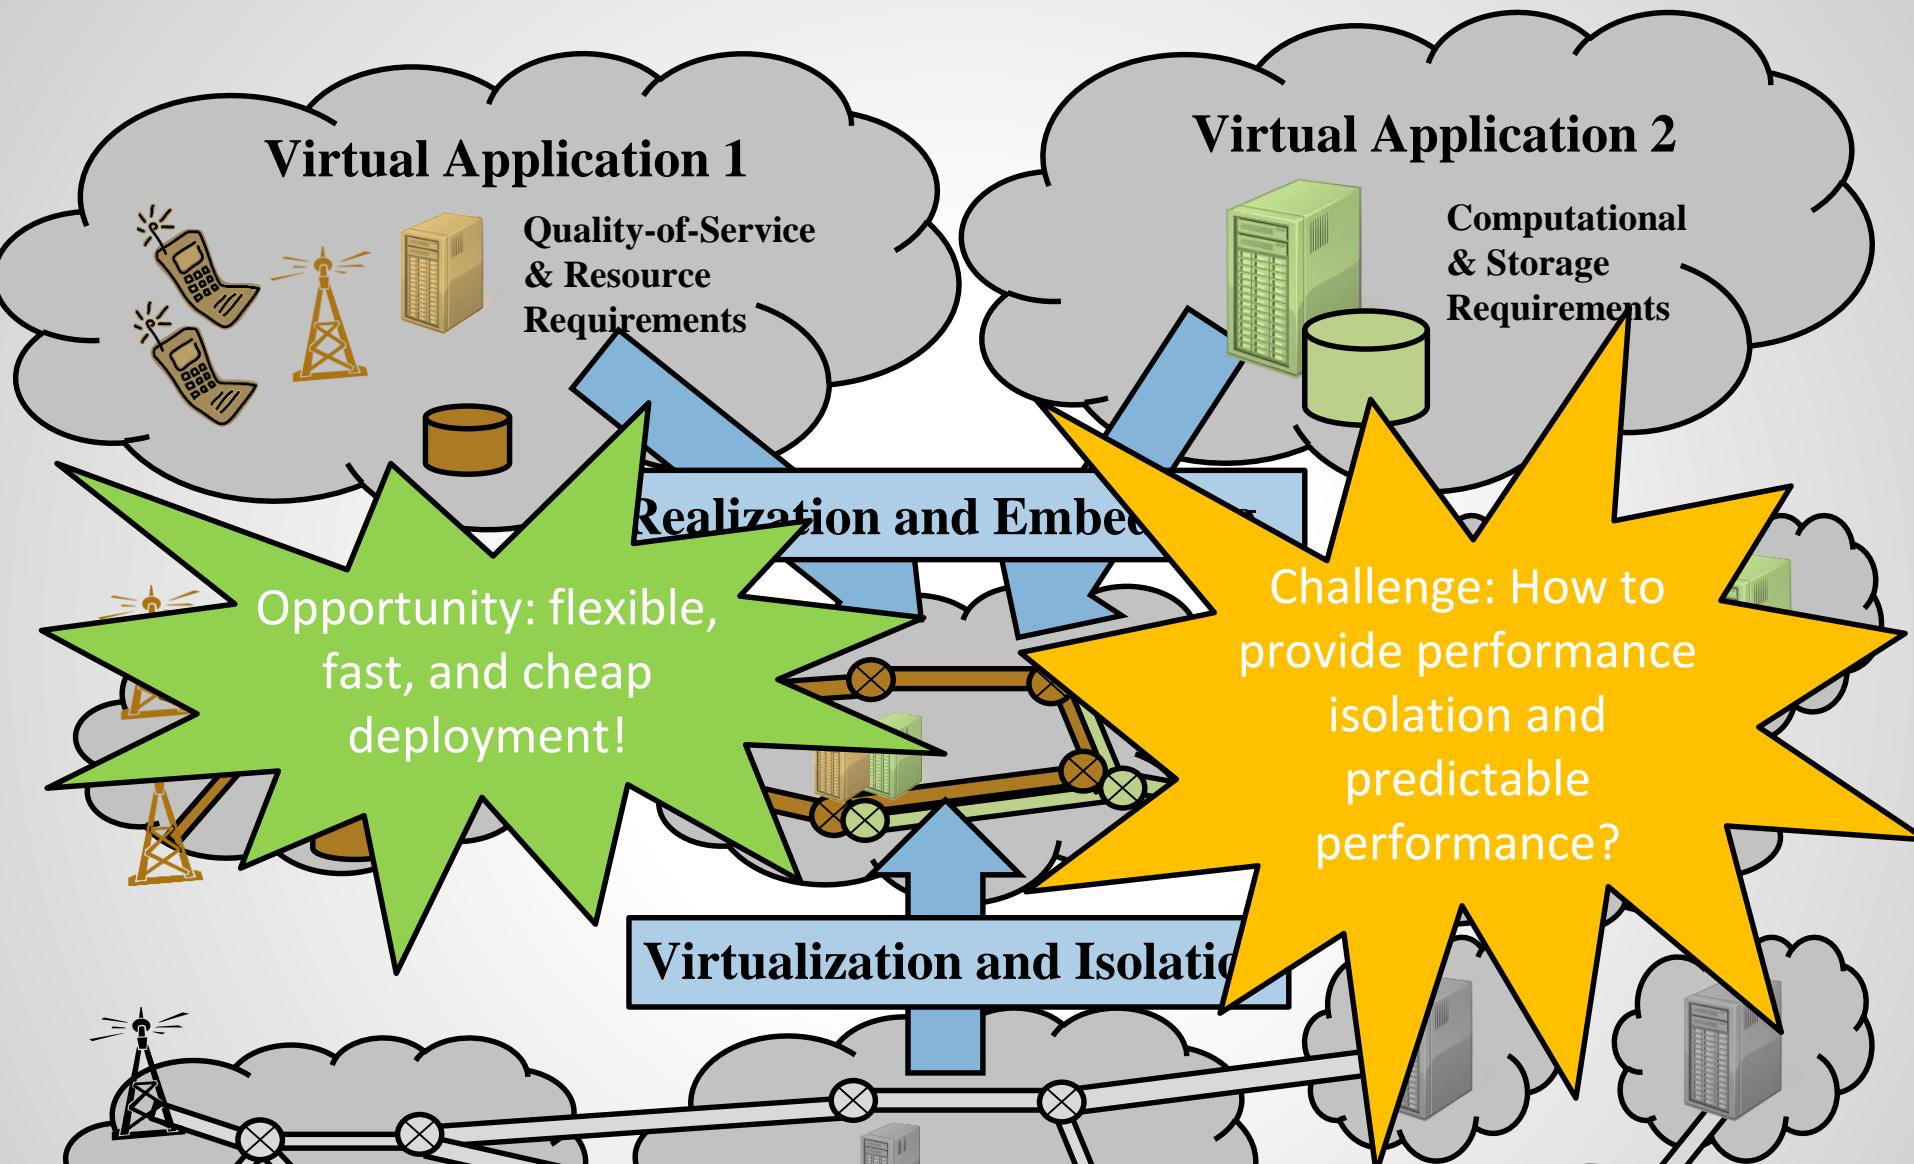
Opportunity: flexible,
fast, and cheap
deployment!

Challenge: How to
provide performance
isolation and
predictable
performance?

Virtualization and Isolation



Opportunity 2: Enable Resource Sharing for Improved Utilization



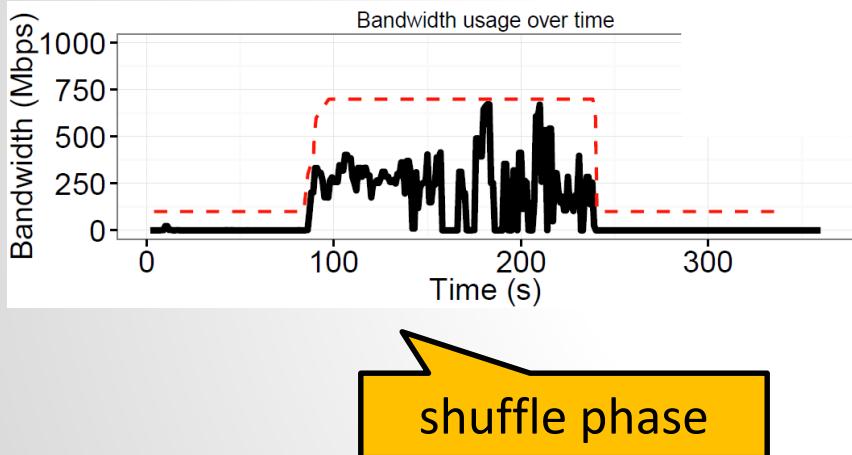
In general: For a predictable application performance, performance isolation needs to be provided along *all* involved resources.

Focus Today: *The Network*

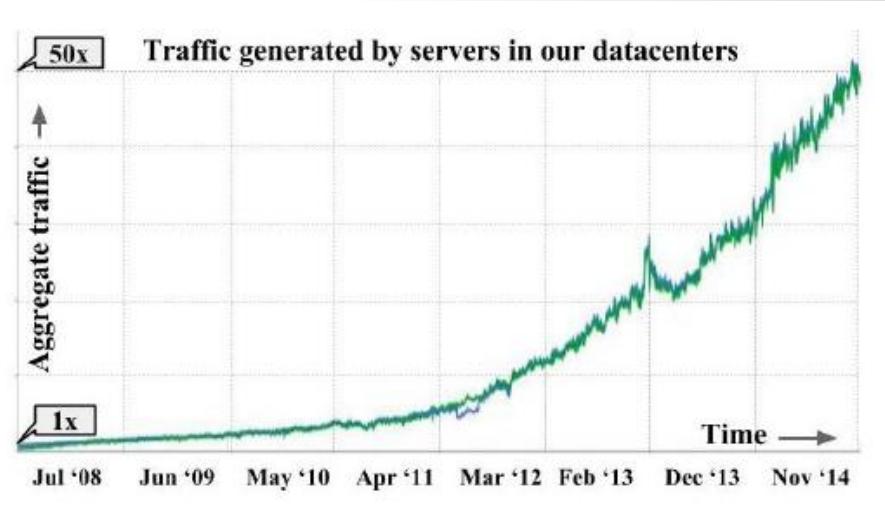
The Network Matters

- Cloud-based applications generate significant network traffic
 - E.g., scale-out databases, streaming, batch processing applications

Example 1: Hadoop Terrasort job



Example 2: Aggregate Server Traffic in Google datacenter



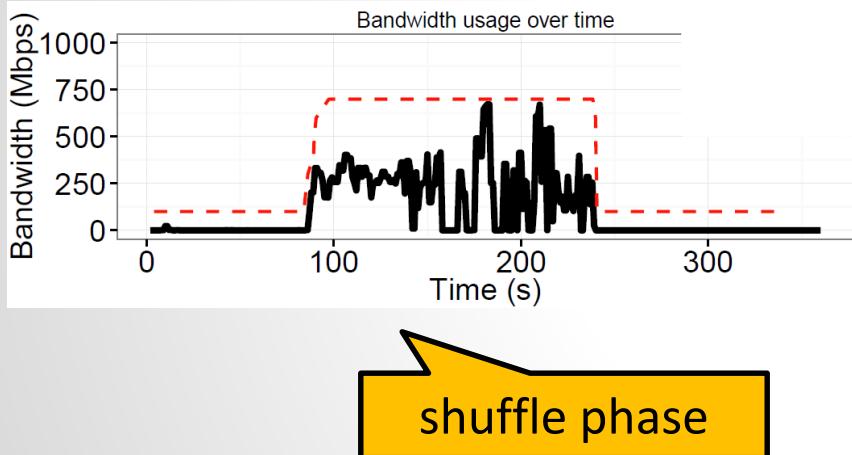
Example 3: More memory-based systems (network becoming bottleneck again)

Jupiter rising @ SIGCOMM 2015

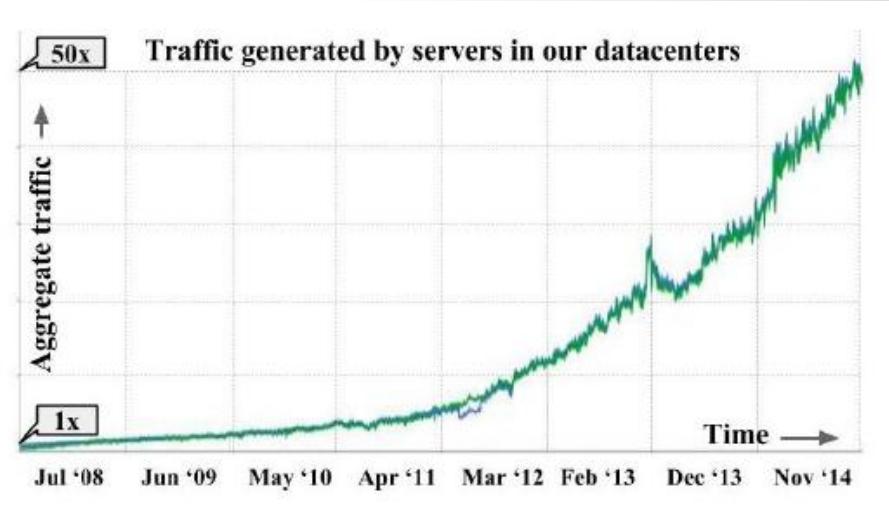
The Network Matters

- Cloud-based applications generate significant network traffic
 - E.g., scale-out databases, streaming, batch processing applications

Example 1: Hadoop Terrasort job



Example 2: Aggregate Server Traffic in Google datacenter



Example 3: More memory-based systems (network becoming bottleneck again)

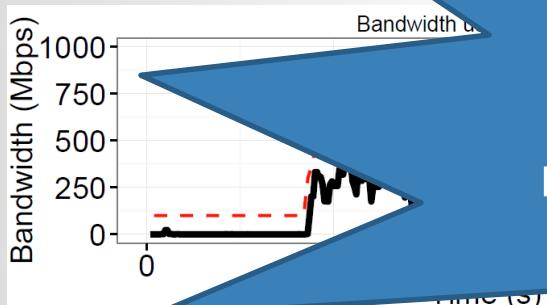
Jupiter rising @ SIGCOMM 2015

As much time is spent on communication: For a predictable application performance, bandwidth resources need to be reserved explicitly.

The Network Matters

- Cloud-based applications generate significant network traffic
 - E.g., scale-out databases, streaming, batch processing applications

Example 1: MapReduce Terrasort



Example 2:

Aggregate Server Traffic

Ideally, communication should
be **local**. And bandwidth
reservations **along short paths!**
An algorithmic problem.



shuffle phase

Example 3: More memory-based systems
(network becoming bottleneck again)

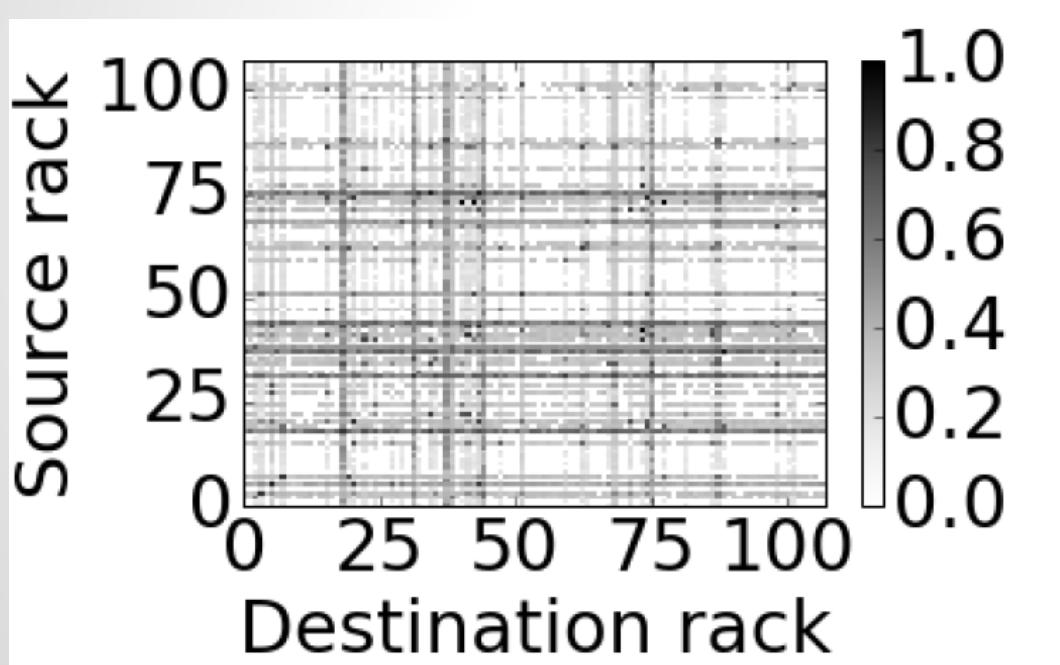
Jupiter rising @ SIGCOMM 2015

As much time is spent on communication: For a predictable application performance, bandwidth resources need to be reserved explicitly.

Structure in Traffic Matrix = Optimization Opportunities

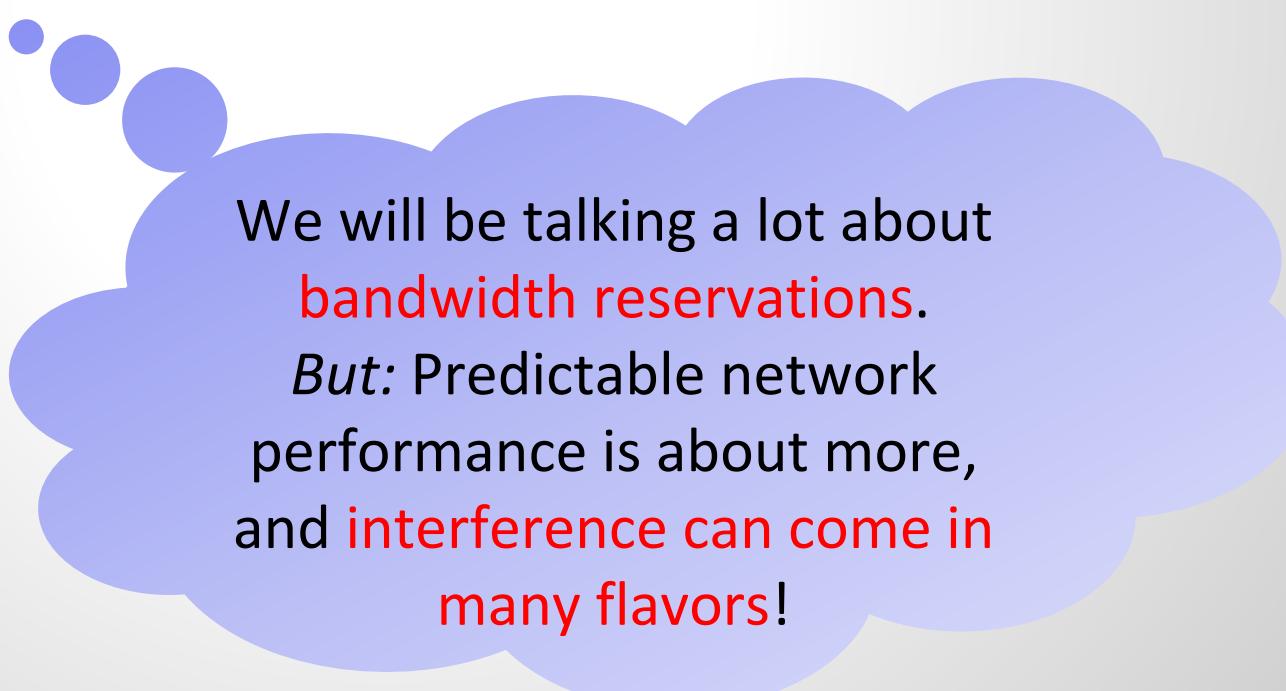
- At the same time, traffic matrices are often **far from random** and uniform, but have a lot of **structure** and are **sparse**

Example 1: Often **little to no traffic** between many racks



Without taking
this structure into
account, some
links may be
overprovisioned
and others
underprovisioned

Focus Today: *The Network*



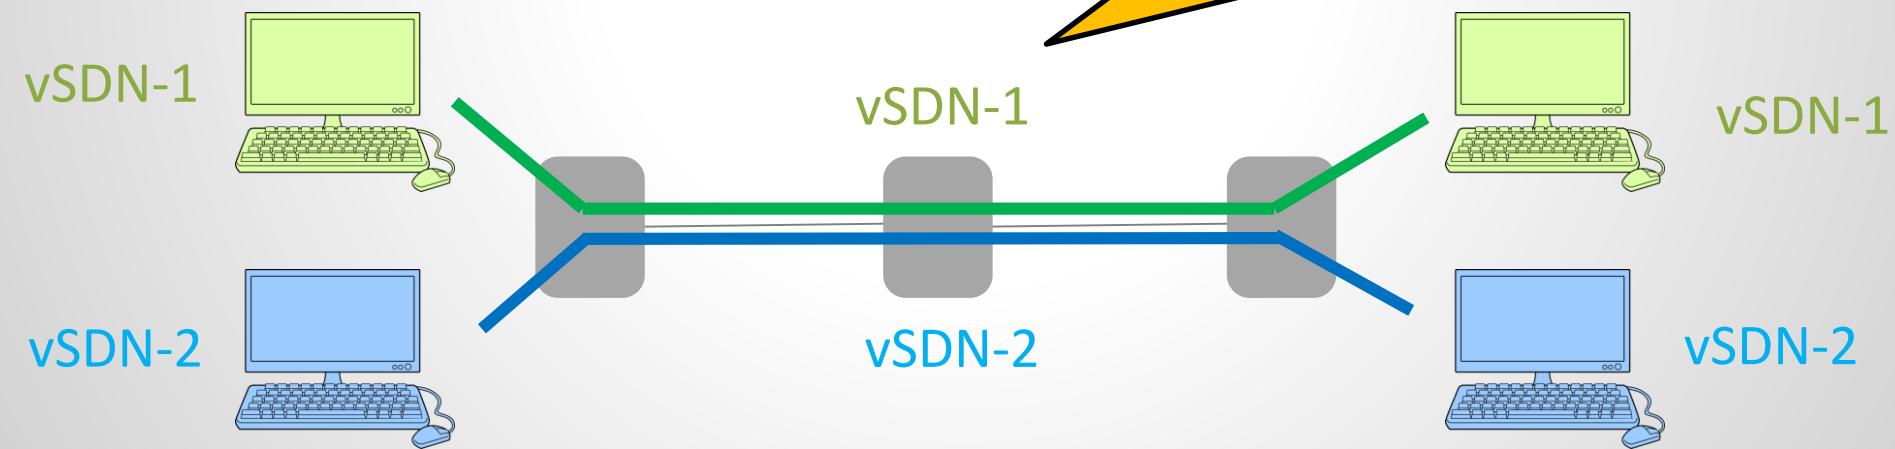
We will be talking a lot about
bandwidth reservations.
But: Predictable network
performance is about more,
and **interference can come in**
many flavors!

<remark>

The Many Faces of Performance Interference

Consider: 2 SDN-based
virtual networks (vSDNs)
sharing physical resources!

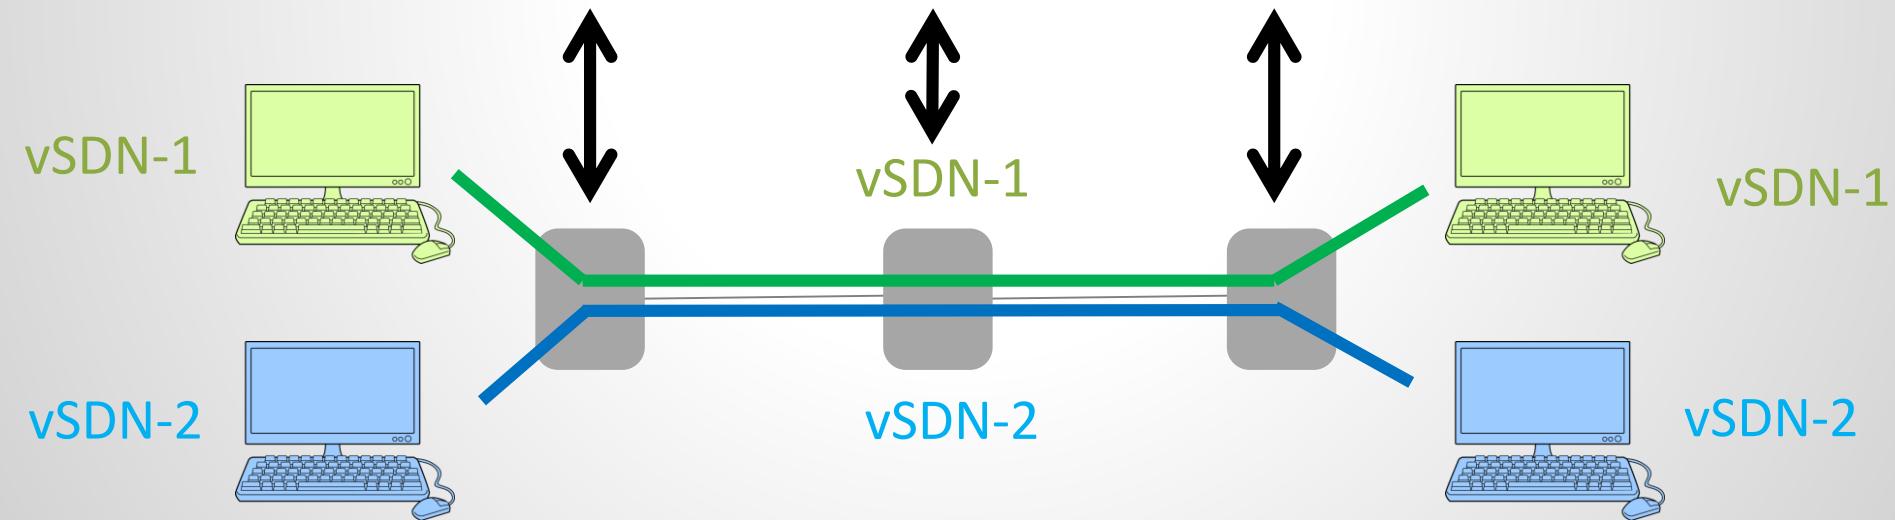
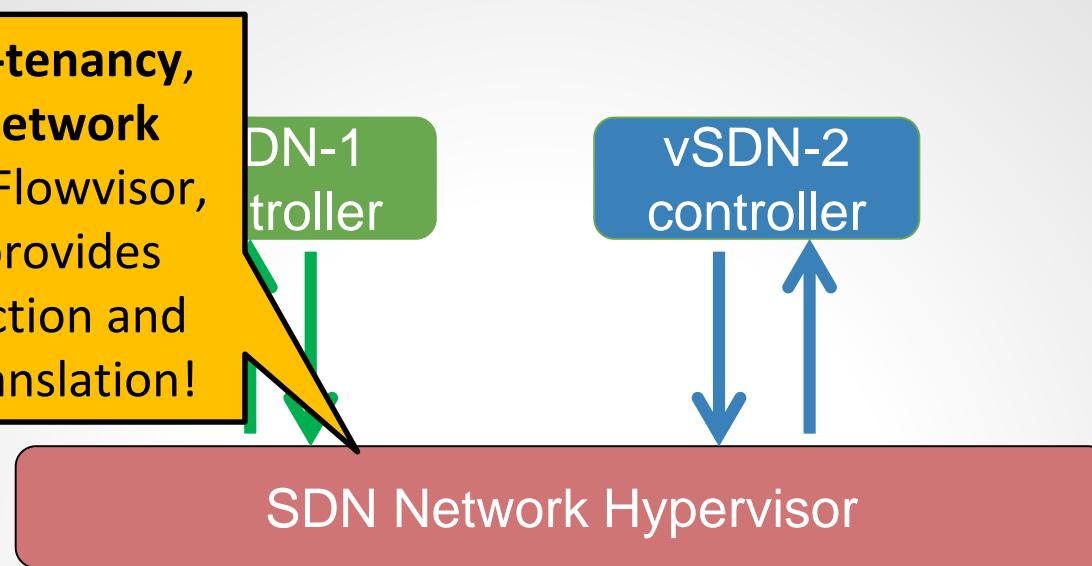
Assume: perfect
performance isolation on
the network!



An Experiment: 2 vSDNs with bw guarantee!

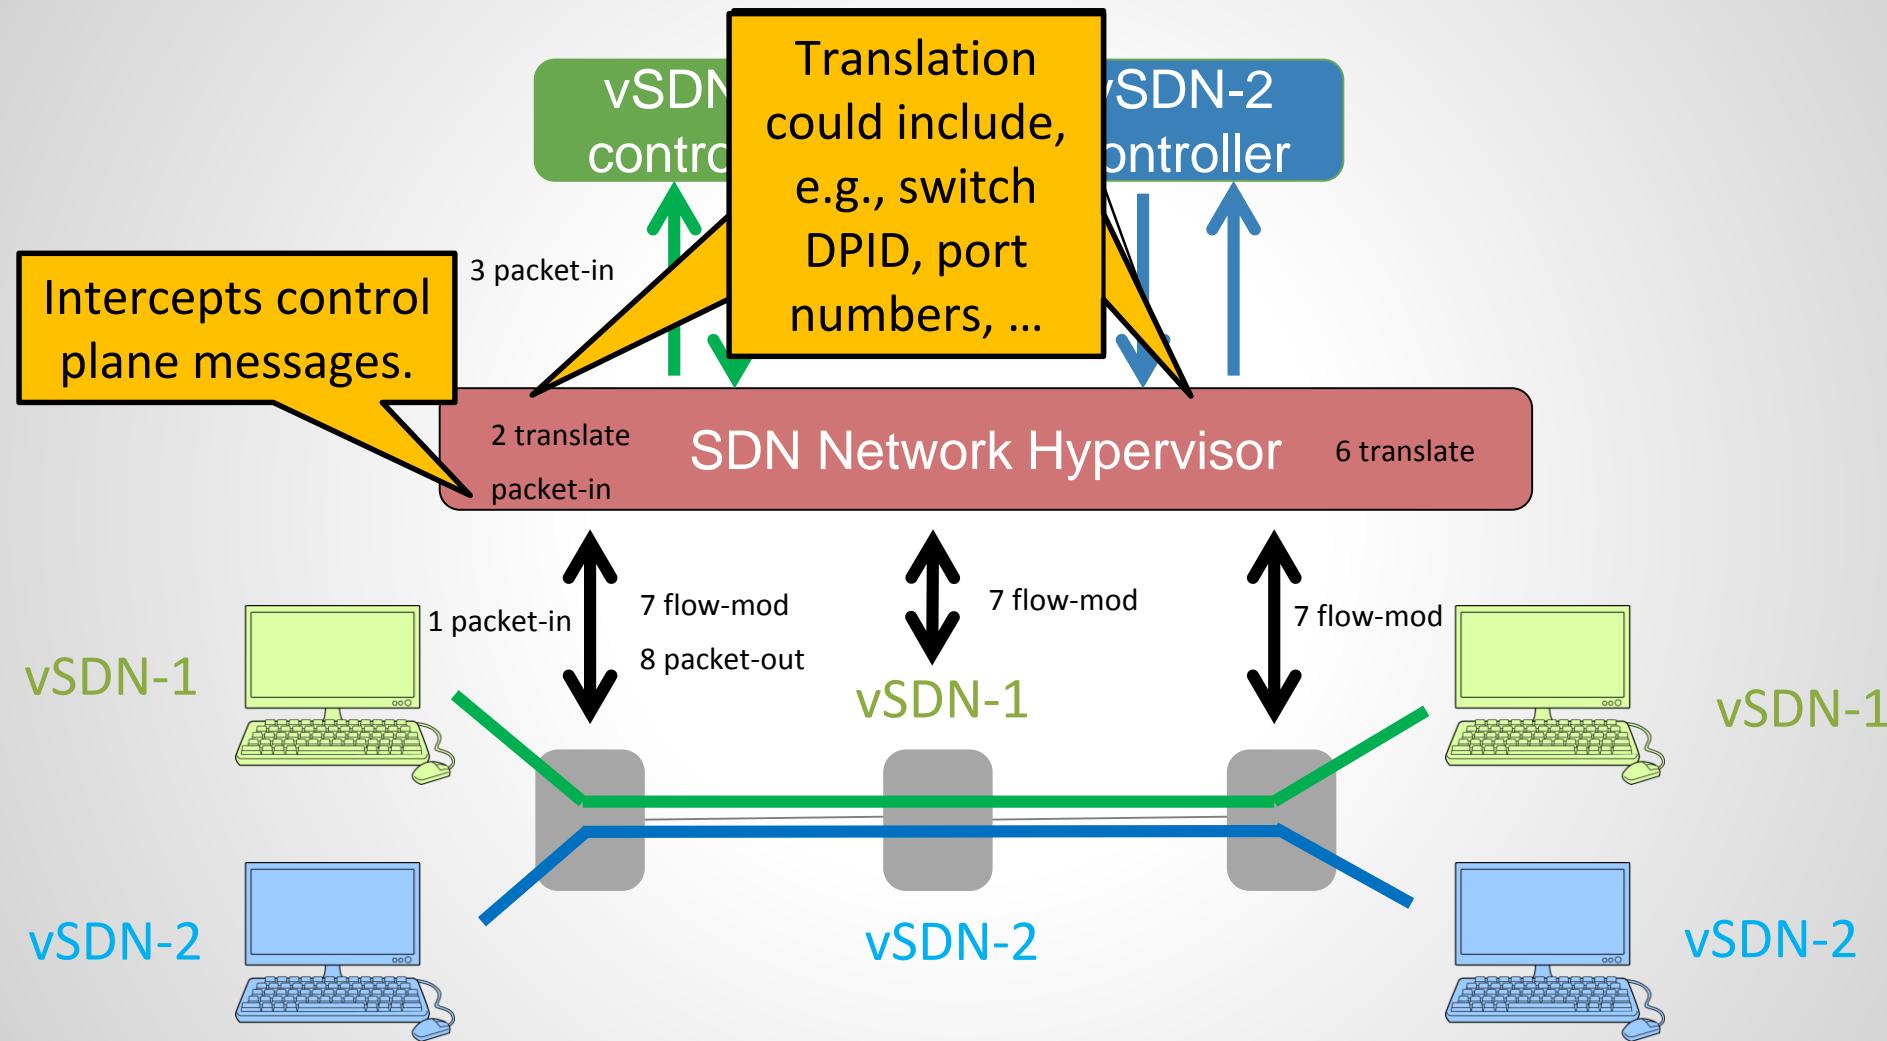
The Many Faces of Performance Interference

To enable **multi-tenancy**, take **existing network hypervisor** (e.g. Flowvisor, OpenVirtex): provides network abstraction and control plane translation!



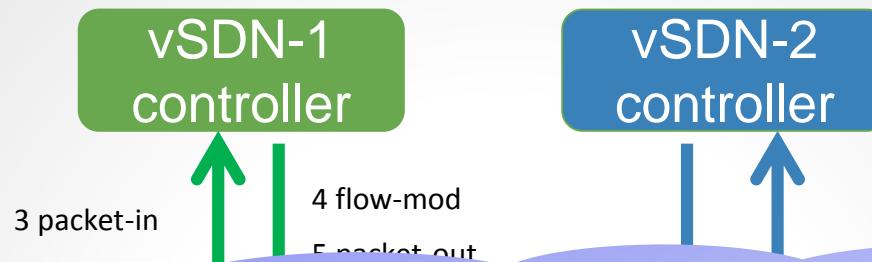
An Experiment: 2 vSDNs with bw guarantee!

The Many Faces of Performance Interference

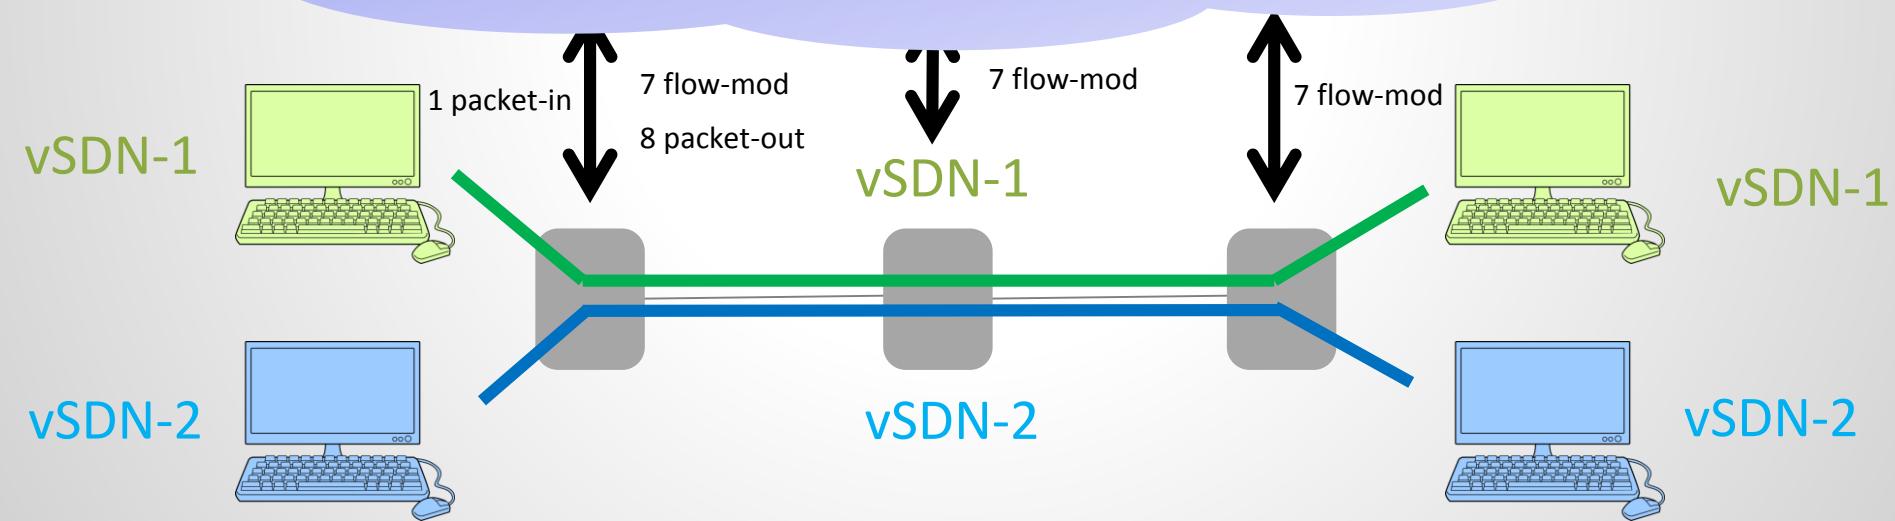


An Experiment: 2 vSDNs with bw guarantee!

The Many Faces of Performance Interference

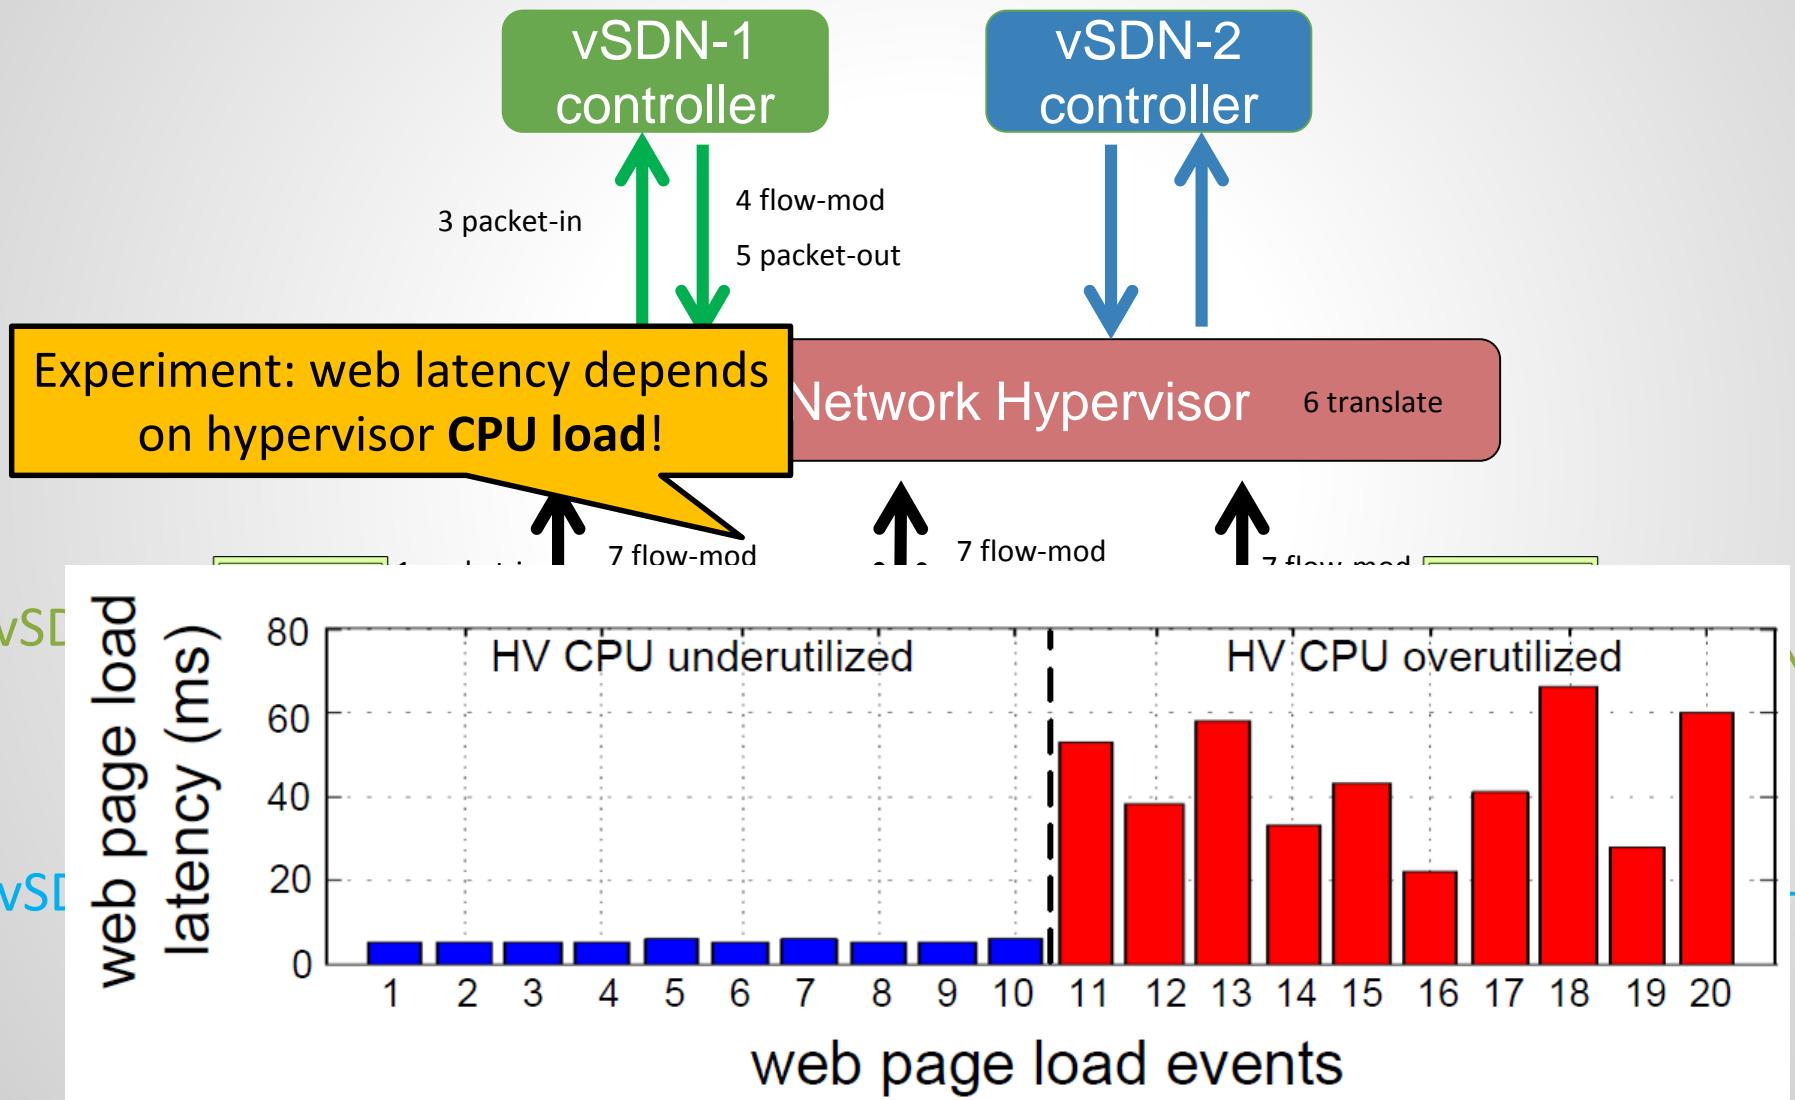


It turns out: the network hypervisor can be source of unpredictable performance!

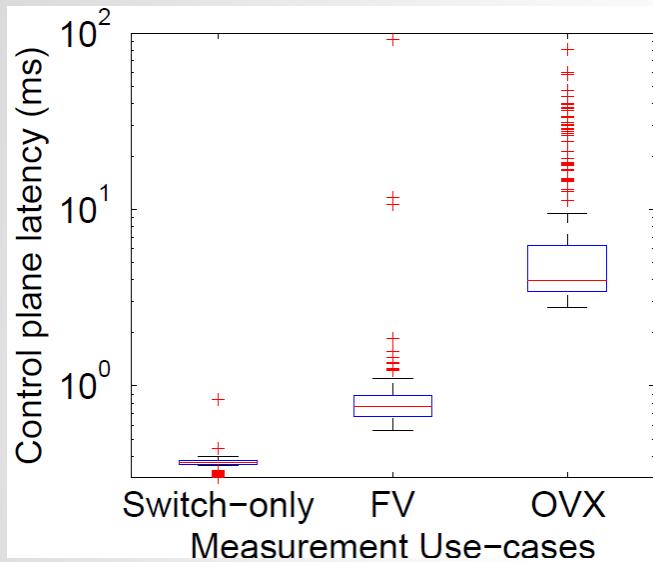


An Experiment: 2 vSDNs with bw guarantee!

The Many Faces of Performance Interference

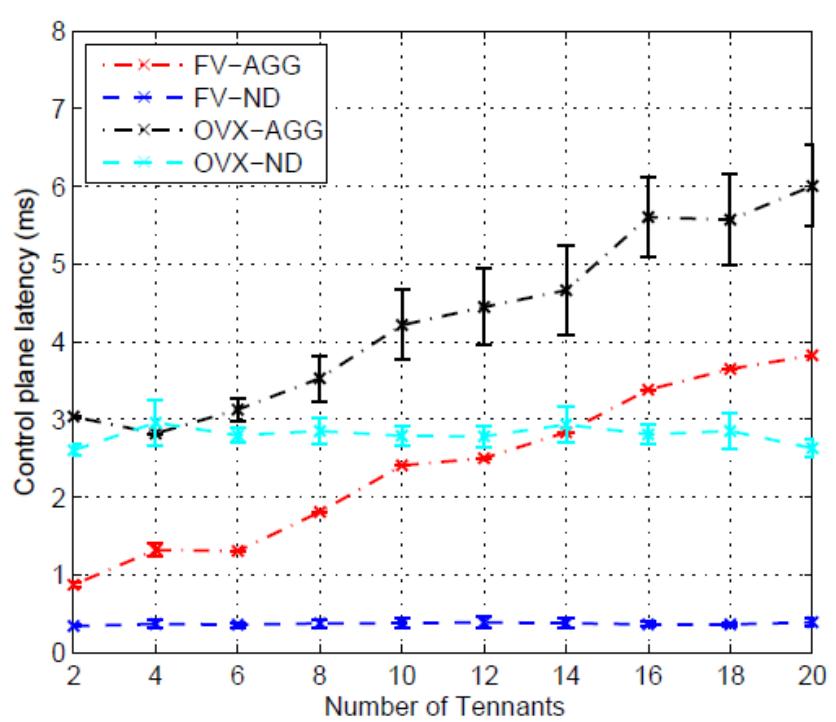


The Many Faces of Performance Interference

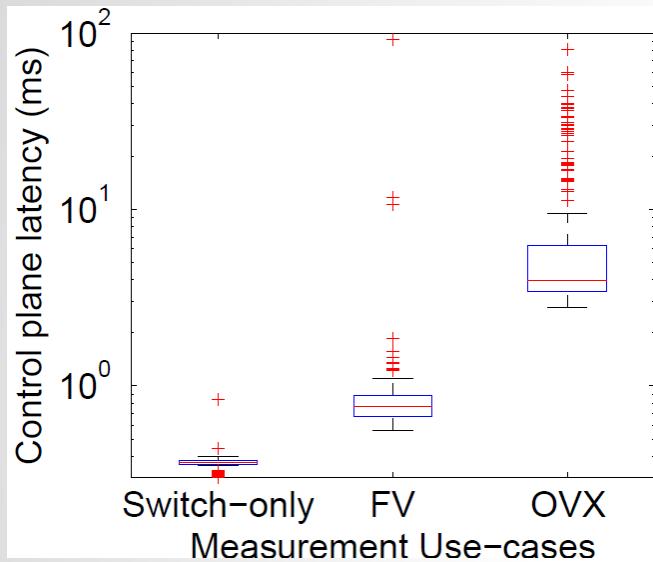


**Performance also depends
on hypervisor type...**
*(multithreaded or not, which version
of Nagle's algorithm, etc.)*

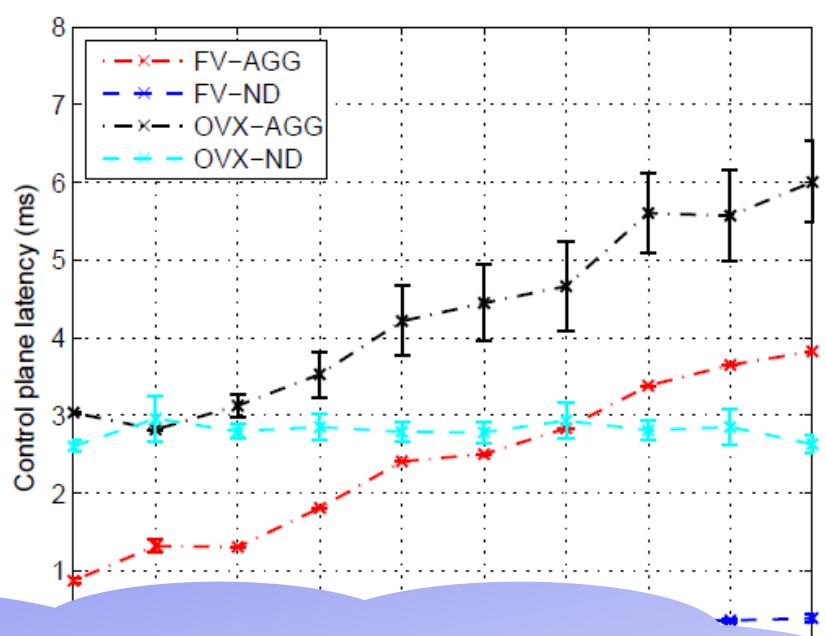
... number of tenants...



The Many Faces of Performance Interference



... number of tenants...



Performance also depends
on hypervisor type

(multithreaded)

Conclusion: for predictable performance,
need to account for all resources!

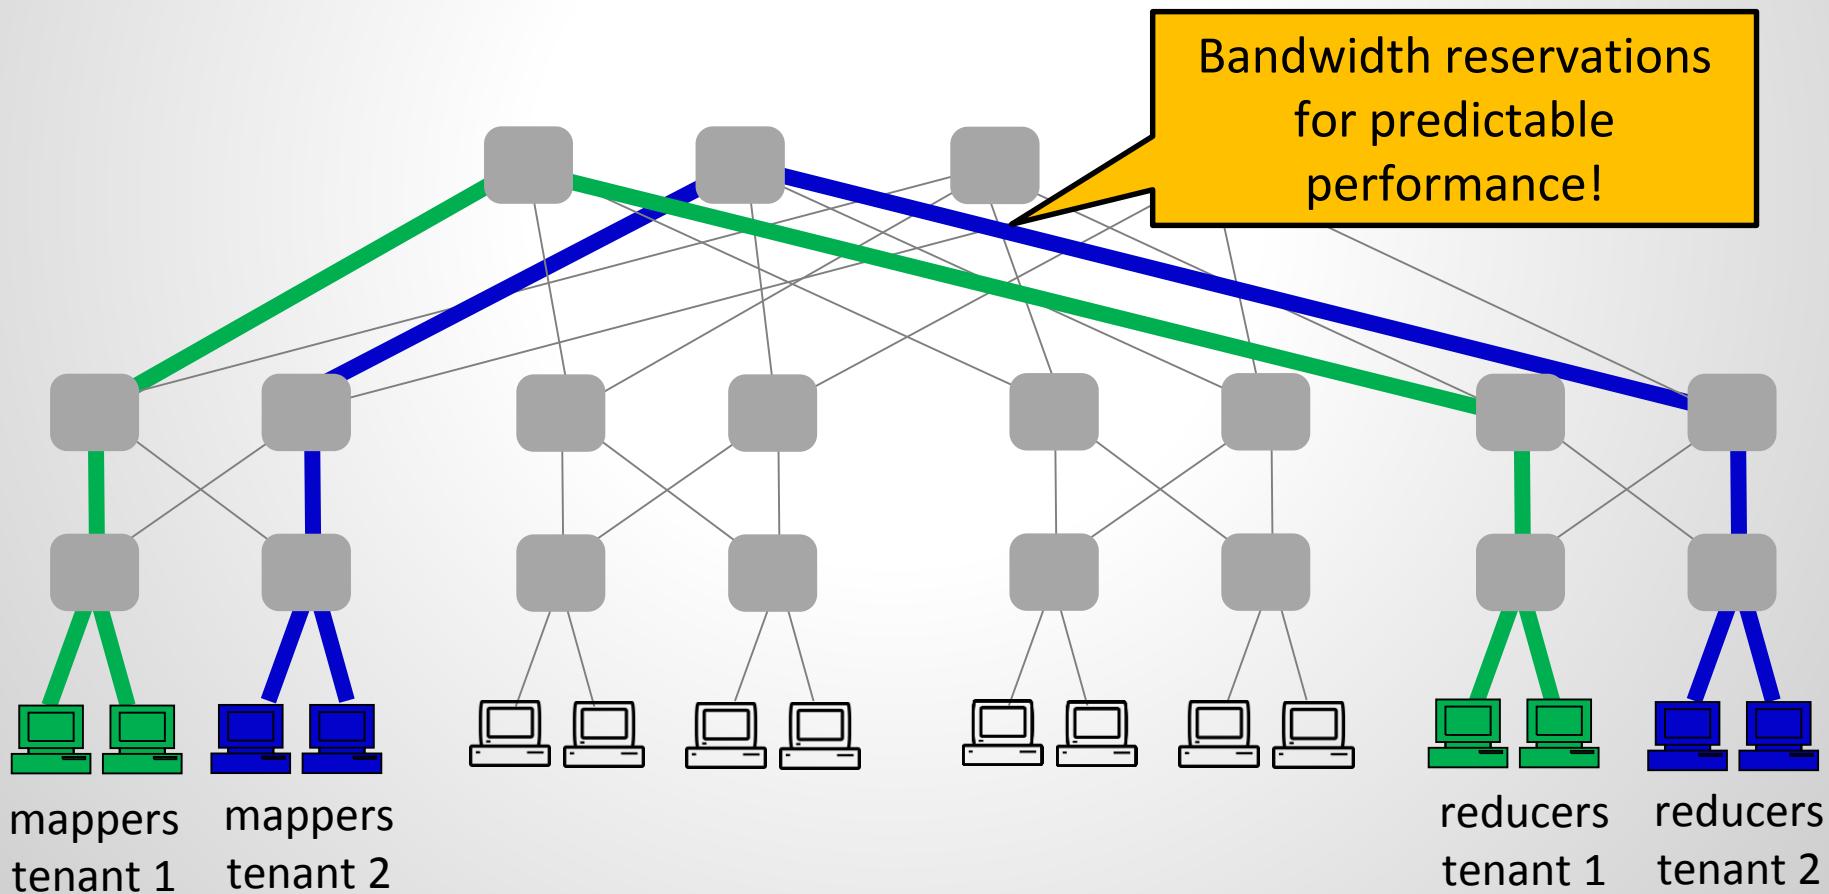
But let us now focus on the network itself.

</remark>

First Algorithmic Challenge:

Keep the traffic local!

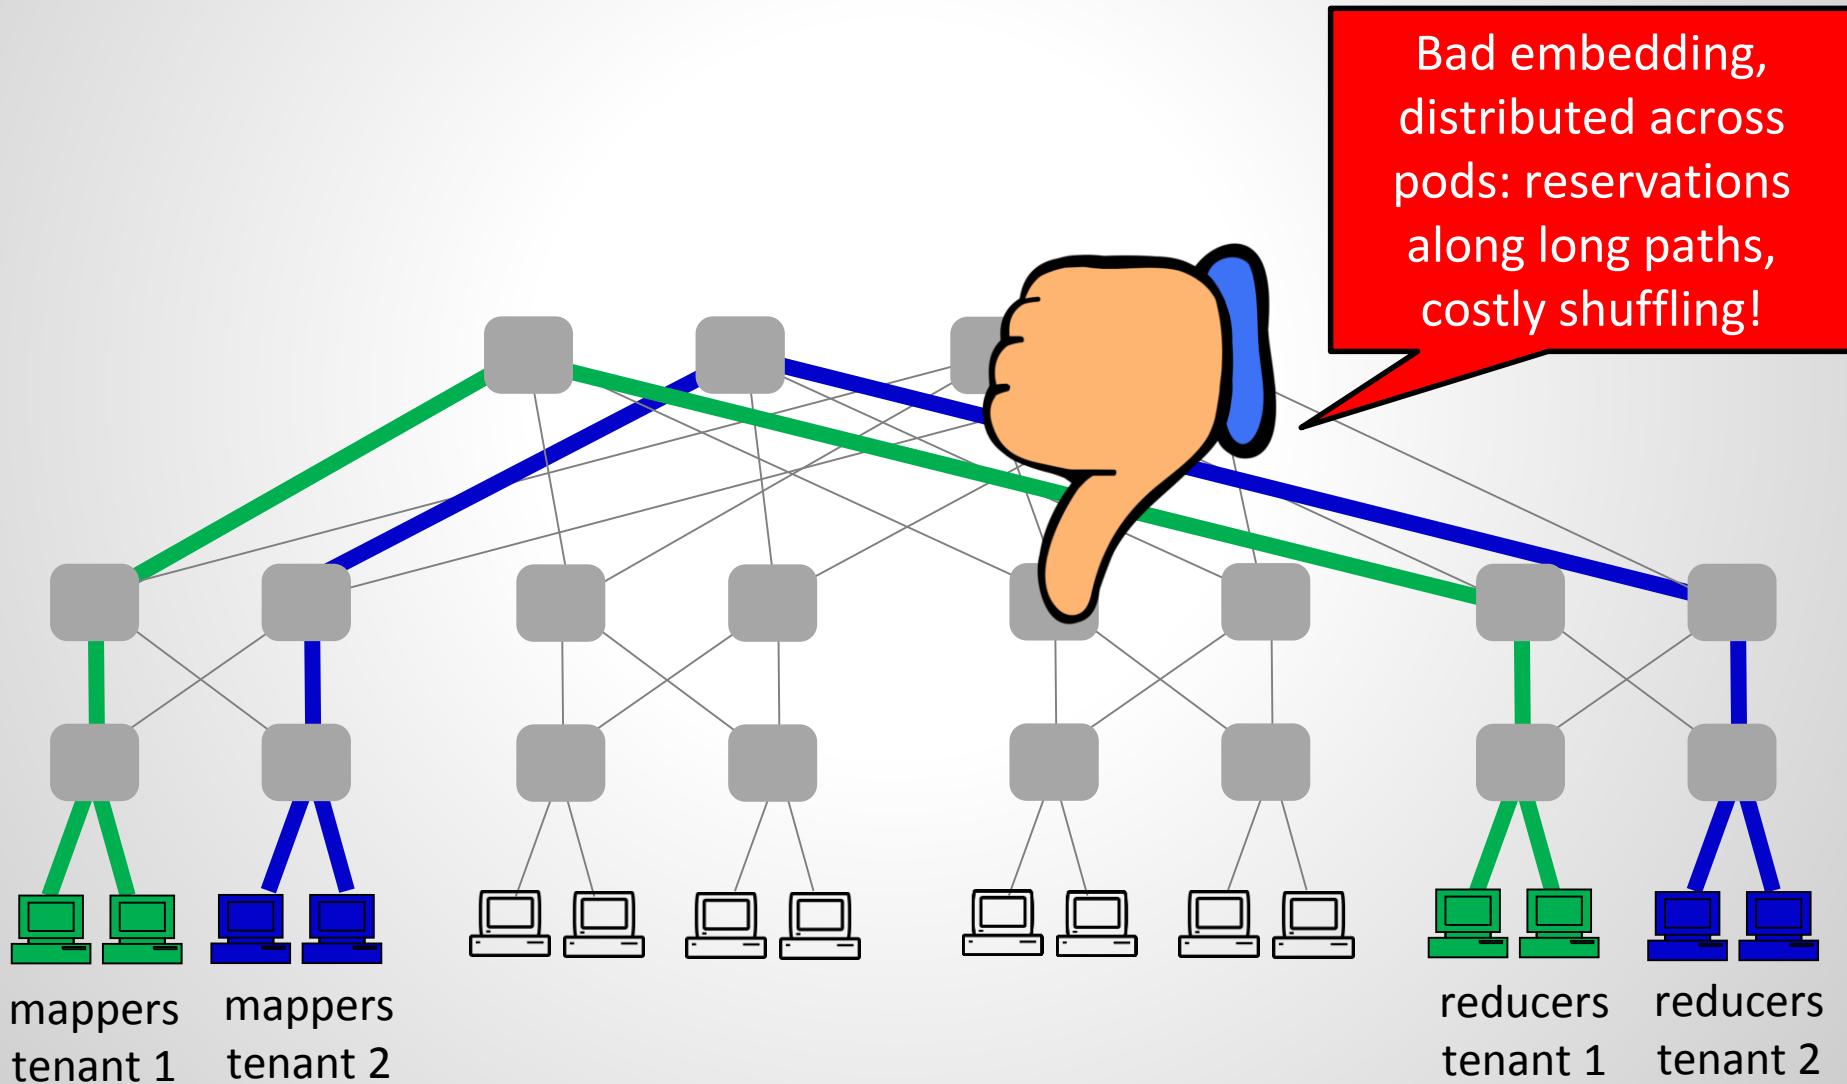
Consider a simple data center hosting two tenants: **green** and **blue**



First Algorithmic Challenge:

Keep the traffic local!

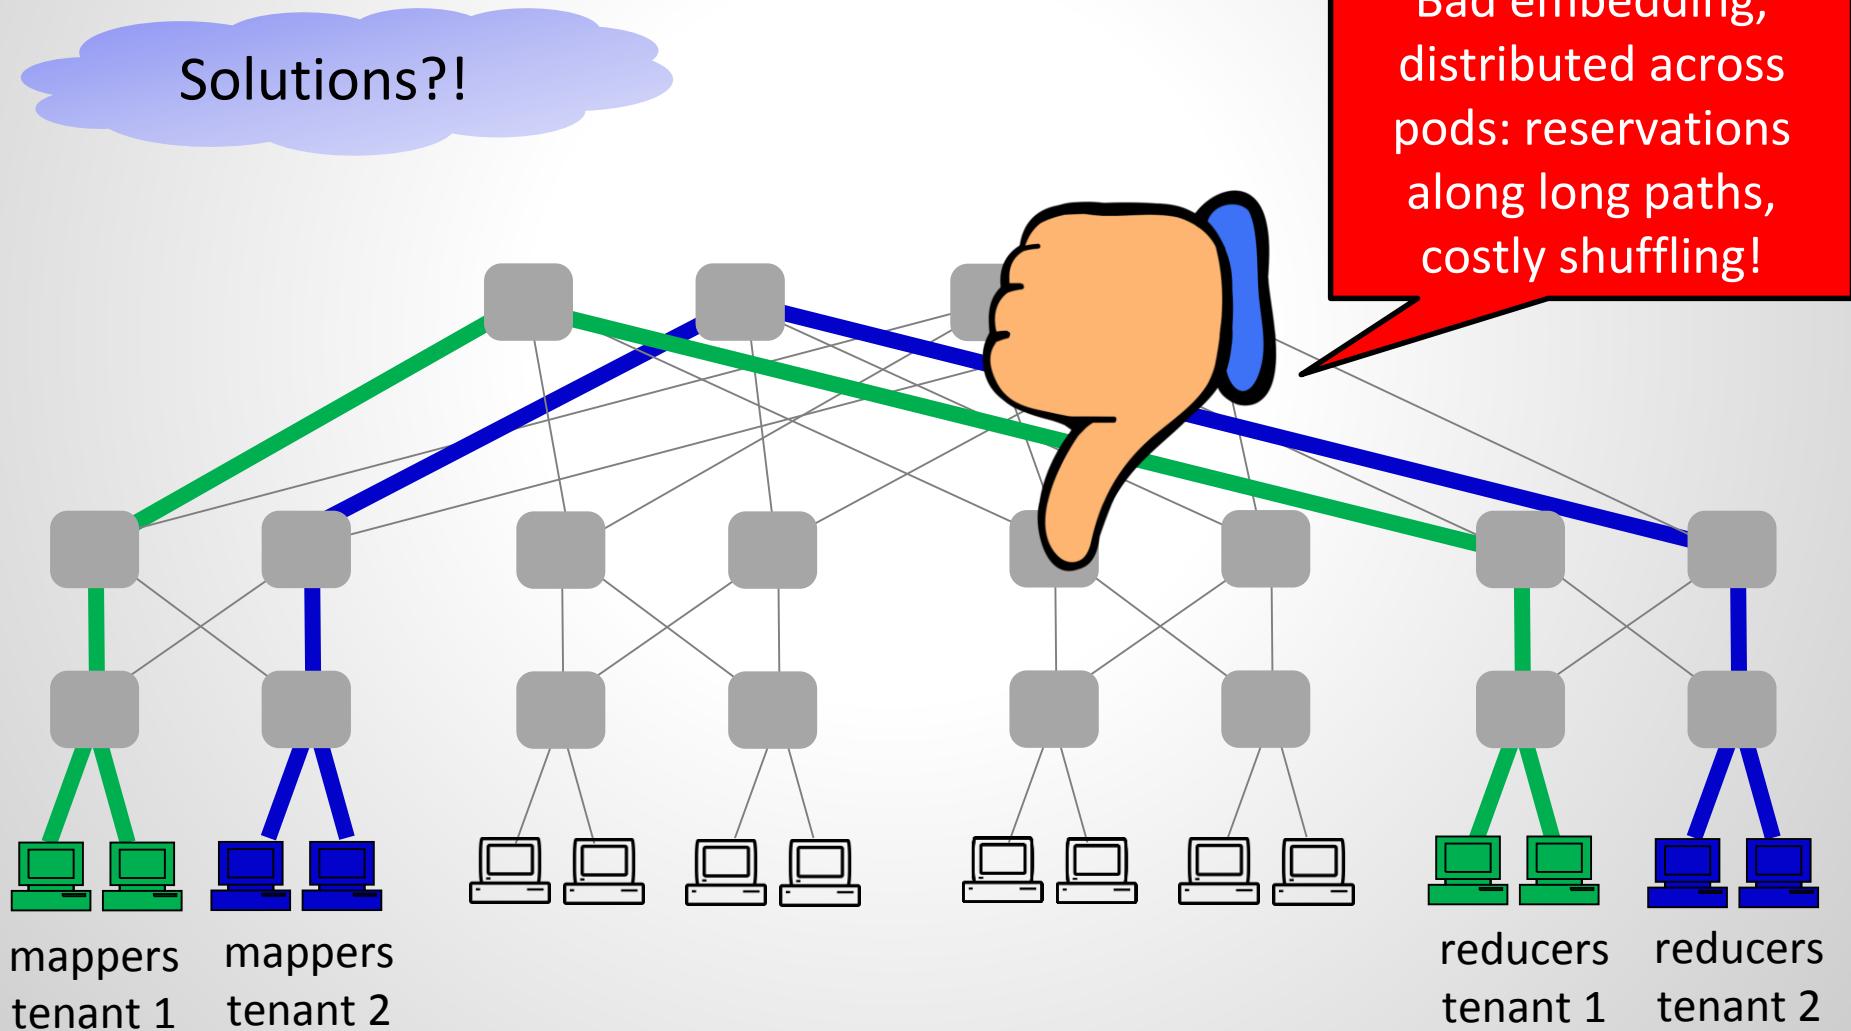
Consider a simple data center hosting two tenants: **green** and **blue**



First Algorithmic Challenge:

Keep the traffic local!

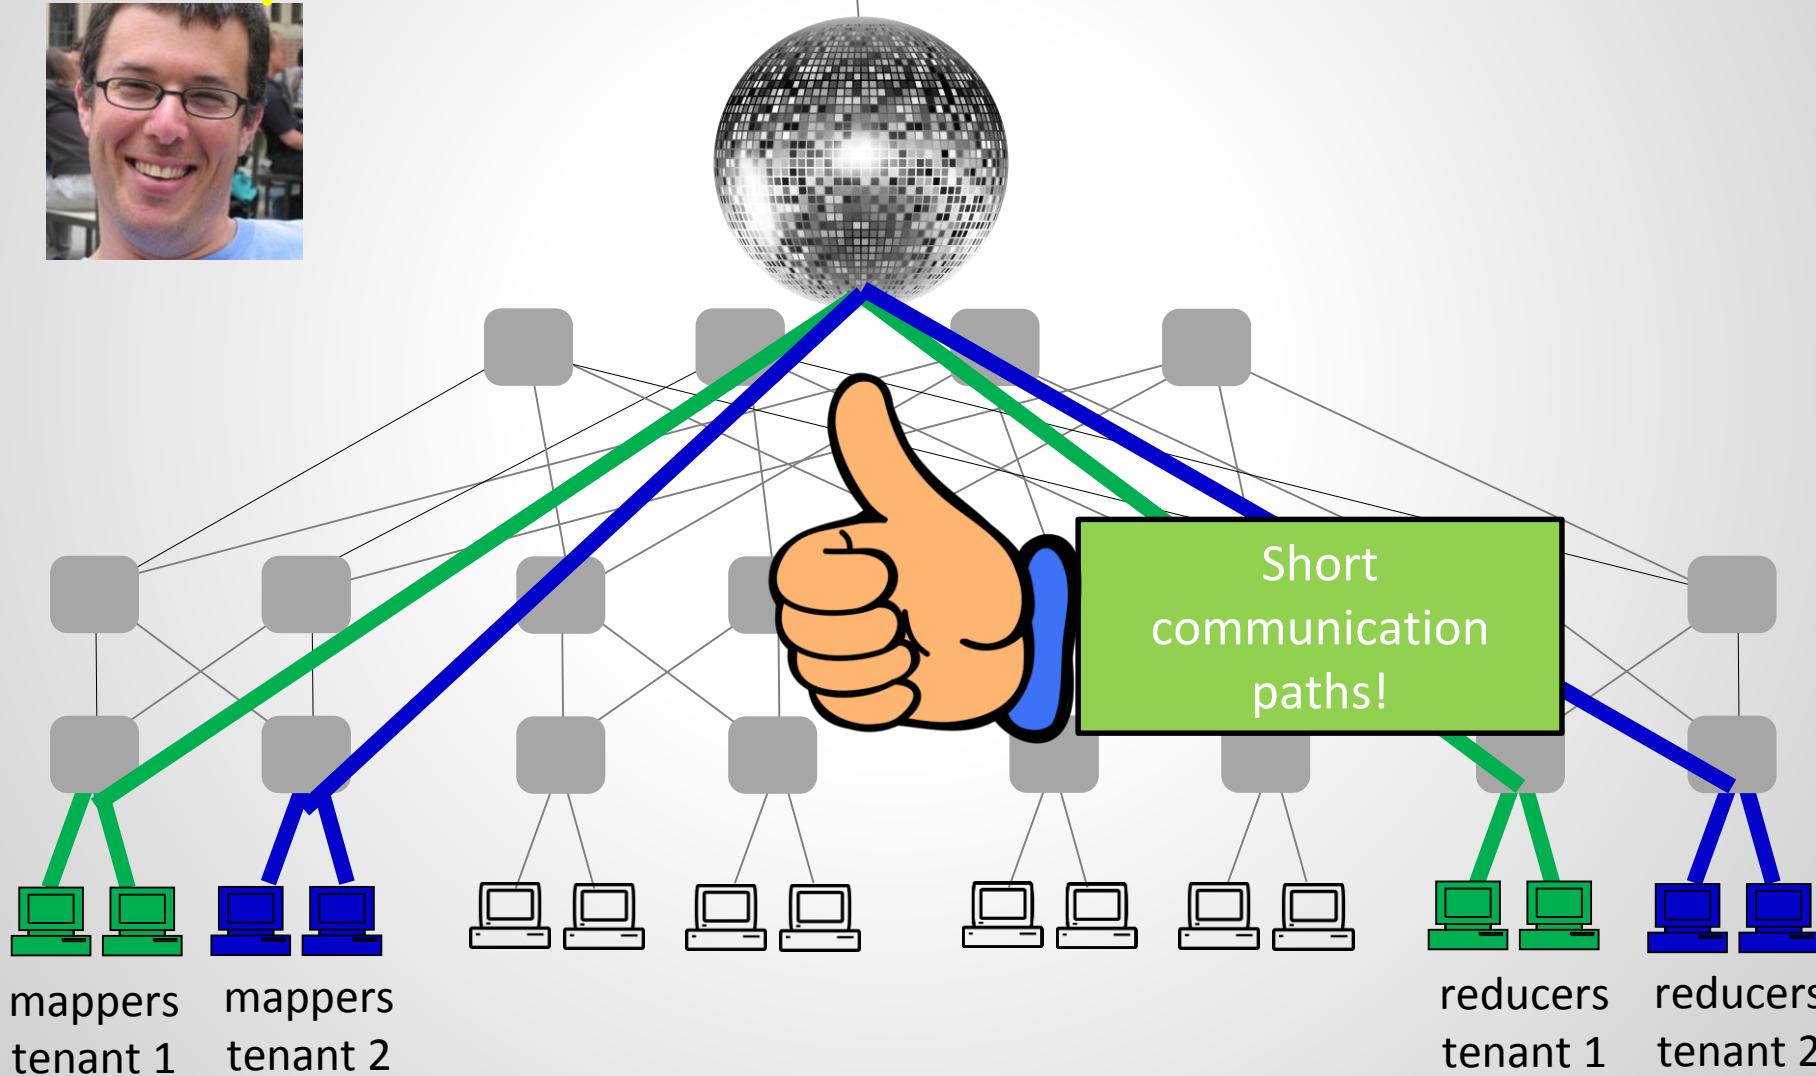
Consider a simple data center hosting two tenants: **green** and **blue**



Solution 1: Adjust the Network

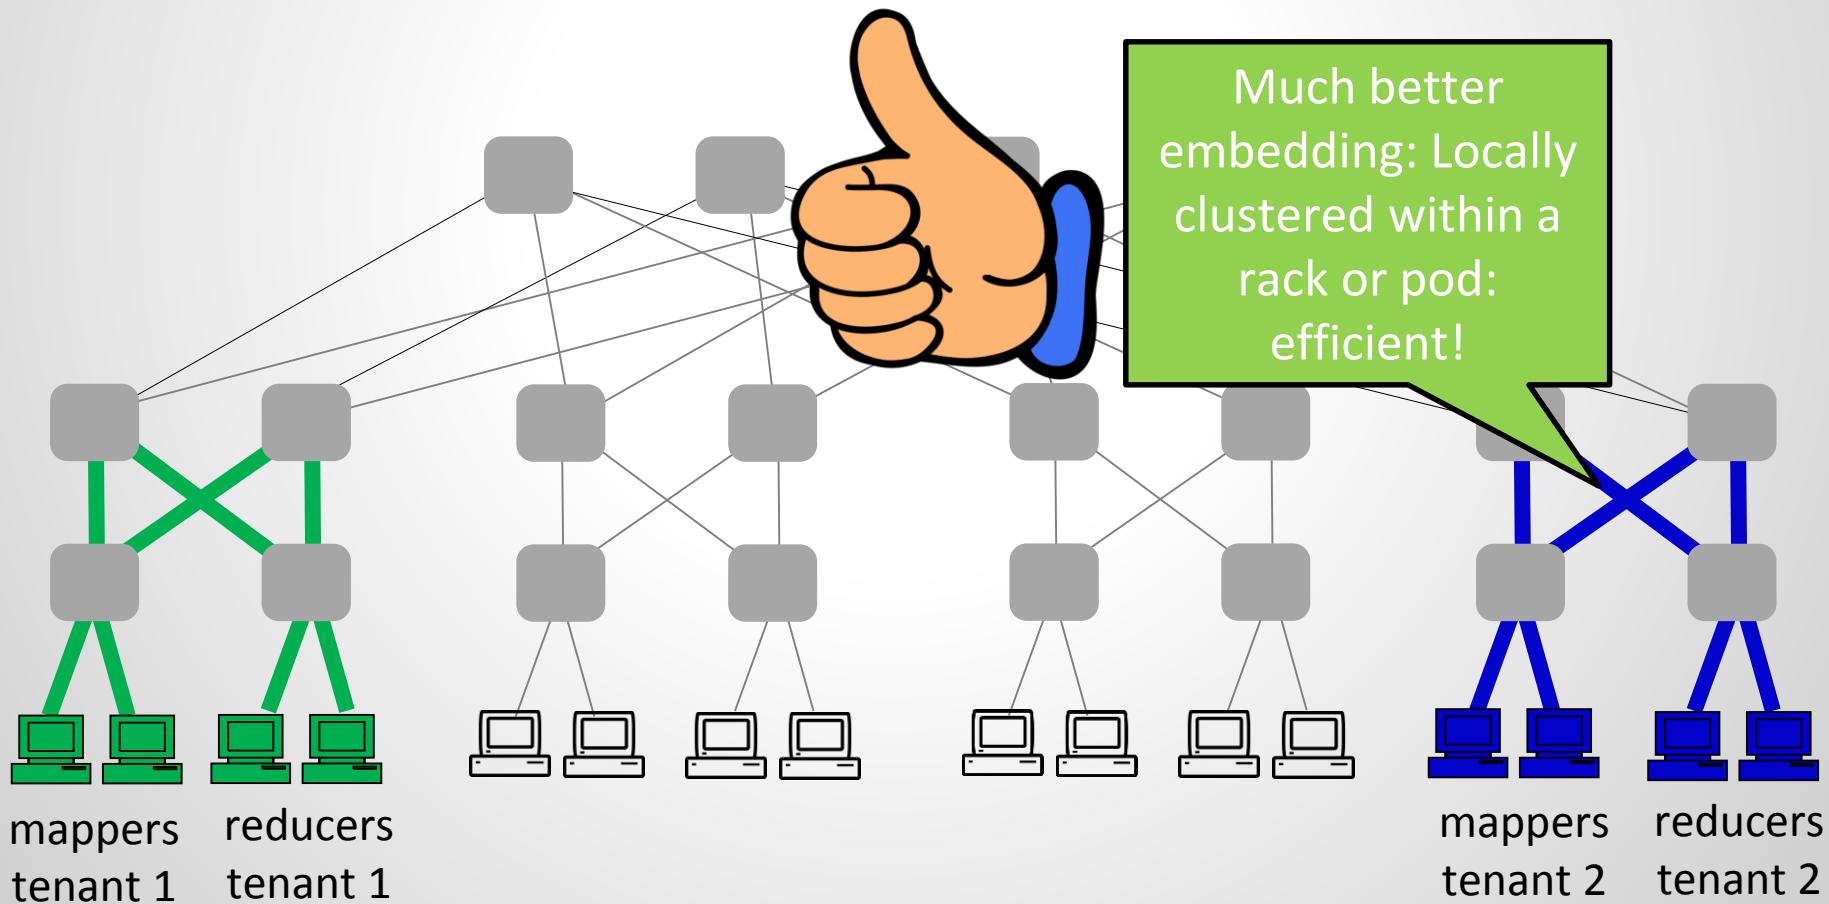
Adjust the network!

Consider a simple data center hosting two tenants: **green** and **blue**



Solution 2: Adjust Embedding

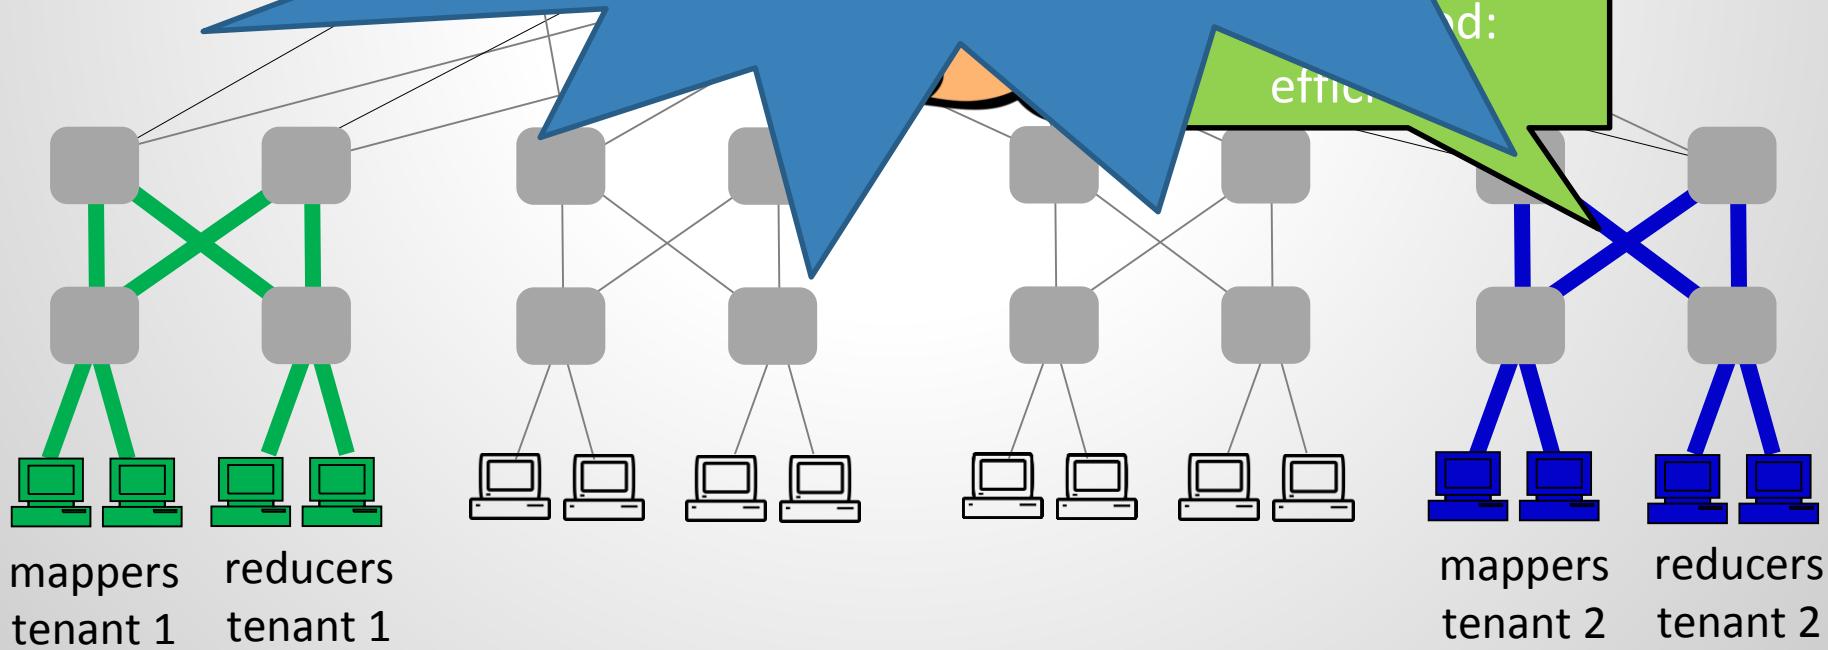
Consider a simple data center hosting two tenants: **green** and **blue**



Solution 2: Adjust Embedding

Consider a simple data center hosting two tenants: green and blue

How to compute a minimal embedding? Known as the Virtual Network Embedding Problem.



Overview

PART I: Static Embeddings

PART II: Reconfiguring Embeddings

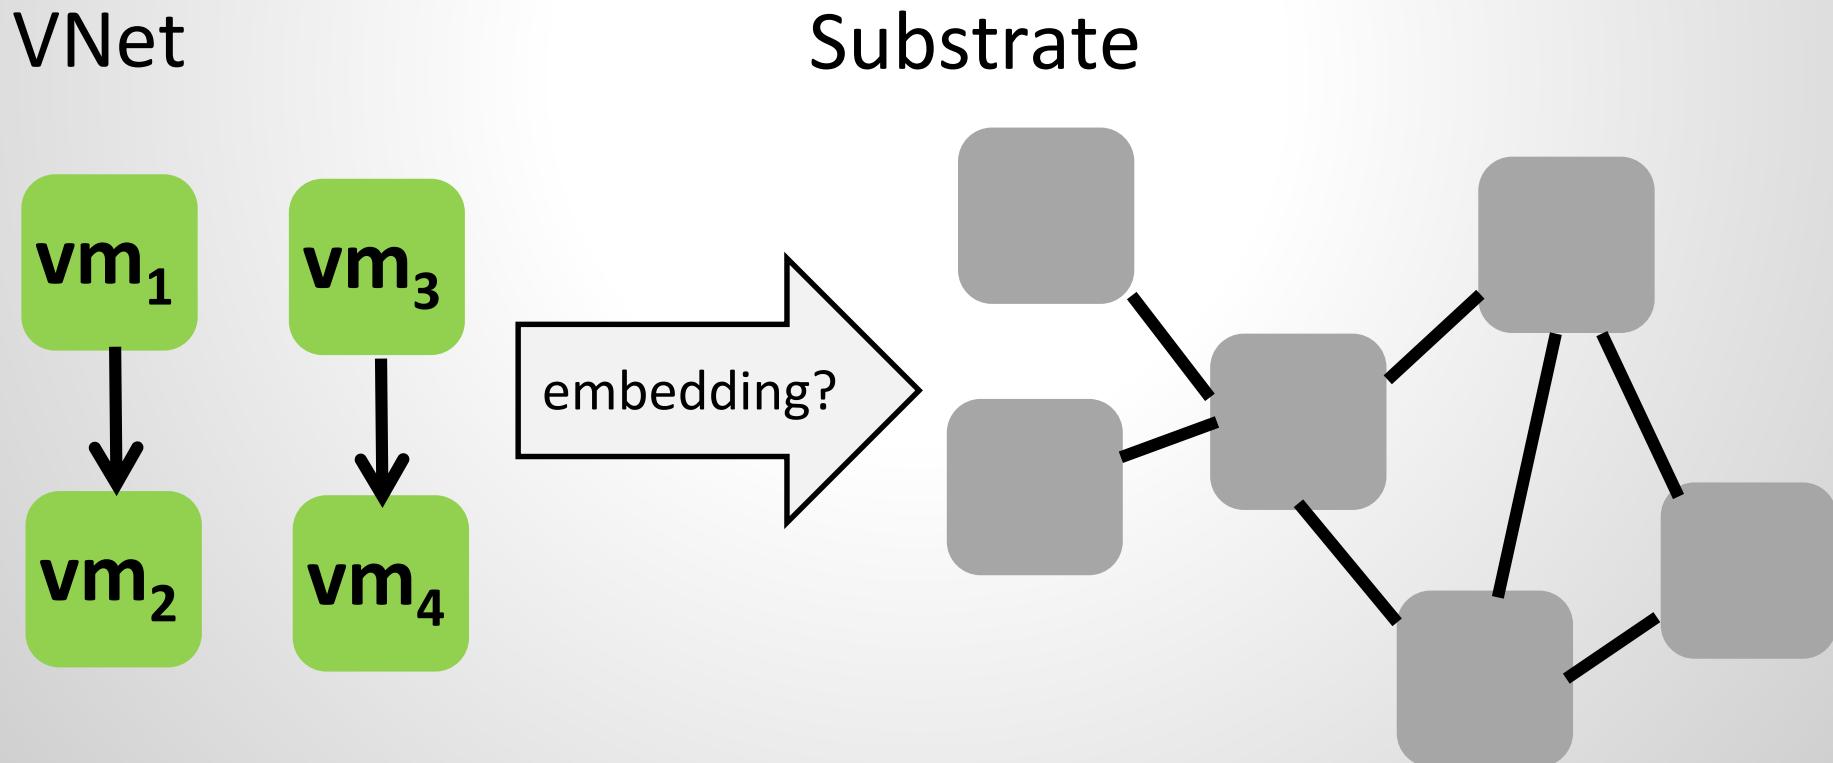
PART III: A request comes seldom alone!

PART I:

Static Embeddings

The Virtual Network Embedding Problem

- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)



The Virtual Network Embedding Problem

- 2 dimensions of flexibility:

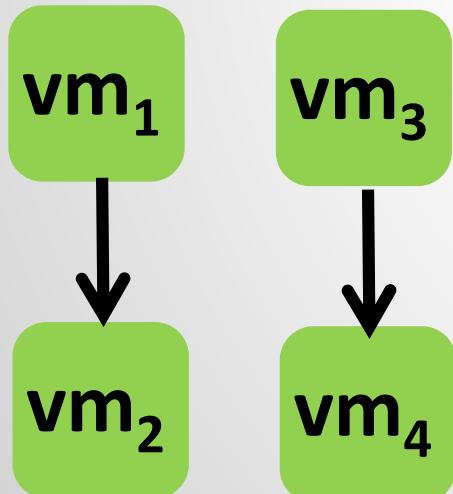
- Mapping of virtual nodes (to physical nodes)

- Mapping of virtual links (to paths)

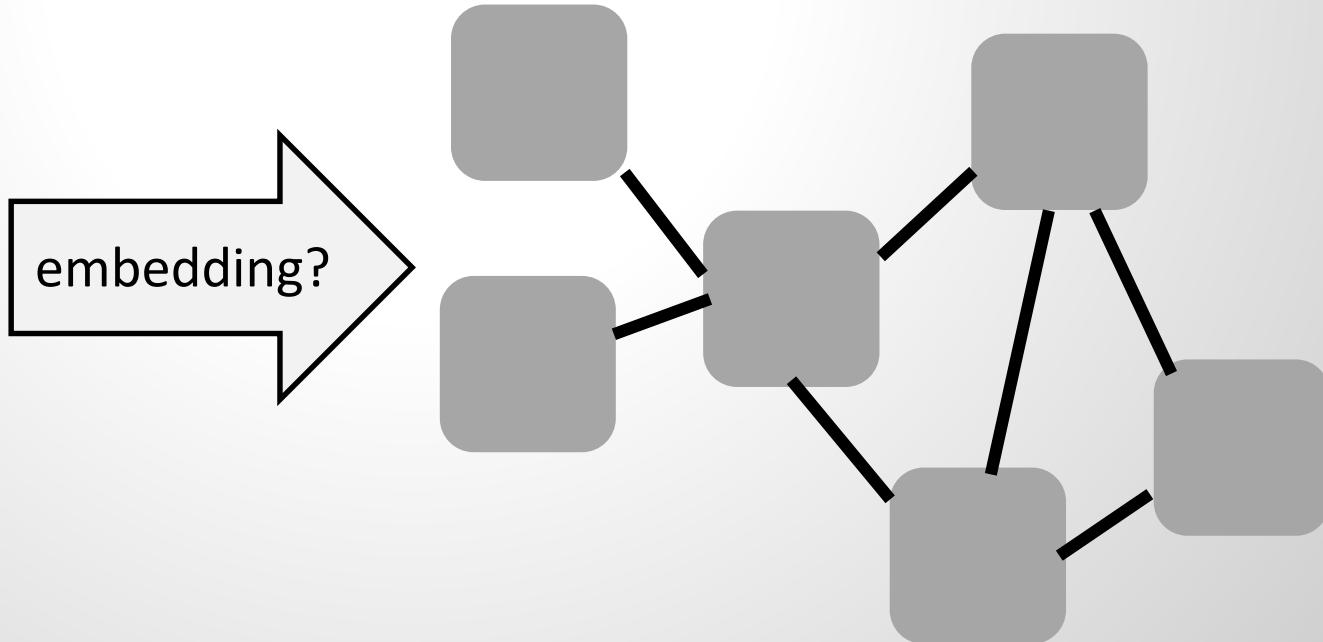
aka “guest graph”

aka “host graph”

VNet



Substrate



The Virtual Network Embedding Problem

- 2 dimensions of flexibility:

- Mapping of virtual nodes (to physical nodes)

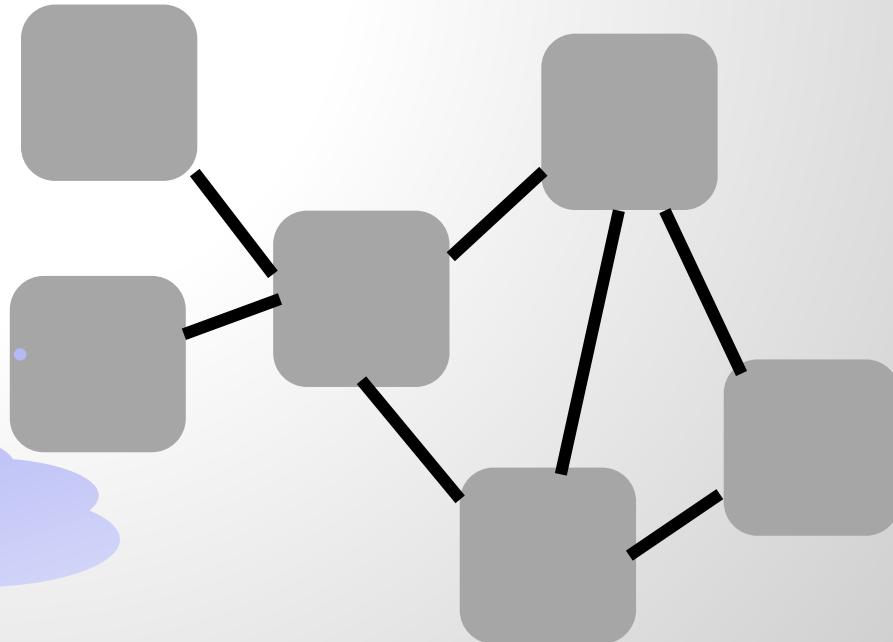
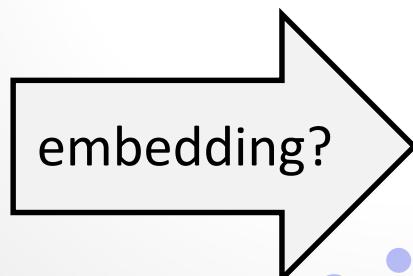
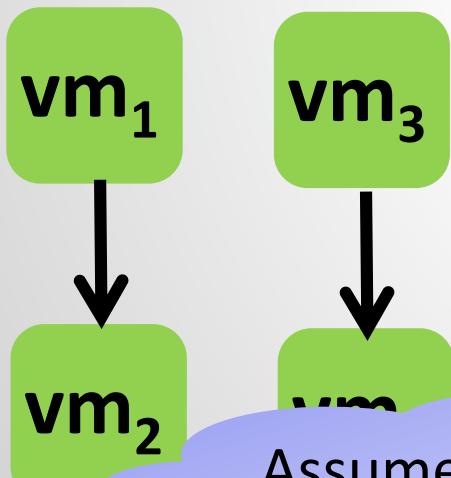
- Mapping of virtual links (to paths)

aka “guest graph”

aka “host graph”

VNet

Substrate

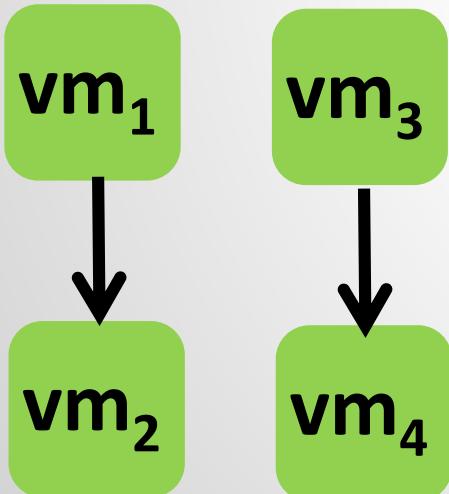


Assume unit demand
and capacity!

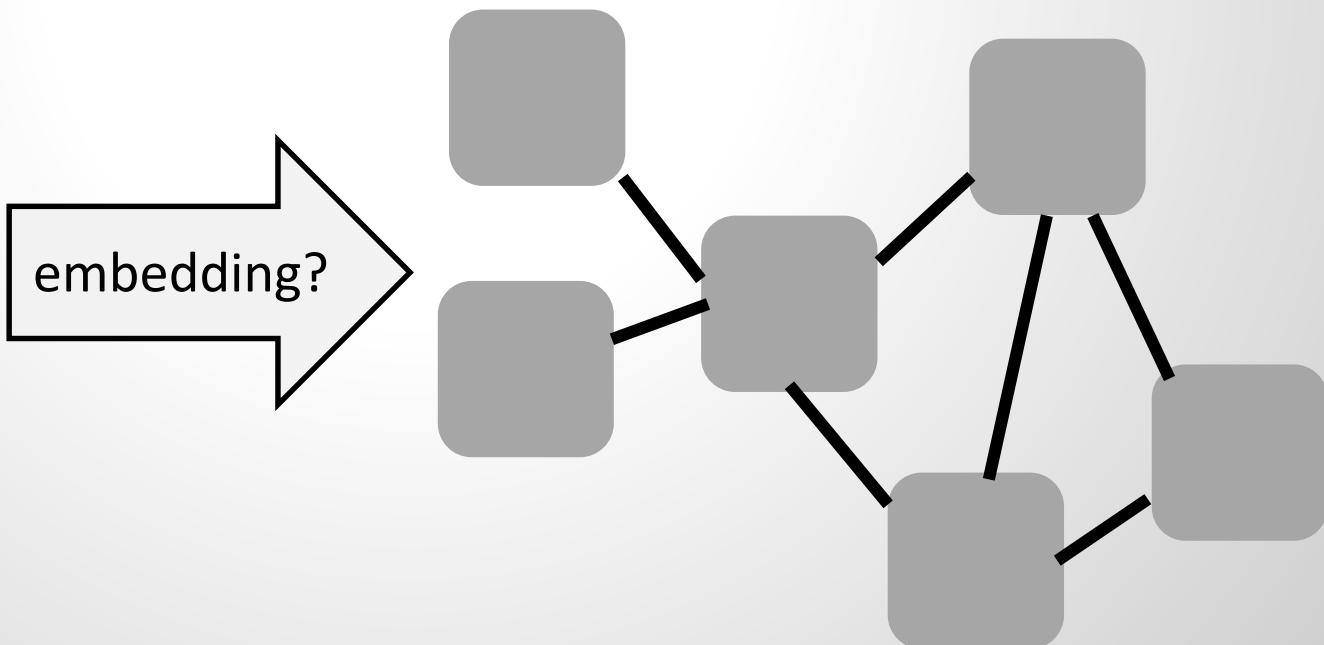
The Virtual Network Embedding Problem

- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)
- Let's start simple: assume node **mappings** are given

VNet



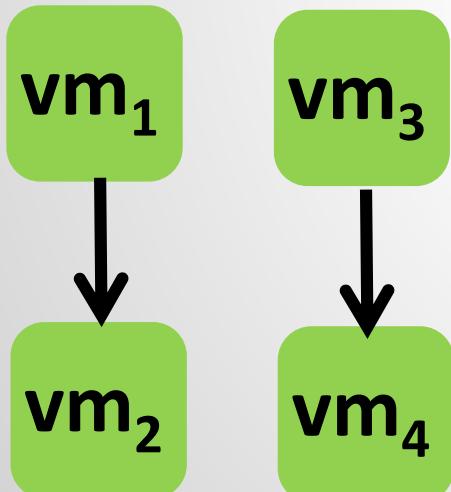
Substrate



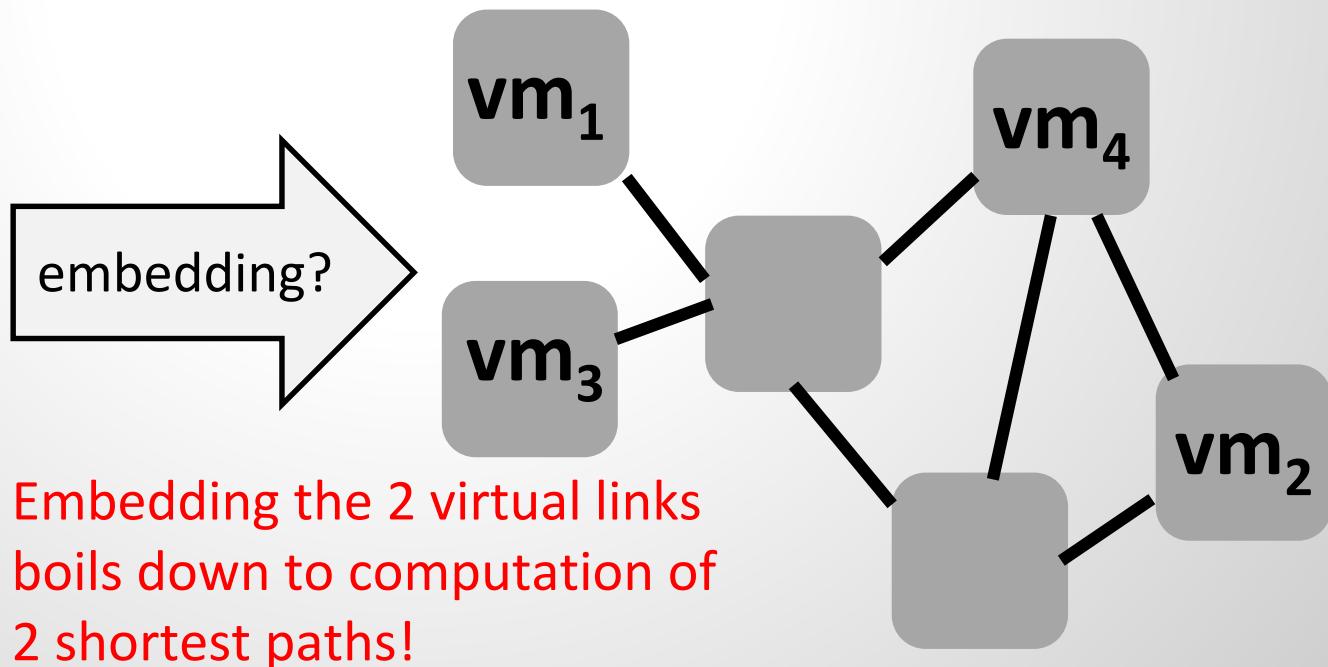
The Virtual Network Embedding Problem

- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)
- Let's start simple: assume node **mappings** are given

VNet



Substrate

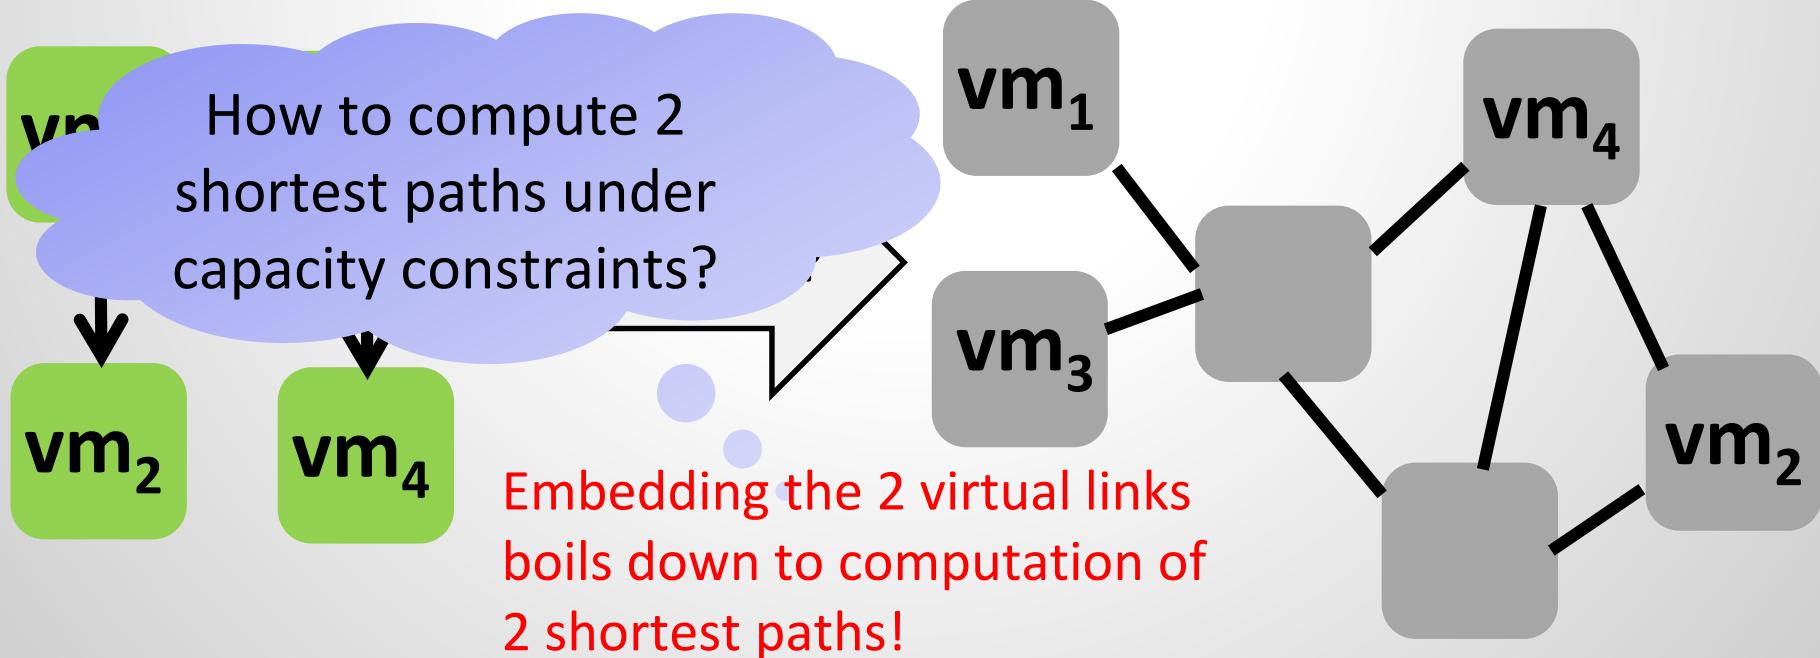


The Virtual Network Embedding Problem

- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)
- Let's start simple: assume node **mappings** are given

VNet

Substrate



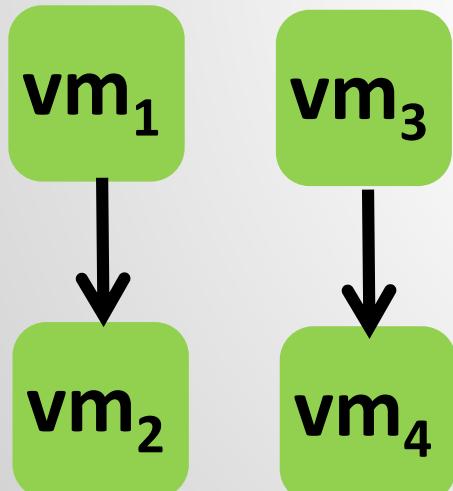
The Virtual Network Embedding Problem

- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)

- Let's start

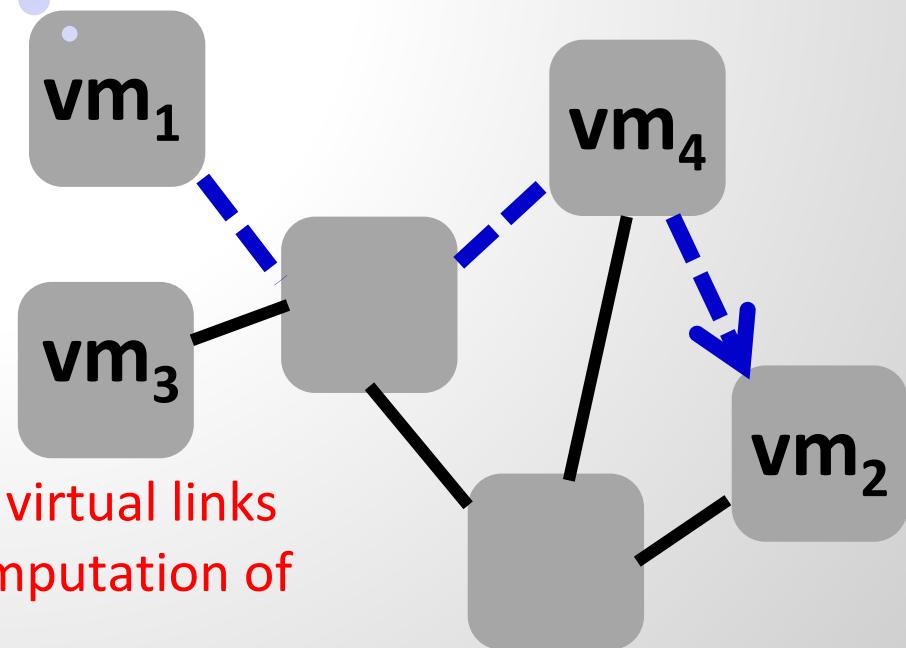
Let's try greedy!
First $vm_1 - vm_2$.

VNet



embedding?

substrate



Embedding the 2 virtual links
boils down to computation of
2 shortest paths!

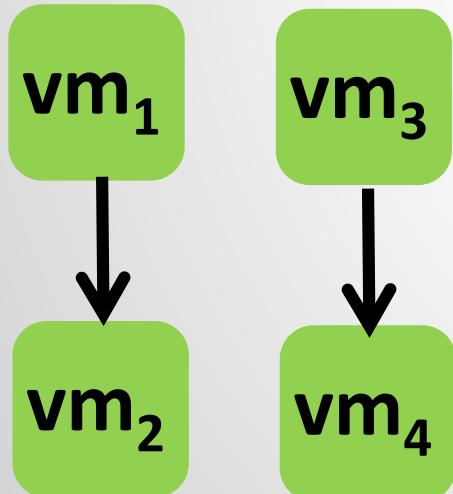
The Virtual Network Embedding Problem

- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)

- Let's start

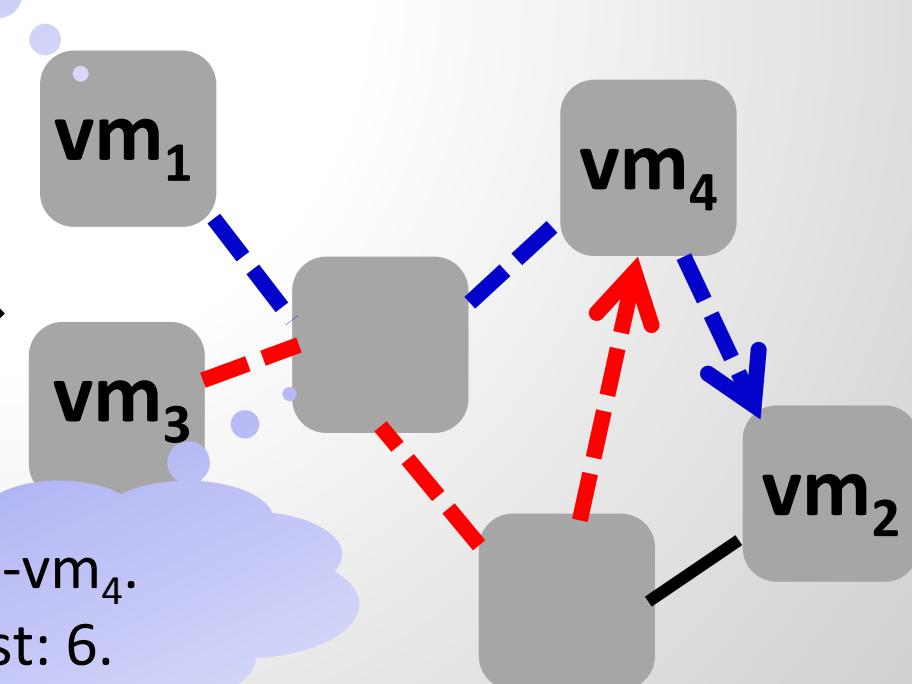
Let's try greedy!
First vm_1 - vm_2 .

VNet



embedding?

substrate

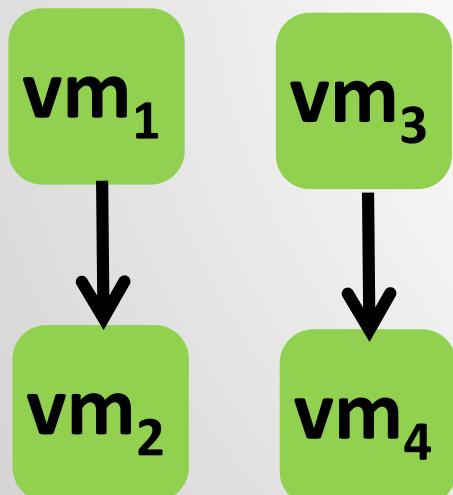


Then vm_3 - vm_4 .
Total cost: 6.

The Virtual Network Embedding Problem

- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)
- Let's start simple: *Virtual network components are given*

VNet

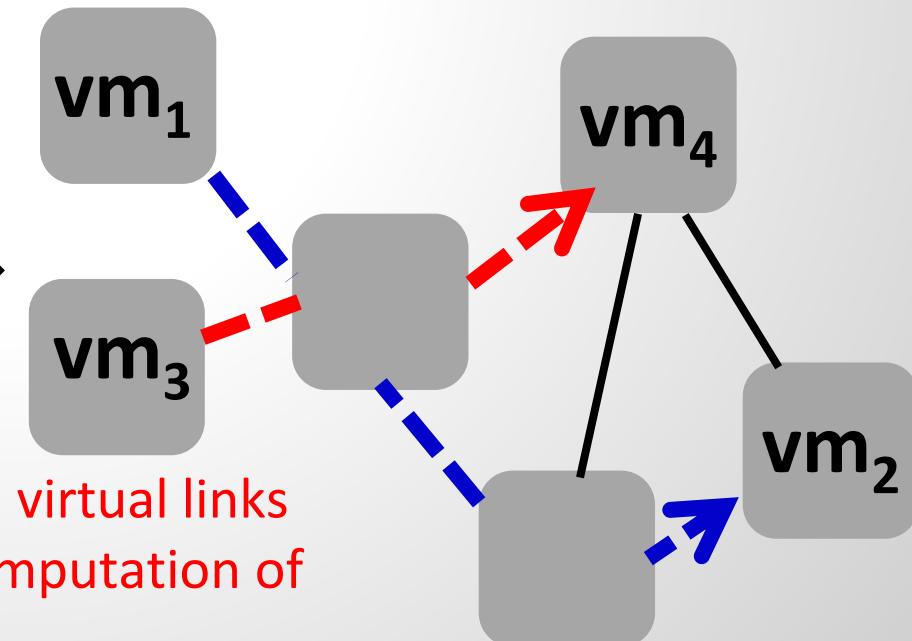


embedding?

Embedding the 2 virtual links
boils down to computation of
2 shortest paths!

A better solution:
cost 5!

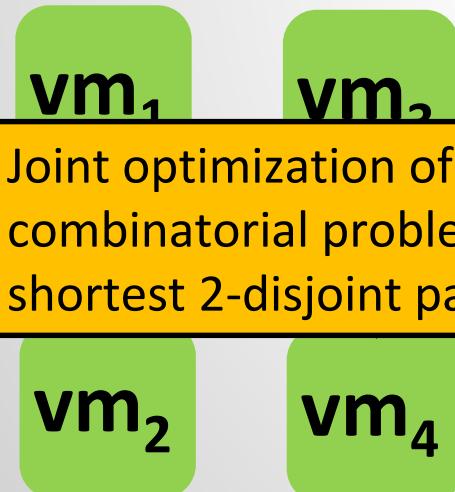
Associate



The Virtual Network Embedding Problem

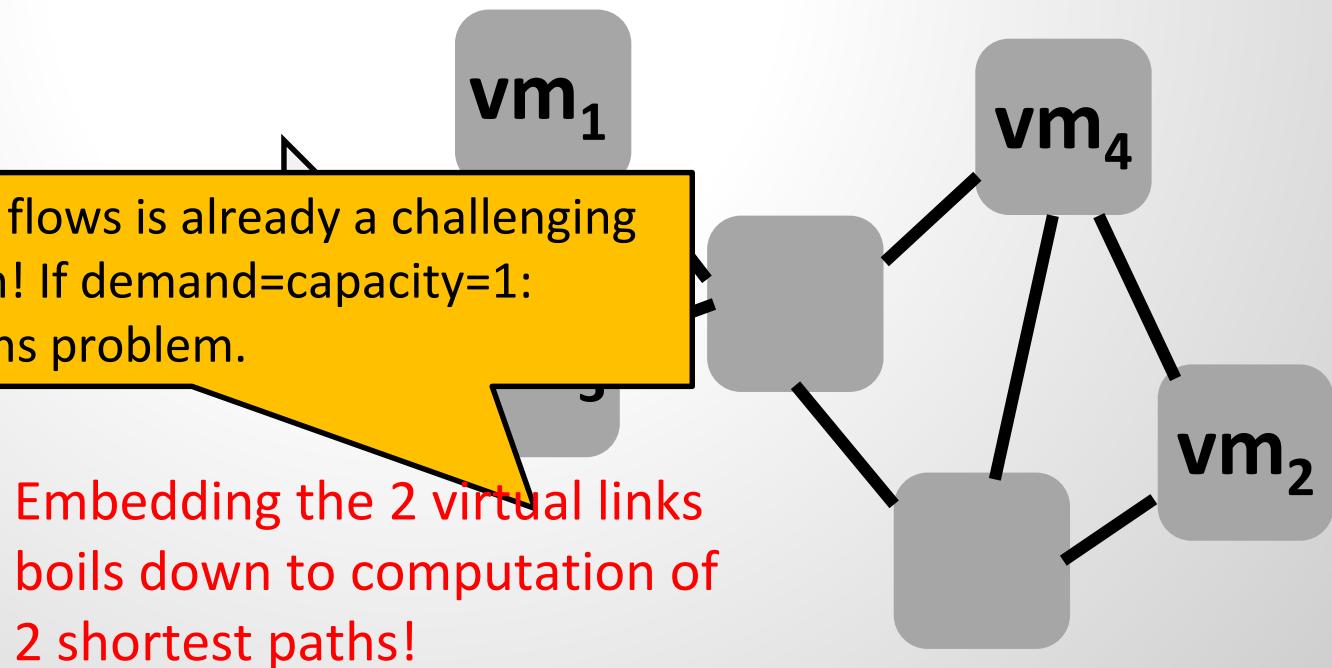
- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)
- Let's start simple: assume node **mappings** are given

VNet



Joint optimization of 2 flows is already a challenging combinatorial problem! If demand=capacity=1: shortest 2-disjoint paths problem.

Substrate



The Virtual Network Embedding Problem

- 2 dimensions of flexibility:
 - Mapping of virtual nodes (to physical nodes)
 - Mapping of virtual links (to paths)
- Let's start simple: assume node **mappings** are given

VNet

Hasn't this problem been
solved a generation ago?!

Joint optimization of 2 flows is already a challenging
combinatorial problem! If demand=capacity=1:
shortest 2-disjoint paths problem.

Embedding the 2 virtual links
boils down to computation of
2 shortest paths!

vm₂

vm₄

vm₁

vm₂

vm₁

vm₄

vm₂

vm₃

vm₃

vm₄

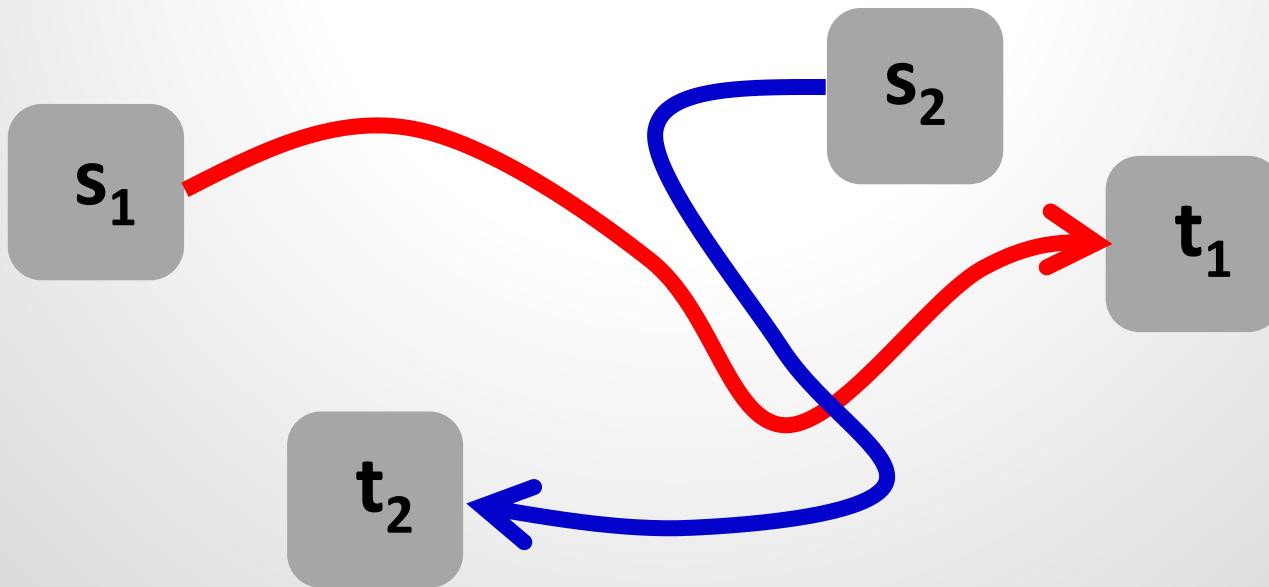
vm₃

vm₂

vm₁

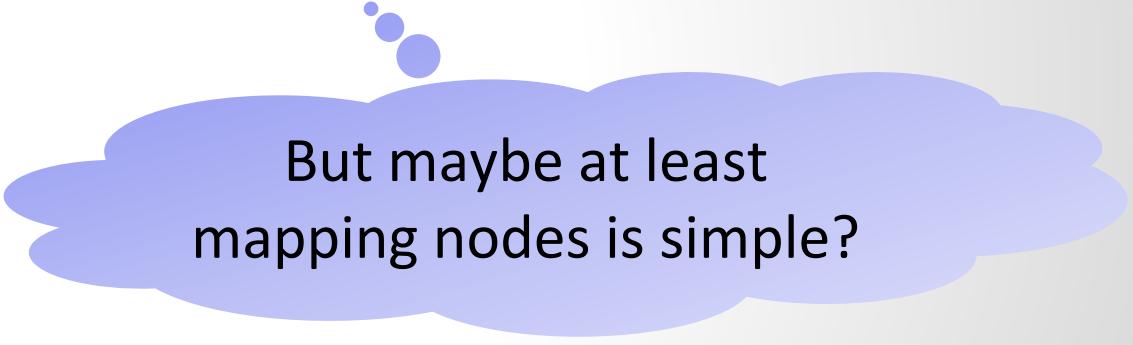
Mapping virtual links: Already hard!

- ❑ Essentially a 2-disjoint shortest paths problem: a deep combinatorial problem
 - ❑ NP-hard on **directed** graphs
 - ❑ For **undirected** graphs:
 - ❑ Feasibility more or less understood: **Robertson&Seymour**
 - ❑ Shortest paths: recent **breakthrough**, first polytime **randomized** algorithm (still slow: a theoretical result)
 - ❑ We are still looking for polytime deterministic algorithms!



Therefore: Mapping Virtual Links is Challenging

Bad news: The Virtual Network Embedding Problem is hard
even if endpoints are already mapped and given.

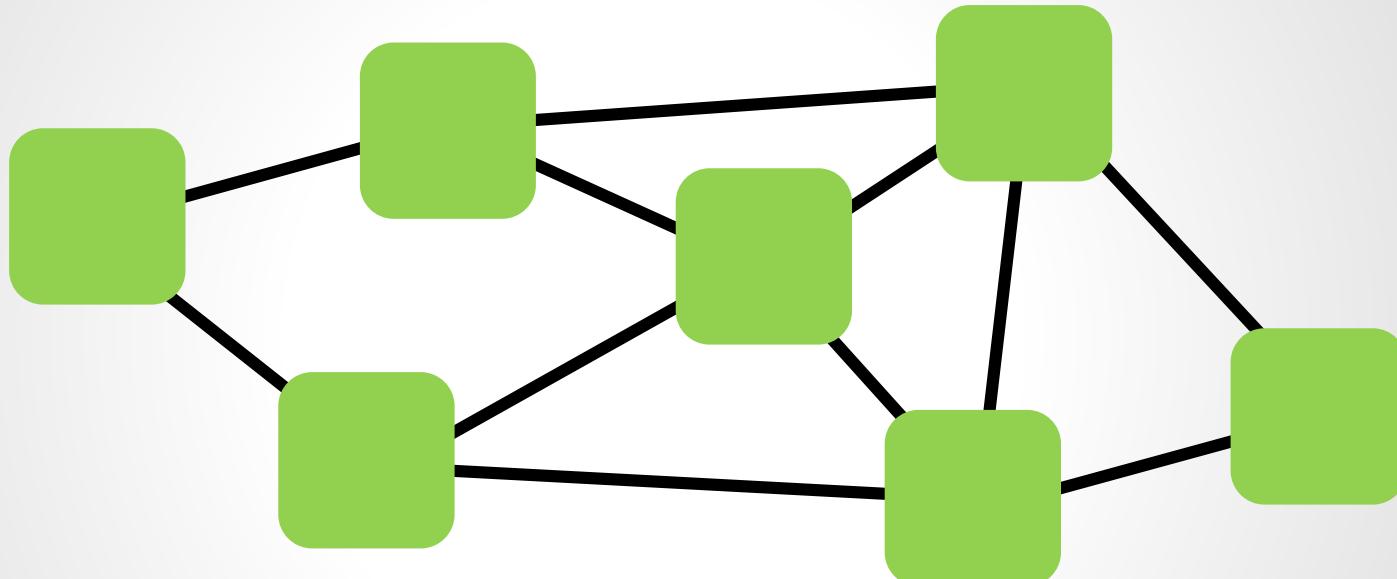


But maybe at least
mapping nodes is simple?

Mapping Virtual Nodes

- Let's start simple again: assume paths are trivial, e.g., the physical network (host graph) **is a line**

Guest



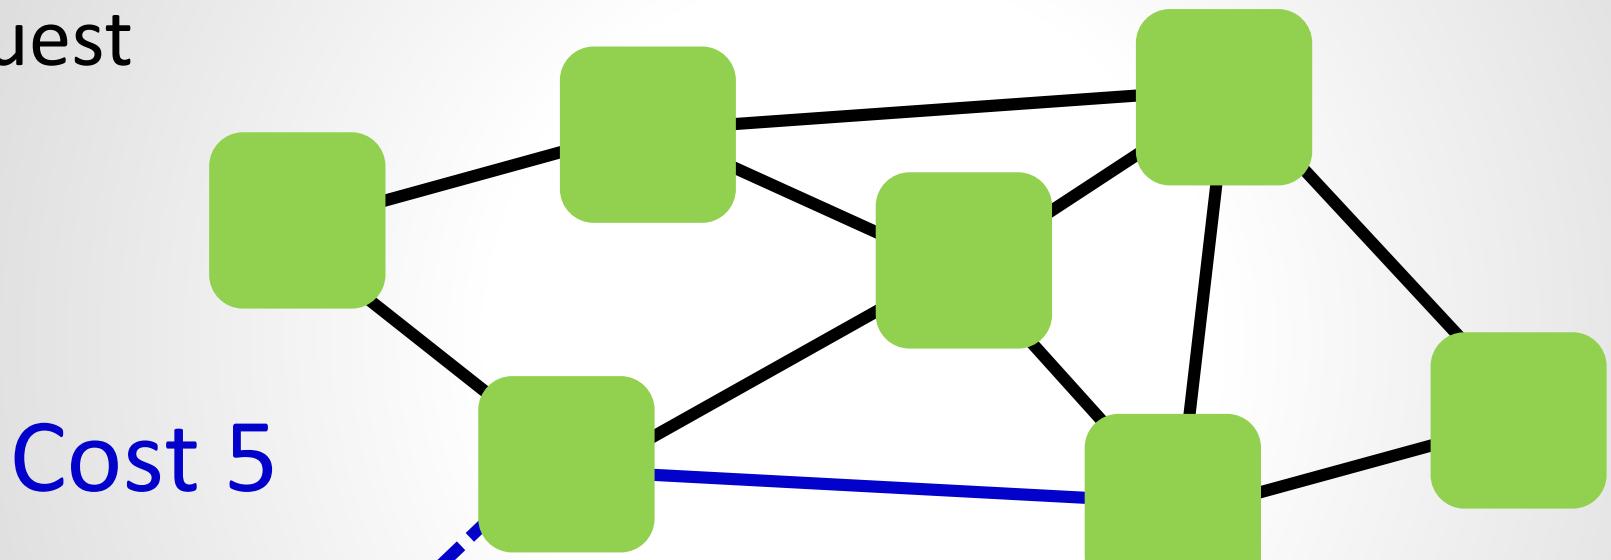
Host



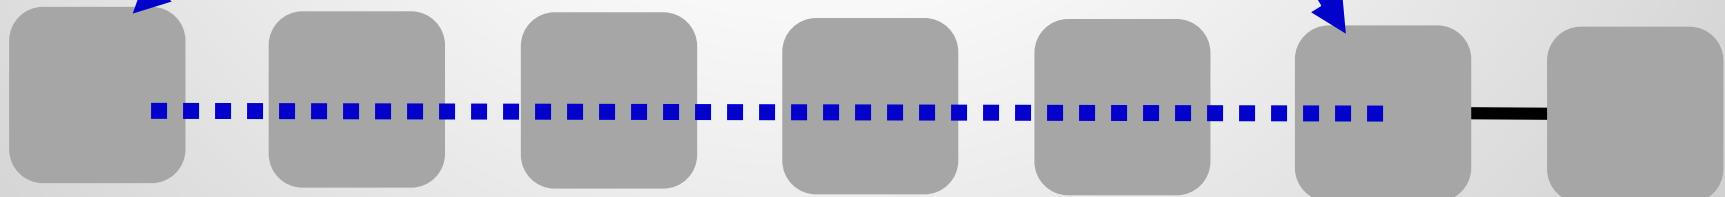
Mapping Virtual Nodes

- Let's start simple again: assume **paths are trivial**, e.g., the physical network (host graph) **is a line**

Guest



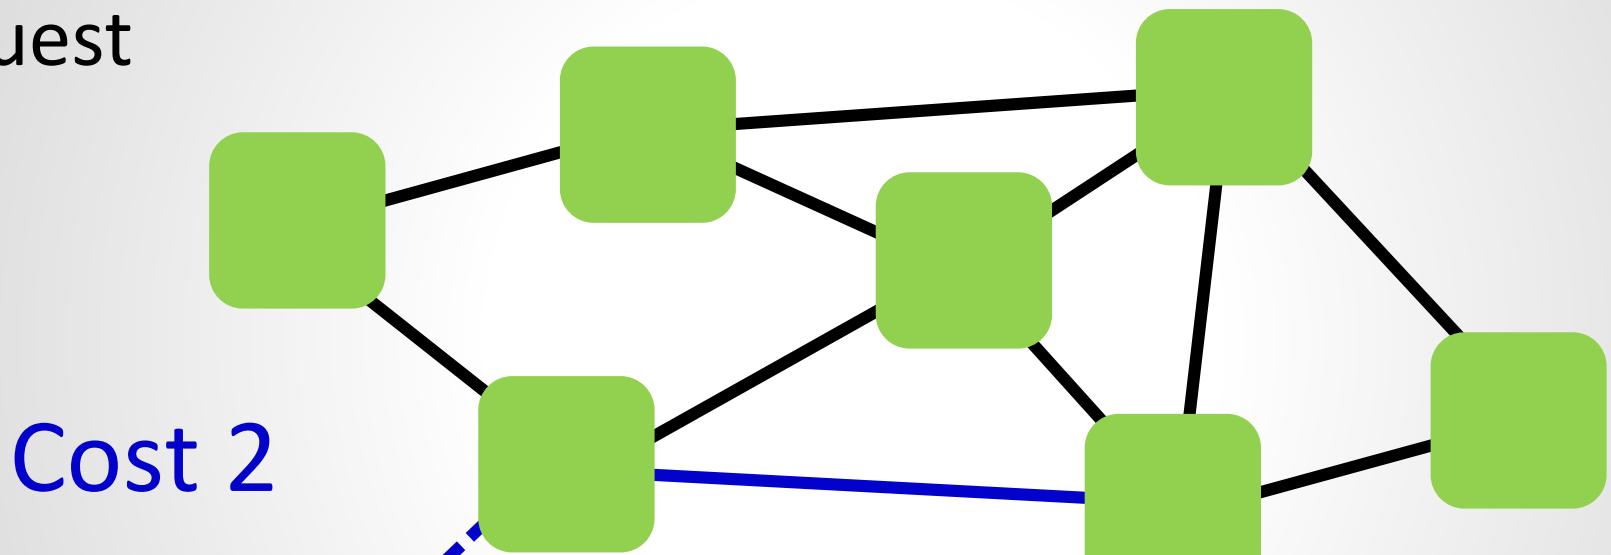
Host



Mapping Virtual Nodes

- Let's start simple again: assume **paths are trivial**, e.g., the physical network (host graph) is a line

Guest

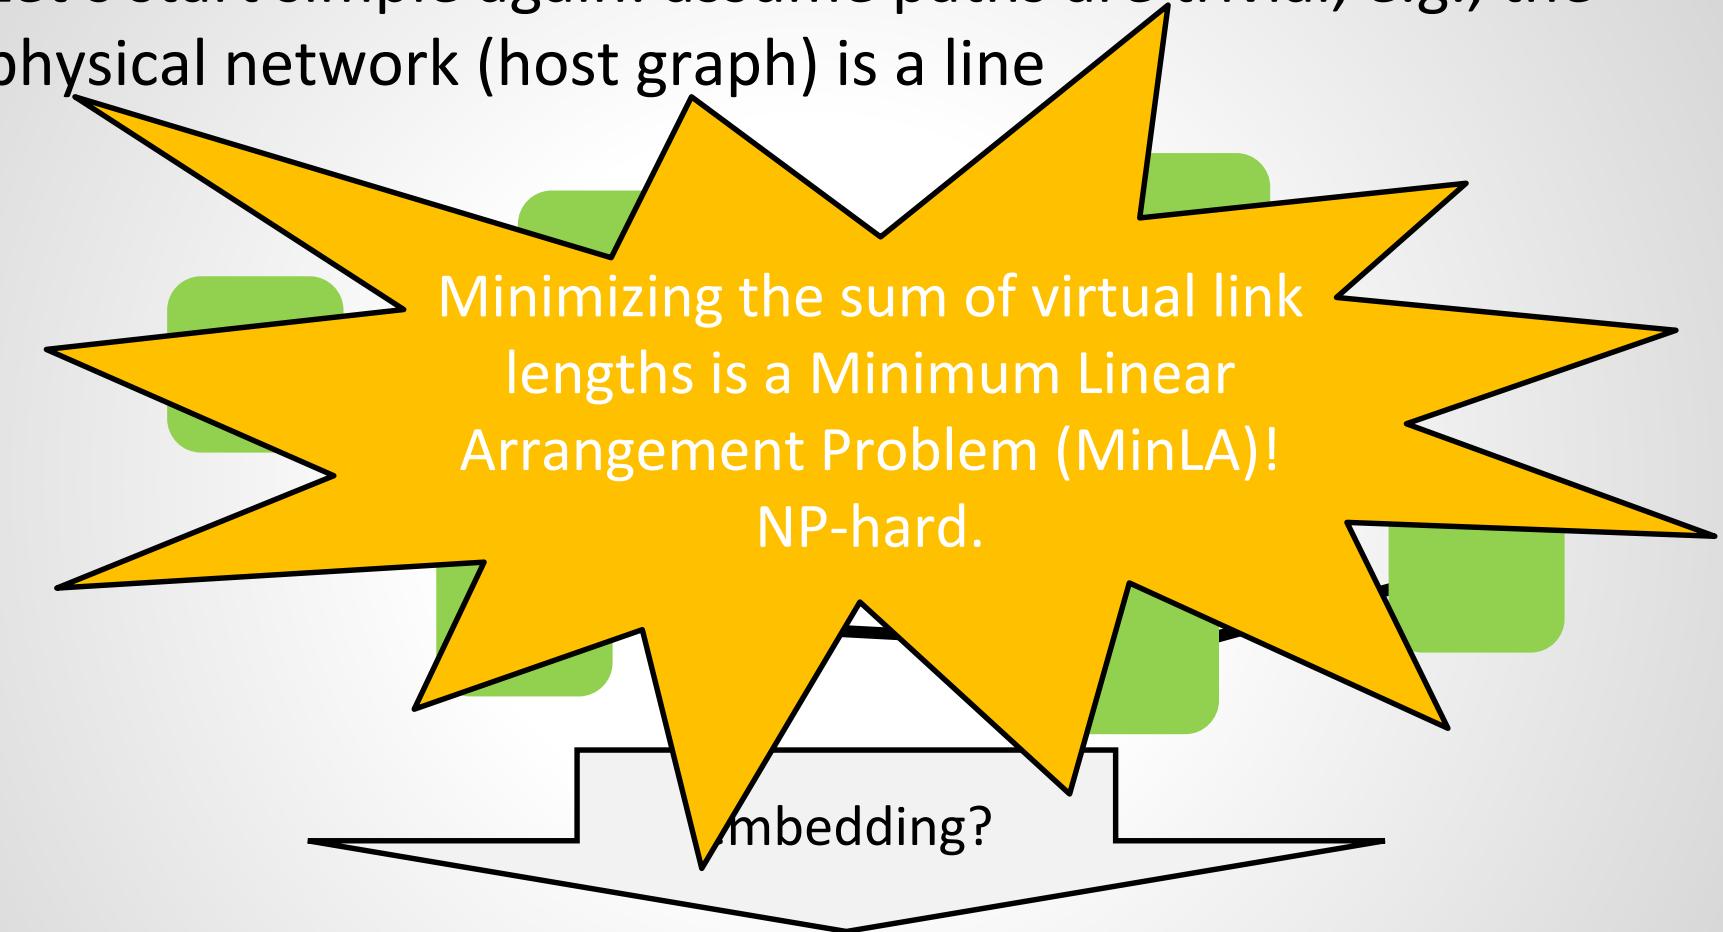


Host



Mapping Virtual Nodes

- Let's start simple again: assume paths are trivial, e.g., the physical network (host graph) is a line



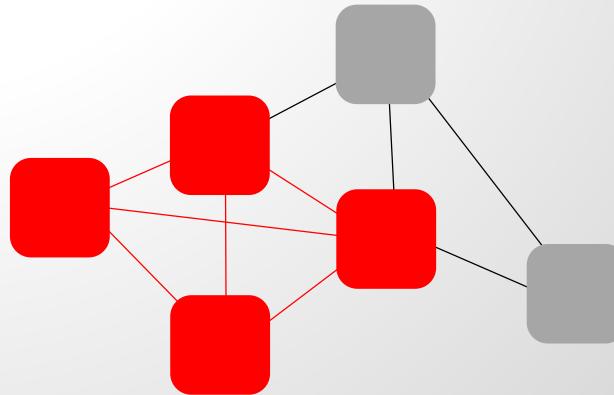
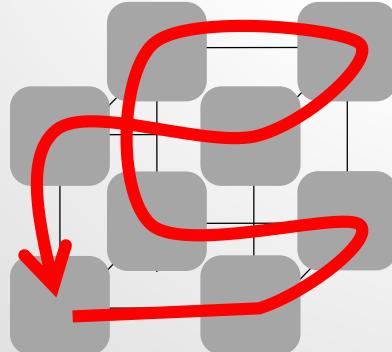
Therefore: VNEP is Hard “in Both Dimensions”!

- We have seen examples that:
 - mapping virtual links is hard (even if nodes are given)
 - mapping virtual nodes is hard (even if links are trivial)
- Remark: the VNEP can also be seen as a generalization of the **Subgraph Isomorphism Problem (SIP)**

Known? Why is SIP NP-hard?

Therefore: VNEP is Hard “in Both Dimensions”!

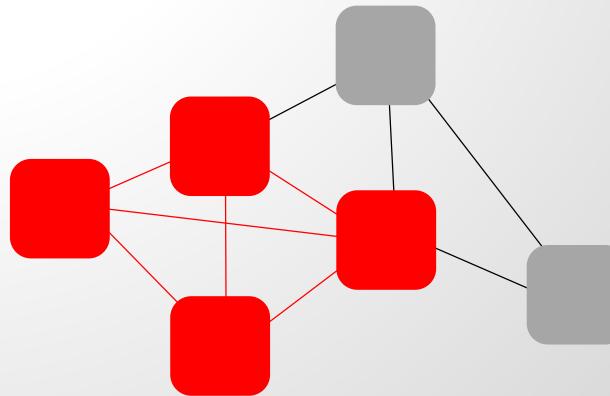
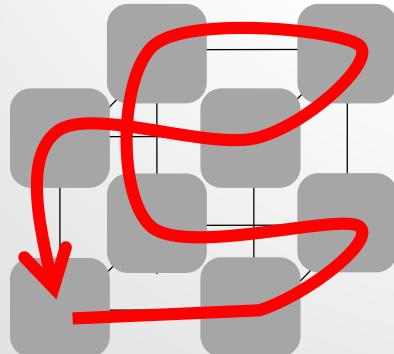
- We have seen examples that:
 - mapping virtual links is hard (even if nodes are given)
 - mapping virtual nodes is hard (even if links are trivial)
- Remark: the VNEP can also be seen as a generalization of the **Subgraph Isomorphism Problem (SIP)**
 - The SIP problem: Given two graphs G, H , determine whether G contains a subgraph that is **isomorphic to H** ?
 - NP-hard: “does G contain an **n -node cycle**?” is a **Hamilton cycle** problem (each node visited exactly once), a solution to “does G contain a **k -clique**?” solves **maximum clique** problem, etc.



Therefore: VNEP is Hard “in Both Dimensions”!

- We have seen examples that:
 - mapping virtual links is hard (even if nodes are given)
 - mapping virtual nodes is hard (even if links are trivial)
 - Remark: the VNEP can also be seen as a generalization of the **Subgraph Isomorphism Problem (SIP)**
 - The SIP problem: Given two graphs G, H , determine whether G contains a subgraph that is isomorphic to H ?
 - NP-hard
- So if SIP is hard, why is VNEP hard?

Hamilton cycle problem (each node visited “?” solves maximum clique problem, etc.)



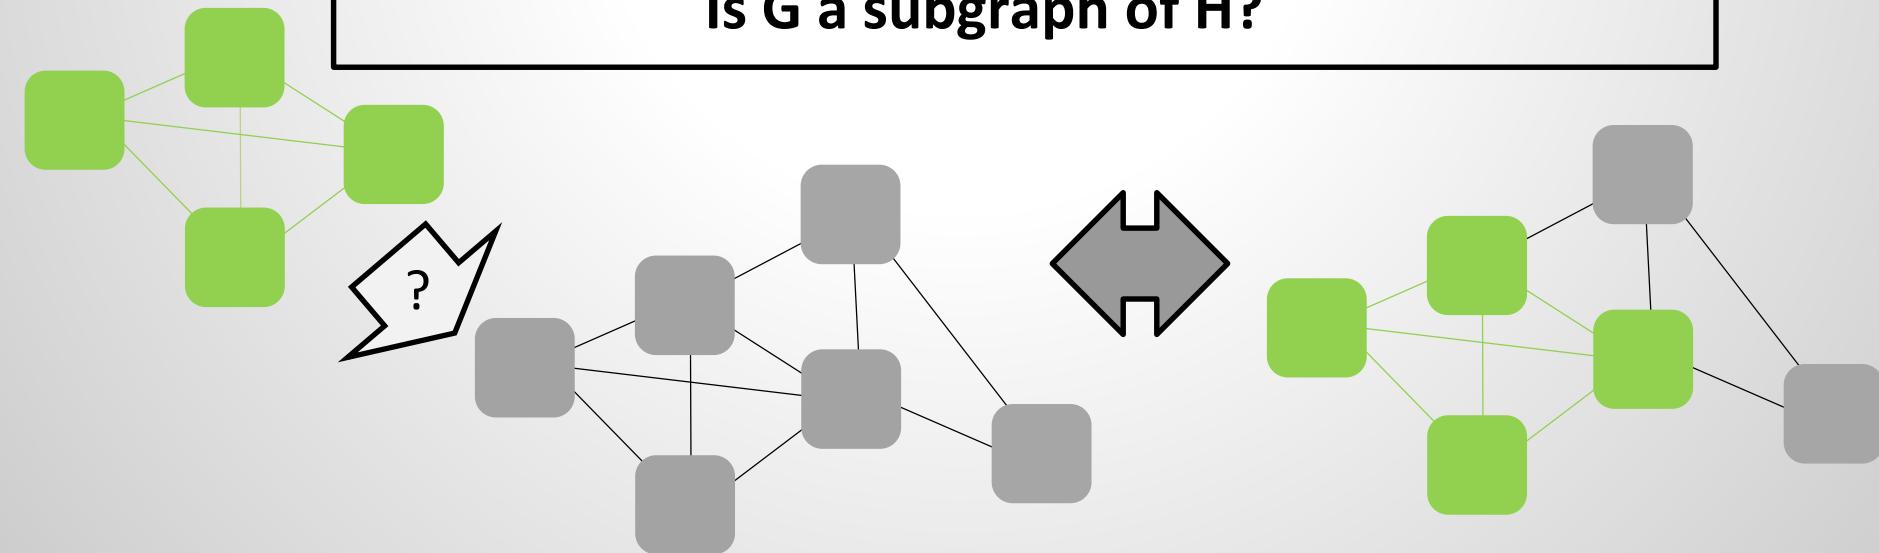
NP-Hardness: From SIP to VNEP

- ❑ Observe: VNEP is a **generalization** of SIP
- ❑ For example:

Can VNet $G=(V,E)$ be embedded in H at cost $|E|$?
(I.e., each virtual edge has **length 1**.)

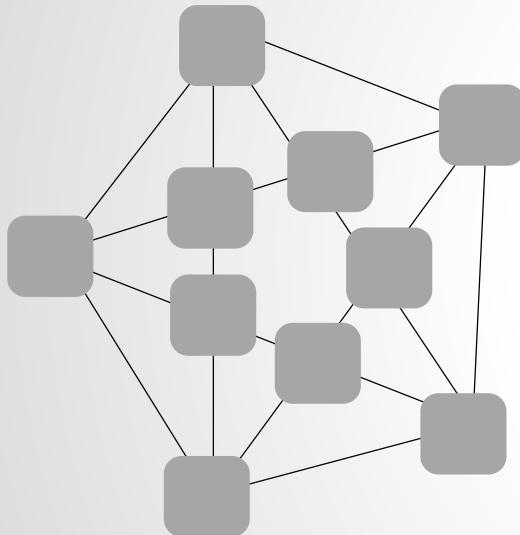


Is G a subgraph of H ?



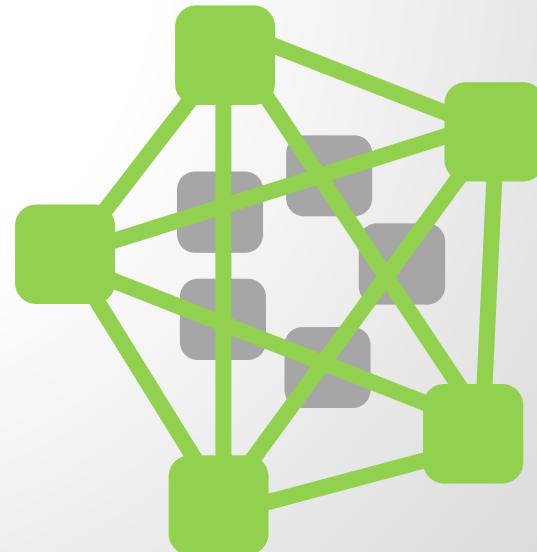
Remark: Graph Minors

Note: It is possible to embed a guest graph G on a host graph H , even though G is *not* a minor of H :



Assume planar host graph H :
 K_5 and $K_{3,3}$ minor-free...

... but it is possible to embed
non-planar guest graph $G=K_5$!



Can we at least formulate a “fast” MIP?

?

Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
 - ❑ Linear objective function (e.g., minimize embedding footprint)
 - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

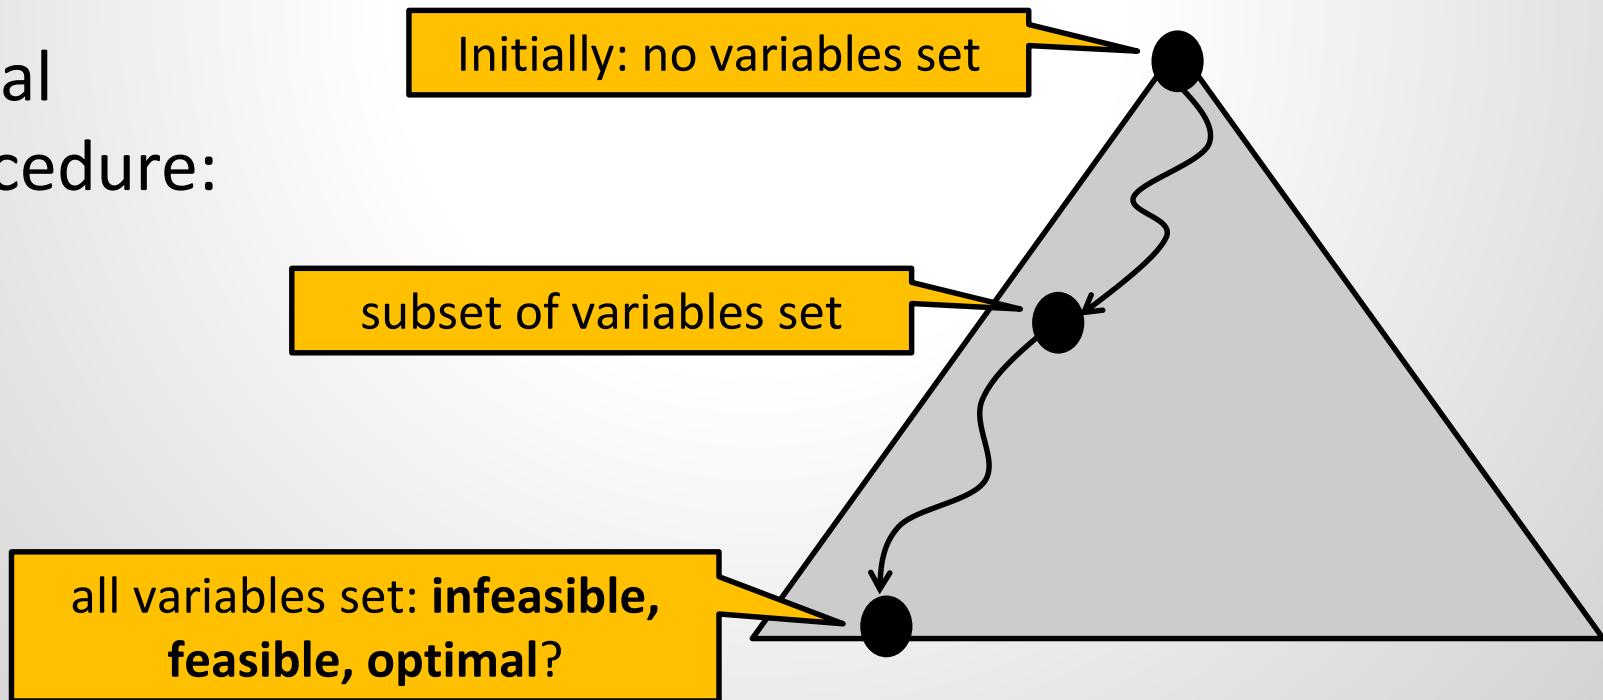
Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
 - ❑ Linear objective function (e.g., minimize)
 - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ One that provides good relaxations!
- ❑ Solved, e.g., with branch-and-bound search tree

Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
 - ❑ Linear objective function (e.g., minimize embedding footprint)
 - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound **search tree**

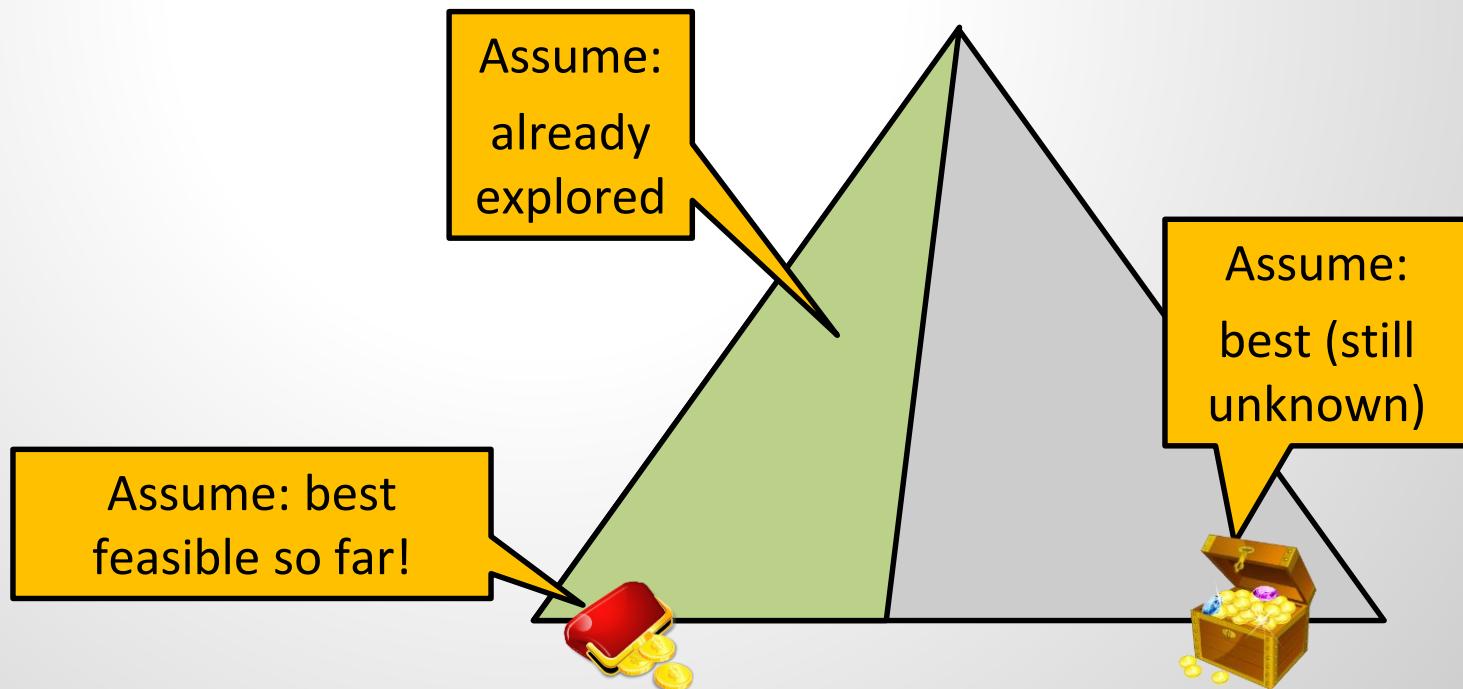
Usual
procedure:



Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
 - ❑ Linear objective function (e.g., minimize embedding footprint)
 - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

Usual
procedure:

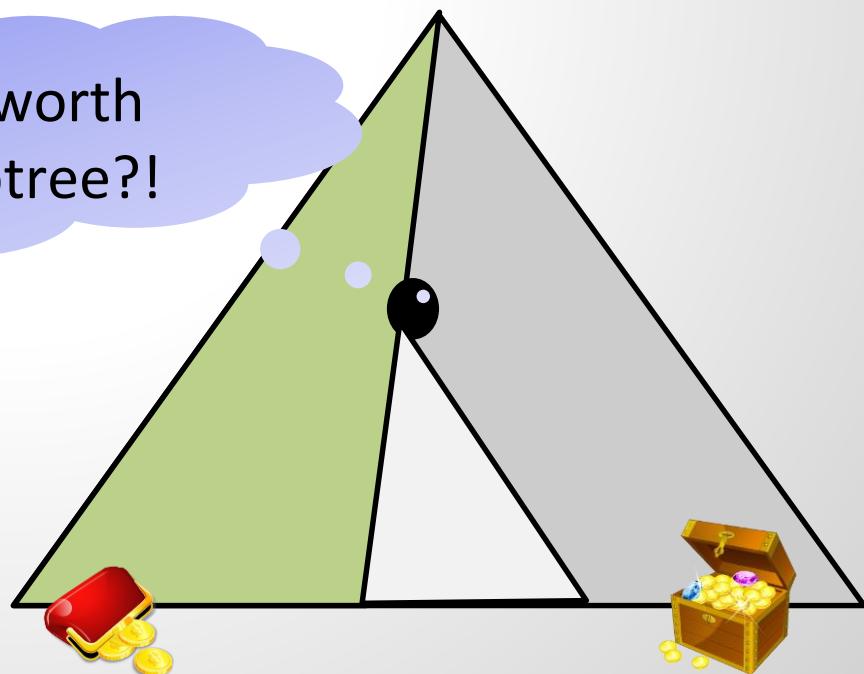


Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
 - ❑ Linear objective function (e.g., minimize embedding footprint)
 - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

Usual
procedure:

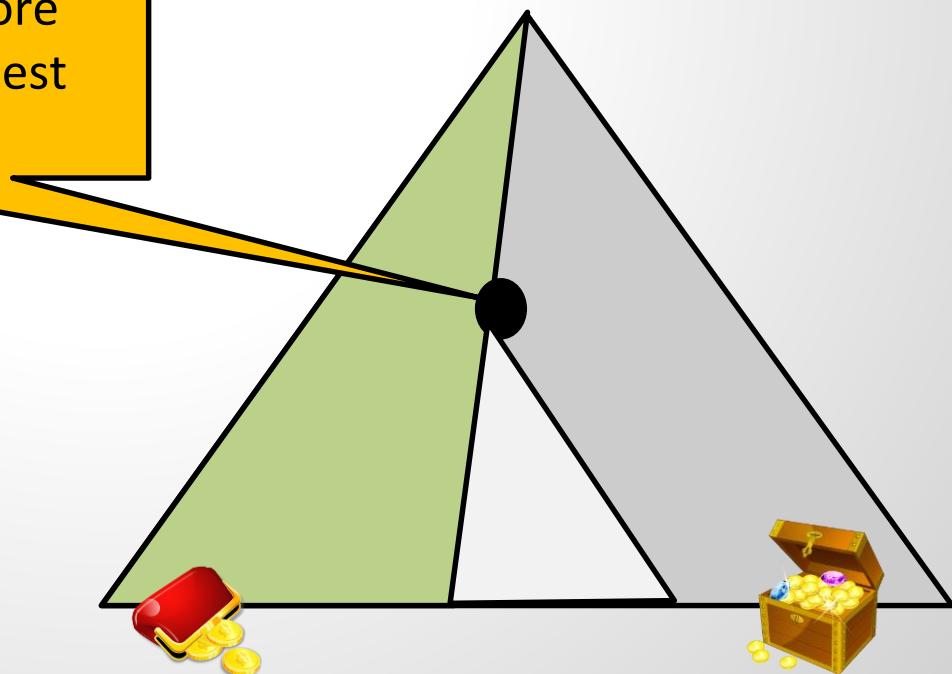
Decide: Is it worth
exploring subtree?!



Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
 - ❑ Linear objective function (e.g., minimize embedding footprint)
 - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

Usual trick: Relax! Solve LP (fast!),
and if **relaxed solution** (more
general!) **not better** then best
solution so far: skip it!

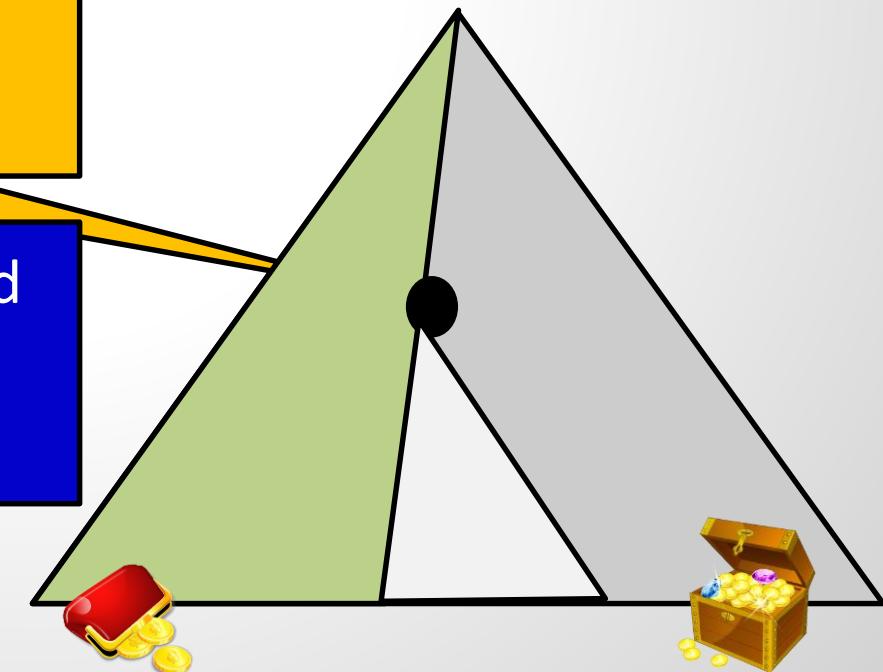


Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
 - ❑ Linear objective function (e.g., minimize embedding footprint)
 - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

Usual trick: Relax! Solve LP (fast!),
and if **relaxed solution** (more
general!) **not better** than best
solution so far: skip it!

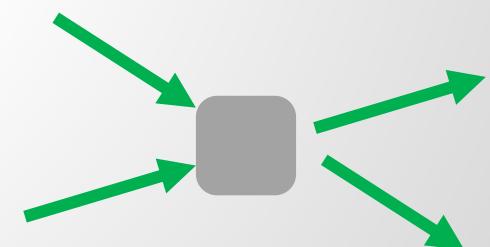
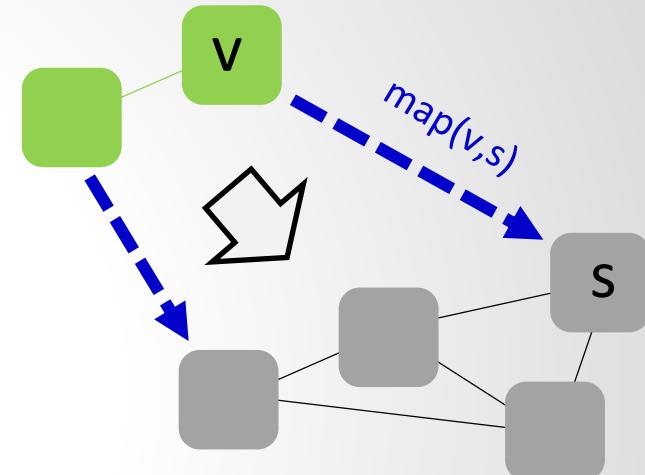
Bottomline: If MIP provides «good relaxations», large parts of the search space can be pruned.



Can we at least formulate a “fast” MIP?

A typical MIP formulation:

- ❑ Introduce **binary variables** $map(v,s)$ to map virtual nodes v to substrate node s
- ❑ Introduce **flow variables** for paths (say **splittable** for now)
- ❑ Ensure **flow conservation**: all flow entering a node must leave the node, unless it is the source or the destination



$$\sum_{u \rightarrow v} f_{uv} = \sum_{v \rightarrow w} f_{vw}$$

In Out

Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b
requested from node s
to node t.

In Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

What does this
formula do and why is
it correct?

Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b
requested from node s
to node t.

In Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

If $\text{map}(s,v)=1$, i.e., s mapped to v:
so **flow starts at v**, and hence
outgoing flow must be larger than
incoming flow (plus b).

Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b requested from node s to node t.

In Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

If $\text{map}(s,v)=0$ and $\text{map}(t,v)=0$, i.e., **v is along the path** from s to t: then we have **flow conservation**: outgoing flow must equal incoming flow (here \geq , objective function will remove unnecessary flow).

Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b
requested from node s
to node t.

In Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

If $\text{map}(t,v)=1$, i.e., t mapped to v: so flow **terminates at node v**: so **no constraint**: minus infinity (but objective function will remove unnecessary flow).

Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b requested from node s to node t.

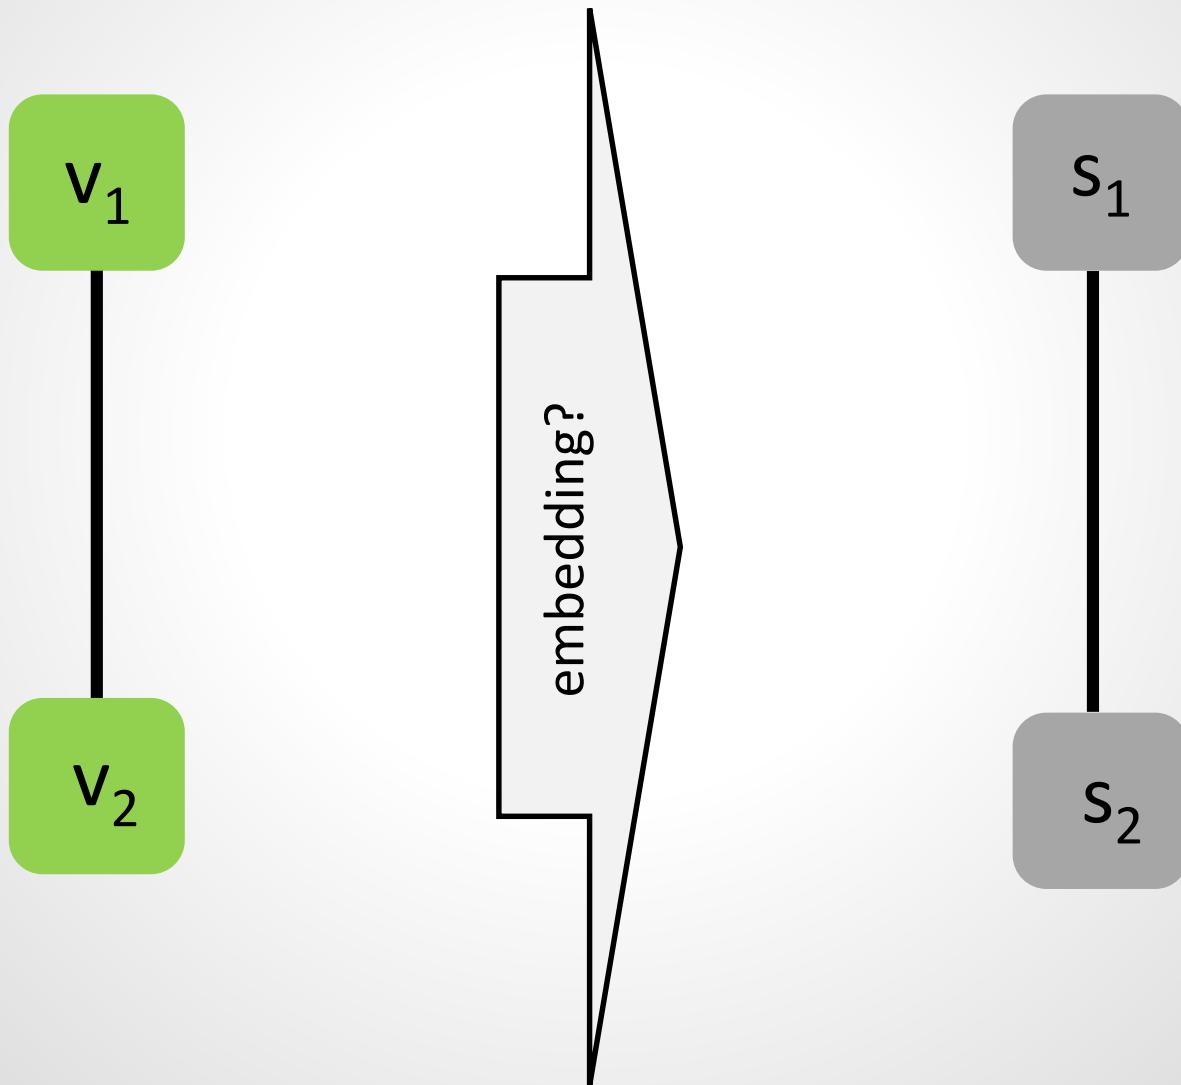
In Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

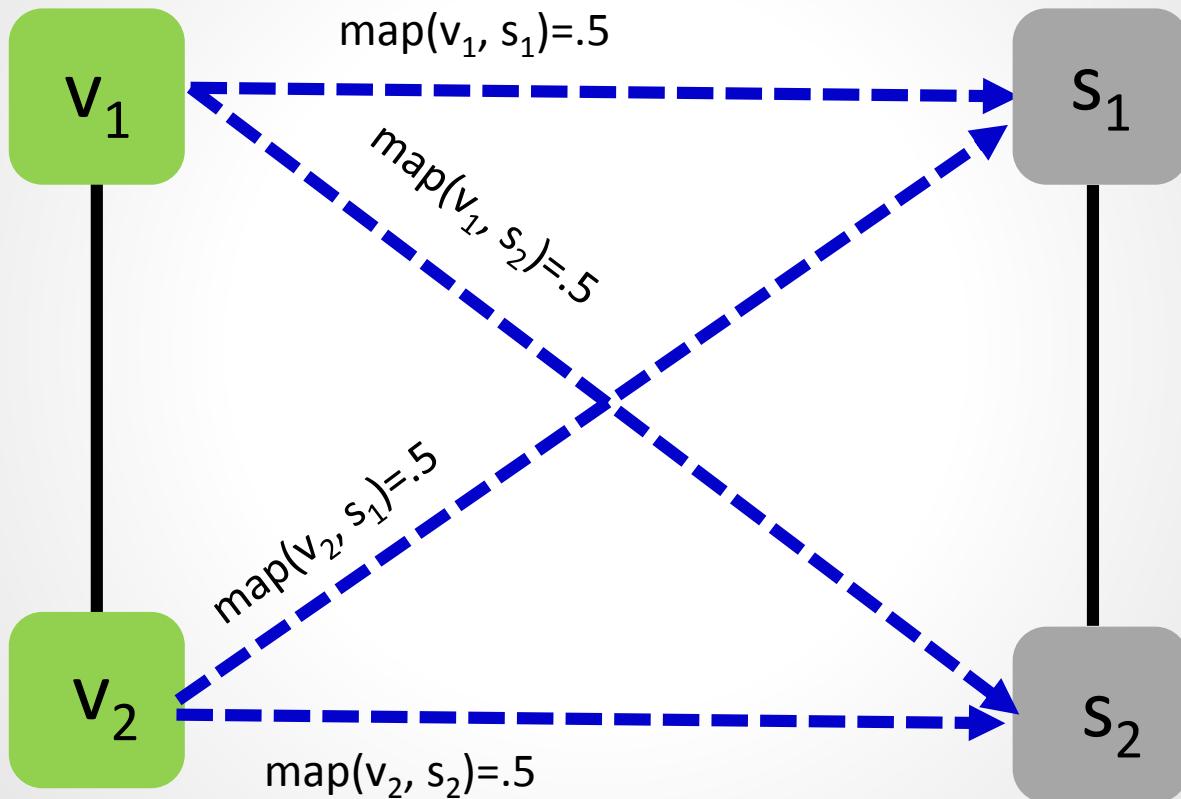
If $\text{map}(t,v)=1$, i.e., t mapped to v: so flow **terminates at node v**: so **no constraint**: minus infinity (but objective function will remove unnecessary flow).

Will such a MIP provide effective pruning?

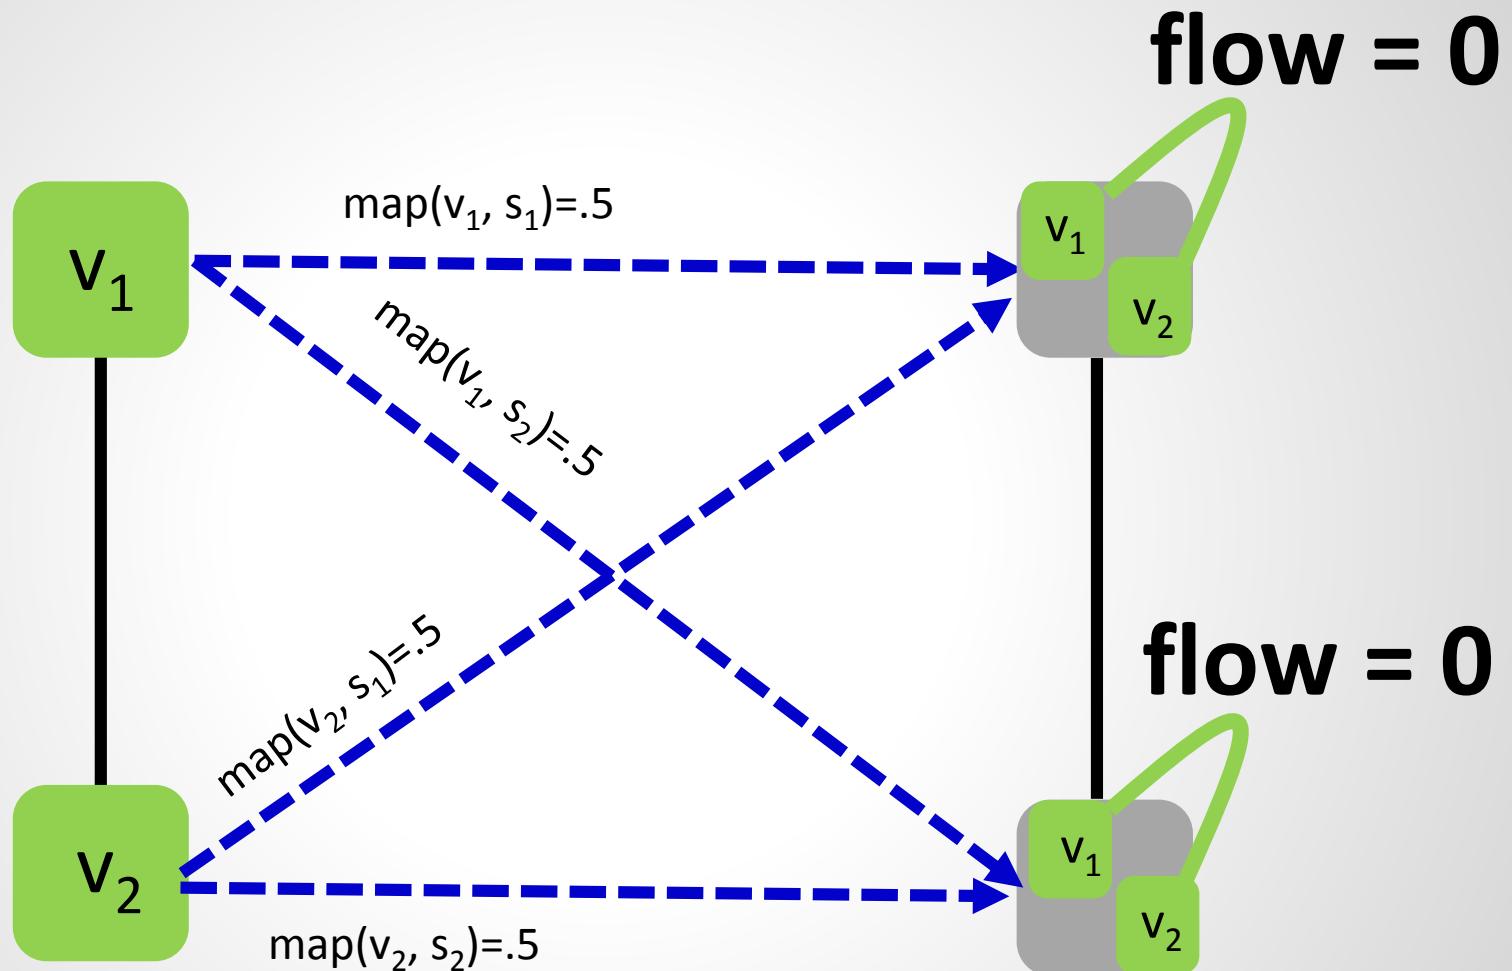
What will happen in this example?



What will happen in this example?



What will happen in this example?



Minimal flow = 0: fulfills flow conservation! Relaxation useless: does not provide any lower bound or indication of good mapping!

What about using randomized rounding?



Recall: classic approximation approach which:

- (i) computes a solution to the linear relaxation of the IP,
- (ii) **decomposes this solution into convex combinations of elementary solutions,**
- and (iii) probabilistically chooses any of the elementary solutions based on their weight.

What about using randomized rounding?

- ❑ Problem 1: relaxed solutions may not be very meaningful
 - ❑ see example for **splittable** flows before
- ❑ Problem 2: also for **unsplittable** flows, if using a standard **Multi-Commodity Flow (MCF)** formulation of VNEP, the **integrality gap** can be huge
 - ❑ Tree-like VNets are still ok
 - ❑ VNets with cycles: randomized rounding not applicable, since problem **not decomposable**



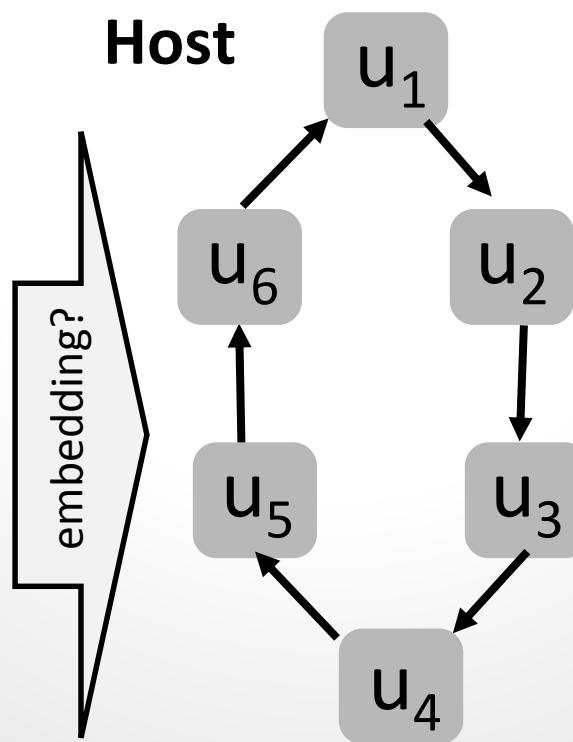
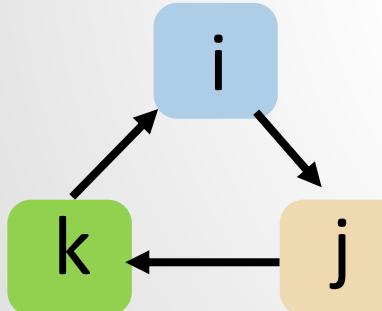
The linear solutions can be decomposed into convex combinations of valid mappings.

Non-Decomposability

- ❑ Relaxations of classic MCF formulation cannot be decomposed into convex combinations of valid mappings (so we need different formulations!)

- ❑ Example:

VNet

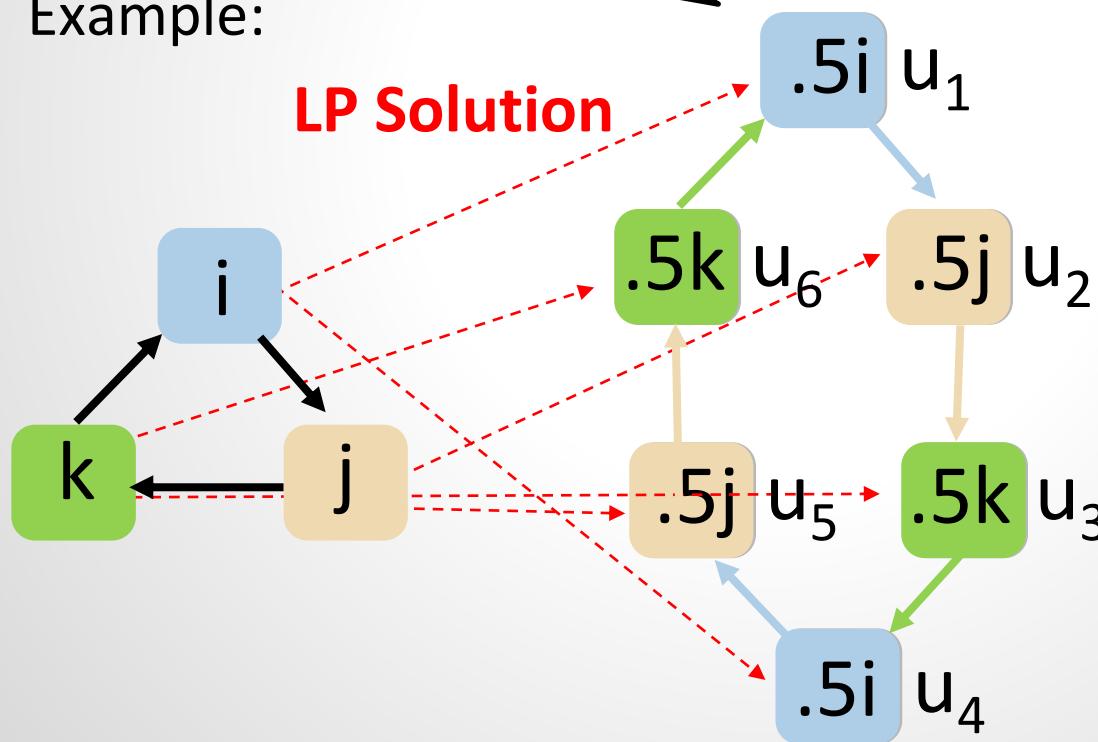


Non-Decomposability

Valid LP solution: virtual node mappings sum to 1 and each virtual node connects to its neighboring node **with half a unit of flow...**

MCE formulation cannot be decomposed
valid mappings (so we need

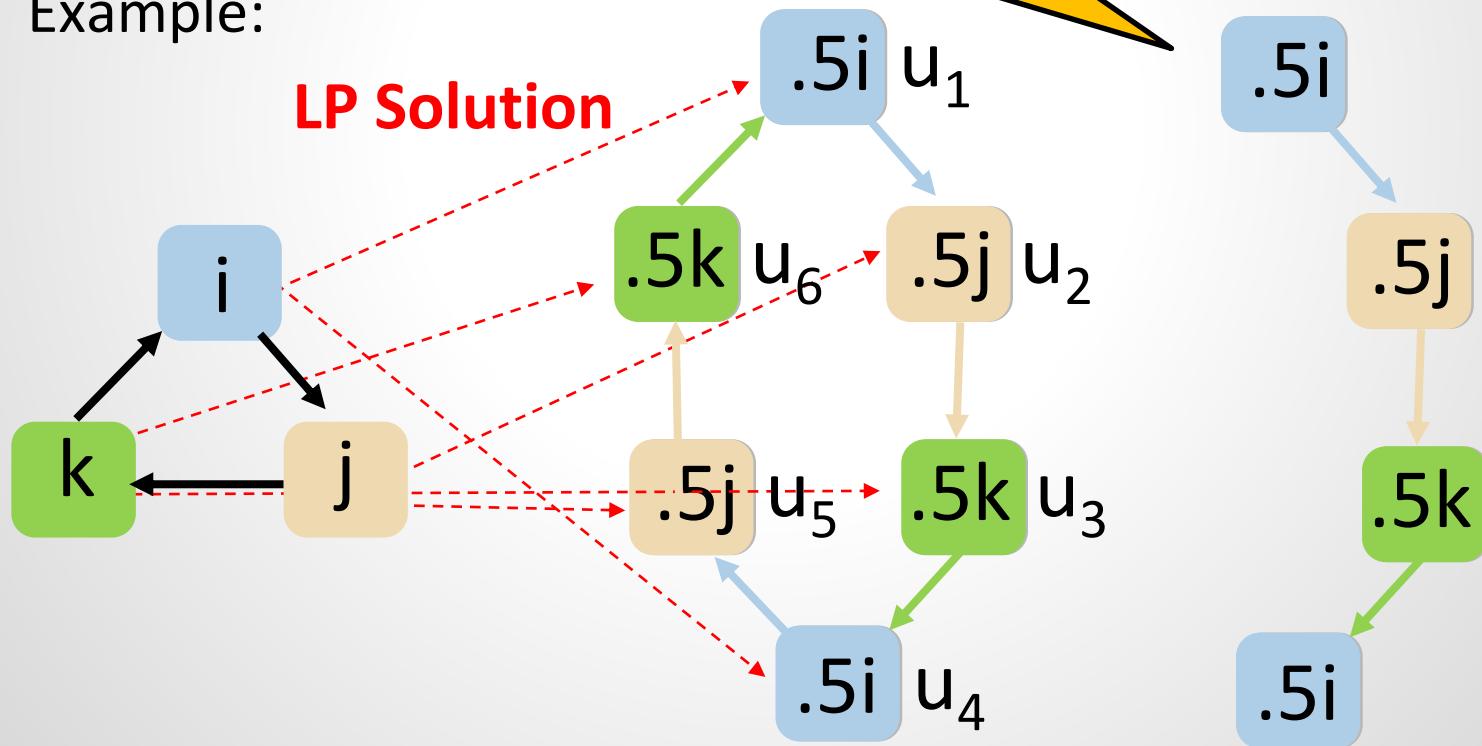
- ❑ Example:



Impossible to decompose and extract **any single valid mapping**. **Intuition:** Node i is mapped to u_1 and the only neighboring node that hosts j is u_2 , so i must be fully mapped on u_1 and j on u_2 . Similarly, k must be mapped on u_3 . But flow of virtual edge (k,i) leaving u_3 only leads to u_4 , so i must be mapped on both u_1 and u_4 . This is **impossible**, even if **capacities are infinite**.

not be decomposed
s (so we need
Partial
Decomposition

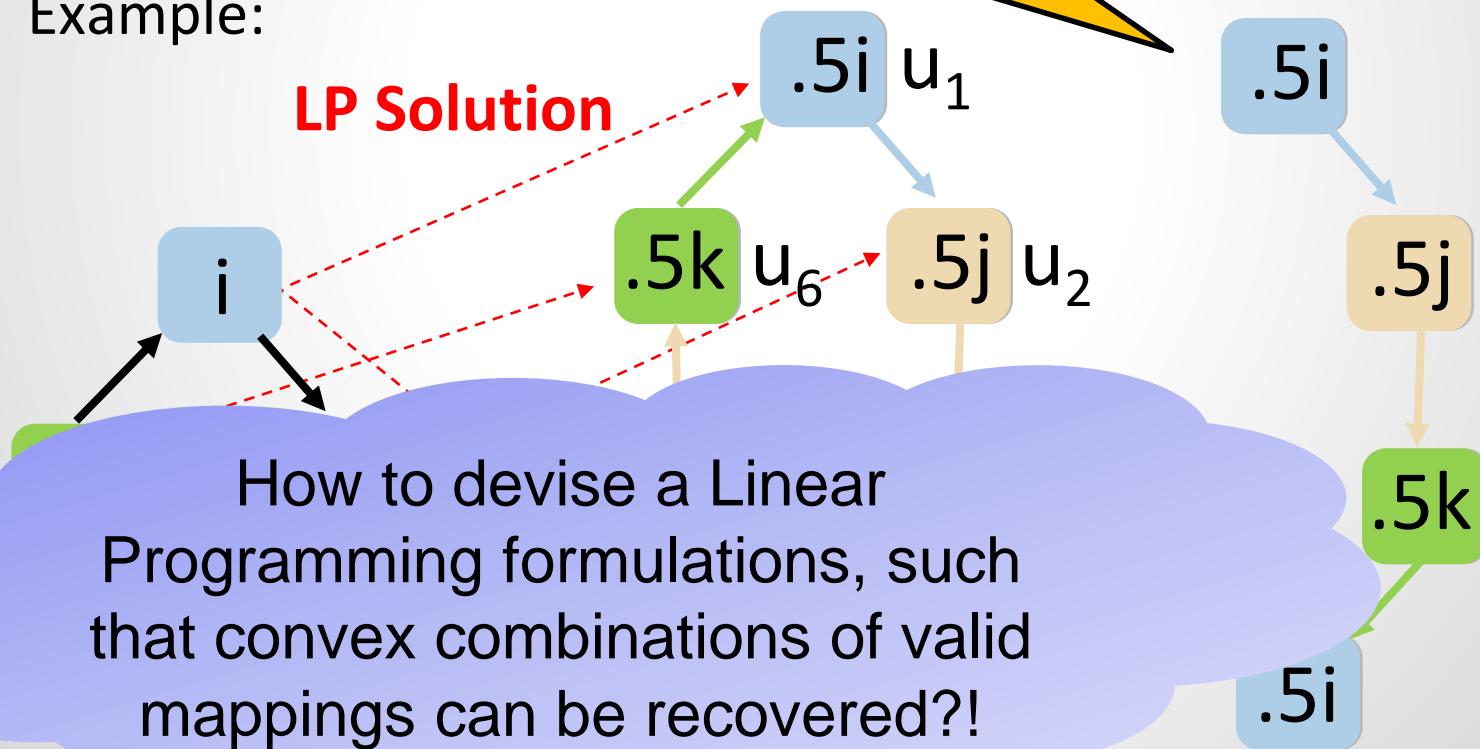
- Example:



Impossible to decompose and extract **any single valid mapping**. **Intuition:** Node i is mapped to u_1 and the only neighboring node that hosts j is u_2 , so i must be fully mapped on u_1 and j on u_2 . Similarly, k must be mapped on u_3 . But flow of virtual edge (k,i) leaving u_3 only leads to u_4 , so i must be mapped on both u_1 and u_4 . This is **impossible**, even if **capacities are infinite**.

not be decomposed
s (so we need
Partial
Decomposition

□ Example:



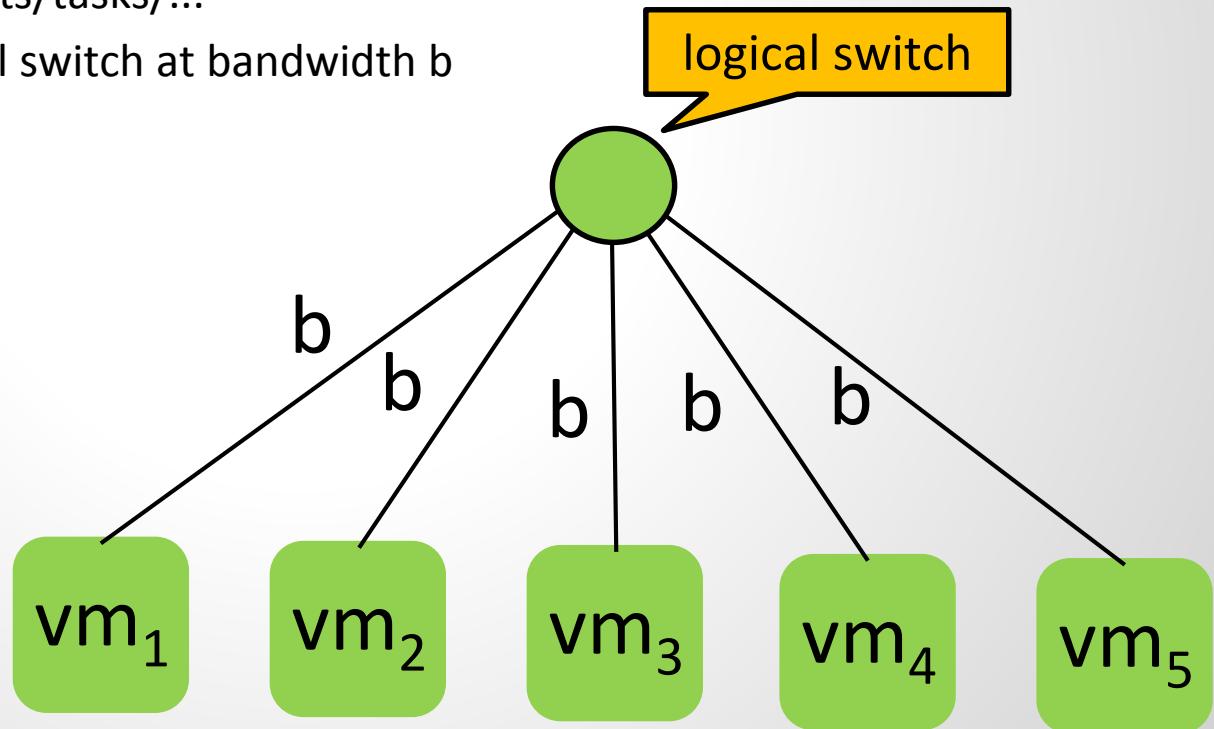
That's all Folks!

That's all Folks!

Wait a minute!
These problems need to be solved!
And they often can, even with guarantees.

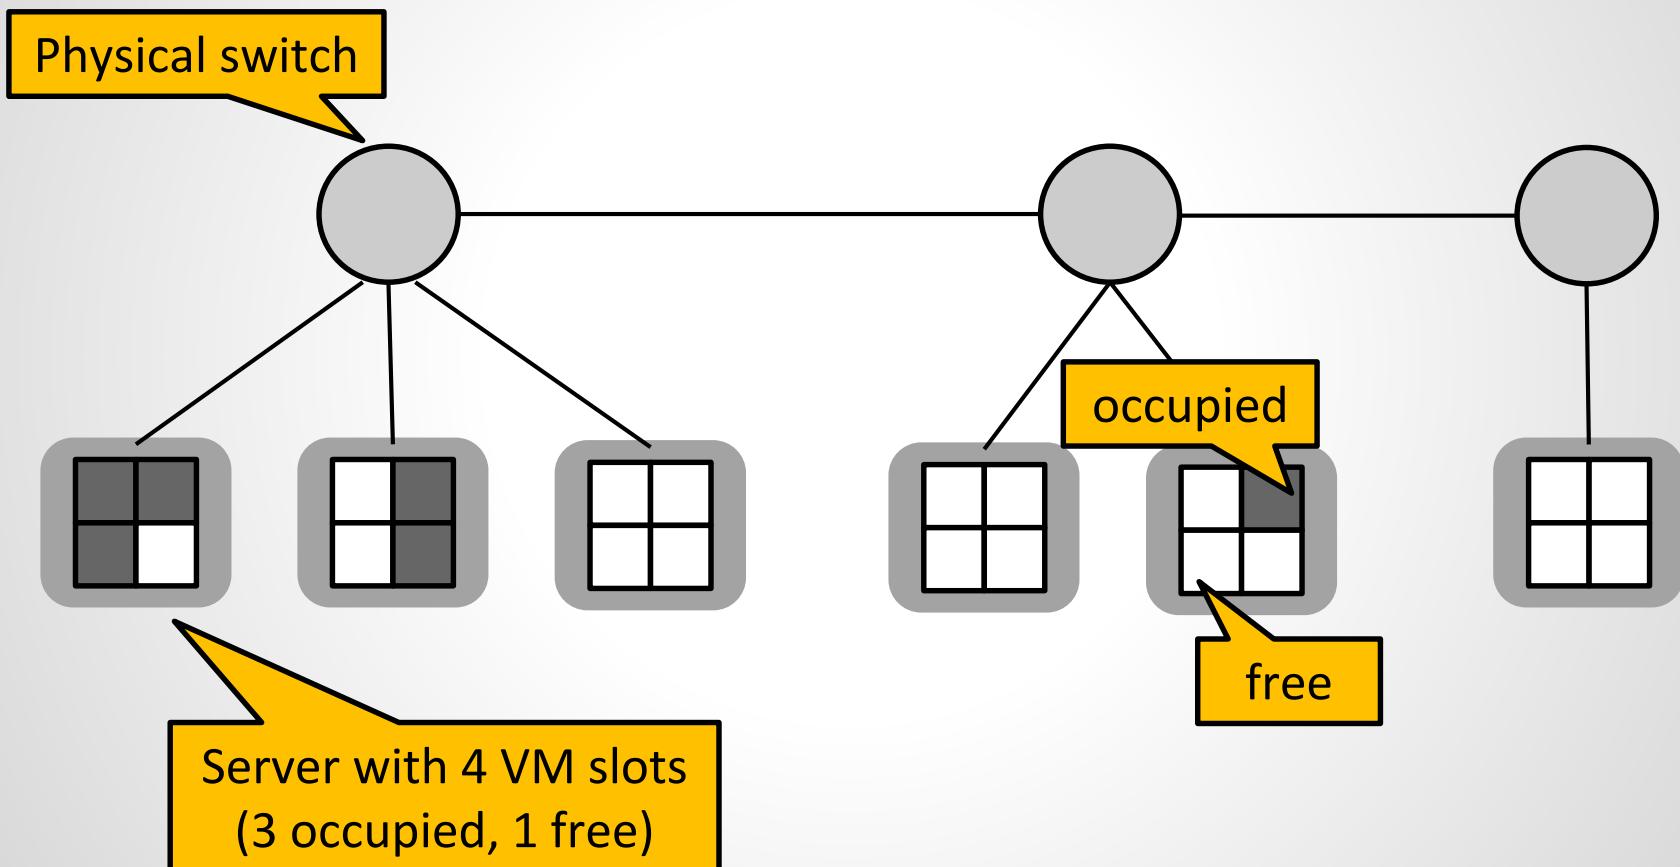
Theory vs Practice: In Practice There is Hope!

- ❑ Guest graphs may not be general graphs, but e.g., virtual clusters: very simple and symmetric, used in context of batch processing
 - ❑ k VMs/compute-units/tasks/...
 - ❑ Connected to virtual switch at bandwidth b



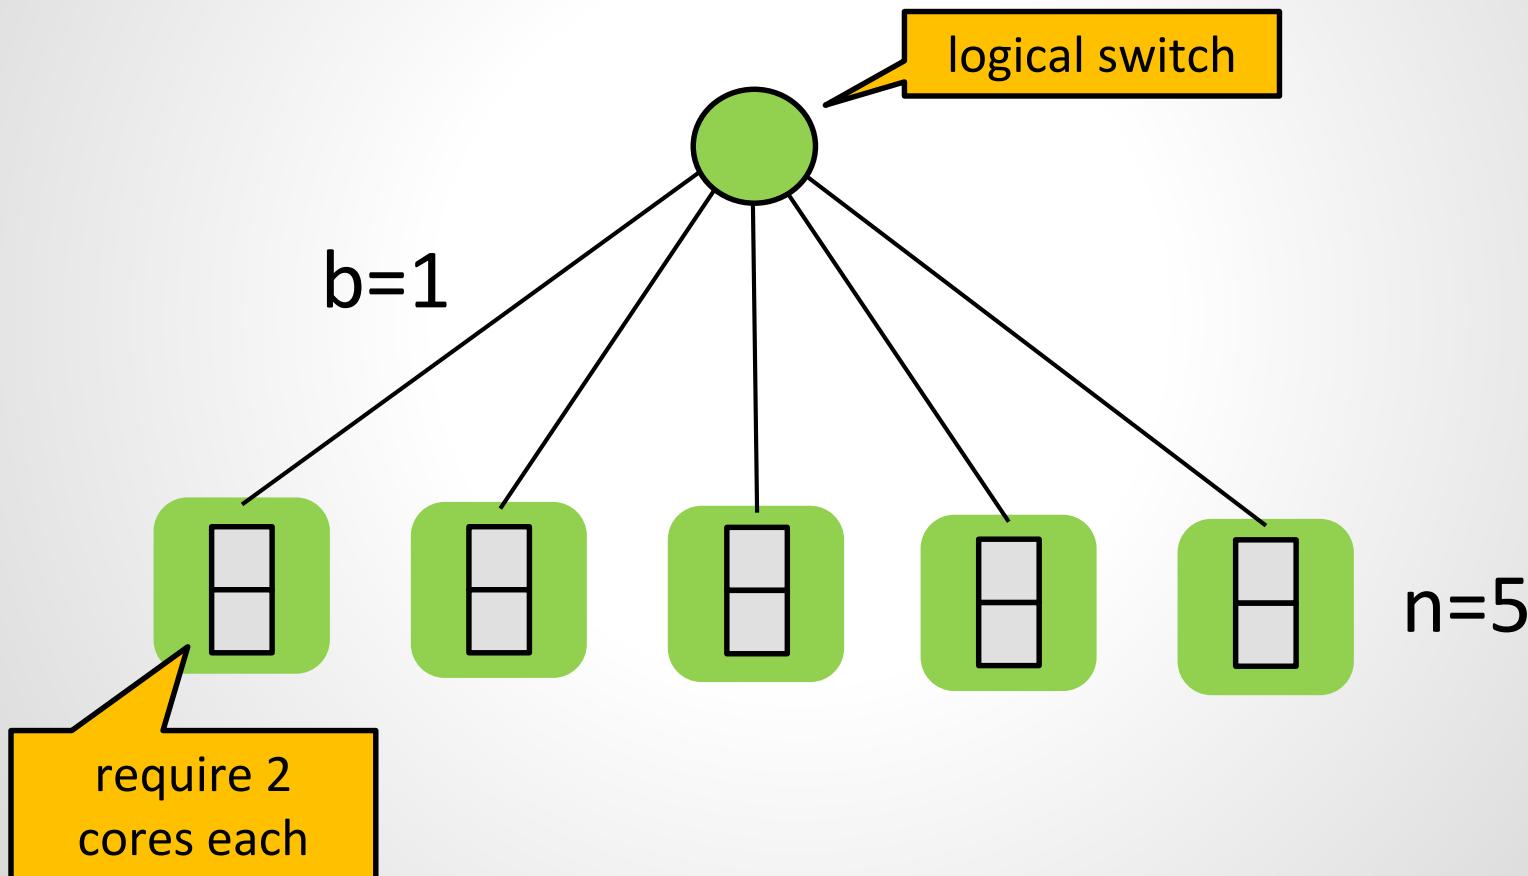
How to Embed a Virtual Cluster?

Consider host graph:

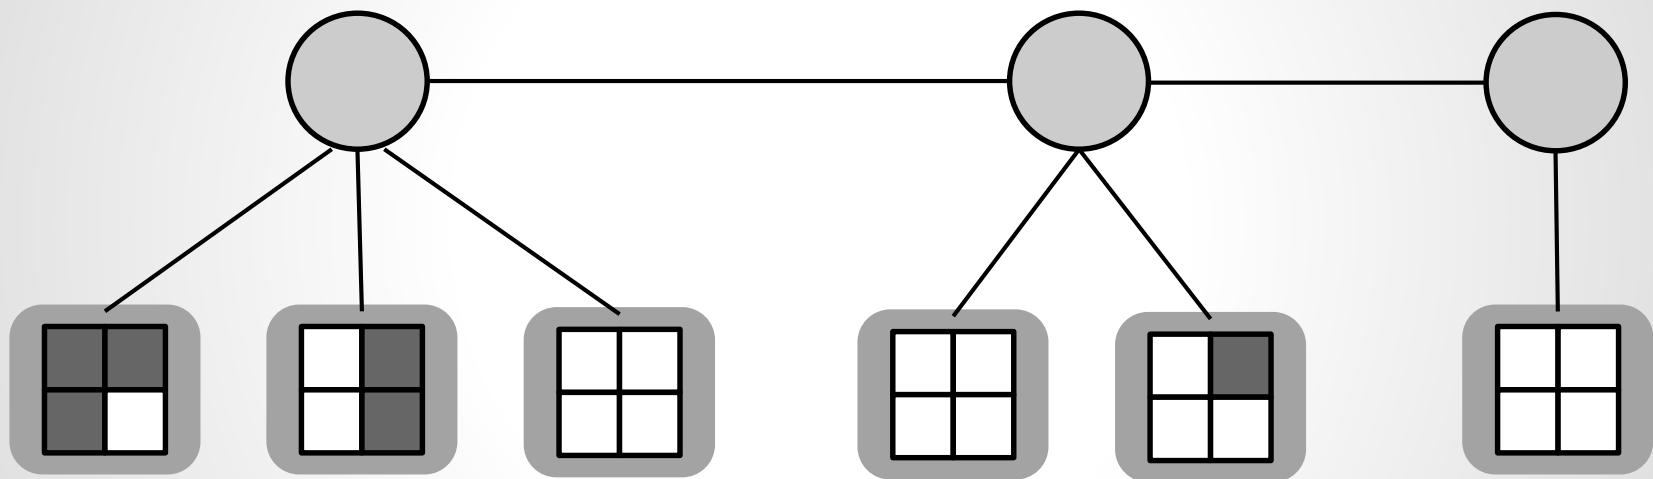


How to Embed a Virtual Cluster?

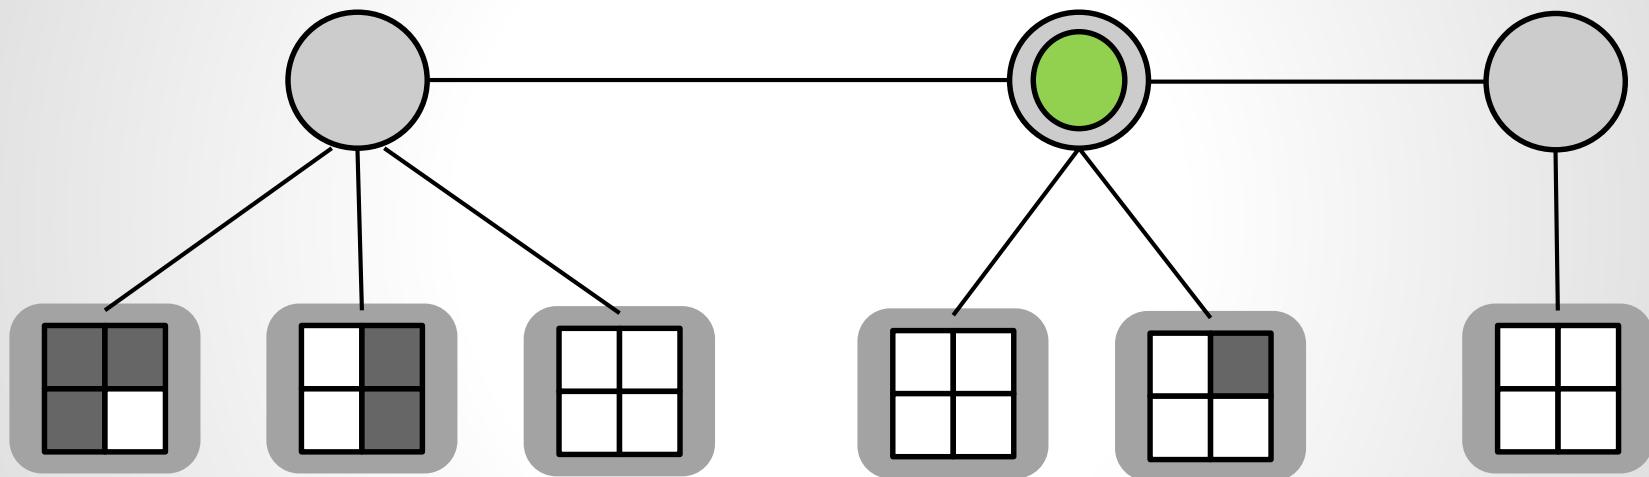
Consider guest graph:



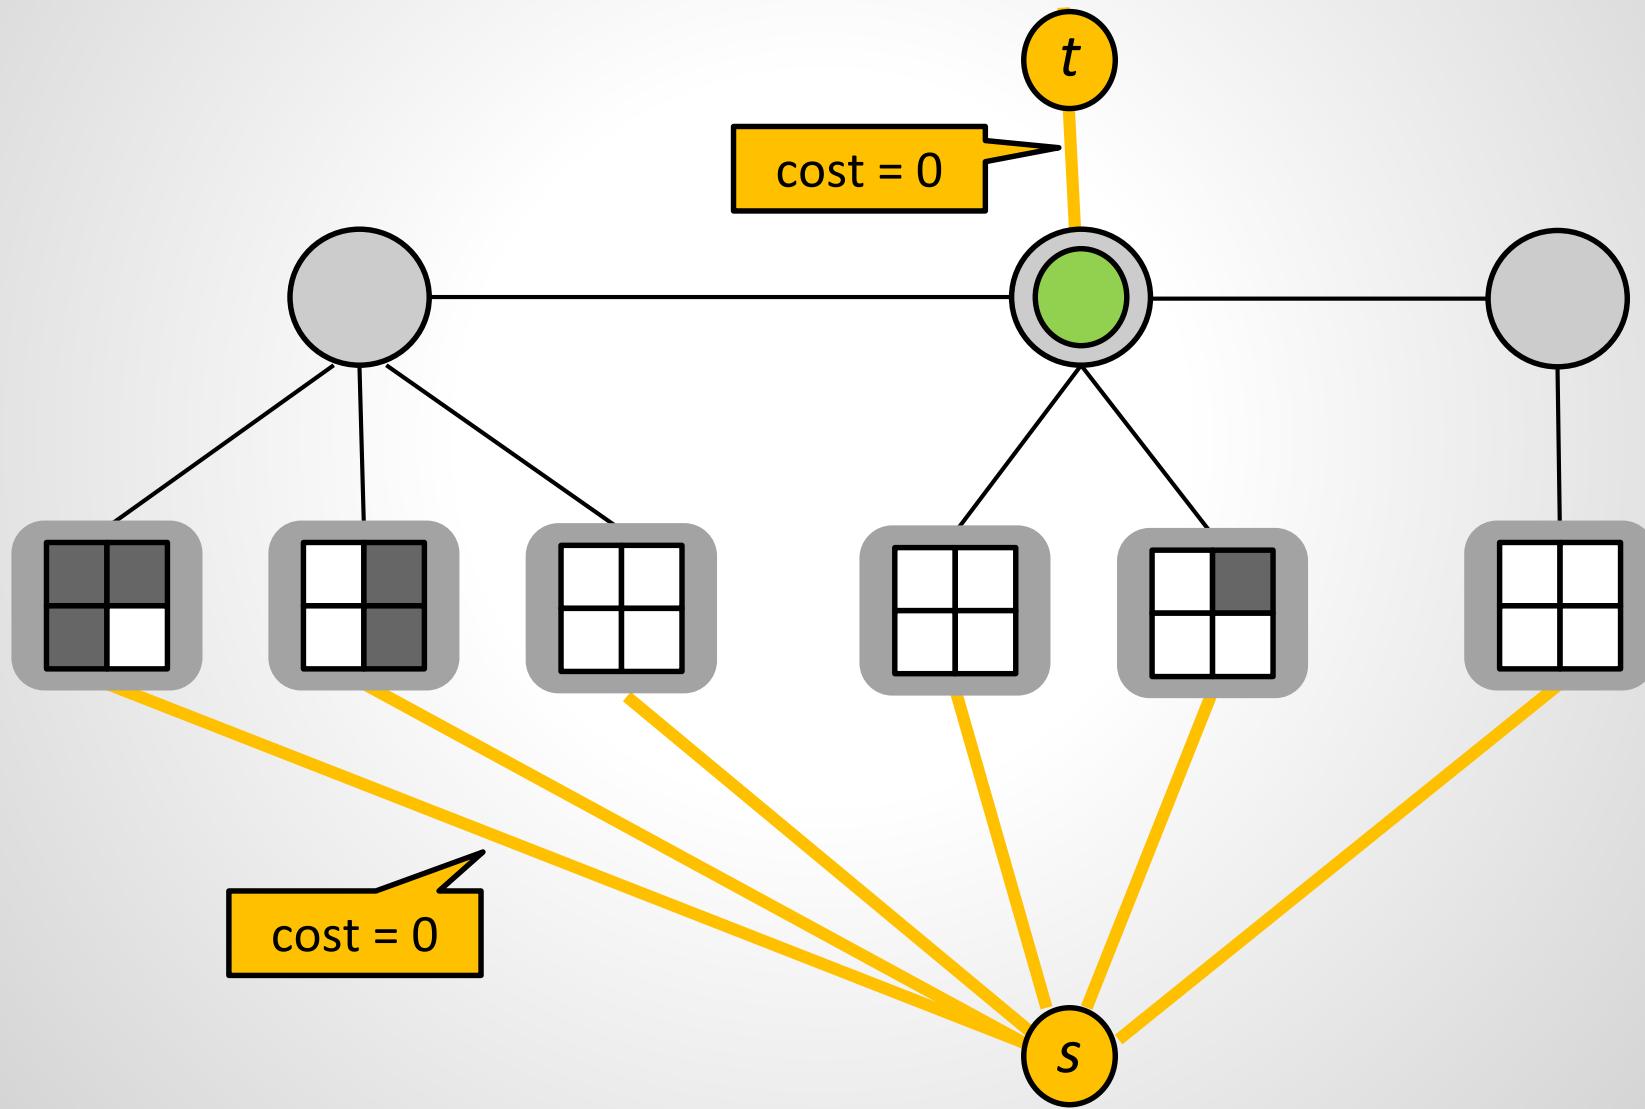
How to Embed a Virtual Cluster?



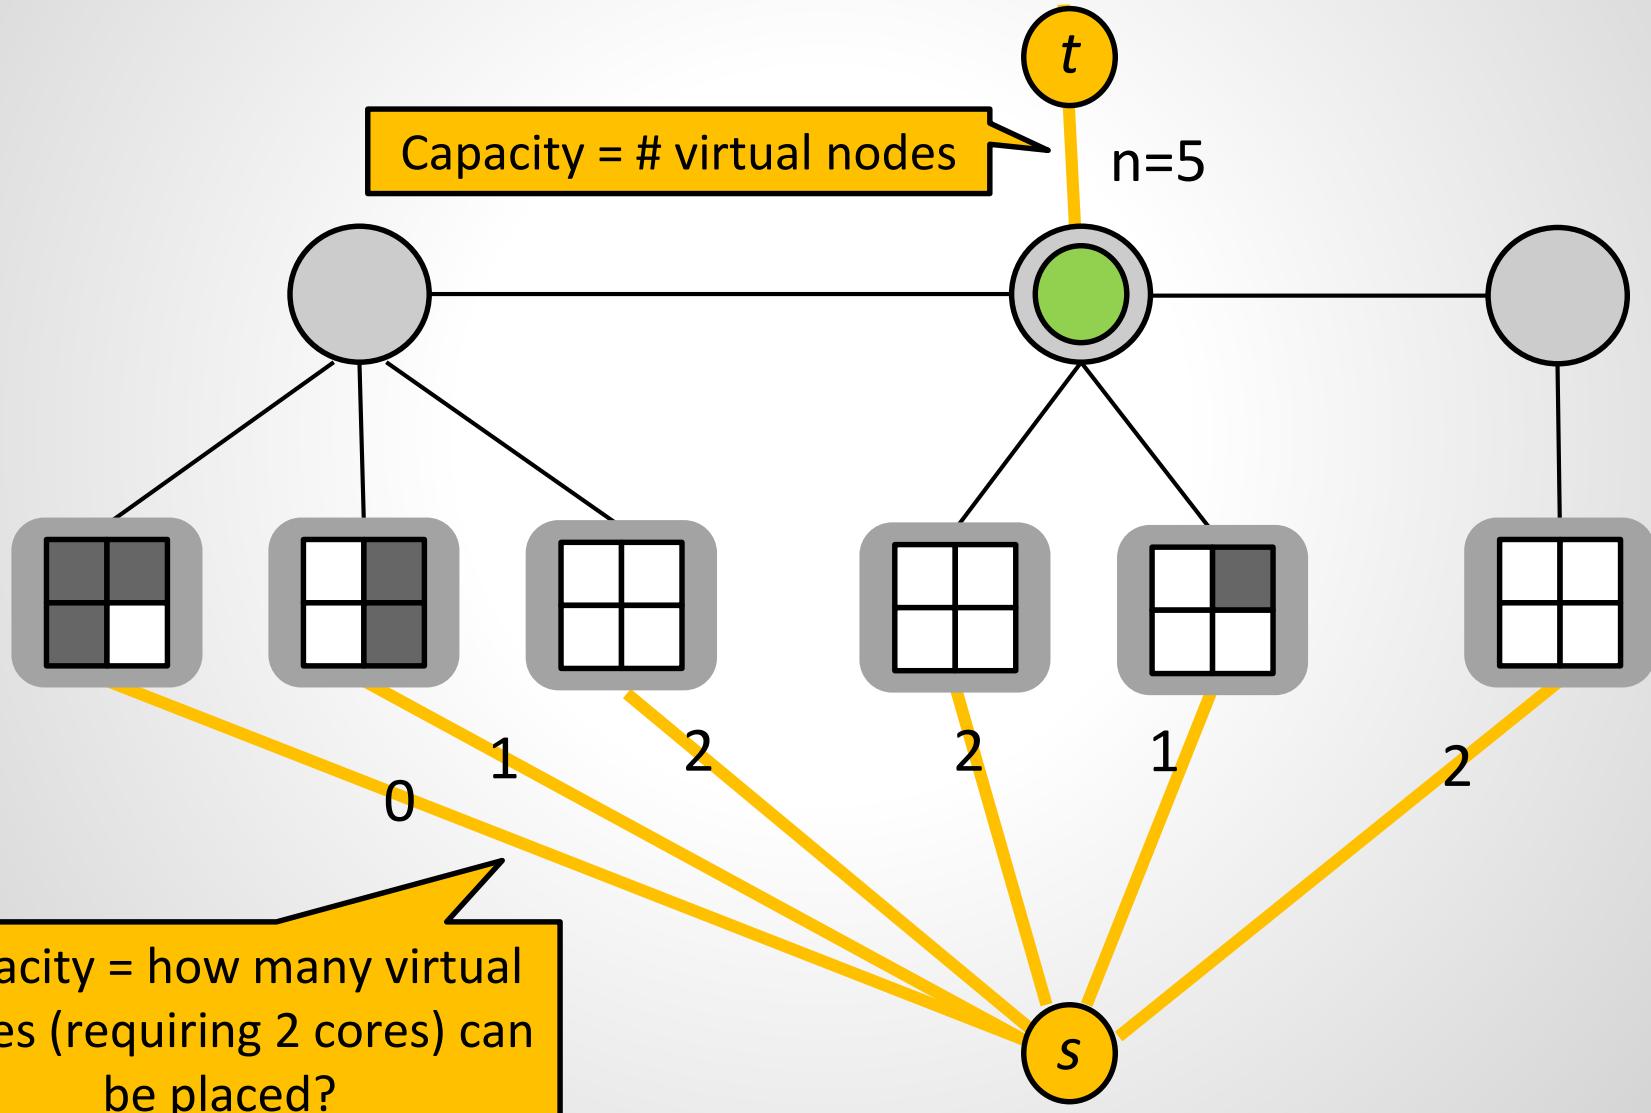
1. Place logical switch (try all options)



1. Place logical switch (try all options)
2. Extend network with artificial source s and sink t



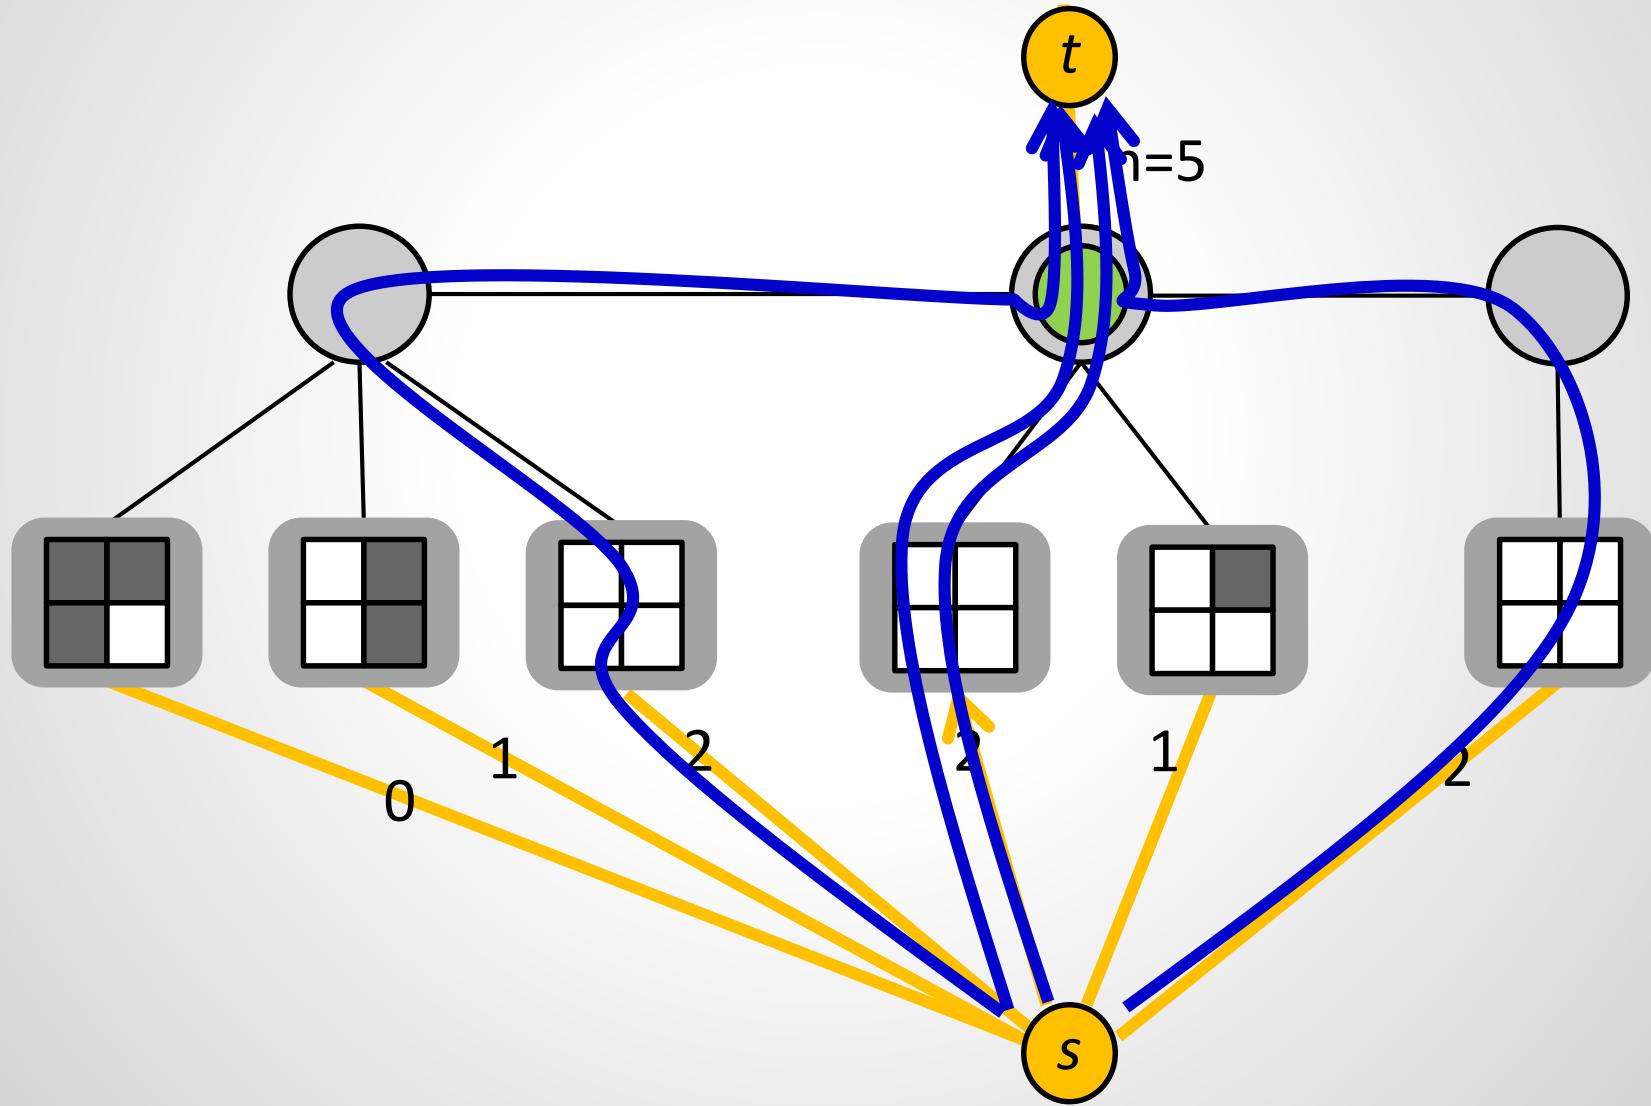
1. Place logical switch (try all options)
2. Extend network with artificial source s and sink t
3. Add capacities (recall that $b=1$, so each virtual node needs one unit of capacity)



Then: Compute min-cost max flow of size n from s to t

(e.g., successive shortest paths): due to capacity

constraints at most size n.

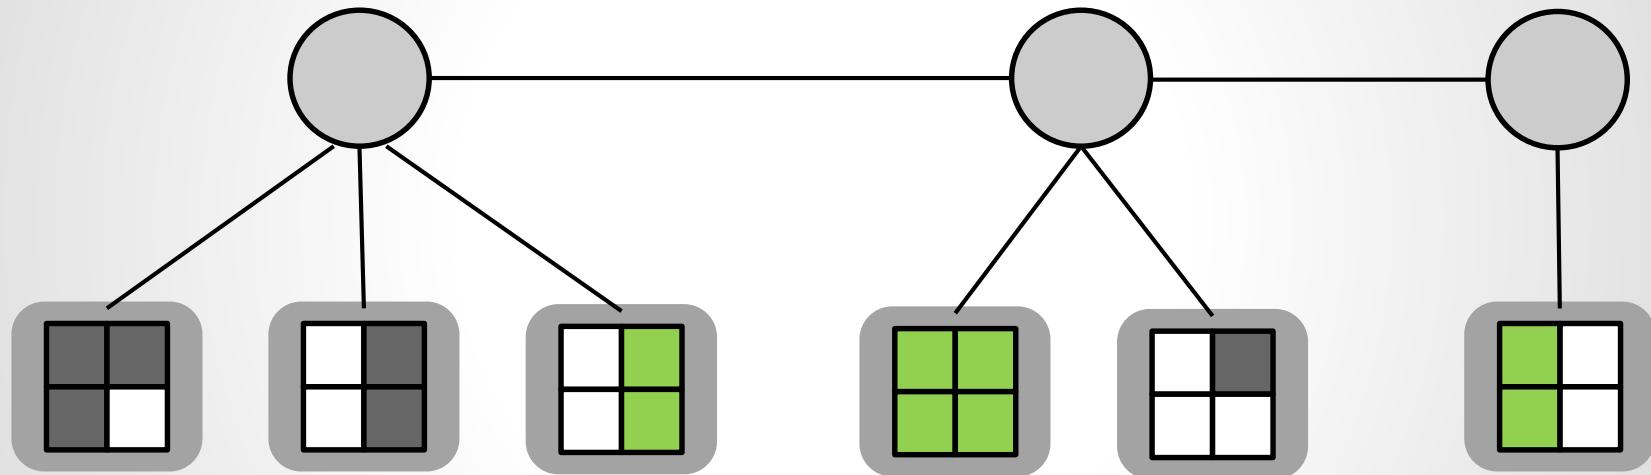


Then: Compute **min-cost max flow** of size n from s to t

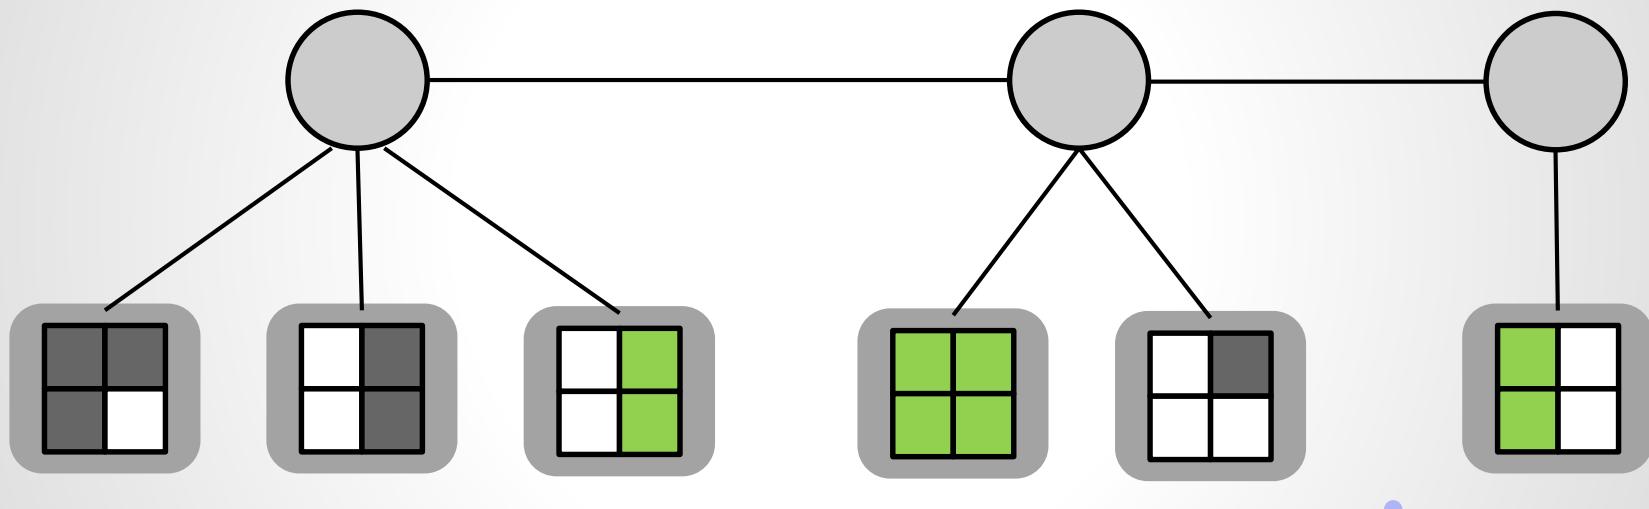
(e.g., **successive shortest paths**): due to capacity

constraints at most size n.

... and **assign virtual nodes (and edges)** accordingly!



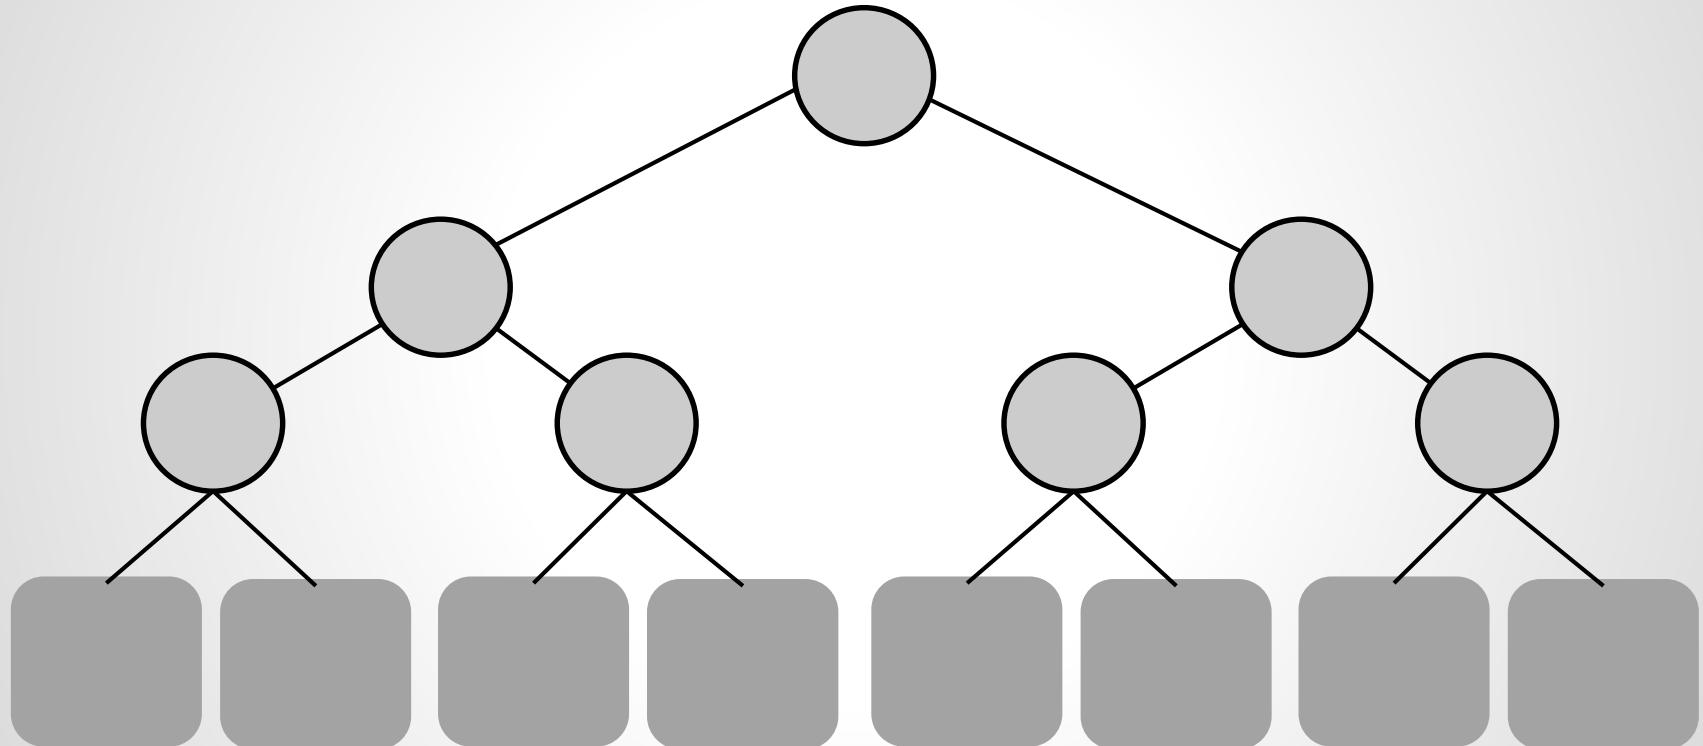
How to Embed a Virtual Cluster?



In fact: this physical network is even ***a tree!***

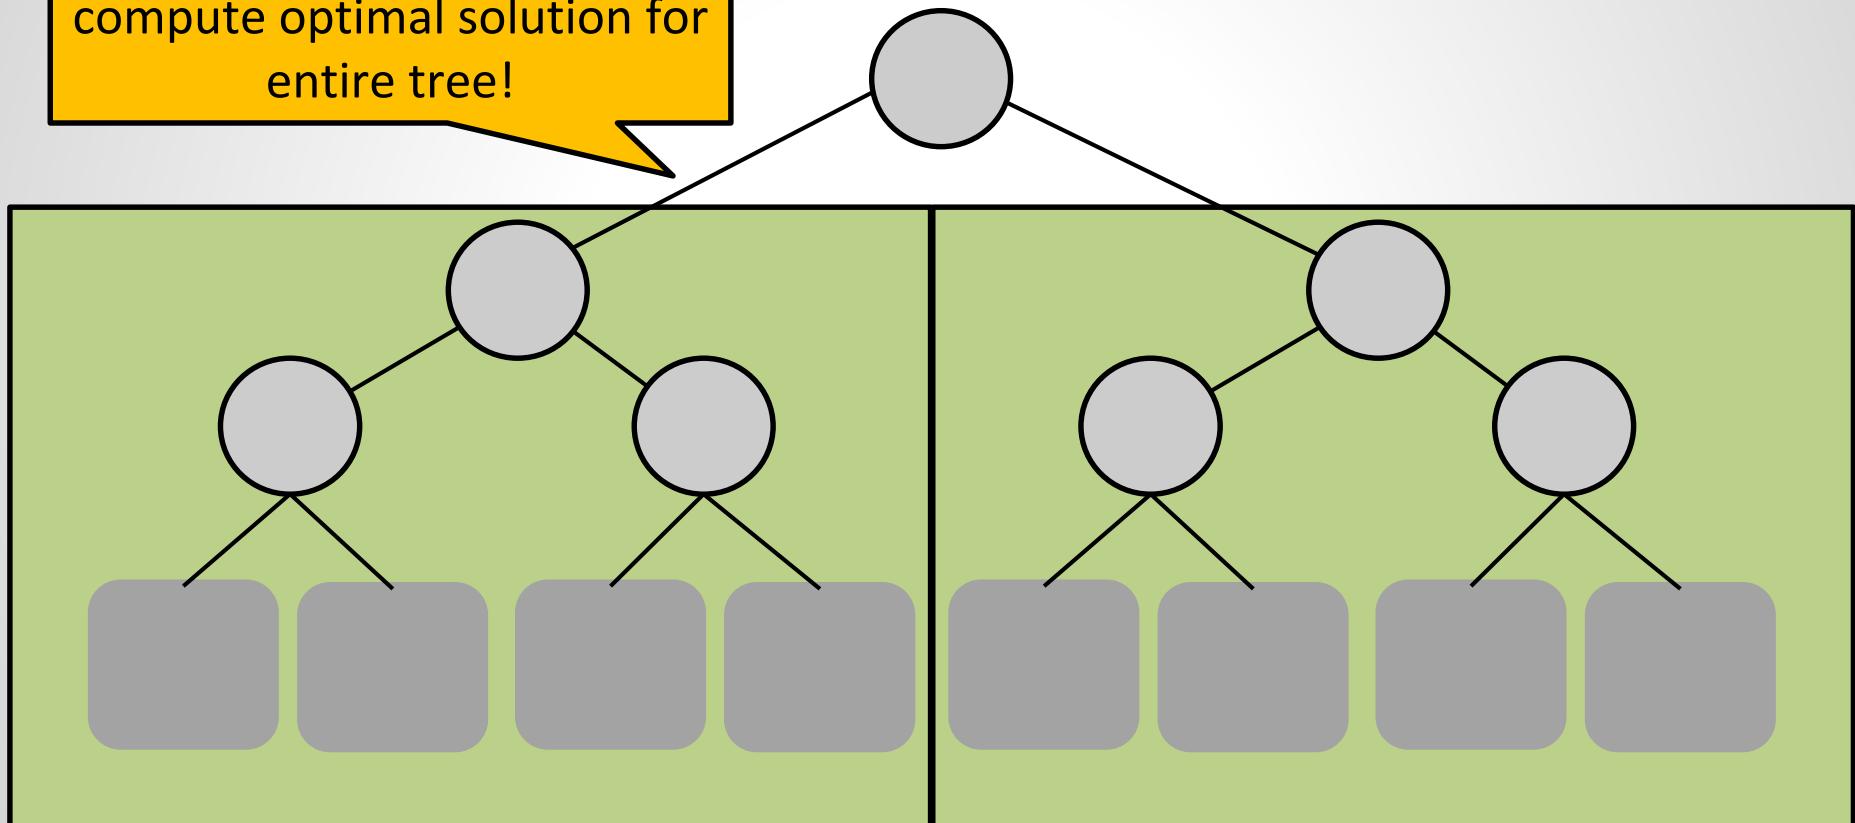
For **trees with servers at leaves**, even simpler algorithms are possible. Ideas?

Dynamic Programming

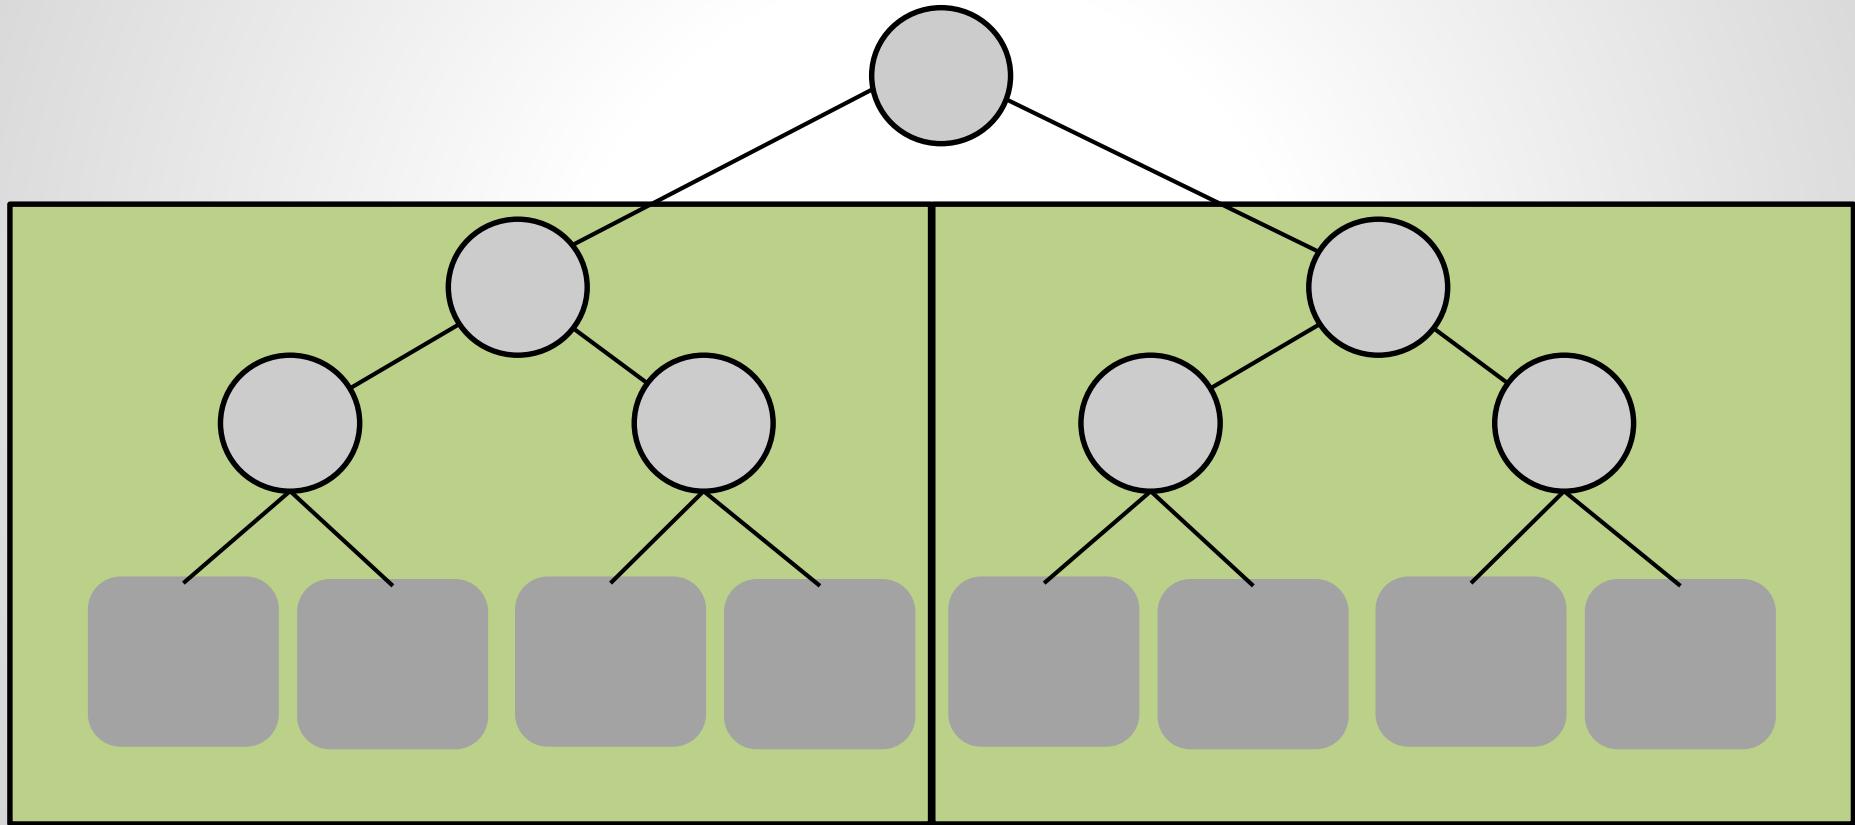


Dynamic Programming

Bottom-up programming:
given optimal solution for
subtrees, can quickly
compute optimal solution for
entire tree!



Dynamic Programming

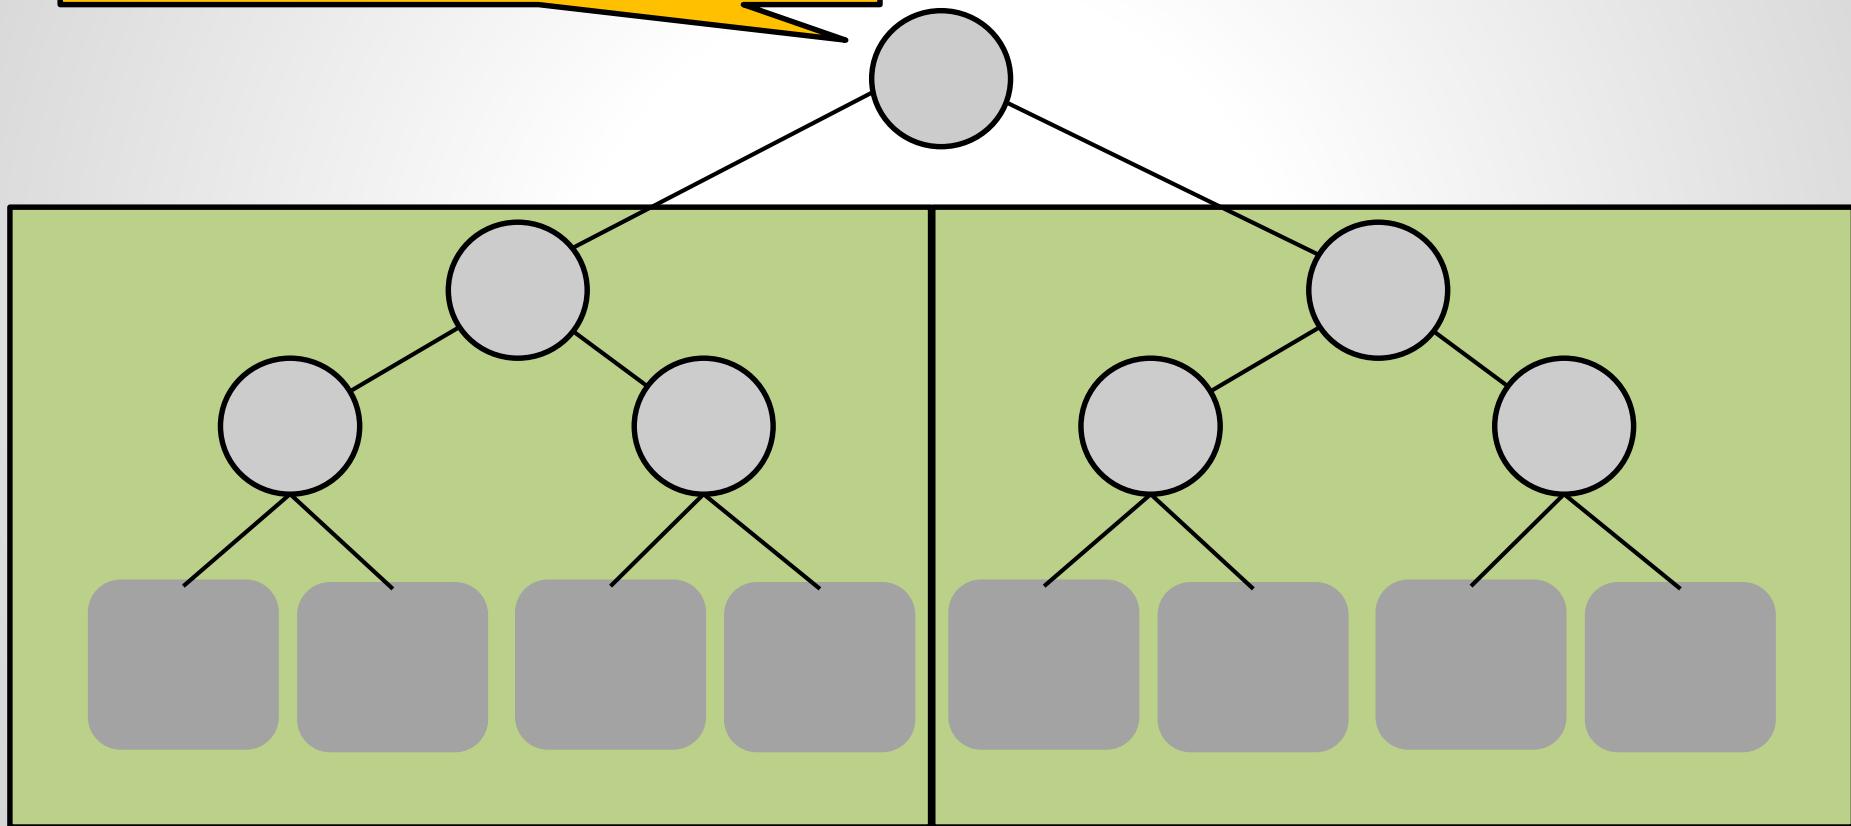


Given **optimal embedding** for $x \in \{0, \dots, n\}$ virtual nodes in left subtree...

Given **optimal embedding** for $x \in \{0, \dots, n\}$ virtual nodes in right subtree...

Dynamic Programming

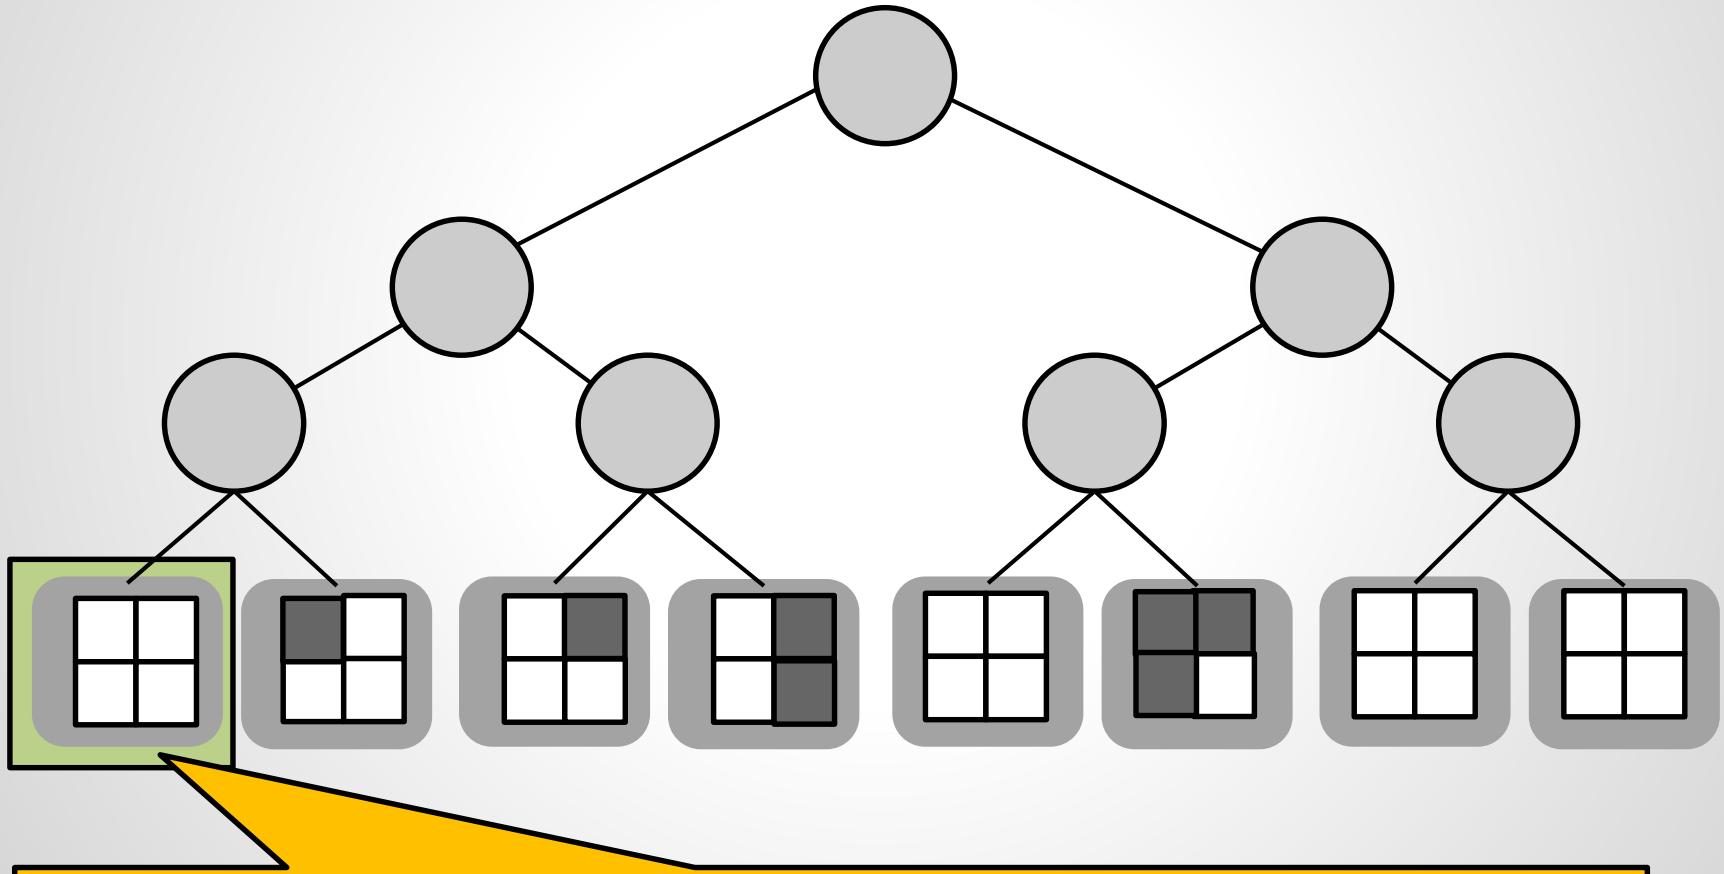
... can compute optimal embedding
of $x \in \{0, \dots, n\}$ virtual nodes in
entire subtree!



Given **optimal embedding** for $x \in \{0, \dots, n\}$ virtual nodes in left subtree...

Given **optimal embedding** for $x \in \{0, \dots, n\}$ virtual nodes in right subtree...

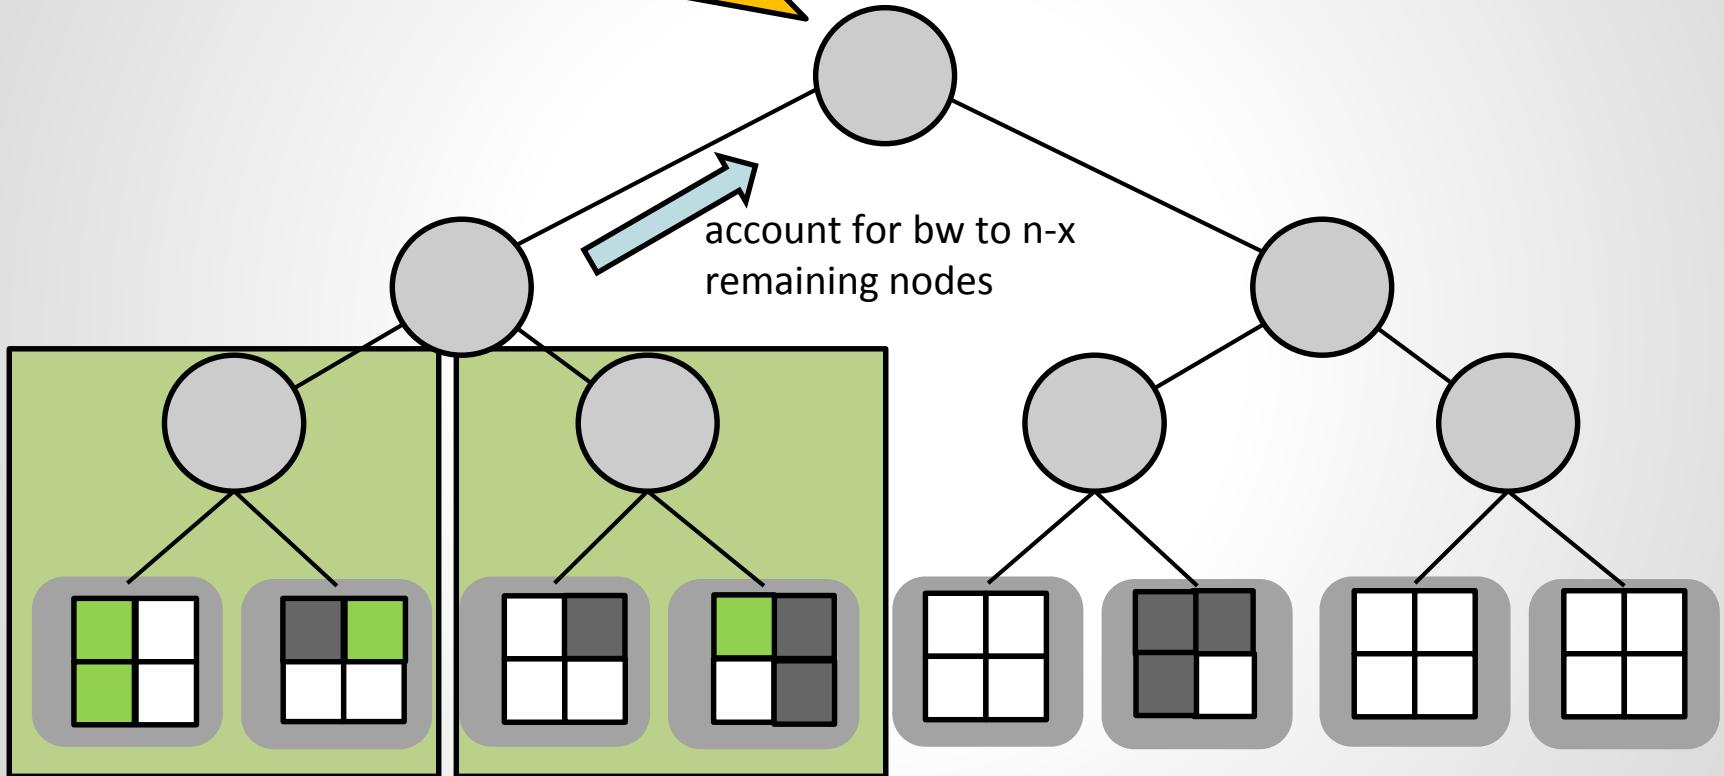
Dynamic Programming



Bottom-up «induction». Leaves easy: either x nodes fit server (cost 0) or not (cost ∞): $\text{opt}[\leq 4] = 0$, $\text{opt}[> 4] = \infty$

Dynamic Programming

$$\text{opt}(T,x) = \min_{0 \leq y \leq x} \{ \text{opt}(\text{left},y) + \text{opt}(\text{right},x-y) + \text{bw}(T,x) \}$$

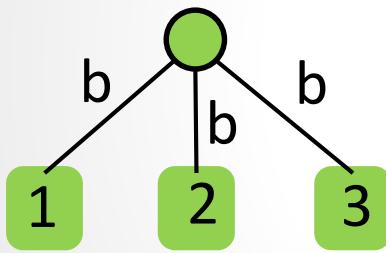


To compute cost of embedding **x nodes** in T , place **y nodes on the left**, **$x-y$ on the right** subtree, and compute cost due to links across root.

Remark on Virtual Cluster Abstraction

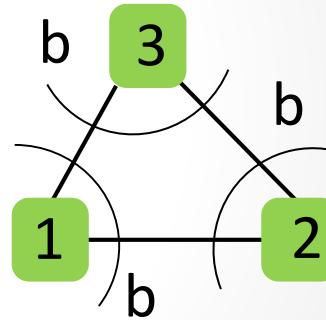
- Two interpretations:

- Logical switch at **unique location**
- Logical switch can be distributed («**Hose model**»)



Logical Switch Variant

vs



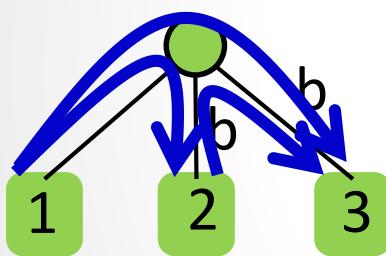
Hose Variant

Aggregated bw in/out
node at most b.

Remark on Virtual Cluster Abstraction

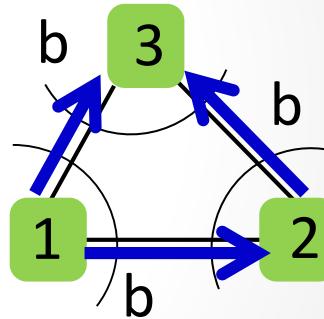
□ Two interpretations:

- Logical switch at **unique location**
- Logical switch can be distributed («**Hose model**»)



Logical Switch Variant

vs



Hose Variant

Can serve the same communication patterns! (A polytope of possible traffic matrices.)

Example:

(1,2): $b/2$

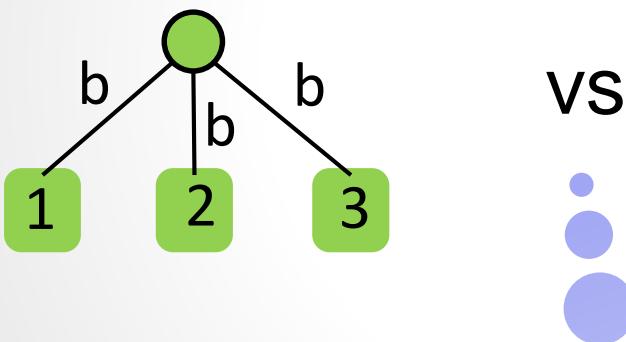
(1,3): $b/2$

(2,3): $b/2$

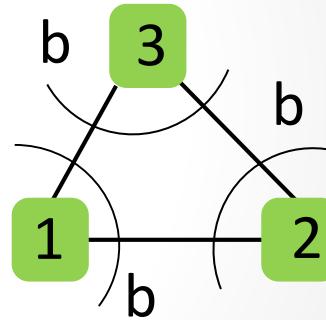
Remark on Virtual Cluster Abstraction

□ Two interpretations:

- Logical switch at **unique location**
- Logical switch can be distributed («**Hose model**»)



vs



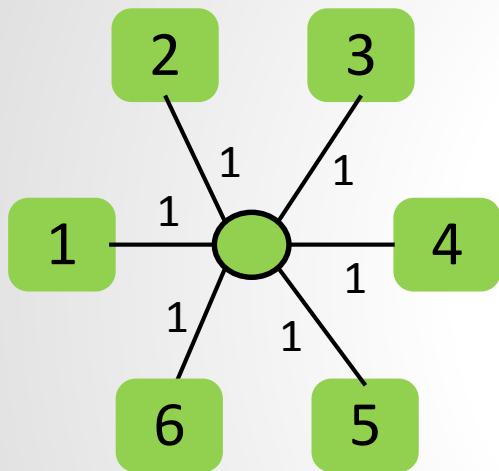
Hose Variant

But **embedding costs** can be different if we do not insist on placing the logical switch explicitly! **Not on trees** though, and **not in uncapacitated networks**: without routing restrictions, optimal routing paths form a tree (**SymG=SymT** a.k.a. **VPN Conjecture**).

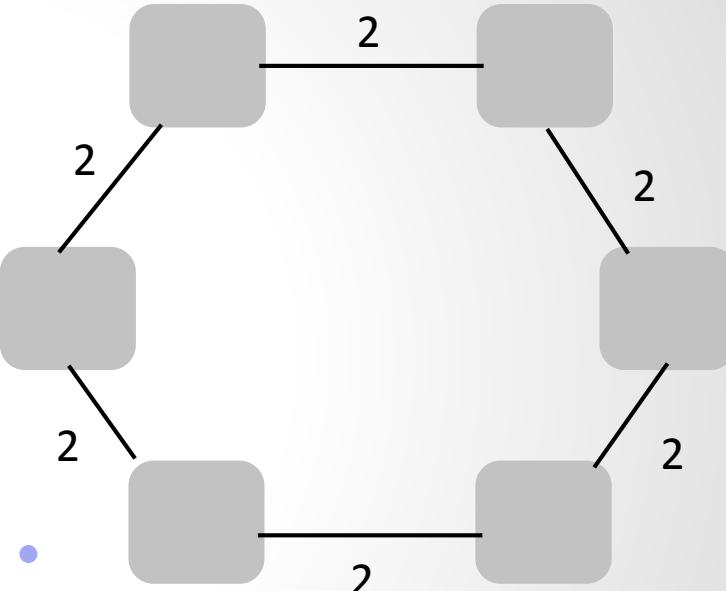
The Benefit of Hose Interpretation

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring

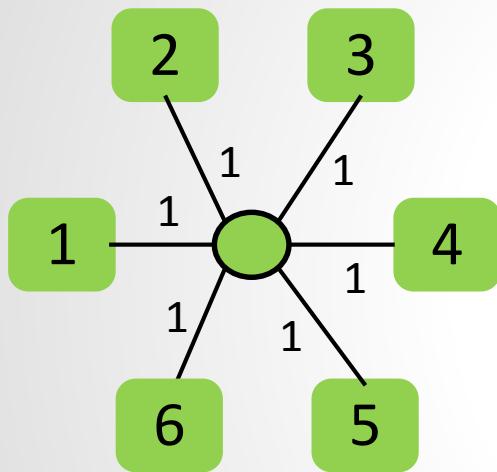


embedding?

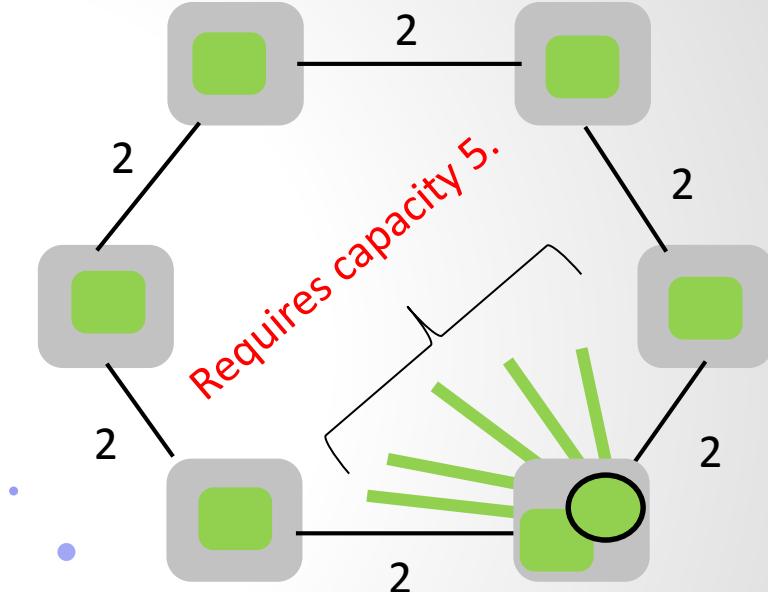
How to embed as a star?

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring



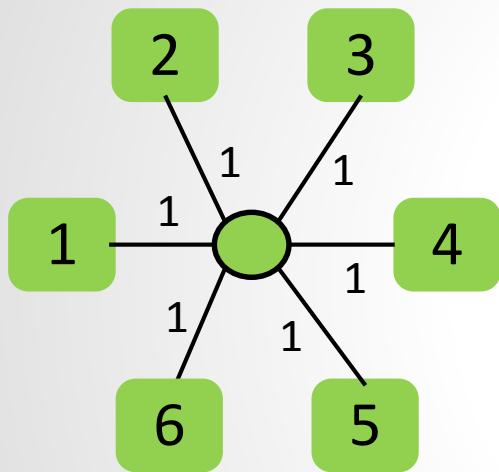
embedding?

How to embed as a star?

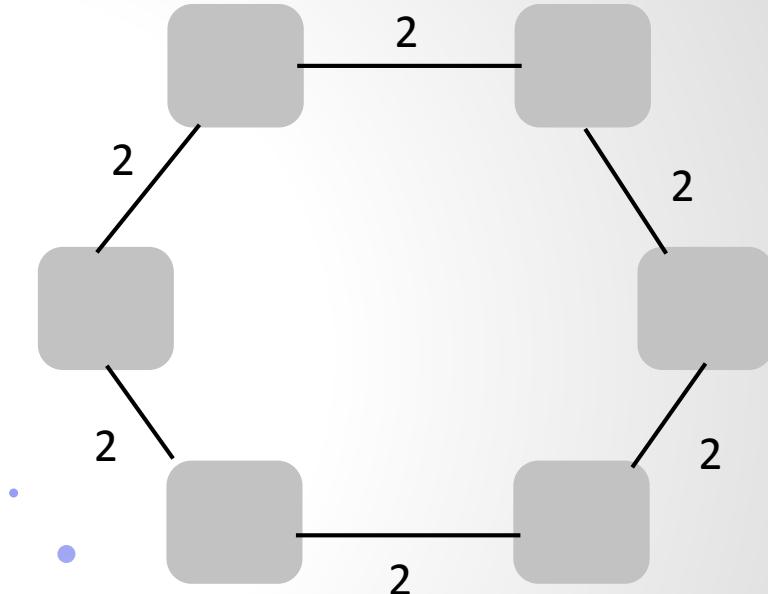
Impossible: need at least 5 units of flow from/to node where *star center is mapped*. However, capacity of incident links is only 4.

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring

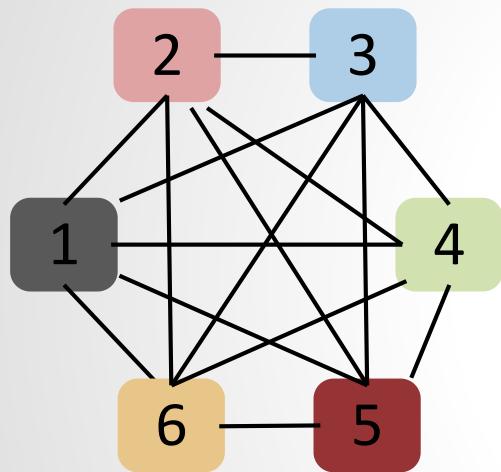


embedding?

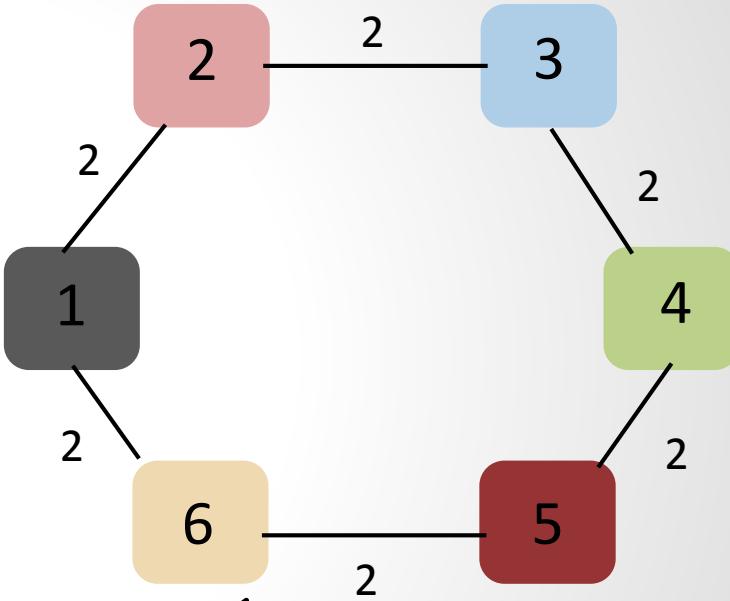
How about hose?

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring

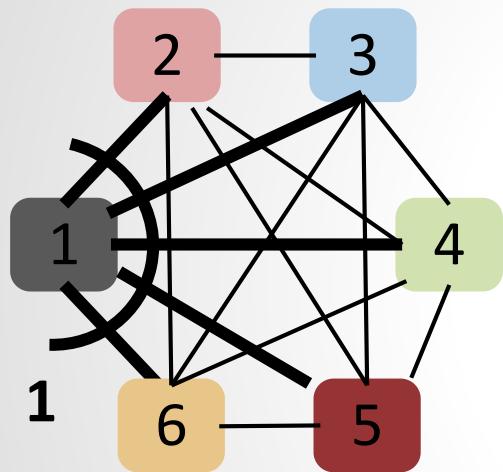


embedding?

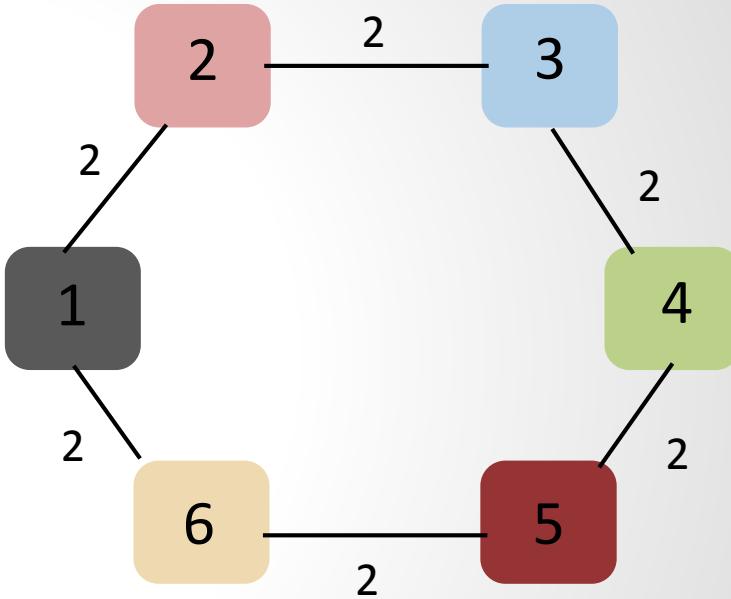
Node mapping! Now: How to
embed these virtual links?

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



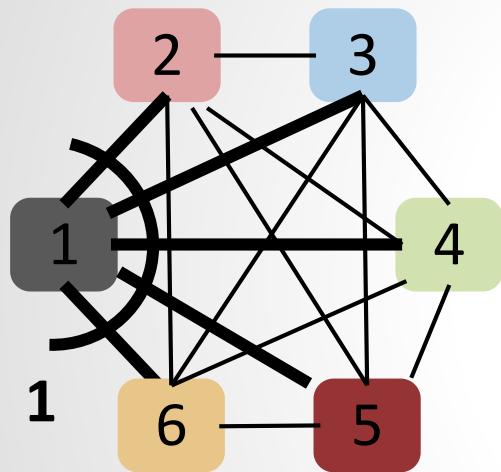
Host Graph: A Ring



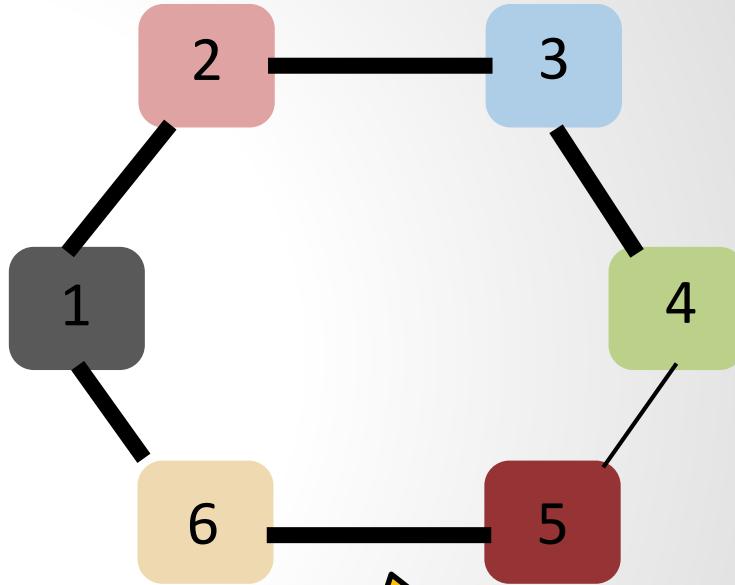
embedding?

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring

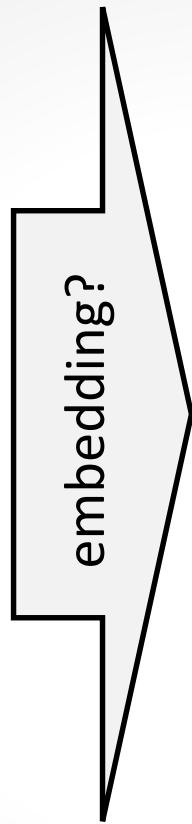
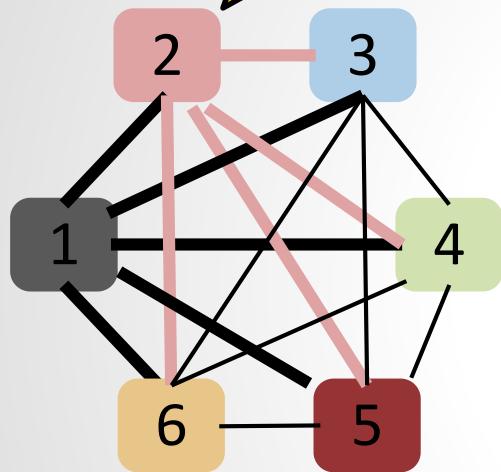


embedding?

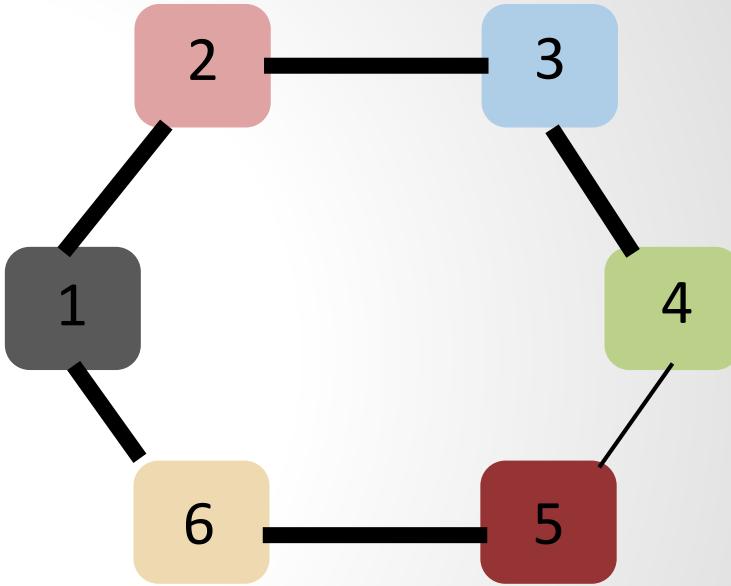
Virtual links from node 1 to $\{2, 3, 4, 5, 6\}$ can be implemented along this route: fulfills capacity constraints under any traffic matrix fulfilling hose specification!

The Benefit of Hose Interpretation

Remaining virtual links to embed for virtual node 2.

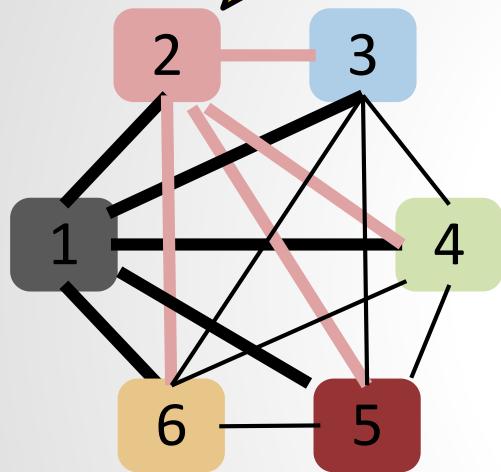


Host Graph: A Ring



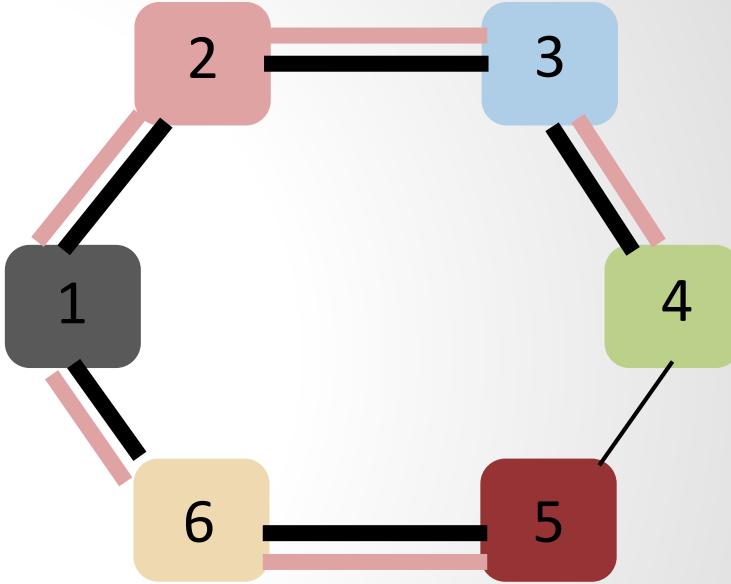
The Benefit of Hose Interpretation

Remaining virtual links to embed for virtual node 2.



embedding?

Host Graph: A Ring



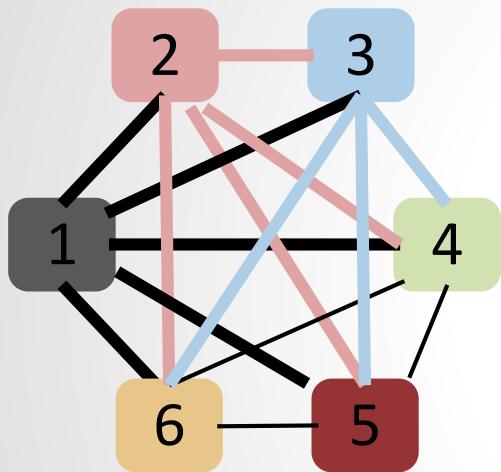
Can be implemented along this route: from node 2, reach nodes {3,4,5,6}.

The Benefit of Hose Interpretation

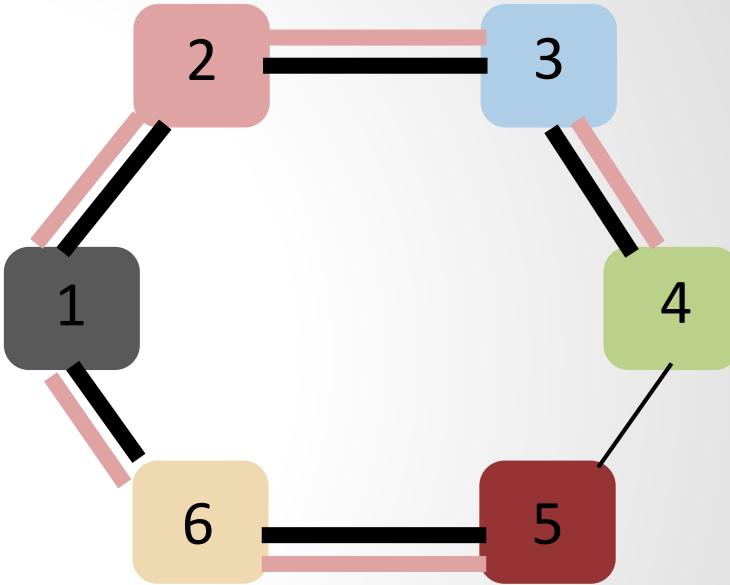
VNet: VC($n=6, b=1$)

Remaining virtual links to embed for virtual node 3.

Host Graph: A Ring



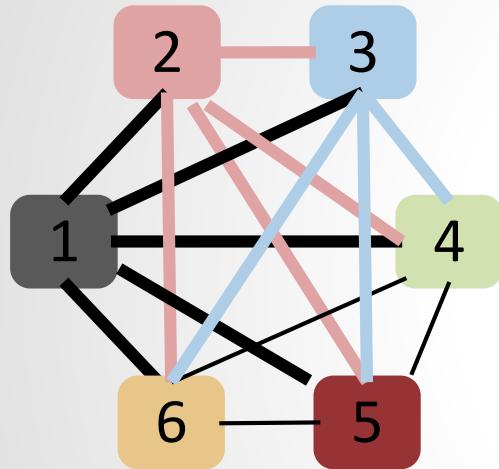
embedding?



The Benefit of Hose Interpretation

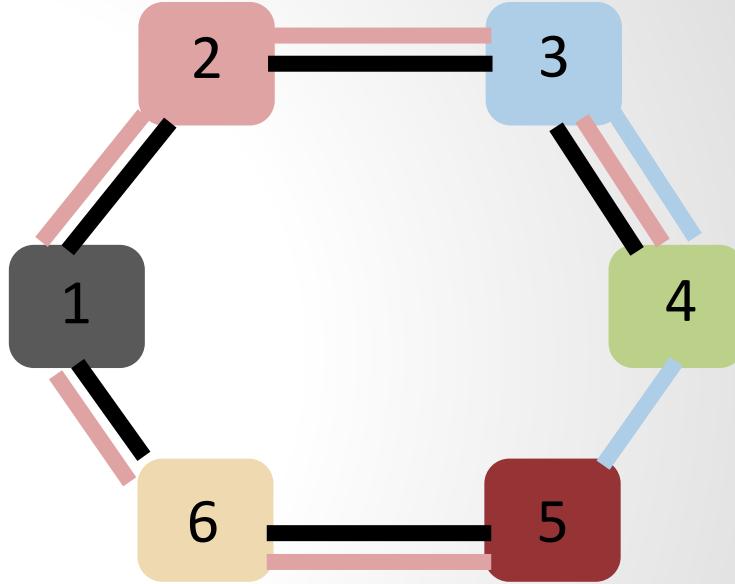
VNet: VC($n=6, b=1$)

Remaining virtual links to embed for virtual node 3.



Host Graph: A Ring

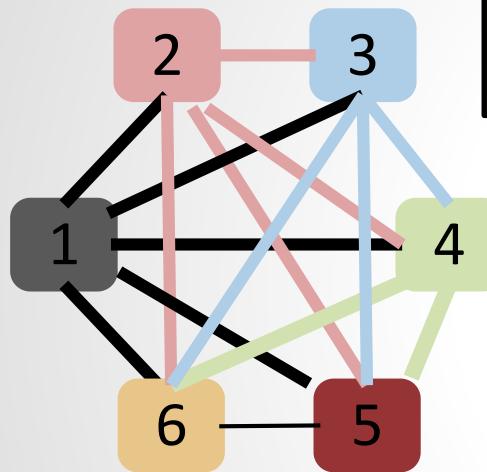
embedding?



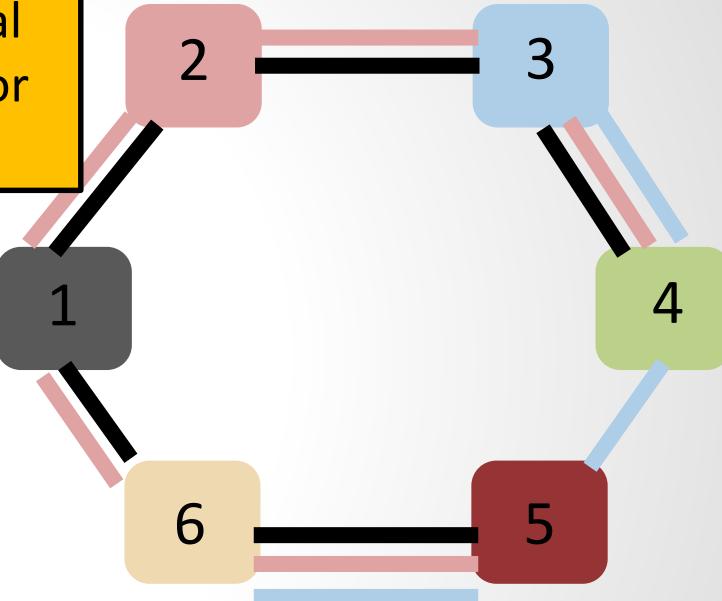
From 3, path reaches {4,5,6}.

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring

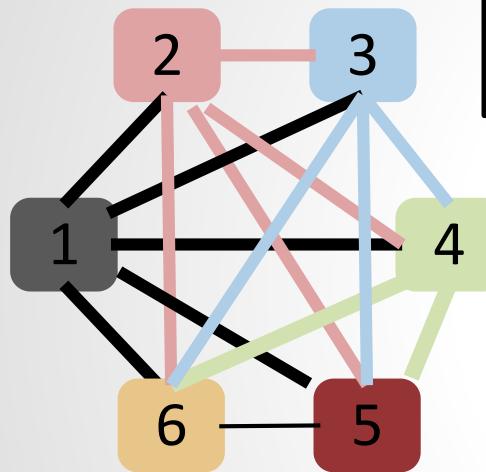


*Remaining virtual
links to embed for
virtual node 4.*

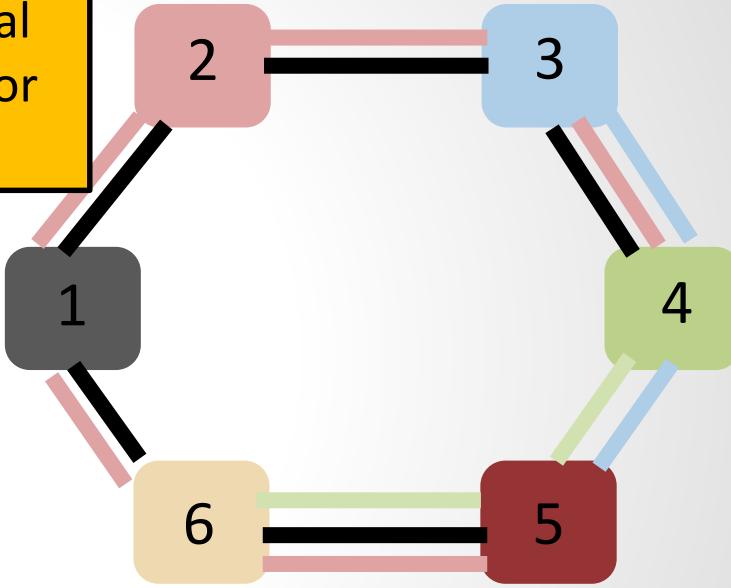
embedding

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring



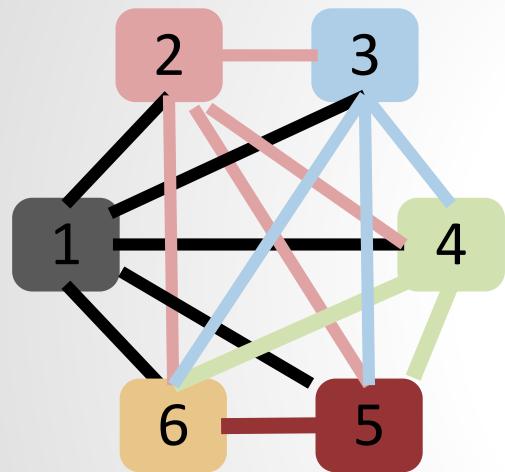
*Remaining virtual
links to embed for
virtual node 4.*

embedding

Route from 4 to {5,6}.

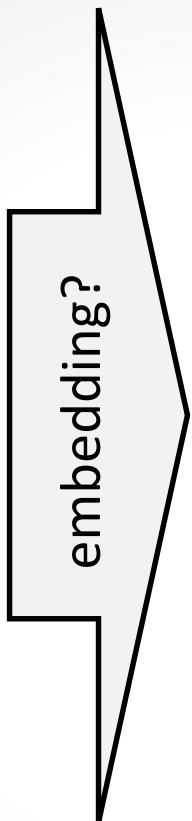
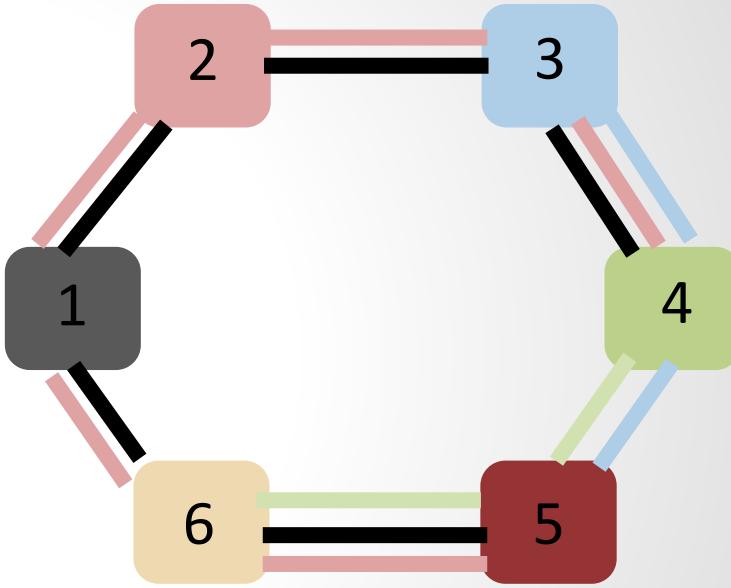
The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



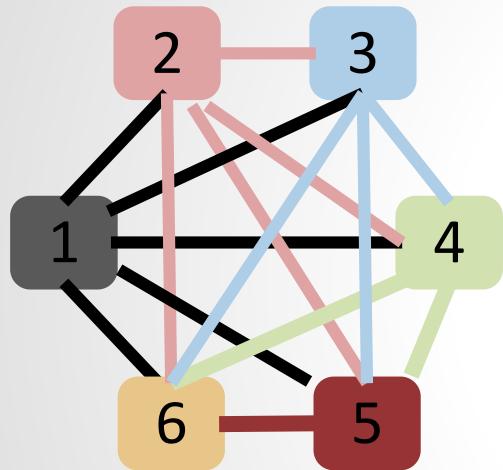
*Remaining
virtual link.*

Host Graph: A Ring

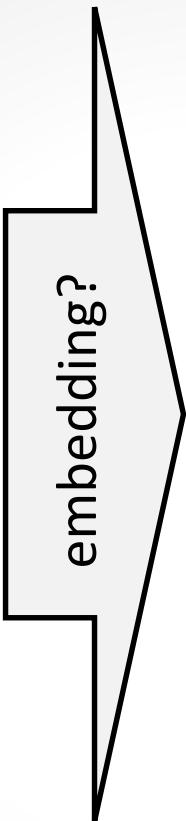
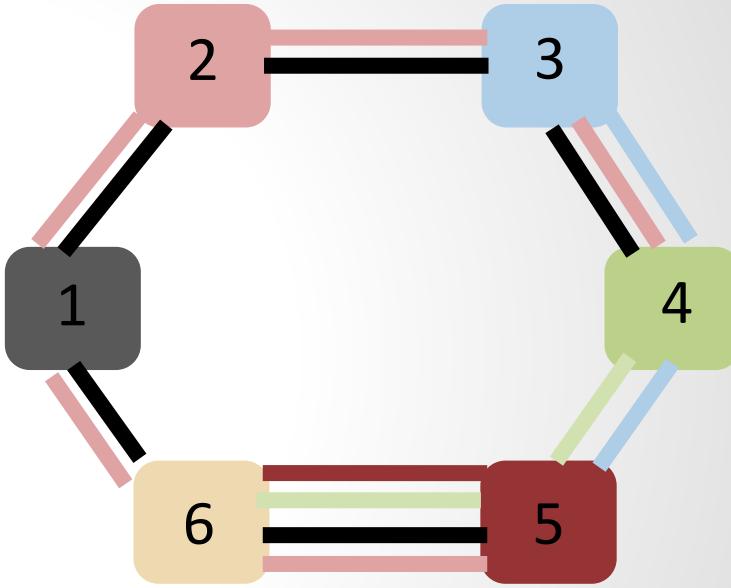


The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



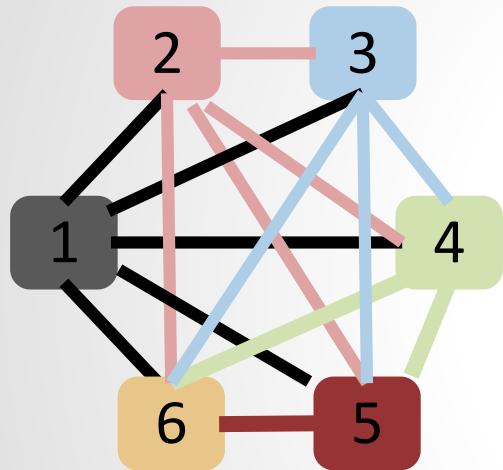
Host Graph: A Ring



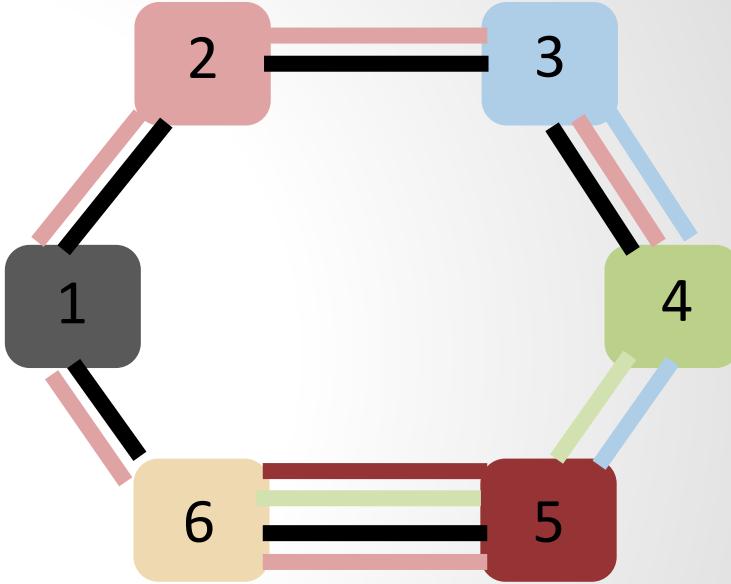
All virtual links mapped to routes!

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring

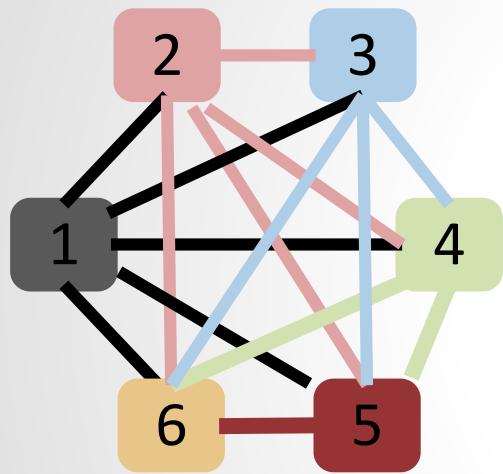


embedding?

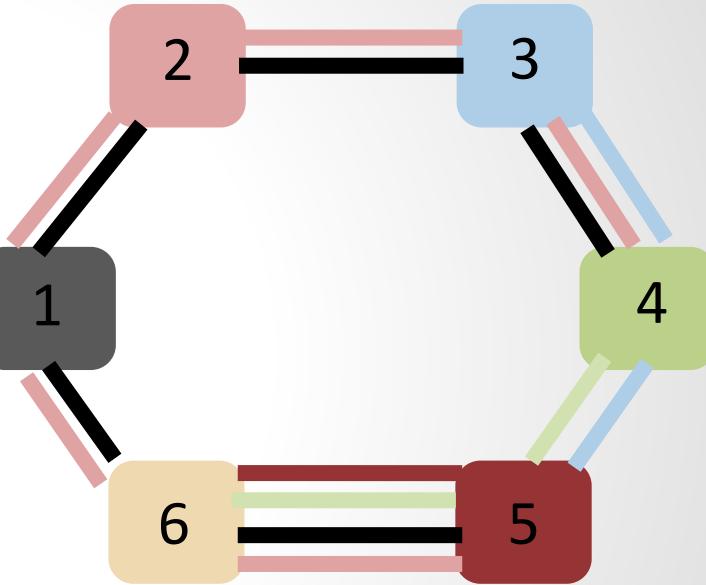
But wait: 5 paths on link {5,6}!
Can demand really be satisfied
given link capacity of 2?!

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring



embedding?

Link $\{5,6\}$ is used by routes: $(1,5), (2,5), (3,6), (4,6), (5,6)$.

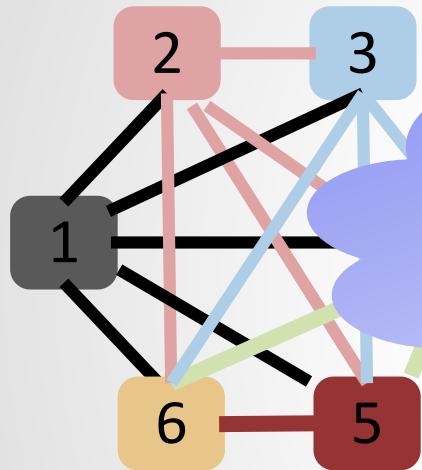
But by definition of the hose model, any traffic matrix M will respect:

$$M_{1,5} + M_{2,5} \leq 1 \text{ and } M_{3,6} + M_{4,6} + M_{5,6} \leq 1.$$

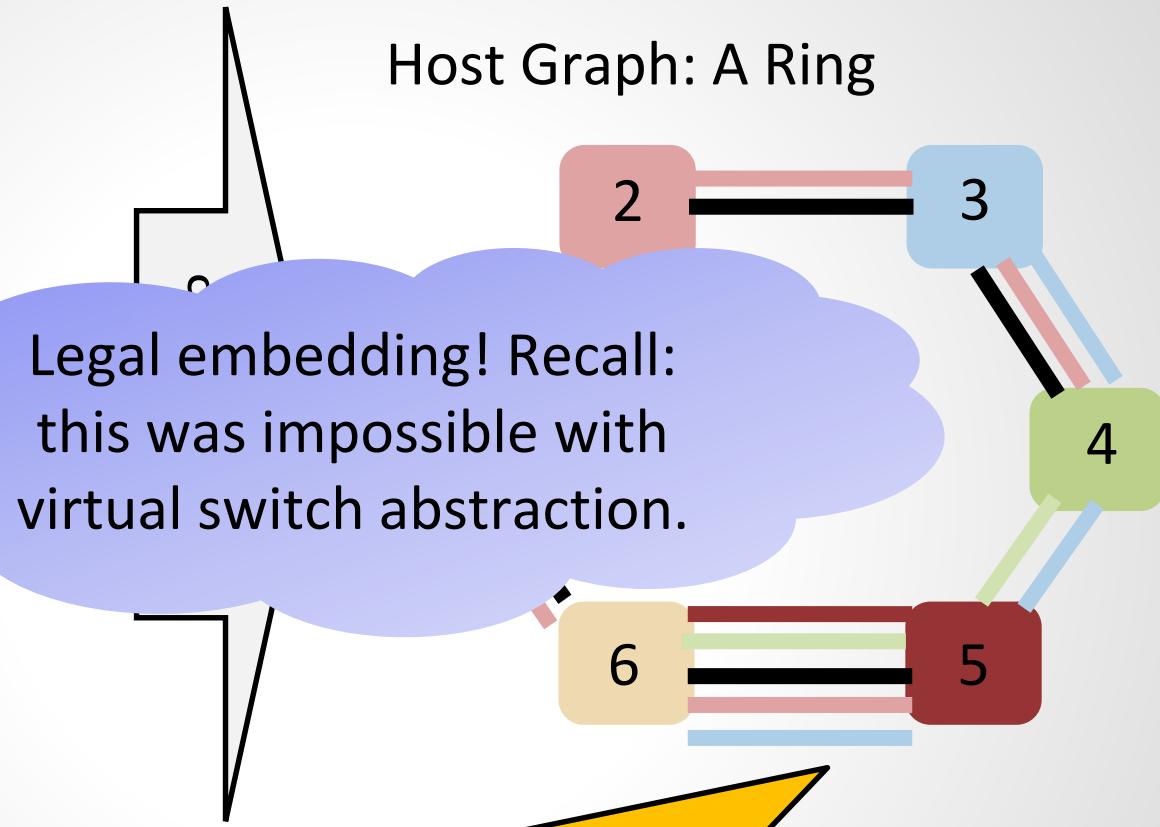
Hence $\sum M_{i,j} \leq 2$ holds!

The Benefit of Hose Interpretation

VNet: VC($n=6, b=1$)



Host Graph: A Ring



Legal embedding! Recall:
this was impossible with
virtual switch abstraction.

Link $\{5,6\}$ is used by routes: $(1,5), (2,5), (3,6), (4,6), (5,6)$.

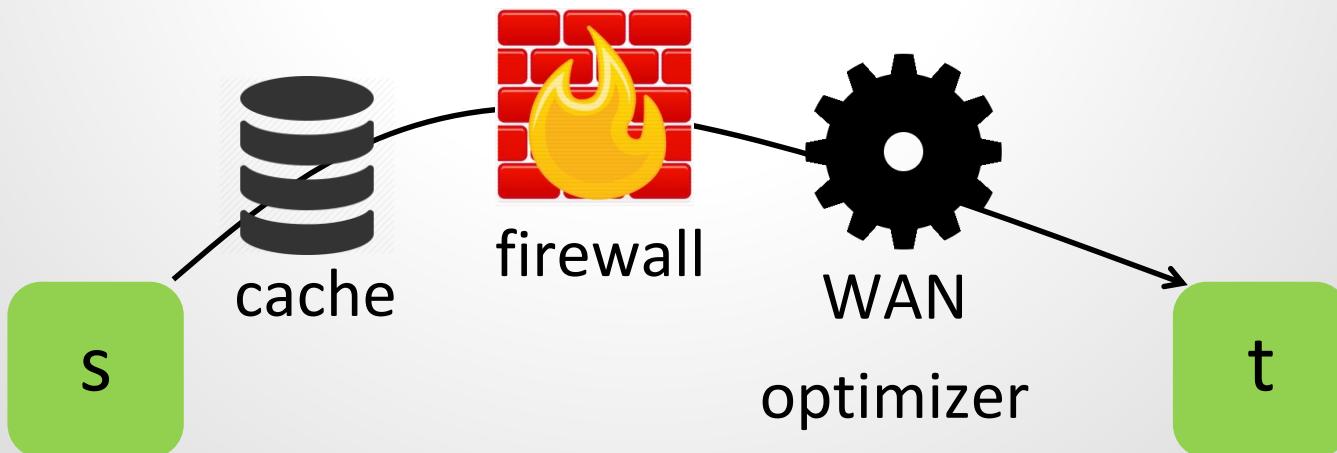
But by definition of the hose model, any traffic matrix M will respect:

$$M_{1,5} + M_{2,5} \leq 1 \text{ and } M_{3,6} + M_{4,6} + M_{5,6} \leq 1.$$

Hence $\sum M_{i,j} \leq 2$ holds!

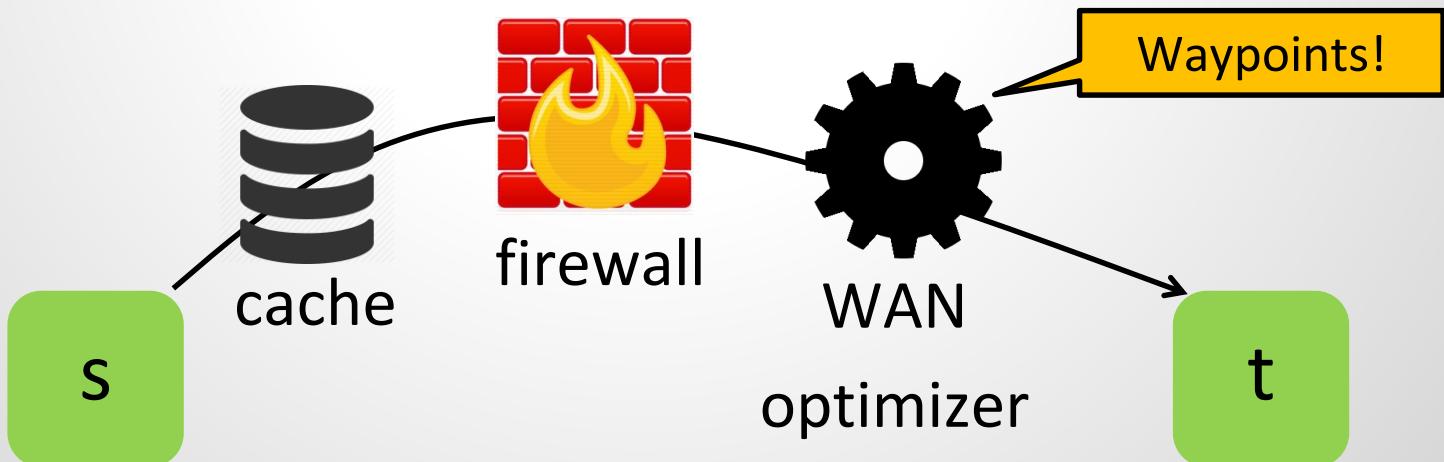
The Many Faces of the VNEP: E.g., Service Chain Embedding

- Similar problems arise in many contexts
- For example, **service chain** embeddings: in a service chain, traffic is steered (e.g., using SDN) through a **sequence of (virtualized) middleboxes** to compose a more complex network service



The Many Faces of the VNEP: E.g., Service Chain Embedding

- Similar problems arise in many contexts
- For example, **service chain** embeddings: in a service chain, traffic is steered (e.g., using SDN) through a **sequence of (virtualized) middleboxes** to compose a more complex network service



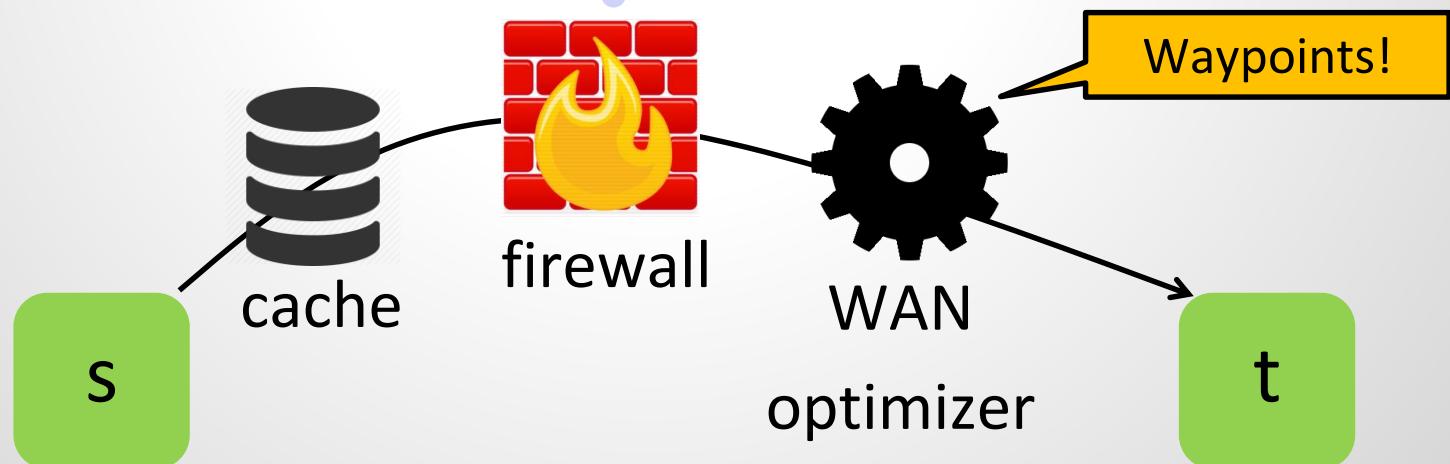
The Many Faces of the VNEP: E.g., Service Chain Embedding

- Similar problems arise in many contexts

Interesting implication: routes from s to t become *walks* (rather than *simple paths*)! How to find shortest walks?

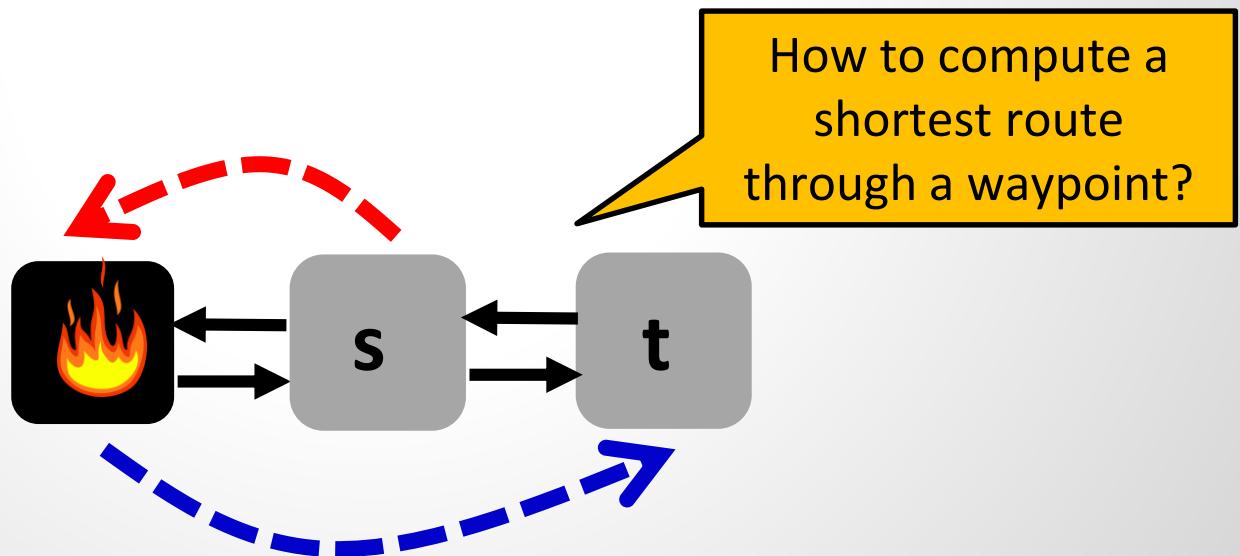
more complex ... service

service
through a
series of nodes
to compose a



Routing Through Waypoints

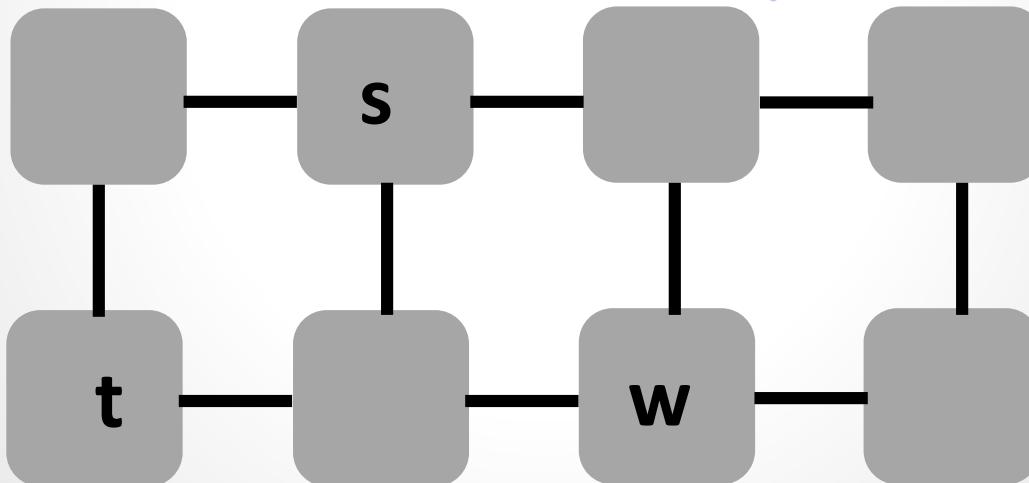
- ❑ Traditionally: routes form **simple paths** (e.g., shortest paths)
- ❑ Novel aspect: routing through **middleboxes** may require more general paths, with loops: **a walk**



Walking Through Waypoints

- Computing shortest routes through waypoints is **non-trivial!**

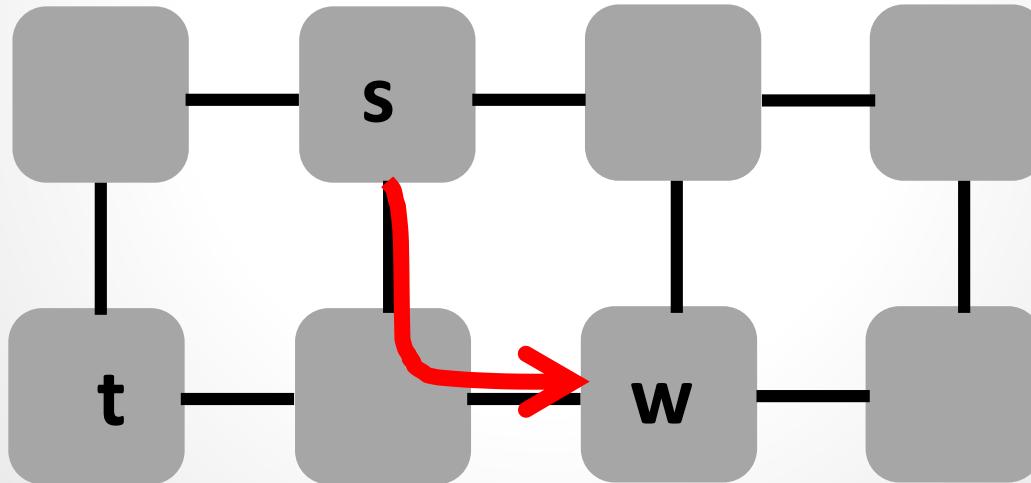
Assume unit capacity and demand for simplicity!



Walking Through Waypoints

- Computing shortest routes through waypoints is **non-trivial!**

Assume unit capacity and demand for simplicity!

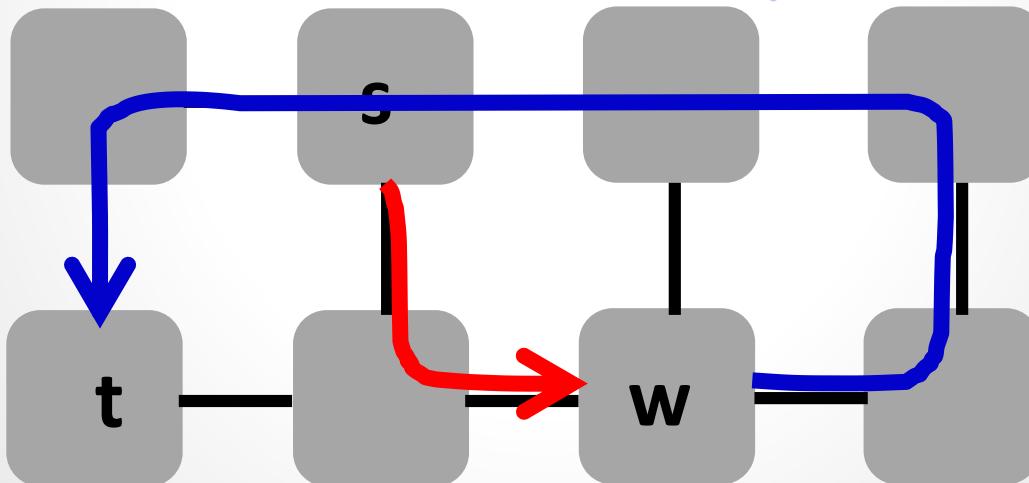


Greedy fails: choose shortest path from s to w ...

Walking Through Waypoints

- Computing shortest routes through waypoints is **non-trivial!**

Assume unit capacity and demand for simplicity!

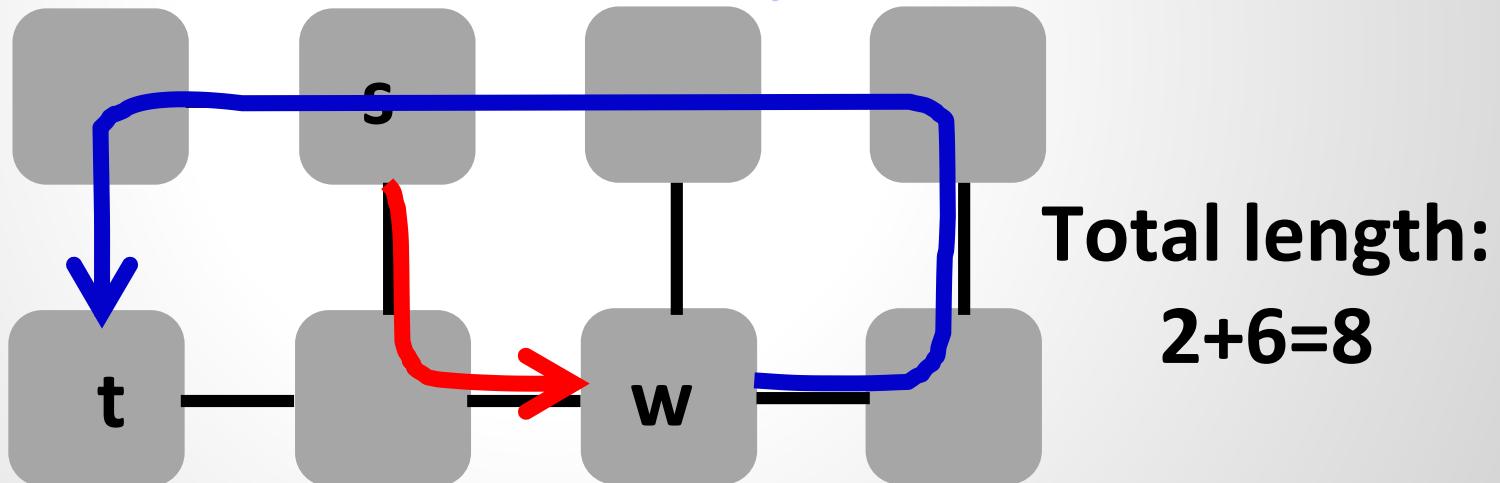


Greedy fails: ... now need long path from w to t

Walking Through Waypoints

- Computing shortest routes through waypoints is **non-trivial!**

Assume unit capacity and demand for simplicity!

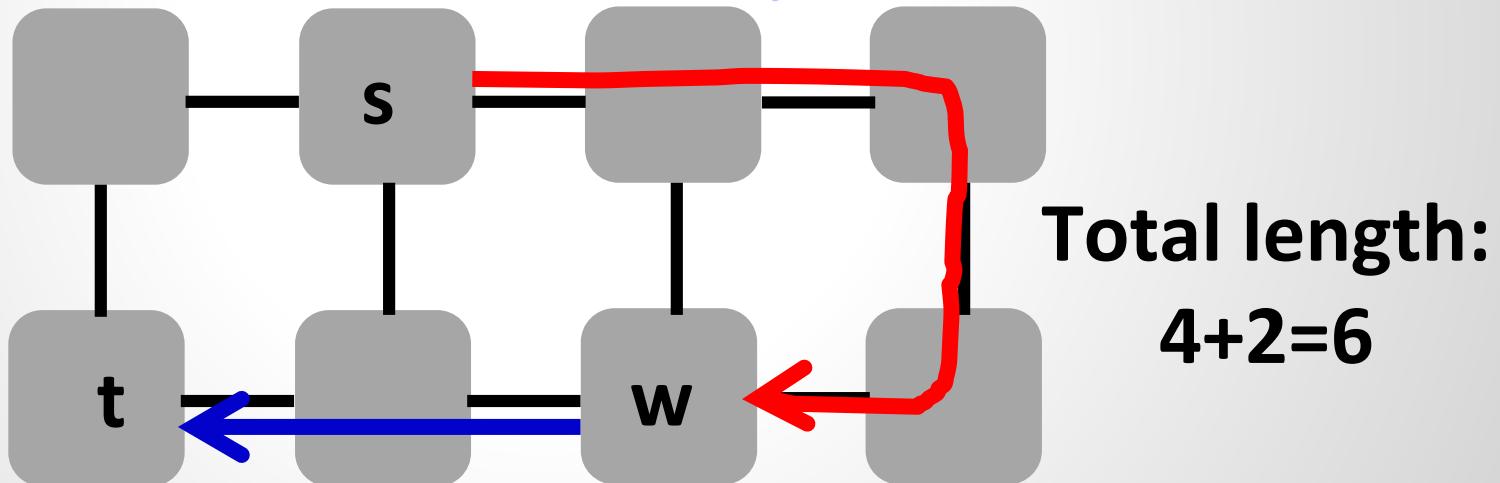


Greedy fails: ... now need long path from w to t

Walking Through Waypoints

- Computing shortest routes through waypoints is **non-trivial!**

Assume unit capacity and demand for simplicity!

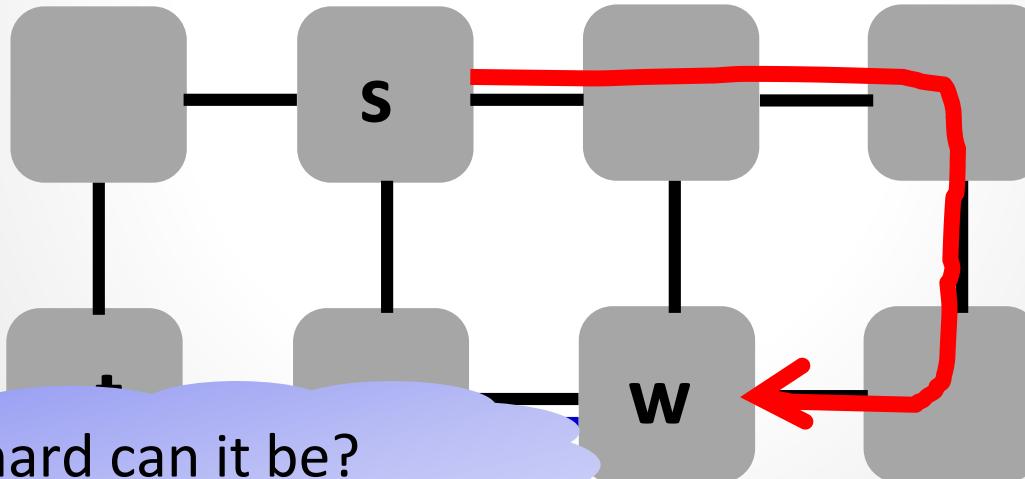


A **better solution**: jointly optimize the two segments!

Walking Through Waypoints

- Computing shortest routes through waypoints is **non-trivial!**

Assume unit capacity and demand for simplicity!

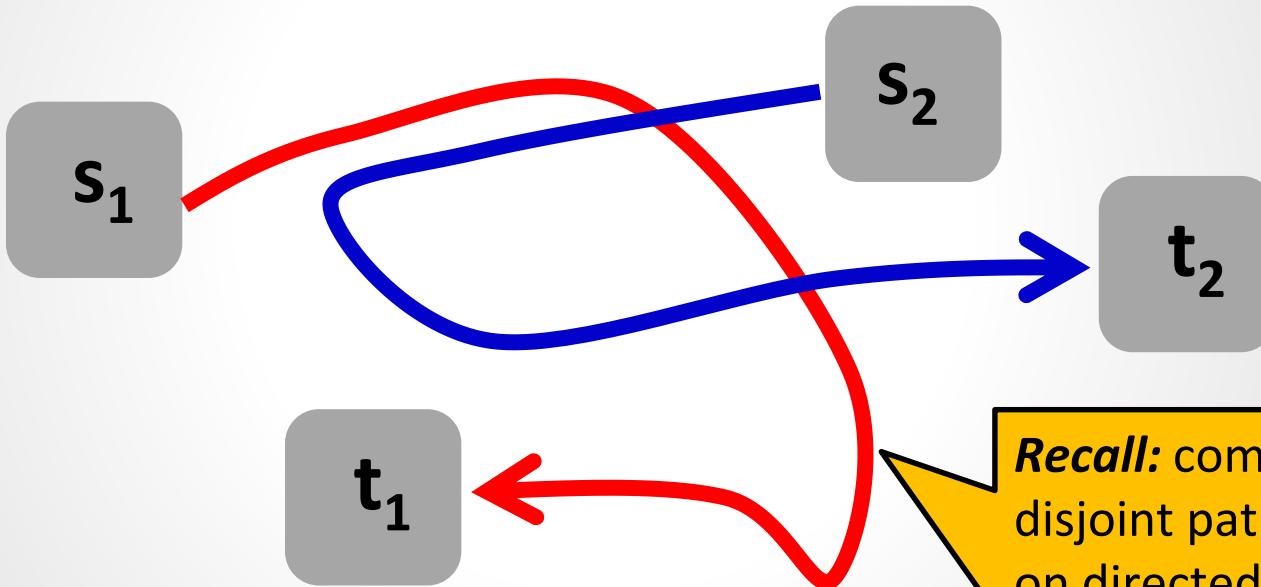


How hard can it be?

A **better solution**: *jointly* optimize the two segments!

NP-hard on Directed Networks: Reduction from Disjoint Paths Problem

Reduction: From joint shortest paths $(s_1, t_1), (s_2, t_2)$
to shortest walk (s, w, t) problem



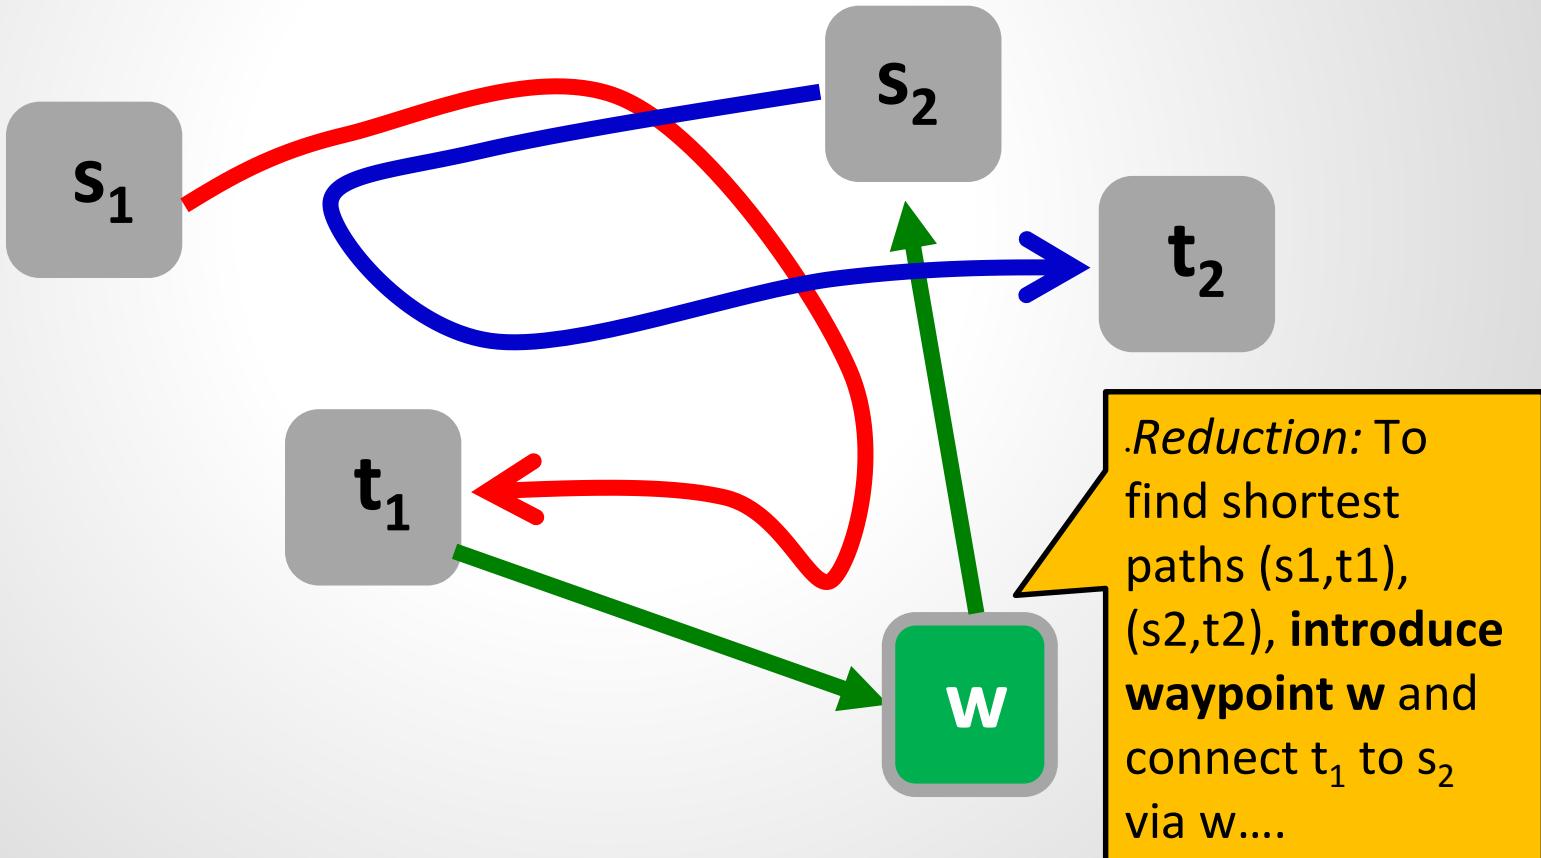
Recall: computing 2-disjoint paths NP-hard on directed graphs.

We show: If waypoint routing was be in P, we could solve it fast.

Contradiction!

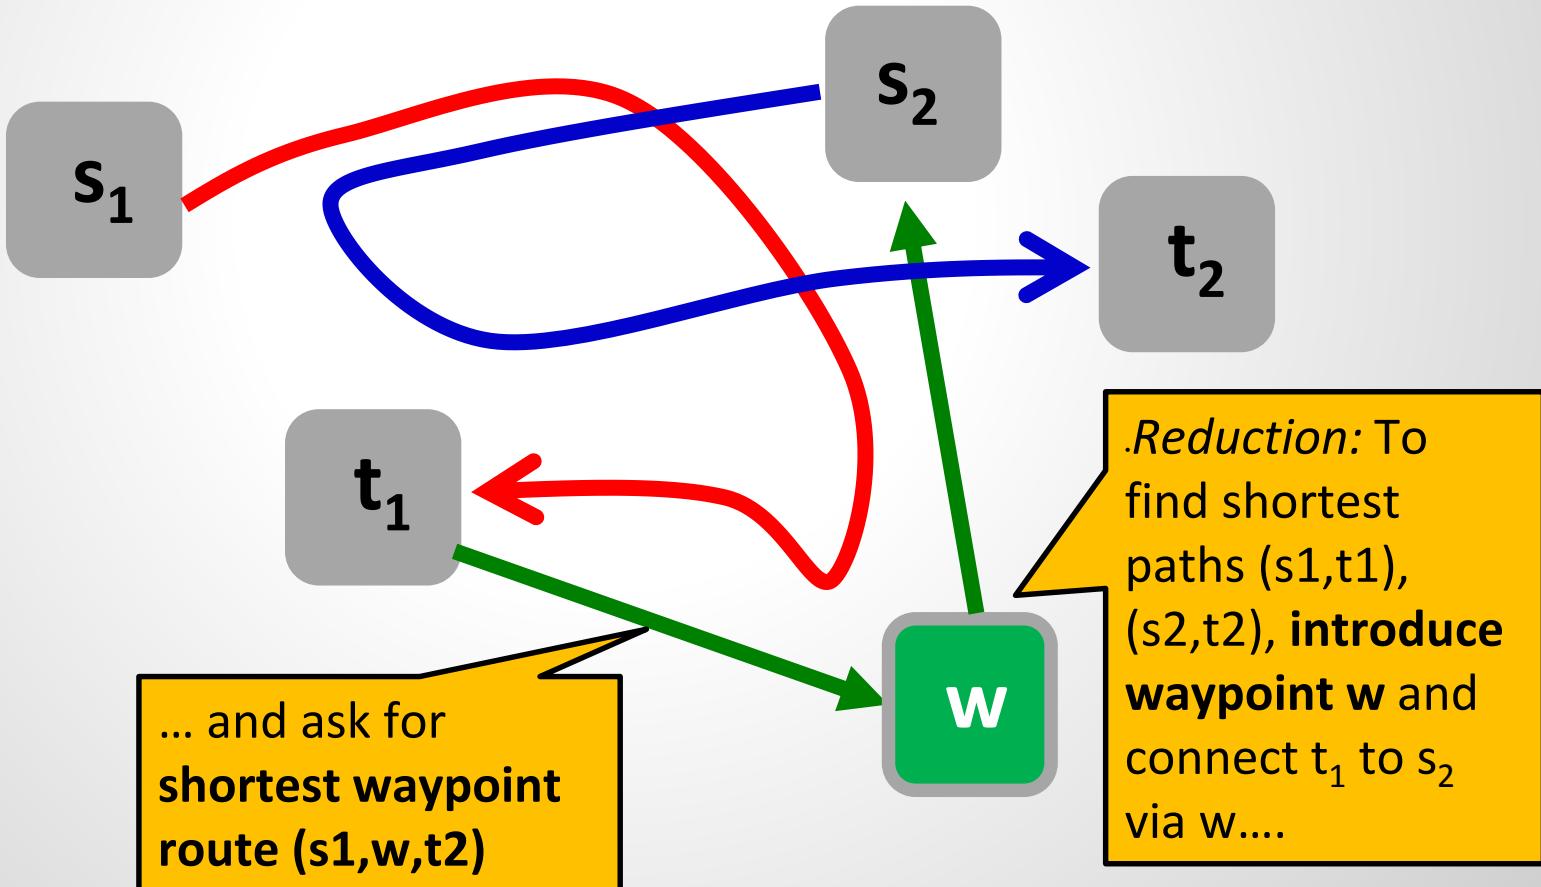
NP-hard on Directed Networks: Reduction from Disjoint Paths Problem

Reduction: From joint shortest paths $(s_1, t_1), (s_2, t_2)$
to shortest walk (s, w, t) problem



NP-hard on Directed Networks: Reduction from Disjoint Paths Problem

Reduction: From joint shortest paths $(s_1, t_1), (s_2, t_2)$
to shortest walk (s, w, t) problem

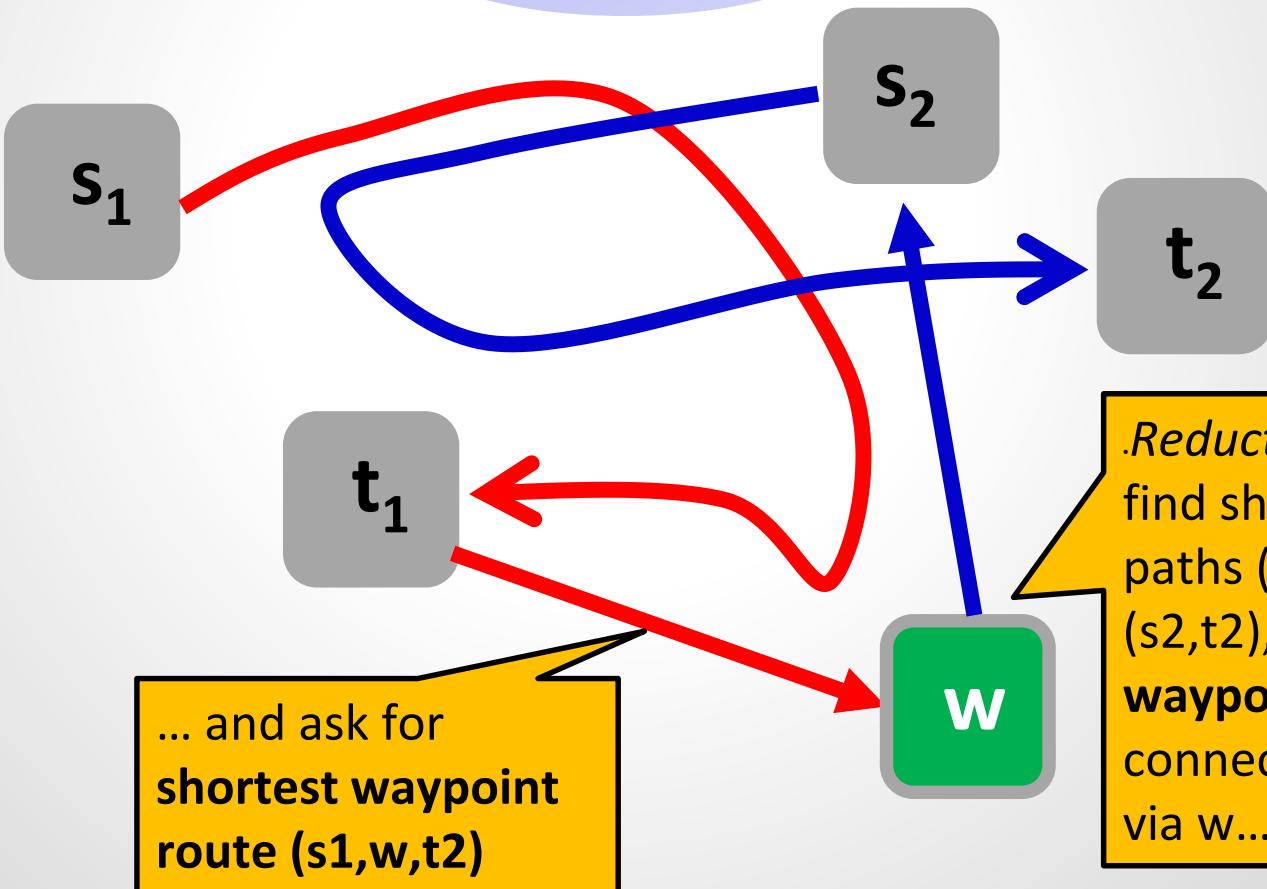


NP-hard on Directed Distances

F The walk (s_1, w, t_2) walk defines a (s_1, t_1) and a (s_2, t_2) path pair before/after the waypoint! Solves original problem: Contradiction!

Reduc

to shortest w.



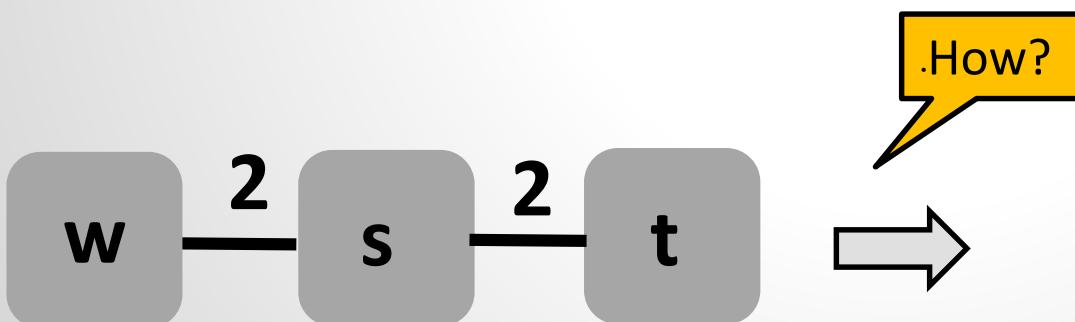
**What about waypoint routes on
undirected networks?**

What about waypoint routes on *undirected* networks?

- Reduction from disjoint paths no longer works: disjoint paths problem **not NP-hard** on undirected networks

What about waypoint routes on *undirected* networks?

- Reduction from disjoint paths no longer works: disjoint paths problem **not NP-hard** on undirected networks
- Indeed, **algorithm exists**: We can reduce **to** edge-disjoint paths to compute a waypoint route!



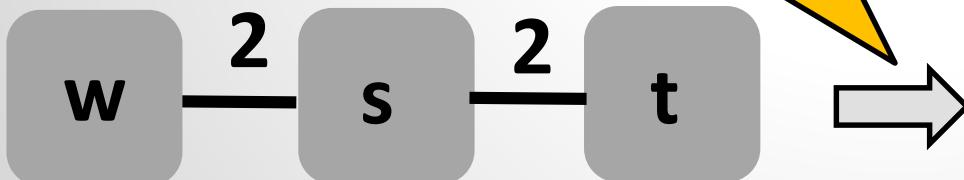
Walks

Edge-Disjoint Paths

What about waypoint routes on *undirected* networks?

- Reduction from disjoint paths no longer works: disjoint paths problem **not NP-hard** on undirected networks
- Indeed, **algorithm exists**: We can reduce **to** edge-disjoint paths ~~to compute a waypoint route!~~

.Replace capacitated
links with parallel links!



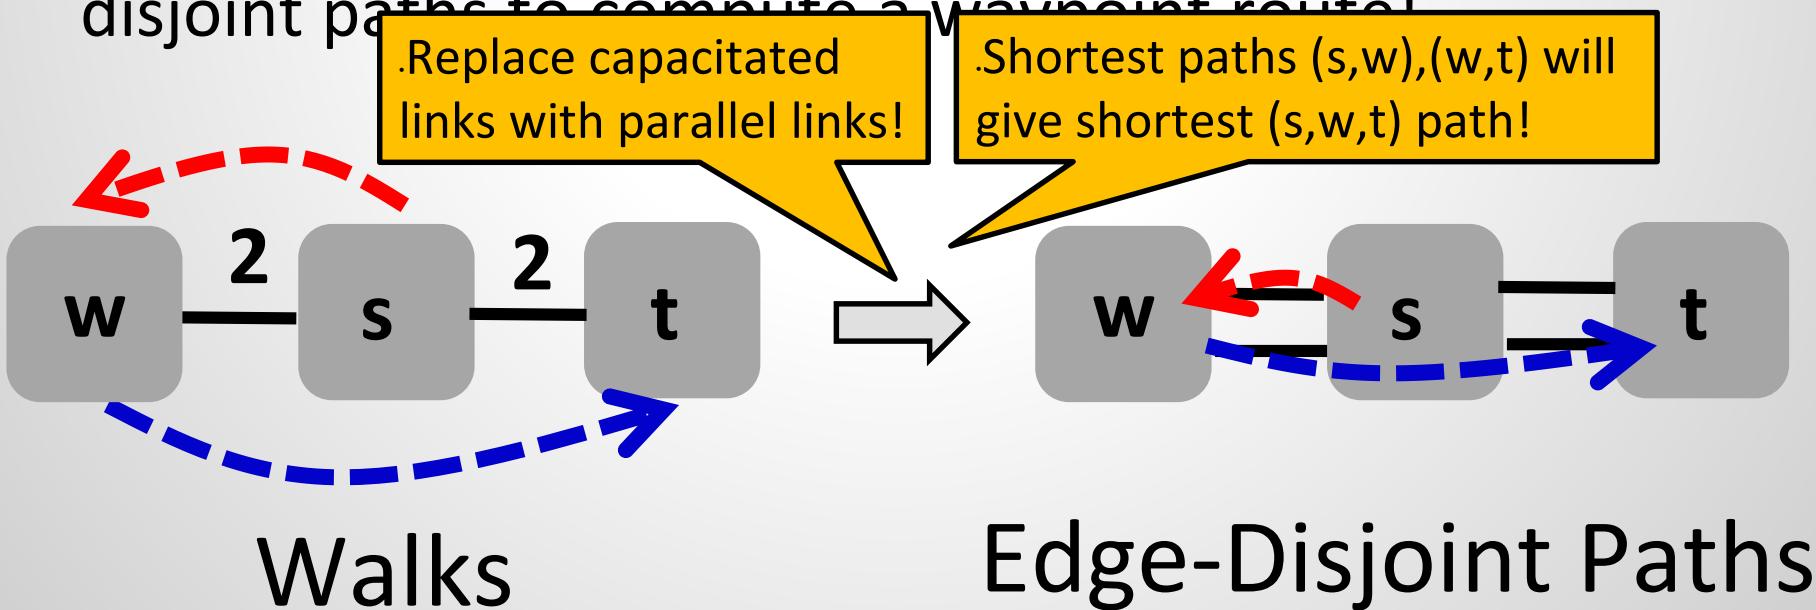
Walks

Edge-Disjoint Paths

What about waypoint routes on *undirected* networks?

- Reduction from disjoint paths no longer works: disjoint paths problem **not NP-hard** on undirected networks

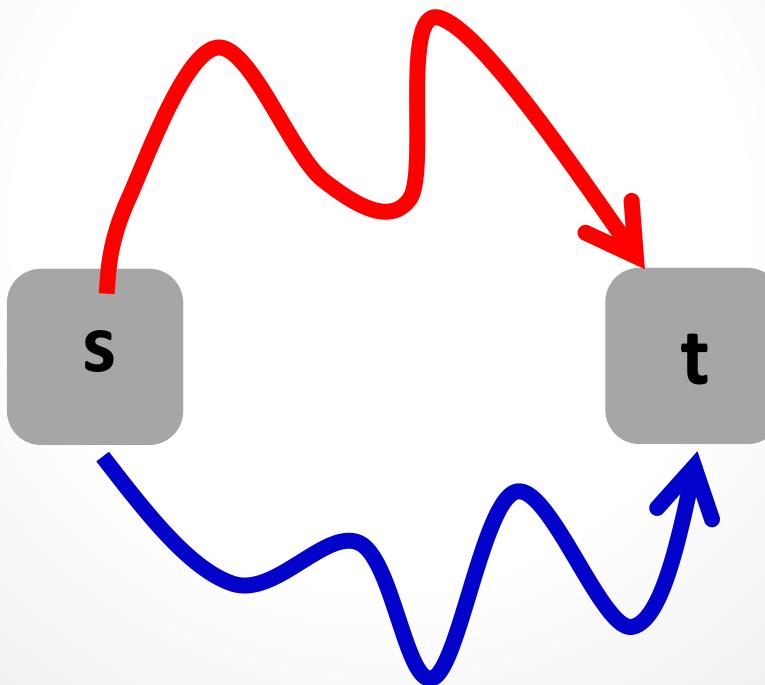
- Indeed, **algorithm exists**: We can reduce to edge-disjoint paths to compute a waypoint route!



Fast and Shortest Waypoint Routing on Undirected Networks: Suurballe's Algorithm

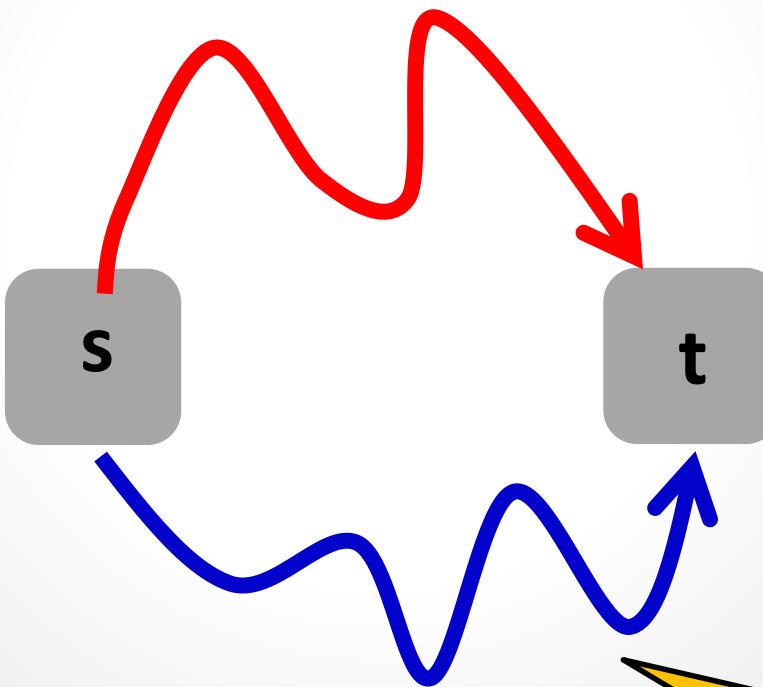
Fast and Shortest Waypoint Routing on Undirected Networks: Suurballe's Algorithm

- Suurballe's algorithm: finds two (edge-)disjoint shortest paths **between same endpoints**:



Fast and Shortest Waypoint Routing on Undirected Networks: Suurballe's Algorithm

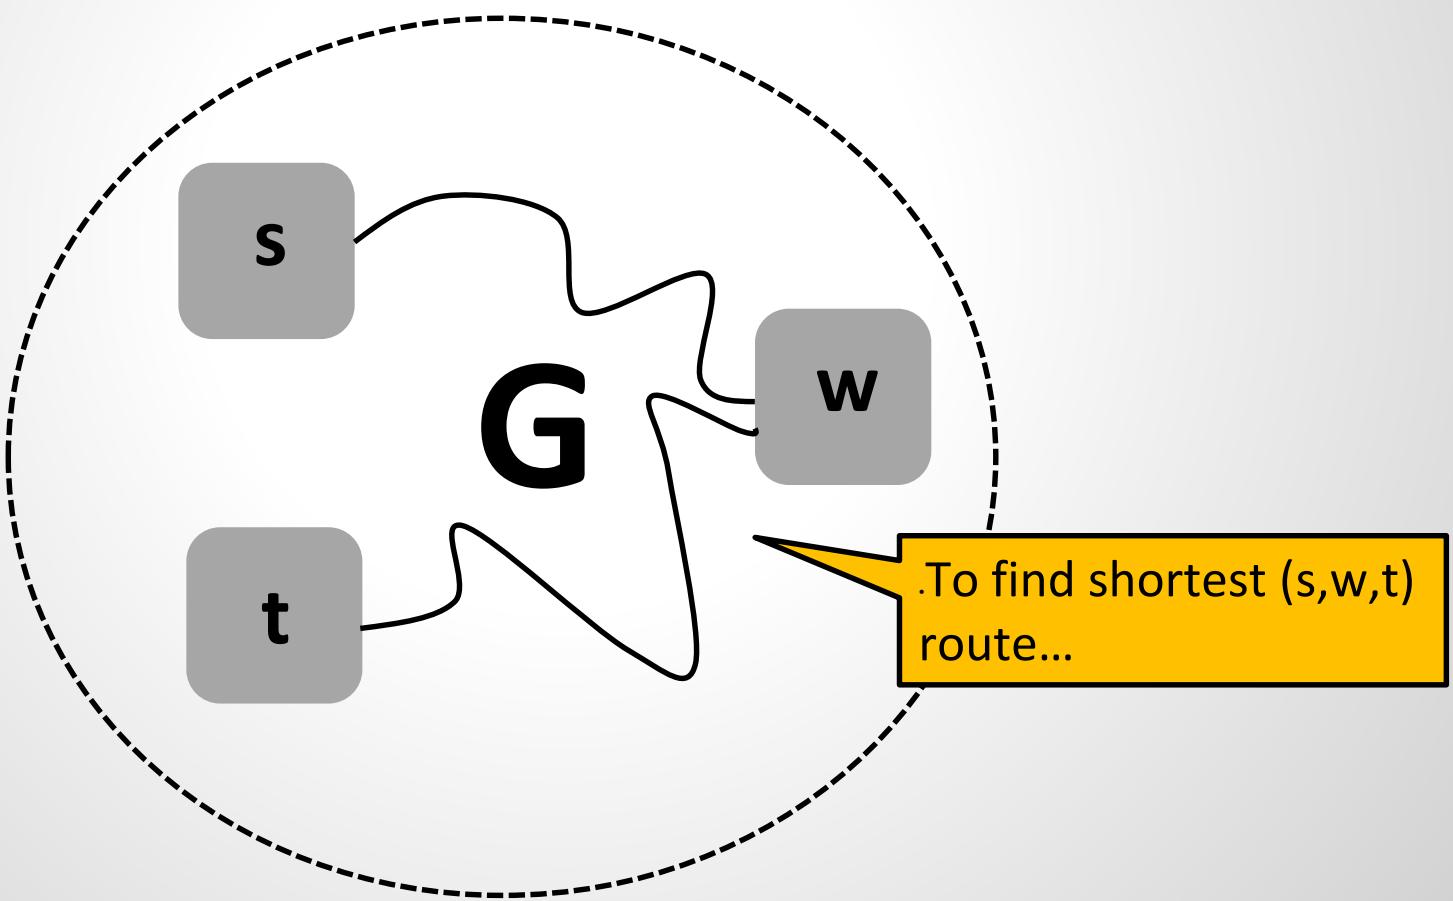
- Suurballe's algorithm: finds two (edge-)disjoint shortest paths **between same endpoints**:



How to compute a shortest (s,w,t) route with this algorithm??

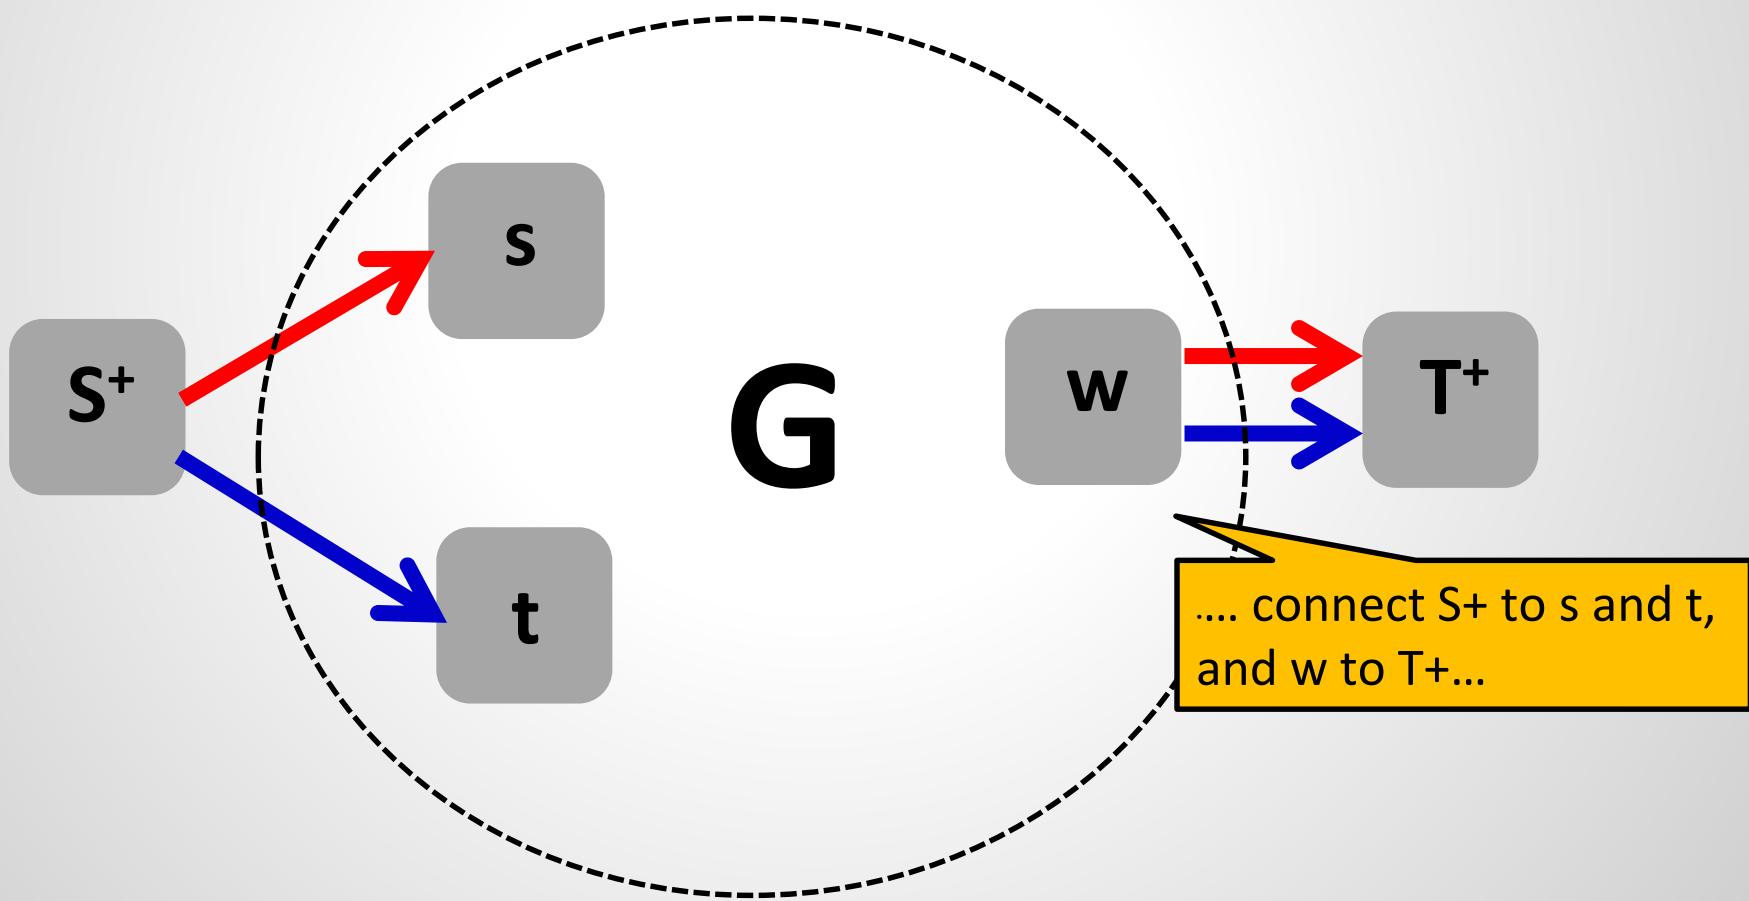
Waypoint Routing on Steroids

- Reduction to Suurballe's algorithm:



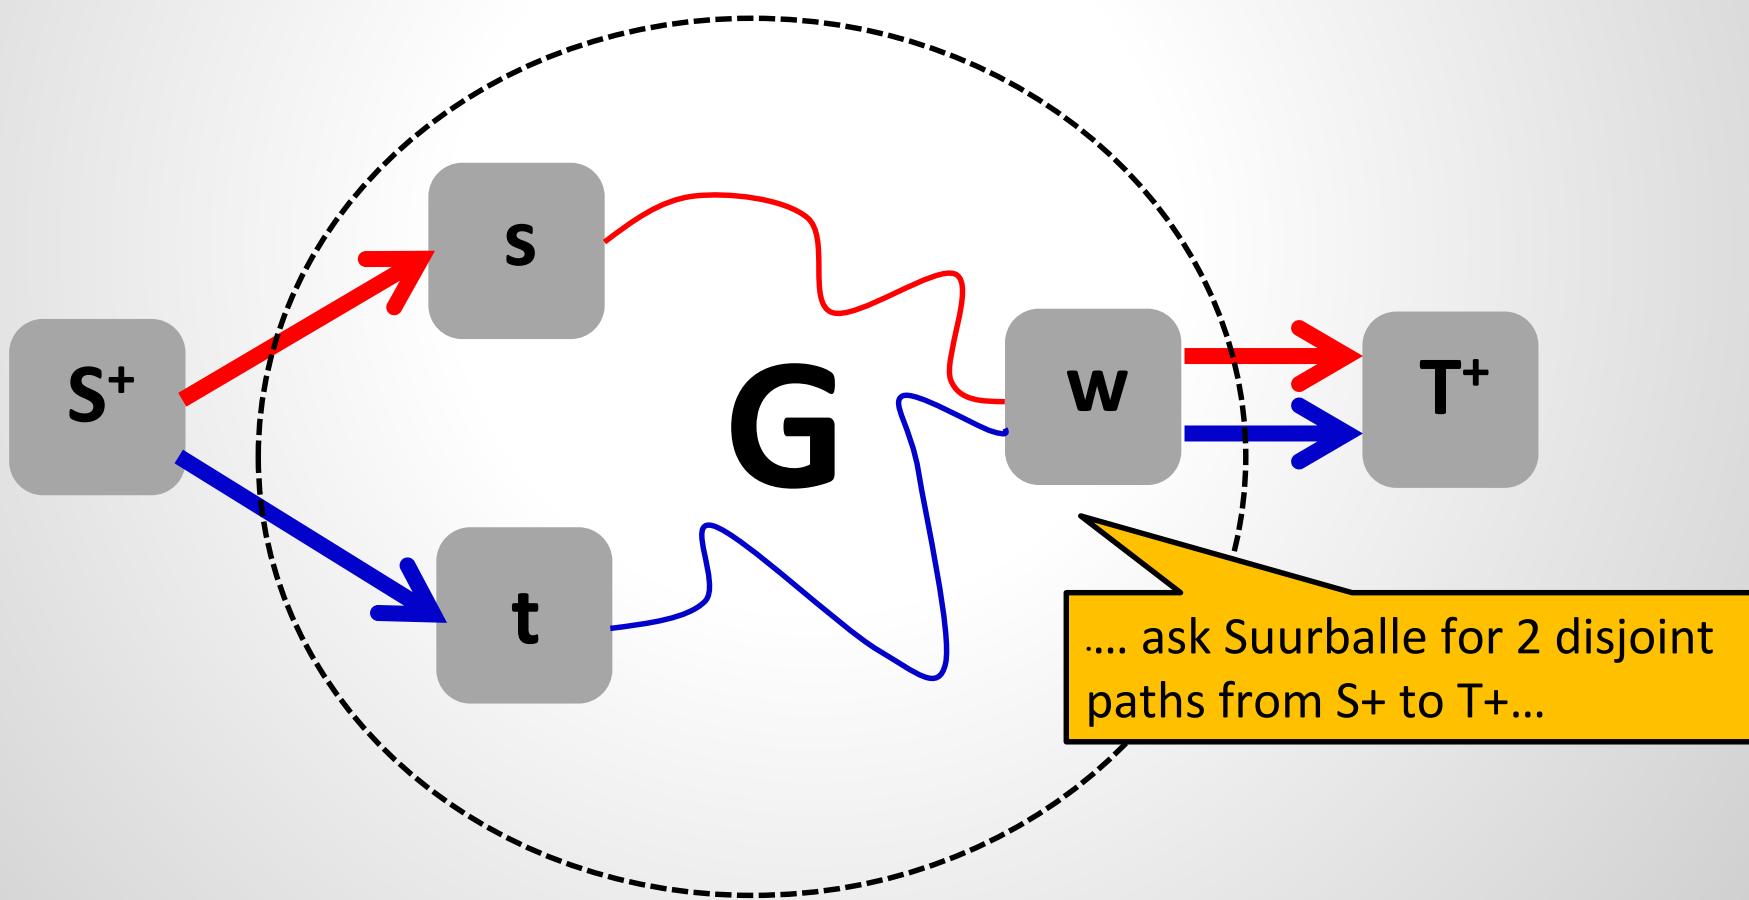
Waypoint Routing on Steroids

- ☐ Reduction to Suurballe's algorithm:



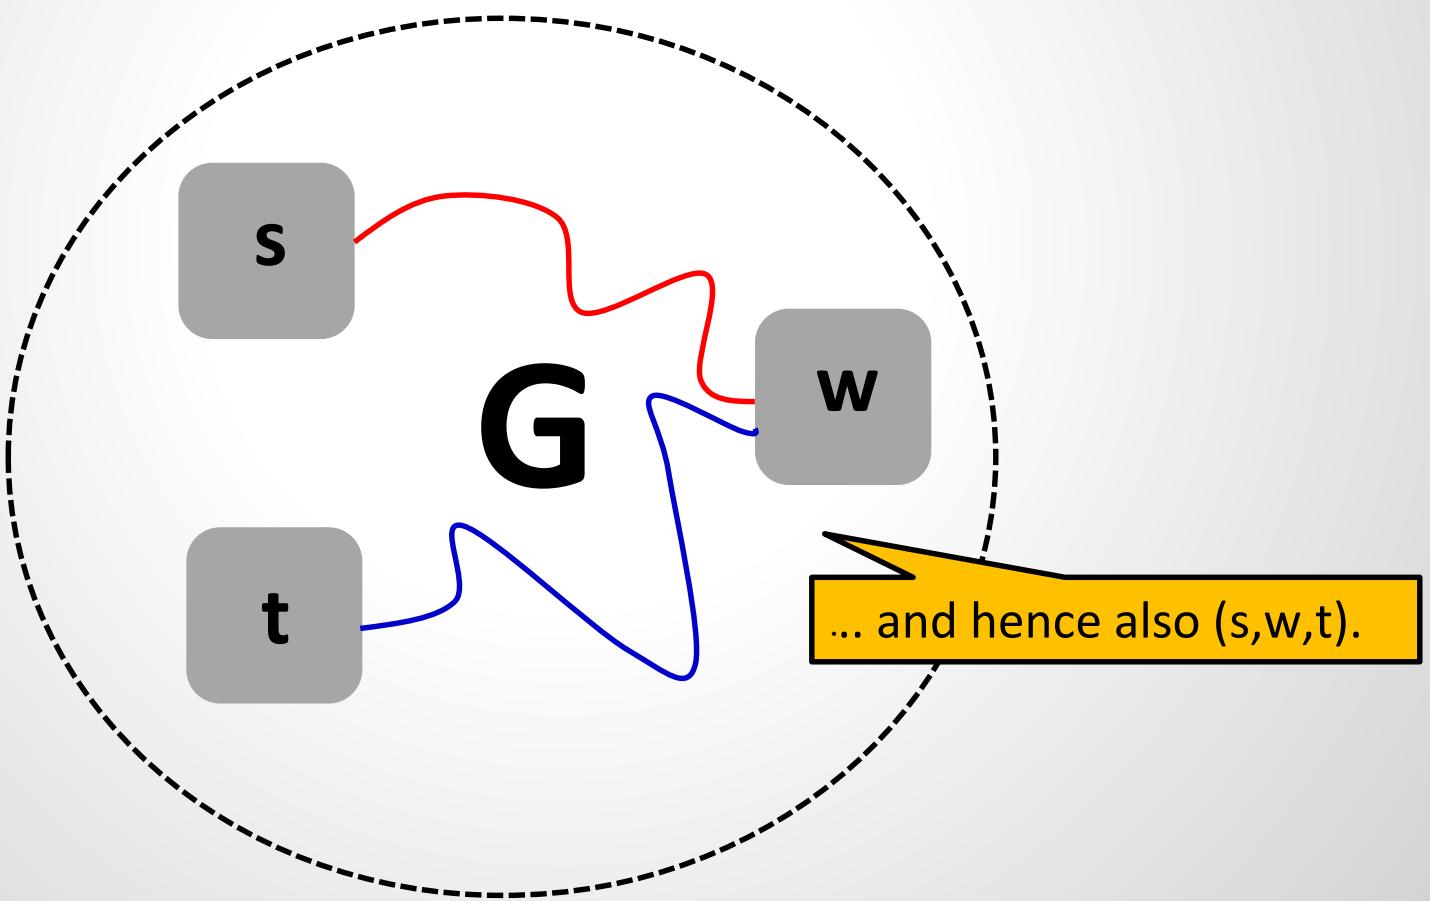
Waypoint Routing on Steroids

- ☐ Reduction to Suurballe's algorithm:



Waypoint Routing on Steroids

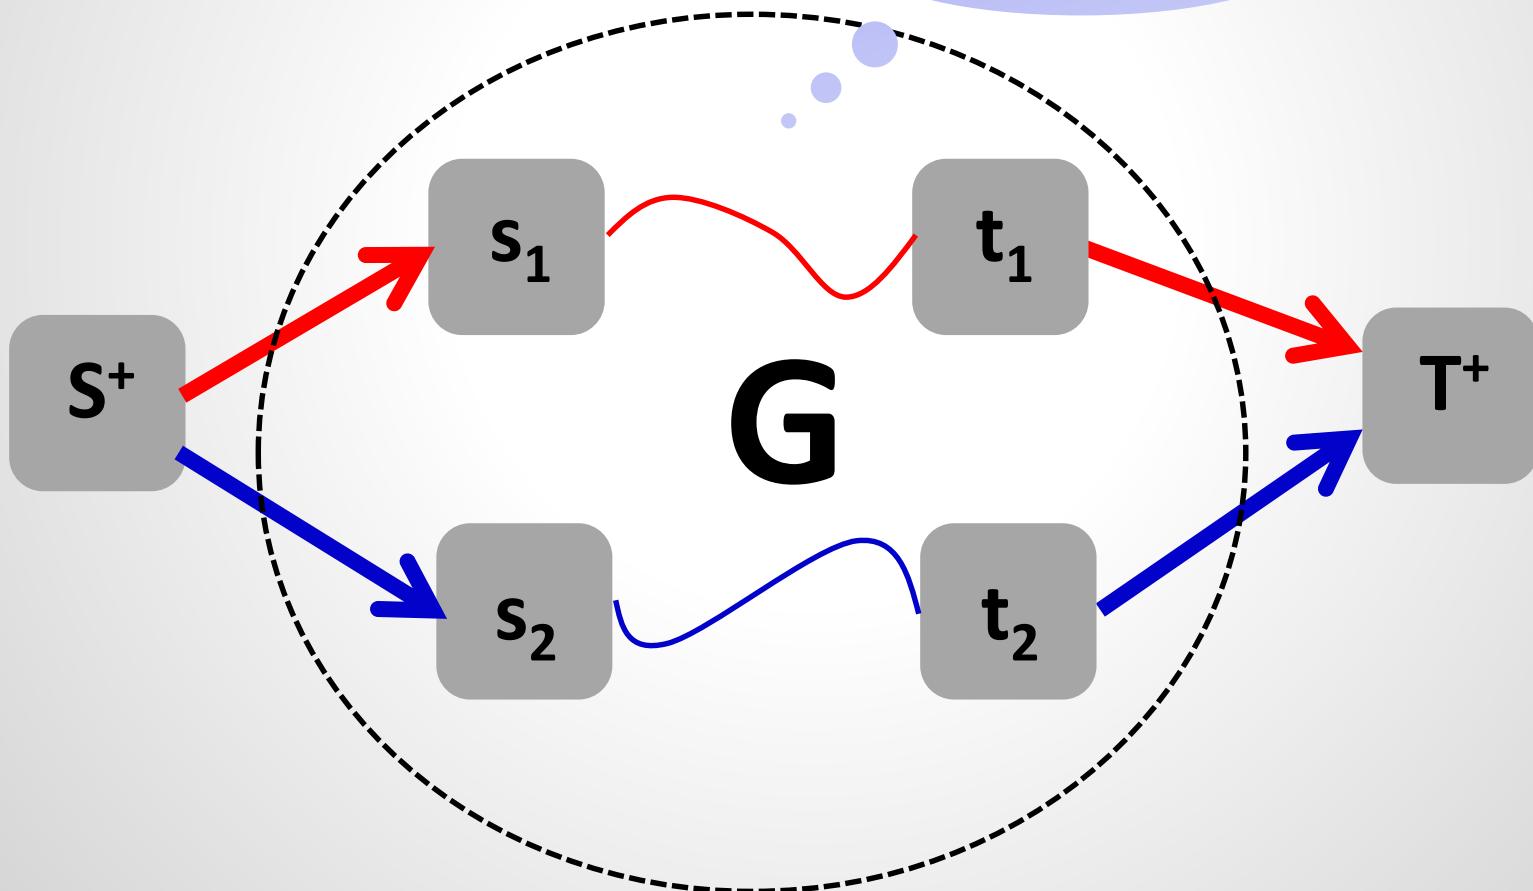
- Reduction to Suurballe's algorithm:



Waypoint Routing on Steroids

☐ Reduction to Suurballe

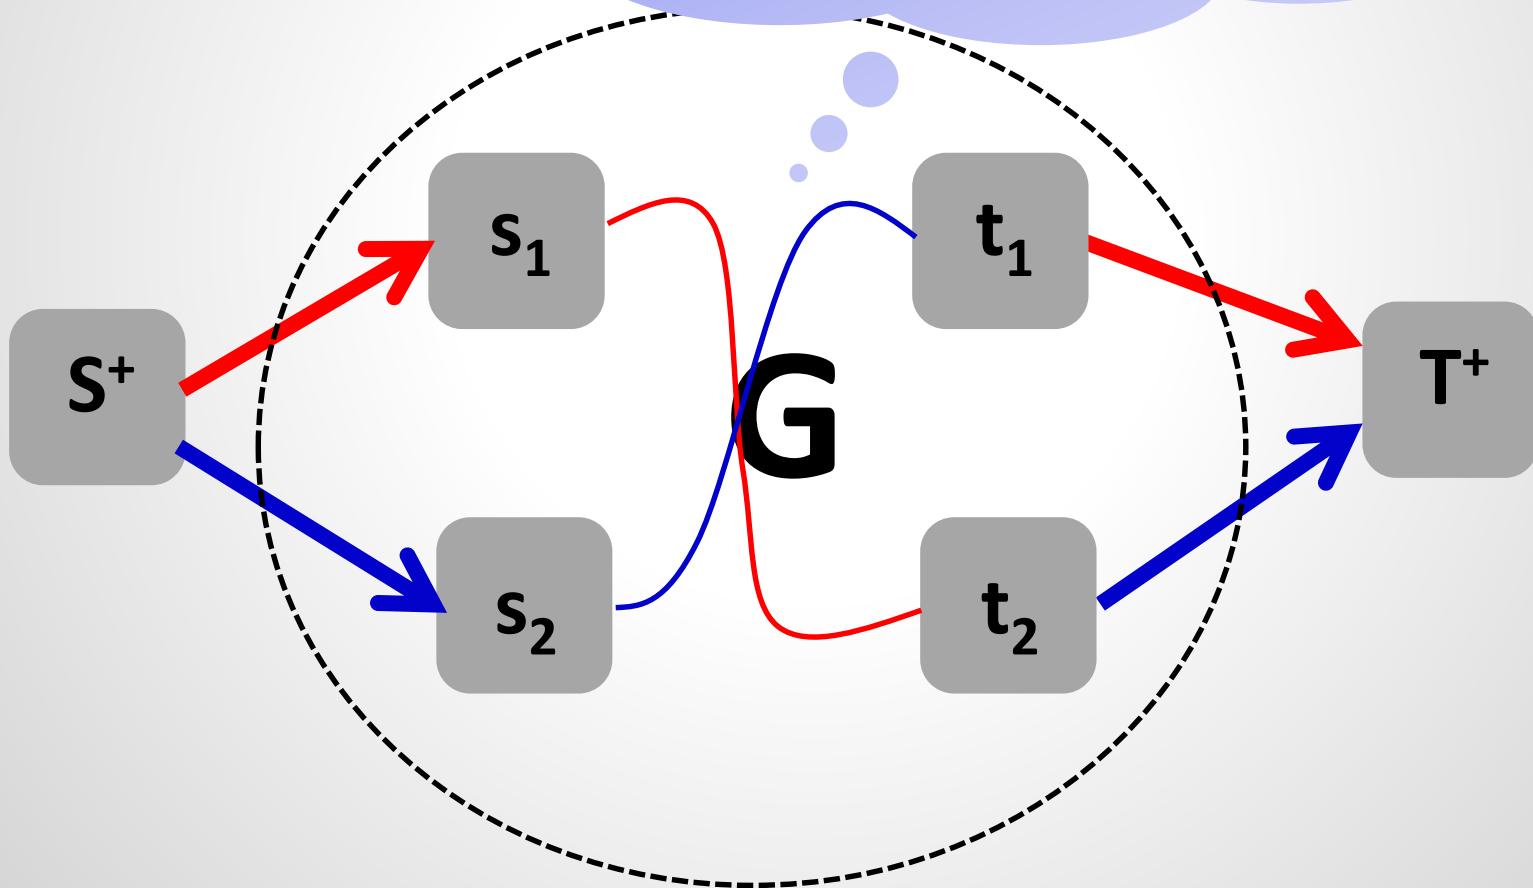
Can't I use Suurballe to efficiently compute disjoint paths as well?!



Waypoint Routing on Steroids

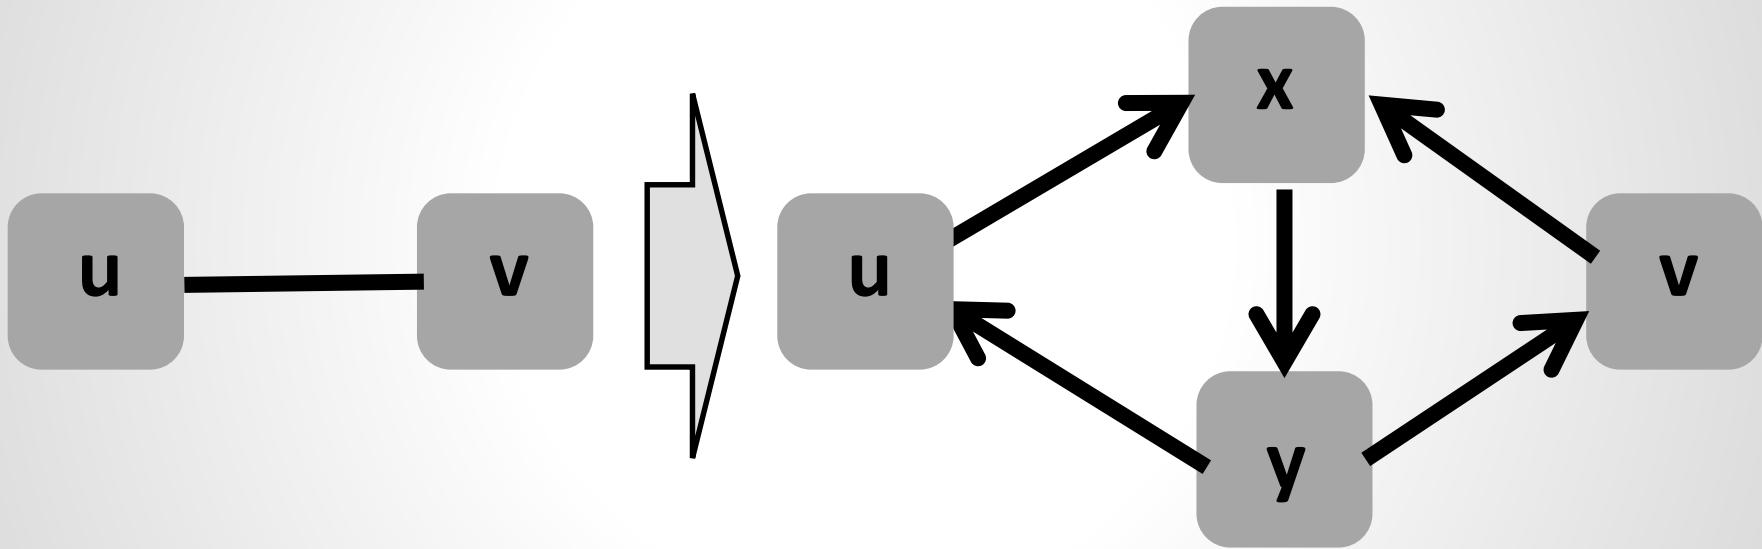
- ☐ Reduction to Suurballe's

No! Solves a much easier problem: 2 routes from $\{s_1, s_2\}$ to $\{t_1, t_2\}$.



Remarks: Under the rug...

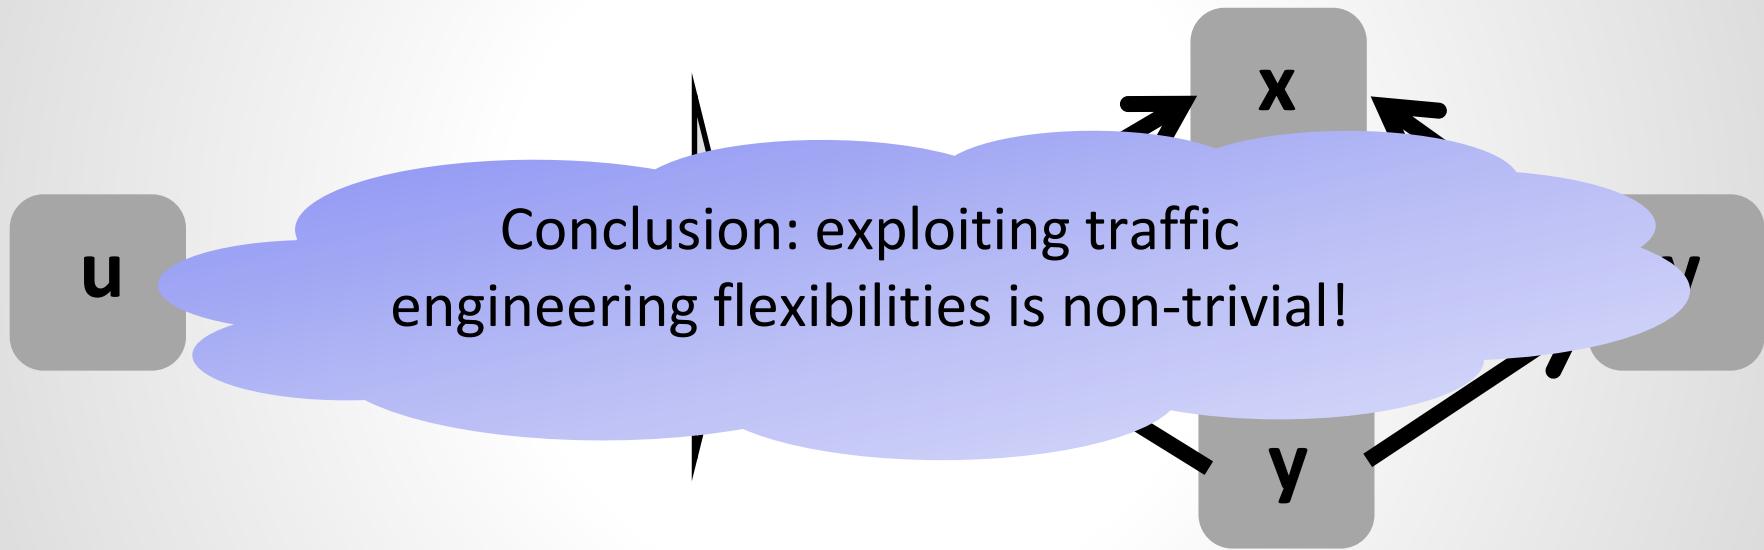
- ❑ Remark 1: Suurballe is actually for **directed substrate graphs**, so need gadget to transform problem in right form:



- ❑ Remark 2: Suurballe: actually vertex disjoint
 - ❑ Suurballe & Tarjan: edge disjoint

Remarks: Under the rug...

- ❑ Remark 1: Suurballe is actually for **directed substrate graphs**, so need gadget to transform problem in right form:



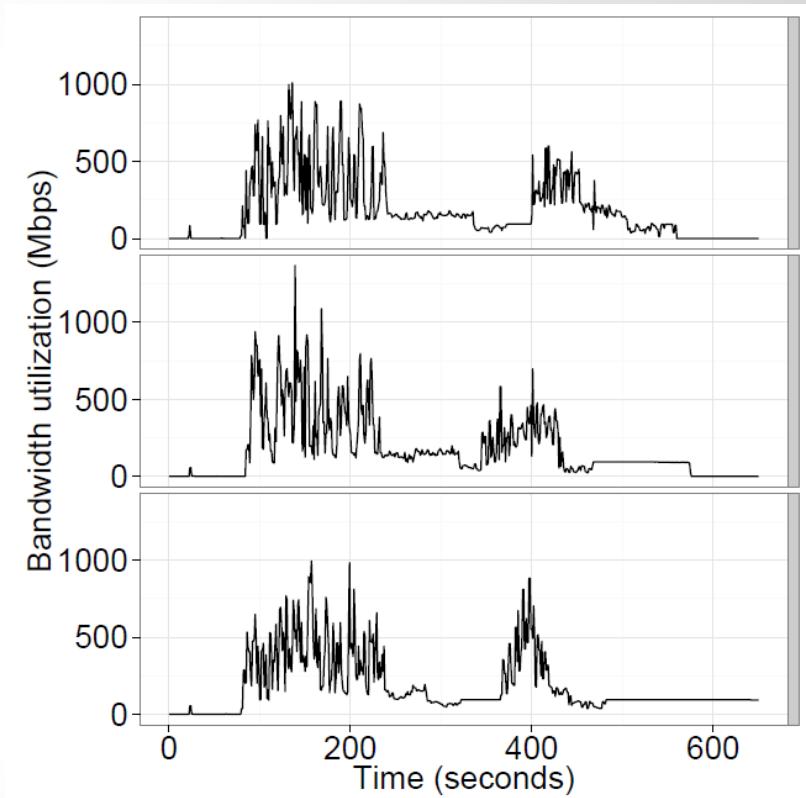
- ❑ Remark 2: Suurballe: actually vertex disjoint
 - ❑ Suurballe & Tarjan: edge disjoint

PART II:

Dynamic Embeddings

Real Communication Patterns Change over Time

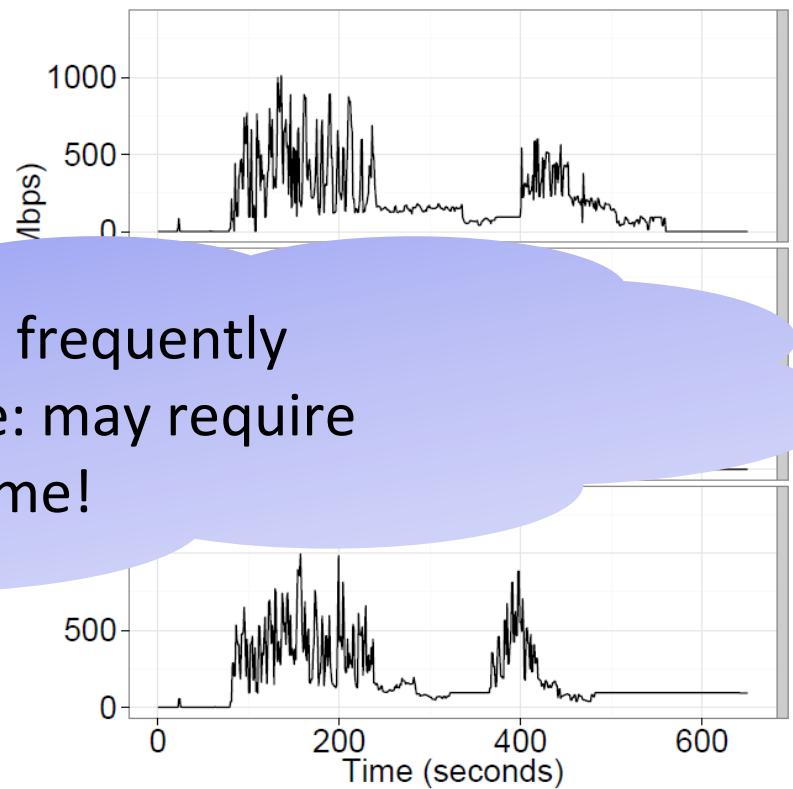
- E.g., **changing bandwidth demand**: map reduce application cycles through phases of high and low bandwidth requirements
- E.g., long-running applications (e.g., streaming) **change in popularity**
- E.g., **job churn**: jobs terminate, new ones arrive
- E.g., large **elephant flow degree**, changing over time (cf ProjecToR at SIGCOMM 2016)



Bandwidth utilization of 3 different runs of the same **TeraSort workload (without interference)**

Real Communication Patterns Change over Time

- E.g., **changing bandwidth demand**: map reduce application cycles through phases of high and low bandwidth requirements



- E.g. Ideally, we want to place frequently communicating nodes close: may require adaptions over time!

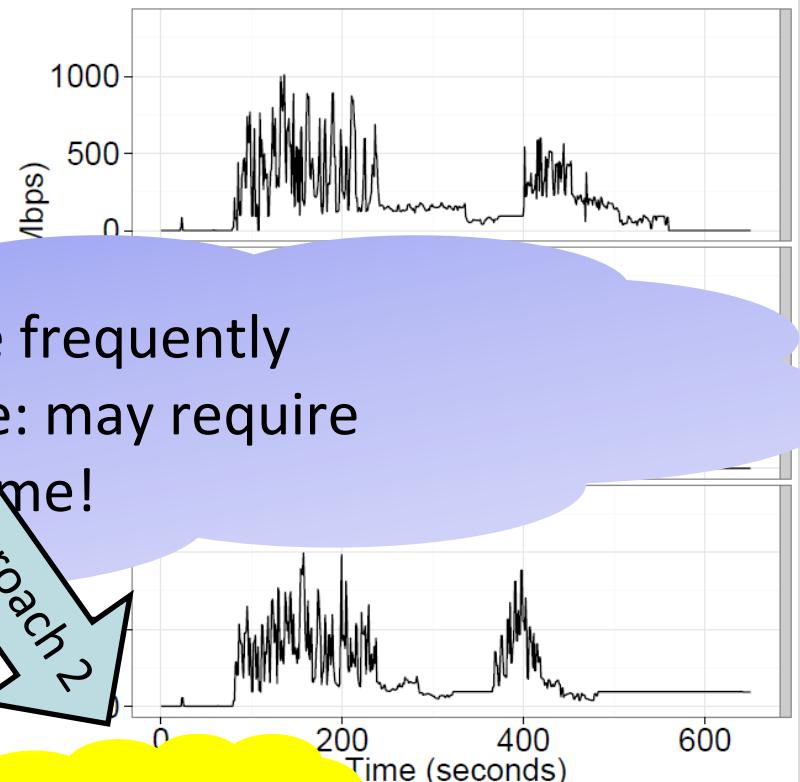
- E.g., **job churn**: new ones arrive

- E.g., large **elephant flow degree**, changing over time (cf ProjectToR at SIGCOMM 2016)

Bandwidth utilization of 3 different runs of the same **TeraSort workload (without interference)**

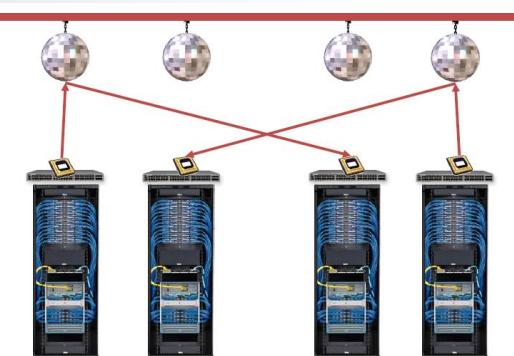
Real Communication Patterns Change over Time

- E.g., **changing bandwidth demand**: map reduce application cycles through phases of high and low bandwidth requirements



Ideally, we want to place frequently communicating nodes close: may require migrations over time!

Adjust the network!
ones arrive



Approach 1

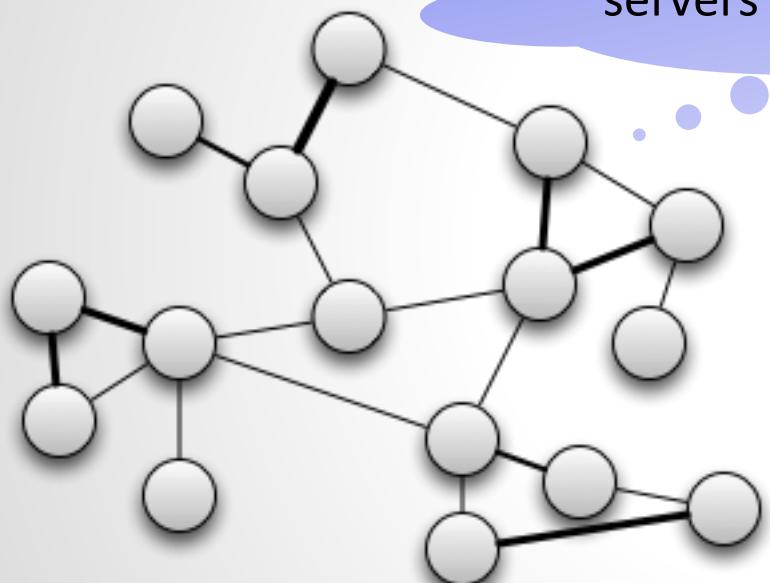
Approach 2

Migrate and collocate!

Utilization of 3 different runs of the same **TeraSort workload (without interference)**

Essentially An Online *RePartitioning* Problem

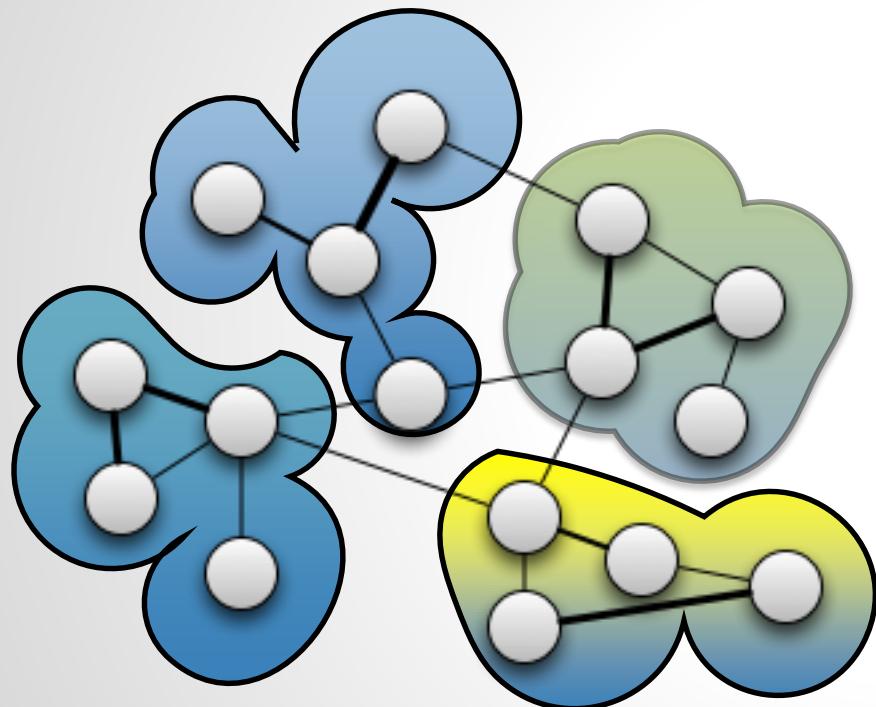
Communication @ time t:



How to embed pattern across $\ell=4$ servers (or racks, pods, etc.) of size $k=4$?

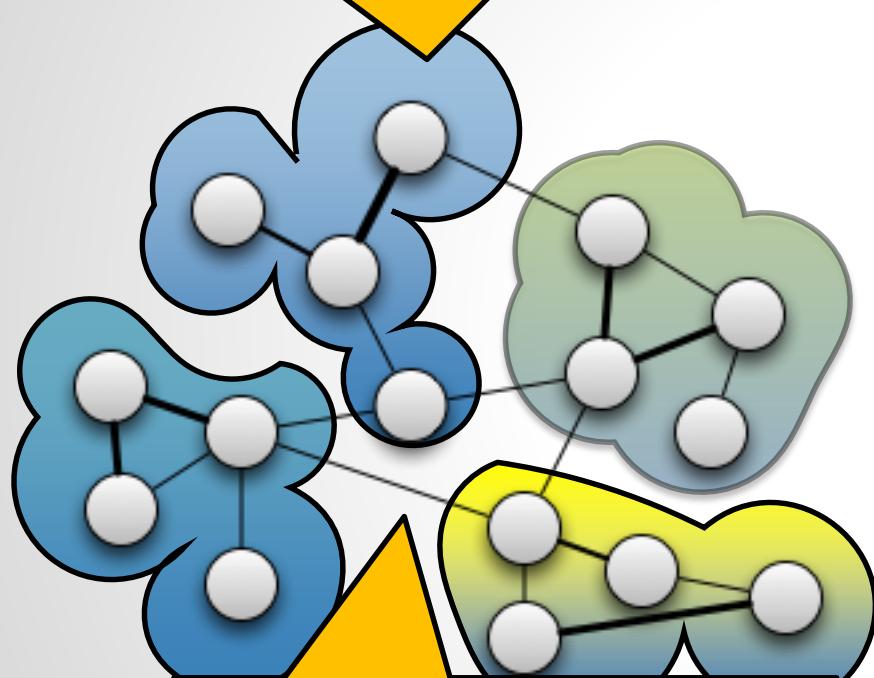
Essentially An Online *RePartitioning* Problem

Communication @ time t:



Essentially An Online *RePartitioning* Problem

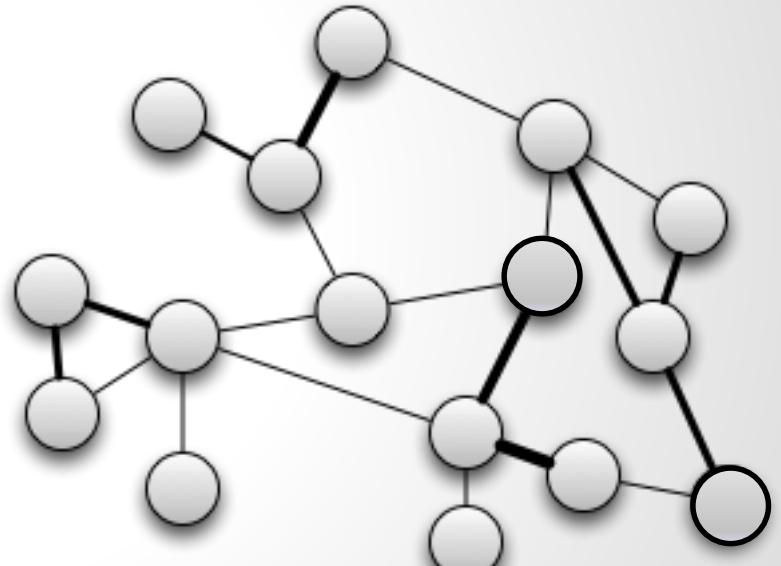
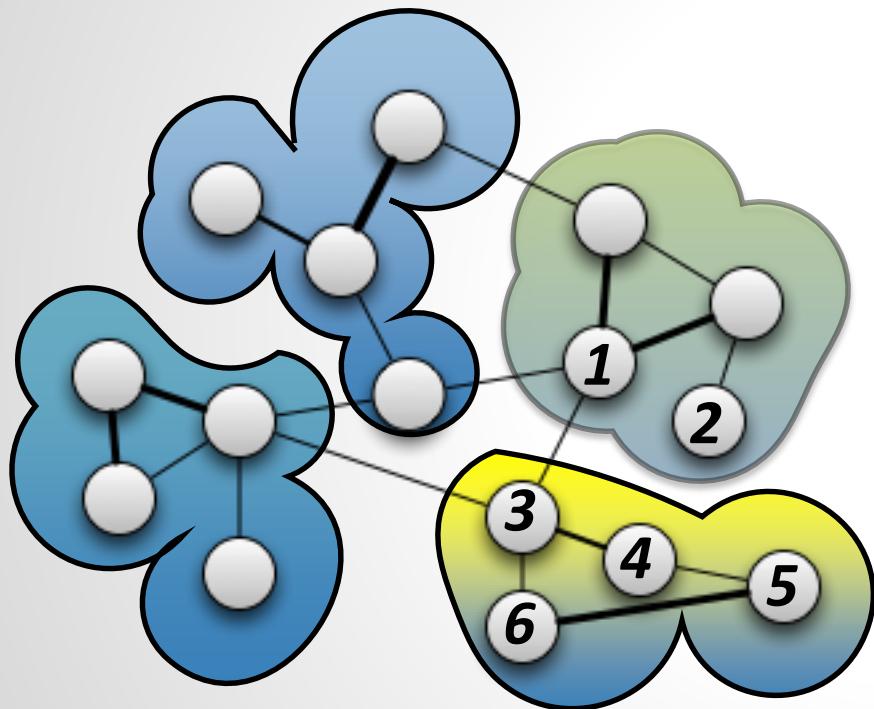
Most communication within cluster (intra-cluster)...



A classic (hard) combinatorial problem!

Essentially An Online *RePartitioning* Problem

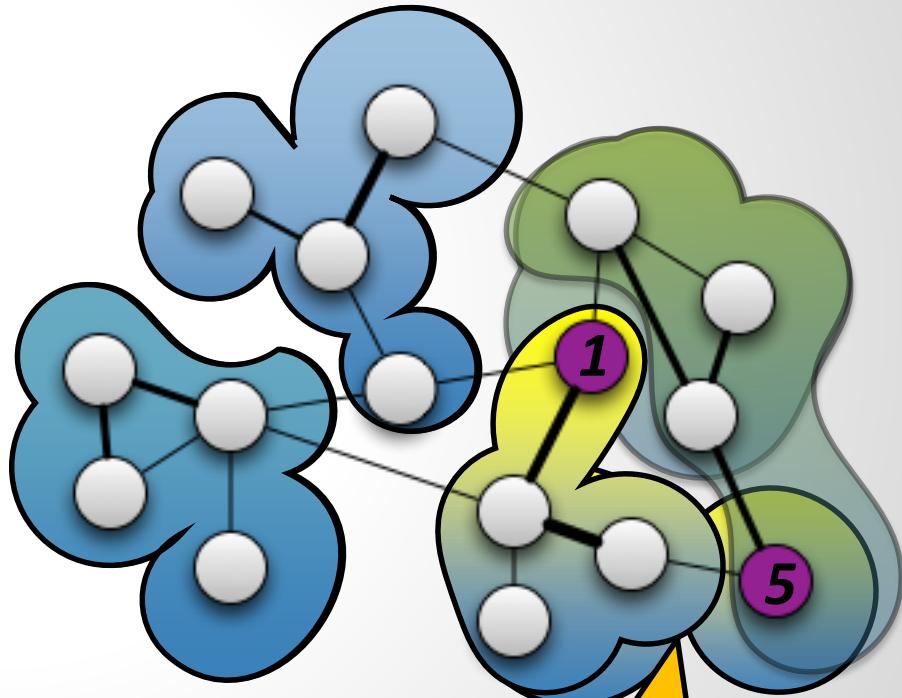
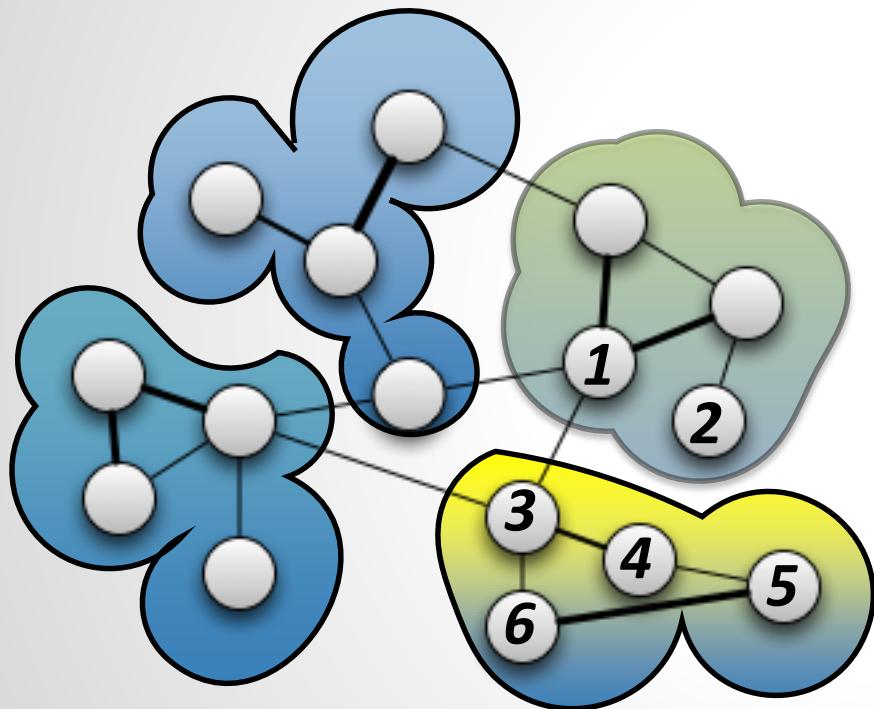
Communication @ time t: Communication @ time t+1:



- ❑ Now assume: changes in communication pattern!
- ❑ E.g., **more** communication (1,3),(3,4),(2,5) but **less** (5,6)

Essentially An Online *RePartitioning* Problem

Communication @ time t: Communication @ time t+1:

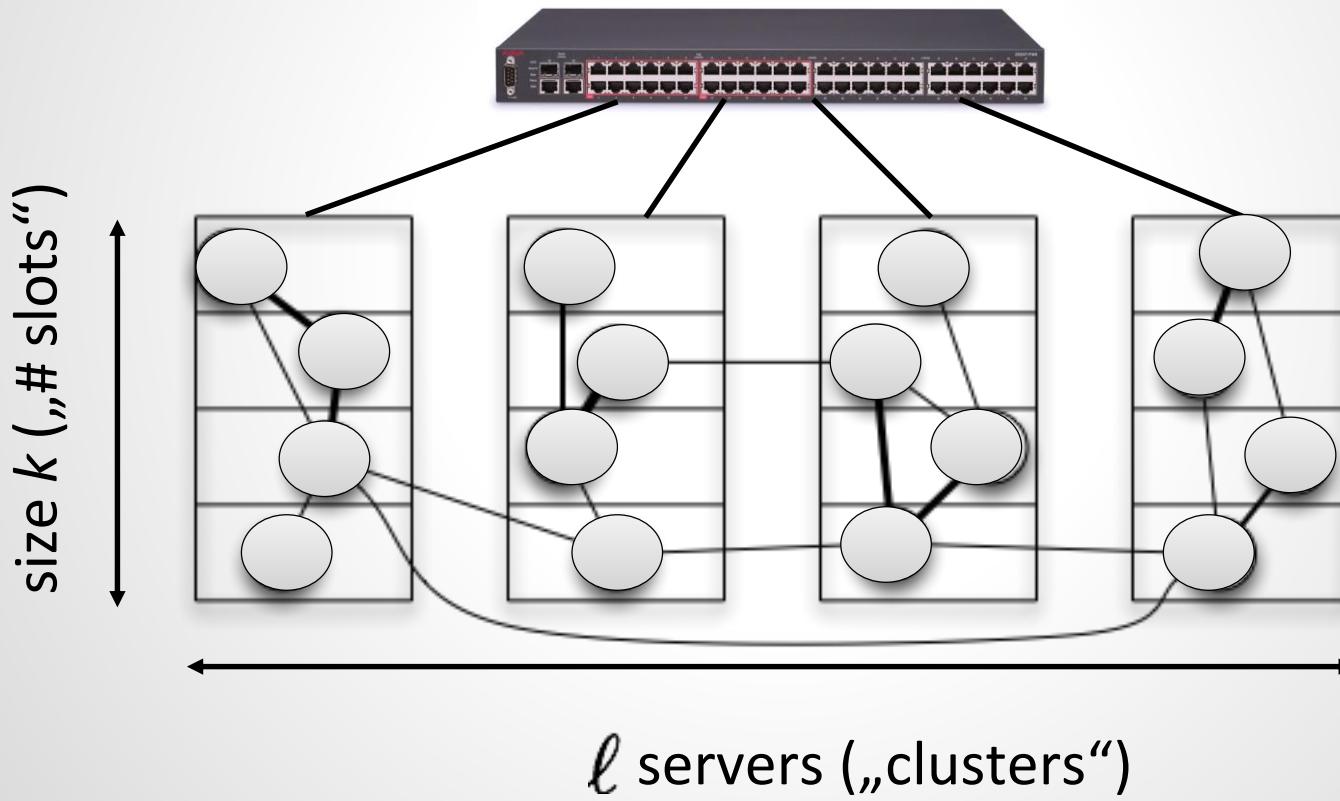


- ❑ Now assume: changes in communication pattern!
- ❑ E.g., more communication (1,3),(3,4),(2,5) but less (5,6)

Makes sense that
nodes 1 and 5
change clusters!

A Simple Model for the Tutorial

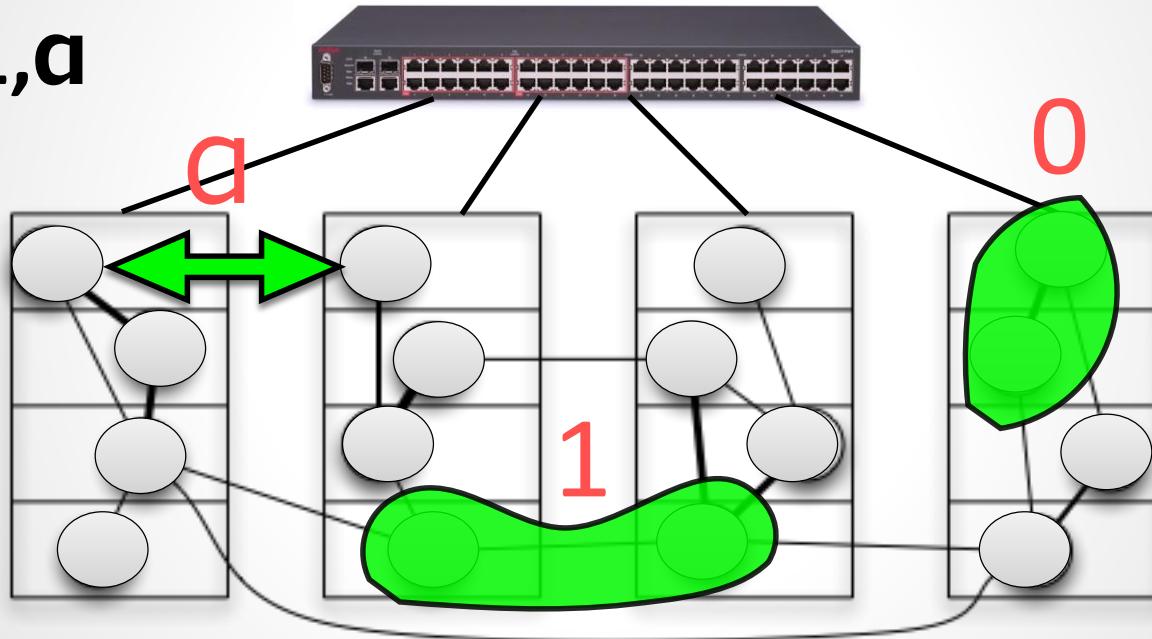
Consider a simple network, e.g., a single switch (e.g., a rack):



A Simple Model for the Tutorial

Consider a simple network, e.g., a single switch (e.g., a rack):

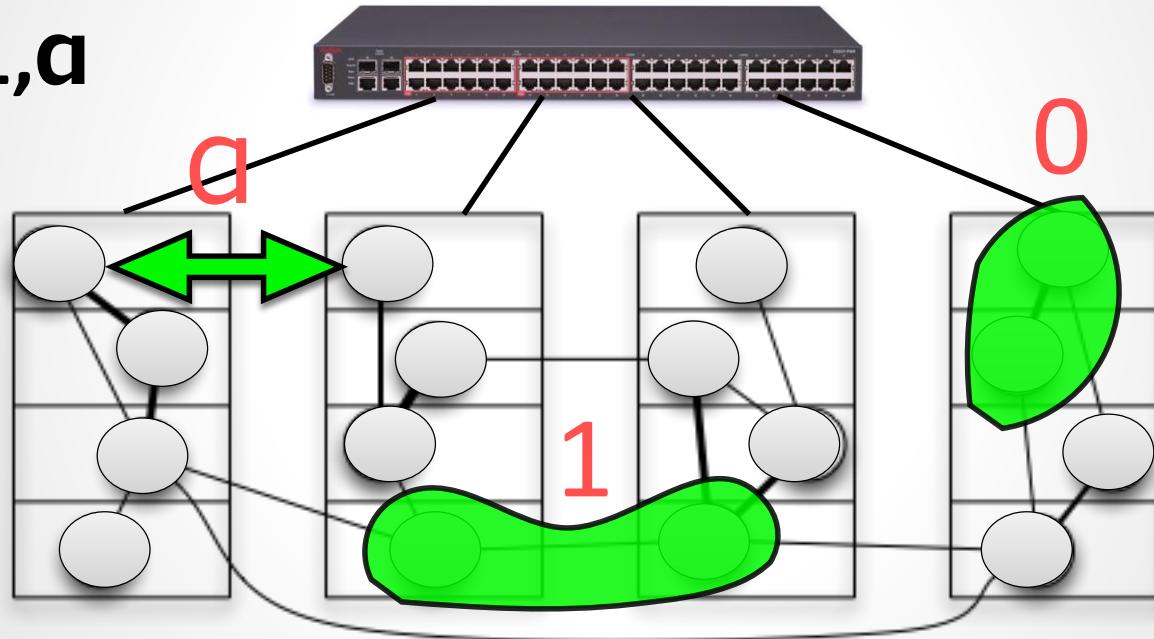
Costs: 0,1,a



A Simple Model for the Tutorial

Consider a simple network, e.g., a single switch (e.g., a rack):

Costs: 0,1,a



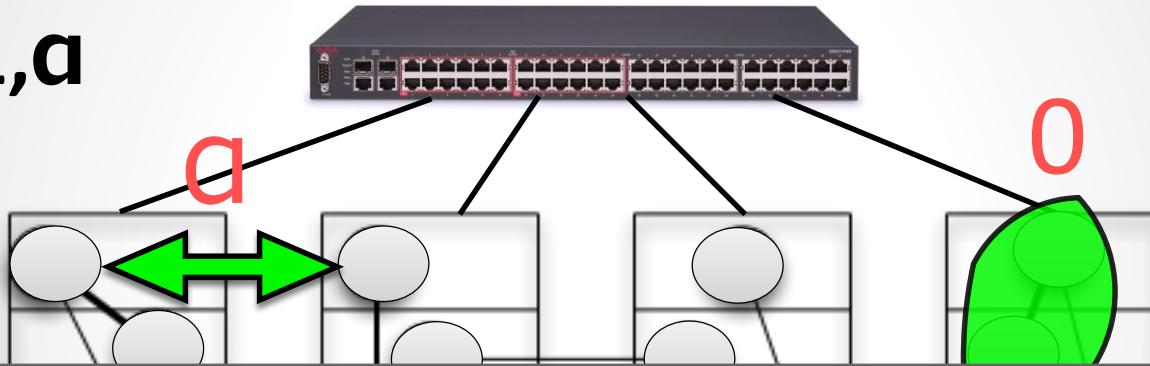
Objective: minimize total communication and migration cost!

More precisely: **competitive ratio** $\rho = \text{cost(ON)}/\text{cost(OPT)}$

A Simple Model for the Tutorial

Consider a simple network, e.g., a single switch (e.g., a rack):

Costs: 0,1,a



Nice: If competitive ratio is low, there is no need to develop any sophisticated prediction models (which may be wrong anyway)! The guarantee holds in the worst-case.

Objective: minimize total communication and migration cost!
More precisely: **competitive ratio** $\rho = \text{cost(ON)}/\text{cost(OPT)}$

“Prediction is difficult,
especially about the future.”



Nils Bohr

torial

(e.g., a rack):

0

C

Nice: If competitive ratio is low, there is no need to develop any sophisticated prediction models (which may be wrong anyway)! The guarantee holds in the worst-case.

Objective: minimize total communication and migration cost!
More precisely: **competitive ratio** $\rho = \text{cost(ON)}/\text{cost(OPT)}$

Adversarial Models

Weak adversary



Strong adversary



- Chooses **request distribution D**
 - Requests **sampled i.i.d.** from D
 - Cannot react to online algo
- Can generate **arbitrary request sequence σ**
 - Knows and can react to online algo

Adversarial Models

Weak adversary



Strong adversary

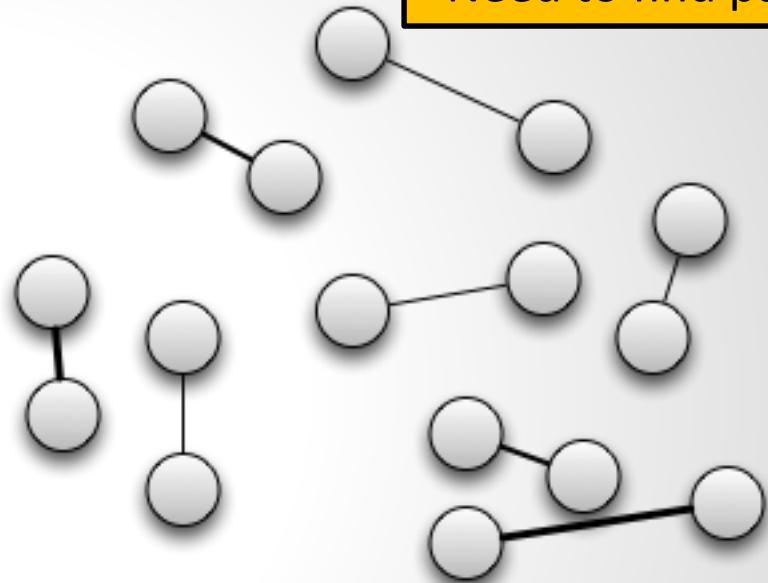
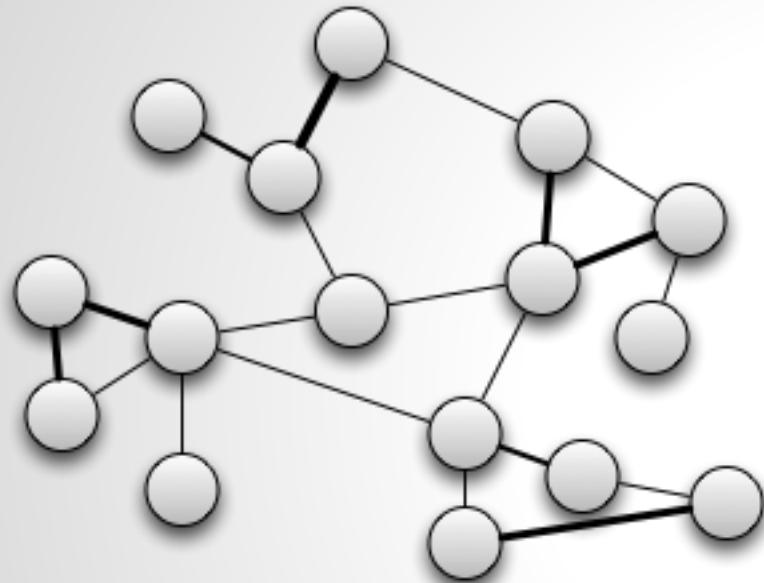


- Chooses **request distribution D**
- Requests **sampled i.i.d.** from D
- Cannot react to online algo
- Can generate **arbitrary request sequence σ**
- Knows and can react to online algo

The Crux: Algorithmic Challenges

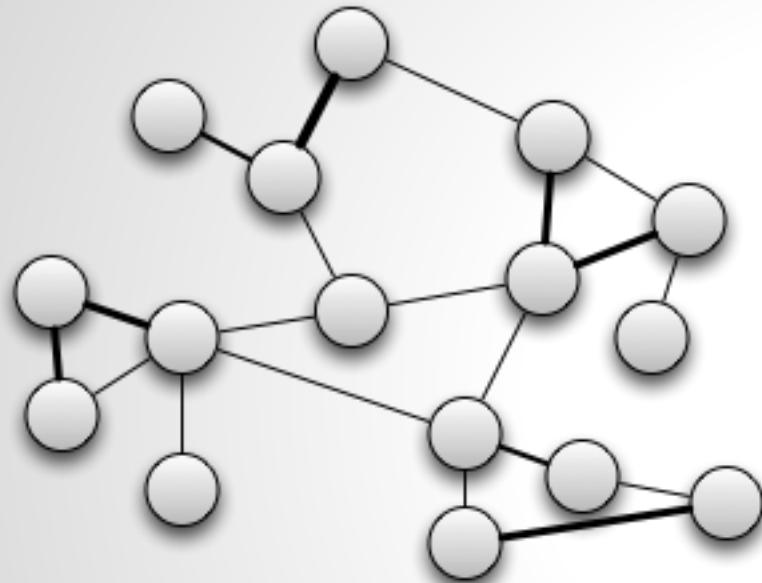
- Do **not know** D resp. σ ahead of time
- Upon each communication request (u, v) :
 - Migrate u and v together? «**Rent-or-buy**»: migration cost should be amortized
 - Migrate where?** u to v , v to u , both to a third cluster?
 - If cluster is full already: what to **evict**?

Example: Special Case $k=2$



Clusters of size 2:
Need to find pairs!

Example: Special Case $k=2$



Clusters of size 2:
Need to find pairs!

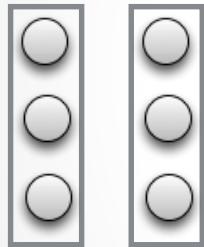
Clusters of size 2: A new
type of online
rematching problem!

It is hard to compete under !



- Assume **two clusters**: for offline algorithm they are of size k...

OFF:

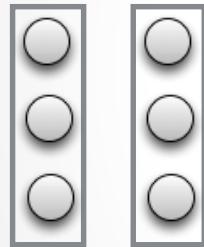


It is hard to compete under !

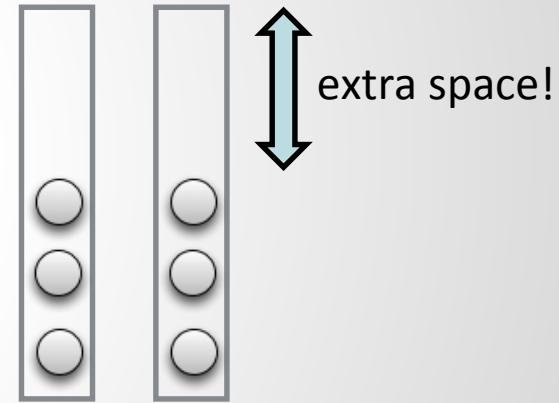


- ❑ Assume **two clusters**: for offline algorithm they are of size k ...
- ❑ ... whereas online algorithm can use clusters of size $2k-1$ even (**augmentation**)!

OFF:



ON:

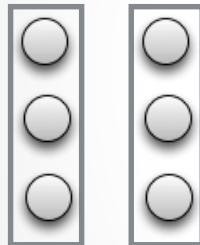


It is hard to compete under !



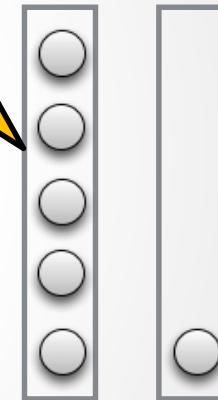
- ❑ Assume **two clusters**: for offline algorithm they are of size k ...
- ❑ ... whereas online algorithm can use clusters of size $2k-1$ even (**augmentation**)!

OFF:



E.g., ON can even collocate all except for one!

ON:

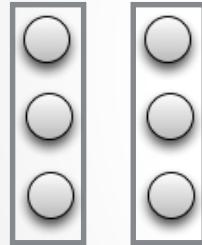


It is hard to compete under !



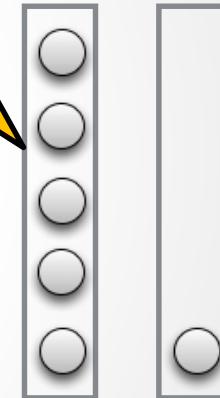
- ❑ Assume **two clusters**: for offline algorithm they are of size k ...
- ❑ ... whereas online algorithm can use clusters of size $2k-1$ even (**augmentation**)!

OFF:



E.g., ON can even collocate all except for one!

ON:

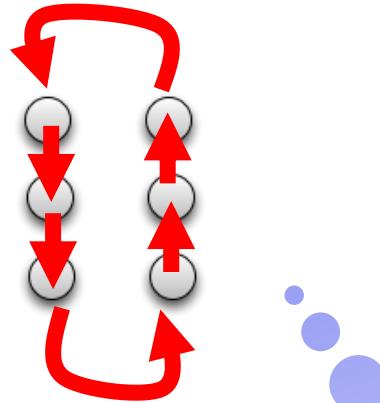


What is the achievable competitive ratio?

It is hard to compete under !



- ❑ Assume two clusters: for offline algorithm they are of size k ...
- ❑ ... whereas online algorithm can use clusters of size $2k-1$ even (augmentation)!



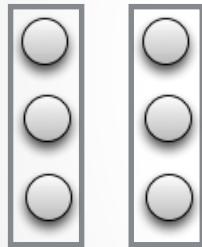
For the sake of lower bound, let us restrict the adversary more: can only ask for node pairs taken from a cyclic order: k pairs (resp. links) in total!

It is hard to compete under !

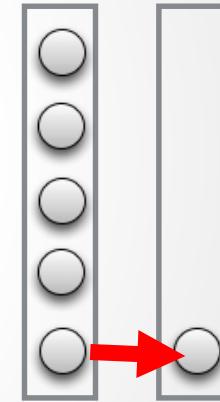


- ❑ Assume two clusters: for offline algorithm they are of size k ...
- ❑ ... whereas online algorithm can use clusters of size $2k-1$ even (augmentation)!

OFF:



ON:



Adversary can always request an inter-cluster link: always exists!



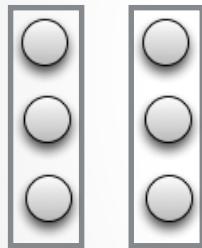
Ouch! Cost 1 for each request.

It is hard to compete under !

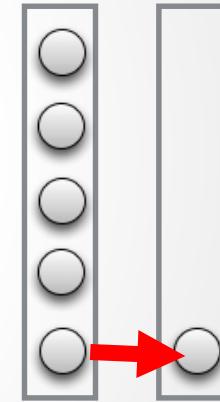


- ❑ Assume two clusters: for offline algorithm they are of size k ...
- ❑ ... whereas online algorithm can use clusters of size $2k-1$ even (augmentation)!

OFF:



ON:



Adversary can always request an inter-cluster link: always exists!



Ouch! Cost 1 for each request.

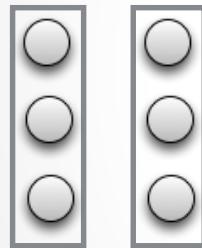
Note: adversarial strategy only depends on ON. So ON cannot learn anything about OFF!

It is hard to compete under !



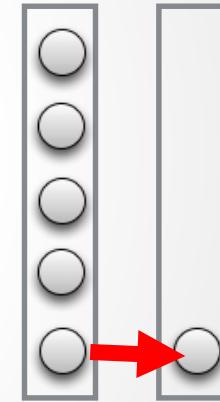
- ❑ Assume two clusters: for offline algorithm they are of size k ...
- ❑ ... whereas online algorithm can use clusters of size $2k-1$ even (augmentation)!

OFF:



What is the cost of OFF?

ON:



Adversary can always request an inter-cluster link: always exists!



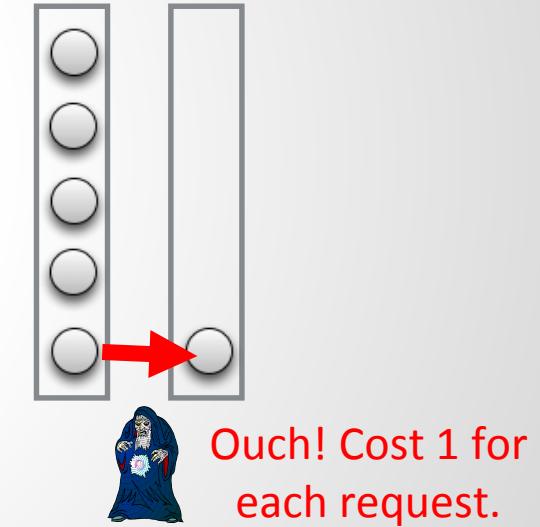
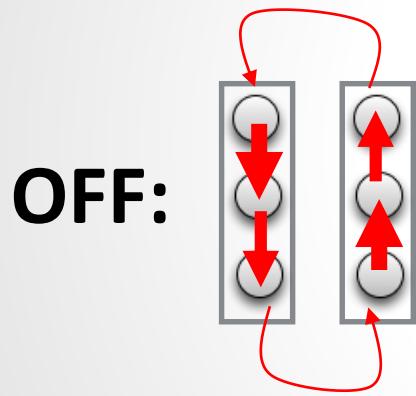
Ouch! Cost 1 for each request.

Note: adversarial strategy only depends on ON. So ON cannot learn anything about OFF!

It is hard to compete under !



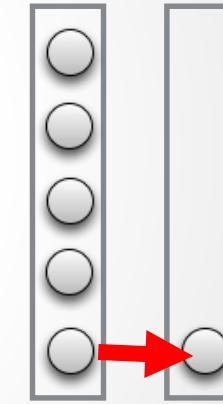
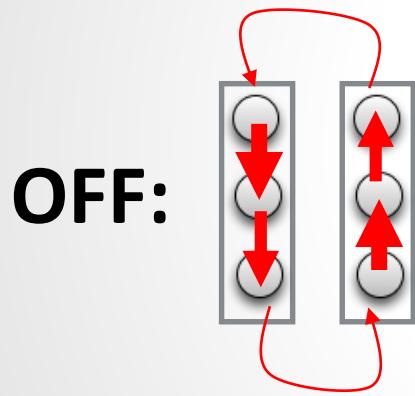
- ❑ Assume two clusters: for offline algorithm they are of size k ...
- ❑ ... whereas online algorithm can use clusters of size $2k-1$ even (augmentation)!



It is hard to compete under !



- Assume two clusters: for offline algorithm they are of size k ...
- ... whereas online algorithm can use clusters of size $2k-1$ even (augmentation)!



Lower bound of $\Omega(k)$ for competitive ratio, despite big augmentation!

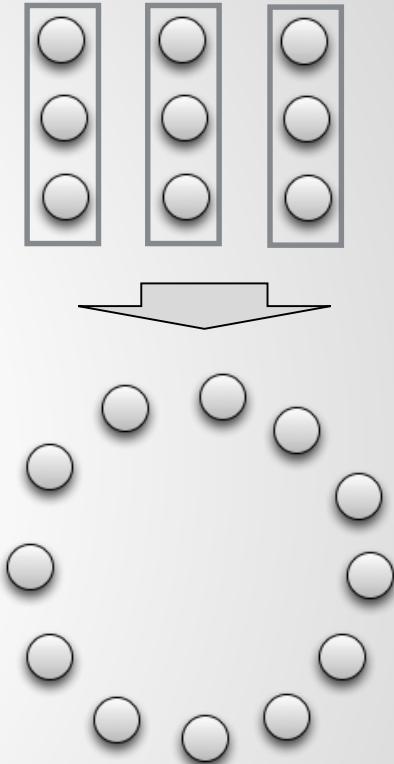
A Simple $O(n^2)$ Upper Bound

At least it does not depend on time! ☺

A Simple $O(n^2)$ Upper Bound

- Algorithm DET:

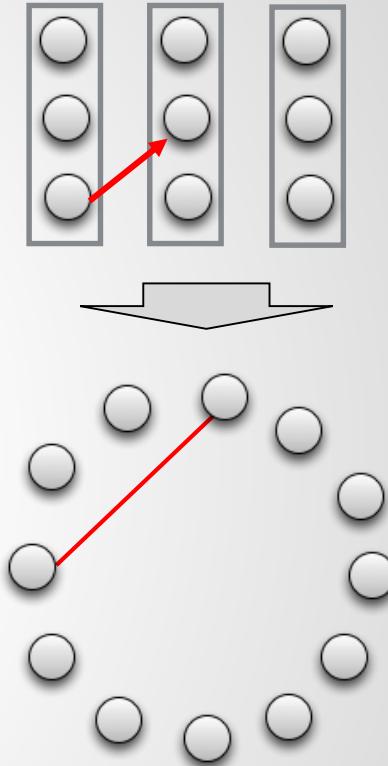
- Based on «growing communication components»
- Cycles through phases
 - Initially in each phase: empty graph of n nodes



A Simple $O(n^2)$ Upper Bound

□ Algorithm DET:

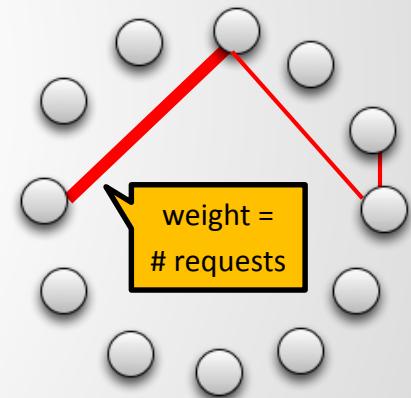
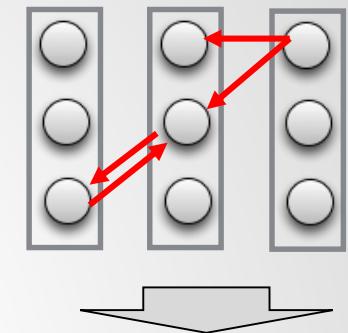
- Based on «**growing communication components**»
- Cycles through **phases**
 - Initially in each phase: **empty graph** of n nodes
 - For each inter-cluster request for ON: **insert edge**



A Simple $O(n^2)$ Upper Bound

□ Algorithm DET:

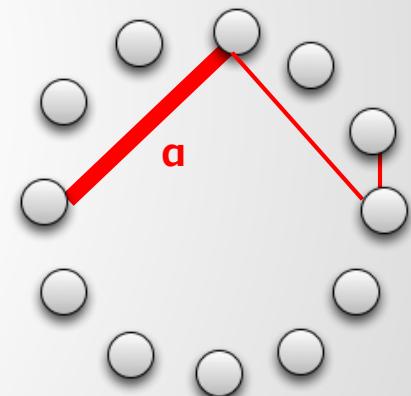
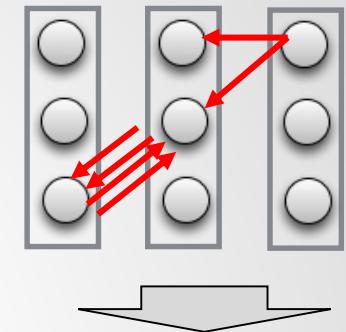
- Based on «growing communication components»
- Cycles through phases
 - Initially in each phase: empty graph of n nodes
 - For each inter-cluster request for ON: insert edge
 - Induces a «communication component»: **edge weight = # requests**



A Simple $O(n^2)$ Upper Bound

□ Algorithm DET:

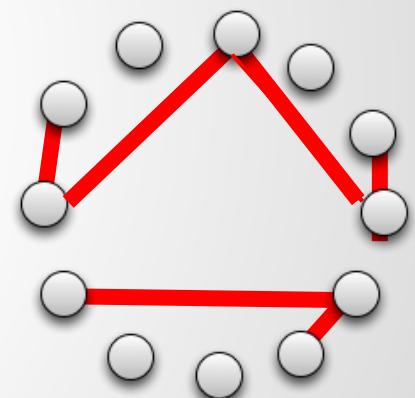
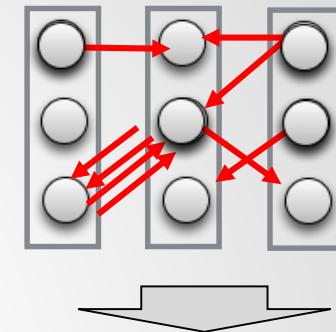
- Based on «growing communication components»
- Cycles through phases
 - Initially in each phase: empty graph of n nodes
 - For each inter-cluster request for ON: insert edge
 - Induces a «communication component»: edge weight = # requests
 - If an edge (u,v) **weight reaches a** , DET
repartitions nodes, so that *all edges which have reached a so far are in same cluster!*



A Simple $O(n^2)$ Upper Bound

□ Algorithm DET:

- Based on «growing communication components»
- Cycles through phases
 - Initially in each phase: empty graph of n nodes
 - For each inter-cluster request for ON: insert edge
 - Induces a «communication component»: edge weight = # requests
 - If an edge (u,v) **weight reaches a** , DET **repartitions nodes**, so that *all edges which have reached a so far are in same cluster!*
 - If this is not possible: **phase ends**

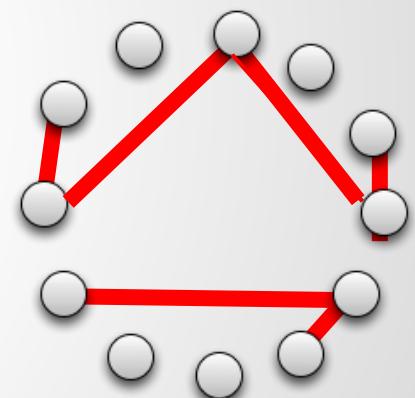
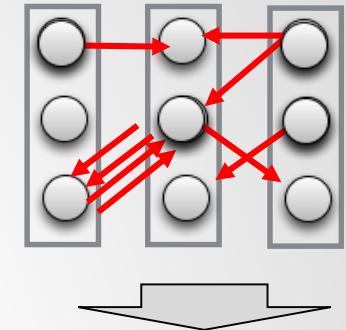


Components cannot be partitioned perfectly (first component alone too large)!

A Simple $O(n^2)$ Upper Bound

□ Algorithm DET:

- Based on «growing communication components»
- Cycles through phases
 - Initially in each phase: empty graph of n nodes
 - For each inter-cluster request for ON: insert edge
 - Induces a «communication component»: edge weight = # requests
 - If an edge (u,v) weight reaches α , DET repartitions nodes, so that *all edges which have reached α so far are in same cluster!*
 - If this is not possible: phase ends

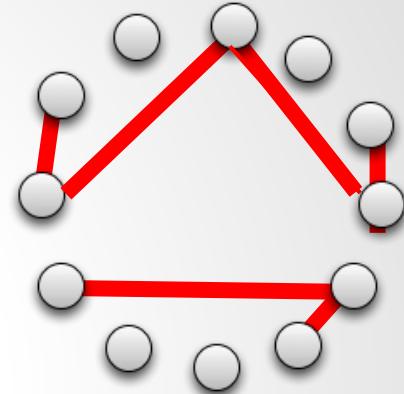


Competitive ratio?

A Simple $O(n^2)$ Upper Bound

- Analysis (costs per phase):

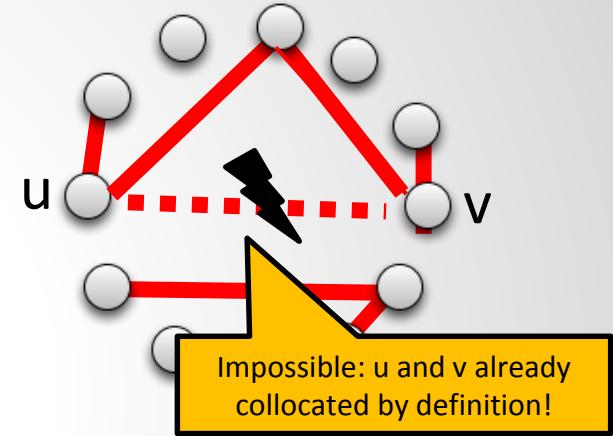
- Observe: edge weights always $\leq \alpha$: once reach α , their endpoints will always be collocated (by algorithm definition)



A Simple $O(n^2)$ Upper Bound

□ Analysis (costs per phase):

- Observe: edge **weights always $\leq \alpha$** : once reach α , their endpoints will always be collocated (by algorithm definition)
- α -edges form a **forest** (so at most n many!): once two nodes (u,v) are connected by a **path of α -edges**, they are in a single cluster and will no longer communicate across clusters



A Simple $O(n^2)$ Upper Bound

□ Analysis (costs per phase):

- Observe: edge **weights always $\leq \alpha$** : once reach α , their endpoints will always be collocated (by algorithm definition)
- α -edges form a **forest** (so at most n many!): once two nodes (u,v) are connected by a **path of α -edges**, they are in a single cluster and will no longer communicate across clusters
- Thus: ON cost per phase:
 - At most 1 reorganization per α -edge (at most $n \alpha$ -edges), so **n times reconfig cost $n \cdot \alpha$** , so $n^2\alpha$
 - Communication cost: at most **α per edge** (at most n^2 many), so also **at most $n^2\alpha$**



A Simple $O(n^2)$ Upper Bound

□ Analysis (costs per phase):

- Observe: edge **weights always $\leq \alpha$** : once reach α , their endpoints will always be collocated (by algorithm definition)
- α -edges form a **forest** (so at most n many!): once two nodes (u,v) are connected by a **path of α -edges**, they are in a single cluster and will no longer communicate across clusters
- Thus: ON cost per phase:
 - At most 1 reorganization per α -edge (at most $n \alpha$ edges), so **n times reconfig cost $n \cdot \alpha$** , so $n^2 \alpha$
 - Communication cost: at most **α per edge** (at most n^2 many), so also **at most $n^2 \alpha$**
- Costs of OFF per phase:
 - If OFF migrates any node, it pays **at least α**
 - If not, it pays communication cost **at least α** : the grown components do not fit clusters (intra-cluster edges only): **definition of «end-of-phase»!**



Upper bound of $O(n^2\alpha/\alpha) = O(n^2)$ for competitive ratio!

Known Results So Far

- ❑ Case $k=2$ („online rematching“): **constant** competitive ratio
- ❑ General case: with a little bit of augmentation: **$O(k \log k)$** possible
 - ❑ Recall $\Omega(k)$ lower bound
 - ❑ Nice: independent of number of clusters!
 - ❑ Practically relevant: # VM slots per server usually small

What about ?



What about ?



- Recall: weak adversary cannot choose request sequence but only the distribution
 - Adversary needs to sample i.i.d. from this distribution
 - Moreover: Adversary knows (deterministic or randomized) «learning» algorithm, i.e., chooses worst distribution

Any ideas?

The Crux: *Joint* Optimization of Efficient Learning *and* Searching

- **Naive idea 1: Take it easy and first learn distribution**
 - Do not move but just sample requests in the beginning: until exact distribution has been **learned whp**
 - Then move to the best location **for good**

The Crux: *Joint Optimization*

Learner vs. Server



Waiting can be very costly: maybe start configuration is very bad and others similarly good: takes long to learn, not competitive! Need to move early on, away from bad locations!

□ Naive idea 1: Take it easy and...

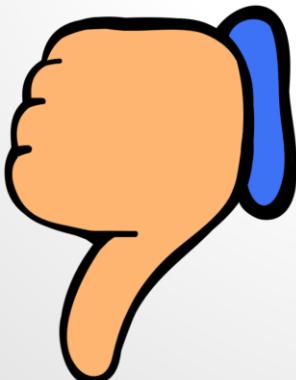
- Do not move but just sample + requests in the beginning: until exact distribution has been learned whp
- Then move to the best location for good

The Crux: *Joint* Optimization of Efficient Learning *and* Searching

- **Naive idea 1: Take it easy and first learn distribution**
 - Do not move but just sample requests in the beginning: until exact distribution has been **learned whp**
 - Then move to the best location **for good**
- **Naive idea 2: Pro-actively always move to the lowest cost configuration seen so far**

The Crux: *Joint* Optimization of Efficient Learning *and* Searching

- **Naive idea 1: Take it easy and first learn distribution**
 - Do not move but just sample requests in the beginning: until exact distribution has been **learned whp**
 - Then move to the best location **for good**
- **Naive idea 2: Pro-actively always move to the lowest cost configuration seen so far**



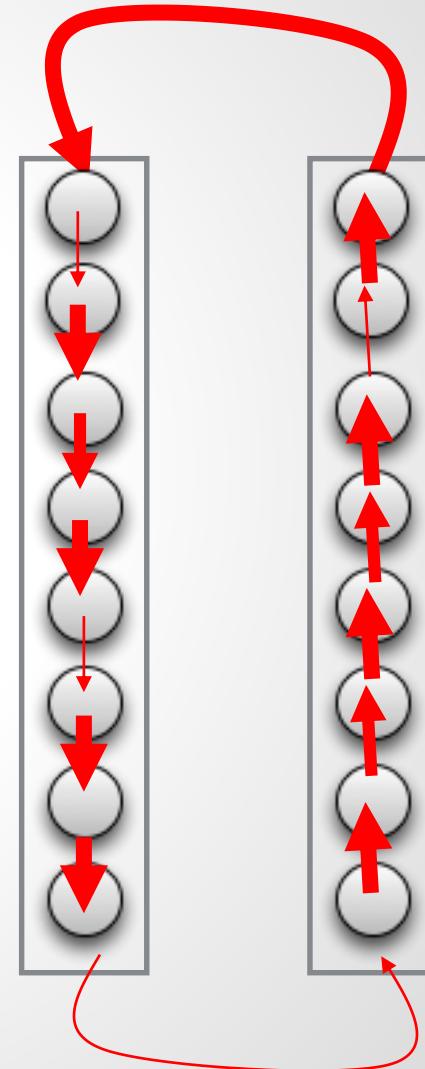
Bad: if requests are uniform at random, you should not move at all!
Migration costs cannot be amortized. Crucial difference to classic distribution learning problems: guessing costs!

The Crux: *Joint* Optimization of Efficient Learning *and* Searching

- **Naive idea 1: Take it easy and first learn distribution**
 - Do not move but just sample requests in the beginning: until exact distribution has been learned
 - Then move ... Only move when it pays off! But e.g., how to differentiate between uniform and „almost uniform“ distribution?
- **Naive idea 2: Always move to lowest cost configuration**
 - Bad, e.g., if requests are distributed uniformly at random: better not to move at all (moving costs cannot be amortized)

Example Learning Algorithm for Ring: Rotate Locally!

- ❑ Mantra of our algorithm: Rotate!
 - ❑ Rotate early, but not too early!
 - ❑ And: rotate **locally**

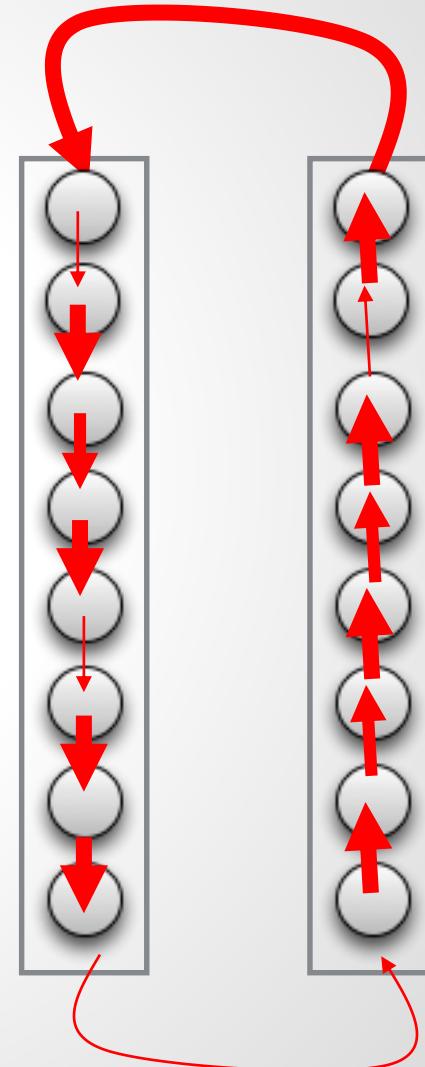


Example Learning Algorithm for Ring: Rotate Locally!

Define conditions for configurations: if met, **never go back** to it (we can afford it w.h.p.: seen enough samples)

Algorithm: Rotate!

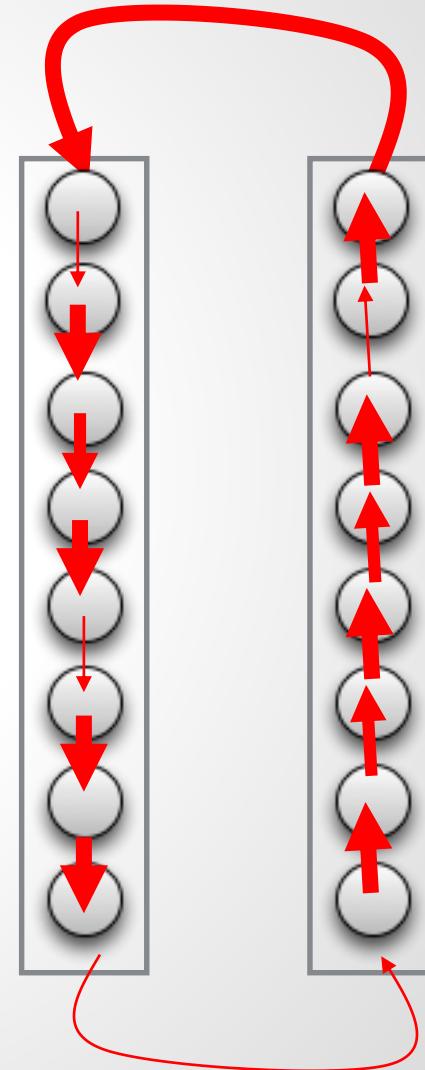
- Rotate early, but not too early!
- And: rotate **locally**



Example Learning Algorithm for Ring: Rotate Locally!

- ❑ Mantra of our algorithm: Rotate!
 - ❑ Rotate early, but not too early!
 - ❑ And: rotate **locally**

If current configuration is eliminated, go to **nearby configuration** (in directed manner: no frequent back and forth)!

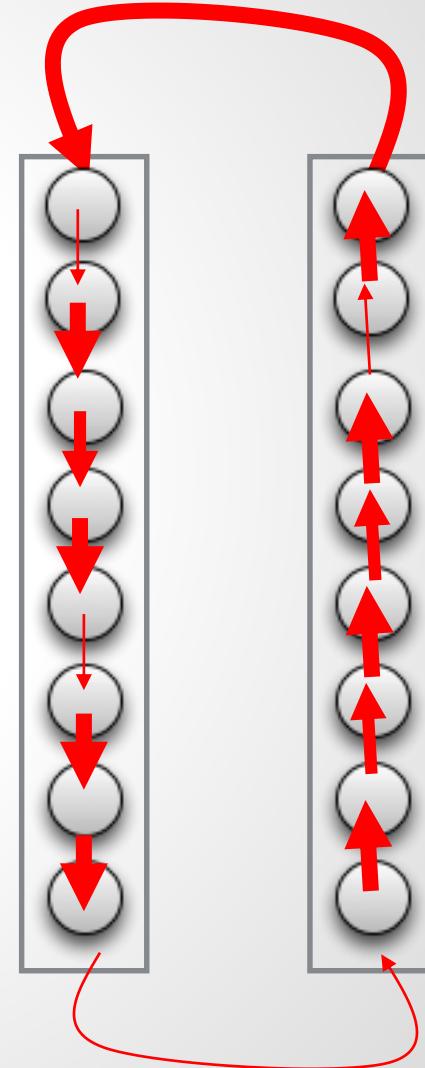


Example Learning Algorithm for Ring: Rotate Locally!

- ❑ Mantra of our algorithm: Rotate!
 - ❑ Rotate early, but not too early!
 - ❑ And: rotate **locally**

If current configuration is eliminated, go to **nearby configuration** (in directed manner: no frequent back and forth)!

Growing radius strategy: allow to move further only once amortized!



Example Learning Algorithm for Ring: Rotate Locally!

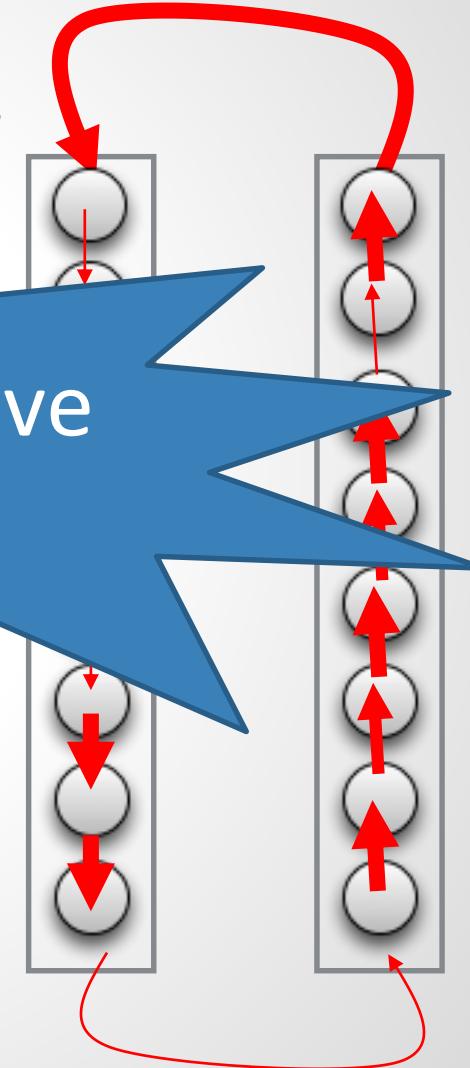
- ❑ Mantra of our algorithm: Rotate!

- ❑ Rotate locally but move easily
- ❑ Append to tail

$\log(n)$ -competitive
w.h.p.

If current configuration is eliminated, go to nearby configuration (in directed manner: no frequent back and forth)!

move to another one amortized!

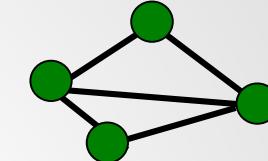
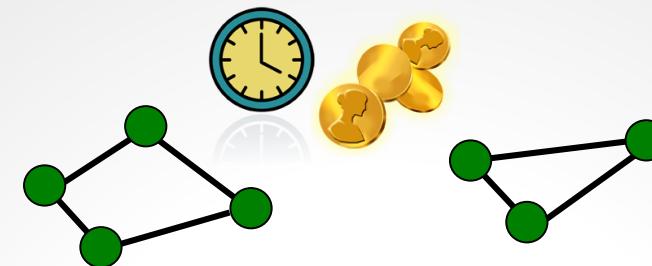


PART III:

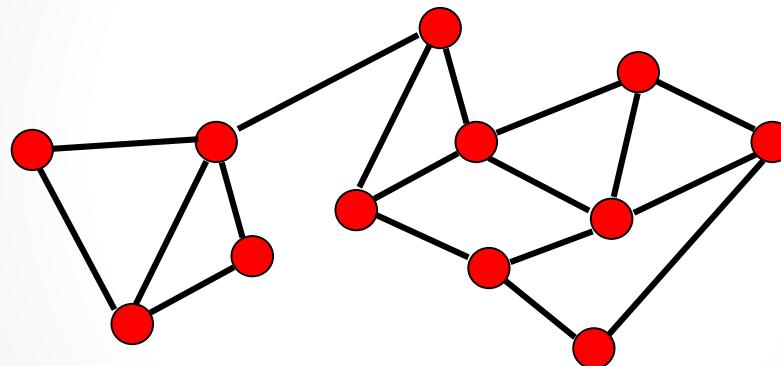
Embeddings over Time

A VNet seldom comes alone!

VNets



Infrastructure

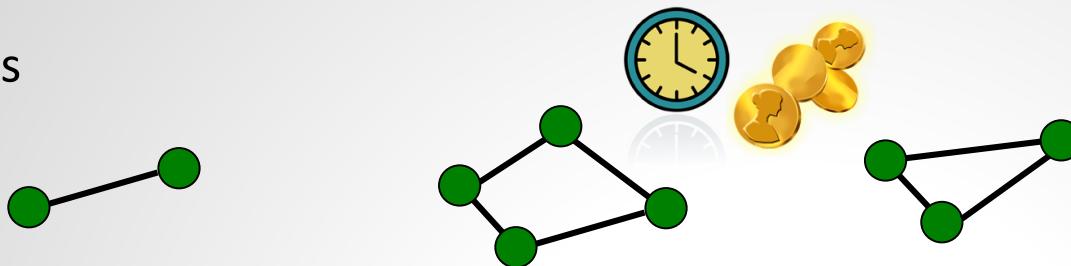


Time

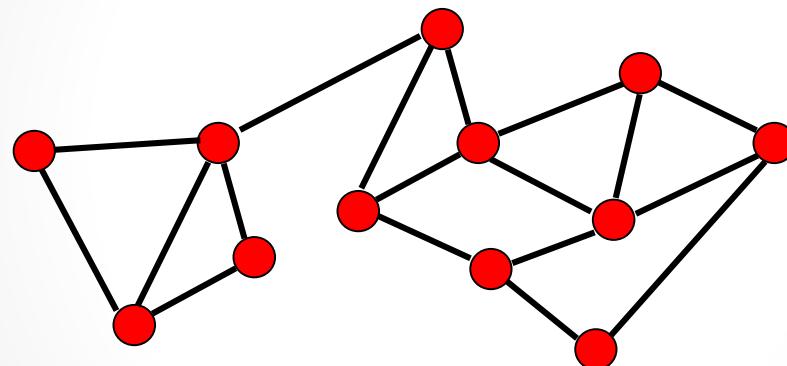
Which ones to accept
and embed?
Admission control!

A VNet seldom comes alone!

VNets



Infrastructure



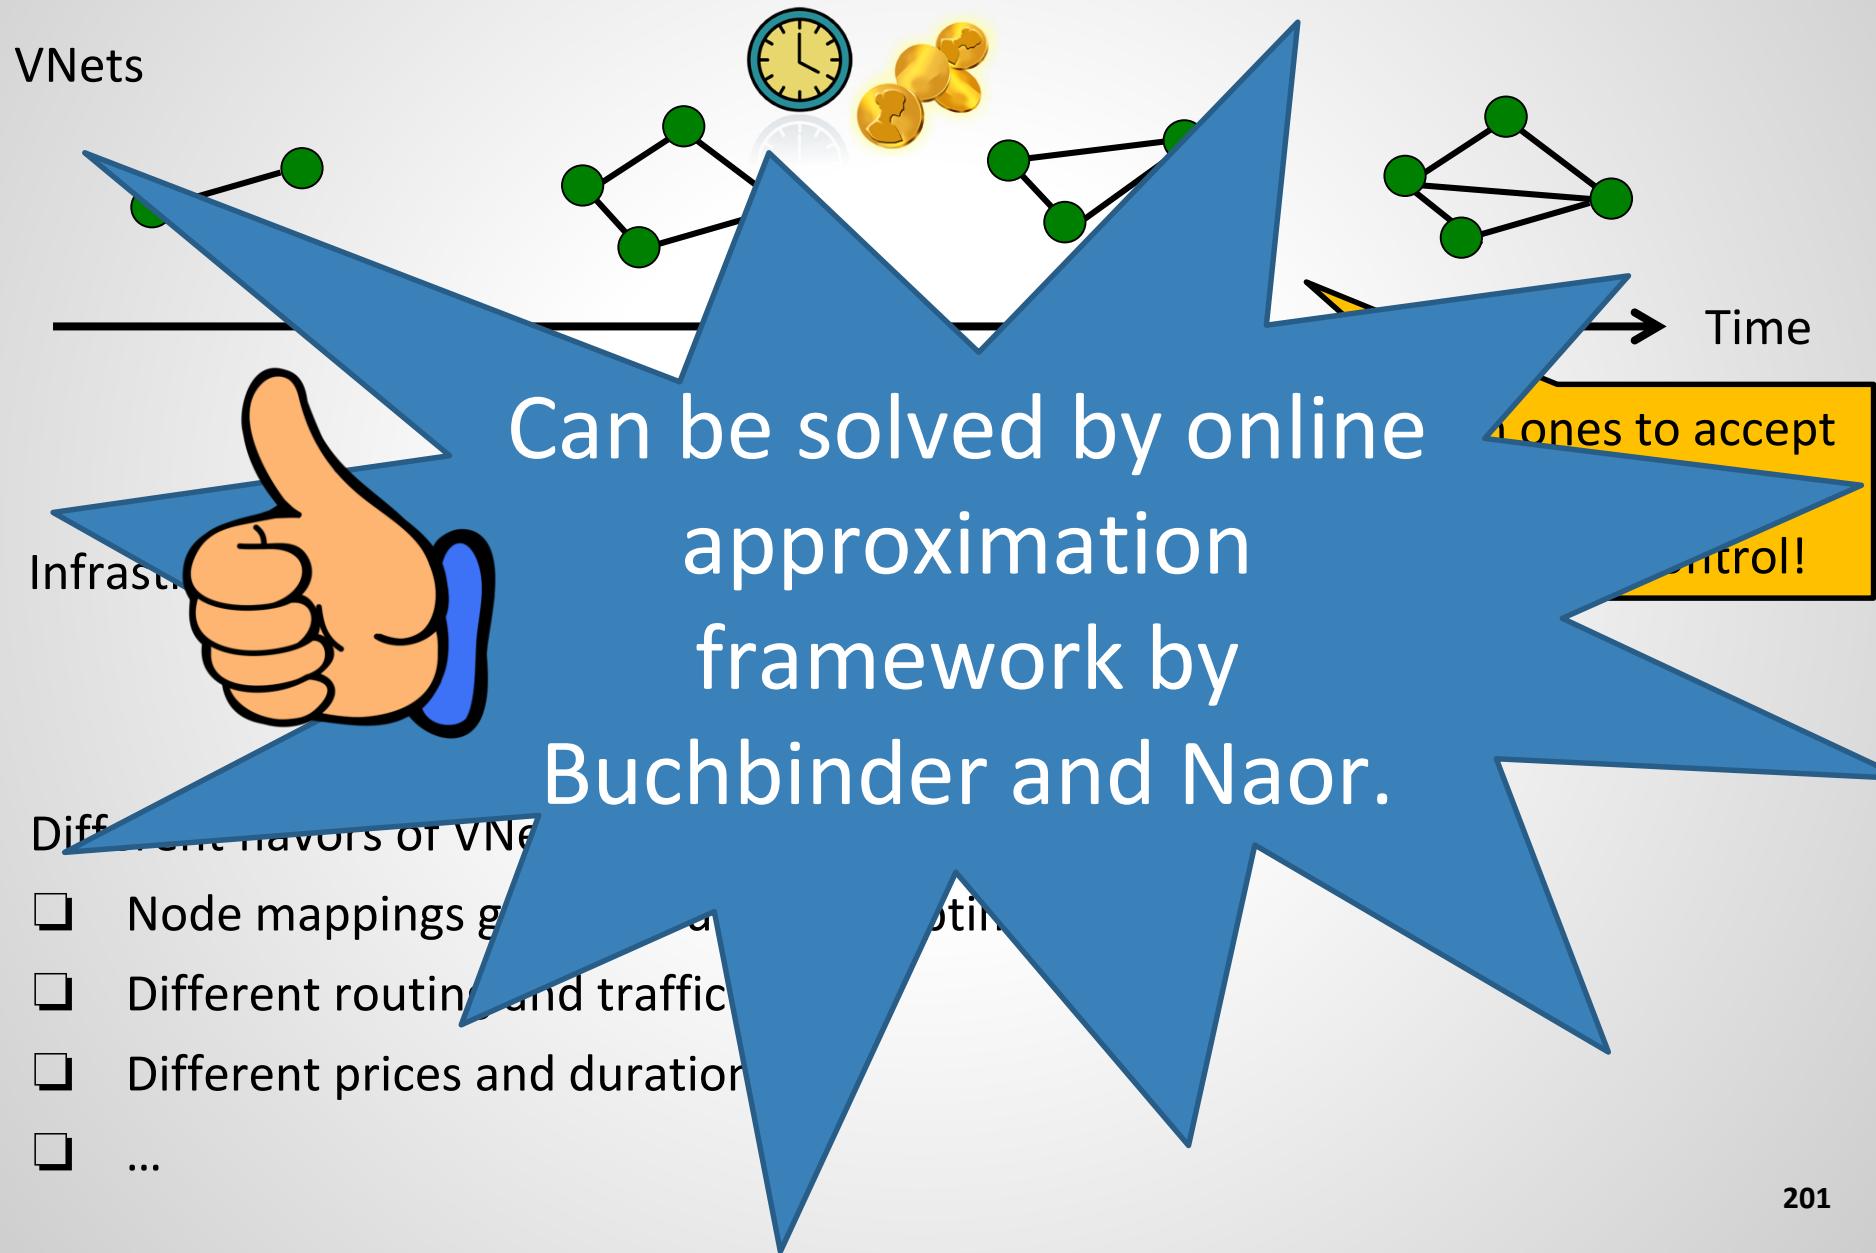
Time

Which ones to accept
and embed?
Admission control!

Different flavors of VNets:

- Node mappings given or subject to optimization
- Different routing and traffic models
- Different prices and durations
- ...

A VNet seldom comes alone!



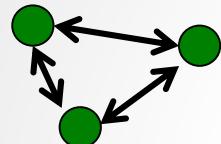
- Node mappings generate a lot of overhead
 - Different routing and traffic patterns
 - Different prices and duration
 - ...

Different VNet Flavors

□ Traffic models

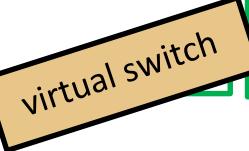
Customer Pipe

Traffic matrix:
Bandwidth per
VM pair (u,v)



Hose Model

Per VM
bandwidth:
polytope of traffic
matrices.



Aggregate Ingress

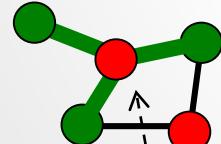
Only ingress
specified: e.g.,
support multicast
etc.



□ Routing models

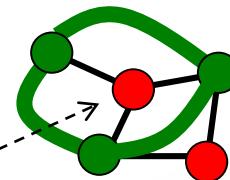
Tree

Steiner tree
embedding



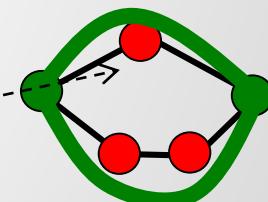
Single Path

Unsplittable
paths



Multi-Path

Splittable paths
(more capacity)



Relay costs: e.g., depending on packet rate

Applying Buchbinder&Naor

| | |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|
| $\begin{aligned} \min Z_j^T \cdot \mathbf{1} + X^T \cdot C & \text{ s.t.} \\ Z_j^T \cdot D_j + X^T \cdot A_j & \geq B_j^T \\ X, Z_j & \geq \mathbf{0} \end{aligned}$ | $\begin{aligned} \max B_j^T \cdot Y_j & \text{ s.t.} \\ A_j \cdot Y_j & \leq C \\ D_j \cdot Y_j & \leq \mathbf{1} \\ Y_j & \geq \mathbf{0} \end{aligned}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.

**Algorithm****Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).Upon the j th round:

1. $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure)
2. If $\gamma(j, \ell) < b_j$ then, (accept)
 - (a) $y_{j,\ell} \leftarrow 1$.
 - (b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

3. Else, (reject)
 - (a) $z_j \leftarrow b_j - \gamma(j, \ell)$.

Applying Buchbinder&Naor

| | |
|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| $\min Z_j^T \cdot \mathbf{1} + X^T \cdot C \text{ s.t.}$ $Z_j^T \cdot D_j + X^T \cdot A_j \geq B_j^T$ $X, Z_j \geq \mathbf{0}$ | $\max B_j^T \cdot Y_j \text{ s.t.}$ $A_j \cdot Y_j \leq C$ $D_j \cdot Y_j \leq \mathbf{1}$ $Y_j \geq \mathbf{0}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.



Formulate the packing
(dual) LP: Maximize profit
(Note: dynamic LP!)

Algorithm

Algorithm 1 The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the j th round:

1. $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure)
2. If $\gamma(j, \ell) < b_j$ then, (accept)
 - (a) $y_{j,\ell} \leftarrow 1$.
 - (b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

3. Else, (reject)
 - (a) $z_j \leftarrow b_j - \gamma(j, \ell)$.

Applying Buchbinder&Naor

| | |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------------------------------------|
| $\begin{aligned} & \min Z_j^T \cdot \mathbf{1} + X^T \cdot C \quad s.t. \\ & Z_j^T \cdot D_j + X^T \cdot A_j \geq B_j^T \\ & X, Z_j \geq \mathbf{0} \end{aligned}$ | $\begin{aligned} & \max B_j^T \cdot Y_j \quad s.t. \\ & A_j \cdot Y_j \leq C \\ & D_j \cdot Y_j \leq 1 \\ & Y_j \geq \mathbf{0} \end{aligned}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.

Algorithm



Algorithm 1 The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the j th round:

- $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure)
 - If $\gamma(j, \ell) < b_j$ then,
 - $y_{j,\ell} \leftarrow 1$.
 - For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)}/c_e} + \frac{1}{w(j,\ell)} \cdot (2^{A_{e,(j,\ell)}/c_e} - 1).$$

- (c) $z_j \leftarrow b_j - \gamma(j, \ell).$

3. Else, (reject)

(a) $z_j \leftarrow 0.$

Applying Buchbinder&Naor

| | |
|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| $\min Z_j^T \cdot \mathbf{1} + X^T \cdot C \text{ s.t.}$ $Z_j^T \cdot D_j + X^T \cdot A_j \geq B_j^T$ $X, Z_j \geq \mathbf{0}$ | $\max B_j^T \cdot Y_j \text{ s.t.}$ $A_j \cdot Y_j \leq C$ $D_j \cdot Y_j \leq \mathbf{1}$ $Y_j \geq \mathbf{0}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.

Algorithm



←

Buchbinder&Naor

Algorithm 1 The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the j th round:

1. $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure)
2. If $\gamma(j, \ell) < b_j$ then, (accept)
 - (a) $y_{j,\ell} \leftarrow 1$.
 - (b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

3. Else, (reject)
 - (a) $z_j \leftarrow b_j - \gamma(j, \ell)$.

Applying Buchbinder&Naor

| | |
|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| $\min Z_j^T \cdot \mathbf{1} + X^T \cdot C \text{ s.t.}$ $Z_j^T \cdot D_j + X^T \cdot A_j \geq B_j^T$ $X, Z_j \geq \mathbf{0}$ | $\max B_j^T \cdot Y_j \text{ s.t.}$ $A_j \cdot Y_j \leq C$ $D_j \cdot Y_j \leq \mathbf{1}$ $Y_j \geq \mathbf{0}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.

**Algorithm****Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).Upon the j th round:

1. $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure) oracle procedure optimal embedding!
2. If $\gamma(j, \ell) < b_j$ then, (accept)
 - (a) $y_{j,\ell} \leftarrow 1$.
 - (b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$
 - (c) $z_j \leftarrow b_j - \gamma(j, \ell)$.
3. Else, (reject)
 - (a) $z_j \leftarrow 0$.

Applying Buchbinder&Naor

| | |
|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| $\min Z_j^T \cdot \mathbf{1} + X^T \cdot C \text{ s.t.}$ $Z_j^T \cdot D_j + X^T \cdot A_j \geq B_j^T$ $X, Z_j \geq \mathbf{0}$ | $\max B_j^T \cdot Y_j \text{ s.t.}$ $A_j \cdot Y_j \leq C$ $D_j \cdot Y_j \leq \mathbf{1}$ $Y_j \geq \mathbf{0}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.

**Algorithm****Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).Upon the j th round:

1. $f_{j,\ell} \leftarrow \arg\min \{ \gamma(j, \ell) : f_{j,\ell} \in \Delta_j \}$ (oracle procedure)
2. If $\gamma(j, \ell) < b_j$ then, (accept)
 - (a) $y_{j,\ell} \leftarrow 1$.
 - (b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

Embedding cost vs profit?

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

- (c) $z_j \leftarrow b_j - \gamma(j, \ell)$.
3. Else, (reject)
 - (a) $z_j \leftarrow 0$.

Applying Buchbinder&Naor

| | |
|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| $\min Z_j^T \cdot \mathbf{1} + X^T \cdot C \text{ s.t.}$ $Z_j^T \cdot D_j + X^T \cdot A_j \geq B_j^T$ $X, Z_j \geq \mathbf{0}$ | $\max B_j^T \cdot Y_j \text{ s.t.}$ $A_j \cdot Y_j \leq C$ $D_j \cdot Y_j \leq \mathbf{1}$ $Y_j \geq \mathbf{0}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.



Algorithm

Algorithm 1 The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the j th round:

1. $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure)

2. If $\gamma(j, \ell) < b_j$ then, (accept)

(a) $y_{j,\ell} \leftarrow 1$.

(b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

(c) $z_j \leftarrow b_j - \gamma(j, \ell)$.

3. Else, (reject)

(a) $z_j \leftarrow 0$.

If cheap: accept and update primal variables
(always feasible solution)

Applying Buchbinder&Naor

| | |
|----------------------------------------------------------------------------------------------------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------|
| $\begin{aligned} \min Z_j^T \cdot \mathbf{1} + X^T \cdot C & \text{ s.t.} \\ Z_j^T \cdot D_j + X^T \cdot A_j & \geq B_j^T \\ X, Z_j & \geq \mathbf{0} \end{aligned}$ | $\begin{aligned} \max B_j^T \cdot Y_j & \text{ s.t.} \\ A_j \cdot Y_j & \leq C \\ D_j \cdot Y_j & \leq \mathbf{1} \\ Y_j & \geq \mathbf{0} \end{aligned}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.



Algorithm

Algorithm 1 The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).

Upon the j th round:

1. $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure)
2. If $\gamma(j, \ell) < b_j$ then, (accept)
 - (a) $y_{j,\ell} \leftarrow 1$.
 - (b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

- (c) $z_j \leftarrow b_j - \gamma(j, \ell)$.
3. Else, (reject)
 - (a) $z_j \leftarrow 0$.

Else reject

Applying Buchbinder&Naor

| | |
|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| $\min Z_j^T \cdot \mathbf{1} + X^T \cdot C \text{ s.t.}$ $Z_j^T \cdot D_j + X^T \cdot A_j \geq B_j^T$ $X, Z_j \geq \mathbf{0}$ | $\max B_j^T \cdot Y_j \text{ s.t.}$ $A_j \cdot Y_j \leq C$ $D_j \cdot Y_j \leq \mathbf{1}$ $Y_j \geq \mathbf{0}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.

**Algorithm****Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).Upon the j th round:

1. $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure) oracle procedure Computationally hard!
2. If $\gamma(j, \ell) < b_j$ then, (accept)
 - (a) $y_{j,\ell} \leftarrow 1$.
 - (b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$
 - (c) $z_j \leftarrow b_j - \gamma(j, \ell)$.
3. Else, (reject)
 - (a) $z_j \leftarrow 0$.

Applying Buchbinder&Naor

| | |
|--------------------------------------------------------------------------------------------------------------------------------|------------------------------------------------------------------------------------------------------------------|
| $\min Z_j^T \cdot \mathbf{1} + X^T \cdot C \text{ s.t.}$ $Z_j^T \cdot D_j + X^T \cdot A_j \geq B_j^T$ $X, Z_j \geq \mathbf{0}$ | $\max B_j^T \cdot Y_j \text{ s.t.}$ $A_j \cdot Y_j \leq C$ $D_j \cdot Y_j \leq \mathbf{1}$ $Y_j \geq \mathbf{0}$ |
| (I) | (II) |

Fig. 1: (I) The primal covering LP. (II) The dual packing LP.

**Algorithm****Algorithm 1** The General Integral (all-or-nothing) Packing Online Algorithm (GIPO).Upon the j th round:

1. $f_{j,\ell} \leftarrow \operatorname{argmin}\{\gamma(j, \ell) : f_{j,\ell} \in \Delta_j\}$ (oracle procedure)
2. If $\gamma(j, \ell) < b_j$ then, (accept)
 - (a) $y_{j,\ell} \leftarrow 1$.
 - (b) For each row e : If $A_{e,(j,\ell)} \neq 0$ do

$$x_e \leftarrow x_e \cdot 2^{A_{e,(j,\ell)} / c_e} + \frac{1}{w(j, \ell)} \cdot (2^{A_{e,(j,\ell)} / c_e} - 1).$$

3. Else, (reject)
 - (a) $z_j \leftarrow b_j - \gamma(j, \ell)$.

Computationally hard!

Use your favorite approximation algorithm! If competitive ratio ρ and approximation r , overall competitive ratio $\rho * r$.

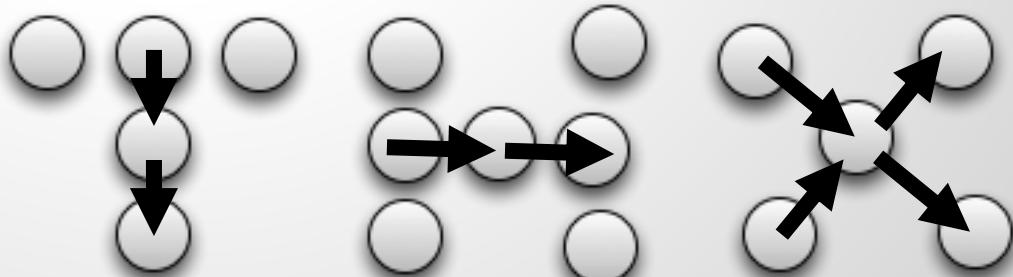
Online VNet Admission Control

Algorithm comes in 2 flavors:

- ❑ Bicriteria guarantees: Obtain **constant fraction** of the optimal benefit while **augmenting resources** by a logarithmic factor.
- ❑ Fractional guarantees resp. limited resource consumption:
The online algorithm achieves a logarithmic competitive ratio without resource augmentation, but either:
 - ❑ may **serve a fraction** of a request (associated benefit is assumed to be the same fraction of the benefit of the request)
 - ❑ the allowed traffic patterns of a request **consumes at most a logarithmic fraction** of every resource (in which case the algorithm rejects the request or fully embeds it)

Summary

- ❑ Predictable performance requires isolation of all resources
- ❑ Static embeddings
- ❑ Reconfiguring embeddings
- ❑ Embeddings over time



Further Reading

Network Hypervisor Performance:

- [Logically Isolated, Actually Unpredictable? Measuring Hypervisor Performance in Multi-Tenant SDNs](#) Arsany Basta, Andreas Blenk, Wolfgang Kellerer, and Stefan Schmid. ArXiv Technical Report, May 2017.

Virtual Network Embedding:

- [Beyond the Stars: Revisiting Virtual Cluster Embeddings](#) Matthias Rost, Carlo Fuerst, and Stefan Schmid. ACM SIGCOMM Computer Communication Review (**CCR**), July 2015.
- [Service Chain and Virtual Network Embeddings: Approximations using Randomized Rounding](#) Matthias Rost and Stefan Schmid. ArXiv Technical Report, April 2016.
- [Charting the Complexity Landscape of Waypoint Routing](#) Saeed Akhoondian Amiri, Klaus-Tycho Foerster, Riko Jacob, and Stefan Schmid. ArXiv Technical Report, May 2017
- [Competitive and Deterministic Embeddings of Virtual Networks](#) Guy Even, Moti Medina, Gregor Schaffrath, and Stefan Schmid. Journal Theoretical Computer Science (**TCS**), Elsevier, 2013.

Online Collocation:

- [Online Balanced Repartitioning](#) Chen Avin, Andreas Loukas, Maciej Pacut, and Stefan Schmid. 30th International Symposium on Distributed Computing (**DISC**), Paris, France, September 2016.
- [Competitive Clustering of Stochastic Communication Patterns on the Ring](#) Chen Avin, Louis Cohen, and Stefan Schmid. 5th International Conference on Networked Systems (**NETYS**), Marrakech, Morocco, May 2017

Network Design:

- [Demand-Aware Network Designs of Bounded Degree](#) Chen Avin, Kaushik Mondal, and Stefan Schmid. ArXiv Technical Report, May 2017.

Further Reading

Network Hypervisor

Hypervisor performance and interference analysis

- [Logically Isolated, Actually Unpredictable? Measuring Hypervisor Performance in Multi-Tenant SDNs](#) Arsany Basta, Andreas Blenk, Wolfgang Kellerer, and Stefan Schmid. ArXiv Technical Report, May 2017.

Virtual Network Embedding:

- [Beyond the Star vs Hose embedding](#) Matthias Rost, SIGCOMM Computer Communication Review (CCR), July 2015. Including decomposability (study of integrality gap)
- [Service Chain and Virtual Network Embeddings: Approximations using Randomized Rounding](#) Matthias Rost and Stefan Schmid. ArXiv Technical Report, April 2016.
- [Charting the Complexity Landscape of Waypoint Routing](#) Saeed Akhoondian Amiri, Klaus-Tycho Foerster, Riko Jacob, and Stefan Schmid. ArXiv Technical Report, May 2017 Relationship to disjoint paths
- [Competitive and Deterministic Embeddings of Virtual Networks](#) Guy Even, Moti Medina, Gregor Schaffrath, and Stefan Schmid. Journal Theoretical Computer Science (TCS), Elsevier, 2013. Online primal-dual

Online Collocation:

- [Online Balanced Repartitioning](#) Chen Avin, Andreas Loukas, Maciej Pacut, and Stefan Schmid. 30th International Symposium on Distributed Computing (DISC), Paris, France. Adaptive VM migration to minimize communication costs
- [Competitive Clustering of Stochastic Communication Patterns on the Ring](#) Chen Avin, Louis Cohen, and Stefan Schmid. 5th International Conference on Networked Systems (NETYS), Marrakech, Morocco, May 2017

Network Design:

- [Demand-Aware Network Designs of Bounded Degree](#) Chen Avin, Kaushik Mondal, and Stefan Schmid. ArXiv Technical Report, May 2017. A model and analysis motivated by ProjectToR