

Synthesis of dependable and self-driving communication networks

Stefan Schmid
@FRIDA 2025

Acknowledgements:



Networks:

Critical Infrastructure

- If networks break, it can have knock-on effects
- For example, Facebook outage in 2021: not only took down their social networking site, but also Instagram, WhatsApp, ...
- ... and their own internal systems, which manage the doors: engineers had to break into their own buildings to bring the network back up

The New York Times

Gone in Minutes, Out for Hours: Outage Shakes Facebook

When apps used by billions of people worldwide blinked out, lives were disrupted, businesses were cut off from customers — and some Facebook employees were locked out of their offices.

 Share full article    884



Facebook's internal communications platform, Workplace, was also taken out, leaving most employees unable to do their jobs. Kelsey McClellan for The New York Times

Credits: Nate Foster

The Challenge: Most Outages due to Human Errors

Human Errors

Countries disconnected

Data Centre ▶ Networks

Google routing blunder sent Japan's Internet dark on Friday

Another big BGP blunder

By Richard Chirgwin 27 Aug 2017 at 22:35

40 SHARE ▾

Last Friday, someone in Google fat-thumbed a border gateway protocol (BGP) advertisement and sent Japanese Internet traffic into a black hole.

The trouble began when The Chocolate Factory "leaked" a big route table to Verizon, the result of which was traffic from Japanese giants like NTT and KDDI was sent to Google on the expectation it would be treated as transit.

Passengers stranded

British Airways' latest Total Inability To Support Upwardness of Planes* caused by Amadeus system outage

Stuck on the ground awaiting a load sheet? Here's why

By Gareth Corfield 19 Jul 2018 at 11:16

109 SHARE ▾



RA flights around the world were canceled as a result of the Amadeus outage

Even 911 affected

Officials: Human error to blame in Minn. 911 outage

According to a press release, CenturyLink told department of public safety that human error by an employee of a third party vendor was to blame for the outage

Aug 16, 2018

Duluth News Tribune

SAINT PAUL, Minn. — The Minnesota Department of Public Safety Emergency Communication Networks division was told by its 911 provider that an Aug. 1 outage was caused by human error.

Even tech-savvy companies struggle:



Slide credits: Nate Foster and Laurent Vanbever


Mainly:
human
errors!

A Reason: Complexity

Especially Under Failures (Policy Compliance)

Example: BGP in

Microsoft
datacenter




A Reason: Complexity

Especially Under Failures (Policy Compliance)

Example: BGP in


Microsoft
datacenter



A Reason: Complexity

Especially Under Failures (Policy Compliance)


Example: BGP in
Microsoft
datacenter



A Reason: Complexity

Especially Under Failures (Policy Compliance)


Example: BGP in
Microsoft
datacenter



A Reason: Complexity

Especially Under Failures (Policy Compliance)


Example: BGP in
Microsoft
datacenter



A Reason: Complexity

Especially Under Failures (Policy Compliance)

Example: BGP in
Microsoft
datacenter




If link (G,X) fails and traffic from G is rerouted via Y and C to X:
X announces (does not block) G and H as it comes from C. (Note: BGP.)

A Reason: Complexity

Especially Under Failures (Policy Compliance)

Example: BGP in
Microsoft
datacenter




If link (G,X) fails and traffic from G is rerouted via Y and C to X:
X announces (does not block) G and H as it comes from C. (Note: BGP.)


Particularly Difficult

Fast Rerouting


Particularly Difficult Local Fast Rerouting



Particularly Difficult Local Fast Rerouting



Particularly Difficult Local Fast Rerouting




Information at Switch for
Local Decision Making?




Information at Switch for Local Decision Making?

- Nodes locally store a forwarding Match -> Action table




Information at Switch for Local Decision Making?

→ The **Packet Header** (e.g., source, destination)




Information at Switch for Local Decision Making?

→ The **Import** of the received packet



Information at Switch for Local Decision Making?

→ Which **incident links** failed



Objective

What-if Analysis & Synthesis

- ... for **robust networks** tolerating many link failures.
- **Verification:** Are the current forwarding rules policy compliant (reachability, waypoint traversal) even under failures?
- **Synthesis:** Can we pre-install local fast failover rules which ensure reachability under multiple failures?
- In general: How **many failures** can be tolerated by static forwarding tables?

Two fundamental

Notions of Resilience

Ideal resilience

Given a k -connected graphs, fast reroute can tolerate *any $k-1$ link failures.*

Perfect resilience

Fast reroute can tolerate any failures as long as the underlying network is *physically connected.*

→ What is the difference? Which is stronger?

A big open challenge


Ideal Resilience

- Given a k-connected network: **how many link failures** can a fast re-routing mechanism tolerate? **Conjecture:** $k-1$.
- Assume: cannot change header, but can match import, src and dst

A big open challenge

Ideal Resilience


- Given a k -connected network: **how many link failures** can a fast re-routing mechanism tolerate? **Conjecture:** $k-1$.
- Assume: cannot change header, but can match import, src and dst



A big open challenge

Ideal Resilience

- Given a k -connected network: **how many link failures** can a fast re-routing mechanism tolerate? **Conjecture:** $k-1$.
- Assume: cannot change header, but can match inport, src and dst



Yes! k disjoint paths: try one after the other, routing *back to source* each time.


A big open challenge

Ideal Resilience

- Given a k -connected network: how many link failures can a fast re-routing mechanism tolerate? Conjecture: $k-1$.
- Assume: cannot change header, but can match import, ~~src~~ and dst




What if I cannot
match source?!
Open conjecture.



State-of-the-Art Approach for Ideal Resilience

Spanning Arborescences




- Fact: k -connected network has **k -arborecence decomposition**
- Basically **disjoint** spanning trees **directed** to destination

State-of-the-Art Approach for Ideal Resilience Spanning Arborescences

Arborescences

1




- Fact: k -connected network has **k -arborescence decomposition**
- Basically **disjoint** spanning trees **directed** to destination

State-of-the-Art Approach for Ideal Resilience Spanning Arborescences

Arborescences

1 2




- Fact: k -connected network has k -arborescence decomposition
- Basically **disjoint** spanning trees **directed** to destination

State-of-the-Art Approach for Ideal Resilience Spanning Arborescences

Arborescences

1 2 3




- Fact: k -connected network has k -arborecence decomposition
- Basically **disjoint** spanning trees **directed** to destination

State-of-the-Art Approach for Ideal Resilience Spanning Arborescences

Arborescences

- 1
- 2
- 3
- 4




- Fact: k -connected network has k -arbrescence decomposition
- Basically **disjoint** spanning trees **directed** to destination

State-of-the-Art Approach for Ideal Resilience Spanning Arborescences

Arborescences

- 1
- 2
- 3
- 4




- Fact: k -connected network has k -arborecence decomposition
- Basically disjoint spanning trees directed to destination


($k/2-1$)-resilient with circular
Arborescence Routing

Arborescences

- 1
- 2
- 3
- 4





($k/2-1$)-resilient with circular
Arborescence Routing



→ Try arborescences **in order**

($k/2-1$)-resilient with circular
Arborescence Routing


Arborescences



→ Try arborescences **in order**

($k/2-1$)-resilient with circular
Arborescence Routing

Arborescences




→ Try arborescences **in order**

($k/2-1$)-resilient with circular
Arborescence Routing

Arborescences

- 1
- 2
- 3
- 4




→ Try arborescences **in order**

→ **$k/2-1$ resilient:** link failure affects at most 2 arborescences

Research Challenges

- Complexity of **verifying** resilience and policy-compliance?
- **Algorithms** for synthesizing resilient fast reroute mechanisms?
- Application to **specific protocols**, like **MPLS** or Segment Routing?




A General Solution: Automation Synthesis with BDDs

- ...> **Binary decision diagrams (BDDs)** allow us to synthesize resilient routings
 - ...> ... or to **repair**
- ...> Attractive: **all solutions**, compactly represented
 - ...> Supports **operator preferences**!
 - ...> Better alternative to e.g. **ILPs**
- ...> Still somewhat **slow**


A General Solution: Automation Synthesis with BDDs

- Binary decision diagrams (BDDs) allow us to synthesize resilient routings
 - ... or to repair
 - Attractive: all solutions, compactly represented
 - Supports operator preferences!
 - Better alternative to e.g. ILPs
 - Still somewhat slow

Network:



BDD 2-resilient routing:s




A General Solution: Automation Synthesis with BDDs


- Binary decision diagrams (BDDs) allow us to synthesize resilient routings
 - ... or to repair
 - Attractive: all solutions, compactly represented
 - Supports operator preferences!
 - Better alternative to e.g. ILPs
 - Still somewhat slow

For specific protocols we can be faster!

Network:




BDD 2-resilient
routing:s



Faster for specific protocol:

MPLS Fast Reroute (FRR)


→ Forwarding based on **top label** of label **stack**



Faster for specific protocol:

MPLS Fast Reroute (FRR)

→ Forwarding based on **top label** of label **stack**




Default
routing of
two flows


Faster for specific protocol:

MPLS Fast Reroute (FRR)

→ Forwarding based on **top label** of label **stack**



Default
routing of
two flows




One failure:
push 30: route
around (v_2, v_3)


Faster for specific protocol:

MPLS Fast Reroute (FRR)

→ Forwarding based on **top label** of label **stack**



Default
routing of
two flows




One failure:
push 30: route
around (v_2, v_3)


Faster for specific protocol:

MPLS Fast Reroute (FRR)


→ Multiple link failures: simply recursive



Original
Routing



One failure:
push 30: route
around (v_2, v_3)



Two failures:
first push 30: route
around (v_2, v_3)

Push recursively
40: route around
 (v_2, v_6)

Faster for specific protocol:


MPLS Fast Reroute (FRR)

- Specific structure of MPLS networks can be exploited for fast what-if analysis: it's a **stack machine**
- Can use the result by **Büchi**: set of all reachable configurations of **pushdown automaton** is regular set
- We hence simply use **Nondeterministic Finite Automata** when reasoning about the pushdown automata
- The resulting regular operations are all **polynomial time**



Julius Richard Büchi
1924-1984
Swiss logician


Example: AalWiNes Tool



Tool: <https://demo.aalwines.cs.aau.dk/>

Youtube: https://www.youtube.com/watch?v=mvXAn9i7_00

Can cover many policies!



Sysadmin responsible for:


- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?
- **Policy:** Is it ensured that traffic from A to B never goes via C?
- **Waypoint enforcement:** Is it ensured that traffic from A to B is always routed via a node C (e.g., intrusion detection system or a firewall)?

... and everything under multiple failures!

Opportunity: Automation for Self-Driving Networks



FT	In-I	In-Label	Out-I	op
τ_{v_1}	m_1	\perp	(v_1, v_2)	$push(10)$
	m_2	\perp	(v_1, v_2)	$push(20)$
τ_{v_2}	(v_1, v_2)	10	(v_2, v_3)	$swap(11)$
	(v_1, v_2)	20	(v_2, v_3)	$swap(21)$
τ_{v_3}	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_2, v_3)	21	(v_3, v_8)	$swap(22)$
τ_{v_4}	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_2, v_3)	21	(v_3, v_8)	$swap(22)$
τ_{v_4}	(v_3, v_4)	12	out_1	pop
	(v_2, v_3)	40	(v_5, v_6)	pop
τ_{v_5}	(v_2, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	30	(v_6, v_7)	$swap(31)$
τ_{v_6}	(v_5, v_6)	61	(v_6, v_7)	$swap(62)$
	(v_5, v_6)	71	(v_6, v_7)	$swap(72)$
τ_{v_7}	(v_6, v_7)	31	(v_7, v_8)	pop
	(v_6, v_7)	62	(v_7, v_8)	$swap(11)$
τ_{v_8}	(v_6, v_7)	72	(v_7, v_8)	$swap(22)$
	(v_5, v_8)	22	out_2	pop
τ_{v_8}	(v_7, v_8)	22	out_2	pop




local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_6)	30	(v_2, v_5)	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
τ'_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$

Router configurations
(Cisco, Juniper, etc.)


Opportunity: Automation for Self-Driving Networks



FT	In-I	In-Label	Out-I	op
τ_{v_1}	m_1	\perp	(v_1, v_2)	$push(10)$
	m_2	\perp	(v_1, v_2)	$push(20)$
τ_{v_2}	(v_1, v_2)	10	(v_2, v_3)	$swap(11)$
	(v_1, v_2)	20	(v_2, v_3)	$swap(21)$
τ_{v_3}	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_2, v_3)	21	(v_3, v_8)	$swap(22)$
	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_2, v_3)	21	(v_3, v_8)	$swap(22)$
τ_{v_4}	(v_3, v_4)	12	out_1	pop
	(v_2, v_3)	40	(v_5, v_6)	pop
τ_{v_5}	(v_2, v_3)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	61	(v_6, v_7)	$swap(62)$
	(v_5, v_6)	71	(v_6, v_7)	$swap(72)$
τ_{v_7}	(v_6, v_7)	31	(v_7, v_3)	pop
	(v_6, v_7)	62	(v_7, v_3)	$swap(11)$
	(v_6, v_7)	72	(v_7, v_8)	$swap(22)$
τ_{v_8}	(v_5, v_6)	22	out_2	pop
	(v_7, v_8)	22	out_2	pop



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_6)	30	(v_2, v_5)	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
τ'_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$




$pX \Rightarrow qXX$

$pX \Rightarrow qYX$

$qY \Rightarrow rYY$

$rY \Rightarrow r$

$rX \Rightarrow pX$




Formal language
which supports
automated analysis

Opportunity: Automation for Self-Driving Networks


What if?!



FT	In-I	In-Label	Out-I	op
τ_{v_1}	m_1	\perp	(v_1, v_2)	$push(10)$
	m_2	\perp	(v_1, v_2)	$push(20)$
τ_{v_2}	(v_1, v_2)	10	(v_2, v_3)	$swap(11)$
	(v_1, v_2)	20	(v_2, v_3)	$swap(21)$
τ_{v_3}	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_2, v_3)	21	(v_3, v_8)	$swap(22)$
	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_2, v_3)	21	(v_3, v_8)	$swap(22)$
τ_{v_4}	(v_3, v_4)	12	out_1	pop
	(v_2, v_3)	40	(v_5, v_6)	pop
τ_{v_5}	(v_2, v_3)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	61	(v_6, v_7)	$swap(62)$
	(v_5, v_6)	71	(v_6, v_7)	$swap(72)$
τ_{v_7}	(v_6, v_7)	31	(v_7, v_3)	pop
	(v_6, v_7)	62	(v_7, v_3)	$swap(11)$
	(v_6, v_7)	72	(v_7, v_8)	$swap(22)$
τ_{v_8}	(v_5, v_8)	22	out_2	pop
	(v_7, v_8)	22	out_2	pop



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_6)	30	(v_2, v_5)	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
τ'_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$




$pX \Rightarrow qXX$

$pX \Rightarrow qYX$

$qY \Rightarrow rYY$

$rY \Rightarrow r$

$rX \Rightarrow pX$







Formal language
which supports
automated analysis

→ Would be nice but synthesis slow.

Opportunity: Automation for Self-Driving Networks




$$\begin{aligned} pX &\Rightarrow qXX \\ pX &\Rightarrow qYX \\ qY &\Rightarrow rYY \\ rY &\Rightarrow r \\ rX &\Rightarrow pX \end{aligned}$$



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_6)	30	(v_2, v_5)	$push(40)$
τ'_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$

Opportunity: Automation for Self-Driving Networks

$$\begin{aligned} pX &\Rightarrow qXX \\ pX &\Rightarrow qYX \\ qY &\Rightarrow rYY \\ rY &\Rightarrow r \\ rX &\Rightarrow pX \end{aligned}$$



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	30	(v_2, v_5)		$push(40)$
τ'_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$

Fast Synthesis: FM+ML

- ...→ **Ideally ML+FM**: guarantees from formal methods, performance from ML
- ...→ For example: synthesize with ML then verify with formal methods
- ...→ Examples: DeepMPLS, DeepBGP, ...
- ...→ ***Self-driving networks!***



Thank you!

[A Survey of Fast-Recovery Mechanisms in Packet-Switched Networks](#)

Marco Chiesa, Andrzej Kamisinski, Jacek Rak, Gabor Retvari, and Stefan Schmid.
IEEE Communications Surveys and Tutorials (**COMST**), 2021.

[AalWiNes: A Fast and Quantitative What-If Analysis Tool for MPLS Networks](#)

Peter Gjøl Jensen, Morten Konggaard, Dan Kristiansen, Stefan Schmid, Bernhard Clemens Schrenk, and Jiri Srba.

16th ACM International Conference on emerging Networking EXperiments and Technologies (**CoNEXT**),
Barcelona, Spain, December 2020.

[A Tight Characterization of Fast Failover Routing: Resiliency to Two Link Failures is Possible](#)

Wenkai Dai, Klaus-Tycho Foerster, and Stefan Schmid.

35th ACM Symposium on Parallelism in Algorithms and Architectures (**SPAA**), Orlando, Florida, USA,
June 2023.

[On the Price of Locality in Static Fast Rerouting](#)

Klaus-Tycho Foerster, Juho Hirvonen, Yvonne-Anne Pignolet, Stefan Schmid, and Gilles Tredan.
52nd IEEE/IFIP International Conference on Dependable Systems and Networks (**DSN**), Baltimore,
Maryland, USA, June 2022.

[SyPer: Synthesis of Perfectly Resilient Local Fast Rerouting Rules for Highly Dependable Networks](#)

Csaba Györgyi, Kim G. Larsen, Stefan Schmid, and Jiri Srba.

IEEE Conference on Computer Communications (**INFOCOM**), Vancouver, Canada, May 2024.

[DeepMPLS: Fast Analysis of MPLS Configurations Using Deep Learning](#)

Fabien Geyer and Stefan Schmid.

IFIP Networking, Warsaw, Poland, May 2019.

[Latte: Improving the Latency of Transiently Consistent Network Update Schedules](#)

Mark Glavind, Niels Christensen, Jiri Srba, and Stefan Schmid.

38th International Symposium on Computer Performance, Modeling, Measurements and Evaluation
(**PERFORMANCE**) and ACM Performance Evaluation Review (**PER**), Milan, Italy, November 2020.

[Model-Based Insights on the Performance, Fairness, and Stability of BBR](#)

Simon Scherrer, Markus Legner, Adrian Perrig, and Stefan Schmid.

ACM Internet Measurement Conference (**IMC**), Nice, France, October 2022.

Slides
available:

