

Revolutionizing Datacenter Networks via Reconfigurable Topologies

Stefan Schmid (TU Berlin)

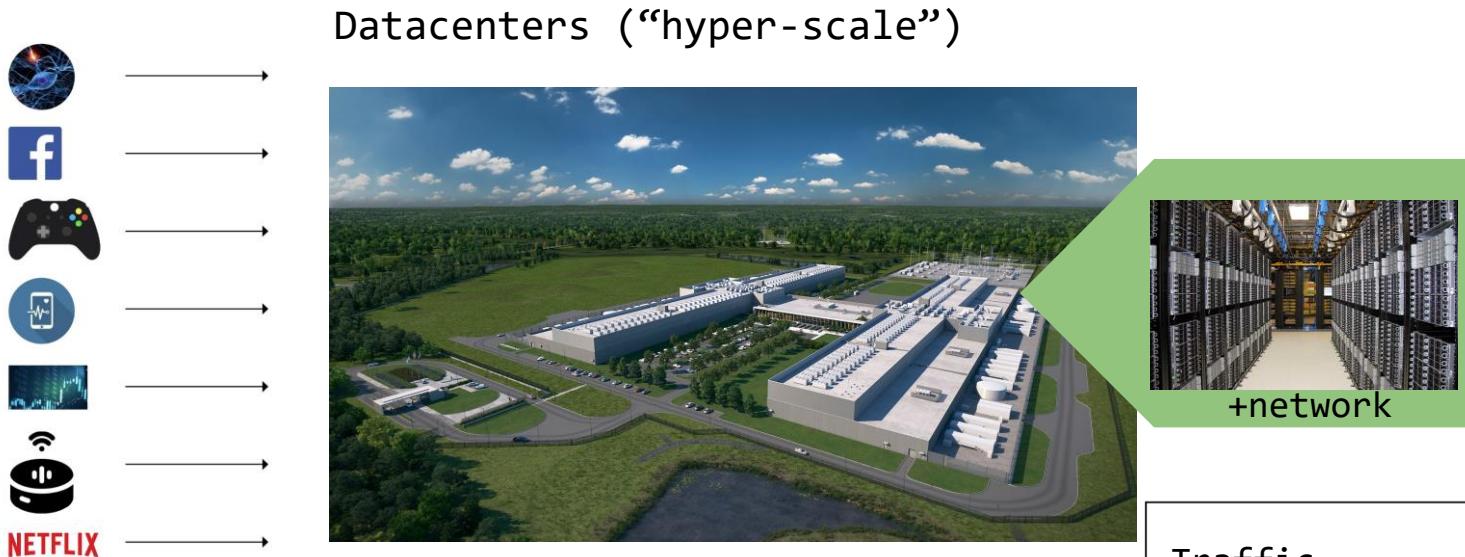
“We cannot direct the wind,
but we can adjust the sails.”

(Folklore)

Acknowledgements:



We live in the age of Distributed Computation



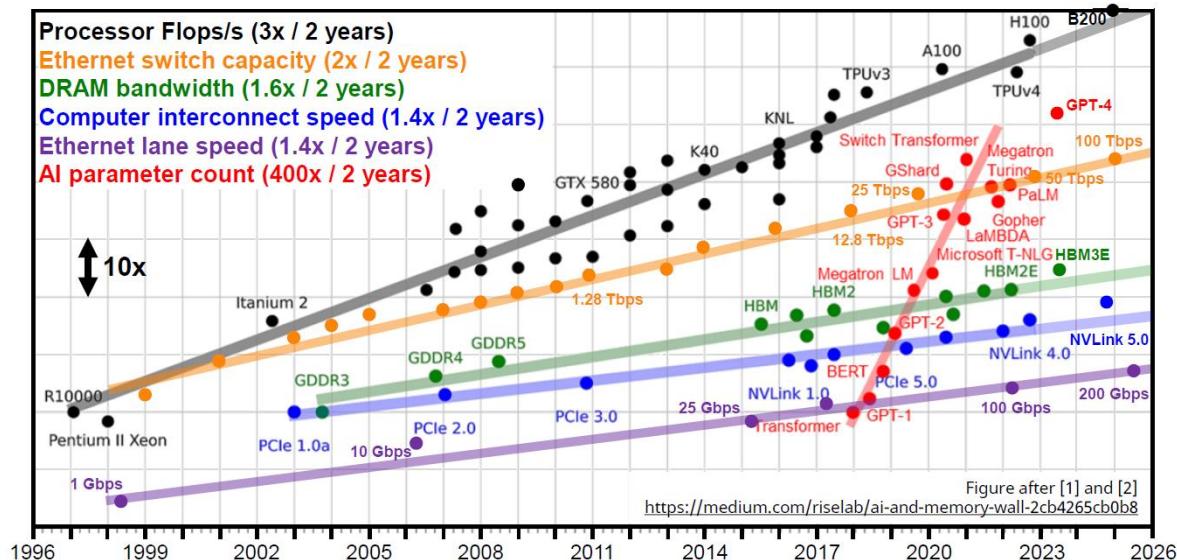
Interconnecting networks:
a **critical infrastructure**
of our digital society.

Traffic
Growth

Source: Facebook

Technological Trends

Increasing Gap Between Compute and Network

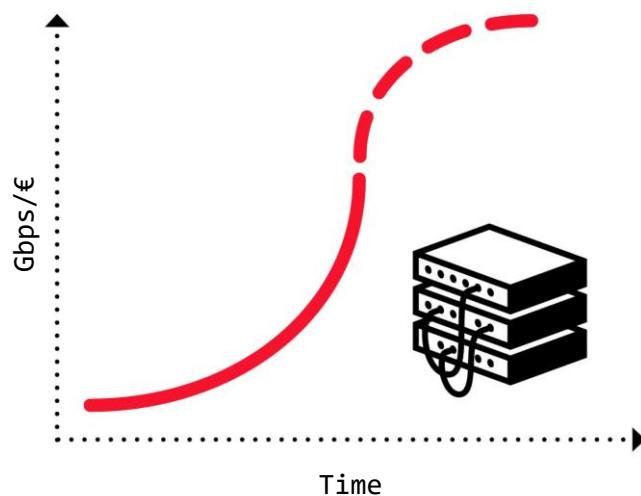


Credits: Nicola Calabretta

The Problem

Huge Infrastructure, Inefficient Use

- Hence: more equipment,
larger networks
- Resource intensive and:
inefficient



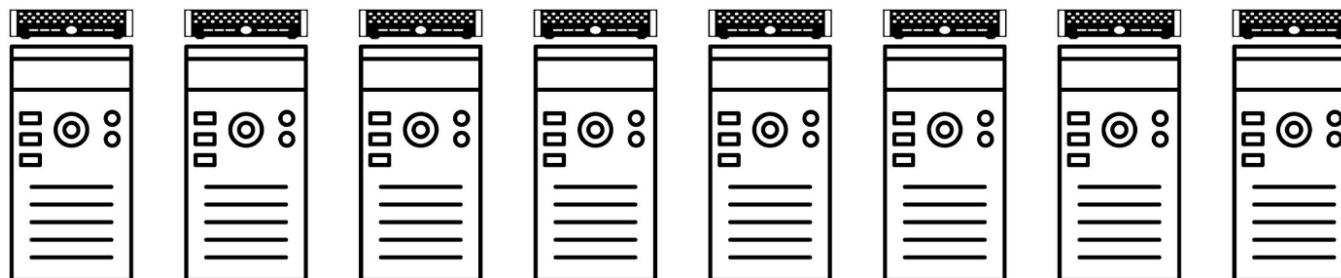
Credits: Paolo Costa, 2019

Annoying for companies,
opportunity for researchers!

An Inefficiency

Fixed and Demand-Oblivious Topology

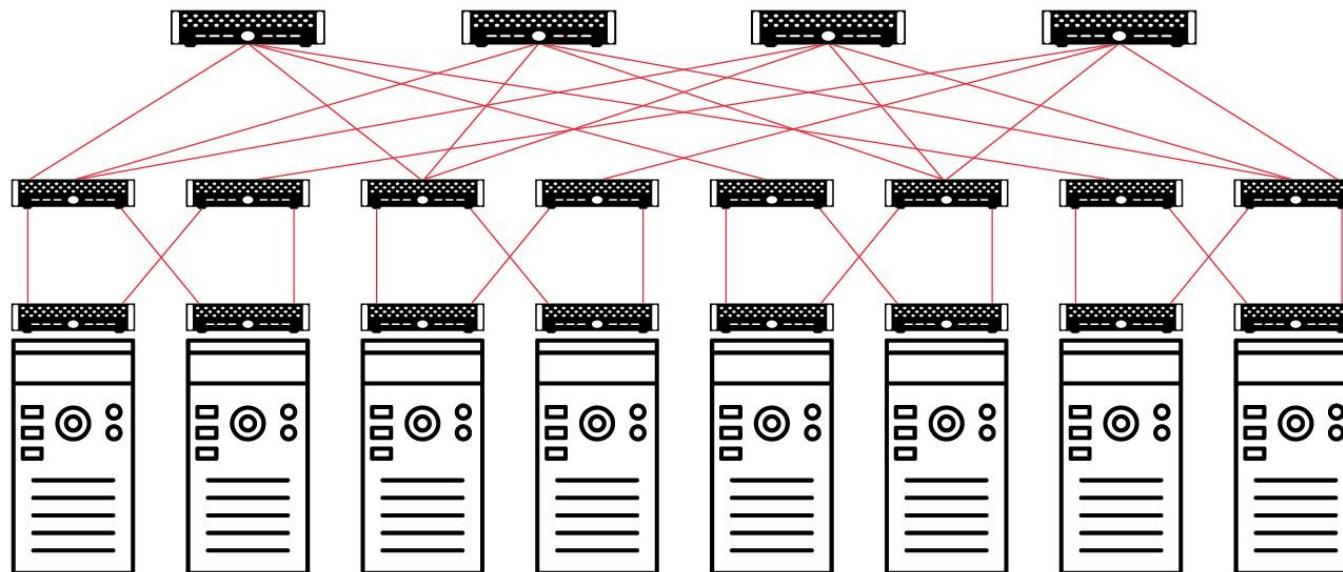
How to interconnect? Focus on this talk: **scale-out network**.



An Inefficiency

Fixed and Demand-Oblivious Topology

- Example: fat-tree topology (**bi-regular**)
 - 2 types of switches: top-of-rack (ToR) connect to hosts, additional switches connecting switches to increase throughput

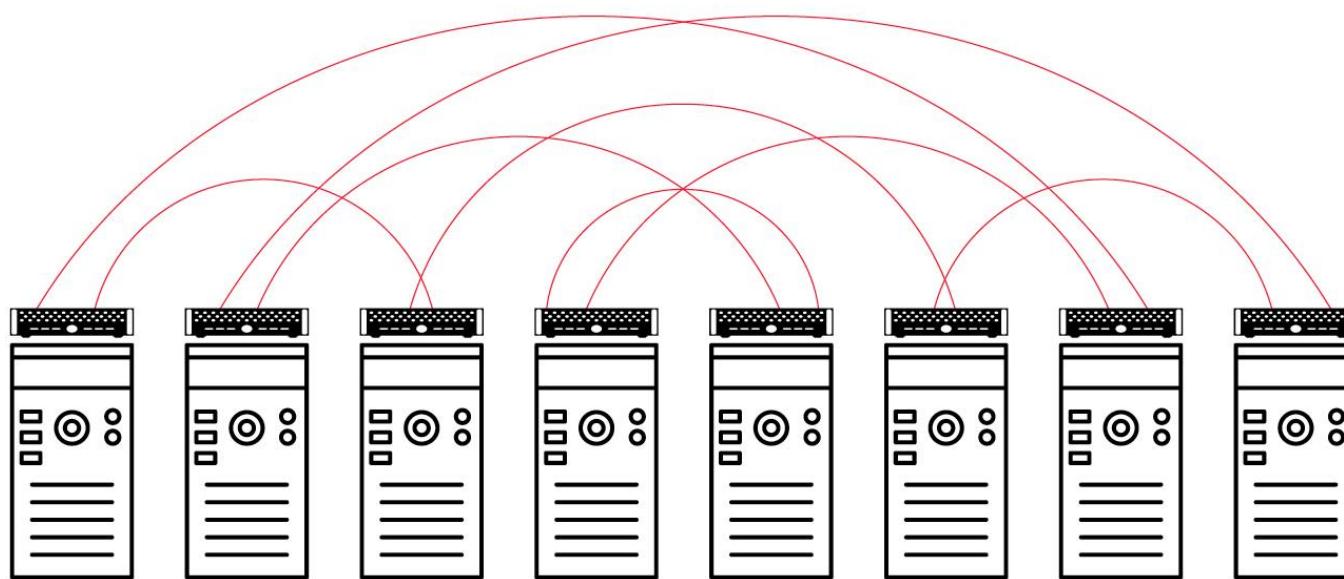


An Inefficiency

Fixed and Demand-Oblivious Topology

→ Example: expander topology (**uni-regular**)

→ Only 1 type of switches:
lower installation and management overheads



An Inefficiency

Fixed and Demand-Oblivious Topology

→ Example: expander topology (**uni-regular**)

→ Only 1 type of switches:
lower installation and management overheads



An Inefficiency

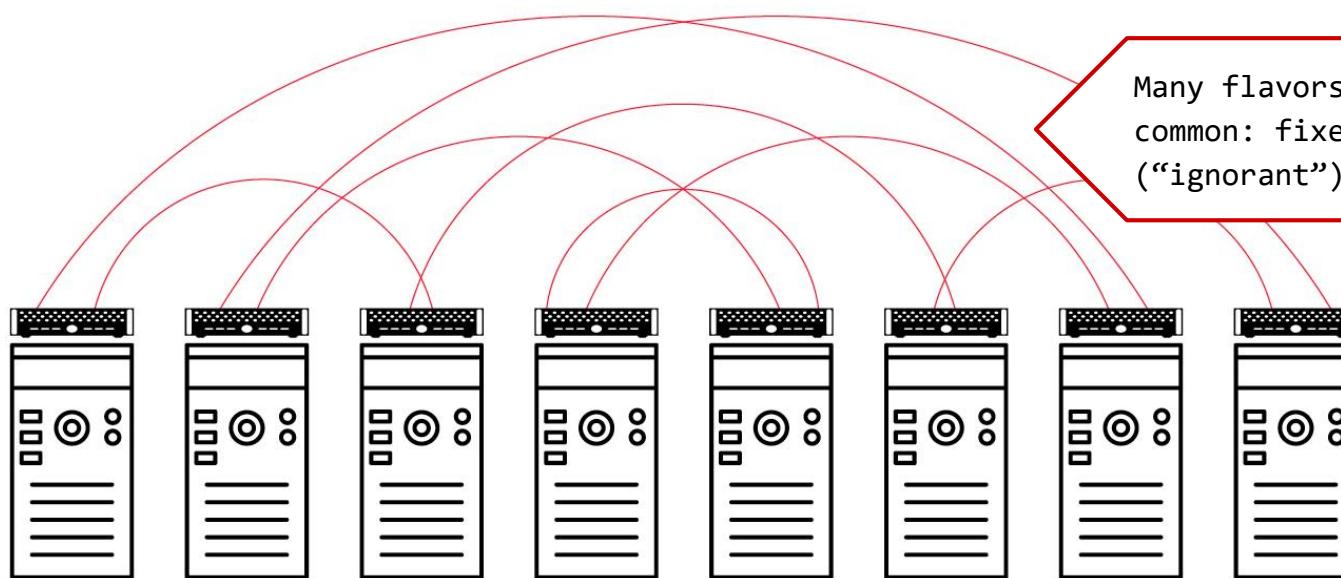
Fixed and Demand-Oblivious Topology



→ Example: expander topology (**uni-regular**)

→ Only 1 type of switches:
lower installation and management overheads

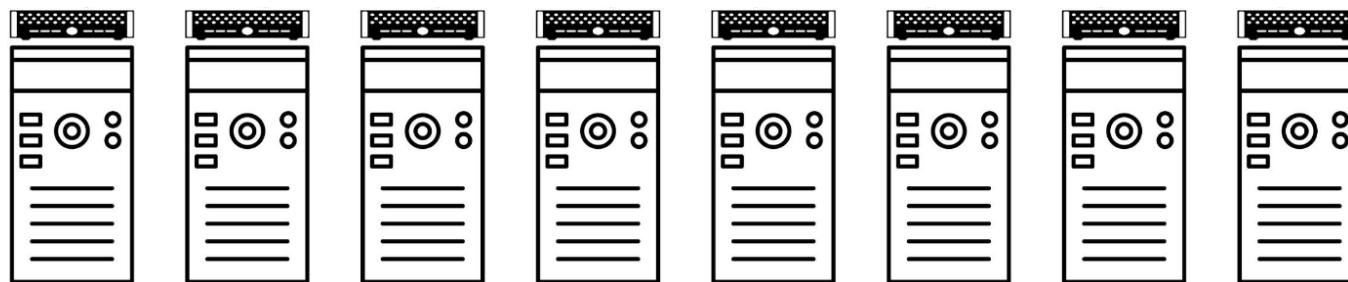
Highway which ignores
actual traffic: **frustrating!**



Many flavors, but in common: fixed and **oblivious** (“ignorant”) to actual demand.

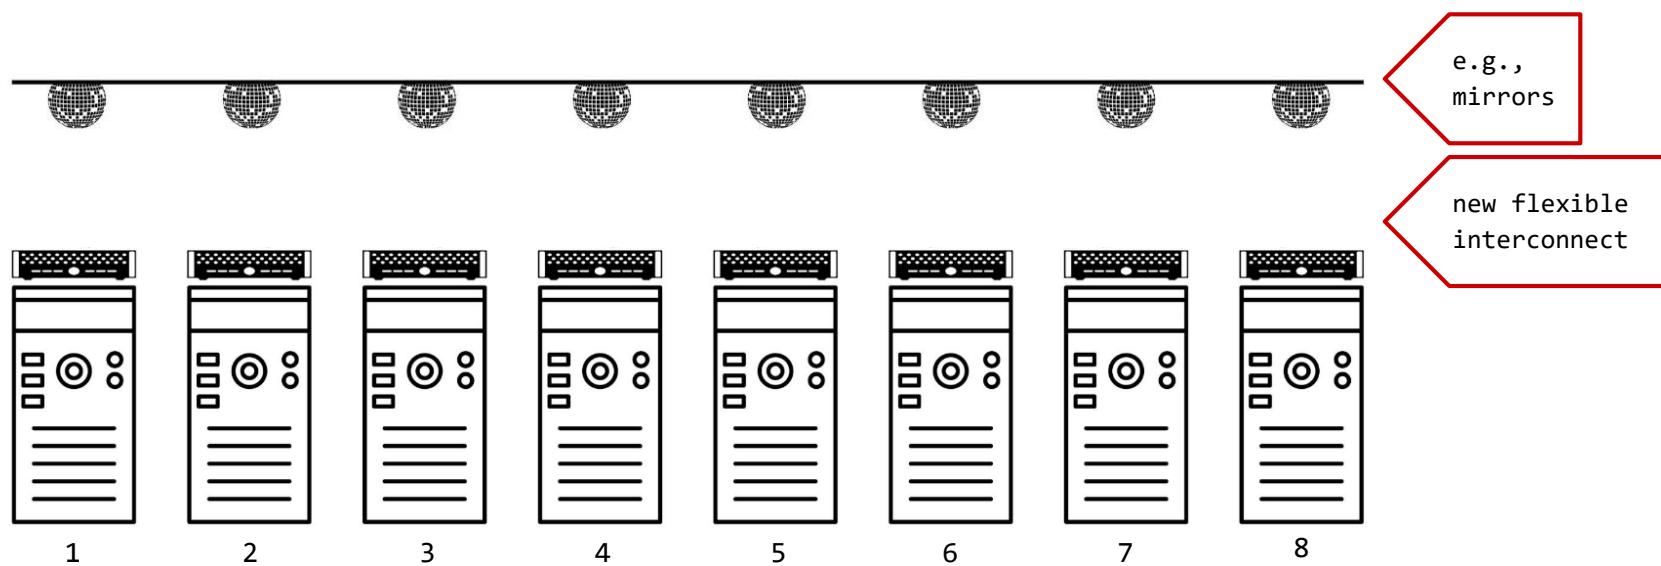
A Vision

Flexible and Demand-Aware Topologies



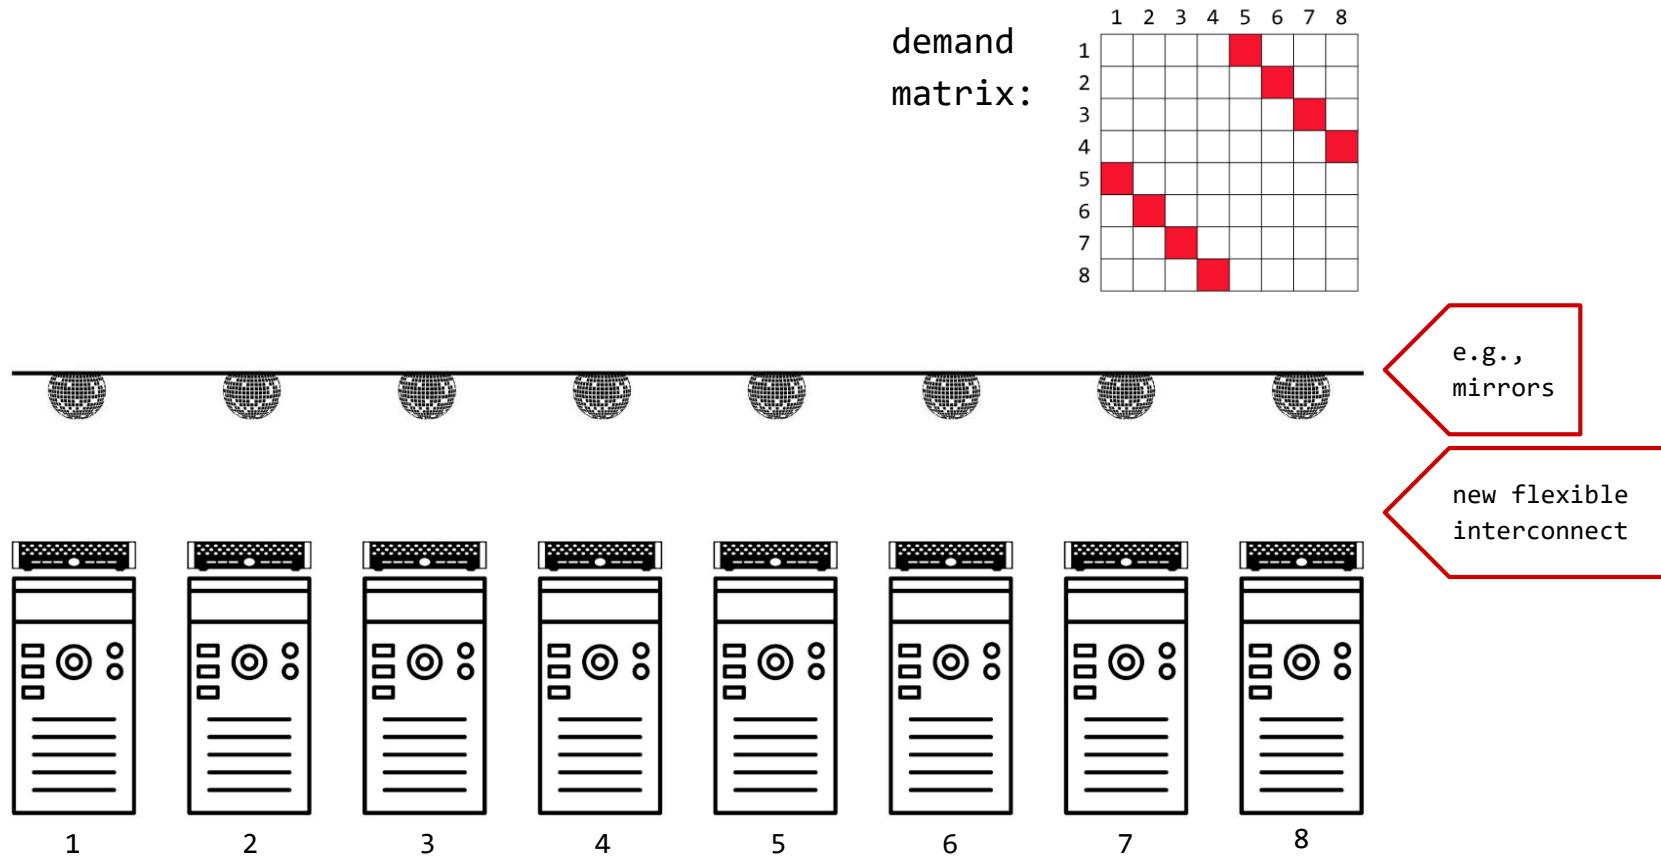
A Vision

Flexible and Demand-Aware Topologies



A Vision

Flexible and Demand-Aware Topologies



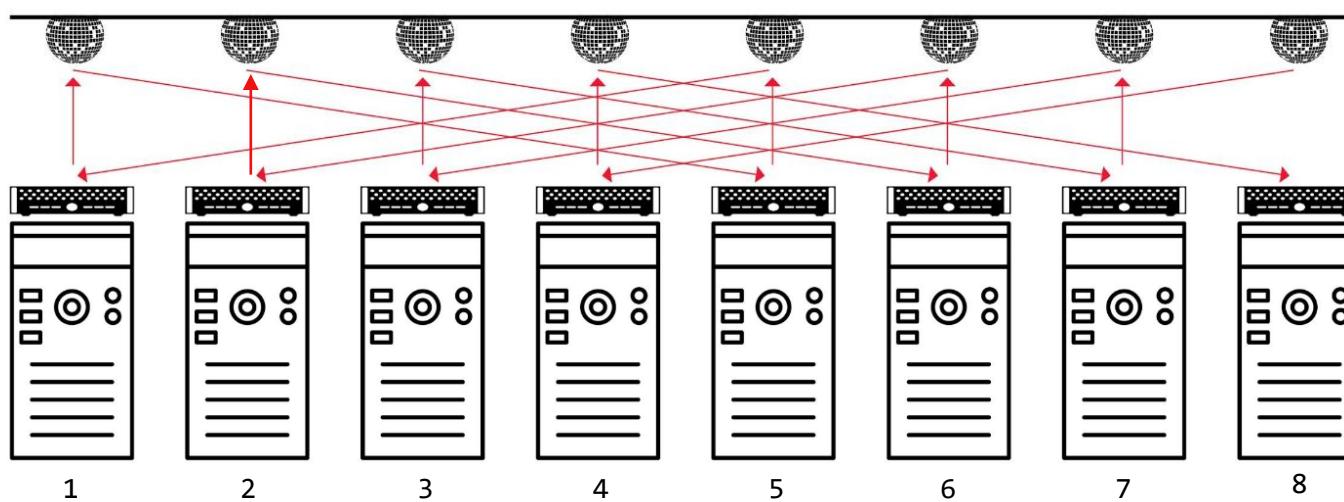
A Vision

Flexible and Demand-Aware Topologies

Matches demand

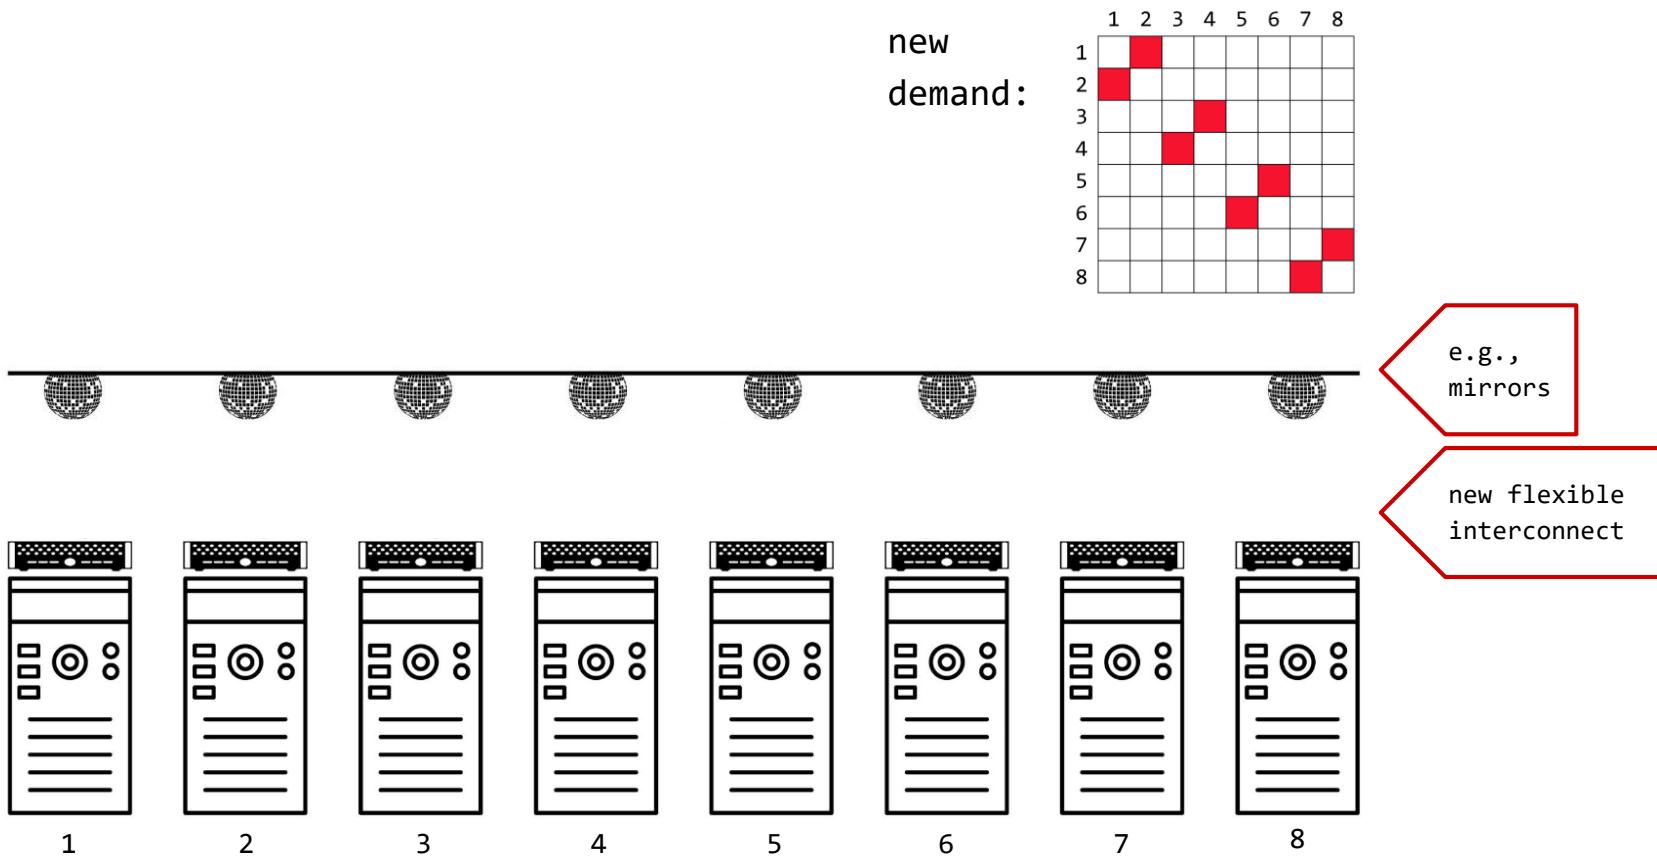
demand
matrix:

1	2	3	4	5	6	7	8
1					1		
2					1	1	
3						1	
4							1
5	1						
6		1					
7			1				
8				1			



A Vision

Flexible and Demand-Aware Topologies



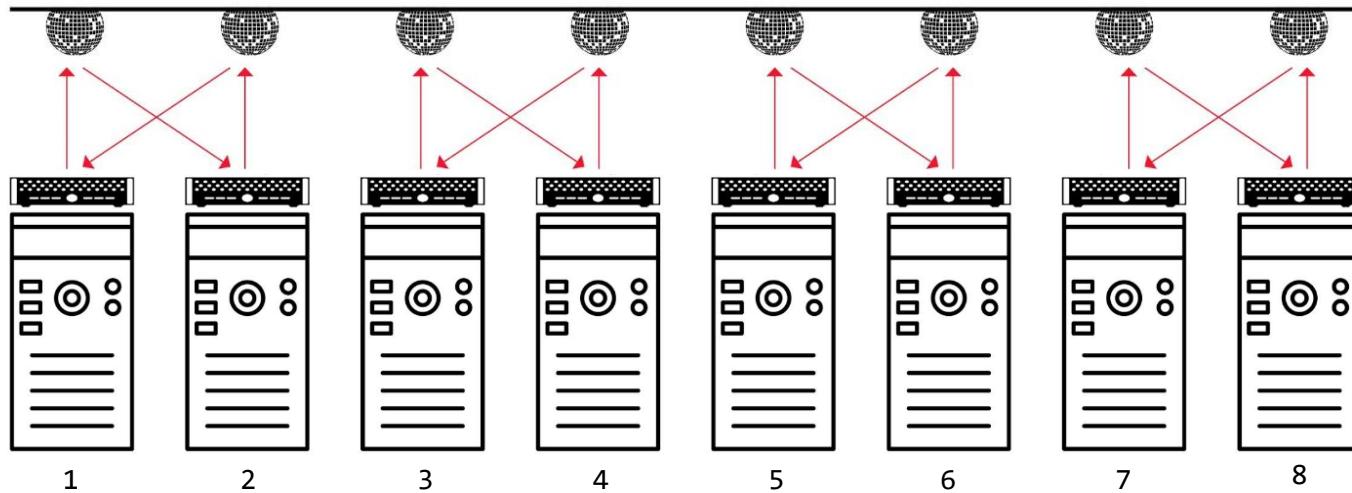
A Vision

Flexible and Demand-Aware Topologies

Matches demand

new
demand:

1	2	3	4	5	6	7	8
1							
2	■						
3							
4		■					
5							
6				■			
7					■		
8						■	



e.g.,
mirrors

new flexible
interconnect

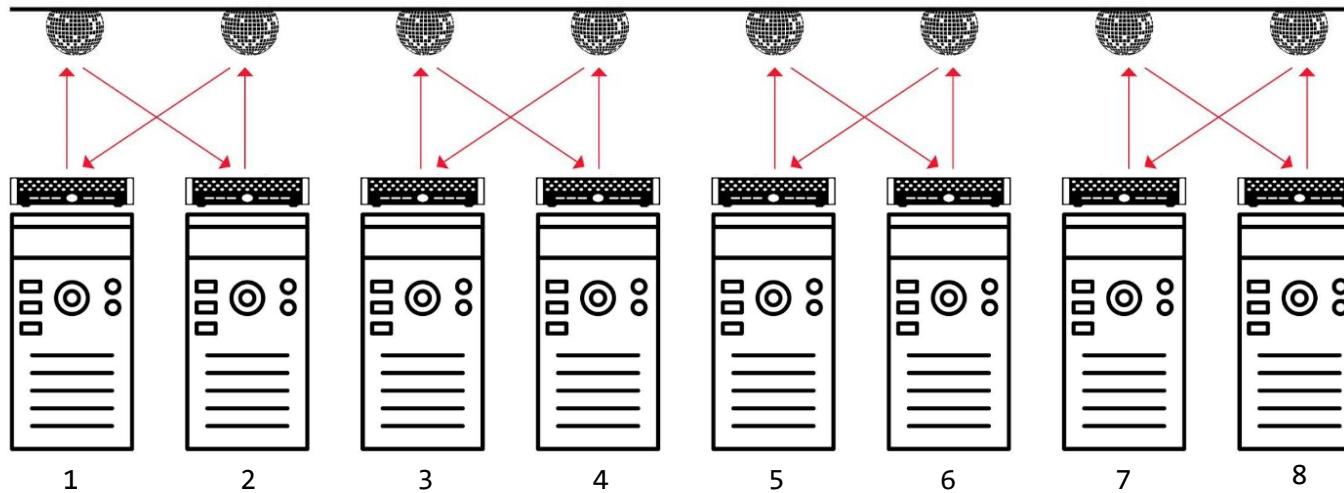
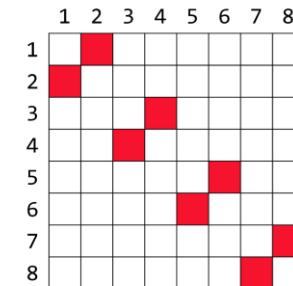
A Vision

Flexible and Demand-Aware Topologies

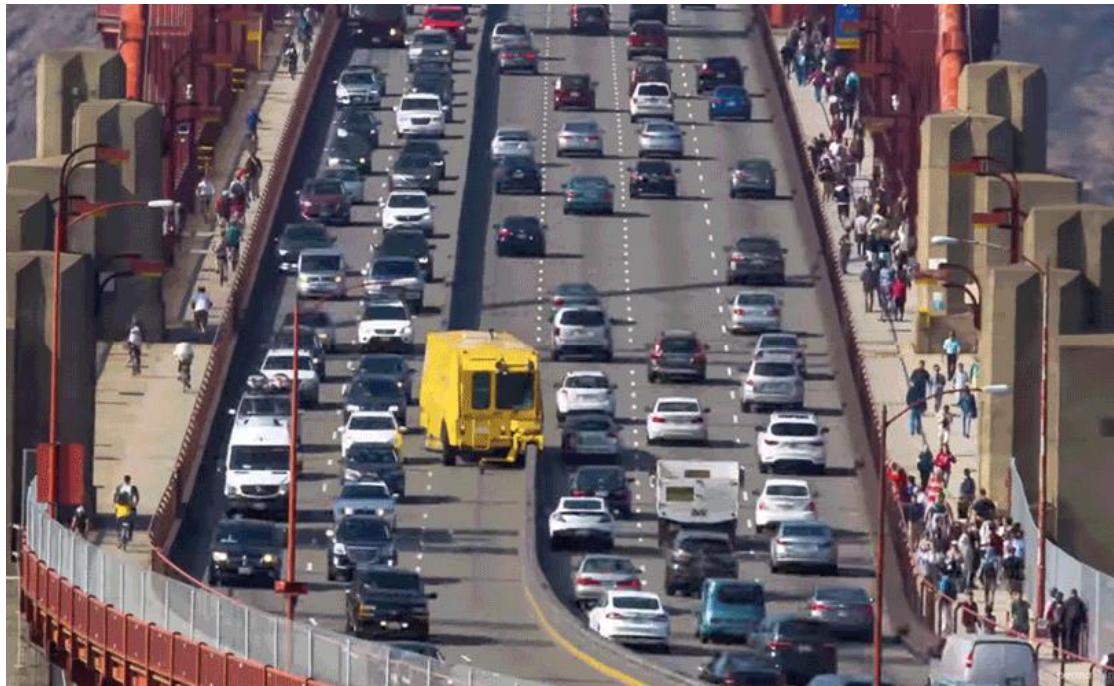


Self-Adjusting
Networks

new
demand:



Analogy



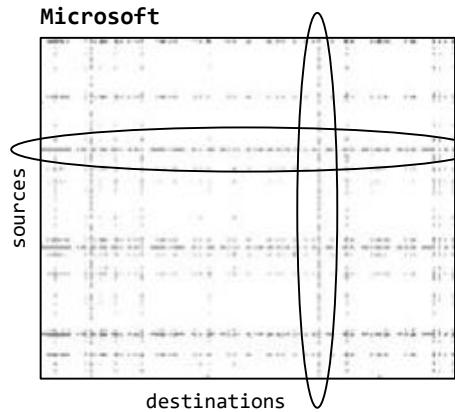
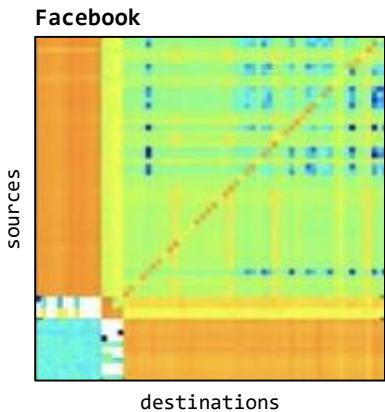
Golden Gate Zipper

The Motivation

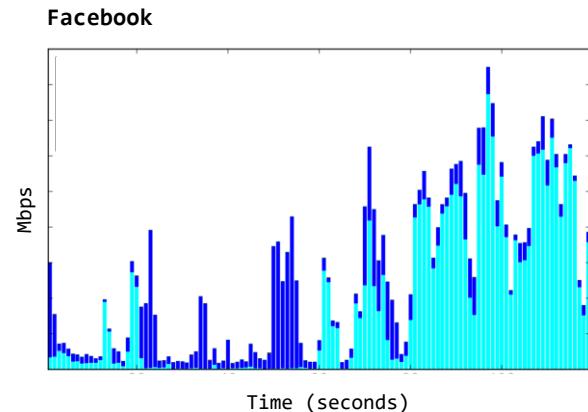
Much Structure in the Demand: Complexity Map

Empirical studies:

traffic matrices **sparse** and **skewed**



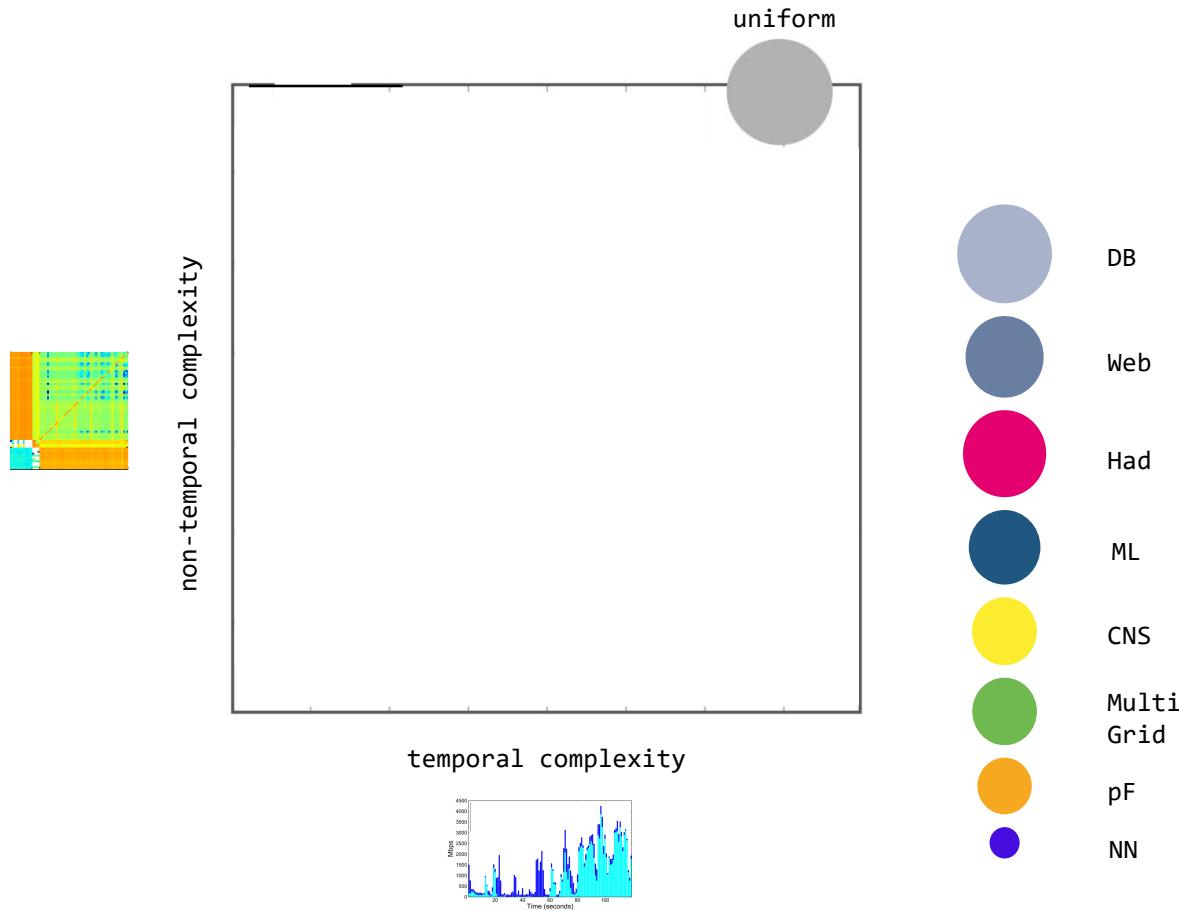
traffic **bursty** over time



The **hypothesis**: can
be exploited.

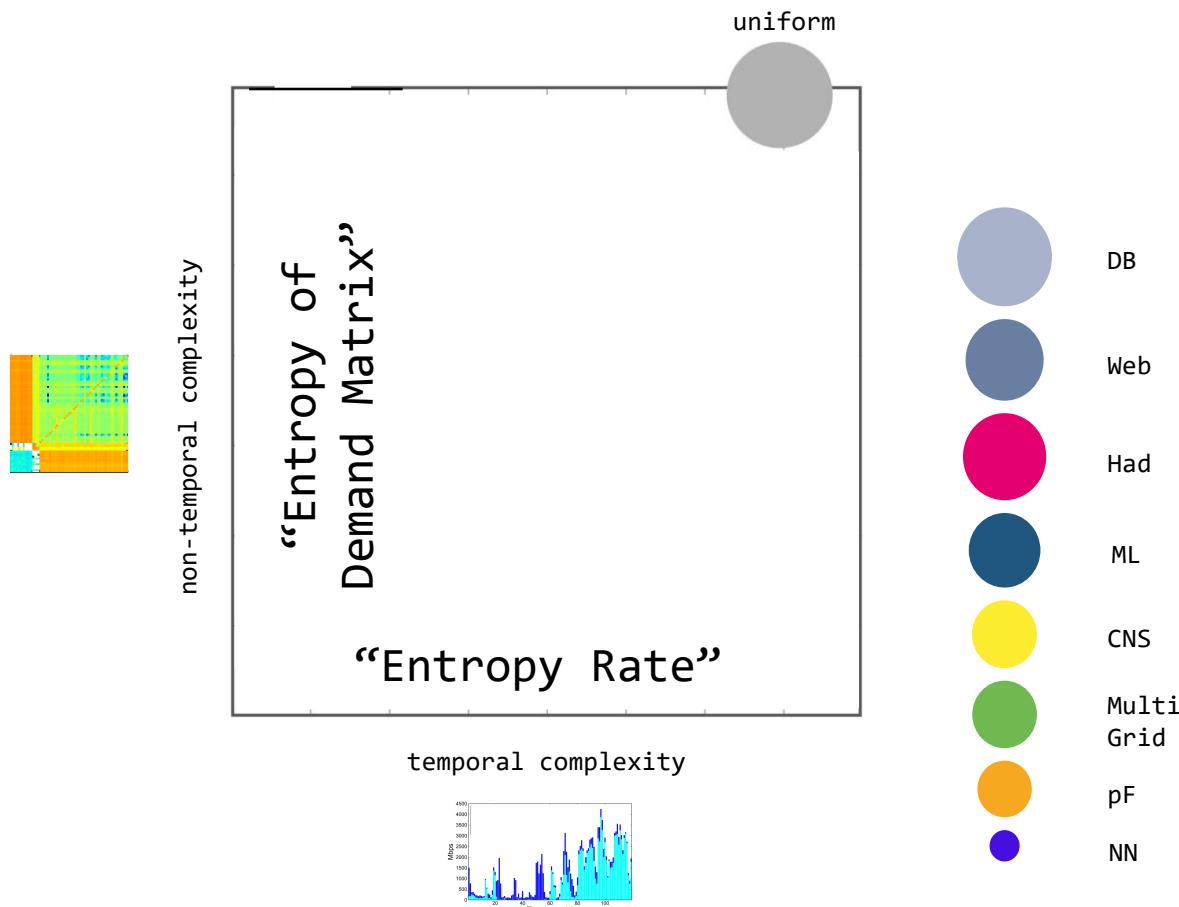
The Motivation

Much Structure in the Demand: Complexity Map



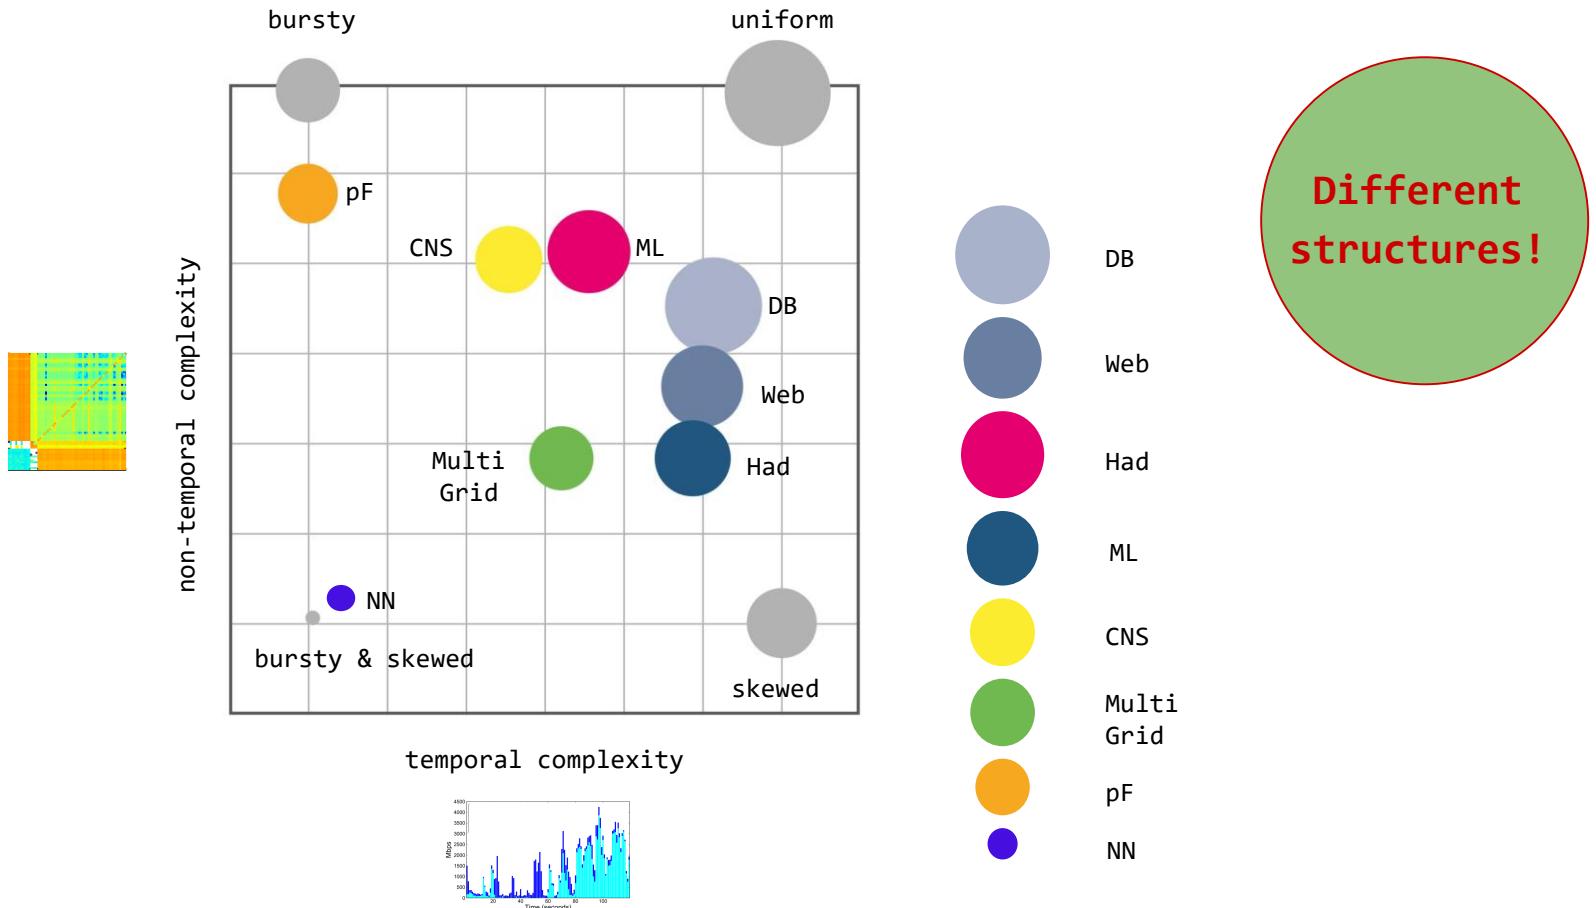
The Motivation

Much Structure in the Demand: Complexity Map



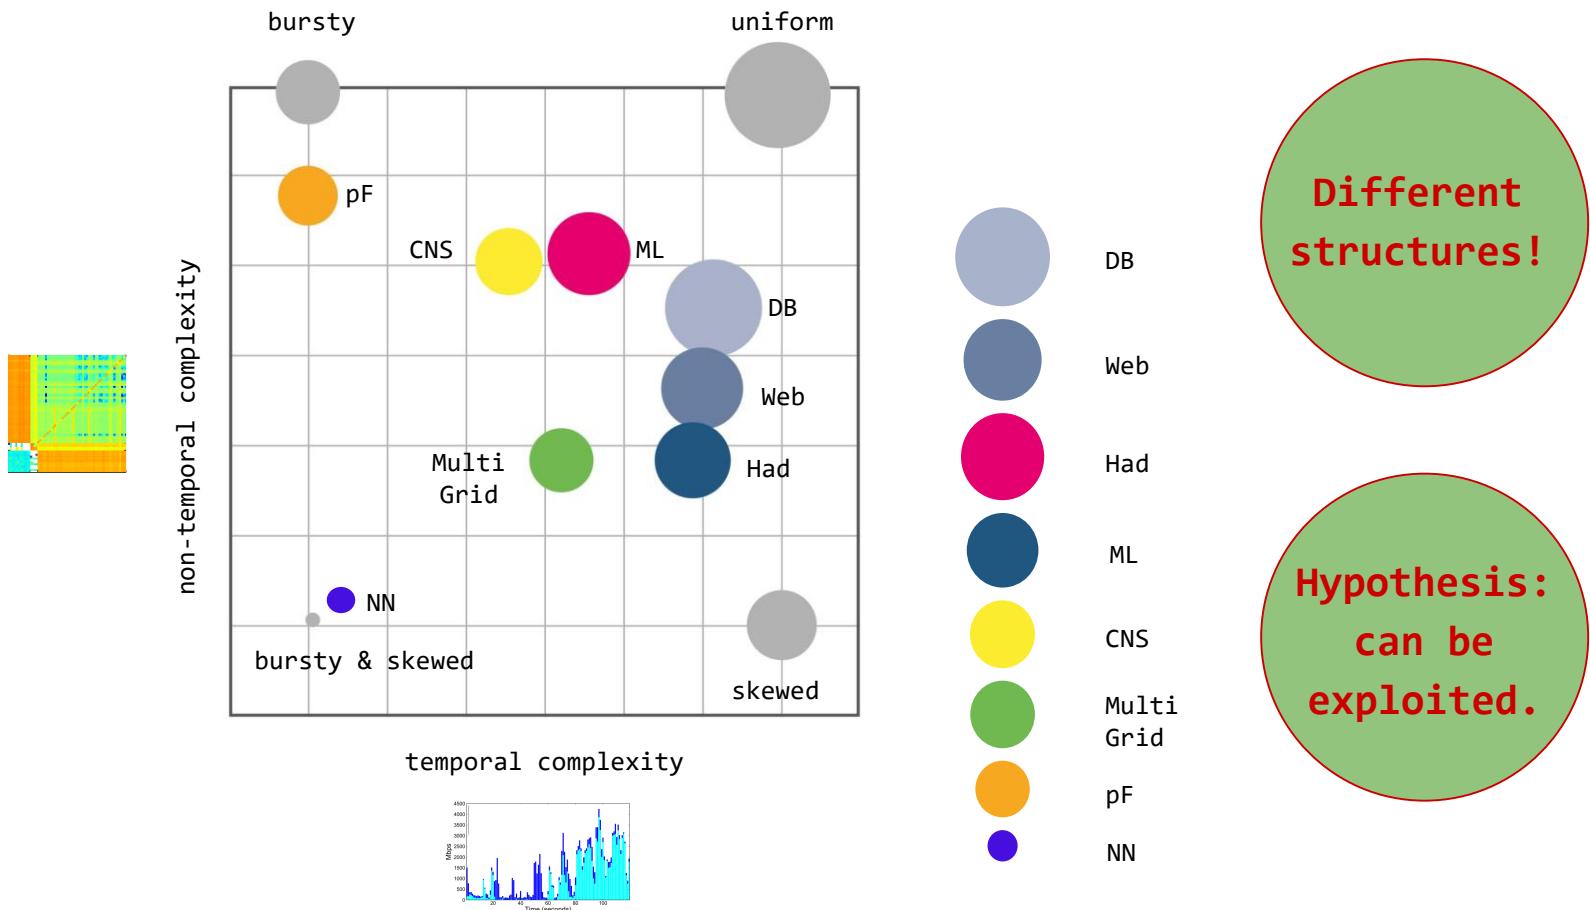
The Motivation

Much Structure in the Demand: Complexity Map



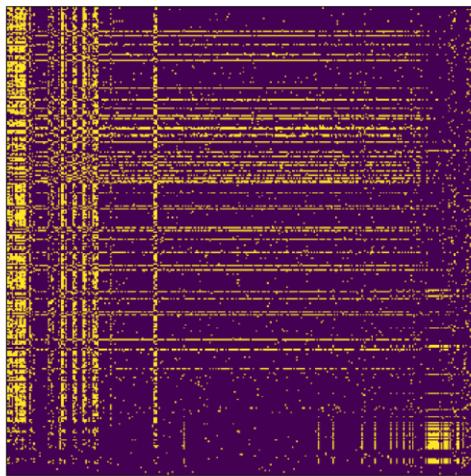
The Motivation

Much Structure in the Demand: Complexity Map

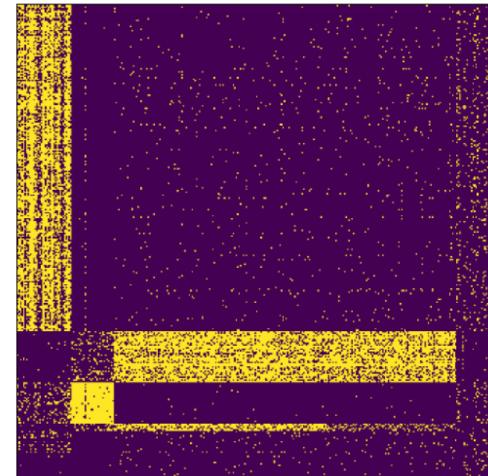


Traffic is also clustered: bi-clustering results

Small Stable Clusters

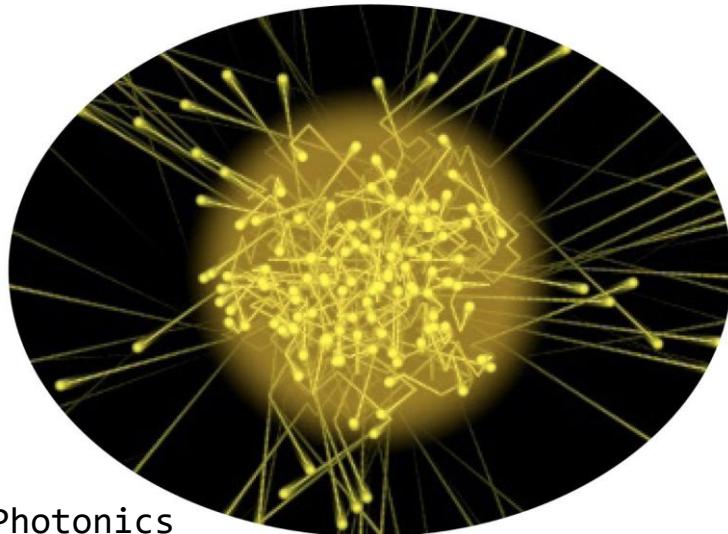


reordering based on
bicluster structure



Opportunity: *exploit* with little reconfigurations!

Sounds Crazy? Emerging Enabling Technology.



H2020:

**“Photronics one of only five
key enabling technologies
for future prosperity.”**

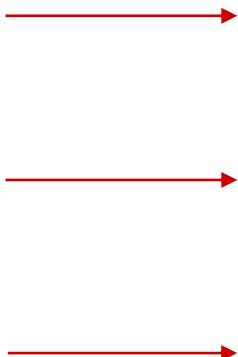
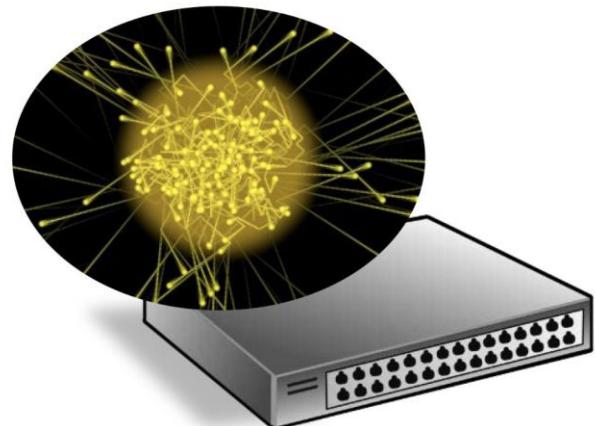
US National Research Council:
**“Photons are the new
Electrons.”**

Enabler

Novel Reconfigurable Optical Switches

→ **Spectrum** of prototypes

- Different sizes, different reconfiguration times
- From our ACM **SIGCOMM** workshop OptSys



Prototype 1



Prototype 2



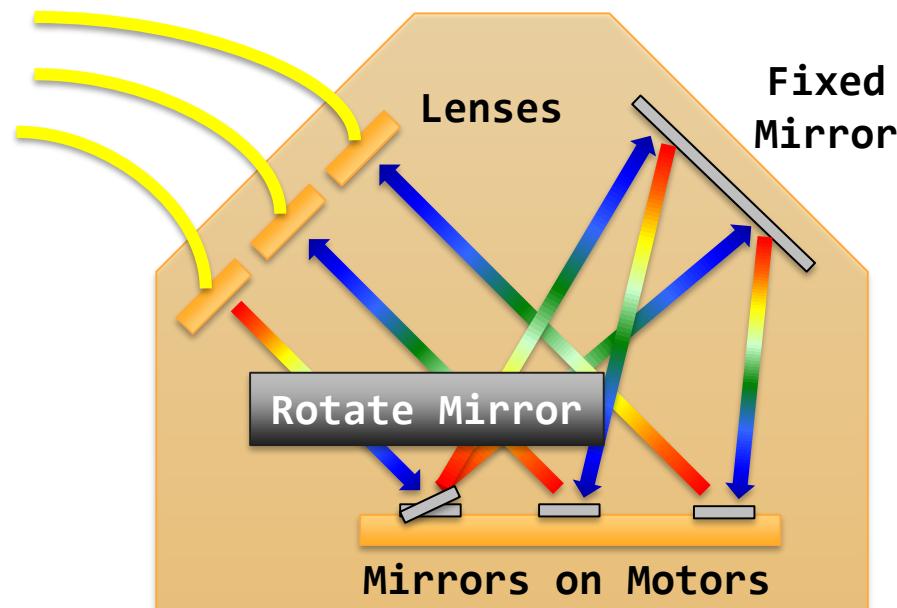
Prototype 3

Changing lambdas (ns)

Example

Optical Circuit Switch

- Optical Circuit Switch rapid adaption of physical layer
 - Based on rotating mirrors



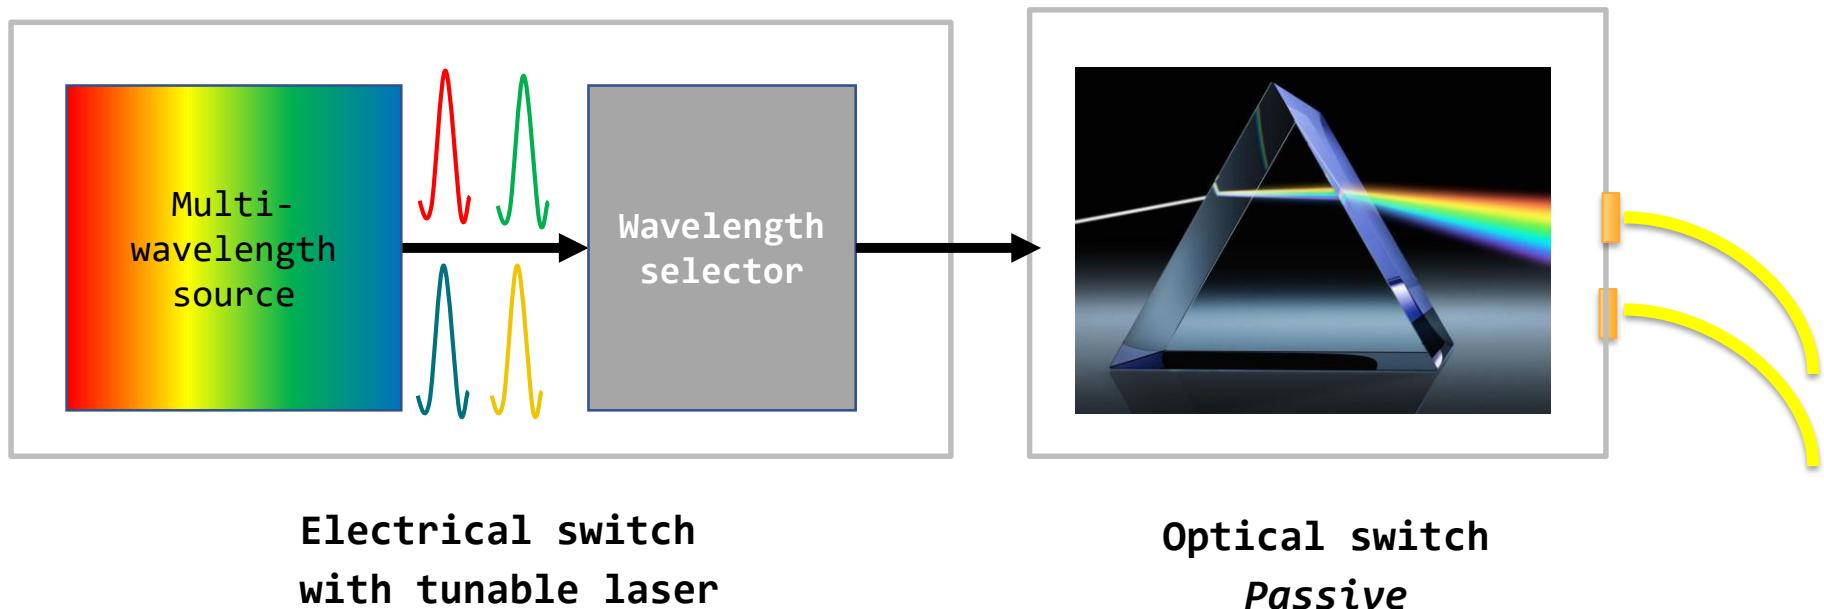
Optical Circuit Switch

By Nathan Farrington, SIGCOMM 2010

Another Example

Tunable Lasers

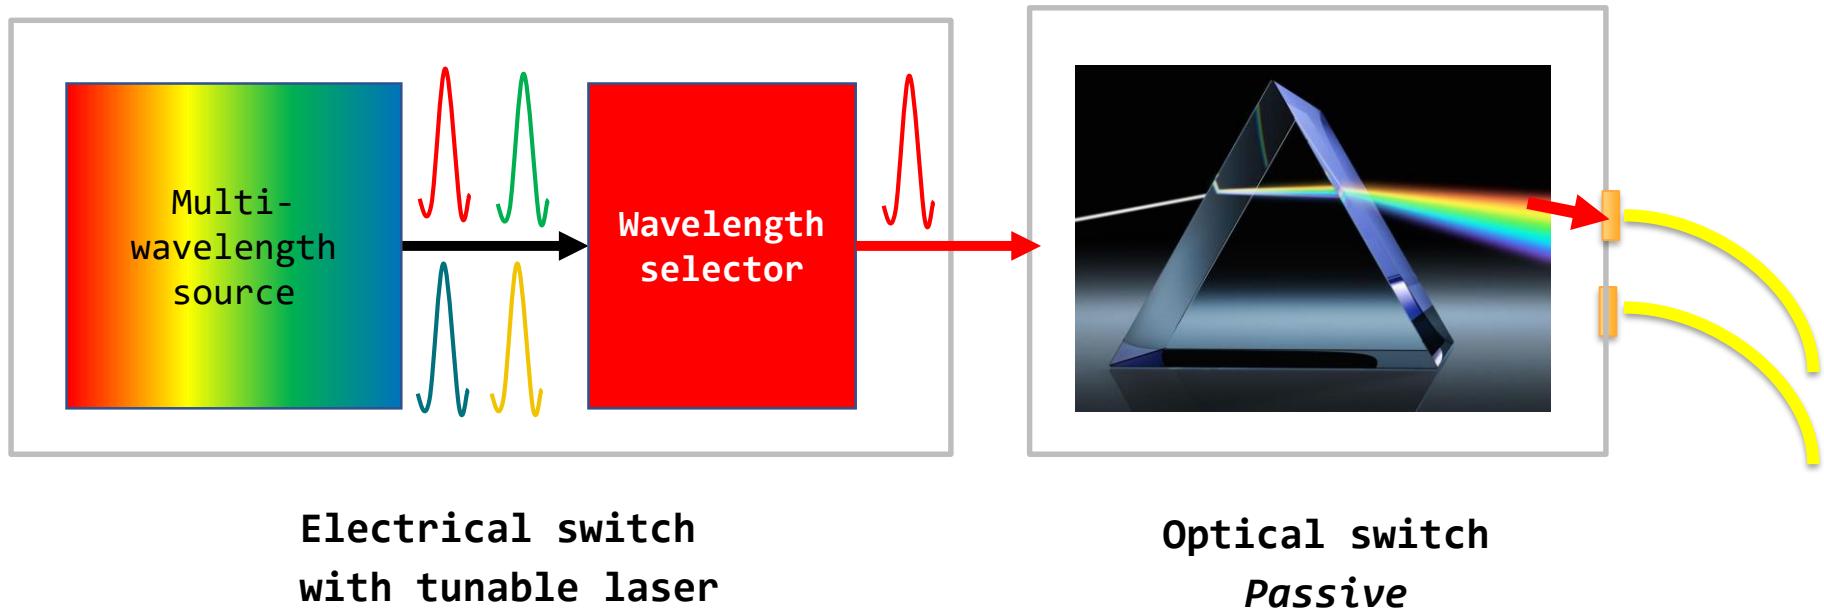
- Depending on wavelength, forwarded differently
- Optical switch is passive



Another Example

Tunable Lasers

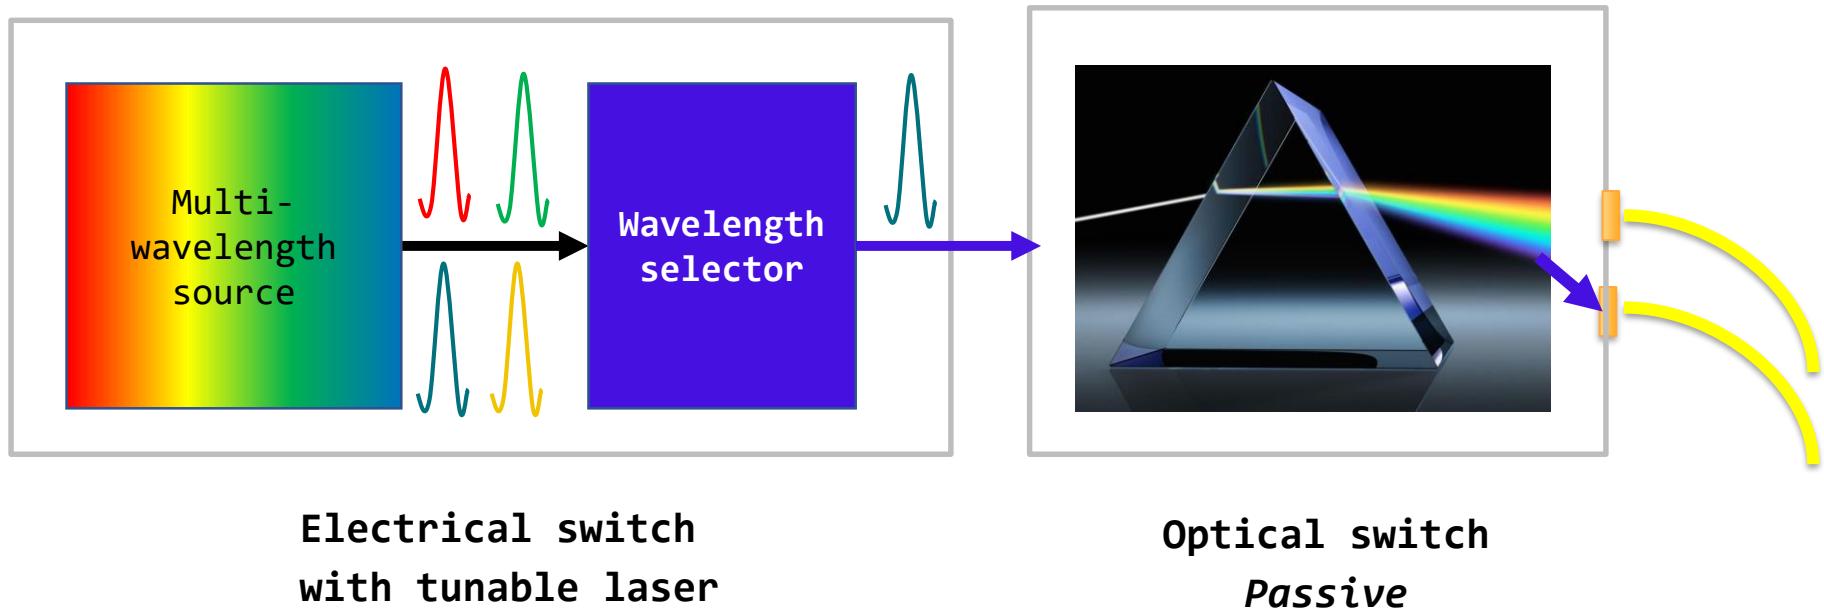
- Depending on wavelength, forwarded differently
- Optical switch is passive



Another Example

Tunable Lasers

- Depending on wavelength, forwarded differently
- Optical switch is passive

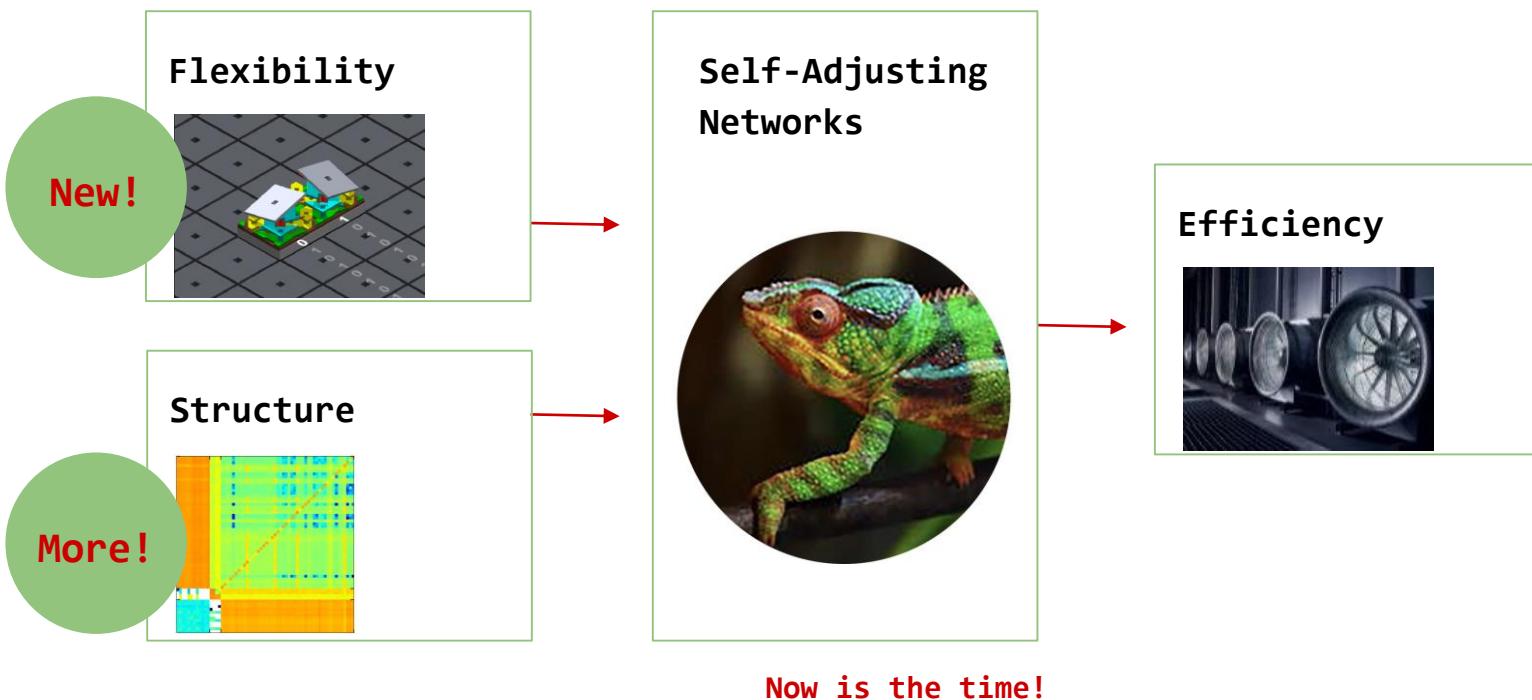


First Deployments

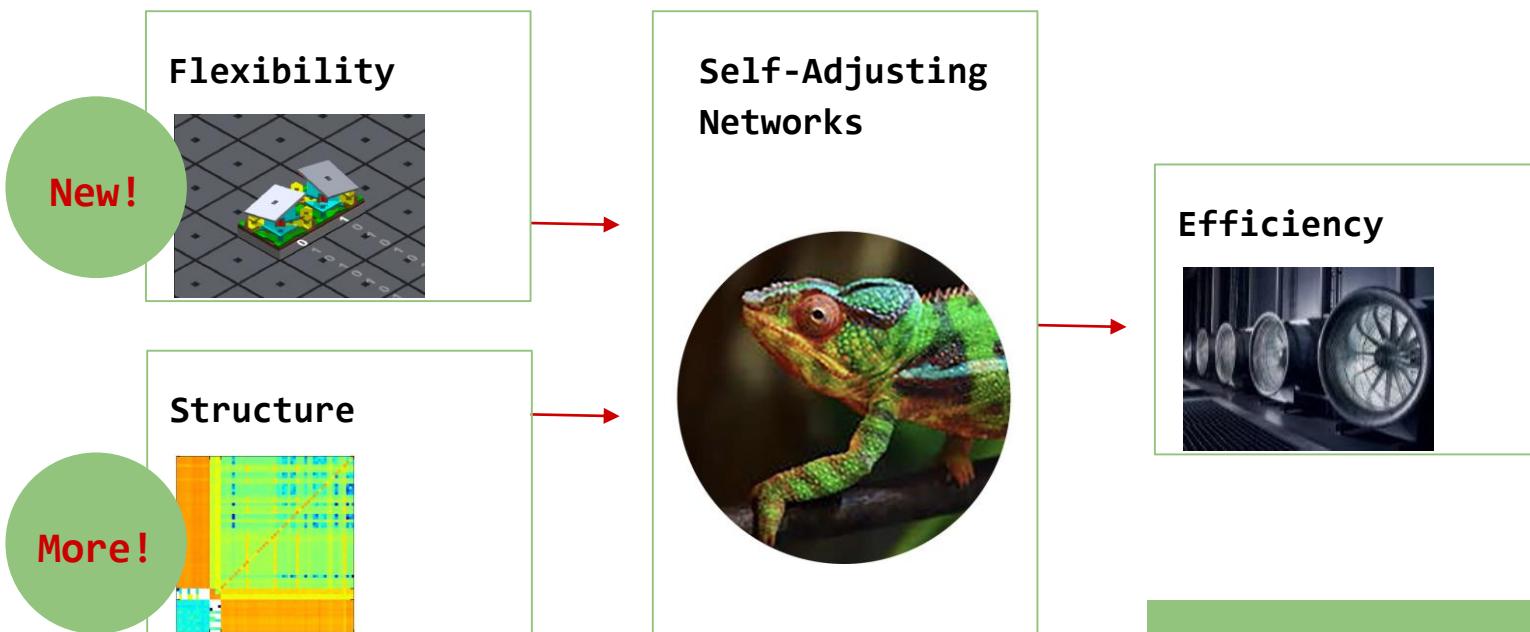
E.g., Google's Datacenter Jupiter



The Big Picture

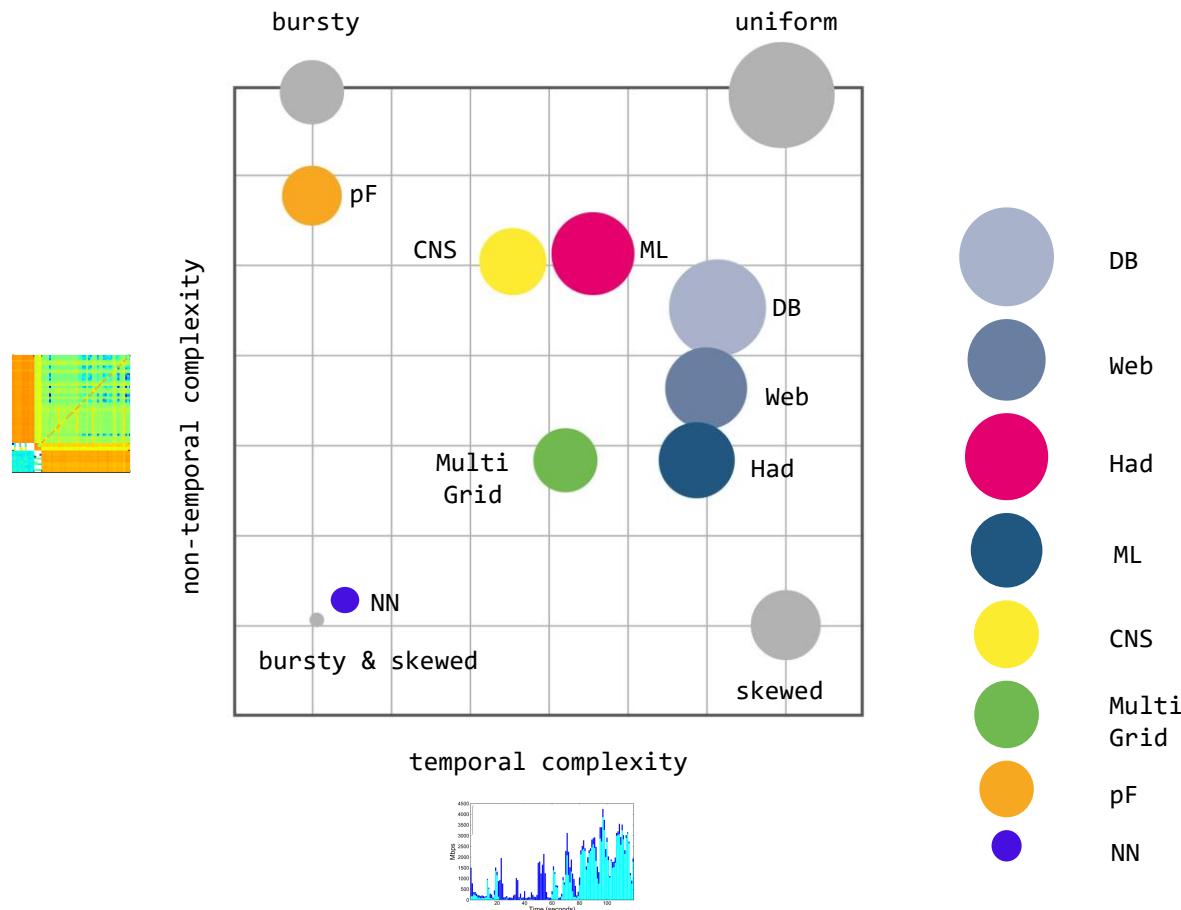


The Big Picture

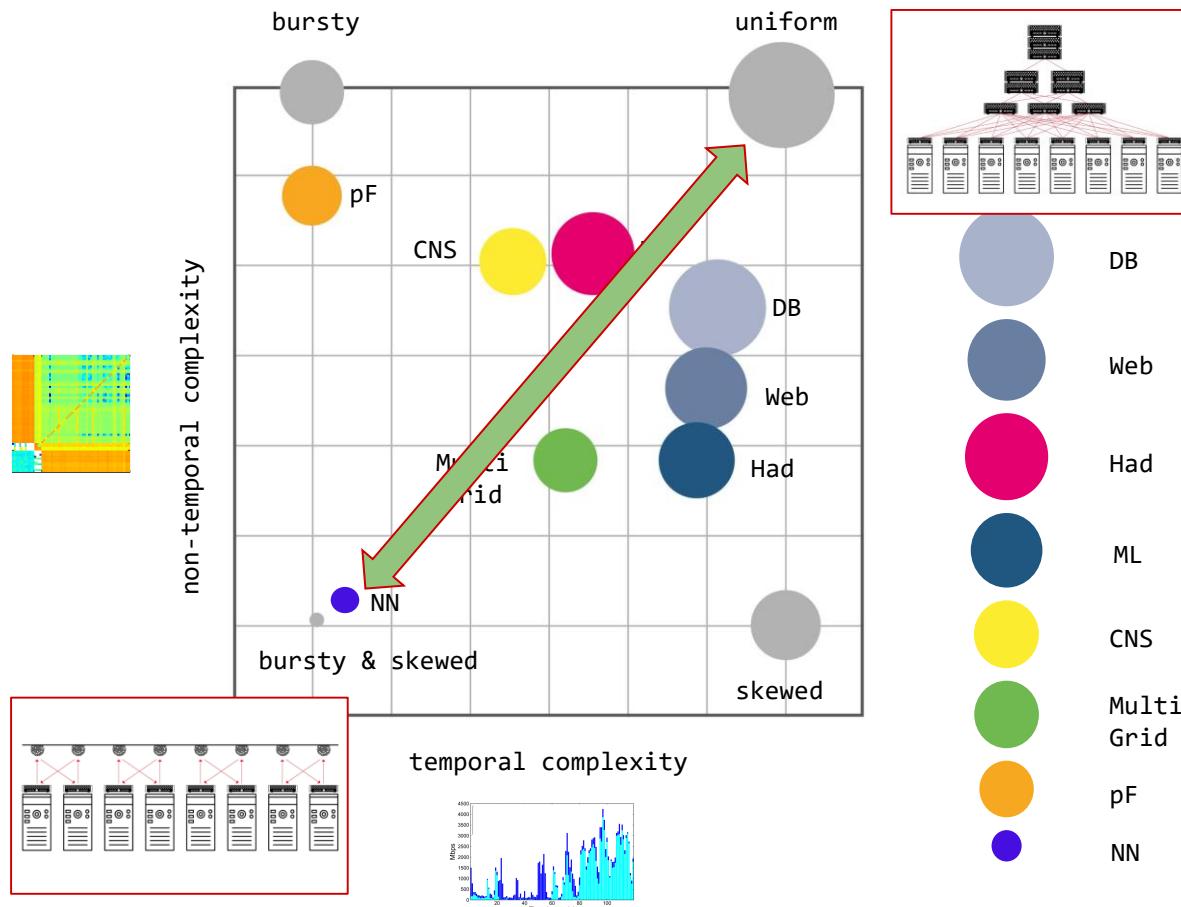


Missing: Foundations of demand-aware, self-adjusting networks.

Potential Gain

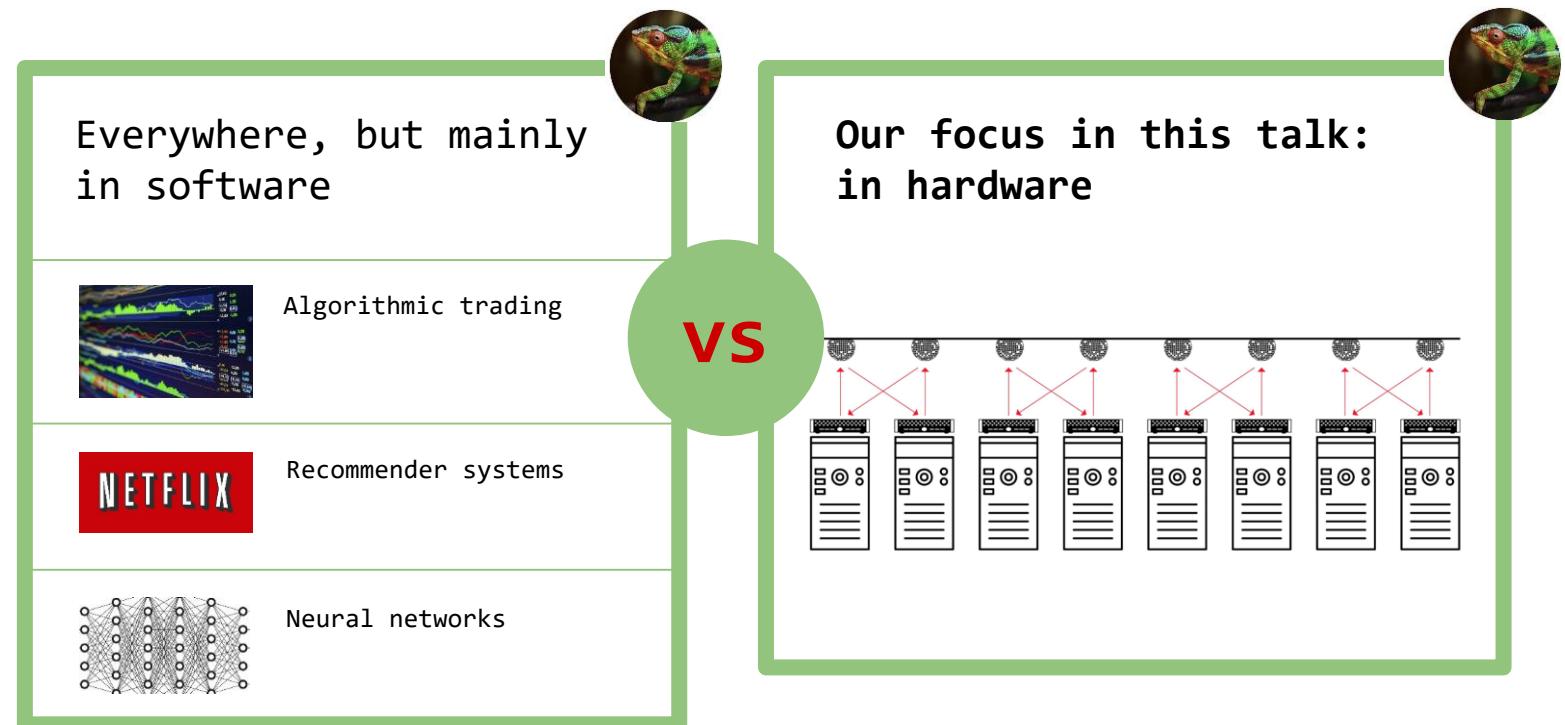


Potential Gain



Unique Position

Demand-Aware, Self-Adjusting Systems



The Natural Question:

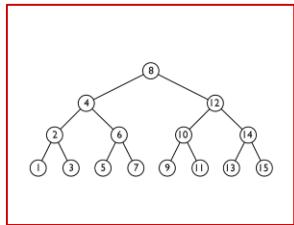
Given This Structure,
What Can Be Achieved?
Metrics and Algorithms?

A first insight: entropy of the demand.

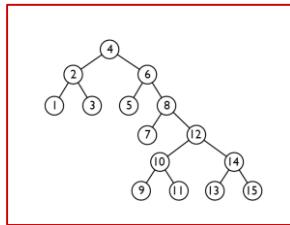
Insight:

Connection to Datastructures

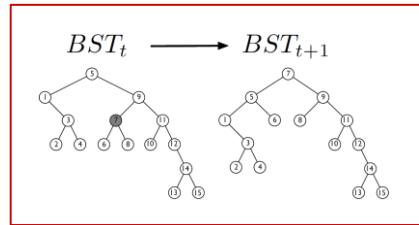
Traditional BST



Demand-aware BST



Self-adjusting BST

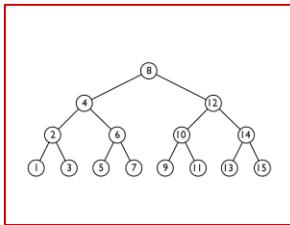


More structure: improved **access cost**

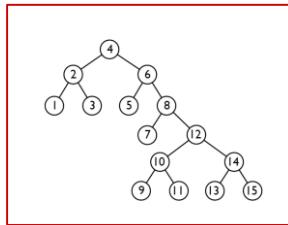
Insight:

Connection to Datastructures & Coding

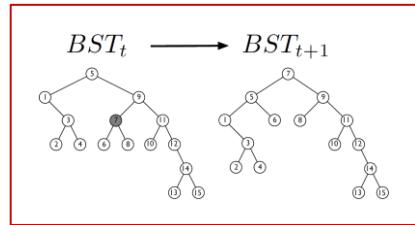
Traditional BST
(Worst-case coding)



Demand-aware BST
(Huffman coding)



Self-adjusting BST
(Dynamic Huffman coding)

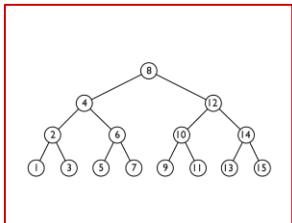


More structure: improved **access cost** / shorter **codes**

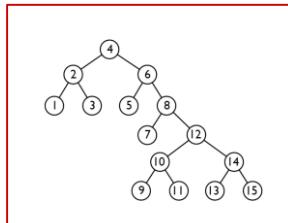
Insight:

Connection to Datastructures & Coding

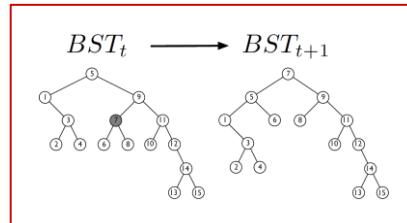
Traditional BST
(Worst-case coding)



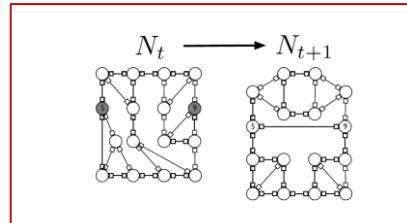
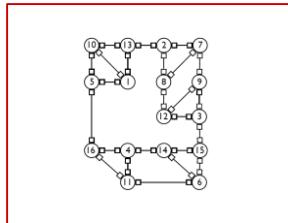
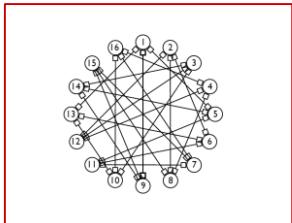
Demand-aware BST
(Huffman coding)



Self-adjusting BST
(Dynamic Huffman coding)



More structure: improved **access cost** / shorter **codes**

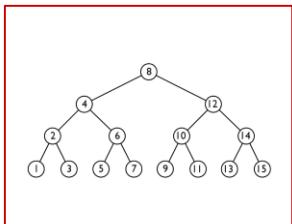


Similar **benefits?**

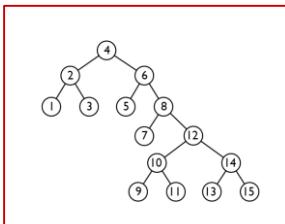
Insight:

Connection to Datastructures & Coding

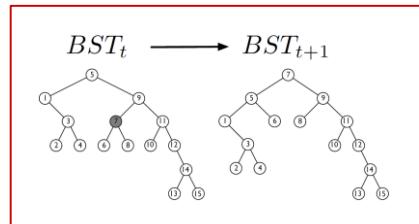
Traditional BST
(Worst-case coding)



Demand-aware BST
(Huffman coding)

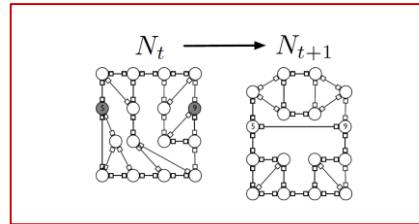
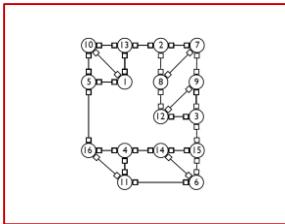
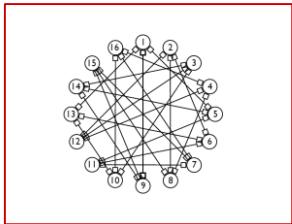


Self-adjusting BST
(Dynamic Huffman coding)



More than
an analogy!

More structure: improved **access cost** / shorter **codes**

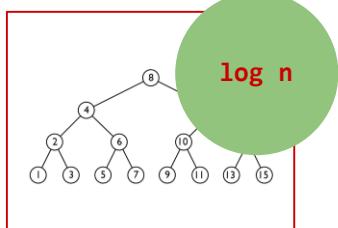


Similar **benefits?**

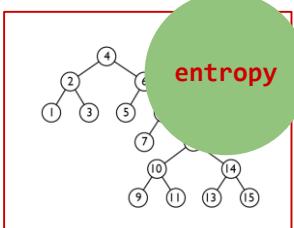
Insight:

Connection to Datastructures & Coding

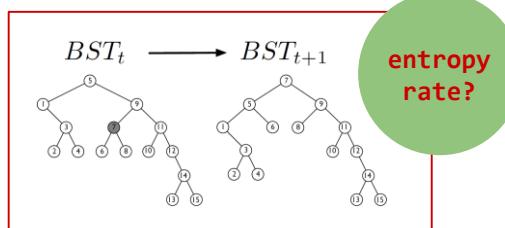
Traditional BST
(Worst-case coding)



Demand-aware BST
(Huffman coding)



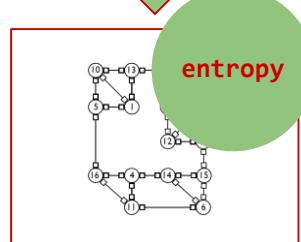
Self-adjusting BST
(Dynamic Huffman coding)



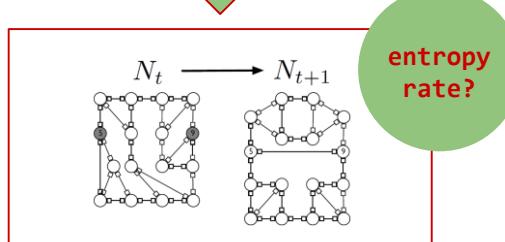
More than
an analogy!



log n



entropy



entropy
rate?

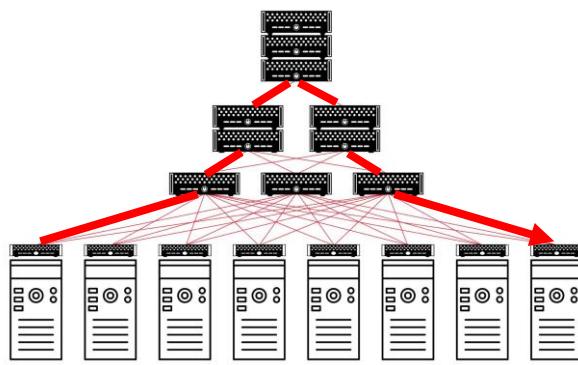
Reduced expected route lengths!

Generalize methodology:
... and transfer
entropy bounds and
algorithms of data-
structures to networks.

First results:
Demand-aware networks
of asymptotically
optimal route lengths.

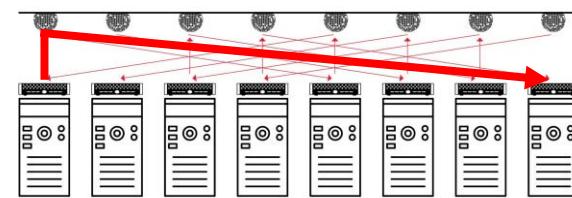
Reality: A Tradeoff

- Self-adjusting networks may be really useful to serve large flows (**elephant flows**): avoiding multi-hop routing



6 hops

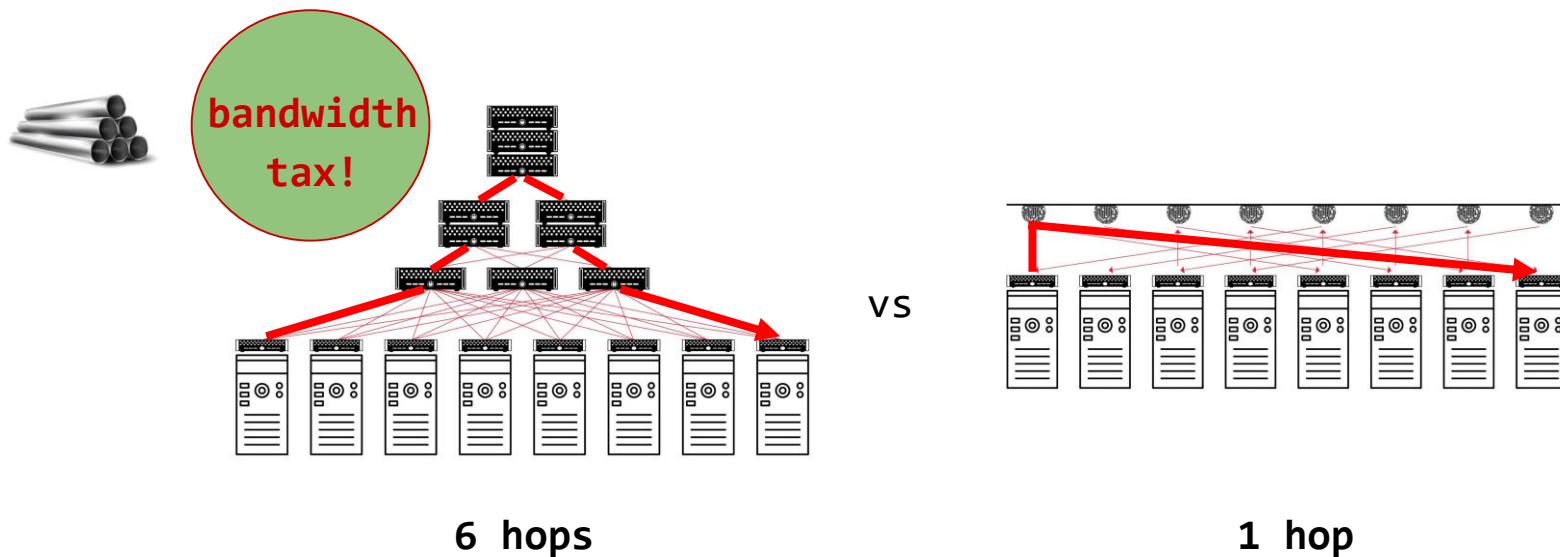
vs



1 hop

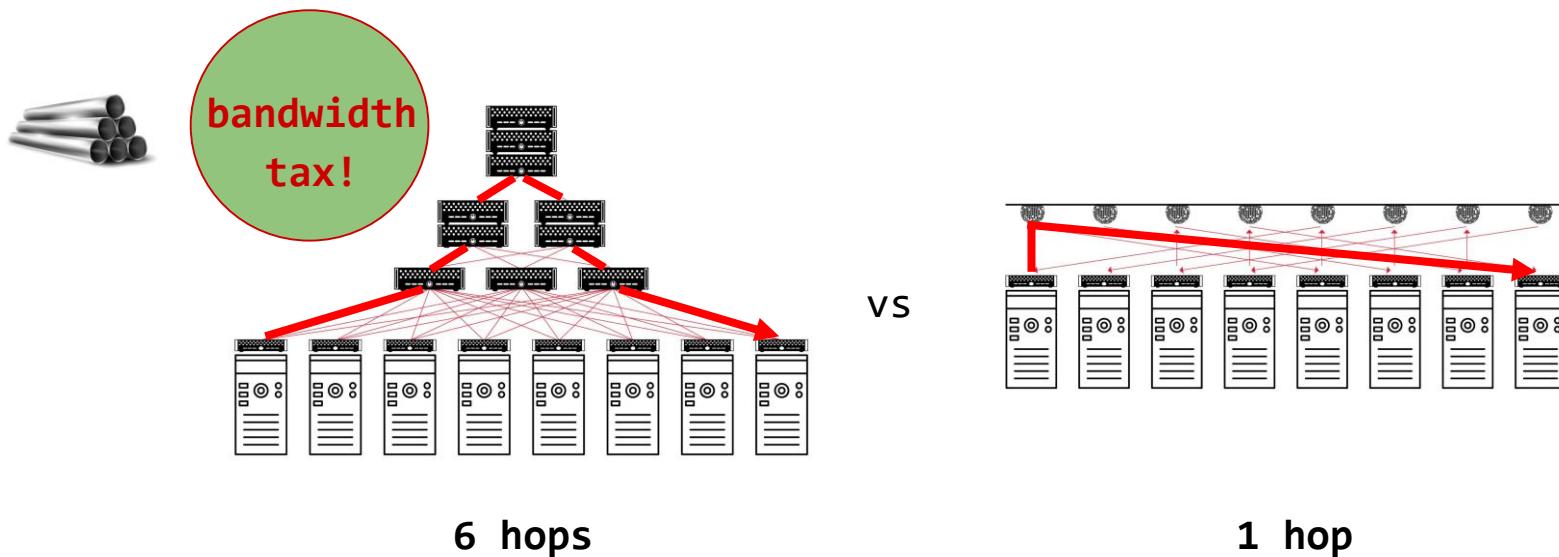
Reality: A Tradeoff

- Self-adjusting networks may be really useful to serve large flows (**elephant flows**): avoiding multi-hop routing



Reality: A Tradeoff

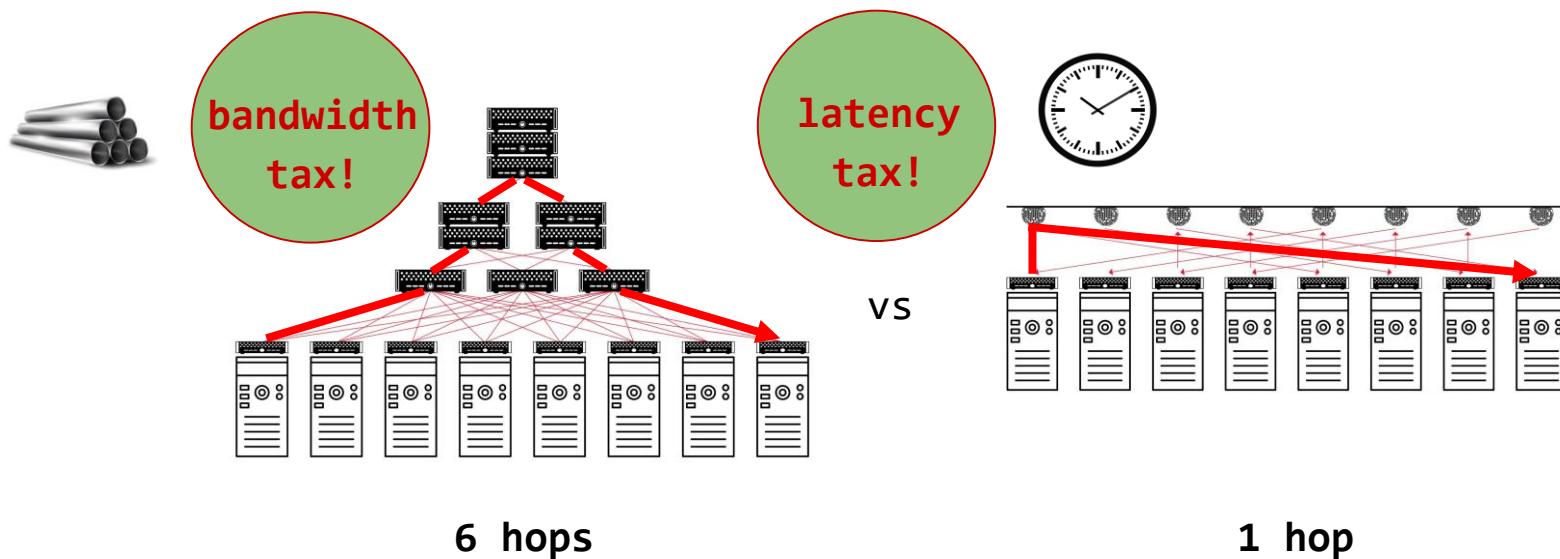
- Self-adjusting networks may be really useful to serve large flows (**elephant flows**): avoiding multi-hop routing



- However, requires optimization and adaption, which **takes time**

Reality: A Tradeoff

- Self-adjusting networks may be really useful to serve large flows (**elephant flows**): avoiding multi-hop routing



- However, requires optimization and adaption, which **takes time**

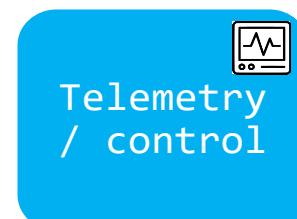
Challenge: Traffic Diversity

Diverse patterns:

- Shuffling/Hadoop:
all-to-all
- All-reduce/ML: **ring** or
tree traffic patterns
 - **Elephant** flows
- Query traffic: skewed
 - **Mice** flows
- Control traffic: does not evolve
but has non-temporal structure

Diverse requirements:

- ML is **bandwidth** hungry,
small flows are **latency-**
sensitive



Opportunity: Tech Diversity

Diverse topology components:

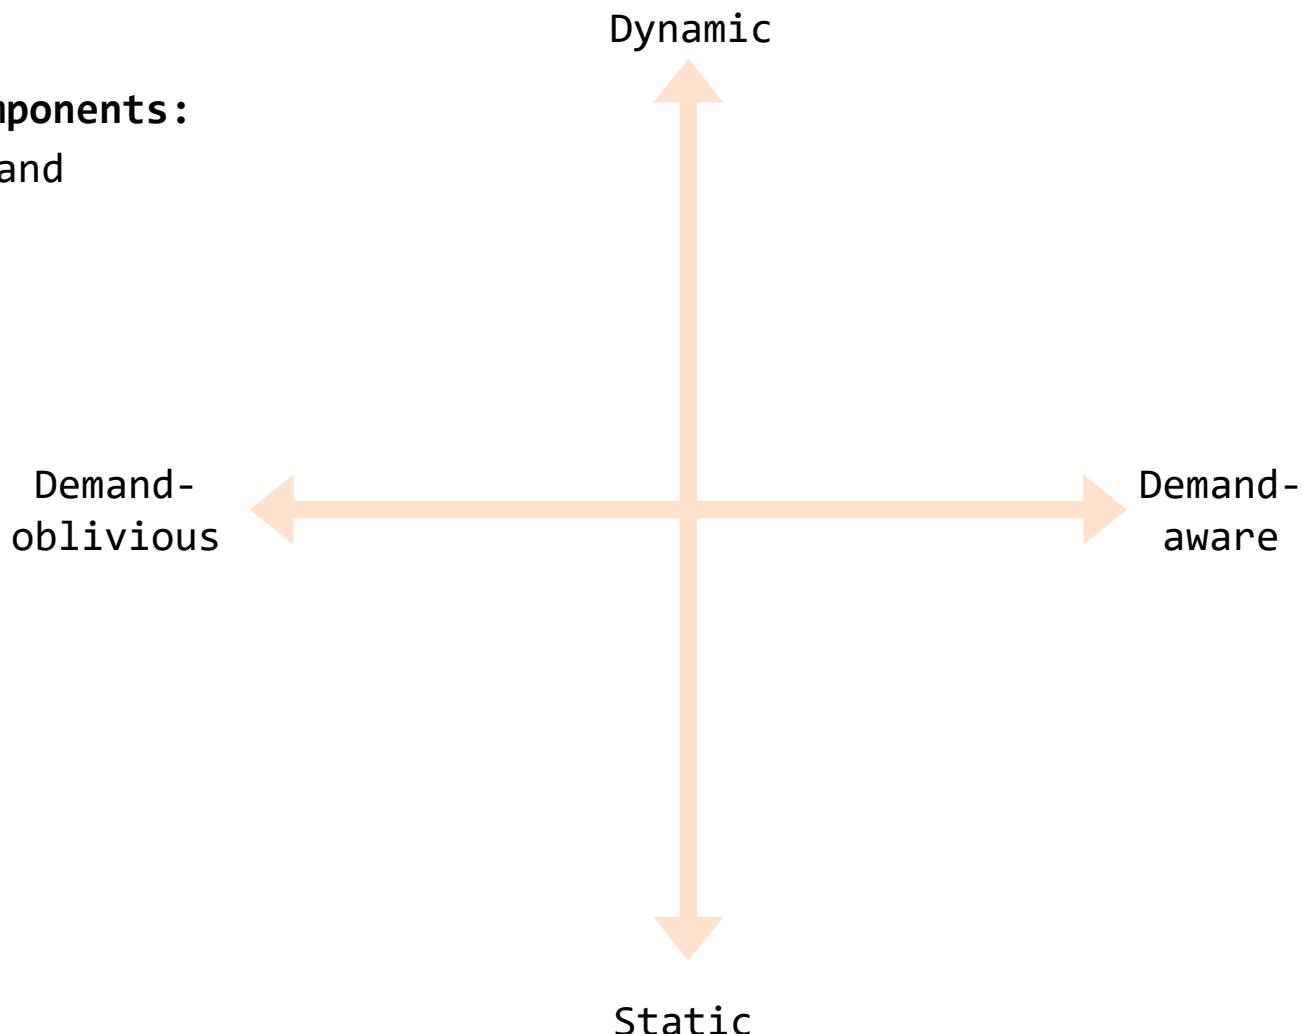
- demand-**oblivious** and
- demand-**aware**



Opportunity: Tech Diversity

Diverse topology components:

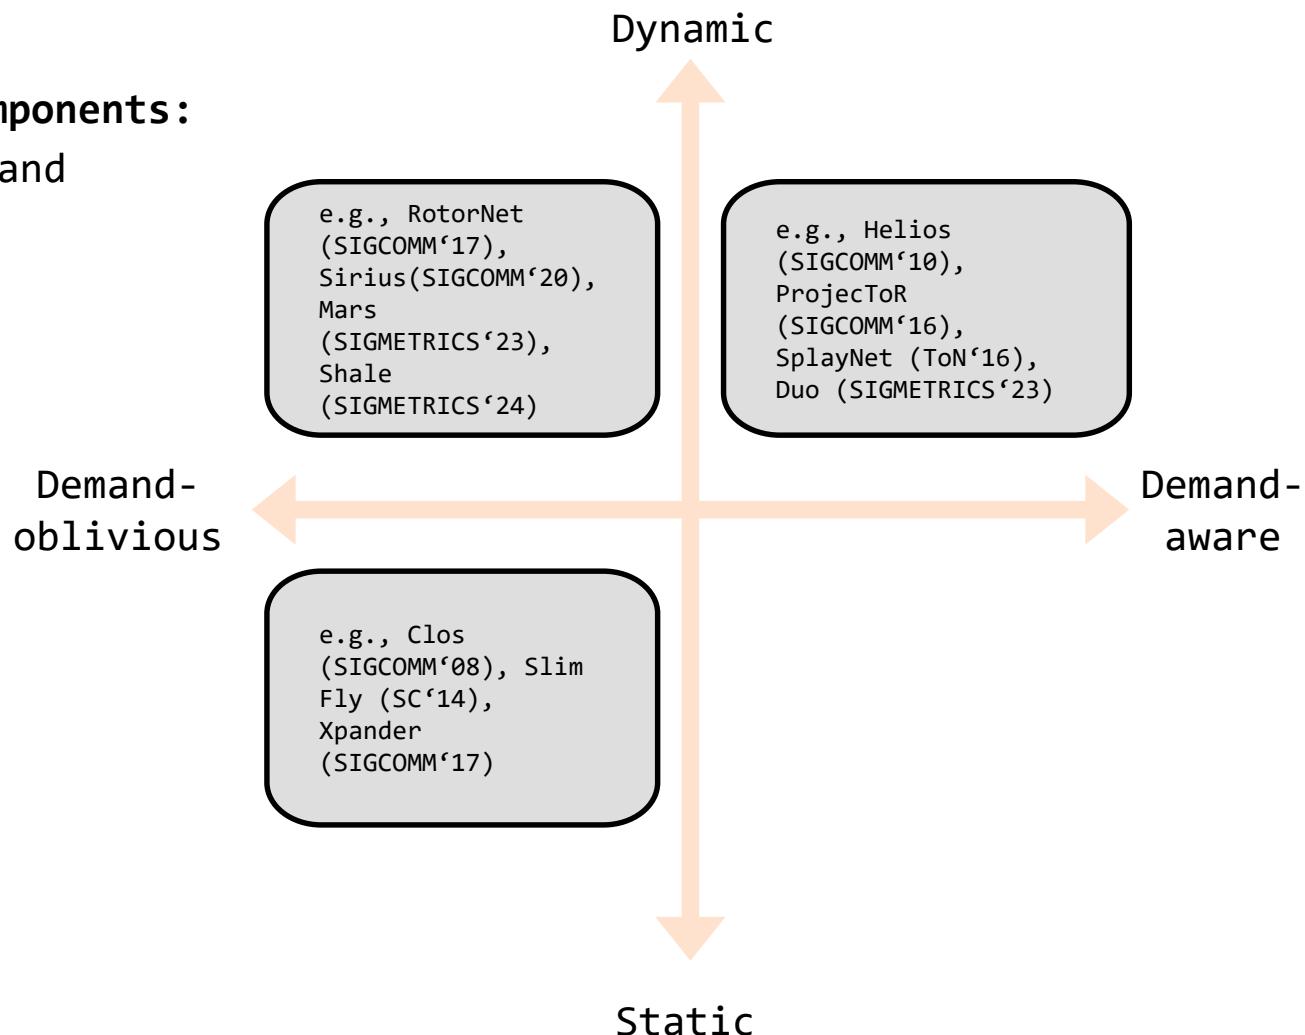
- demand-**oblivious** and
demand-**aware**
- static vs dynamic



Opportunity: Tech Diversity

Diverse topology components:

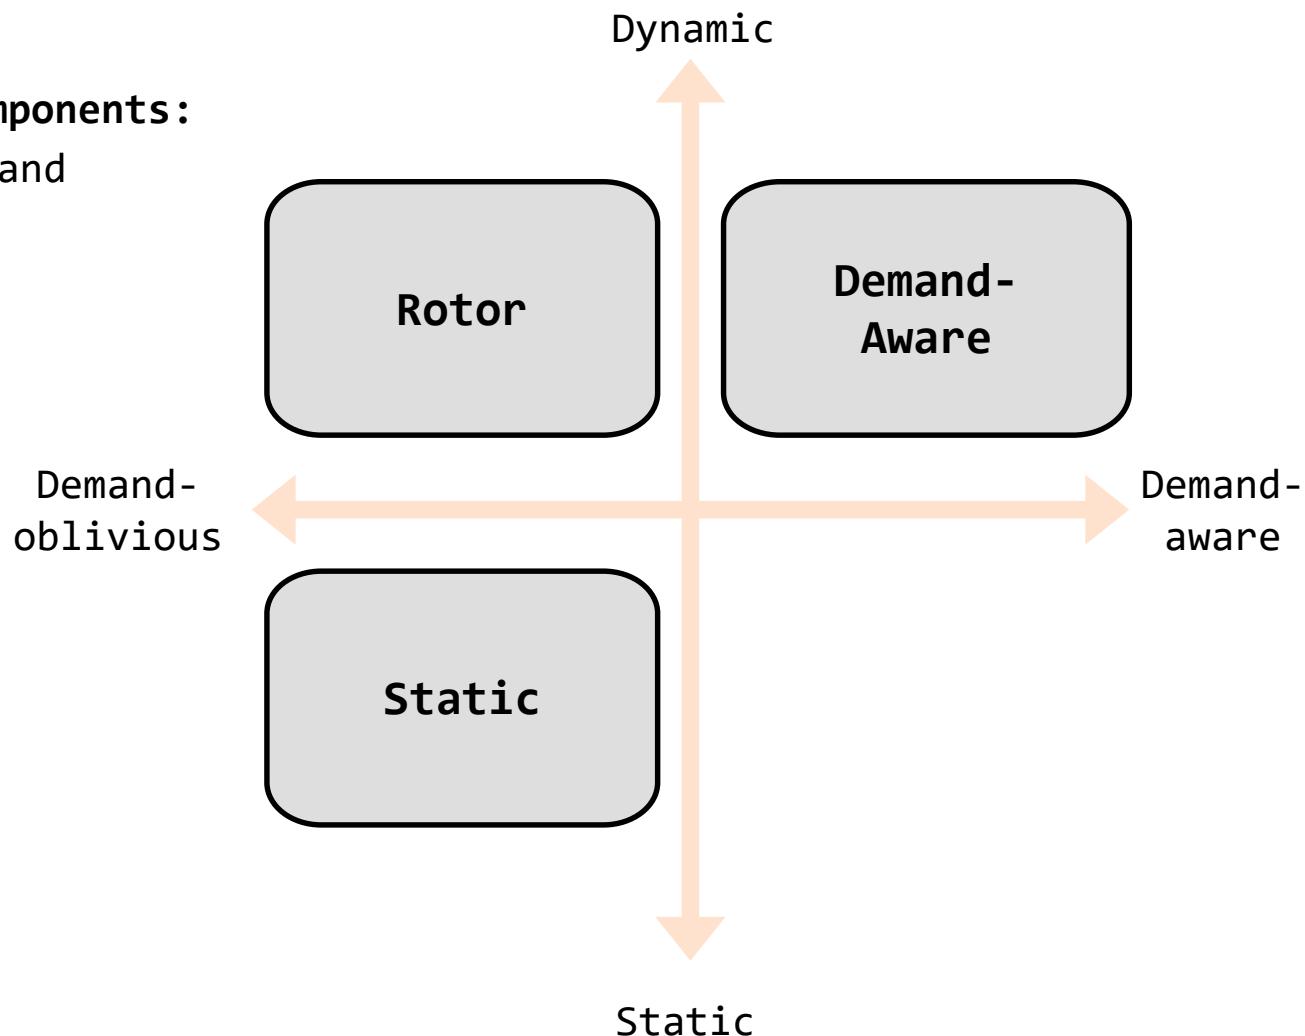
- demand-**oblivious** and
demand-**aware**
- static vs dynamic



Opportunity: Tech Diversity

Diverse topology components:

- demand-**oblivious** and
demand-**aware**
- static vs dynamic



Opportunity: Tech Diversity

Diverse topology components:

- demand-**oblivious** and
demand-**aware**
- static vs dynamic



Demand-
oblivious

Dynamic

Demand-
Aware

Rotor

Static

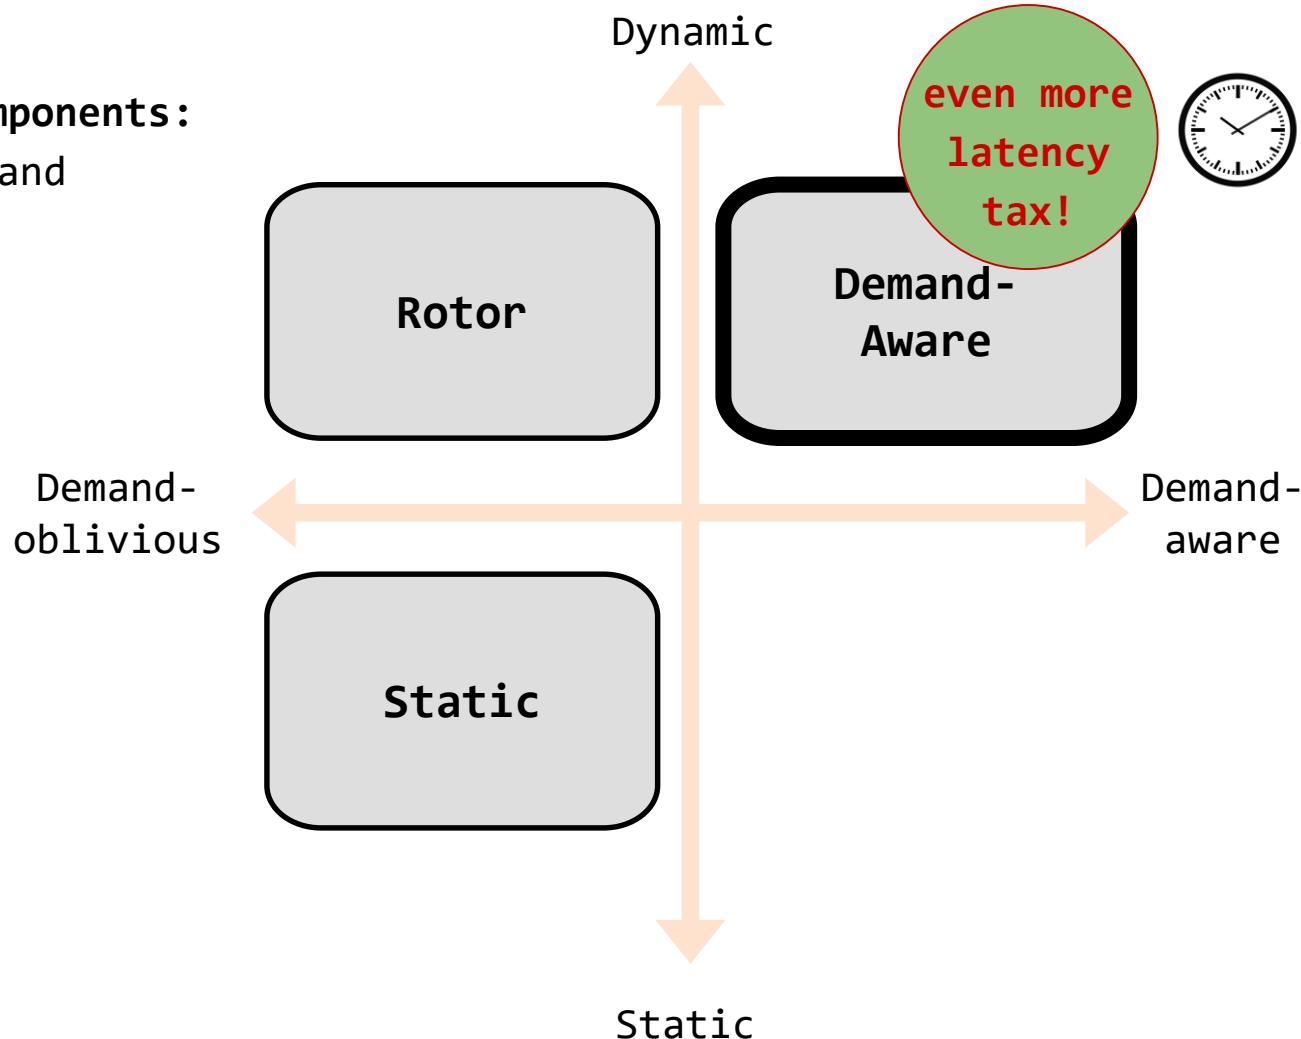
Demand-
aware

Static

Opportunity: Tech Diversity

Diverse topology components:

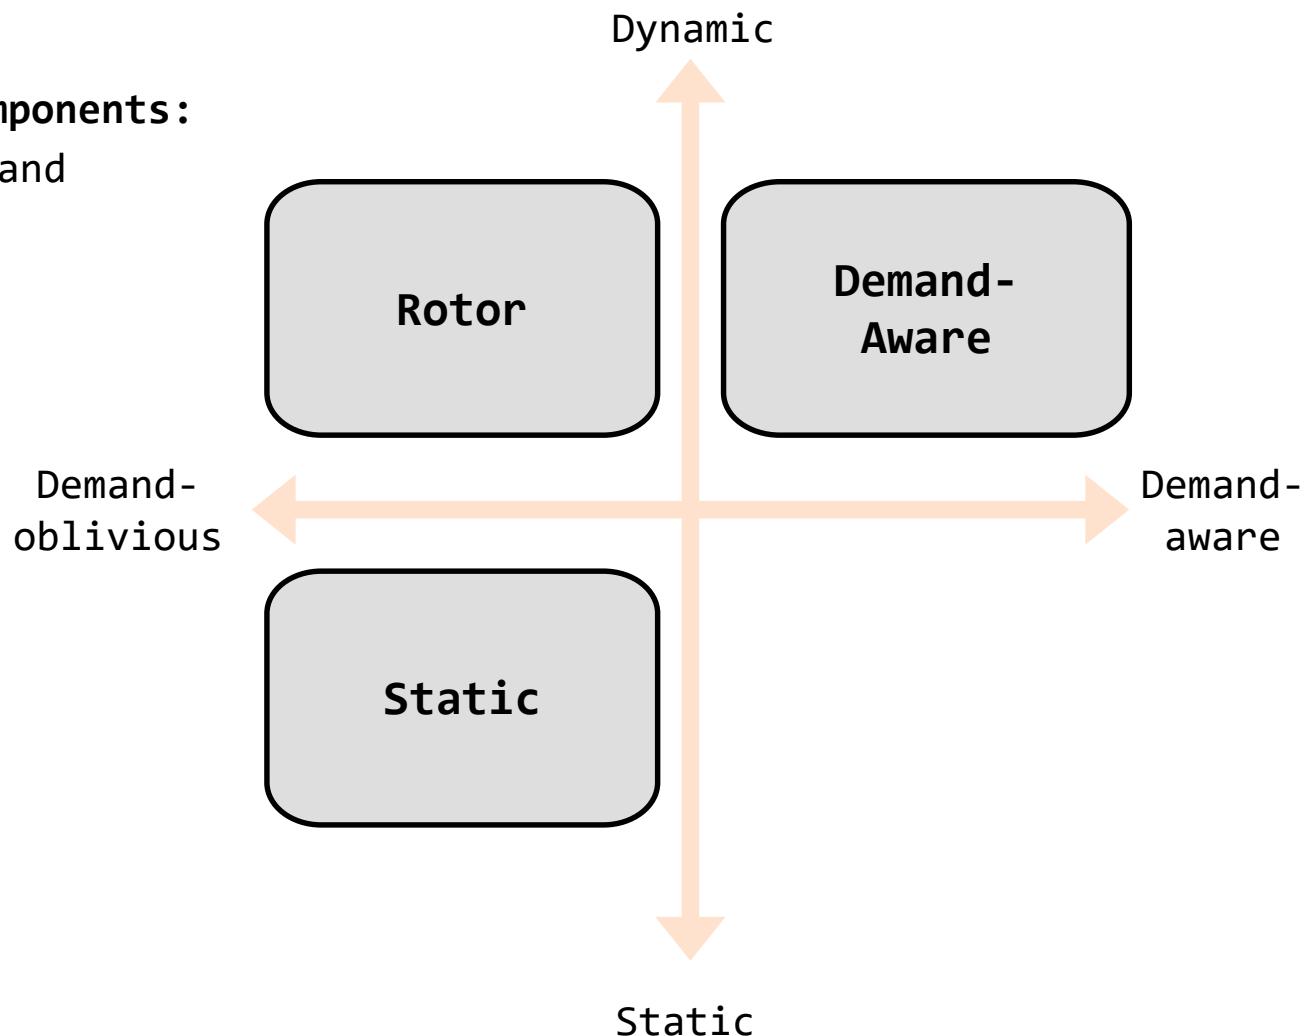
- demand-**oblivious** and demand-**aware**
- static vs dynamic



Opportunity: Tech Diversity

Diverse topology components:

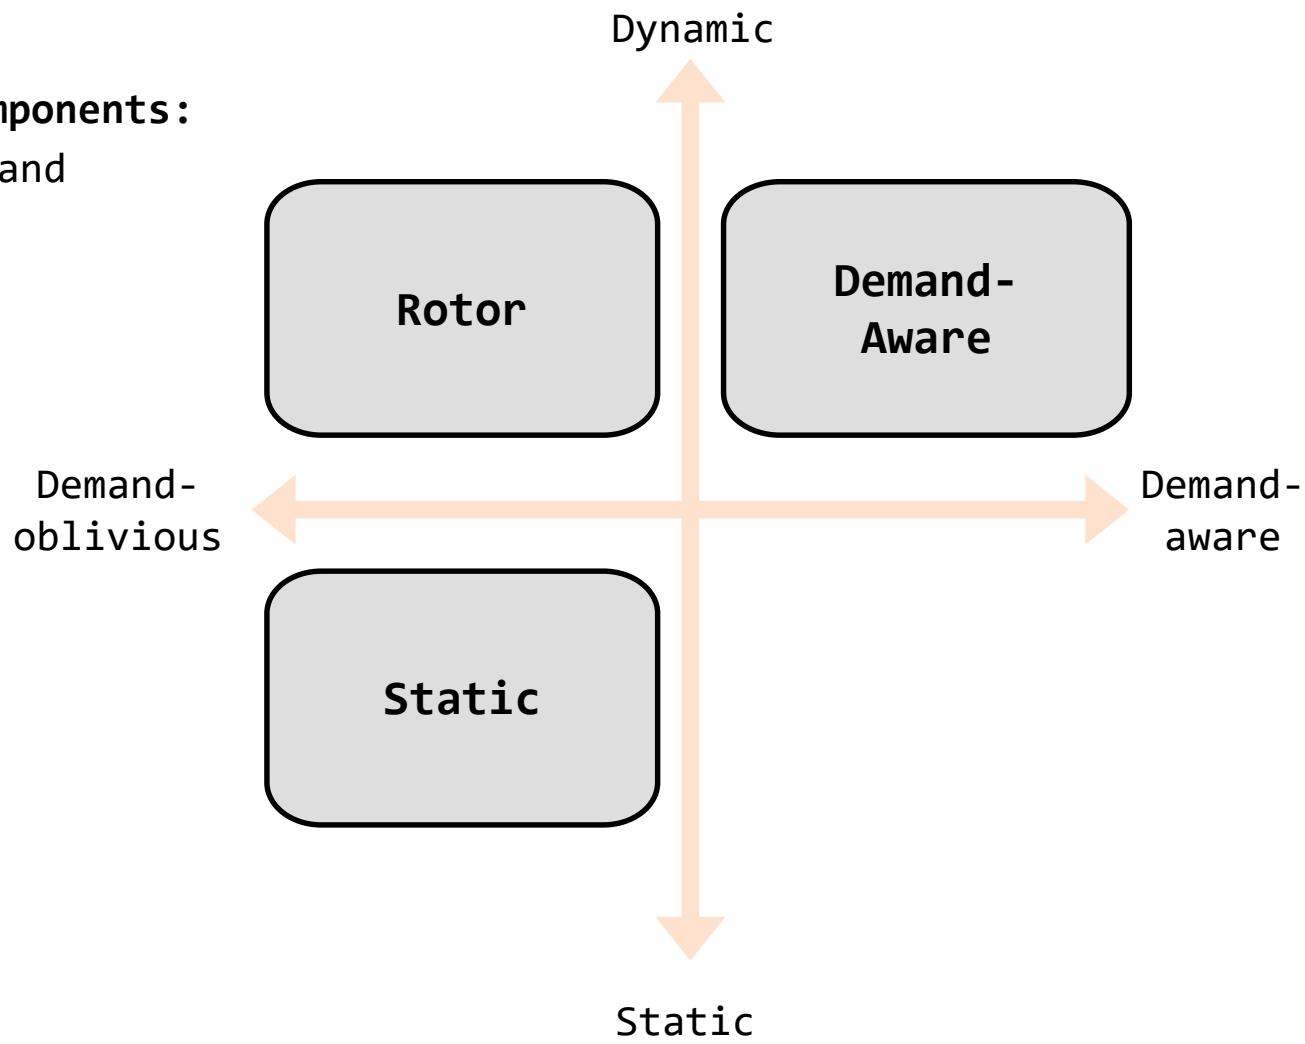
- demand-**oblivious** and
demand-**aware**
- static vs dynamic



Opportunity: Tech Diversity

Diverse topology components:

- demand-**oblivious** and
demand-**aware**
- static vs dynamic



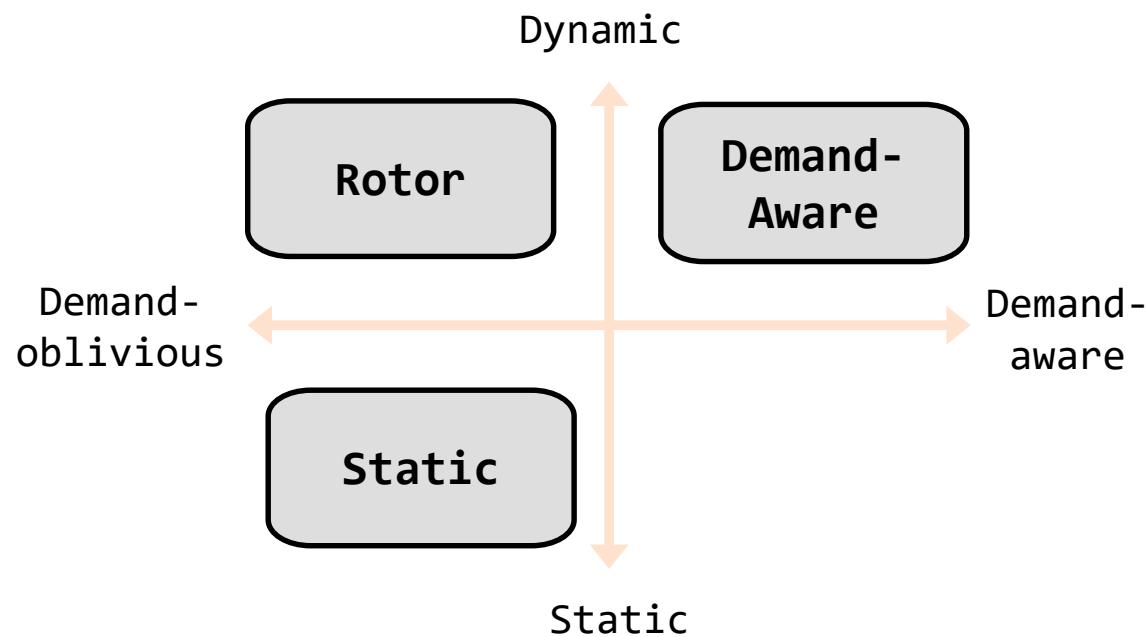
As always in CS:
It depends...

Examples:

Match or Mismatch?

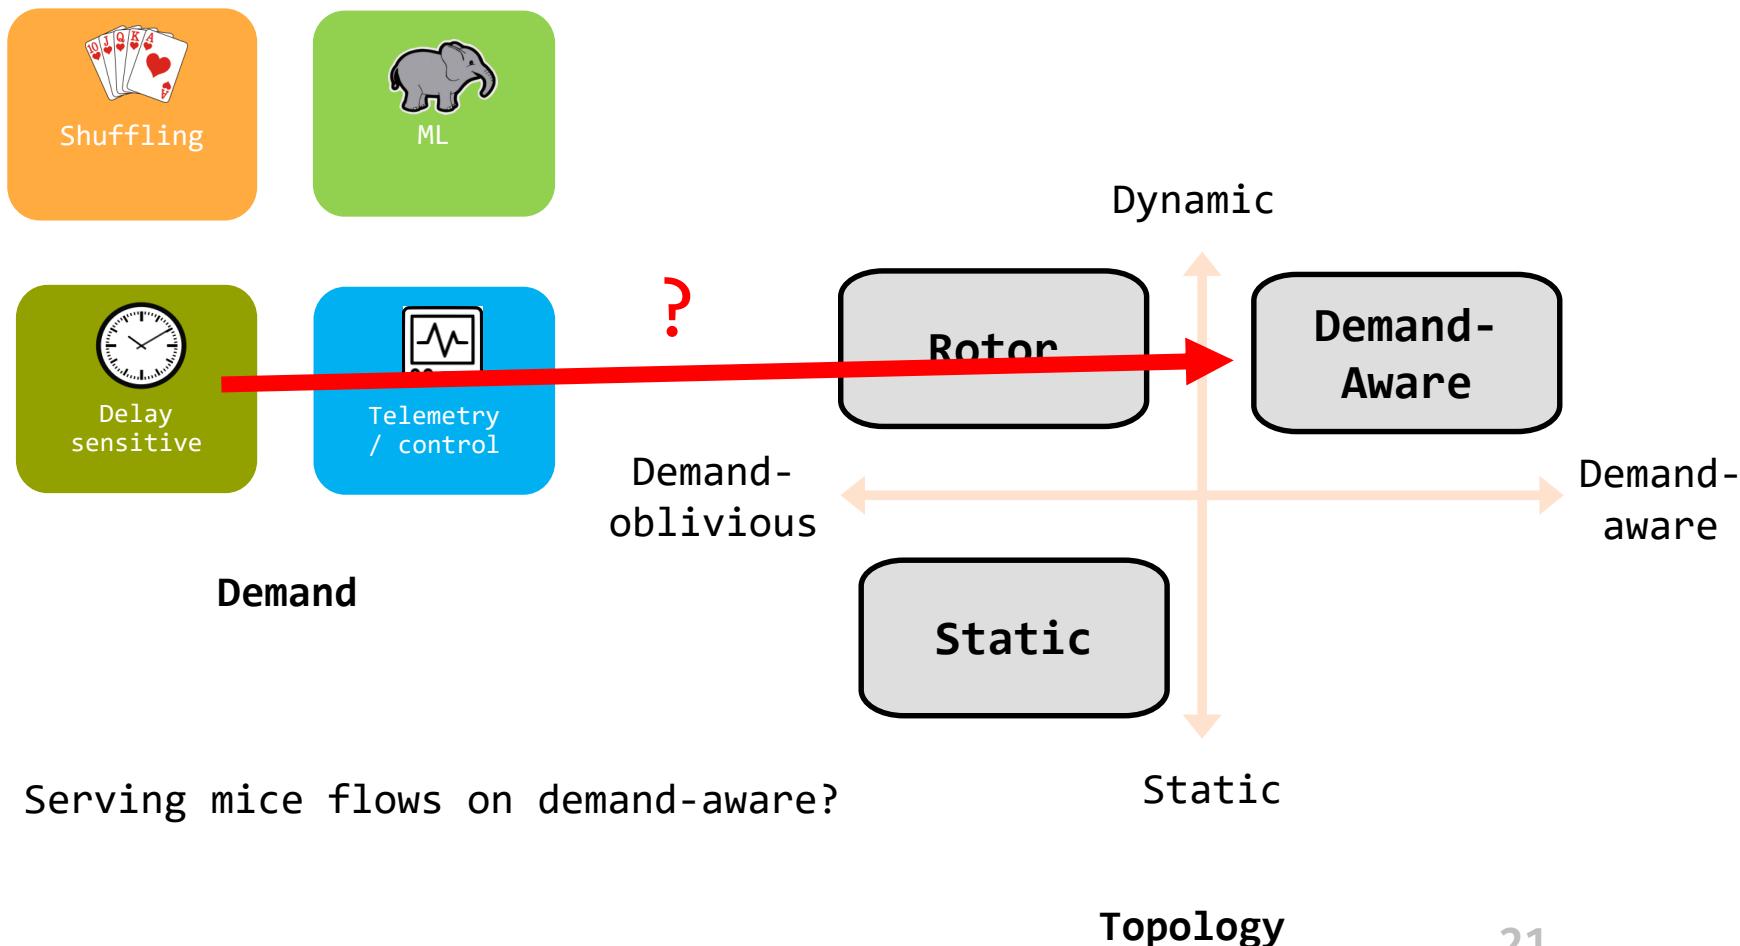


Demand

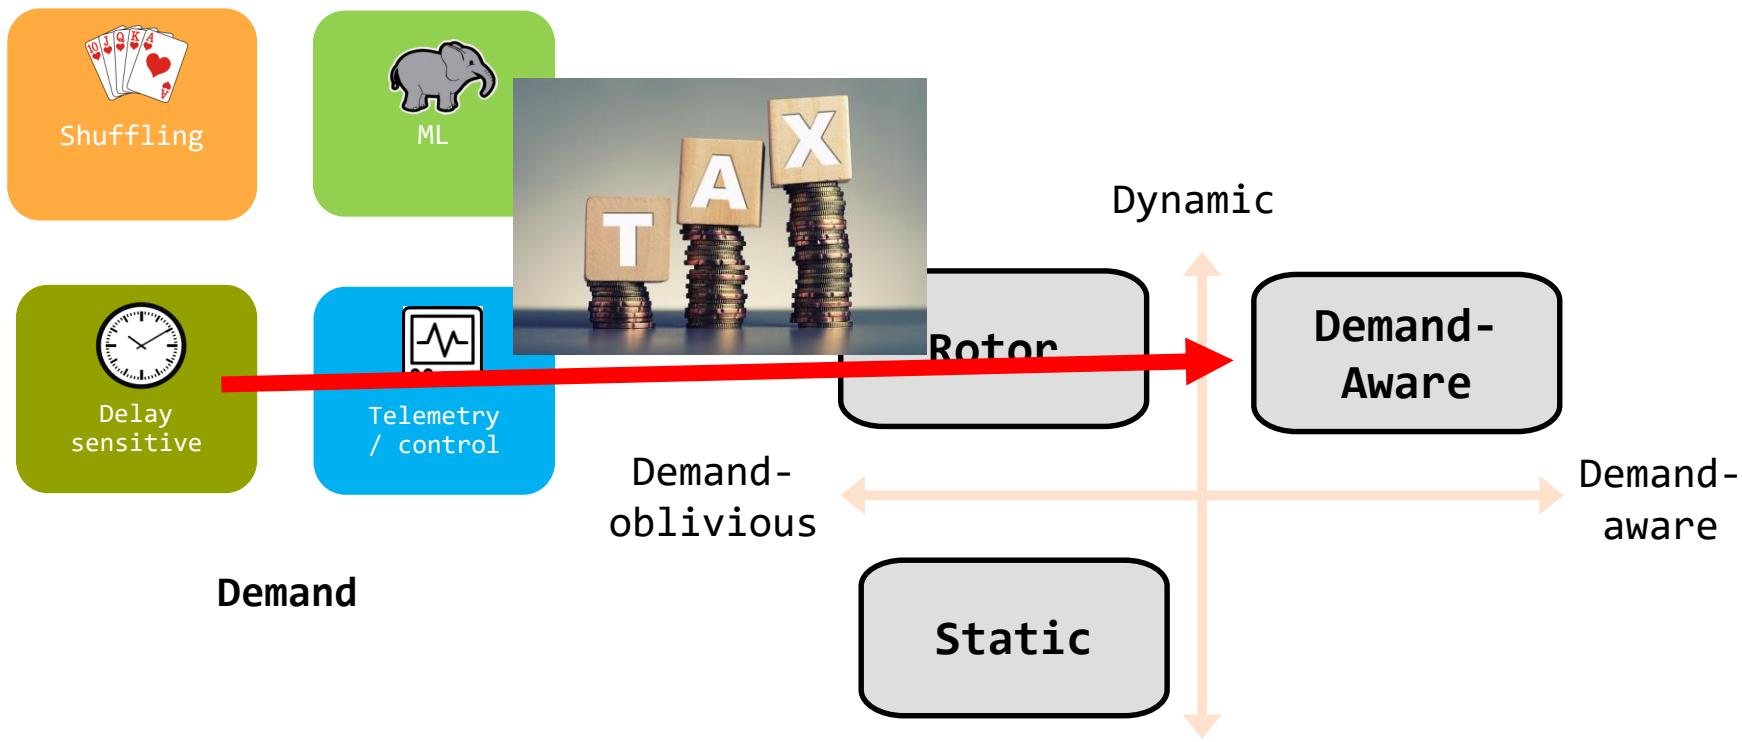


Examples:

Match or Mismatch?



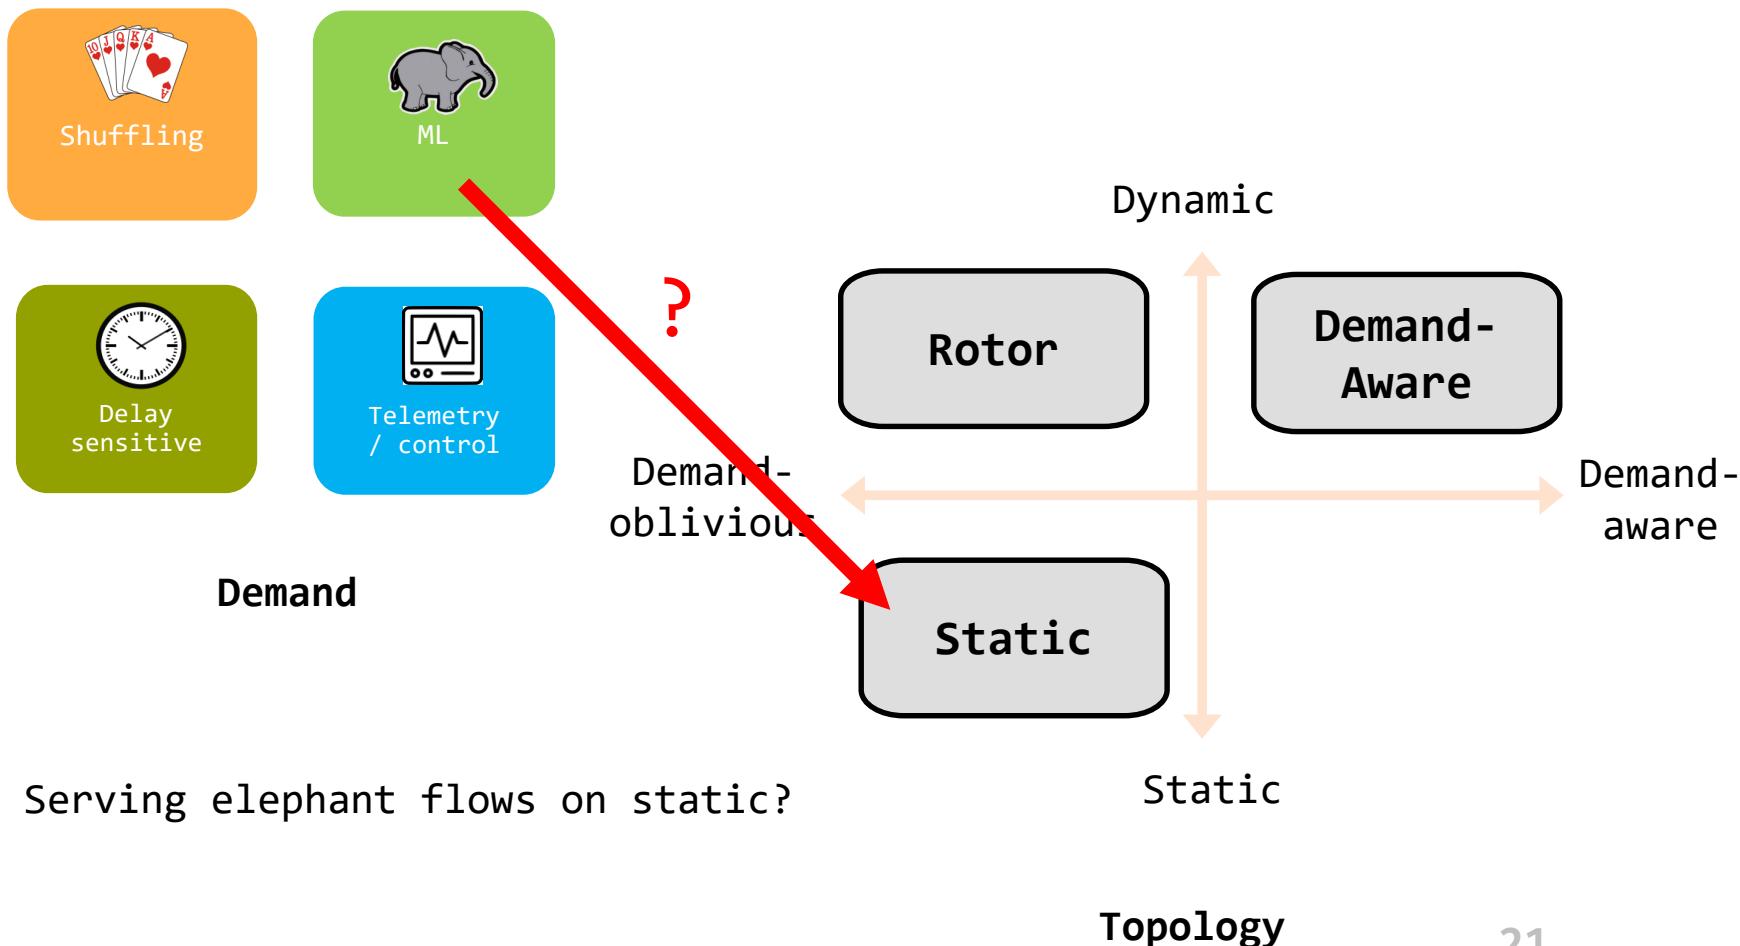
Examples: Match or Mismatch?



Serving mice flows on demand-aware?
Bad idea! Latency tax.

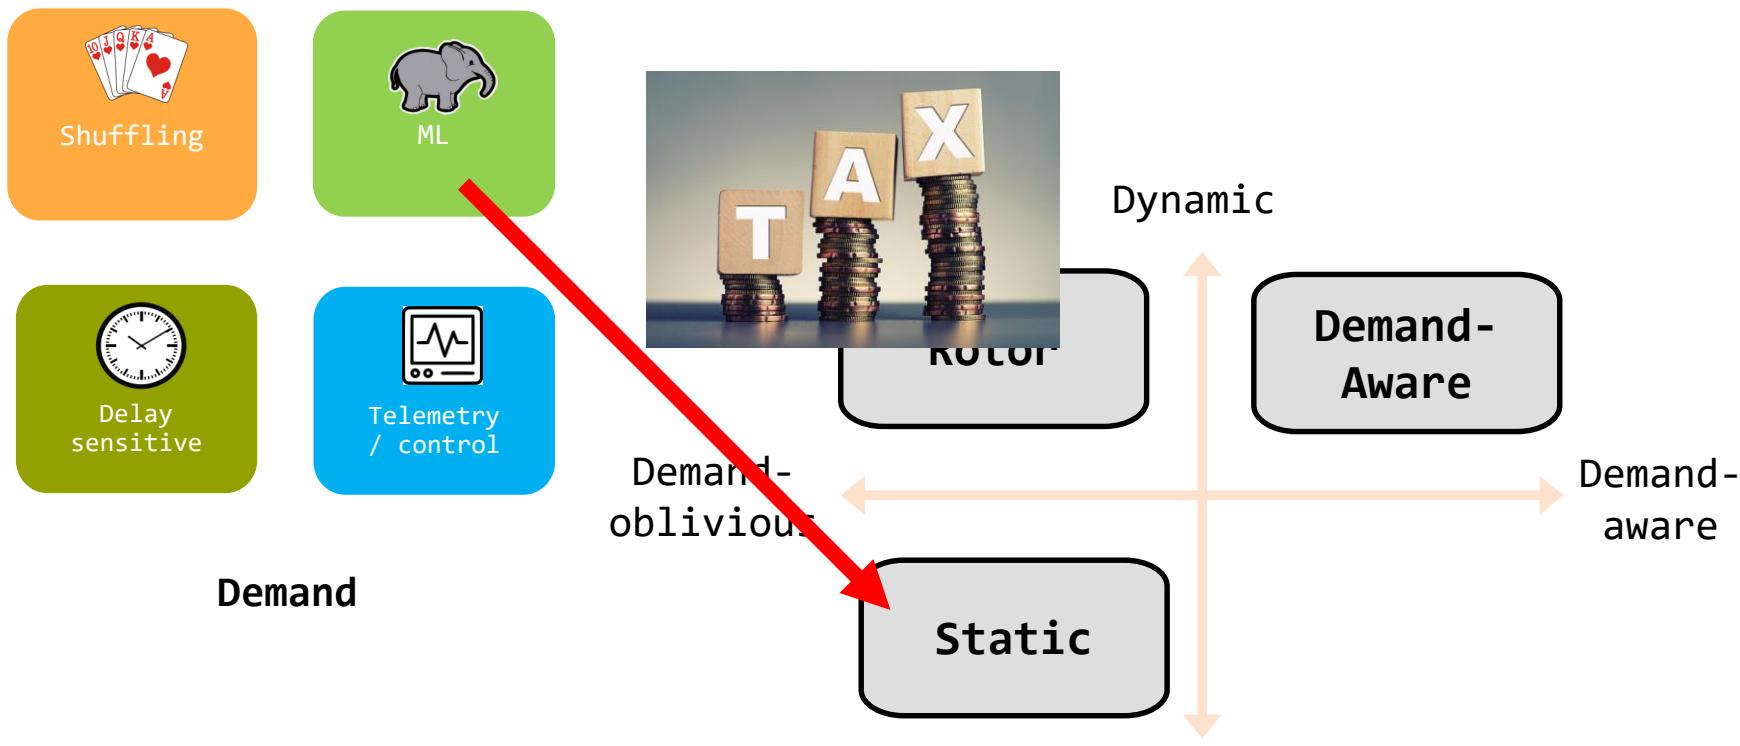
Examples:

Match or Mismatch?



Examples:

Match or Mismatch?



Serving elephant flows on static?
Bad idea! Bandwidth tax.

Examples: Match or Mismatch?



Demand

Demand-
oblivious

Dynamic

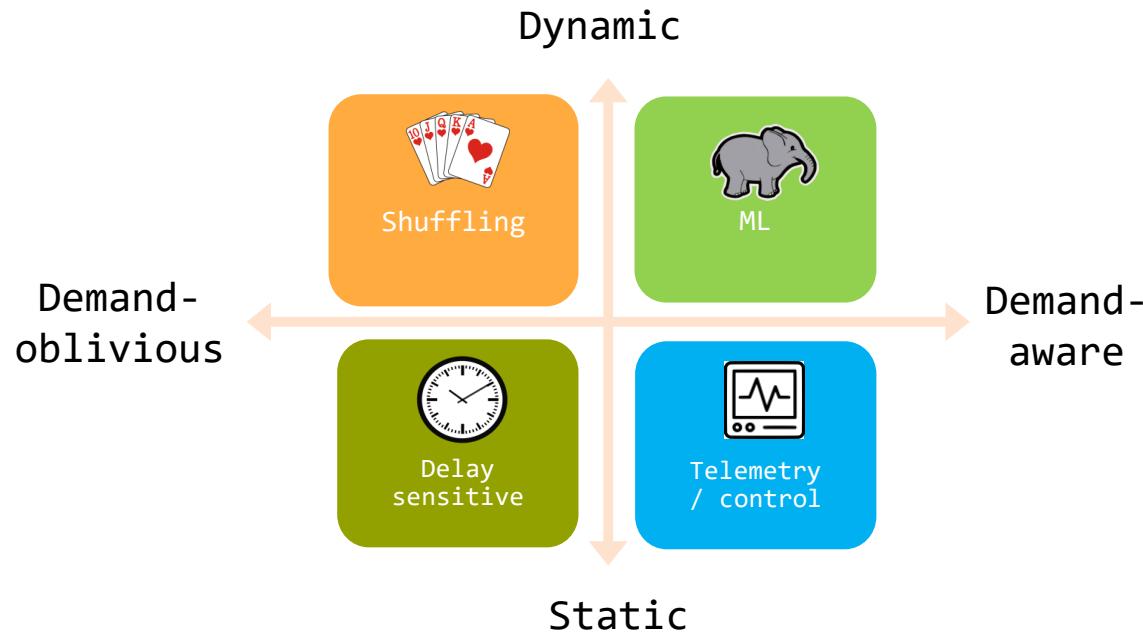
Demand-
aware

Static

Topology

Serving elephant flows on static?
Bad idea! Bandwidth tax.

Conceptual Solution



Conceptually ideal solution:

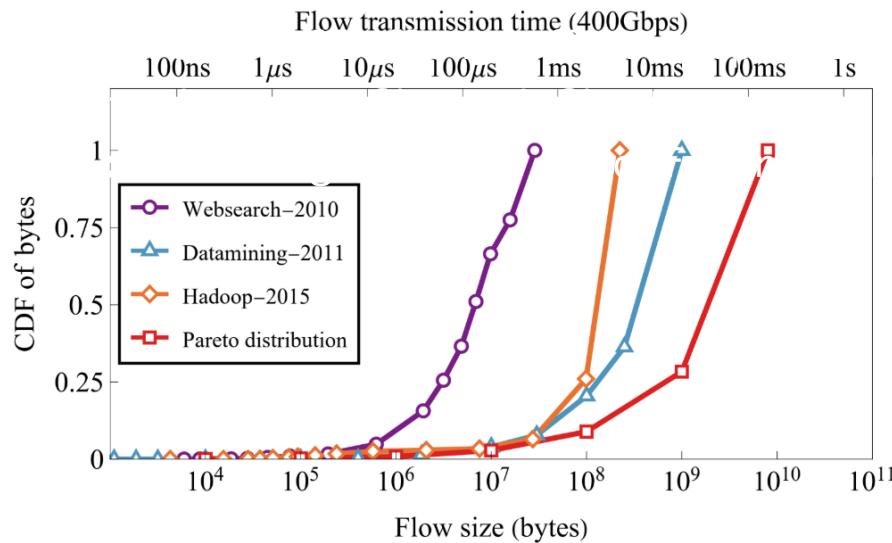
*Cerberus** serves traffic on the “best topology”!

Flow Size Matters

On what should topology type depend? We argue: **flow size.**

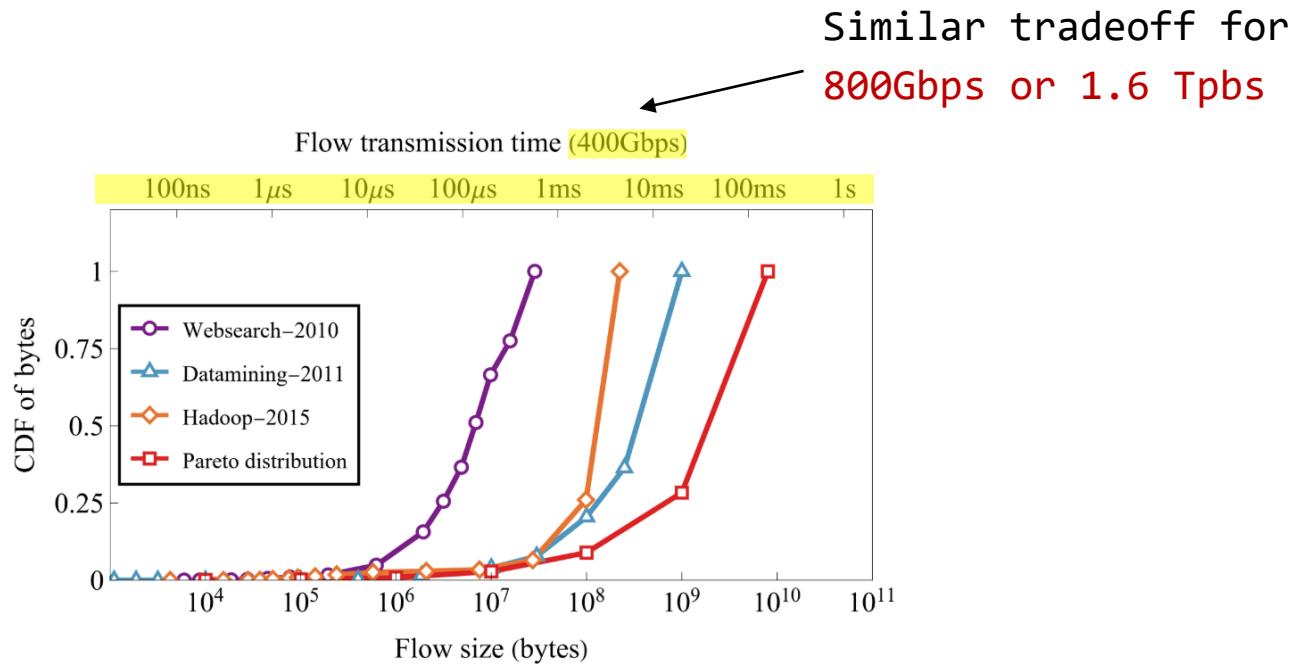
Flow Size Matters

On what should topology type depend? We argue: **flow size**.



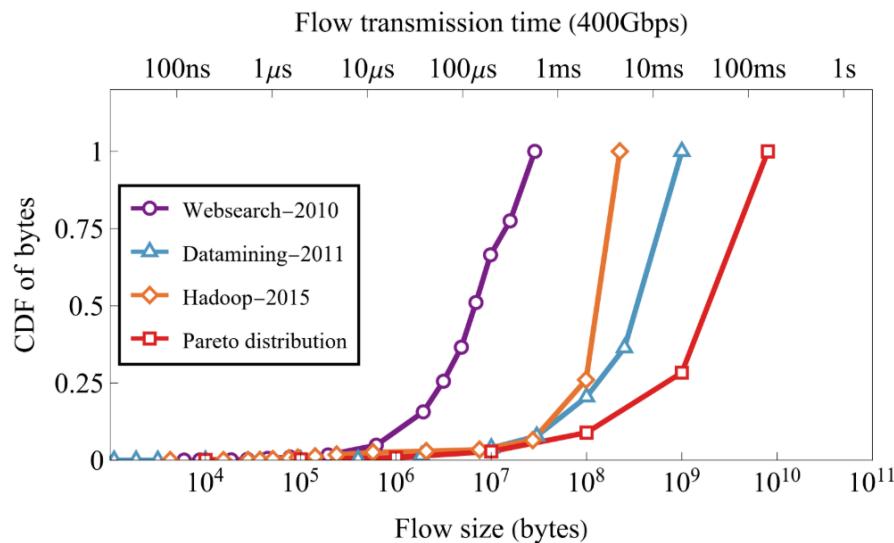
→ **Observation 1:** Different apps have different flow size distributions.

Flow Size Matters



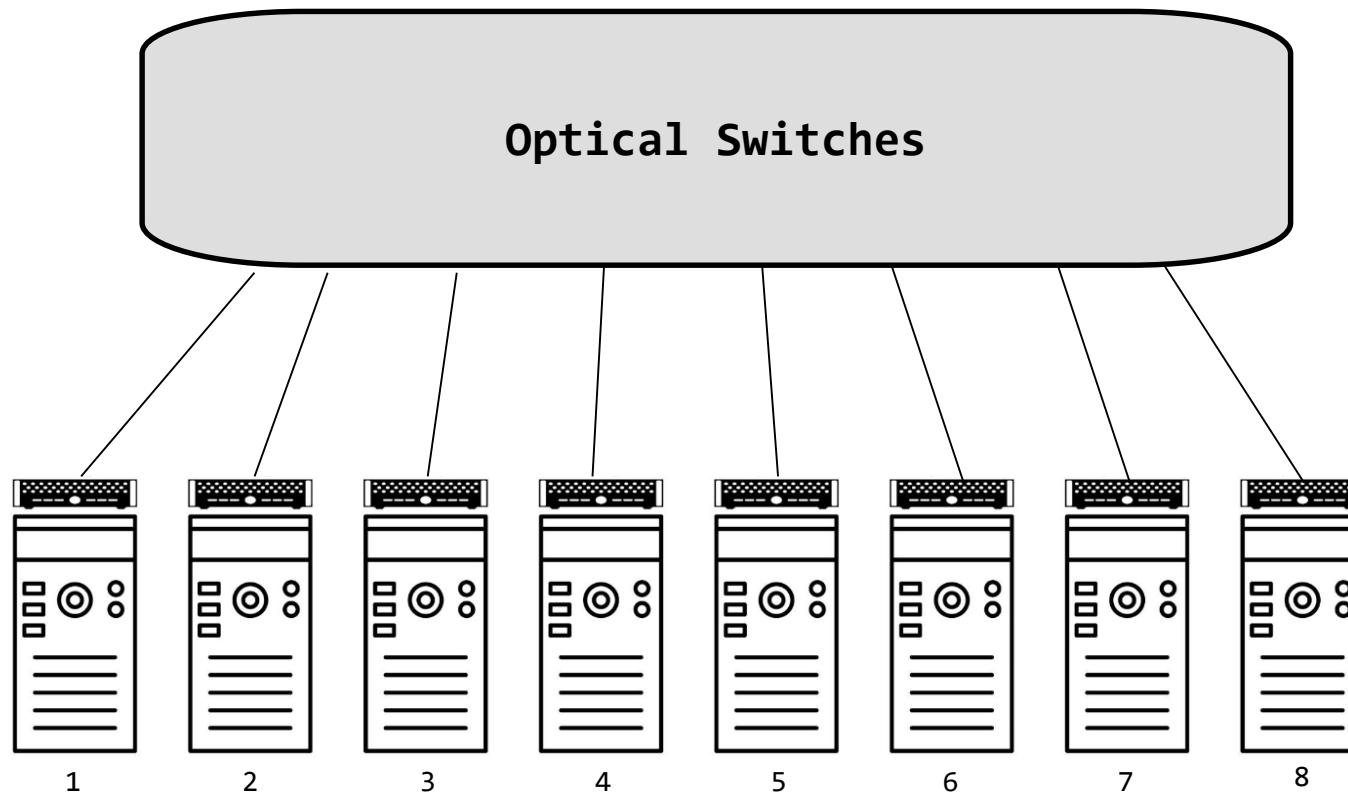
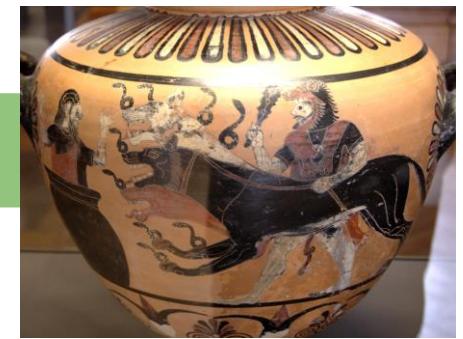
- **Observation 1:** Different apps have different flow size distributions.
- **Observation 2:** The transmission time of a flow depends on its **size**.

Flow Size Matters

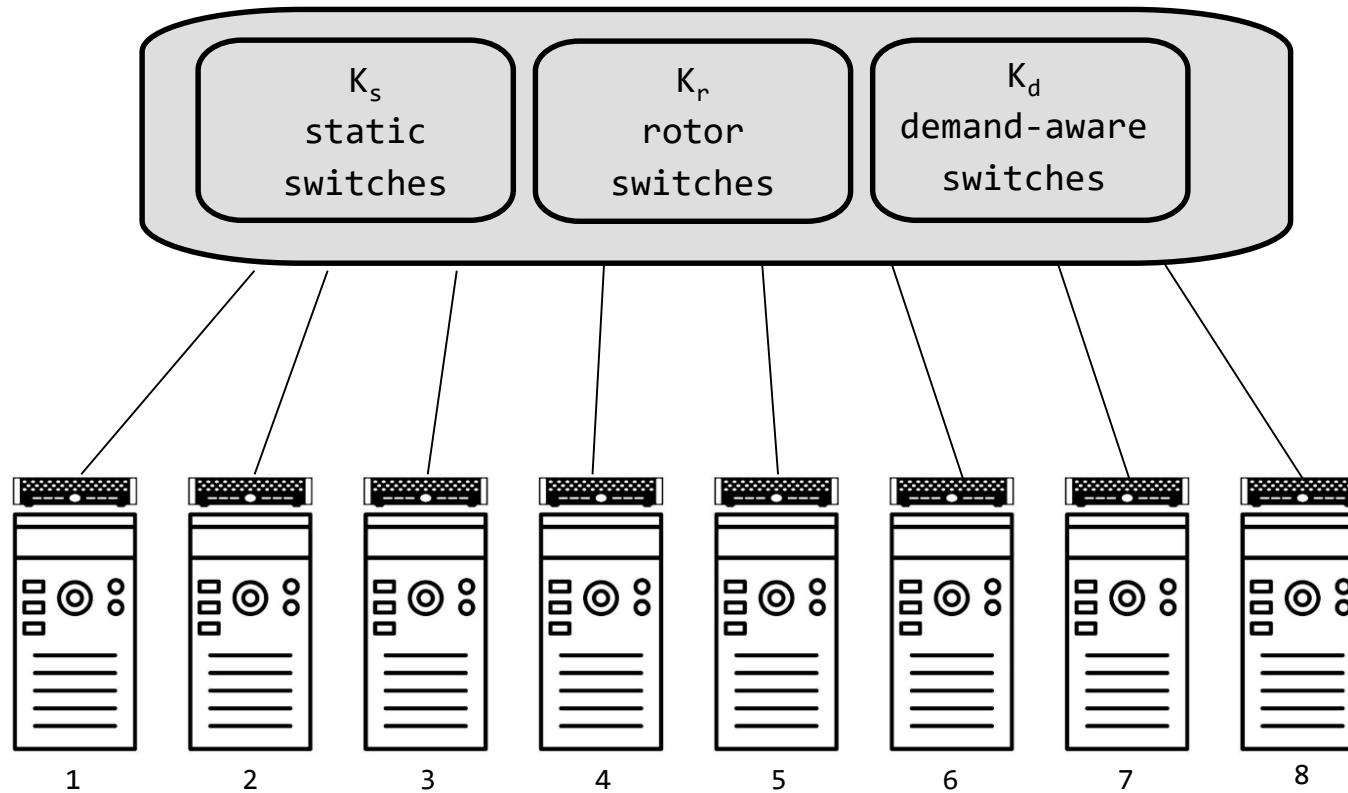
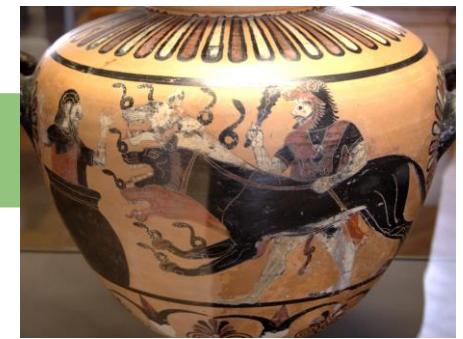


- **Observation 1:** Different apps have different flow size distributions.
- **Observation 2:** The transmission time of a flow depends on its **size**.
- **Observation 3:** For small flows, **flow completion time suffers** if network needs to be **reconfigured** first.
- **Observation 4:** For large flows, reconfiguration time may **amortize**.

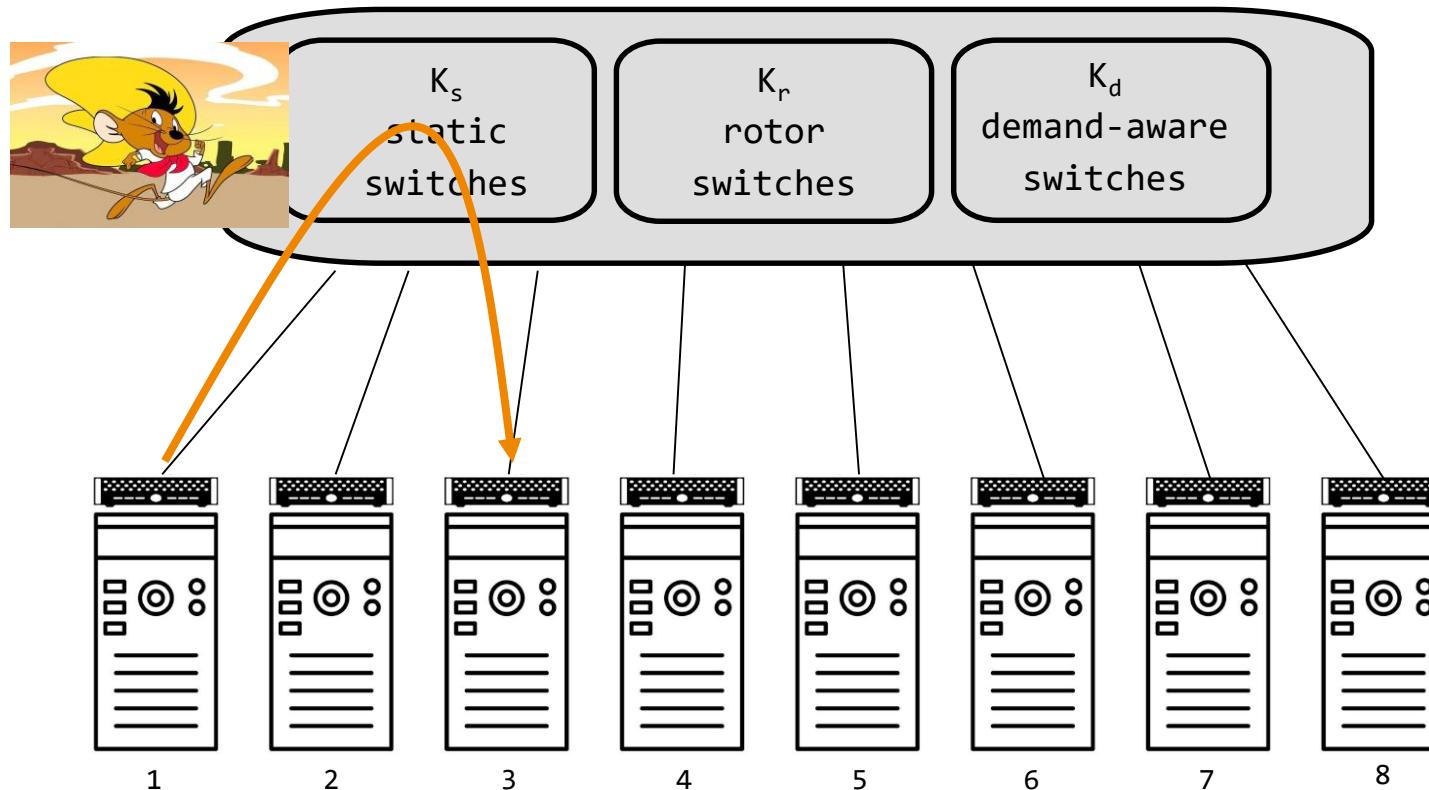
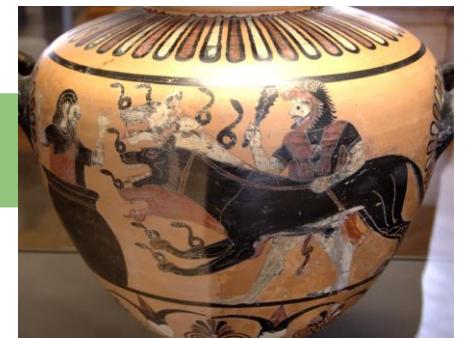
Cerberus



Cerberus

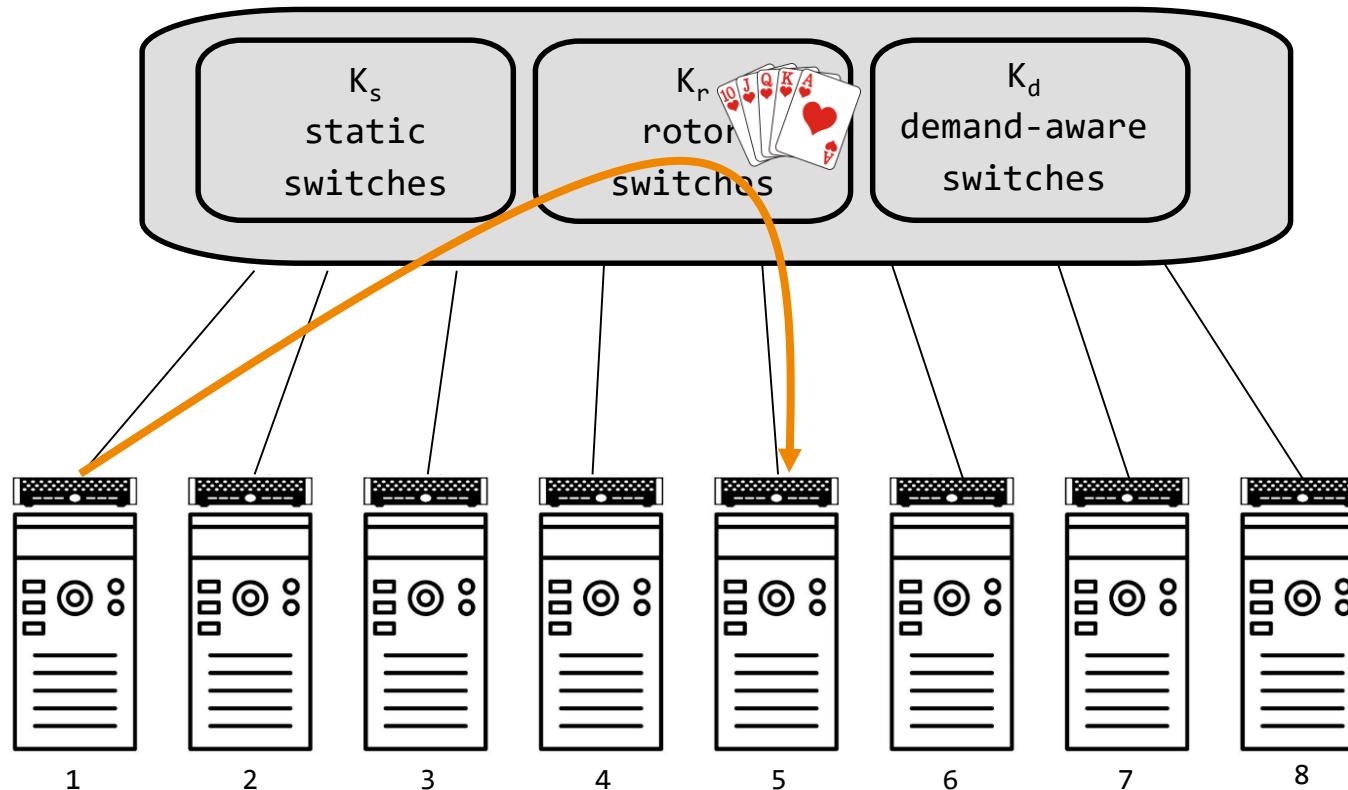
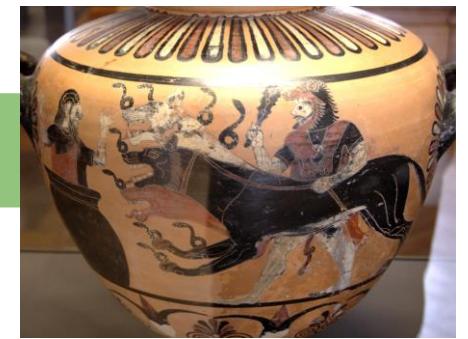


Cerberus



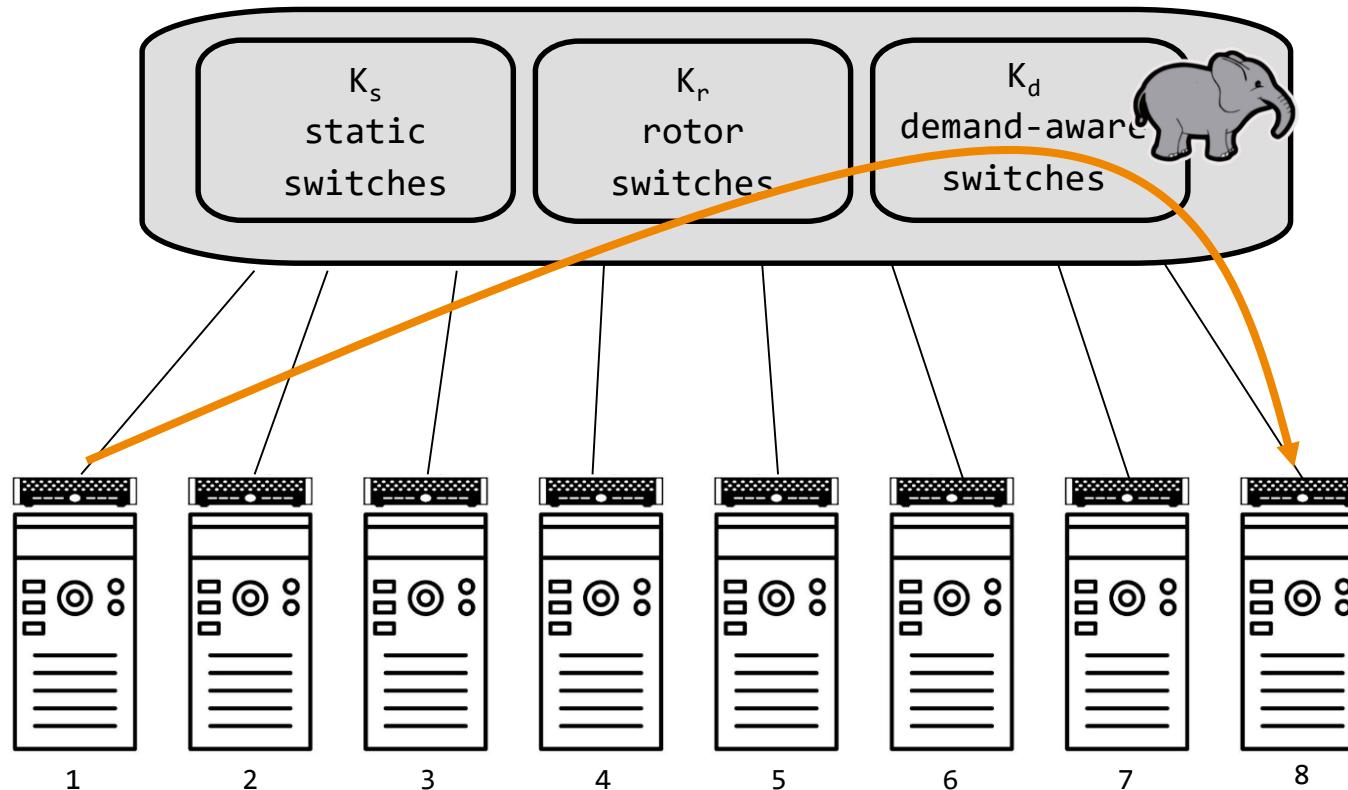
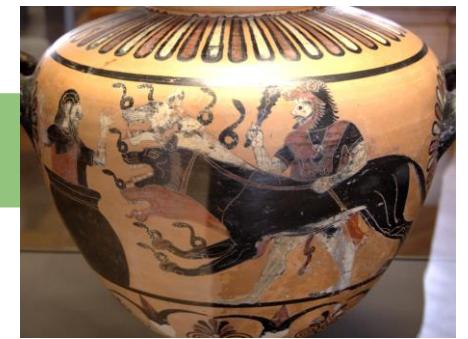
Scheduling: Small flows go via static switches...

Cerberus



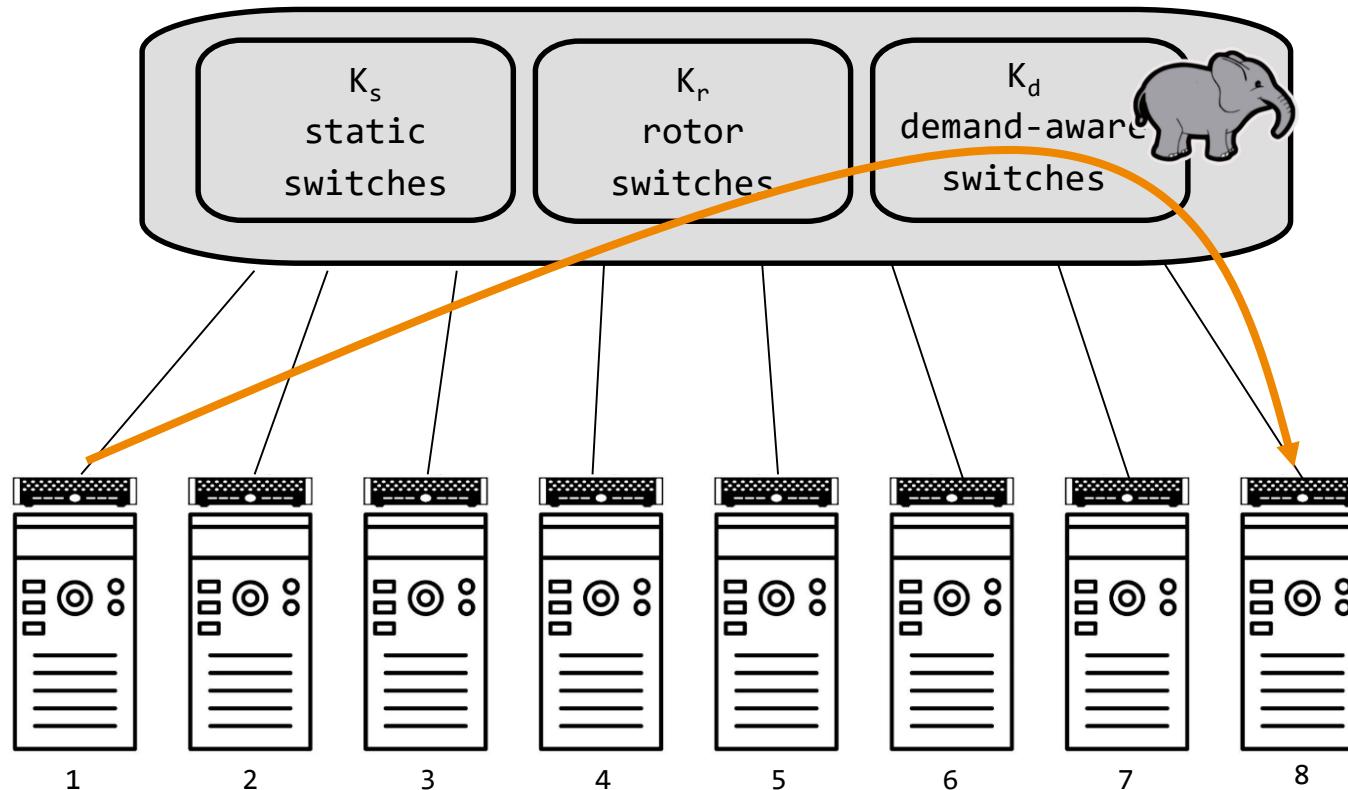
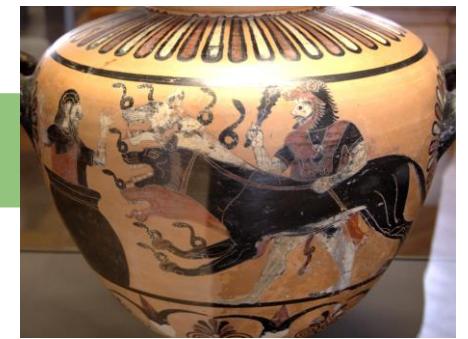
Scheduling: ... **medium flows** via rotor switches...

Cerberus



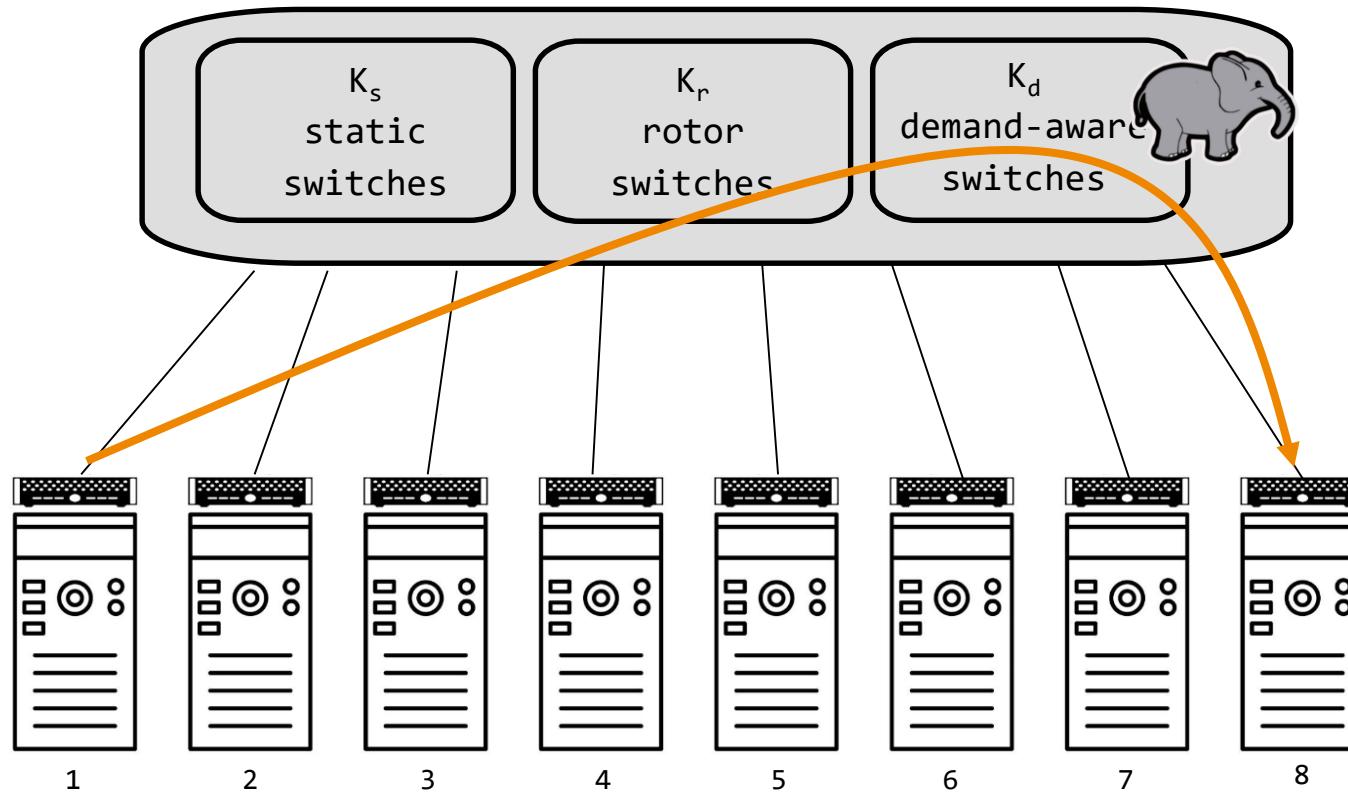
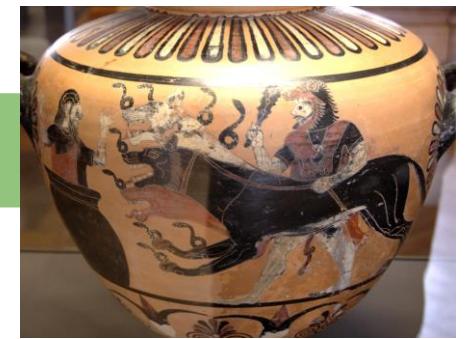
Scheduling: ... and **large flows** via demand-aware switches
(if one available, otherwise via rotor).

Cerberus



How good is it? Open problem. But there are bounds.

Cerberus



How good is it? Open problem. But there are bounds.

How to realize such an architecture?! Scalable control plane?

Throughput of RDCNs?

T

Demand Matrix

	1	2	3	4	5	6	7	8
1								
2								
3								
4								
5								
6								
7								
8								

Metric: throughput
of a demand matrix...

Abdu et al., SC 2016
Namyar et al., SIGCOMM 2021

Throughput of RDCNs?

Demand Matrix

T

	1	2	3	4	5	6	7	8
1								
2	■							
3			■					
4		■						
5				■				
6					■			
7						■		
8							■	

$\times \theta(T)$

Metric: throughput
of a **demand matrix**...

... is the maximal scale
down **factor** by which
traffic is **feasible**
 $0 \leq \theta(T) \leq 1.$

Abdu et al., SC 2016
Namyar et al., SIGCOMM 2021

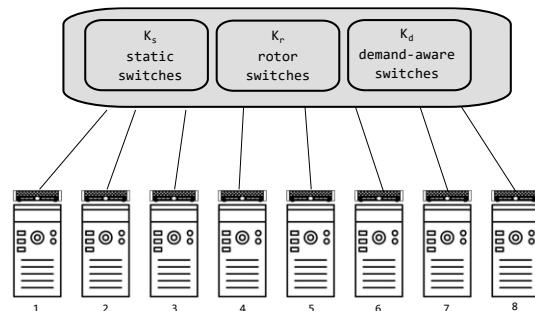
Throughput of RDCNs?

Demand Matrix

T

1	2	3	4	5	6	7	8
2							
3							
4							
5							
6							
7							
8							

$\times \theta(T) \Rightarrow$



Metric: throughput
of a **demand matrix**...

... is the maximal scale
down **factor** by which
traffic is **feasible**
 $0 \leq \theta(T) \leq 1.$

Throughput of network θ^* :
worst case T

Abdu et al., SC 2016
Namyar et al., SIGCOMM 2021

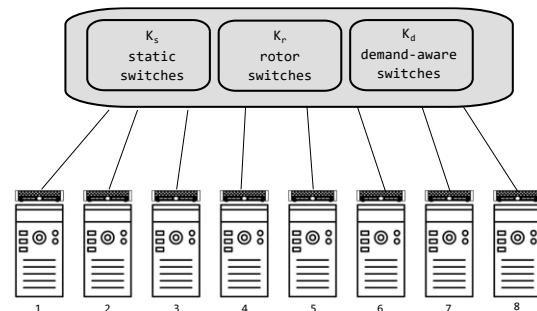
Throughput of RDCNs?

Demand Matrix

1	2	3	4	5	6	7	8
2							
3							
4							
5							
6							
7							
8							

T

$\times \theta(T) \Rightarrow$



Metric: throughput
of a demand matrix...

... is the maximal scale
down **factor** by which
traffic is **feasible**
 $0 \leq \theta(T) \leq 1$.

Worst T for
different
networks?

Throughput of network θ^* :
worst case T

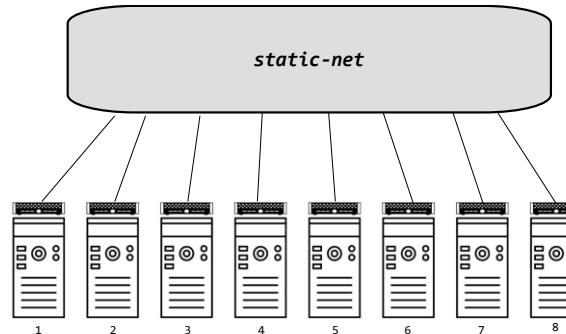
Abdu et al., SC 2016
Namyar et al., SIGCOMM 2021

Throughput: Expander

T

Demand Matrix

	1	2	3	4	5	6	7	8
1								
2								
3								
4								
5								
6								
7								
8								



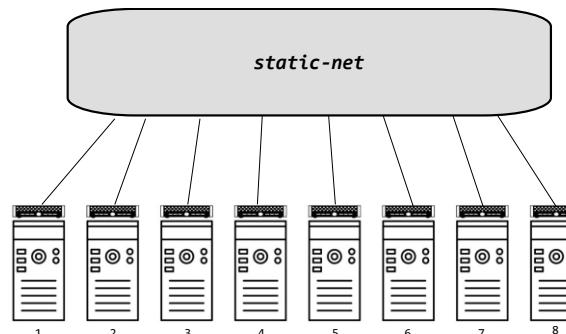
Permutation matrix
is the **worst demand**

Throughput: Expander

T

Demand Matrix

	1	2	3	4	5	6	7	8
1								
2	1							
3		1						
4			1					
5				1				
6					1			
7						1		
8							1	



Permutation matrix
is the **worst demand**

$$\theta^* \leq \frac{1}{\text{epl}(G(k))}$$

Expected path length

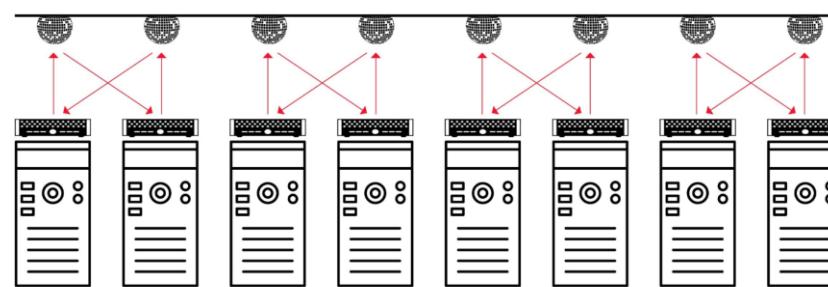


Bandwidth tax

Throughput: Demand-Aware

Demand Matrix
 T

	1	2	3	4	5	6	7	8
1								
2	1							
3		1						
4			1					
5				1				
6					1			
7						1		
8							1	

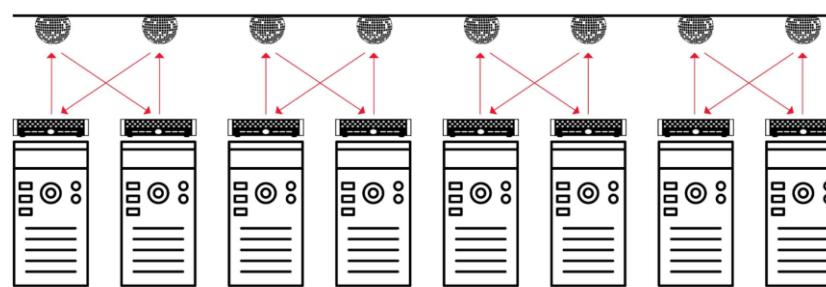


Permutation matrix
is the **best demand** ☺

Throughput: Demand-Aware

Demand Matrix
 T

1	2	3	4	5	6	7	8
1							
2	■						
3		■					
4			■				
5				■			
6					■		
7						■	
8							■



Permutation matrix
is the **best demand** ☺

Demand-aware performs poorly for **unstructured demand**.

Throughput formula is a function of Latency tax.

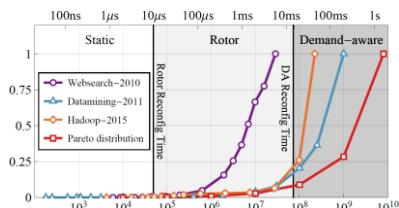
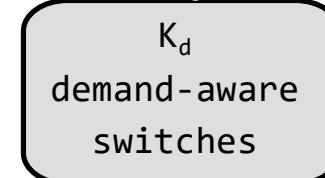
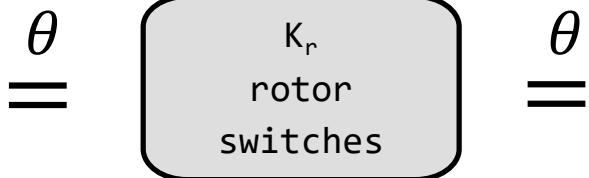
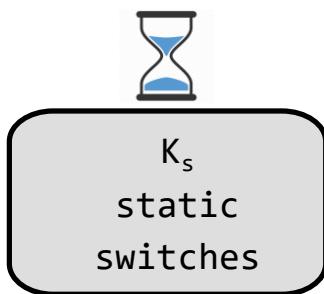
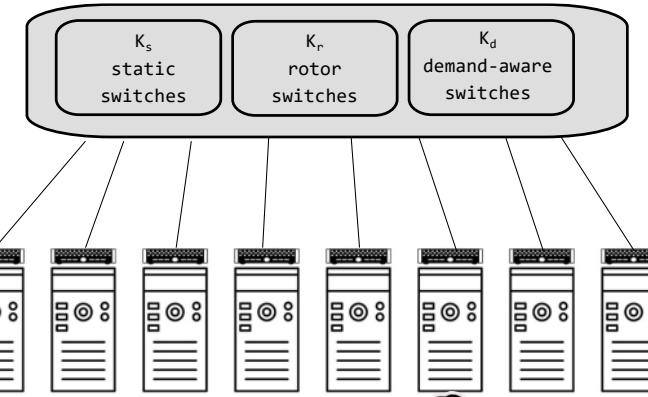


Throughput: Cerberus

Demand Matrix

	1	2	3	4	5	6	7	8
1								
2								
3								
4								
5								
6								
7								
8								

T



Cerberus **balances optimally** across switch types.

Throughput depends on both:



Bandwidth tax + Latency tax

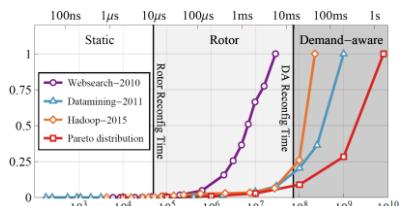
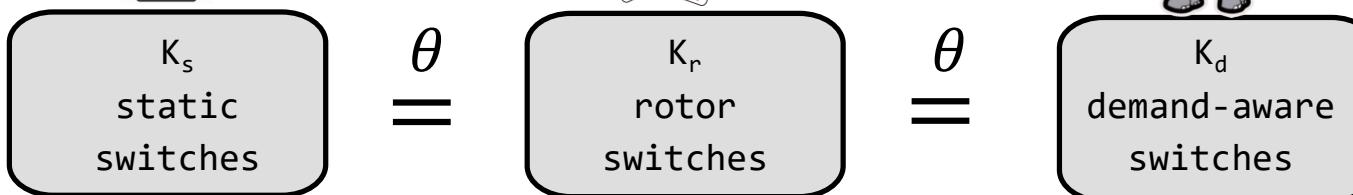
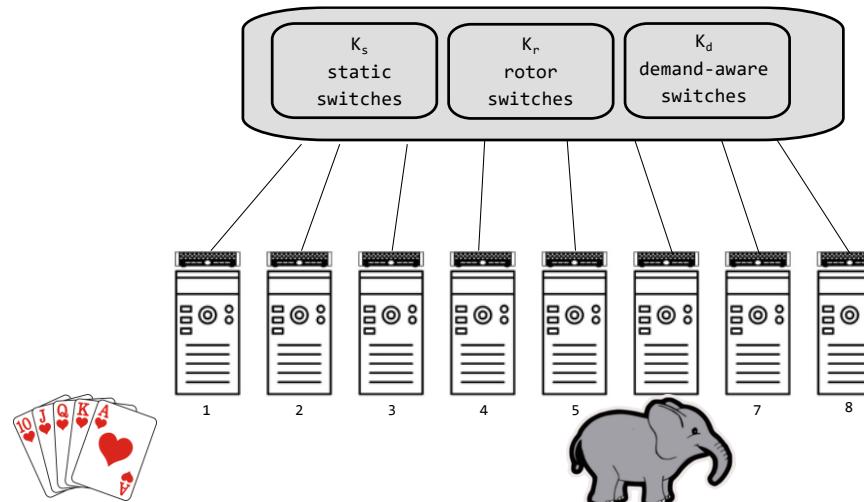


Throughput: Cerberus

Demand Matrix

T

	1	2	3	4	5	6	7	8
1								
2								
3								
4								
5								
6								
7								
8								



Cerberus **balances optimally** across switch types.

Throughput depends on both:



Bandwidth tax + Latency tax

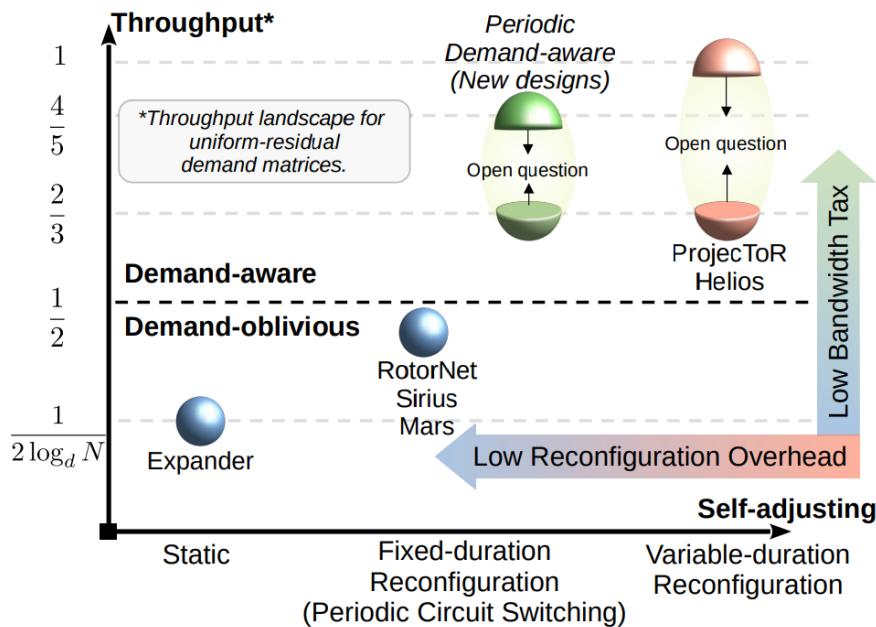


But theoretical **problem open**: no exact formula known yet.

Throughput:

Many More Open Questions

- Throughput bounds for many designs not fully understood yet



Addanki et al., arXiv 2025:
<https://arxiv.org/pdf/2405.20869.pdf>

Addanki et al., Vermillion:
<https://arxiv.org/pdf/2504.09892.pdf>

How to support such
dynamic networks on
other layers?

More Challenges:

Network Layer?

- *ECMP* reconvergence?! Benefits of *Valiant* routing?
- How to avoid packet *reorderings*? *RDMA* network cards don't like them!
- Routing in hybrid networks: segregated vs *non-segregated*?
- First ideas: *Local* routing! Techniques from dynamic P2P systems?

Duo: A High-Throughput Reconfigurable Datacenter Network Using Local Routing and Control

JOHANNES ZERWAS, TUM School of Computation, Information and Technology, Technical University of Munich, Germany

CSABA GYÖRGYI, University of Vienna and ELTE Eötvös Loránd University, Austria and Hungary

ANDREAS BLENK, Siemens AG, Germany

STEFAN SCHMID, TU Berlin & Fraunhofer SIT, Germany

CHEN AVIN, Ben-Gurion University, Israel

The performance of many cloud-based applications critically depends on the capacity of the underlying datacenter network. A particularly innovative approach to improve the throughput in datacenters is enabled by emerging optical technologies, which allow to dynamically adjust the physical network topology, both in an oblivious or demand-aware manner. However, such topology engineering, i.e., the operation and control of dynamic datacenter networks, is considered complex and currently comes with restrictions and overheads.

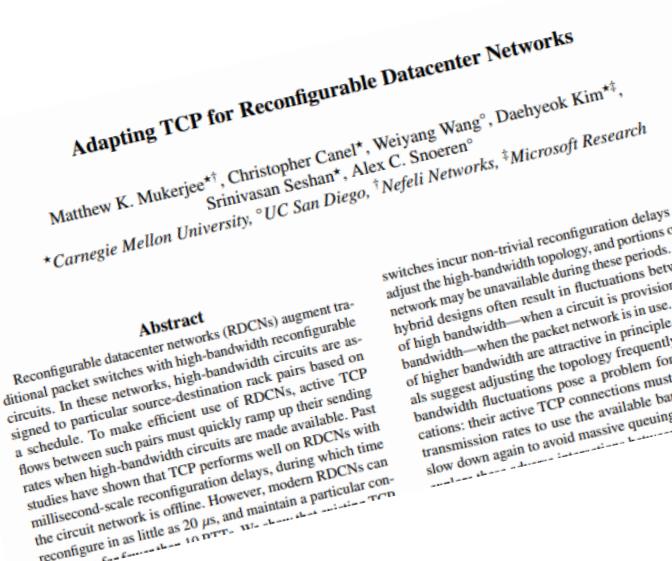
We present Duo, a novel demand-aware reconfigurable rack-to-rack datacenter network design realized with a simple and efficient control plane. Duo is based on the well-known de Bruijn topology (implemented using a small number of optical circuit switches) and the key observation that this topology can be enhanced using dynamic ("opportunistic") links between its nodes.

In contrast to previous systems, Duo has several desired features: i) It makes effective use of the network

More Challenges:

Congestion Control?

- First ideas for quickly reacting TCP: *ReTCP*, *PowerTCP*, ...
- Or better completely different approach? Even centralized?!



switches incur non-trivial reconfiguration delays while they adjust the high-bandwidth topology, and portions of the circuit network may be unavailable during these periods. Hence, such hybrid designs often result in fluctuations between periods of high bandwidth—when a circuit is provisioned—and low bandwidth—when the packet network is in use. While periods of higher bandwidth are attractive in principle, recent proposals suggest adjusting the topology frequently. The resulting bandwidth fluctuations pose a problem for end-host applications: their active TCP connections must rapidly increase transmission rates to use the available bandwidth and then slow down again to avoid massive queuing. In this paper, we



stringent performance requirements are introduced by today's trend of resource disaggregation in datacenters where fast access to remote resources (e.g., GPUs or memory) is pivotal for the overall system performance [36]. Building systems with strict performance requirements is especially challenging under heavy traffic patterns as they are commonly observed

More Challenges: Buffering?

ABM: Active Buffer Management in Datacenters

Vamsi Addanki^{*}
TU Berlin
Stefan Schmid
TU Berlin
Maria Apostolaki^{*}
Princeton University
Laurent Vanbever
ETH Zurich
Manya Ghobadi
MIT

ABSTRACT

Today's network devices share buffer across queues to avoid drops during transient congestion and absorb bursts. As the buffer-per-bandwidth-unit in datacenter decreases, the need for optimal buffer utilization becomes more pressing. Typical devices use a hierarchical packet admission control scheme: First, a Buffer Management (BM) scheme decides the maximum length per queue at the device level and then an Active Queue Management (AQM) scheme decides which packets will be admitted at the queue level. Unfortunately, the lack of cooperation between the two control schemes leads to: (i) harmful interference across queues, due to the lack of isolation; (ii) increased queuing delay, due to the obliviousness to the per-queue drain time; and (iii) thus unpredictable burst tolerance. To overcome these limitations, we propose ABM, Active Buffer Management which incorporates insights from both BM and AQM. Concretely, ABM accounts for both total buffer occupancy (typically used by BM) and queue drain time (typically used by AQM). We analytically prove that ABM provides isolation, bounded buffer drain time and achieves predictable burst tolerance without sacrificing throughput. We empirically find that ABM improves the 99th percentile FCT for short flows by up to 94% compared to the

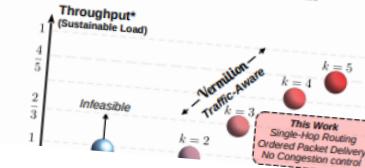
Figure 1: BM and AQM are orthogonal in their goals, and the hierarchical scheme fundamentally limits the burst absorption capabilities of the buffer.

Vermilion: A Traffic-Aware Reconfigurable Optical Interconnect with Formal Throughput Guarantees

Vamsi Addanki TU Berlin Giannis Patronas NVIDIA Paraskevas Bakopoulos NVIDIA	Chen Avin Ben-Gurion University of the Negev Dimitris Syrivelis NVIDIA Ilias Marinou NVIDIA	Goran Dario Knabe TU Berlin Nikos Terzenidis NVIDIA Stefan Schmid TU Berlin
---	--	--

ABSTRACT

The increasing gap between datacenter traffic volume and the capacity of electrical switches has driven the development of reconfigurable network designs utilizing optical circuit switching. Recent advancements, particularly those featuring periodic fixed-duration reconfigurations, have achieved practical end-to-end delays of just a few microseconds. However, current designs rely on multi-hop



Case Study: PowerTCP

Existing congestion control algorithms based on either

- State (“**voltage**”) like BDP, queue length, loss, e.g.:
 - DCTCP: uses ECN/loss
 - Swift: RTT
 - HPCC: inflight packets
- Gradient (“**current**”) like reaction to queue length change
 - Timely: RTT-gradient based

Case Study: PowerTCP

Existing congestion control algorithms based on either

- State (“**voltage**”) like BDP, queue length, loss, e.g.:
 - DCTCP: uses ECN/loss
 - Swift: RTT
 - HPCC: inflight packets
 - Gradient (“**current**”) like reaction to queue length change
 - Timely: RTT-gradient based
- 
- ☺ Can achieve near-zero queue equilibrium
☹ Slow reaction

Case Study: PowerTCP

Existing congestion control algorithms based on either

- State (“**voltage**”) like BDP, queue length, loss, e.g.:

- DCTCP: uses ECN/loss

- Swift: RTT

- HPCC: inflight packets

- Gradient (“**current**”) like reaction to queue length change

- Timely: RTT-gradient based

} ☺ Fast reaction
☹ No equilibrium

Case Study: PowerTCP

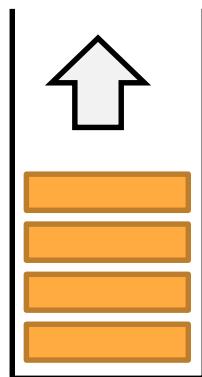
Existing congestion control algorithms based on either

- State (“**voltage**”) like BDP, queue length, loss, e.g.:
 - DCTCP: uses ECN/loss
 - Swift: RTT
 - HPCC: inflight packets
- Gradient (“**current**”) like reaction to queue length change
 - Timely: RTT-gradient based

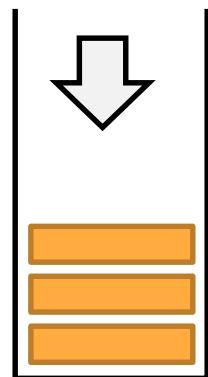
Limitation: using only one of the two may miss useful information for fine-grained adaptions!

Limitation of SOTA

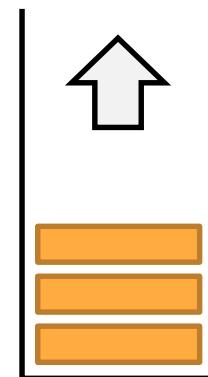
→ Consider a queue which may be in three different states:



1
growing



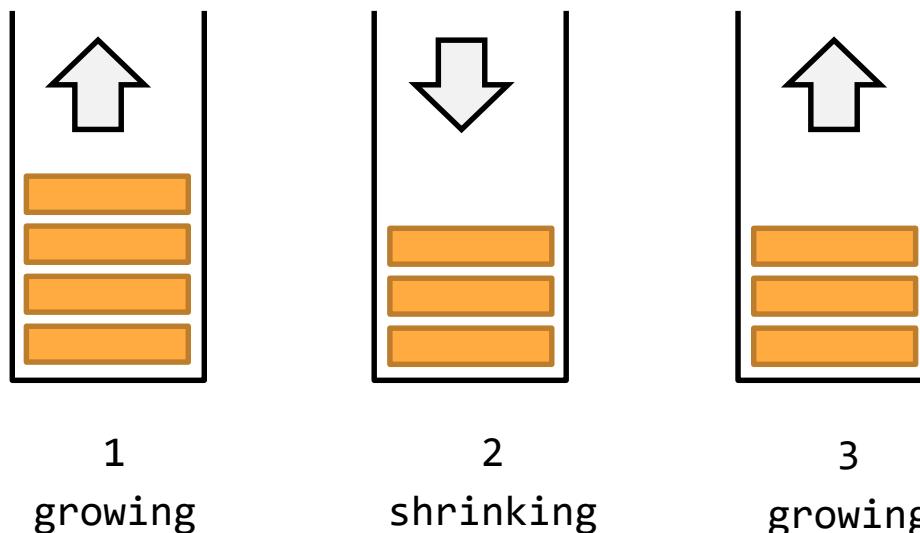
2
shrinking



3
growing

Limitation of SOTA

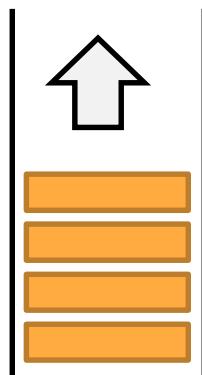
→ Consider a queue which may be in three different states:



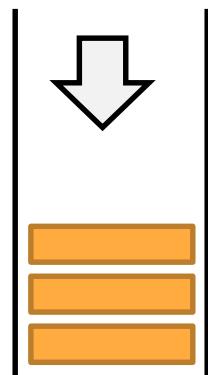
2 and 3: impossible to
distinguish for voltage-based CCA

Limitation of SOTA

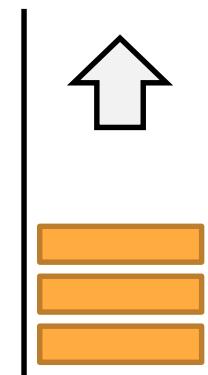
→ Consider a queue which may be in three different states:



1
growing



2
shrinking

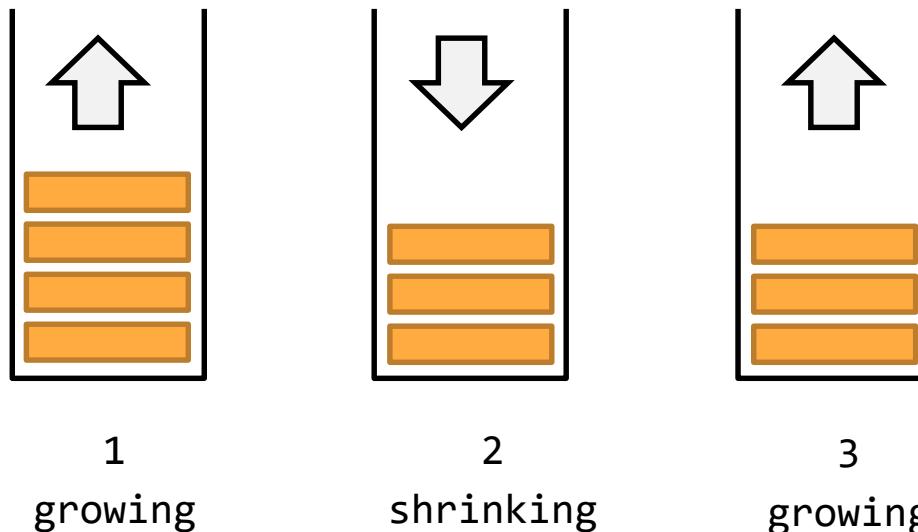


3
growing

1 and 3: impossible to
distinguish for current-based CC

Limitation of SOTA

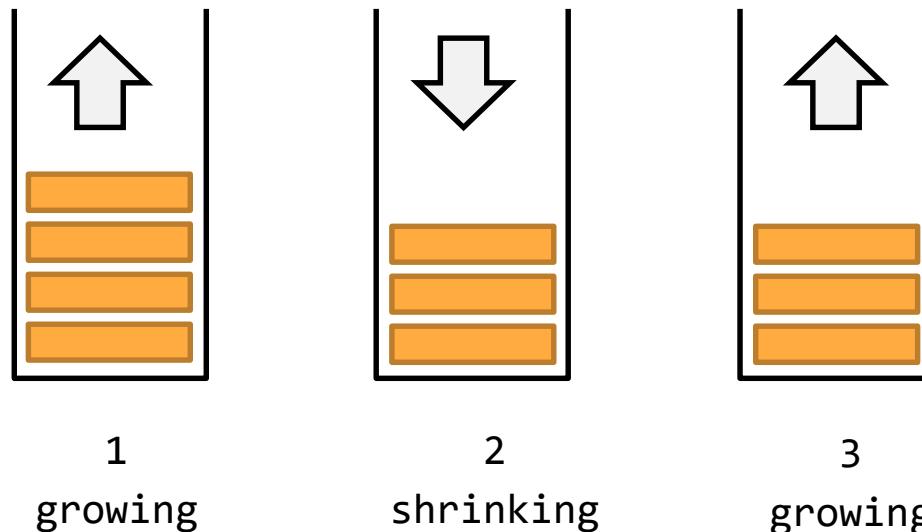
→ Consider a queue which may be in three different states:



We need both: Power ($Voltage \times Current$)

Limitation of SOTA

→ Consider a queue which may be in three different states:



We need both: Power (Voltage x Current)

Inspired: **POWERTCP**

More benefits of optical & reconfigurable switching

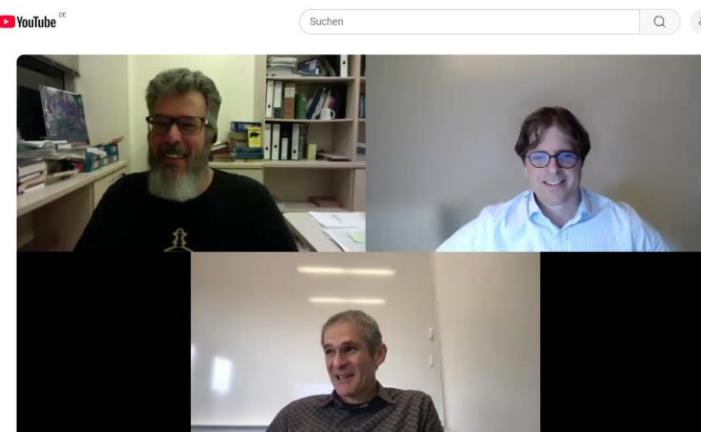
So far: focus on throughput performance.

Benefit 1:

Evolving Datacenters

- Reconfigurable datacenter networks naturally support *heterogeneous* network elements
- And therefore also *incremental* hardware upgrades

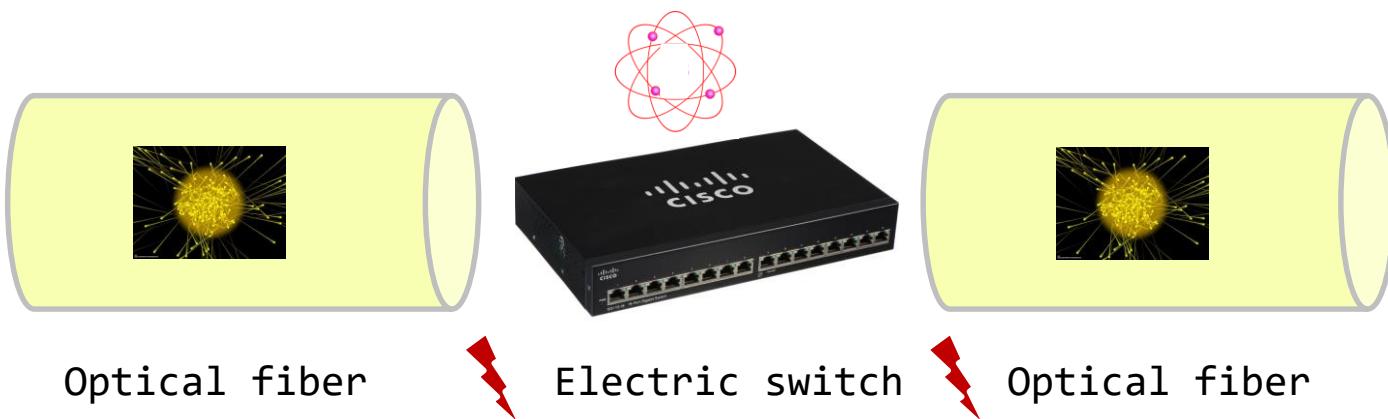
See interview with Amin Vahdat, Google in CACM:
<https://www.youtube.com/watch?v=Ixv1gu8ETA>



Benefit 2:

Energy and Latency

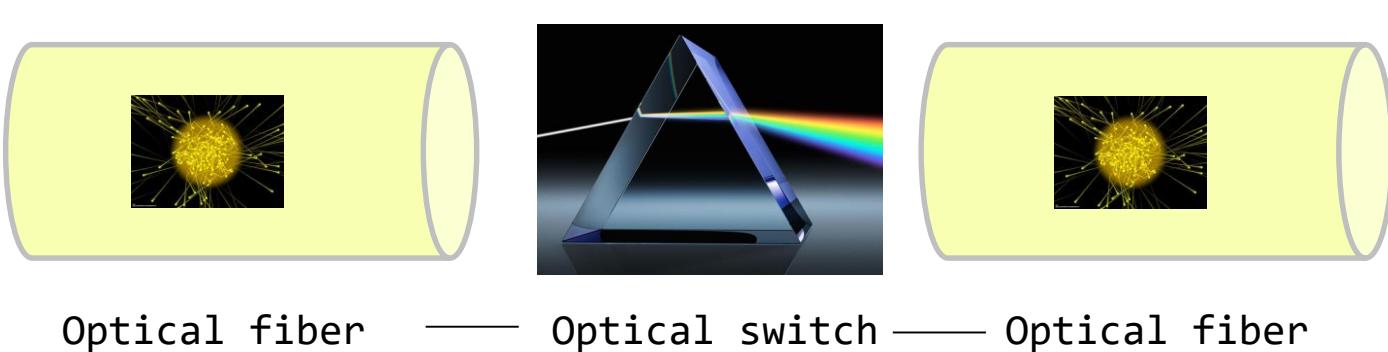
- No need to *convert* photons in fiber to electrons in switch (and back)
- Can save *energy* and reduce *latency* (in addition to enabling almost unlimited throughput)



Benefit 2:

Energy and Latency

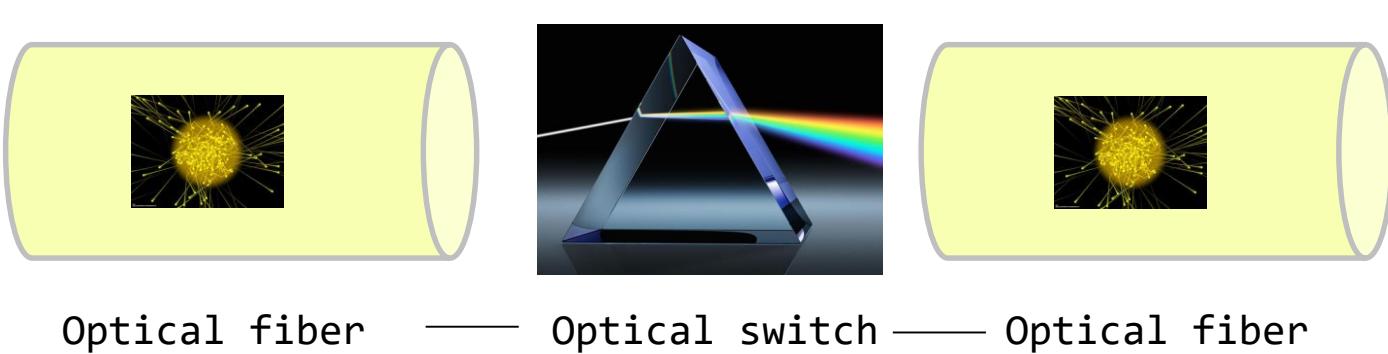
- No need to *convert* photons in fiber to electrons in switch (and back)
- Can save *energy* and reduce *latency* (in addition to enabling almost unlimited throughput)



Benefit 2:

Energy and Latency

- No need to *convert* photons in fiber to electrons in switch (and back)
- Can save *energy* and reduce *latency* (in addition to enabling almost unlimited throughput)

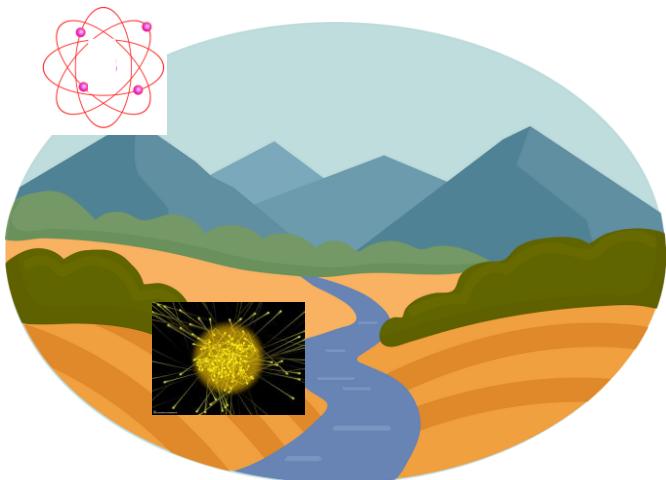


- Interesting for emerging *distributed datacenters*!

Benefit 3:

Resilience

Floodings in South Germany destroyed much electrical network infrastructure



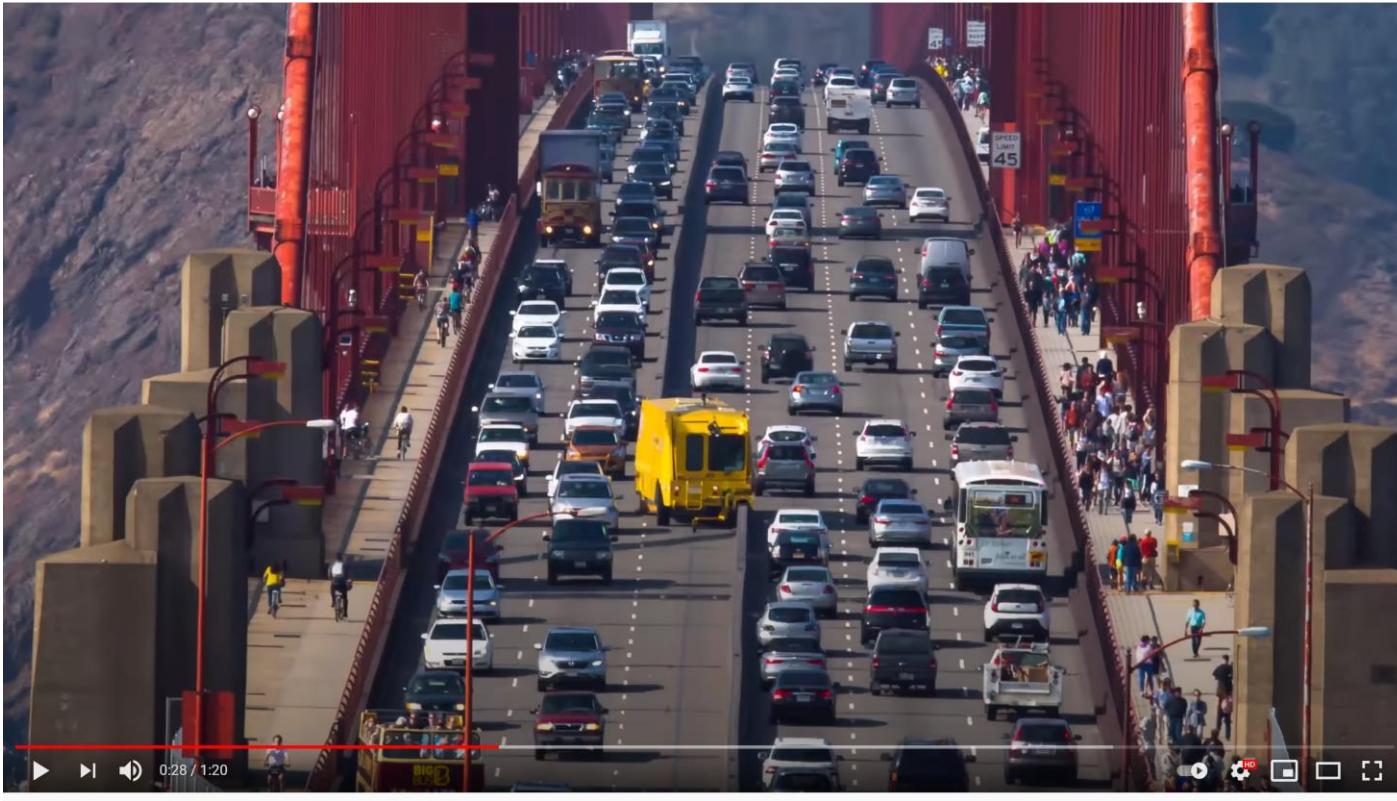
Solution: deploy optical infrastructure (in valleys) and electrical *on hills* where safe?

Conclusion

- Opportunity: *structure* in demand and *reconfigurable* networks
- So far: tip of the iceberg
- Many challenges
 - Optimal design depends on traffic pattern
 - How to *measure/predict* traffic?
 - Impact on other *Layers*?
 - *Scalable control* plane
 - *Application-specific* self-adjusting networks?
- Many more *opportunities* for optical networks



Thank you! Questions?



Slides
available
here:



Online Video Course

Invitation to
Self-Adjusting Networks
A short video course

Before demand:

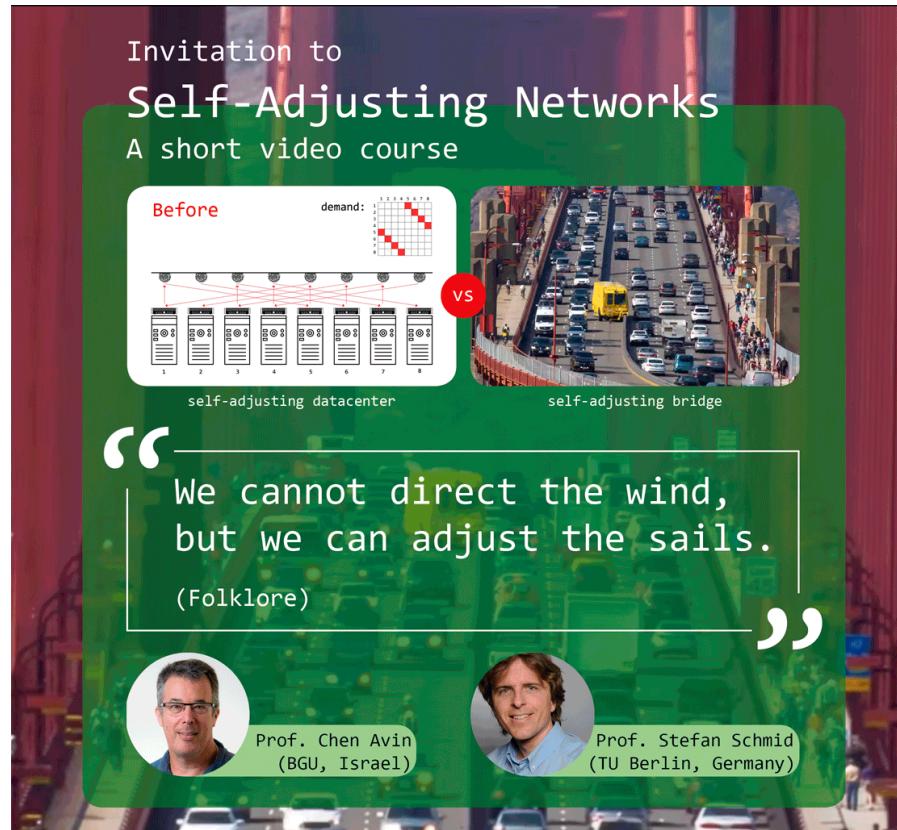
1	2	3	4	5	6	7	8
1	✓	✓	✓	✓	✓	✓	✓
2	✓	✓	✓	✓	✓	✓	✓
3	✓	✓	✓	✓	✓	✓	✓
4	✓	✓	✓	✓	✓	✓	✓
5	✓	✓	✓	✓	✓	✓	✓
6	✓	✓	✓	✓	✓	✓	✓
7	✓	✓	✓	✓	✓	✓	✓
8	✓	✓	✓	✓	✓	✓	✓

self-adjusting datacenter self-adjusting bridge

“ We cannot direct the wind,
but we can adjust the sails.
(Folklore) ”

Prof. Chen Avin
(BGU, Israel)

Prof. Stefan Schmid
(TU Berlin, Germany)



erc

<https://self-adjusting.net/course>



YouTube Interview & CACM

Check out our **YouTube interviews**
on Reconfigurable Datacenter Networks:



[Revolutionizing Datacenter Networks via Reconfigurable Topologies](#)

Chen Avin and Stefan Schmid.

Communications of the ACM (CACM), 2025.

Watch here: <https://www.youtube.com/@self-adjusting-networks-course>



Websites

SELF-ADJUSTING NETWORKS
RESEARCH ON SELF-ADJUSTING DEMAND-AWARE NETWORKS

Project Overview Team Publications Contact Us

AdjustNet

Breaking new ground with demand-aware self-adjusting networks

Our Vision:
Flexible and Demand-Aware Topologies

This site provides an overview of our ongoing research on the foundations of self-adjusting networks.

MARCH 17, 2020

WEBSITE LAUNCHED!

Download Slides

<http://self-adjusting.net/>
Project website



TRACE COLLECTION
WAN AND DC NETWORK TRACES

Publication Team Download Traces Contact Us

The following table lists the traces used in the publication: *On the Complexity of Traffic Traces and Implications*
To reference this website, please use: bibtex

File Name	Source Information	Type	Lines	Size	Download
exact_BoxLib_MultiGrid_C_Large_1024.csv	High Performance Computing Traces	Traces	17,947,800	151.3 MB	Download
exact_BoxLib_CNS_NoSpec_Large_1024.csv	High Performance Computing Traces	Traces	11,108,068	9.3 MB	Download
cesar_Nekbone_1024.csv	High Performance Computing Traces	Traces	21,745,229	184.0 MB	Download

<https://trace-collection.net/>
Trace collection website



June'25 CACM Article

Revolutionizing Datacenter Networks via Reconfigurable Topologies

CHEN AVIN, is a Professor at Ben-Gurion University of the Negev, Beersheva, Israel

STEFAN SCHMID, is a Professor at TU Berlin, Berlin, Germany

With the popularity of cloud computing and data-intensive applications such as machine learning, datacenter networks have become a critical infrastructure for our digital society. Given the explosive growth of datacenter traffic and the slowdown of Moore's law, significant efforts have been made to improve datacenter network performance over the last decade. A particularly innovative solution is reconfigurable datacenter networks (RDCNs): datacenter networks whose topologies dynamically change over time, in either a demand-oblivious or a demand-aware manner. Such dynamic topologies are enabled by recent optical switching technologies and stand in stark contrast to state-of-the-art datacenter network topologies, which are fixed and oblivious to the actual traffic demand. In particular, reconfigurable demand-aware and "self-adjusting" datacenter networks are motivated empirically by the significant spatial and temporal structures observed in datacenter communication traffic. This paper presents an overview of reconfigurable datacenter networks. In particular, we discuss the motivation for such reconfigurable architectures, review the technological enablers, and present a taxonomy that classifies the design space into two dimensions: static vs. dynamic and demand-oblivious vs. demand-aware. We further present a formal model and discuss related research challenges. Our article comes with complementary video interviews in which three leading experts, Manya Ghobadi, Amin Vahdat, and George Papen, share with us their perspectives on reconfigurable datacenter networks.

KEY INSIGHTS

- Datacenter networks have become a critical infrastructure for our digital society, serving explosively growing communication traffic.
- Reconfigurable datacenter networks (RDCNs) which can adapt their topology dynamically, based on innovative optical switching technologies, bear the potential to improve datacenter network performance, and to simplify datacenter planning and operations.
- Demand-aware dynamic topologies are particularly interesting because of the significant spatial and temporal structures observed in real-world traffic, e.g., related to distributed machine learning.
- The study of RDCNs and self-adjusting networks raises many novel technological and research challenges related to their design, control, and performance.

References (1)

[Revolutionizing Datacenter Networks via Reconfigurable Topologies](#)

Chen Avin and Stefan Schmid.
Communications of the ACM (**CACM**), 2025.

[Cerberus: The Power of Choices in Datacenter Topology Design \(A Throughput Perspective\)](#)

Chen Griner, Johannes Zerwas, Andreas Blenk, Manya Ghobadi, Stefan Schmid, and Chen Avin.
ACM **SIGMETRICS** and ACM Performance Evaluation Review (**PER**), Mumbai, India, June 2022.

[Mars: Near-Optimal Throughput with Shallow Buffers in Reconfigurable Datacenter Networks](#)

Vamsi Addanki, Chen Avin, and Stefan Schmid.
ACM **SIGMETRICS** and ACM Performance Evaluation Review (**PER**), Orlando, Florida, USA, June 2023.

[Duo: A High-Throughput Reconfigurable Datacenter Network Using Local Routing and Control](#)

Johannes Zerwas, Csaba Györgyi, Andreas Blenk, Stefan Schmid, and Chen Avin.
ACM **SIGMETRICS** and ACM Performance Evaluation Review (**PER**), Orlando, Florida, USA, June 2023.

[On the Complexity of Traffic Traces and Implications](#)

Chen Avin, Manya Ghobadi, Chen Griner, and Stefan Schmid.
ACM **SIGMETRICS** and ACM Performance Evaluation Review (**PER**), Boston, Massachusetts, USA, June 2020.

[Toward Demand-Aware Networking: A Theory for Self-Adjusting Networks \(Editorial\)](#)

Chen Avin and Stefan Schmid.
ACM SIGCOMM Computer Communication Review (**CCR**), October 2018.

[Credence: Augmenting Datacenter Switch Buffer Sharing with ML Predictions](#)

Vamsi Addanki, Maciej Pacut, and Stefan Schmid.
21st USENIX Symposium on Networked Systems Design and Implementation (**NSDI**), Santa Clara, California, USA, April 2024.

[PowerTCP: Pushing the Performance Limits of Datacenter Networks](#)

Vamsi Addanki, Oliver Michel, and Stefan Schmid.
19th USENIX Symposium on Networked Systems Design and Implementation (**NSDI**), Renton, Washington, USA, April 2022.

[TCP's Third Eye: Leveraging eBPF for Telemetry-Powered Congestion Control](#)

Jörn-Thorben Hinz, Vamsi Addanki, Csaba Györgyi, Theo Jepsen, and Stefan Schmid.
SIGCOMM Workshop on eBPF and Kernel Extensions (**eBPF**), Columbia University, New York City, New York, USA, September 2023.

References (2)

[ABM: Active Buffer Management in Datacenters](#)

Vamsi Addanki, Maria Apostolaki, Manya Ghobadi, Stefan Schmid, and Laurent Vanbever.
ACM SIGCOMM, Amsterdam, Netherlands, August 2022.

[ExRec: Experimental Framework for Reconfigurable Networks Based on Off-the-Shelf Hardware](#)

Johannes Zerwas, Chen Avin, Stefan Schmid, and Andreas Blenk.
16th ACM/IEEE Symposium on Architectures for Networking and Communications Systems (ANCS), Virtual Conference, December 2021.

[Demand-Aware Network Design with Minimal Congestion and Route Lengths](#)

Chen Avin, Kaushik Mondal, and Stefan Schmid.
IEEE/ACM Transactions on Networking (TON), 2022.

[A Survey of Reconfigurable Optical Networks](#)

Matthew Nance Hall, Klaus-Tycho Foerster, Stefan Schmid, and Ramakrishnan Durairajan.
Optical Switching and Networking (OSN), Elsevier, 2021.

[SplayNet: Towards Locally Self-Adjusting Networks](#)

Stefan Schmid, Chen Avin, Christian Scheideler, Michael Borokhovich, Bernhard Haeupler, and Zvi Lotker.
IEEE/ACM Transactions on Networking (TON), Volume 24, Issue 3, 2016.

.

Bonus Material



Hogwarts Stair

Question:

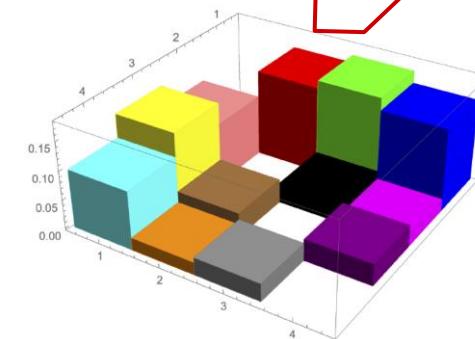
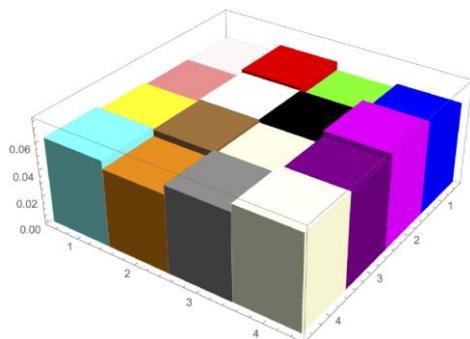
How to Quantify
such “Structure”
in the Demand?

Intuition

Which demand has more structure?

→ Traffic matrices of two different distributed
ML applications

→ GPU-to-GPU



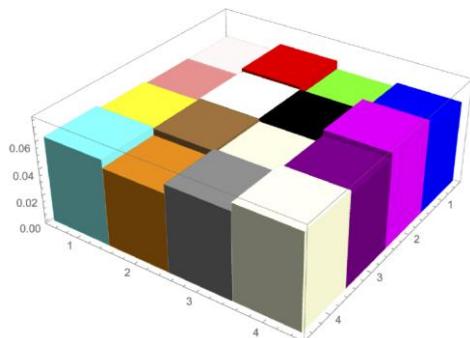
Color = communication pair

Intuition

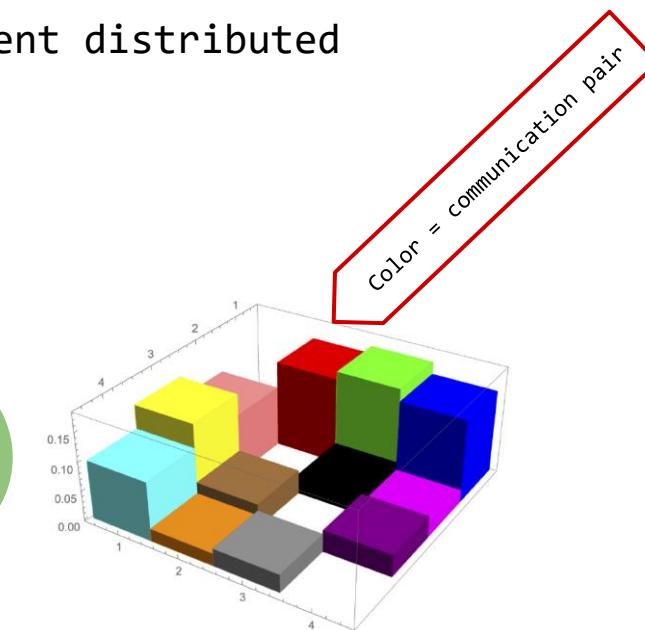
Which demand has more structure?

→ Traffic matrices of two different distributed
ML applications

→ GPU-to-GPU



More uniform



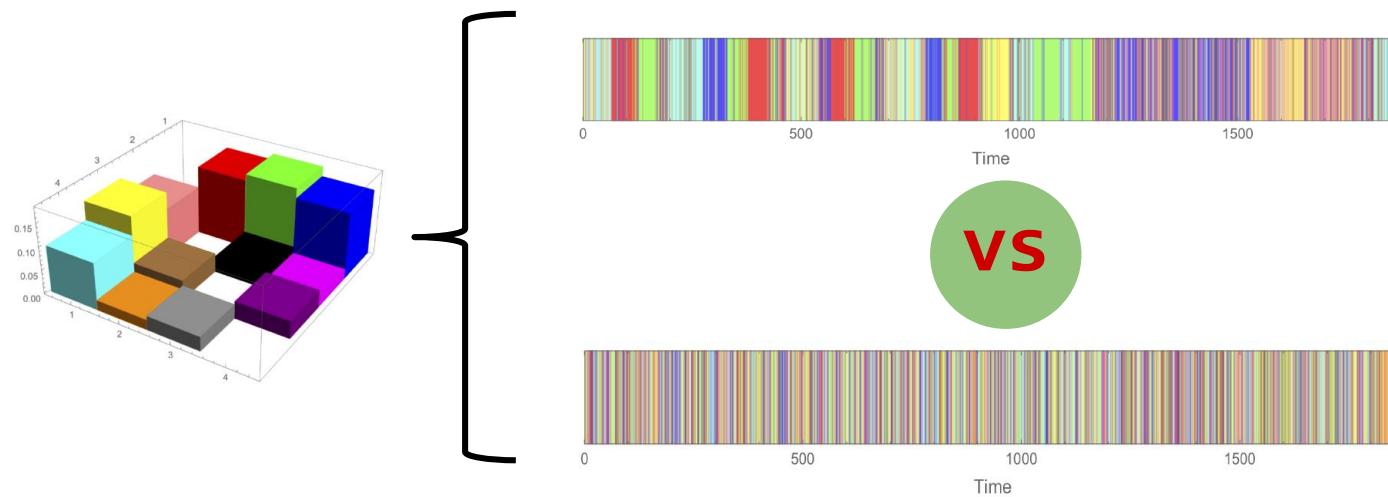
More structure

Color = communication pair

Intuition

Spatial vs temporal structure

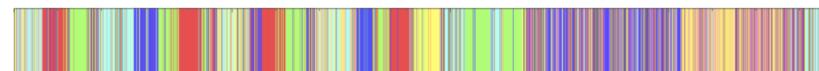
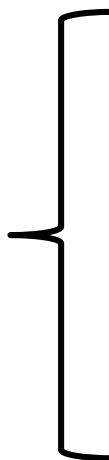
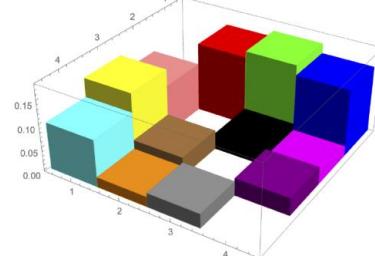
- Two different ways to generate same traffic matrix:
 - Same non-temporal structure
- Which one has more structure?



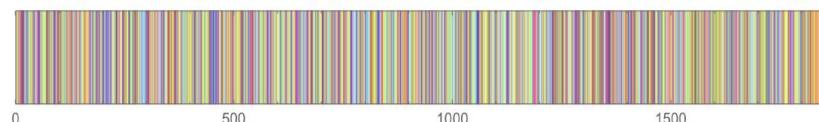
Intuition

Spatial vs temporal structure

- Two different ways to generate same traffic matrix:
 - Same non-temporal structure
- Which one has more structure?



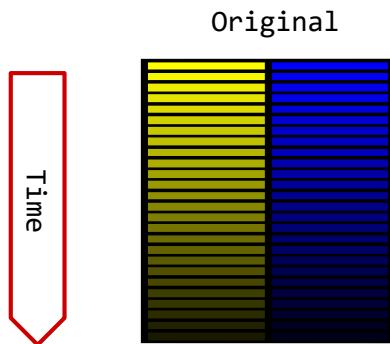
VS



Systematically?

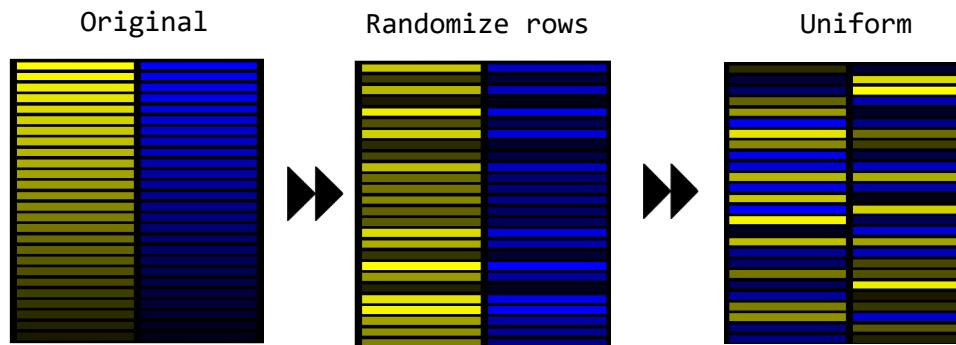
Trace Complexity

Information-Theoretic Approach
“Shuffle&Compress”



Trace Complexity

Information-Theoretic Approach
“Shuffle&Compress”

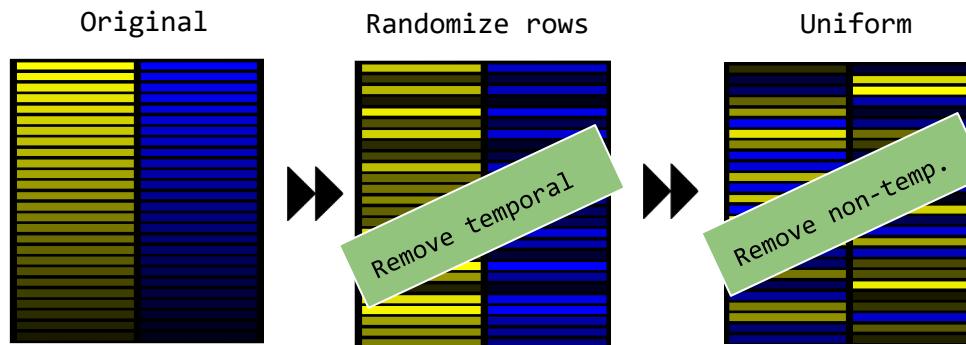


Increasing complexity (systematically randomized)

More structure (compresses better)

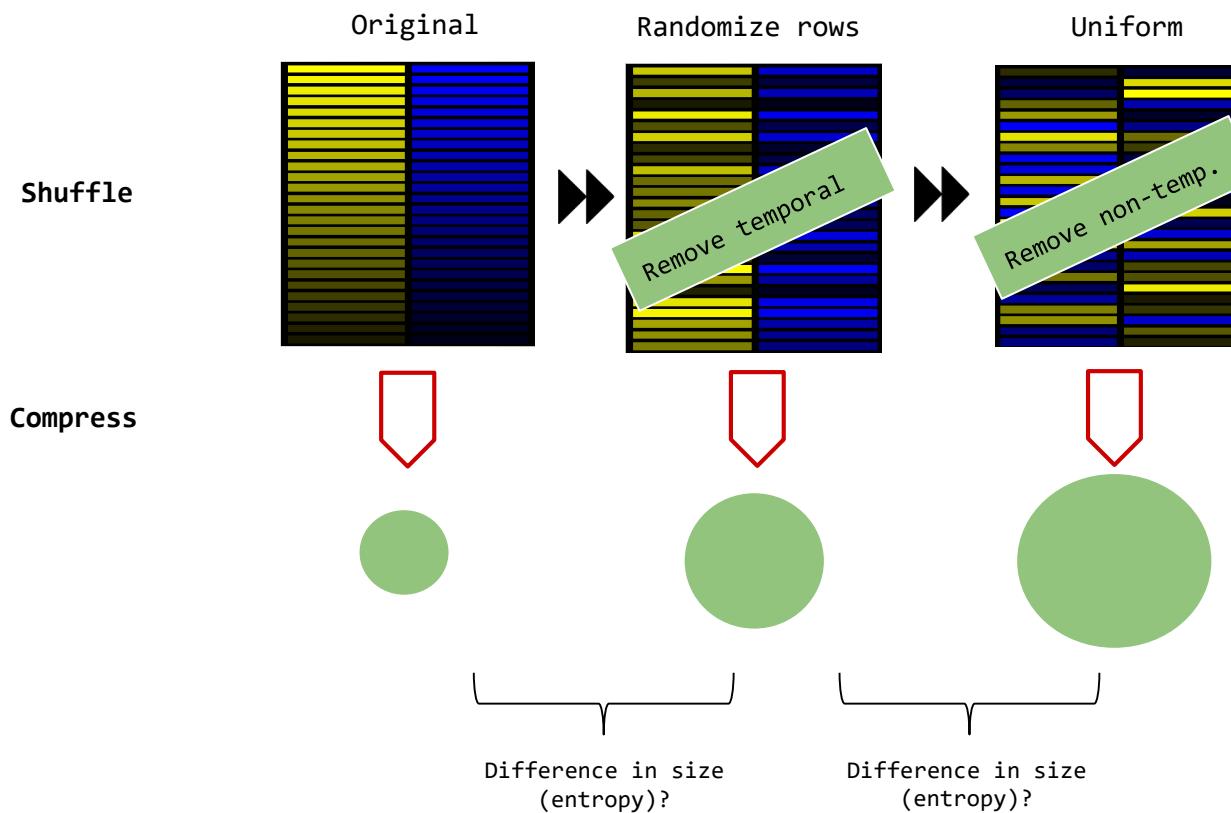
Trace Complexity

Information-Theoretic Approach
“Shuffle&Compress”



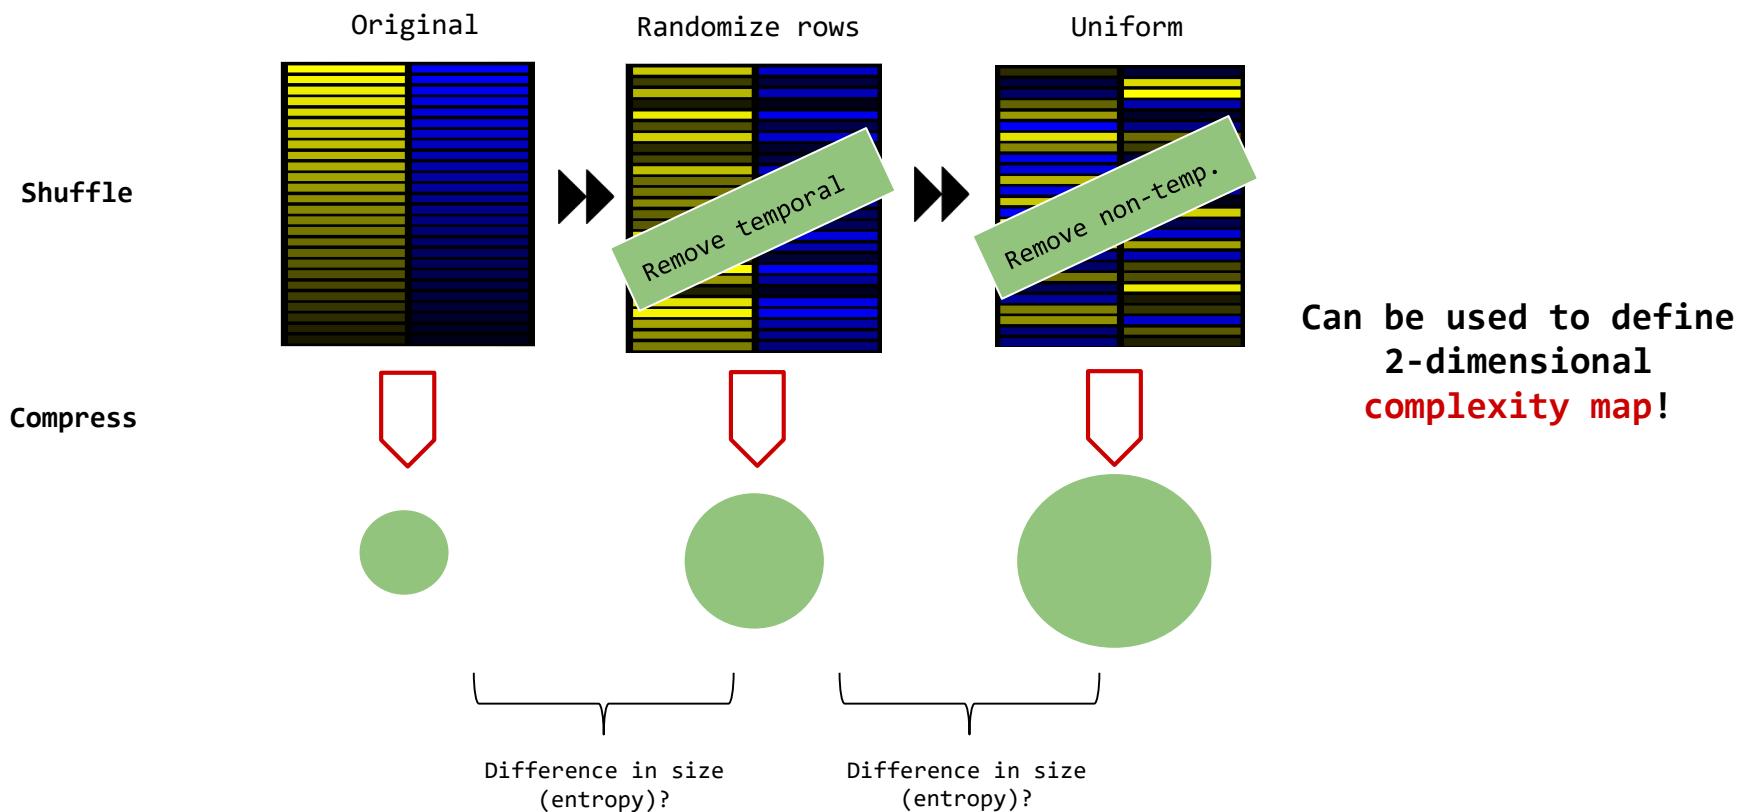
Trace Complexity

Information-Theoretic Approach
“Shuffle&Compress”



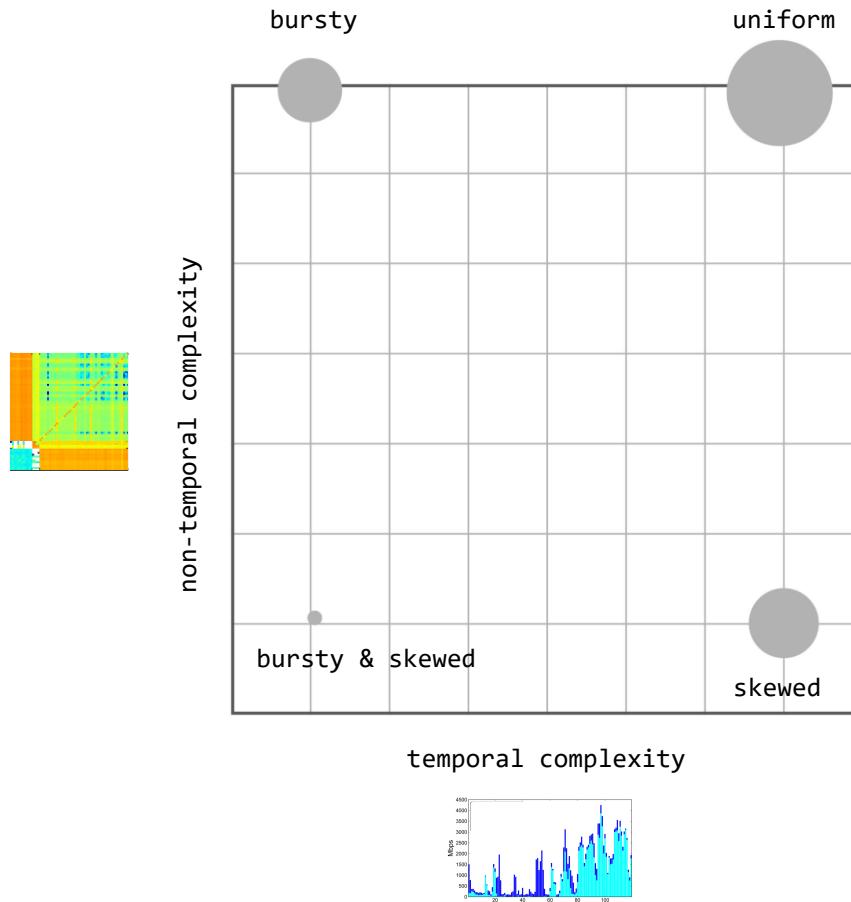
Trace Complexity

Information-Theoretic Approach
“Shuffle&Compress”



Our Methodology

Complexity Map

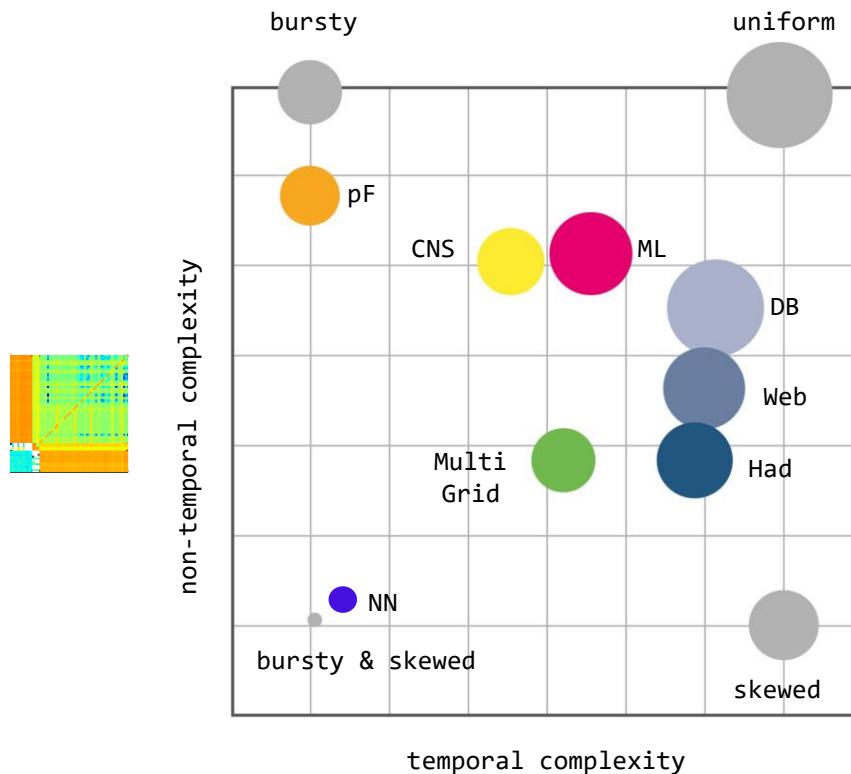


No structure

Our approach: iterative
randomization and
compression of trace to
identify dimensions of
structure.

Our Methodology

Complexity Map



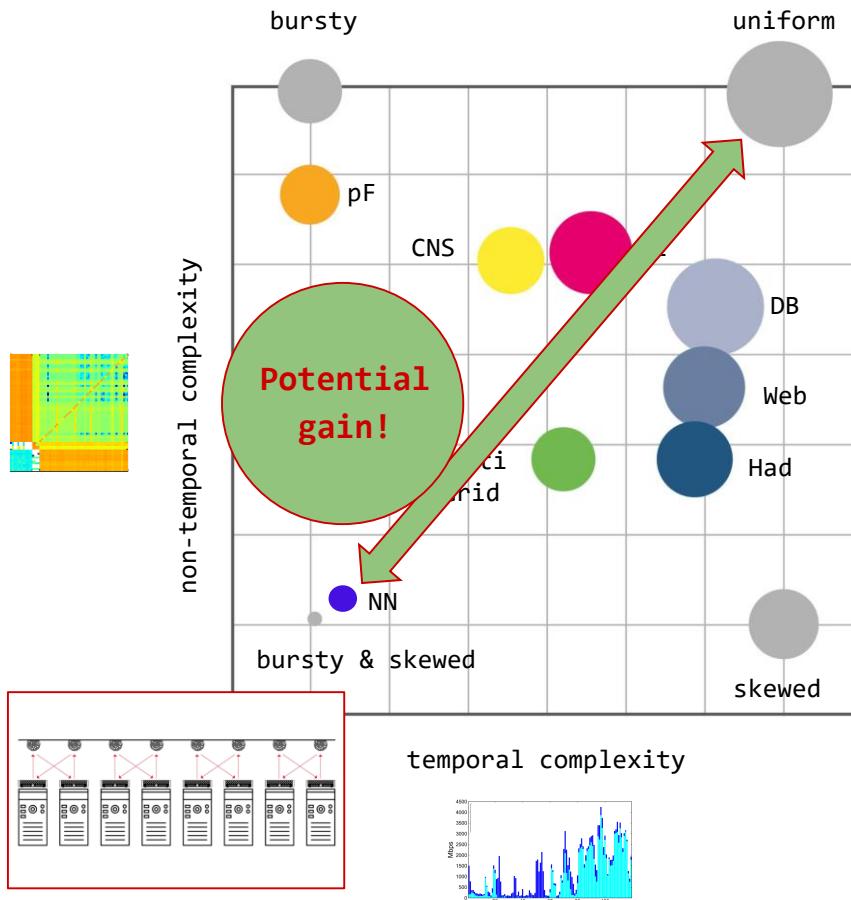
No structure

Our approach: iterative
randomization and
compression of trace to
identify dimensions of
structure.

Different
structures!

Our Methodology

Complexity Map



Our approach: iterative randomization and compression of trace to identify dimensions of structure.

Different structures!

Further Reading

ACM SIGMETRICS 2020

On the Complexity of Traffic Traces and Implications

CHEN AVIN, School of Electrical and Computer Engineering, Ben Gurion University of the Negev, Israel

MANYA GHOBADI, Computer Science and Artificial Intelligence Laboratory, MIT, USA

CHEN GRINER, School of Electrical and Computer Engineering, Ben Gurion University of the Negev, Israel

STEFAN SCHMID, Faculty of Computer Science, University of Vienna, Austria

This paper presents a systematic approach to identify and quantify the types of structures featured by packet traces in communication networks. Our approach leverages an information-theoretic methodology, based on iterative randomization and compression of the packet trace, which allows us to systematically remove and measure dimensions of structure in the trace. In particular, we introduce the notion of *trace complexity* which approximates the entropy rate of a packet trace. Considering several real-world traces, we show that trace complexity can provide unique insights into the characteristics of various applications. Based on our approach, we also propose a traffic generator model able to produce a synthetic trace that matches the complexity levels of its corresponding real-world trace. Using a case study in the context of datacenters, we show that insights into the structure of packet traces can lead to improved demand-aware network designs: datacenter topologies that are optimized for specific traffic patterns.

CCS Concepts: • Networks → Network performance evaluation; Network algorithms; Data center networks; • Mathematics of computing → Information theory;

Additional Key Words and Phrases: trace complexity, self-adjusting networks, entropy rate, compress, complexity map, data centers

ACM Reference Format:

Chen Avin, Manya Ghobadi, Chen Griner, and Stefan Schmid. 2020. On the Complexity of Traffic Traces and Implications. *Proc. ACM Meas. Anal. Comput. Syst.* 4, 1, Article 20 (March 2020), 29 pages. <https://doi.org/10.1145/3379486>

1 INTRODUCTION

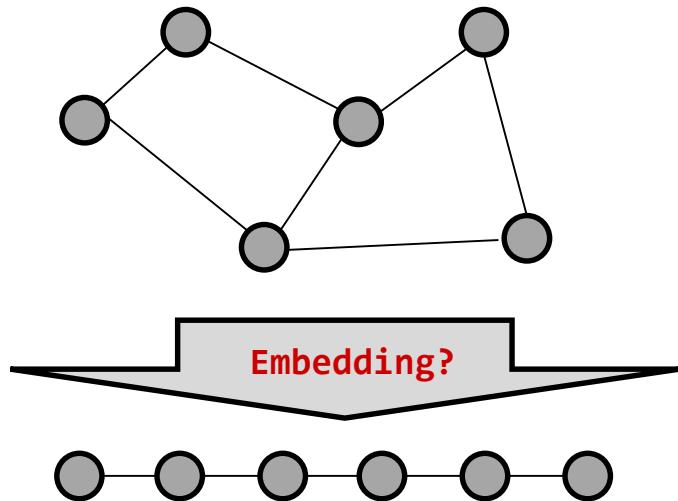
Packet traces collected from networking applications, such as datacenter traffic, have been shown to feature much *structure*: datacenter traffic matrices are sparse and skewed [16, 39], exhibit

20

Related Problem: Remember Bernardetta's Talk

Virtual Network Embedding Problem (VNEP)

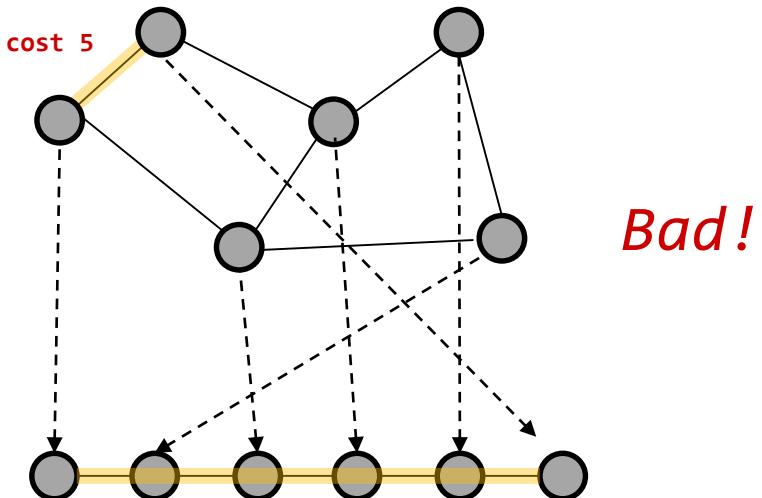
Example $\Delta=2$: A Minimum Linear
Arrangement (**MLA**) Problem
→ Minimizes sum of virtual
edges



Related Problem: Remember Bernardetta's Talk

Virtual Network Embedding Problem (VNEP)

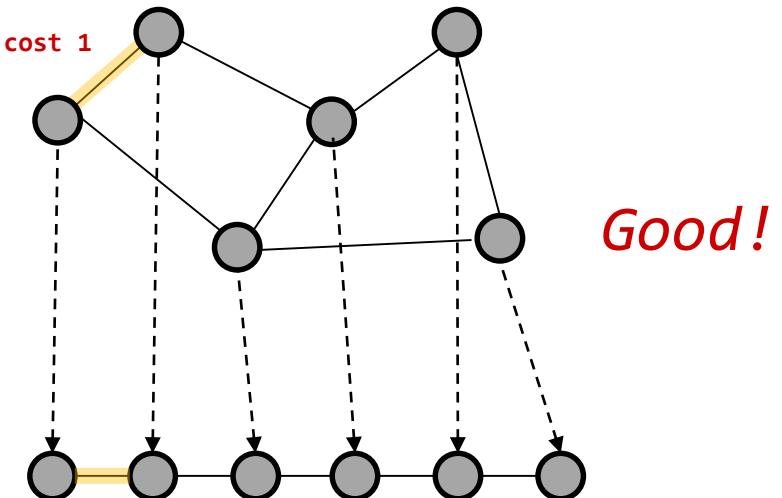
Example $\Delta=2$: A Minimum Linear Arrangement (MLA) Problem
→ Minimizes sum of virtual edges



Related Problem: Remember Bernardetta's Talk

Virtual Network Embedding Problem (VNEP)

Example $\Delta=2$: A Minimum Linear Arrangement (MLA) Problem
→ Minimizes sum of virtual edges



Related Problem: Remember Bernardetta's Talk

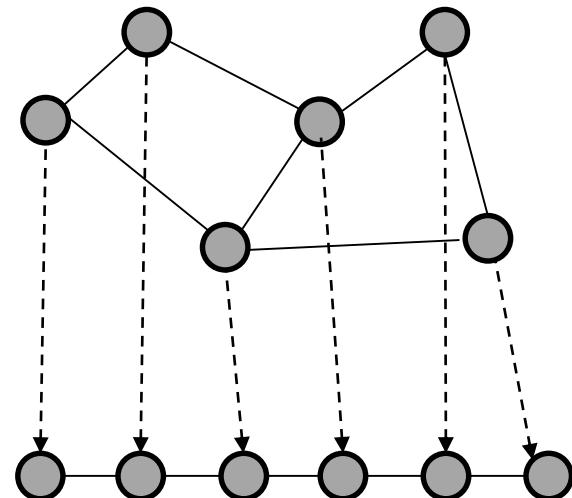
Virtual Network Embedding Problem (VNEP)

Example $\Delta=2$: A Minimum Linear Arrangement (MLA) Problem

→ Minimizes sum of virtual edges

MLA is **NP-hard**

→ ... and so is our problem!



Related Problem: Remember Bernardetta's Talk

Virtual Network Embedding Problem (VNEP)

Example $\Delta=2$: A Minimum Linear Arrangement (MLA) Problem

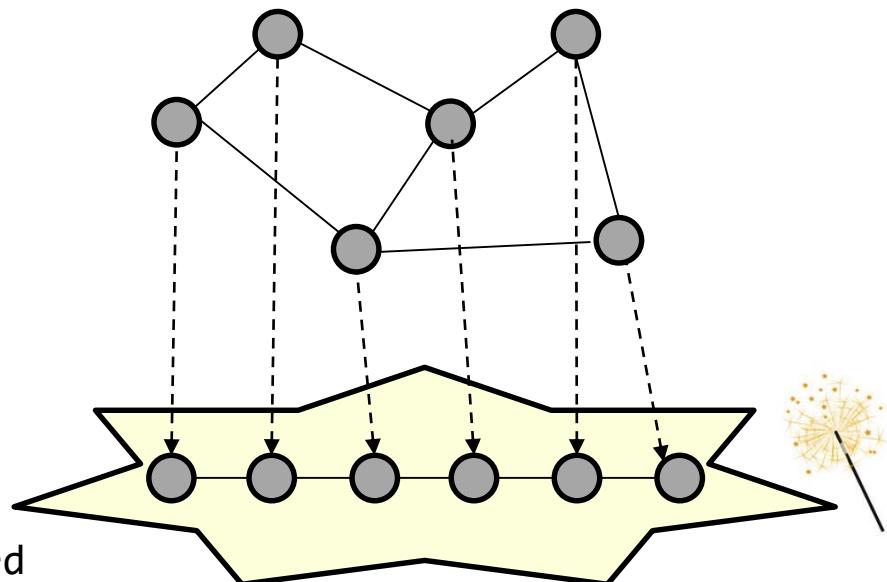
→ Minimizes sum of virtual edges

MLA is **NP-hard**

→ ... and so is our problem!

But what about $\Delta > 2$?

→ Embedding problem still hard
→ But we have a new **degree of freedom!**



Related Problem: Remember Bernardetta's Talk

Virtual Network Embedding Problem (VNEP)

Example $\Delta=2$: A Minimum Linear Arrangement (MLA) Problem

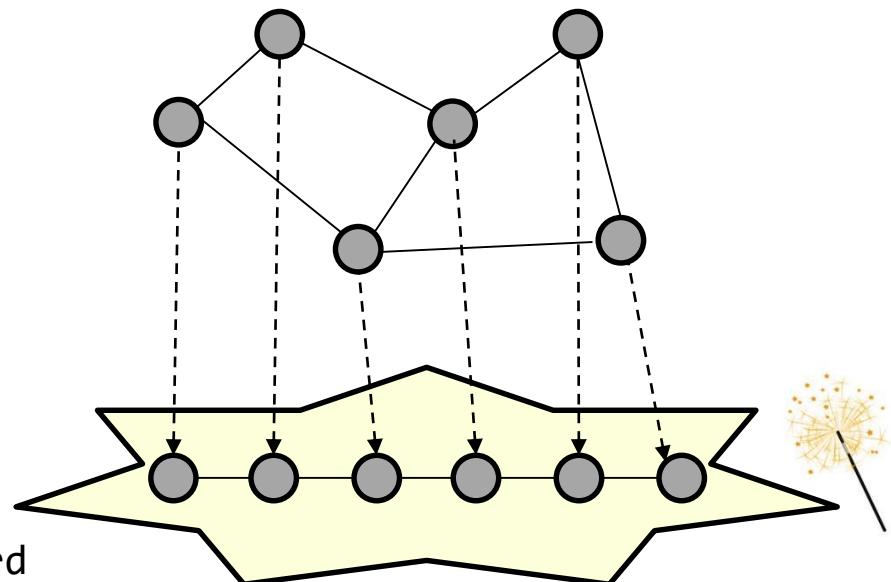
→ Minimizes sum of virtual edges

MLA is **NP-hard**

→ ... and so is our problem!

But what about $\Delta > 2$?

→ Embedding problem still hard
→ But we have a new **degree of freedom!**



Simplifies problem?!

Another Related Problem

Low Distortion Spanners

- ...> Classic problem: find *sparse*, *distance-preserving* (low-distortion) spanner of a graph
- ...> But:
 - ...> Spanners aim at low distortion *among all pairs*; in our case, we are only interested in the **local distortion**, 1-hop communication neighbors
 - ...> We allow *auxiliary edges* (not a subgraph): similar to geometric spanners
 - ...> We require *constant degree*

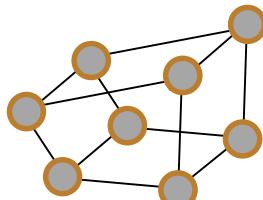
From Spanners to DANs

An Algorithm

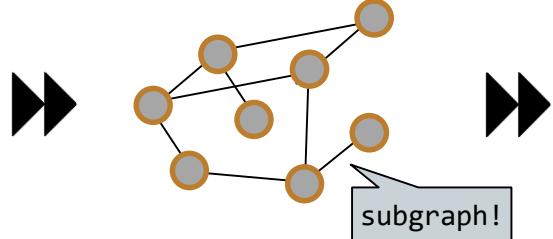
→ Yet, can leverage the connection to spanners sometimes!

Theorem: If demand matrix is regular and uniform, and if we can find a constant distortion, linear sized (i.e., constant, sparse) spanner for this request graph: then we can design a constant degree DAN providing an optimal expected route length (i.e., $O(H(X/Y)+H(Y/X))$).

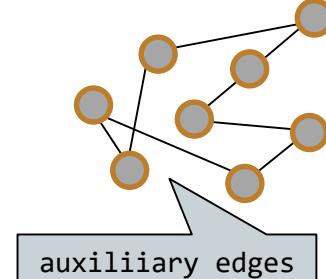
r-regular and
uniform demand:



Sparse, *irregular*
(*constant*) spanner:



Constant degree
optimal DAN (ERL
at most *Log r*):



From Spanners to DANs

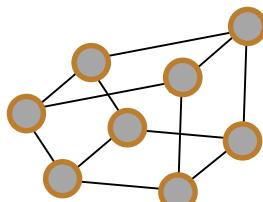
An Algorithm

→ Yet, can leverage the connection to spanners sometimes!

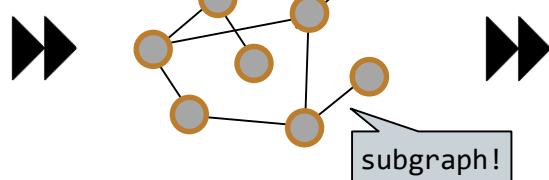
Theorem: If demand matrix is regular and uniform, and if we can find a constant distortion, linear sized (i.e., constant, sparse) spanner for this request graph: then we can design a constant degree DAN providing an optimal expected route length (i.e., $O(H(X/Y)+H(Y/X))$).

Our degree reduction trick again!

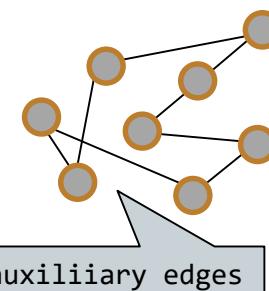
r-regular and
uniform demand:



Sparse, irregular
(constant) spanner:

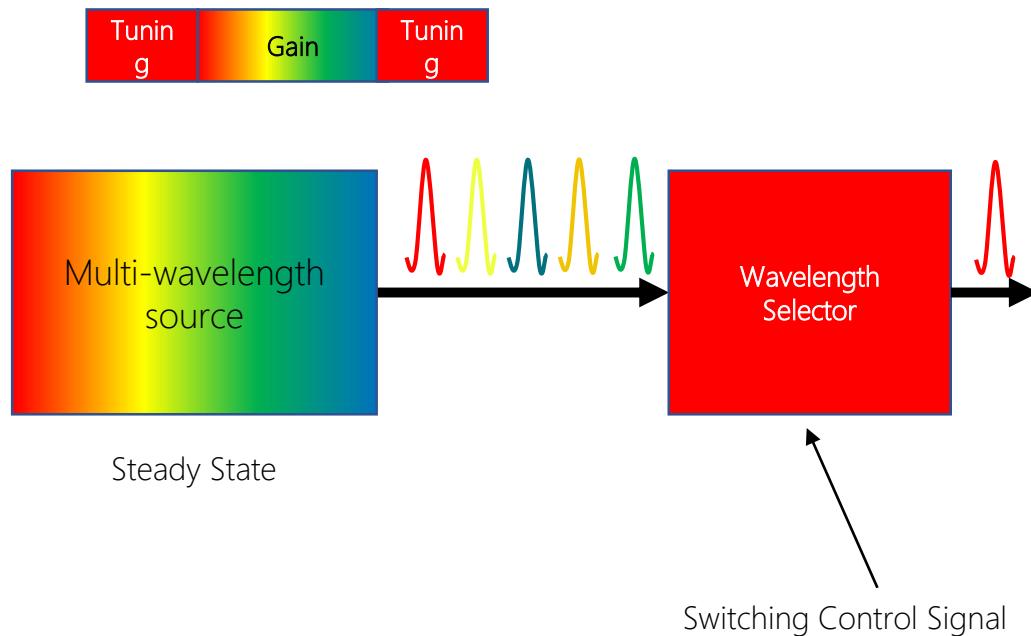


Constant degree
optimal DAN (ERL
at most $\log r$):

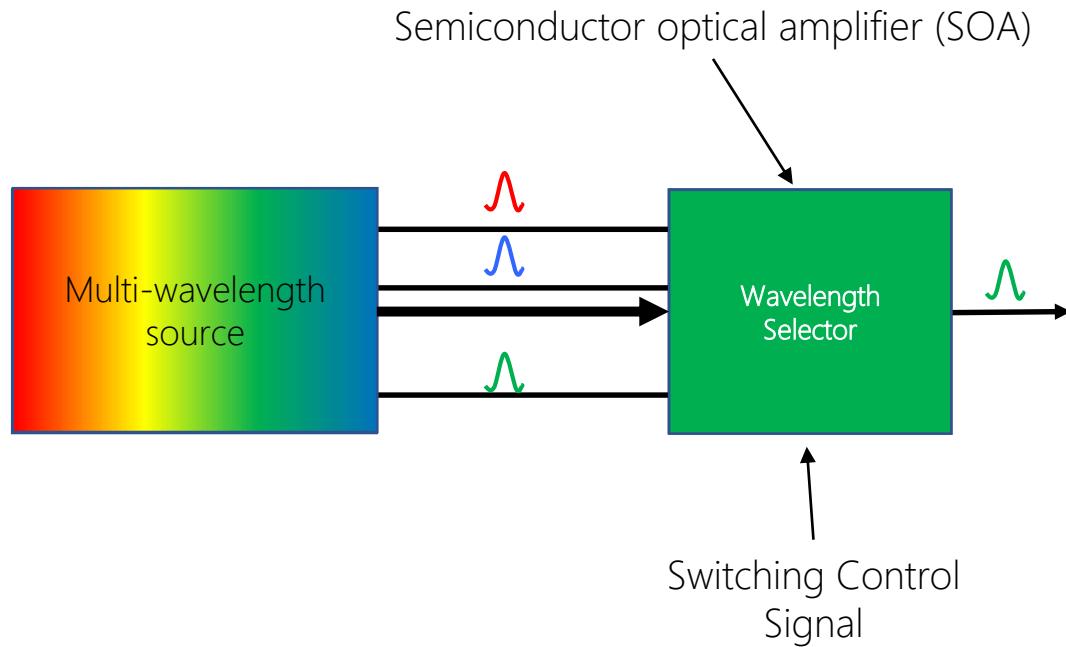


Idea

Disaggregated Laser



Example Design



Sirius also implemented other designs
(details in the paper)