

Toward Self-* Networks

Stefan Schmid (Uni Vienna)



A Great Time to Be a Networking Researcher!



Rhone and Arve Rivers,
Switzerland

Credits: George Varghese.

Flexibilities: Along 3 Dimensions



Passau, Germany

Inn, Donau, Ilz

Flexibilities: Along 3 Dimensions

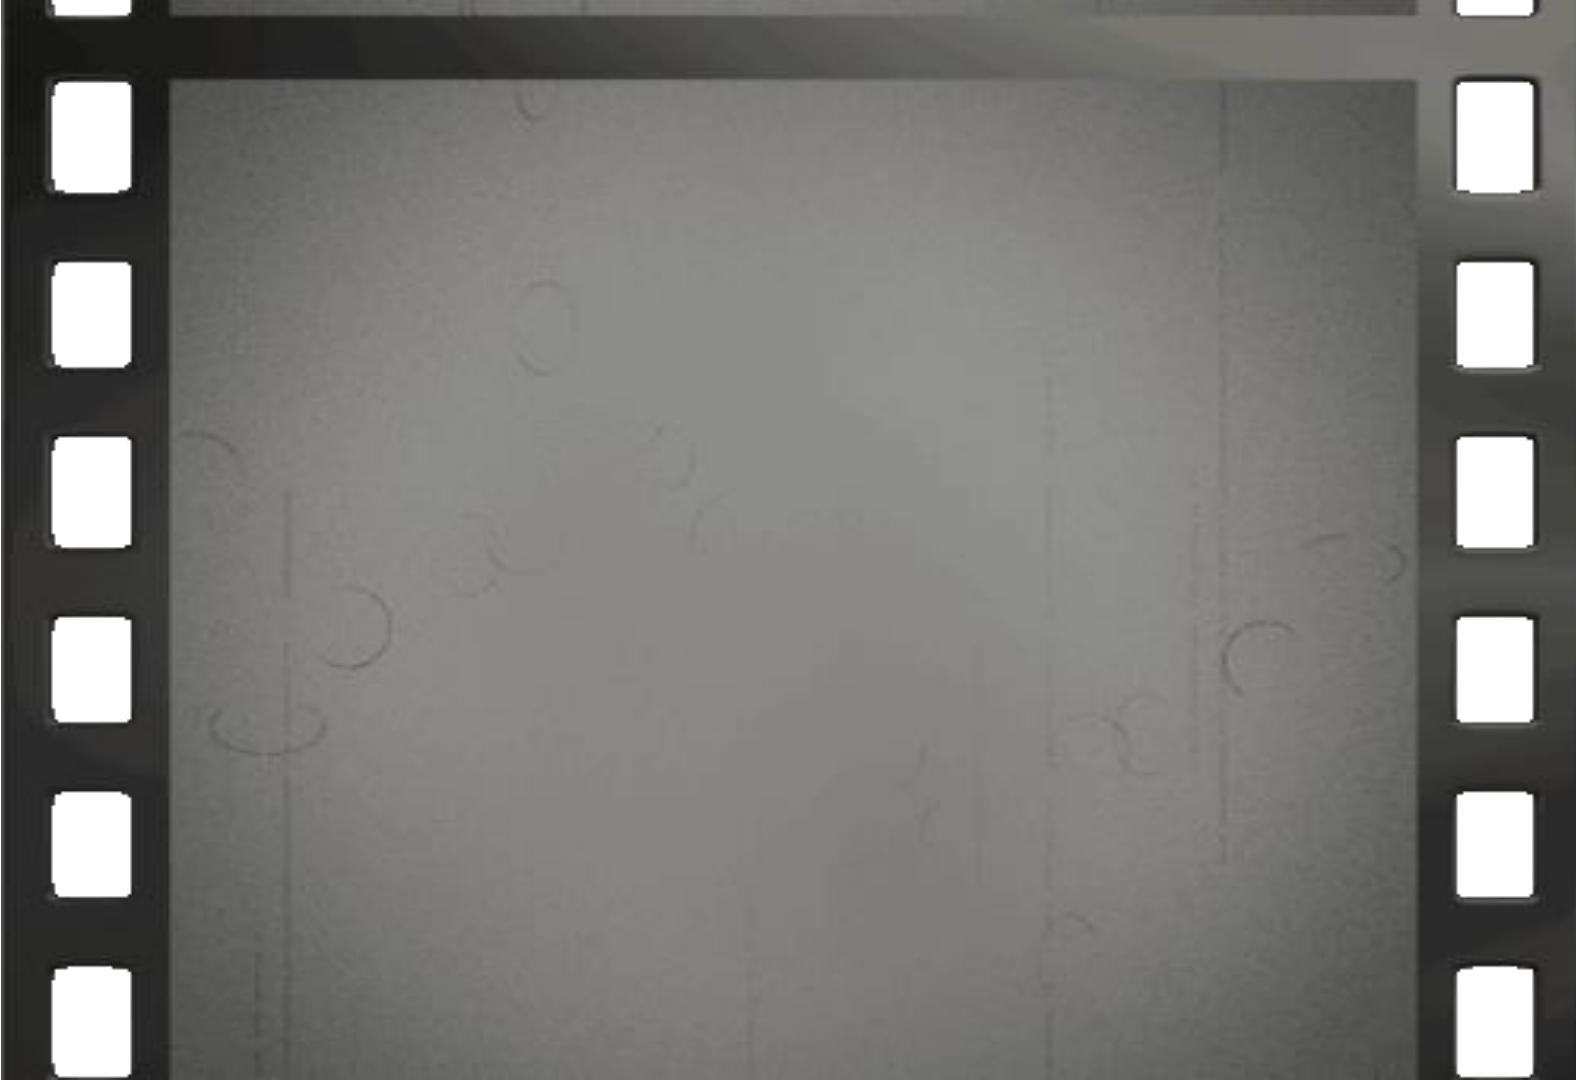


Passau, Germany

Inn, Donau, Ilz

Flexibilities: Along 3 Dimensions





Rewinding the clock of the Internet...

Shortest path routing only

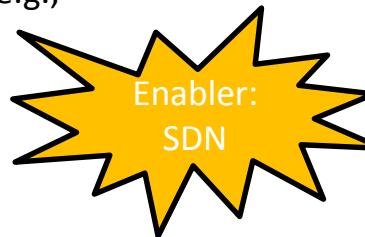
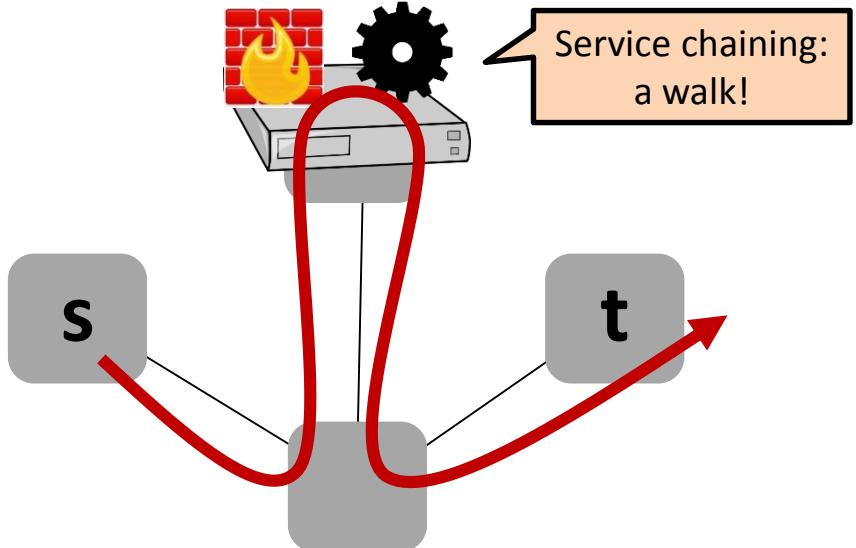
Indirect control: via weights only

Proprietary, blackbox implementations

Difficult and slow innovation

Opportunity: Flexible Routing

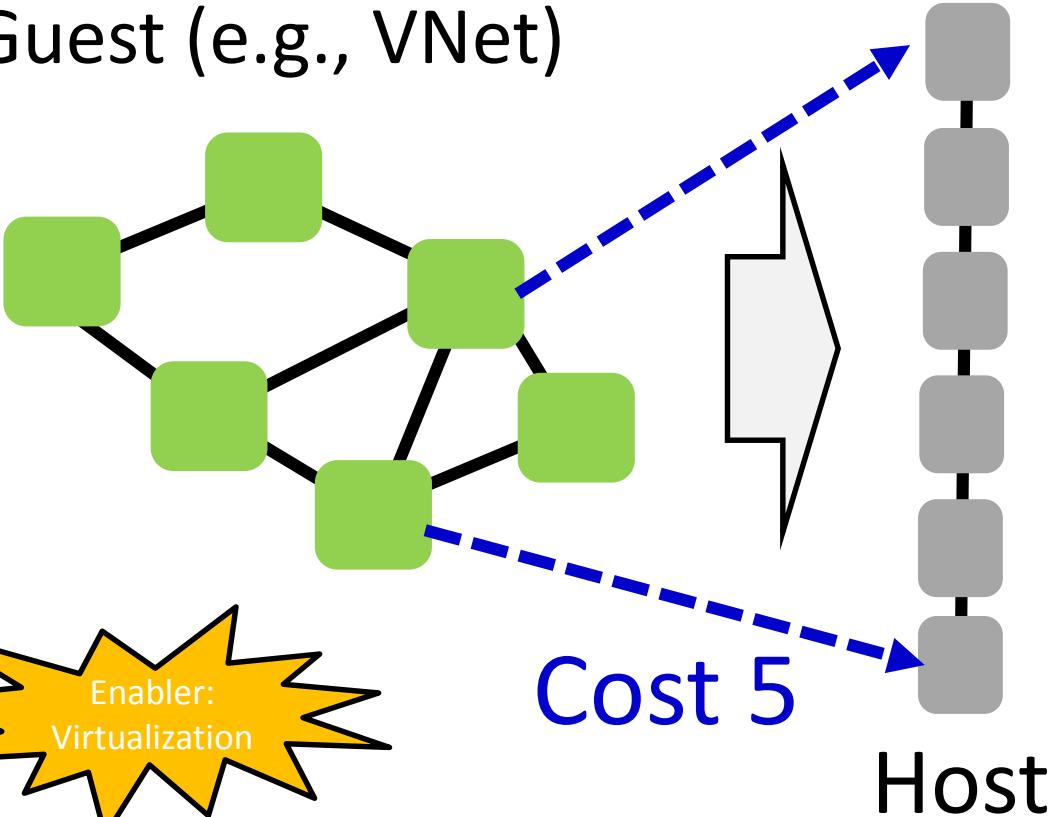
- Direct control over paths
 - Traditionally: indirect control *via weights* based on which *shortest paths* are computed
- More general routes
 - Beyond shortest paths, even *beyond „paths“*
 - E.g., steer traffic through (virtualized) middleboxes to compose new services like *service chains* („walk“)
- General match-action
 - SDN allows to match *L2-L4* and route, e.g., HTTP traffic differently (e.g., to cache)



Charting the Algorithmic Complexity
of Waypoint Routing. Amiri et al. ACM
SIGCOMM CCR, 2018.

Opportunity: Flexible Embedding

Guest (e.g., VNet)

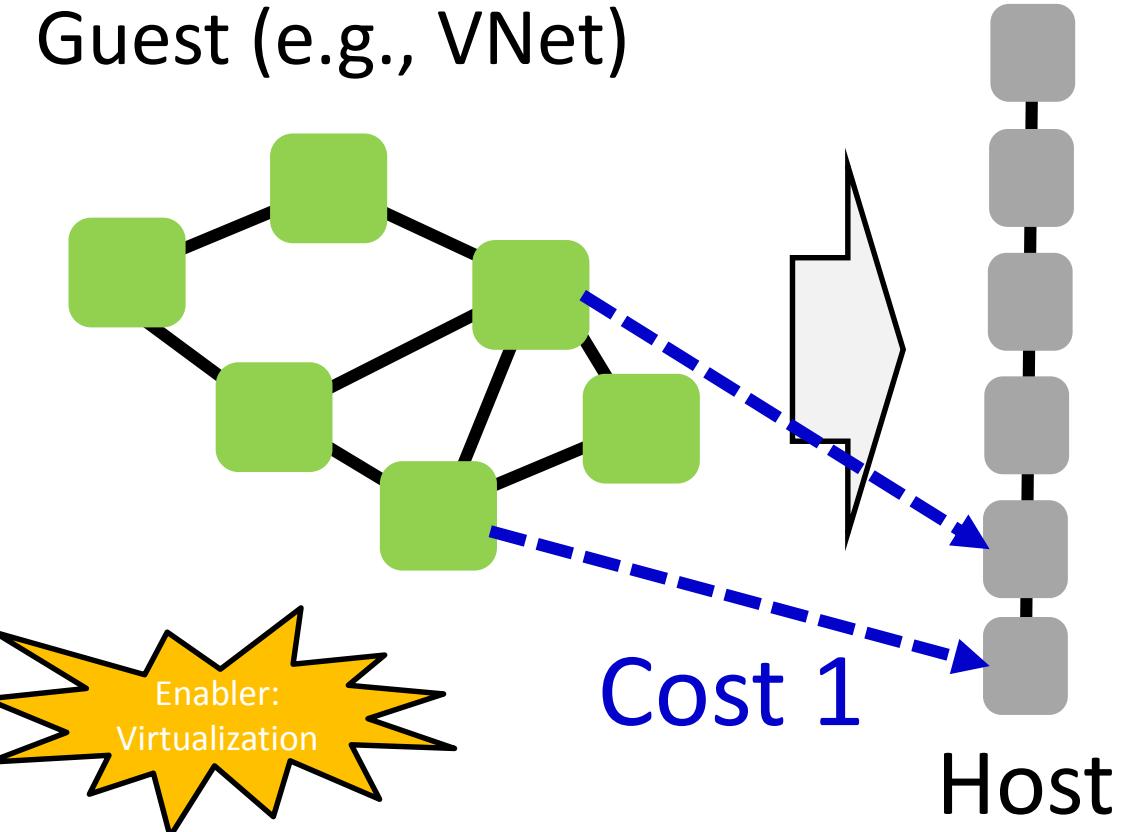


- Flexibly **allocate** (virtualized) network functions or **map** (virtualized) communication partners...
- ... to improve **utilization**, minimize **latency** and **load**, etc.

Charting the Complexity Landscape of
Virtual Network Embeddings. Rost et
al. IFIP Networking, 2018.

Opportunity: Flexible Embedding

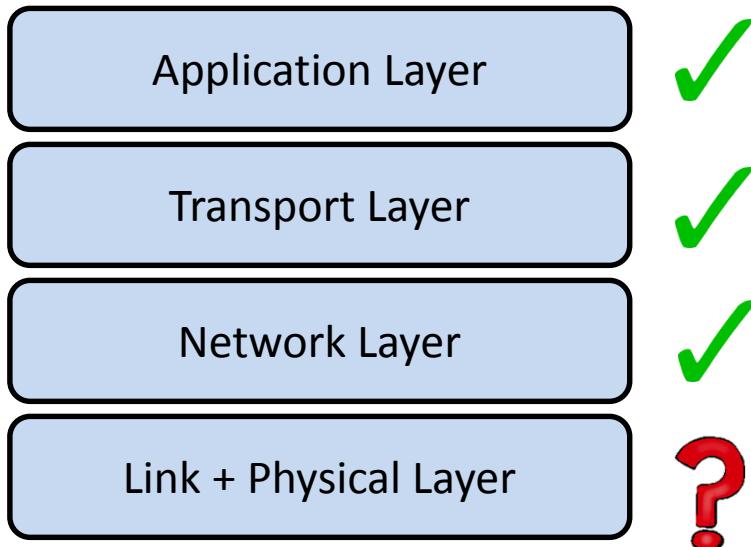
Guest (e.g., VNet)



- Flexibly **allocate** (virtualized) network functions or **map** (virtualized) communication partners...
- ... to improve **utilization**, minimize **latency** and **load**, etc.

Charting the Complexity Landscape of
Virtual Network Embeddings. Rost et
al. IFIP Networking, 2018.

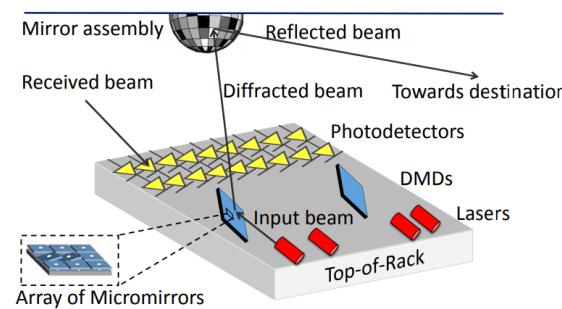
The Internet: Capable of Change on All Layers!



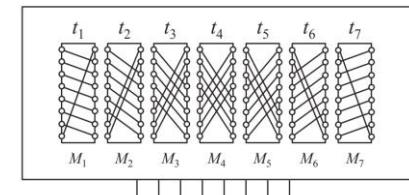
The Internet: Capable of Change on All Layers!

- Application Layer ✓
- Transport Layer ✓
- Network Layer ✓
- Link + Physical Layer ✓

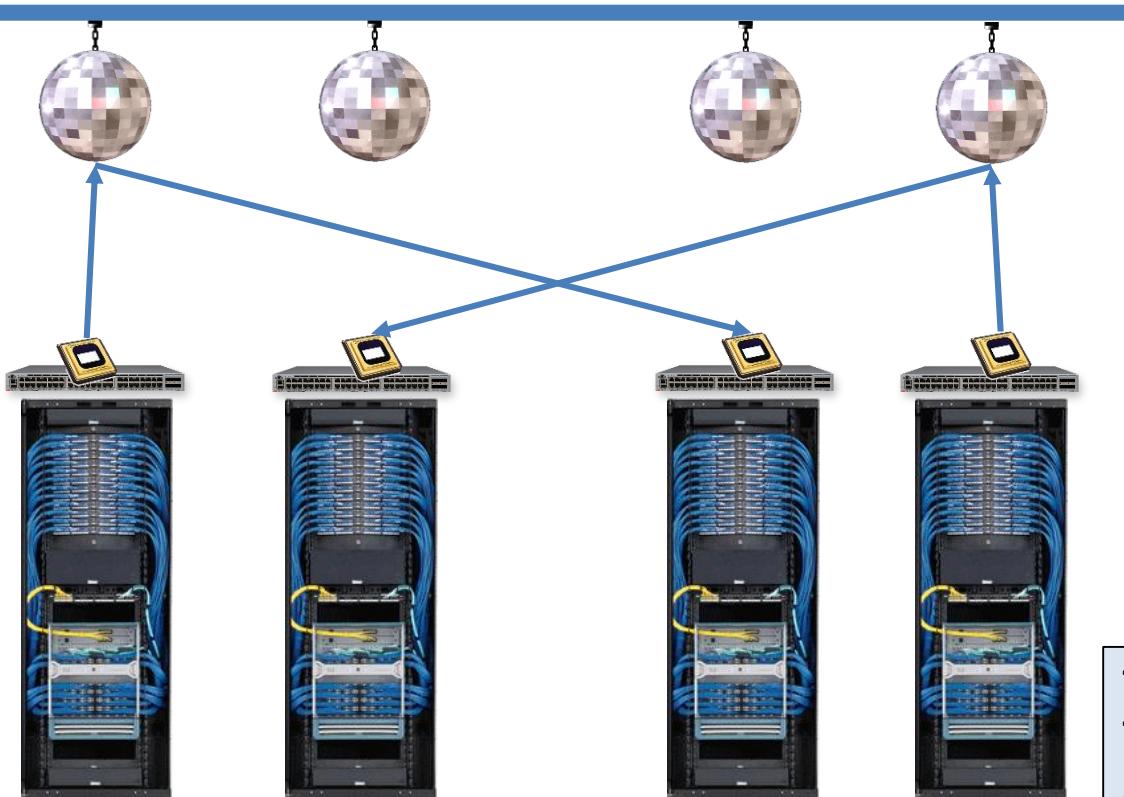
Based on **free-space optics**, **60GHz**, **optical circuit switches**, **movable antennas** and **mirrors**, etc.



Rotor switch



Opportunity: Flexible Topology Programming

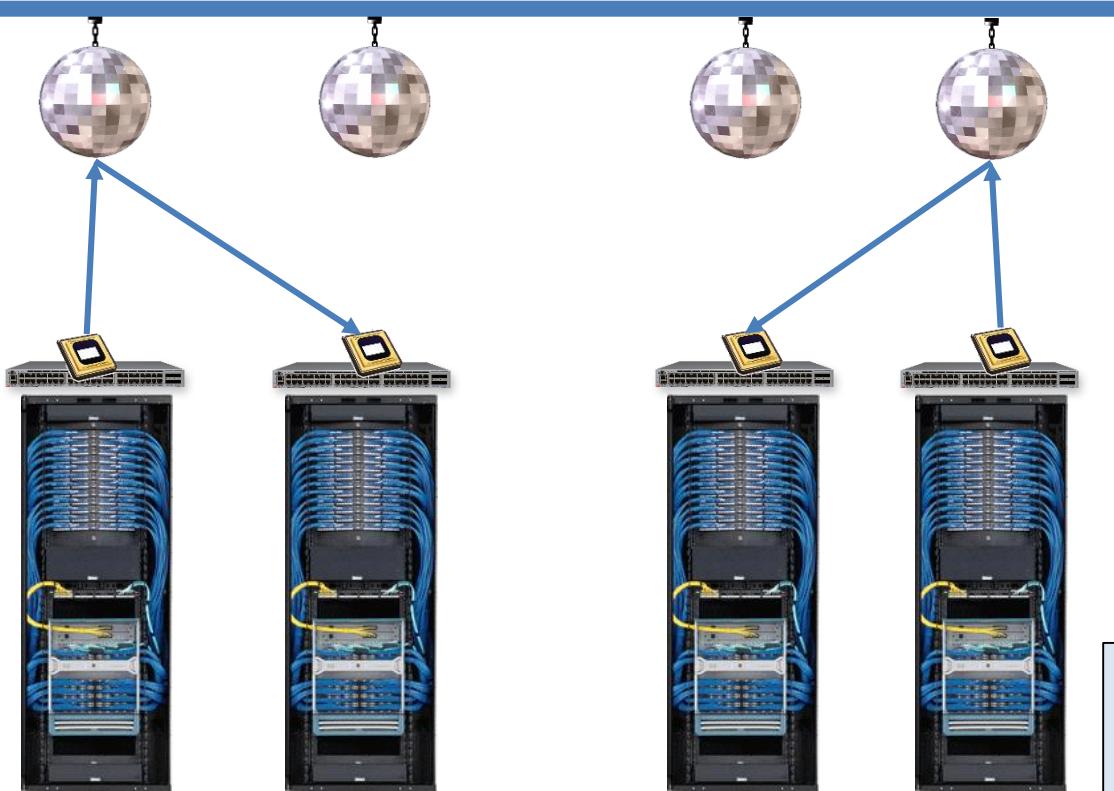


- **Reconfigure** networks towards needs

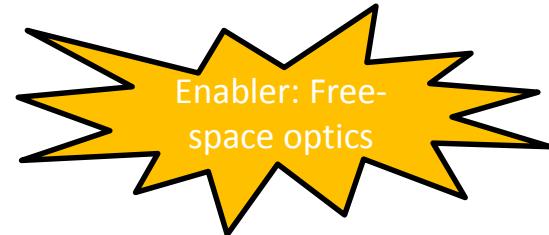


Toward Demand-Aware Networking: A Theory for Self-Adjusting Networks.
Avin et al. ACM SIGCOMM CCR, 2018.

Opportunity: Flexible Topology Programming



- **Reconfigure** networks towards needs



Toward Demand-Aware Networking: A Theory for Self-Adjusting Networks.
Avin et al. ACM SIGCOMM CCR, 2018.

Timeline

Reconfiguration time: from milliseconds **to microseconds** (and decentralized).

-
- 2009
 - *Flyways* [51]: Steerable antennas (narrow beamwidth at 60 GHz [78]) to serve hotspots
 - 2010
 - *Helios* [33]/*c-Through* [98, 99]: Hybrid switch architecture, maximum matching (Edmond's algorithm [30]), single-hop reconfigurable connections ($O(10)ms$ reconfiguration time).
 - *Proteus* [21, 89]: k reconfigurable connections per ToR, multi-hop path stitching, multi-hop reconfigurable connections (weighted b -matching [69], edge-exchanges for connectivity [72], wavelength assignment via edge-coloring [67] on multigraphs)
 - 2011
 - Extension of *Flyways* [51] to better handle practical concerns such as stability and interference for 60GHz links, along with greedy heuristics for dynamic link placement [45]
 - 2012
 - *Mirror Mirror on the ceiling* [106]: 3D-beamforming (60 Ghz wireless), signals bounce off the ceiling
 - 2013
 - *Mordia* [31, 32, 77]: Traffic matrix scheduling, matrix decomposition (Birkhoff-von-Neumann (BvN) [18, 97]), fiber ring structure with wavelengths ($O(10)\mu s$ reconfiguration time)
 - *SplayNets* [6, 76, 82]: Fine-grained and online reconfigurations in the spirit of self-adjusting datastructures (all links are reconfigurable), aiming to strike a balance between short route lengths and reconfiguration costs
 - 2014
 - *REACToR* [56]: Buffer burst of packets at end-hosts until circuit provisioned, employs [77]
 - *Firefly* [14]: Combination of Free Space Optics and Galvo/switchable mirrors (small fan-out)
 - 2015
 - *Solstice* [57]: Greedy perfect matching based hybrid scheduling heuristic that outperforms BvN [77]
 - Designs for optical switches with a reconfiguration latency of $O(10)ns$ [3]
 - 2016
 - *ProjectToR* [39]: Distributed Free Space Optics with digital micromirrors (high fan-out) [38] (Stable Matching [26]), goal of (starvation-free) low latency
 - *Eclipse* [95, 96]: $(1 - 1/e^{(1-\varepsilon)})$ -approximation for throughput in traffic matrix scheduling (single-hop reconfigurable connections, hybrid switch architecture), outperforms heuristics in [57]
 - 2017
 - *DAN* [7, 8, 11, 12]: Demand-aware networks based on reconfigurable links only and optimized for a demand snapshot, to minimized average route length and/or minimize load
 - *MegaSwitch* [23]: Non-blocking circuits over multiple fiber rings (stacking rings in [77] doesn't suffice)
 - *Rotornet* [63]: Oblivious cyclical reconfiguration w. selector switches [64] (Valiant load balancing [94])
 - *Tale of Two Topologies* [105]: Convert locally between Clos [24] topology and random graphs [87, 88]
 - 2018
 - *DeepConf* [81]/*xWeaver* [102]: Machine learning approaches for topology reconfiguration
 - 2019
 - Complexity classifications for weighted average path lengths in reconfigurable topologies [34, 35, 36]
 - *ReNet* [13] and *Push-Down-Trees* [9] providing statically and dynamically optimal reconfigurations
 - *DisSplayNets* [75]: fully decentralized *SplayNets*
 - *Opera* [60]: Maintaining expander-based topologies under (oblivious) reconfiguration

Survey of Reconfigurable Data Center Networks. Foerster and Schmid. SIGACT News, 2019.

Opportunity



Great **optimization opportunities**



In principle flexibilities can be exploited „fast“: **open interfaces** (***bring your own algorithm!***)



Also easier to **collect data**: programmable networks, **telemetry**

Challenge



Operating networks may become more **complex**, e.g.: **traversal** of firewall not mapped to „edge“



Exploiting **online optimization** at runtime hard at **human** time scale: difficult algorithmic problems



Modelling can be **more difficult** too: new components like **hypervisor** can affect performance

Challenge: Model vs Reality

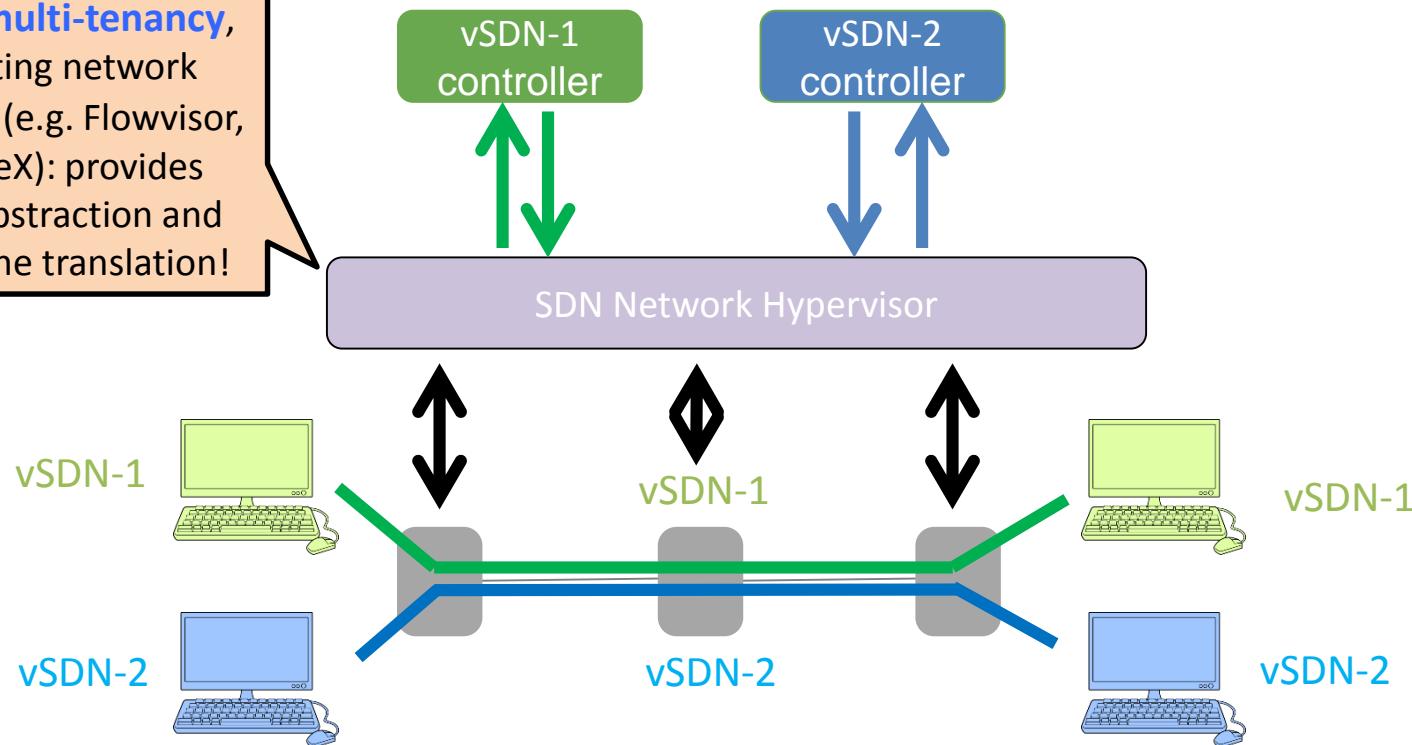
You: I invented a great new algorithm to route and embed service chains at low resource cost and providing minimal bandwidth guarantees!

Boss: So can I promise our customers a predictable performance?

You: hmm...

Recall Andi Blenk's Talk

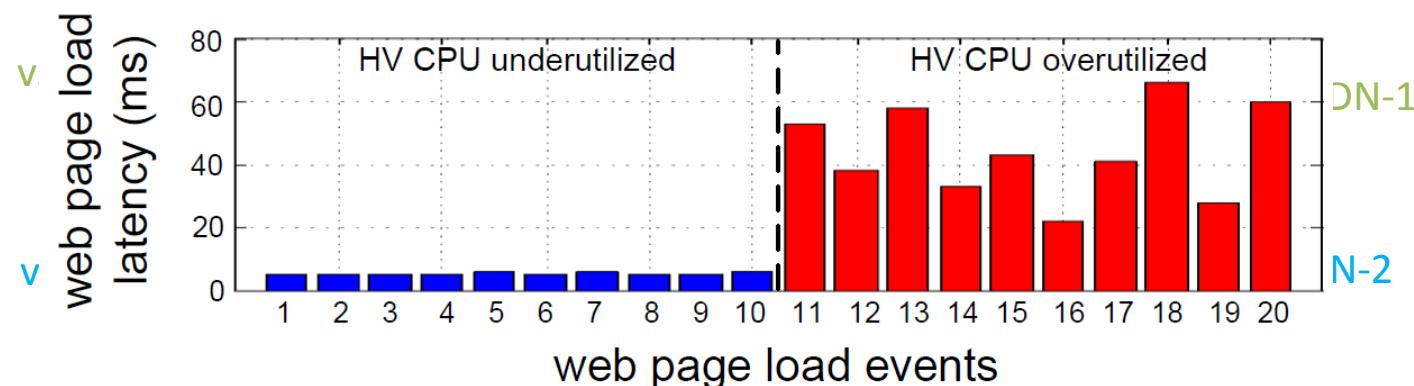
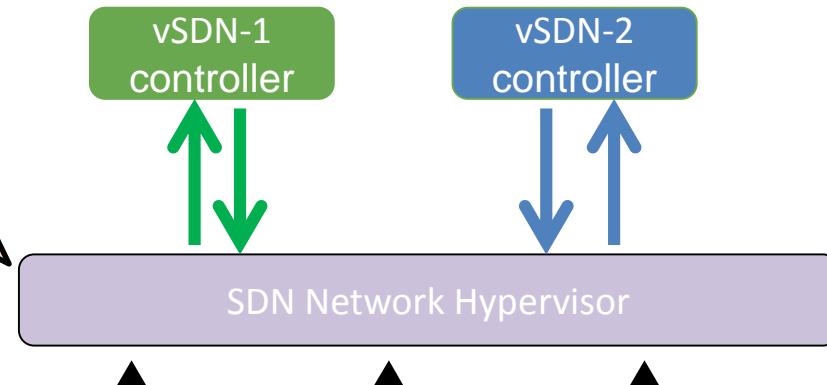
To enable **multi-tenancy**, take existing network **hypervisor** (e.g. Flowvisor, OpenVirtex): provides network abstraction and control plane translation!



An Experiment: 2 vSDNs with bw guarantee!

Recall Andi Blenk's Talk

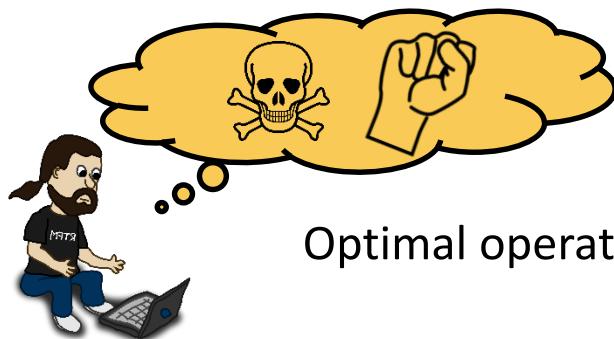
To enable **multi-tenancy**, take existing network **hypervisor** (e.g. Flowvisor, OpenVirtex): provides network abstraction and control plane translation!



An Experiment: 2 vSDNs with bw guarantee!

First Conclusions

- Exploiting network flexibilities is non-trivial, especially if **fine-grained** and **fast** reactions are desired
- Also modelling such networked systems is challenging: details of **interference**, **demand**, etc. will only be available at runtime



Optimal operation of flexible networks ***too complex for humans.***



Let's give up control: self-* networks!

Self-observing, self-adjusting, self-repairing, “self-driving”, ...

It's about
automation!

Roadmap

- Opportunities of self-* networks
 - Example 1: Demand-aware, self-adjusting networks
 - Example 2: Self-repairing networks
- Challenges of designing self-* networks



Roadmap

- Opportunities of self-* networks
 - **Example 1: Demand-aware, self-adjusting networks**
 - Example 2: Self-repairing networks
- Challenges of designing self-* networks

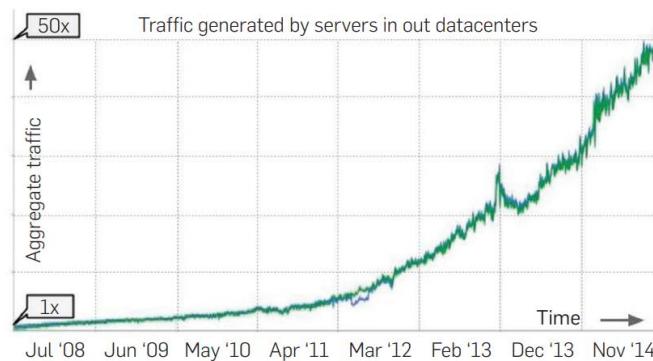


Why Demand-Aware...?

Case study: datacenter networks

Explosive Growth of Demand...

Batch processing, web services,
distributed ML, ...: *data-centric applications* are distributed and interconnecting network is *critical*



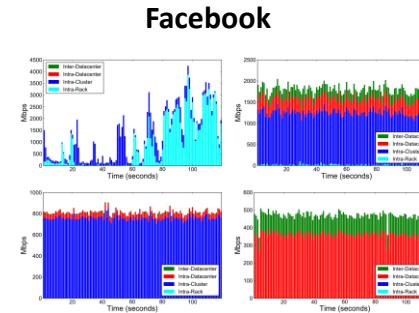
Source: Jupiter Rising. SIGCOMM 2015.

Aggregate server traffic in
Google's datacenter fleet

... But Much Structure!

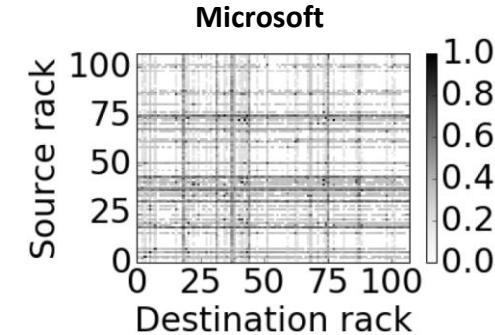
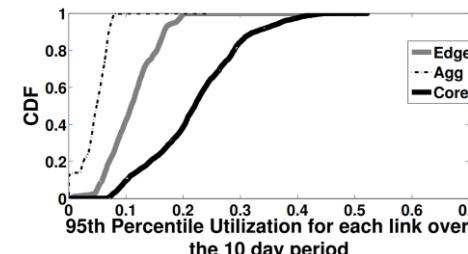


Spatial (*sparse!*) and temporal *locality*



Inside the Social Network's
(Datacenter) Network @
SIGCOMM 2015

Benson et al.

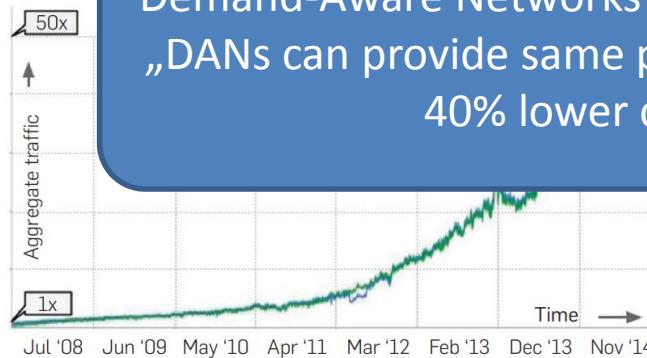


ProjecToR @ SIGCOMM 2016

Understanding Data Center Traffic
Characteristics @ WREN 2009

Explosive Growth of Demand...

Batch processing, web services,
distributed ML, ...: ***data-centric applications*** are distributed and
interconnecting network is ***critical***

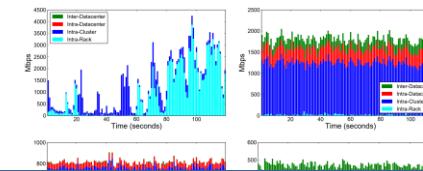


Source: Jupiter Rising. SIGCOMM 2015.

Aggregate server traffic in
Google's datacenter fleet

... But Much Structure!

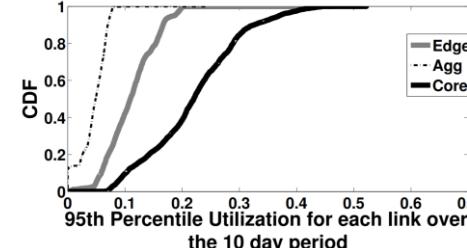
Facebook



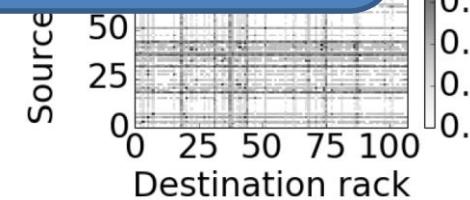
Small Varying and

Demand-Aware Networks (DANs) can exploit this structure by adapting to it:
„DANs can provide same performance as demand-oblivious networks at 25-
40% lower costs.“ Firefly, SIGCOMM CCR, 2014.

BENSON ET AL.



Understanding Data Center Traffic
Characteristics @ WREN 2009



ProjecToR @ SIGCOMM 2016

Datacenter Networks

Traditionally: demand-**oblivious**:



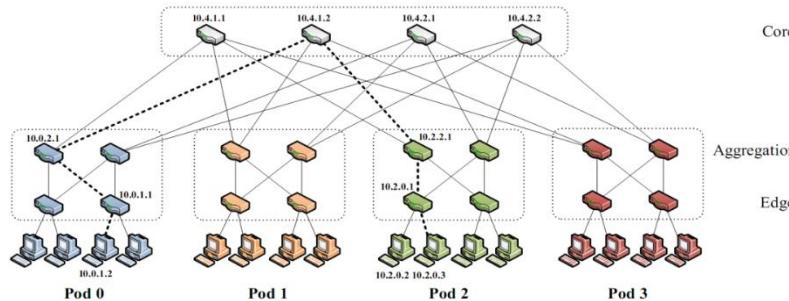
Datacenter Networks

Traditionally: demand-**oblivious**:



Datacenter Networks

Traditional datacenter network



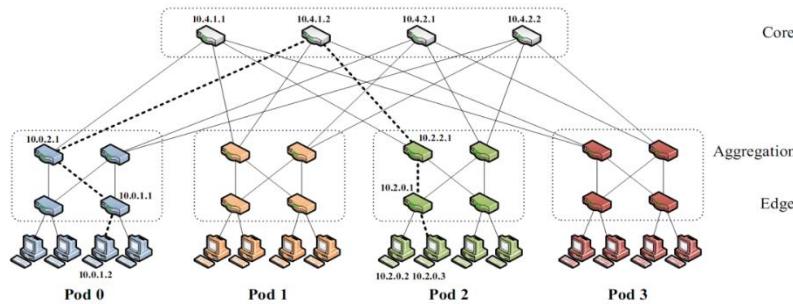
Traditionally: demand-**oblivious**:



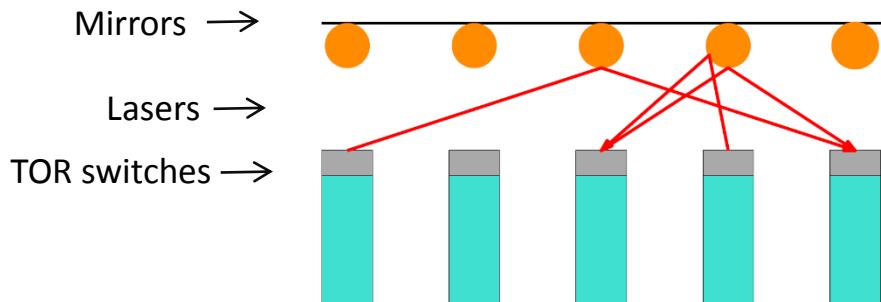
- Usually optimized **for the “worst-case”** (all-to-all communication)
- Example, fat-tree topologies: provide **full bisection bandwidth**

Datacenter Networks

Traditional datacenter network



Reconfigurable datacenter network

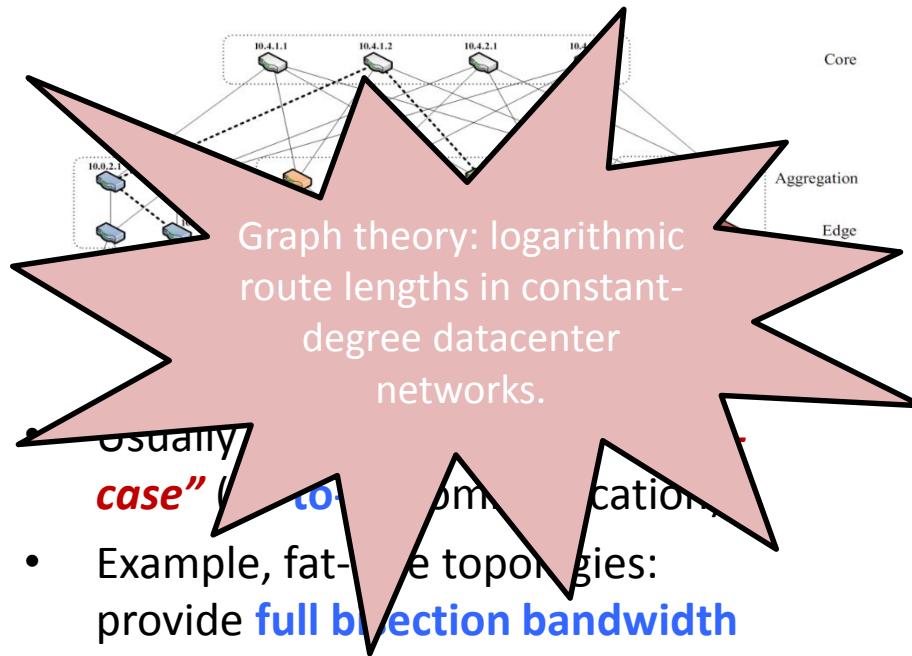


- Usually optimized *for the “worst-case”* (all-to-all communication)
- Example, fat-tree topologies: provide **full bisection bandwidth**

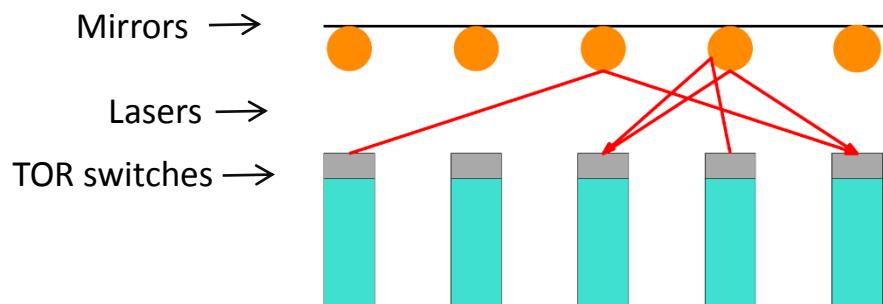
- Optimized *toward the workload* it serves (e.g., **route length**)
- Statically or *even dynamically*

Datacenter Networks

Traditional datacenter network



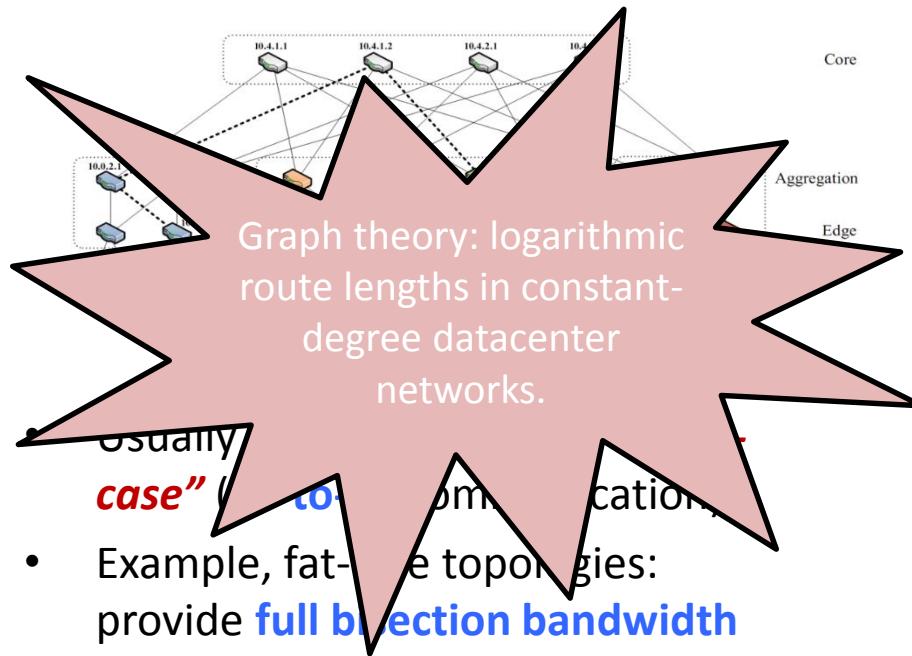
Reconfigurable datacenter network



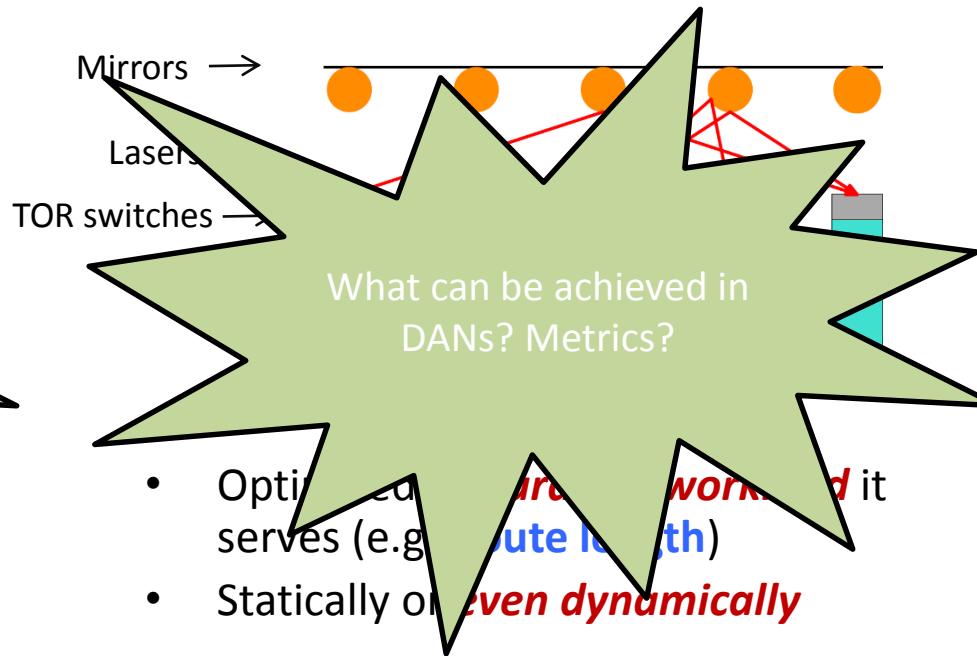
- Optimized *toward the workload* it serves (e.g., *route length*)
- Statically or *even dynamically*

Datacenter Networks

Traditional datacenter network



Reconfigurable datacenter network



DAN Design: New Types of Problems

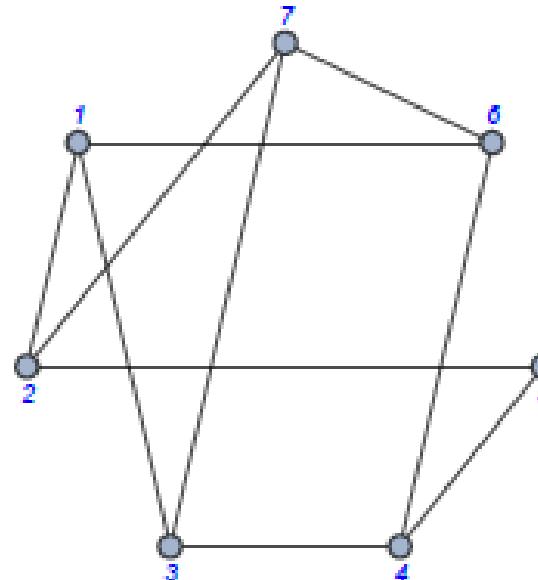
Input: Workload

Destinations

		1	2	3	4	5	6	7
		0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
Sources		1	$\frac{2}{65}$	$\frac{13}{65}$	$\frac{65}{65}$	$\frac{65}{65}$	$\frac{65}{65}$	$\frac{65}{65}$
1	2	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
2	3	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$
3	4	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	0	$\frac{1}{13}$
4	5	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0
5	6	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0	0	0
6	7	$\frac{2}{65}$	0	0	0	0	0	$\frac{3}{65}$
7		$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0

design

Output: DAN



Demand matrix: joint distribution

... of *constant degree* (scalability)

DAN Design: New Types of Problems

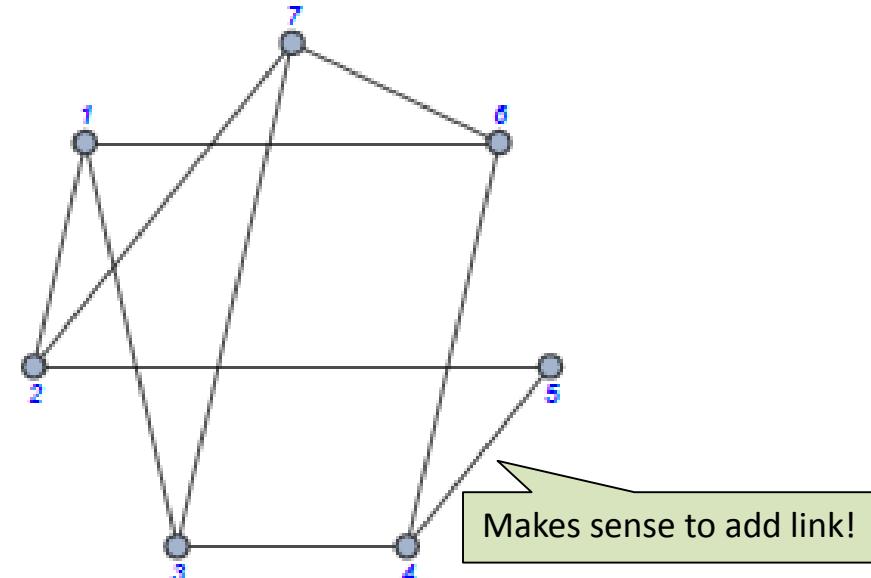
Input: Workload

Destinations

Sources	1	2	3	4	5	6	7
1	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
2	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$
3	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{13}{65}$	0
4	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0
5	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0	0	0
6	$\frac{2}{65}$	0	0	0	0	0	$\frac{3}{65}$
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0

design

Output: DAN



Demand matrix: joint distribution

... of *constant degree* (scalability)

DAN Design: New Types of Problems

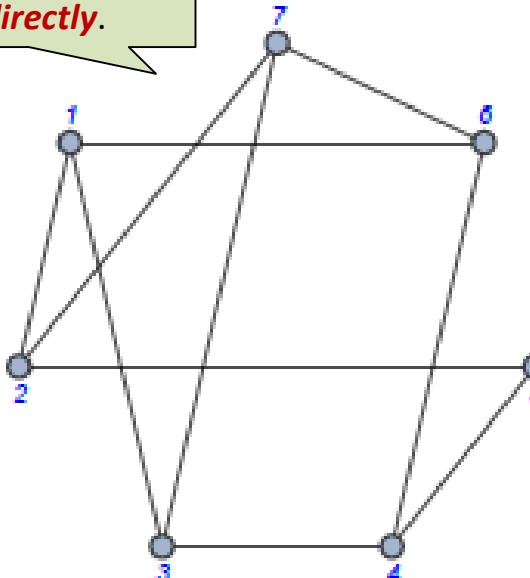
Input: Workload

Destinations								
		1 communicates to many.						
		6	7					
1	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$	
2	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$	
3	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	0	$\frac{1}{13}$	
4	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0	
5	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0	0	0	
6	$\frac{2}{65}$	0	0	0	0	0	$\frac{3}{65}$	
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0	

Bounded degree: route
to 7 *indirectly*.

Output: DAN

design



Demand matrix: joint distribution

... of *constant degree* (scalability)

DAN Design: New Types of Problems

Input: Workload

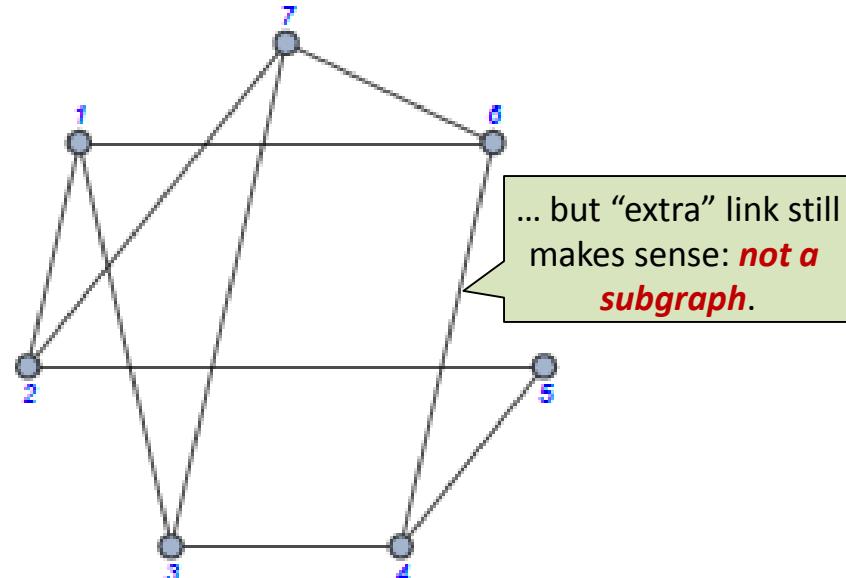
Destinations

Sources	1	2	3	4	5	6	7
1	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
2	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$
3	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	0	$\frac{1}{13}$
4	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0
5	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0		
6	$\frac{2}{65}$	0	0	0	0		
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0

design

Demand matrix: joint distribution

Output: DAN



... of **constant degree** (scalability)

More Formally: DAN Design Problem

Input:

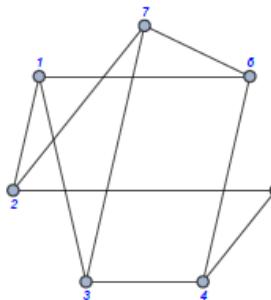
$\mathcal{D}[p(i,j)]$: joint **distribution**, Δ

		Y						
		1	2	3	4	5	6	7
X	1	0	<u>2</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>2</u>	<u>3</u>
	2	<u>2</u>	0	<u>1</u>	0	0	0	<u>2</u>
3	<u>1</u>	<u>1</u>	0	<u>2</u>	0	0	0	<u>1</u>
	13	65	65	65	13	13	65	65
4	<u>1</u>	0	<u>2</u>	0	<u>4</u>	0	0	0
	65	65	65	65	65	65	65	65
5	<u>1</u>	0	<u>3</u>	<u>4</u>	0	0	0	0
	65	65	65	65	65	65	65	65
6	<u>2</u>	0	0	0	0	0	0	<u>3</u>
	65	65	65	65	65	65	65	65
7	<u>3</u>	<u>2</u>	<u>1</u>	0	0	<u>3</u>	0	0
	65	65	13	65	65	65	65	65



Output:

N: DAN



Bounded degree
 $\Delta=3$

Objective:

Expected Path Length (EPL):

Demand-weighted route length

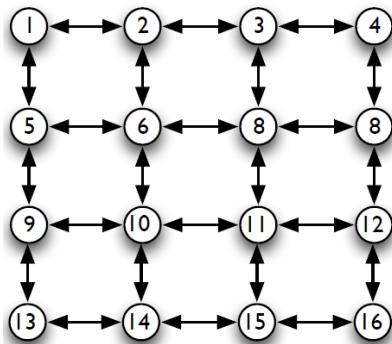
$$\text{EPL}(\mathcal{D}, N) = \sum_{(u,v) \in \mathcal{D}} p(u, v) \cdot d_N(u, v)$$

Path length **on DAN N.**

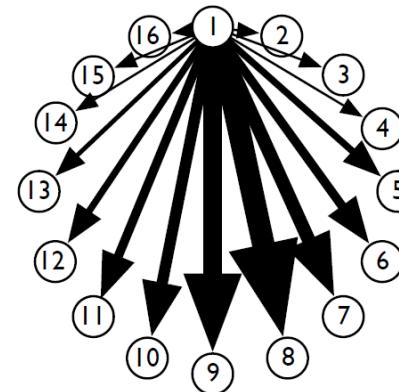
Frequency

Sometimes, DANs can be much better!

Example 1: low-degree demand



Example 2: high-degree but *skewed* demand



- Already low degree: degree-4 DAN can serve this *at cost 1*.
- If sufficiently skewed: constant-degree DAN can serve it at cost *O(1)*

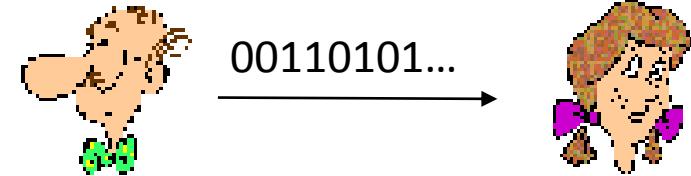
So on what does it depend?

So on what does it depend?



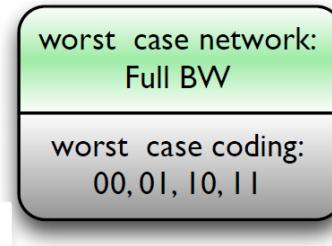
We argue (but still don't know!): on the
“entropy” of the demand!





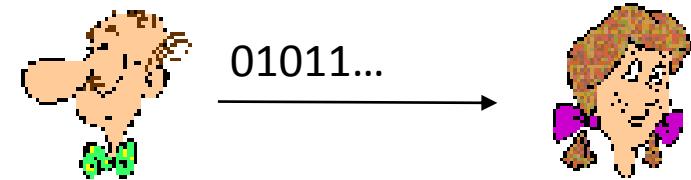
An Analogy to Coding

if demand **arbitrary** and **unknown**



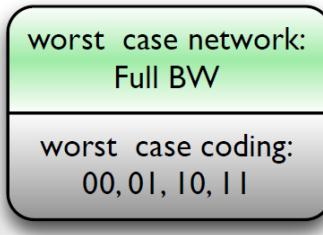
log diameter

log # bits / symbol



An Analogy to Coding

if demand **arbitrary** and **unknown**



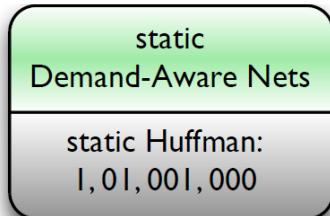
log diameter

log # bits / symbol



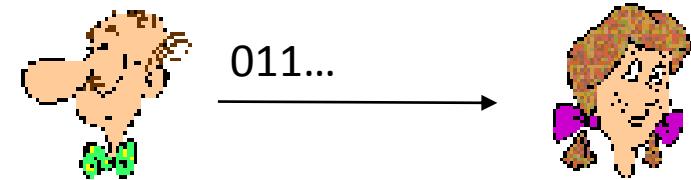
DAN!

if demand **known** and **fixed**



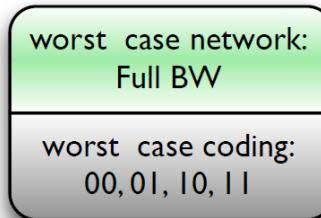
entropy?

entropy / symbol



An Analogy to Coding

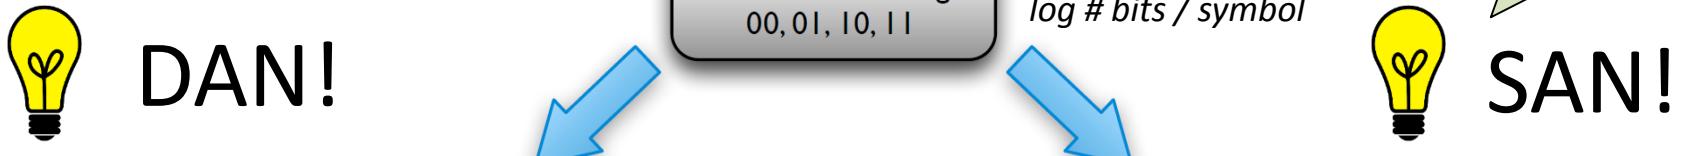
if demand **arbitrary** and **unknown**



log diameter

Dynamic DANs:
Aka. **Self-Adjusting Networks (SANs)!**

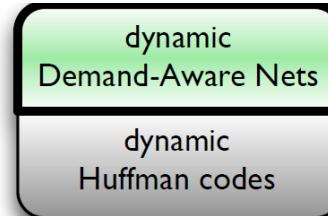
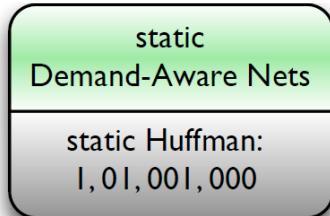
log # bits / symbol



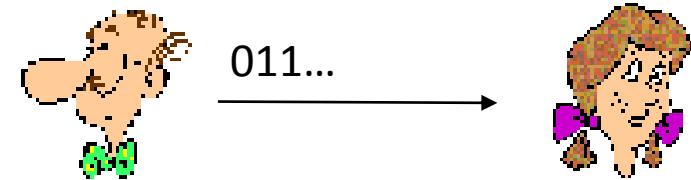
if demand **known** and **fixed**

entropy?

entropy / symbol

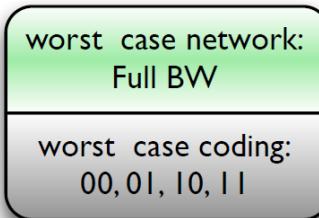


if demand **unknown** but **reconfigurable**



An Analogy to Coding

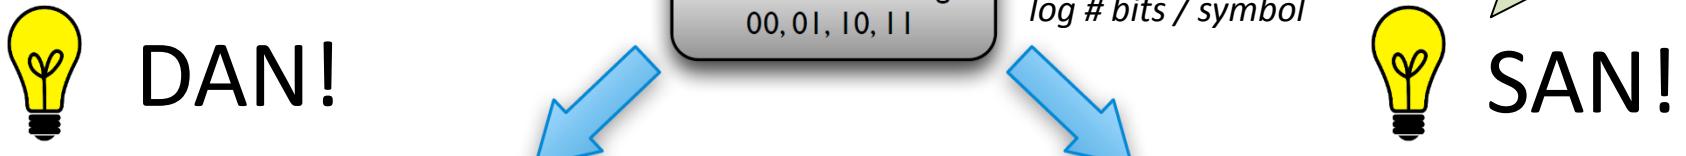
if demand **arbitrary** and **unknown**



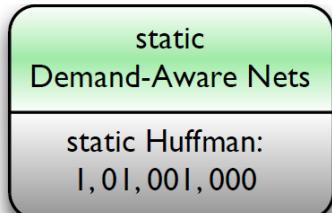
log diameter

Dynamic DANs:
Aka. **Self-Adjusting Networks (SANs)!**

log # bits / symbol

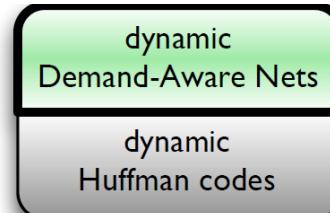


if demand **known** and **fixed**



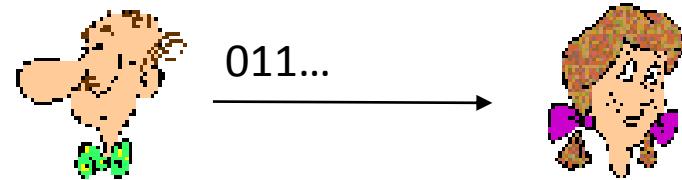
Can exploit
spatial locality!

if demand **unknown** but **reconfigurable**

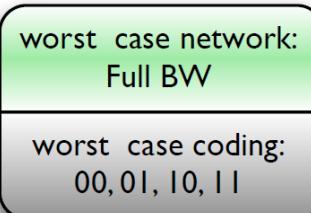


Additionally exploit
temporal locality!

An Analogy to Coding



if demand **arbitrary** and **unknown**



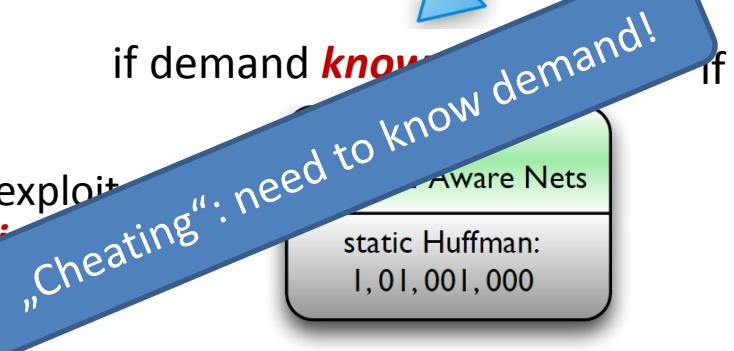
log diameter

Dynamic DANs:
Aka. **Self-Adjusting Networks (SANs)!**

log # bits / symbol



if demand **know**



if demand **unknown** but **regular**

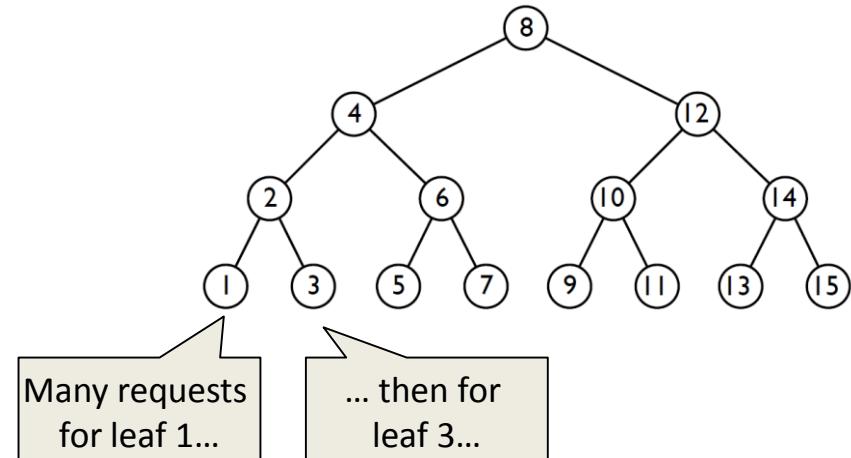
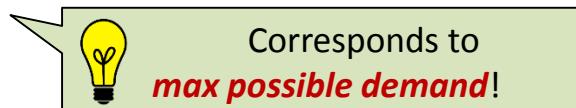


Need online algorithms!

Additionally exploit **temporal locality!**

Analogous to *Datastructures*: Oblivious...

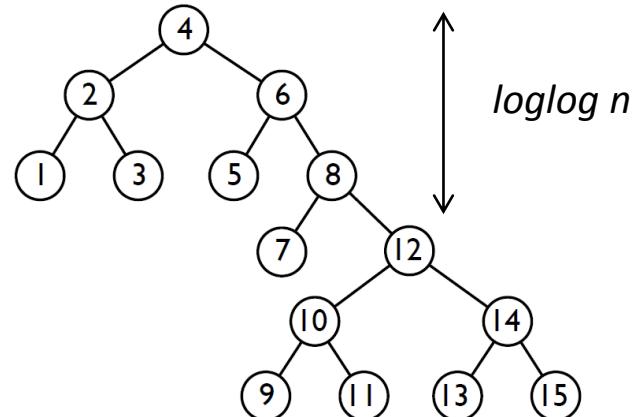
- Traditional, **fixed** BSTs do not rely on any assumptions on the demand
- Optimize for the **worst-case**
- Example **demand**:
 $1, \dots, 1, 3, \dots, 3, 5, \dots, 5, 7, \dots, 7, \dots, \log(n), \dots, \log(n)$
 $\longleftrightarrow \longleftrightarrow \longleftrightarrow \longleftrightarrow \longleftrightarrow \quad \longleftrightarrow$
many many many many *many*
- Items stored at **$O(\log n)$** from the root, **uniformly** and **independently** of their frequency



... Demand-Aware ...

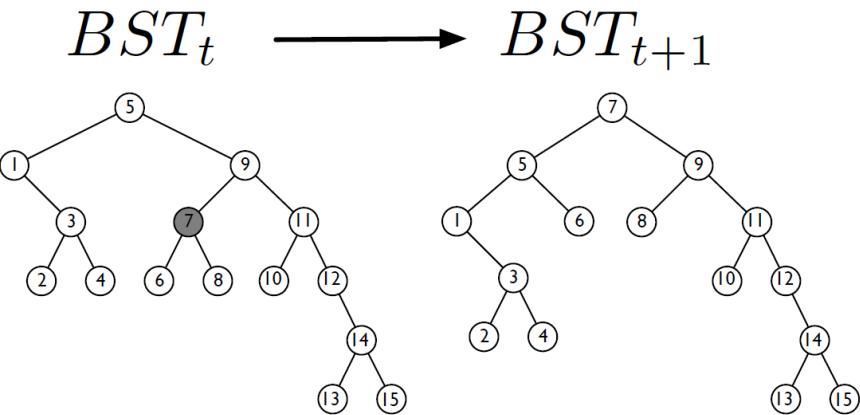
- **Demand-aware fixed** BSTs can take advantage of *spatial locality* of the demand
- E.g.: place frequently accessed elements close to the root
- E.g., **Knuth/Mehlhorn/Tarjan** trees
- Recall example **demand**:
 $1, \dots, 1, 3, \dots, 3, 5, \dots, 5, 7, \dots, 7, \dots, \log(n), \dots, \log(n)$
 - Amortized cost $O(\log \log n)$

 Amortized cost corresponds
to *empirical entropy of demand!*



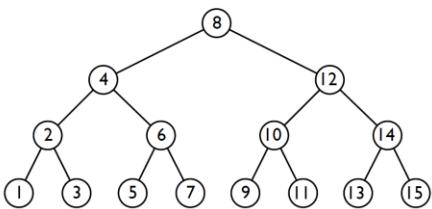
... Self-Adjusting!

- Demand-aware reconfigurable BSTs can additionally take advantage of *temporal locality*
- By moving accessed element to the root: amortized cost is *constant*, i.e., $O(1)$
 - Recall example demand:
 $1, \dots, 1, 3, \dots, 3, 5, \dots, 5, 7, \dots, 7, \dots, \log(n), \dots, \log(n)$

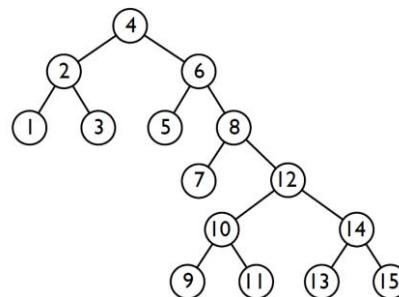


Datastructures

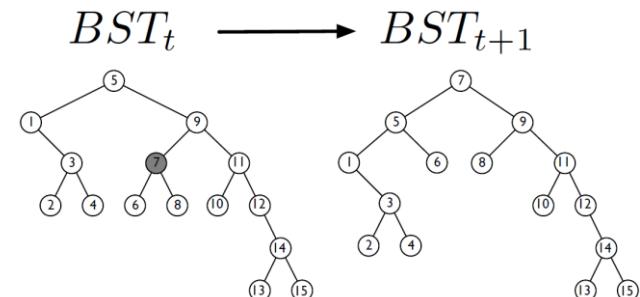
Oblivious



Demand-Aware



Self-Adjusting



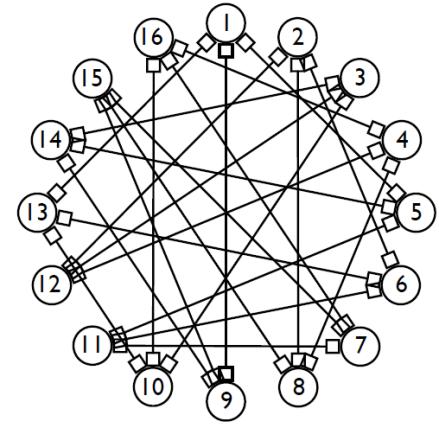
Lookup
 $O(\log n)$

Exploit **spatial locality**:
empirical entropy $O(\log \log n)$

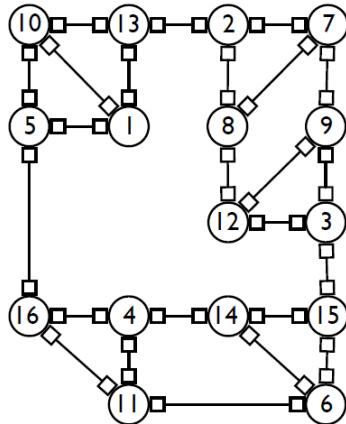
Exploit **temporal locality** as well:
 $O(1)$

Analogously for Networks

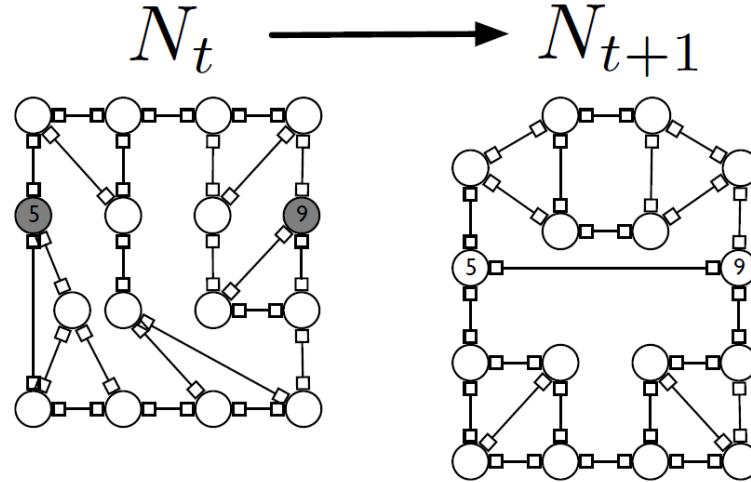
Oblivious



DAN



SAN



Const degree
(e.g., **expander**):
route lengths **$O(\log n)$**

Exploit **spatial locality**

Exploit **temporal locality** as well

SPEED
LIMIT
?

Intuition: Entropy Lower Bound



Lower Bound Idea: Leverage Coding or Datastructure

Destinations		1	2	3	4	5	6	7
Sources	1	0 65	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
	2	$\frac{2}{65}$	0 65	$\frac{1}{65}$	0 0	0 0	0 0	$\frac{2}{65}$
	3	$\frac{1}{13}$	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 0	0 0	$\frac{1}{13}$
	4	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 0	$\frac{4}{65}$	0 0	0 0
	5	$\frac{1}{65}$	0 65	$\frac{3}{65}$	$\frac{4}{65}$	0 0	0 0	0 0
	6	$\frac{2}{65}$	0 65	0 0	0 0	0 0	0 0	$\frac{3}{65}$
	7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0 0	0 0	$\frac{3}{65}$	0 0

- DAN just for a *single (source) node 1*: cannot do better than Δ -ary **Huffman tree** (or a **biased BST**) for its destinations

- How good can this tree be?



Entropy lower bound on EPL known for binary trees, e.g. **Mehlhorn** 1975 for BST

Lower Bound Idea: Leverage Coding or Data Structure

Destinations		1	2	3	4	5	6	7
Sources	1	0 65	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
	2	$\frac{2}{65}$	0 65	$\frac{1}{65}$	0 0	0 0	0 0	$\frac{2}{65}$
3	$\frac{1}{13}$	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 0	0 0	0 0	$\frac{1}{13}$
4	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 0	$\frac{4}{65}$	0 0	0 0	0 0
5	$\frac{1}{65}$	0 65	$\frac{3}{65}$	$\frac{4}{65}$	0 0	0 0	0 0	0 0
6	$\frac{2}{65}$	0 65	0 0	0 0	0 0	0 0	$\frac{3}{65}$	0 0
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0 0	0 0	$\frac{3}{65}$	0 0	0 0

An optimal “ego-tree”
for this source!

- DAN just for a **single (source) node 1**: cannot do better than Δ -ary **Huffman tree** (or a **biased BST**) for its destinations

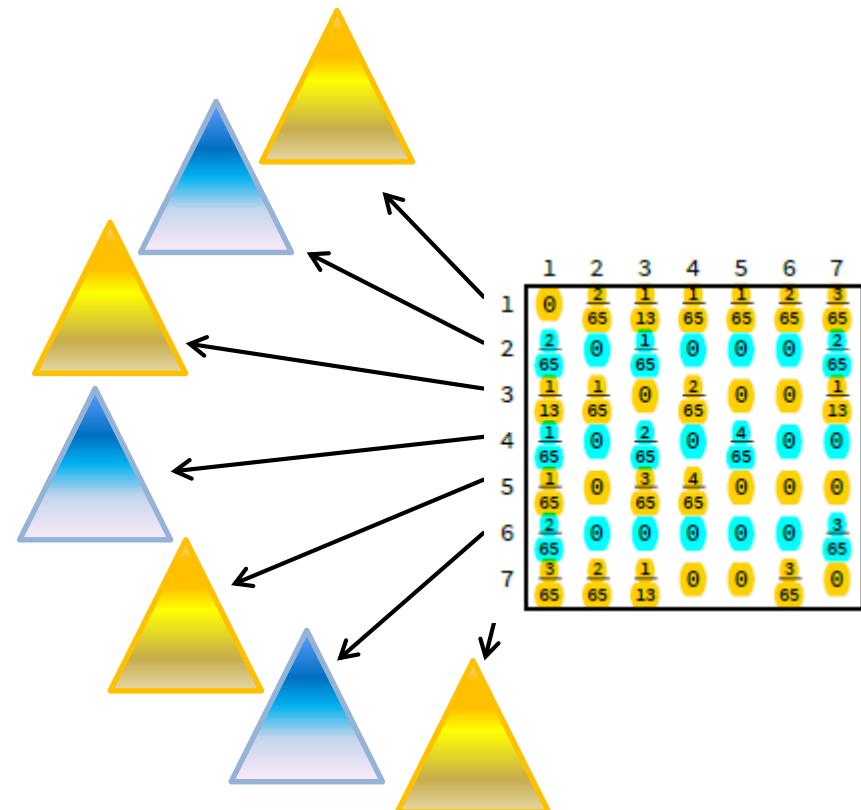
- How good can this tree be?



Entropy lower bound on EPL known for binary trees, e.g. **Mehlhorn** 1975 for BST

So: Entropy of the *Entire* Demand

- Proof idea ($EPL = \Omega(H_{\Delta}(Y|X))$):
 - sources
 - destinations
 - entropy
 - degree
- Compute **ego-tree** for each source node
- Take **union** of all **ego-trees**
- Violates **degree restriction** but valid lower bound



Entropy of the *Entire* Demand: Sources and Destinations

Do this in **both dimensions**:

$$\text{EPL} \geq \Omega(\max\{\mathcal{H}_\Delta(Y|X), \mathcal{H}_\Delta(X|Y)\})$$

$\Omega(\mathcal{H}_\Delta(X Y))$						
1	2	3	4	5	6	7
1	0 $\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
2	$\frac{2}{65}$	0 $\frac{1}{65}$	0 $\frac{1}{65}$	0 $\frac{1}{65}$	0 $\frac{1}{65}$	$\frac{2}{65}$
3	$\frac{1}{13}$	$\frac{1}{65}$	0 $\frac{2}{65}$	0 $\frac{1}{65}$	0 $\frac{1}{65}$	$\frac{1}{13}$
4	$\frac{1}{65}$	0 $\frac{2}{65}$	0 $\frac{1}{65}$	0 $\frac{1}{65}$	$\frac{4}{65}$	0 $\frac{0}{65}$
5	$\frac{1}{65}$	0 $\frac{3}{65}$	$\frac{4}{65}$	0 $\frac{1}{65}$	0 $\frac{0}{65}$	0 $\frac{0}{65}$
6	$\frac{2}{65}$	0 $\frac{0}{65}$	0 $\frac{0}{65}$	0 $\frac{0}{65}$	0 $\frac{0}{65}$	$\frac{3}{65}$
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0 $\frac{0}{65}$	0 $\frac{0}{65}$	0 $\frac{0}{65}$

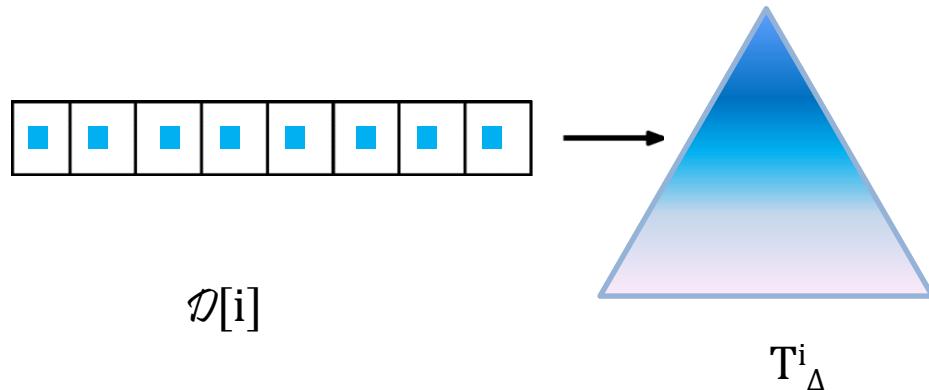
\mathcal{D}

Intuition: Reaching Entropy Limit in Datacenters



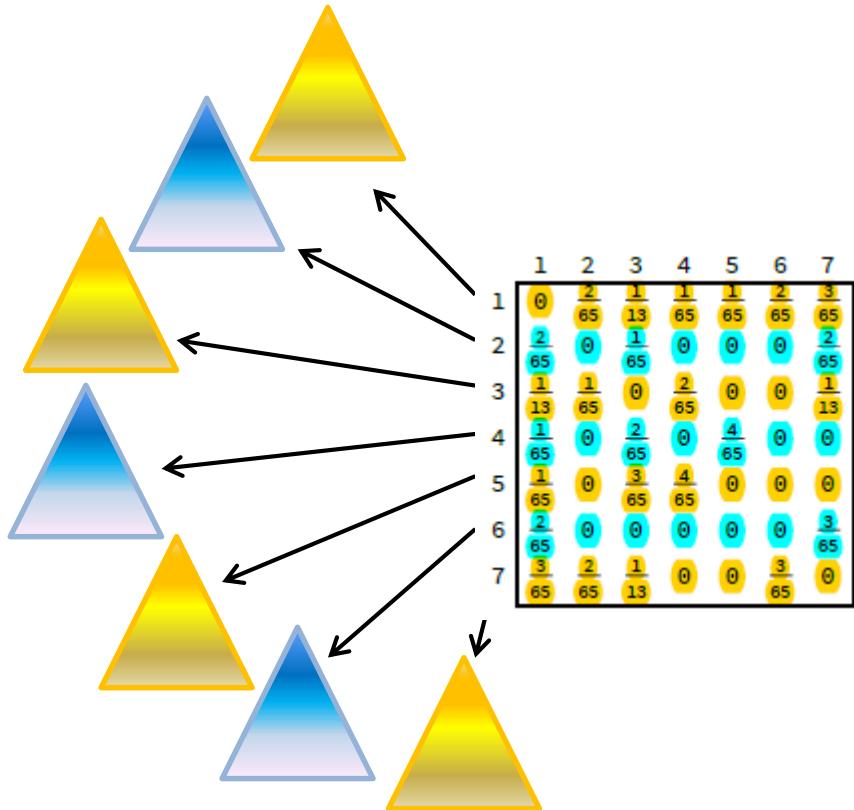
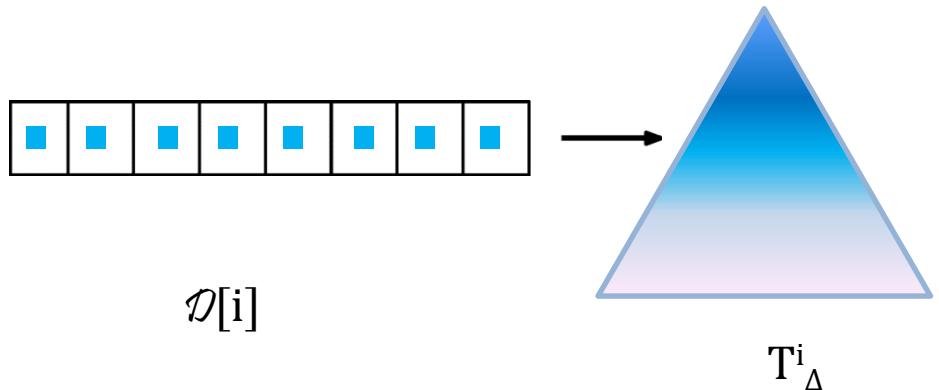
Ego-Trees Revisited

- ego-tree: optimal tree for a row (= given source)



Ego-Trees Revisited

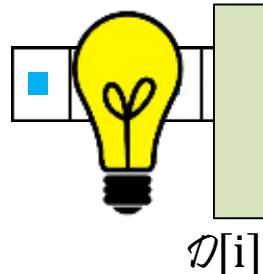
- ego-tree: optimal tree for a row (= given source)



Can we merge the trees **without distortion** and **keep degree low**?

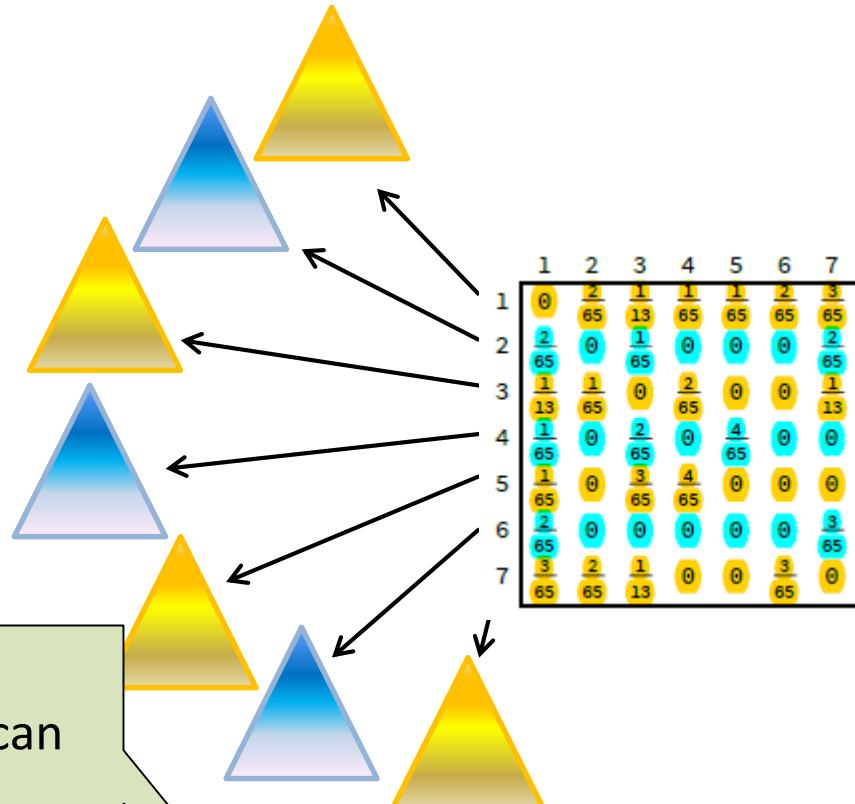
Ego-Trees Revisited

- ego-tree: optimal tree for a row (= given source)



For **sparse demands** yes:
enough *low-degree nodes* which can
serve as “**helper nodes**”!

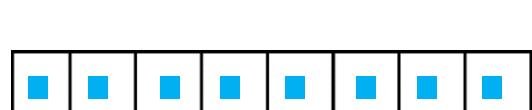
T_{Δ}^i



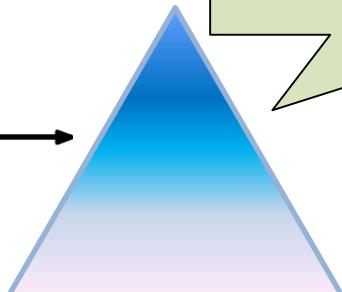
Can we merge the trees **without distortion** and **keep degree low**?

Ego-Trees Revisited

- ego-tree: optimal tree for a row (= given source)

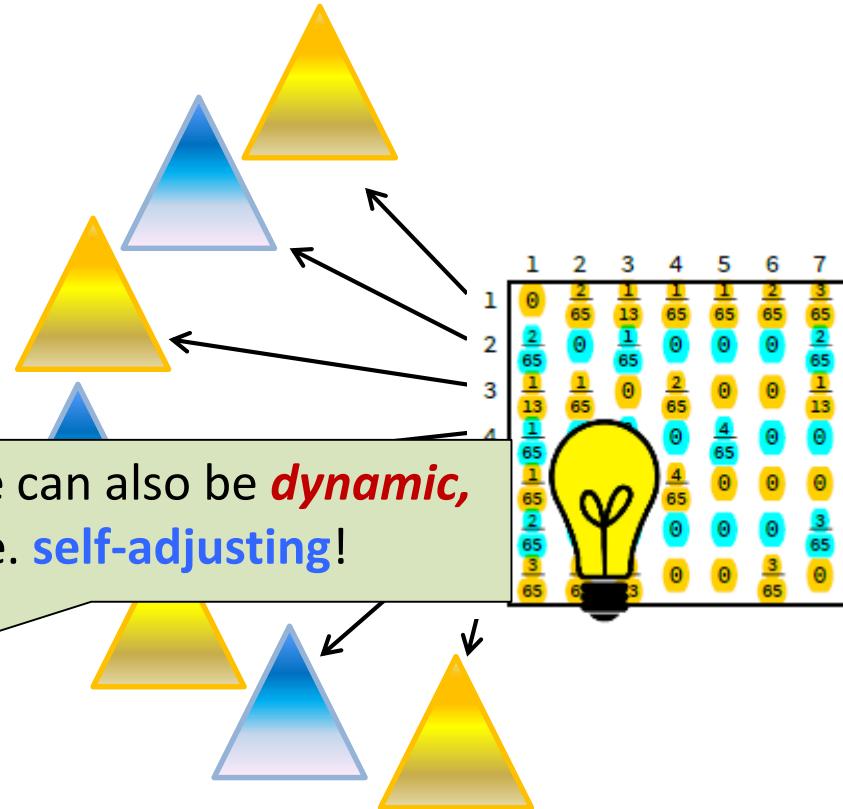


$\mathcal{D}[i]$



T_{Δ}^i

Ego-tree can also be *dynamic*,
i.e. *self-adjusting*!

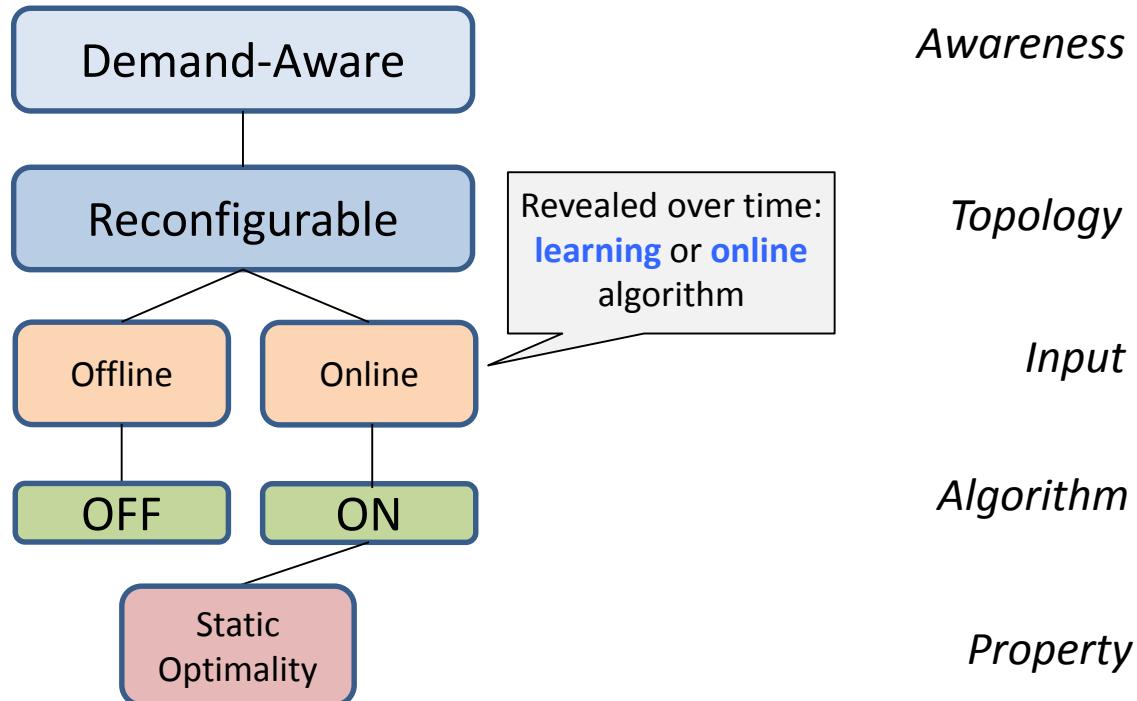


Can we merge the trees *without distortion* and *keep degree low*?

Other metrics for self-adjusting networks?

A Taxonomy: Reconfigurable Networks

Static Optimality:
“Not worse than static
which knows demand
ahead of time!”
 $\rho = \text{Cost(ON)}/\text{Cost(STAT*)}$
is constant.



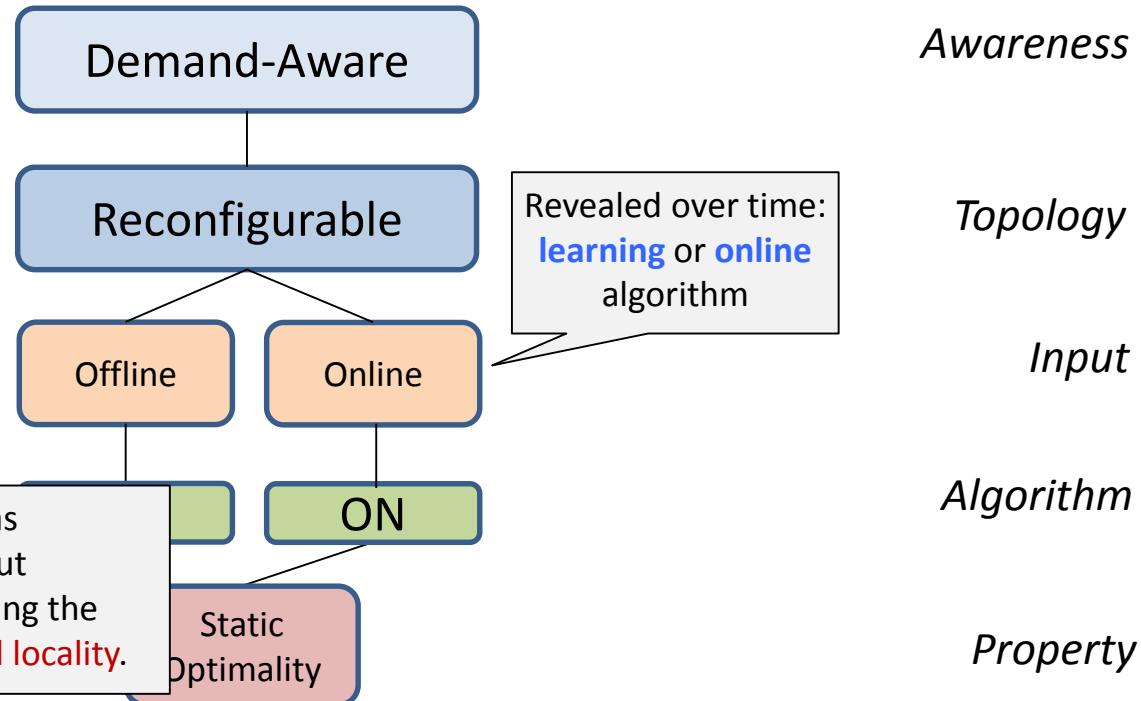
A Taxonomy: Reconfigurable Networks

Static Optimality:

"Not worse than static
which knows demand
ahead of time!"

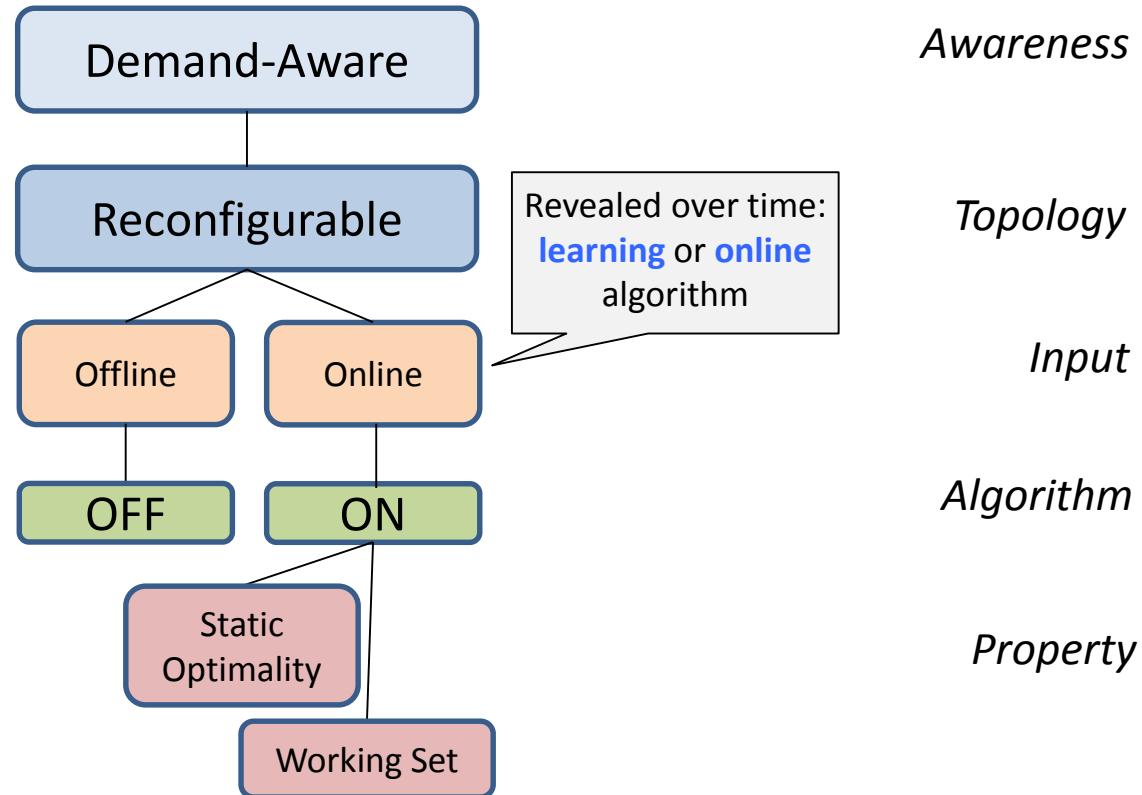
$\rho = \text{Cost(ON)}/\text{Cost(STAT*)}$
is constant.

Note: may be $<<1$. ON has
advantage of adjusting, but
the **disadvantage** of not knowing the
workload. E.g. if much **temporal locality**.



A Taxonomy: Reconfigurable Networks

Working Set Property:
“*Topological distance* between nodes
proportional to how recently they communicated!”



A Taxonomy: Reconfigurable Networks

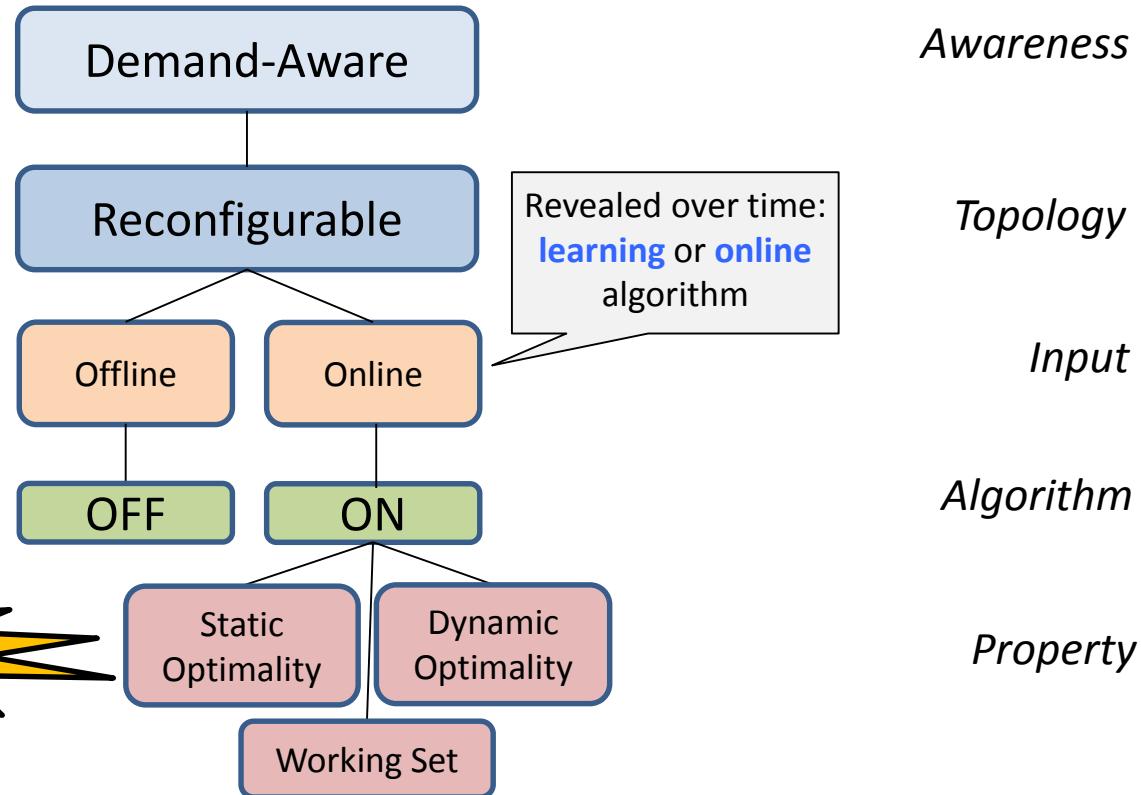
Dynamic Optimality:

“No worse than an offline algorithm which *knows the sequence!*”

$\rho = \text{Cost(ON)}/\text{Cost(OFF*)}$
is constant.

Always ≥ 1 .

The holy grail!



So: How *much* structure/entropy is there?



How to **measure** it?
How to distinguish between **temporal**
and **non-temporal** structure?
More **tricky**!

Often only intuitions in the literature...

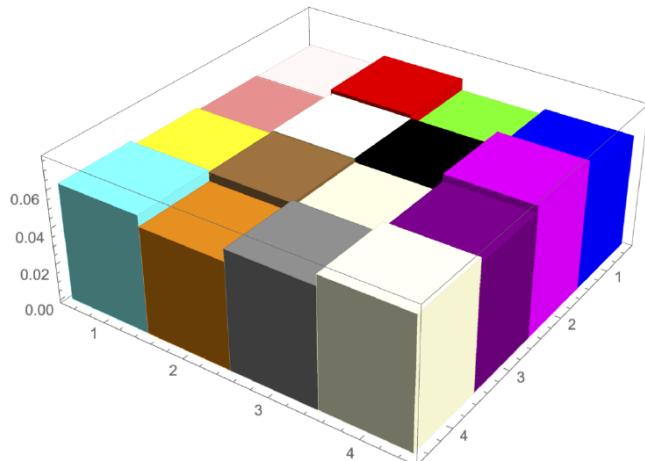
“less than 1% of the rack pairs account for 80% of the total traffic”

“only a few ToRs switches are hot and most of their traffic goes to a few other ToRs”

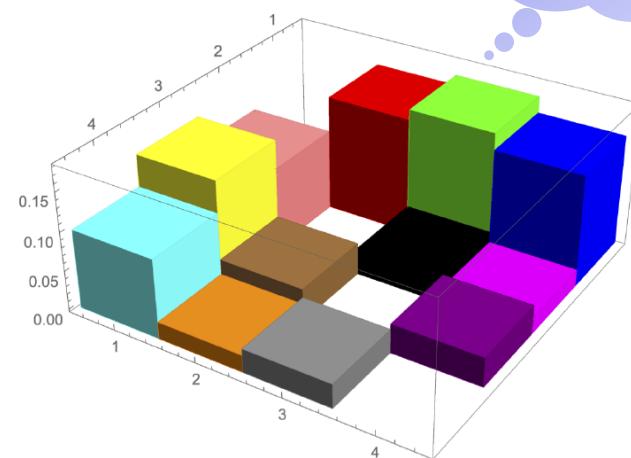
“over 90% bytes flow in elephant flows”

... and it *is* intuitive!

Non-temporal Structure



VS



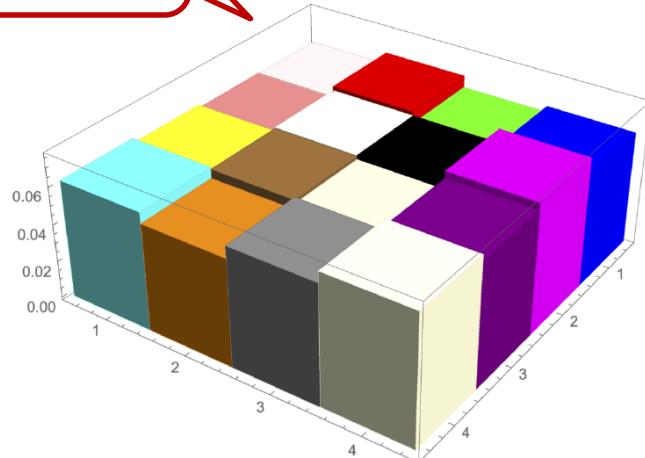
Color =
comm. pair

Traffic matrix of two different **distributed ML** applications (GPU-to-GPU):
Which one has *more structure*?

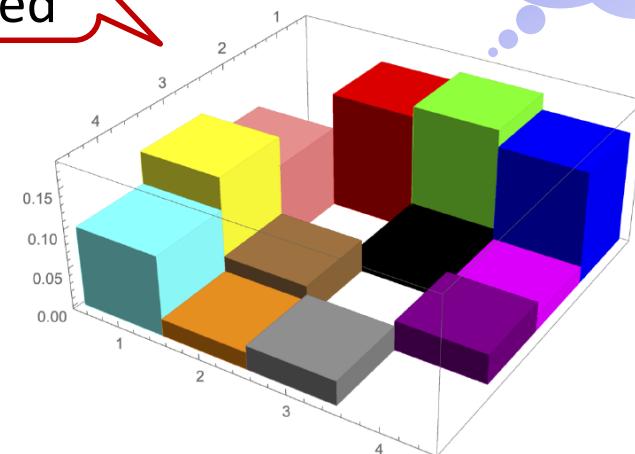
... and it *is* intuitive!

Non-temporal Structure

More uniform



More skewed

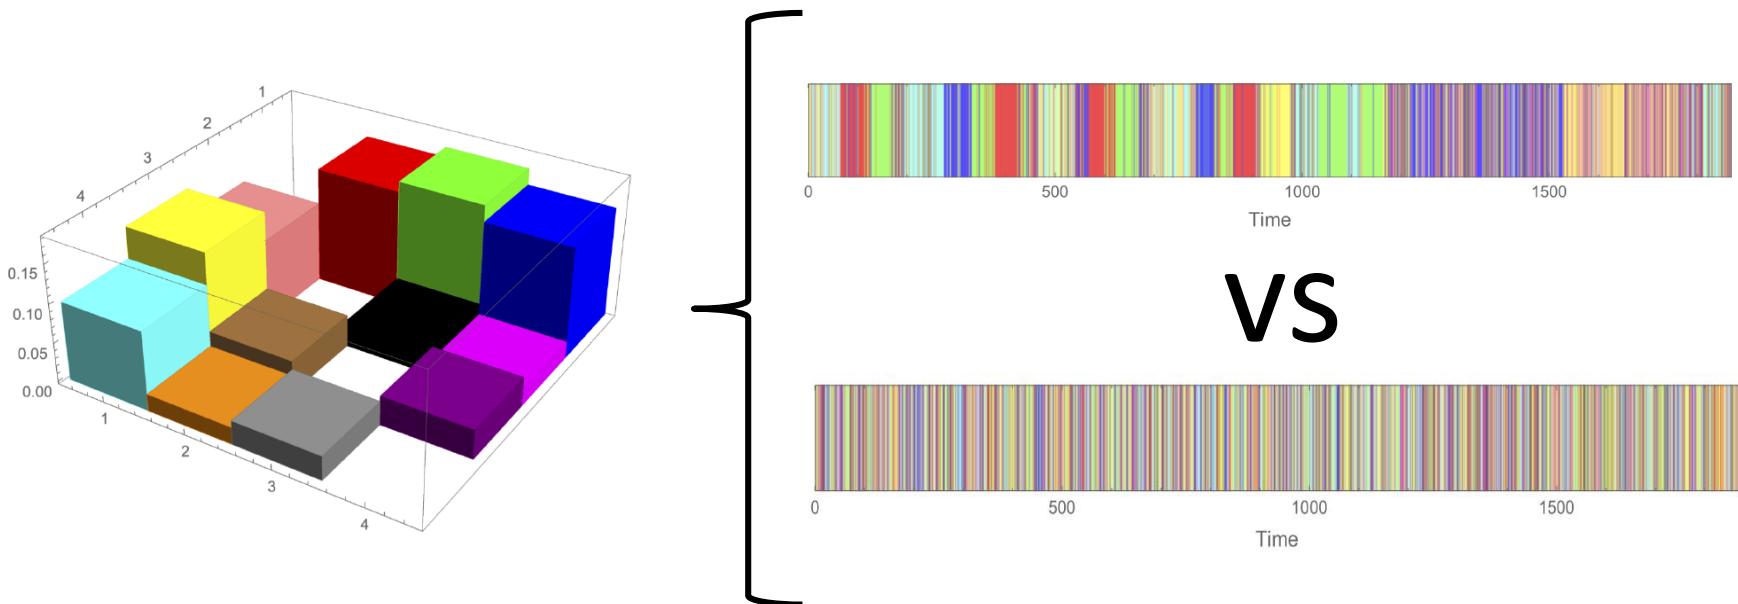


VS

Traffic matrix of two different **distributed ML** applications (GPU-to-GPU):
Which one has *more structure*?

... and it *is* intuitive!

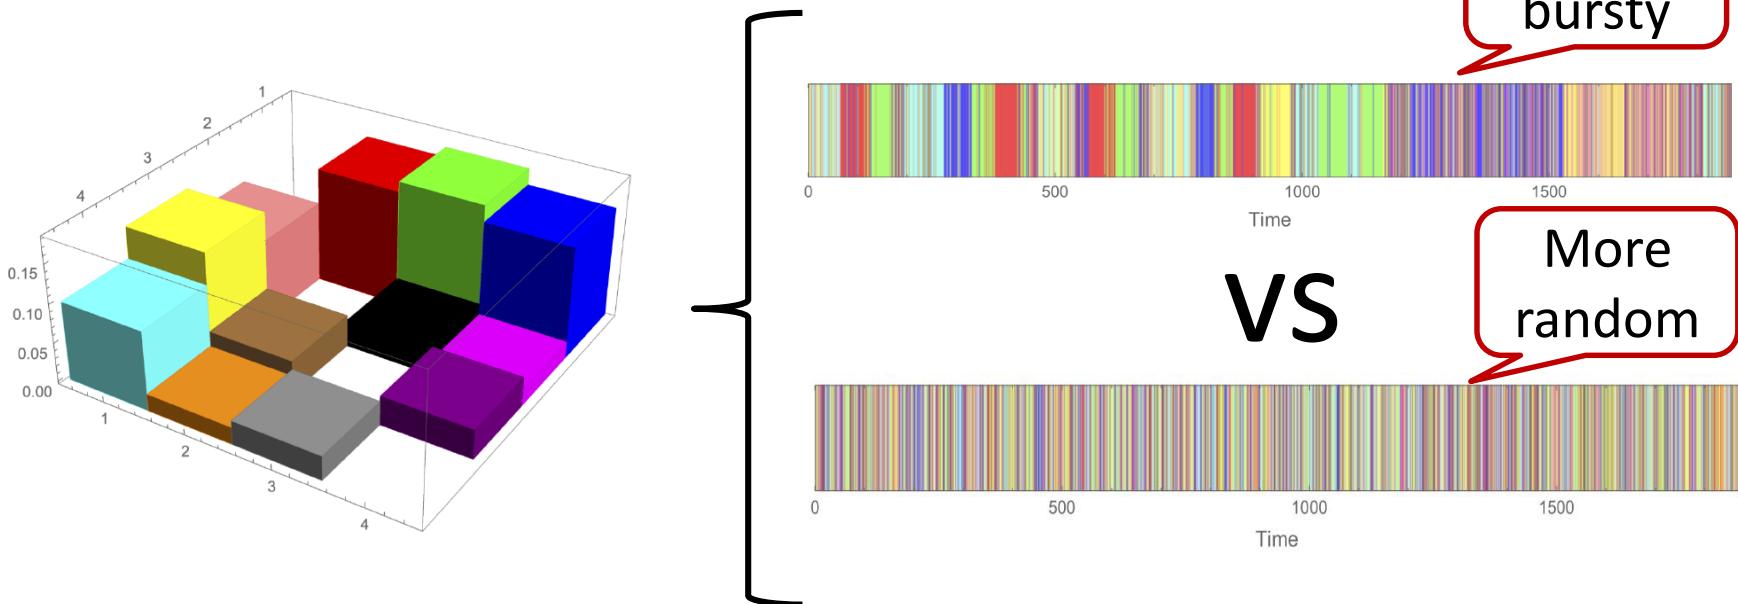
Temporal Structure



Two different ways to generate *same traffic matrix* (same non-temporal structure):
Which one has *more structure*?

... and it *is* intuitive!

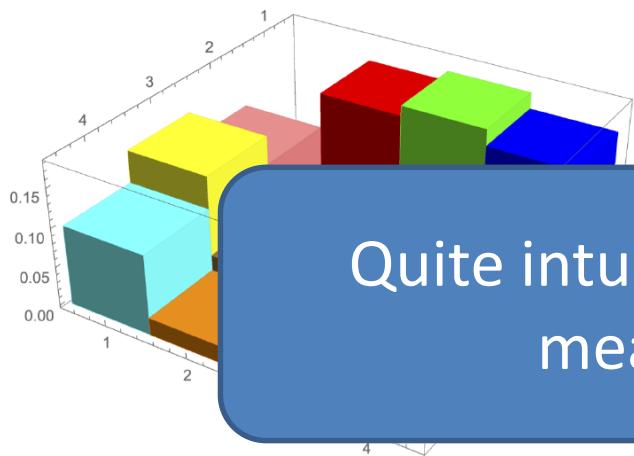
Temporal Structure



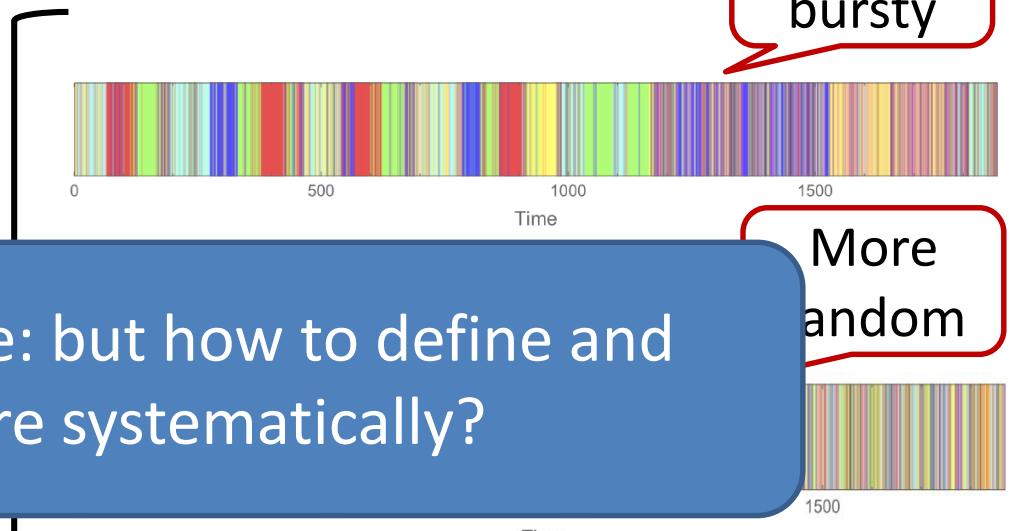
Two different ways to generate *same traffic matrix* (same non-temporal structure):
Which one has *more structure*?

... and it *is* intuitive!

Temporal Structure



Quite intuitive: but how to define and measure systematically?



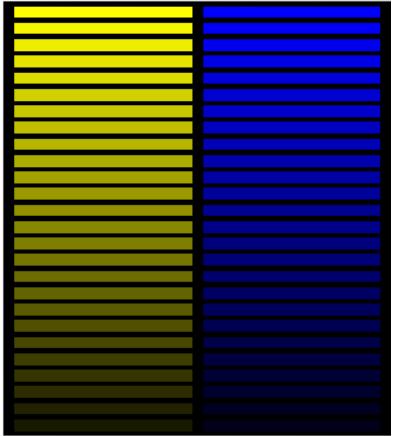
Two different ways to generate *same traffic matrix* (same non-temporal structure):
Which one has *more structure*?

The Trace Complexity

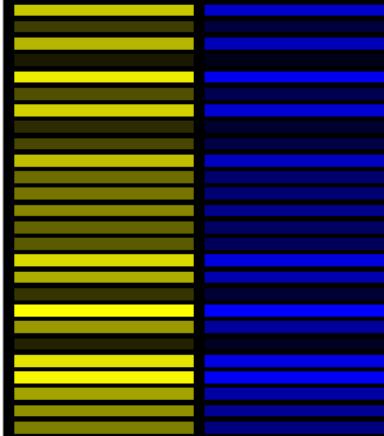
- An **information-theoretic**: what is the entropy (rate) of a traffic trace?
- Systematic „**shuffle&compress**“ methodology
 - Remove structure by iterative *randomization*
 - Difference of compression *before and after* randomization: structure

The Trace Complexity

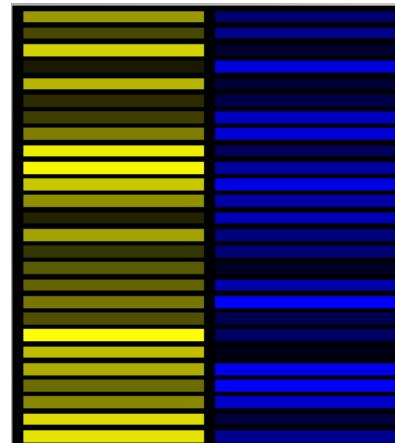
Original src-dst trace



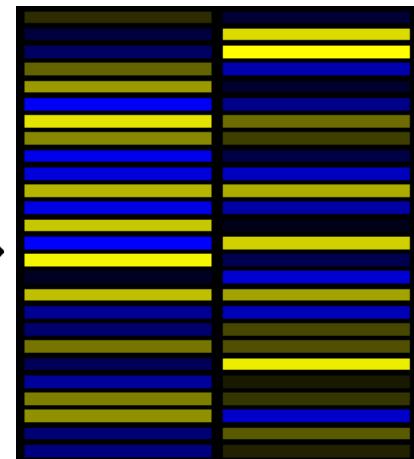
Randomize rows



Randomized columns



Uniform trace

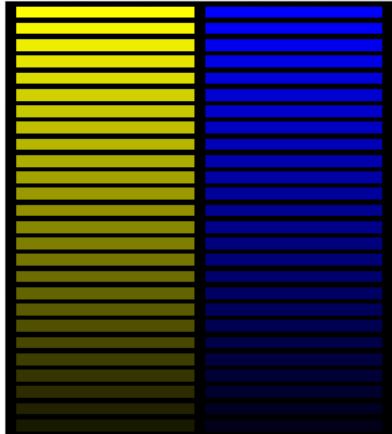


Increasing complexity (systematically randomized)

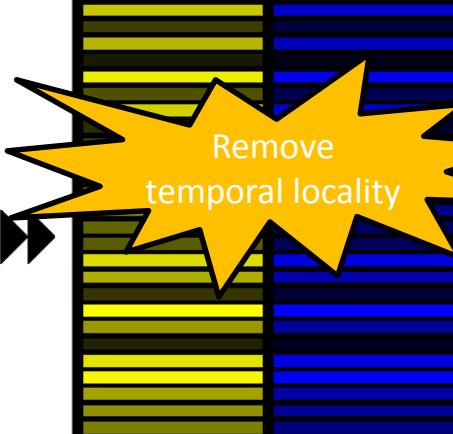
More structure (compresses better)

The Trace Complexity

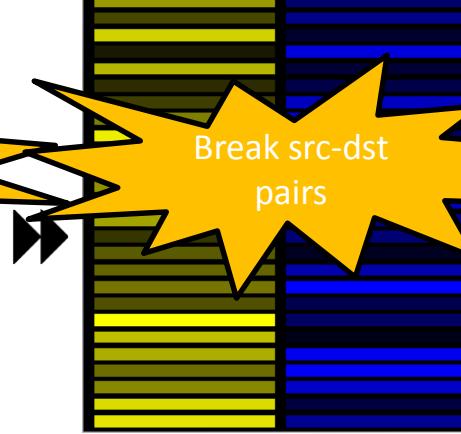
Original src-dst trace



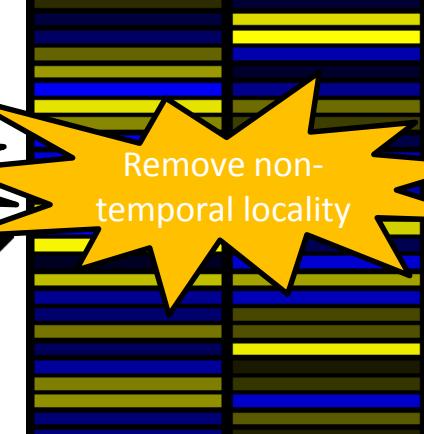
Randomize rows



Randomized columns



Uniform trace



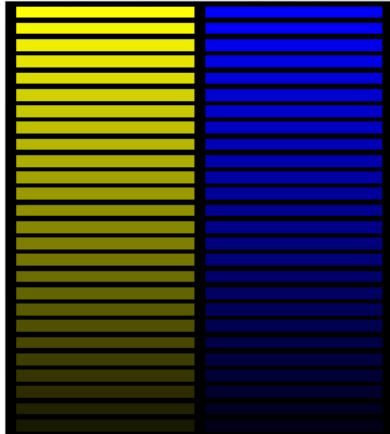
Difference in
compression?

Difference in
compression?

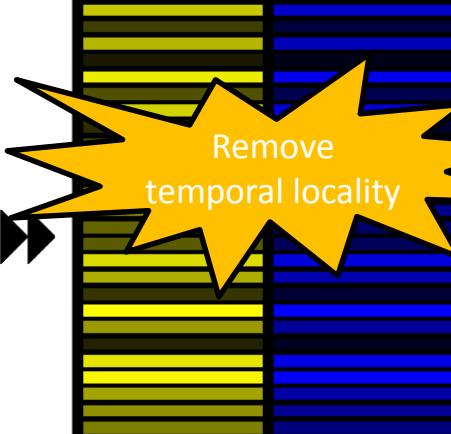
Difference in
compression?

The Trace Complexity

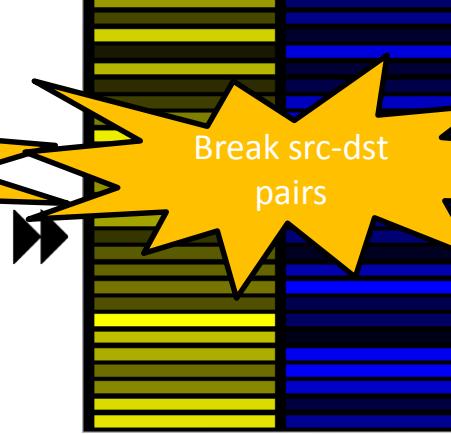
Original src-dst trace



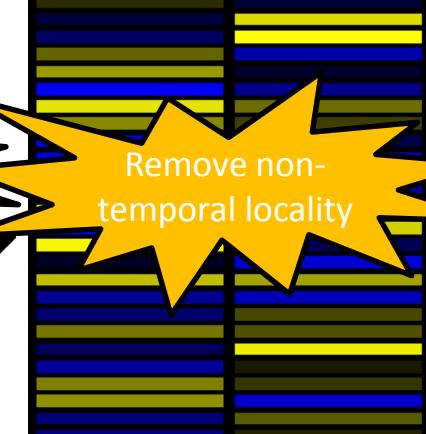
Randomize rows



Randomized columns



Uniform trace



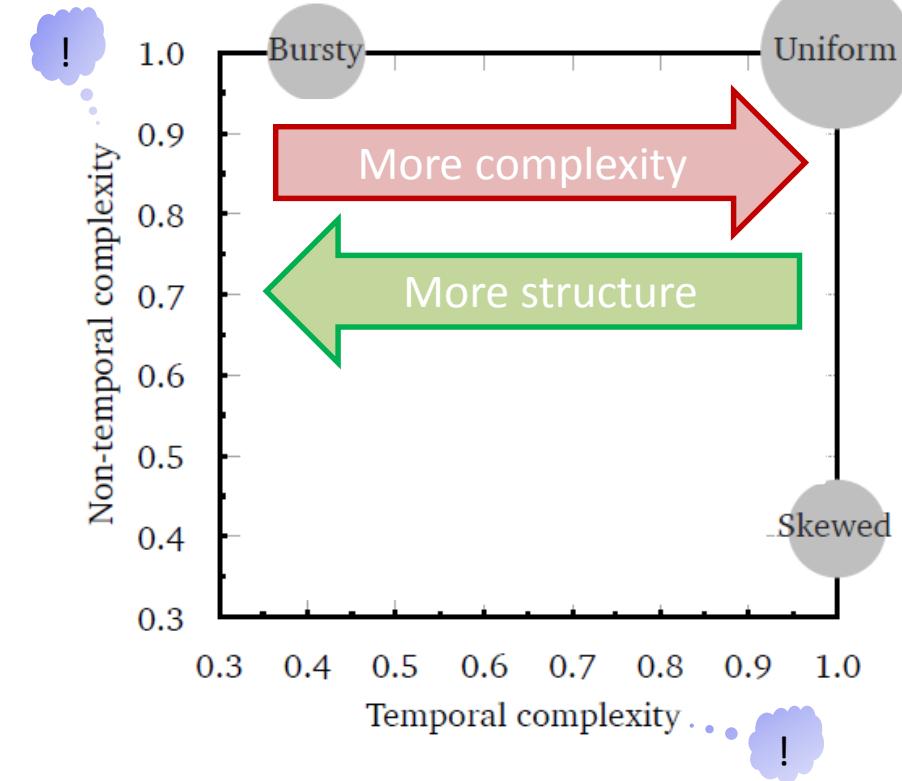
Difference in
compression?

Difference in
compression?

Difference in
compression?

Can be used to define a „complexity map“!

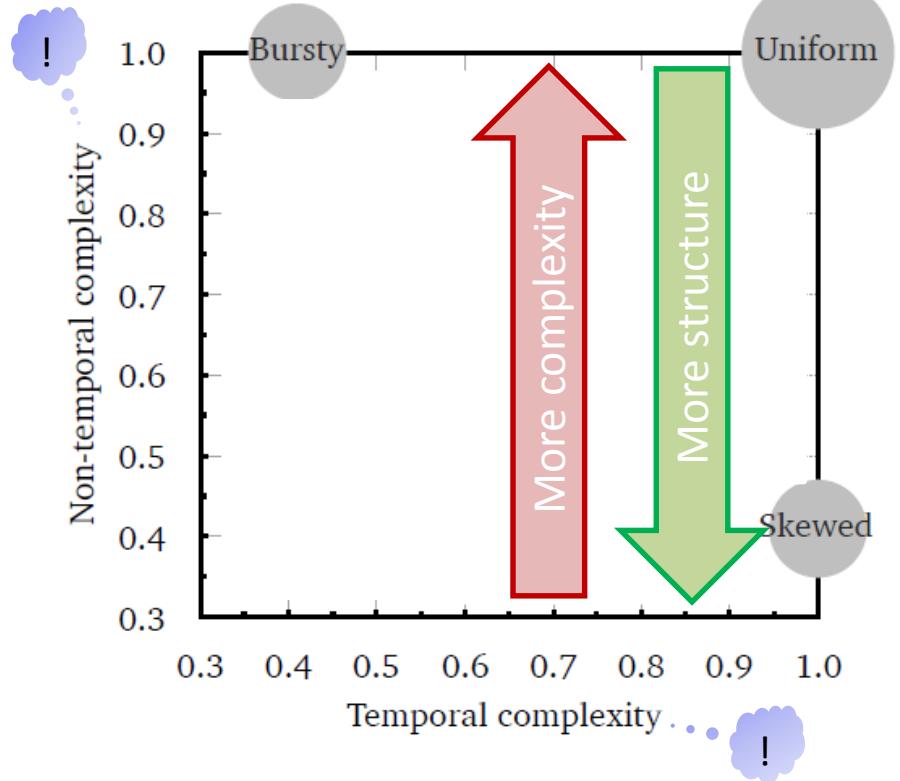
The Complexity Map



Complexity Map: Entropy („complexity“) of traffic traces.

Measuring the Complexity of Packet Traces.
Avin, Ghobadi, Griner, Schmid. ArXiv 2019.

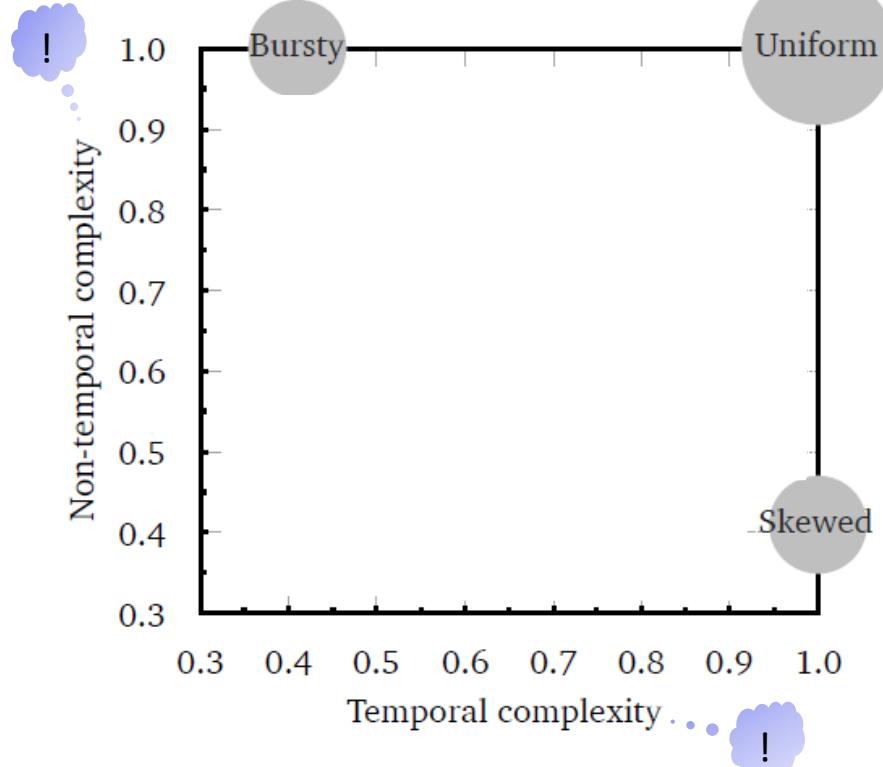
The Complexity Map



Complexity Map: Entropy („complexity“) of traffic traces.

Measuring the Complexity of Packet Traces.
Avin, Ghobadi, Griner, Schmid. ArXiv 2019.

The Complexity Map

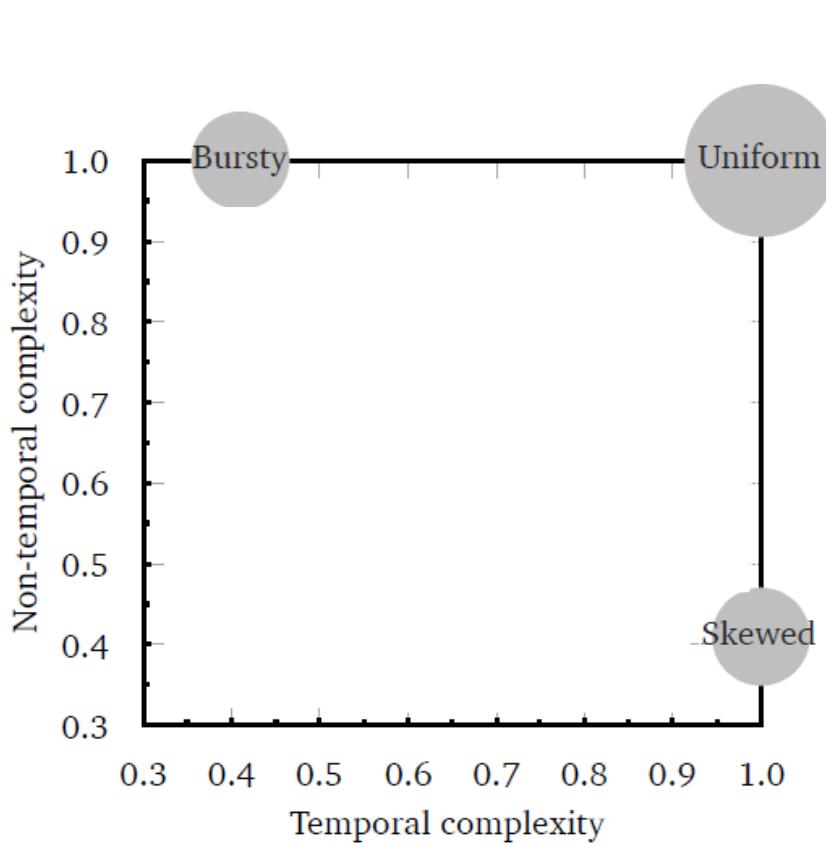


Complexity Map: Entropy
("complexity") of traffic traces.

Size = product
of entropy

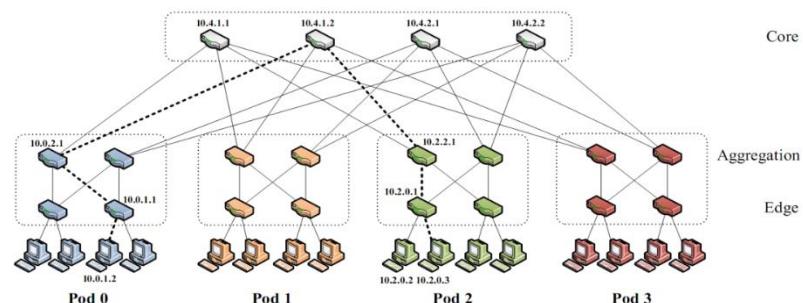
Measuring the Complexity of Packet Traces.
Avin, Ghobadi, Griner, Schmid. ArXiv 2019.

The Complexity Map

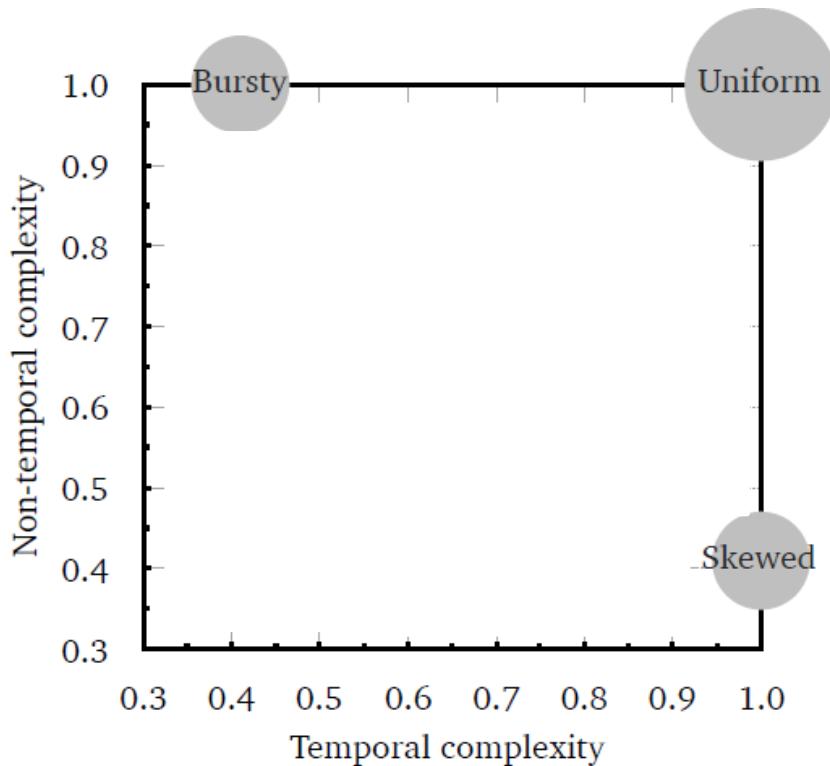


Uniform: Today's
datacenters

- Traditional networks are optimized **for the “worst-case” (all-to-all)** communication traffic)
- Example, fat-tree topologies: provide **full bisection bandwidth**

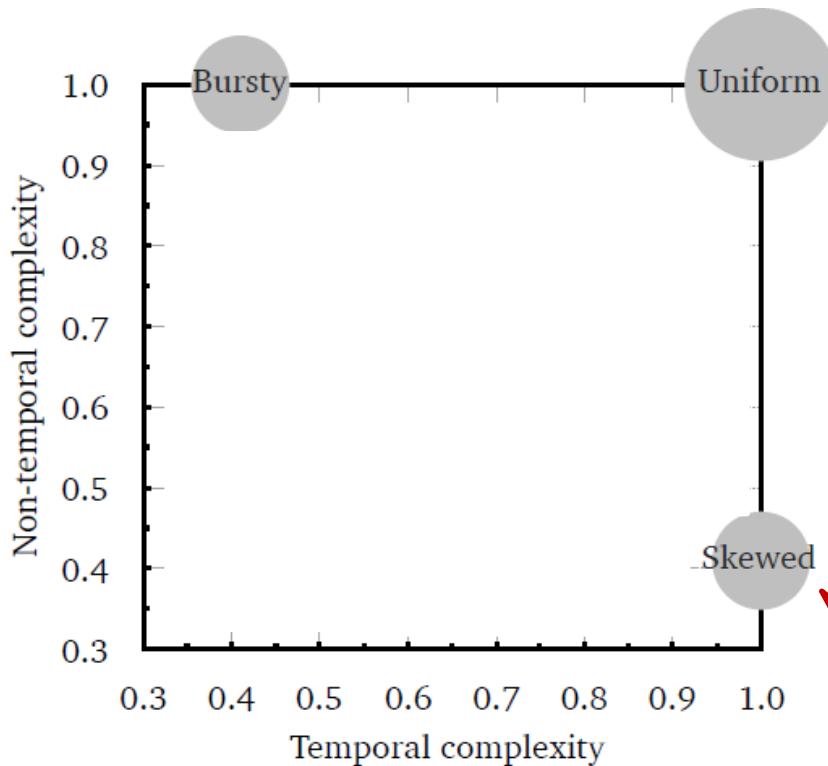


The Complexity Map



Good in the worst case ***but***:
cannot leverage different
temporal and **non-temporal**
structures of traffic traces!

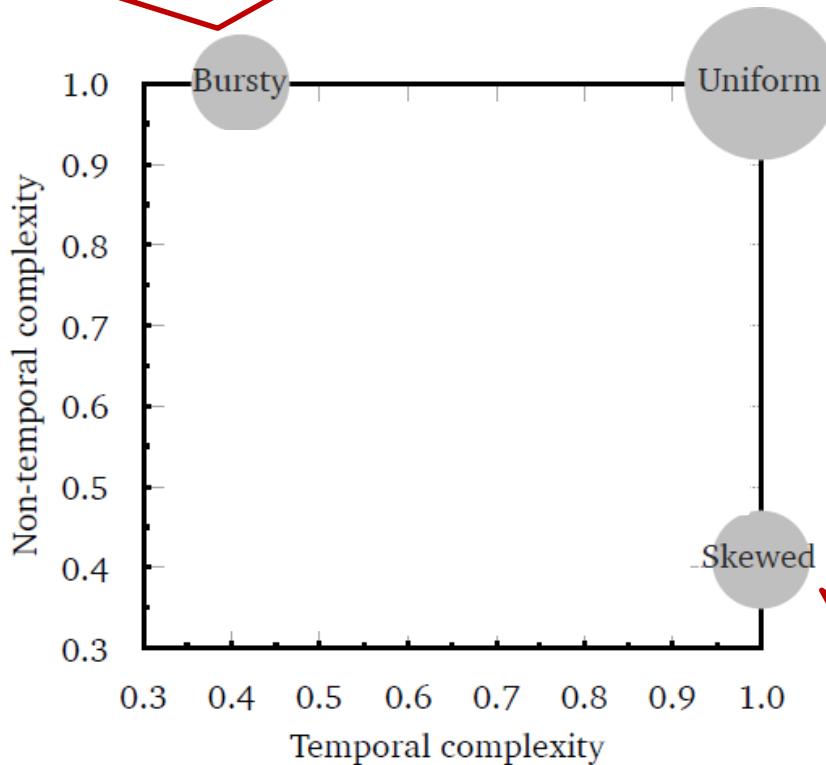
The Complexity Map



Good in the worst case ***but:***
cannot leverage different
temporal and **non-temporal**
structures of traffic traces!

Non-temporal structure could
be exploited already with ***static***
demand-aware networks!

To exploit **temporal** structure,
need ***adaptive demand-aware***
("self-adjusting") networks.

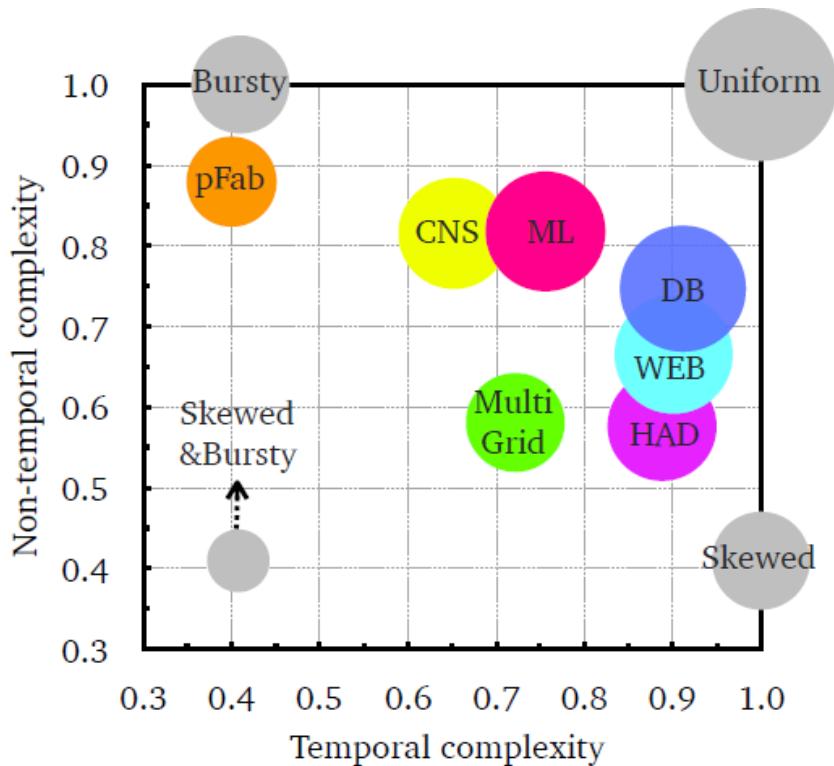


Complexity Map

Good in the worst case ***but:***
cannot leverage different
temporal and **non-temporal**
structures of traffic traces!

Non-temporal structure could
be exploited already with ***static***
demand-aware networks!

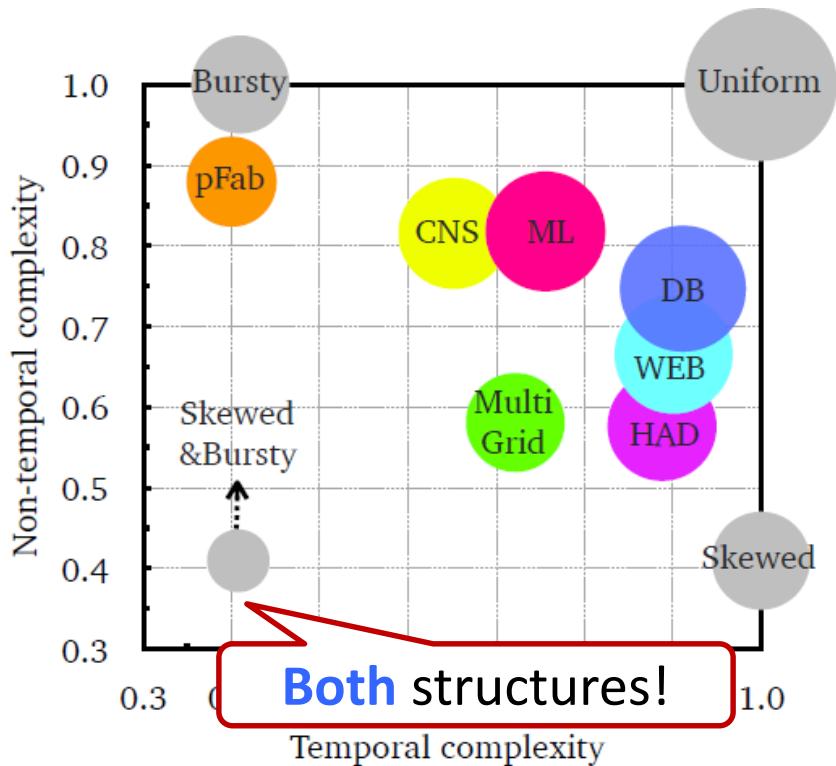
The Complexity Map



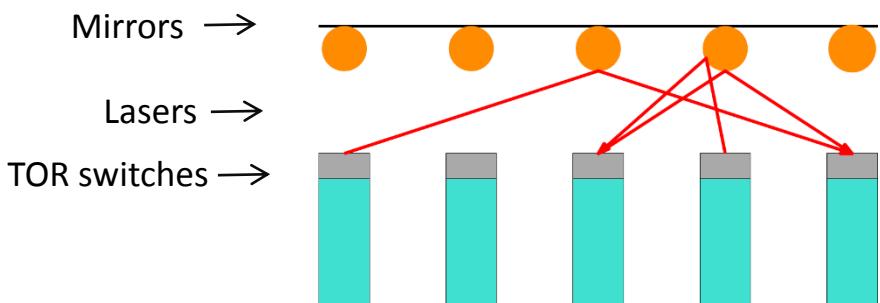
Observation: different applications feature quite significant (and different!) **temporal** and **non-temporal** structures.

- Facebook clusters: DB, WEB, HAD
- HPC workloads: CNS, Multigrid
- Distributed Machine Learning (ML)
- Synthetic traces like pFabric

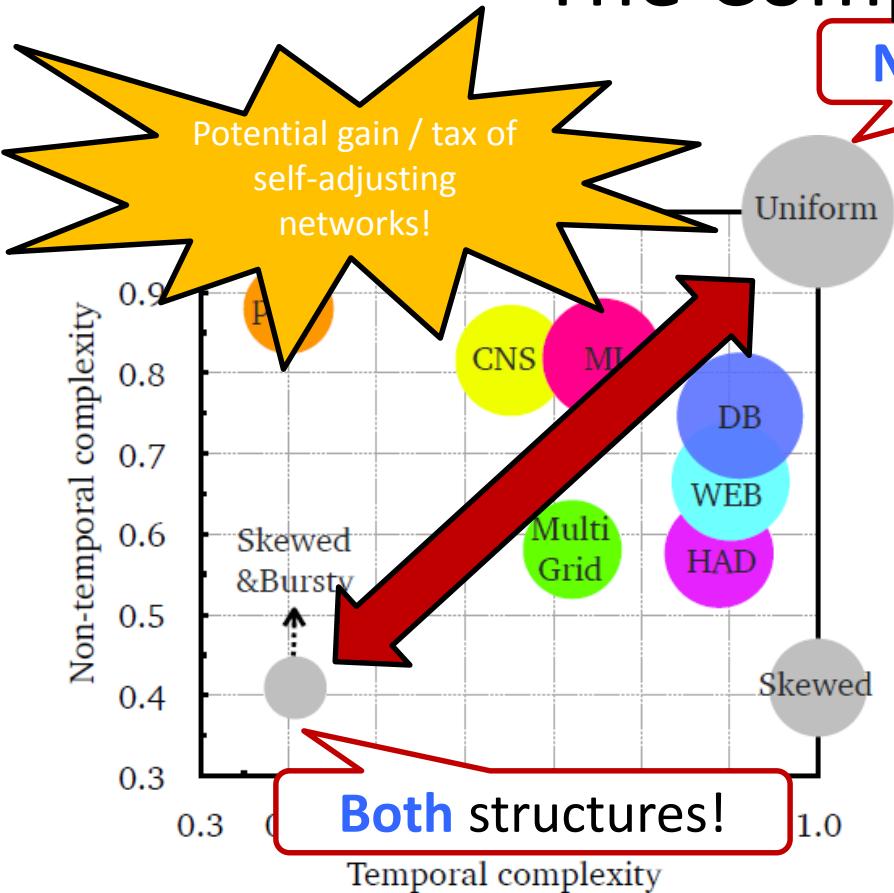
The Complexity Map



Goal: Design **self-adjusting networks** which leverage **both** dimensions of structure!

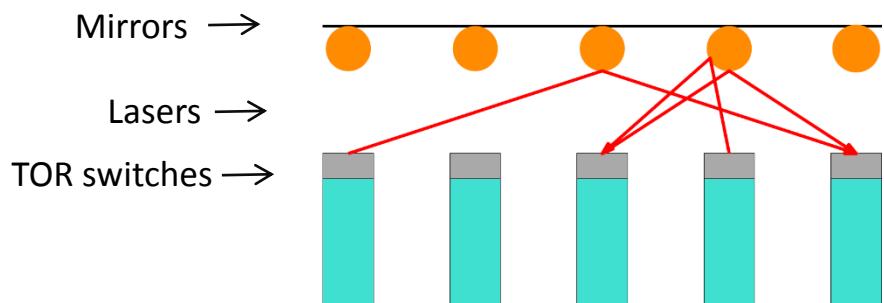


The Complexity Map



No structure!

Goal: Design **self-adjusting networks** which leverage **both** dimensions of structure!



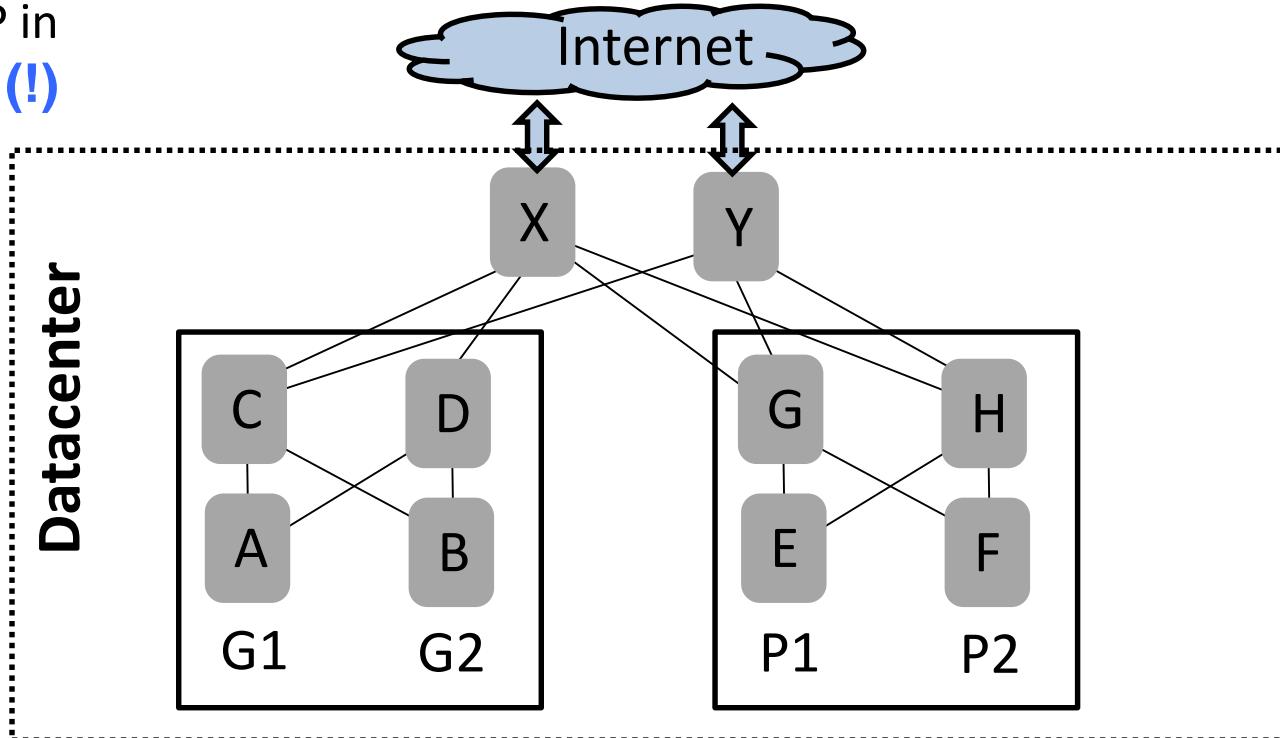
Roadmap

- Opportunities of self-* networks
 - Example 1: Demand-aware, self-adjusting networks
 - **Example 2: Self-repairing networks**
- Challenges of designing self-* networks



Reasoning About Failures is Hard

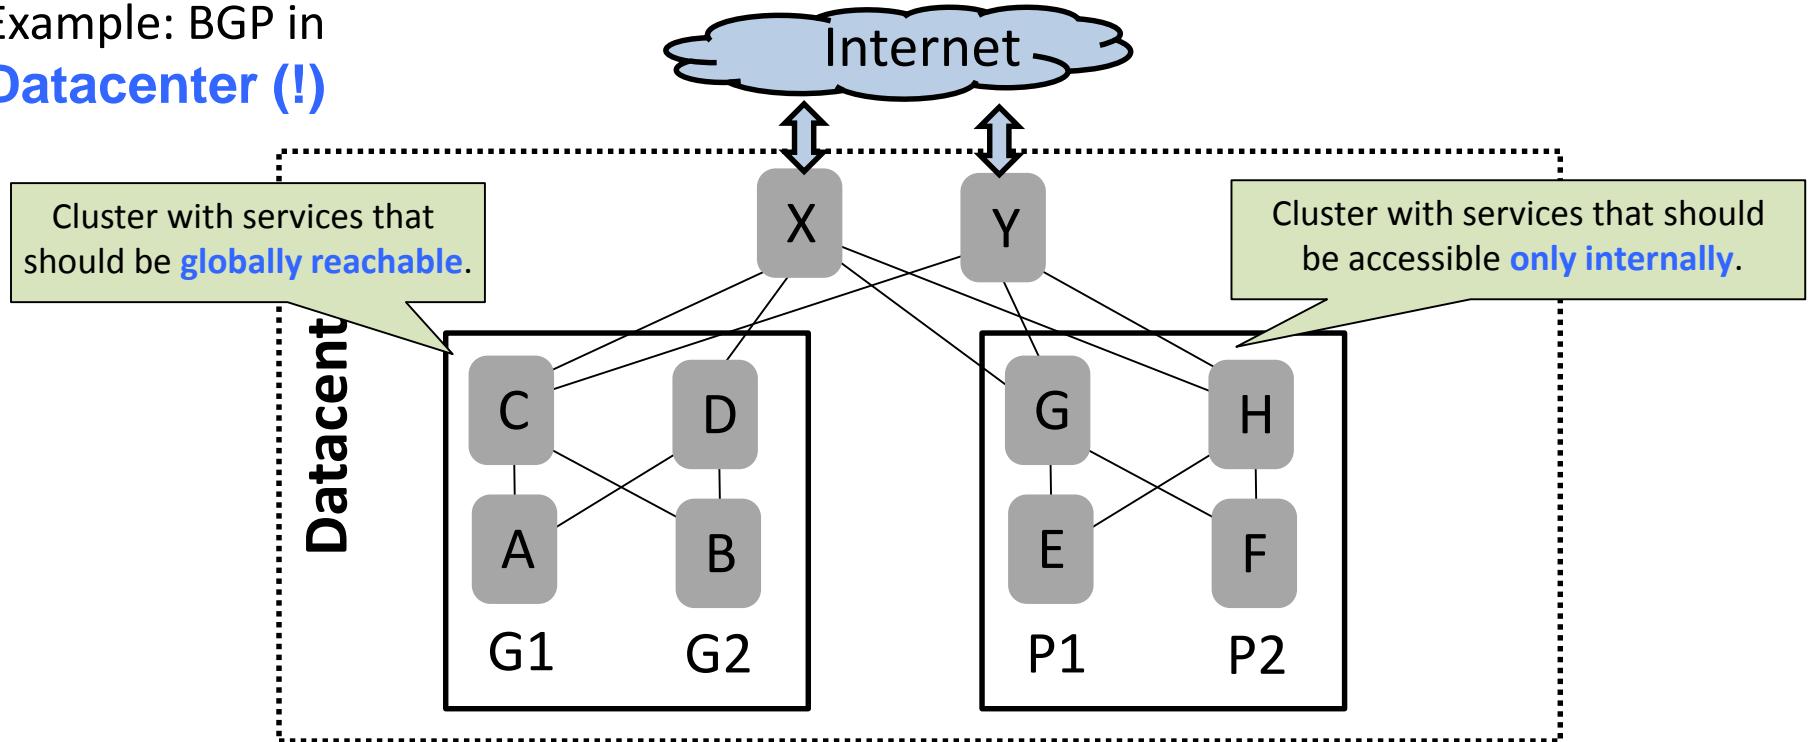
Example: BGP in
Datacenter (!)



Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

Reasoning About Failures is Hard

Example: BGP in
Datacenter (!)



Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

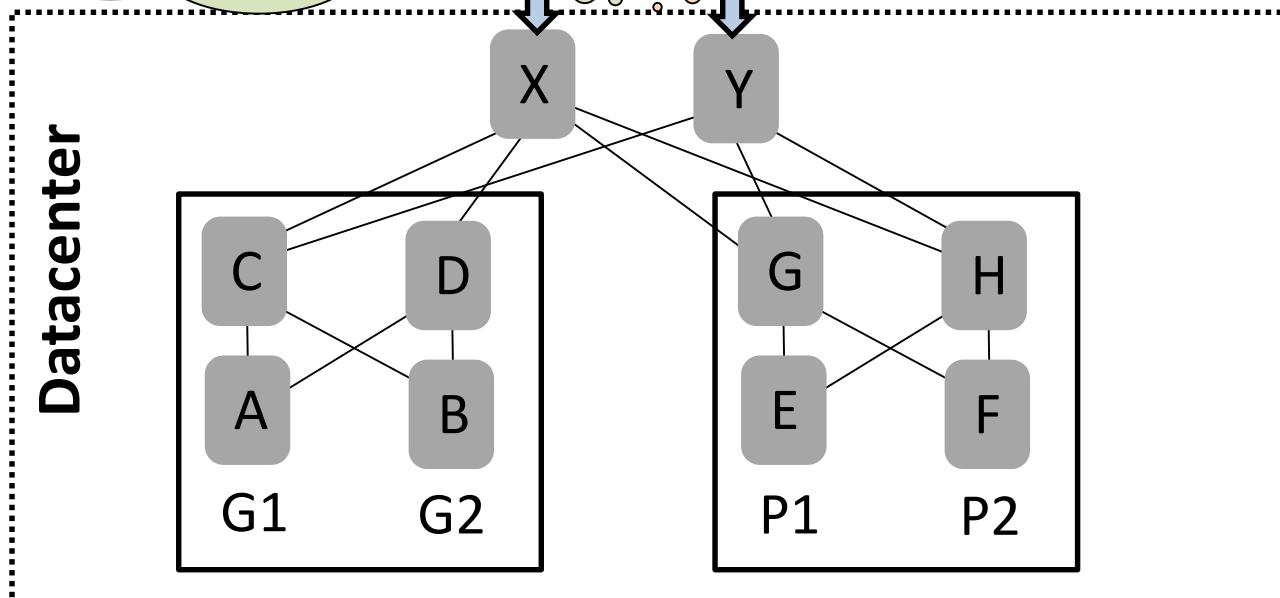
Reasoning About Failures is Hard

Example:

Datacenter

X and Y **announce** to Internet what is from G^* (prefix).

X and Y **block** what is from P^* .



Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

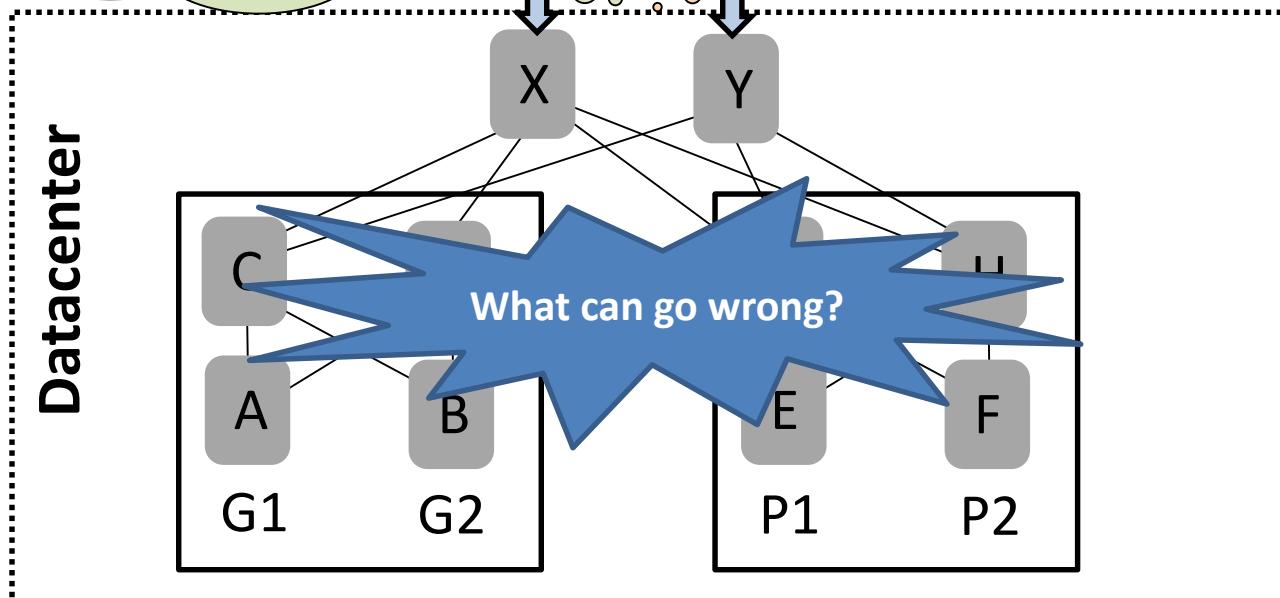
Reasoning About Failures is Hard

Example:

Datacenter

X and Y **announce** to Internet what is from G^* (prefix).

X and Y **block** what is from P^* .



Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

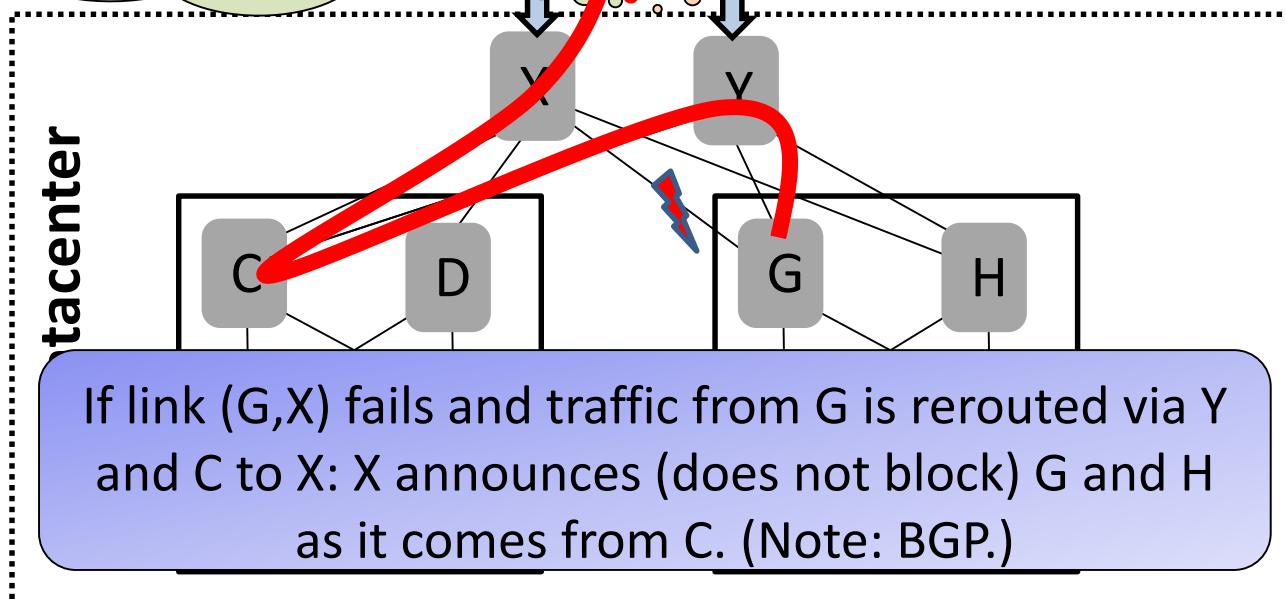
Reasoning About Failures is Hard

Example:

Datacenter

X and Y **announce** to Internet what is from G^* (prefix).

X and Y **block** what is from P^* .



Credits: Beckett et al. (SIGCOMM 2016): Bridging Network-wide Objectives and Device-level Configurations.

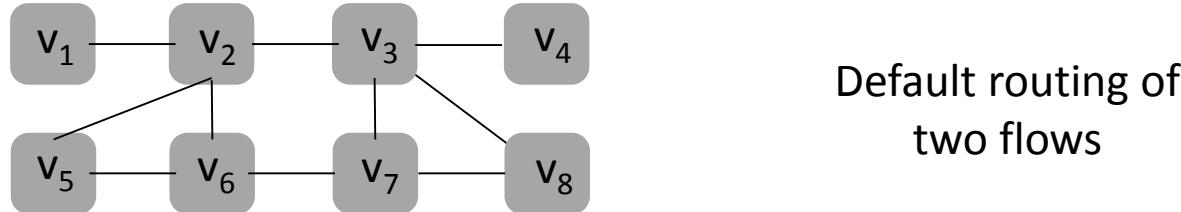
Managing Complex Networks is Hard for Humans



Another Case for Automation!

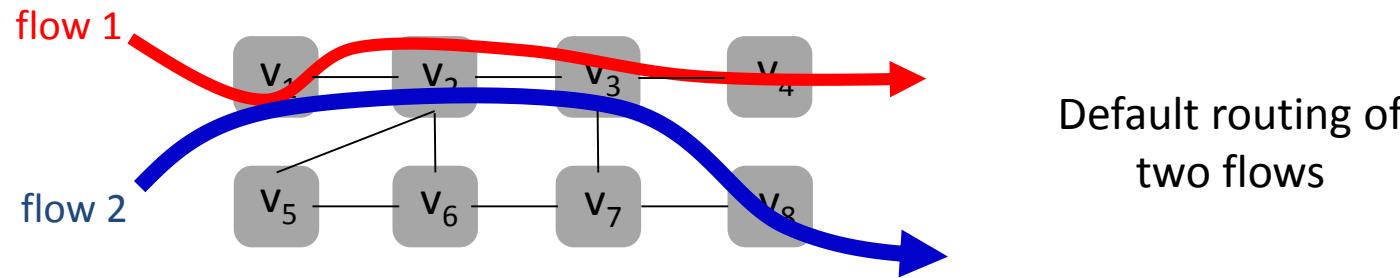
Case Study: Self-Repairing MPLS Networks

- MPLS: forwarding based on **top label** of label **stack**



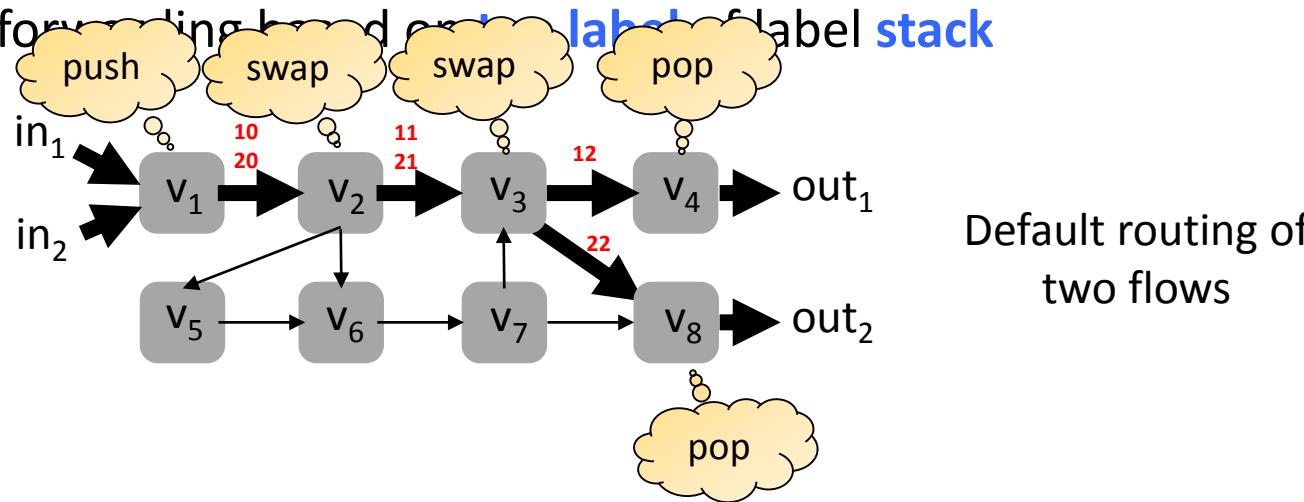
Case Study: Self-Repairing MPLS Networks

- MPLS: forwarding based on **top label** of label **stack**



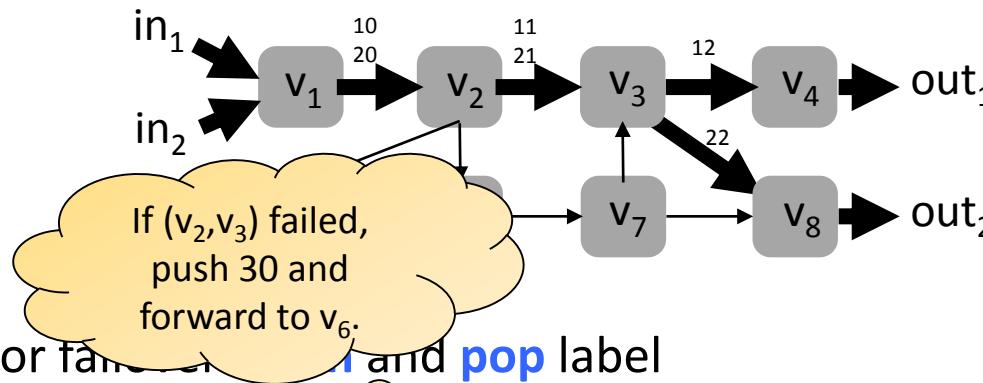
Case Study: Self-Repairing MPLS Networks

- MPLS: forwarding based on **label stack**

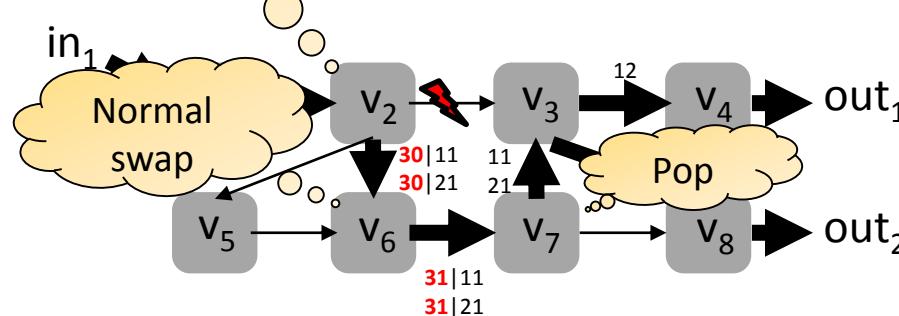


Fast Reroute Around 1 Failure

- MPLS: forwarding based on **top label** of label **stack**



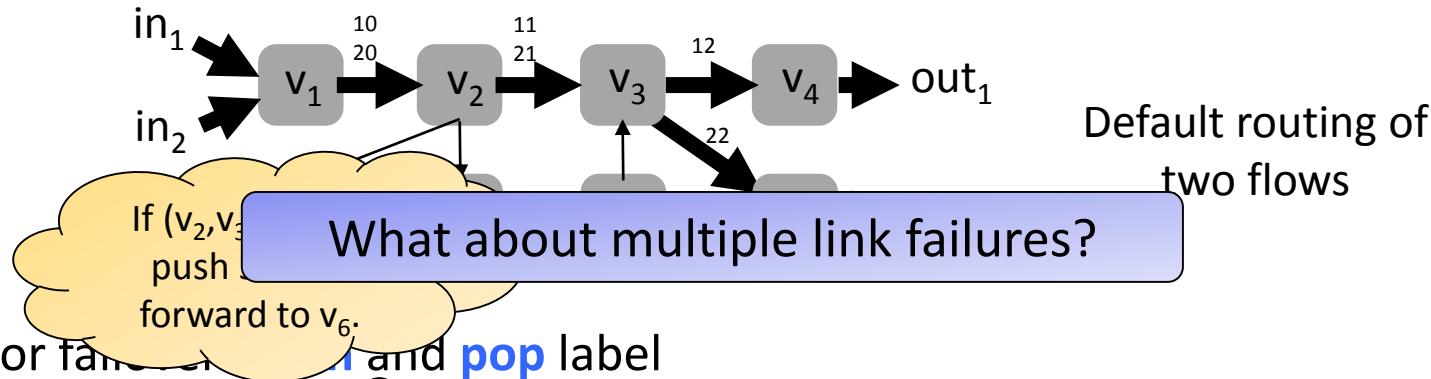
- For failure, **swap** and **pop** label



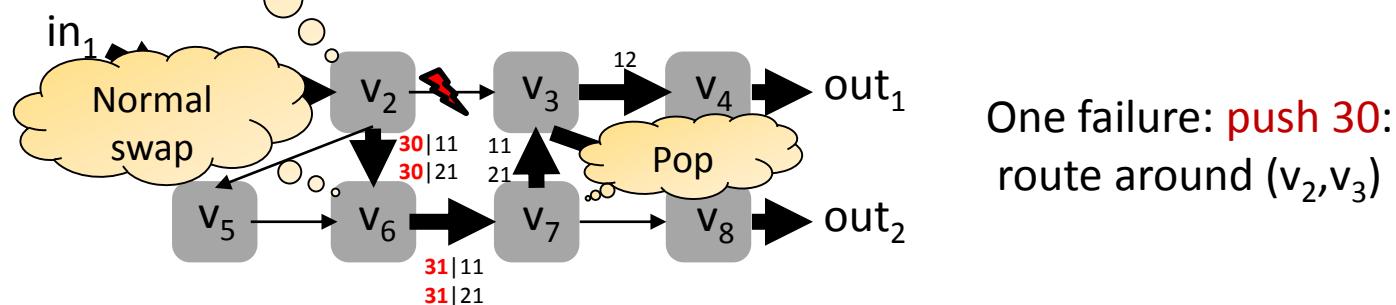
One failure: **push 30:** route around (v_2, v_3)

Fast Reroute Around 1 Failure

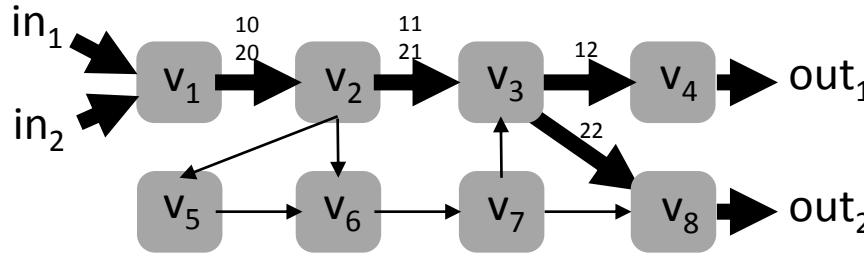
- MPLS: forwarding based on **top label** of label **stack**



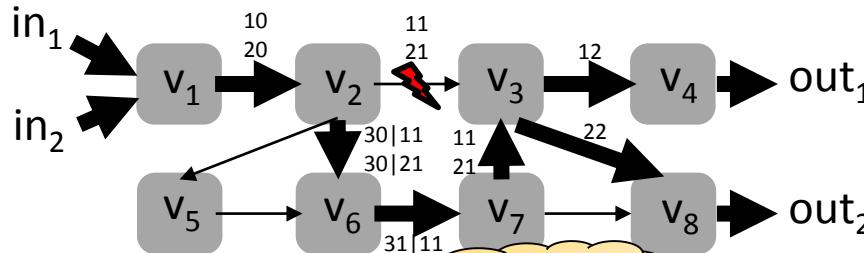
- For failure, **swap** and **pop** label



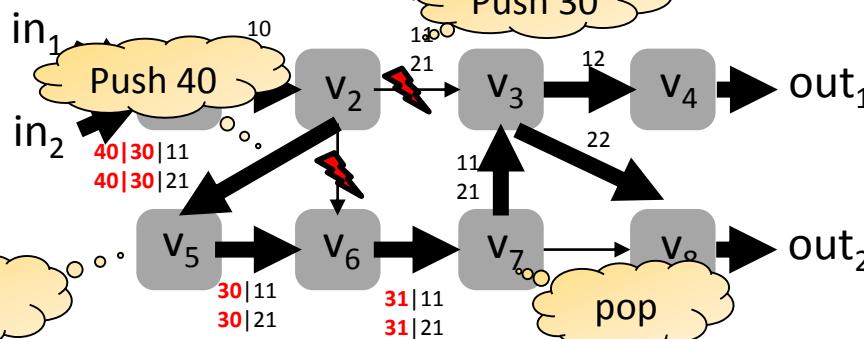
2 Failures: Push *Recursively*



Original Routing



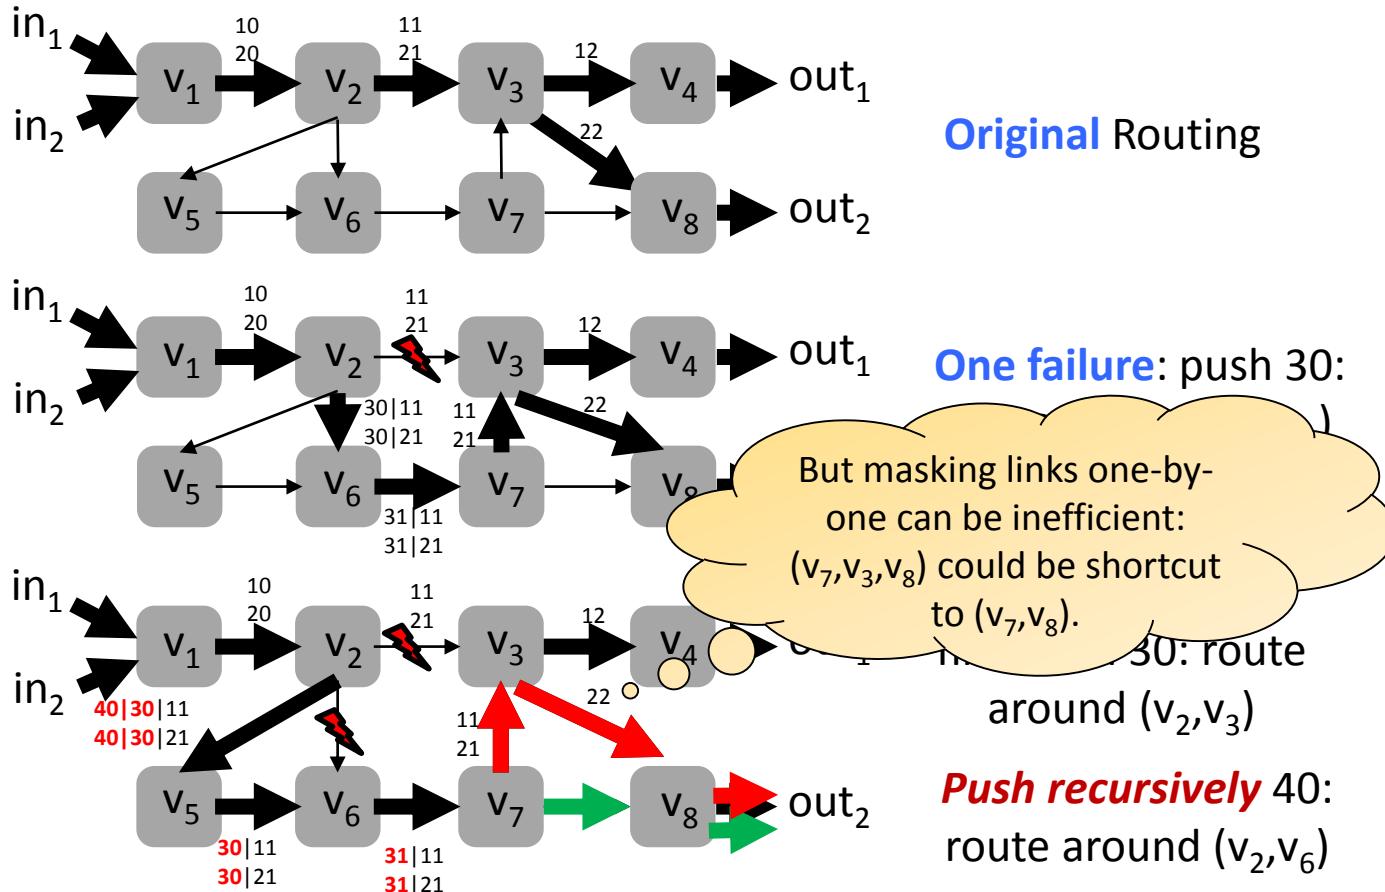
One failure: push 30:
route around (v_2, v_3)



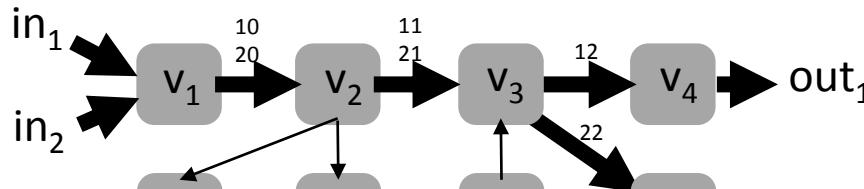
Two failures:
first push 30: route
around (v_2, v_3)

Push recursively 40:
route around (v_2, v_6)

2 Failures: Push *Recursively*

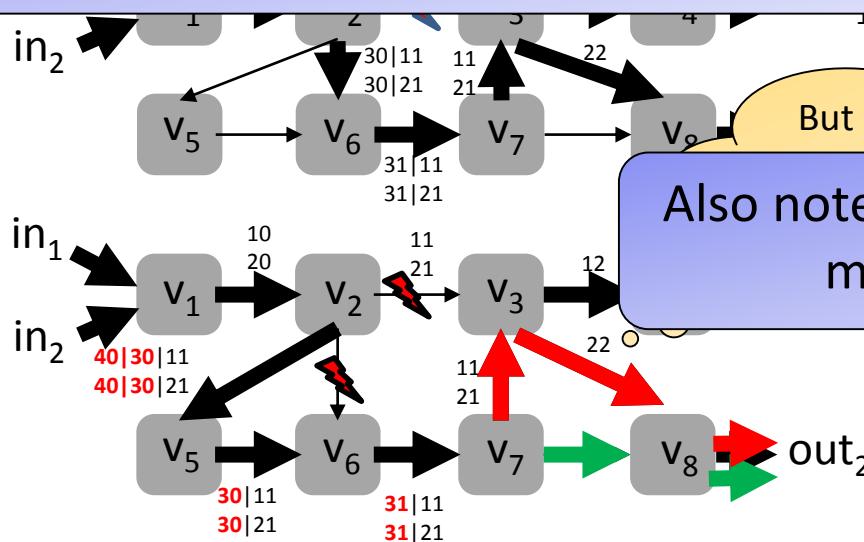


2 Failures: Push *Recursively*



Original Routing

More efficient but also more complex:
Cisco does **not recommend** using this option!



One failure: push 30:

But masking links one-by-

Also note: due to push, **header size**
may grow arbitrarily!

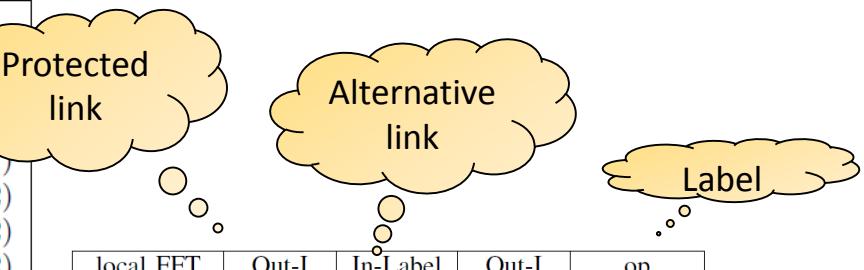
around (v_2, v_3)

Push recursively 40:
route around (v_2, v_6)

Reasoning About Low-Level Rules is Hard

FT	In-I	In-Label	Out-I	op
τ_{v_1}	in_1	\perp	(v_1, v_2)	$push(1)$
	in_2	\perp	(v_1, v_2)	$push(1)$
τ_{v_2}	(v_1, v_2)	10	(v_2, v_3)	$swap(1)$
	(v_1, v_2)	20	(v_2, v_3)	$swap(21)$
τ_{v_3}	(v_2, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_2, v_3)	21	(v_3, v_8)	$swap(22)$
τ_{v_4}	(v_7, v_3)	11	(v_3, v_4)	$swap(12)$
	(v_7, v_3)	21	(v_3, v_8)	$swap(22)$
τ_{v_5}	(v_3, v_4)	12	out_1	pop
	(v_2, v_5)	40	(\dots, \dots)	pop
τ_{v_6}	(v_5, v_6)	71	(v_6, v_7)	$push(1)$
	(v_6, v_7)	31	(v_7, v_3)	$push(31)$
τ_{v_7}	(v_6, v_7)	62	(v_7, v_3)	pop
	(v_6, v_7)	72	(v_7, v_8)	$swap(11)$
τ_{v_8}	(v_3, v_8)	22	out_2	$swap(22)$
	(v_7, v_8)	22	out_2	pop
				pop

Flow Table



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_6)	30	(v_2, v_5)	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
τ'_{v_2}	(v_2, v_3)	11	(v_2, v_6)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$

Failover Tables

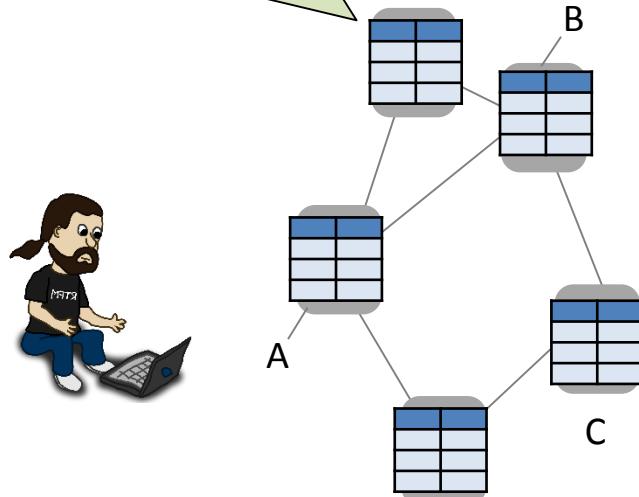
Tables for our example

MPLS Tunnels in Today's ISP Networks

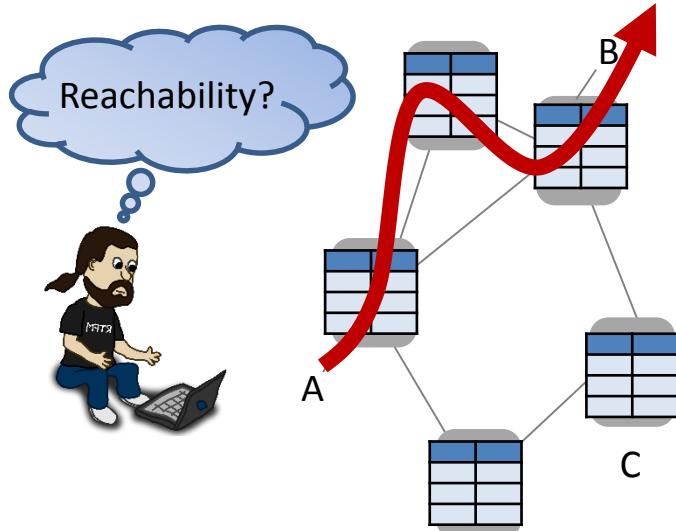


Responsibilities of a Sysadmin

Routers and switches store list of **forwarding rules**, and conditional **failover rules**.



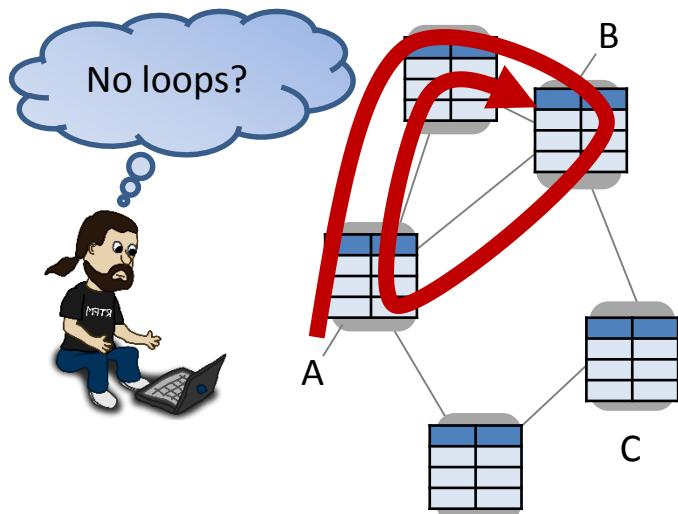
Responsibilities of a Sysadmin



Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?

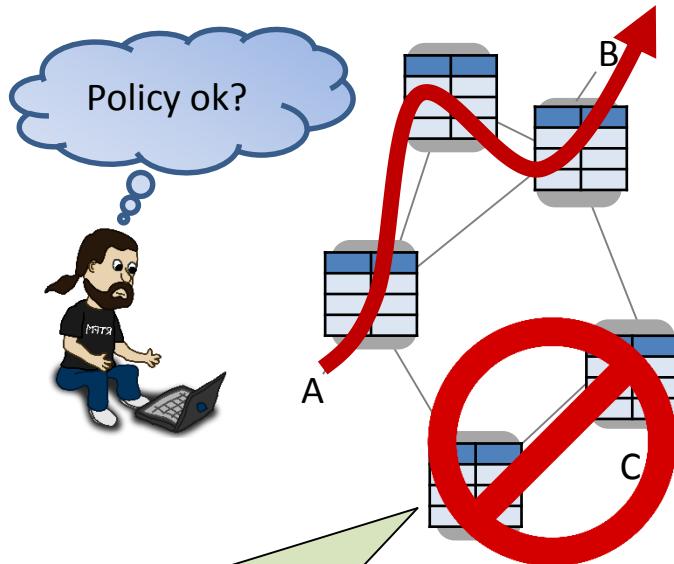
Responsibilities of a Sysadmin



Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?

Responsibilities of a Sysadmin

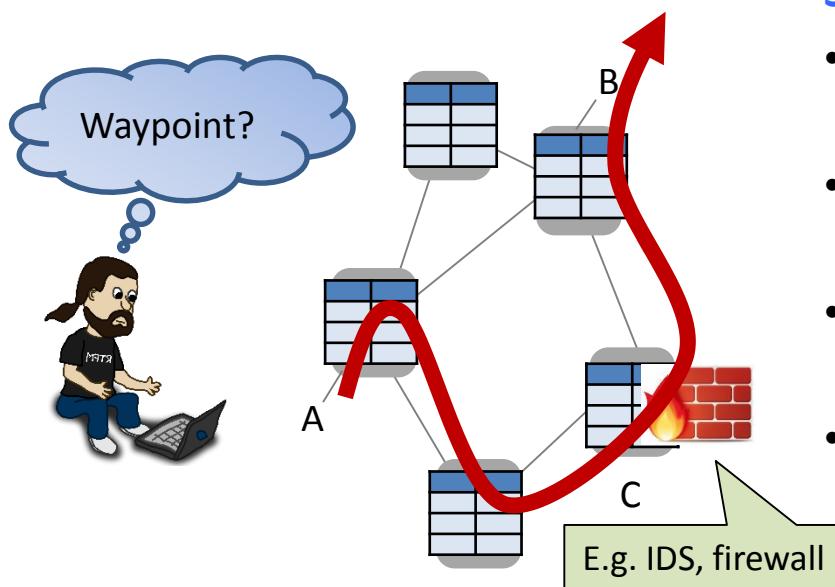


E.g. **NORDUnet**: no traffic via Iceland (expensive!). Or no traffic through *route reflectors*.

Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?
- **Policy:** Is it ensured that traffic from A to B never goes via C?

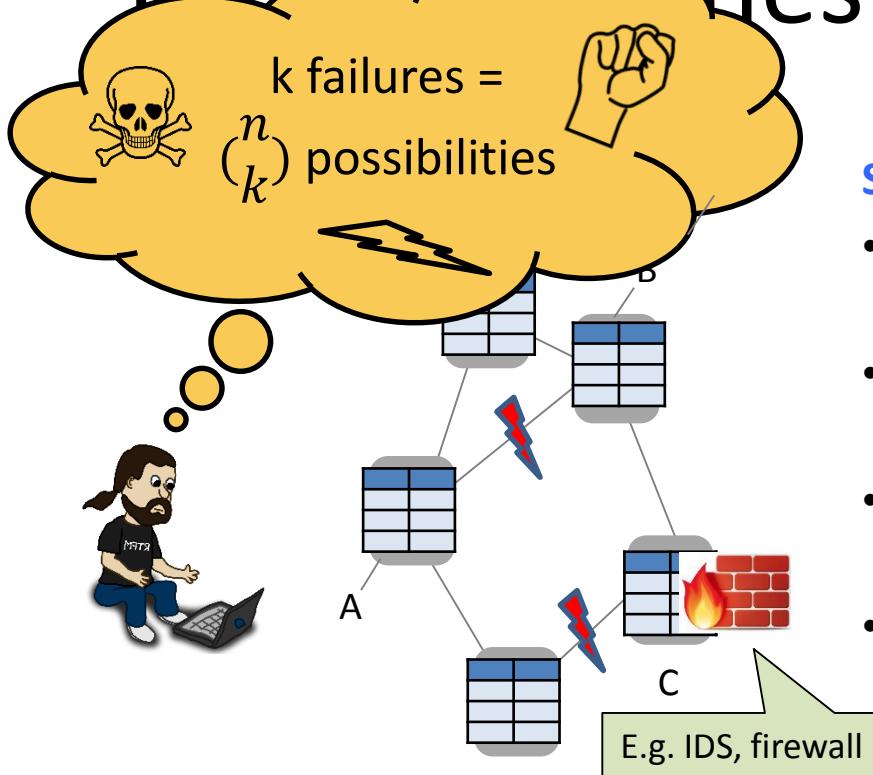
Responsibilities of a Sysadmin



Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?
- **Policy:** Is it ensured that traffic from A to B never goes via C?
- **Waypoint enforcement:** Is it ensured that traffic from A to B is always routed via a node C?

Responsibilities of a Sysadmin



Sysadmin responsible for:

- **Reachability:** Can traffic from ingress port A reach egress port B?
- **Loop-freedom:** Are the routes implied by the forwarding rules loop-free?
- **Policy:** Is it ensured that traffic from A to B never goes via C?
- **Waypoint enforcement:** Is it ensured that traffic from A to B is always routed via a node C?

... and everything even under multiple failures?!

Can we automate such tests
or even self-repair?

Can we automate such tests or even self-repair?



Yes! Encouraging: sometimes even ***fast***:
What-if Analysis Tool for MPLS and SR

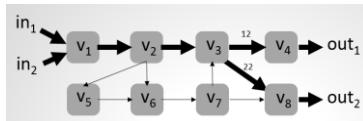
Leveraging Automata-Theoretic Approach



What if...?!

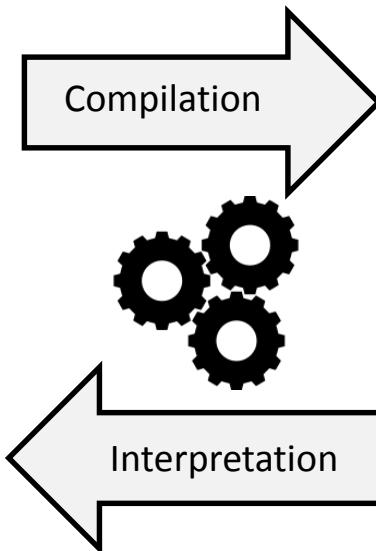


FT	In-I	In-Label	Out-I	op
τ_{v_1}	m_1	\perp	(v_1, v_2)	$push(10)$
	m_2	\perp	(v_1, v_2)	$push(20)$
τ_{v_2}	(v_1, v_2)	10	(v_2, v_3)	$swap(11)$
	(v_1, v_2)	20	(v_2, v_3)	$swap(21)$
τ_{v_3}	(v_2, v_3)	10	(v_1, v_2)	$swap(12)$
	(v_2, v_3)	21	(v_2, v_3)	$swap(22)$
	(v_7, v_3)	11	(v_3, v_1)	$swap(12)$
τ_{v_4}	(v_3, v_4)	12	out_1	pop
τ_{v_5}	(v_2, v_5)	40	(v_5, v_6)	pop
τ_{v_6}	(v_2, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	61	(v_6, v_7)	$swap(62)$
	(v_5, v_7)	71	(v_7, v_5)	$swap(72)$
τ_{v_7}	(v_6, v_7)	31	(v_7, v_5)	$pop(30)$
	(v_6, v_7)	62	(v_7, v_5)	$swap(11)$
	(v_6, v_7)	72	(v_7, v_5)	$swap(22)$
τ_{v_8}	(v_3, v_8)	22	out_2	pop
	(v_7, v_8)	22	out_2	pop



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_5)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_5)	30	(v_2, v_5)	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
$\tau_{v_2}^f$	(v_2, v_3)	11	(v_2, v_5)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$

MPLS configurations,
Segment Routing etc.



$pX \Rightarrow qXX$

$pX \Rightarrow qYX$

$qY \Rightarrow rYY$

$rY \Rightarrow r$

$rX \Rightarrow pX$

Pushdown Automaton
and Prefix Rewriting
Systems Theory

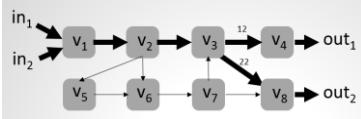
Leveraging Automata

Use cases: Sysadmin **issues queries**
to test certain properties, or do it
on a **regular basis** automatically!

What if...?!



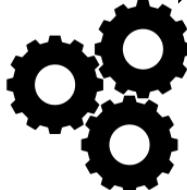
FT	In-I	In-Label	Out-I	op
τ_{v_1}	m_1	\perp	(v_1, v_2)	$push(10)$
	m_2	\perp	(v_1, v_2)	$push(20)$
τ_{v_2}	(v_1, v_2)	10	(v_2, v_3)	$swap(11)$
	(v_1, v_2)	20	(v_2, v_3)	$swap(21)$
τ_{v_3}	(v_2, v_3)	\perp	(v_1, v_2)	$swap(12)$
	(v_2, v_3)	21	(v_2, v_3)	$swap(22)$
	(v_7, v_3)	11	(v_3, v_4)	$swap(12)$
τ_{v_4}	(v_3, v_4)	21	(v_3, v_4)	$swap(22)$
	(v_3, v_4)	12	out_1	pop
τ_{v_5}	(v_2, v_5)	40	(v_5, v_6)	pop
τ_{v_6}	(v_2, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	30	(v_6, v_7)	$swap(31)$
	(v_5, v_6)	61	(v_6, v_7)	$swap(62)$
τ_{v_7}	(v_1, v_7)	\perp	(v_7, v_8)	$swap(72)$
	(v_6, v_7)	31	(v_7, v_8)	$pop(30)$
	(v_6, v_7)	62	(v_7, v_8)	$swap(11)$
τ_{v_8}	(v_6, v_7)	72	(v_7, v_8)	$swap(22)$
	(v_3, v_8)	22	out_2	pop
	(v_7, v_8)	22	out_2	pop



local FFT	Out-I	In-Label	Out-I	op
τ_{v_2}	(v_2, v_3)	11	(v_2, v_5)	$push(30)$
	(v_2, v_3)	21	(v_2, v_6)	$push(30)$
	(v_2, v_5)	30	(v_2, v_5)	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
$\tau_{v_2}^f$	(v_2, v_3)	11	(v_2, v_5)	$swap(61)$
	(v_2, v_3)	21	(v_2, v_6)	$swap(71)$
	(v_2, v_6)	61	(v_2, v_5)	$push(40)$
	(v_2, v_6)	71	(v_2, v_5)	$push(40)$

MPLS configurations,
Segment Routing etc.

Compilation



Interpretation

$$pX \Rightarrow qXX$$

$$pX \Rightarrow qYX$$

$$qY \Rightarrow rYY$$

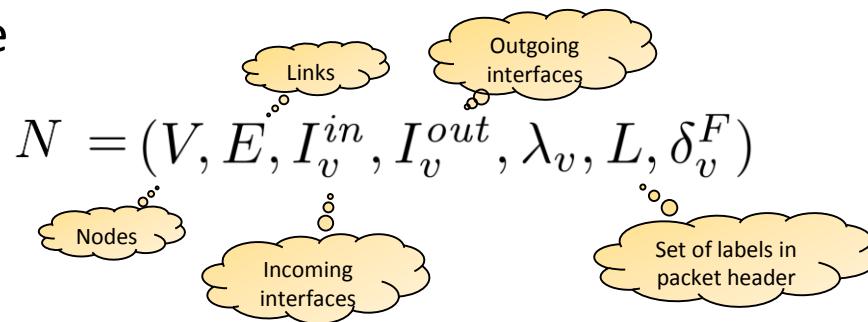
$$rY \Rightarrow r$$

$$rX \Rightarrow pX$$

Pushdown Automaton
and Prefix Rewriting
Systems Theory

Mini-Tutorial: A Network Model

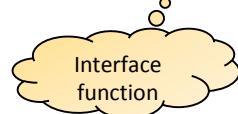
- Network: a 7-tuple



Mini-Tutorial: A Network Model

- Network: a 7-tuple

$$N = (V, E, I_v^{in}, I_v^{out}, \lambda_v, L, \delta_v^F)$$



Interface function: maps outgoing interface to next hop node and incoming interface to previous hop node

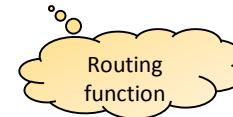
$$\lambda_v : I_v^{in} \cup I_v^{out} \rightarrow V$$

That is: $(\lambda_v(in), v) \in E$ and $(v, \lambda_v(out)) \in E$

Mini-Tutorial: A Network Model

- Network: a 7-tuple

$$N = (V, E, I_v^{in}, I_v^{out}, \lambda_v, L, \delta_v^F)$$



Routing function: for each set of **failed links** $F \subseteq E$, the routing function

$$\delta_v^F : I_v^{in} \times L^* \rightarrow 2^{(I_v^{out} \times L^*)}$$

defines, for all **incoming interfaces** and packet **headers**, **outgoing interfaces** together with **modified headers**.

Routing in Network

Packet routing sequence can be represented using **sequence of tuples**:

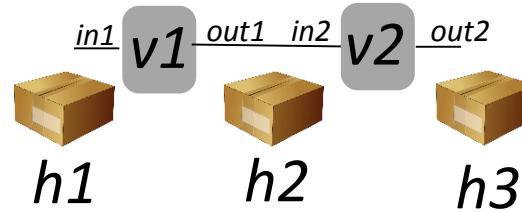


- Example: **routing** (in)finite sequence of tuples

$$(v_1, in_1, h_1, out_1, h_2, F_1),$$

$$(v_2, in_2, h_2, out_2, h_3, F_2),$$

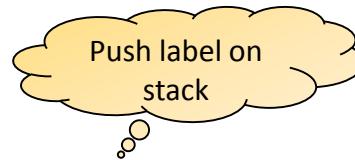
...



Example Rules:

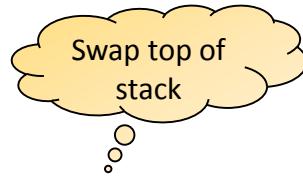
Regular Forwarding on Top-Most Label

Push:



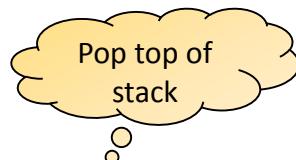
$$(v, \text{in})\ell \rightarrow (v, \text{out}, 0)\ell'\ell \text{ if } \tau_v(\text{in}, \ell) = (\text{out}, \text{push}(\ell'))$$

Swap:



$$(v, \text{in})\ell \rightarrow (v, \text{out}, 0)\ell' \text{ if } \tau_v(\text{in}, \ell) = (\text{out}, \text{swap}(\ell'))$$

Pop:



$$(v, \text{in})\ell \rightarrow (v, \text{out}, 0) \text{ if } \tau_v(\text{in}, \ell) = (\text{out}, \text{pop})$$

Example Failover Rules

Failover-Push:

Enumerate all
rerouting options

$$(v, \text{out}, i)\ell \rightarrow (v, \text{out}', i + 1)\ell'\ell \text{ for every } i, 0 \leq i < k, \\ \text{where } \pi_v(\text{out}, \ell) = (\text{out}', \text{push}(\ell'))$$

Failover-Swap:

$$(v, \text{out}, i)\ell \rightarrow (v, \text{out}', i + 1)\ell' \text{ for every } i, 0 \leq i < k, \\ \text{where } \pi_v(\text{out}, \ell) = (\text{out}', \text{swap}(\ell')),$$

Failover-Pop:

$$(v, \text{out}, i)\ell \rightarrow (v, \text{out}', i + 1) \text{ for every } i, 0 \leq i < k, \\ \text{where } \pi_v(\text{out}, \ell) = (\text{out}', \text{pop}).$$

Example rewriting sequence:

$$(v_1, \text{in}_1)h_1\perp \rightarrow (v_1, \text{out}, 0)h\perp \rightarrow (v_1, \text{out}', 1)h'\perp \rightarrow (v_1, \text{out}'', 2)h''\perp \rightarrow \dots \rightarrow (v_1, \text{out}_1, i)h_2\perp$$

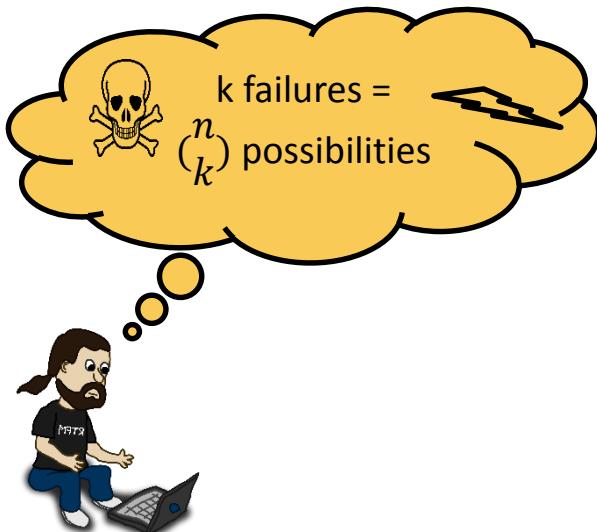
Try default

Try first backup

Try second backup

A Complex and Big Formal Language!

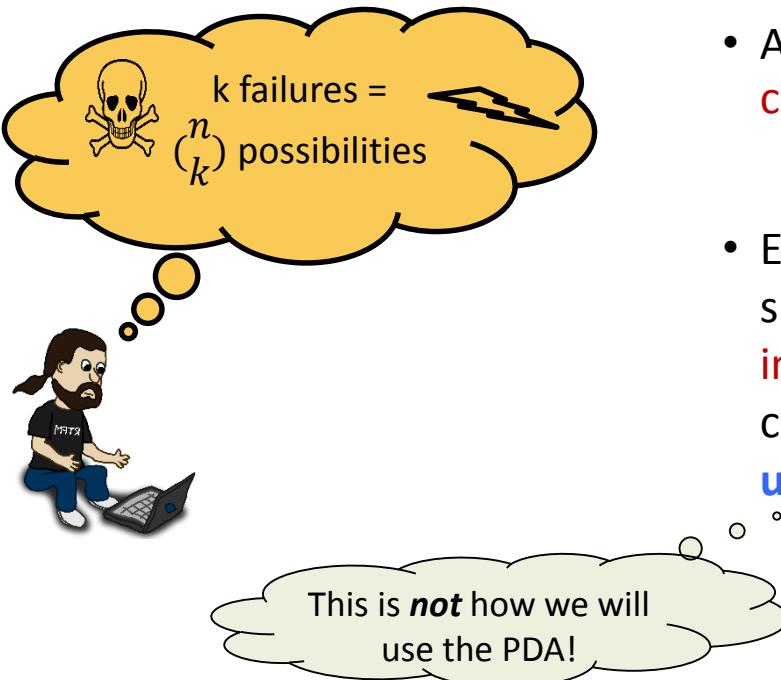
Why Polynomial Time?!



- Arbitrary number k of failures: How can I avoid **checking all $\binom{n}{k}$ many options?**!
- Even if we reduce to **push-down automaton**: simple operations such as **emptiness testing** or **intersection on Push-Down Automata (PDA)** is computationally non-trivial and sometimes even **undecidable**!

A Complex and Big Formal Language!

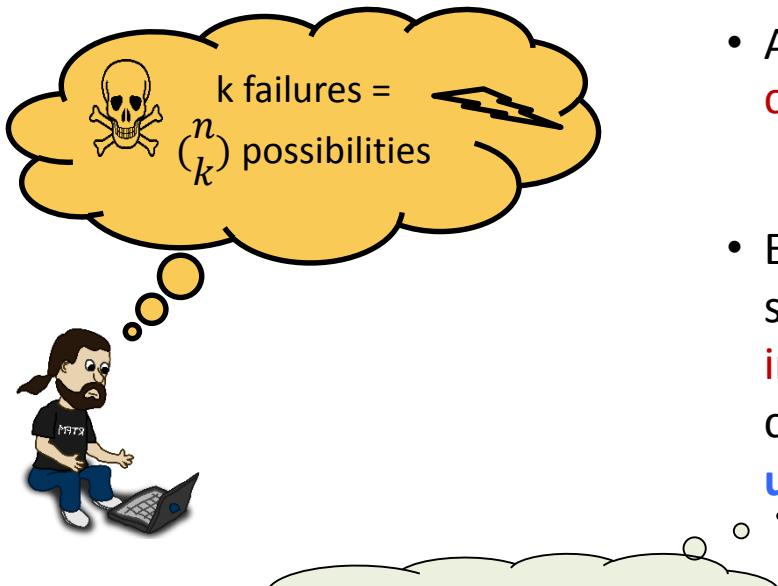
Why Polynomial Time?!



- Arbitrary number k of failures: How can I avoid checking all $\binom{n}{k}$ many options?!
- Even if we reduce to **push-down automaton**: simple operations such as **emptiness testing** or **intersection on Push-Down Automata (PDA)** is computationally non-trivial and sometimes even **undecidable**!

A Complex and Big Formal Language!

Why Polynomial Time?!



- Arbitrary number k of failures: How can I avoid checking all $\binom{n}{k}$ many options?!
- Even if we reduce to **push-down automaton**: simple operations such as **emptiness testing** or **intersection on Push-Down Automata (PDA)** is computationally non-trivial and sometimes even **undecidable**!

The words in our language are sequences of pushdown stack symbols, not the labels of transitions.

Time for Automata Theory (from Switzerland)!

- Classic result by **Büchi** 1964: the set of all reachable configurations of a pushdown automaton a is **regular set**
- Hence, we can operate only on **Nondeterministic Finite Automata (NFAs)** when reasoning about the pushdown automata
- The resulting **regular operations** are all **polynomial time**
 - Important result of **model checking**



Julius Richard Büchi

1924-1984

Swiss logician

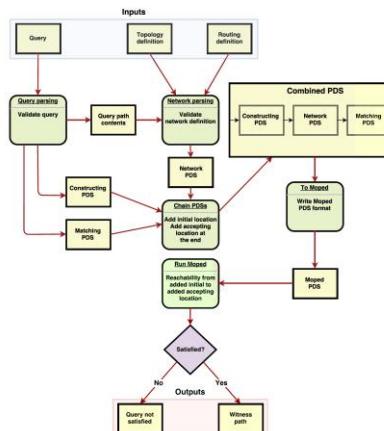
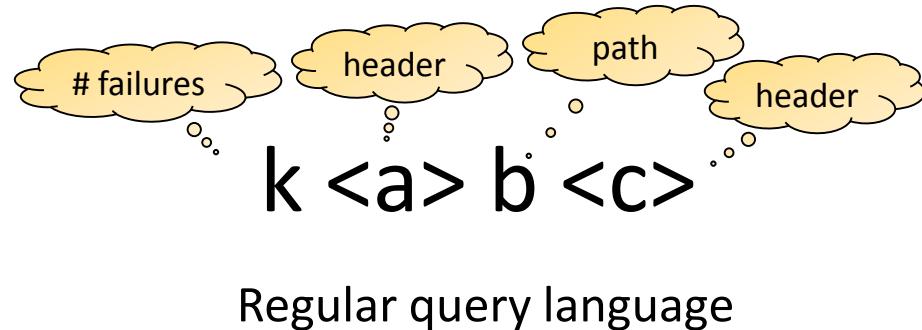
Tool and Query Language

Part 1: Parses query and constructs Push-Down System (PDS)

- In Python 3

Part 2: Reachability analysis of constructed PDS

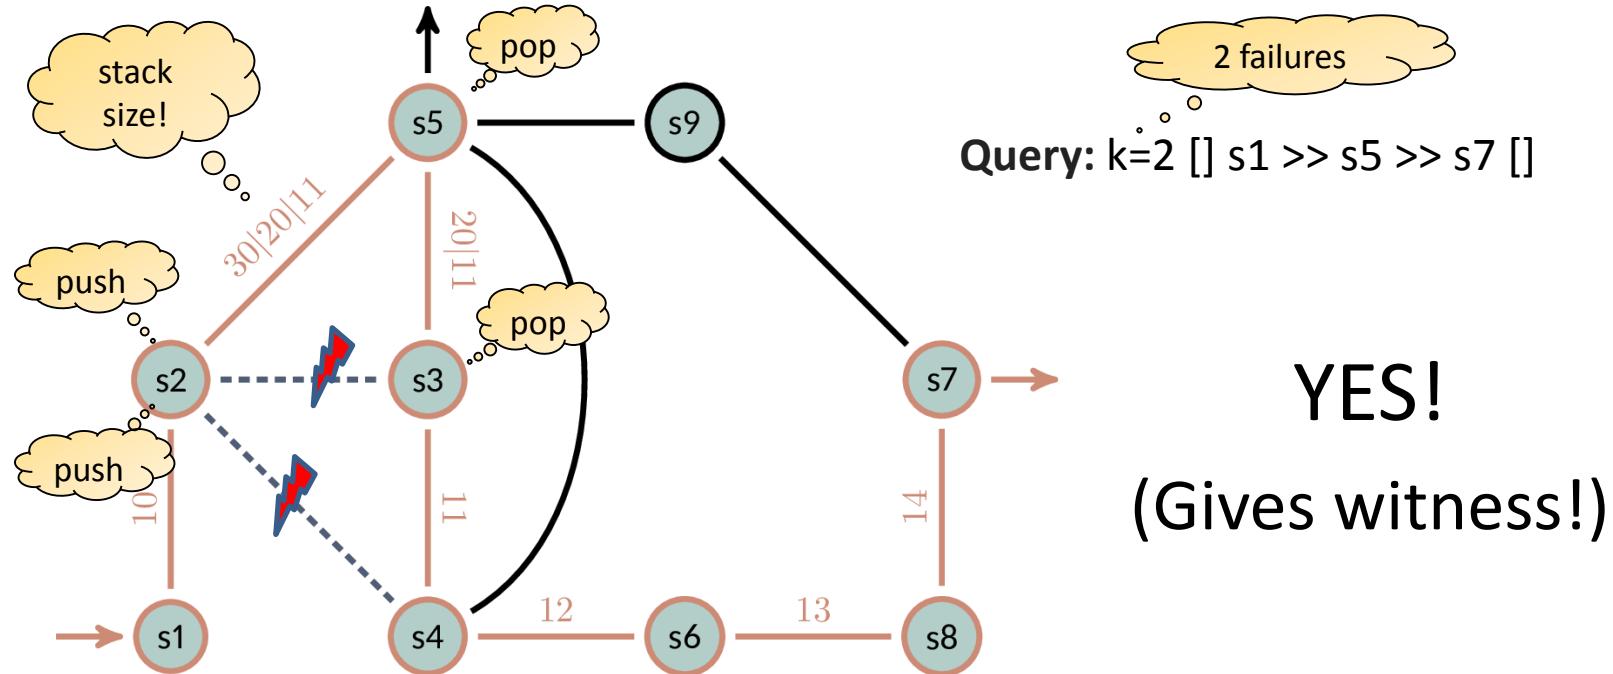
- Using **Moped** tool



query processing flow

Example: Traversal Testing With 2 Failures

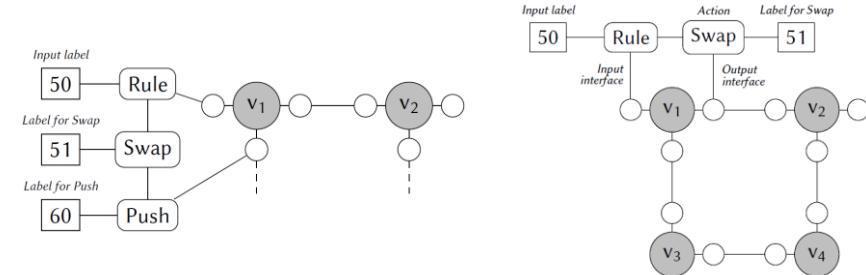
Traversal test with $k=2$: Can traffic starting with [] go through s_5 , under up to $k=2$ failures?



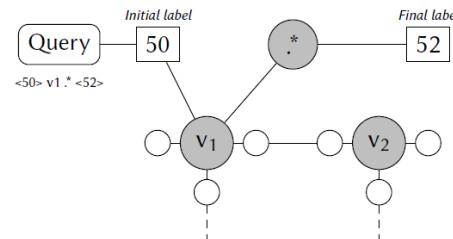
Formal methods are nice (give guarantees!)... But what about ML...?!

Speed Up Further and Synthesize: Deep Learning (s. talk by Fabien Geyer)

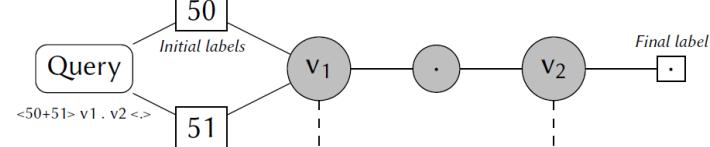
- Yes sometimes **without losing guarantees**
- Extend **graph-based neural networks**
- **Predict** counter-examples and **fixes**



Network topologies and MPLS rules

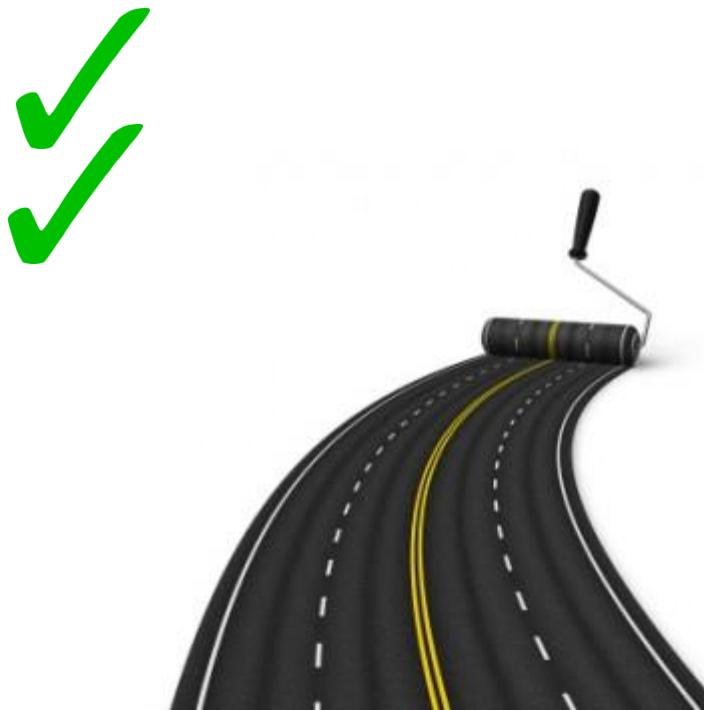


Network topologies and query



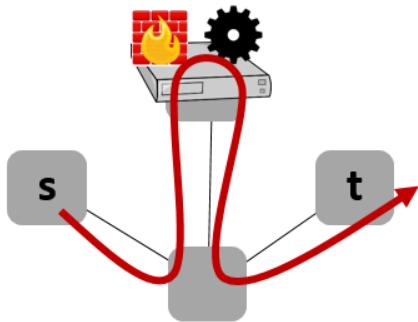
Roadmap

- Opportunities of self-* networks
 - Example 1: Demand-aware, self-adjusting networks
 - Example 2: Self-repairing networks
- Challenges of designing self-* networks

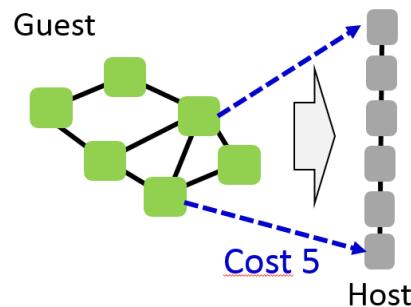


Challenge 1: Hard Problems

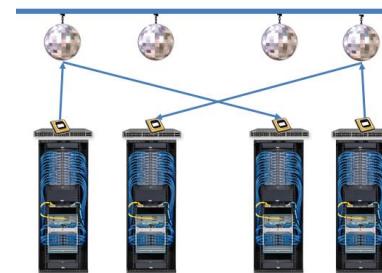
- Optimization problems are often **NP-hard**: hard *even for computers!*



Waypoint routing:
disjoint paths

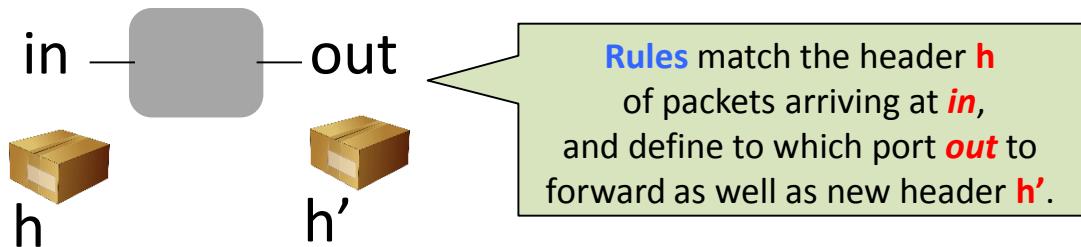


Embedding:
Minimum Lin. Arrangement



Topology design:
Graph spanners

It can get worse...: intractable!



(Simplified) MPLS rules:

prefix rewriting

$$in \times L \rightarrow out \times OP$$

where **OP** = {swap, push, pop}

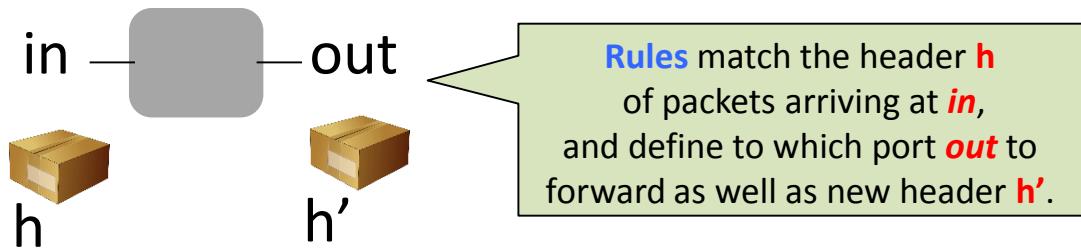
Rules of general networks (e.g., SDN):

VS

arbitrary header rewriting

$$in \times L^* \rightarrow out \times L^*$$

It can get worse...: intractable!



(Simplified) MPLS rules:

prefix rewrite

Polynomial time
 $in \xrightarrow{L} out \times OP$

where $OP = \{swap, push, pop\}$

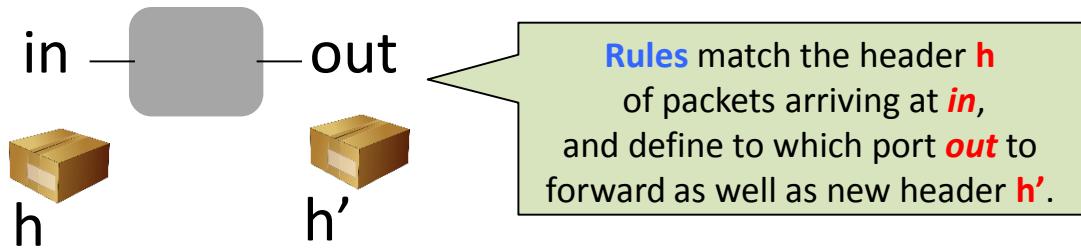
VS

Rules of general networks (e.g., SDN):

arbitrary headers

Undecidable!
 $in \xrightarrow{L} out \times L^*$

It can get worse....: intractable!



(Simplifying)

problem

What is a good tradeoff between generality and performance?

Polynomial

where $LOP = \{swap, push, pop\}$

$in \rightarrow out \times L^*$

Challenge 2: Realizing Limits?

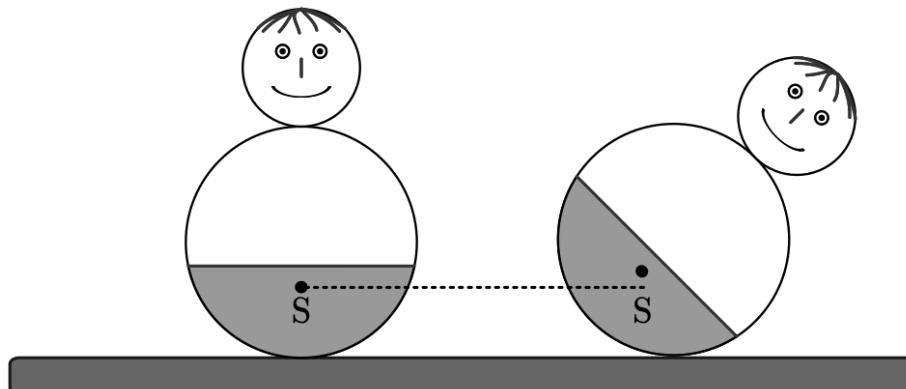
- Can a self-* network realize its **limits**?
- E.g., when quality of **input data** is not good enough?
- When to hand over to human? Or **fall back** to „safe/oblivious mode“?
- Can we learn from self-driving **cars**?



Challenge 3: Self-Stabilization

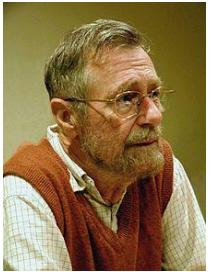
- Could be an attractive property of self-* network!

A **self-stabilizing** system guarantees that it *reconverges to a desirable configuration* or state, *from any initial state*.



„Stehaufmännchen“

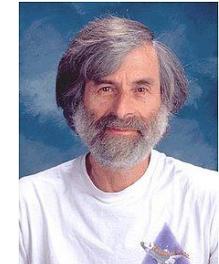
Self-Stabilization



Self-stabilizing algorithms pioneered by **Dijkstra** (1973): for example **self-stabilizing mutual exclusion**.

“I regard this as Dijkstra’s most brilliant work. Self-stabilization is a very important concept in **fault tolerance**.”

Leslie **Lamport** (PODC 1983)

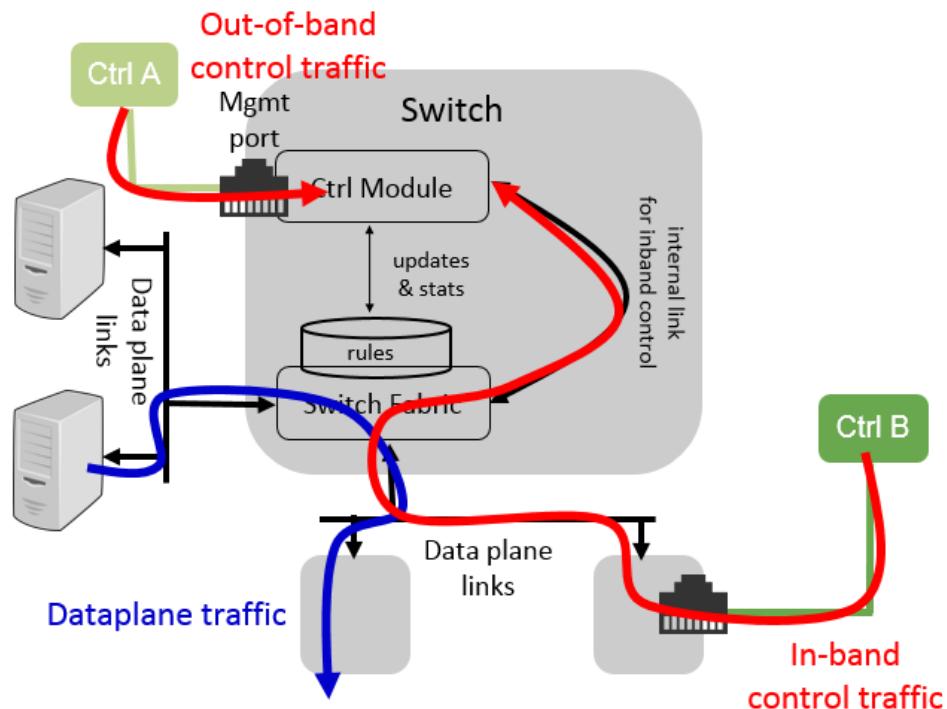


Some notable works by **Perlman** toward self-stabilizing Internet, e.g., **self-stabilizing spanning trees**.

Yet, many protocols in the Internet are *not* self-stabilizing. Much need for future work.

E.g., Self-Stabilizing SDN Control?

- Distributed SDN control plane which **self-organizes management** of switches?
- Especially challenging: **inband control** (how to distinguish traffic?)

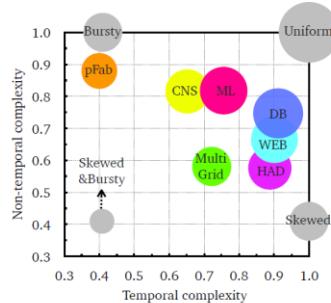
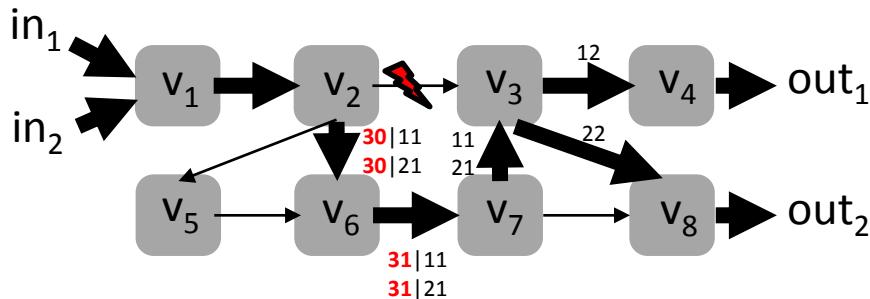


Challenge 4: Uncertainties

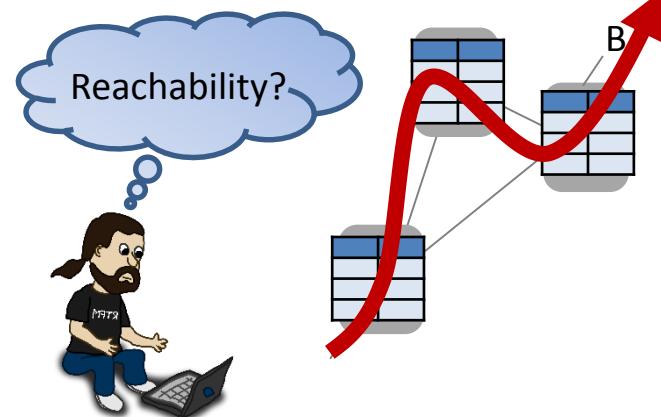
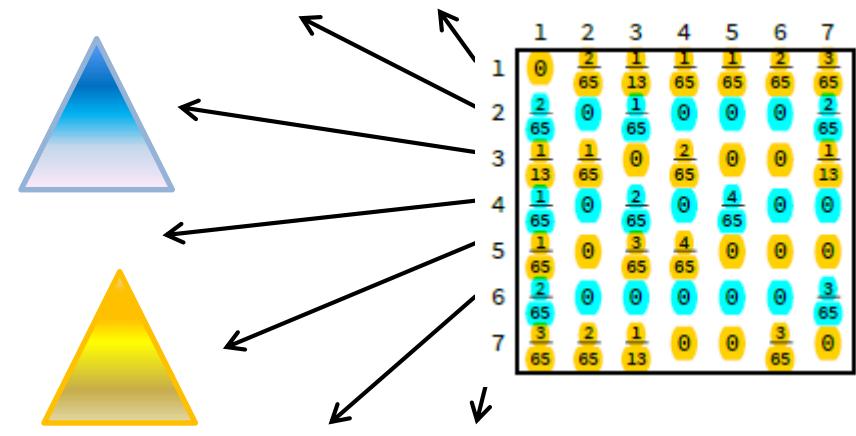
- How to deal with **uncertainties**?
- How to maintain flexibilities?
- Use of principles from robotics? E.g., **empowerment**?

Conclusion

- **Flexibilities** in networks: great opportunities for **optimization** and **automation**
- **Demand-aware** and **self-adjusting** networks: beating the routing lower bounds of oblivious networks, ***reaching entropy bounds***
- Potential of **self-repairing** networks, self-stabilizing networks, etc.
- Much work ahead: ***tradeoff*** generality vs efficiency? How to self-monitor and **fall-back** if needed? Use of **formal methods** and ML?



Thank you! Questions?



Further Reading

Flexibilities and Complexity

[On The Impact of the Network Hypervisor on Virtual Network Performance](#)

Andreas Blenk, Arsany Basta, Wolfgang Kellerer, and Stefan Schmid.

IFIP Networking, Warsaw, Poland, May 2019.

[Adaptable and Data-Driven Softwarized Networks: Review, Opportunities, and Challenges](#) (Invited Paper)

Wolfgang Kellerer, Patrick Kalmbach, Andreas Blenk, Arsany Basta, Martin Reisslein, and Stefan Schmid.

Proceedings of the IEEE (PIEEE), 2019.

[Efficient Distributed Workload \(Re-\)Embedding](#)

Monika Henzinger, Stefan Neumann, and Stefan Schmid.

ACM/IFIP **SIGMETRICS/PERFORMANCE**, Phoenix, Arizona, USA, June 201

[Parametrized Complexity of Virtual Network Embeddings: Dynamic & Linear Programming Approximations](#)

Matthias Rost, Elias Döhne, and Stefan Schmid.

ACM SIGCOMM Computer Communication Review (**CCR**), January 2019.

[Charting the Complexity Landscape of Virtual Network Embeddings](#) (Best Paper Award)

Matthias Rost and Stefan Schmid.

IFIP Networking, Zurich, Switzerland, May 2018.

[Tomographic Node Placement Strategies and the Impact of the Routing Model](#)

Yvonne Anne Pignolet, Stefan Schmid, and Gilles Tredan.

ACM **SIGMETRICS**, Irvine, California, USA, June 2018. hmid.

ACM/IEEE Symposium on Architectures for Networking and Communications Systems (**ANCS**), Ithaca, New York, USA, July 2018.

Further Reading

Demand-Aware and Self-Adjusting Networks

[Survey of Reconfigurable Data Center Networks: Enablers, Algorithms, Complexity](#)

Klaus-Tycho Foerster and Stefan Schmid.

SIGACT News, June 2019.

[Toward Demand-Aware Networking: A Theory for Self-Adjusting Networks \(Editorial\)](#)

Chen Avin and Stefan Schmid.

ACM SIGCOMM Computer Communication Review (**CCR**), October 2018.

[Demand-Aware Network Design with Minimal Congestion and Route Lengths](#)

Chen Avin, Kaushik Mondal, and Stefan Schmid.

38th IEEE Conference on Computer Communications (**INFOCOM**), Paris, France, April 2019.

Documents: paper [pdf](#), bibtex [bib](#)

[Distributed Self-Adjusting Tree Networks](#)

Bruna Peres, Otavio Augusto de Oliveira Souza, Olga Goussevskaia, Chen Avin, and Stefan Schmid.

38th IEEE Conference on Computer Communications (**INFOCOM**), Paris, France, April 2019.

[Efficient Non-Segregated Routing for Reconfigurable Demand-Aware Networks](#)

Thomas Fenz, Klaus-Tycho Foerster, Stefan Schmid, and Anaïs Villedieu.

IFIP Networking, Warsaw, Poland, May 2019.

[DaRTree: Deadline-Aware Multicast Transfers in Reconfigurable Wide-Area Networks](#)

Long Luo, Klaus-Tycho Foerster, Stefan Schmid, and Hongfang Yu.

IEEE/ACM International Symposium on Quality of Service (**IWQoS**), Phoenix, Arizona, USA, June 2019.

[Demand-Aware Network Designs of Bounded Degree](#)

Chen Avin, Kaushik Mondal, and Stefan Schmid.

31st International Symposium on Distributed Computing (**DISC**), Vienna, Austria, October 2017.

[SplayNet: Towards Locally Self-Adjusting Networks](#)

Stefan Schmid, Chen Avin, Christian Scheideler, Michael Borokhovich, Bernhard Haeupler, and Zvi Lotker.

IEEE/ACM Transactions on Networking (**TON**), Volume 24, Issue 3, 2016. Early version: IEEE **IPDPS** 2013.

[Characterizing the Algorithmic Complexity of Reconfigurable Data Center Architectures](#)

Klaus-Tycho Foerster, Monia Ghobadi, and Stefan Schmid.

ACM/IEEE Symposium on Architectures for Networking and Communications Systems (**ANCS**), Ithaca, New York, USA, July 2018.

Further Reading

Self-Repairing Networks

[P-Rex: Fast Verification of MPLS Networks with Multiple Link Failures](#)

Jesper Stenbjerg Jensen, Troels Beck Krogh, Jonas Sand Madsen, Stefan Schmid, Jiri Srba, and Marc Tom Thorgersen.

14th International Conference on emerging Networking EXperiments and Technologies (**CoNEXT**), Heraklion, Greece, December 2018.

[Polynomial-Time What-If Analysis for Prefix-Manipulating MPLS Networks](#)

Stefan Schmid and Jiri Srba.

37th IEEE Conference on Computer Communications (**INFOCOM**), Honolulu, Hawaii, USA, April 2018.

[Renaissance: A Self-Stabilizing Distributed SDN Control Plane](#)

Marco Canini, Iosif Salem, Liron Schiff, Elad Michael Schiller, and Stefan Schmid.

38th IEEE International Conference on Distributed Computing Systems (**ICDCS**), Vienna, Austria, July 2018.

[Empowering Self-Driving Networks](#)

Patrick Kalmbach, Johannes Zerwas, Peter Babarczi, Andreas Blenk, Wolfgang Kellerer, and Stefan Schmid.

ACM SIGCOMM 2018 Workshop on Self-Driving Networks (**SDN**), Budapest, Hungary, August 2018.

[DeepMPLS: Fast Analysis of MPLS Configurations using Deep Learning](#)

Fabien Geyer and Stefan Schmid.

[IFIP Networking](#), Warsaw, Poland, May 2019.