

# Fast Re-Routing in Networks: On the Complexity of Perfect Resilience

Matthias Bentert (TU Berlin)

Esra Ceylan (TU Berlin + ISTA)

Valentin Huebner (ISTA)

Stefan Schmid (TU Berlin)

Jiri Srba (Aalborg University)

Acknowledgements:



Networks:

# Critical Infrastructure

- If networks break, it can have knock-on effects
- For example, Facebook outage in 2021: not only took down their social networking site, but also Instagram, WhatsApp, ...
- ... and their own internal systems, which manage the doors: engineers had to break into their own buildings to bring the network back up

The New York Times

## *Gone in Minutes, Out for Hours: Outage Shakes Facebook*

When apps used by billions of people worldwide blinked out, lives were disrupted, businesses were cut off from customers — and some Facebook employees were locked out of their offices.

 Share full article    884



Facebook's internal communications platform, Workplace, was also taken out, leaving most employees unable to do their jobs. Kelsey McClellan for The New York Times

Credits: Nate Foster

# The Challenge: Most Outages due to Human Errors

## Human Errors

### Countries disconnected

Data Centre ▶ Networks

#### Google routing blunder sent Japan's Internet dark on Friday

Another big BGP blunder

By Richard Chirgwin 27 Aug 2017 at 22:35

40 SHARE ▾

Last Friday, someone in Google fat-thumbed a border gateway protocol (BGP) advertisement and sent Japanese Internet traffic into a black hole.

The trouble began when The Chocolate Factory "leaked" a big route table to Verizon, the result of which was traffic from Japanese giants like NTT and KDDI was sent to Google on the expectation it would be treated as transit.

### Passengers stranded

#### British Airways' latest Total Inability To Support Upwardness of Planes\* caused by Amadeus system outage

Stuck on the ground awaiting a load sheet? Here's why

By Gareth Corfield 19 Jul 2018 at 11:16

109 SHARE ▾



RA flights around the world were canceled as a result of the Amadeus outage

### Even 911 affected

#### Officials: Human error to blame in Minn. 911 outage

According to a press release, CenturyLink told department of public safety that human error by an employee of a third party vendor was to blame for the outage

Aug 16, 2018

Duluth News Tribune

SAINT PAUL, Minn. — The Minnesota Department of Public Safety Emergency Communication Networks division was told by its 911 provider that an Aug. 1 outage was caused by human error.

Even tech-savvy companies struggle:



Slide credits: Nate Foster and Laurent Vanbever


Mainly:  
human  
errors!

# A Reason: Complexity

Especially Under Failures (Policy Compliance)

Example: BGP in

Microsoft  
datacenter




# A Reason: Complexity

Especially Under Failures (Policy Compliance)

Example: BGP in


Microsoft  
datacenter



# A Reason: Complexity

Especially Under Failures (Policy Compliance)


Example: BGP in  
**Microsoft**  
datacenter



# A Reason: Complexity

Especially Under Failures (Policy Compliance)


Example: BGP in  
**Microsoft**  
datacenter



# A Reason: Complexity

Especially Under Failures (Policy Compliance)


Example: BGP in  
**Microsoft**  
datacenter



# A Reason: Complexity

Especially Under Failures (Policy Compliance)

Example: BGP in  
Microsoft  
datacenter




If link (G,X) fails and traffic from G is rerouted via Y and C to X:  
X announces (does not block) G and H as it comes from C. (Note: BGP.)

# A Reason: Complexity

Especially Under Failures (Policy Compliance)

Example: BGP in  
Microsoft  
datacenter




If link (G,X) fails and traffic from G is rerouted via Y and C to X:  
X announces (does not block) G and H as it comes from C. (Note: BGP.)

# Vision to Deal with Complexity: Automation




FT	In-I	In-Label	Out-I	op
$\tau_{v_1}$	$m_1$	$\perp$	$(v_1, v_2)$	$push(10)$
	$m_2$	$\perp$	$(v_1, v_2)$	$push(20)$
$\tau_{v_2}$	$(v_1, v_2)$	20	$(v_2, v_3)$	$swap(11)$
	$(v_1, v_2)$	20	$(v_2, v_3)$	$swap(21)$
$\tau_{v_3}$	$(v_2, v_3)$	11	$(v_3, v_4)$	$swap(12)$
	$(v_2, v_3)$	21	$(v_3, v_8)$	$swap(22)$
	$(v_2, v_3)$	11	$(v_3, v_4)$	$swap(12)$
	$(v_2, v_3)$	21	$(v_3, v_8)$	$swap(22)$
$\tau_{v_4}$	$(v_3, v_4)$	12	$out_1$	$pop$
	$(v_2, v_3)$	40	$(v_5, v_6)$	$pop$
$\tau_{v_5}$	$(v_2, v_6)$	30	$(v_6, v_7)$	$swap(31)$
	$(v_5, v_6)$	30	$(v_6, v_7)$	$swap(31)$
	$(v_5, v_6)$	61	$(v_6, v_7)$	$swap(62)$
	$(v_5, v_6)$	71	$(v_6, v_7)$	$swap(72)$
$\tau_{v_7}$	$(v_6, v_7)$	31	$(v_7, v_3)$	$pop$
	$(v_6, v_7)$	62	$(v_7, v_3)$	$swap(11)$
	$(v_6, v_7)$	72	$(v_7, v_8)$	$swap(22)$
$\tau_{v_8}$	$(v_5, v_8)$	22	$out_2$	$pop$
	$(v_7, v_8)$	22	$out_2$	$pop$




local FFT	Out-I	In-Label	Out-I	op
$\tau_{v_2}$	$(v_2, v_3)$	11	$(v_2, v_6)$	$push(30)$
	$(v_2, v_3)$	21	$(v_2, v_6)$	$push(30)$
	$(v_2, v_6)$	30	$(v_2, v_5)$	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
$\tau'_{v_2}$	$(v_2, v_3)$	11	$(v_2, v_6)$	$swap(61)$
	$(v_2, v_3)$	21	$(v_2, v_6)$	$swap(71)$
	$(v_2, v_6)$	61	$(v_2, v_5)$	$push(40)$
	$(v_2, v_6)$	71	$(v_2, v_5)$	$push(40)$

Router **configurations**  
(Cisco, Juniper, etc.)


# Vision to Deal with Complexity: Automation



FT	In-I	In-Label	Out-I	op
$\tau_{v_1}$	$m_1$	$\perp$	$(v_1, v_2)$	$push(10)$
	$m_2$	$\perp$	$(v_1, v_2)$	$push(20)$
$\tau_{v_2}$	$(v_1, v_2)$	10	$(v_2, v_3)$	$swap(11)$
	$(v_1, v_2)$	20	$(v_2, v_3)$	$swap(21)$
$\tau_{v_3}$	$(v_2, v_3)$	11	$(v_3, v_4)$	$swap(12)$
	$(v_2, v_3)$	21	$(v_3, v_8)$	$swap(22)$
	$(v_7, v_3)$	11	$(v_3, v_4)$	$swap(12)$
	$(v_7, v_3)$	21	$(v_3, v_8)$	$swap(22)$
$\tau_{v_4}$	$(v_3, v_4)$	12	$out_1$	$pop$
	$(v_2, v_5)$	40	$(v_5, v_6)$	$pop$
$\tau_{v_5}$	$(v_2, v_6)$	30	$(v_6, v_7)$	$swap(31)$
	$(v_5, v_6)$	30	$(v_6, v_7)$	$swap(31)$
	$(v_5, v_6)$	61	$(v_6, v_7)$	$swap(62)$
	$(v_5, v_6)$	71	$(v_6, v_7)$	$swap(72)$
$\tau_{v_7}$	$(v_6, v_7)$	31	$(v_7, v_3)$	$pop$
	$(v_6, v_7)$	62	$(v_7, v_3)$	$swap(11)$
	$(v_6, v_7)$	72	$(v_7, v_8)$	$swap(22)$
$\tau_{v_8}$	$(v_5, v_8)$	22	$out_2$	$pop$
	$(v_7, v_8)$	22	$out_2$	$pop$



local FFT	Out-I	In-Label	Out-I	op
$\tau_{v_2}$	$(v_2, v_3)$	11	$(v_2, v_6)$	$push(30)$
	$(v_2, v_3)$	21	$(v_2, v_6)$	$push(30)$
	$(v_2, v_6)$	30	$(v_2, v_5)$	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
$\tau'_{v_2}$	$(v_2, v_3)$	11	$(v_2, v_6)$	$swap(61)$
	$(v_2, v_3)$	21	$(v_2, v_6)$	$swap(71)$
	$(v_2, v_6)$	61	$(v_2, v_5)$	$push(40)$
	$(v_2, v_6)$	71	$(v_2, v_5)$	$push(40)$




$pX \Rightarrow qXX$   
 $pX \Rightarrow qYX$   
 $qY \Rightarrow rYY$   
 $rY \Rightarrow r$   
 $rX \Rightarrow pX$

**Formal language** which  
 supports **automated  
verification**  
 („what-if analysis“)


# Vision to Deal with Complexity: Automation




FT	In-I	In-Label	Out-I	op
$\tau_{v_1}$	$m_1$	$\perp$	$(v_1, v_2)$	$push(10)$
	$m_2$	$\perp$	$(v_1, v_2)$	$push(20)$
$\tau_{v_2}$	$(v_1, v_2)$	10	$(v_2, v_3)$	$swap(11)$
	$(v_1, v_2)$	20	$(v_2, v_3)$	$swap(21)$
$\tau_{v_3}$	$(v_2, v_3)$	11	$(v_3, v_4)$	$swap(12)$
	$(v_2, v_3)$	21	$(v_3, v_8)$	$swap(22)$
	$(v_2, v_3)$	11	$(v_3, v_4)$	$swap(12)$
	$(v_2, v_3)$	21	$(v_3, v_8)$	$swap(22)$
$\tau_{v_4}$	$(v_3, v_4)$	12	$out_1$	$pop$
	$(v_2, v_3)$	40	$(v_5, v_6)$	$pop$
$\tau_{v_5}$	$(v_2, v_3)$	30	$(v_6, v_7)$	$swap(31)$
	$(v_5, v_6)$	30	$(v_6, v_7)$	$swap(31)$
	$(v_5, v_6)$	61	$(v_6, v_7)$	$swap(62)$
	$(v_5, v_6)$	71	$(v_6, v_7)$	$swap(72)$
$\tau_{v_7}$	$(v_6, v_7)$	31	$(v_7, v_8)$	$pop$
	$(v_6, v_7)$	62	$(v_7, v_8)$	$swap(11)$
	$(v_6, v_7)$	72	$(v_7, v_8)$	$swap(22)$
$\tau_{v_8}$	$(v_5, v_6)$	22	$out_2$	$pop$
	$(v_5, v_6)$	22	$out_2$	$pop$



local FFT	Out-I	In-Label	Out-I	op
$\tau_{v_2}$	$(v_2, v_3)$	11	$(v_2, v_6)$	$push(30)$
	$(v_2, v_3)$	21	$(v_2, v_6)$	$push(30)$
	$(v_2, v_6)$	30	$(v_2, v_5)$	$push(40)$
global FFT	Out-I	In-Label	Out-I	op
$\tau'_{v_2}$	$(v_2, v_3)$	11	$(v_2, v_6)$	$swap(61)$
	$(v_2, v_3)$	21	$(v_2, v_6)$	$swap(71)$
	$(v_2, v_6)$	61	$(v_2, v_5)$	$push(40)$
	$(v_2, v_6)$	71	$(v_2, v_5)$	$push(40)$



$pX \Rightarrow qXX$   
 $pX \Rightarrow qYX$   
 $qY \Rightarrow rYY$   
 $rY \Rightarrow r$   
 $rX \Rightarrow pX$



Or even  
*generate/fix?*

**Formal language** which  
 supports **automated**  
**verification**  
 („what-if analysis“)


→ Generation aka. *Synthesis*

Information at Switch for  
**Local Decision Making?**




# Information at Switch for Local Decision Making?

- Nodes locally store a forwarding Match -> Action table




# Information at Switch for Local Decision Making?

→ The **Packet Header** (e.g., source, destination)




# Information at Switch for Local Decision Making?

→ The **Import** of the received packet



# Information at Switch for Local Decision Making?

→ Which **incident links** failed



Two fundamental

# Notions of Resilience

## Ideal resilience

Given a  $k$ -connected graphs, fast reroute can tolerate *any  $k-1$  link failures.*

## Perfect resilience

Fast reroute can tolerate any failures as long as the underlying network is *physically connected.*

What is the difference? Which is stronger?

Two fundamental

# Notions of Resilience

## Ideal resilience

Given a  $k$ -connected graphs, fast reroute can tolerate *any  $k-1$  link failures.*

## Perfect resilience

Fast reroute can tolerate any failures as long as the underlying network is *physically connected.*


- Our focus: *verification and synthesis of perfect resilience*
- We do have insights in verification of ideal resilience
- Synthesis of ideal resilience: *open problem* (harder)

# Our Results

- *Synthesis* of Perfect Resilience is *in P*. (Non-constructive! Big open problem: *synthesis of perfect resilience*.)
- *Verification* of Perfect Resilience is *coNP-complete*. Even if the input graph is *planar* and for simple “*skipping* rules”.
- *Verification and synthesis* problems for Perfect Resilience using *in-port oblivious* are in *Linear time* (constructive).

# Our Results

- *Synthesis* of Perfect Resilience is *in P*. (Non-constructive! Big open problem: *synthesis of perfect resilience*.)
- *Verification* of Perfect Resilience is *coNP-complete*. Even if the input graph is *planar* and for simple “*skipping* rules”.
- *Verification and synthesis* problems for Perfect Resilience using *in-port oblivious* are in *Linear time* (constructive).



There exists a perfectly resilient in-port oblivious local (or skipping) routing for any given target node if and only if all simple *cycles* of  $G$  have length *at most 3*. Can be computed with a modified BFS.

# Our Results

From *Robertson and Seymour*'s result on the *minor-stability* property for the graphs for which perfect resilience is not impossible (known from *Foerster et al.*). Whether a graph contains a *forbidden minor* can be checked in polynomial time.

- *Synthesis* of Perfect Resilience is *in P*. (Non-constructive! Big open problem: *synthesis of perfect resilience*.)
- *Verification* of Perfect Resilience is *coNP-complete*. Even if the input graph is *planar* and for simple “*skipping* rules”.
- *Verification and synthesis* problems for Perfect Resilience using *in-port oblivious* are in *Linear time* (constructive).

# Our Results

- *Synthesis* of Perfect Resilience is *in P*. (Non-constructive! Big open problem: *synthesis of perfect resilience*.)
- *Verification* of Perfect Resilience is *coNP-complete*. Even if the input graph is *planar* and for simple “*skipping* rules”.
- *Verification and synthesis* problems for Perfect Resilience using *in-port oblivious* are in *Linear time* (constructive).



# CoNP Completeness

Recall *coNP-complete*: complement of the language is *NP-complete*

→ Alternative to show hardness: reduce from coNP-hard problem

**The problem is in coNP:**

- We can *guess a failure* scenario where a source node is connected to the target but where the routing tables *create a forwarding Loop*
- The given failure scenario can be checked in polynomial time

**coNP-hardness:**

- By a reduction from *3-Sat with exactly 3 literals*
- The constructed skipping routing is perfectly resilient if and only if the original instance of 3-Sat is *not satisfiable*

# Future Work

- Ideal resilience conjecture
- Constructive proof that perfect resilience synthesis is in P

# Thank you and References

## [A Survey of Fast-Recovery Mechanisms in Packet-Switched Networks](#)

Marco Chiesa, Andrzej Kaminski, Jacek Rak, Gabor Retvari, and Stefan Schmid.  
IEEE Communications Surveys and Tutorials (**COMST**), 2021.

## [AalWiNes: A Fast and Quantitative What-If Analysis Tool for MPLS Networks](#)

Peter Gjøl Jensen, Morten Konggaard, Dan Kristiansen, Stefan Schmid, Bernhard Clemens Schrenk, and Jiri Srba.

16th ACM International Conference on emerging Networking EXperiments and Technologies (**CoNEXT**), Barcelona, Spain, December 2020.

## [A Tight Characterization of Fast Failover Routing: Resiliency to Two Link Failures is Possible](#)

Wenkai Dai, Klaus-Tycho Foerster, and Stefan Schmid.

35th ACM Symposium on Parallelism in Algorithms and Architectures (**SPAA**), Orlando, Florida, USA, June 2023.

## [On the Price of Locality in Static Fast Rerouting](#)

Klaus-Tycho Foerster, Juho Hirvonen, Yvonne-Anne Pignolet, Stefan Schmid, and Gilles Tredan.  
52nd IEEE/IFIP International Conference on Dependable Systems and Networks (**DSN**), Baltimore, Maryland, USA, June 2022.

## [SyPer: Synthesis of Perfectly Resilient Local Fast Rerouting Rules for Highly Dependable Networks](#)

Csaba Györgyi, Kim G. Larsen, Stefan Schmid, and Jiri Srba.

IEEE Conference on Computer Communications (**INFOCOM**), Vancouver, Canada, May 2024.

## [DeepMPLS: Fast Analysis of MPLS Configurations Using Deep Learning](#)

Fabien Geyer and Stefan Schmid.

IFIP Networking, Warsaw, Poland, May 2019.

## [Latte: Improving the Latency of Transiently Consistent Network Update Schedules](#)

Mark Glavind, Niels Christensen, Jiri Srba, and Stefan Schmid.

38th International Symposium on Computer Performance, Modeling, Measurements and Evaluation (**PERFORMANCE**) and ACM Performance Evaluation Review (**PER**), Milan, Italy, November 2020.

## [Model-Based Insights on the Performance, Fairness, and Stability of BBR](#)

Simon Scherrer, Markus Legner, Adrian Perrig, and Stefan Schmid.

ACM Internet Measurement Conference (**IMC**), Nice, France, October 2022.

Slides  
available:

