



universität
wien



Fast Rerouting Against Dynamic Failures: 2-Resilience via Ear- Decomposition and Planarity

Wenkai Dai (TU Berlin, Germany & University of Vienna, Austria), Klaus-Tycho Foerster (TU Dortmund, Germany), and Stefan Schmid (TU Berlin, Germany)




29th Conference on Principles of Distributed Systems




Motivation: Handling Link-Failures on Control-Plane Is Slow

- Link **failures** are **unavoidable** so mechanism to **route around**
- Traditionally, control-plane handles link failures by recomputing routing tables (e.g., OSPF)
- **Slow**, unacceptable for critical applications




Motivation: Handling Link-Failures on Control-Plane Is Slow

- Link **failures** are **unavoidable** so mechanism to **route** around
- Traditionally, control-plane handles link failures by recomputing routing tables (e.g., OSPF)
- **Slow**, unacceptable for critical applications




Motivation: Handling Link-Failures on Control-Plane Is Slow

- Fast reroute (FRR) on data-plane handles failures locally



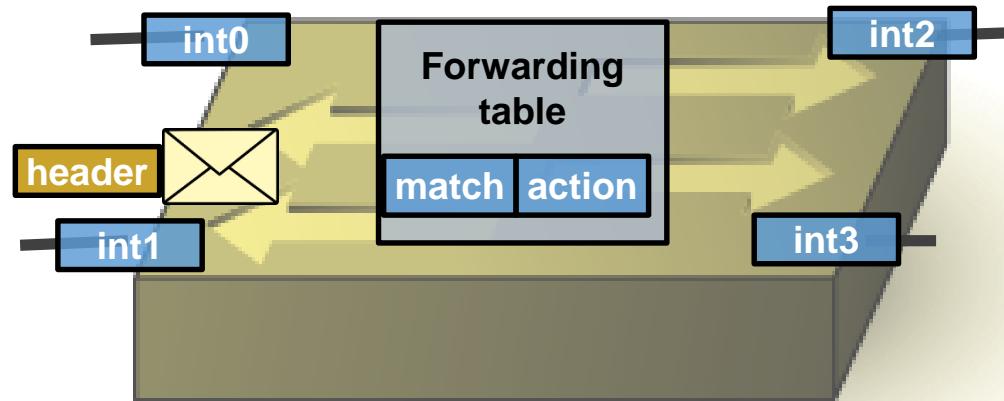
Fast Rerouting (FRR) on Data-Plane

- Local failover without invoking control plane, by pre-installing alternative paths
- Conditional forwarding rules, only depending on local failures, no information about failures downstream
- Challenge: How to determine these pre-installing reroute rules



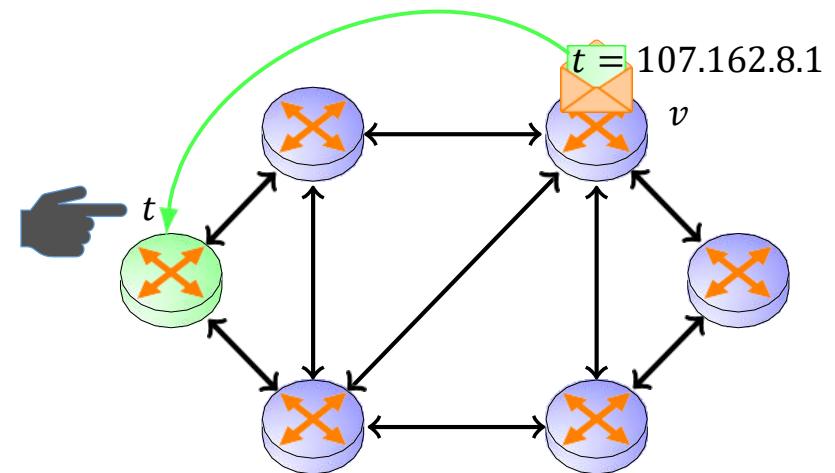
Basic Routing Functions Only Use Local Information

- Network is a connected undirected graph $G=(V,E)$
- Packet **forwarding** can be model as **a function**



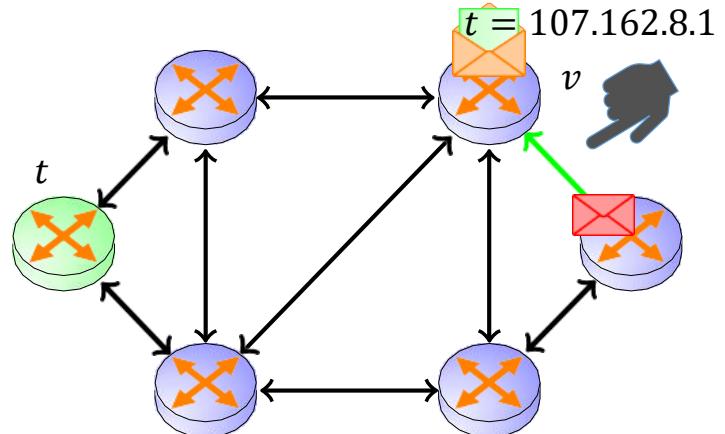
Basic Routing Functions Only Use Local Information

- Network is a connected undirected graph $G=(V,E)$
- **Basic** routing function at node v may only match on:
 - Destination t



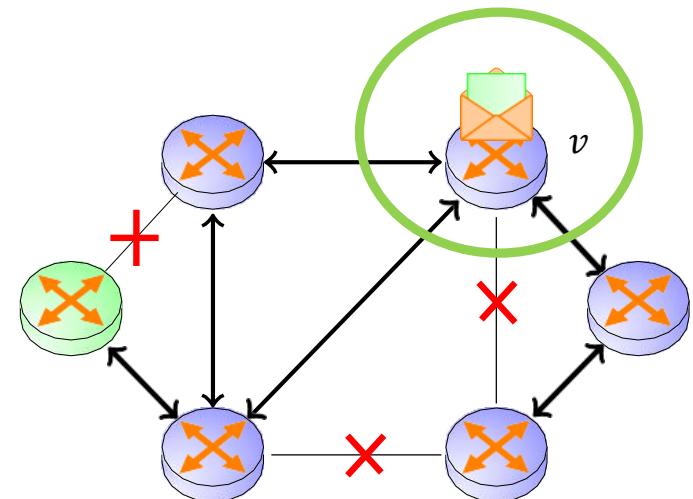
Basic Routing Functions Only Use Local Information

- Network is a connected undirected graph $G=(V,E)$
- **Basic** routing function at node v may only match on:
 - Destination t
 - Incoming port from $E(v) \cup \{\perp\}$




Basic Routing Functions Only Use Local Information

- Network is a connected undirected graph $G=(V,E)$
- **Basic** routing function at node v may only match on:
 - Destination t
 - Incoming port from $E(v) \cup \{\perp\}$
 - Incident edge failures $F \cap E(v)$




Basic Routing Functions Only Use Local Information

- Network is a connected undirected graph $G=(V,E)$
- **Basic** routing function at node v may only match on:
 - Destination t
 - Incoming port from $E(v) \cup \{\perp\}$
 - Incident edge failures $F \cap E(v)$
- To determine an outgoing link in $E(v) - F$
- We know, e.g., how to achieve 1-resilience [Kwong et al. 2011], [Feigenbaum et al. 2012]



Our Focus: Deterministic Forwarding Rules

- Only consider **deterministic** routing functions
- Randomized routing cons:
 - packet recording
 - long paths?
 - device limitations



Spectrum of Models



Basic Routing

Incident
links

Per
destination

Incoming
port

Routing
Function

Outgoing
port

Source-matched Routing

Per
source

Incident
links

Per
destination

Incoming
port

Routing
Function

Outgoing
port

Optional: Header Bit-rewriting

Packet header
rewritable-bits

Inputs of above
routing, e.g., basic

Routing
Function

Outgoing
port

Packet header
rewritable-bits

Spectrum of Models



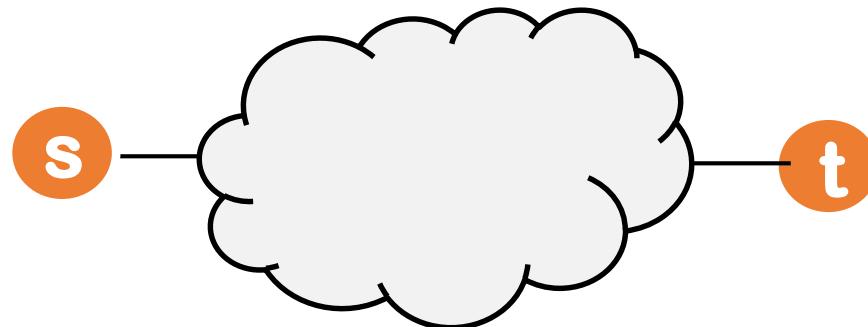
- **The more we match, the larger the tables**
- **Source-matched: $O(n^2)$ (for source and destinations), times failure scenarios (?)**
- **Bits-rewriting: even more overheads (and where to write in header?)**

Problem: Computing a k -Resilient Routing Scheme

- Given an undirected graph $G=(V,E)$ and a destination $t \in V$
- A routing scheme is called k -Resilient ($1 \leq k \leq |E|$) :
 - if any $v \in V$ can reach a destination $t \in V$ **as long as $v-t$ still connected in $G-F$** , for k arbitrary link failures
 $F \subseteq E(G)$
- Goal: Compute a k -resilient routing scheme for given $G=(V,E)$, $t \in V$ and k

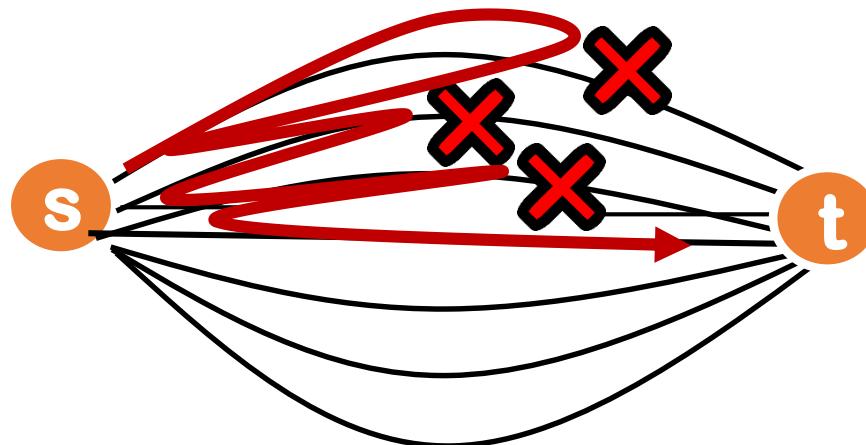
The Model Matters

- Assume k-connected graph: how to tolerate $k-1$ failures?
- If I can match source, import and destination?



The Model Matters

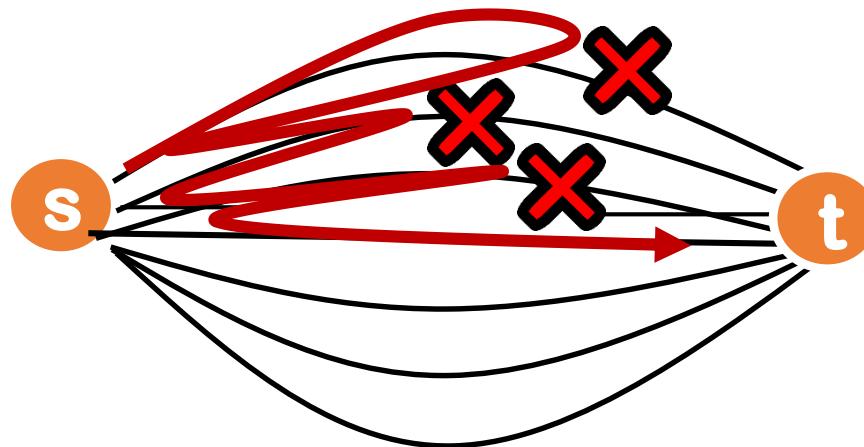
- Assume k-connected graph: how to tolerate k-1 failures?
- If I can match source, import and destination?



- Try one path after the other!

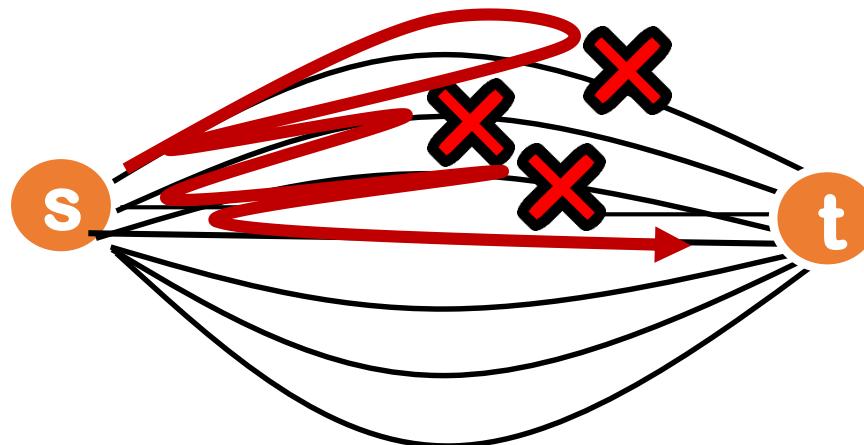
The Model Matters

- Assume k-connected graph: how to tolerate k-1 failures?
- If I can match source, import and destination?



The Model Matters


- Assume k-connected graph: how to tolerate $k-1$ failures?
- If I can match source, import and destination?



- Long-standing conjecture (that it still works)!

Related Works (Static Link Failures)

- Fast failover is a widely studied area, see recent survey
 - Chiesa et al. (2021): *A Survey of Fast Recovery Mechanisms in Packet-Switched Networks*
- Source-matched routing
 - Latest result: *Dai et al. (SPAA' 23): A Tight Characterization of Fast Failover Routing: Resiliency to Two Link Failures is Possible*



The Model of Static Link Failure Is Unrealistic

- Previous research considers **static** link failures, i.e.,
 - **Simultaneous** and **permanent** (*unrecoverable*) Failures

The Model of Static Link Failure Is Unrealistic

- Previous research considers **static** link failures, i.e.,
 - **Simultaneous** and **permanent** (*unrecoverable*) Failures

Reality:

- Links fail *over the time*
- And can recover!

P. Gill et al. Understanding network failures in data centers:
measurement, analysis, and implications (Sigcomm), 2011.

The Model of Static Link Failure Is Unrealistic

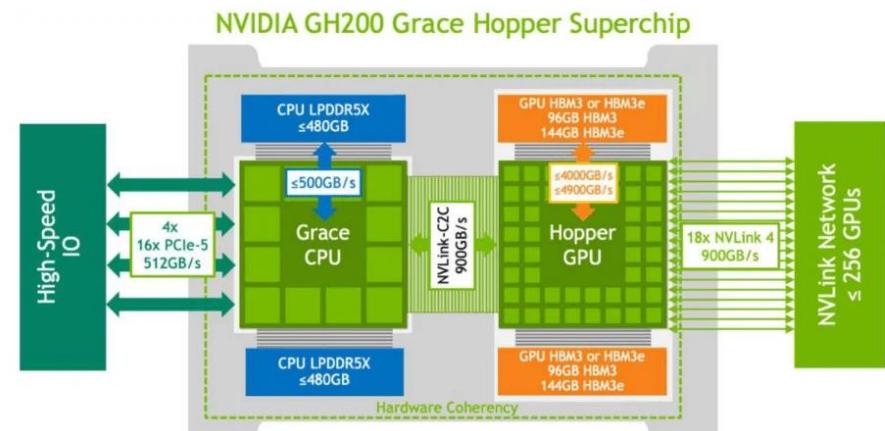
- Previous research considers **static** link failures, i.e.,
 - **Simultaneous** and **permanent (unrecoverable)** Failures

Reality:

- Links fail *over the time*
- And can recover!

P. Gill et al. Understanding network failures in data centers: measurement, analysis, and implications

E.g., NVLink vs Copper



The Novel Models for Link Failures

Challenges: devise
failover reroute rules on
dynamic graphs?

Semi-Dynamic Link Failure

Links F fail *non-simultaneously*,
but *still permanently*.

Dynamic Link Failure

Links F fail *non-simultaneously*,
and *non-permanently*.

First insights: Dai et al. On the Resilience of Fast
Failover Routing Against Dynamic Link Failures (IFIP
Networking), 2025.

The Novel Models for Link Failures

Challenges: devise
failover reroute rules on
dynamic graphs?

Semi-Dynamic Link Failure

Links F fail *non-simultaneously*,
but *still permanently*.

Dynamic Link Failure

Links F fail *non-simultaneously*,
and *non-permanently*.

Algorithm that works for semi-dynamic also works for dynamic!

First insights: Dai et al. On the Resilience of Fast
Failover Routing Against Dynamic Link Failures (IFIP
Networking), 2025.

Related Previous Results and Our Main Contributions


Failure type	Per-destination	Per source	Incoming port	Incident links	k -edge-connected graph	Graph Restrictions	Packet rewriting bits	Resiliency	Ref.
static	X		X	X	$k = 2$	planar		No 2-resilient	[Chiesa et al. 2016]
static	X	X	X	X		general		No 3-resilient	[Dai et al. 2023]
static	X	X	X	X		general		≤ 2 -resilient	[Dai et al. 2023]
static	X		X	X		general	1 bit	≥ 2 -resilient?	Open
static	X	X	X	X		general		≥ 2 -resilient?	Open
dynamic	X		X	X	$k = 2$	planar	0	No 2-resilient	[This Talk]
dynamic	X		X	X		general	$\log k$	No k -resilient	[This Talk]
dynamic	X		X	X		planar	1 or 2	≤ 2 -resilient	[This Talk]

Previous results (highlighted in blue)

Our results (highlighted in green)

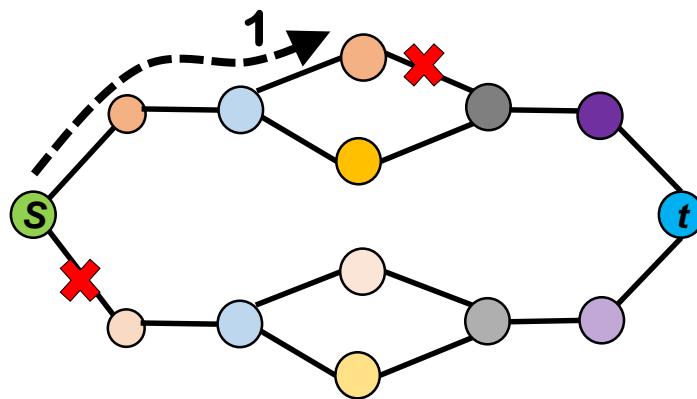
In particular:

- A 2-resilient **basic** routing scheme against **dynamic** failures by **rewriting 1-bit (tight)** on **planar** graphs, **computable** within $O(n + m)$ per destination
- A 2-resilient **source-matched** routing against **dynamic** failures on **planar** graphs **without** rewriting is **impossible**



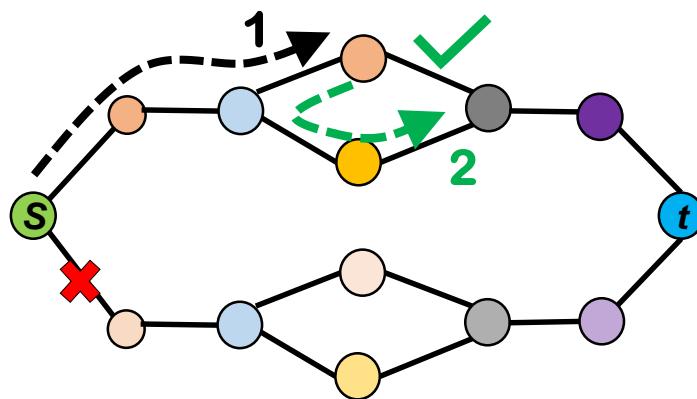
Impossibility: Via Case Distinction

- Impossible to have 2-resilient *source-matched routing* against dynamic failures *without rewriting bits* even on planar graph
 - In the example, every node must use **link-circular routing** (clockwise or counter-clockwise)
 - Depending on order, fail dynamically: *indistinguishable*




Impossibility: Via Case Distinction

- Impossible to have 2-resilient *source-matched routing* against dynamic failures *without rewriting bits* even on planar graph
 - In the example, every node must use **link-circular routing** (clockwise or counter-clockwise)
 - Depending on order, fail dynamically: *indistinguishable*




Impossibility: Via Case Distinction

- Impossible to have 2-resilient *source-matched routing* against dynamic failures *without rewriting bits even on planar graph*
 - In the example, every node must use **link-circular routing** (clockwise or counter-clockwise)
 - Depending on order, fail dynamically: *indistinguishable*





Dynamic 2-Resilient, 2-Bit Rewriting Algorithm: Via Ear Decomposition

- We can focus on 2-connected graphs wlog
- Ear decomposition of G is a sequence of subgraphs $G = P_0 \cup P_1 \dots P_k$,
 - P_0 is a cycle including a fixed node (destination) t
 - Each P_i is an ear (path or cycle) with only its endpoints in $P_0 \cup P_1 \dots P_{i-1}$: endpoints of each ear belong to earlier ears in the sequence but internal vertices of each ear do not




Directed Routing on Faces

- Idea: Can traverse the **boundary** in the **anti-clockwise and clockwise** direction (depending on failures)





- **Right** routing on faces
(clockwise)




- **left** routing on faces
(anti-clockwise)

Can do this for each ear...

- For each ear, left/right routing traverse it in anti-clockwise/clockwise direction




- Right routing on ears (clockwise)



- left routing on ears (anti-clockwise)

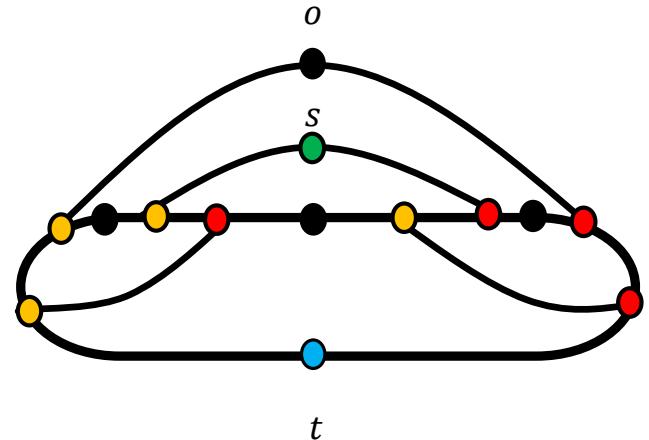
Idea: Precompute important nodes: left/right-most nodes for each phase

- For each face, $\exists P_i$ (**ear**) that shares **common edges**
- Precompute: **left-most** (yellow-colored) and **right-most** (red-colored)



Idea: Precompute important nodes: left/right-most nodes for each phase

- For each face, $\exists P_i$ (**ear**) that shares **common edges**
- Precompute: **left-most** (yellow-colored) and **right-most** (red-colored)
- **Complexity:** compute ear-decomposition of 2-edge-connected planar graph and left/right routes on ears and faces, left/right-most nodes: **time** $O(n + m)$ per destination

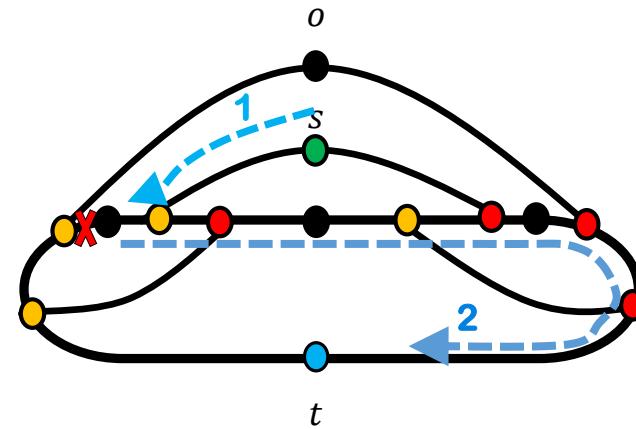


Idea of the algorithm

- Route packets **through ears** in a "waterfall manner" towards destination. Upon failure, **change direction**.
- If every ear contains **at most one failure**, the packet eventually reaches t
- If a **second failure occurs on the same ear**, routing switches to face-based mode (**1 bit**) on the **adjacent faces**
- The algorithm first attempts both directions (now using **1 more bit**)
- Will **reach a right/left-most node**


How to Route by Rewriting One Bit (Ear-Based Routing)

- Starting at the **original node**, initially $\beta = 0$
- When bit $\beta = 0$, routing on ears,
 - **first left and then right direction** upon any failure
 - **Successful**, if seeing ≤ 1 failures for ear-based routing
-




How to Route by Rewriting One Bit (Ear-Based Routing)

- Starting at the **original node**, initially $\beta = 0$
- When bit $\beta = 0$, routing on ears,
 - **first left and then right direction** upon any failure
 - **Successful**, if seeing ≤ 1 failures for ear-based routing
 - Otherwise, when hitting the **second** failure, let $\beta = 1$




How to Route by Rewriting One Bit (Face-Based Routing)

- When hitting the **second** failure (u, v) , $\beta = 1$
- Right routing on the face f_i including (v, u) until stopping the **right-most** (red) node ($->3$)
- Otherwise, seeing another failure (a, b) during the right routing on f_i ($->3$)



How to Route by Rewriting One Bit (Face-Based Routing)

- When hitting the **second** failure (u, v) , $\beta = 1$
- Right routing on the face f_i including (v, u) until stopping the **right-most** (red) node ($\rightarrow 3$)
- Otherwise, seeing another failure (a, b) during the right routing on f_i ($\rightarrow 3$)
- Left routing on the face f_j including (b, a) until stopping the **left-most** (yellow) node ($\rightarrow 4$)
- Now, let $\beta = 0$ to **restart the ear-based routing** ($\rightarrow 5$), always successful (left/right-most node implies left/right routing on the ear)





Thank you!

Failure type	Per-destination	Per source	Incoming port	Incident links	k -edge-connected graph	Graph Restrictions	Packet rewriting bits	Resiliency	Ref.
static	X		X	X	$k = 2$	planar		No 2-resilient	[Chiesa et al. 2016]
static	X	X	X	X		general		No 3-resilient	[Dai et al. 2023]
static	X	X	X	X		general		≤ 2 -resilient	[Dai et al. 2023]
static	X		X	X		general	1 bit	≥ 2 -resilient?	Open
static	X	X	X	X		general		≥ 2 -resilient?	Open
dynamic	X		X	X	$k = 2$	planar	0	No 2-resilient	[This Talk]
dynamic	X		X	X		general	$\log k$	No k -resilient	[This Talk]
dynamic	X		X	X		planar	1 or 2	≤ 2 -resilient	[This Talk]

Previous results (highlighted with a red border)

Our results (highlighted with a green border)