

Principles of Distributed Network Design

Stefan Schmid ***et al.***, mainly Chen Avin (BGU, Israel)



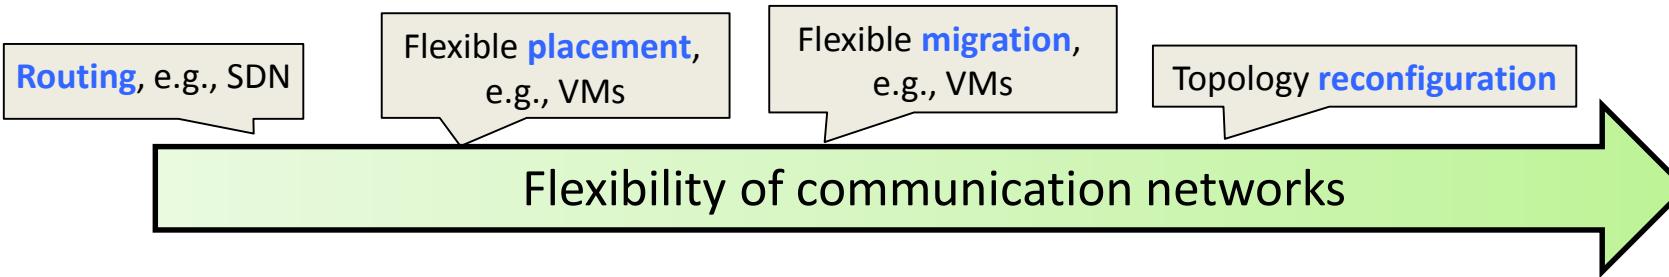
Flexible Distributed Systems: A Great Time To Be a Researcher!



Rhone and Arve Rivers,
Switzerland

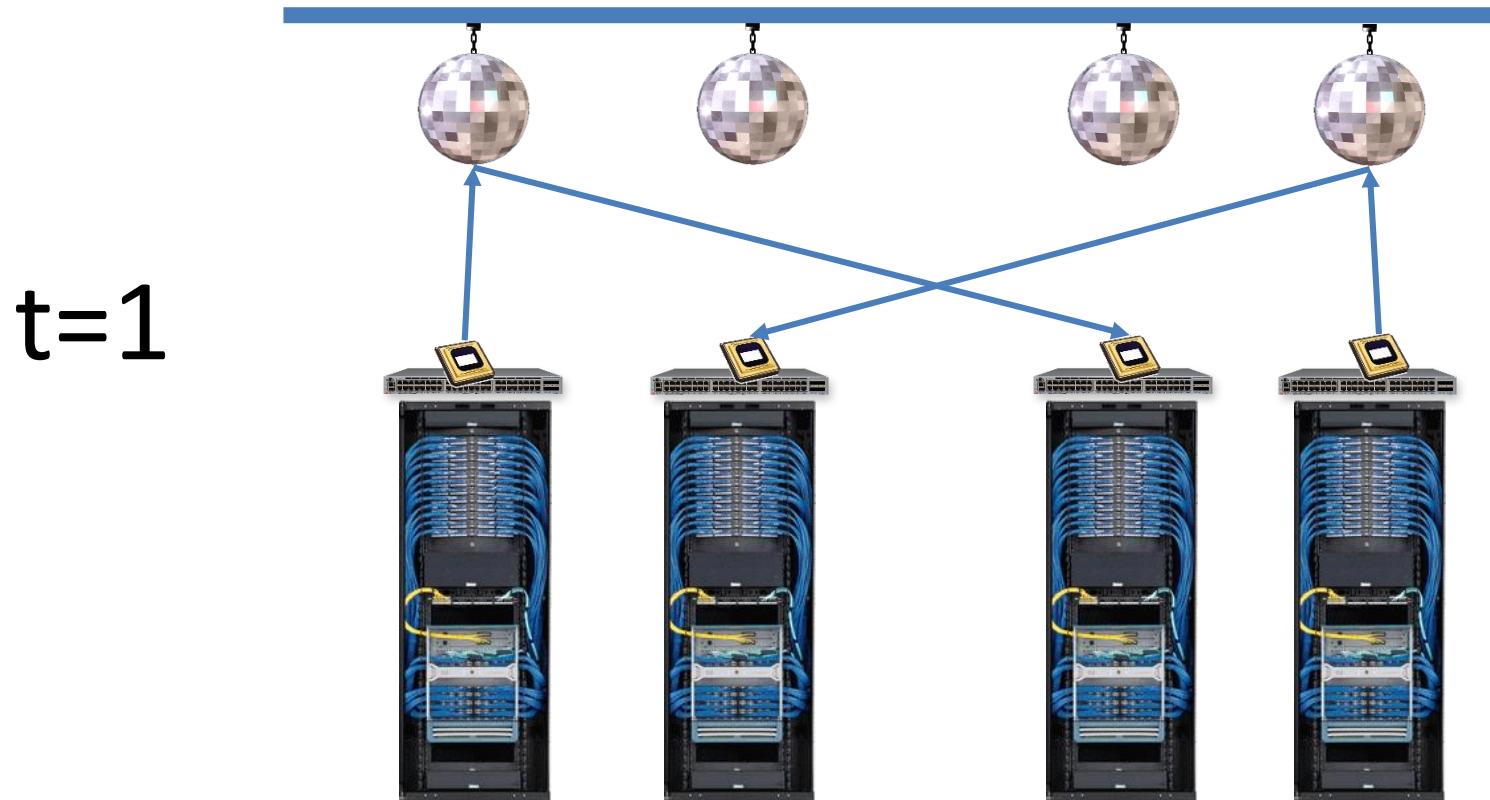
Credits: George
Varghese.

*The new
frontier!*



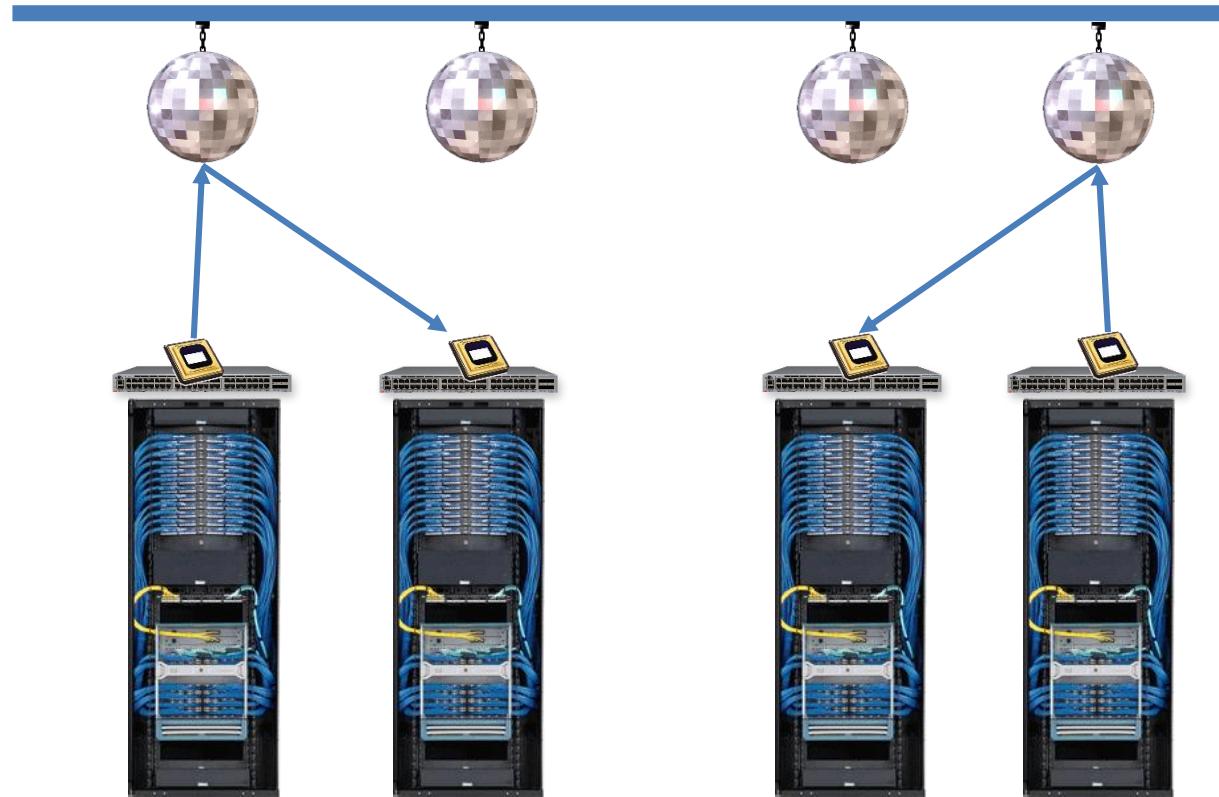
Technological Motivation

Example: Free-Space Optics (*ProjecToR*)



Example: Free-Space Optics (*ProjecToR*)

$t=2$



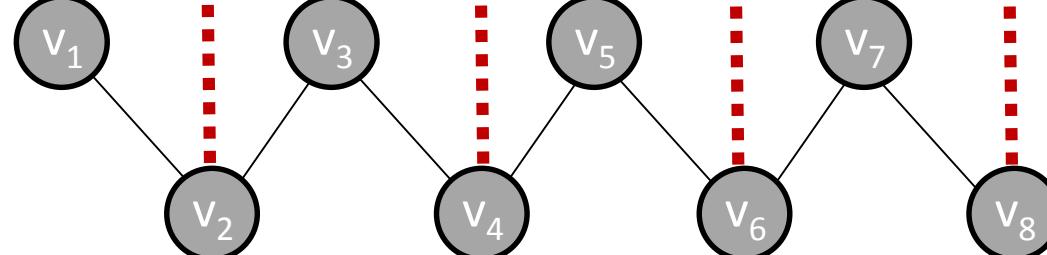
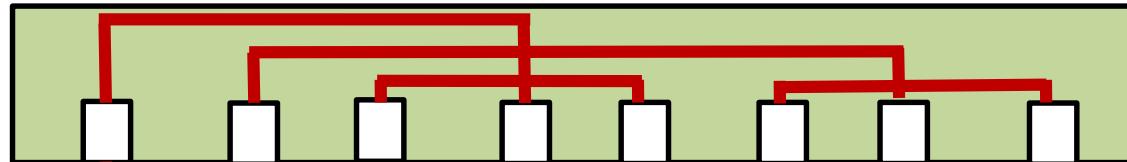
Example: Reconfigurable Optical Switches (*Helios*, *c-Through*, etc.)

Matching!

Dynamic topology:
optical switch
(e.g. matching)

$t=1$

Static topology:
electric



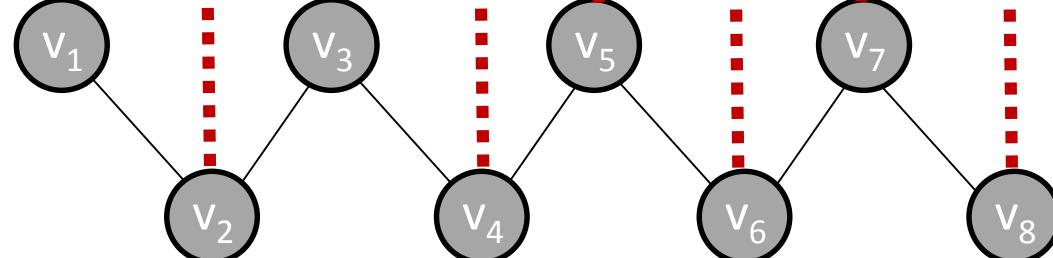
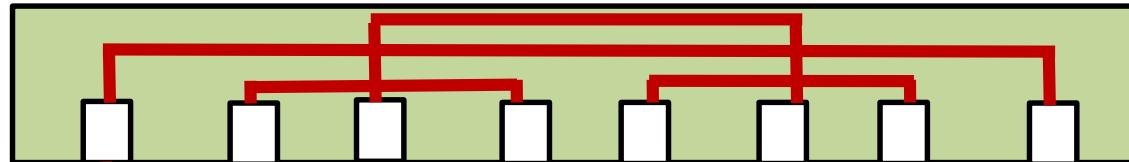
Example: Reconfigurable Optical Switches (*Helios*, *c-Through*, etc.)

Matching!

Dynamic topology:
optical switch
(e.g. matching)

$t=2$

Static topology:
electric



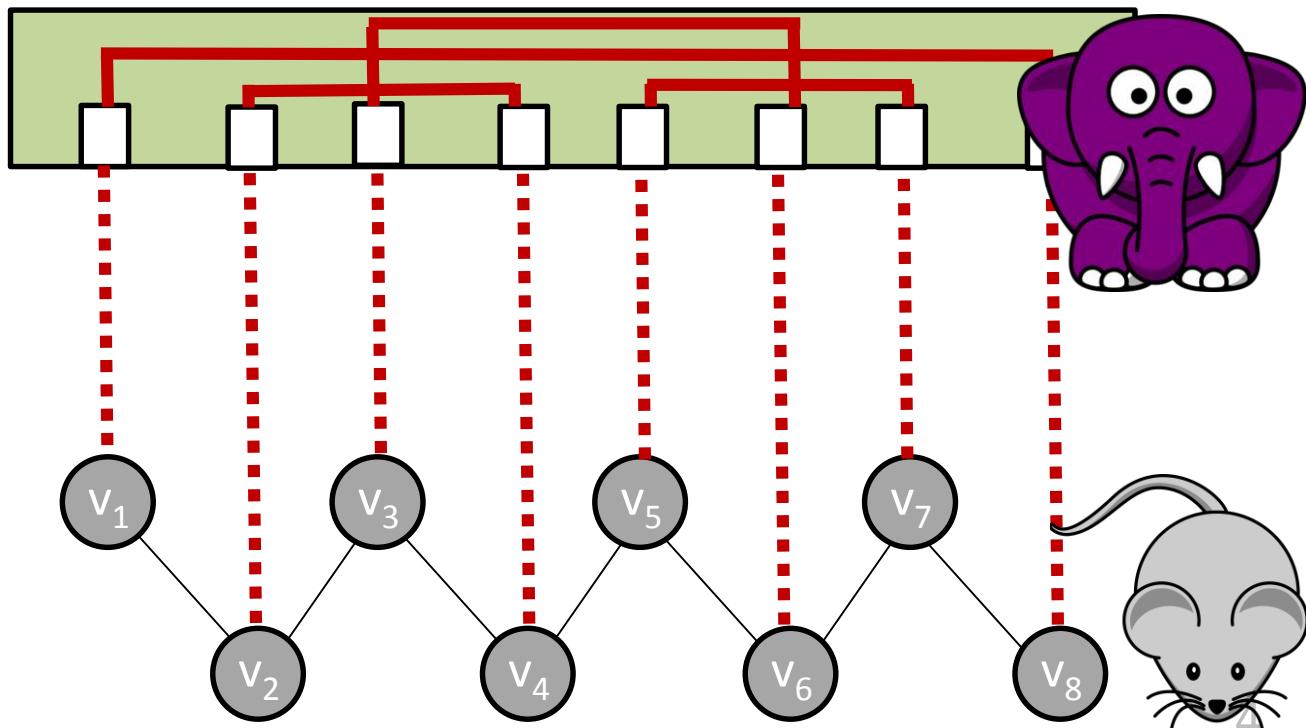
Example: Reconfigurable Optical Switches (*Helios*, *c-Through*, etc.)

Matching!

Dynamic topology:
optical switch
(e.g. matching)

$t=2$

Static topology:
electric



Further Reading

Free-Space Optics

- Ghobadi et al., “**Projector**: Agile reconfigurable data center interconnect,” SIGCOMM 2016.
- Hamedazimi et al. “**Firefly**: A reconfigurable wireless data center fabric using free-space optics,” CCR 2014.

Optical Circuit Switches

- Farrington et al. “**Helios**: a hybrid electrical/optical switch architecture for modular data centers,” CCR 2010.
- Mellette et al. “**Rotornet**: A scalable, low-complexity, optical datacenter network,” SIGCOMM 2017.
- Farrington et al. “Integrating microsecond circuit switching into the data center,” SIGCOMM 2013.
- Liu et al. “Circuit switching under the radar with reactor.,” NSDI 2014

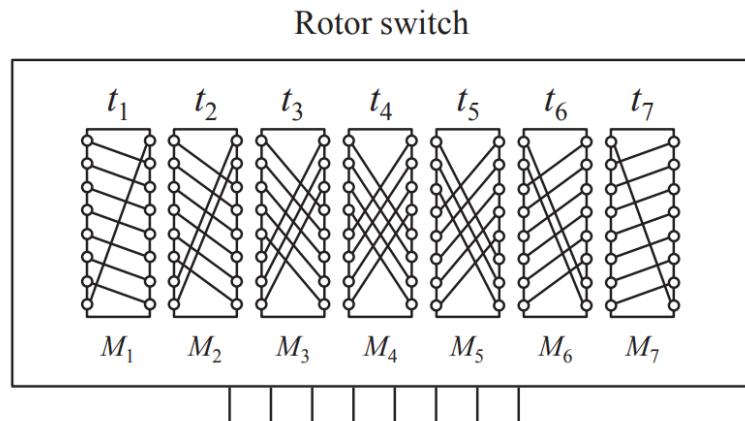
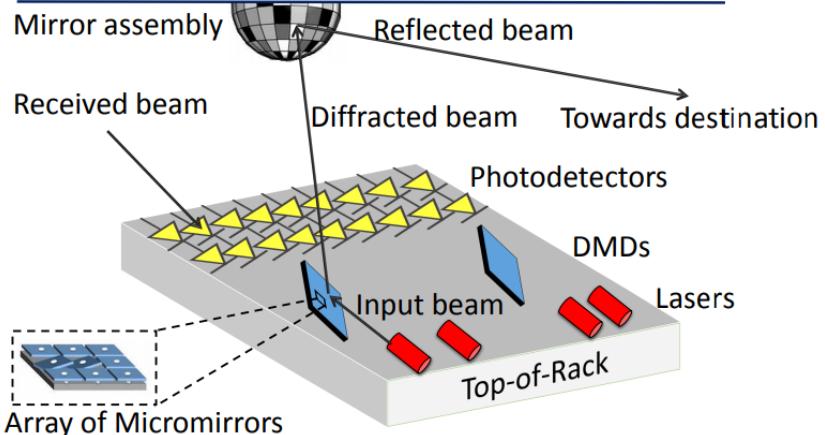
Movable Antennas

- Halperin et al. “Augmenting data center networks with multi-gigabit wireless links,” SIGCOMM 2011.

60GHz Wireless Communication

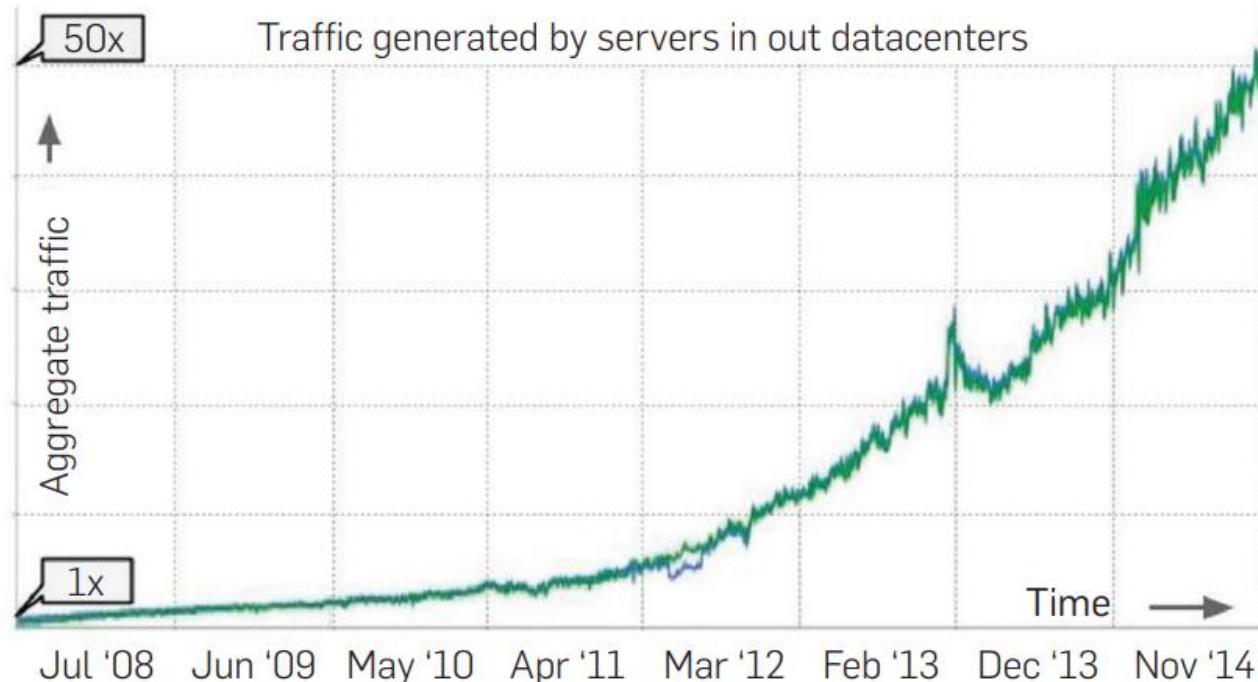
- Zhou et al. “Mirror mirror on the ceiling: Flexible wireless links for data centers,” CCR 2012.
- Kandula et al. “**Flyways** to de-congest data center networks,” 2009.

Etc.!



Empirical Motivation

Data-Centric Applications: Growing Traffic...



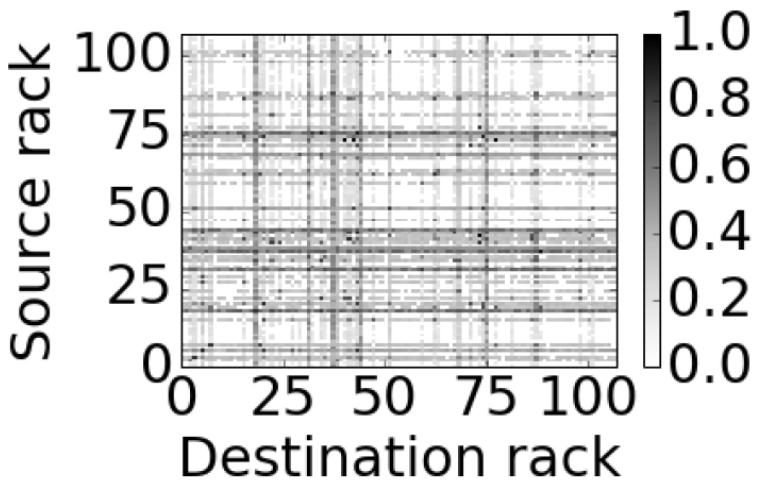
Aggregate server traffic in
Google's datacenter fleet

Source: Jupiter Rising.

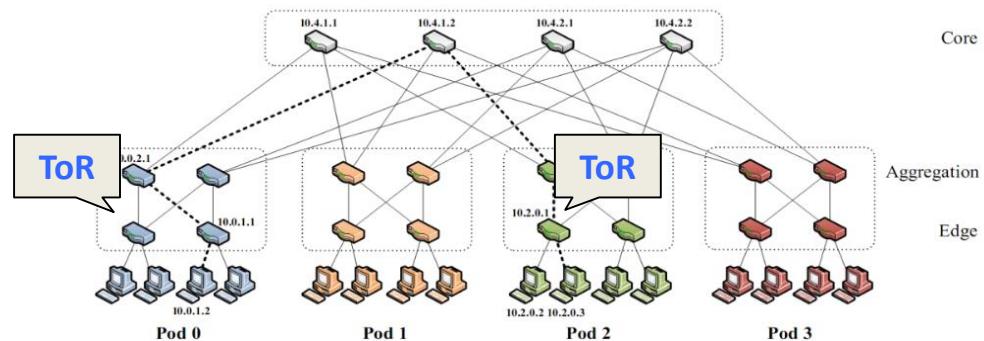
SIGCOMM 2015.

... But Much Structure!

Heatmap rack-to-rack
traffic:



Clos topology and ToR
switches:



The Idea

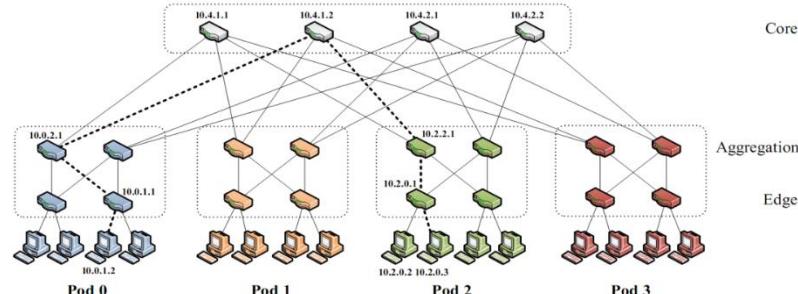


Traditional Networks

Demand-Oblivious

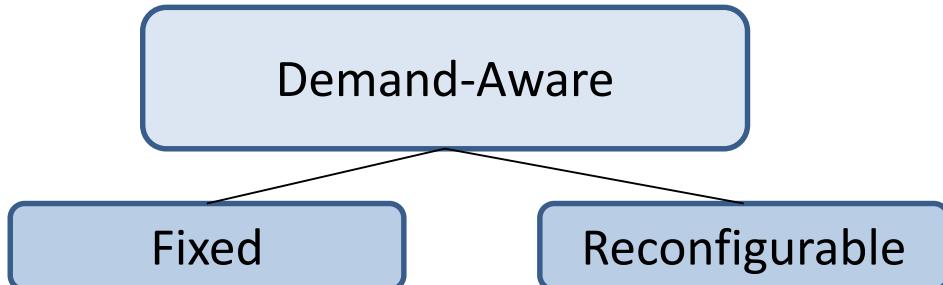
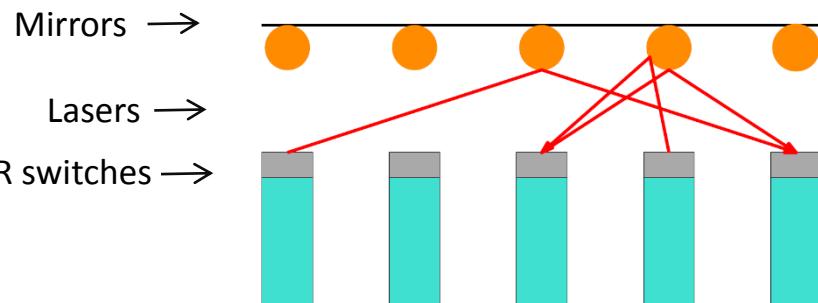
Fixed

- Usually optimized for the “worst-case” (**all-to-all** communication)
- Example, fat-tree topologies: provide **full bisection bandwidth**
- Lower bounds and hard **trade-offs**, e.g., degree vs diameter



Demand-Aware Networks

- **DAN**: Demand-Aware Network
 - Statically optimized **toward the demand**
- **SAN**: Self-Adjusting Network
 - **Dynamically optimized toward the (time-varying) demand**



Roadmap

- Vision and Motivation
- An analogy: coding and datastructures
- Principles of Demand-Aware Network (DAN) Designs
- Principles of Self-Adjusting Network (SAN) Designs
- Principles of Decentralized Approaches



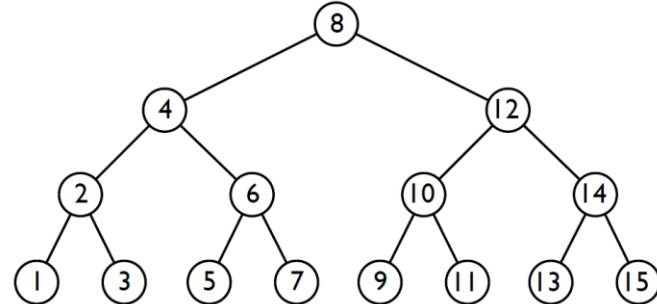
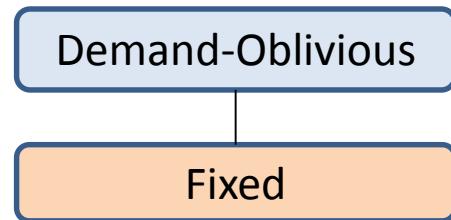
Roadmap

- Vision and Motivation
- **An analogy: coding and datastructures**
- Principles of Demand-Aware Network (DAN) Designs
- Principles of Self-Adjusting Network (SAN) Designs
- Principles of Decentralized Approaches



Analogous to *Datastructures*: Oblivious...

- Traditional, **fixed** BSTs do not rely on any assumptions on the demand
- Optimize for the **worst-case**



Analogous to *Datastructures*: Oblivious...

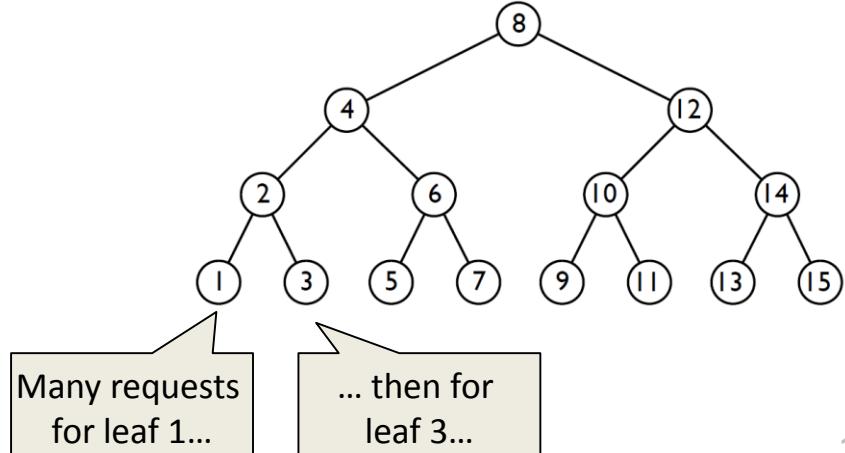
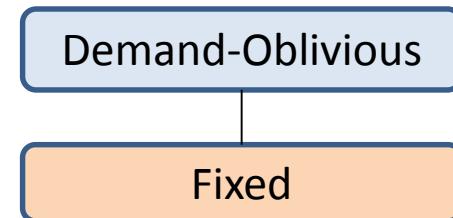
- Traditional, **fixed** BSTs do not rely on any assumptions on the demand

- Optimize for the **worst-case**

- Example **demand**:

$1, \dots, 1, 3, \dots, 3, 5, \dots, 5, 7, \dots, 7, \dots, \log(n), \dots, \log(n)$
 $\longleftrightarrow \longleftrightarrow \longleftrightarrow \longleftrightarrow \longleftrightarrow \quad \leftarrow \qquad \qquad \qquad \qquad \qquad \text{many} \quad \text{many} \quad \text{many} \quad \text{many} \quad \text{many}$

- Items stored at **$O(\log n)$** from the root, **uniformly** and **independently** of their frequency



Analogous to *Datastructures*: Oblivious...

- Traditional, **fixed** BSTs do not rely on any assumptions on the demand

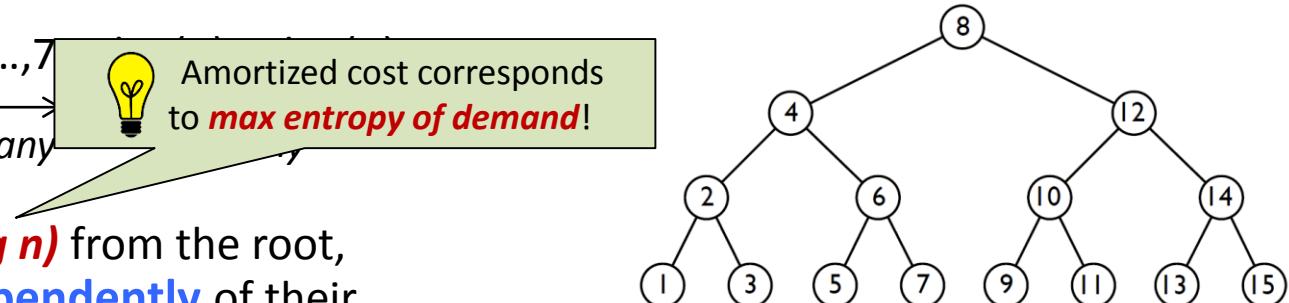
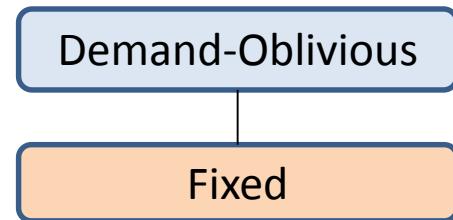
- Optimize for the **worst-case**

- Example **demand**:

1,...,1,3,...,3,5,...,5,7,...,7
↔ ↔ ↔ ↔ ↔
many many many many

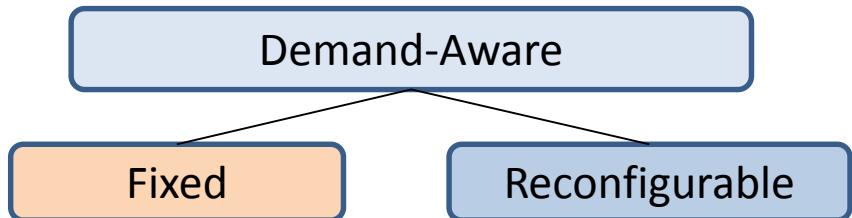
Amortized cost corresponds to **max entropy of demand!**

- Items stored at **$O(\log n)$** from the root, **uniformly** and **independently** of their frequency



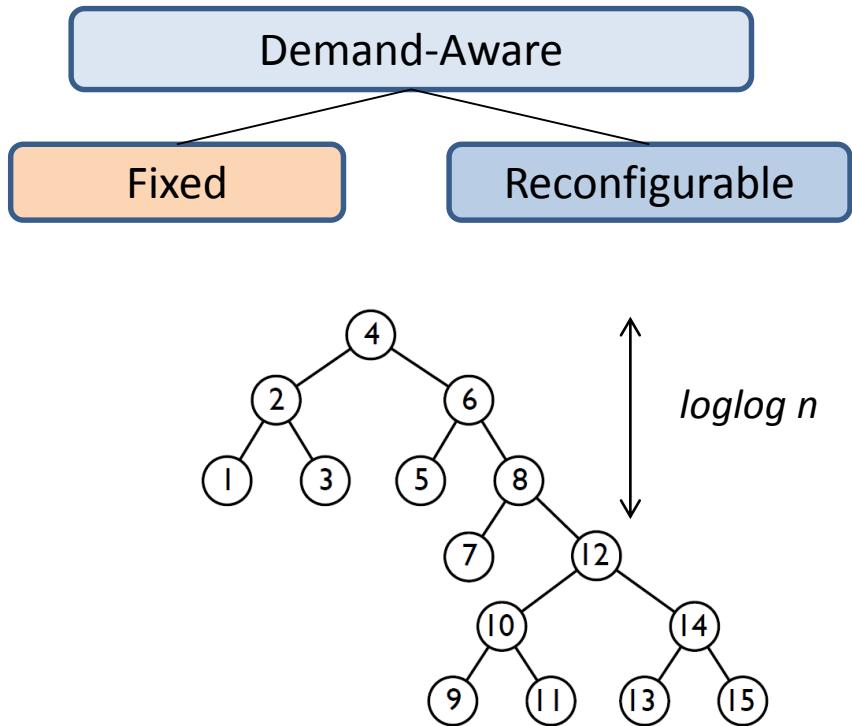
... Demand-Aware ...

- **Demand-aware fixed** BSTs can take advantage of *spatial locality* of the demand
- E.g.: place frequently accessed elements close to the root
- E.g., Knuth/Mehlhorn/Tarjan trees



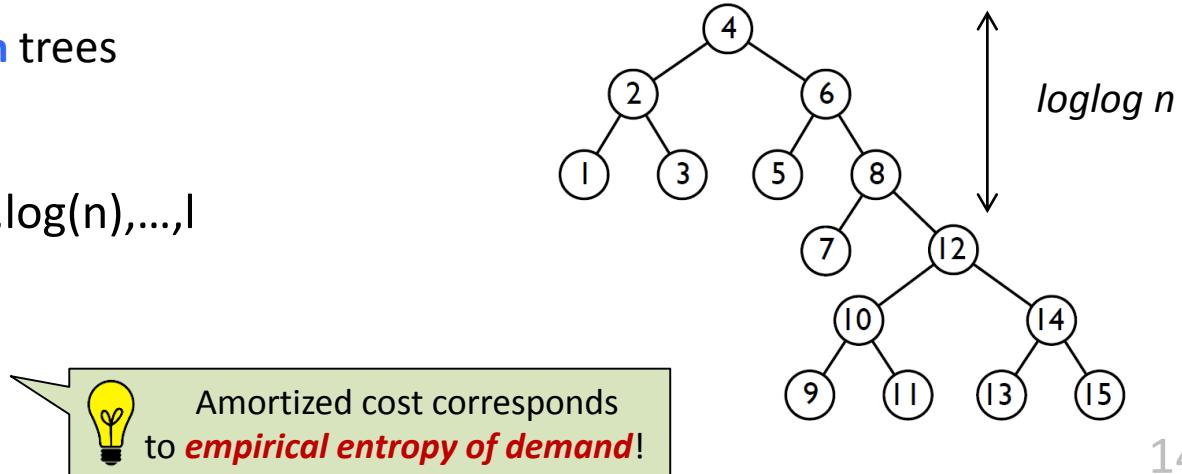
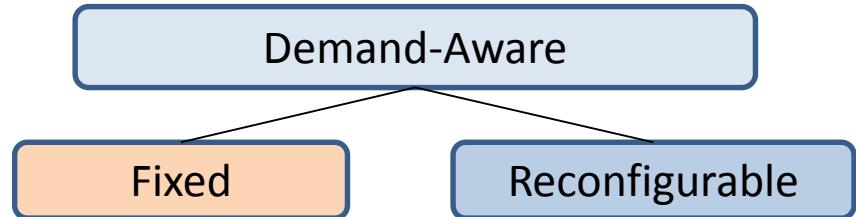
... Demand-Aware ...

- **Demand-aware fixed** BSTs can take advantage of *spatial locality* of the demand
- E.g.: place frequently accessed elements close to the root
- E.g., Knuth/Mehlhorn/Tarjan trees
- Recall example **demand**:
 $1, \dots, 1, 3, \dots, 3, 5, \dots, 5, 7, \dots, 7, \dots, \log(n), \dots, \log(n)$
 - Amortized cost $O(\log \log n)$



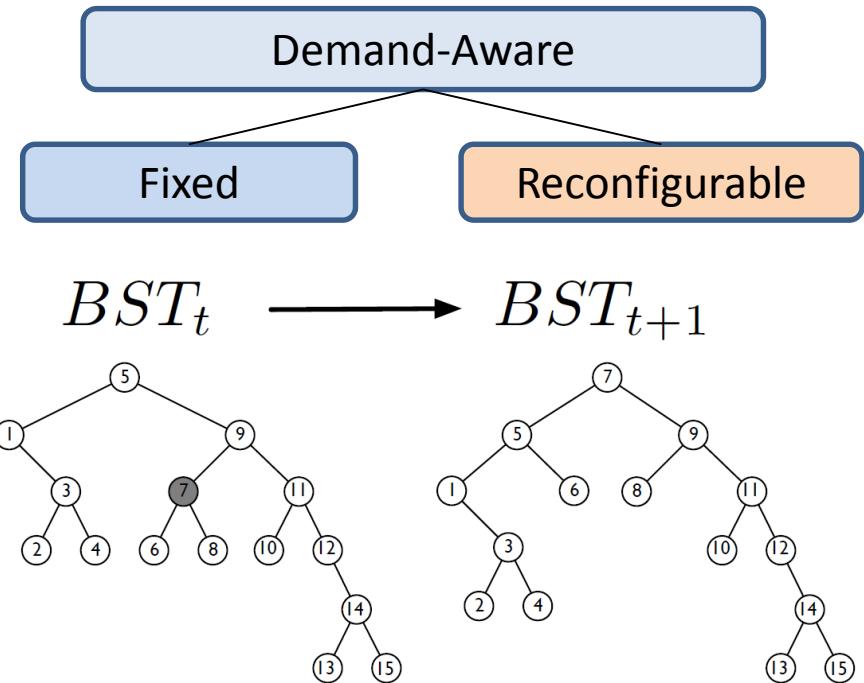
... Demand-Aware ...

- **Demand-aware fixed** BSTs can take advantage of *spatial locality* of the demand
- E.g.: place frequently accessed elements close to the root
- E.g., Knuth/Mehlhorn/Tarjan trees
- Recall example **demand**:
 $1, \dots, 1, 3, \dots, 3, 5, \dots, 5, 7, \dots, 7, \dots, \log(n), \dots, \log(n)$
 - Amortized cost $O(\log \log n)$



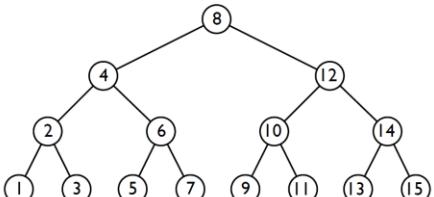
... Self-Adjusting!

- Demand-aware reconfigurable BSTs can additionally take advantage of *temporal locality*
- By moving accessed element to the root: amortized cost is *constant*, i.e., $O(1)$
 - Recall example demand:
 $1, \dots, 1, 3, \dots, 3, 5, \dots, 5, 7, \dots, 7, \dots, \log(n), \dots, \log(n)$
- Self-adjusting BSTs e.g., useful for implementing *caches* or garbage collection

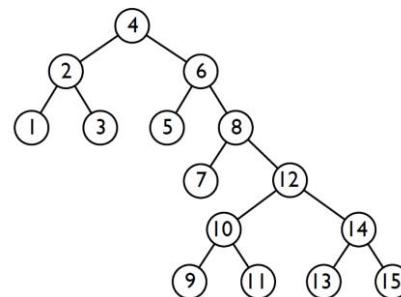


Datastructures

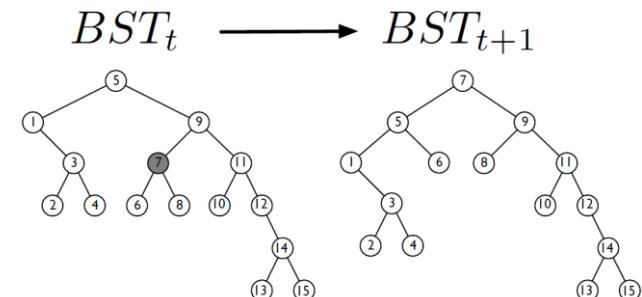
Oblivious



Demand-Aware



Self-Adjusting



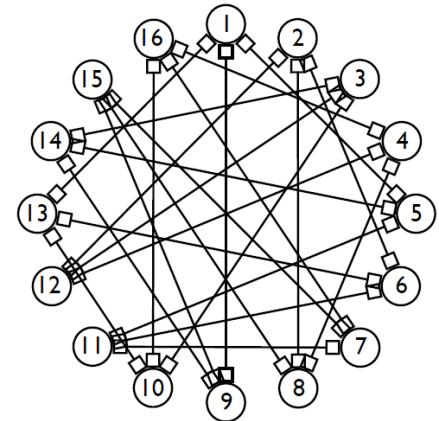
Lookup $O(\log n)$

Exploit **spatial locality**:
empirical entropy $O(\log \log n)$

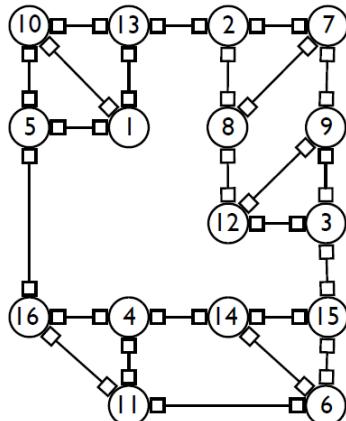
Exploit **temporal locality** as well:
 $O(1)$

Analogously for Networks

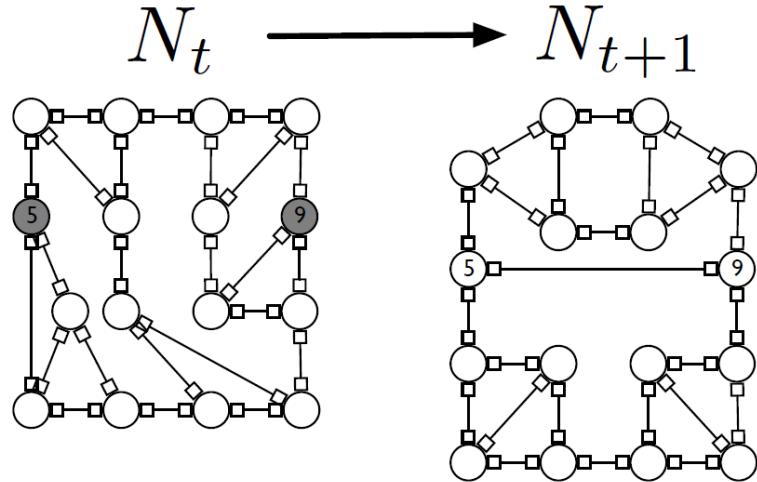
Oblivious



DAN



SAN



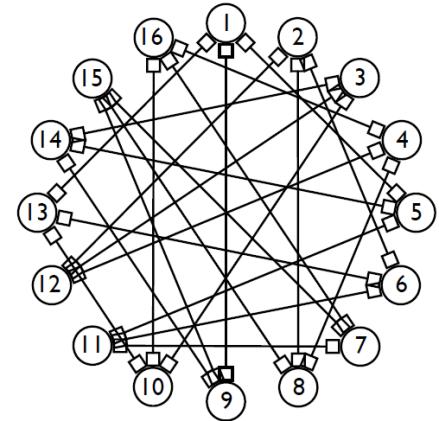
Const degree
(e.g., **expander**):
route lengths $O(\log n)$

Exploit **spatial locality**: Route
lengths depend on
conditional entropy of demand

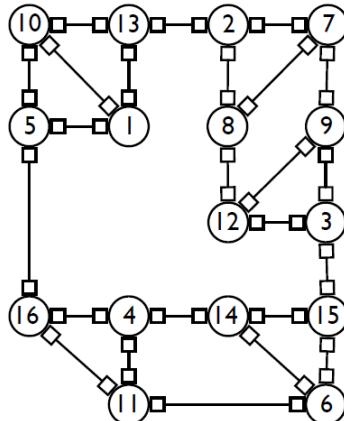
Exploit **temporal locality** as well

Analogously for Networks

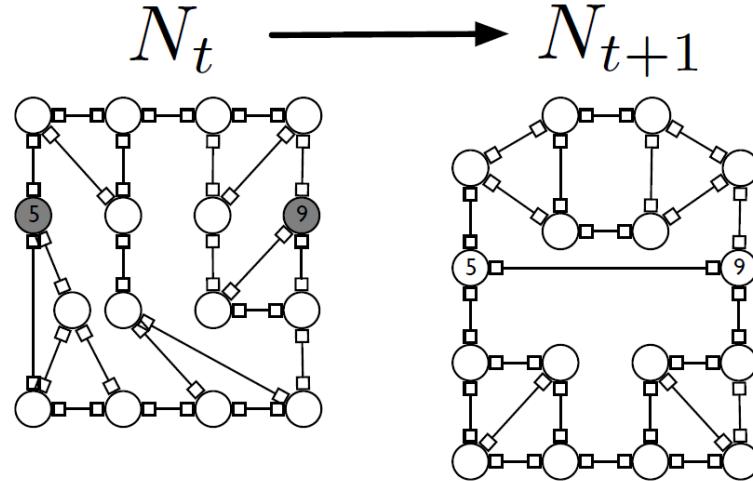
Oblivious



DAN



SAN



Const degree
(e.g., **expander**):
route lengths **$O(\log n)$**

Exploit **spatial locality**: Route
lengths depend on
conditional entropy of demand

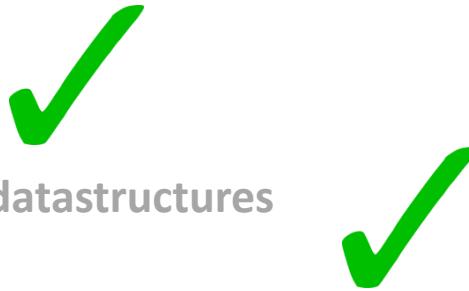
Exploit **temporal locality** as well

How useful are DANs/SANs?

As always in computer science (e.g., also in coding, in self-adjusting datastructures, etc.): it depends! ☺

Roadmap

- Vision and Motivation
- An analogy: coding and datastructures
- **Principles of Demand-Aware Network (DAN) Designs**
- Principles of Self-Adjusting Network (SAN) Designs
- Principles of Decentralized Approaches



The DAN Design Problem

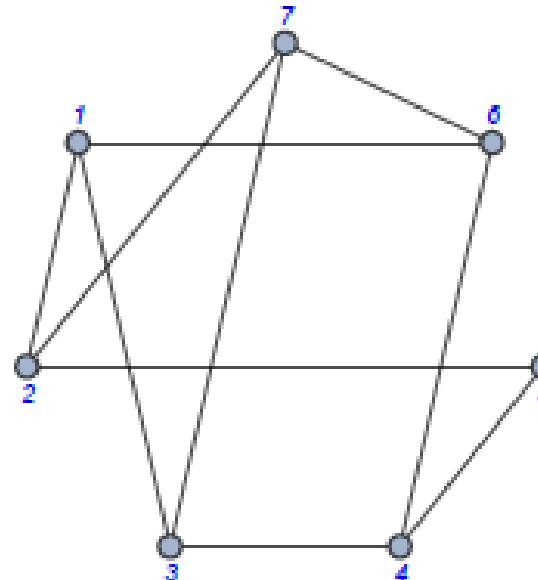
Input: Workload

Output: DAN

Destinations

Sources	1	2	3	4	5	6	7
1	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
2	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$
3	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	0	$\frac{1}{13}$
4	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0
5	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0	0	0
6	$\frac{2}{65}$	0	0	0	0	0	$\frac{3}{65}$
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0

design



Demand matrix: joint distribution

... of constant degree (**scalability**)

The DAN Design Problem

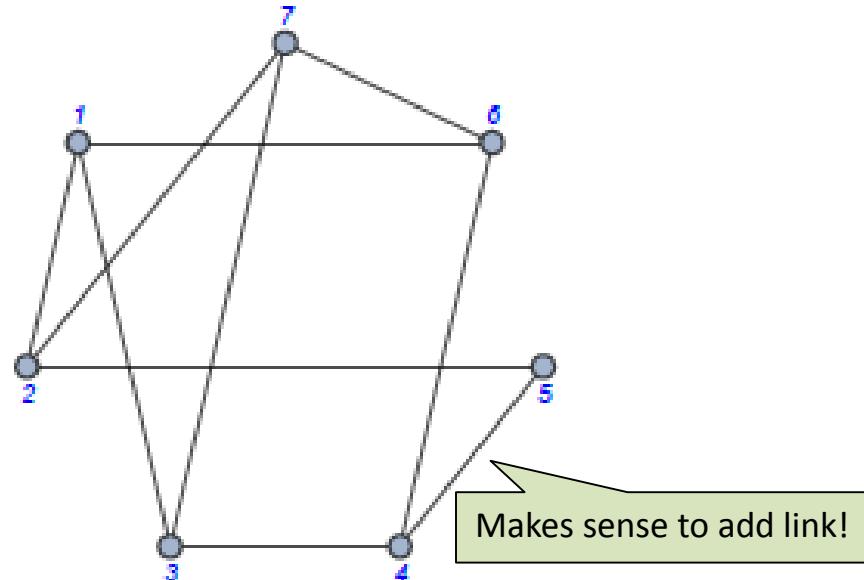
Input: Workload

Output: DAN

Destinations

Sources	1	2	3	4	5	6	7
1	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
2	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$
3	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{13}{65}$	0
4	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0
5	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0	0	0
6	$\frac{2}{65}$	0	0	0	0	0	$\frac{3}{65}$
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0

design



Demand matrix: joint distribution

... of constant degree (**scalability**)

The DAN Design Problem

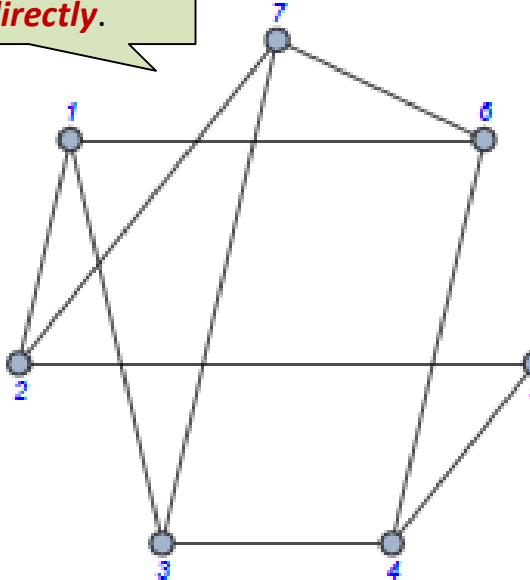
Input: Workload

Destinations								
		1 communicates to many.						
		6	7					
1	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$	
2	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$	
3	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	0	$\frac{1}{13}$	
4	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0	
5	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0	0	0	
6	$\frac{2}{65}$	0	0	0	0	0	$\frac{3}{65}$	
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0	

Output: DAN

Bounded degree: route
to 7 *indirectly*.

design



Demand matrix: joint distribution

... of constant degree (**scalability**)

The DAN Design Problem

Input: Workload

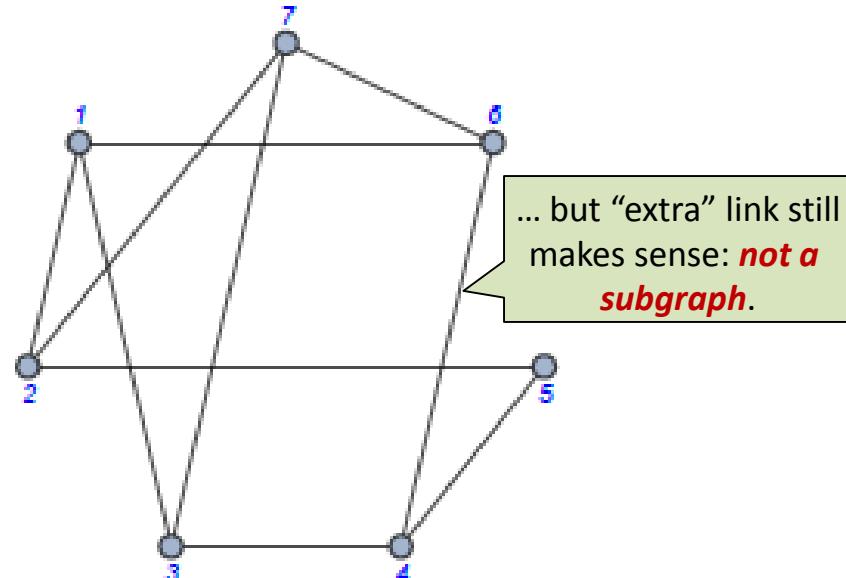
Output: DAN

Destinations

Sources	1	2	3	4	5	6	7
1	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
2	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$
3	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	0	$\frac{1}{13}$
4	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0
5	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0		
6	$\frac{2}{65}$	0	0	0	0		
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0

design

4 and 6 don't communicate...



Demand matrix: joint distribution

... of constant degree (**scalability**)

Case Study: Expected Route Length

Shorter paths: smaller bandwidth footprint, lower latency, less energy, ...

More Formally: DAN Design Problem

Input:

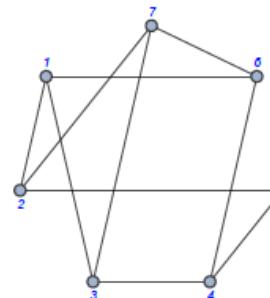
$\mathcal{D}[p(i, j)]$: joint distribution, Δ

		Y						
		1	2	3	4	5	6	7
X	1	0	<u>2</u>	<u>1</u>	<u>1</u>	<u>1</u>	<u>2</u>	<u>3</u>
	2	<u>2</u>	0	<u>1</u>	0	0	0	<u>2</u>
	3	<u>1</u>	<u>1</u>	0	<u>2</u>	0	0	<u>1</u>
	4	<u>1</u>	0	<u>2</u>	0	<u>4</u>	0	0
	5	<u>1</u>	0	<u>3</u>	<u>4</u>	0	0	0
	6	<u>2</u>	0	0	0	0	0	<u>3</u>
	7	<u>3</u>	<u>2</u>	<u>1</u>	0	0	<u>3</u>	0



Output:

N: DAN



Bounded degree $\Delta = 3$

Expected Path Length (EPL): Basic measure of efficiency

$$\text{EPL}(\mathcal{D}, N) = E_{\mathcal{D}}[d_N(\cdot, \cdot)] = \sum_{(u,v) \in \mathcal{D}} p(u, v) \cdot d_N(u, v)$$

Path length *on DAN*.

Frequency

The Goal: Bounded Network Design (BND)

- **Inputs:** Communication distribution $\mathcal{D}[p(i,j)]_{nxn}$ and a maximum degree Δ .
- **Output:** A Demand Aware Network $N \in N_\Delta$ s.t.

$$\text{BND}(\mathcal{D}, \Delta) = \min_{N \in N_\Delta} \text{EPL}(\mathcal{D}, N)$$

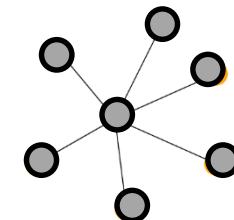
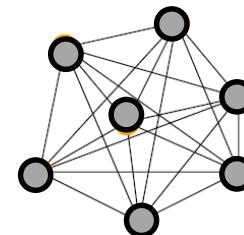
Belong to some graph family:
bounded-degree, but e.g. also **local routability**!

Examples

- What about BND for $\Delta = n$?

Examples

- What about BND for $\Delta = n$?
- Easy: e.g., **clique** and **star** have **constant** EPL but unbounded degree.

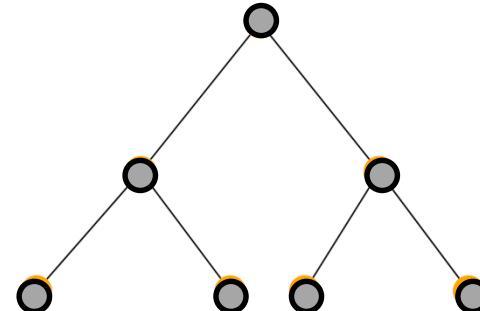


Some Insights

- What about $\Delta = 3$?

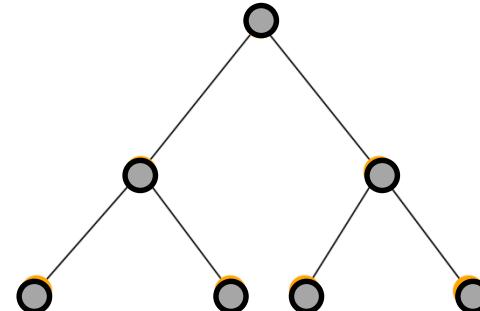
Some Insights

- What about $\Delta = 3$?
 - E.g., complete binary tree
 - $d_N(u,v) \leq 2 \log n$
 - Can we do **better** than **log n**?



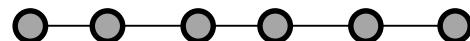
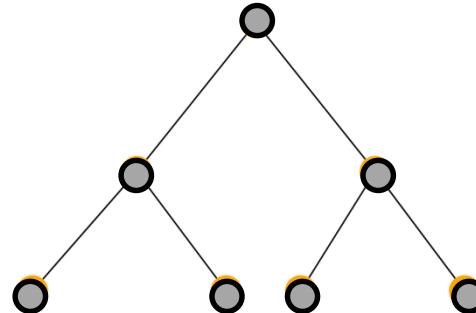
Some Insights

- What about $\Delta = 3$?
 - E.g., complete binary tree
 - $d_N(u,v) \leq 2 \log n$
 - Can we do **better** than **log n**?
- What about $\Delta = 2$?



Some Insights

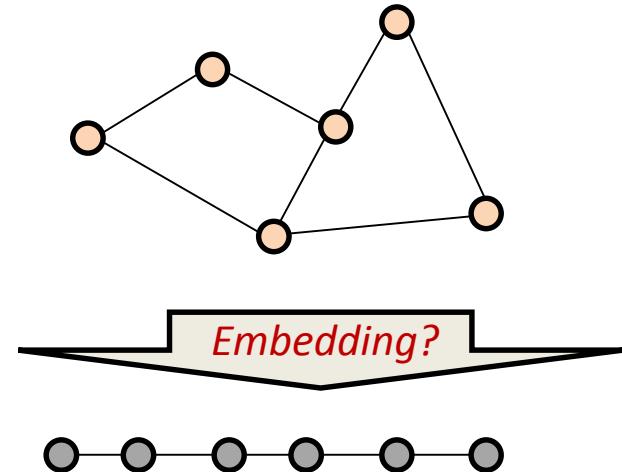
- What about $\Delta = 3$?
 - E.g., complete binary tree
 - $d_N(u,v) \leq 2 \log n$
 - Can we do **better** than **log n**?
- What about $\Delta = 2$?
 - E.g., set of lines and cycles



How hard is it to design a DAN?

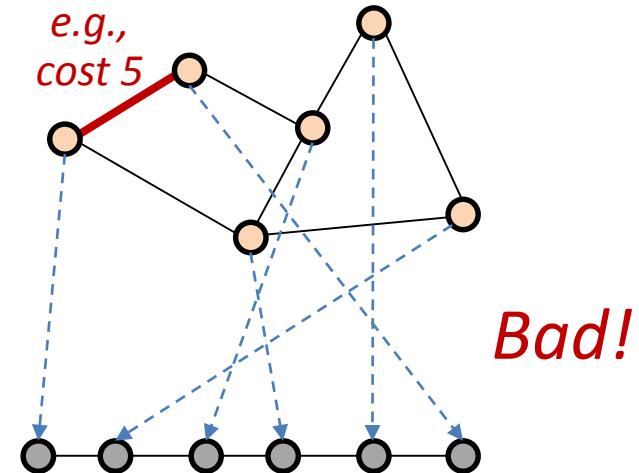
DAN design can be NP-hard

- Example $\Delta = 2$: A Minimum Linear Arrangement (MLA) problem
 - A “Virtual Network Embedding Problem”, VNEP
 - Minimize sum of lengths of virtual edges



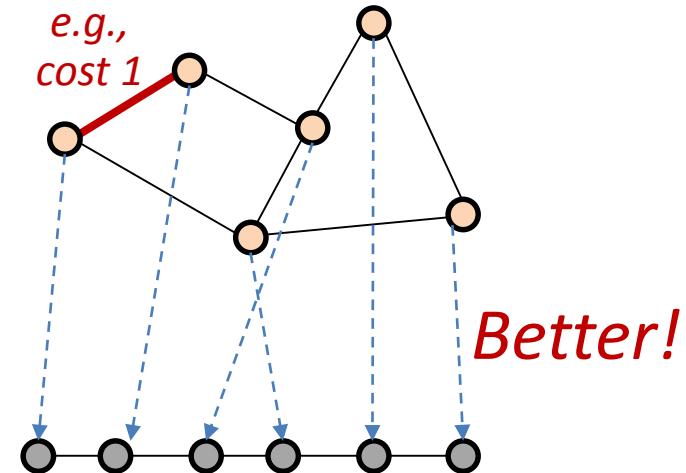
DAN design can be NP-hard

- Example $\Delta = 2$: A Minimum Linear Arrangement (MLA) problem
 - A “Virtual Network Embedding Problem”, VNEP
 - Minimize sum of lengths of virtual edges



DAN design can be NP-hard

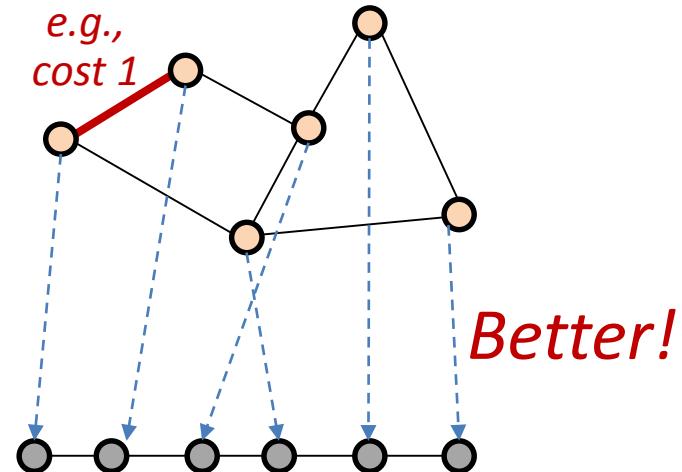
- Example $\Delta = 2$: A Minimum Linear Arrangement (MLA) problem
 - A “Virtual Network Embedding Problem”, VNEP
 - Minimize sum of lengths of virtual edges



DAN design can be NP-hard

- Example $\Delta = 2$: A Minimum Linear Arrangement (MLA) problem
 - A “Virtual Network Embedding Problem”, VNEP
 - Minimizing the lengths of virtual edges

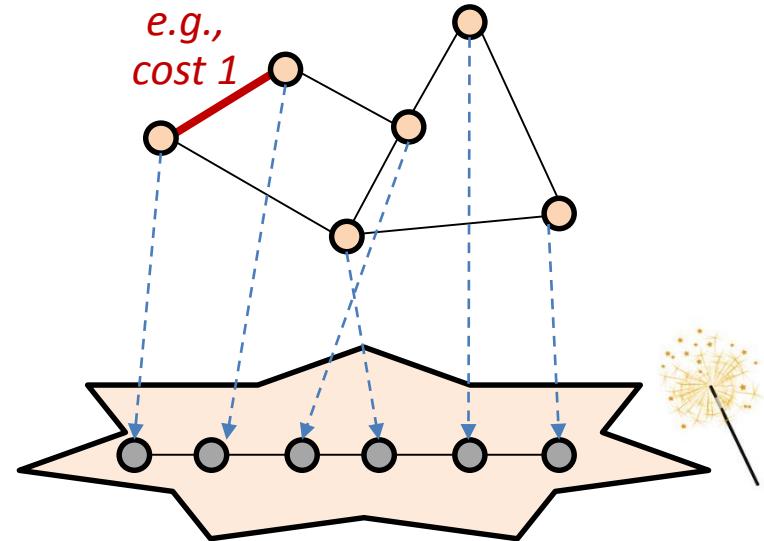
NP-hard, and so is DAN design.



DAN design can be NP-hard

- Example $\Delta = 2$: A Minimum Linear Arrangement (MLA) problem
 - A “Virtual Network Embedding Problem”, VNEP
 - Minimizing the lengths of virtual edges
- But what about > 2 ? Embedding problem still hard, but we have an additional degree of freedom:

Do topological flexibilities make problem easier or harder?!



A new knob for optimization!

Also Related To...:

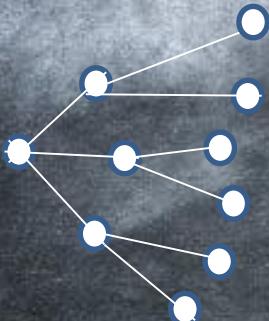
- Sparse, distance-preserving (low-distortion) **spanners**
- But:
 - Spanners aim at low distortion among ***all pairs***; in our case, we are only interested in the ***local distortion***, 1-hop communication neighbors
 - We allow **auxiliary edges** (not a subgraph): similar to **geometric spanners**
 - We require ***constant degree***

Expected Path Length in Traditional Networks?

Theorem (Traditional Networks):

Each network with n nodes and max degree $\Delta > 2$ must have a diameter of at least $\log(n)/\log(\Delta-1)-1$.

Proof.



In k steps, reach at most $1 + \sum \Delta(\Delta - 1)^i$

$$1 \quad \Delta \quad \Delta(\Delta-1)$$

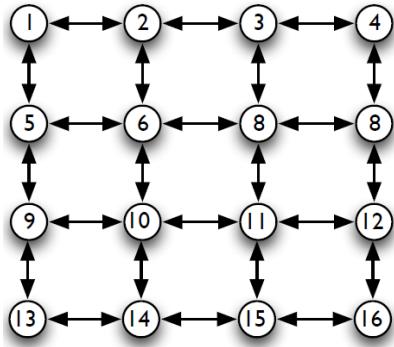
Corollary: Constant-degree graphs have at least logarithmic diameter.

Example: Clos, Bcube, Xpander.

Can DANs do better?

Yes, constant-degree DANs can!

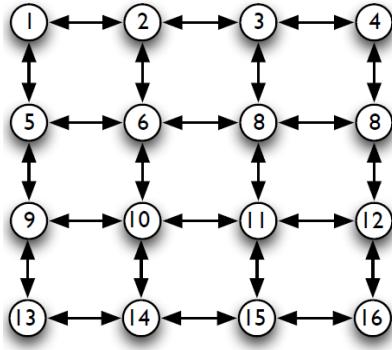
- Example 1: **demand**



- Oblivious design diameter $O(\log n)$
(e.g., Δ -ary tree)
- But constant-degree DAN can serve it at cost $O(1)$.

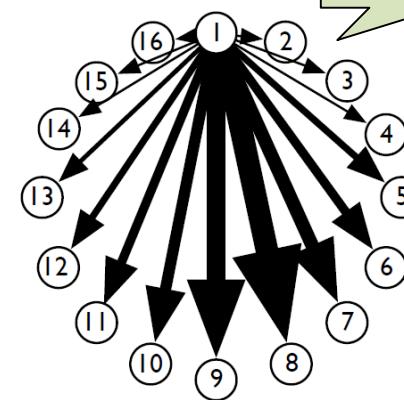
Yes, constant-degree DANs can!

- Example 1: **demand**

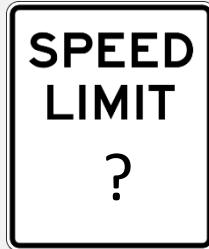


- Oblivious design diameter $O(\log n)$ (e.g., Δ -ary tree)
- But constant-degree DAN can serve it at cost $O(1)$.

- Example 2: **demand**



- Oblivious design: diameter $O(\log n)$ (e.g., Δ -ary tree)
- But constant-degree DAN can serve it at cost $O(1)$ (e.g., **Huffman** tree for node 1)



Limitations of DANs: Lower Bounds on EPL



Lower Bound Idea: Leverage Coding or Datastructure!

Destinations		1	2	3	4	5	6	7
Sources	1	0 65	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
	2	$\frac{2}{65}$	0 65	$\frac{1}{65}$	0 0	0 0	0 0	$\frac{2}{65}$
	3	$\frac{1}{13}$	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 0	0 0	$\frac{1}{13}$
	4	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 65	$\frac{4}{65}$	0 0	0 0
	5	$\frac{1}{65}$	0 65	$\frac{3}{65}$	$\frac{4}{65}$	0 0	0 0	0 0
	6	$\frac{2}{65}$	0 65	0 0	0 0	0 0	0 0	$\frac{3}{65}$
	7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0 0	0 0	$\frac{3}{65}$	0 0

- Consider **(source) node 1**: best Δ -ary tree we can build for this source is **Huffman tree** for its destinations
[0,1/65,1/13,1/65,1/65,2/65,3/65]
 - resp. **Knuth/Mehlhorn/Tarjan** tree if search property required
- How good can this tree be?

Lower Bound Idea: Leverage Coding or Datastructure!

Destinations		1	2	3	4	5	6	7
Sources	1	0 65	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
	2	$\frac{2}{65}$	0 65	$\frac{1}{65}$	0 0	0 0	0 0	$\frac{2}{65}$
	3	$\frac{1}{13}$	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 0	0 0	$\frac{1}{13}$
	4	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 65	$\frac{4}{65}$	0 0	0 0
	5	$\frac{1}{65}$	0 65	$\frac{3}{65}$	$\frac{4}{65}$	0 0	0 0	0 0
	6	$\frac{2}{65}$	0 65	0 0	0 0	0 0	0 0	$\frac{3}{65}$
	7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0 0	0 0	$\frac{3}{65}$	0 0

- Consider **(source) node 1**: best Δ -ary tree we can build for this source is **Huffman tree** for its destinations
[0,1/65,1/13,1/65,1/65,2/65,3/65]
 - resp. **Knuth/Mehlhorn/Tarjan** tree if search property required
- How good can this tree be?



Entropy lower bound known on EPL known for binary trees, e.g. **Mehlhorn** 1975 for BST

Lower Bound Idea: Leverage Coding or Datastructure!

Destinations		1	2	3	4	5	6	7
Sources	1	0 65	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
	2	$\frac{2}{65}$	0 65	$\frac{1}{65}$	0 0	0 0	0 0	$\frac{2}{65}$
	3	$\frac{1}{13}$	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 0	0 0	$\frac{1}{13}$
	4	$\frac{1}{65}$	0 65	$\frac{2}{65}$	0 65	$\frac{4}{65}$	0 0	0 0
	5	$\frac{1}{65}$	0 65	$\frac{3}{65}$	$\frac{4}{65}$	0 0	0 0	0 0
	6	$\frac{2}{65}$	0 65	0 0	0 0	0 0	0 0	$\frac{3}{65}$
	7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0 0	0 0	$\frac{3}{65}$	0 0

- Consider **(source) node 1**: best Δ -ary tree we can build for this source is **Huffman tree** for its destinations [0,1/65,1/13,1/65]
 - resp. **Knuth/Morris-Pratt** property requires
- An optimal **ego-tree** for this source!
- How good can this tree be?



Entropy lower bound known on EPL known for binary trees, e.g. **Mehlhorn** 1975 for BST

Entropy lower bound for Binary Search Trees (BST):

Let $H(p)$ be the entropy of the frequency distribution $p = (p_1, p_2, \dots, p_n)$. Let T be an optimal binary search tree built for the above frequency distribution. Then $EPL(p, T) \geq H(p)/\log(3)$.

Corollary: Can generalize it easily to other degrees Δ .

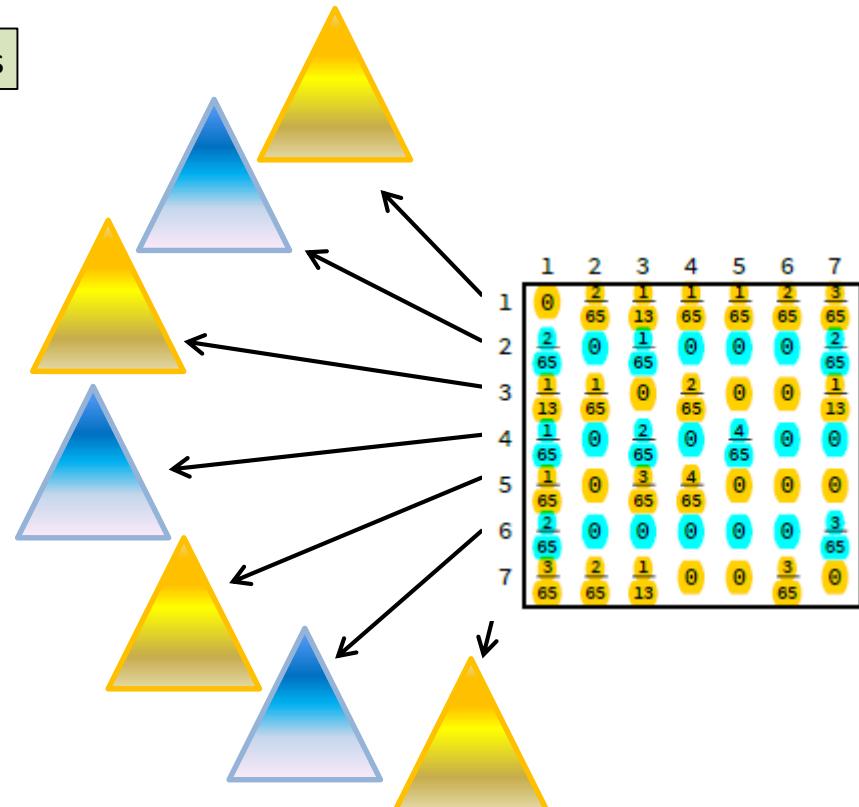
An Entropy Lower Bound (Sources)

- Proof idea ($EPL = \Omega(H_{\Delta}(Y|X))$):

- Consider ***union*** of all **ego-trees**

- Violates ***degree restriction***
but valid lower bound

sources destinations



Lower Bound: Sources + Destinations

Do this in **both dimensions**:

$$\text{EPL} \geq \Omega(\max\{\mathcal{H}_\Delta(Y|X), \mathcal{H}_\Delta(X|Y)\})$$

						$\Omega(\mathcal{H}_\Delta(X Y))$
						$\Omega(\mathcal{H}_\Delta(Y X))$
1	2	3	4	5	6	7
1	0 65	2 65	1 13	1 65	1 65	2 65 65
2	2 65	0 65	1 65	0 0	0 0	2 65
3	1 13	1 65	0 65	2 65	0 0	1 13
4	1 65	0 65	2 65	0 65	4 65	0 0
5	1 65	0 65	3 65	4 65	0 0	0 0
6	2 65	0 65	0 13	0 0	0 0	3 65
7	3 65	2 65	1 13	0 0	3 65	0 0

\mathcal{D}

Lower Bound: Summary

Considering union of n trees,

$$\begin{aligned} \text{EPL}(\mathcal{D}, N_{\Delta}) &\geq \sum_{i=1}^n p(i) \text{EPL}(\mathcal{D}[i], T_{\Delta}^i) \\ &\geq \sum_{i=1}^n p(i) (H_{\Delta}(Y | X=i)) \\ &= \Omega(H_{\Delta}(Y | X)) \end{aligned}$$

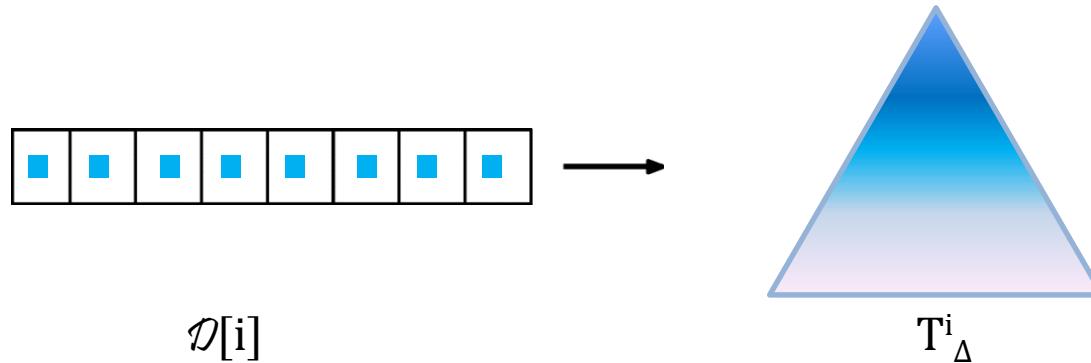
Marginal distribution of source i

Optimal tree

Can DANs Reach The Entropy Speed Limit? Upper Bounds on EPL



Ego-Trees Revisited



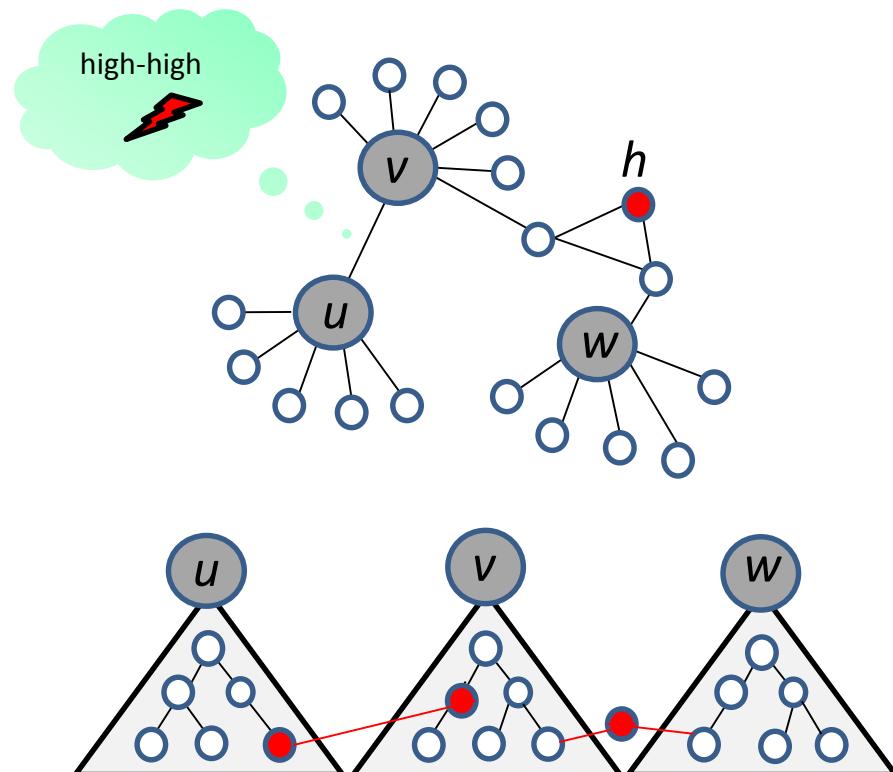
(Tight) Upper Bounds: Algorithm Idea

- Idea: construct **per-node optimal tree**

- BST (e.g., Mehlhorn)
- Huffman tree
- Splay tree (!)

- Take **union** of trees but reduce degree

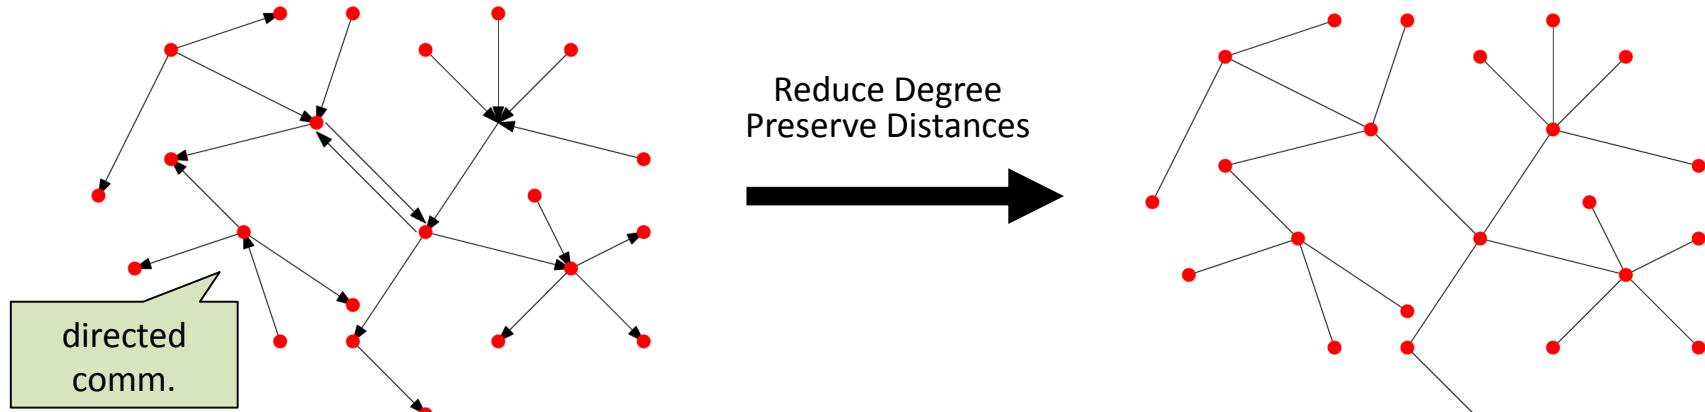
- E.g., in **sparse distribution**: leverage **helper** nodes between two “large” (i.e., high-degree) nodes



How to Reduce Degree?

Example: Tree Distributions

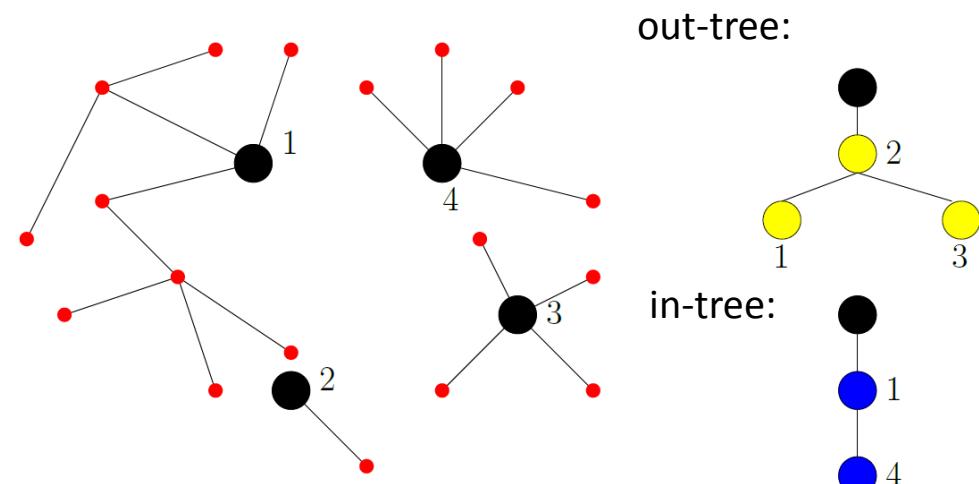
Theorem: Let \mathcal{D} be such that $G_{\mathcal{D}}$ is a tree (ignoring the edge direction). It is possible to generate a DAN with maximum degree 8, such that, $EPL(\mathcal{D}, N) \leq O(H(Y|X) + H(X|Y))$.



Tree Distributions

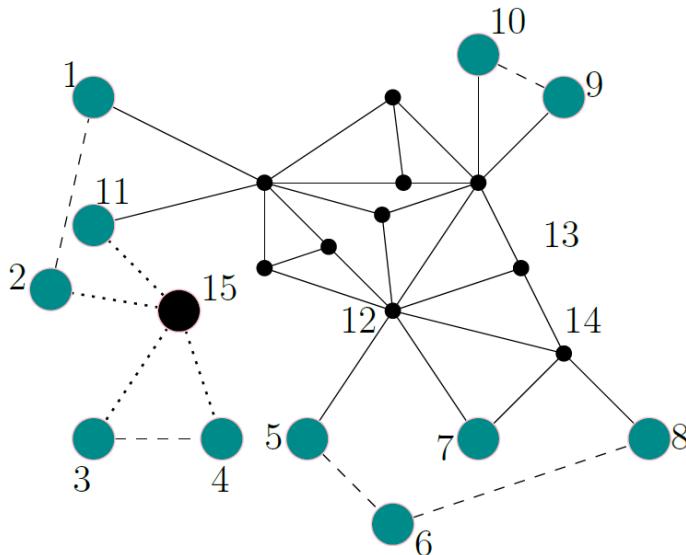
Proof idea:

- Make tree **rooted** and directed: gives **parent-child** relationship
- Arrange the outgoing edges (to children) of each node in a **binary (Huffman) tree**
- Repeat for the incoming edges: make another **another binary (Huffman) tree** with incoming edges from children
- Analysis
 - Can **appear in at most 4 trees**: in&out own tree and in&out tree of parent (parent-child helps to avoid many “children trees”)
 - **Degree** at most 6
 - Huffman trees maintain **distortion**: proportional to **conditional entropy**



Generalize to Arbitrary Sparse Distributions

Demand graph:

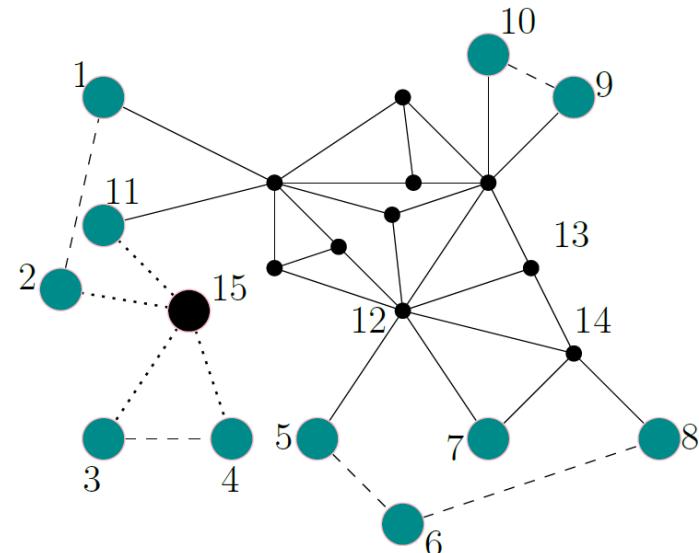
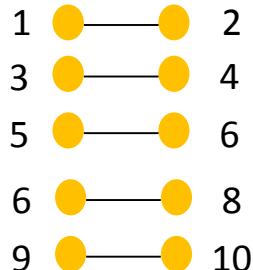


Demand

Sparse Distributions: Construction

Demand graph:

Low-low edges:
remove from
guest, add to DAN

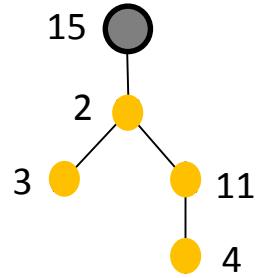


DAN

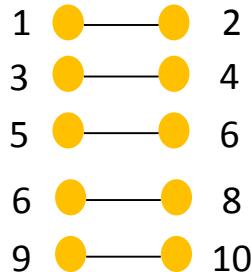
Demand

Sparse Distributions: Construction

High degree nodes with low degree neighbors: need binarization

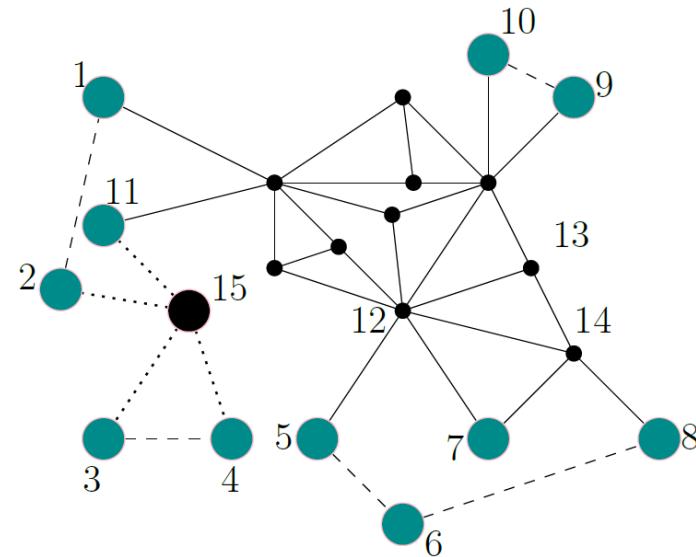


Low-low edges:
remove from
guest, add to DAN



DAN

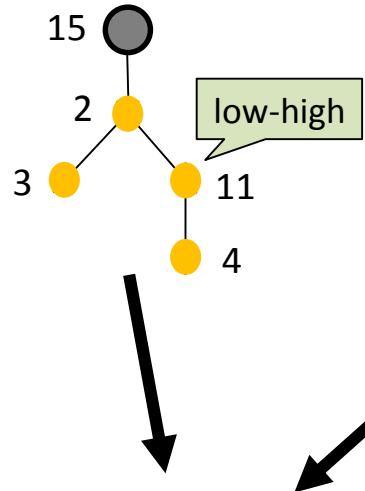
Demand graph:



Demand

Sparse Distributions: Construction

High degree nodes with low degree neighbors: need binarization

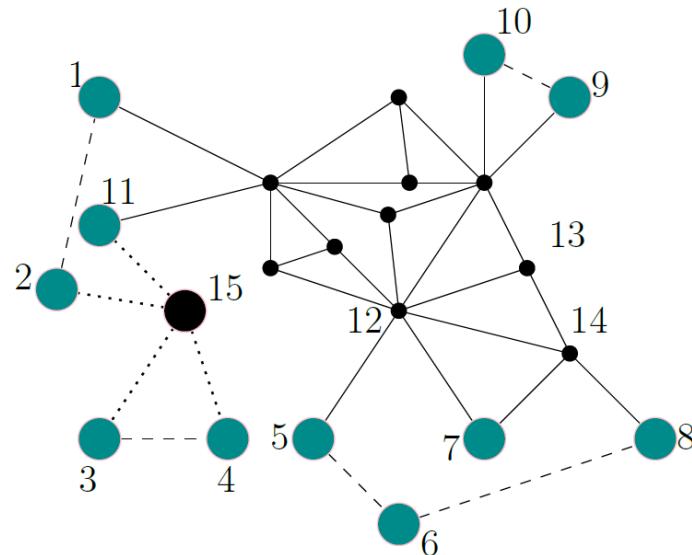


Low-low edges:
remove from
guest, add to DAN

1	2
3	4
5	6
6	8
9	10

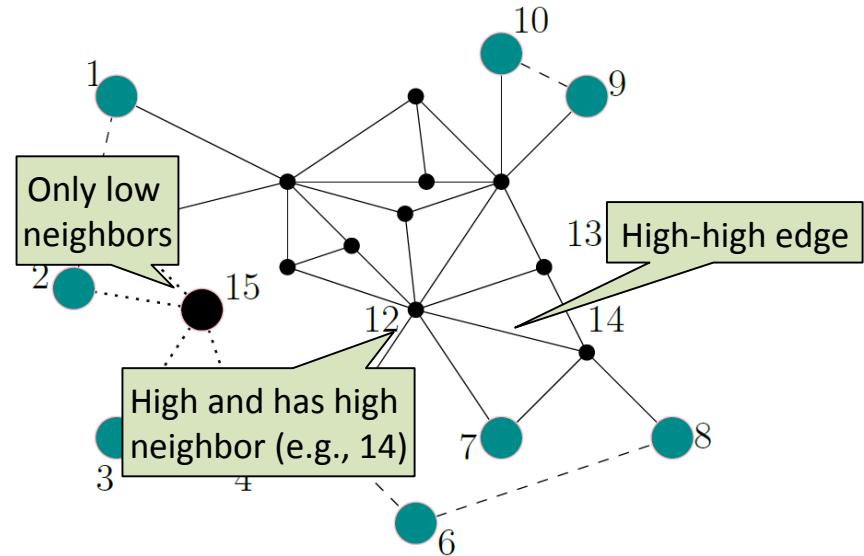
Remove from guest,
add to DAN!

Demand graph:



Sparse Distributions: Construction

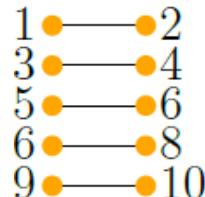
- Find **low** degree nodes
 - Half of the nodes of lowest degree: “below twice average degree”
- Find high degree nodes having only low degree neighbors (e.g., 15 but not 12):
 - Create optimal **binary tree** with low degree neighbors
- Put the **low-low** edges and the **binary tree** into DAN and remove from demand
- Mark **high-high** edges
 - Put (any) **low degree** nodes in between (e.g., 1 or 2): one is enough so distance increased by **+1**
- Now high degree nodes have only low degree neighbors: make **tree** again



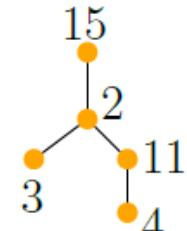
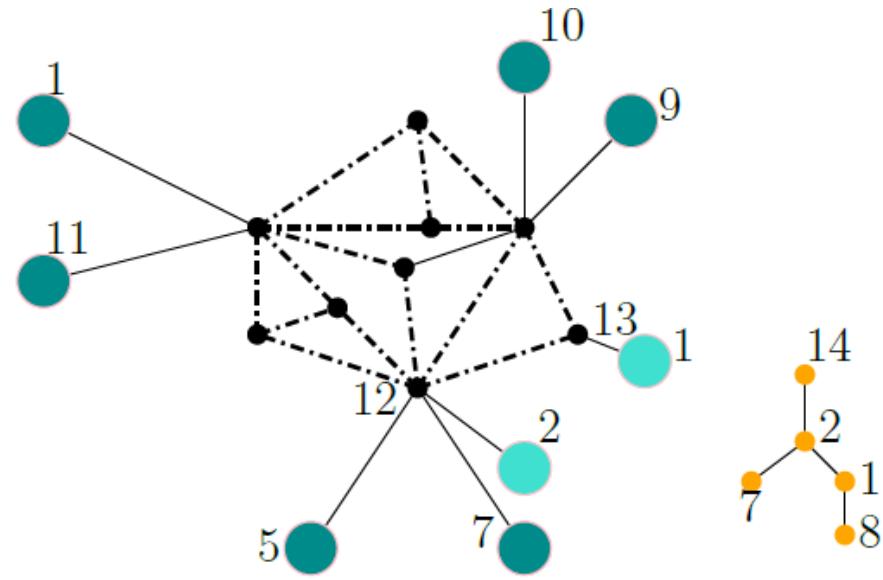
Example Illustrated

- Find **low** degree nodes
- Mark **low-low** edges
- Find **high** degree nodes with only **low** degree neighbors (e.g., 15)
- Make **binary tree** for them
- Add **low** degree node (e.g., 1 and 2) between **high-high** edge (e.g., 12-14, e.g., 14 has two high-degree neighbors 12 and 13)
- Now high nodes have only low neighbors as well, so make **tree** again (at 12-14)

DAN:



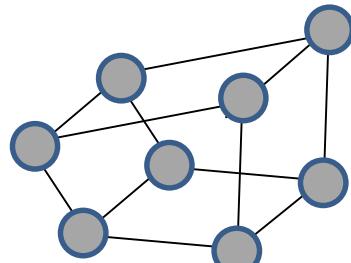
Demand:



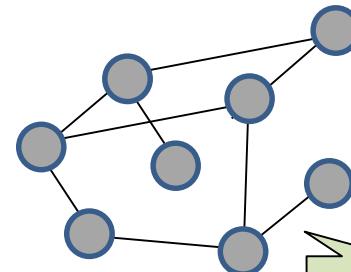
Regular and Uniform Distribution: Leveraging The Connection to Spanners

Theorem: If request distribution \mathcal{D} is **regular and uniform**, and if we can find a constant distortion, linear sized (i.e., **constant, sparse**) spanner for this request graph: can design a constant degree DAN providing an **optimal expected path length** (i.e., $O(H(X|Y)+H(Y|X))$).

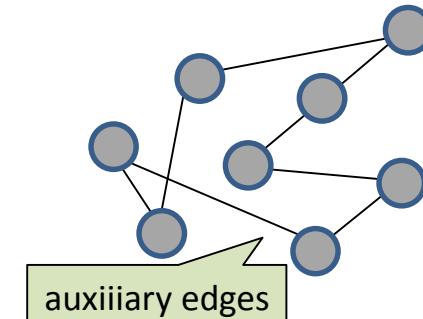
***r-regular and uniform
request:***



***Sparse, irregular
(constant) spanner:***



***Constant degree optimal
DAN (EPL at most **log r**):***

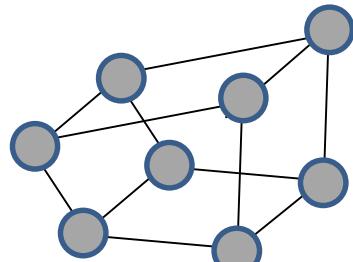


Regular and Uniform Distribution: Leveraging The Connection to Spanners

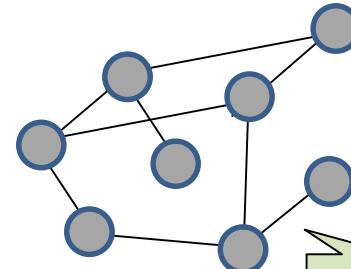
Theorem: If request distribution \mathcal{D} is **regular and uniform**, and if we can find a constant distortion, linear sized (i.e., **constant, sparse**) spanner for this request graph: can design a constant degree DAN providing an **optimal** ~~constant degree DAN~~ ($H(Y|X) \leq H(Y|X)$).

Optimal: in r -regular graphs,
conditional entropy is $\log r$.

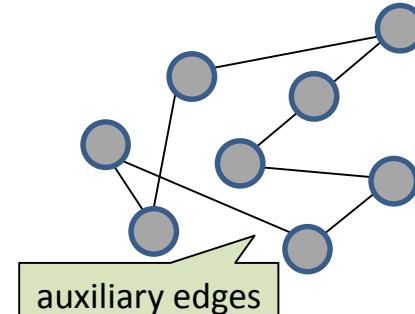
***r*-regular and uniform
request:**



**Sparse, irregular
(constant) spanner:**



**Constant deg.
optimal
DAN (EPL at most $\log r$):**

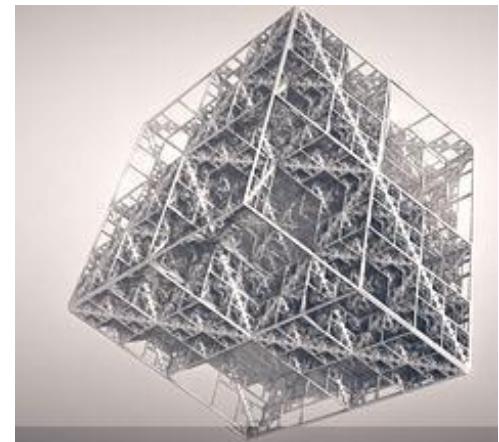


Reduction from Spanner

- Proof technique: degree-reduction again, this time **from sparse spanner** (before: from sparse demand graph)

- Consequences: optimal DAN designs for

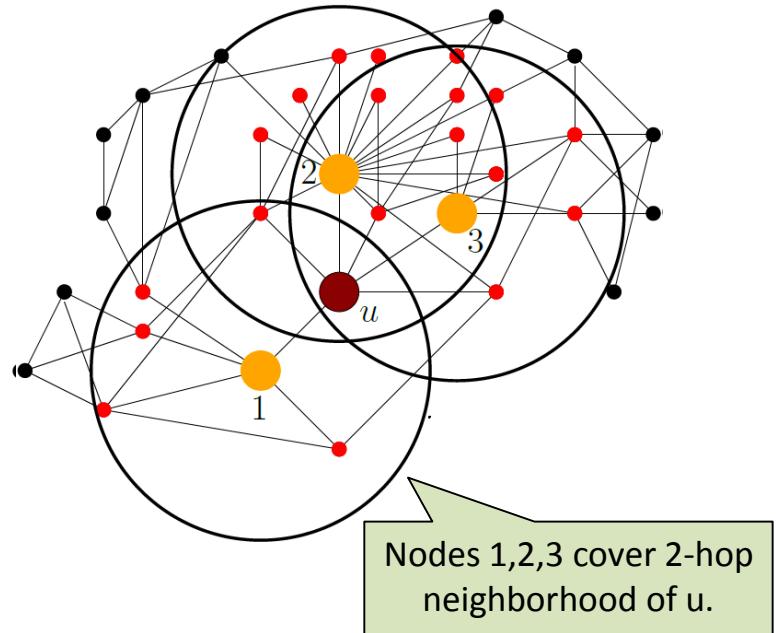
- **Hypercubes** (with $n \log n$ edges) Has sparse 3-spanner.
- **Chordal** graphs Has sparse $O(1)$ -spanner.
- Trivial: graphs with polynomial degree (**dense graphs**)



A dense graph

Another Example: Demands of Locally-Bounded Doubling Dimension

- LDD: G_{\emptyset} has a **Locally-bounded Doubling Dimension** (LDD) iff all 2-hop neighbors are covered by 1-hop neighbors of just λ nodes
 - Note: care only about **2-neighborhood**
- Formally, $B(u, 2) \subseteq \bigcup_{i=1}^{\lambda} B(v_i, 1)$
- Challenge: can be of **high degree**!



DAN for Locally-Bounded Doubling Dimension

Lemma: There exists a sparse 9-(subgraph)spanner for LDD.

This *implies optimal DAN*: still focus on regular and uniform!

Def. (ε -net): A subset V' of V is a **ε -net** for a graph $G = (V, E)$ if

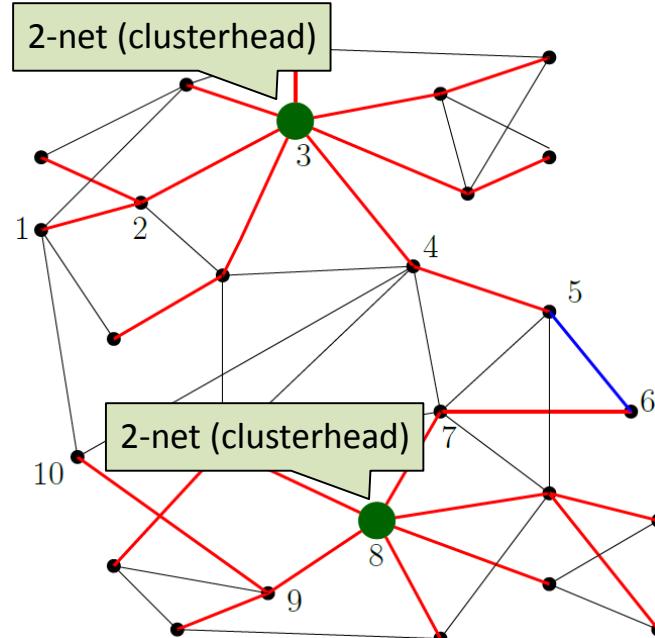
- V' sufficiently “**independent**”: for every $u, v \in V'$, $d_G(u, v) > \varepsilon$
- “**dominating**” V : for each $w \in V$, \exists at least one $u \in V'$ such that, $d_G(u, w) \leq \varepsilon$

9-Spanner for LDD (= optimal DAN)

Simple algorithm:

1. Find a 2-net

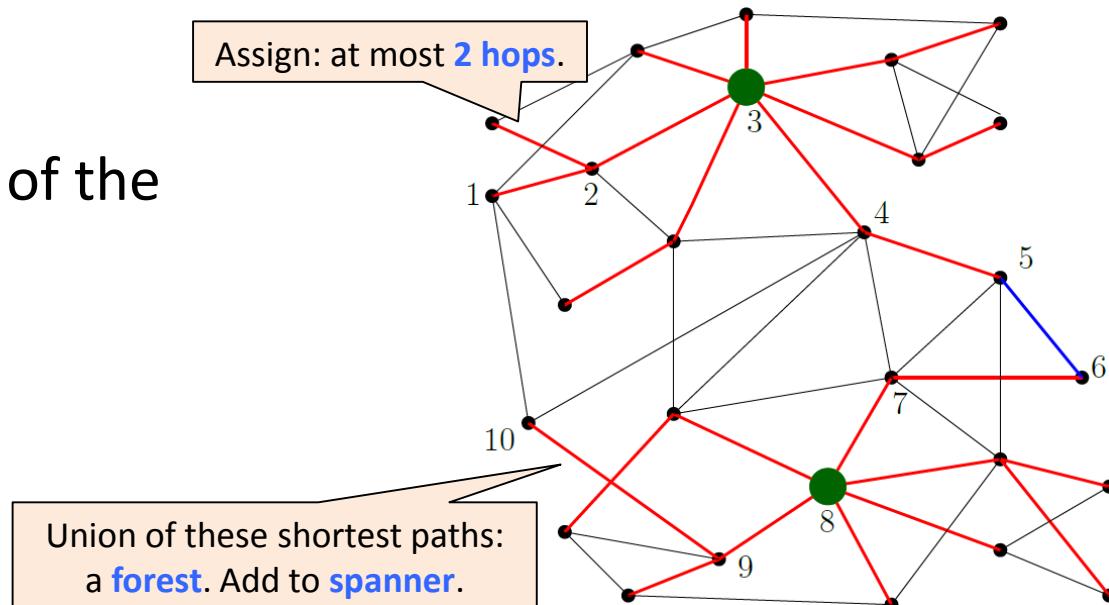
Easy: Select nodes into 2-net **one-by-one** in decreasing (remaining) degrees, **remove 2-neighborhood**. Iterate.



9-Spanner for LDD (= optimal DAN)

Simple algorithm:

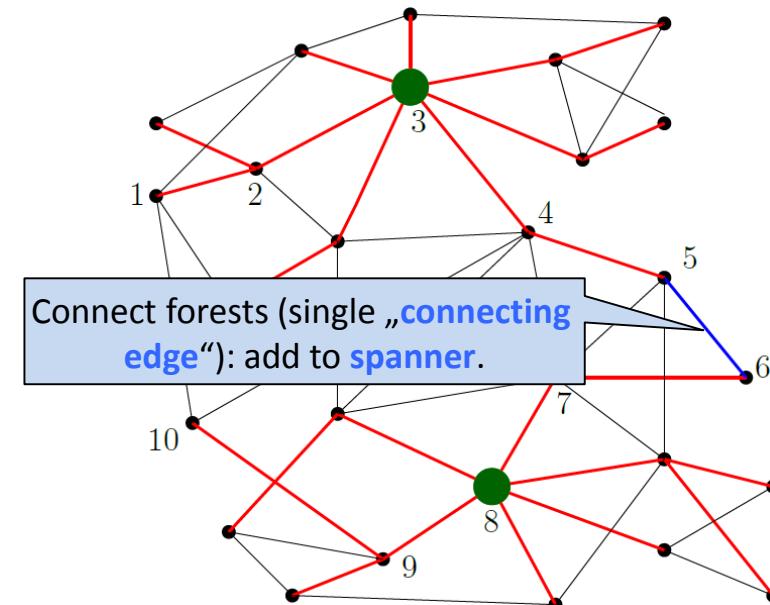
1. Find a 2-net
2. Add nodes to one of the closest 2-net nodes



9-Spanner for LDD (= optimal DAN)

Simple algorithm:

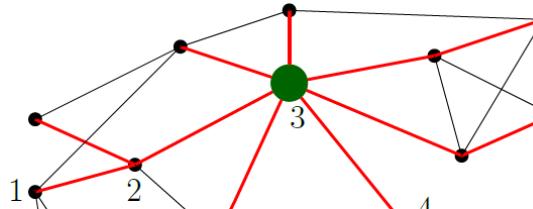
1. Find a 2-net
2. Add nodes to one of the closest 2-net nodes
3. Join two clusters if there are edges in between



9-Spanner for LDD (= optimal DAN)



Distortion 9: Detour via clusterheads:
 $u, ch(u), x, y, ch(v), v$



1. Build a 2-net
2. Add nodes to one of the closest 2-net nodes

3. Join two clusters
edges in between



Sparse: Spanner only includes forest (sparse) plus “connecting edges”: but since in a locally doubling dimension graph the number of cluster heads at distance 5 is bounded, only a small number of neighboring clusters will communicate.



Further Reading

Demand-Aware Network Designs of Bounded Degree

Chen Avin, Kaushik Mondal, and Stefan Schmid.

31st International Symposium on Distributed Computing (**DISC**),

Vienna, Austria, October 2017.

So how useful are entropy-proportional DANs?

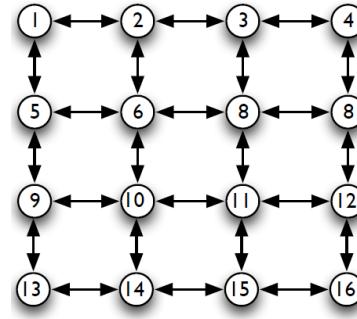
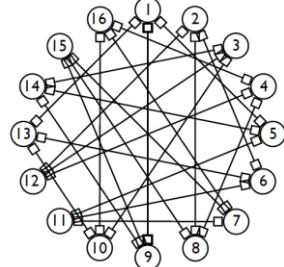
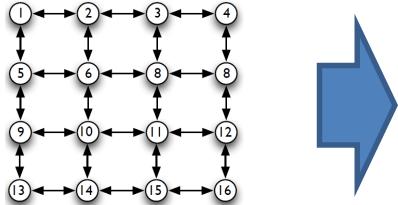
It depends...

- Demand-oblivious network: no information about matrix: **maximum uncertainty (entropy)**, so EPL = entropy = $\Omega(\log n)$...
 - ... even if the demand matrix has a very low **actual entropy**
- **Actual entropy** depends on ***spatial locality*** of communication

	1	2	3	4	5	6	7
1	0	$\frac{2}{65}$	$\frac{1}{13}$	$\frac{1}{65}$	$\frac{1}{65}$	$\frac{2}{65}$	$\frac{3}{65}$
2	$\frac{2}{65}$	0	$\frac{1}{65}$	0	0	0	$\frac{2}{65}$
3	$\frac{1}{13}$	$\frac{1}{65}$	0	$\frac{2}{65}$	0	0	$\frac{1}{13}$
4	$\frac{1}{65}$	0	$\frac{2}{65}$	0	$\frac{4}{65}$	0	0
5	$\frac{1}{65}$	0	$\frac{3}{65}$	$\frac{4}{65}$	0	0	0
6	$\frac{2}{65}$	0	0	0	0	0	$\frac{3}{65}$
7	$\frac{3}{65}$	$\frac{2}{65}$	$\frac{1}{13}$	0	0	$\frac{3}{65}$	0

Example 1: 2-dim Grid

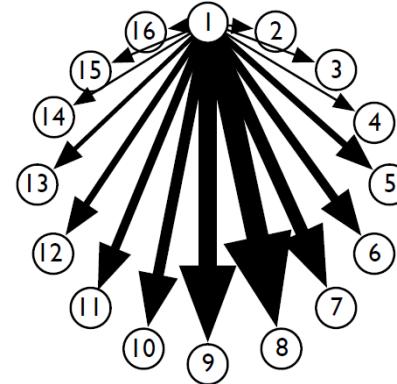
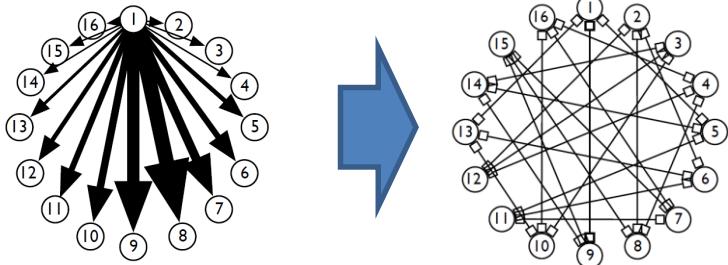
Low **spatial locality**: conditional entropy
less than two (i.e., $O(1)$)



But **embedding** a 2-dim grid demand graph on a const-degree expander would result in
 $\Omega(\log n)$ path lengths

Example 2: Weighted Star

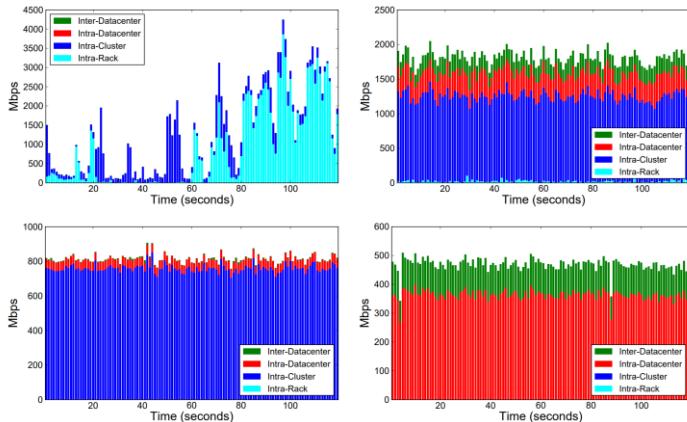
Low **spatial locality**: conditional entropy can be ***very small*** if skew is large



But **embedding** a weighted star demand graph on a const-degree expander would result in **$\Omega(\log n)$** path lengths

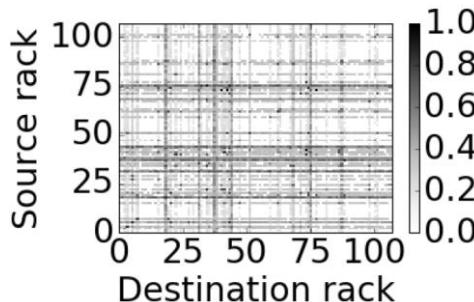
Many Empirical Studies Confirm Spatial Locality

Facebook



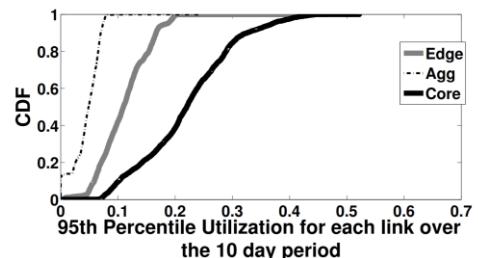
Inside the Social Network's (Datacenter)
Network @ SIGCOMM 2015

Microsoft



ProjectToR @ SIGCOMM 2016

Benson et al.



Understanding Data Center Traffic
Characteristics @ WREN 2009

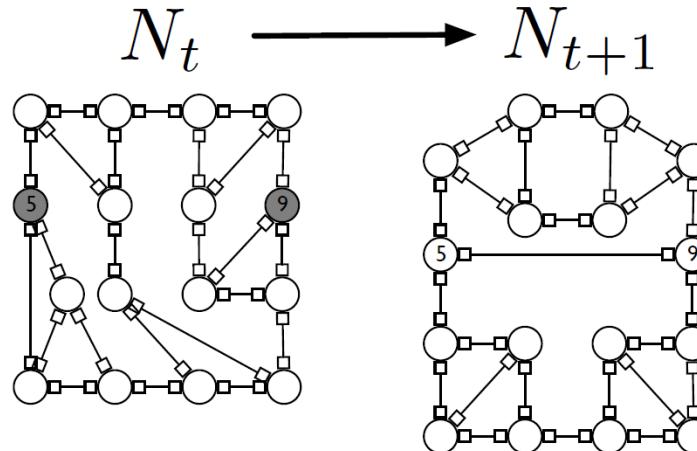
Roadmap

- Vision and Motivation
- An analogy: coding and datastructures
- Principles of Demand-Aware Network (DAN) Designs
- **Principles of Self-Adjusting Network (SAN) Designs**
- Principles of Decentralized Approaches



What Are the Objectives and Metrics for SANs?

Model: A Cost-Benefit Tradeoff



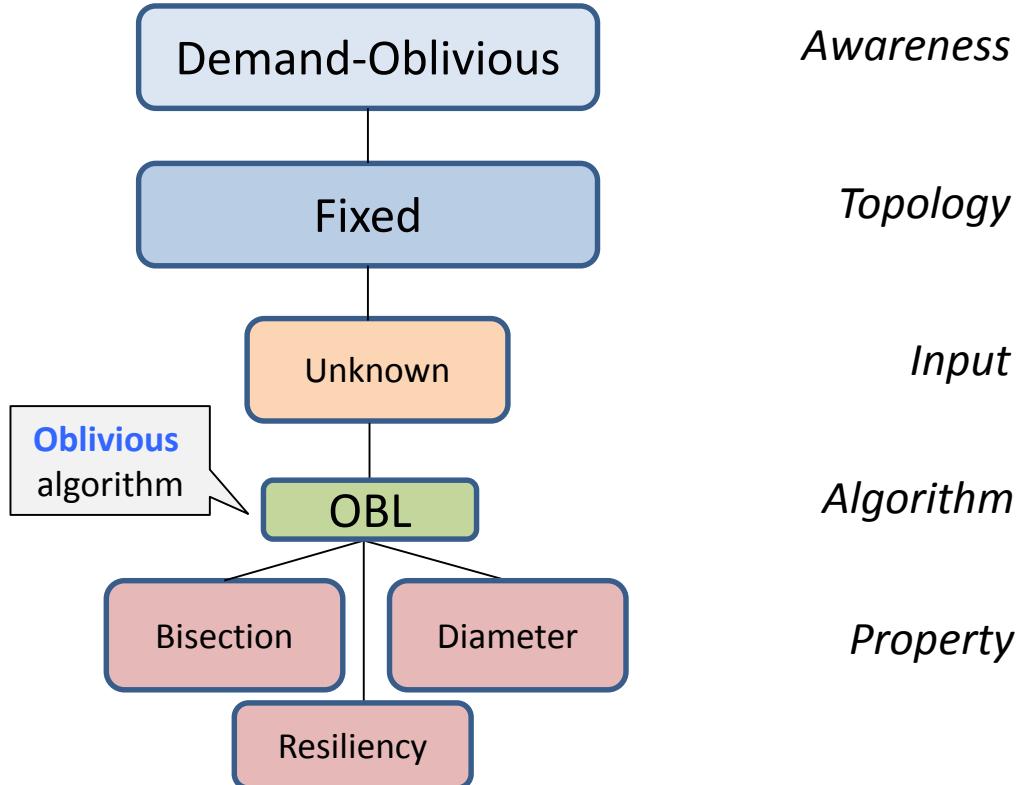
Basic question:
How often to reconfigure?



Short routes
High reconfiguration cost

Low reconfiguration cost
Long routes

A Taxonomy: Traditional Networks



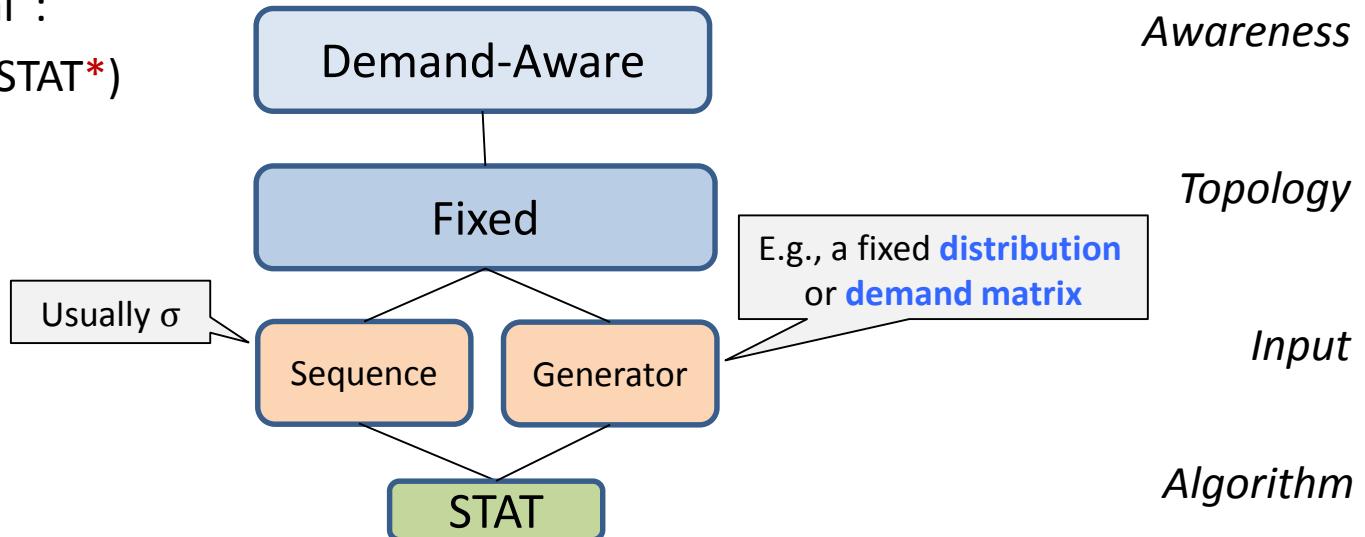
A Taxonomy: Static DANs

Objective:

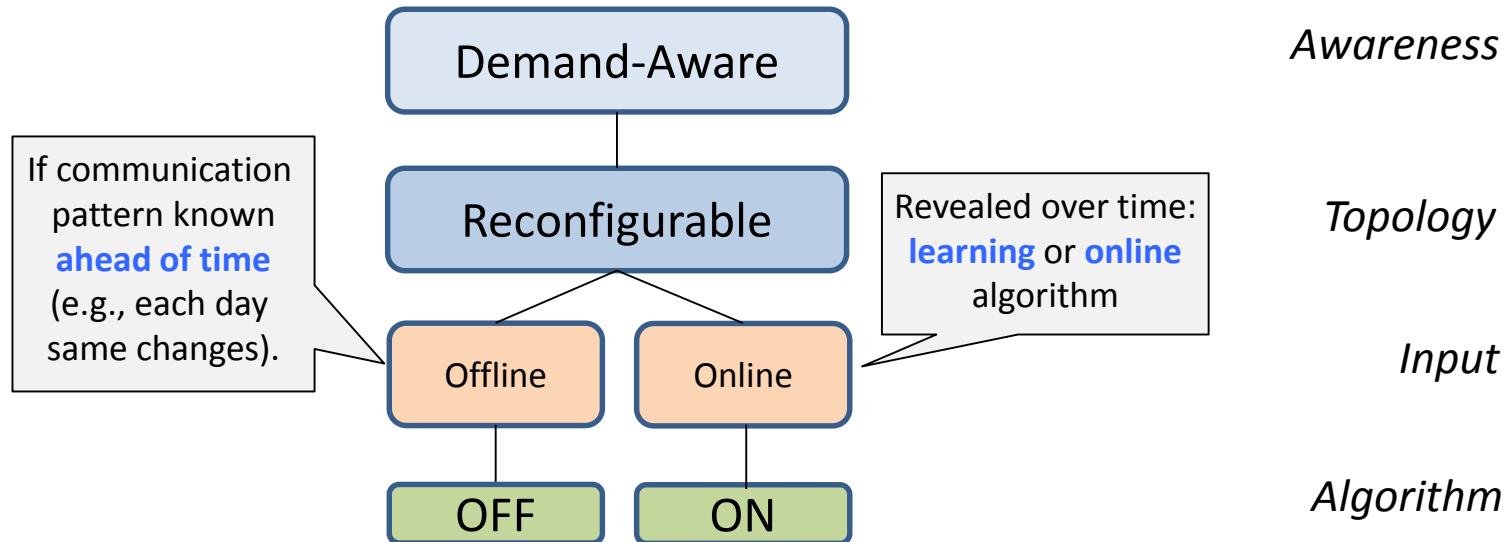
“Close to optimal”:

$$\rho = \text{Cost(STAT)}/\text{Cost(STAT}^*)$$

is small.



A Taxonomy: Reconfigurable Networks

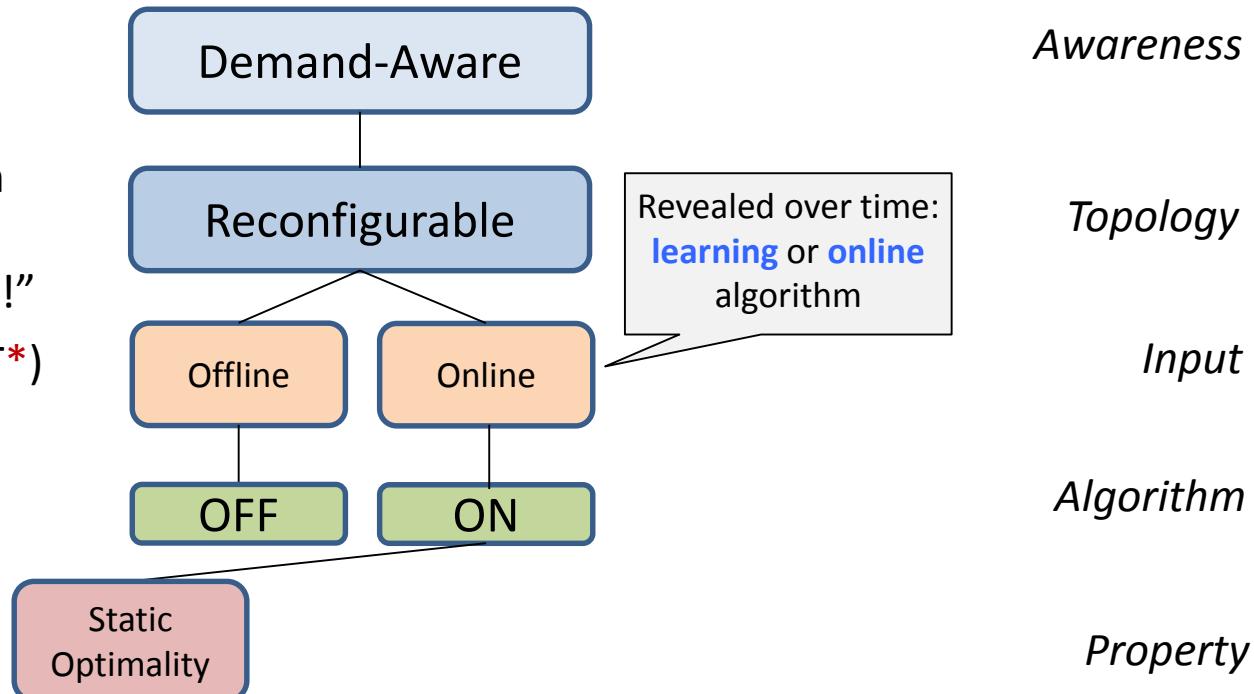


A Taxonomy: Reconfigurable Networks

Static Optimality:

“Don’t be worse than static which knows demand ahead of time!”

$\rho = \text{Cost(ON)}/\text{Cost(STAT*)}$ is constant.

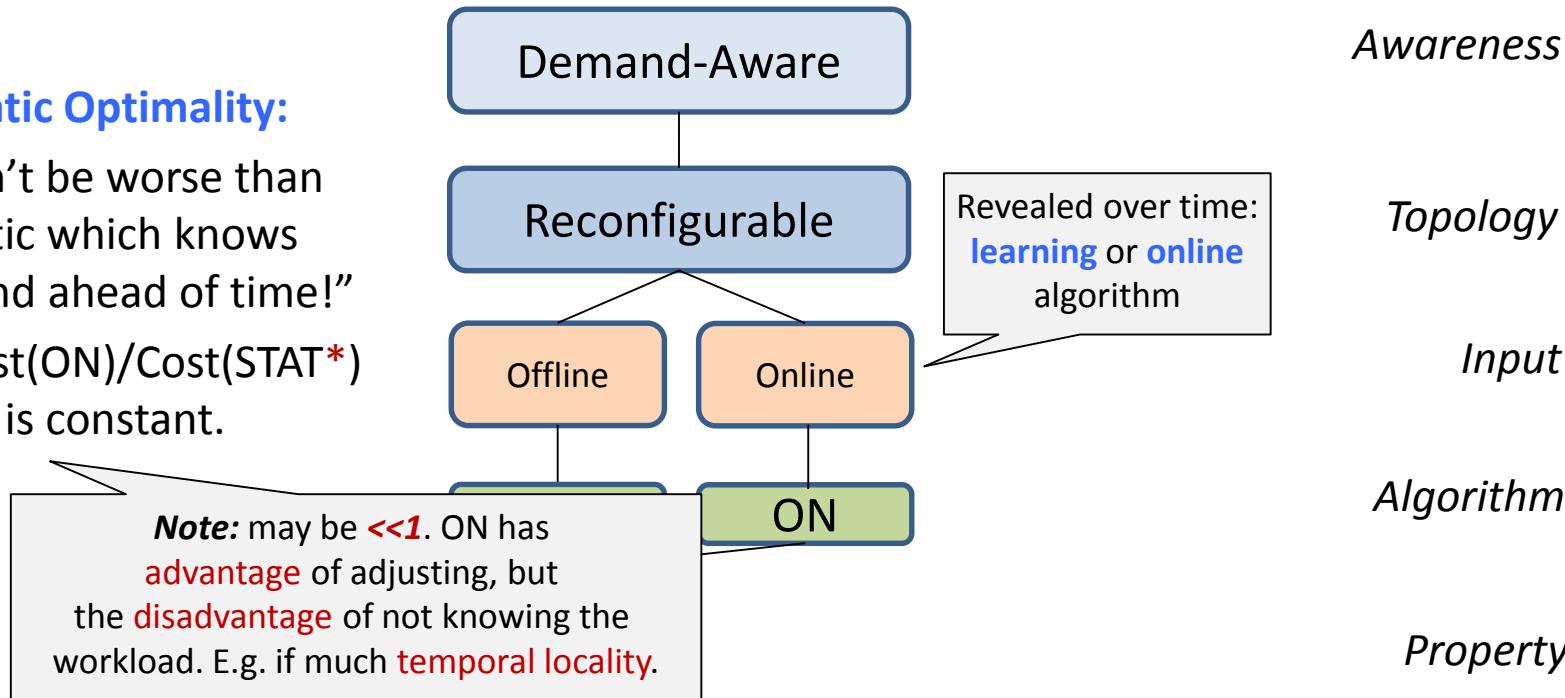


A Taxonomy: Reconfigurable Networks

Static Optimality:

“Don’t be worse than static which knows demand ahead of time!”

$\rho = \text{Cost(ON)}/\text{Cost(STAT*)}$ is constant.



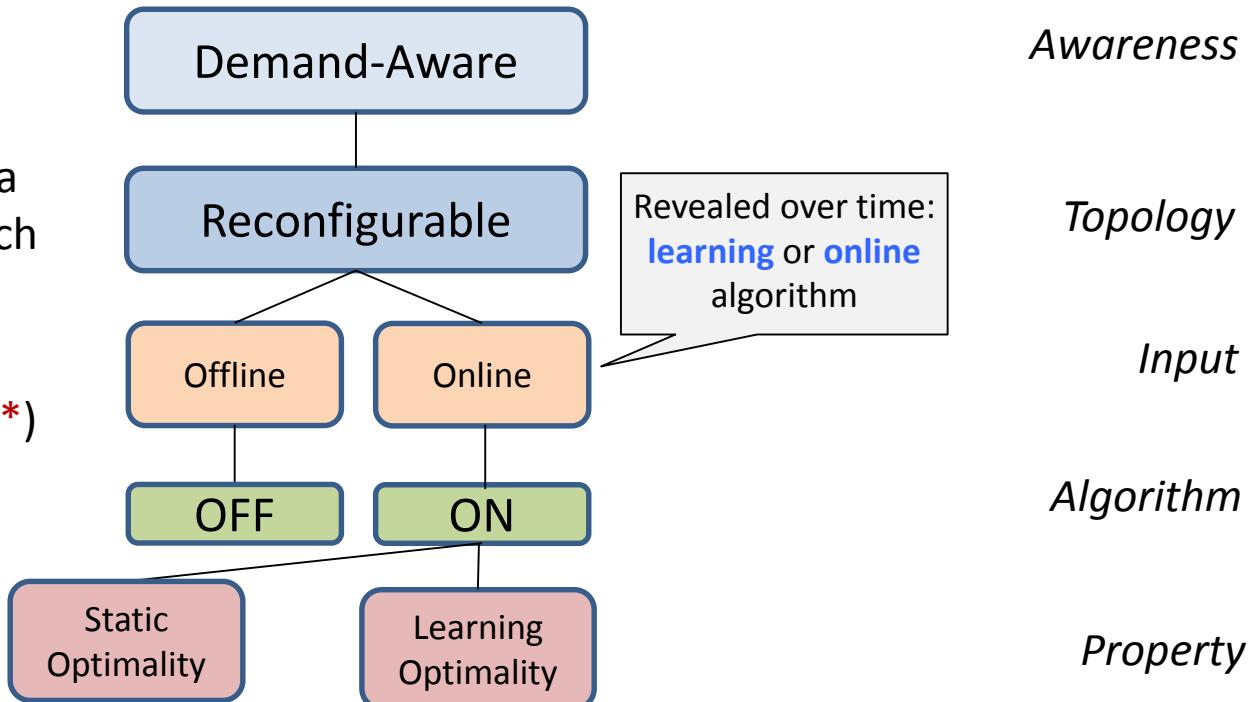
A Taxonomy: Reconfigurable Networks

Learning Optimality:

“Don’t be worse than a dynamic algorithm which knows the (fixed) generator!”

$\rho = \text{Cost(ON)}/\text{Cost(GEN*)}$
is constant.

Always ≥ 1 .



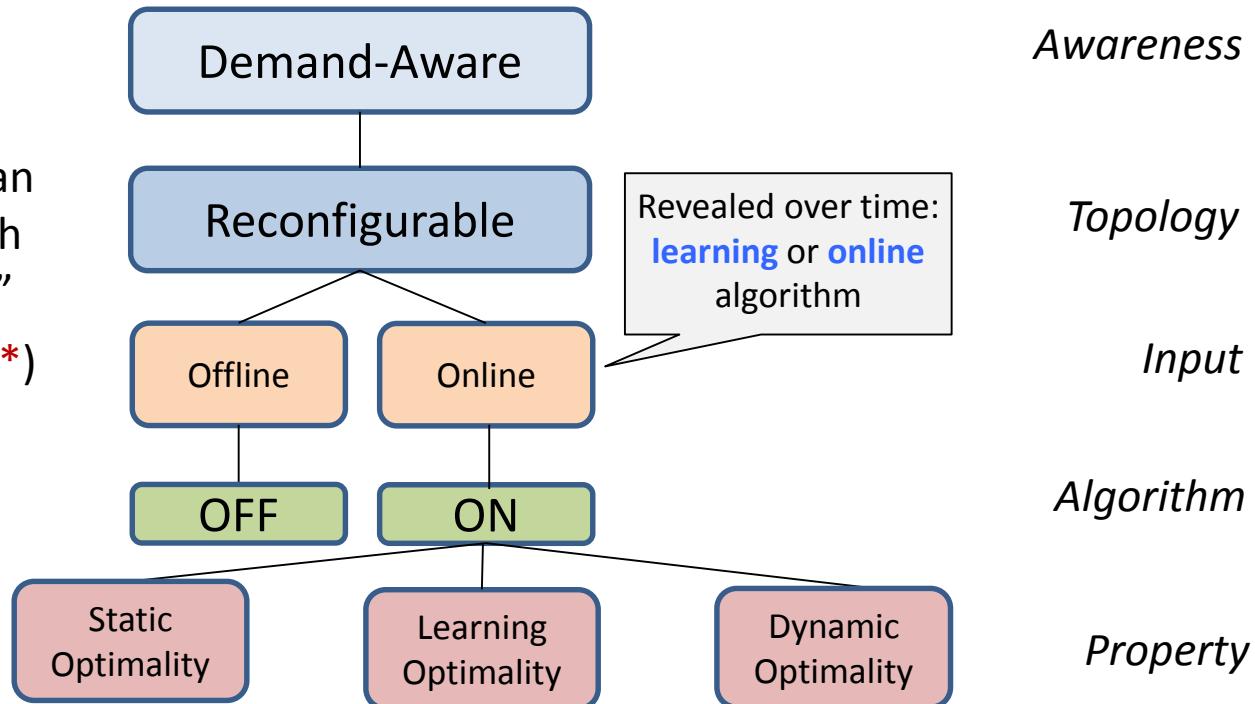
A Taxonomy: Reconfigurable Networks

Dynamic Optimality:

“Don’t be worse than an offline algorithm which knows the sequence!”

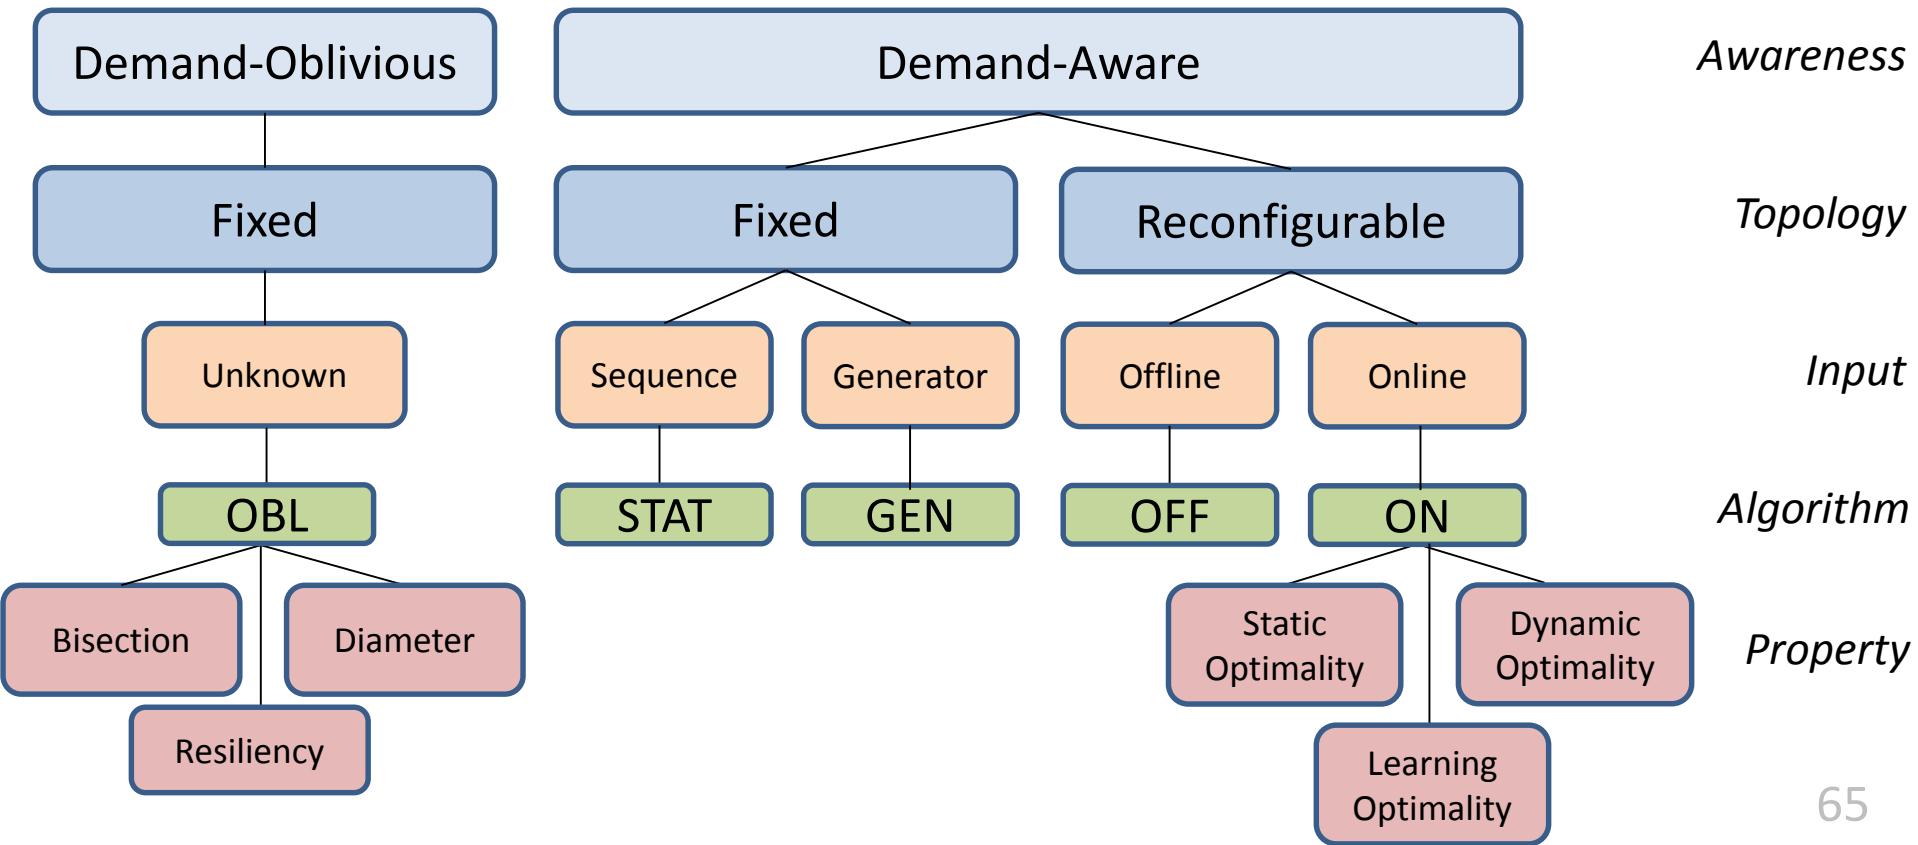
$\rho = \text{Cost(ON)}/\text{Cost(OFF*)}$
is constant.

Always $>= 1$.



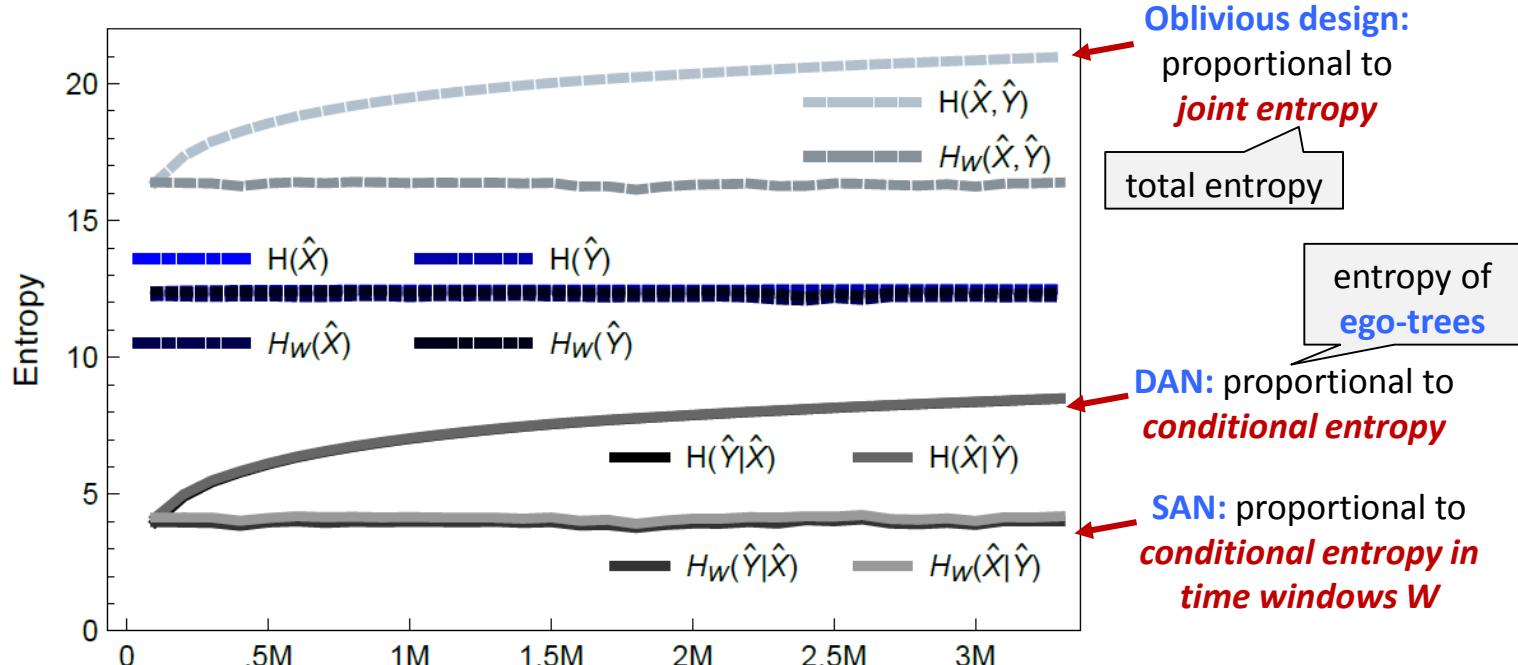
Taxonomy

Toward Demand-Aware Networking: A Theory for
Self-Adjusting Networks. ArXiv 2018.



When are SANs better than DANs?

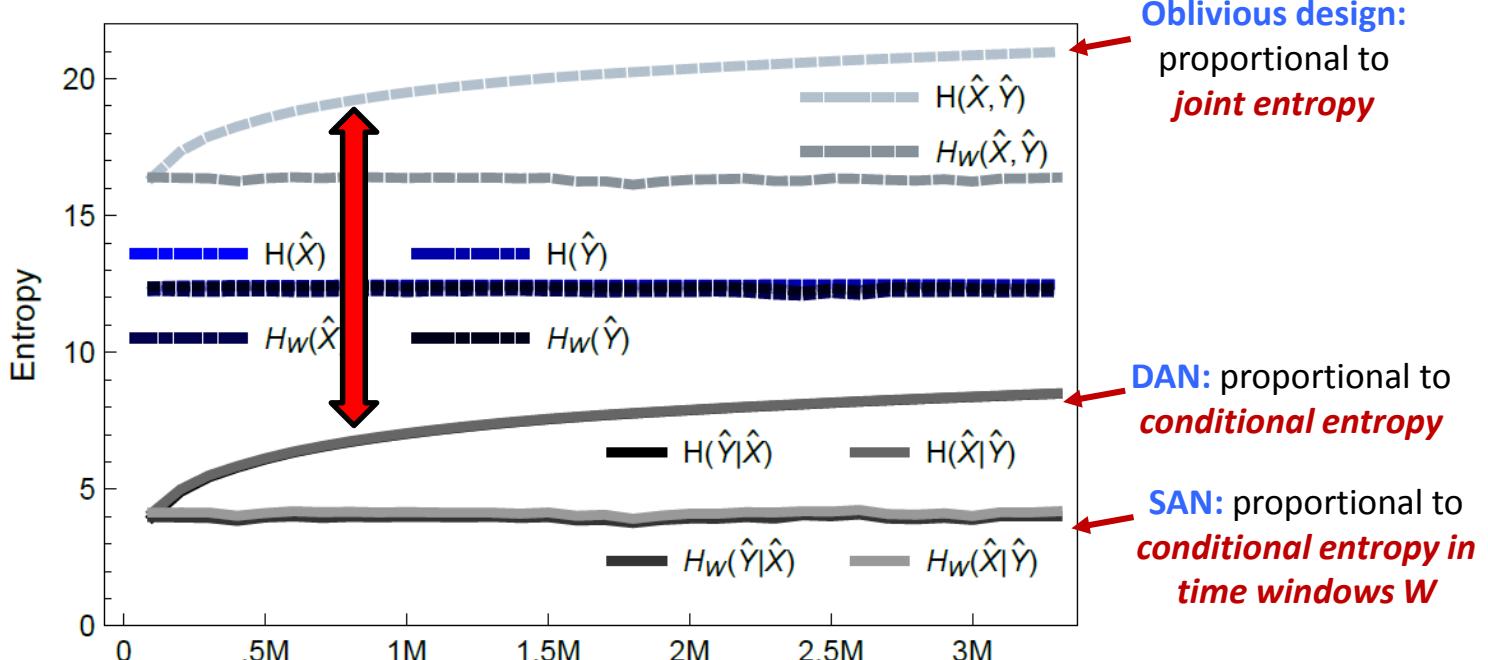
If There is Much Temporal Locality



Entropy measures in Facebook's workload
(3 million requests)

If There is Much Temporal Locality

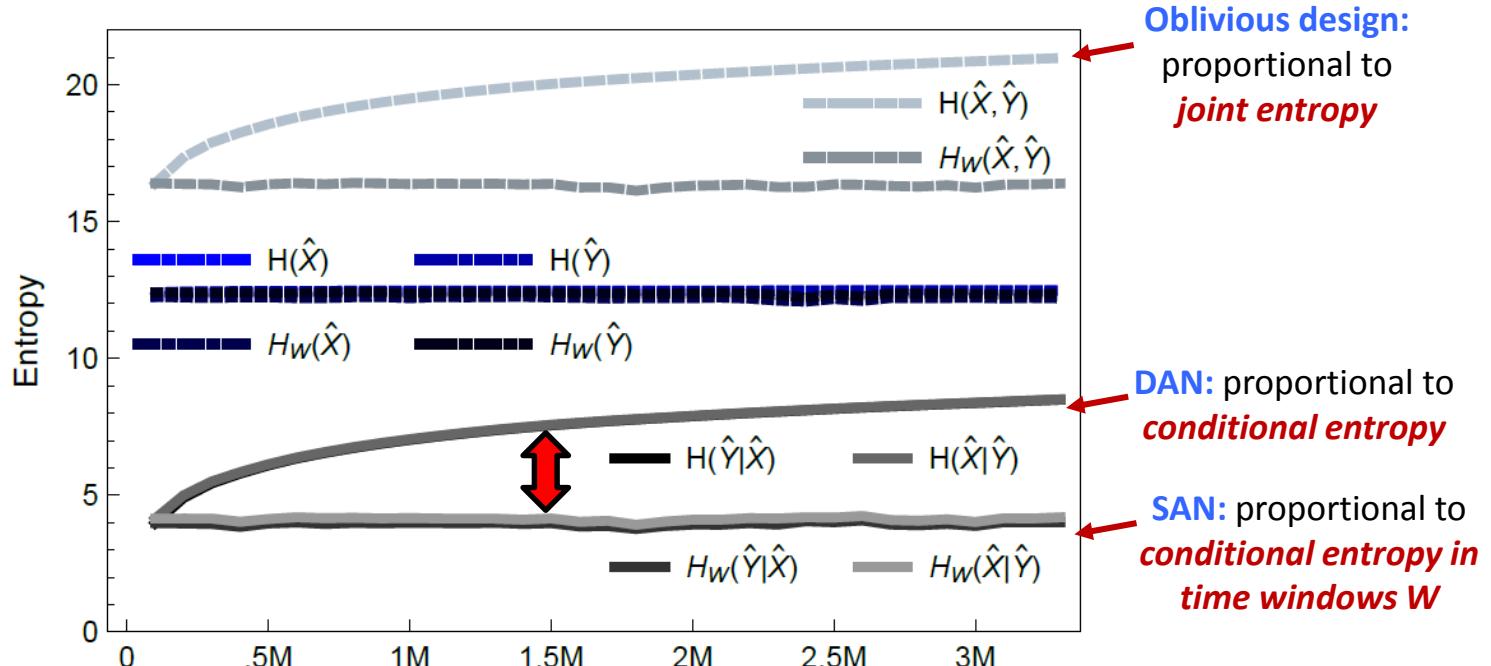
Benefit of
DAN



Entropy measures in Facebook's workload
(3 million requests)

If There is Much Temporal Locality

Benefit of SAN
(if we change every
 $W=100k$ requests)



Entropy measures in Facebook's workload
(3 million requests)

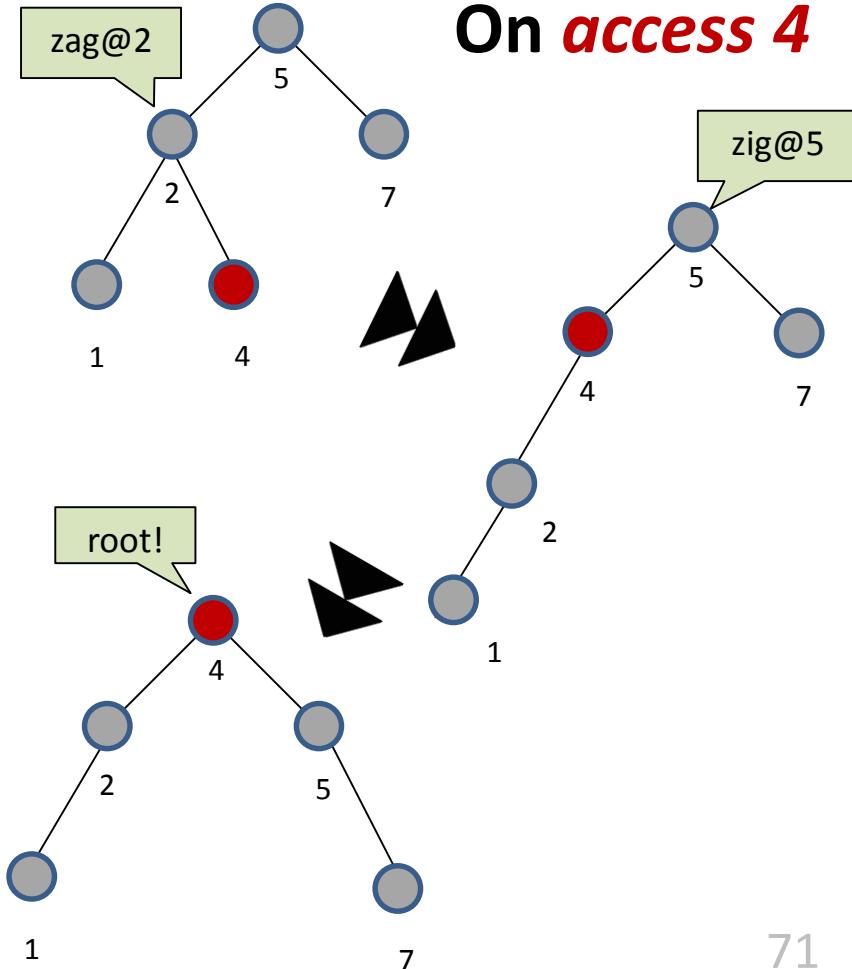
How to Design SANs?



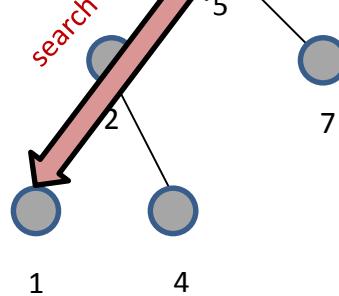
Inspiration from **self-adjusting
datastructures** again!

The Classic Self-Adjusting Datastructure: Splay Tree

- A Binary Search Tree (**BST**)
- Inspired by “**move-to-front**”: move **to root!**
- Self-adjustment: **zig**, **zigzag**, **zigzag**
 - Maintains **search property**
- Many nice properties
 - **Static optimality**, **working set**, (static,dynamic) **fingers**, ...

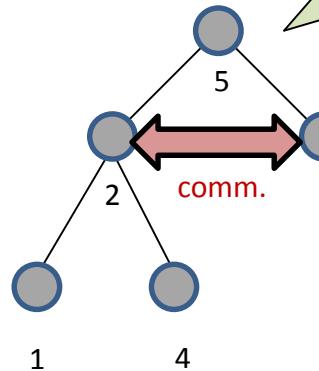


A Simple Idea: Generalize Splay Tree To *SplayNet*



Splay Tree

vs

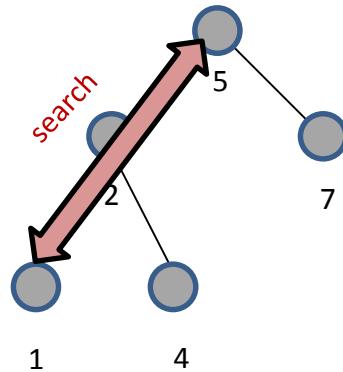


SplayNet

BST is nice for networks:
local (greedy) search!

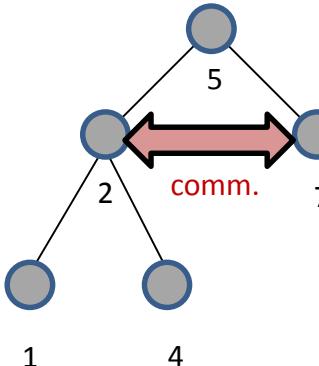
A Simple Idea: Generalize Splay Tree To *SplayNet*

But how?



Splay Tree

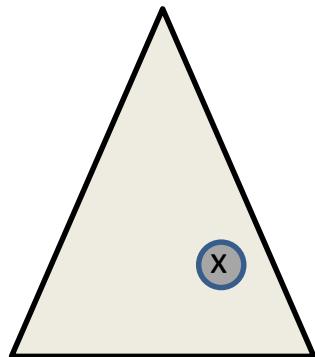
vs



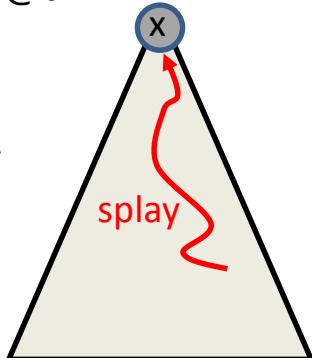
SplayNet

SplayNet: A Simple Idea

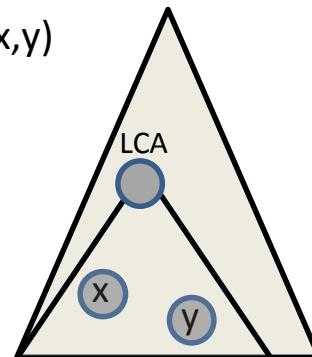
@t: access x



@t+1



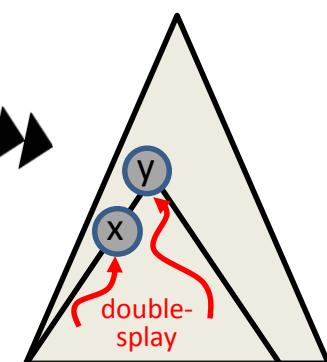
@t: comm
(x,y)



@t+1



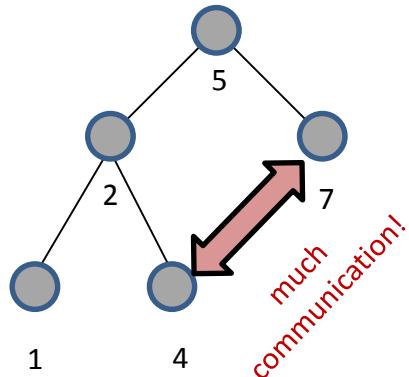
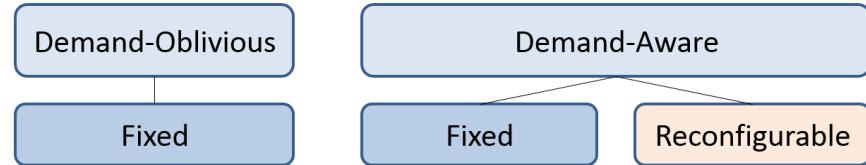
double-splay



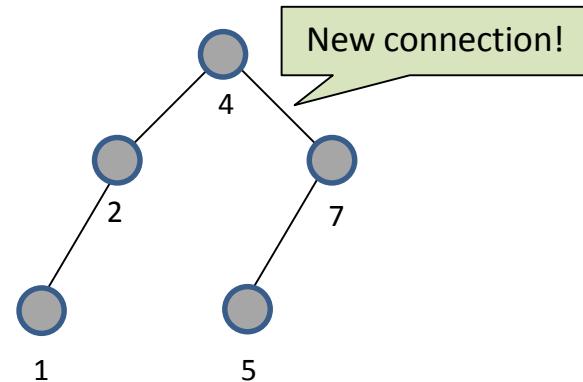
Splay Tree

SplayNet

Example



adjust



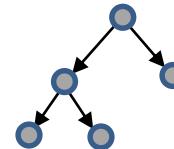
Challenges: How to minimize reconfigurations?

How to keep network locally routable?

Properties of SplayNets

- **Statically optimal** if demand comes from a *product distribution*
 - Product distribution: entropy equals conditional entropy, i.e., $H(X)+H(Y)=H(X|Y)+H(Y|X)$
- Converges to optimal static topology in
 - **Multicast scenario**: requests come from a BST as well
 - **Cluster scenario**: communication only *within* interval
 - **Laminated scenario** : communication is „non-crossing matching“

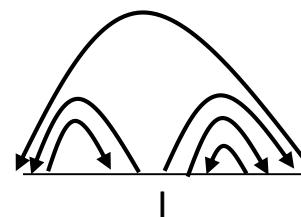
Multicast Scenario



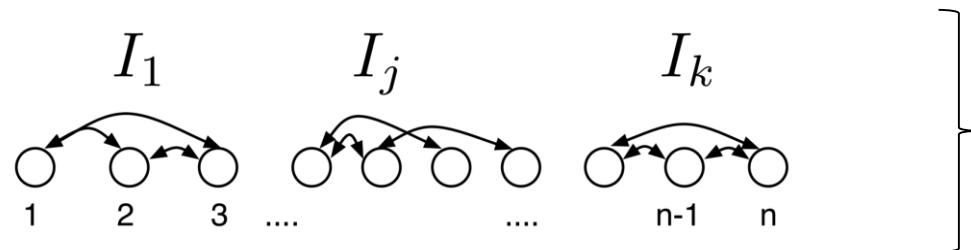
Cluster Scenario



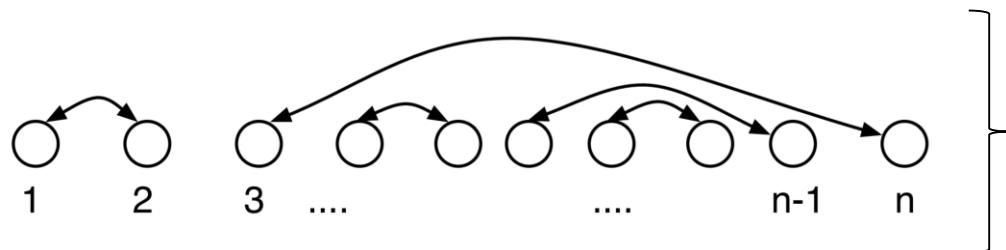
Laminated Scenario



More Specifically



Cluster scenario: SplayNet will converge to state where **path between cluster nodes** only includes cluster nodes



Non-crossing matching scenario: SplayNet will converge to state where all communication pairs are adjacent

Remark: Fast Static SplayNet

Theorem: Optimal static SplayNet can be computed in polynomial-time (**dynamic programming**)

- Unlike unordered tree?

1. Define: flow out of interval I

$$W_I(v) = \sum_{u \in I} w(u, v) + w(v, u)$$

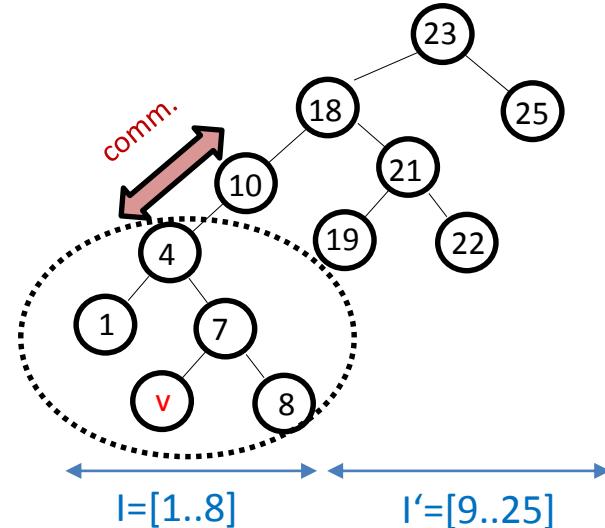
2. Cost of a given tree T_I on I :

$$\text{Cost}(T_I, W_I) = [\sum_{u,v \in I} (d(u, v) + 1)w(u, v)] + D_I * W_I$$

(D_I distances of nodes in I from root of T_I)

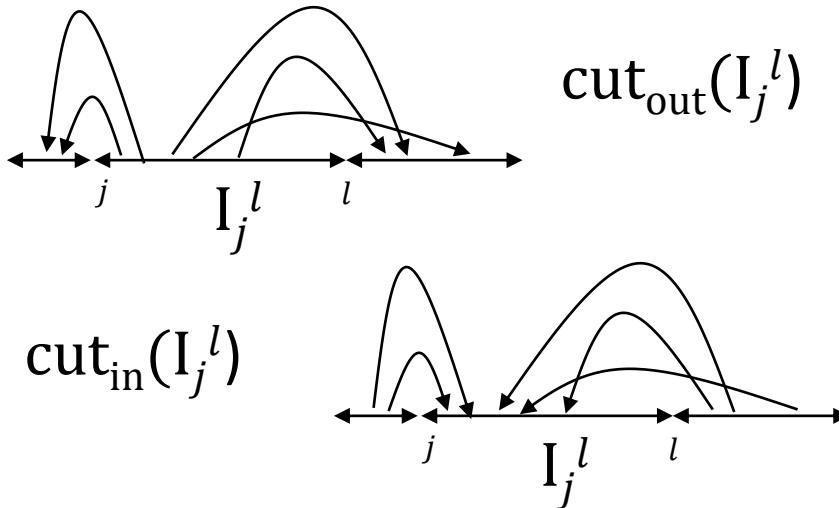
3. Dynamic program over intervals

Choose optimal root and add dist to root



Remark: Improved Lower Bounds

Interval Cuts Bound



$$\text{Cost} = \Omega(\max_i \min_{j,l} H(\text{cut}_{\text{in}}(I_j^l)))$$

$$\text{Cost} = \Omega(\max_i \min_{j,l} H(\text{cut}_{\text{out}}(I_j^l)))$$

If much goes over same cut:
leads to long paths!

Edge Expansion Bound

- Let cut $W(S)$ be weight of edges in cut (S, S') for a given S
- Define a distribution $w_S(u)$ according to the weights to all possible nodes v :

$$w_S(u) = \sum_{\substack{(u,v) \in E(S, \bar{S}) \\ u \in S}} w(u, v) / W(S)$$

- Define entropy of cut and $\text{src}(S), \text{dst}(S)$ distributions accordingly:
- Conductance entropy is lower bound:

$$\varphi_H(S) = W(S) (H(\text{src}(S)) + H(\text{dst}(S)))$$

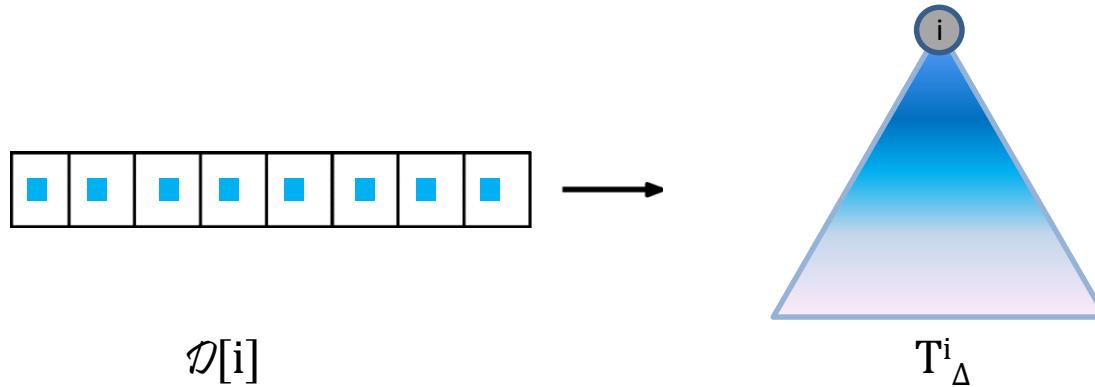
$$\Omega(\phi_H(\mathcal{R}(\sigma)))$$

For SplayNets: stronger lower bounds than conditional entropy.

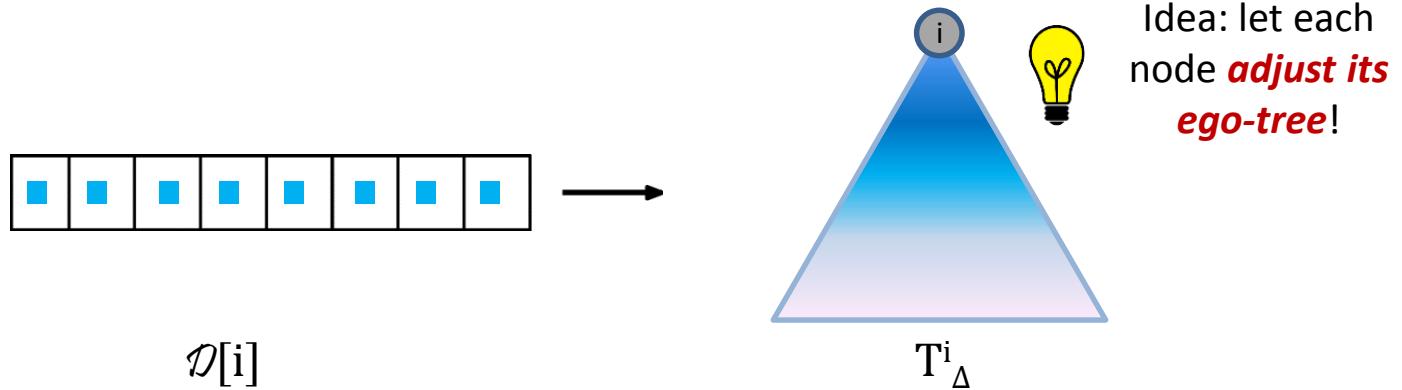
Further Reading

SplayNet: Towards Locally Self-Adjusting Networks
Stefan Schmid, Chen Avin, Christian Scheideler, Michael Borokhovich, Bernhard Haeupler, and Zvi Lotker.
IEEE/ACM Transactions on Networking (**TON**), Volume 24, Issue 3, 2016.

Better Idea: Back to Ego-Trees!

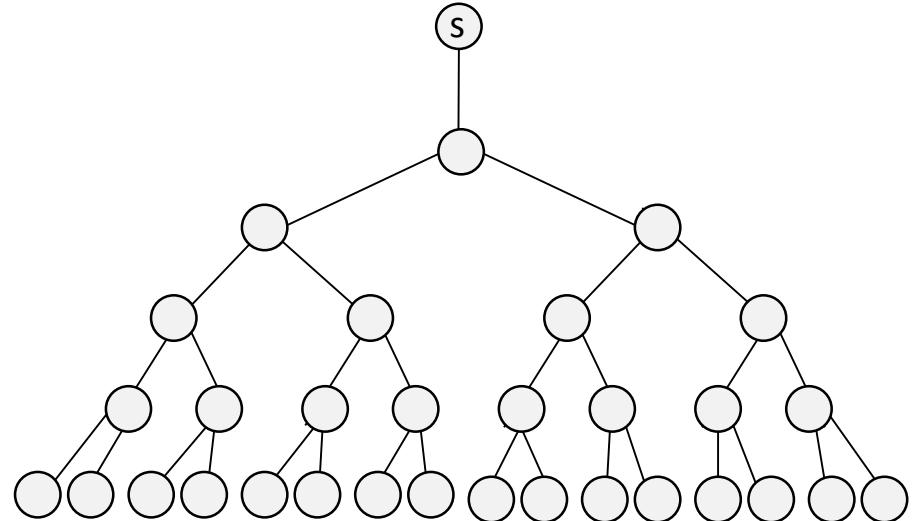


Better Idea: Back to Ego-Trees!



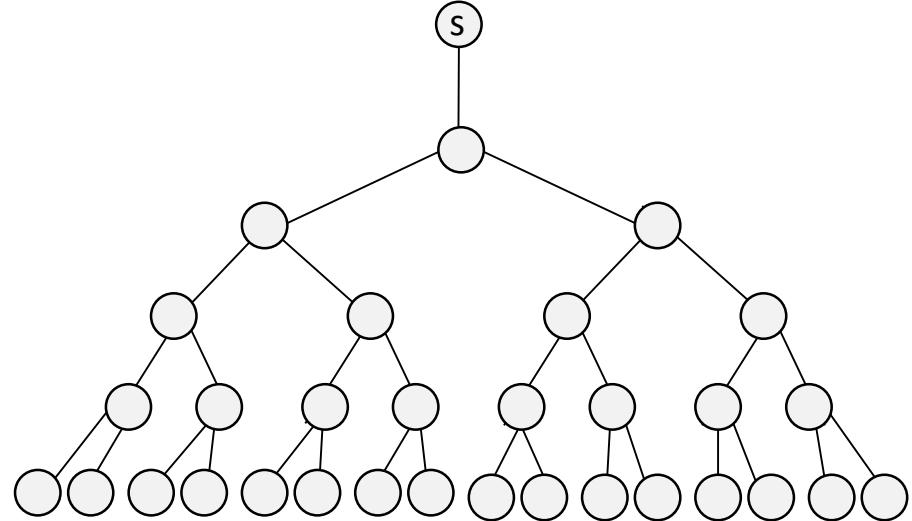
A Balanced Self-Adjusting Tree: Push-Down Tree

- **Push-down tree:** a self-adjusting complete tree
- ***Dynamically optimal***
- Not ordered: requires **a map**



A Balanced Self-Adjusting Tree: Push-Down Tree

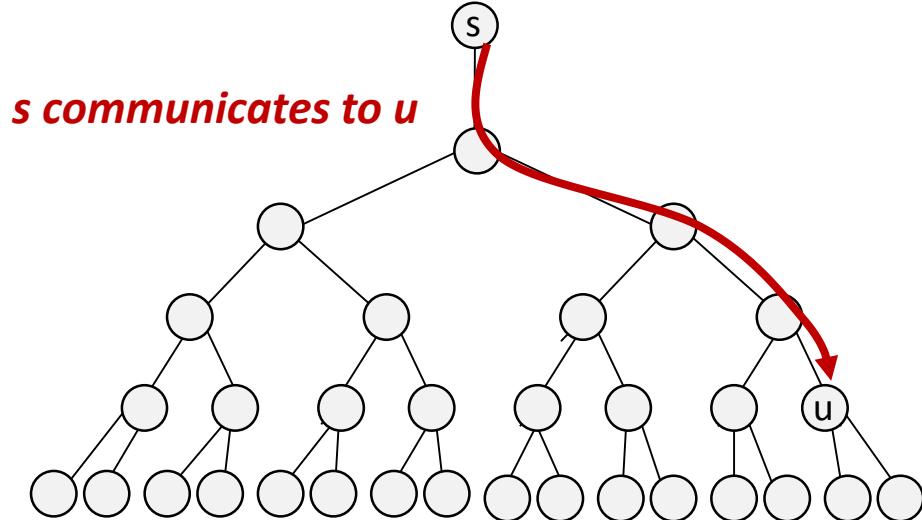
- **Push-down tree:** a self-adjusting complete tree
- ***Dynamically optimal***
- Not ordered: requires **a map**



A useful dynamic property: **Most-Recently Used (MRU)**!
Similar to **Working Set Property**: more recent communication Partners closer to source.

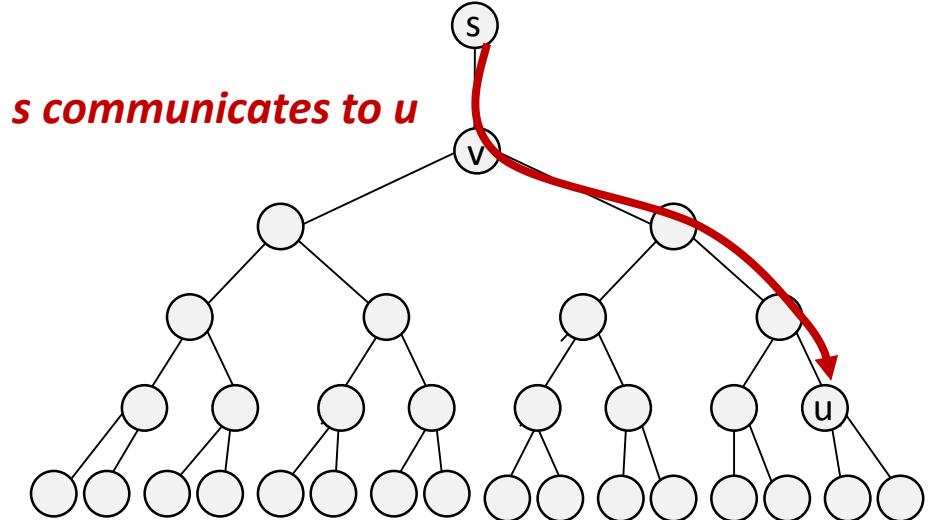
A Balanced Self-Adjusting Tree: Push-Down Tree

- **Push-down tree:** a self-adjusting complete tree
- *Dynamically optimal*
- Not ordered: requires **a map**



A Balanced Self-Adjusting Tree: Push-Down Tree

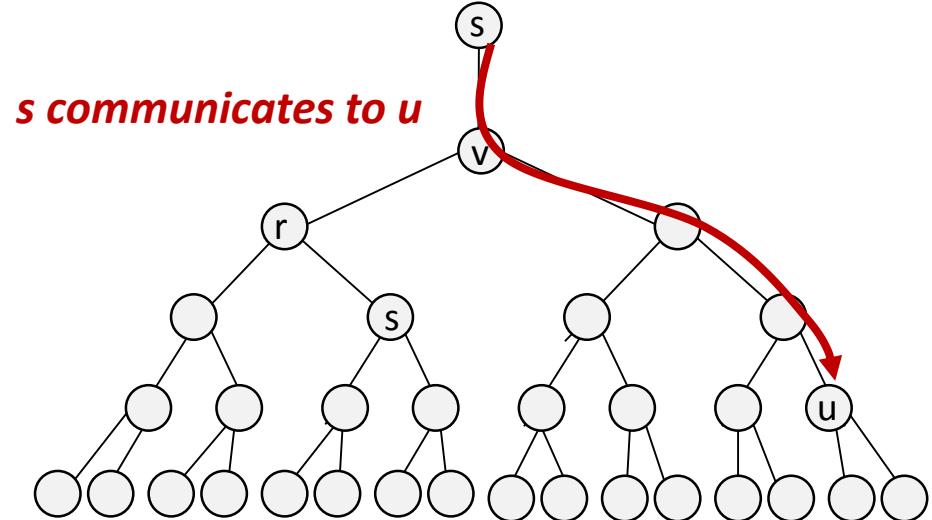
- **Push-down tree:** a self-adjusting complete tree
- **Dynamically optimal**
- Not ordered: requires **a map**



Strict MRU requires: move u to root! But how? Cannot swap with v: v no longer MRU!

A Balanced Self-Adjusting Tree: Push-Down Tree

- **Push-down tree:** a self-adjusting complete tree
- **Dynamically optimal**
- Not ordered: requires **a map**



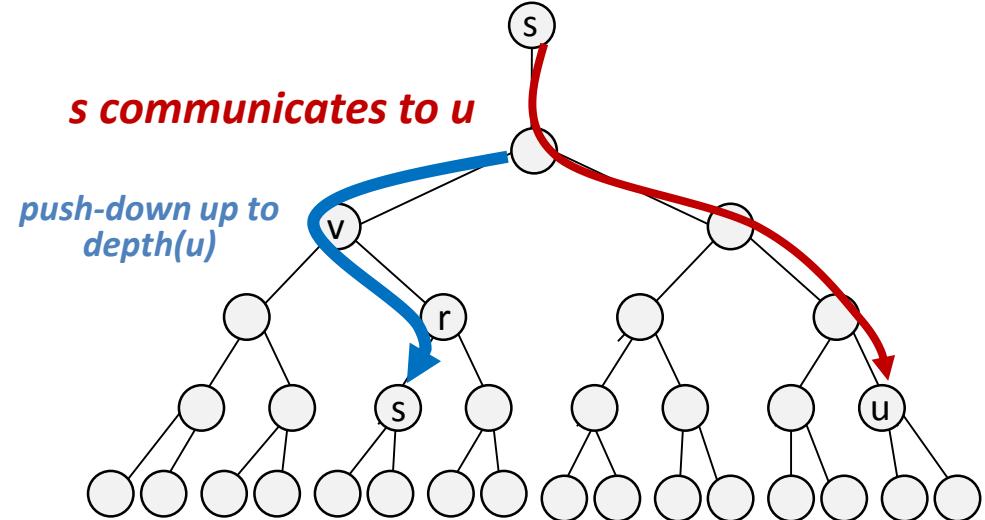
Strict MRU requires: move u to root! But how? Cannot swap with v: v no longer MRU!



Idea: Push v down, in a balanced manner, up to depth(u): left-right-left-right („rotate-push“)

A Balanced Self-Adjusting Tree: Push-Down Tree

- **Push-down tree:** a self-adjusting complete tree
- **Dynamically optimal**
- Not ordered: requires **a map**



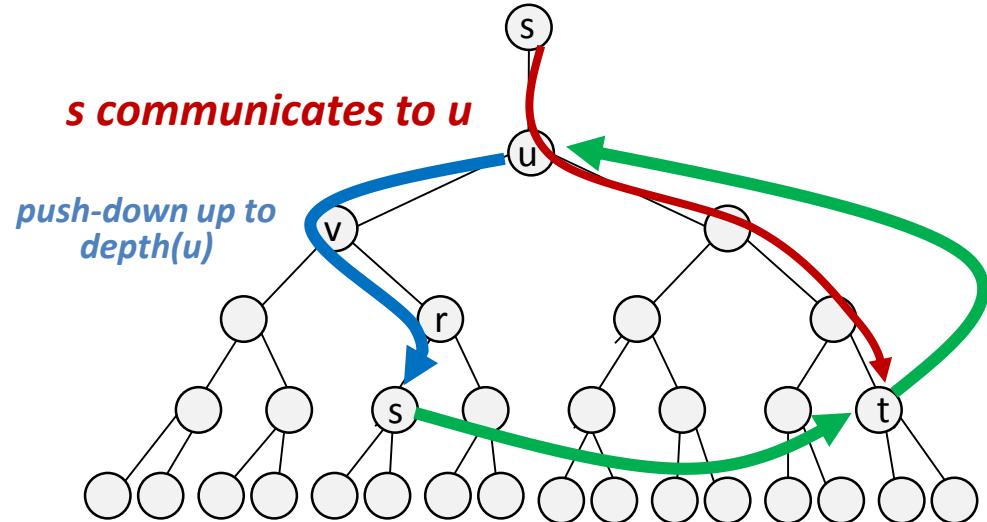
Strict MRU requires: move u to root! But how? Cannot swap with v: v no longer MRU!



Idea: Push v down, in a balanced manner, up to depth(u): left-right-left-right („rotate-push“)

A Balanced Self-Adjusting Tree: Push-Down Tree

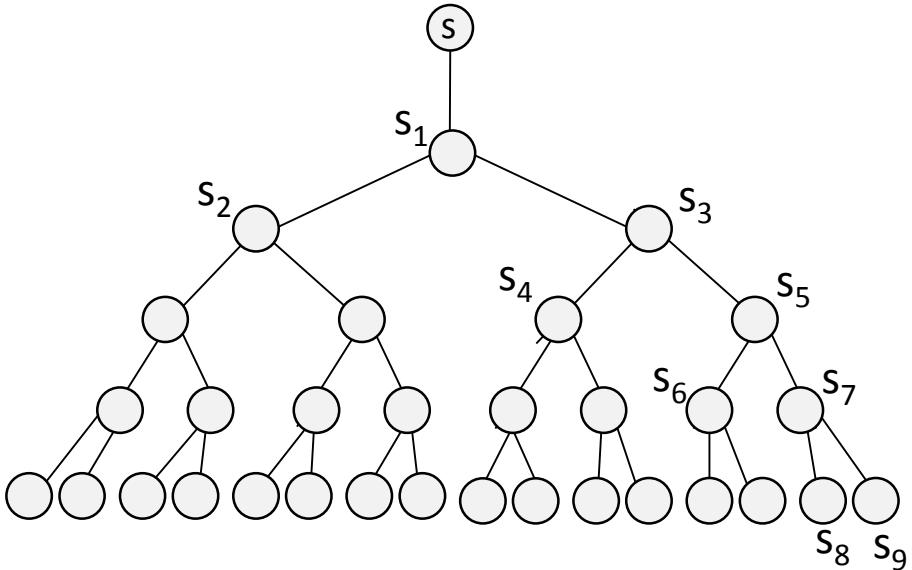
- **Push-down tree:** a self-adjusting complete tree
- **Dynamically optimal**
- Not ordered: requires **a map**



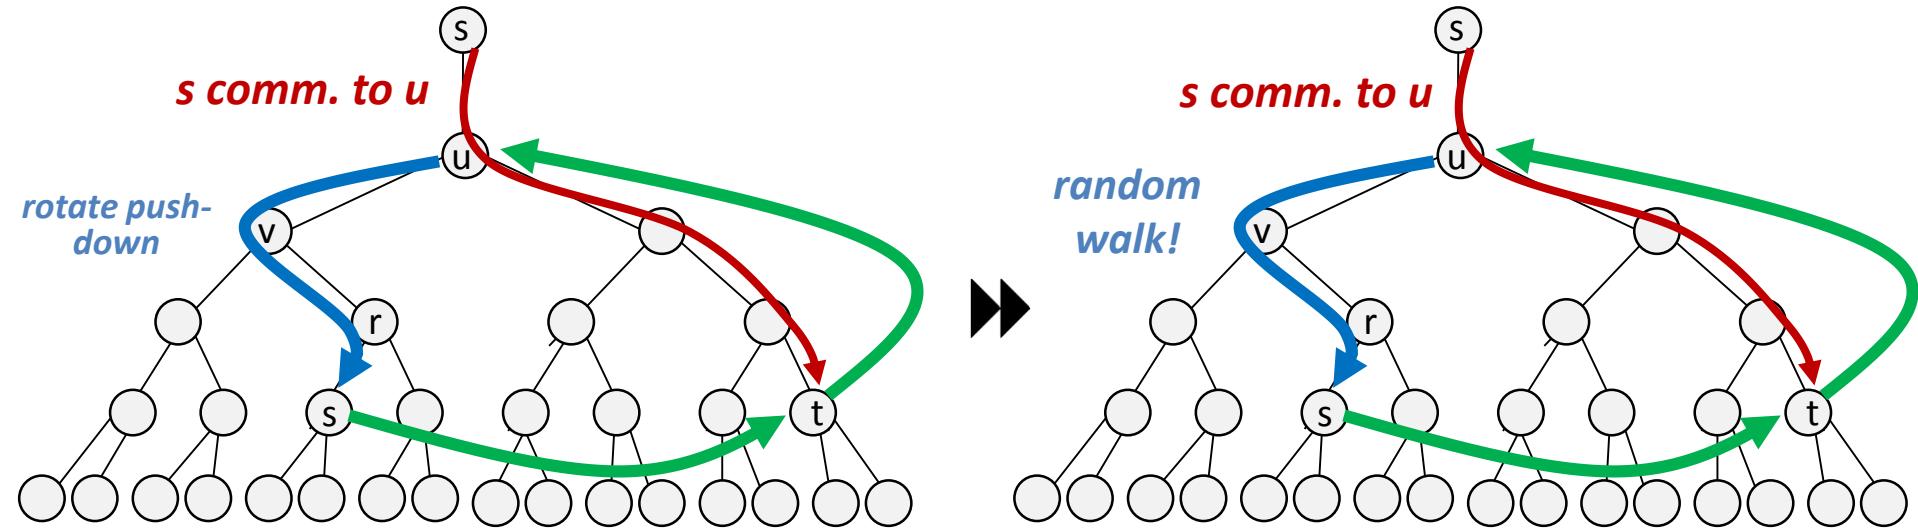
*Then: promote *u* to available root, and
t to *u*: at original depth!*

Remarks

- Unfortunately, alternating push-down does ***not maintain MRU*** (working set) property
- Tree can ***degrade***, e.g.: sequence of requests from level 4,1,2,1,3,1,4,1



Solution: Random Walk



*At least maintains approximate
working set / MRU!*

Further Reading

Push-Down Trees: Optimal Self-Adjusting Complete Trees

Chen Avin, Kaushik Mondal, and Stefan Schmid.

ArXiv Technical Report, July 2018.

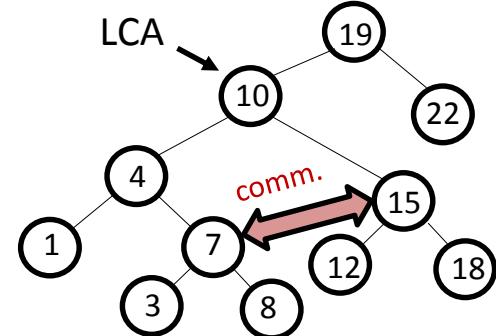
Roadmap

- Vision and Motivation
- An analogy: coding and datastructures
- Principles of Demand-Aware Network (DAN) Designs
- Principles of Self-Adjusting Network (SAN) Designs
- **Principles of Decentralized Approaches**



A “Simple” Decentralized Solution: Distributed SplayNet (*DiSplayNet*)

- SplayNet attractive: ordered BST supports **local routing**
 - Nodes **maintain three ranges**: interval of left subtree, right subtree, upward
- If communicate (frequently): **double-splay** toward LCA
- Challenge: **concurrency**!
 - Access Lemma of splay trees no longer works: **potential function** does not „**telescope**“ anymore: a concurrently rising node may push down another rising node again

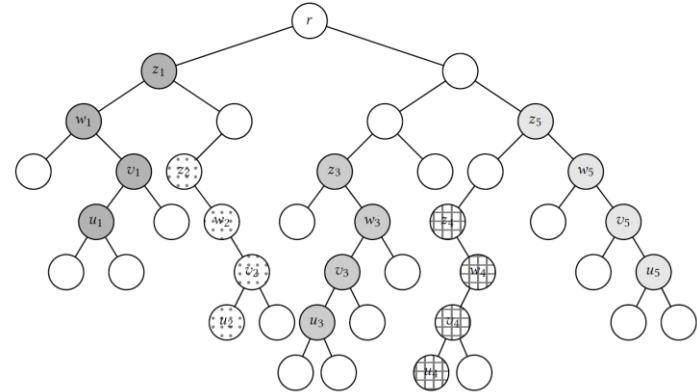


SplayNet

DiSplayNet: Challenges

- DiSplayNet: Rotations (zig,zigzag,zigzag) are **concurrent**
- To avoid conflict: distributed computation of **independent clusters**
- Still challenging:

	1	2	3	4	5	6	7	8	...	$i - 6$	$i - 5$	$i - 4$	$i - 3$	$i - 2$	$i - 1$	i
σ_1	✓	✓	✓	✓	-	-	-	-	...	-	-	-	-	-	-	-
σ_2	-	X	X	X	✓	✓	✓	-	...	-	-	-	-	-	-	-
...
σ_{m-1}	-	-	-	-	-	-	-	-	...	✓	✓	-	-	-	-	-
σ_m	-	-	-	-	-	-	-	-	...	X	X	✓	✓	✓	✓	-



	1	2	3	...	i	$i + 1$	$i + 2$	$i + 3$	$i + 4$	$i + 5$	$i + 6$...	j	...	k
s_1	✓	✓	✓	...	✓	✓	✓	✓	✓	✓	...	-	✓	...	-
d_1	✓	✓	✓	...	✓	✓	✓	✓	✓	✓	...	-	✓	...	-
s_2	-	✓	✓	...	✓	✓	✓	✓	-	-	...	-	-	...	-
d_2	-	✓	✓	...	✓	✓	X	✓	-	-	...	-	-	...	-
s_3	-	-	✓	...	X	X	X	✓	X	X	...	✓	...	-	-
d_3	-	-	✓	...	X	X	X	X	X	X	...	✓	...	-	-

Sequential SplayNet: requests **one after another**

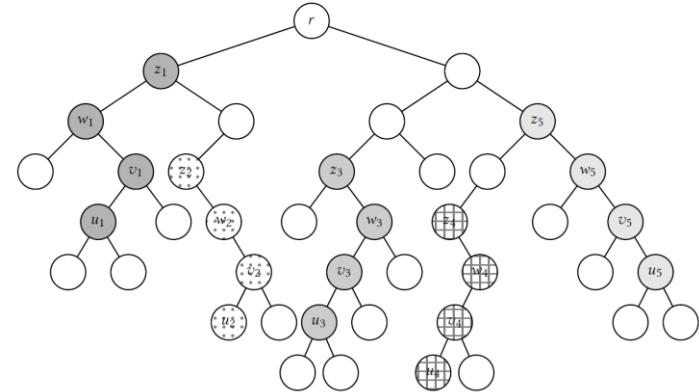
DiSplayNet: Analysis more challenging: potential function sum no longer **telescopic**. One request can “push-down” another.

DiSplayNet: Challenges

- DiSplayNet: Rotations (zig,zigzag,zigzag) are **concurrent**
- To avoid conflict: distributed computation of **independent clusters**
- Still challenging: **Telescopic: max potential drop**

	1	2	3	4	5	6	7	8	...	$i - 6$	$i - 5$	$i - 4$	$i - 3$	$i - 2$	$i - 1$	i
σ_1	✓	✓	✓	✓				-	...	-	-	-	-	-	-	-
σ_2	-	X	X	X	✓	✓	✓	-	...	-	-	-	-	-	-	-
...
σ_{m-1}	-	-	-	-	-	-	-	-	...	✓	✓				-	-
σ_m	-	-	-	-	-	-	-	-	...	X	X	✓	✓	✓	✓	-

Sequential SplayNet: requests **one after another**



	1	2	3	...	i	$i + 1$	$i + 2$	$i + 3$	$i + 4$	$i + 5$	$i + 6$...	j	...	k
s_1	✓	✓	✓	...	✓	✓	✓	✓	✓	✓	...	-	✓	...	-
d_1	✓	✓	✓	...	✓	✓	✓	✓	✓	✓	...	-	✓	...	-
s_2	-	✓	✓	...	✓	✓	✓	✓	-	-	...	-	-	...	-
d_2	-	✓	✓	...	✓	✓	X	✓	-	-	...	-	-	...	-
s_3	-	-	✓	...	X	X	X	✓	X	X	...	✓	...	-	-
d_3	-	-	✓	...	X	X	X	X	X	X	...	✓	...	-	-

DiSplayNet: Analysis more challenging: potential function sum no longer **telescopic**. One request can “push-down” another.

Further Reading

Brief Announcement: Distributed SplayNets

Bruna Peres, Olga Goussevskaia, Stefan Schmid, and Chen Avin.

31st International Symposium on Distributed Computing (**DISC**), Vienna,
Austria, October 2017.

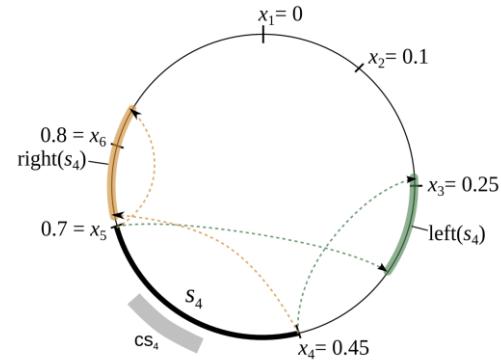
Roadmap

- Vision and Motivation
- An analogy: coding and datastructures
- Principles of Demand-Aware Network (DAN) Designs
- Principles of Self-Adjusting Network (SAN) Designs
- Principles of Decentralized Approaches



Many Open Questions

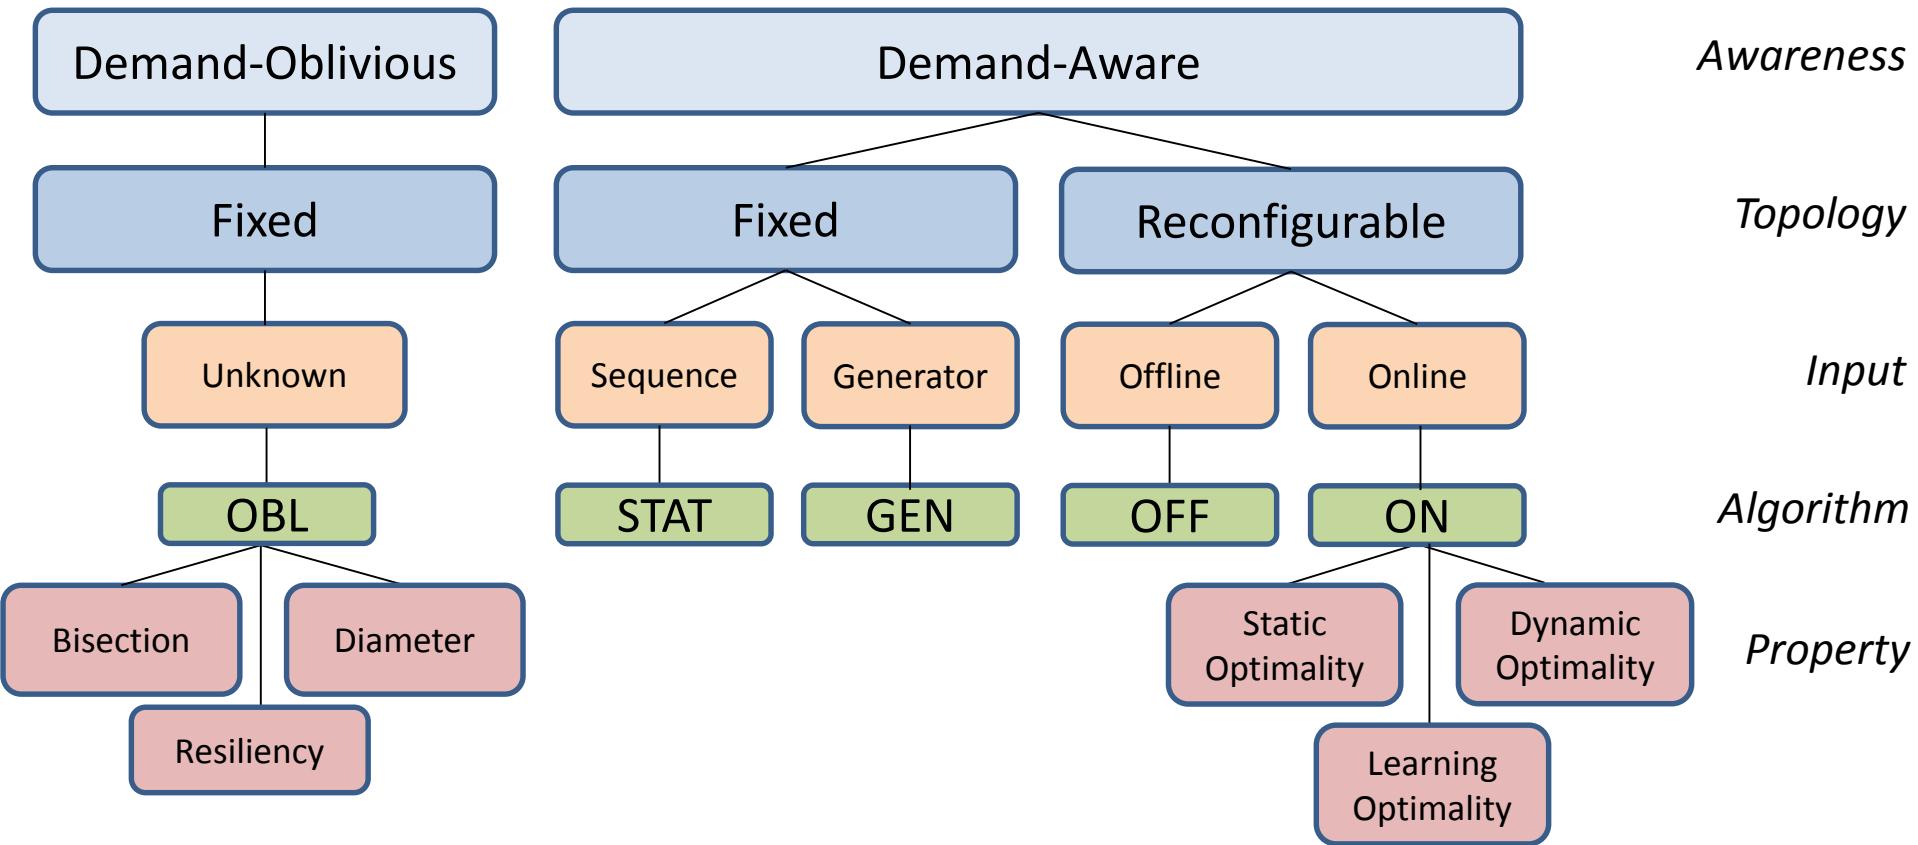
- E.g., robust demand-aware networks?
 - First idea: **rDAN**, based on continuous-discrete approach and Shannon-Fano-Elias coding
- Serving dense communication patterns
- Distributed version
- ...



rDAN: Toward Robust Demand-Aware Network Designs. Chen Avin, Alexandr Hercules, Andreas Loukas, and Stefan Schmid. Information Processing Letters (IPL), Elsevier, 2018.

Uncharted Landscape!

Toward Demand-Aware Networking: A Theory for
Self-Adjusting Networks. ArXiv 2018.



Conclusion

- Networked systems become reconfigurable
- First techniques emerging (e.g., ego-trees)
- How much it helps depends on spatial and temporal locality
- Entropy seems to be a useful metric

*Thank you!
Question?*

Further Reading

[Toward Demand-Aware Networking: A Theory for Self-Adjusting Networks](#)

Chen Avin and Stefan Schmid.

ArXiv Technical Report, July 2018.

[Demand-Aware Network Designs of Bounded Degree](#)

Chen Avin, Kaushik Mondal, and Stefan Schmid.

31st International Symposium on Distributed Computing (**DISC**), Vienna, Austria, October 2017.

[Push-Down Trees: Optimal Self-Adjusting Complete Trees](#)

Chen Avin, Kaushik Mondal, and Stefan Schmid.

ArXiv Technical Report, July 2018.

[Online Balanced Repartitioning](#)

Chen Avin, Andreas Loukas, Maciej Pacut, and Stefan Schmid.

30th International Symposium on Distributed Computing (**DISC**), Paris, France, September 2016.

[rDAN: Toward Robust Demand-Aware Network Designs](#)

Chen Avin, Alexandr Hercules, Andreas Loukas, and Stefan Schmid.

Information Processing Letters (**IPL**), Elsevier, 2018.

[SplayNet: Towards Locally Self-Adjusting Networks](#)

Stefan Schmid, Chen Avin, Christian Scheideler, Michael Borokhovich, Bernhard Haeupler, and Zvi Lotker.

IEEE/ACM Transactions on Networking (**TON**), Volume 24, Issue 3, 2016. Early version: IEEE **IPDPS** 2013.

[Characterizing the Algorithmic Complexity of Reconfigurable Data Center Architectures](#)

Klaus-Tycho Foerster, Monia Ghobadi, and Stefan Schmid.

ACM/IEEE Symposium on Architectures for Networking and Communications Systems (**ANCS**), Ithaca, New York, USA, July 2018.

[Charting the Complexity Landscape of Virtual Network Embeddings](#)

Matthias Rost and Stefan Schmid. **IFIP Networking**, Zurich, Switzerland, May 2018.

Chapter 72: [Overlay Networks for Peer-to-Peer Networks](#)

Andrea Richa, Christian Scheideler, and Stefan Schmid.

Handbook on Approximation Algorithms and Metaheuristics (**AAM**), 2nd Edition, 2017.