

# Performance and Security Isolation in Softwarized Networks: Advances and Challenges

**Stefan Schmid**

Faculty of Computer Science  
University of Vienna

# *Softwarized Networks:*

## It's a great time to be a networking researcher!



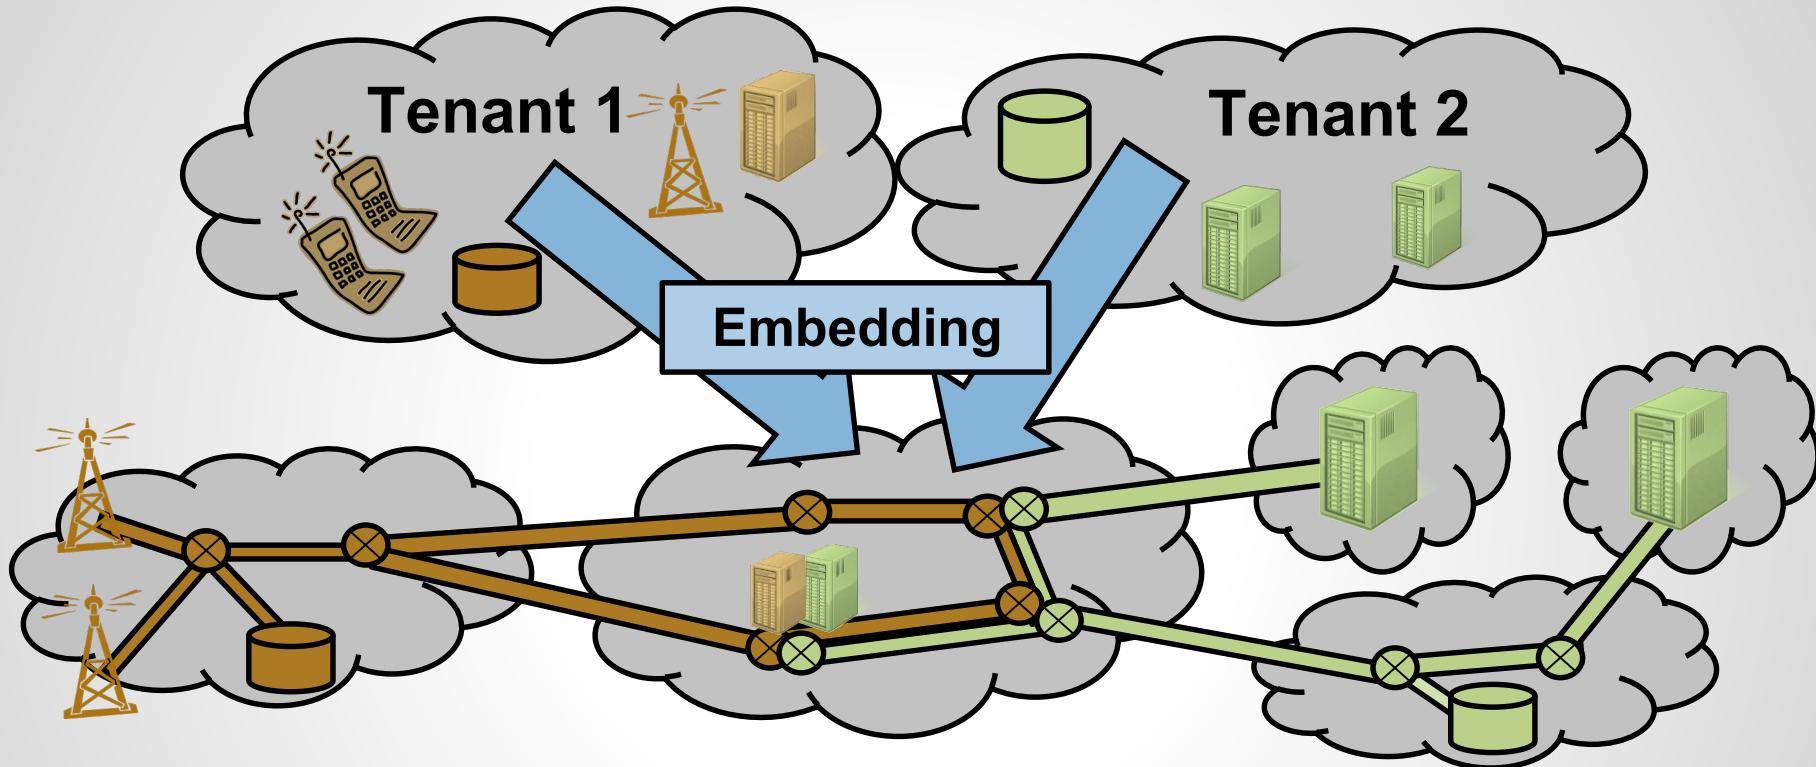
Rhone and Arve Rivers,  
Switzerland

Credits: George Varghese.

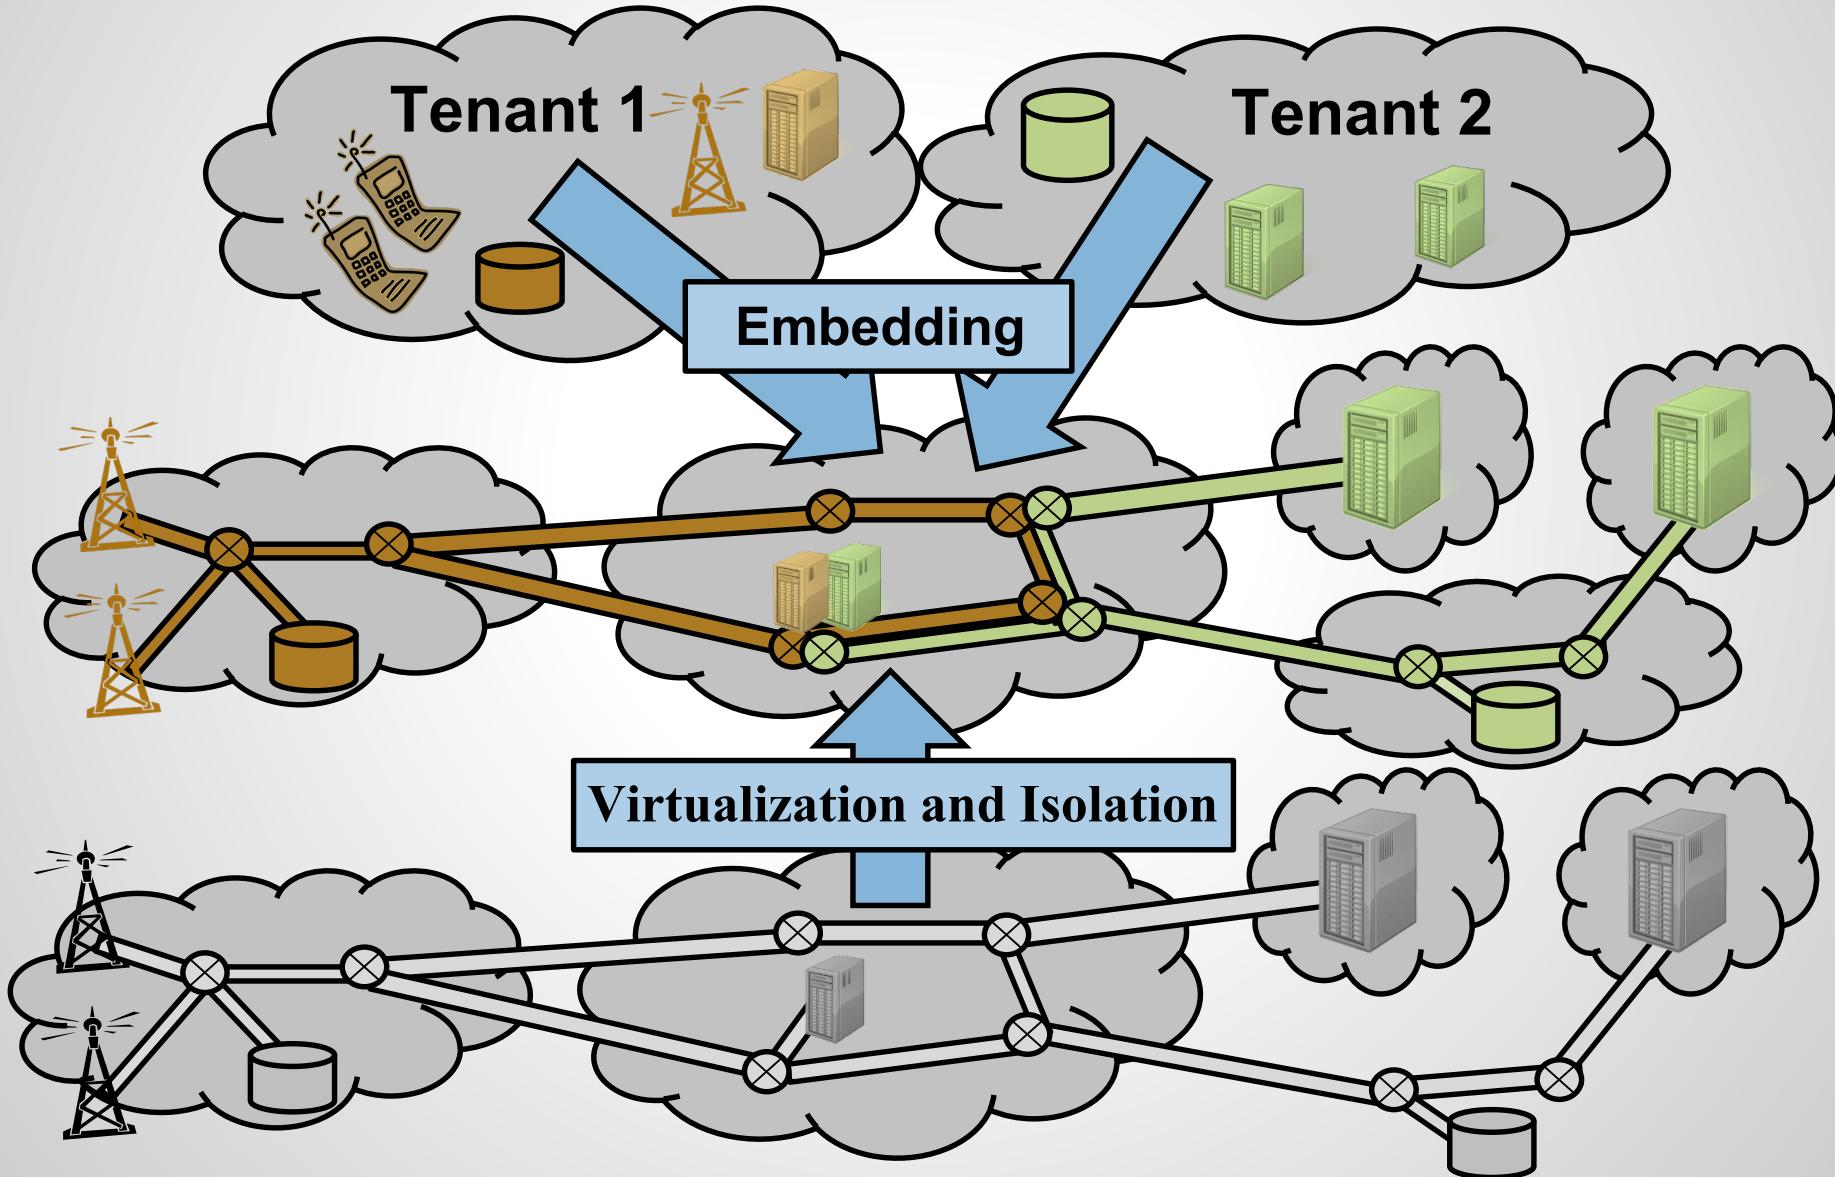
# The (At Least) 3 Dimensions of Network Flexibility



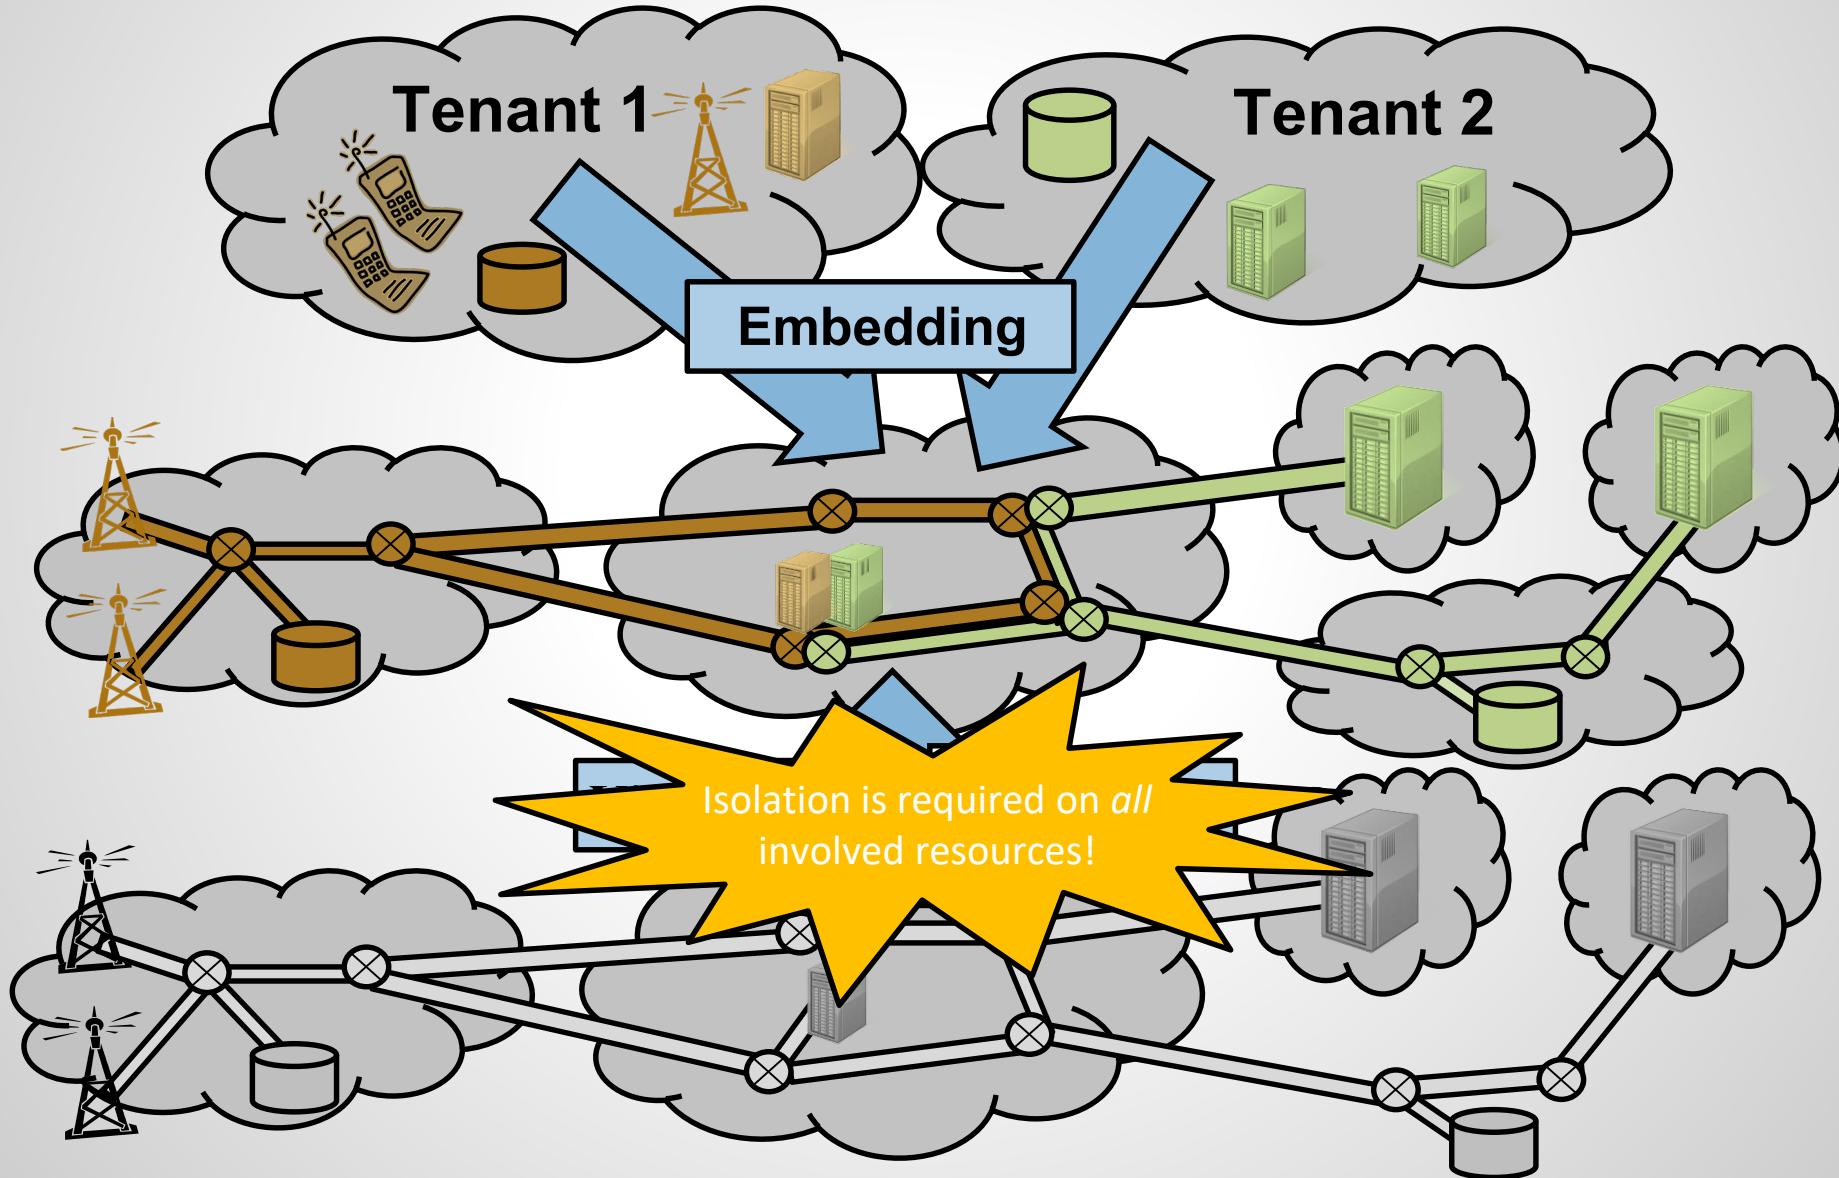
# Opportunity: Improved Sharing of Physical Network Infrastructure



# Challenges: Isolation (and Embedding)



# Challenges: Isolation (and Embedding)



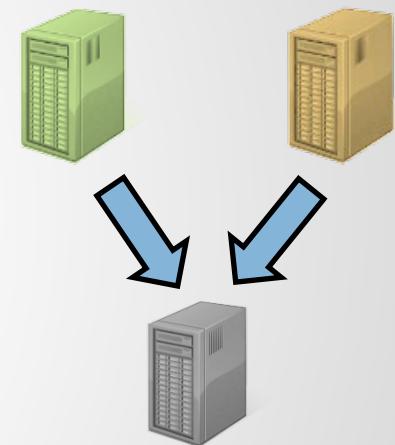
# Two Flavors of Isolation

- ❑ Logical isolation

- ❑ E.g., prevent from communication, no need to coordinate name/address space, etc.
  - ❑ Relevant for **security**

- ❑ Performance isolation

- ❑ E.g., prevent resource interference, ensure SLAs, make it appear like a dedicated infrastructure
  - ❑ Relevant for **quality-of-service**



We'll consider both in this talk!

# Invitation: A Roadtrip Through The Opportunities and Challenges of Network Isolation

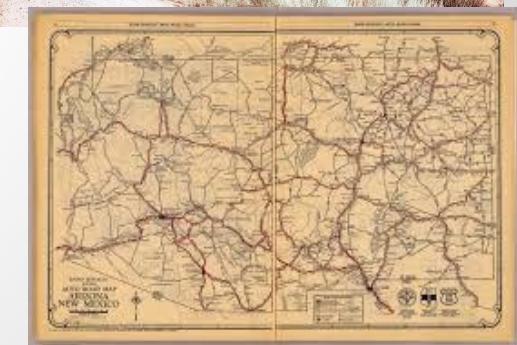
- ❑ Opportunities
  - ❑ Algorithmic opportunities
  - ❑ Technological opportunities
- ❑ Challenges
  - ❑ Modelling challenges
  - ❑ Security challenges
- ❑ A perspective how AI can improve slicing efficiency and security



Road map 1927: Arizona and New Mexico

# Invitation: A Roadtrip Through The Opportunities and Challenges of Network Isolation

- ❑ Opportunities
  - ❑ Algorithmic opportunities
  - ❑ Technological opportunities
- ❑ Challenges
  - ❑ Modelling challenges
  - ❑ Security challenges
- ❑ A perspective how AI can improve slicing efficiency and security



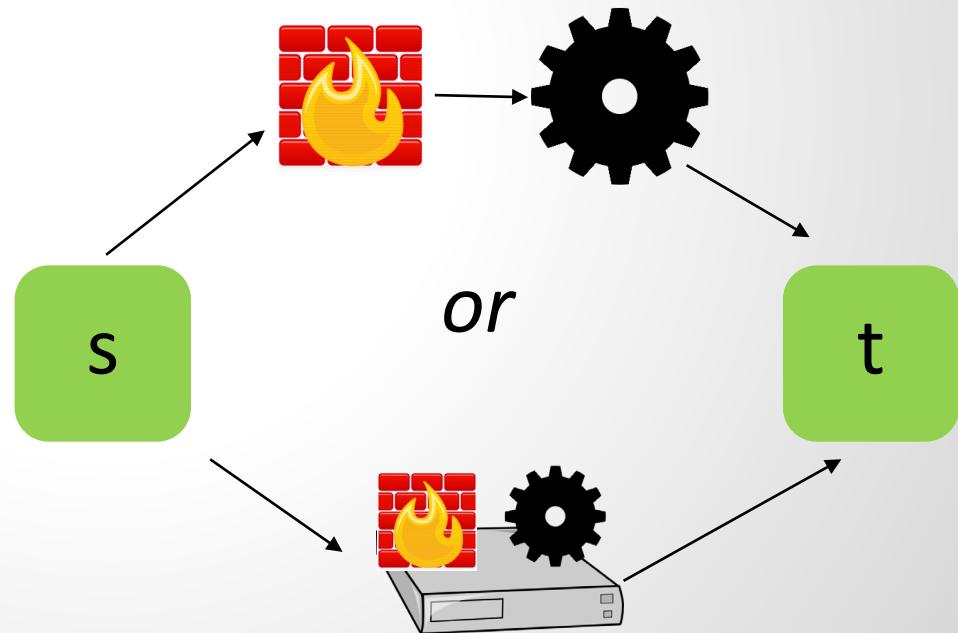
Road map 1927: Arizona and New Mexico

# Opportunity: Define and Flexibly Allocate Complex Services



Steer traffic through network functions to compose complex service chains

More complex requests:  
allowing for alternatives and  
different decompositions

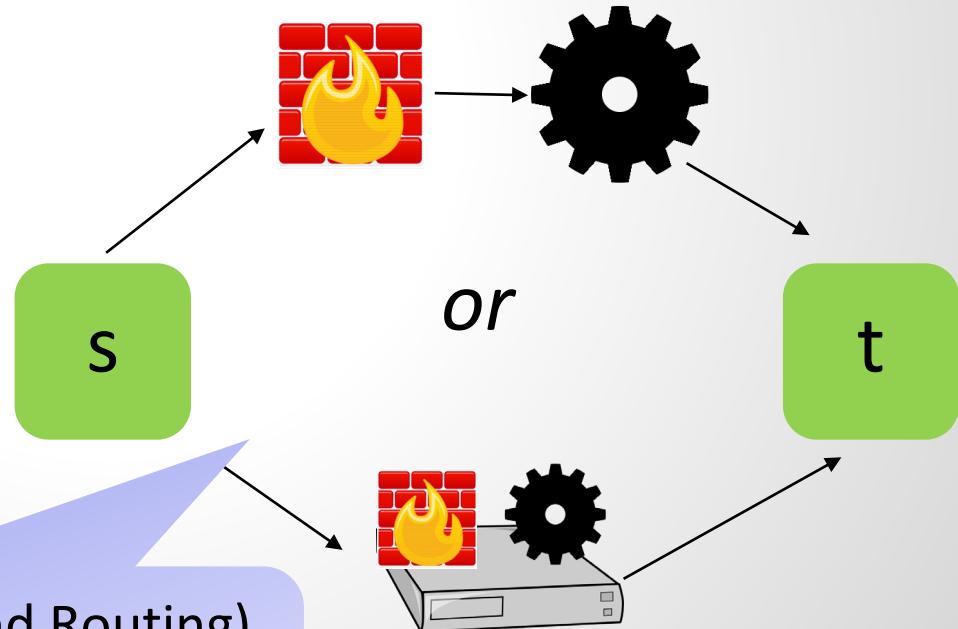


# Opportunity: Define and Flexibly Allocate Complex Services



Steer traffic through network functions to compose complex service chains

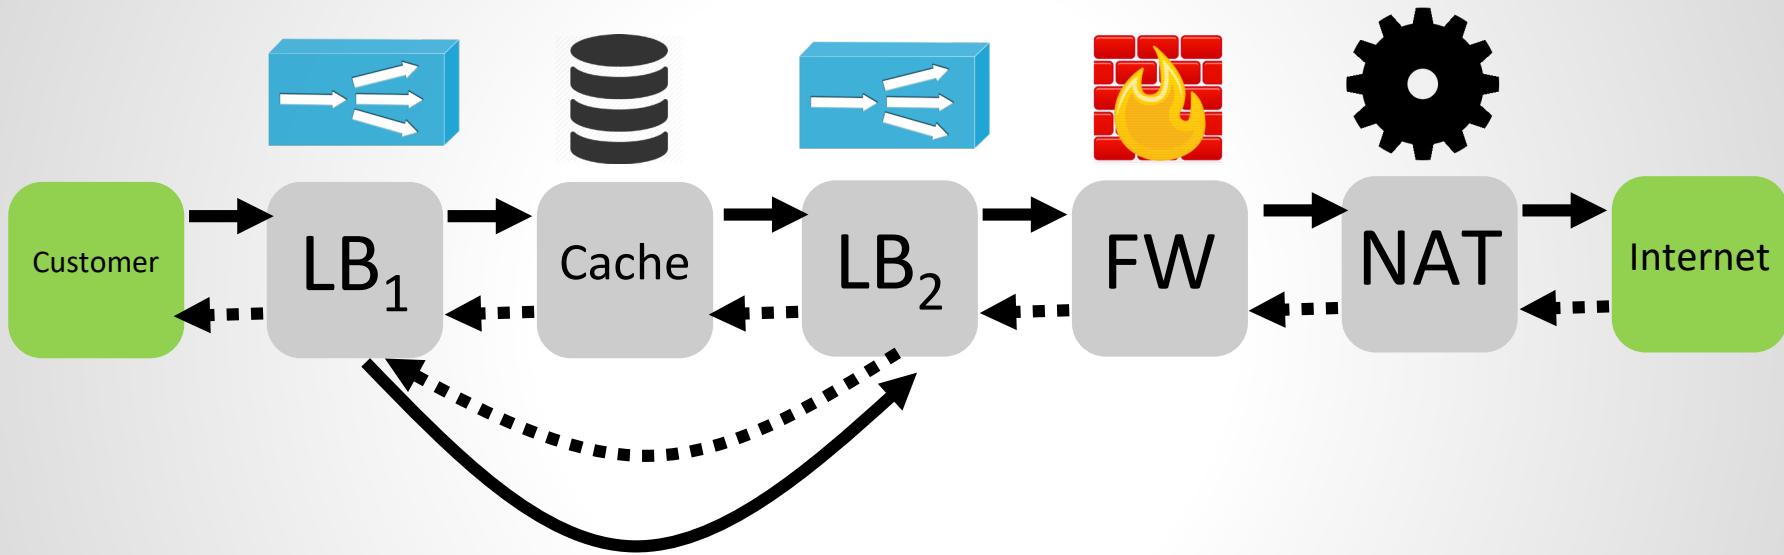
More complex requests:  
allowing for alternatives and  
different decompositions



Known as PR (Processing and Routing)  
Graph: allows to model different  
choices and implementations!

# More Complex Service Chains

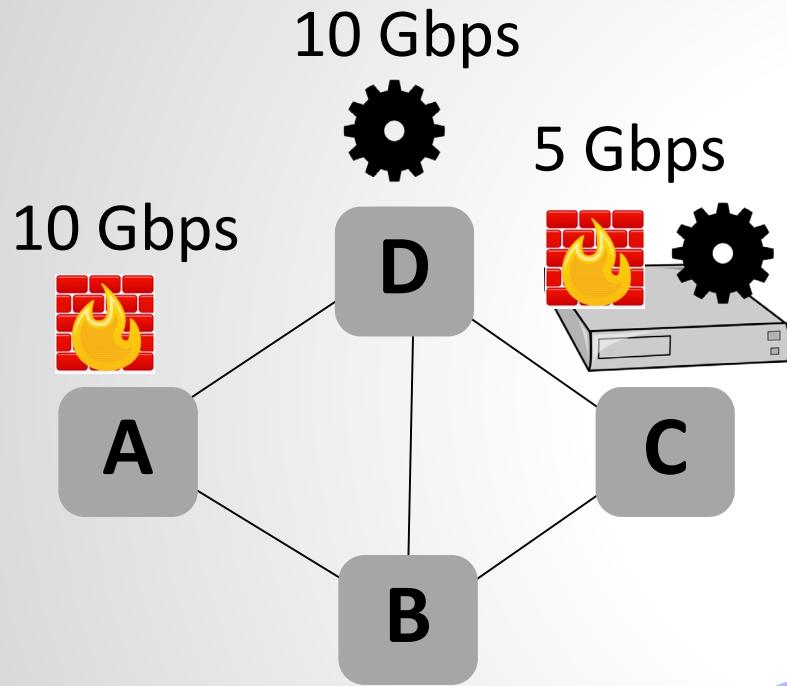
IETF Draft: Service chain for mobile operators



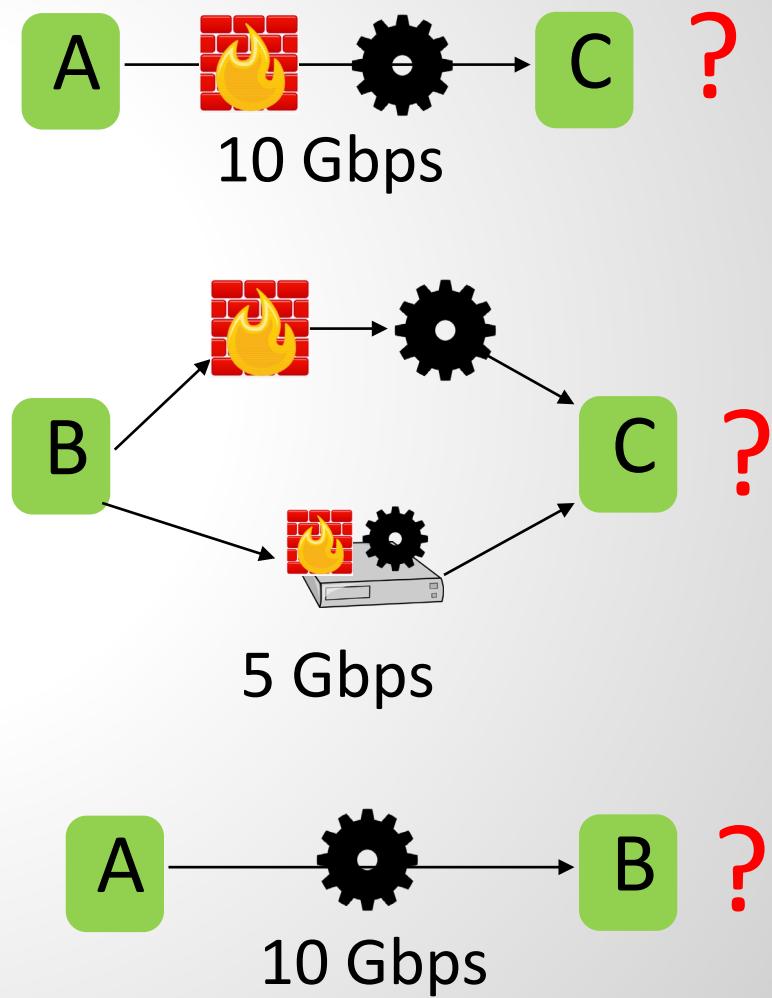
- Load-balancers are used to route (parts of) the traffic through cache

# Algorithmic Challenges: Admission Control and Embedding

Substrate:



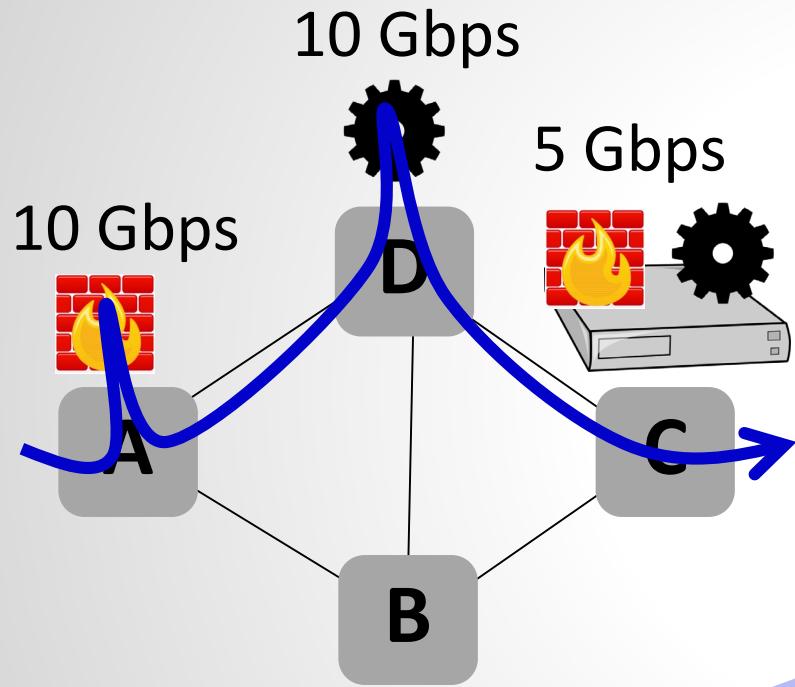
Requests:



Which ones can be  
admitted and embedded?

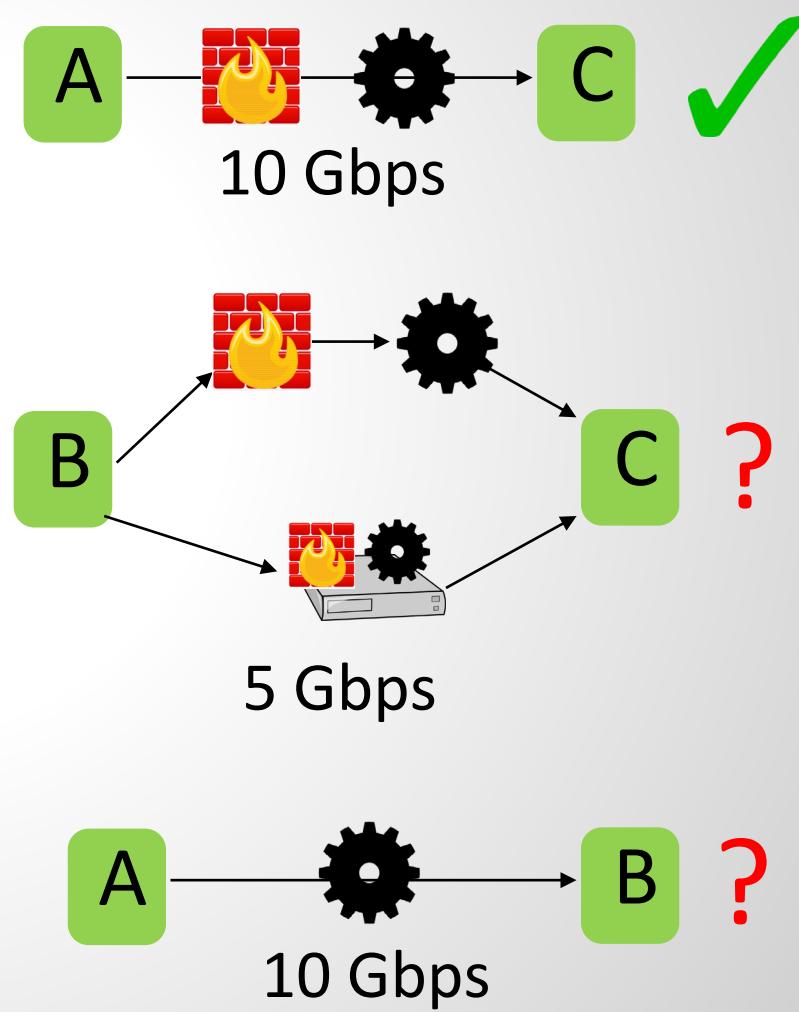
# Algorithmic Challenges: Admission Control and Embedding

Substrate:



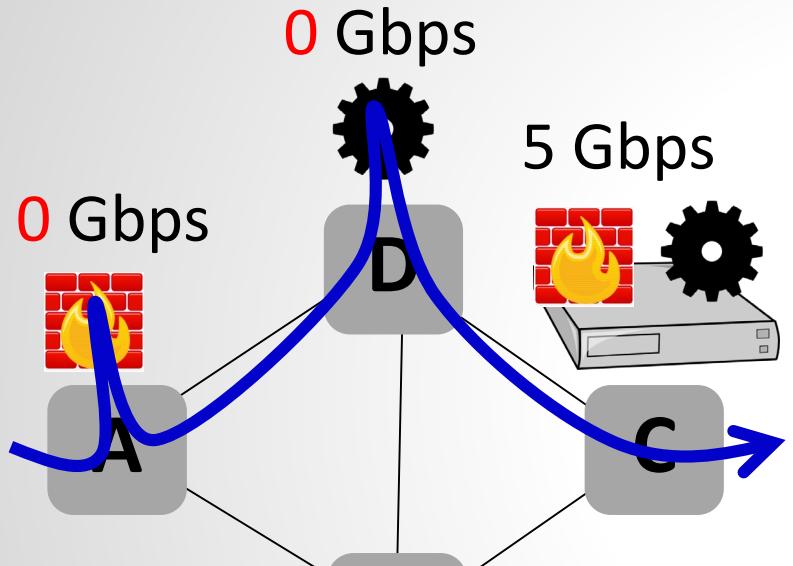
Which ones can be  
admitted and embedded?

Requests:

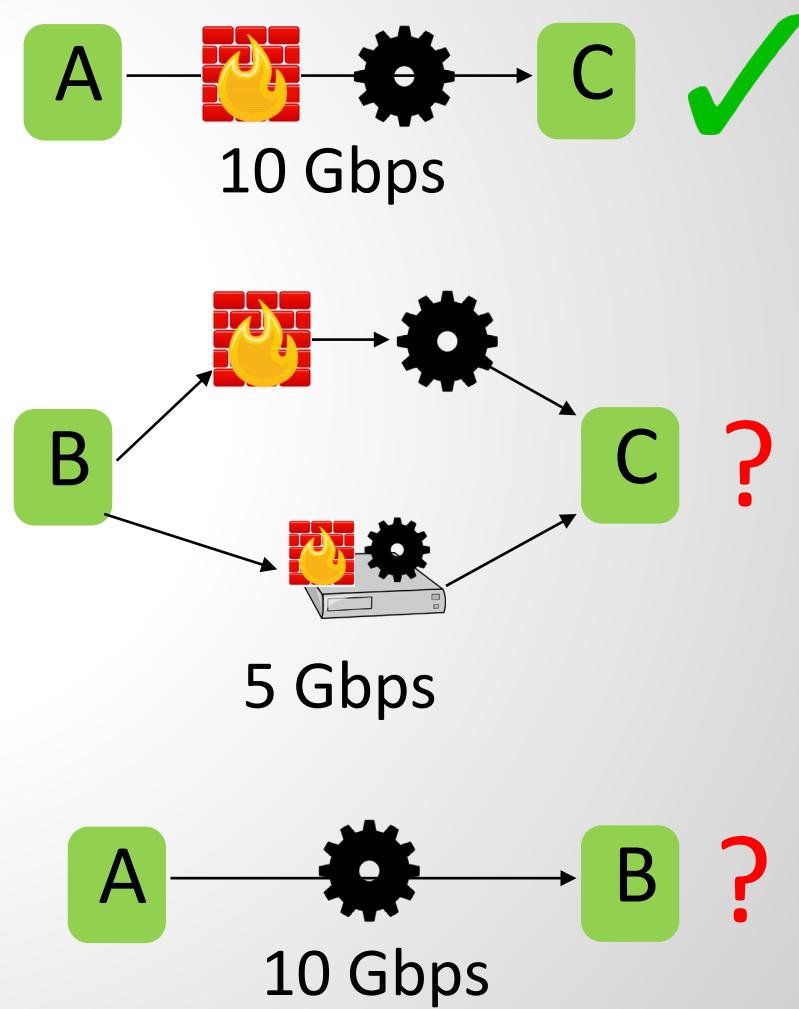


# Algorithmic Challenges: Admission Control and Embedding

Substrate:



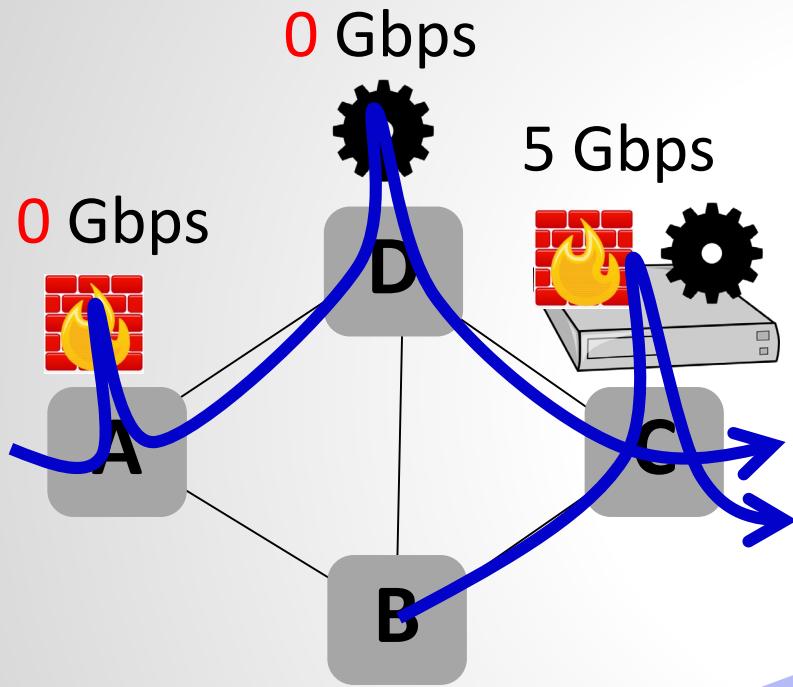
Requests:



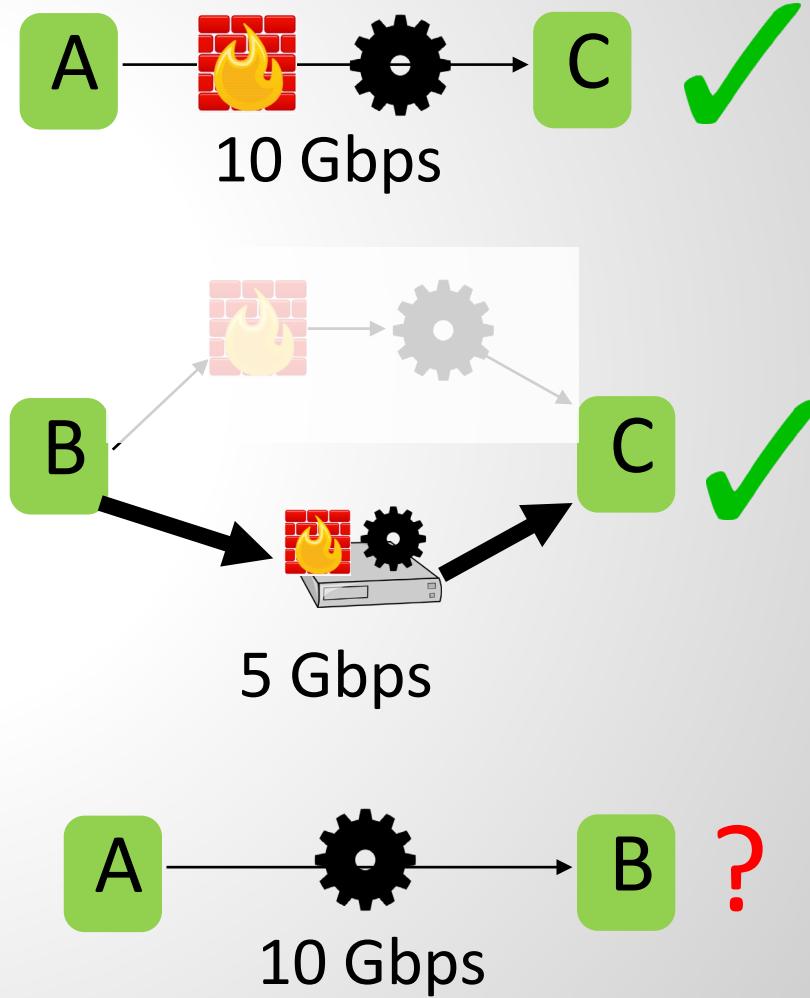
Deduct resources for  
performance isolation!

# Algorithmic Challenges: Admission Control and Embedding

Substrate:



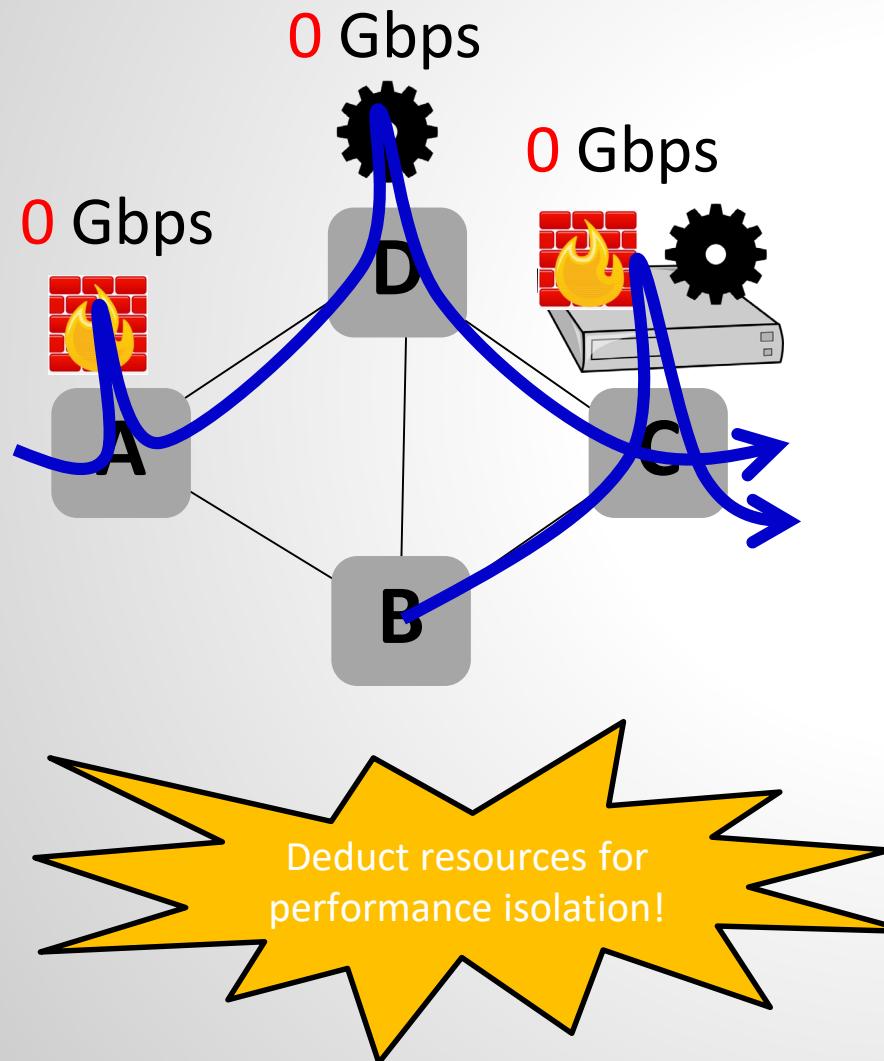
Requests:



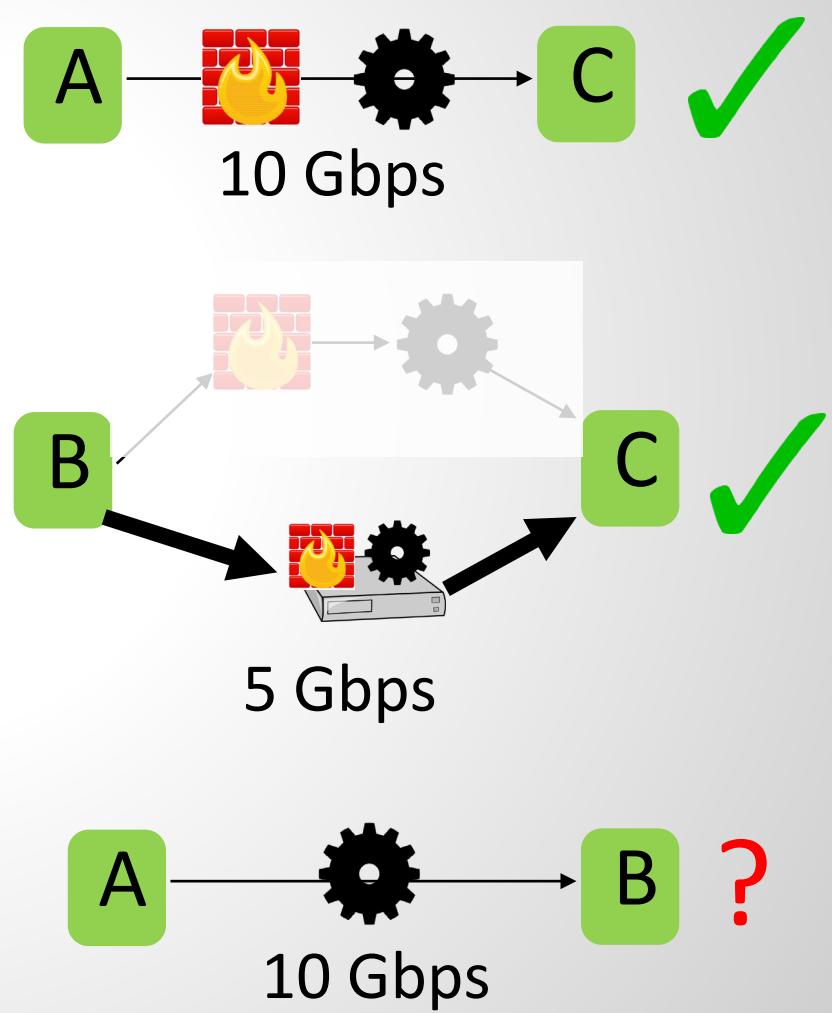
Which ones can be  
admitted and embedded?

# Algorithmic Challenges: Admission Control and Embedding

Substrate:

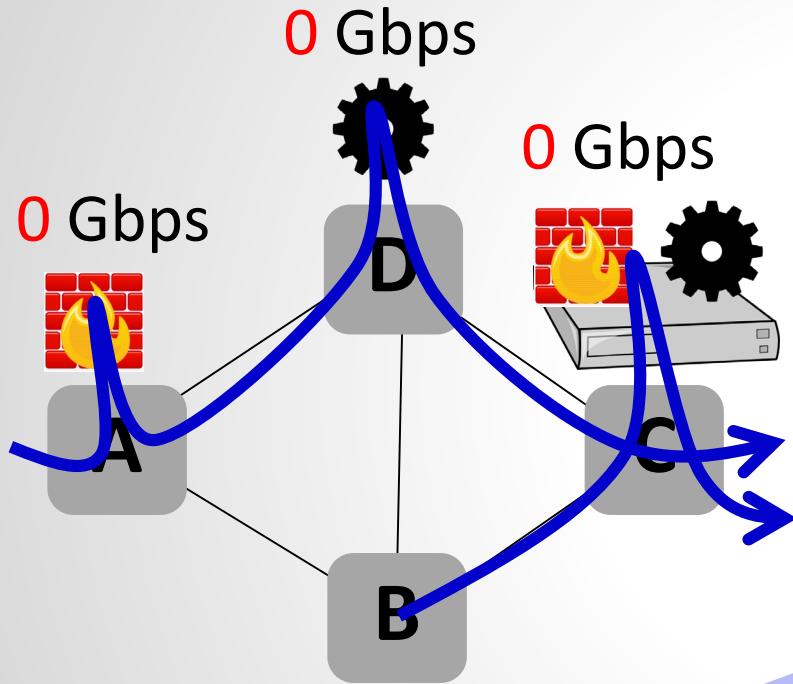


Requests:



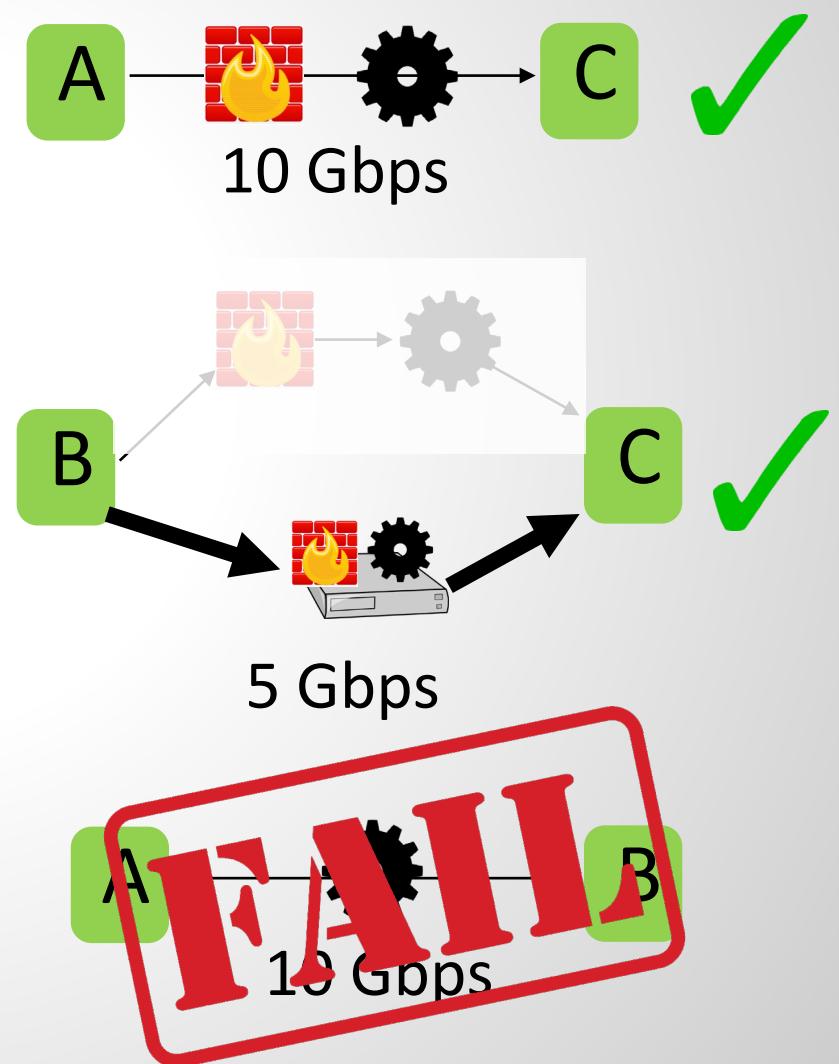
# Algorithmic Challenges: Admission Control and Embedding

Substrate:



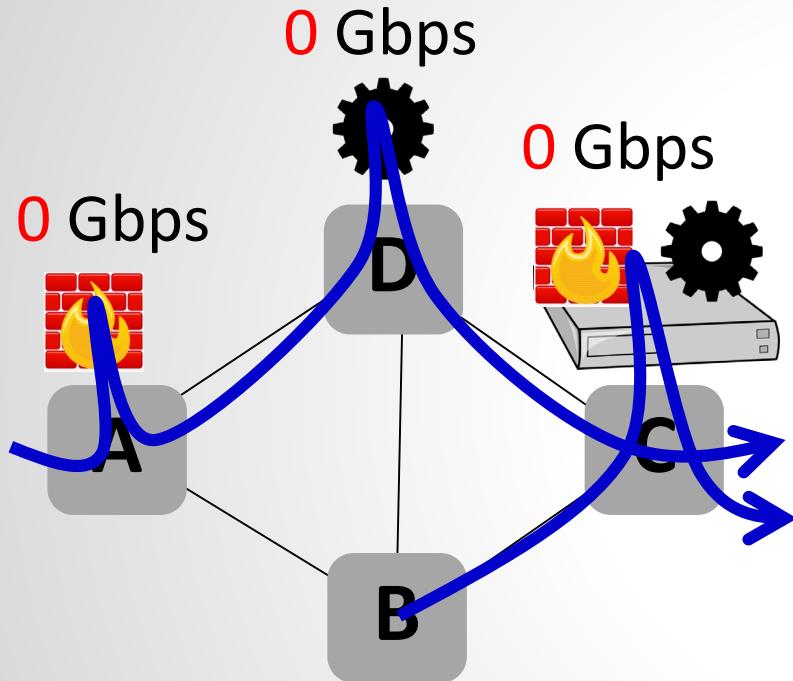
Which ones can be admitted and embedded?

Requests:

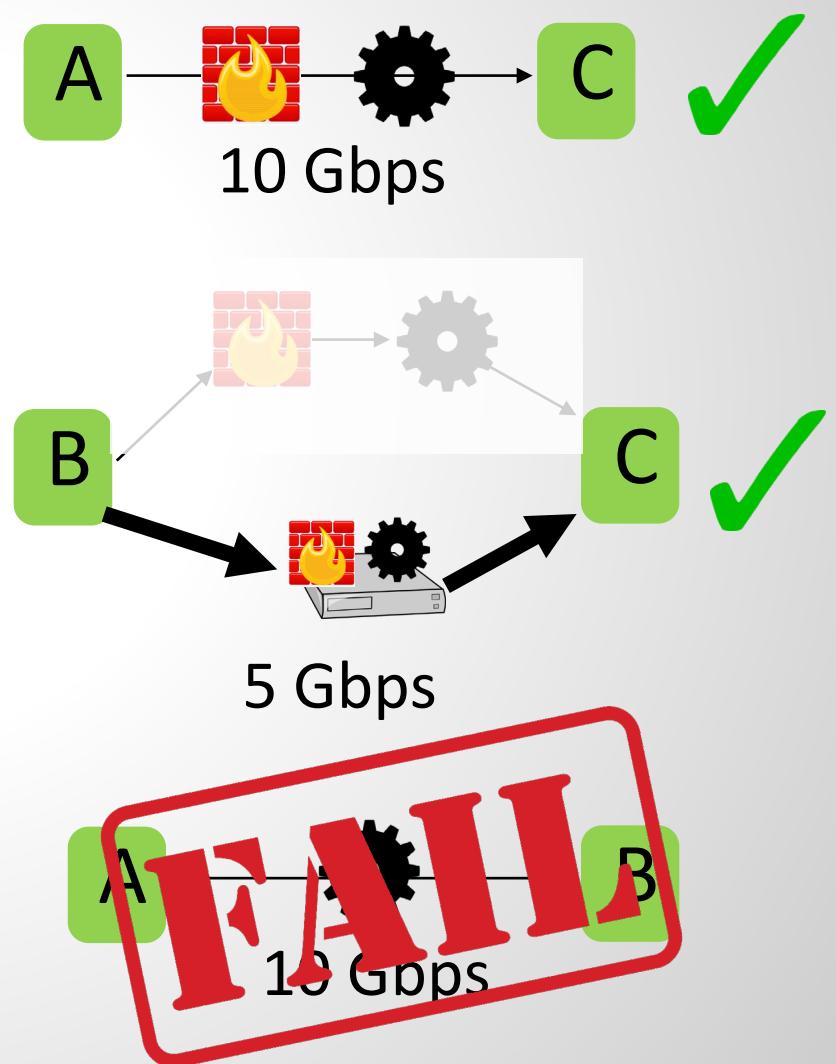


# Algorithmic Challenges: Admission Control and Embedding

Substrate:



Requests:

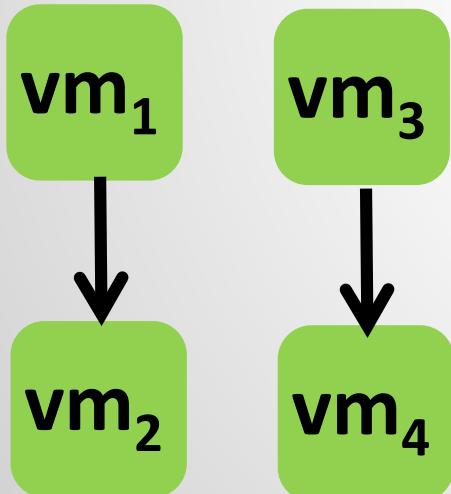


Essentially a **virtual network**  
embedding problem!

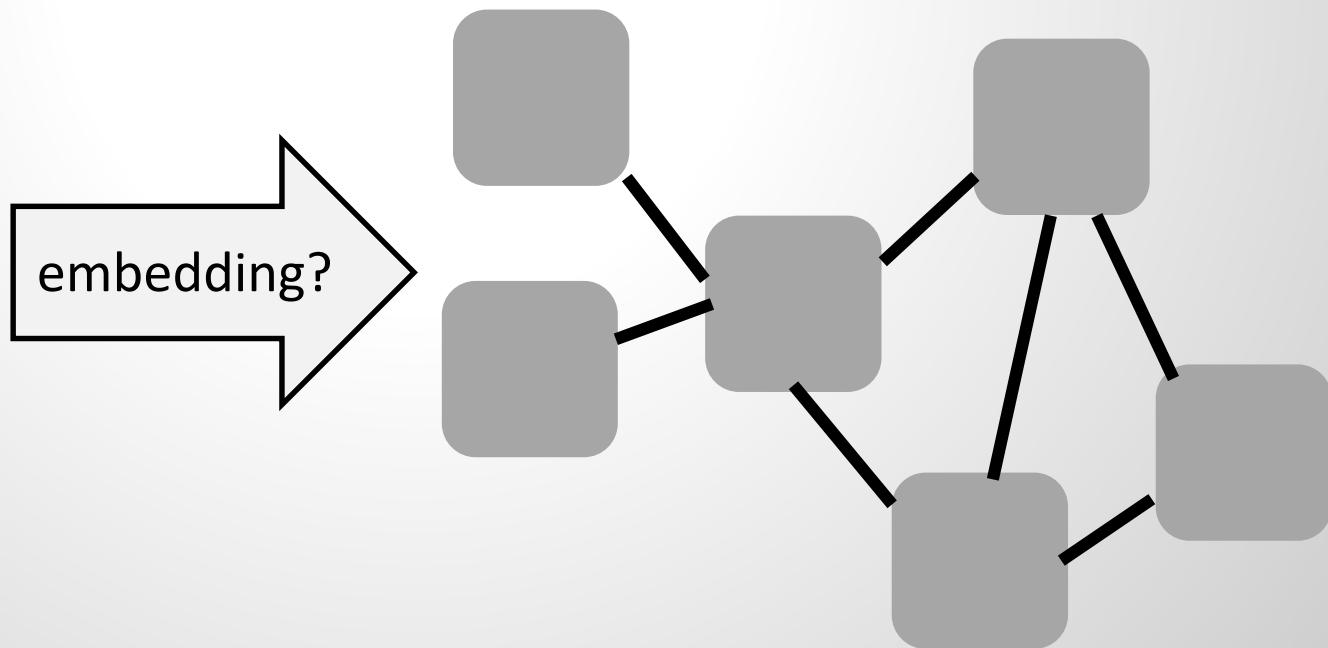
# The Virtual Network Embedding Problem

- ❑ A fundamental resource allocation problem
- ❑ 2 dimensions of flexibility:
  - ❑ Mapping of virtual **nodes** (to physical nodes)
  - ❑ Mapping of virtual **links** (to paths)

VNet



Substrate



# The Virtual Network Embedding Problem

- ❑ A fundamental resource allocation problem
- ❑ 2 dimensions of flexibility:

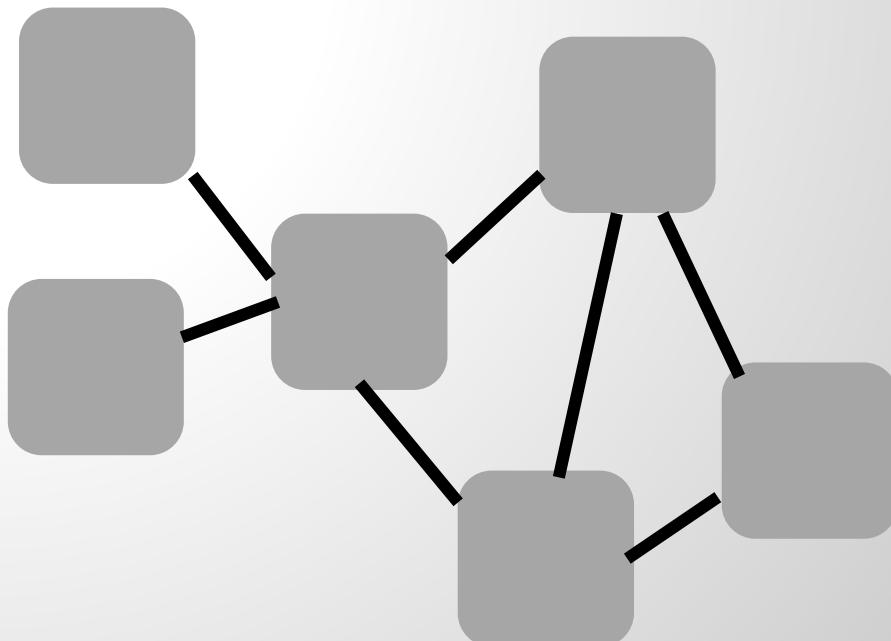
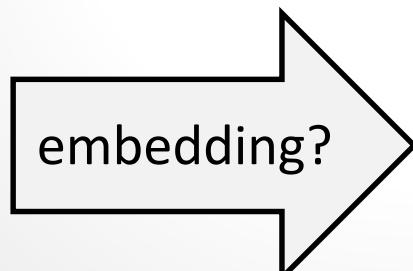
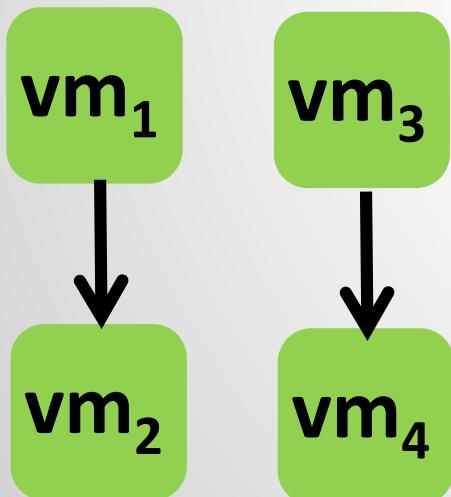
❑ aka “guest graph”

nodes (to physical nodes)

aka “host graph”

VNet

Substrate



# The Virtual Network Embedding Problem

- ❑ A fundamental resource allocation problem
- ❑ 2 dimensions of flexibility:

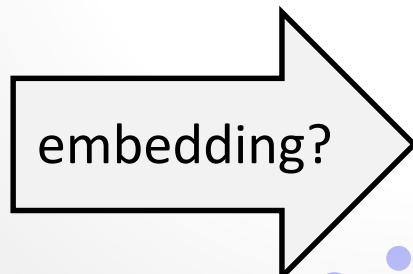
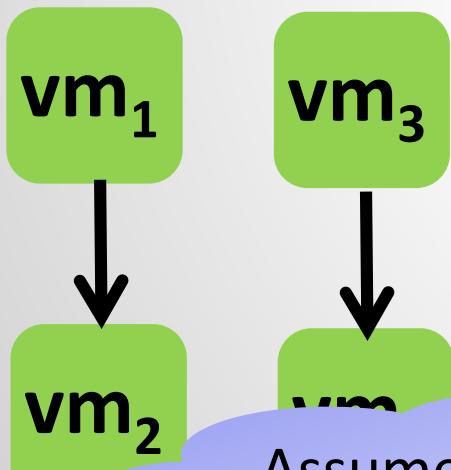
❑ aka “guest graph”

nodes (to physical nodes)

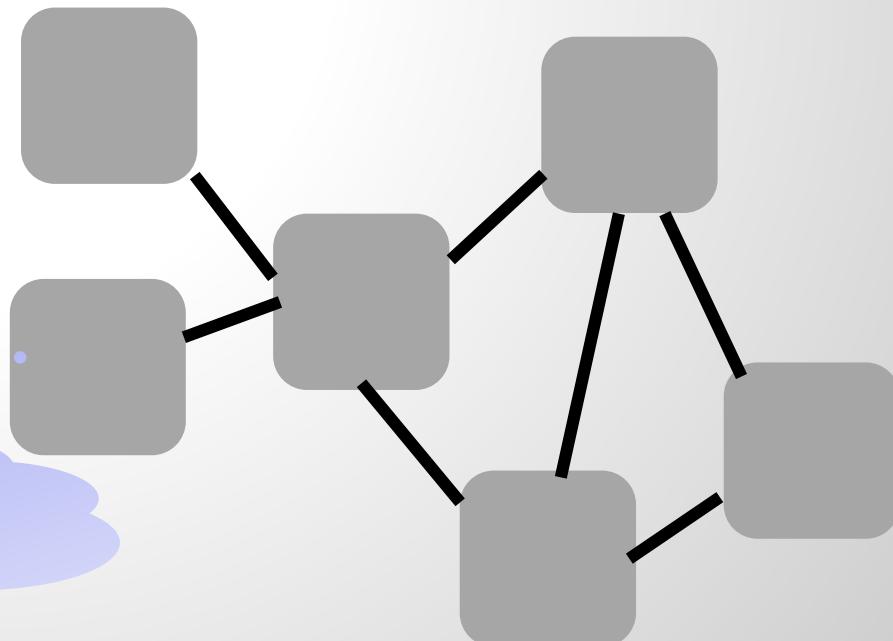
❑ aka “host graph”

VNet

Substrate



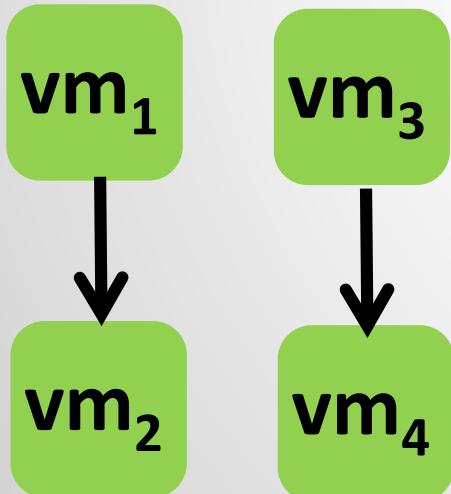
Assume unit demand  
and capacity!



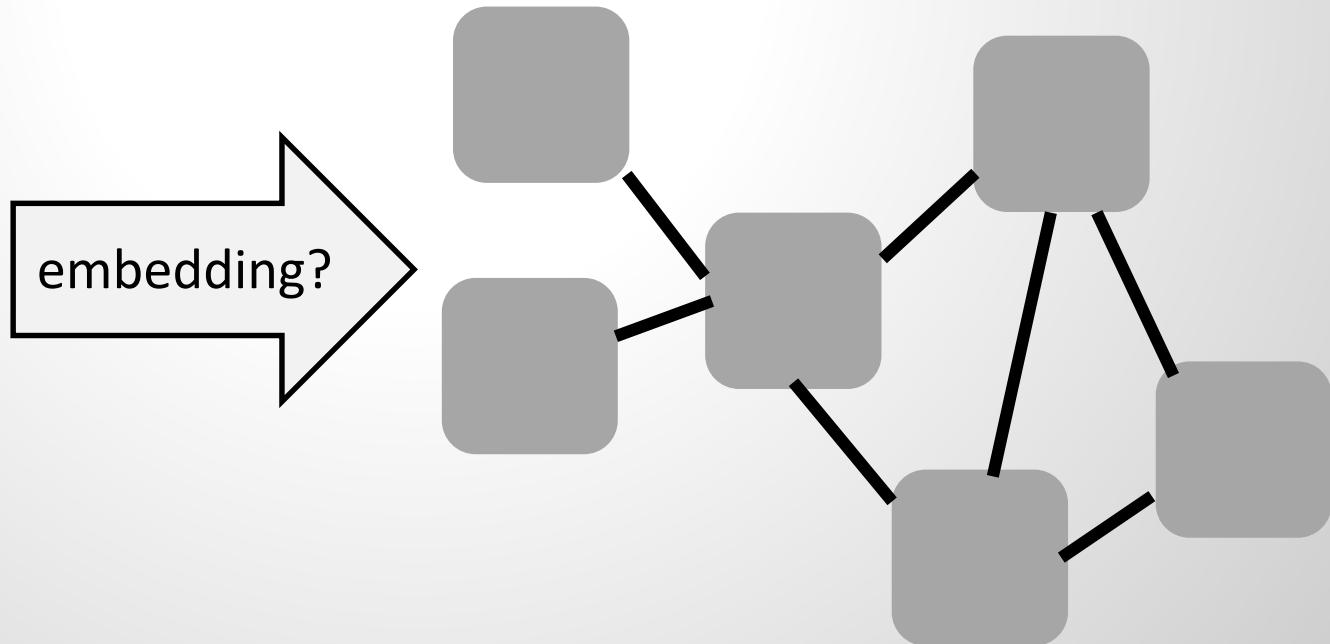
# The Virtual Network Embedding Problem

- Let's start simple: assume node **mappings** are given

VNet



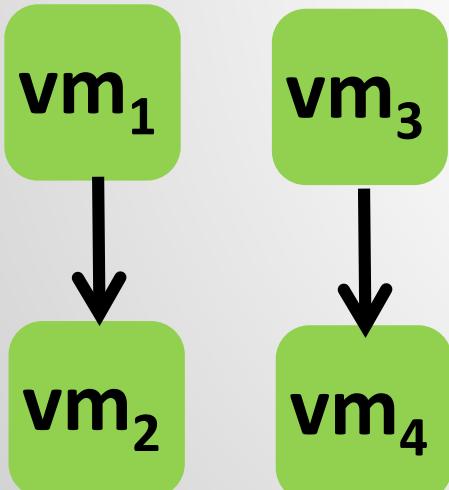
Substrate



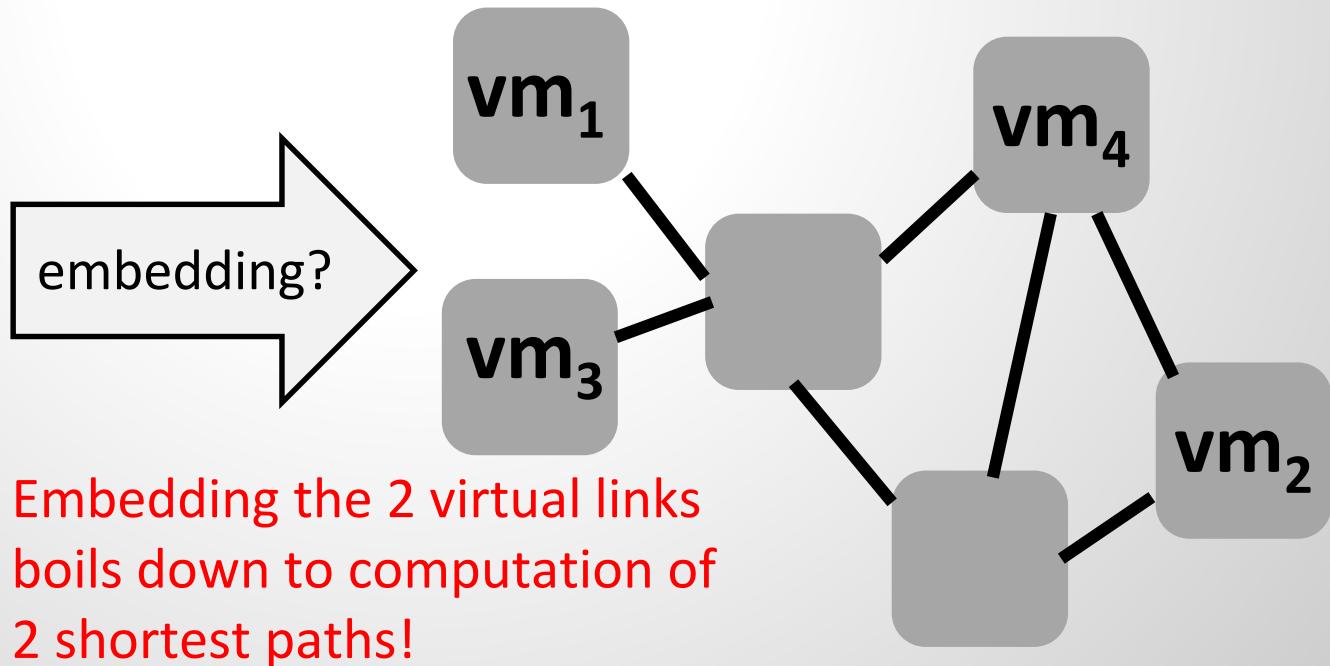
# The Virtual Network Embedding Problem

- Let's start simple: assume node **mappings** are given

VNet



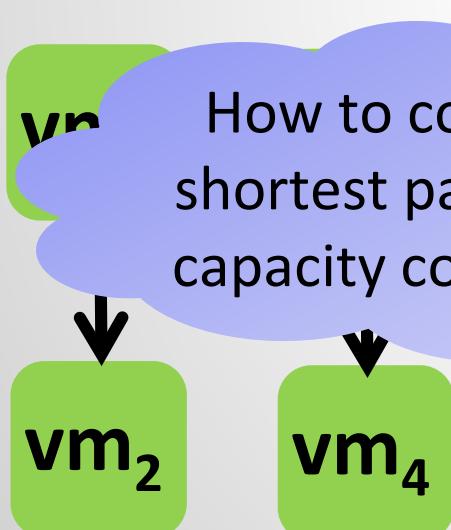
Substrate



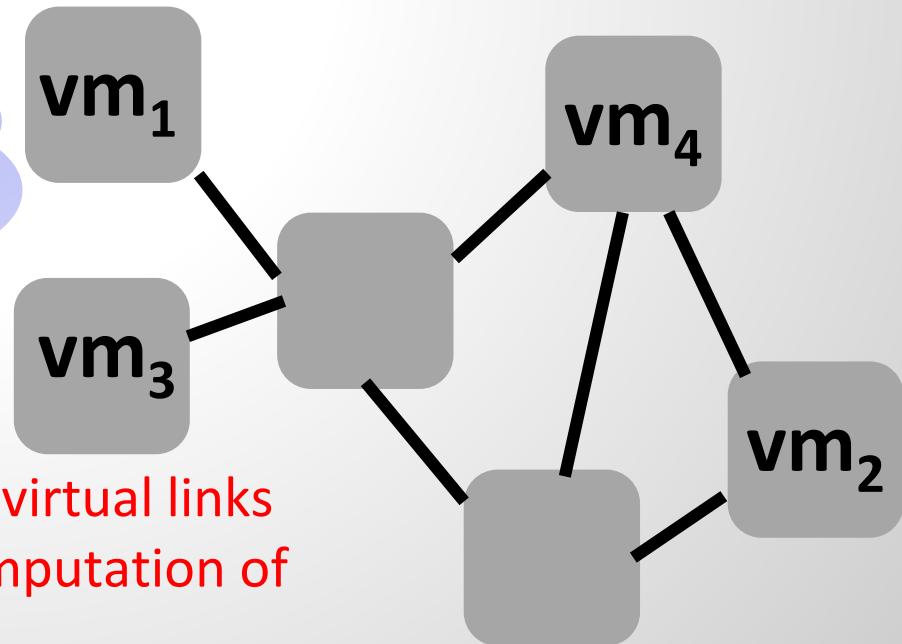
# The Virtual Network Embedding Problem

- Let's start simple: assume node **mappings** are given

VNet



Substrate

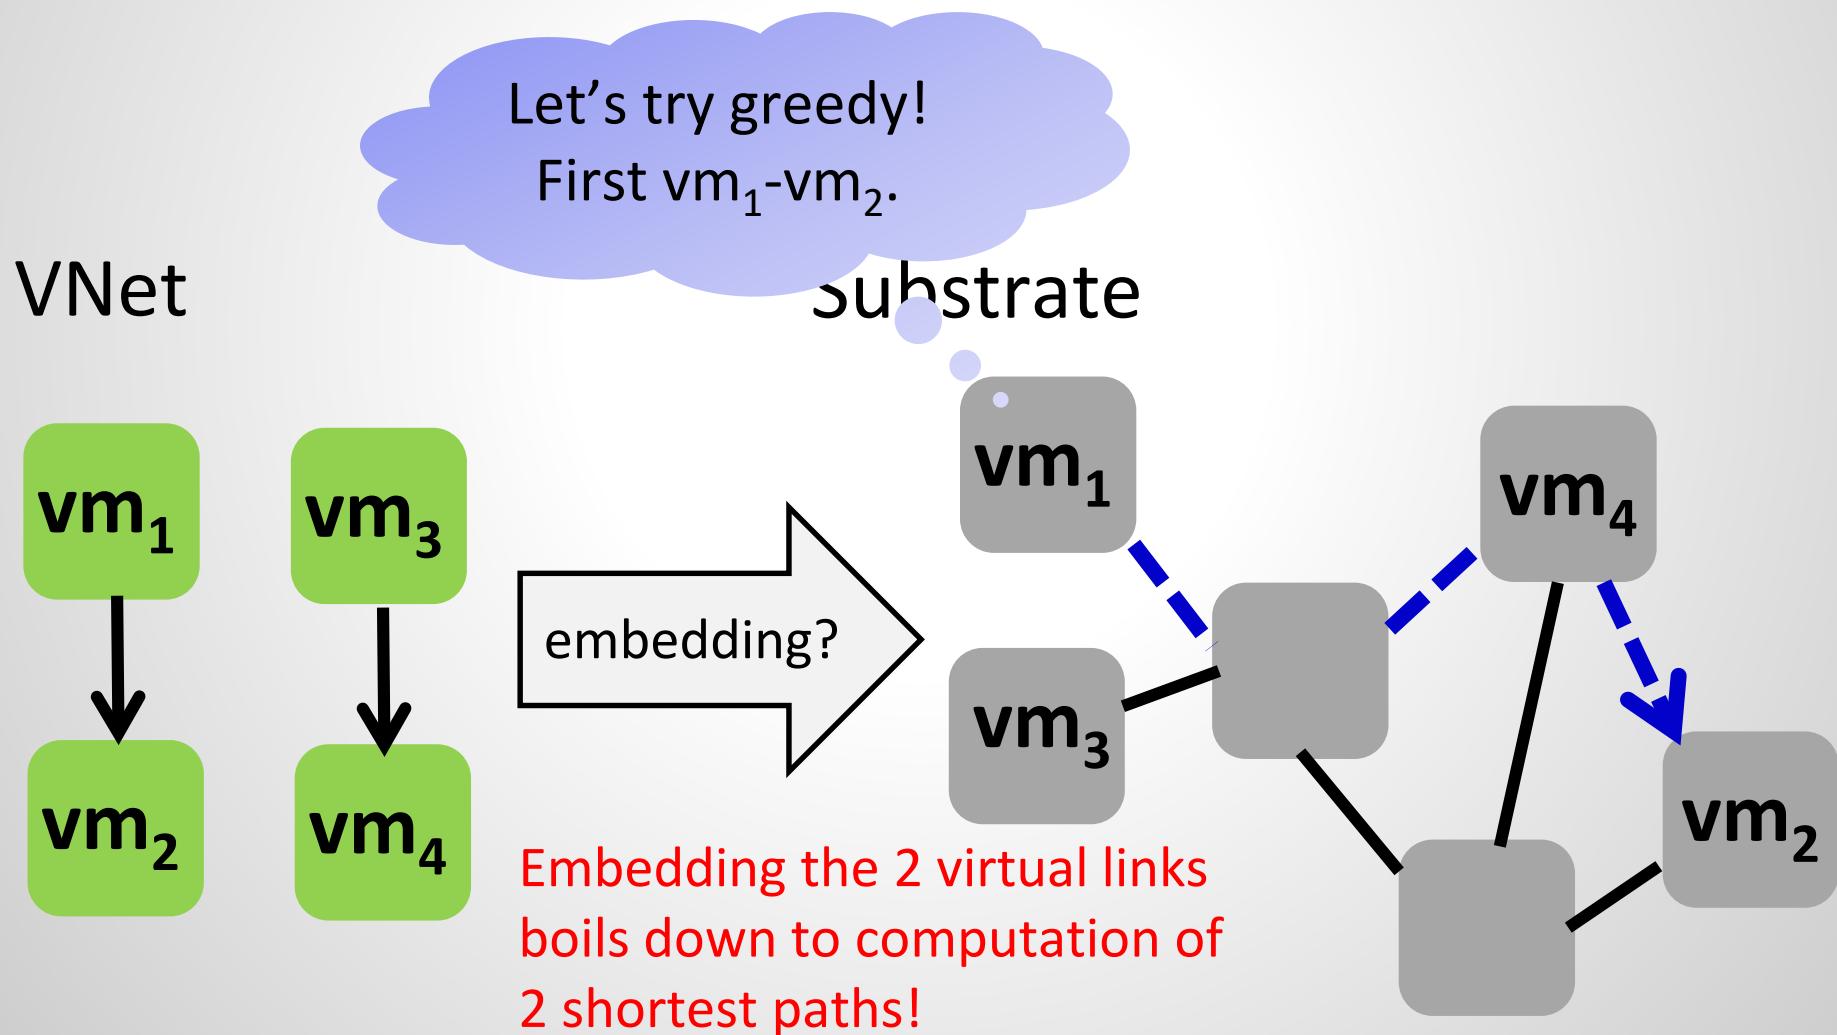


How to compute 2 shortest paths under capacity constraints?

Embedding the 2 virtual links  
boils down to computation of  
2 shortest paths!

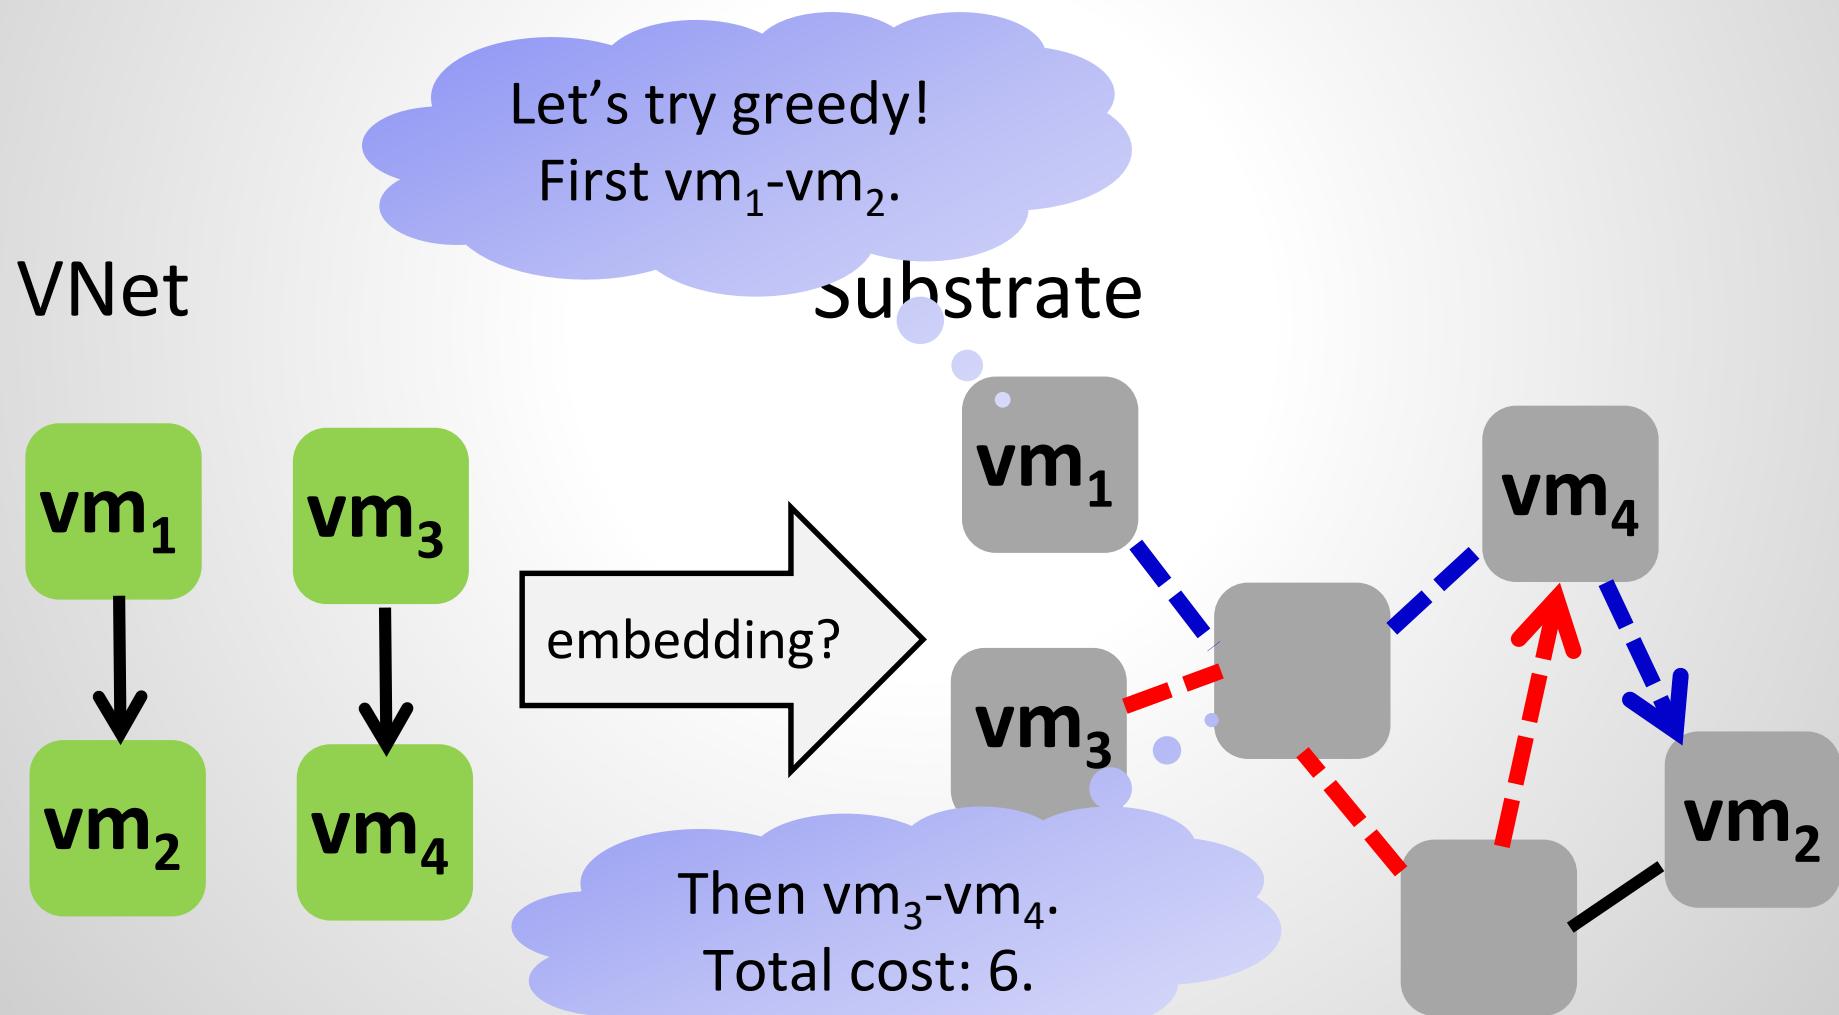
# The Virtual Network Embedding Problem

- Let's start simple: assume node **mappings** are given



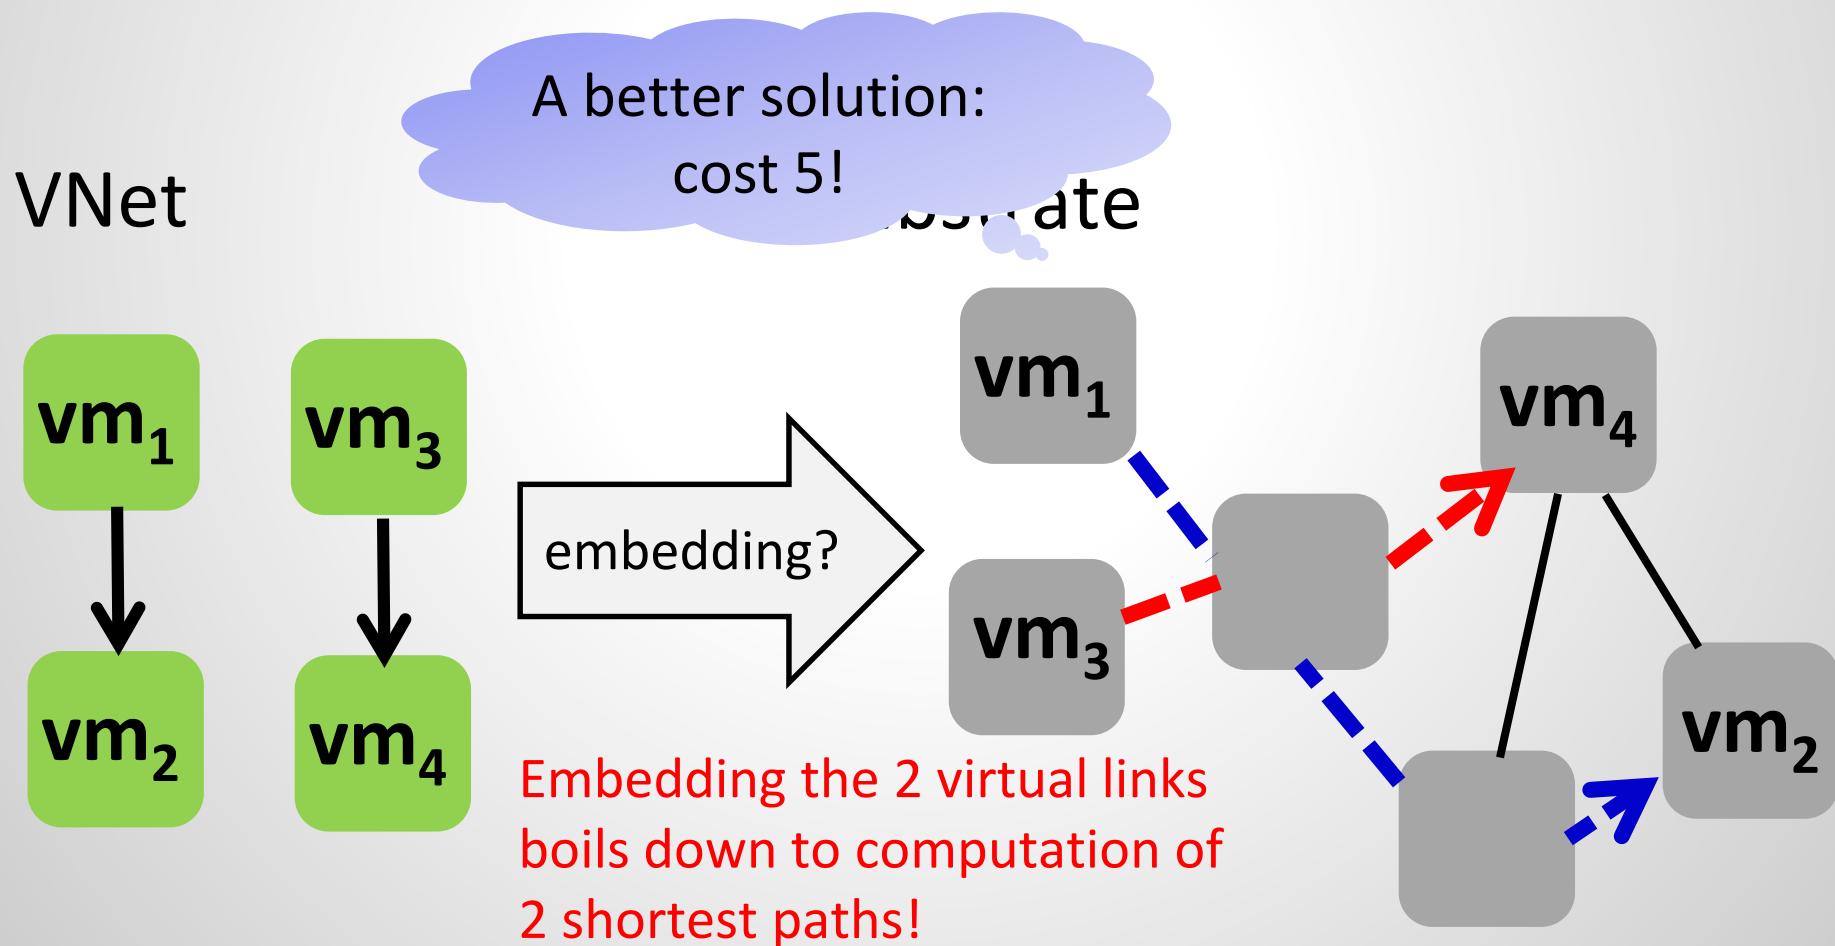
# The Virtual Network Embedding Problem

- Let's start simple: assume node **mappings** are given



# The Virtual Network Embedding Problem

- Let's start simple: assume node **mappings** are given



# The Virtual Network Embedding Problem

- Let's start simple: assume node **mappings** are given

VNet

Substrate

Joint optimization of 2 flows is already a challenging combinatorial problem! If demand=capacity=1:  
shortest 2-disjoint paths problem.

**vm<sub>2</sub>**

**vm<sub>4</sub>**

**vm<sub>1</sub>**

**vm<sub>3</sub>**

**vm<sub>4</sub>**

**vm<sub>2</sub>**

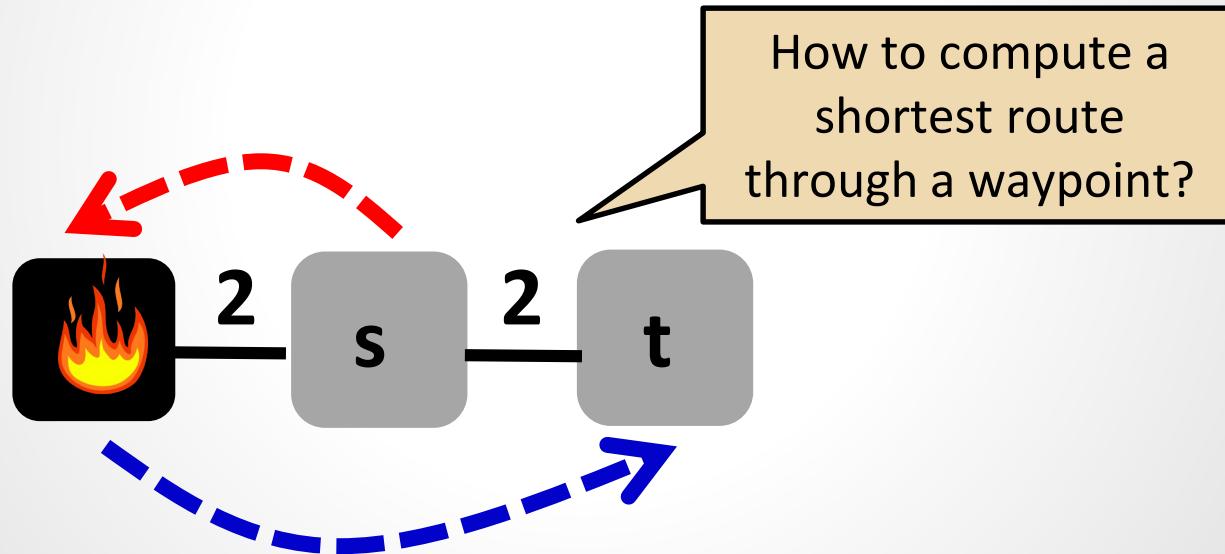
Embedding the 2 virtual links  
boils down to computation of  
2 shortest paths!

# Therefore: Mapping Virtual Links is Challenging

Bad news: The Virtual Network Embedding Problem is hard  
**even if endpoints are already mapped** and given.

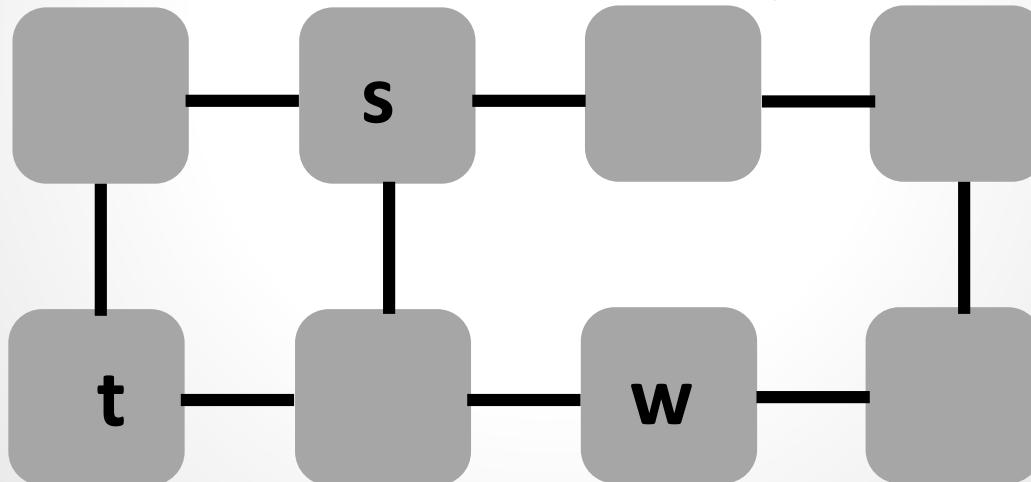
# Remark: Also Hard to Route 1 Waypoint!

Steering traffic through a single network function / **middlebox**:  
**a walk**



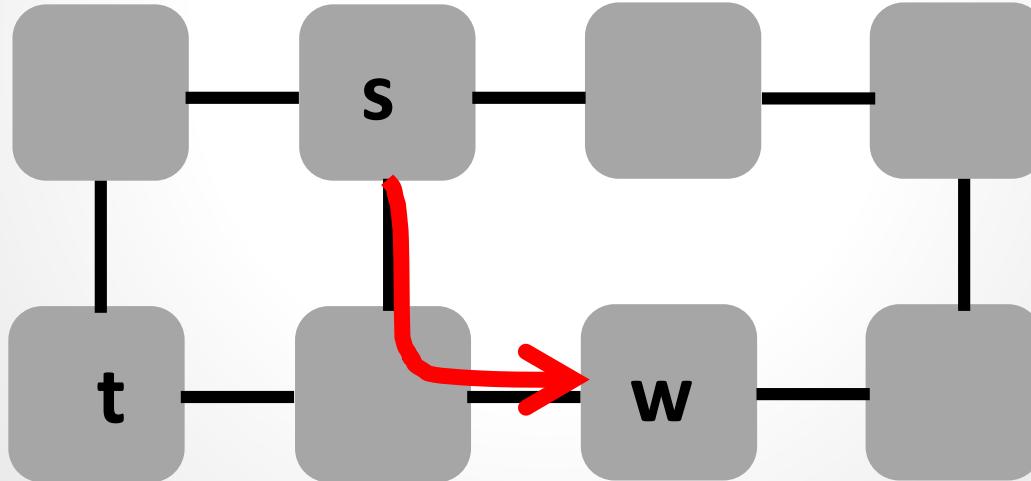
# Comuting A Shortest Walk Through A *Single* Given Waypoint is Non-Trivial!

Assume unit capacity and  
demand for simplicity!



# Comuting A Shortest Walk Through A *Single* Given Waypoint is Non-Trivial!

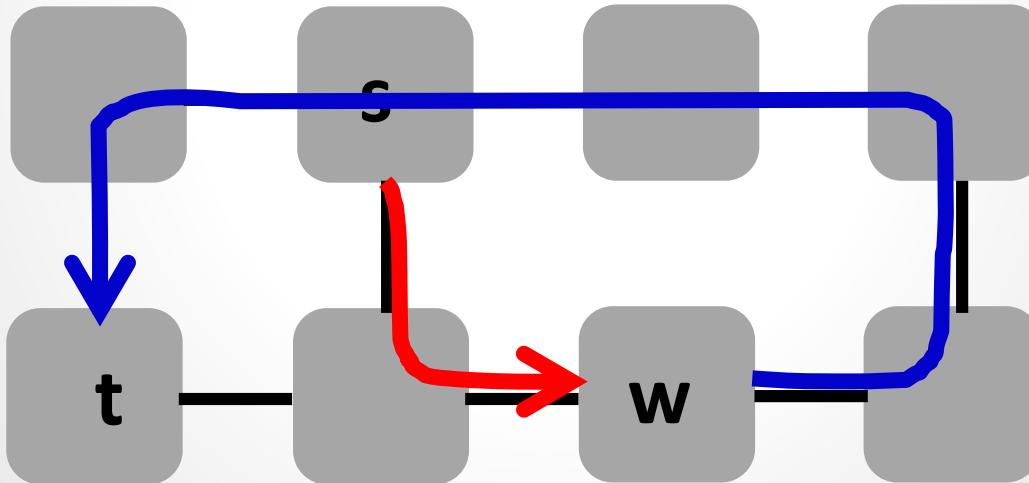
Assume unit capacity and  
demand for simplicity!



*Greedy fails:* choose shortest path from  $s$  to  $w$ ...

# Comuting A Shortest Walk Through A *Single* Given Waypoint is Non-Trivial!

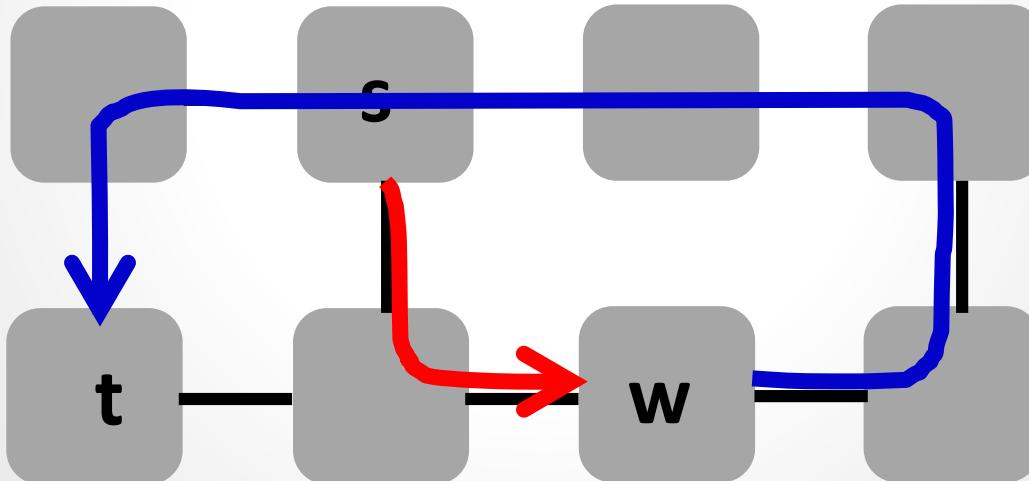
Assume unit capacity and  
demand for simplicity!



*Greedy fails:* ... now need long path from  $w$  to  $t$

# Comuting A Shortest Walk Through A *Single* Given Waypoint is Non-Trivial!

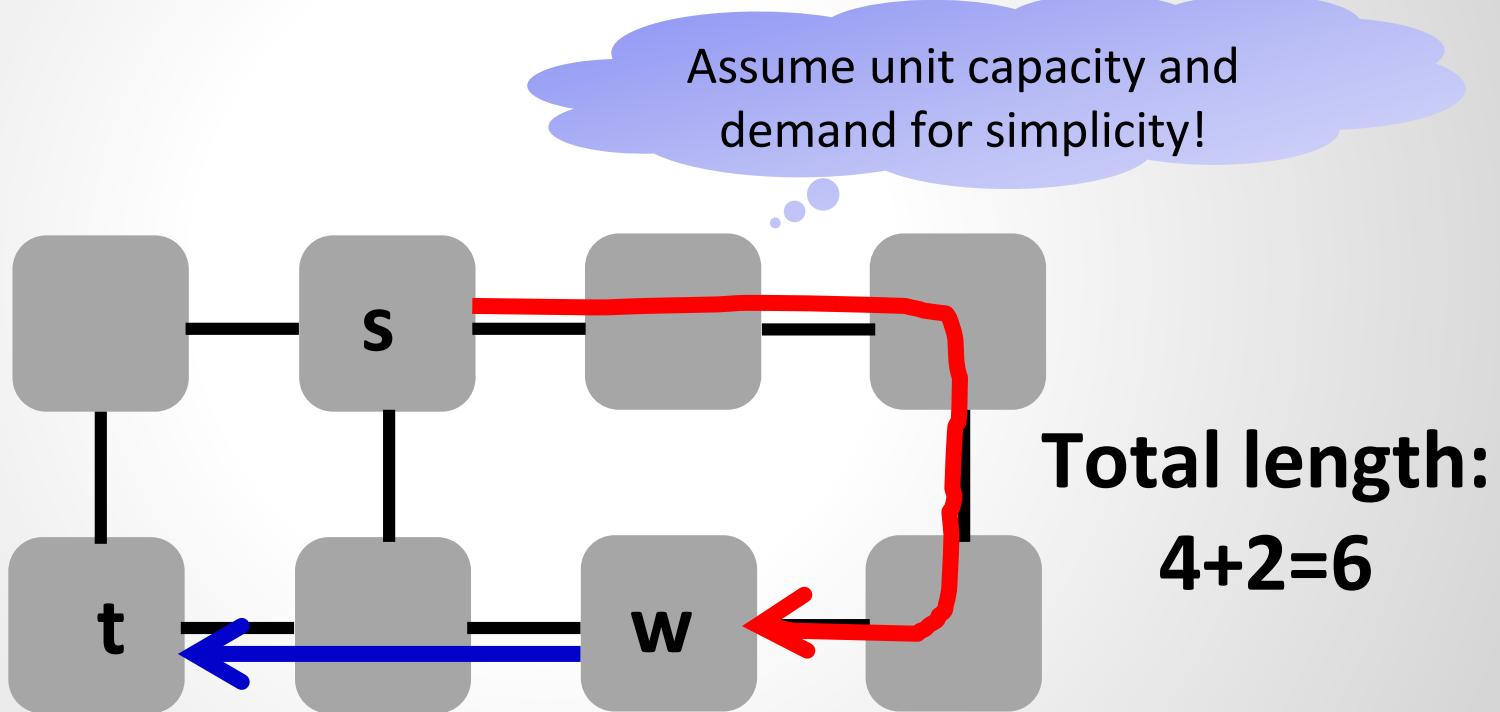
Assume unit capacity and demand for simplicity!



Total length:  
 $2+6=8$

*Greedy fails:* ... now need long path from w to t

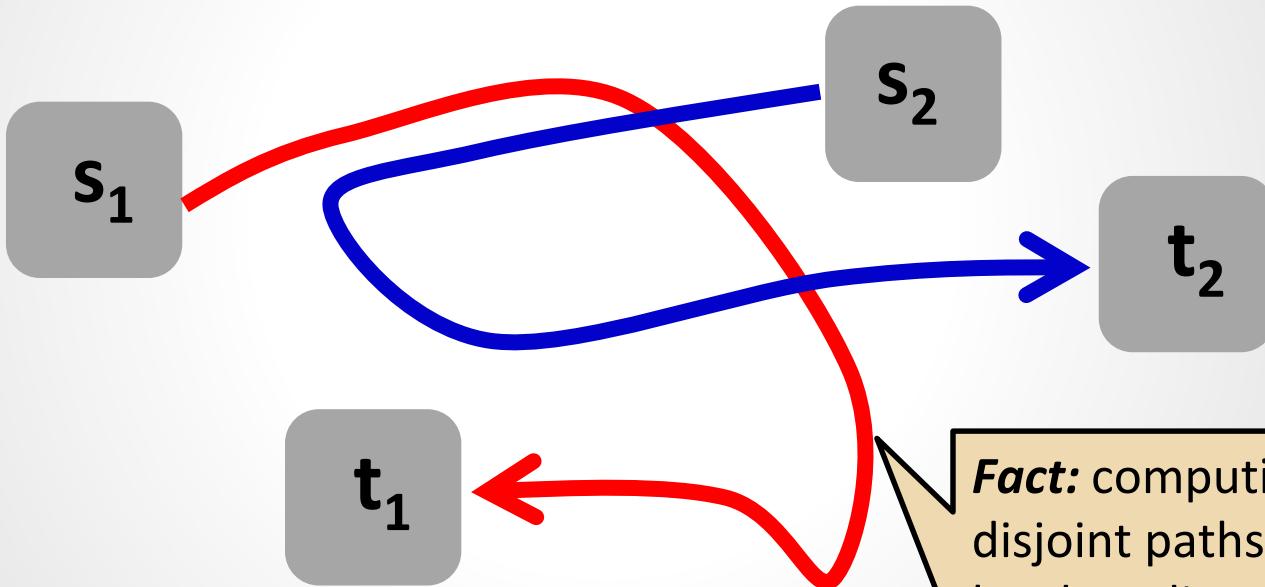
# Comuting A Shortest Walk Through A *Single* Given Waypoint is Non-Trivial!



A *better solution*: jointly optimize the two segments!

# NP-hard on *Directed Networks*: Reduction from Disjoint Paths Problem

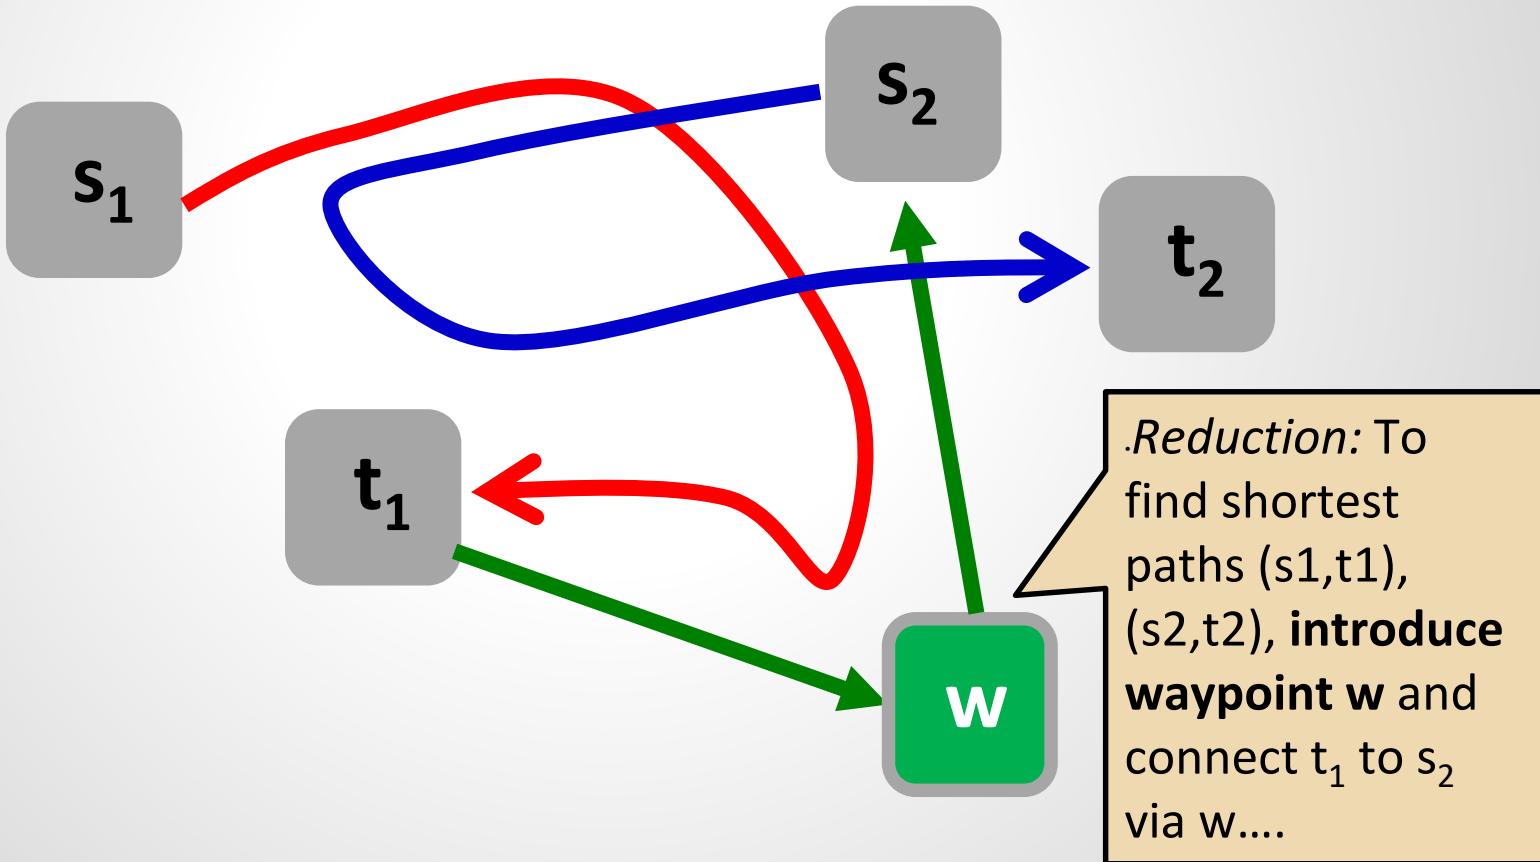
**Reduction:** From joint shortest paths  $(s_1, t_1), (s_2, t_2)$   
to shortest walk  $(s, w, t)$  problem



**Fact:** computing 2-disjoint paths (2DP) is NP-hard on directed graphs.  
**We show:** If waypoint routing was in P, we could solve 2DP fast.  
**Contradiction!**

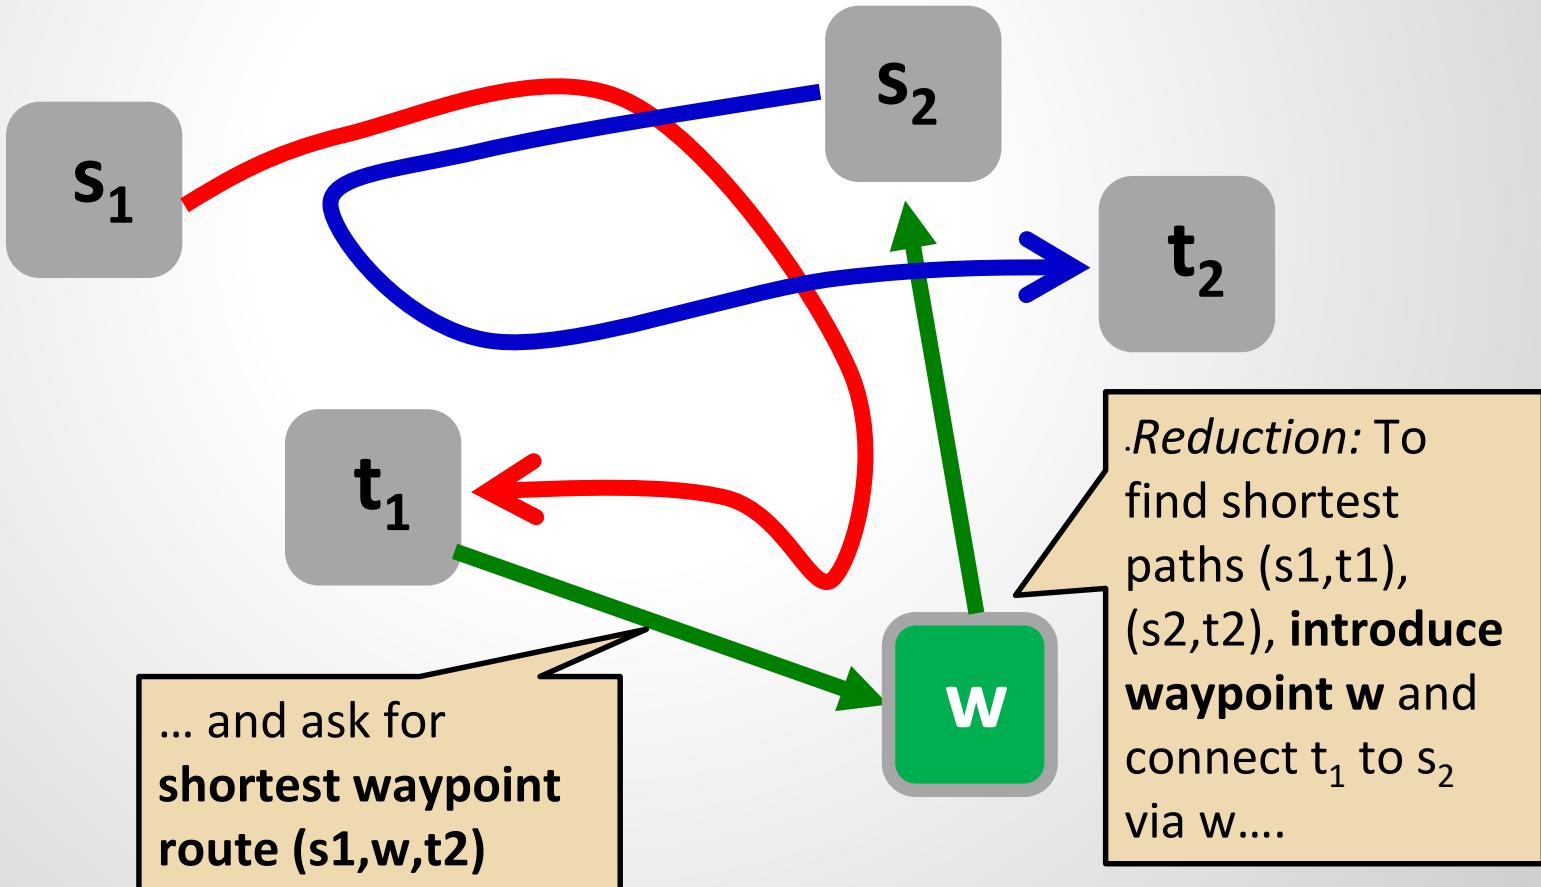
# NP-hard on *Directed Networks*: Reduction from Disjoint Paths Problem

**Reduction:** From joint shortest paths  $(s_1, t_1), (s_2, t_2)$   
to shortest walk  $(s, w, t)$  problem



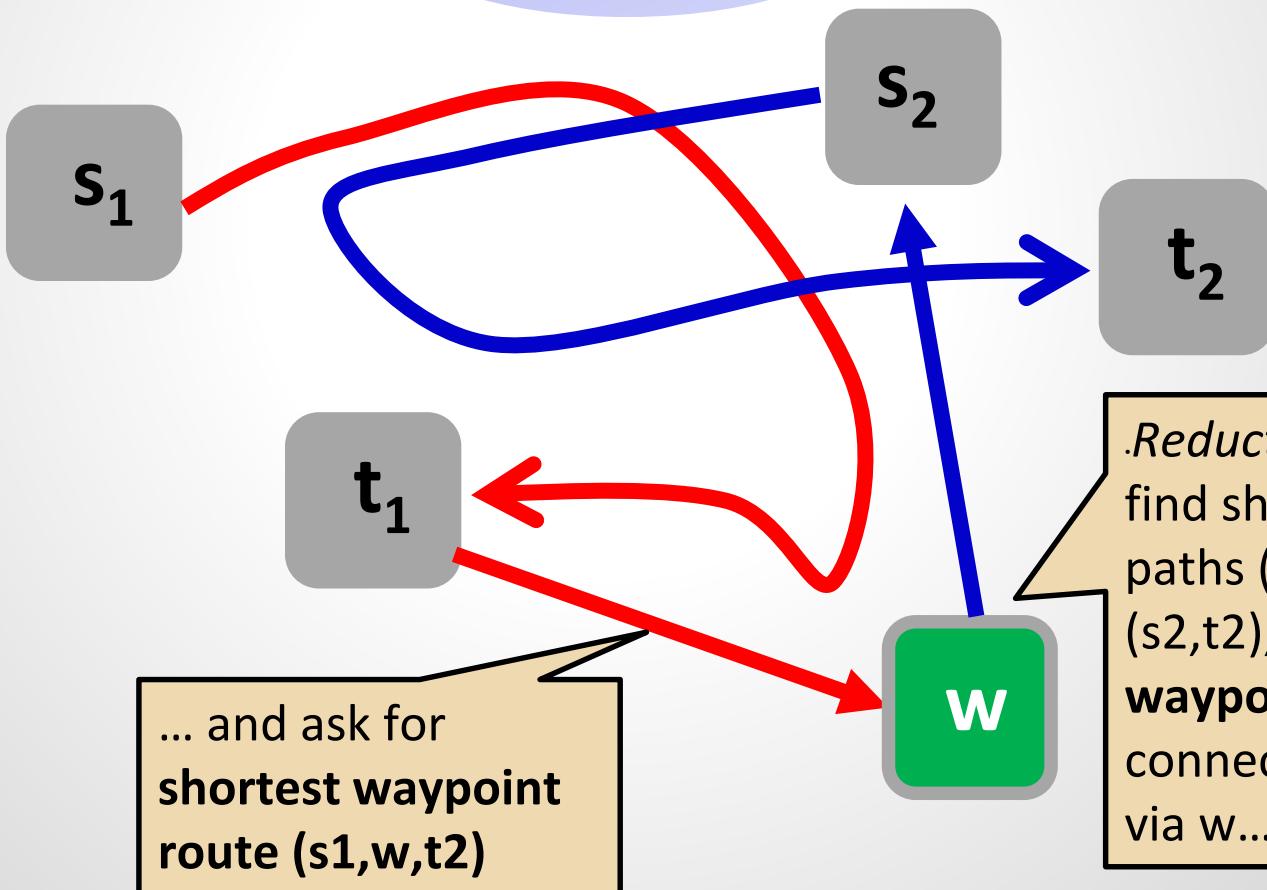
# NP-hard on *Directed Networks*: Reduction from Disjoint Paths Problem

Reduction: From joint shortest paths  $(s_1, t_1), (s_2, t_2)$   
to shortest walk  $(s, w, t)$  problem

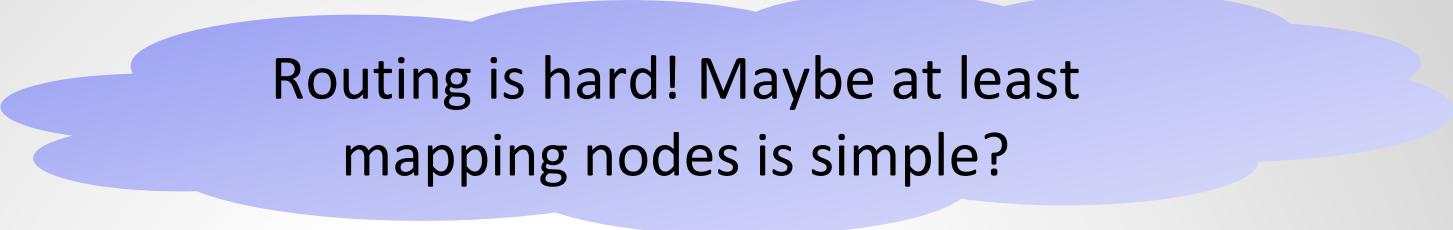


# NP-hard on Directed Networks

F The walk  $(s_1, w, t_2)$  walk defines a  $(s_1, t_1)$  and a  $(s_2, t_2)$  path pair before/after the waypoint! Solves original problem:  
**Reduction to shortest w.** **Contradiction!**



# Mapping Virtual Nodes

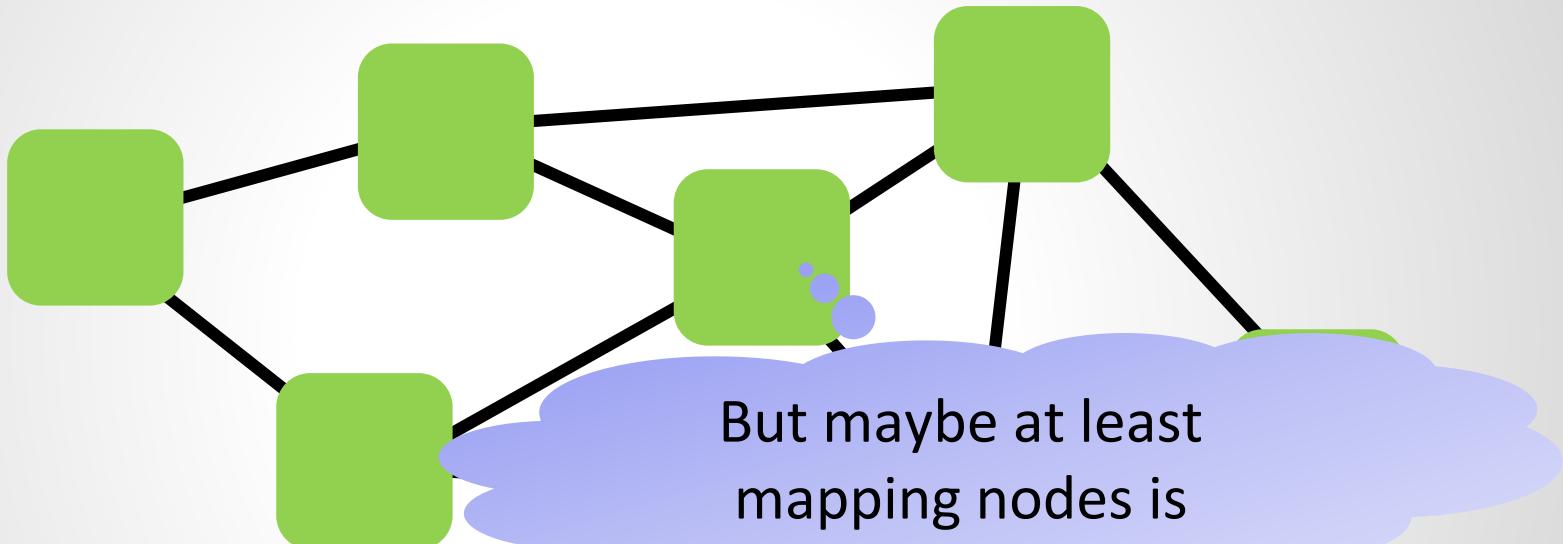


Routing is hard! Maybe at least  
mapping nodes is simple?

# Mapping Virtual Nodes

- Let's start simple again: assume paths are trivial, e.g., the physical network (host graph) **is a line**

Guest

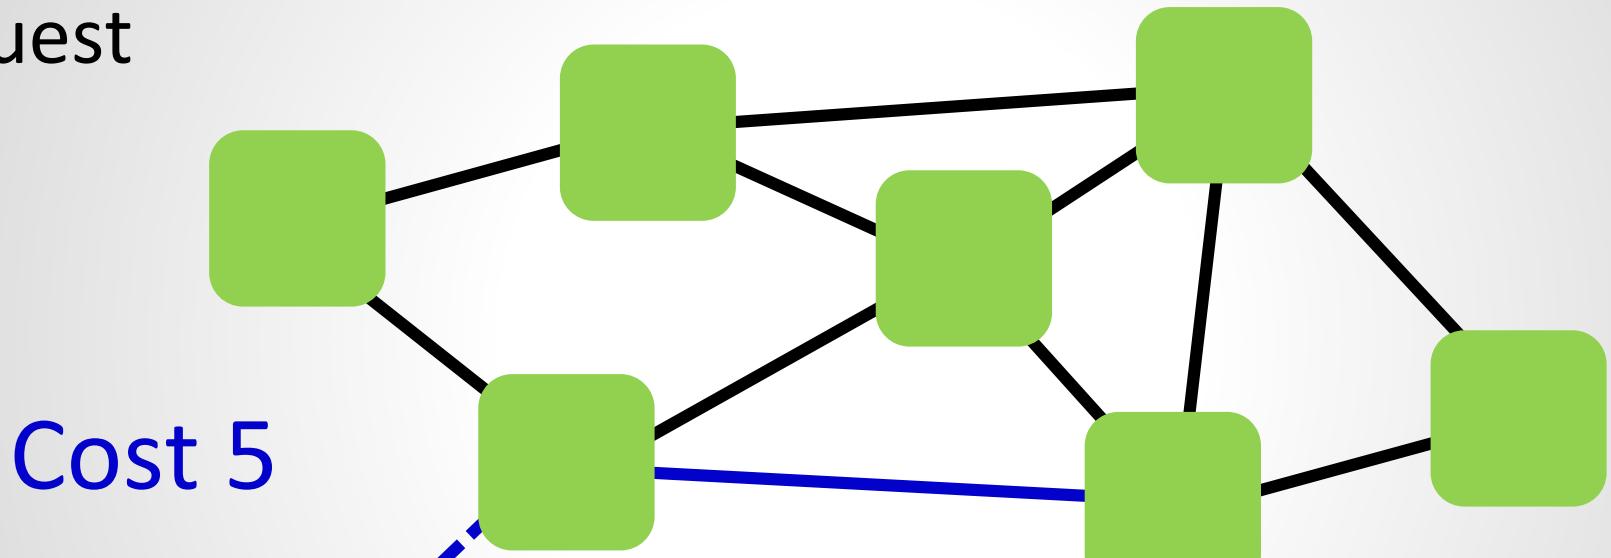


Host

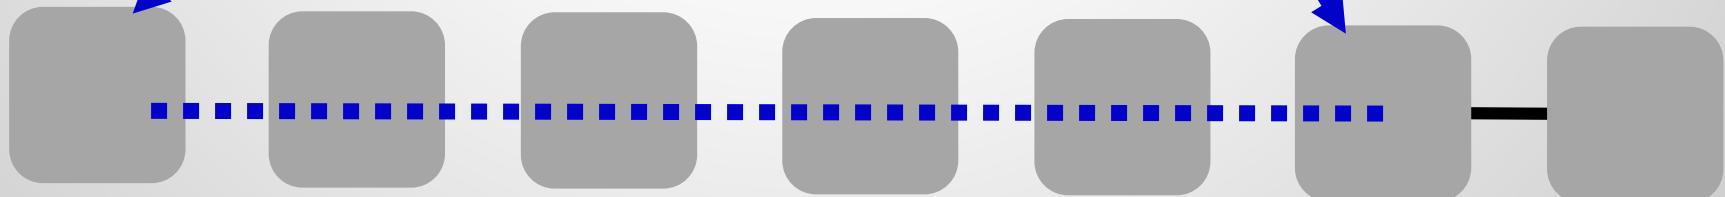
# Mapping Virtual Nodes

- Let's start simple again: assume **paths are trivial**, e.g., the physical network (host graph) **is a line**

Guest



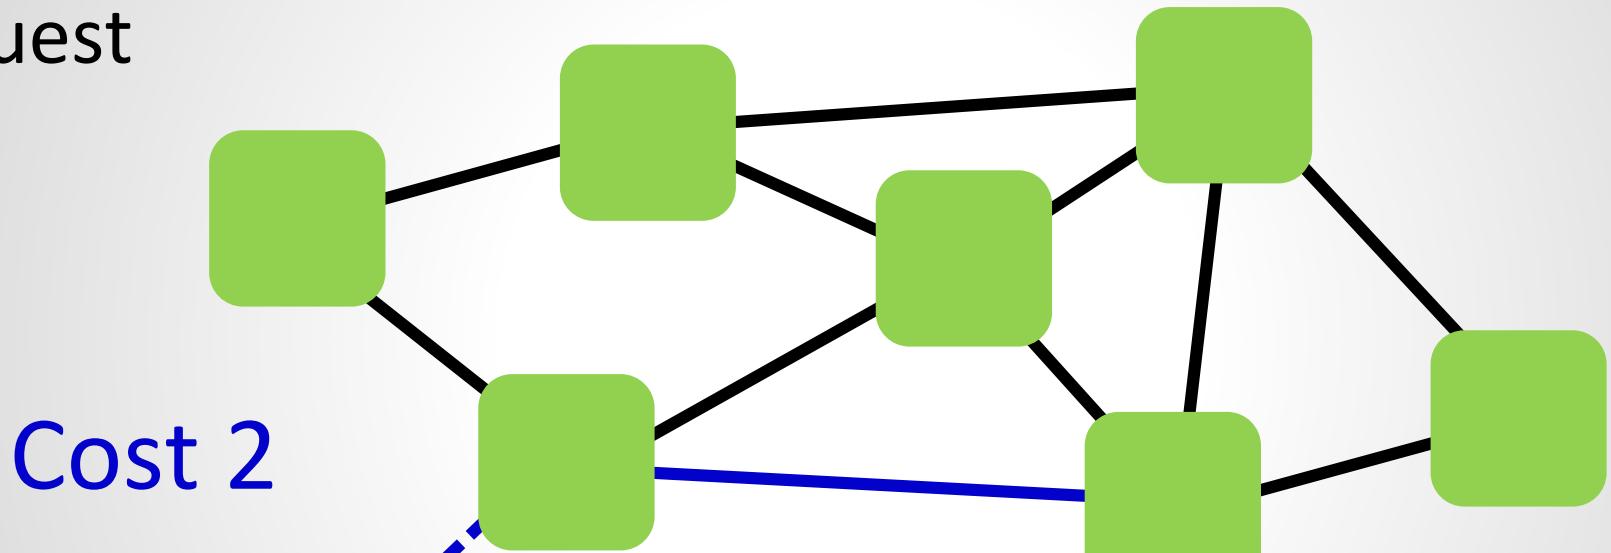
Host



# Mapping Virtual Nodes

- Let's start simple again: assume **paths are trivial**, e.g., the physical network (host graph) is a line

Guest

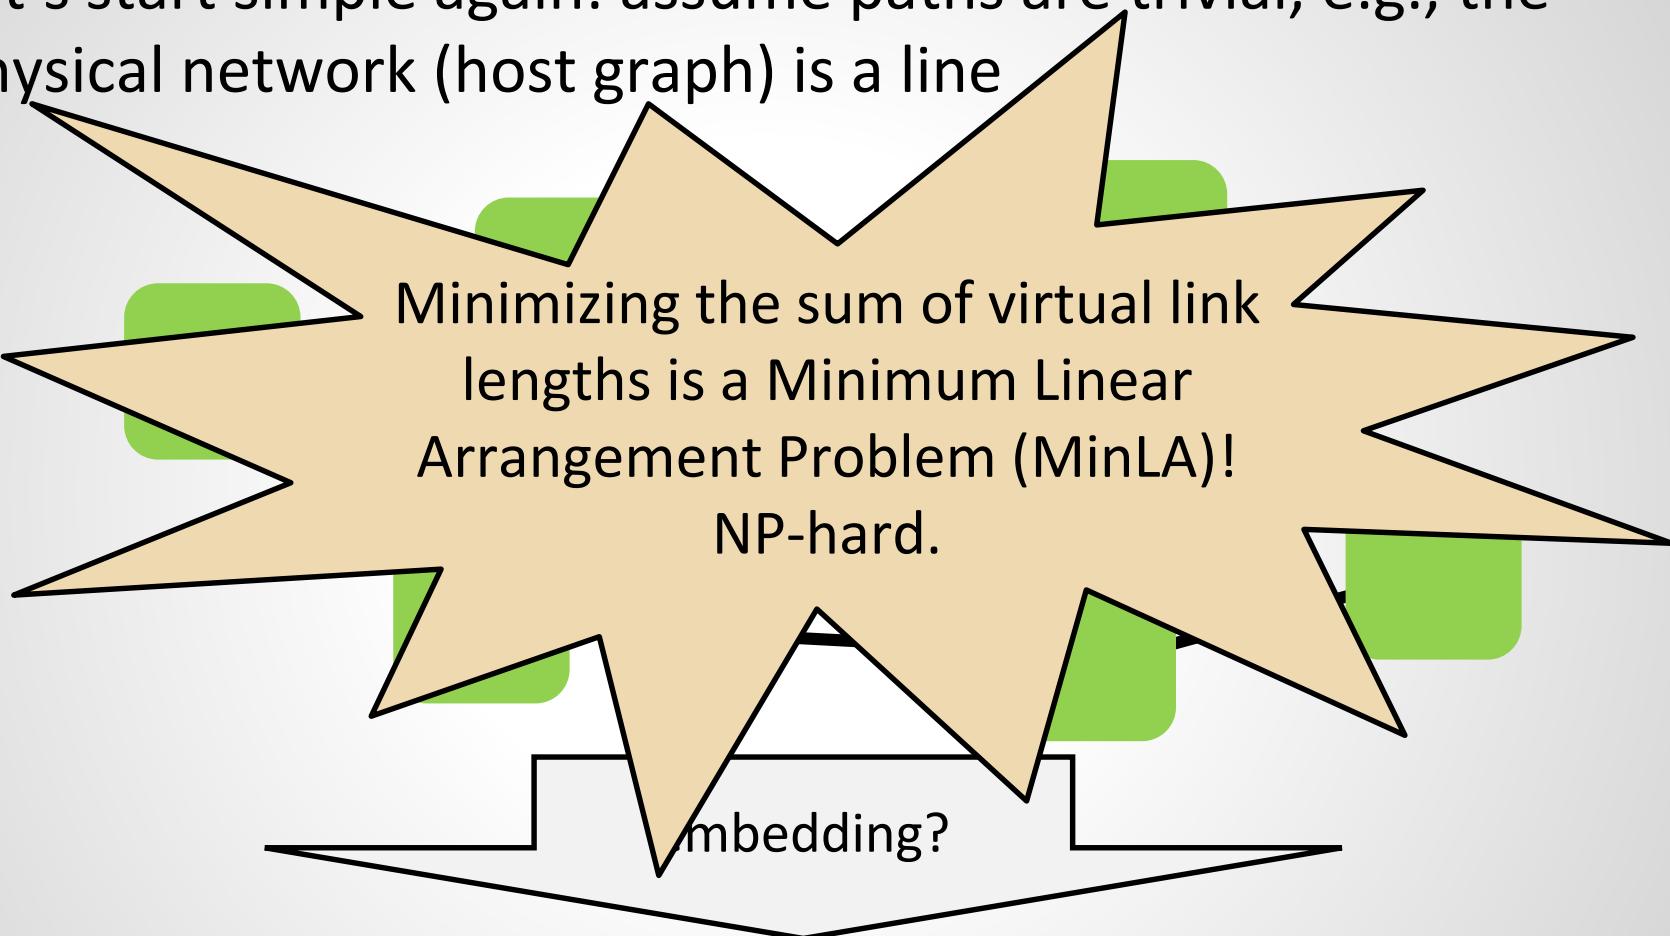


Host



# Mapping Virtual Nodes

- Let's start simple again: assume paths are trivial, e.g., the physical network (host graph) is a line



Minimizing the sum of virtual link lengths is a Minimum Linear Arrangement Problem (MinLA)!  
NP-hard.

embedding?



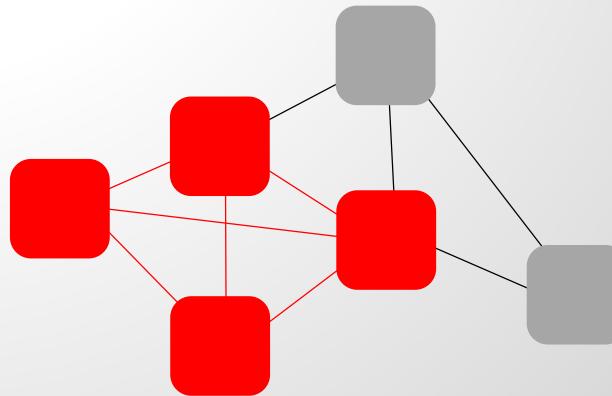
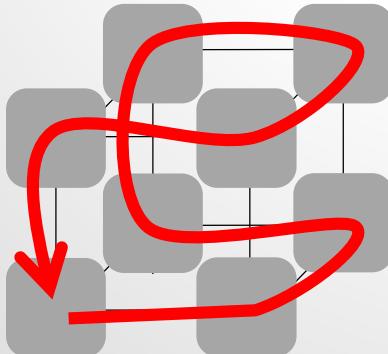
# Therefore: VNEP is Hard “in Both Dimensions”!

- We have seen examples that:
  - mapping virtual links is hard (even if nodes are given)
  - mapping virtual nodes is hard (even if links are trivial)
- Remark: the VNEP can also be seen as a generalization of the **Subgraph Isomorphism Problem (SIP)**

Known? Why is SIP NP-hard?

# Therefore: VNEP is Hard “in Both Dimensions”!

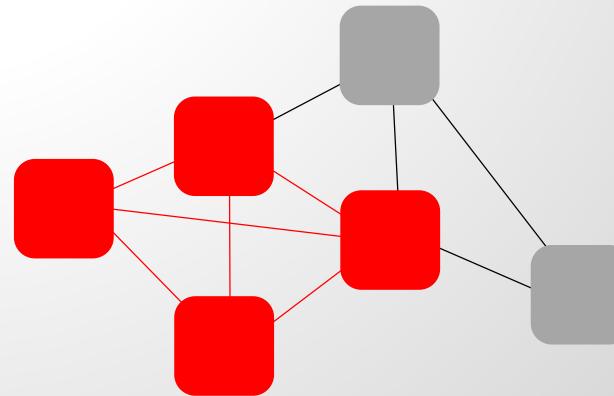
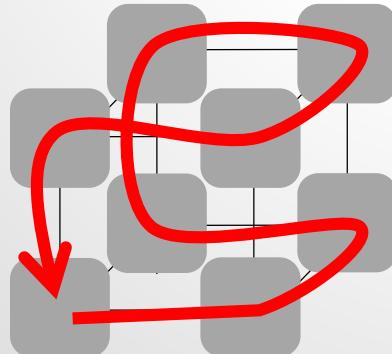
- We have seen examples that:
  - mapping virtual links is hard (even if nodes are given)
  - mapping virtual nodes is hard (even if links are trivial)
- Remark: the VNEP can also be seen as a generalization of the **Subgraph Isomorphism Problem (SIP)**
  - The SIP problem: Given two graphs  $G, H$ , determine whether  $G$  contains a subgraph that is **isomorphic to  $H$** ?
  - NP-hard: “does  $G$  contain an  **$n$ -node cycle**?” is a **Hamilton cycle** problem (each node visited exactly once), a solution to “does  $G$  contain a  **$k$ -clique**?” solves **maximum clique** problem, etc.



# Therefore: VNEP is Hard “in Both Dimensions”!

- We have seen examples that:
  - mapping virtual links is hard (even if nodes are given)
  - mapping virtual nodes is hard (even if links are trivial)
- Remark: the VNEP can also be seen as a generalization of the **Subgraph Isomorphism Problem (SIP)**
  - The SIP problem: Given two graphs  $G, H$ , determine whether  $G$  contains a subgraph that is **isomorphic to  $H$** ?
  - A “Hamilton cycle” problem (each node visited “”) solves **maximum clique** problem, etc.

So if SIP is hard, why is  
VNEP hard?



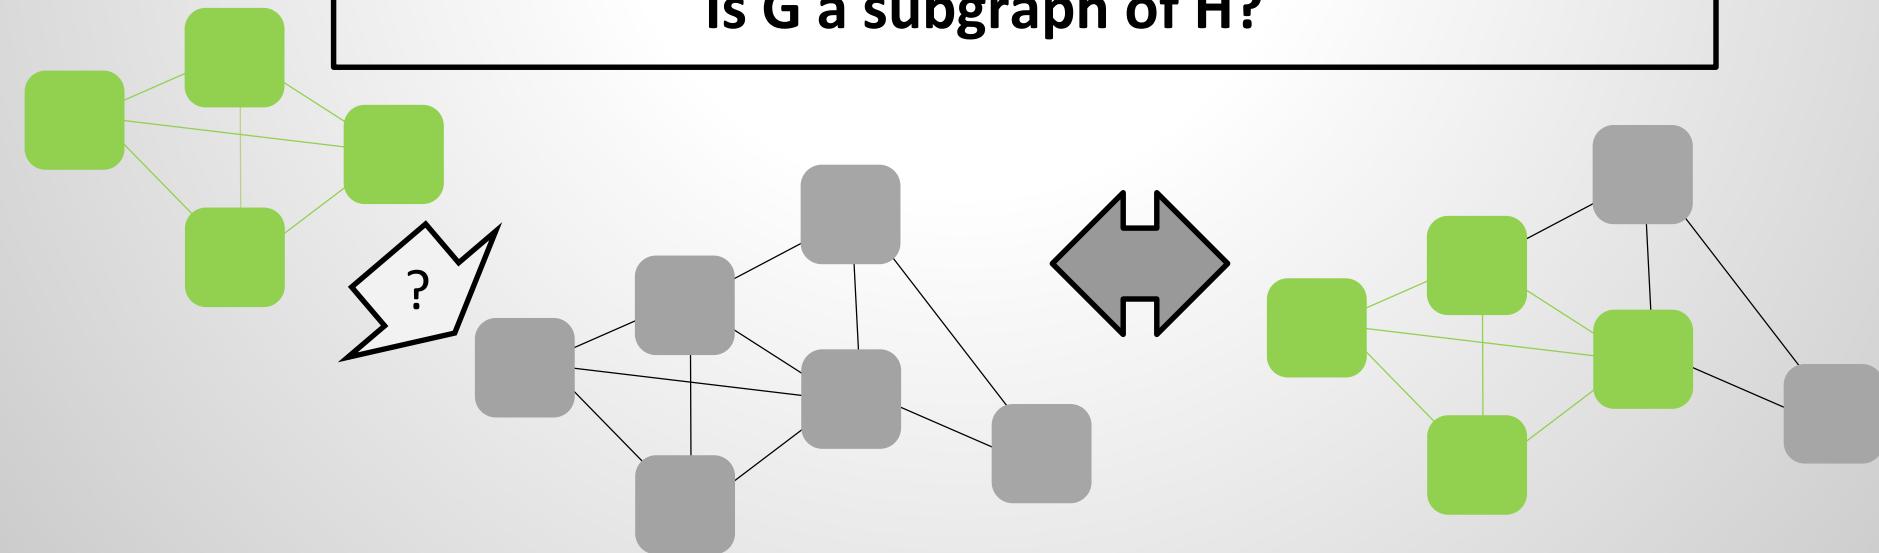
# NP-Hardness: From SIP to VNEP

- Observe: VNEP is a **generalization** of SIP
- For example:

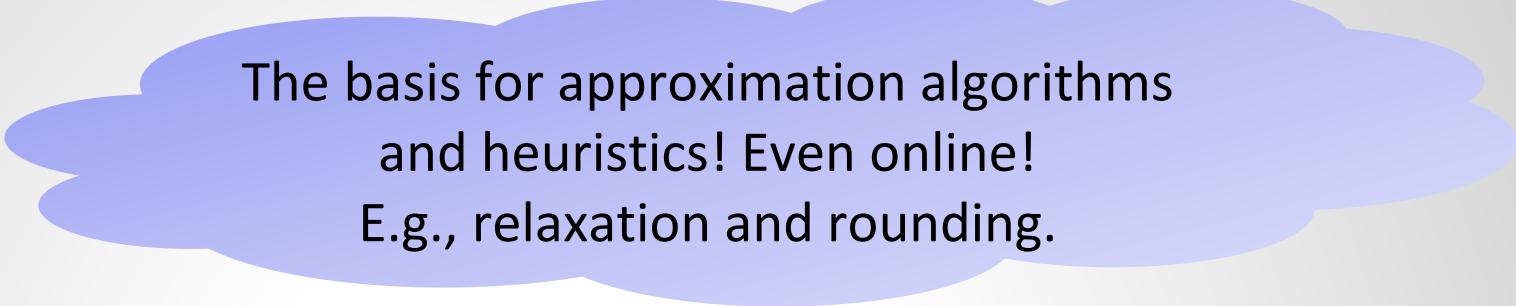
Can VNet  $G=(V,E)$  be embedded in  $H$  at cost  $|E|$ ?  
(I.e., each virtual edge has **length 1**.)



Is  $G$  a subgraph of  $H$ ?



# Can we at least formulate a “fast” MIP?



The basis for approximation algorithms  
and heuristics! Even online!  
E.g., relaxation and rounding.

# Can we at least formulate a “fast” MIP?



?

# Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
  - ❑ Linear objective function (e.g., minimize embedding footprint)
  - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

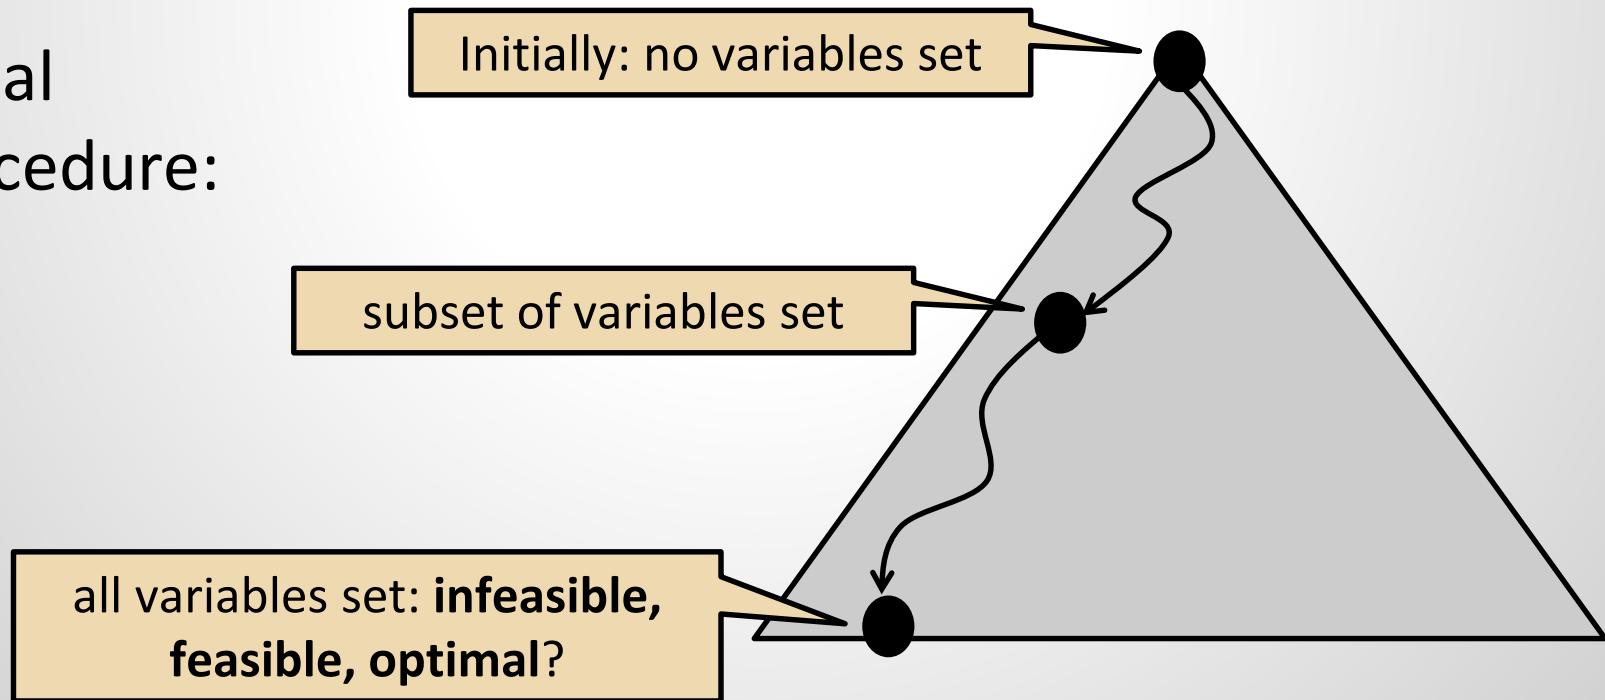
# Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
  - ❑ Linear objective function (e.g., minimize)
  - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ One that provides good relaxations!
- ❑ Solved, e.g., with branch-and-bound search tree

# Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
  - ❑ Linear objective function (e.g., minimize embedding footprint)
  - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound **search tree**

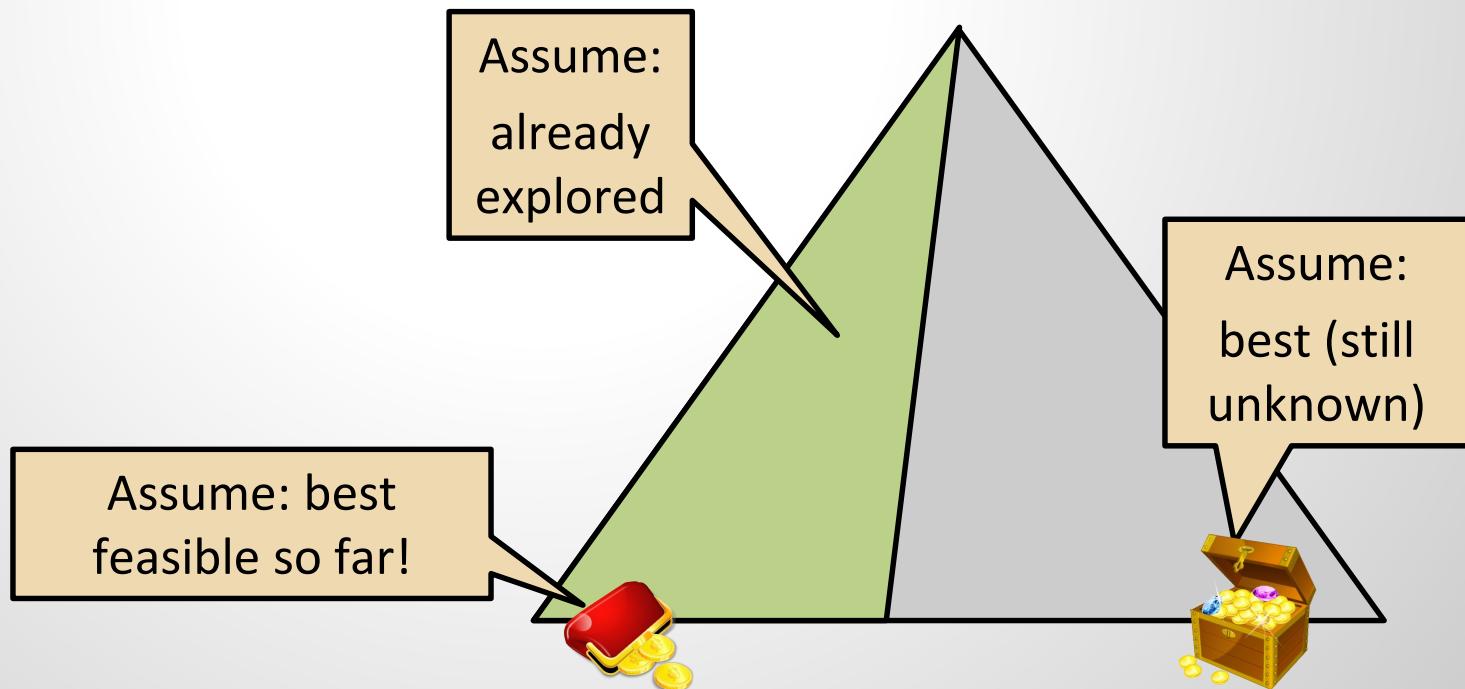
Usual  
procedure:



# Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
  - ❑ Linear objective function (e.g., minimize embedding footprint)
  - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

Usual  
procedure:

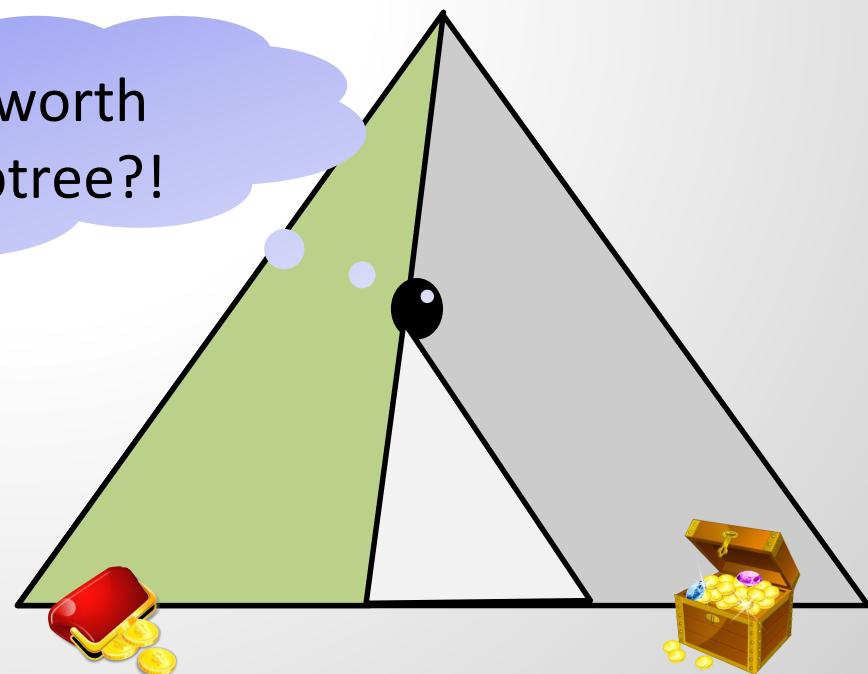


# Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
  - ❑ Linear objective function (e.g., minimize embedding footprint)
  - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

Usual  
procedure:

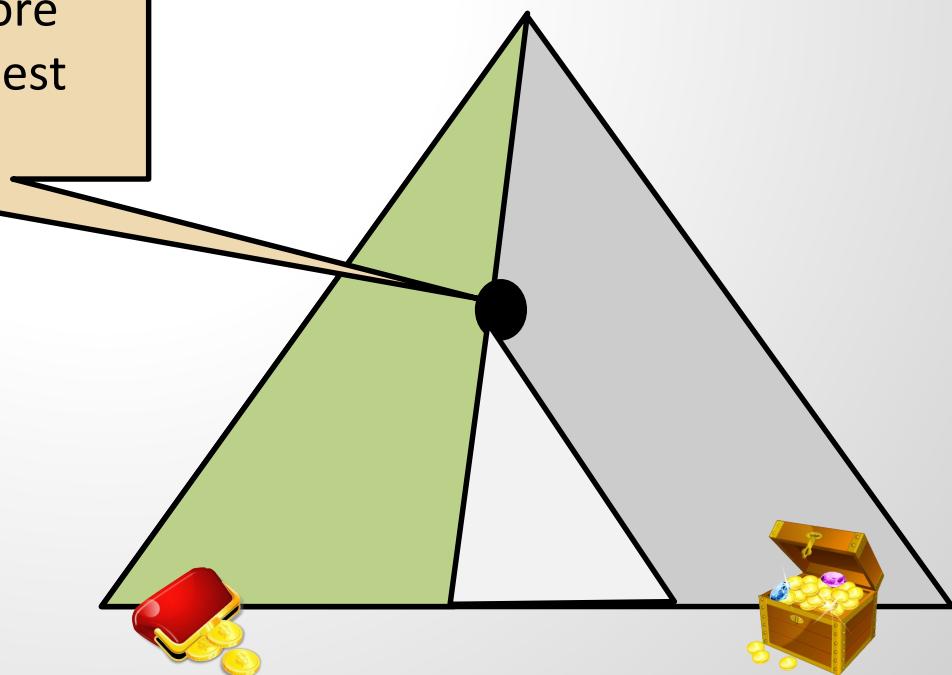
Decide: Is it worth  
exploring subtree?!



# Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
  - ❑ Linear objective function (e.g., minimize embedding footprint)
  - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

Usual trick: Relax! Solve LP (fast!),  
and if **relaxed solution** (more  
general!) **not better** then best  
solution so far: skip it!

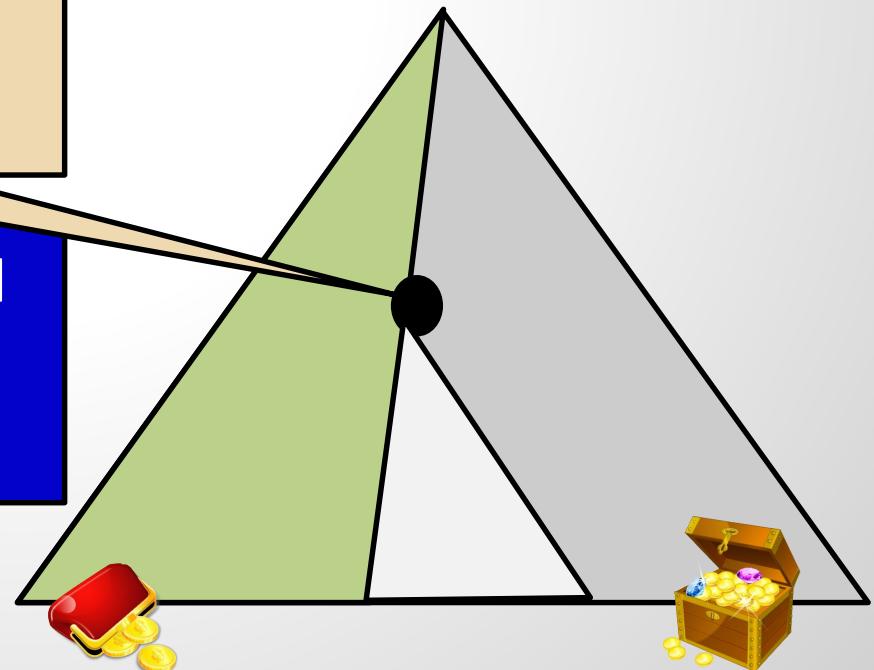


# Can we at least formulate a “fast” MIP?

- ❑ Recall: Mixed Integer Program (MIP)
  - ❑ Linear objective function (e.g., minimize embedding footprint)
  - ❑ Linear constraints (e.g., do not violate capacity constraints)
- ❑ Solved, e.g., with branch-and-bound search tree

Usual trick: Relax! Solve LP (fast!),  
and if **relaxed solution** (more  
general!) **not better** then best  
solution so far: skip it!

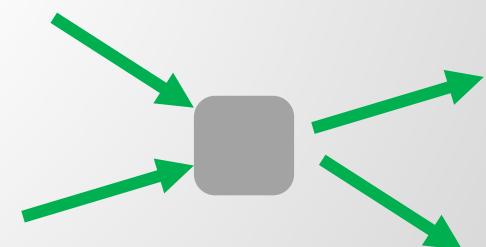
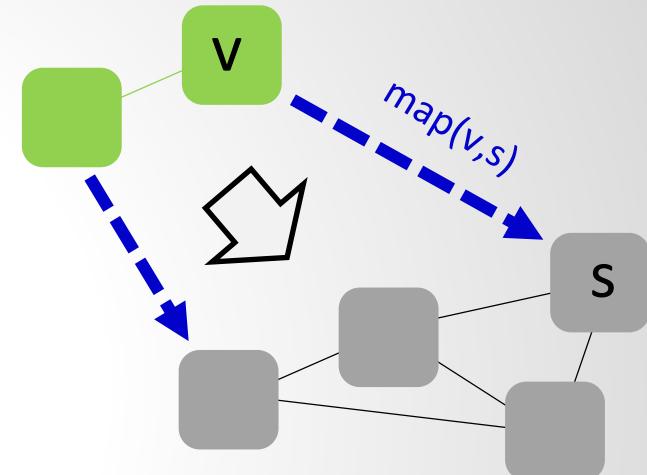
Bottomline: If MIP provides «good relaxations», large parts of the search space can be pruned.



# Can we at least formulate a “fast” MIP?

A typical MIP formulation:

- ❑ Introduce **binary variables**  $map(v,s)$  to map virtual nodes  $v$  to substrate node  $s$
- ❑ Introduce **flow variables** for paths (say **splittable** for now)
- ❑ Ensure **flow conservation**: all flow entering a node must leave the node, unless it is the source or the destination



$$\sum_{u \rightarrow v} f_{uv} = \sum_{v \rightarrow w} f_{vw}$$

In                      Out

# Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b requested from node s to node t.

In      Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

What does this formula do and why is it correct?

# Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b  
requested from node s  
to node t.

In      Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

If  $\text{map}(s,v)=1$ , i.e., s mapped to v:  
so **flow starts at v**, and hence  
outgoing flow must be larger than  
incoming flow (plus b).

# Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b requested from node s to node t.

In      Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

If  $\text{map}(s,v)=0$  and  $\text{map}(t,v)=0$ , i.e., **v is along the path** from s to t: then we have **flow conservation**: outgoing flow must equal incoming flow (here  $\geq$ , objective function will remove unnecessary flow).

# Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b  
requested from node s  
to node t.

In      Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

If  $\text{map}(t,v)=1$ , i.e., t mapped to v: so flow **terminates at node v**: so **no constraint**: minus infinity (but objective function will remove unnecessary flow).

# Can we at least formulate a “fast” MIP?

We get constraints like:

Assume bandwidth b requested from node s to node t.

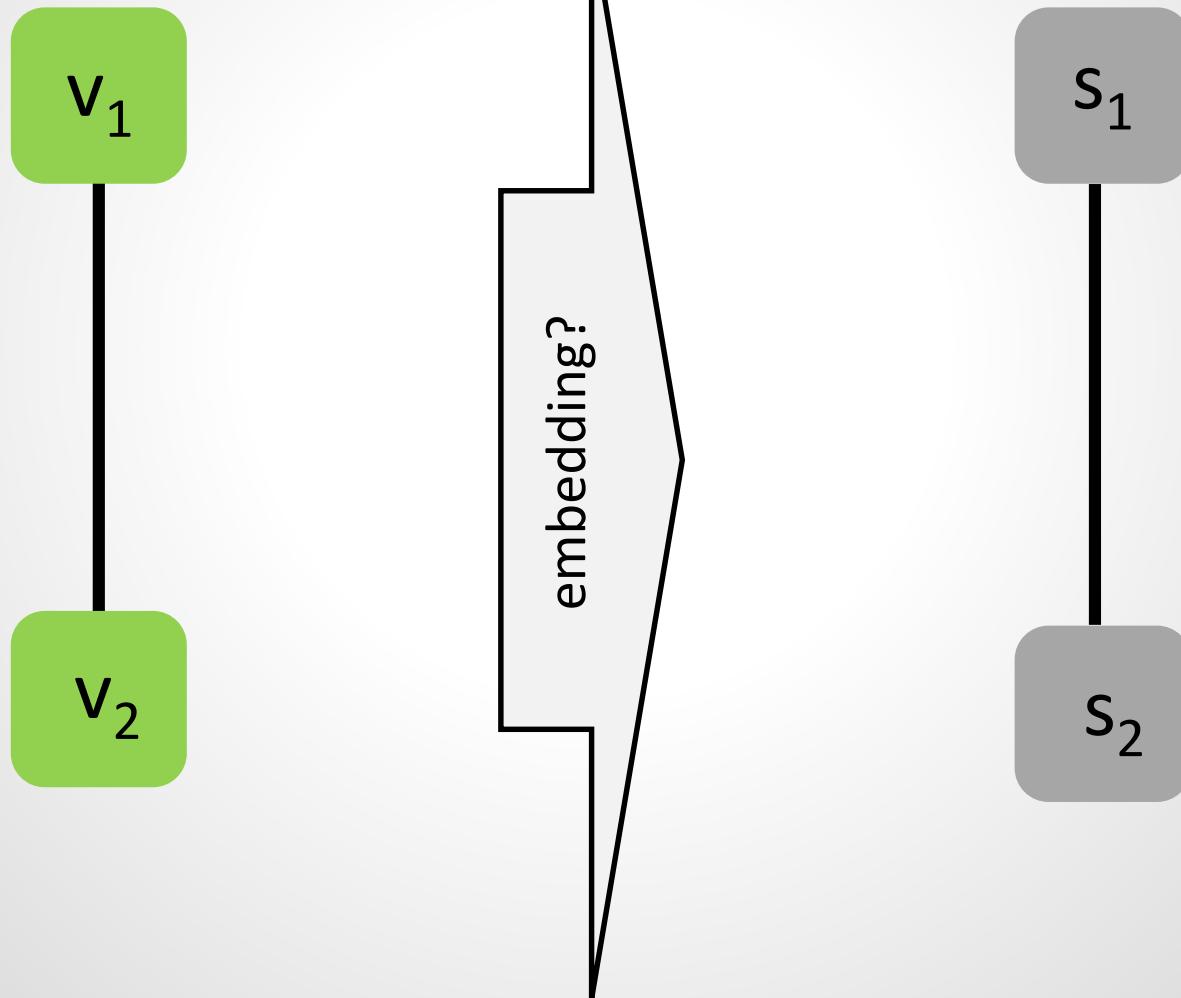
In      Out

$$\forall v: \sum_u f_{uv} - f_{vu} \geq \text{map}(s,v) * b - \text{map}(t,v) * \infty$$

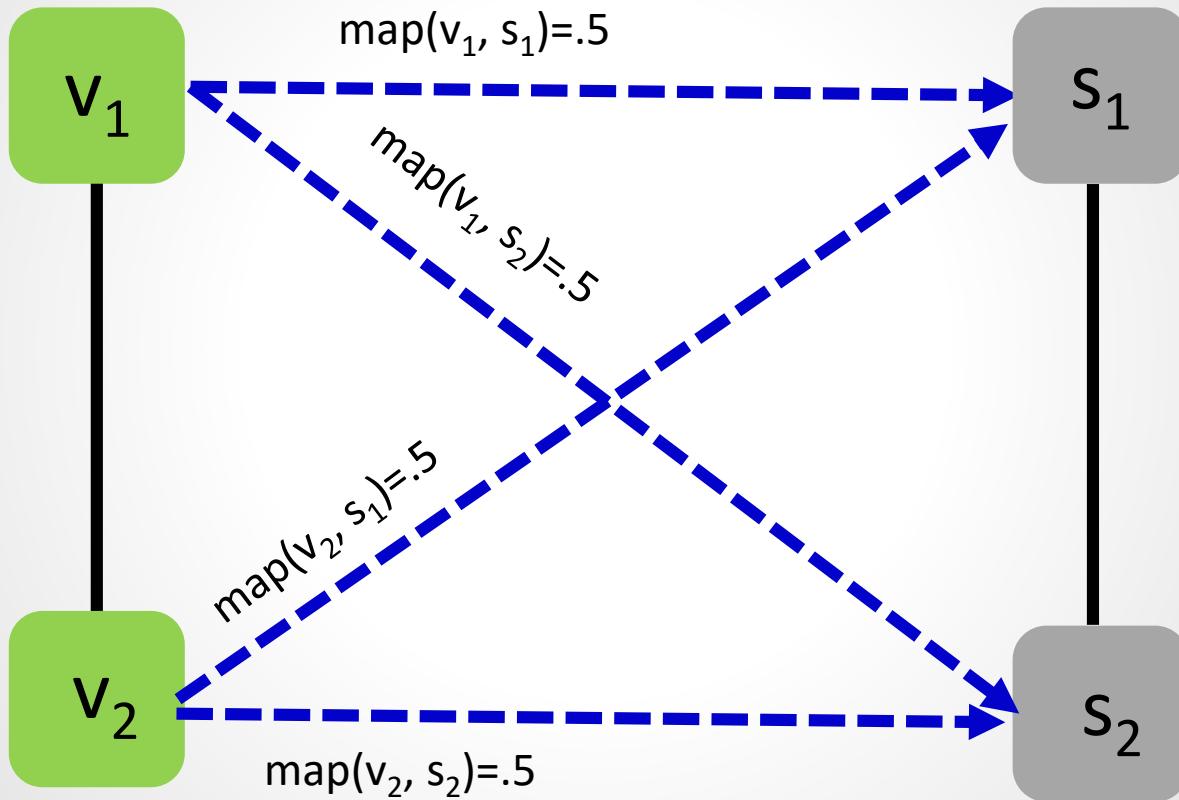
If  $\text{map}(t,v)=1$ , i.e., t mapped to v: so flow **terminates at node v**: so **no constraint**: minus infinity (but objective function will remove unnecessary flow).

Will such a MIP provide effective pruning?

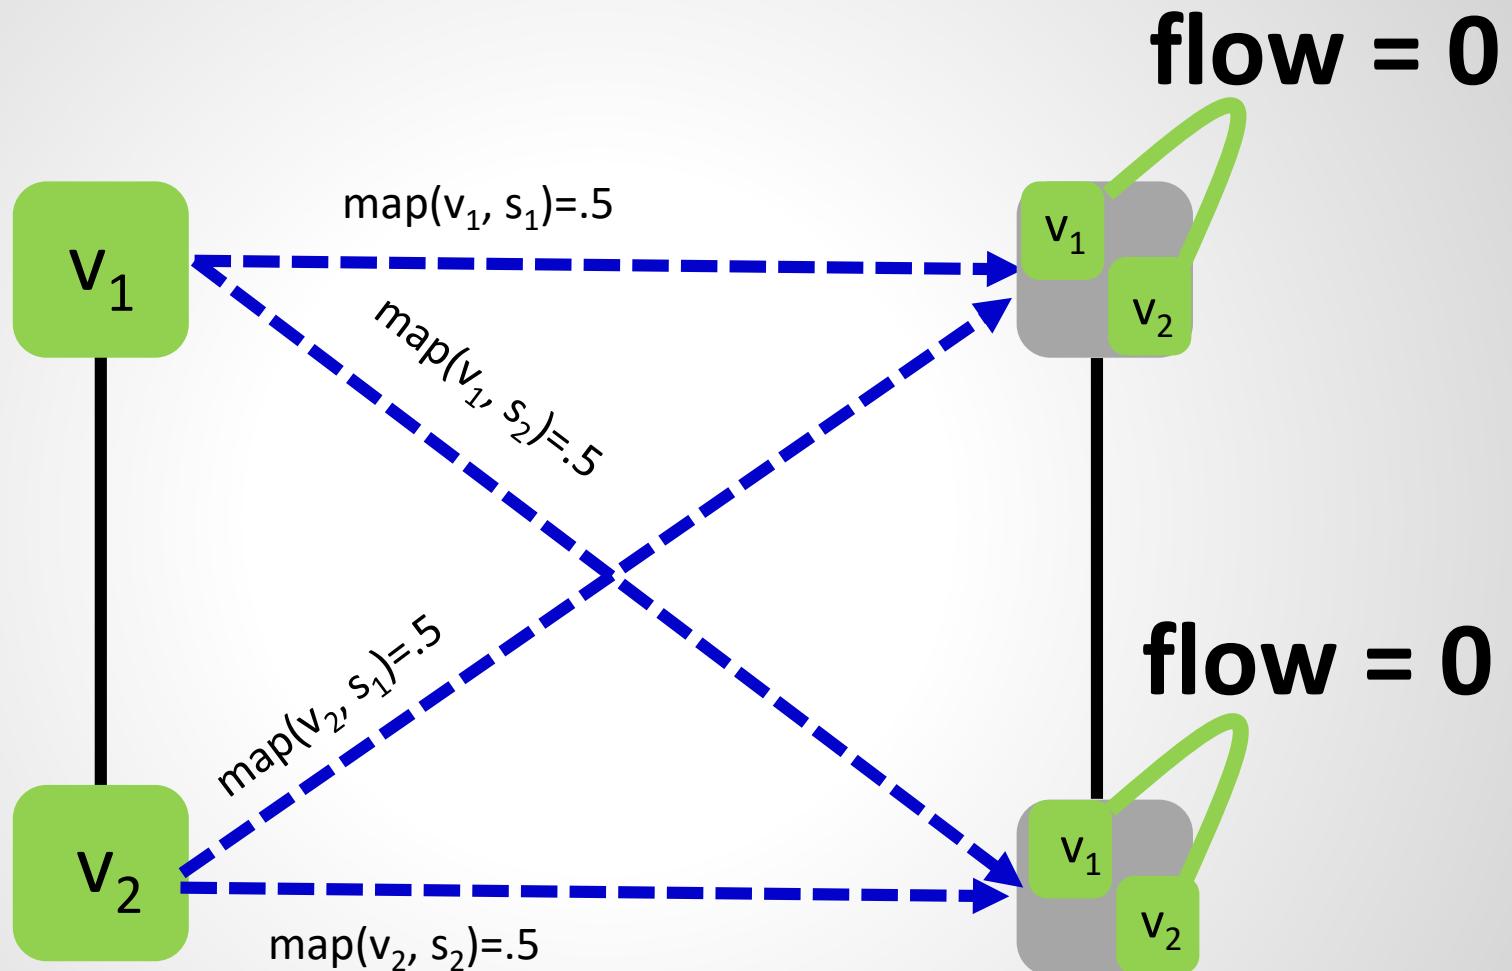
# What will happen in this example?



# What will happen in this example?



# What will happen in this example?



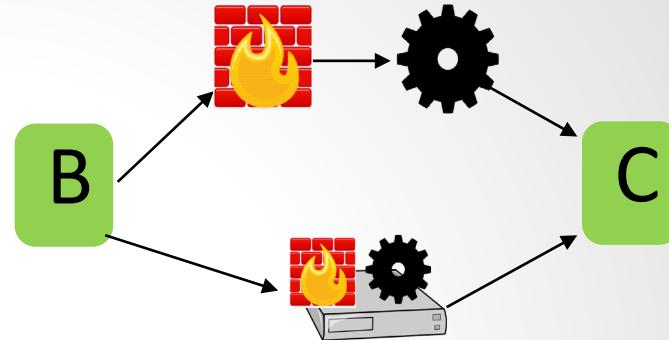
Minimal flow = 0: fulfills flow conservation! Relaxation useless: does not provide any lower bound or indication of good mapping!

*That's all Folks!*

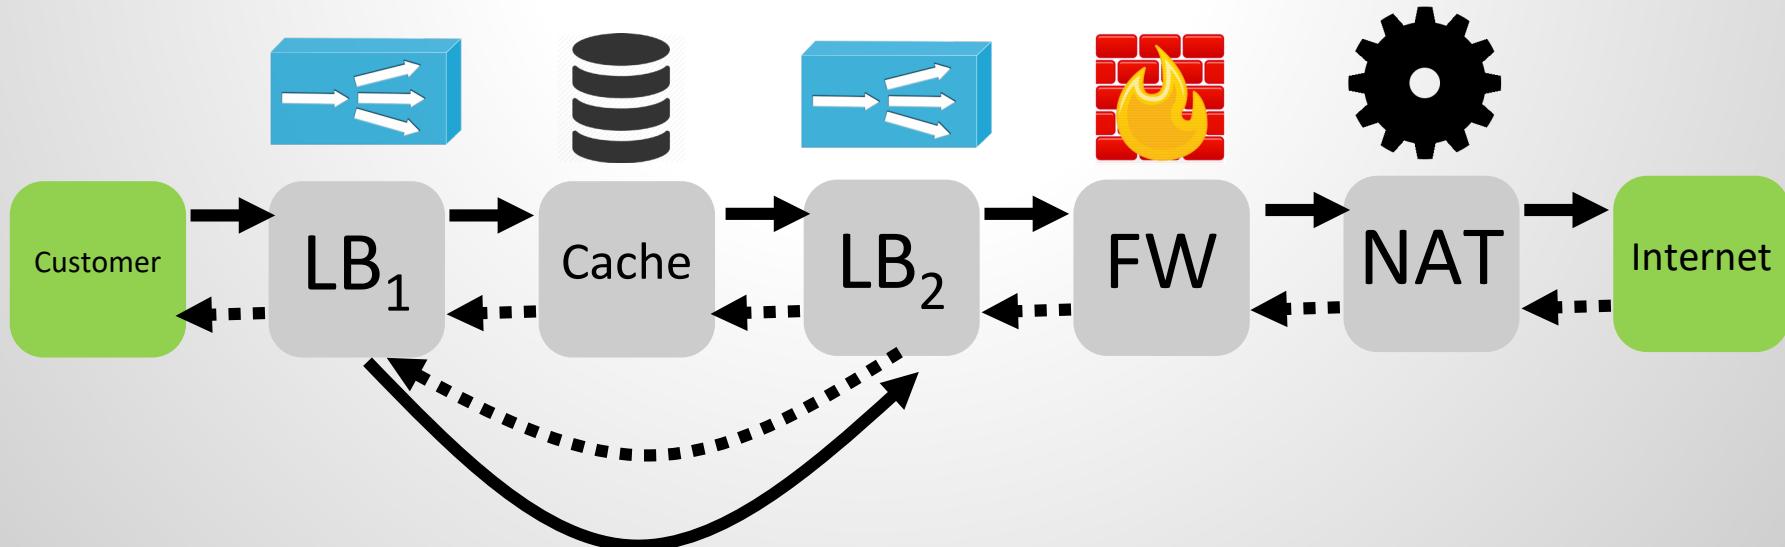
Wait a minute!  
These problems need to be solved!  
And they often can, even with guarantees.

# Theory vs Practice: In Practice There is Hope!

In practice, requests may have more structure:



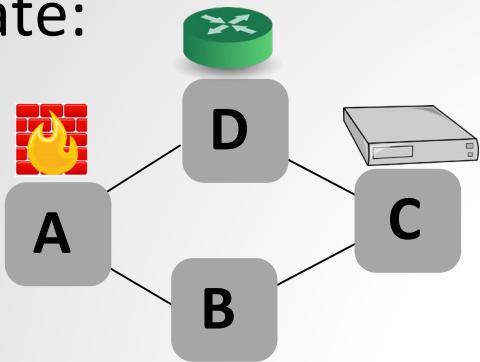
Even this beast:



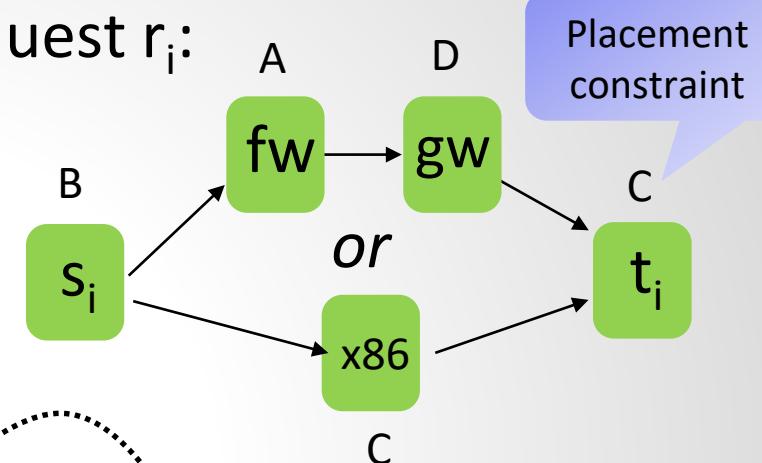
# But In Theory There is Hope Too: Approximations!

## Product graphs and randomized rounding

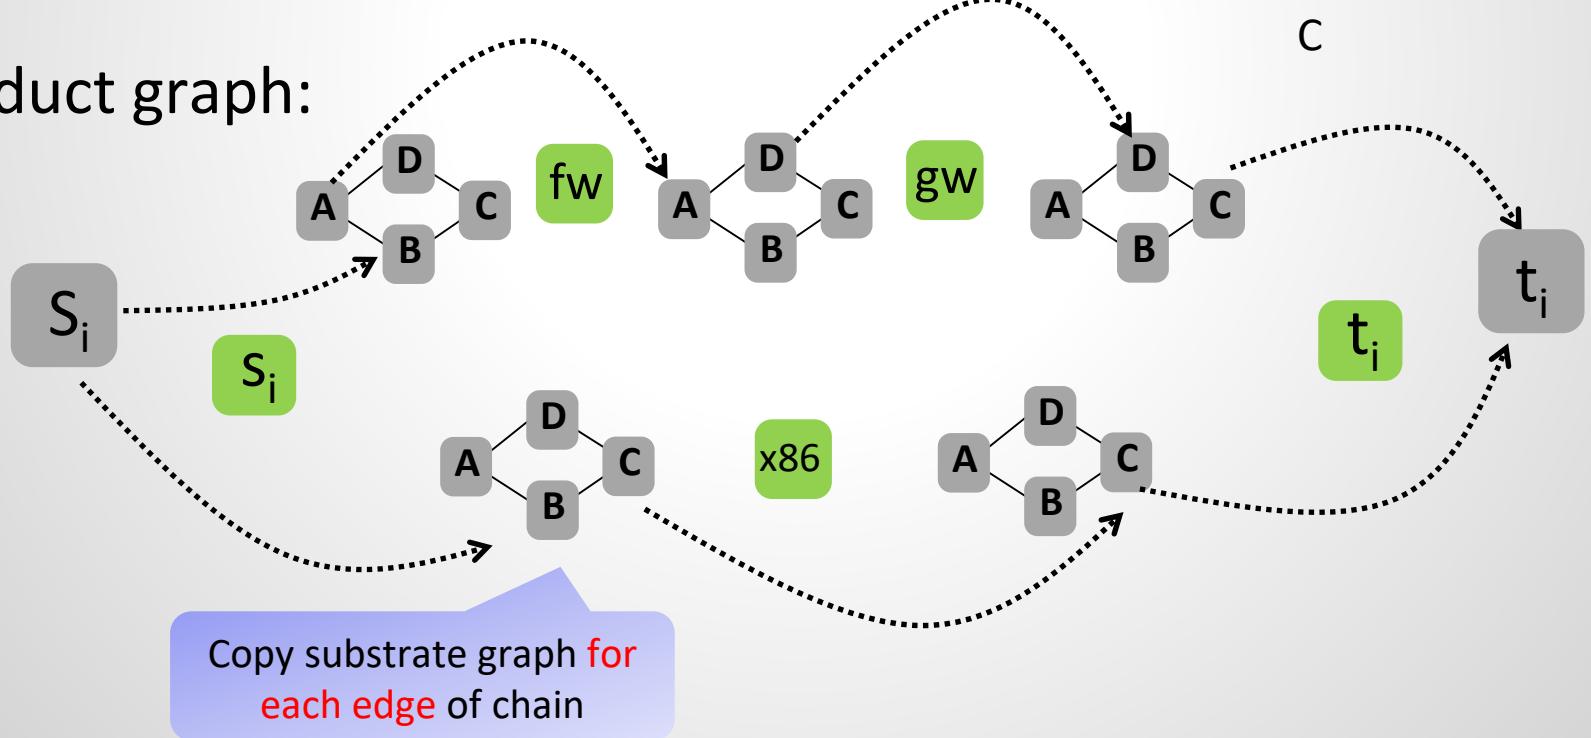
Substrate:



$i^{\text{th}}$  request  $r_i$ :



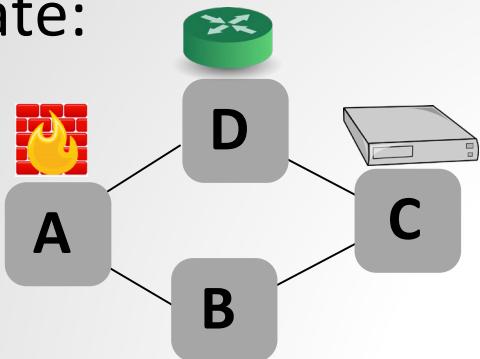
Product graph:



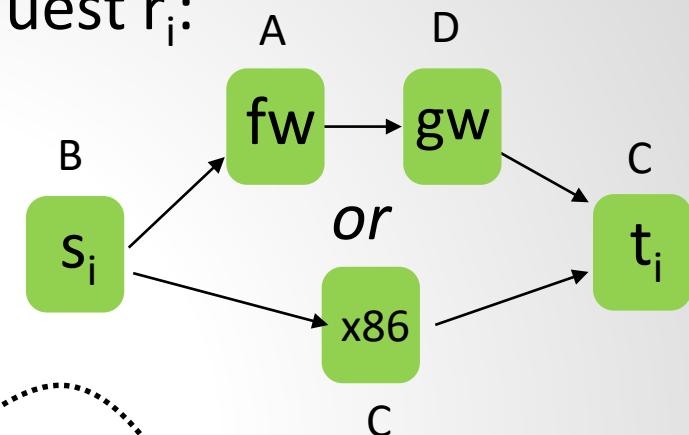
# But In Theory There is Hope Too: Approximations!

## Product graphs and randomized rounding

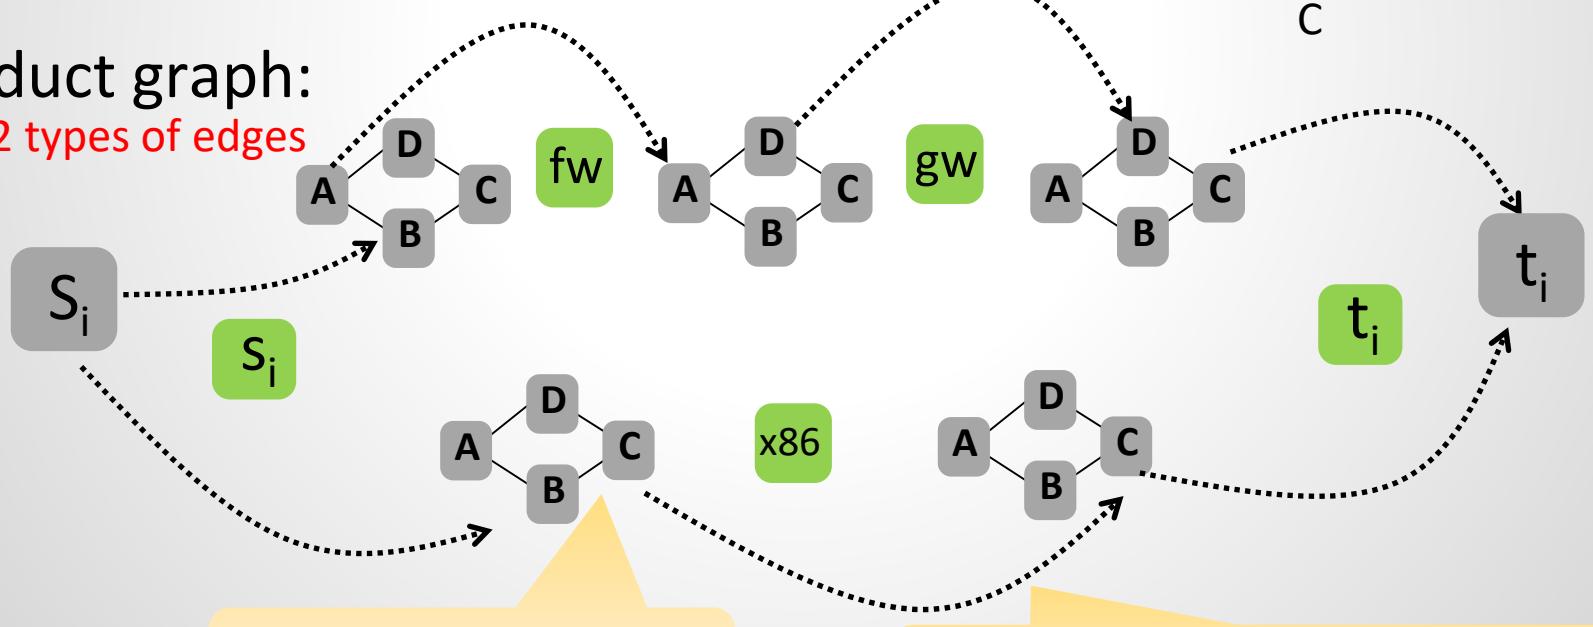
Substrate:



$i^{\text{th}}$  request  $r_i$ :



Product graph:  
with 2 types of edges



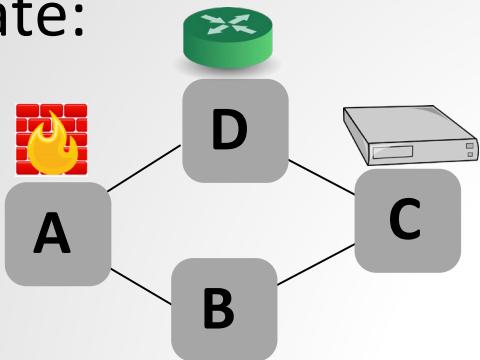
**Routing edge:** graph edge  
on same layer

**Processing edge:** processing happens on C:  
connect C to C in next layer!

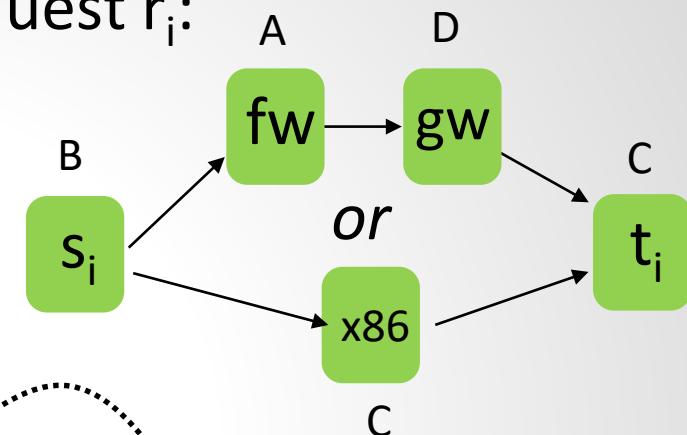
# But In Theory There is Hope Too: Approximations!

## Product graphs and randomized rounding

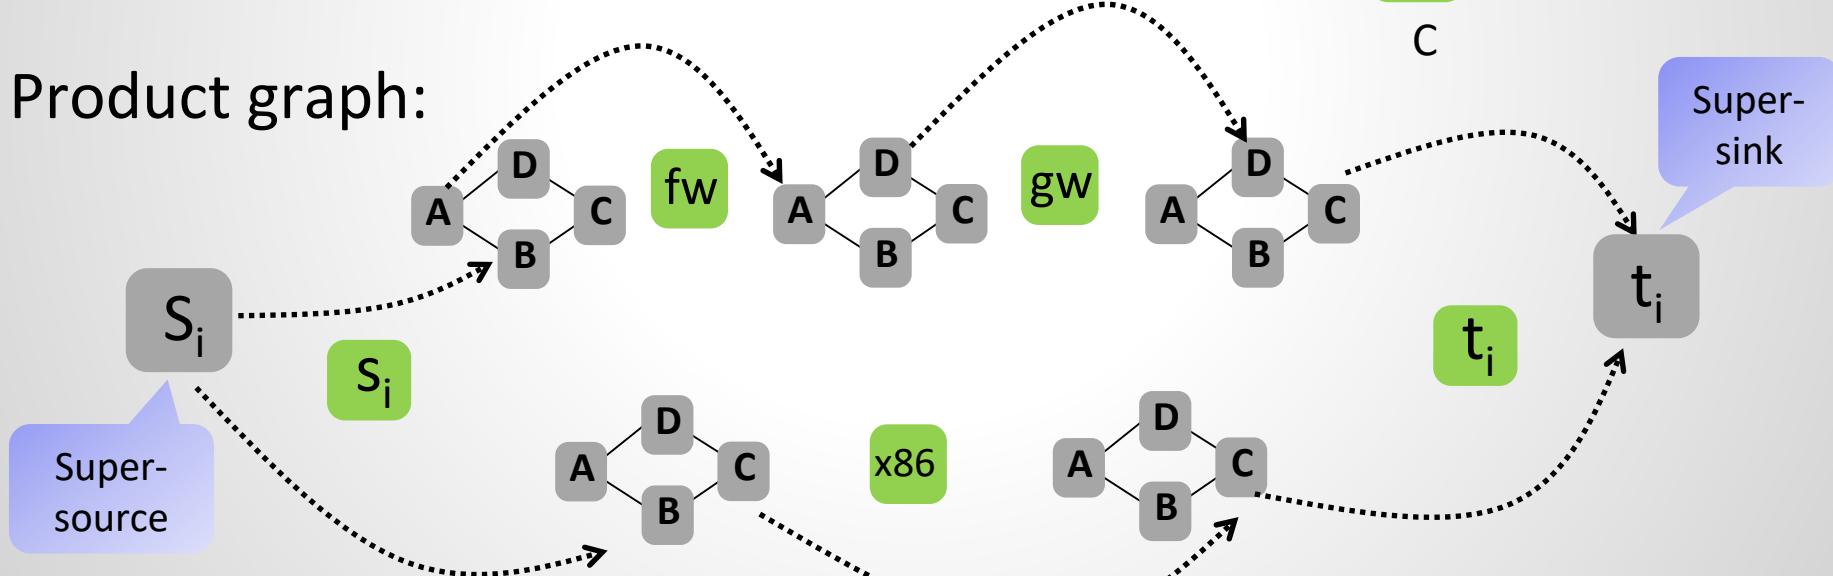
Substrate:



$i^{\text{th}}$  request  $r_i$ :



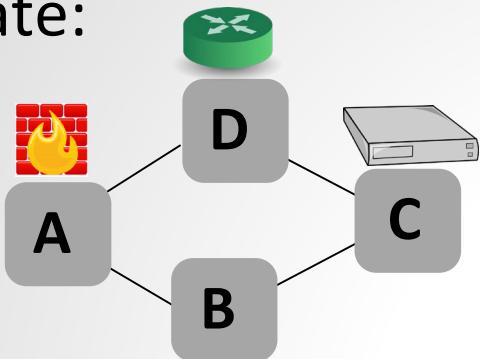
Product graph:



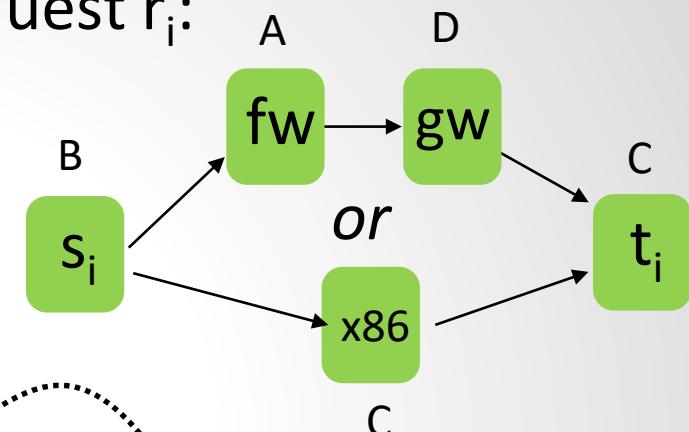
# But In Theory There is Hope Too: Approximations!

## Product graphs and randomized rounding

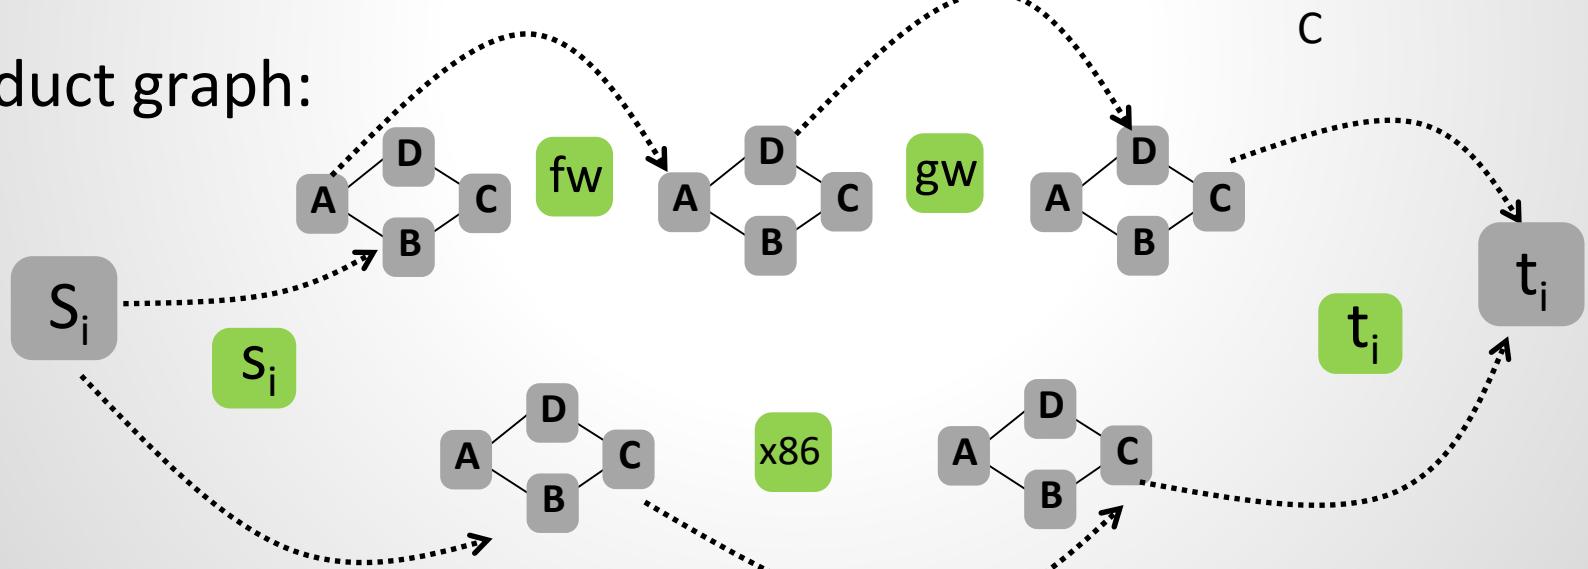
Substrate:



$i^{\text{th}}$  request  $r_i$ :



Product graph:

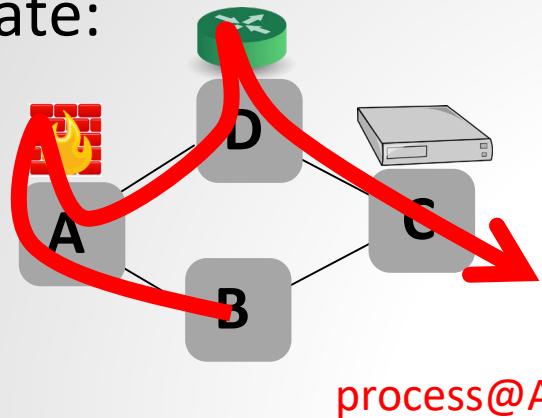


Any  $(s_i, t_i)$  flow presents a route of the request  $r_i$ !

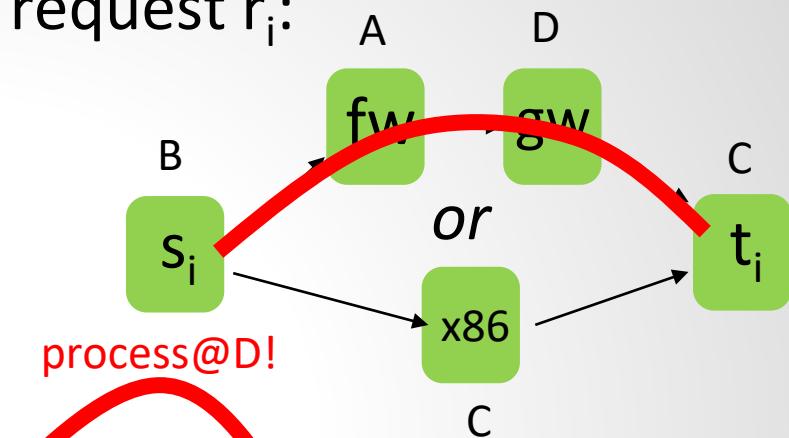
# But In Theory There is Hope Too: Approximations!

## Product graphs and randomized rounding

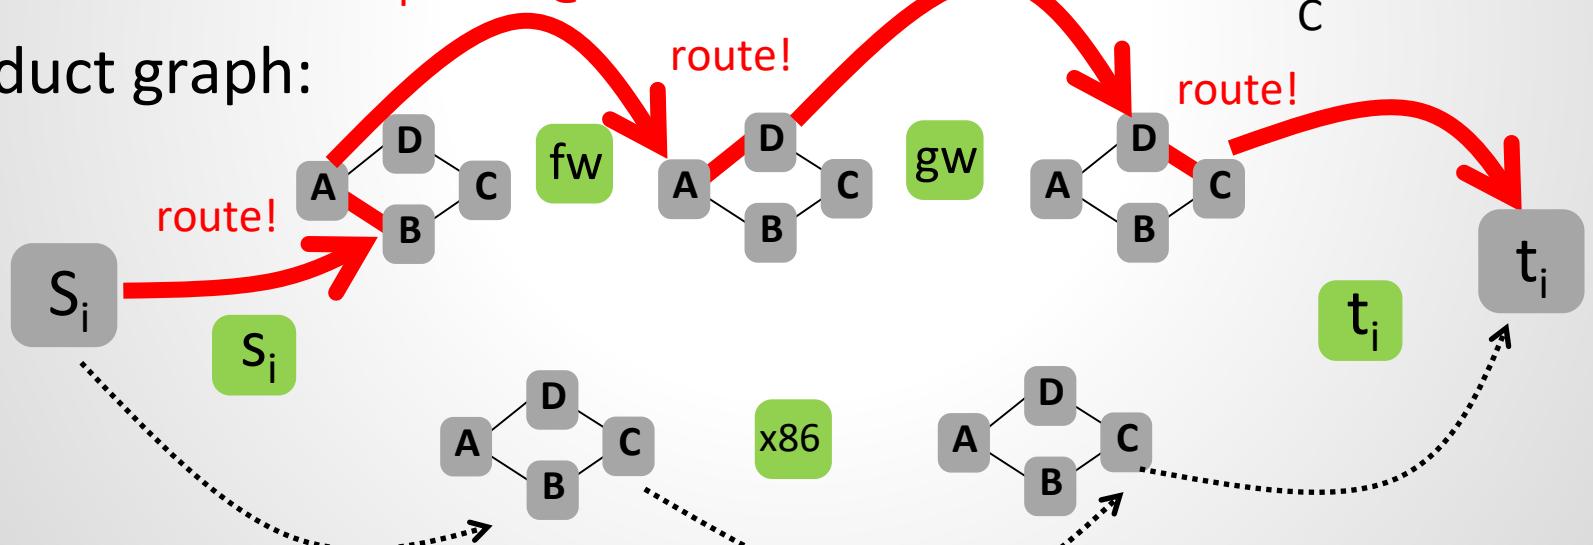
Substrate:



$i^{\text{th}}$  request  $r_i$ :



Product graph:

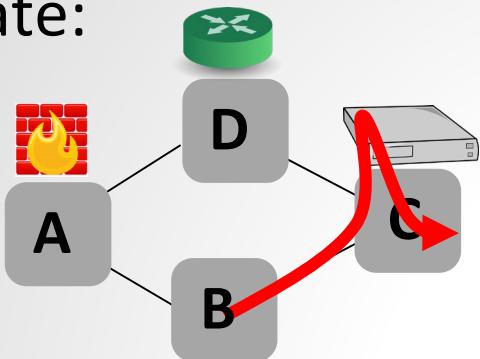


Any  $(s_i, t_i)$  flow presents a route of the request  $r_i$ !

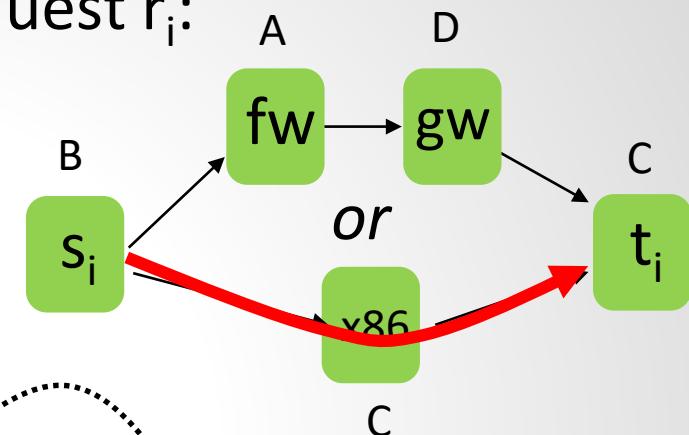
# But In Theory There is Hope Too: Approximations!

## Product graphs and randomized rounding

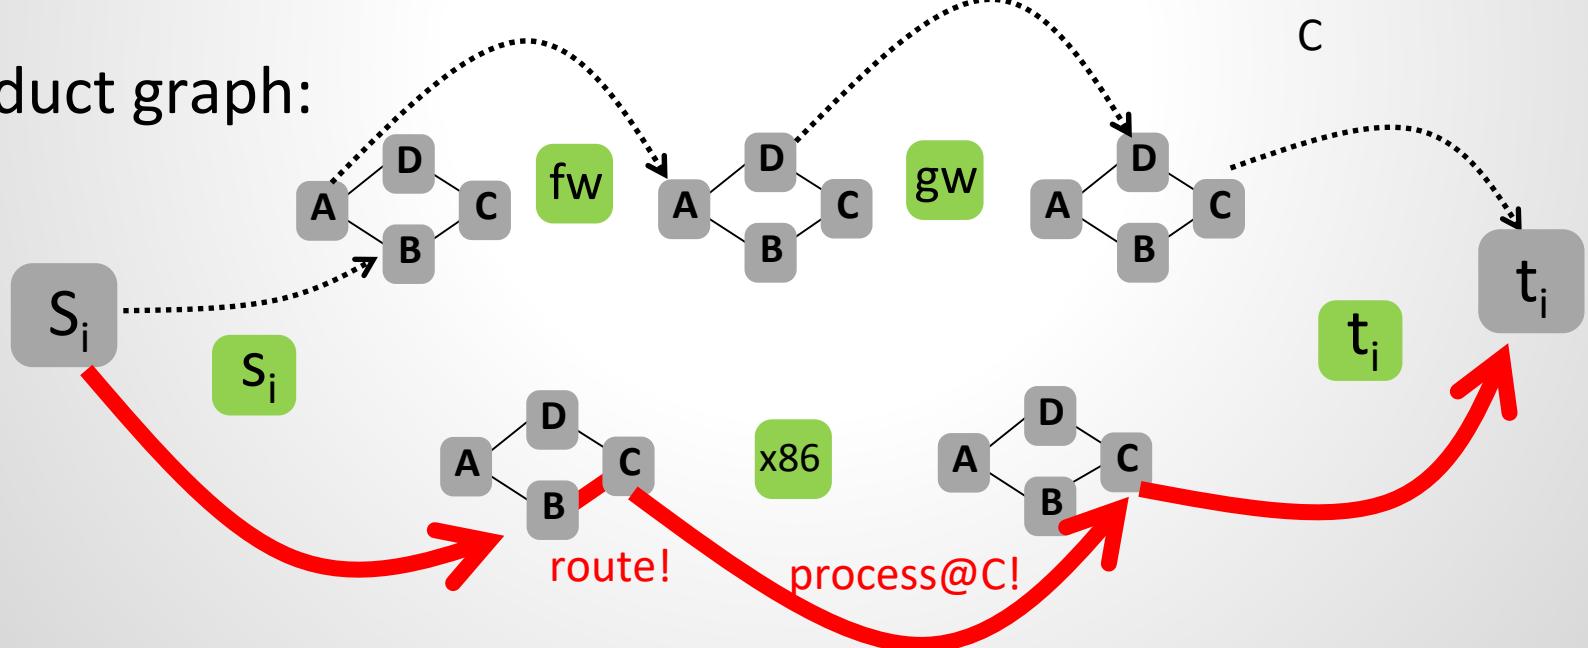
Substrate:



$i^{\text{th}}$  request  $r_i$ :



Product graph:



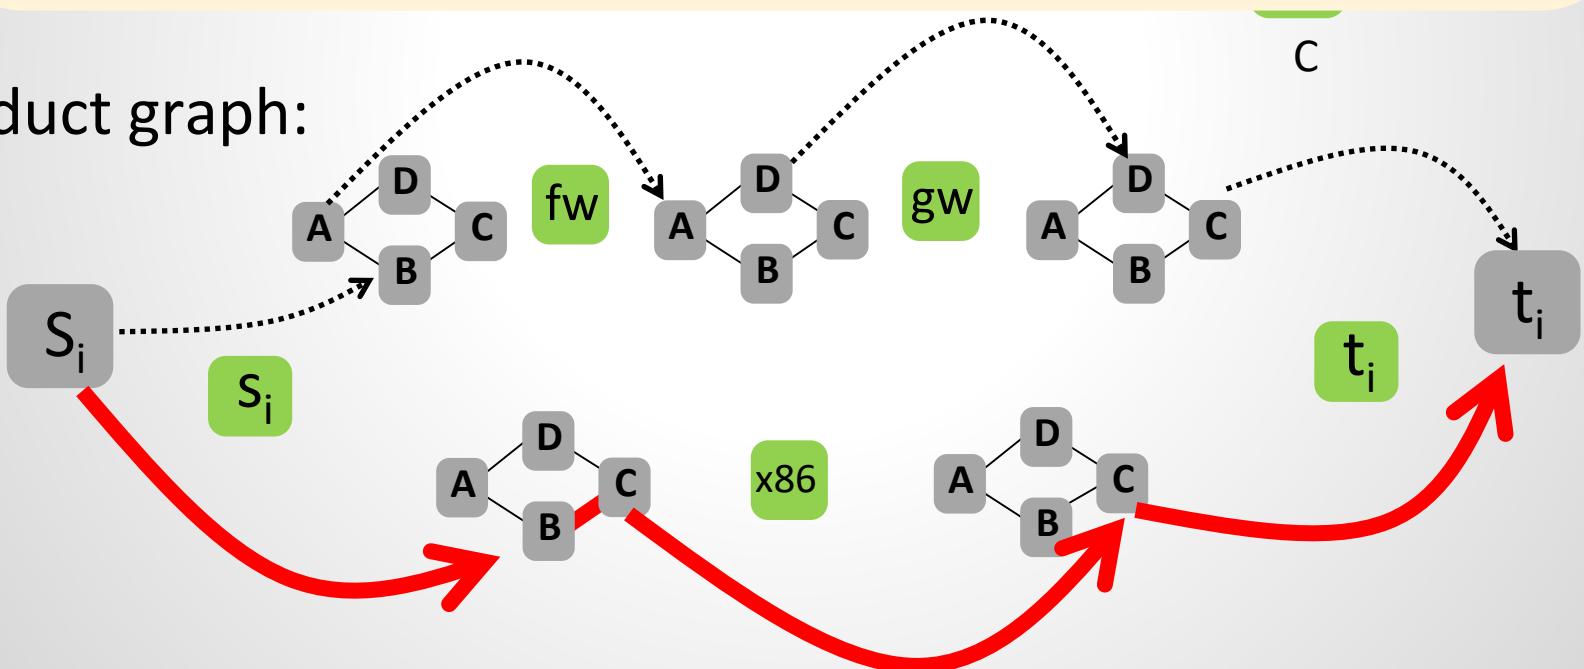
Any  $(s_i, t_i)$  flow presents a route of the request  $r_i$ !

# But In Theory There is Hope Too: Approximations!

Product graphs and randomized rounding

This problem can be solved using mincost unsplittable multi-commodity flow (approximation) algorithms (e.g., randomized rounding).

Product graph:



Any  $(s_i, t_i)$  flow presents a route of the request  $r_i$ !

# Invitation: A Roadtrip Through The Opportunities and Challenges of Network Isolation

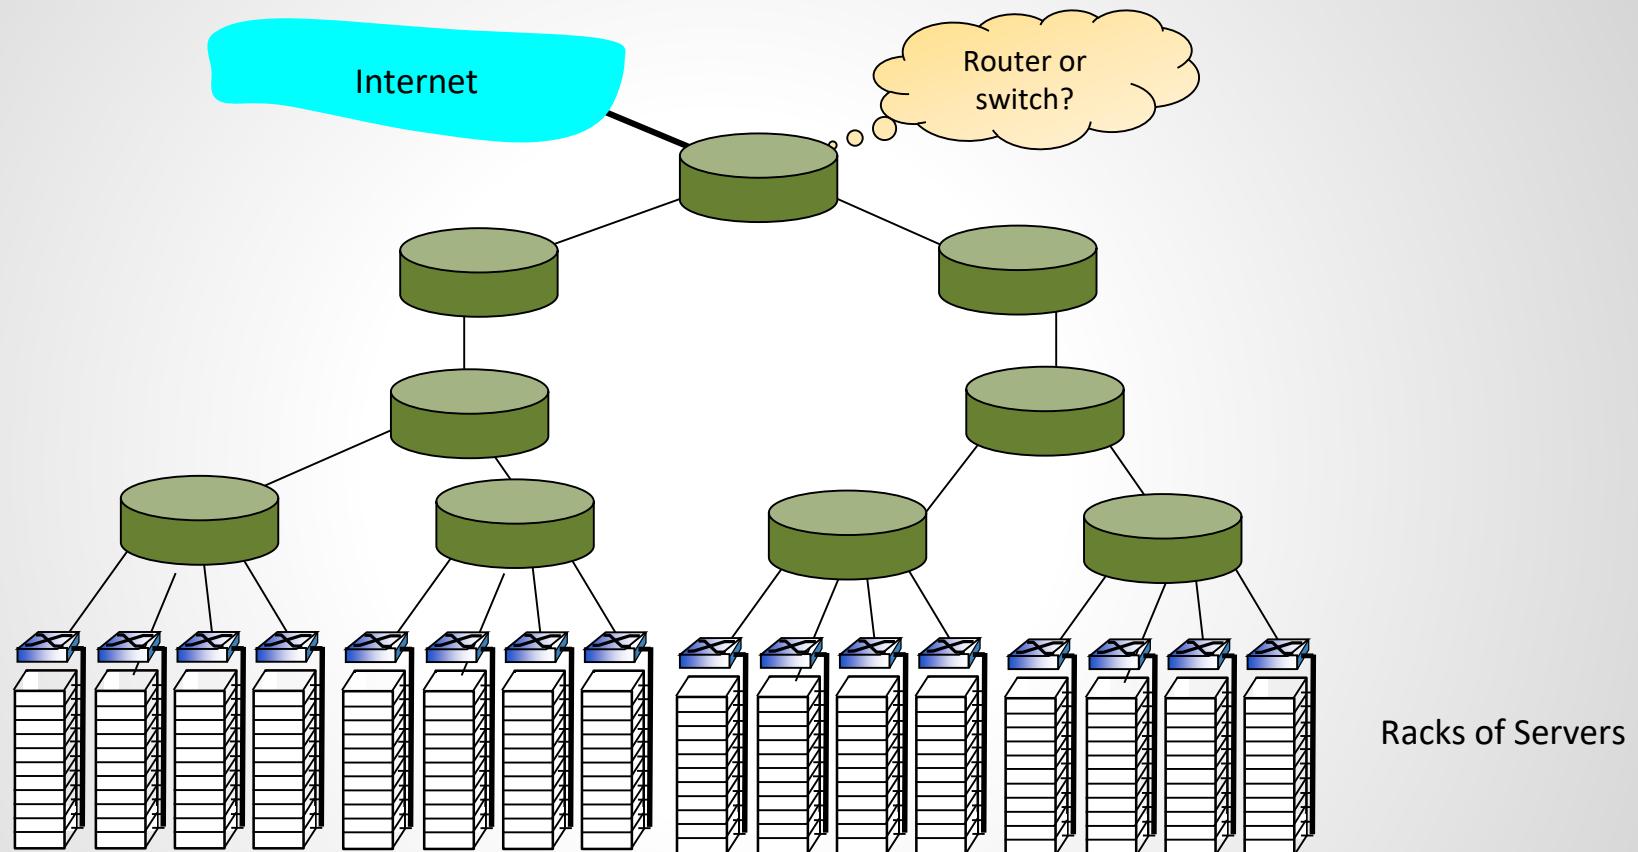
- ❑ Opportunities
  - ❑ Algorithmic opportunities
  - ❑ **Technological opportunities**
- ❑ Challenges
  - ❑ Modelling challenges
  - ❑ Security challenges
- ❑ A perspective how AI can improve slicing efficiency and security



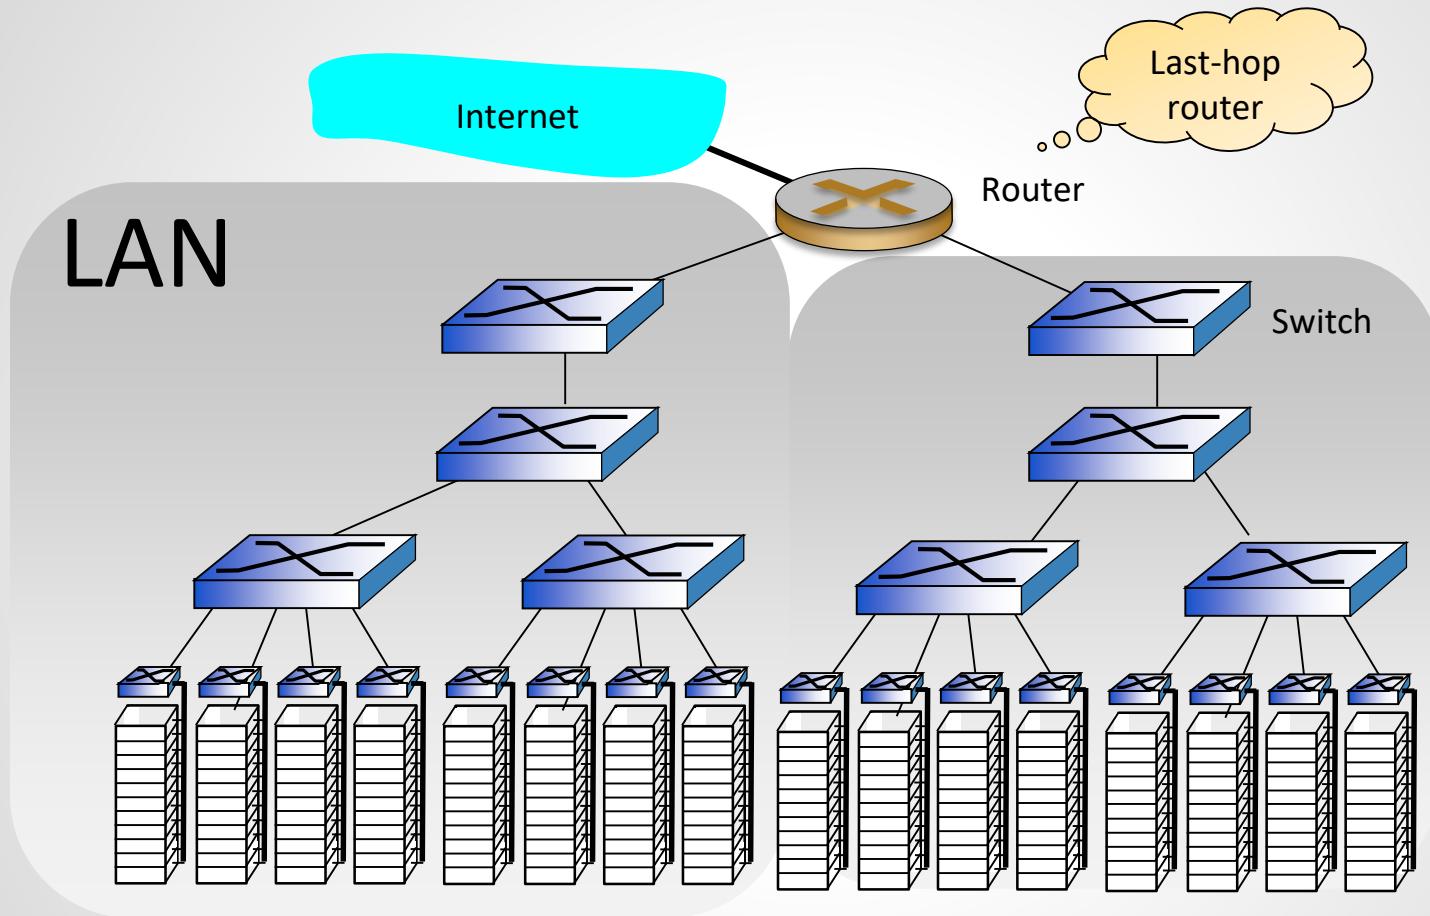
Road map 1927: Arizona and New Mexico

# Example: Isolation in Datacenter

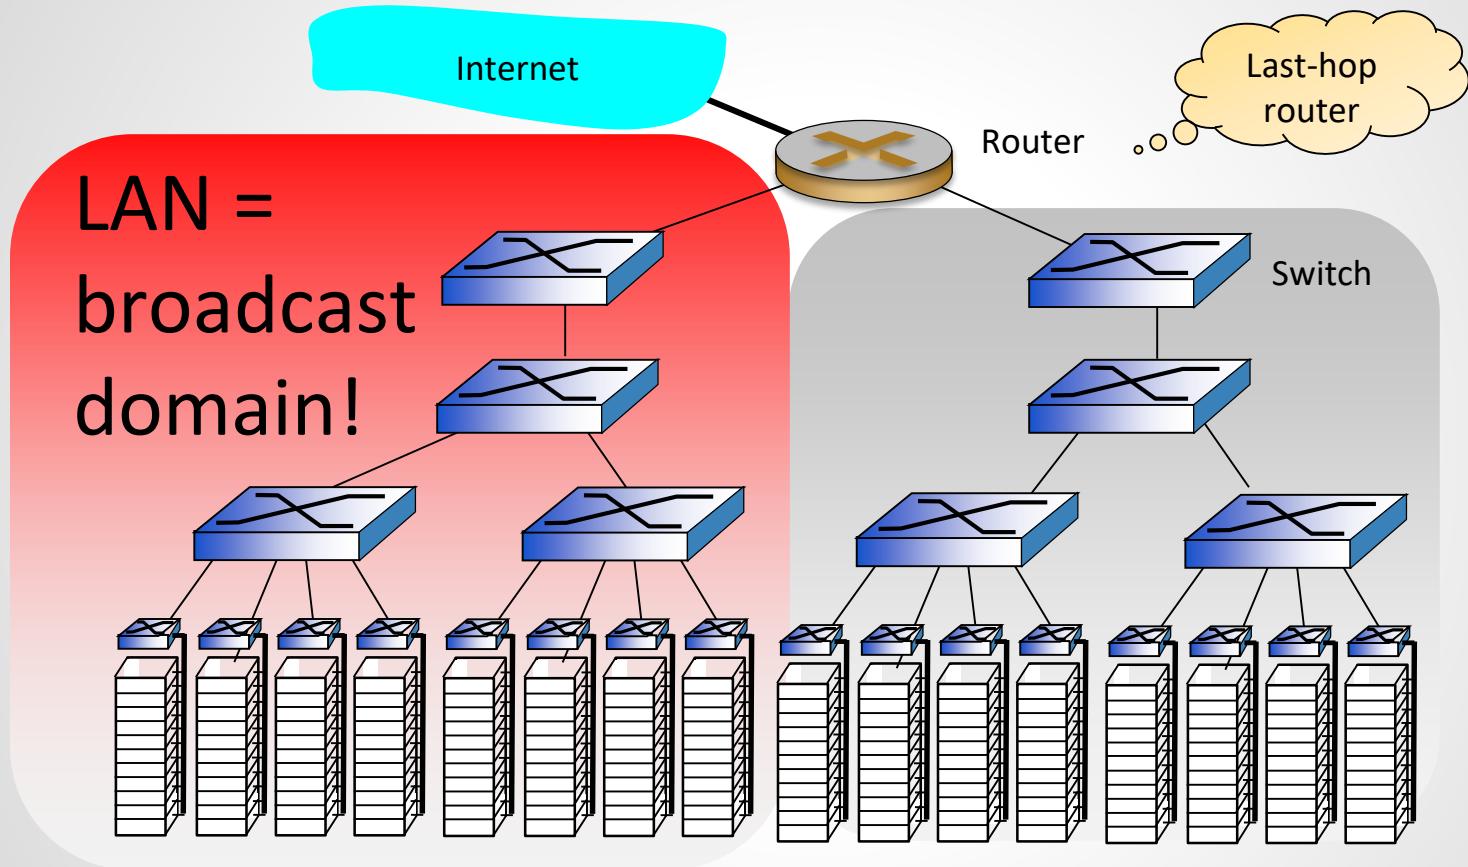
Tradeoff, traditionally:



# Architecture #1



# Architecture #1

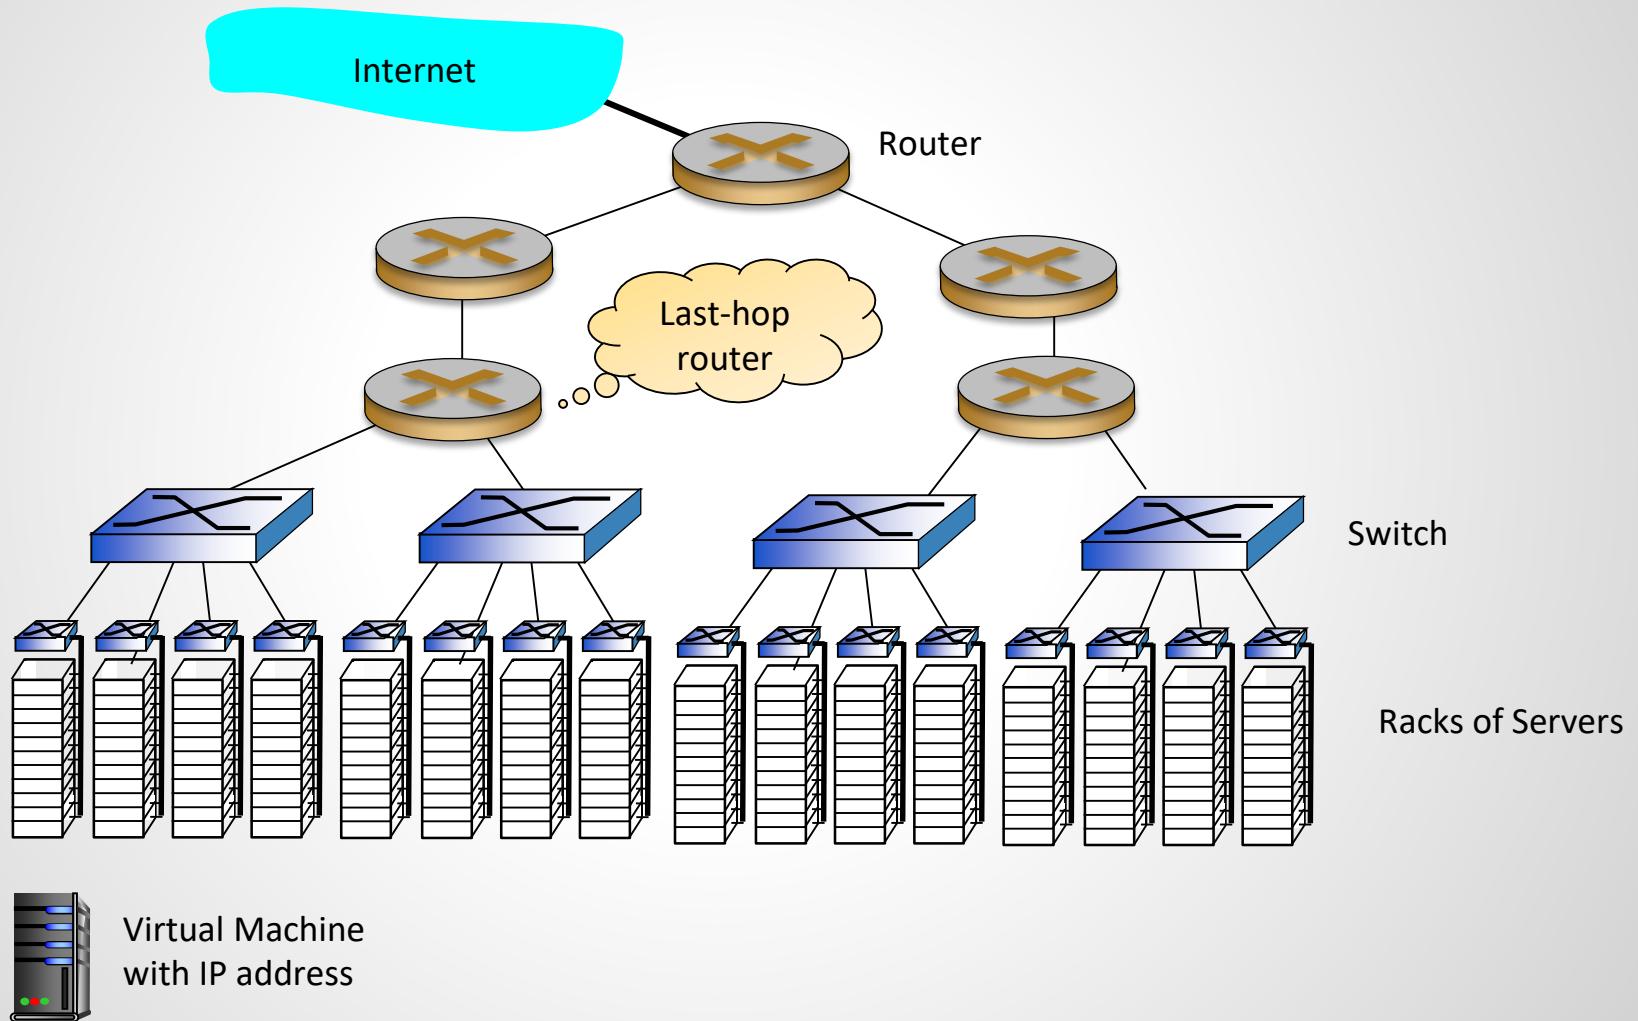


Virtual Machine  
with IP address

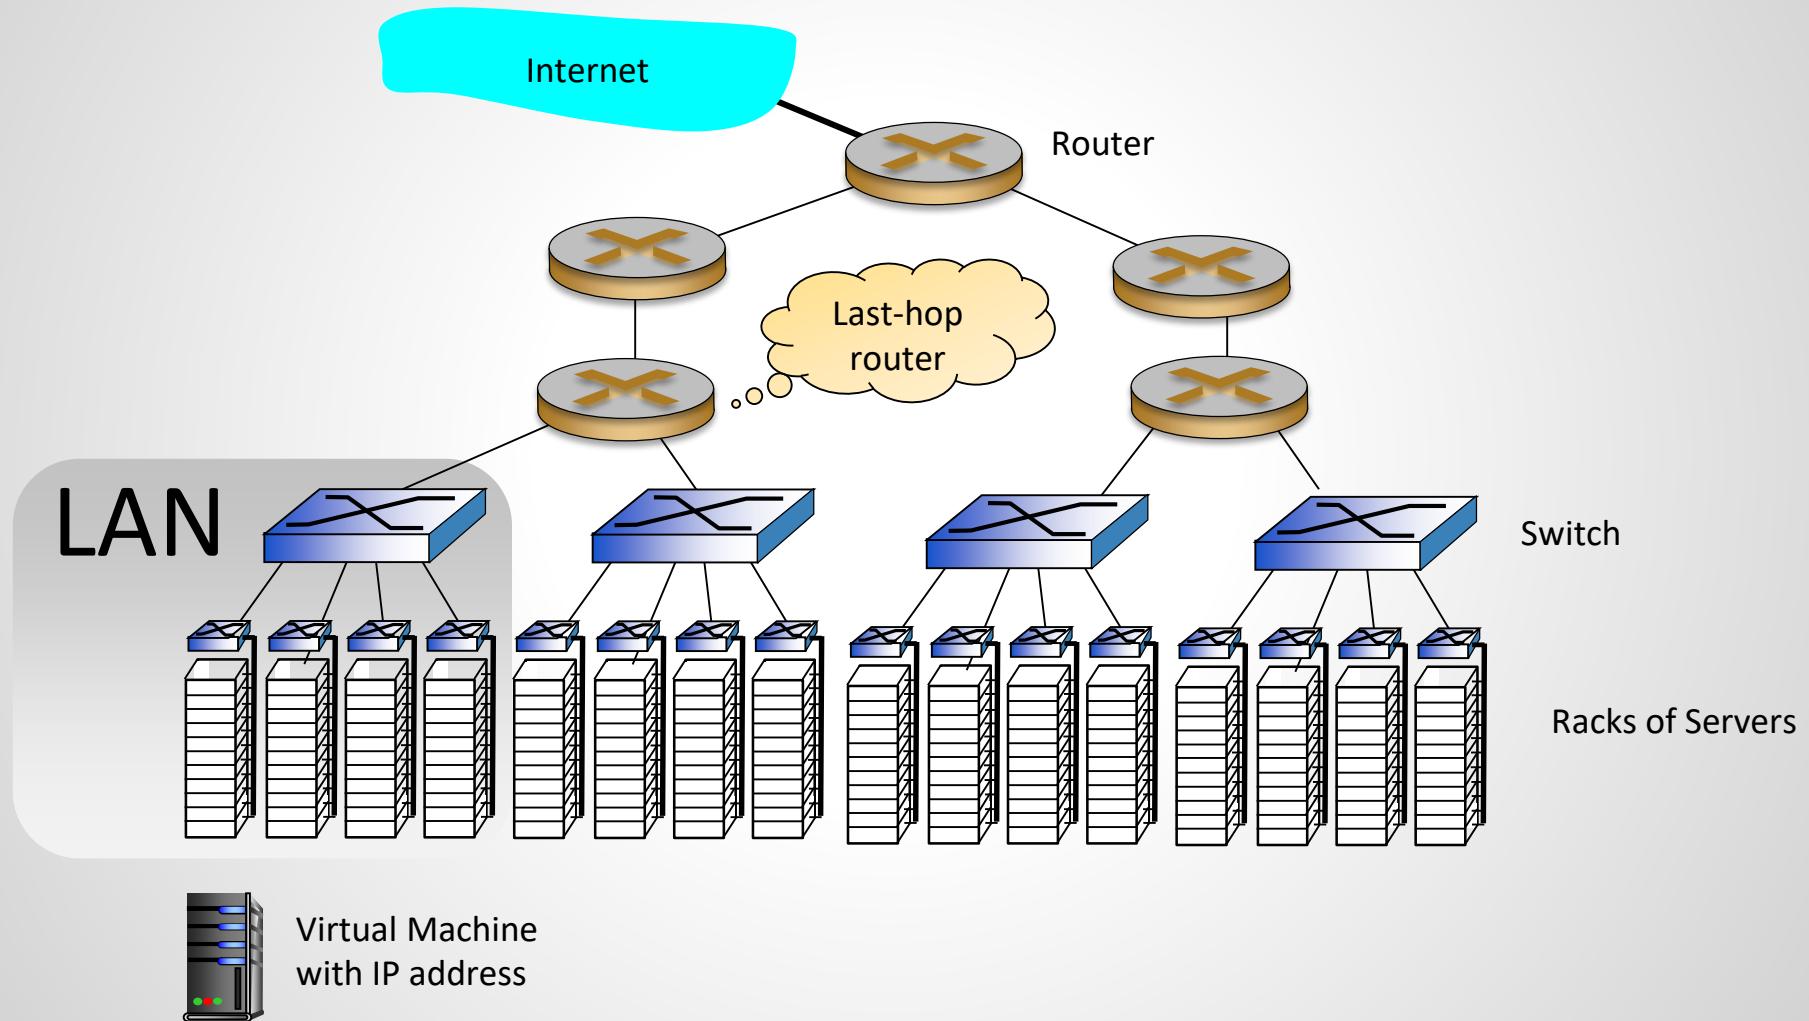
No need to change IP!

A large LAN: High mobility...  
... but high overhead due to learning  
and broadcasting.

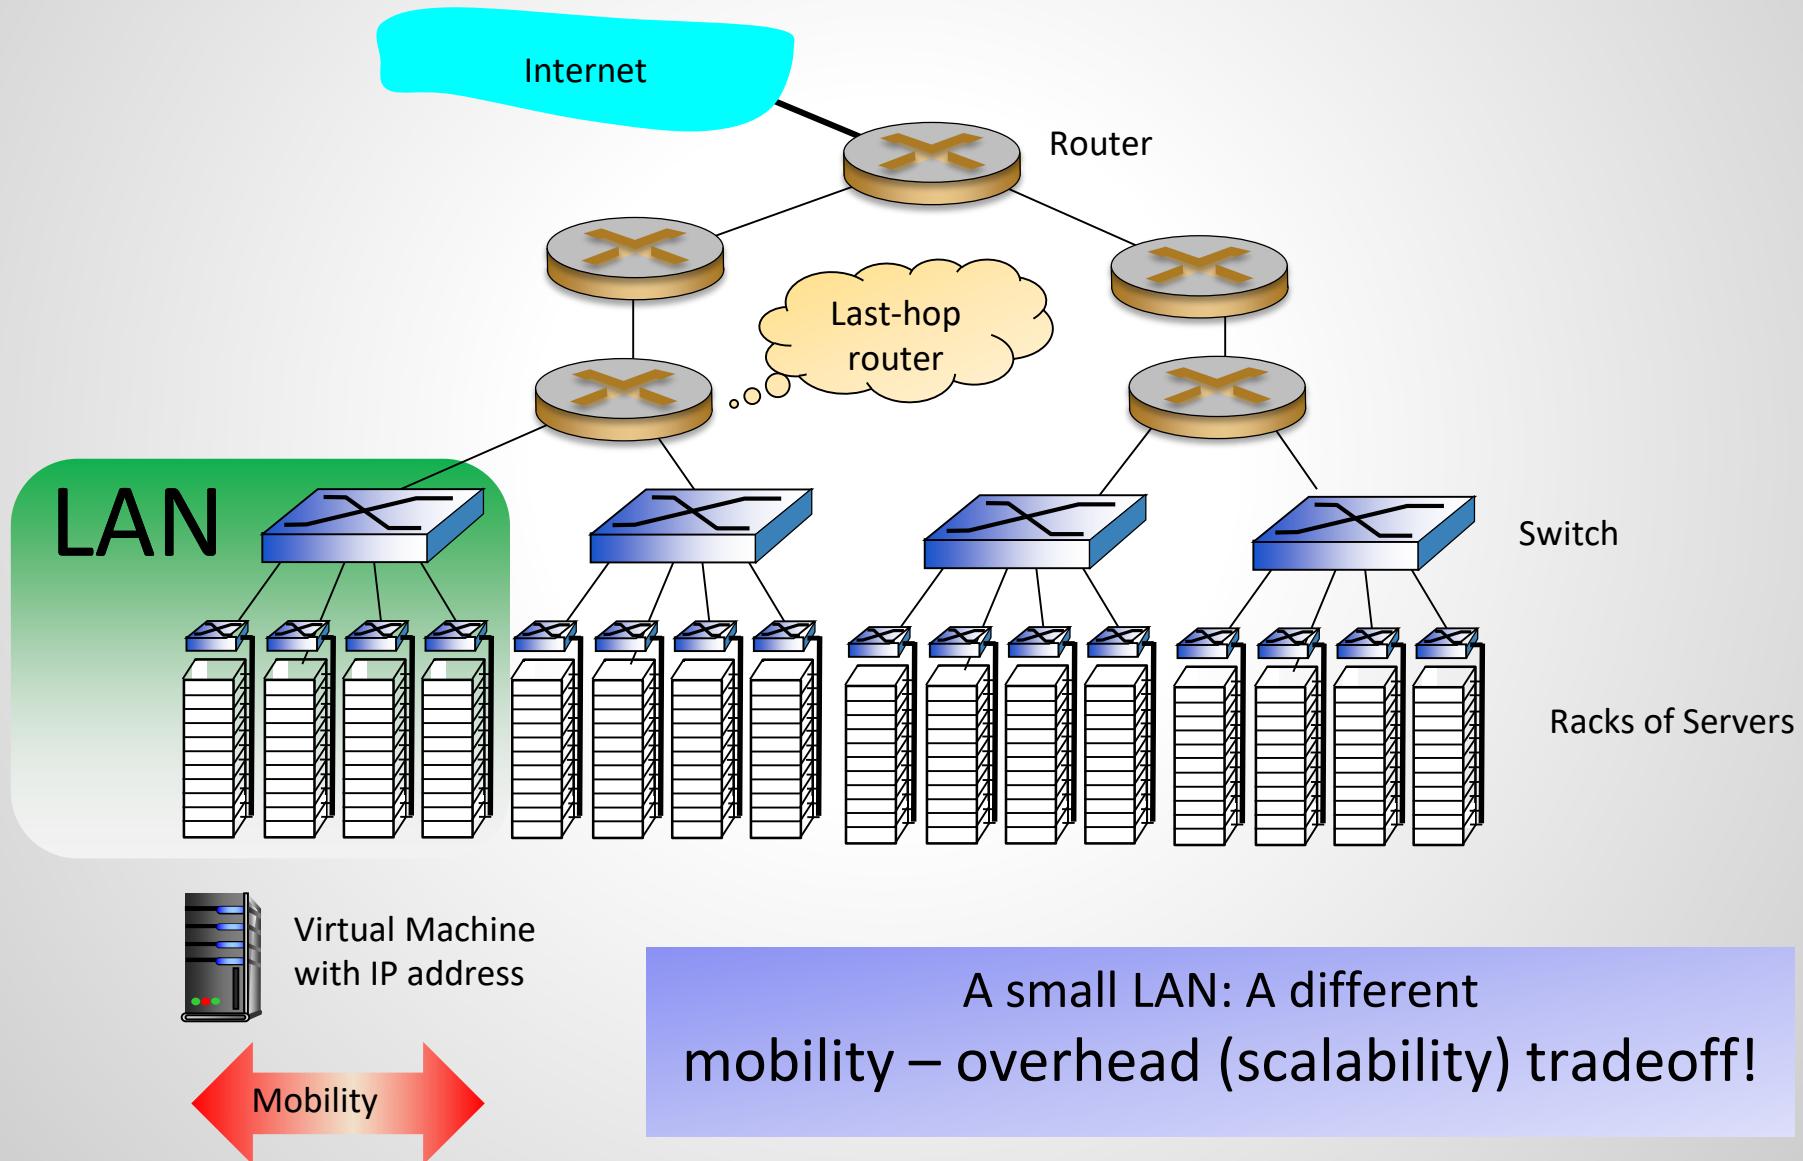
# Architecture #2



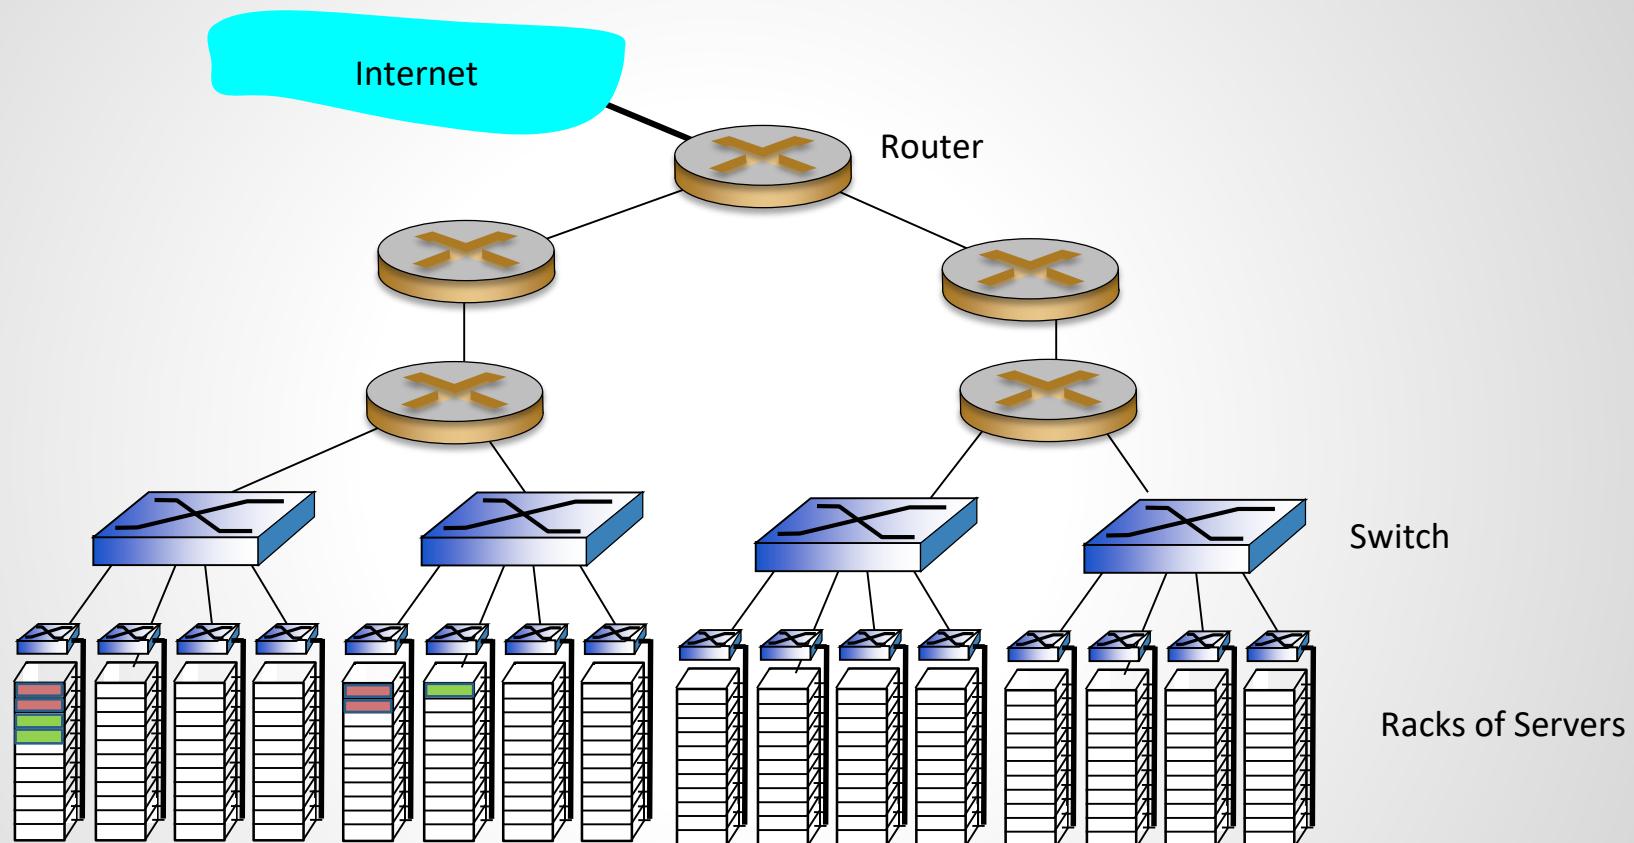
# Architecture #2



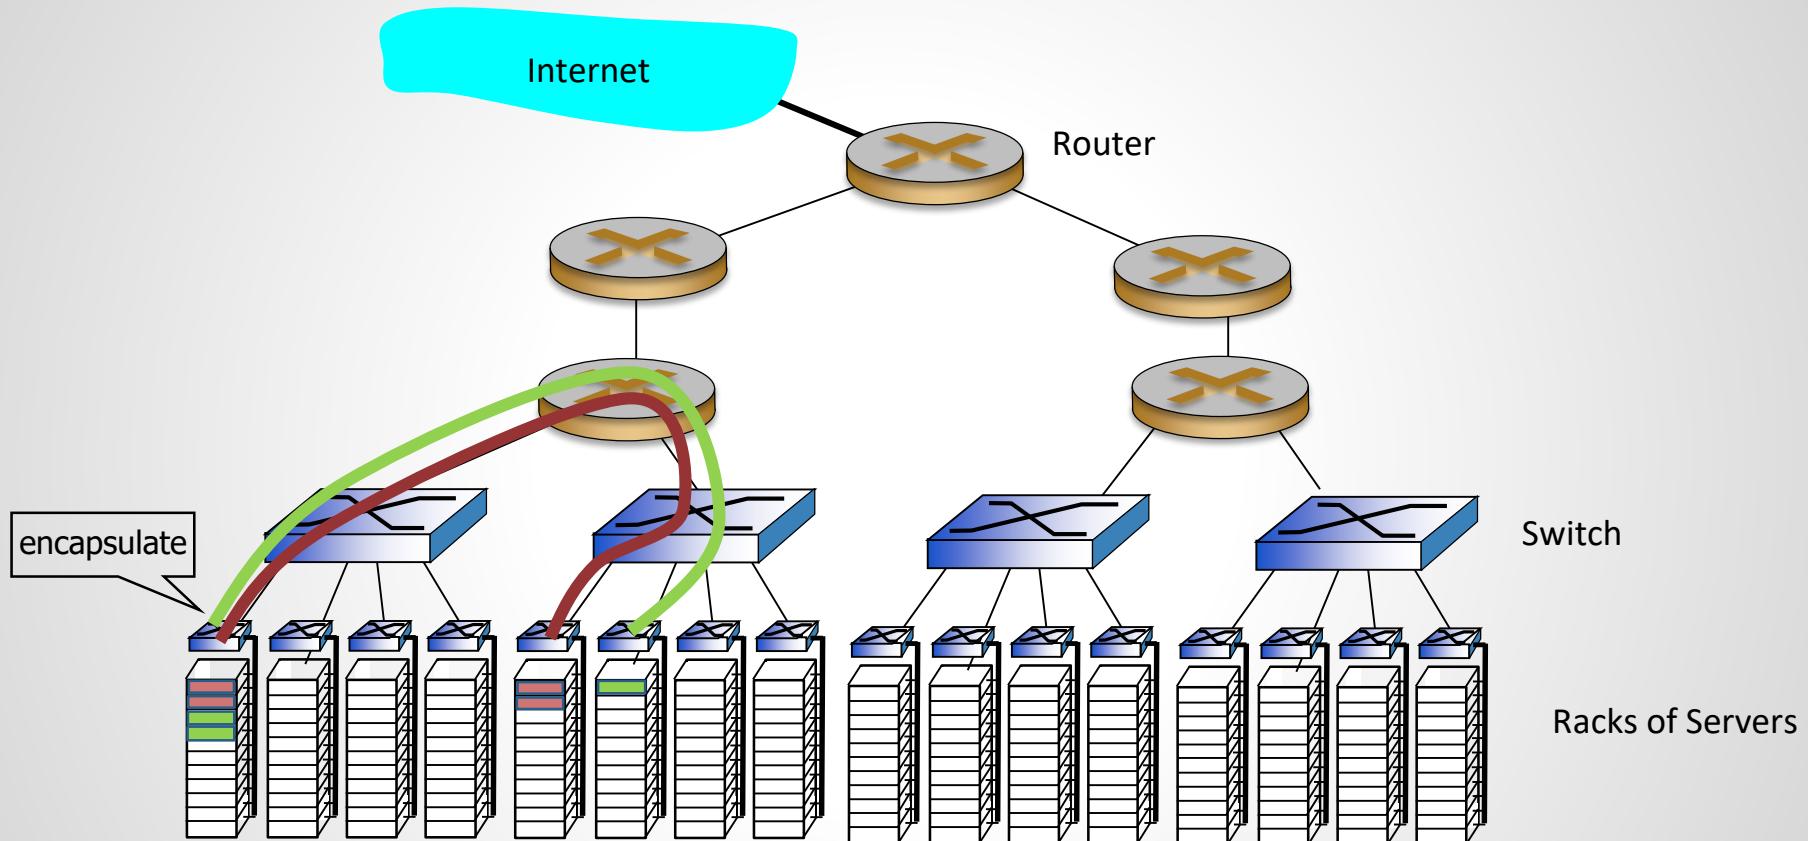
# Architecture #2



# Virtualization Technologies: Isolation of Tenants



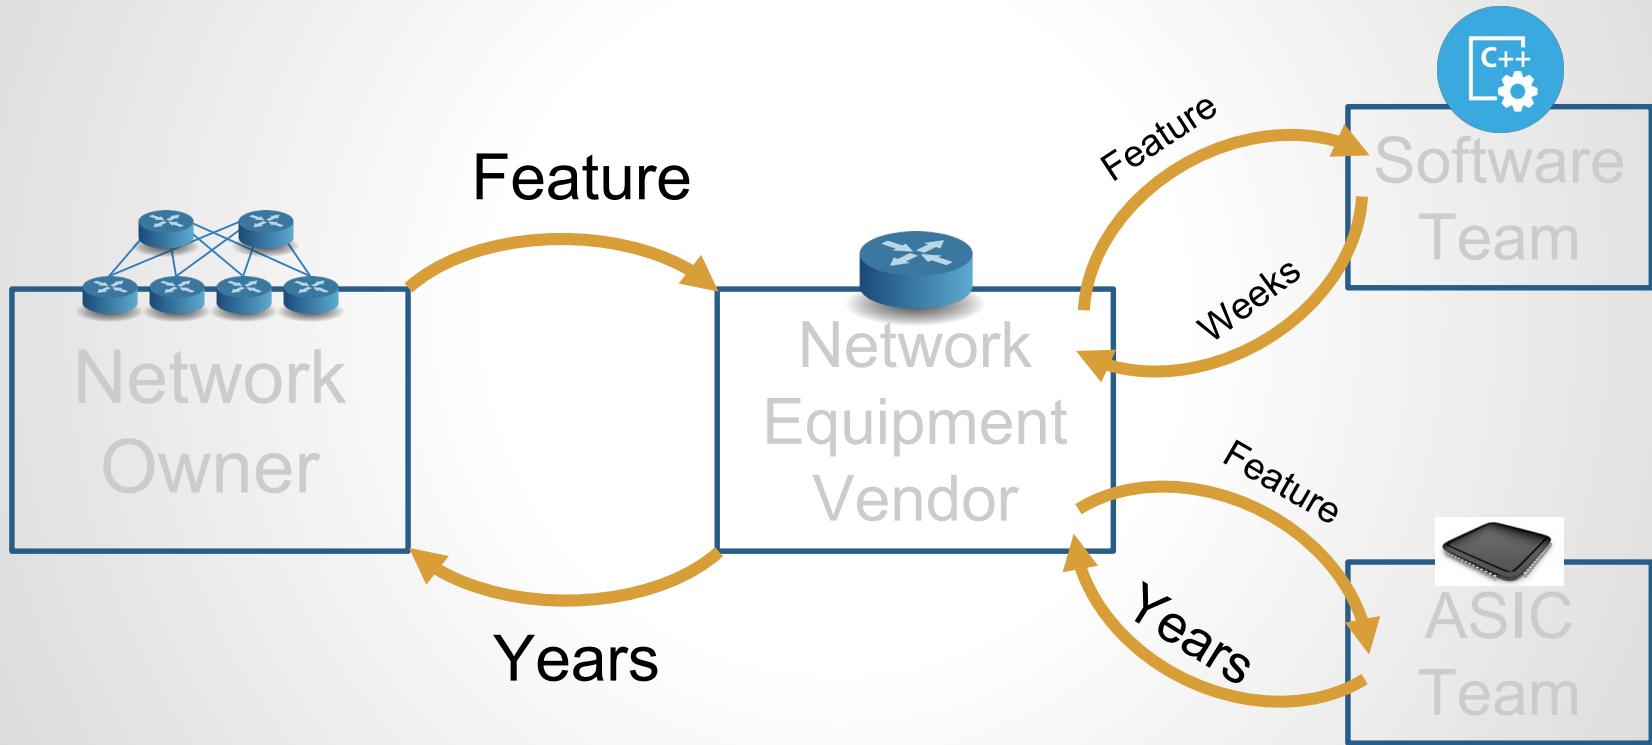
# Virtualization Technologies: Isolation of Tenants



Network virtualization: VLANs, VxLANs, tunneling, ... *or SDN!*

# In the Past, Introducing Virtualization Technologies Took Years

Example: VxLAN



# In the Past: Slow Innovation

Operator says:

I need extended VTP  
(VLAN Trunking  
Protocol) / a 3rd  
spanport etc. !

Vendor's answer:

Buy one of these!



# In the Past: Slow Innovation

Operator says:



I need  
something  
better than STP  
for my data-  
center...

Vendor's answer:



We don't  
have that!

# Opportunity: Softwarization, e.g., Programmable Dataplanes

## Innovation at the Speed of Software Development

The screenshot shows a YouTube video player interface. At the top, there's a search bar and a magnifying glass icon. Below the video frame, there are playback controls: a play button, a progress bar indicating 8:15 / 36:37, and other standard video controls.

The video content itself features a slide with the following text and diagram:

Network systems will be programmed “top-down”

*“This is precisely how you must process packets”*

Diagram illustrating the architecture of a Programmable Switch:

- A **Switch OS** box is connected to a **Driver** box.
- The **Driver** box is connected to a **Programmable Switch**.
- Two vertical arrows point downwards from the **Switch OS** and **Driver** boxes to the **Programmable Switch**.

On the left side of the slide, there is a small video thumbnail showing a man in a grey t-shirt standing and speaking.

Below the video frame, the video title is "Why Does the Internet Need a Programmable Forwarding Plane with Nick McKeown" and the upload date is "8.984 Aufrufe • 19.01.2017". To the right of the video controls, there are social sharing icons: likes (110), dislikes (3), share (TEILEN), save (SPEICHERN), and more (three dots).

<https://www.youtube.com/watch?v=zR88NIg3n3g>

# Invitation: A Roadtrip Through The Opportunities and Challenges of Network Isolation

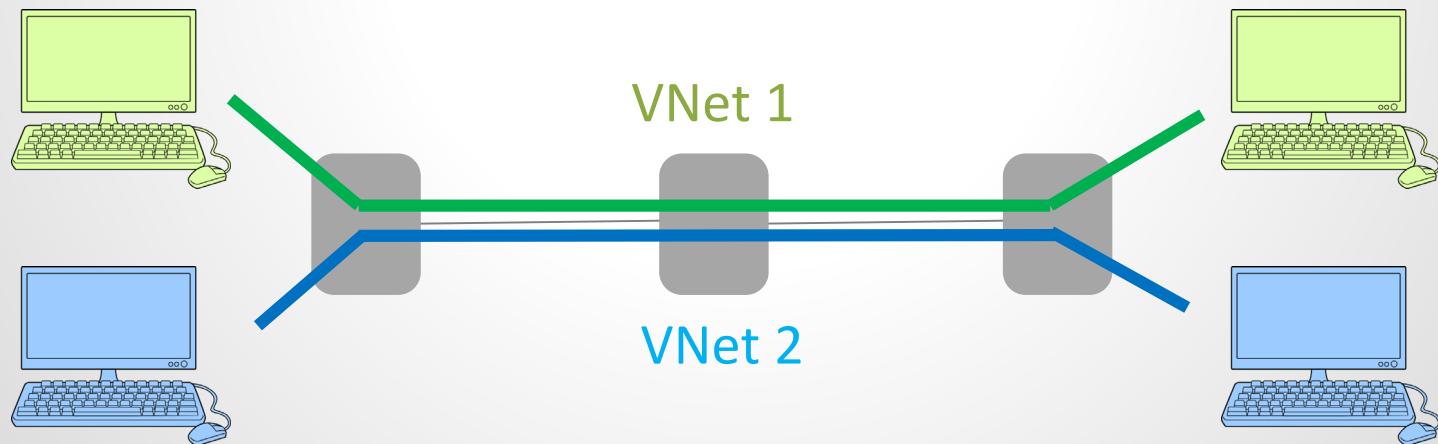
- ❑ Opportunities
  - ❑ Algorithmic opportunities
  - ❑ Technological opportunities
- ❑ Challenges
  - ❑ **Modelling challenges**
  - ❑ Security challenges
- ❑ A perspective how AI can improve slicing efficiency and security



Road map 1927: Arizona and New Mexico

# Models: Mind the Gap!

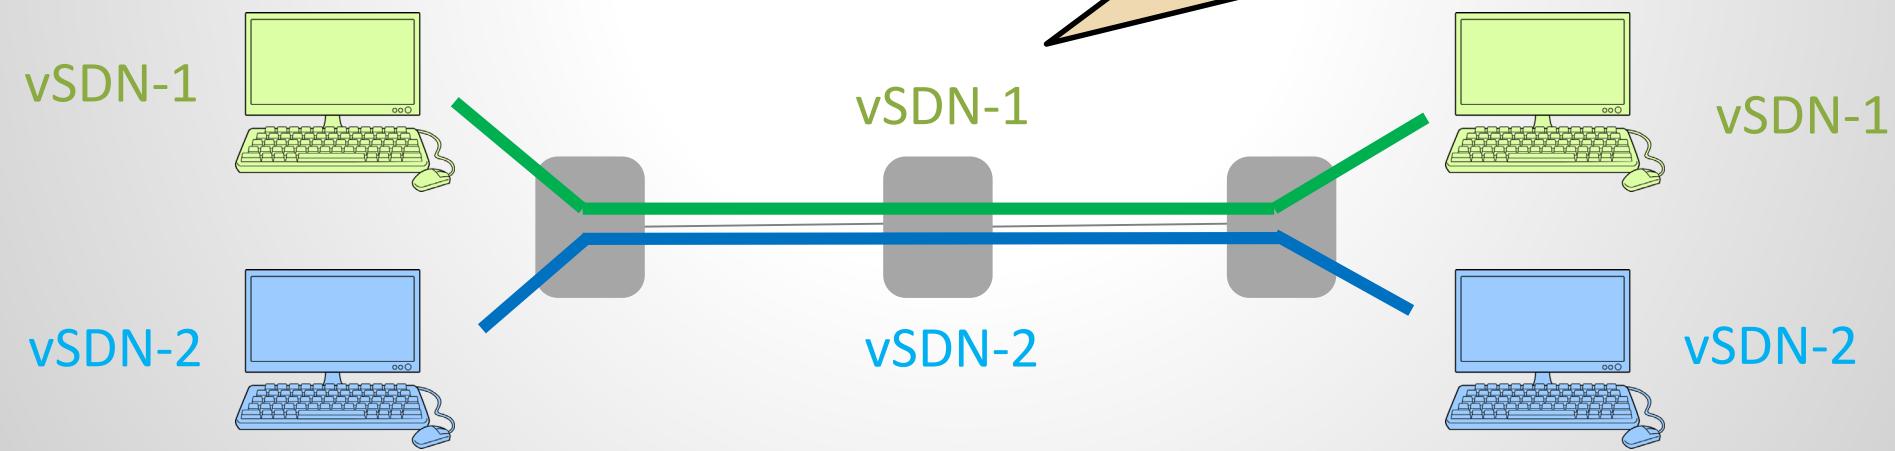
In theory land: bandwidth reservation for virtual networks = predictable performance



# Models: Mind the Gap!

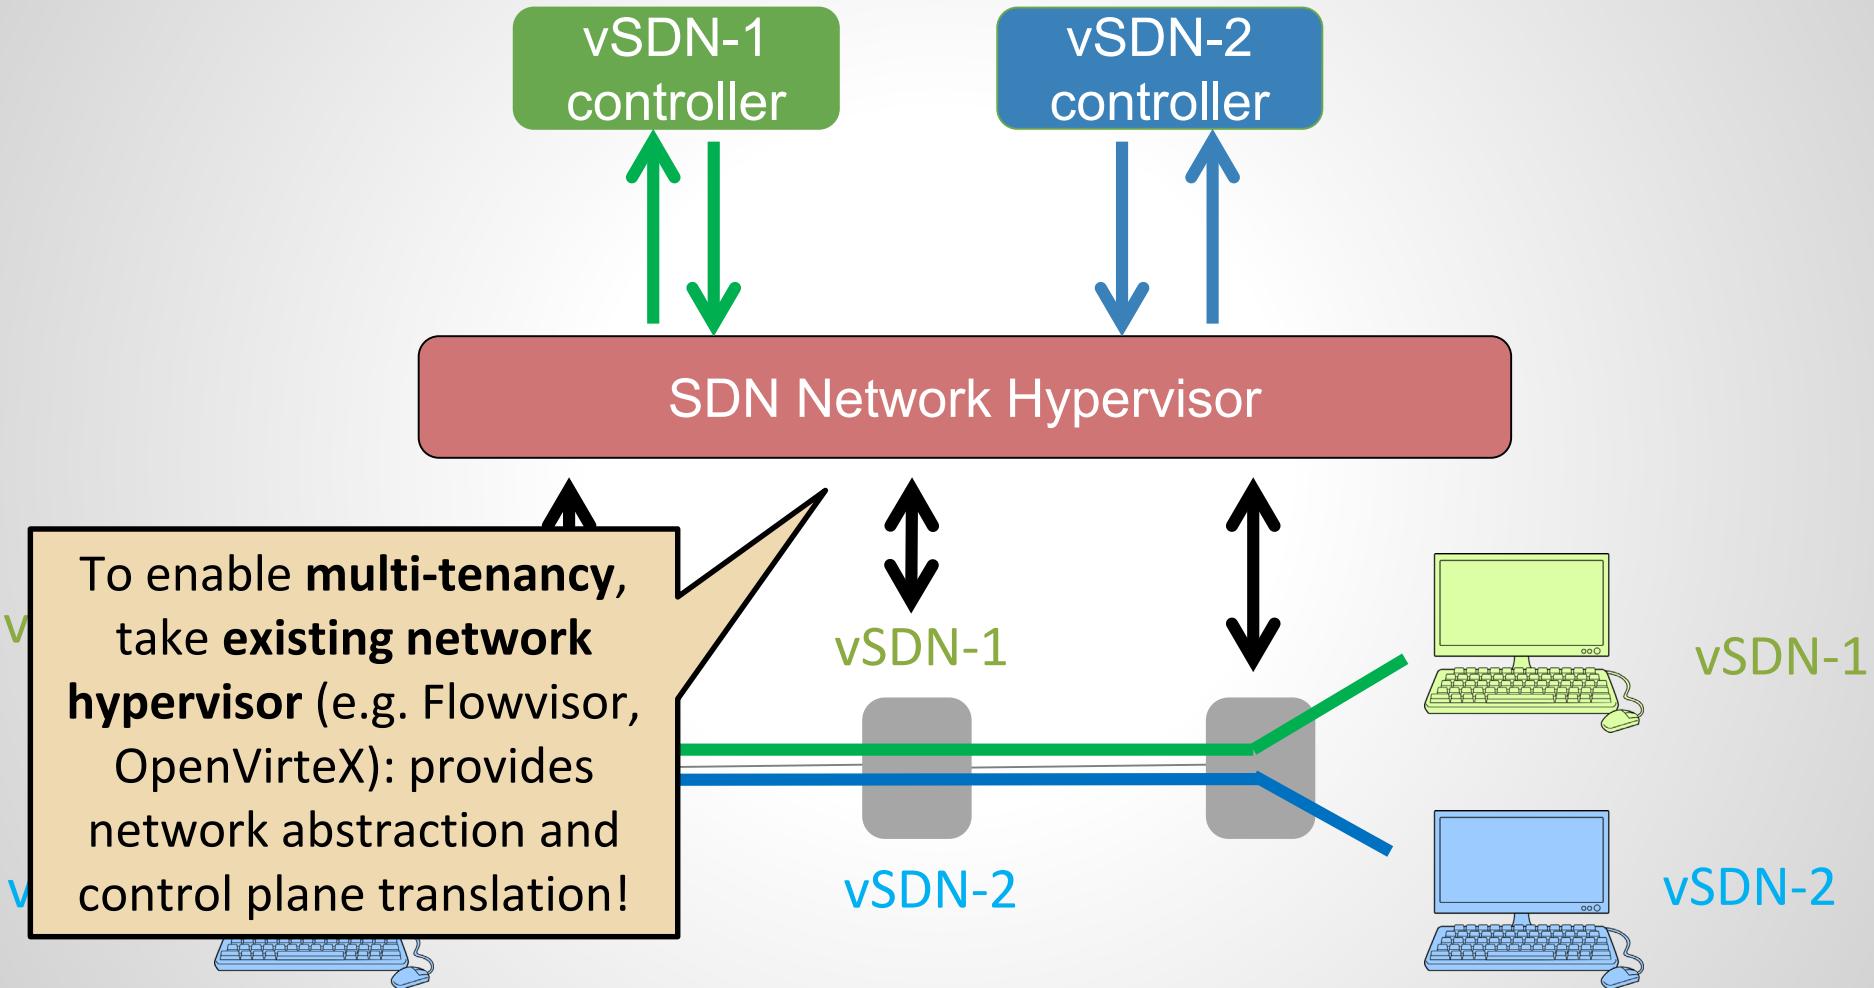
Consider: 2 SDN-based  
**virtual networks (vSDNs)**  
sharing physical resources!

Assume: perfect  
performance isolation on  
the network!



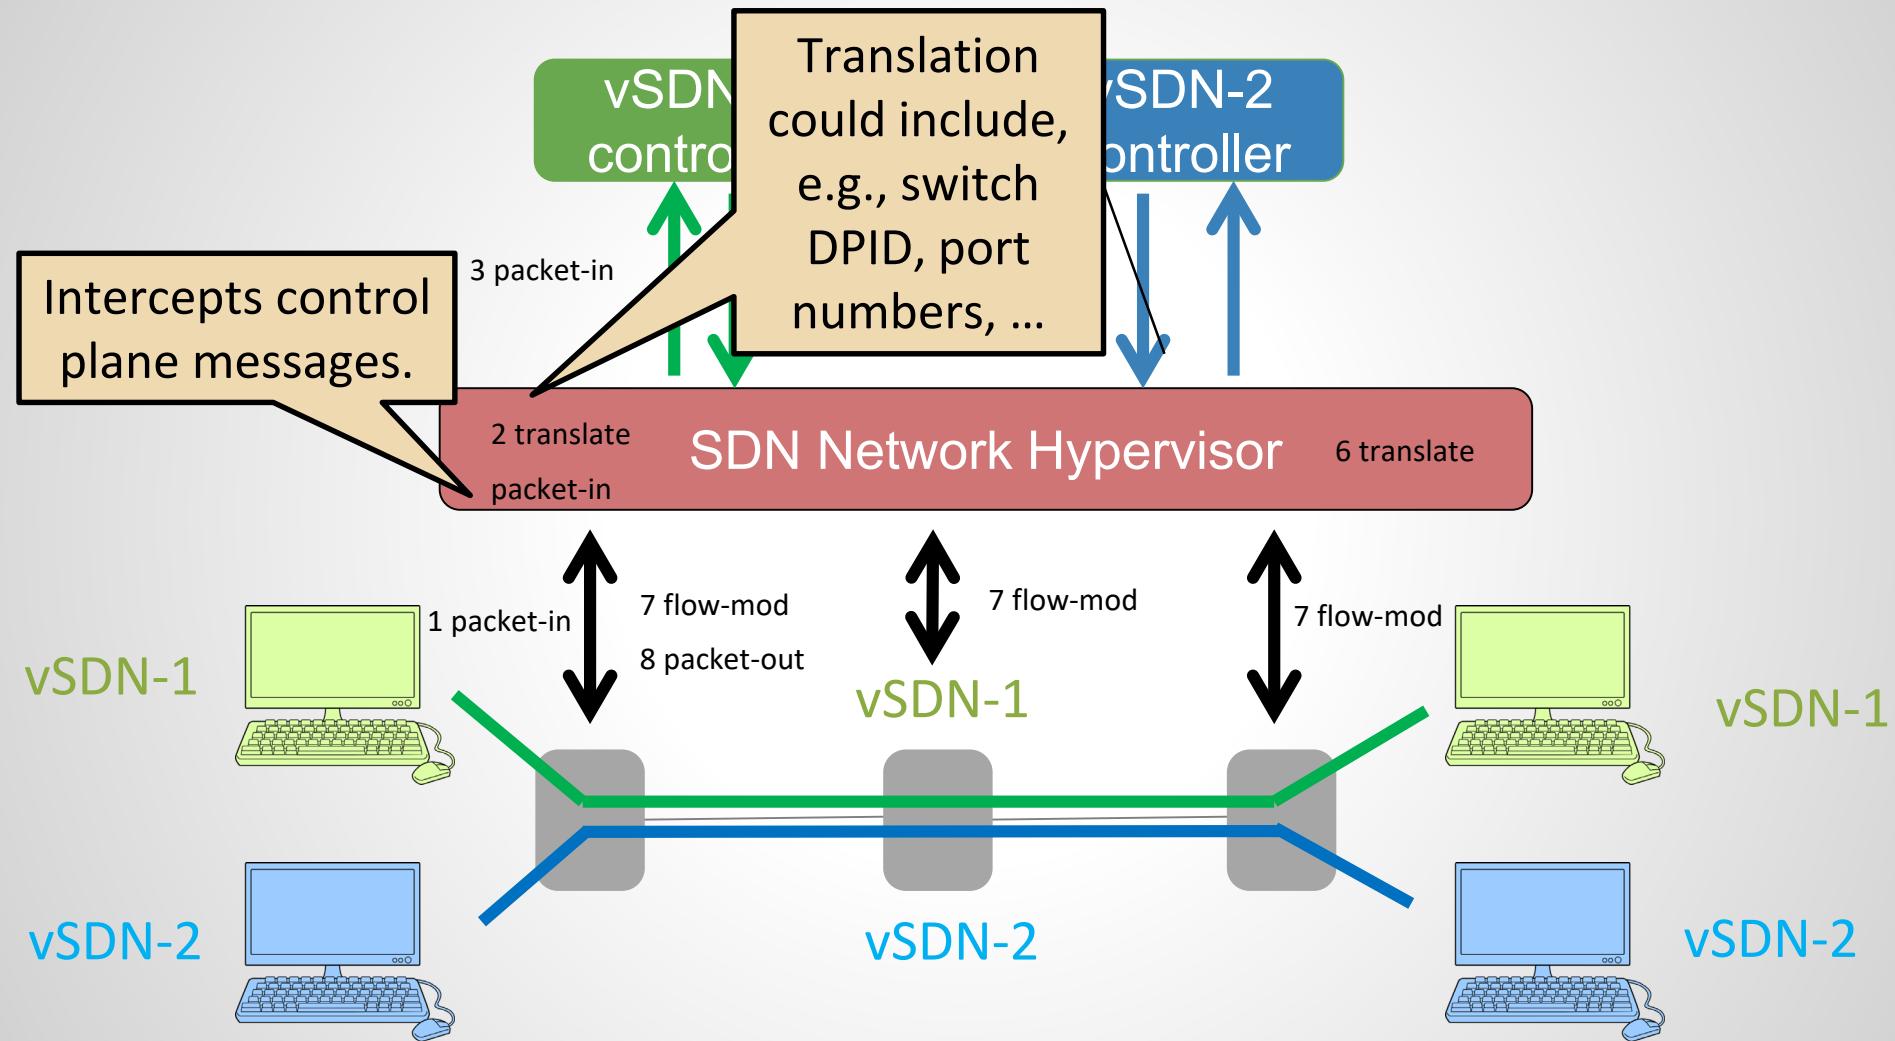
Realization: Virtual networks based on SDN

# Models: Mind the Gap!



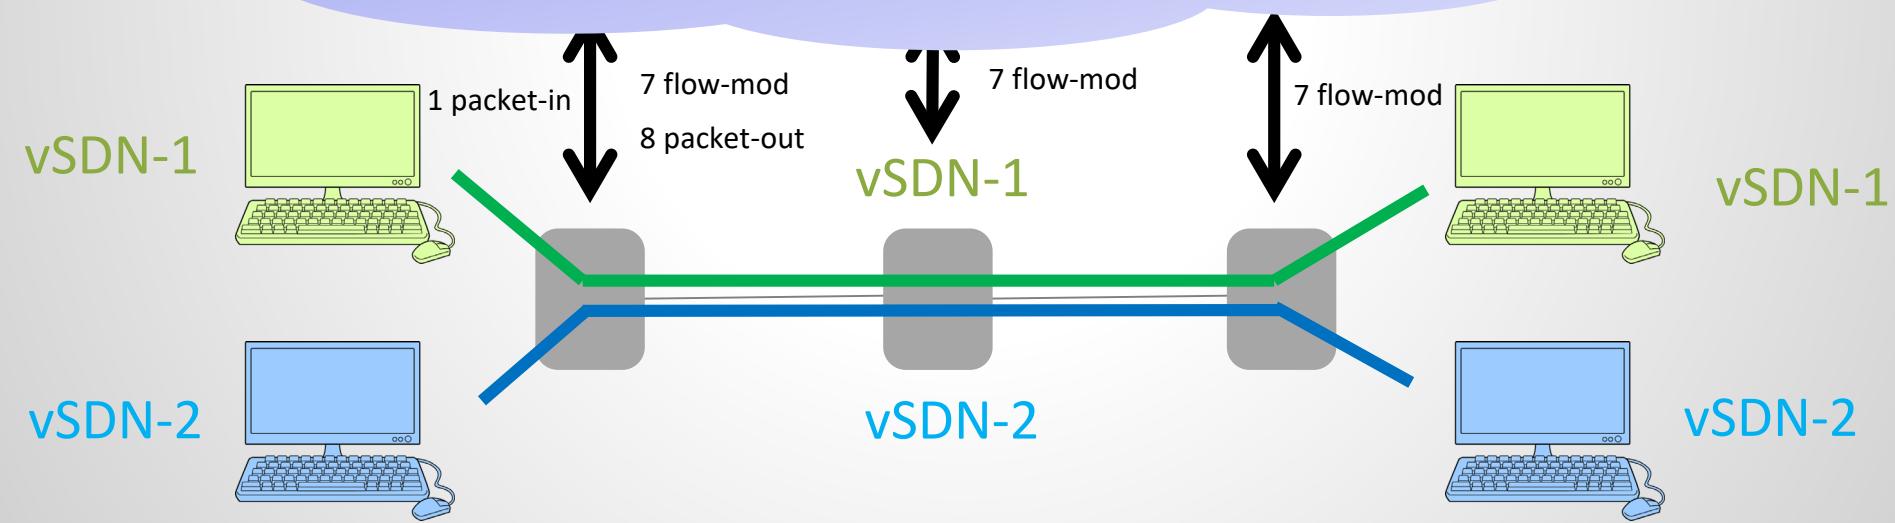
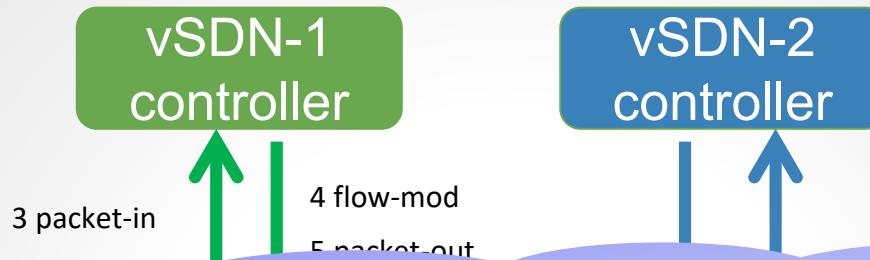
Realization: Virtual networks based on SDN

# Models: Mind the Gap!



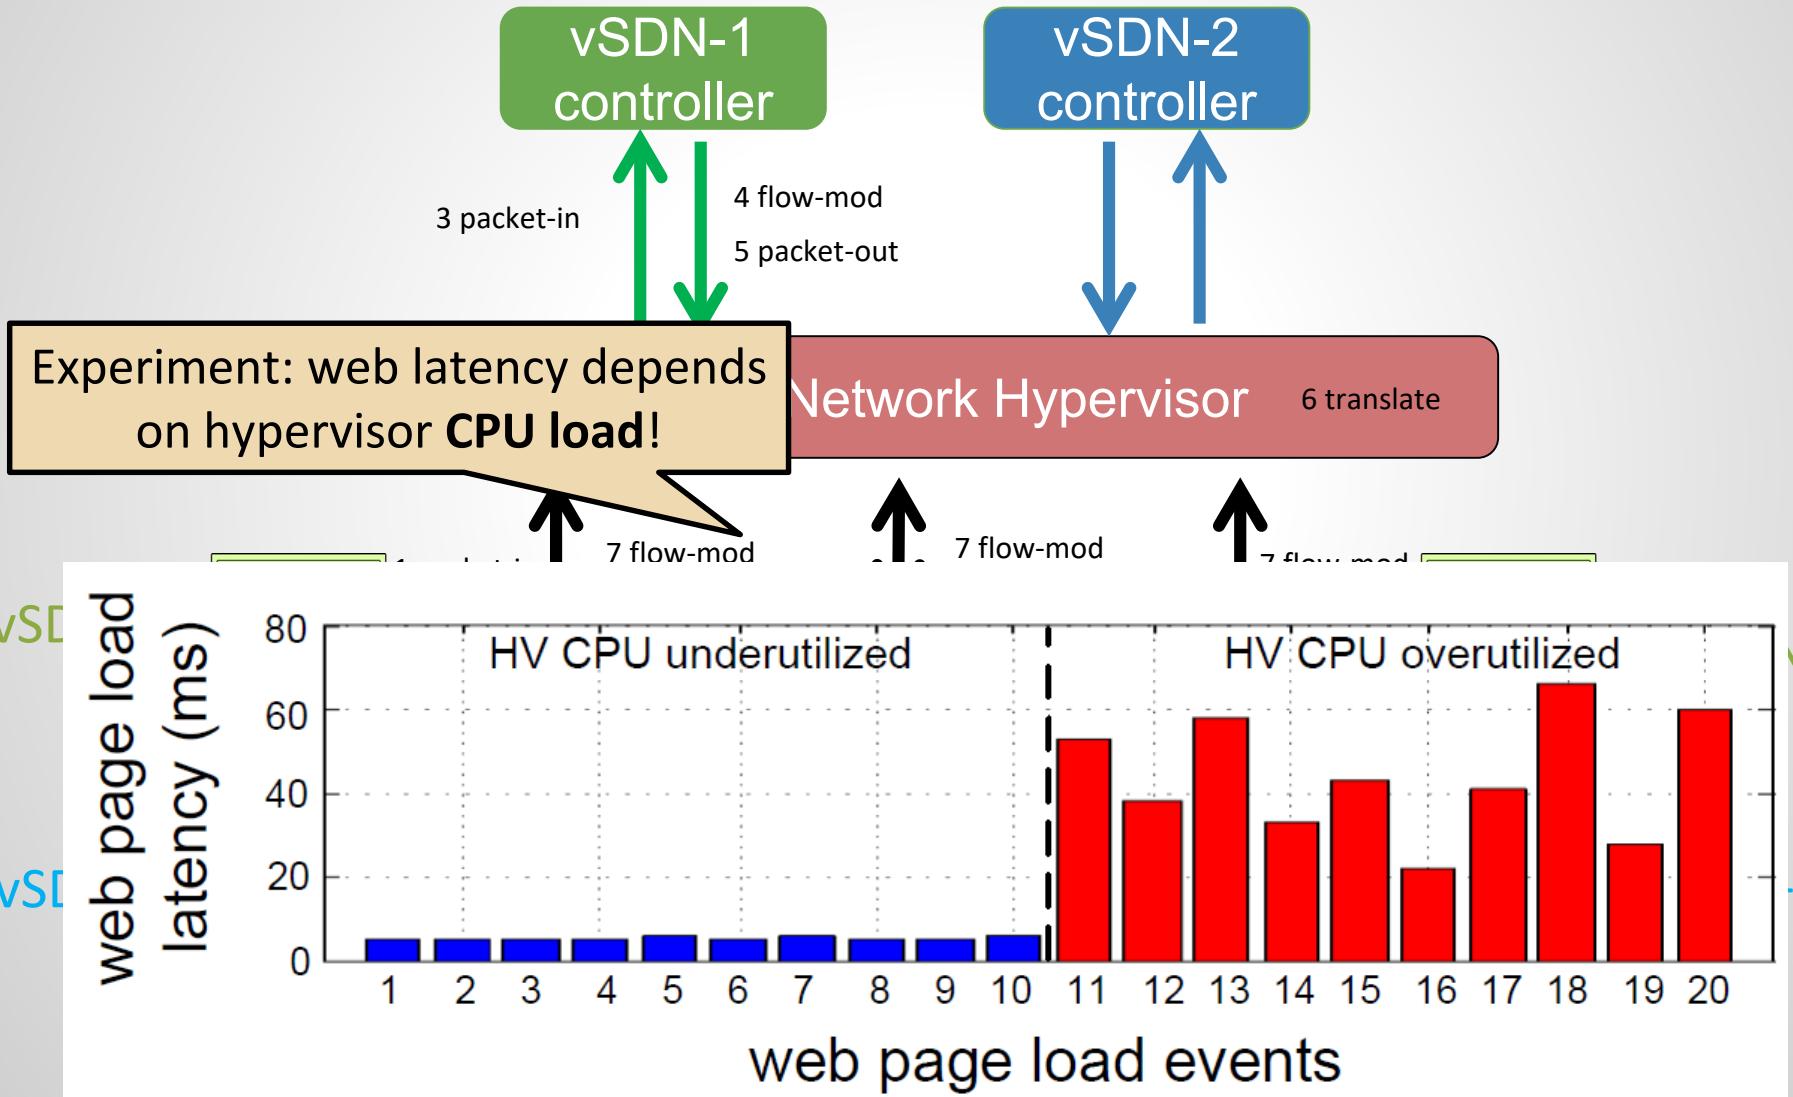
Realization: Virtual networks based on SDN

# Models: Mind the Gap!

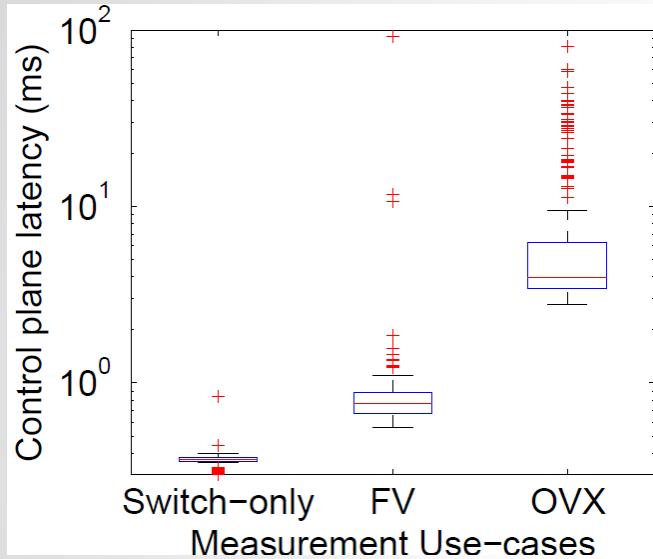


Realization: Virtual networks based on SDN

# Models: Mind the Gap!

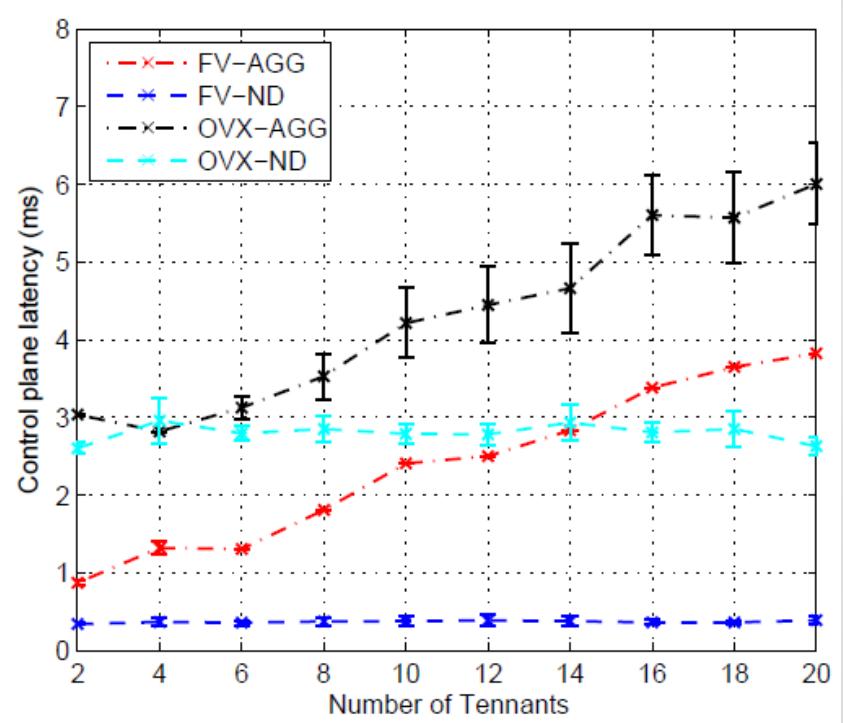


# The Many Faces of Performance Interference



**Performance also depends  
on hypervisor type...**  
*(multithreaded or not, which version  
of Nagle's algorithm, etc.)*

**... number of tenants...**



**It's complex!**

# Invitation: A Roadtrip Through The Opportunities and Challenges of Network Isolation

- ❑ Opportunities
  - ❑ Algorithmic opportunities
  - ❑ Technological opportunities
- ❑ Challenges
  - ❑ Algorithmic challenges
  - ❑ Modelling challenges
  - ❑ **Security challenges**
- ❑ A perspective how AI can improve slicing efficiency and security



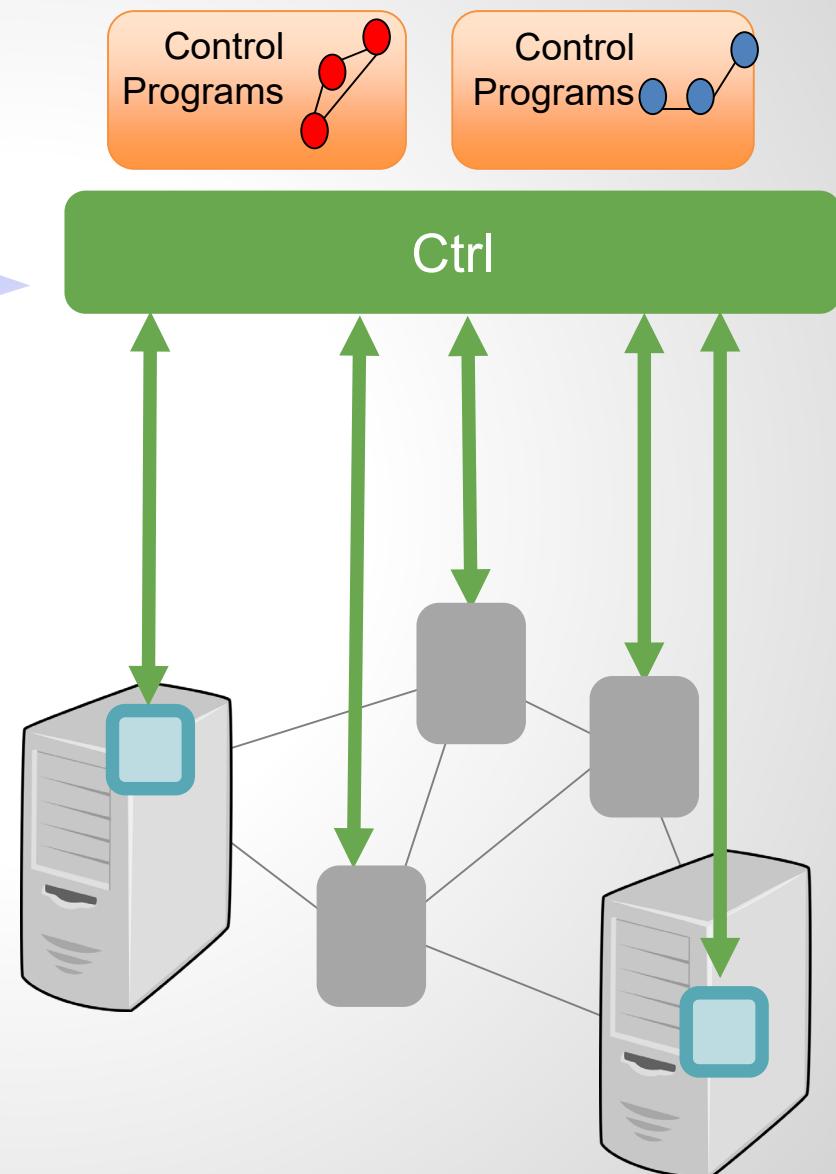
Road map 1927: Arizona and New Mexico

# Programmable and Virtualized Networks

Networked systems:

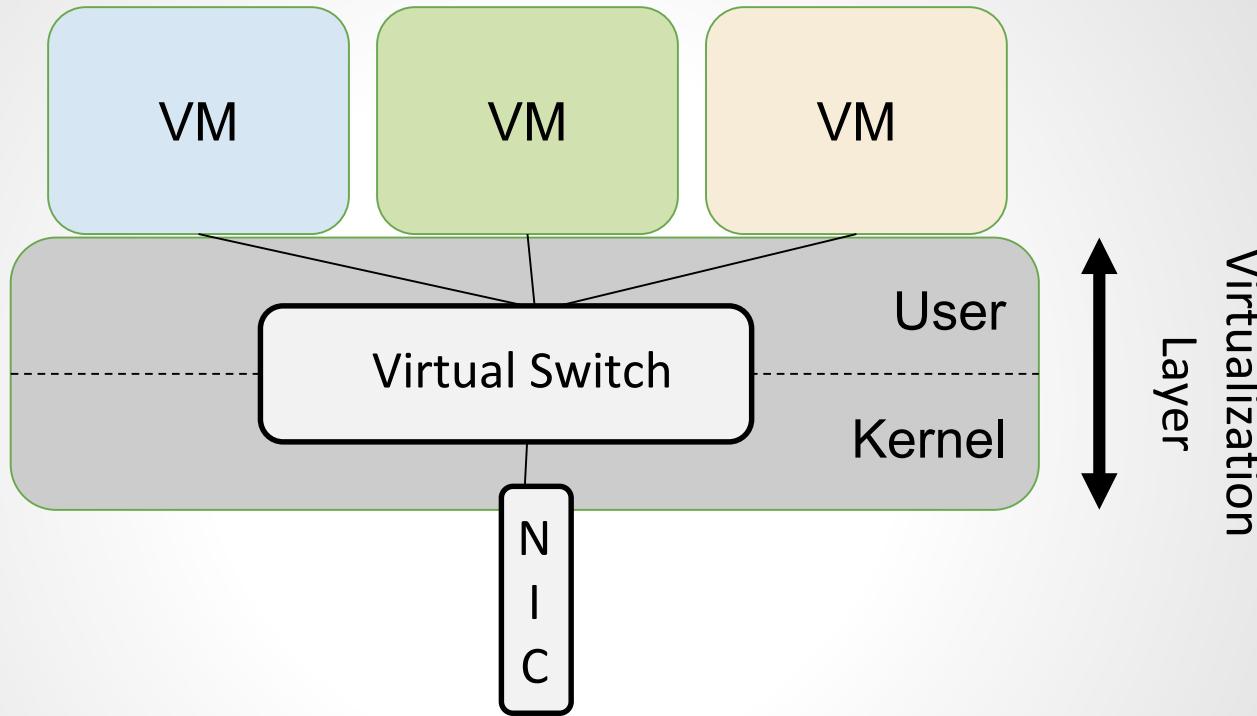
Increasingly  
centralized

Increasingly  
virtualized



Challenge: security!

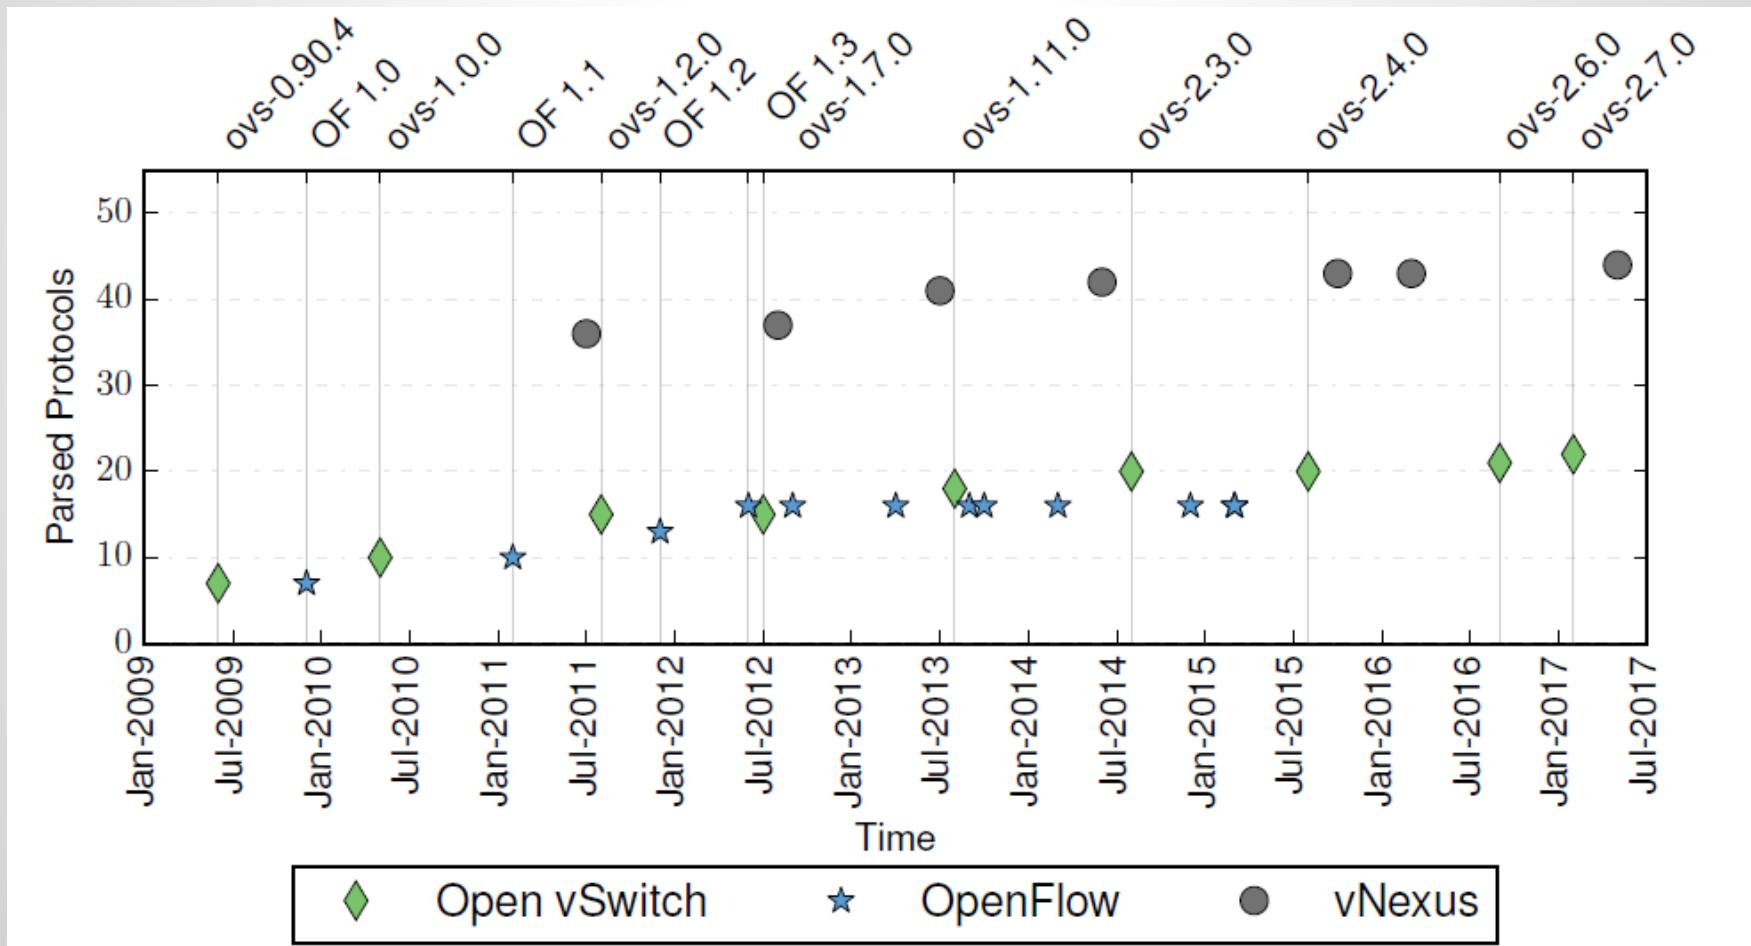
# Potential New Attack Surface: Virtual Switches



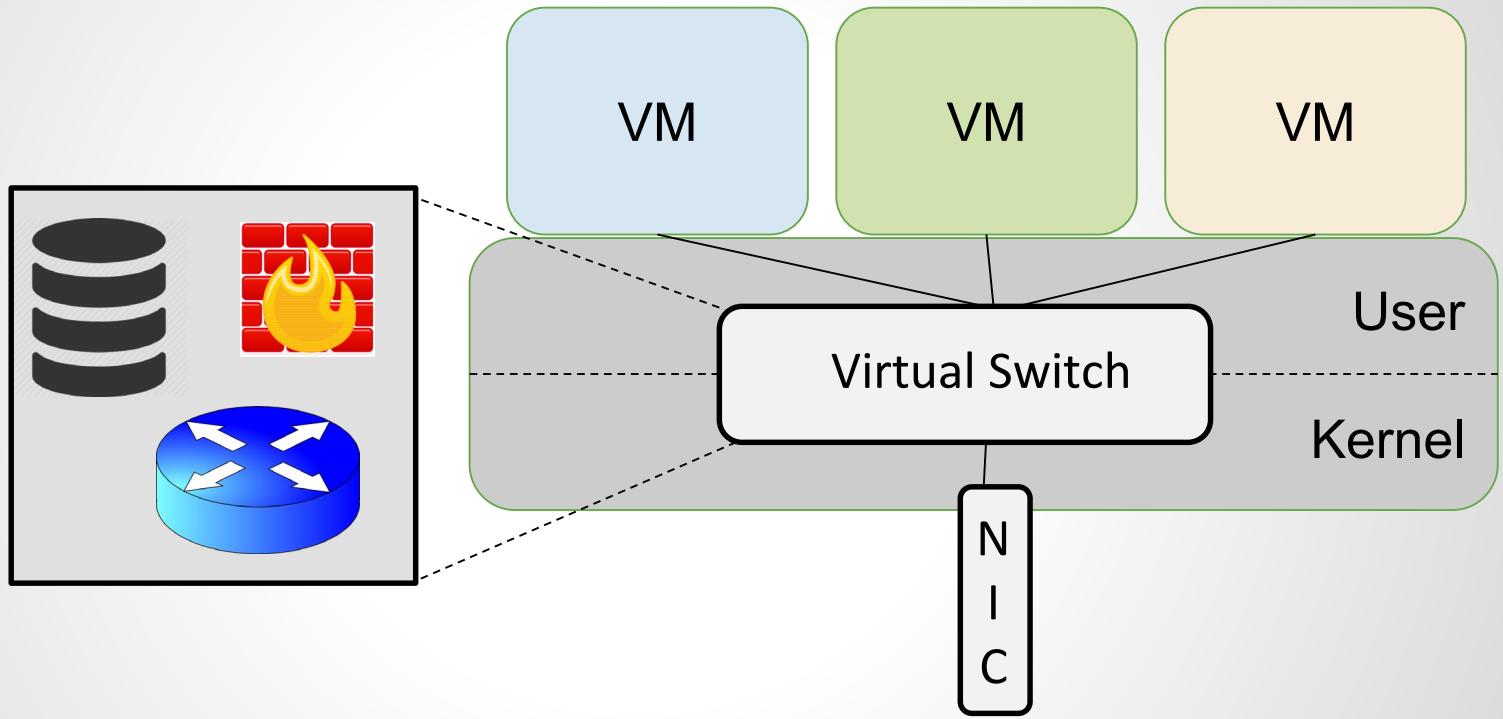
Virtual switches reside in the **server's virtualization layer** (e.g., Xen's Dom0). Goal: provide connectivity and isolation.

# Increasing Complexity: *# Parsed Protocols*

Number of parsed high-level protocols constantly increases:



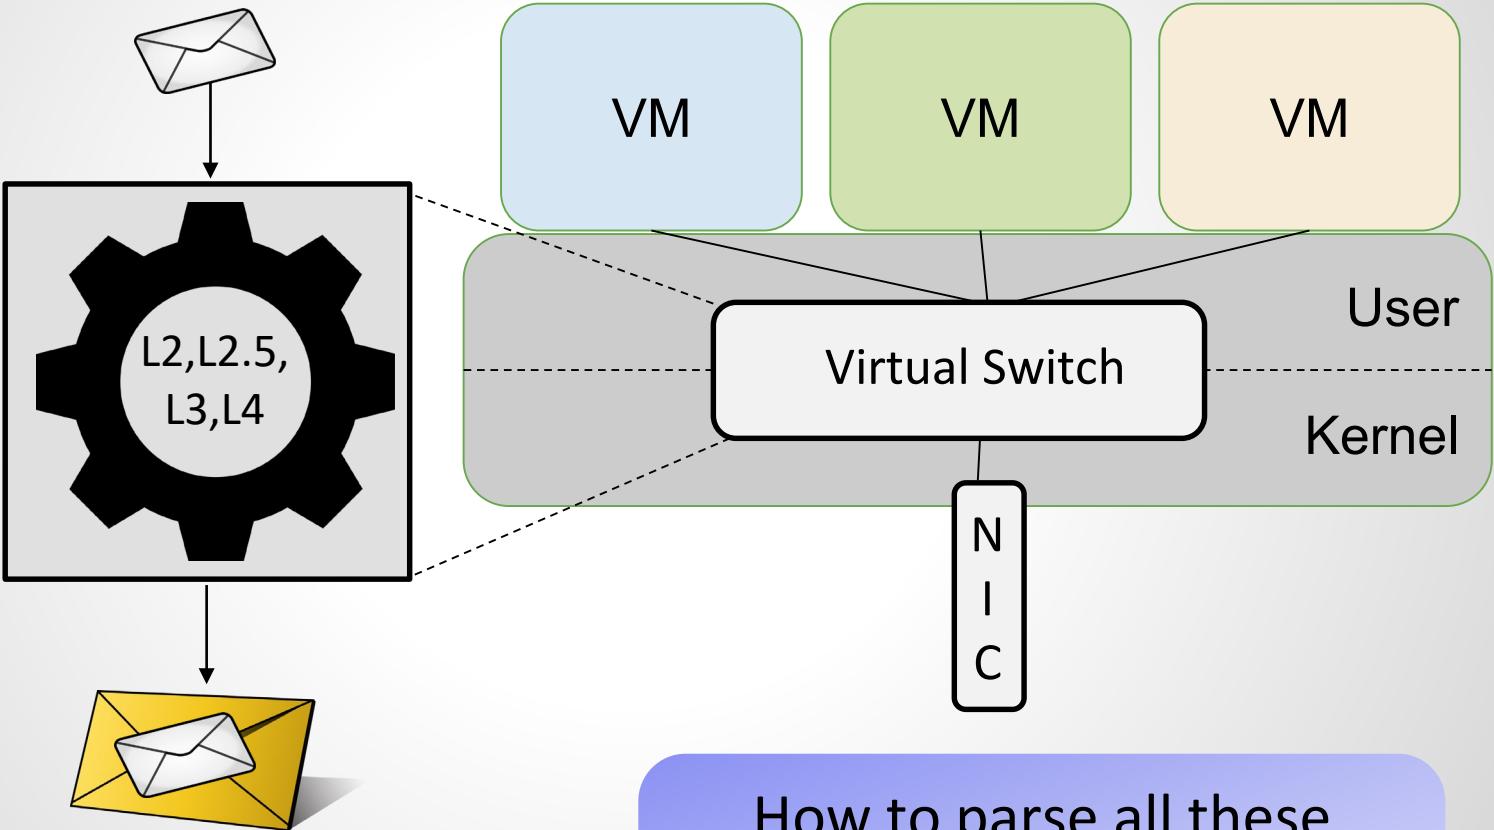
# Increasing Complexity: *Introduction of middlebox functionality*



**Increasing workloads** and advancements in network virtualization drive virtual switches to **implement middlebox functions** such as load-balancing, DPI, firewalls, etc.

# Increasing Complexity: *Unified Packet Parsing*

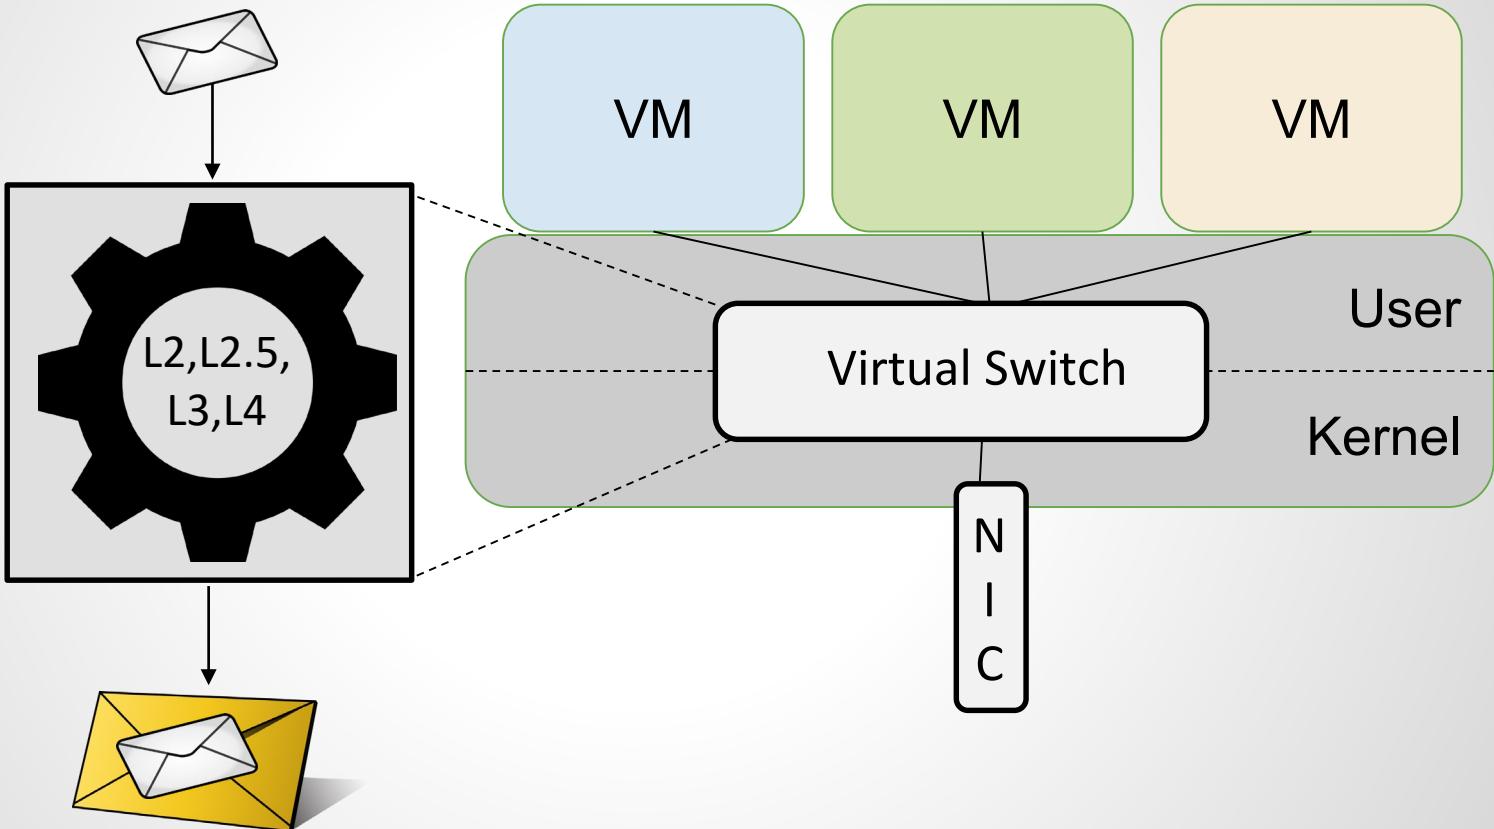
Ethernet  
LLC  
VLAN  
MPLS  
IPv4  
ICMPv4  
TCP  
UDP  
ARP  
SCTP  
IPv6  
ICMPv6  
IPv6 ND  
GRE  
LISP  
VXLAN  
PBB  
IPv6 EXT HDR  
TUNNEL-ID  
IPv6 ND  
IPv6 EXT HDR  
IPv6HOPOPTS  
IPv6ROUTING  
IPv6Fragment  
IPv6DESTOPT  
IPv6ESP  
IPv6 AH  
RARP  
IGMP



How to parse all these  
protocols without lowering  
forwarding performance?!

# Increasing Complexity: *Unified Packet Parsing*

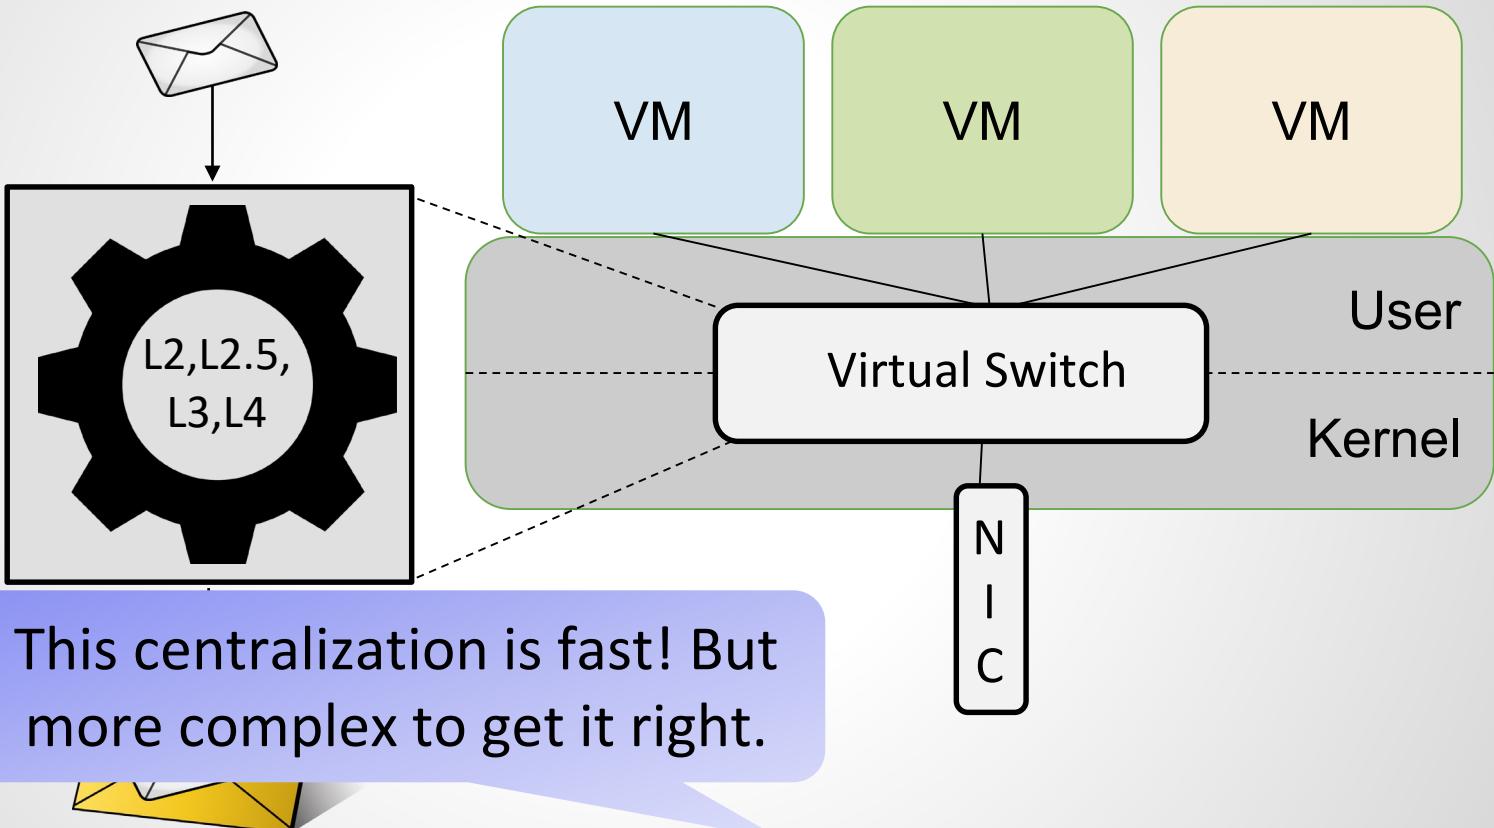
Ethernet  
LLC  
VLAN  
MPLS  
IPv4  
ICMPv4  
TCP  
UDP  
ARP  
SCTP  
IPv6  
ICMPv6  
IPv6 ND  
GRE  
LISP  
VXLAN  
PBB  
IPv6 EXT HDR  
TUNNEL-ID  
IPv6 ND  
IPv6 EXT HDR  
IPv6HOPOPTS  
IPv6ROUTING  
IPv6Fragment  
IPv6DESTOPT  
IPv6ESP  
IPv6 AH  
RARP  
IGMP



Unified packet parsing allows parse more and more protocols efficiently: **in a single pass!**

# Increasing Complexity: *Unified Packet Parsing*

Ethernet  
LLC  
VLAN  
MPLS  
IPv4  
ICMPv4  
TCP  
UDP  
ARP  
SCTP  
IPv6  
ICMPv6  
IPv6 ND  
GRE  
LISP  
VXLAN  
PBB  
IPv6 EXT HDR  
TUNNEL-ID  
IPv6 ND  
IPv6 EXT HDR  
IPv6HOPOPTS  
IPv6ROUTING  
IPv6Fragment  
IPv6DESOPT  
IPv6ESP  
IPv6 AH  
RARP  
IGMP



Unified packet parsing allows parse more and more protocols efficiently: **in a single pass!**

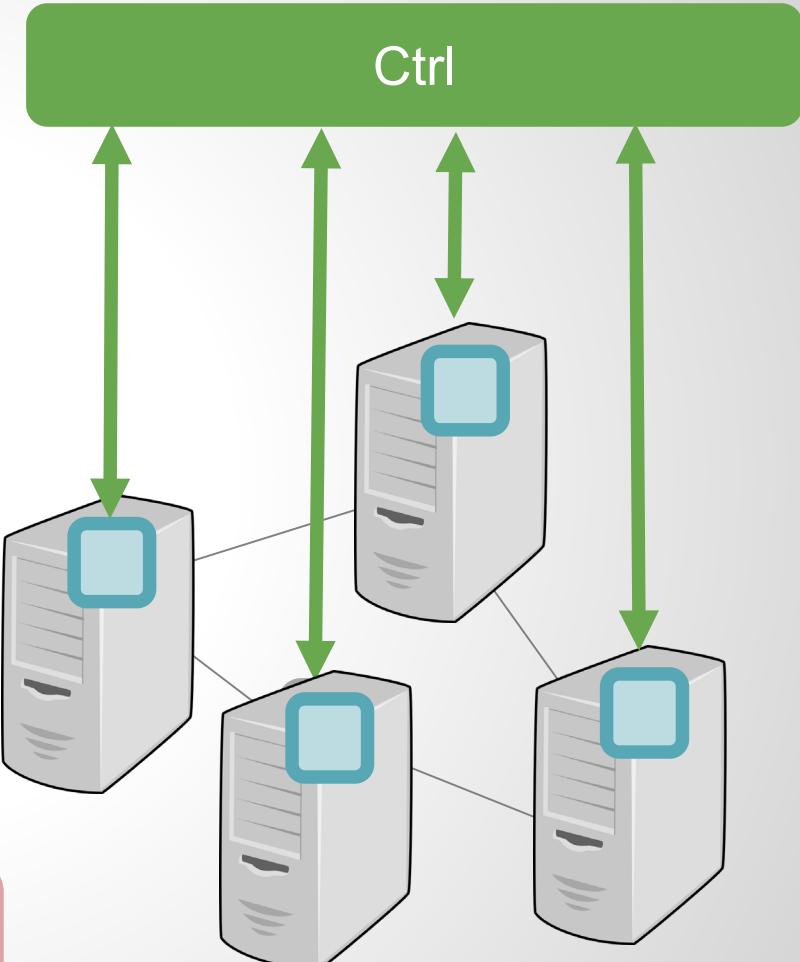
# Complexity: The Enemy of Security!

- Data plane security not well-explored (in general, not only virtualized): most security research on control plane

- Two conjectures:

1. Virtual switches increase the attack surface.

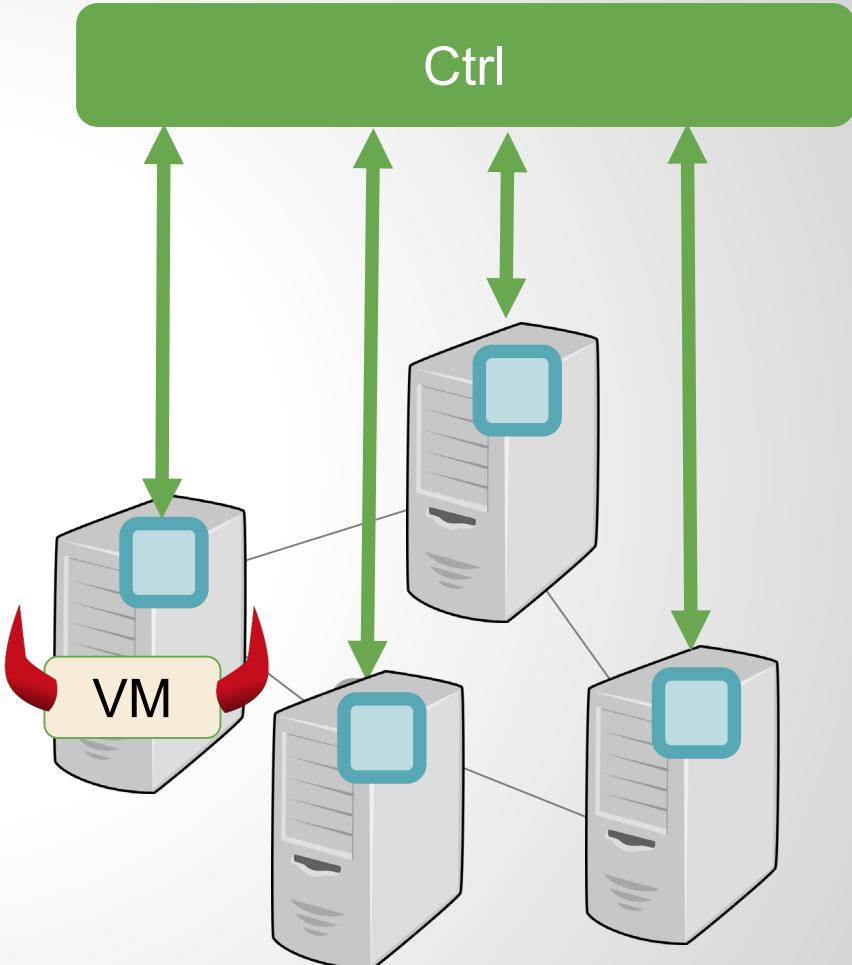
2. Impact of attack larger than with traditional data planes.



# The Attack Surface: Closer...

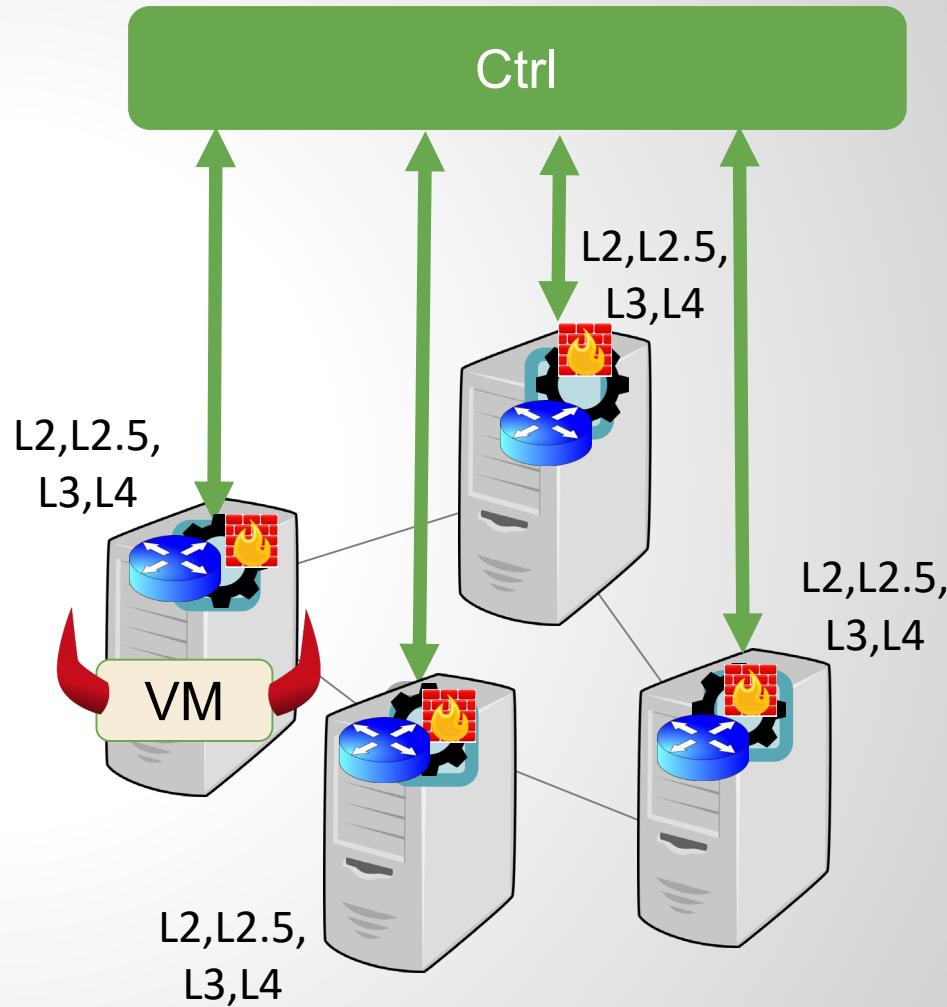
Attack surface **becomes closer**:

- ❑ Packet parser typically integrated into the code base of virtual switch
- ❑ **First component** of the virtual switch to process network packets it receives from the network interface
- ❑ May process **attacker-controlled packets!**



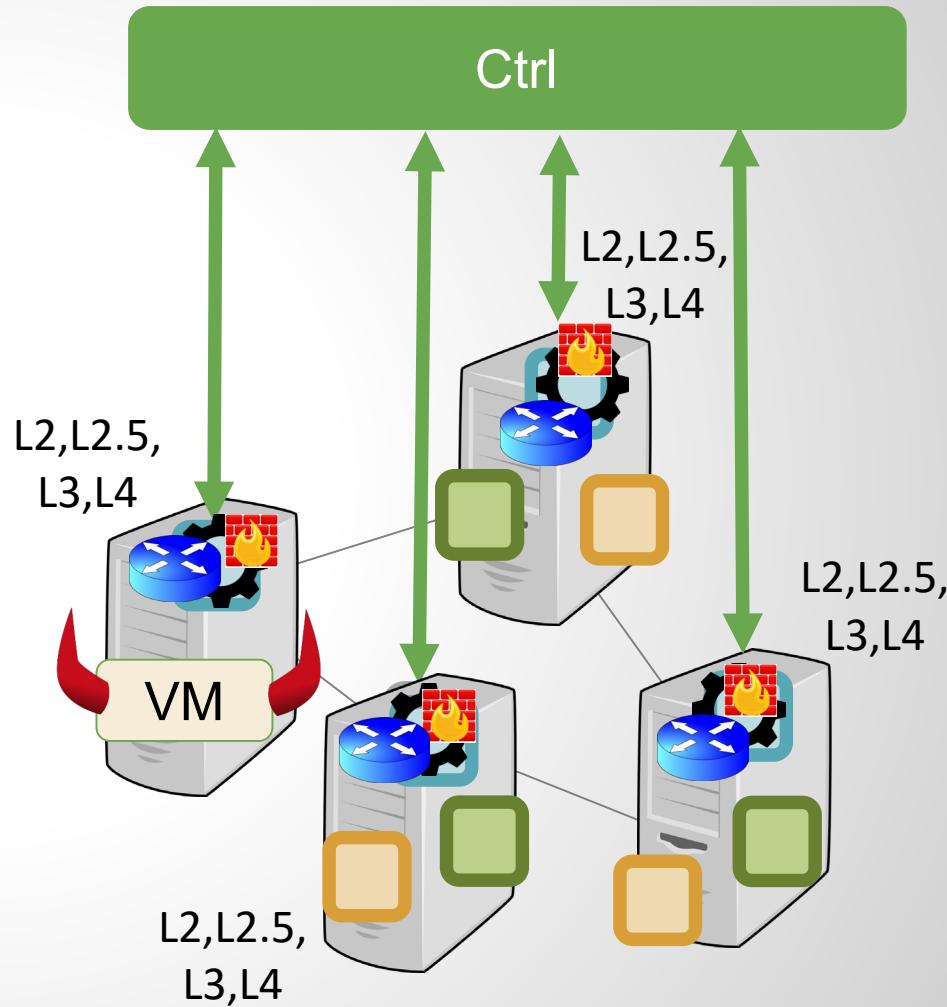
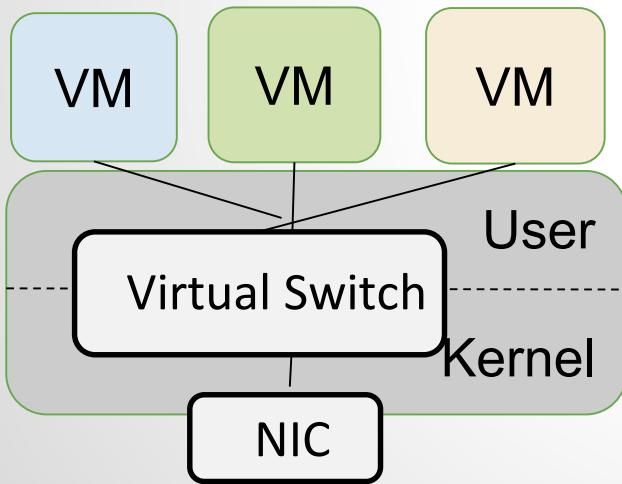
# The Attack Surface: ... More Complex ...

Ethernet	PBB
LLC	IPv6 EXT HDR
VLAN	TUNNEL-ID
MPLS	IPv6 ND
IPv4	IPv6 EXT HDR
ICMPv4	IPv6HOPOPTS
TCP	IPv6ROUTING
UDP	IPv6Fragment
ARP	IPv6DESTOPT
SCTP	IPv6ESP
IPv6	IPv6 AH
ICMPv6	RARP
IPv6 ND	IGMP
GRE	
LISP	
VXLAN	



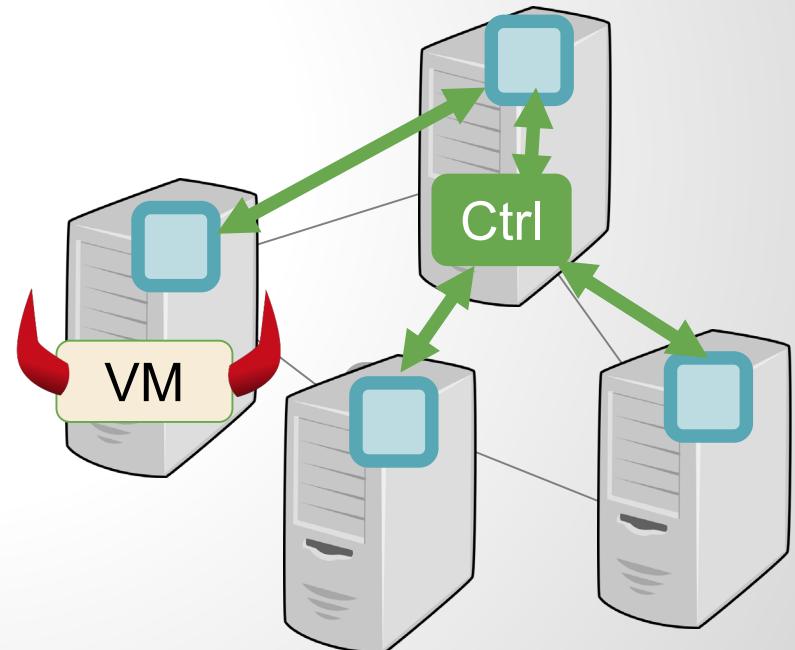
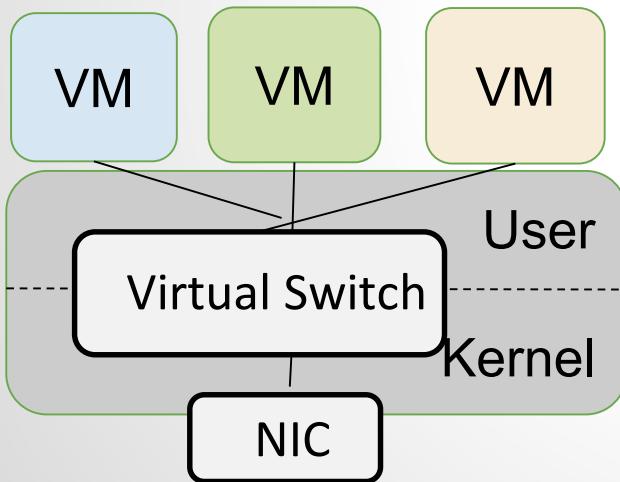
# ... Elevated Privileges and Collocation ...

- Collocated (at least partially) with **hypervisor's Dom0 kernel space**, guest VMs, image management, block storage, identity management, ...



# ... Elevated Privileges and Collocation ...

- ❑ Collocated (at least partially) with **hypervisor's Dom0 kernel** space, guest VMs, image management, block storage, identity management, ...

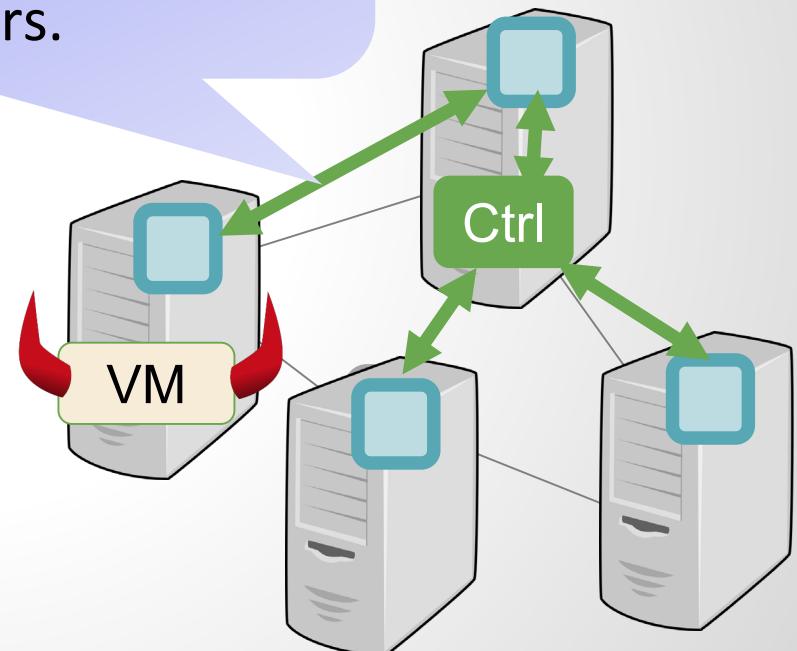
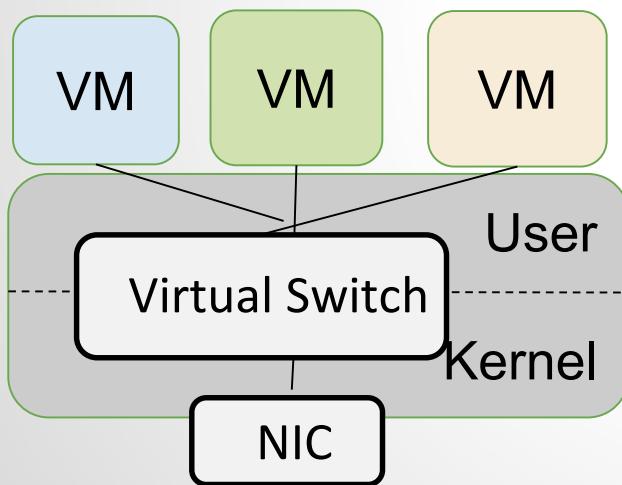


- ❑ ... the **controller** itself.

# ... Centralization ...

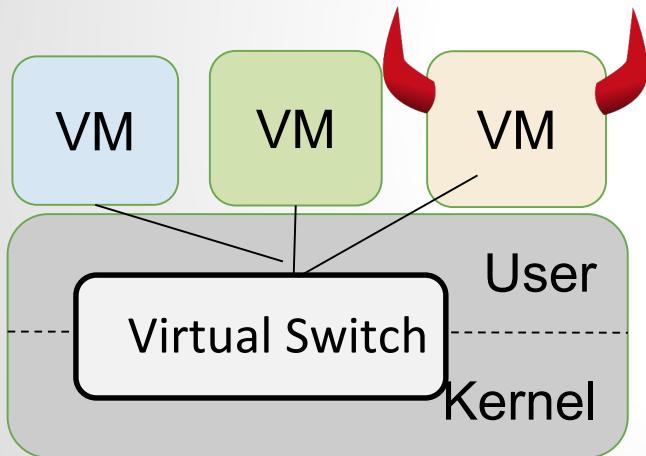
- ❑ Collocate management services with **hypervisor** (in host space), get access to management interface, identity management, ...

Available communication channels to (SDN/Openstack) controller!  
Controller needs to be reachable from all servers.



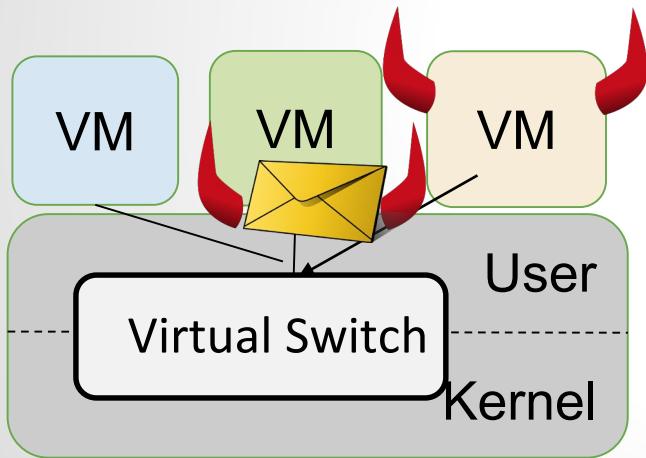
- ❑ ... the **controller** itself.

# Larger Impact: Case Study OVS



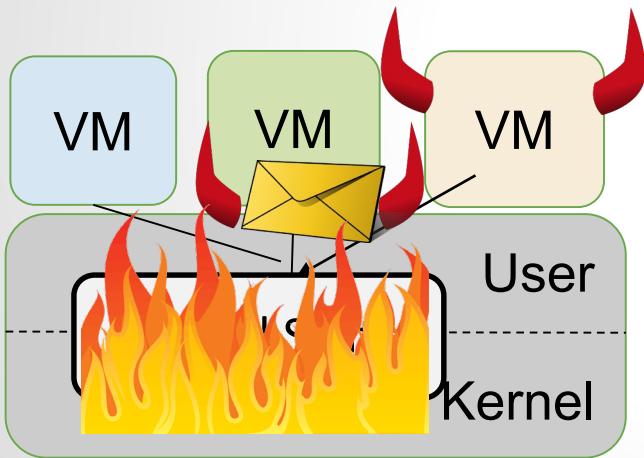
1. Rent a VM in the cloud (**cheap**)

# Larger Impact: Case Study OVS



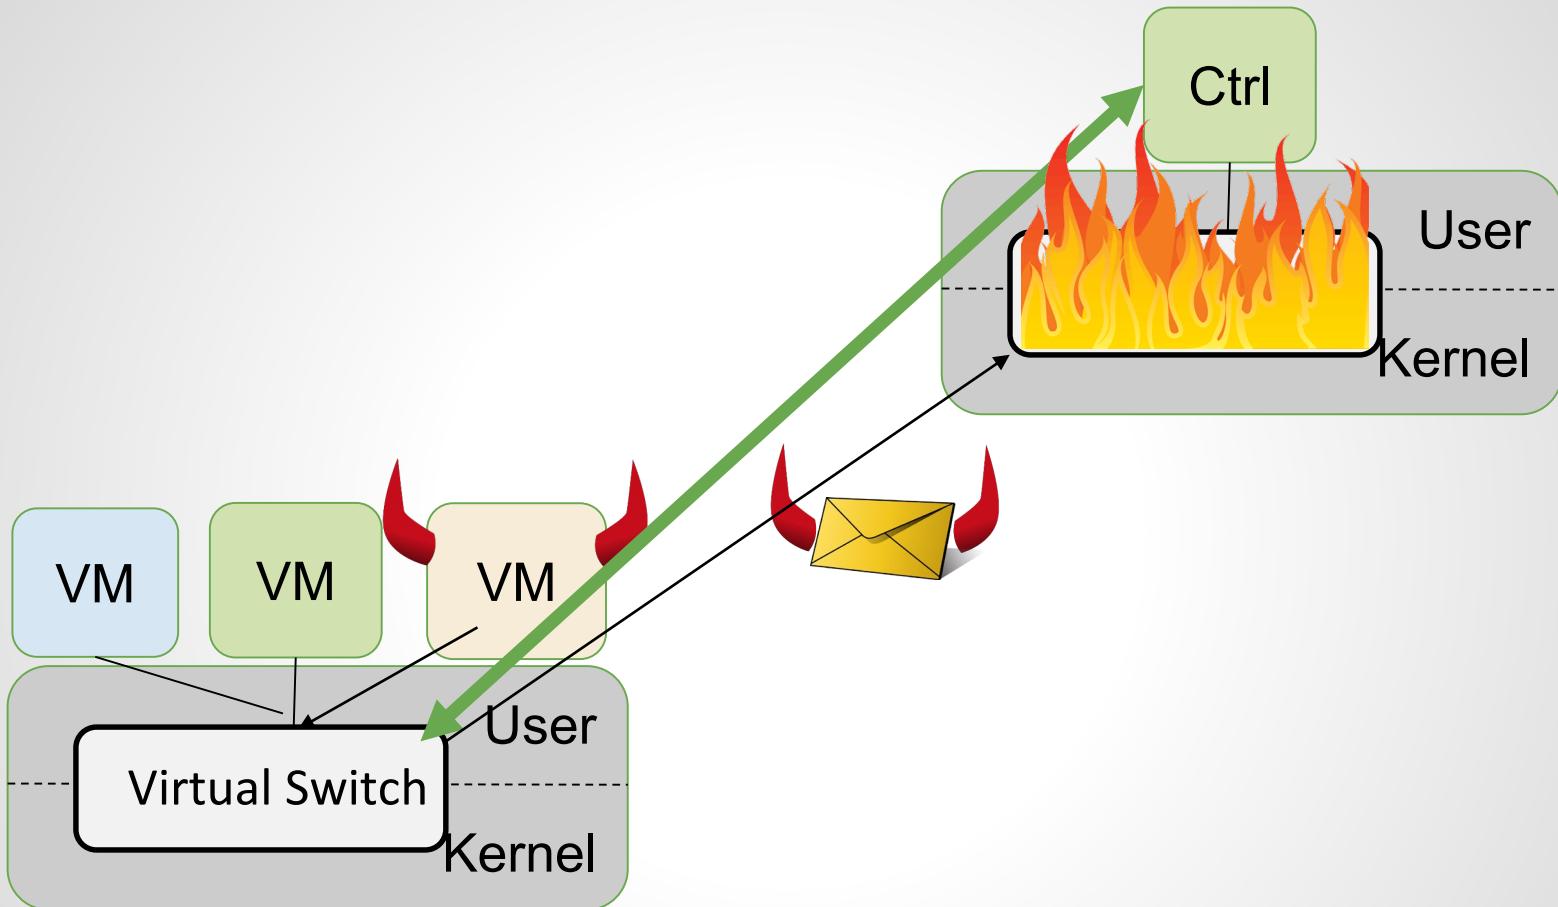
2. Send **malformed MPLS packet** to virtual switch (**unified parser** parses label stack packet **beyond the threshold**)

# Larger Impact: Case Study OVS



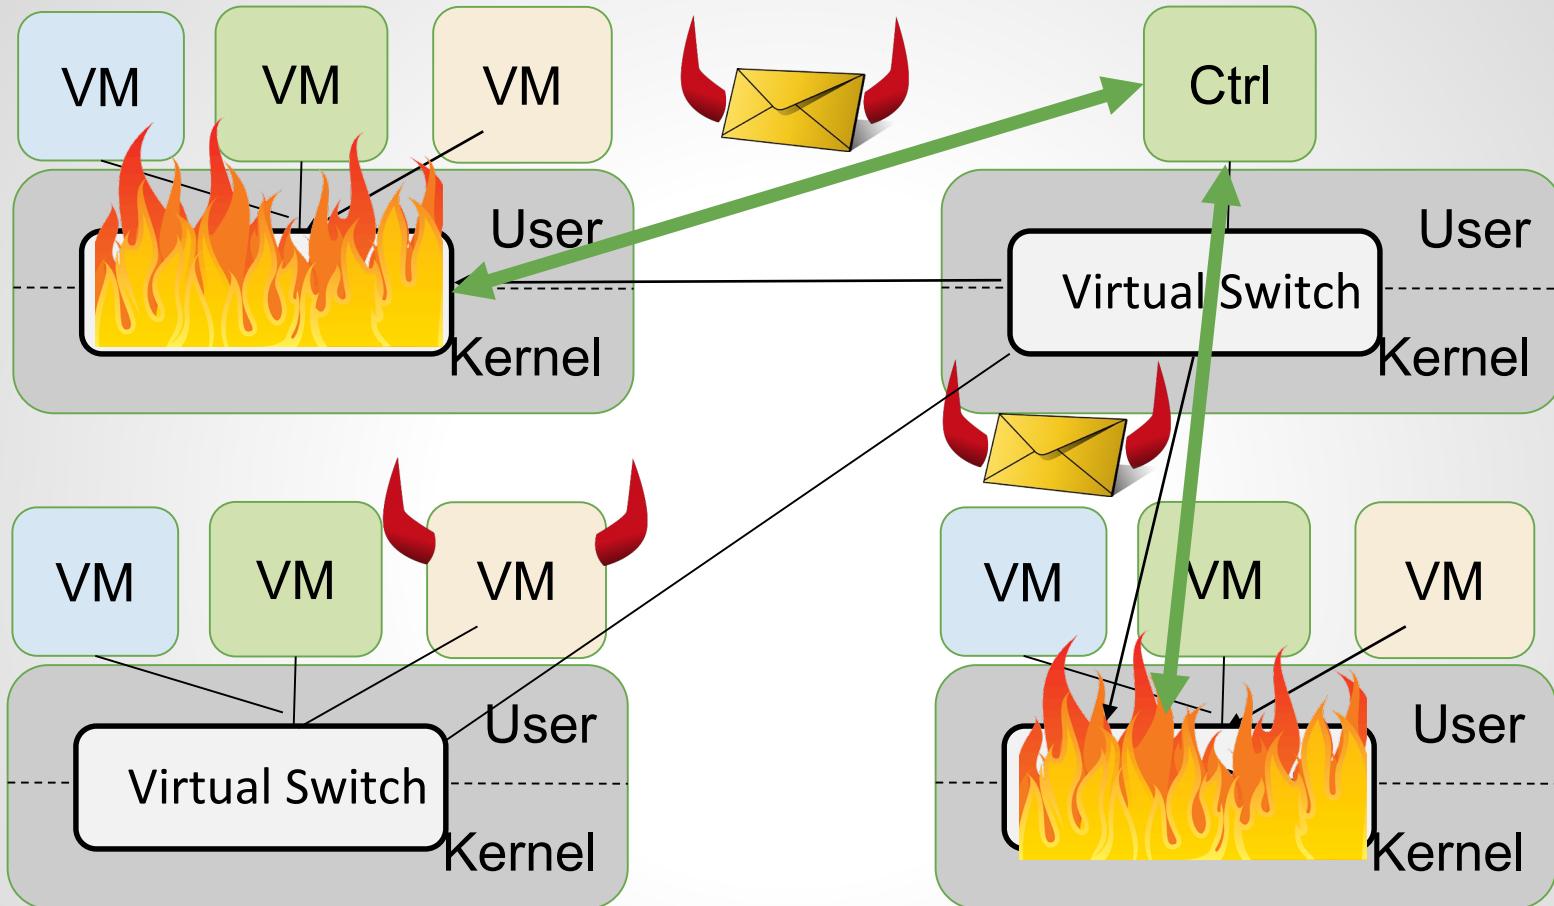
3. **Stack buffer overflow** in (unified) MPLS parsing code:  
enables **remote code execution**

# Larger Impact: Case Study OVS



4. Send malformed packet to server (virtual switch) where controller is located (use **existing communication channel**)

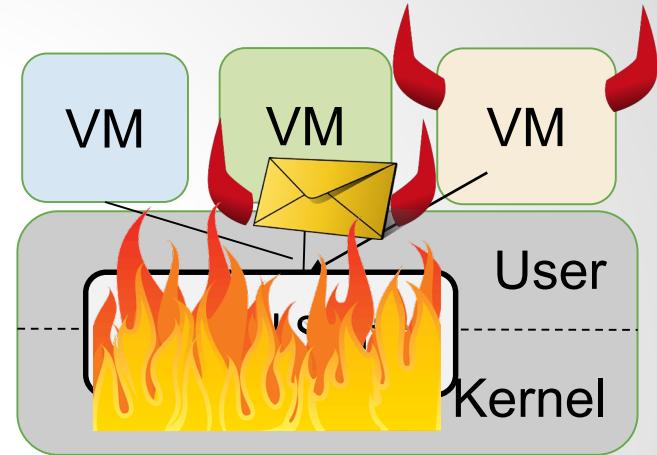
# Larger Impact: Case Study OVS



5. Spread

# A New Threat Model

- ❑ Limited skills required
  - ❑ Use standard fuzzer to find crashes
  - ❑ Construct malformed packet
  - ❑ Build ROP chain
- ❑ Limited resources
  - ❑ rent a VM in the cloud
- ❑ No physical access needed

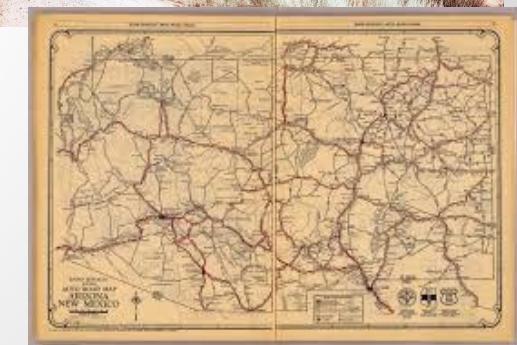


No need to be a state-level attacker to compromise the dataplane (and beyond)!

Similar problems in NFV: need even more complex parsing/processing. And are often built on top of OvS.

# Invitation: A Roadtrip Through The Opportunities and Challenges of Network Isolation

- Opportunities
  - Algorithmic opportunities
  - Technological opportunities
- Challenges
  - Modelling challenges
  - Security challenges
- A perspective how AI can improve slicing efficiency and security



Road map 1927: Arizona and New Mexico

# A Case for AI?

1. Modelling virtualized and softwarized systems is complex

- ❑ E.g., recall our SDN setup
- ❑ Often, many algorithms and parameters involved
- ❑ Wireless/radio components likely to increase complexity

2. In practice, ‘optimal’ resource sharing typically achieved with statistical multiplexing

- ❑ Requires data: the more the better the statistics and hence the efficiency

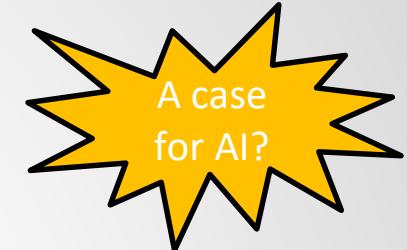
3. Resource allocation algorithms are often executed repeatedly

- ❑ E.g., routing, embedding, switching...

# A Case for AI?

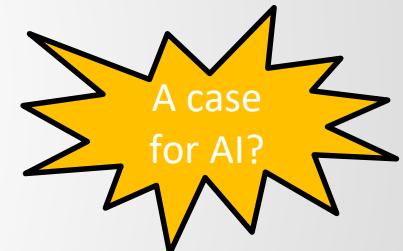
1. Modelling virtualized and softwarized systems is complex

- ❑ E.g., recall our SDN setup
- ❑ Often, many algorithms and parameters involved
- ❑ Wireless/radio components likely to increase complexity



2. In practice, 'optimal' resource sharing typically achieved with statistical multiplexing

- ❑ Requires data: the more the better the statistics and hence the efficiency



3. Resource allocation algorithms are often executed repeatedly

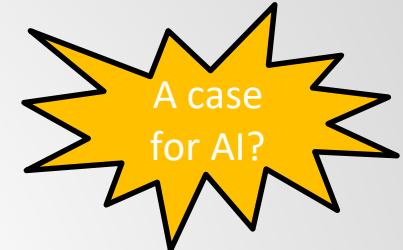
- ❑ E.g., routing, embedding, switching...



# A Case for AI?

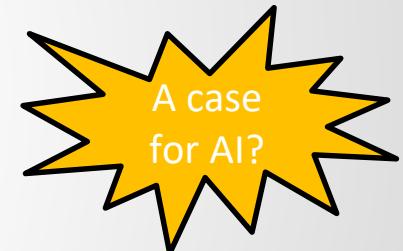
1. Modelling virtualized and softwarized systems is complex

- ❑ E.g., recall our SDN setup
- ❑ Often, many algorithms and parameters involved
- ❑ Wireless/radio components likely to increase complexity



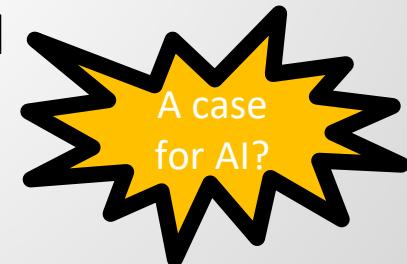
2. In practice, 'optimal' resource sharing typically achieved with statistical multiplexing

- ❑ Requires data: the more the better the statistics and hence the efficiency

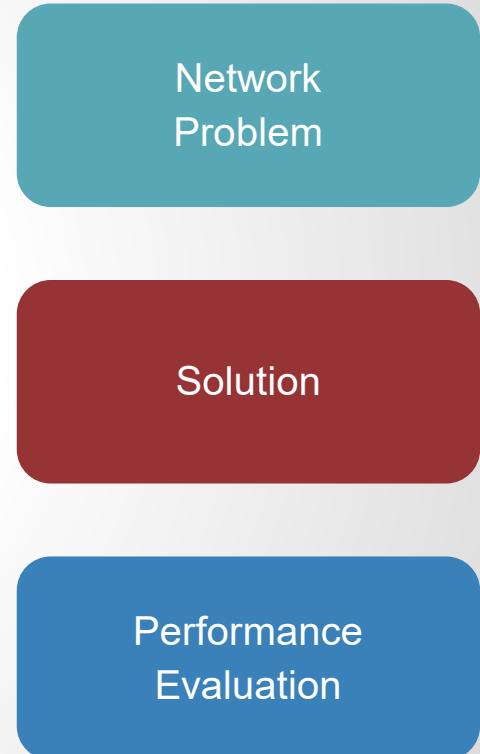
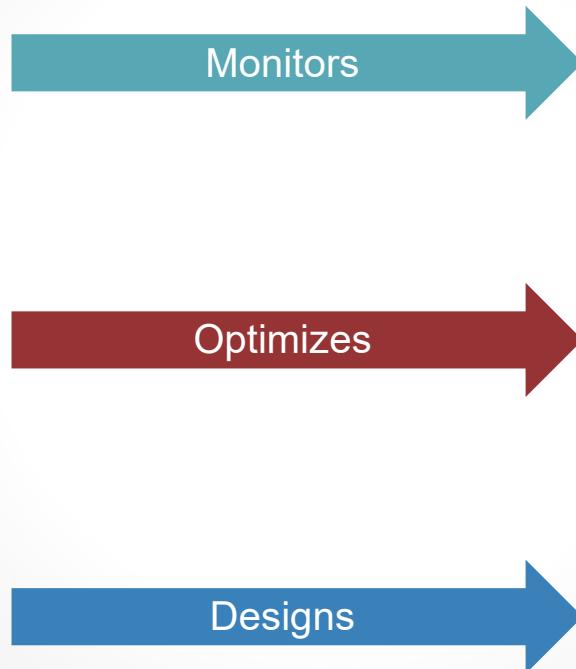


3. Resource allocation algorithms are often executed repeatedly

- ❑ E.g., routing, embedding, switching...

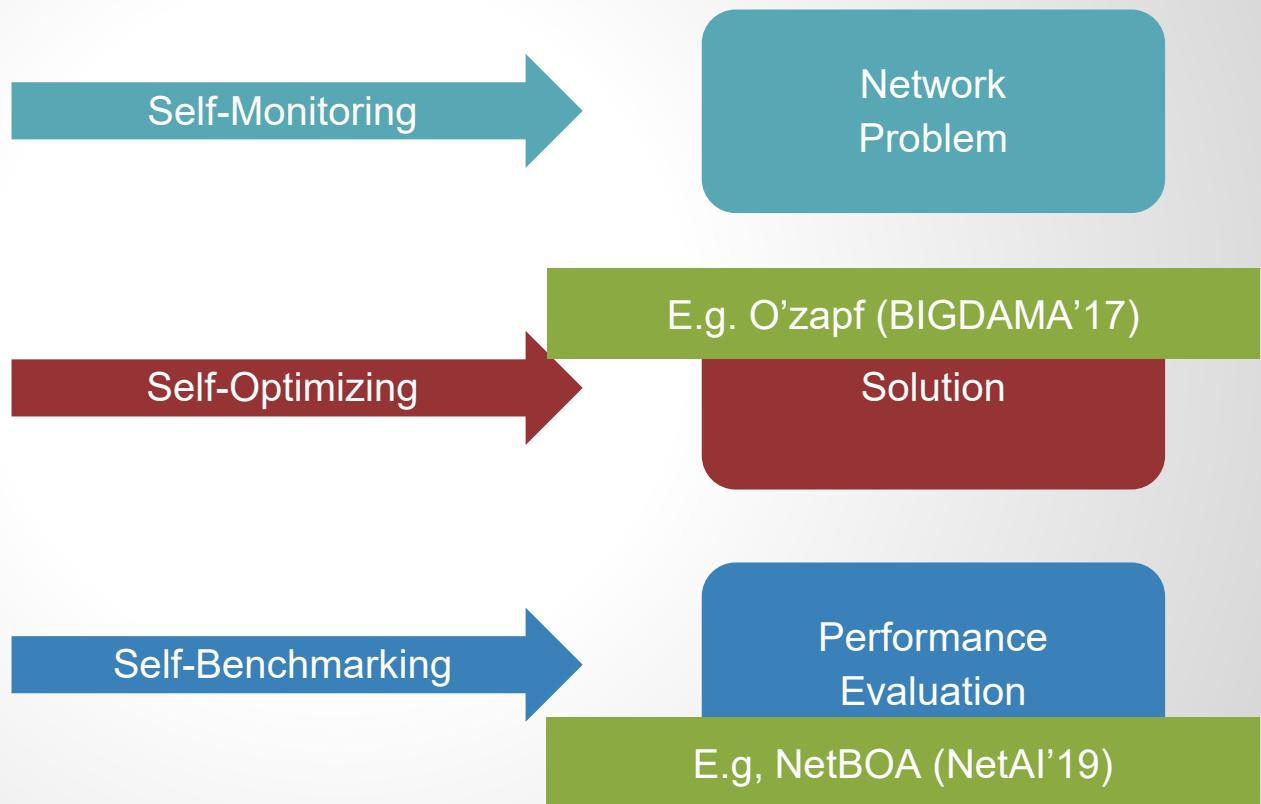


# Today's Approach to Operate Networks

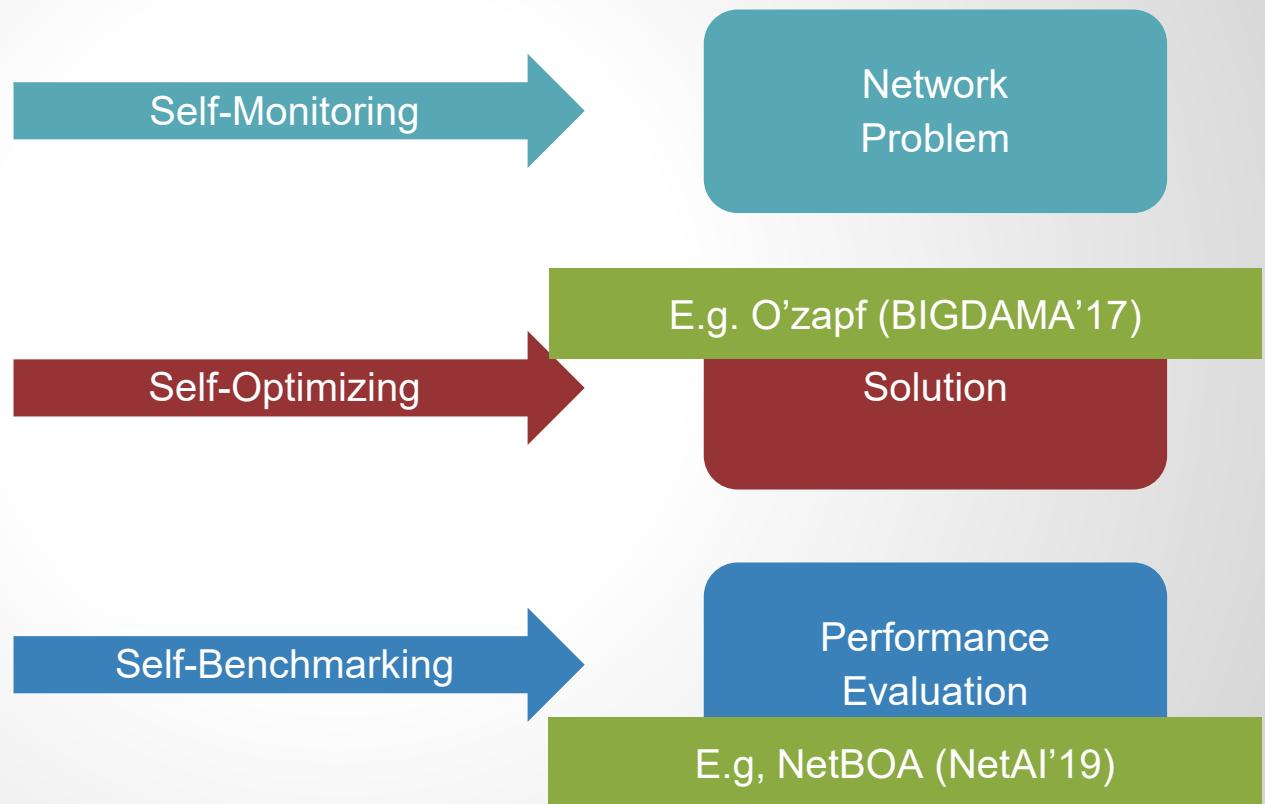


With more complex networks: need for automation!

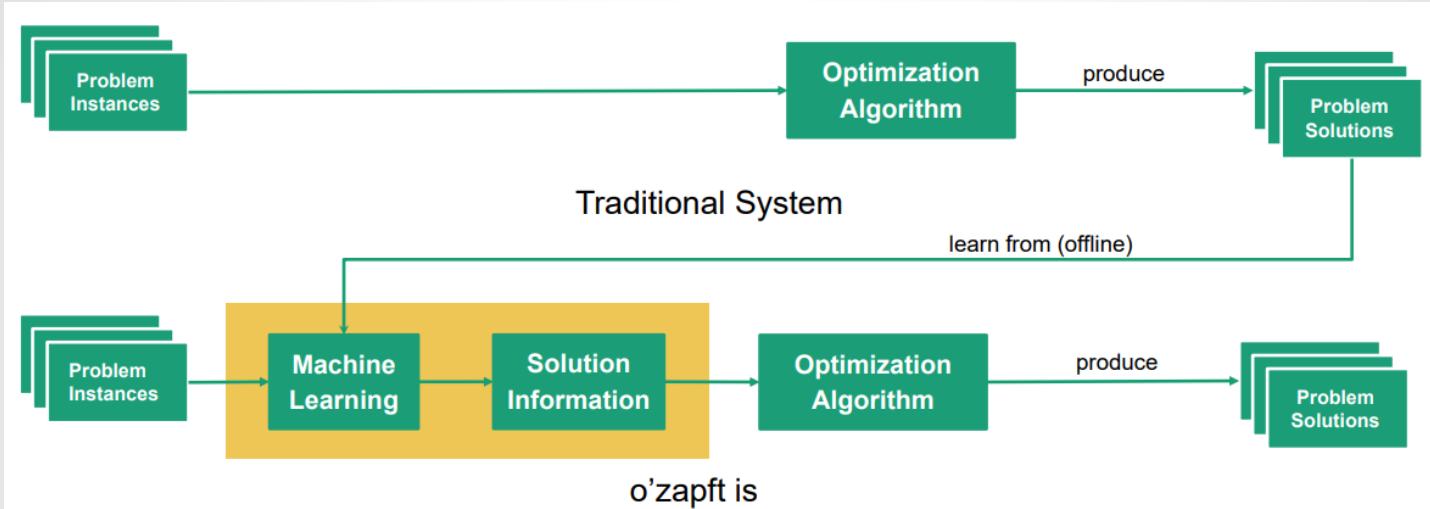
# What Self-Driving Networks Could Do?



# What Self-Driving Networks Could Do?



# Example: Data-Driven Algorithms

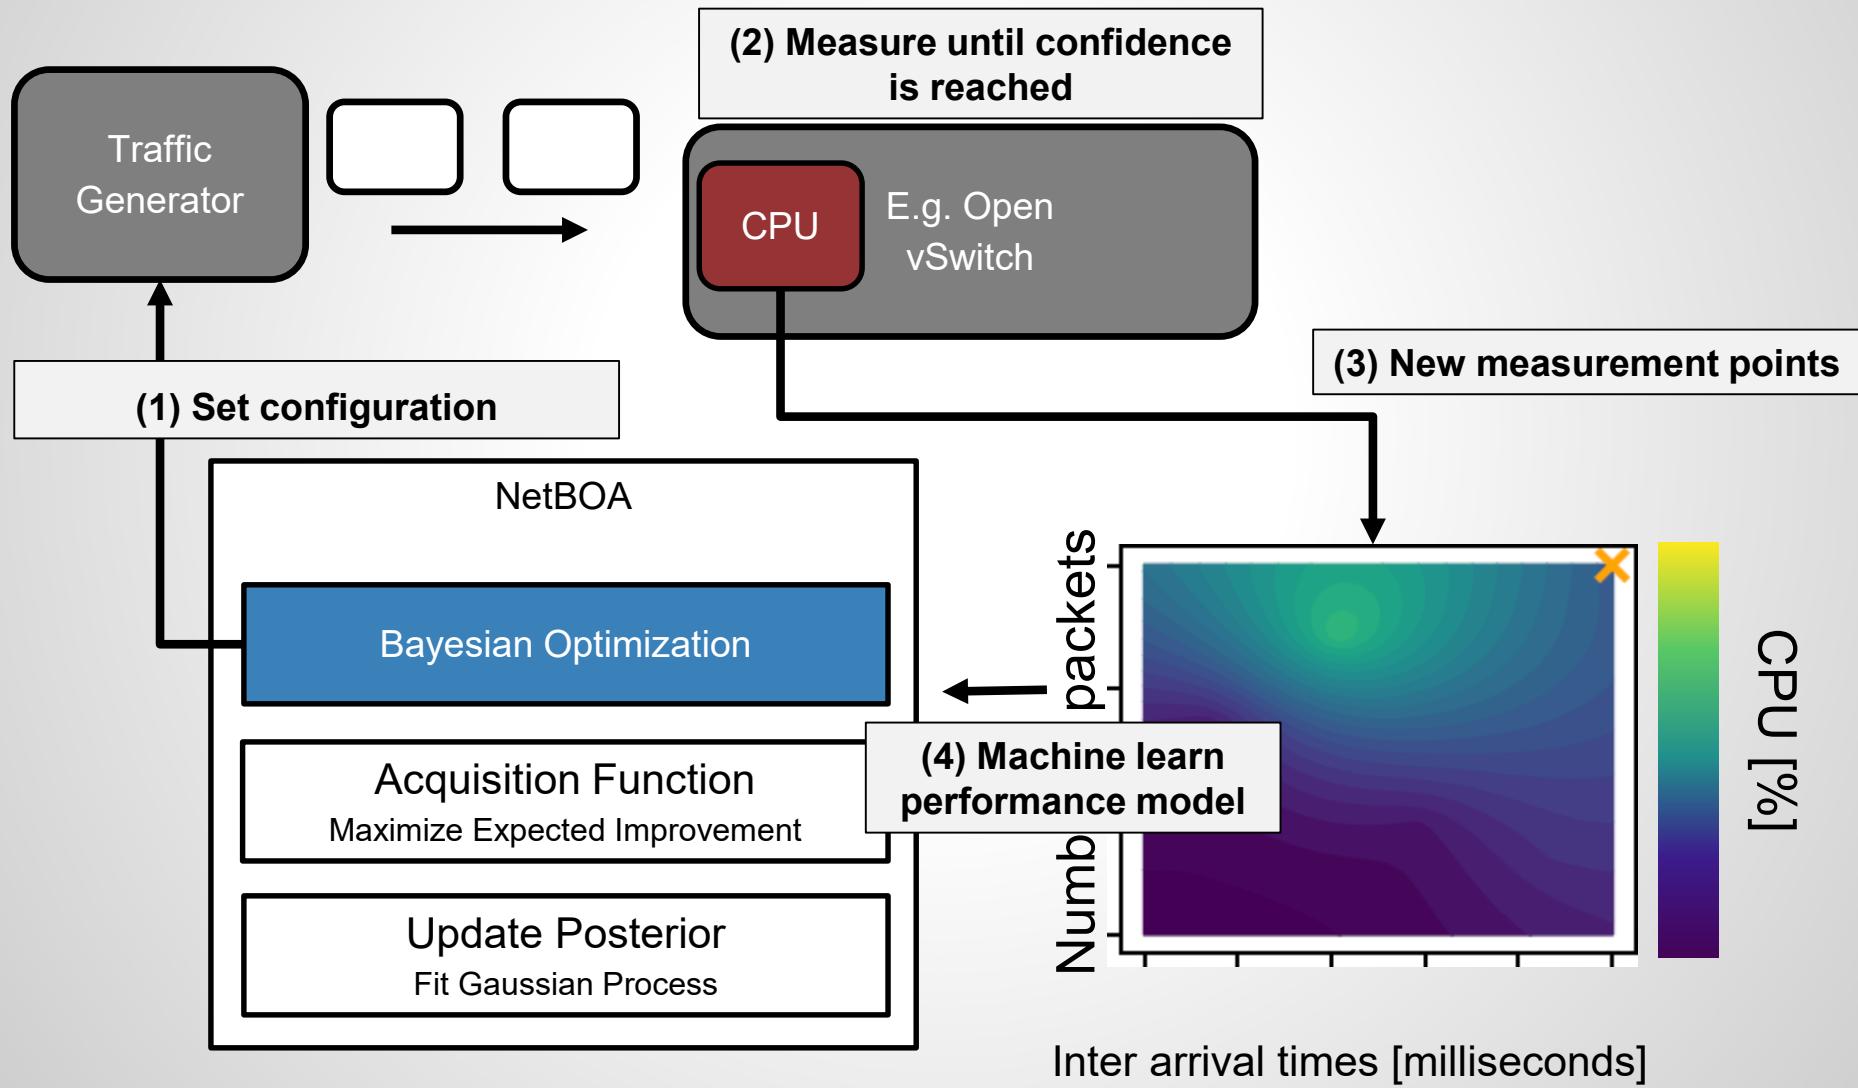


Can we learn from past solutions?

- E.g., to speed up future solutions?

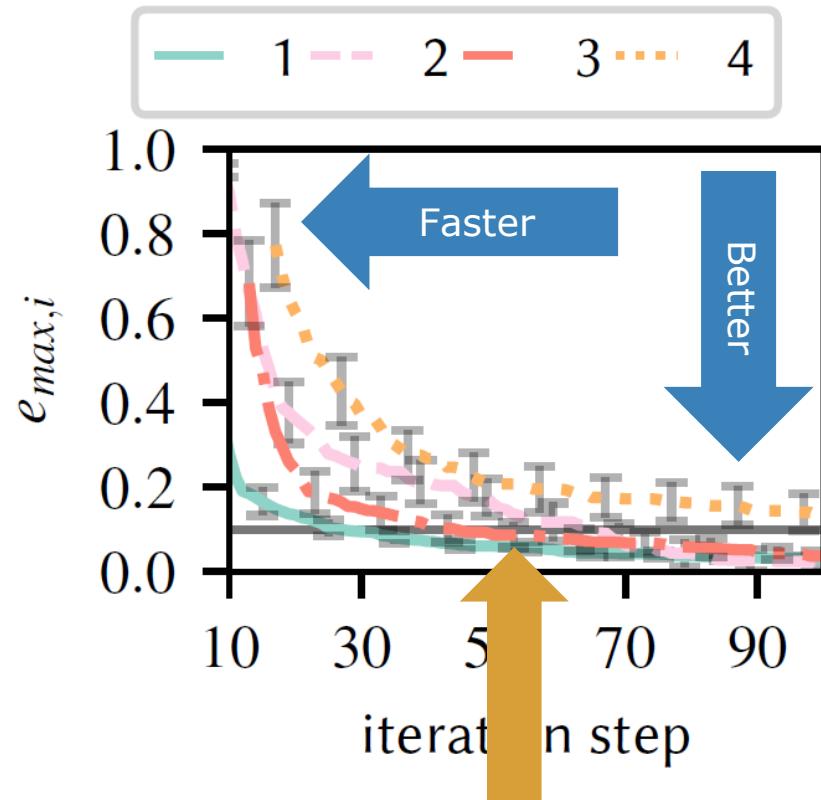
# Example: NetBOA

## Automated Learning of “Bad Inputs”

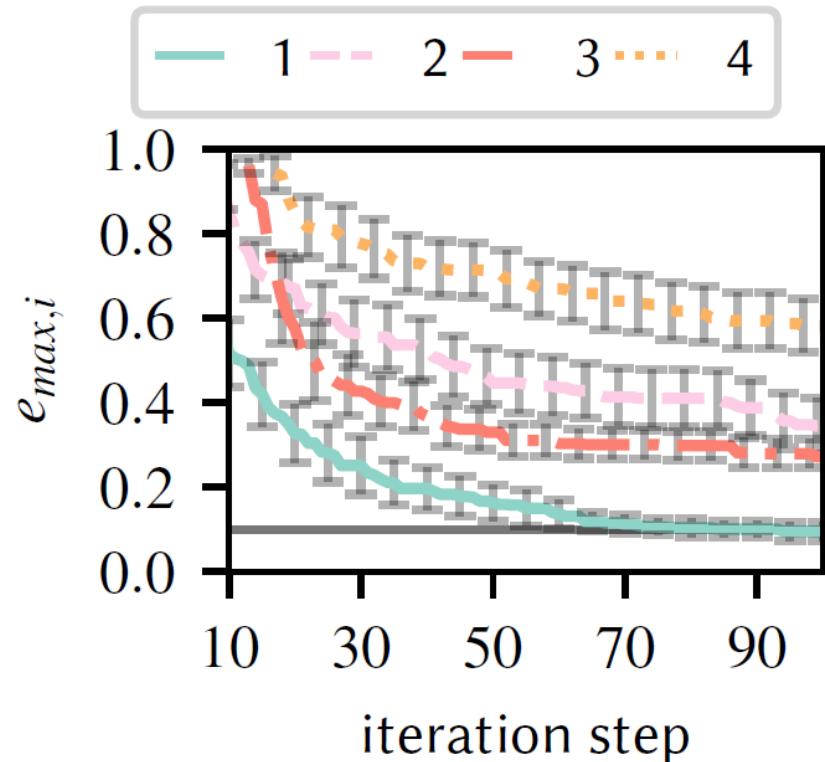


# NetBOA vs Random Search

**NetBOA**

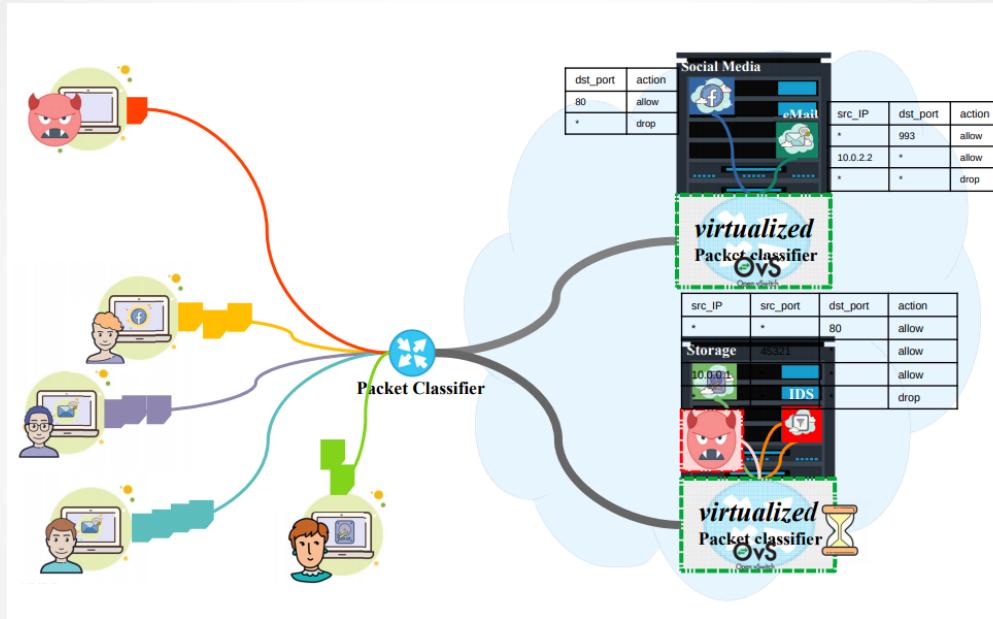


**Random Search**



**24 % higher CPU utilization**

# May Also Be Exploited: Algorithmic Complexity Attacks



E.g., automated learning of bad inputs to **packet classifier**

- ❑ E.g., difficult regular expressions
- ❑ Severely affects performance of OvS
- ❑ Can result in denial-of-service

# Challenges of AI-Based and Self-Driving Networks

- How much control are we willing to give away?
- Can a self-\* network realize its **limits**?
- E.g., when quality of **input data** is not good enough?
- When to hand over to human? Or **fall back** to „safe/oblivious mode“?
- Can we learn from self-driving **cars**?



# Conclusion

- ❑ Programmability and virtualization: algorithmic opportunities but also challenges
  - ❑ E.g.,: faster innovation, flexibilities in resource allocation, etc.
  - ❑ But, e.g.: **performance isolation** needs to be ensured across all involved resources, resulting resource allocation **problems hard** (open: good LP formulations, accounting for **latencies, derandomization**, special graphs, etc.)
- ❑ Security: more opportunities and challenges
  - ❑ Also faster **innovation**, but also new **attack surface** and potentially high impact
- ❑ AI opens interesting new opportunities
  - ❑ To deal with algorithmic complexities
  - ❑ To deal with modelling complexities
  - ❑ To find performance weaknesses
  - ❑ But also new challenges: how much control can we **give away**?

# References

AI Approaches	<p><a href="#">On The Impact of the Network Hypervisor on Virtual Network Performance</a> Andreas Blenk, Arsany Basta, Wolfgang Kellerer, and Stefan Schmid. <b>IFIP Networking</b>, Warsaw, Poland, May 2019.</p> <p><a href="#">Waypoint Routing in Special Networks</a> Saeed Akhoondian Amiri, Klaus-Tycho Foerster, Riko Jacob, Mahmoud Parham, and Stefan Schmid. <b>IFIP Networking</b>, Zurich, Switzerland, May 2018.</p> <p><a href="#">Walking Through Waypoints</a> Saeed Akhoondian Amiri, Klaus-Tycho Foerster, and Stefan Schmid. 13th Latin American Theoretical Informatics Symposium (<b>LATIN</b>), Buenos Aires, Argentina, April 2018.</p> <p><a href="#">Charting the Algorithmic Complexity of Waypoint Routing</a> Saeed Akhoondian Amiri, Klaus-Tycho Foerster, Riko Jacob, and Stefan Schmid. ACM SIGCOMM Computer Communication Review (<b>CCR</b>), 2018.</p> <p><a href="#">Virtual Network Embedding Approximations: Leveraging Randomized Rounding</a> Matthias Rost and Stefan Schmid. IEEE/ACM Transactions on Networking (<b>TON</b>), 2019.</p> <p><a href="#">Parametrized Complexity of Virtual Network Embeddings: Dynamic &amp; Linear Programming Approximations</a> Matthias Rost, Elias Döhne, and Stefan Schmid. ACM SIGCOMM Computer Communication Review (<b>CCR</b>), January 2019.</p> <p><a href="#">Charting the Complexity Landscape of Virtual Network Embeddings</a> (Best Paper Award) Matthias Rost and Stefan Schmid. <b>IFIP Networking</b>, Zurich, Switzerland, May 2018.</p> <p><a href="#">Competitive and Deterministic Embeddings of Virtual Networks</a> Guy Even, Moti Medina, Gregor Schaffrath, and Stefan Schmid. Journal Theoretical Computer Science (<b>TCS</b>), Elsevier, 2013.</p> <p><a href="#">MTS: Bringing Multi-Tenancy to Virtual Switches</a> Kashyap Thimmaraju, Saad Hermak, Gabor Retvari, and Stefan Schmid. USENIX Annual Technical Conference (<b>ATC</b>), Renton, Washington, USA, July 2019.</p> <p><a href="#">Taking Control of SDN-based Cloud Systems via the Data Plane</a> Kashyap Thimmaraju, Bhargava Shastry, Tobias Fiebig, Felicitas Hetzelt, Jean-Pierre Seifert, Anja Feldmann, and Stefan Schmid. ACM Symposium on SDN Research (<b>SOSR</b>), Los Angeles, California, USA, March 2018.</p> <p><a href="#">NetBOA: Self-Driving Network Benchmarking</a> Johannes Zerwas, Patrick Kalmbach, Laurenz Henkel, Gabor Retvari, Wolfgang Kellerer, Andreas Blenk, and Stefan Schmid. ACM SIGCOMM Workshop on Network Meets AI &amp; ML (<b>NetAI</b>), Beijing, China, August 2019.</p> <p><a href="#">o'zapft is: Tap Your Network Algorithm's Big Data!</a> Andreas Blenk, Patrick Kalmbach, Stefan Schmid, and Wolfgang Kellerer. ACM SIGCOMM 2017 Workshop on Big Data Analytics and Machine Learning for Data Communication Networks (<b>Big-DAMA</b>), Los Angeles, California, USA, August 2017.</p> <p><a href="#">Tuple Space Explosion: A Denial-of-Service Attack Against a Software Packet Classifier</a> Levente Csikor, Dinil Mon Divakaran, Min Suk Kang, Attila Korosi, Balazs Sonkoly, David Haja, Dimitrios Pezaros, Stefan Schmid, and Gabor Retvari. ACM CoNEXT, Orlando, Florida, USA, December 2019.</p>
---------------	--